



search dblp

CVPR 2022

> Home > Conferences and Workshops > CVPR

Dagstuhl

- IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2022, New Orleans, LA, USA, June 18-24, 2022. IEEE 2022, ISBN 978-1-6654-6946-3
- Meina Zhang, Yingying Fang, Guoxi Ni, Tieyong Zeng: **Pixel screening based intermediate correction for blind deblurring.** 1-9
- Elijah Cole, Xuan Yang, Kimberly Wilber, Oisin Mac Aodha, Serge J. Belongie: **When Does Contrastive Visual Representation Learning Work?** 1-10
- Dengpan Fu, Dongdong Chen, Hao Yang, Jianmin Bao, Lu Yuan, Lei Zhang, Houqiang Li, Fang Wen, Dong Chen: **Large-Scale Pre-training for Person Re-identification with Noisy Labels.** 1-11
- Yunhui Guo, Xudong Wang, Yubei Chen, Stella X. Yu: **Clipped Hyperbolic Classifiers Are Super-Hyperbolic Classifiers.** 1-10
- Yunhui Guo, Haoran Guo, Stella X. Yu: **CO-SNE: Dimensionality Reduction and Visualization for Hyperbolic Data.** 11-20
- Jinyu Cai, Jicong Fan, Wenzhong Guo, Shiping Wang, Yunhe Zhang, Zhao Zhang: **Efficient Deep Embedded Subspace Clustering.** 21-30
- Jiexi Yan, Lei Luo, Chenghao Xu, Cheng Deng, Heng Huang: **Noise Is Also Useful: Negative Correlation-Steered Latent Contrastive Learning.** 31-40
- Kun-Peng Ning, Xun Zhao, Yu Li , Sheng-Jun Huang: **Active Learning for Open-set Annotation.** 41-49
- Theodoros Tsiligkaridis, Jay Roberts: **Understanding and Increasing Efficiency of Frank-Wolfe Adversarial Training.** 50-59
- Kezhi Kong, Guohao Li , Mucong Ding, Zuxuan Wu, Chen Zhu, Bernard Ghanem , Gavin Taylor, Tom Goldstein: **Robust Optimization as Data Augmentation for Large-scale Graphs.** 60-69
- Sihao Yu, Jiafeng Guo, Ruqing Zhang, Yixing Fan, Zizhen Wang, Xueqi Cheng: **A Re-Balancing Strategy for Class-Imbalanced Classification Based on Instance Difficulty.** 70-79
- Bingyuan Liu, Ismail Ben Ayed, Adrian Galdran , Jose Dolz: **The Devil is in the Margin: Margin-based Label Smoothing for Network Calibration.** 80-88
- Guoliang Lin, Hanlu Chu, Hanjiang Lai: **Towards Better Plasticity-Stability Trade-off in Incremental Learning: A Simple Linear Connector.** 89-98
- Rishabh Tiwari, KrishnaTeja Killamsetty, Rishabh K. Iyer, Pradeep Shenoy: **GCR: Gradient Coreset based Replay Buffer Selection for Continual Learning.** 99-108
- Qingsen Yan, Dong Gong , Yuhang Liu, Anton van den Hengel, Javen Qinfeng Shi : **Learning Bayesian Sparse Networks with Full Experience Replay for Continual Learning.** 109-118
- Daniel Grzech, Mohammad Farid Azampour, Ben Glocker, Julia A. Schnabel, Nassir Navab, Bernhard Kainz , Loïc Le Folgoc: **A variational Bayesian method for similarity learning in non-rigid image registration.** 119-128
- Yadong Ding, Yu Wu , Chengyue Huang, Siliang Tang, Yi Yang, Longhui Wei, Yuetong Zhuang, Qi Tian: **Learning to Learn by Jointly Optimizing Neural Architecture and Weights.** 129-138
- Zifeng Wang, Zizhao Zhang, Chen-Yu Lee, Han Zhang, Ruoxi Sun, Xiaoqi Ren, Guolong Su, Vincent Perot, Jennifer G. Dy, Tomas Pfister: **Learning to Prompt for Continual Learning.** 139-149
- Mengqi Xue , Haofei Zhang , Jie Song, Mingli Song: **Meta-attention for ViT-backed Continual Learning.** 150-159
- Vitor Guizilini, Rares Ambrus, Dian Chen, Sergey Zakharov, Adrien Gaidon: **Multi-Frame Self-Supervised Depth with Transformers.** 160-170

SPARQL queries

[run query for this table of contents](#)

- frequent references
 - [highly cited publications](#)
- or build your own?

Refine list

showing all ?? records

- Zhen Wang, Liu Liu, Yiqun Duan, Yajing Kong, Dacheng Tao:
Continual Learning with Lifelong Vision Transformer. 171-181
- Jianfeng Wang, Thomas Lukasiewicz:
Rethinking Bayesian Deep Learning Methods for Semi-Supervised Volumetric Medical Image Segmentation. 182-190
- Yawei Li , Kamil Adamczewski, Wen Li, Shuhang Gu, Radu Timofte , Luc Van Gool:
Revisiting Random Channel Pruning for Neural Network Compression. 191-201
- Huayi Tang, Yong Liu:
Deep Safe Multi-view Clustering: Reducing the Risk of Clustering Performance Degradation Caused by View Increase. 202-211
- Jongin Lim, Sangdoo Yun, Seulki Park, Jin Young Choi:
Hypergraph-Induced Semantic Triplet Loss for Deep Metric Learning. 212-222
- Prateek Munjal, Nasir Hayat, Munawar Hayat, Jamshid Sourati, Shadab Khan:
Towards Robust and Reproducible Active Learning using Neural Networks. 223-232
- Jiulong Liu, Zhaoqiang Liu:
Non-Iterative Recovery from Nonlinear Observations using Generative Models. 233-243
- Minyoung Kim:
Gaussian Process Modeling of Approximate Inference Errors for Variational Autoencoders. 244-253
- Kwang In Kim:
Robust Combination of Distributed Gradients Under Adversarial Perturbations. 254-263
- Lan Wang, Vishnu Naresh Boddegi:
Do learned representations respect causal relationships? 264-274
- Rafid Mahmood, James Lucas, David Acuna, Daiqing Li, Jonah Philion, José M. Álvarez, Zhiding Yu, Sanja Fidler, Marc T. Law:
How Much More Data Do I Need? Estimating Requirements for Downstream Tasks. 275-284
- Magzhan Gabidolla, Miguel Á. Carreira-Perpiñán:
Pushing the Envelope of Gradient Boosting Forests via Globally-Optimized Oblique Trees. 285-294
- Dian Chen, Dequan Wang, Trevor Darrell, Sayna Ebrahimi:
Contrastive Test-Time Adaptation. 295-305
- Paritosh Mittal, Yen-Chi Cheng , Maneesh Singh, Shubham Tulsiani:
AutoSDF: Shape Priors for 3D Completion, Reconstruction and Generation. 306-315
- Shikun Li , Xiaobo Xia, Shimeng Ge, Tongliang Liu :
Selective-Supervised Contrastive Learning with Noisy Labels. 316-325
- Yufei Guo, Xinyi Tong , Yuanpei Chen, Liwen Zhang, Xiaode Liu, Zhe Ma, Xuhui Huang:
RecDis-SNN: Rectifying Membrane Potential Distribution for Directly Training Spiking Neural Networks. 326-335
- M. Saquib Sarfraz, Marios Koulakis, Constantin Seibold, Rainer Stiefelhagen:
Hierarchical Nearest Neighbor Graph Embedding for Efficient Dimensionality Reduction. 336-345
- Yikai Wang, Xinwei Sun , Yanwei Fu :
Scalable Penalized Regression for Noise Detection in Learning with Noisy Labels. 346-355
- Xiran Fan, Chun-Hao Yang , Baba C. Vemuri :
Nested Hyperbolic Spaces for Dimensionality Reduction and Hyperbolic NN Design. 356-365
- Ivor J. A. Simpson , Sara Vicente, Neill D. F. Campbell :
Learning Structured Gaussians to Approximate Deep Ensembles. 366-374
- Ruoyu Wang, Mingyang Yi, Zhitang Chen, Shengyu Zhu:
Out-of-distribution Generalization with Causal Invariant Transformations. 375-385
- Tom Ryder, Chen Zhang, Ning Kang, Shifeng Zhang:
Split Hierarchical Variational Compression. 386-395
- Iordanis Fostiroopoulos, Barry W. Boehm:
Implicit Feature Decoupling with Depthwise Quantization. 396-405
- Jurijs Nazarovs, Zhichun Huang, Songwong Tasneeyapant, Rudrasis Chakraborty, Vikas Singh:
Understanding Uncertainty Maps in Vision with Statistical Testing. 406-416
- Anh-Dzung Doan , Michele Sasdelli , David Suter , Tat-Jun Chin:
A Hybrid Quantum-Classical Algorithm for Robust Fitting. 417-427
- Paul Roetzer, Paul Swoboda, Daniel Cremers , Florian Bernard:
A Scalable Combinatorial Solver for Elastic Geometrically Consistent 3D Shape Matching. 428-438
- Ahmed Abbas , Paul Swoboda:
FastDOG: Fast Discrete Optimization on GPU. 439-449
- Vladimir Chikin, Mikhail Antukh:
Data-Free Network Compression via Parametric Non-uniform Mixed Precision Quantization. 450-459

- Huu Le, Rasmus Kjær Høier , Che-Tsung Lin , Christopher Zach:
AdaSTE: An Adaptive Straight-Through Estimator to Train Binary Neural Networks. 460-469
- Sanjeev Muralikrishnan , Siddhartha Chaudhuri, Noam Aigerman, Vladimir G. Kim, Matthew Fisher, Niloy J. Mitra:
GLASS: Geometric Latent Augmentation for Shape Spaces. 470-479
- Matteo Spallanzani, Gian Paolo Leonardi, Luca Benini:
Training Quantised Neural Networks with STE Variants: the Additive Noise Annealing Algorithm. 470-479
- Nuo Xu, Jianlong Chang, Xing Nie, Chunlei Huo, Shiming Xiang, Chunhong Pan:
AME: Attention and Memory Enhancement in Hyper-Parameter Optimization. 480-489
- Christina Baek, Ziyang Wu, Kwan Ho Ryan Chan, Tianjiao Ding, Yi Ma, Benjamin D. Haeffele:
Efficient Maximal Coding Rate Reduction by Variational Forms. 490-498
- Marvin Eisenberger, Aysim Toker, Laura Leal-Taixé, Florian Bernard, Daniel Cremers :
A Unified Framework for Implicit Sinkhorn Differentiation. 499-508
- Yidong Chen, Chen Li, Zhonghua Lu :
Computing Wasserstein-\$p\$ Distance Between Images with Linear Cost. 509-518
- Natacha Kuete Meli, Florian Mannel
- Nourhan Bayasi, Ghassan Hamarneh, Rafeef Garbi:
BoosterNet: Improving Domain Generalization of Deep Neural Nets using Culpability-Ranked Features. 528-538
- Dong-Hwan Jang, Sanghyeok Chu, Joonhyuk Kim, Bohyung Han:
Pooling Revisited: Your Receptive Field is Suboptimal. 539-548
- Jiajing Chen, Burak Kakillioglu, Huantao Ren, Senem Velipasalar:
Why Discard if You can Recycle?: A Recycling Max Pooling Module for 3D Point Cloud Analysis. 549-557
- Mu Hu, Junyi Feng, Jiashen Hua, Baisheng Lai, Jianqiang Huang, Xiaojin Gong, Xiansheng Hua:
Online Convolutional Reparameterization. 558-567
- Xiaohan Ding, Honghao Chen, Xiangyu Zhang, Jungong Han, Guiguang Ding:
RepMLPNet: Hierarchical Vision MLP with Re-parameterized Locality. 568-577
- Tao Huang, Shan You, Bohan Zhang, Yuxuan Du, Fei Wang, Chen Qian, Chang Xu:
DyRep: Bootstrapping Training with Dynamic Re-parameterization. 578-587
- Tianlong Chen, Zhenyu Zhang, Yihua Zhang, Shiyu Chang, Sijia Liu, Zhangyang Wang:
Quarantine: Sparsity Can Uncover the Trojan Attack Trigger for Free. 588-599
- Anil Kag, Venkatesh Saligrama :
Condensing CNNs with Partial Differential Equations. 600-609
- Shaojie Bai, Zhengyang Geng, Yash Savani, J. Zico Kolter:
Deep Equilibrium Optical Flow Estimation. 610-620
- Matan Atzmon, Koki Nagano, Sanja Fidler, Sameh Khamis, Yaron Lipman:
Frame Averaging for Equivariant Shape Space Learning. 621-631
- Gee-Sern Hsu, Chun-Hung Tsai, Hung-Yi Wu:
Dual-Generator Face Reenactment. 632-640
- Rongzhen Zhao, Jian Li, Zhenzhi Wu:
Convolution of Convolution: Let Kernels Spatially Collaborate. 641-650
- Matthias Wödlinger, Jan Kotera , Jan Xu, Robert Sablatnig:
SASIC: Stereo Image Compression with Latent Shifts and Stereo Attention. 651-660
- Michael Schelling, Pedro Hermosilla, Timo Ropinski :
RADU: Ray-Aligned Depth Update Convolutions for ToF Data Denoising. 661-670
- Utkarsh Singhal, Yifei Xing, Stella X. Yu:
Co-domain Symmetry for Complex-Valued Deep Learning. 671-680
- Tong Yu, Ruslan Khalitov, Lei Cheng, Zhirong Yang:
Paramixer: Parameterizing Mixing Links in Sparse Factors Works Better than Dot-Product Self-Attention. 681-690
- Huanyu Wang, Junjie Liu , Xin Ma, Yang Yong, Zhenhua Chai, Jianxin Wu:
Compressing Models with Few Samples: Mimicking then Replacing. 691-700
- Raymond A. Yeh, Yuan-Ting Hu, Zhongzheng Ren, Alexander G. Schwing:
Total Variation Optimization Layers for Computer Vision. 701-711
- Vinit Veerendraveer Singh, Chandra Kambhamettu:
AIM: an Auto-Augmenter for Images and Meshes. 712-721
- George Yiasemis, Jan-Jakob Sonke, Clarisa Sánchez, Jonas Teuwen:
Recurrent Variational Network: A Deep Learning Inverse Problem Solver applied to the task of Accelerated MRI Reconstruction. 722-731
- Nicolas Donati, Etienne Corman, Maks Ovsjanikov:
Deep orientation-aware functional maps: Tackling symmetry issues in Shape Matching. 732-741

- Jingqi Huang, Yue Ning, Dong Nie, Linan Guan, Xiping Jia:
Weakly-supervised Metric Learning with Cross-Module Communications for the Classification of Anterior Chamber Angle Images. 742-752
- Lei Huang, Yi Zhou, Tian Wang, Jie Luo, Xianglong Liu:
Delving into the Estimation Shift of Batch Normalization in a Network. 753-762
- Fanqing Lin, Brian L. Price, Tony R. Martinez:
Generalizing Interactive Backpropagating Refinement for Dense Prediction Networks. 763-772
- Wenshuo Li, Hanting Chen, Jianyuan Guo , Ziyang Zhang, Yunhe Wang:
Brain-inspired Multilayer Perceptron with Spiking Neurons. 773-783
- Koushik Biswas, Sandeep Kumar, Shilpak Banerjee , Ashish Kumar Pandey:
Smooth Maximum Unit: Smooth Activation Function for Deep Networks using Smoothing Maximum Technique. 784-793
- Mannat Singh, Laura Gustafson, Aaron Adcock, Vinicius de Freitas Reis, Bugra Gedik, Raj Prateek Kosaraju, Dhruv Mahajan, Ross B. Girshick, Piotr Dollár, Laurens van der Maaten:
Revisiting Weakly Supervised Pre-Training of Visual Perception Models. 794-804
- Xuran Pan, Chunjiang Ge, Rui Lu, Shiji Song, Guanfu Chen, Zeyi Huang, Gao Huang:
On the Integration of Self-Attention and Convolution. 805-815
- Jianyuan Guo , Yehui Tang, Kai Han, Xinghao Chen, Han Wu, Chao Xu, Chang Xu, Yunhe Wang:
Hire-MLP: Vision MLP via Hierarchical Rearrangement. 816-826
- Benjamin Naoto Chiche, Arnaud Woiselle, Joana Frontera-Pons , Jean-Luc Starck :
Stable Long-Term Recurrent Video Super-Resolution. 827-836
- Aming Wu, Cheng Deng:
Single-Domain Generalized Object Detection in Urban Scene via Cyclic-Disentangled Self-Distillation. 837-846
- Anlin Zheng, Yuang Zhang, Xiangyu Zhang, Xiaojuan Qi, Jian Sun:
Progressive End-to-End Object Detection in Crowded Scenes. 847-856
- Ajay Jain, Ben Mildenhall, Jonathan T. Barron, Pieter Abbeel, Ben Poole:
Zero-Shot Text-Guided Object Generation with Dream Fields. 857-866
- Mingjin Zhang, Rui Zhang, Yuxiang Yang, Haichen Bai , Jing Zhang, Jie Guo :
ISNet: Shape Matters for Infrared Small Target Detection. 867-876
- Yi-Nan Chen, Hang Dai, Yong Ding:
Pseudo-Stereo for Monocular 3D Object Detection in Autonomous Driving. 877-887
- Tu Zheng, Yifei Huang, Yang Liu, Wenjian Tang, Zheng Yang, Deng Cai, Xiaofei He:
CLRNet: Cross Layer Refinement Network for Lane Detection. 888-897
- Yanan Zhang , Jiaxin Chen, Di Huang:
CAT-Det: Contrastively Augmented Transformer for Multimodal 3D Object Detection. 898-907
- Yu-Jhe Li, Jinyung Park, Matthew O'Toole, Kris Kitani:
Modality-Agnostic Learning for Radar-Lidar Fusion in Vehicle Detection. 908-917
- Yanbin Hao, Hao Zhang, Chong-Wah Ngo, Xiangnan He:
Group Contextualization for Video Recognition. 918-928
- Suchen Wang, Yueqi Duan, Henghui Ding , Yap-Peng Tan, Kim-Hui Yap, Junsong Yuan:
Learning Transferable Human-Object Interaction Detector with Natural Language Supervision. 929-938
- Gongjie Zhang, Zhipeng Luo, Yingchen Yu, Kaiwen Cui, Shijian Lu:
Accelerating DETR Convergence via Semantic-Aligned Matching. 939-948
- Jialian Wu, Sudhir Yaram, Hui Liang, Tian Lan, Junsong Yuan, Jayan Eledath , Gérard G. Medioni:
Efficient Video Instance Segmentation via Tracklet Query and Proposal. 949-958
- Zhaozheng Chen, Tan Wang, Xiongwei Wu, Xian-Sheng Hua, Hanwang Zhang, Qianru Sun:
Class Re-Activation Maps for Weakly-Supervised Semantic Segmentation. 959-968
- Siyue Yu , Jimin Xiao, Bingfeng Zhang , Eng Gee Lim:
Democracy Does Matter: Comprehensive Feature Mining for Co-Salient Object Detection. 969-978
- Jinheng Xie, Jianfeng Xiang, Junliang Chen, Xianxu Hou, Xiaodong Zhao, Linlin Shen:
C² AM: Contrastive learning of Class-agnostic Activation Map for Weakly Supervised Object Localization and Semantic Segmentation. 979-988
- Ayan Kumar Bhunia, Subhadeep Koley , Abdullah Faiz Ur Rahman Khilji, Aneeshan Sain , Pinaki Nath Chowdhury, Tao Xiang, Yi-Zhe Song :
Sketching without Worrying: Noise-Tolerant Sketch-Based Image Retrieval. 989-998
- Hao Li, Tianwen Fu, Jifeng Dai , Hongsheng Li , Gao Huang, Xizhou Zhu:
AutoLoss-Zero: Searching Loss Functions from Scratch for Generic Tasks. 999-1008
- Jihwan Park , Seungjun Lee, Hwan Heo, Hyeong Kyu Choi, Hyunwoo J. Kim:
Consistency Learning via Decoding Path Augmentation for Transformers in Human Object Interaction Detection. 1009-1018

- Hanyu Xuan, Zhiliang Wu , Jian Yang, Yan Yan, Xavier Alameda-Pineda:
A Proposal-based Paradigm for Self-supervised Sound Source Localization in Videos. 1019-1028
- Canjie Luo, Lianwen Jin, Jingdong Chen:
SimAN: Exploring Self-Supervised Representation Learning of Scene Text via Similarity-Aware Normalization. 1029-1038
- Shangbang Long, Siyang Qin, Dmitry Panteleev, Alessandro Bissacco, Yasuhisa Fujii, Michalis Raptis:
Towards End-to-End Unified Scene Text Detection and Layout Analysis. 1039-1049
- Xinqian Gu, Hong Chang, Bingpeng Ma, Shutao Bai, Shiguang Shan, Xilin Chen:
Clothes-Changing Person Re-identification with RGB Modality Only. 1050-1059
- Qing Lian, Peiliang Li, Xiaozhi Chen :
MonoJSG: Joint Semantic and Geometric Cost Volume for Monocular 3D Object Detection. 1060-1069
- Jiaqi Gu , Bojian Wu , Lubin Fan, Jianqiang Huang, Shen Cao, Zhiyu Xiang, Xian-Sheng Hua:
Homography Loss for Monocular 3D Object Detection. 1070-1079
- Xuyang Bai, Zeyu Hu, Xinge Zhu, Qingqiu Huang, Yilun Chen, Hongbo Fu, Chiew-Lan Tai:
TransFusion: Robust LiDAR-Camera Fusion for 3D Object Detection with Transformers. 1080-1089
- Ruihang Chu, Xiaoqing Ye, Zhengzhe Liu, Xiao Tan, Xiaojuan Qi, Chi-Wing Fu, Jiaya Jia:
TWIST: Two-Way Inter-label Self-Training for Semi-supervised 3D Instance Segmentation. 1090-1099
- Haiyang Wang, Shaoshuai Shi, Ze Yang , Rongyao Fang, Qi Qian, Hongsheng Li , Bernt Schiele , Liwei Wang:
RBGNet: Ray-based Grouping for 3D Object Detection. 1100-1109
- Yanwei Li, Xiaojuan Qi, Yukang Chen, Liwei Wang, Zeming Li, Jian Sun, Jiaya Jia:
Voxel Field Fusion for 3D Object Detection. 1110-1119
- Yurong You , Katie Luo, Cheng Perng Phoo, Wei-Lun Chao, Wen Sun, Bharath Hariharan, Mark E. Campbell, Kilian Q. Weinberger:
Learning to Detect Mobile Objects from LiDAR Scans Without Labels. 1120-1130
- David Schinagl, Georg Krispel, Horst Possegger , Peter M. Roth, Horst Bischof :
OccAM's Laser: Occlusion-based Attribution Maps for 3D Object Detectors on LiDAR Data. 1131-1140
- Yichun Shen, Wanli Jiang, Zhen Xu, Rundong Li, Junghyun Kwon:
Confidence Propagation Cluster: Unleash Full Potential of Object Detectors. 1141-1151
- Sijie Zhu, Mubarak Shah, Chen Chen:
TransGeo: Transformer Is All You Need for Cross-view Image Geo-localization. 1152-1161
- Yongjian Deng , Hao Chen, Hai Liu, Youfu Li :
A Voxel Graph CNN for Object Classification with Event Cameras. 1162-1171
- Dongchen Lu, Dongmei Li, Yali Li, Shengjin Wang:
OSKDet: Orientation-sensitive Keypoint Localization for Rotated Object Detection. 1172-1182
- Yang You, Zelin Ye, Yujing Lou , Chengkun Li, Yong-Lu Li, Lizhuang Ma, Weiming Wang, Cewu Lu:
Canonical Voting: Towards Robust Oriented Bounding Box Detection in 3D Scenes. 1183-1192
- Jiaxing Huang, Dayan Guan, Aoran Xiao, Shijian Lu, Ling Shao:
Category Contrast for Unsupervised Domain Adaptation in Visual Tasks. 1193-1204
- Xiaohua Zhai, Alexander Kolesnikov, Neil Houlsby, Lucas Beyer:
Scaling Vision Transformers. 1204-1213
- Yihong Sun, Adam Kortylewski, Alan L. Yuille:
Amodal Segmentation through Out-of-Task and Out-of-Distribution Generalization with a Bayesian Model. 1205-1214
- Xingzhe He, Bastian Wandt, Helge Rhodin:
GANSeg: Learning to Segment by Unsupervised Hierarchical Image Generation. 1215-1225
- Anirud Thyagarajan, Benjamin Ummenhofer, Prashant Laddha, Om Ji Omer, Sreenivas Subramoney:
Segment-Fusion: Hierarchical Context Fusion for Robust 3D Semantic Segmentation. 1226-1235
- Liulei Li, Tianfei Zhou, Wenguan Wang , Jianwu Li, Yi Yang:
Deep Hierarchical Semantic Segmentation. 1236-1247
- Yifan Zhang, Bo Pang, Cewu Lu:
Semantic Segmentation by Early Region Proxy. 1248-1258
- Shubhankar Borse , Hyojin Park, Hong Cai, Debasmit Das, Risheek Garrepalli, Fatih Porikli:
Panoptic, Instance and Semantic Relations: A Relational Context Encoder to Enhance Panoptic Segmentation. 1259-1269
- Zhiqi Li, Wenhui Wang, Enze Xie, Zhiding Yu, Anima Anandkumar, José M. Álvarez, Ping Luo, Tong Lu:
Panoptic SegFormer: Delving Deeper into Panoptic Segmentation with Transformers. 1270-1279
- Bowen Cheng, Ishan Misra, Alexander G. Schwing, Alexander Kirillov, Rohit Girdhar:
Masked-attention Mask Transformer for Universal Image Segmentation. 1280-1289
- Xi Chen, Zhiyan Zhao, Yilei Zhang, Manni Duan, Donglian Qi, Hengshuang Zhao:
FocalClick: Towards Practical Interactive Image Segmentation. 1290-1299

- Tiancheng Shen, Yuechen Zhang, Lu Qi, Jason Kuen, Xingyu Xie, Jianlong Wu, Zhe Lin, Jiaya Jia: **High Quality Segmentation for Ultra High-resolution Images.** 1300-1309
- Wenwen Pan, Haonan Shi, Zhou Zhao, Jieming Zhu, Xiuqiang He, Zhigeng Pan, Lianli Gao, Jun Yu, Fei Wu, Qi Tian: **Wnet: Audio-Guided Video Object Segmentation via Wavelet-Based Cross- Modal Denoising Networks.** 1310-1321
- Mingxing Li, Li Hu, Zhiwei Xiong, Bang Zhang, Pan Pan, Dong Liu: **Recurrent Dynamic Embedding for Video Object Segmentation.** 1322-1331
- Kai Xu, Angela Yao: **Accelerating Video Object Segmentation with Compressed Video.** 1332-1341
- Kwanyong Park, Sanghyun Woo, Seoung Wug Oh, In So Kweon, Joon-Young Lee: **Per-Clip Video Object Segmentation.** 1342-1351
- Zihui Lin, Tianyu Yang, Maomao Li, Ziyu Wang, Chun Yuan, Wenhao Jiang, Wei Liu: **SWEM: Towards Real-Time Video Object Segmentation with Sequential Weighted Expectation-Maximization.** 1352-1362
- Hanyuan Liu, Chengze Li, Xuetong Liu, Tien-Tsin Wong: **Neural Recognition of Dashed Curves with Gestalt Law of Continuity.** 1363-1372
- Ziqiang Xu, Chunyan Xu, Zhen Cui, Xiangwei Zheng, Jian Yang: **CVNet: Contour Vibration Network for Building Extraction.** 1373-1381
- Jinsheng Wang, Yinchen Ma, Shaofei Huang, Tianrui Hui, Fei Wang, Chen Qian, Tianzhu Zhang: **A Keypoint-based Global Association Network for Lane Detection.** 1382-1391
- Mengyang Pu, Yaping Huang, Yuming Liu, Qingji Guan, Haibin Ling: **EDTER: Edge Detection with Transformer.** 1392-1402
- Yining Hong, Kaichun Mo, Li Yi, Leonidas J. Guibas, Antonio Torralba, Joshua B. Tenenbaum, Chuang Gan: **Fixing Malfunctional Objects With Learned Physical Simulation and Functional Prediction.** 1403-1413
- Aoxiang Fan, Jiayi Ma, Xin Tian, Xiaoguang Mei, Wei Lin: **Coherent Point Drift Revisited for Non-rigid Shape Matching and Registration.** 1414-1424
- Tianchen Zhao, Niansong Zhang, Xuefei Ning, He Wang, Li Yi, Yu Wang: **CodedVTR: Codebook-based Sparse Voxel Transformer with Geometric Guidance.** 1425-1434
- Rishabh Singh, Pranav Gupta, Pradeep Shenoy, Ravikiran Sarvadevabhatla : **FLOAT: Factorized Learning of Object Attributes for Improved Multi-object Multi-part Scene Parsing.** 1435-1445
- Hong-Xing Yu, Jiajun Wu, Li Yi: **Rotationally Equivariant 3D Object Detection.** 1446-1454
- Zhiqin Chen, Kangxue Yin, Sanja Fidler: **AUV-Net: Learning Aligned UV Maps for Texture Transfer and Synthesis.** 1455-1464
- Hongsuk Choi, Gyeongsik Moon, JoonKyu Park, Kyoung Mu Lee: **Learning to Estimate Robust 3D Human Mesh from In-the-Wild Crowded Scenes.** 1465-1474
- Georgios Pavlakos, Jitendra Malik, Angjoo Kanazawa: **Human Mesh Recovery from Multiple Shots.** 1475-1485
- JoonKyu Park, Yeonguk Oh, Gyeongsik Moon, Hongsuk Choi, Kyoung Mu Lee: **HandOccNet: Occlusion-Robust 3D Hand Mesh Estimation Network.** 1486-1495
- Thiendo Alldieck, Mihai Zanfir, Cristian Sminchisescu: **Photorealistic Monocular 3D Reconstruction of Humans Wearing Clothing.** 1496-1505
- Ayush Tewari, Mallikarjun B. R., Xingang Pan, Ohad Fried, Maneesh Agrawala , Christian Theobalt : **Disentangled3D: Learning a 3D Generative Model with Disentangled Geometry and Appearance from Monocular Images.** 1506-1515
- Keyu Wu, Yifan Ye, Lingchen Yang, Hongbo Fu, Kun Zhou, Youyi Zheng: **NeuralHDHair: Automatic High-fidelity Hair Modeling from a Single Image Using Implicit Neural Representations.** 1516-1525
- Shivam Duggal, Deepak Pathak: **Topologically-Aware Deformation Fields for Single-View 3D Reconstruction.** 1526-1536
- Rahul Dey, Vishnu Naresh Boddeti: **Generating Diverse 3D Reconstructions from a Single Occluded Face Image.** 1537-1547
- Daniel Rebain, Mark J. Matthews , Kwang Moo Yi, Dmitry Lagun, Andrea Tagliasacchi: **LOLNeRF: Learn from One Look.** 1548-1557
- Yida Wang, David Joseph Tan, Nassir Navab, Federico Tombari: **Learning Local Displacements for Point Cloud Completion.** 1558-1567
- Andra Petrovai, Sergiu Nedevschi : **Exploiting Pseudo Labels in a Self-Supervised Learning Framework for Improved Monocular Depth Estimation.** 1568-1578
- Yunpeng Zhang, Wenzhao Zheng, Zheng Zhu, Guan Huang, Dalong Du, Jie Zhou, Jiwen Lu: **Dimension Embeddings for Monocular 3D Object Detection.** 1579-1588

- Shengyi Qian, Linyi Jin, Chris Rockwell, Siyi Chen, David F. Fouhey:
Understanding 3D Object Articulation in Internet Videos. 1589-1599
- Vaishakh Patil, Christos Sakaridis , Alexander Liniger, Luc Van Gool:
P3Depth: Monocular Depth Estimation with a Piecewise Planarity Prior. 1600-1611
- Kehan Wang, Jia Zheng , Zihan Zhou:
Neural Face Identification in a 2D Wireframe Projection of a Manifold Object. 1612-1621
- Naiyu Gao, Fei He, Jian Jia, Yanhu Shan, Haoyang Zhang, Xin Zhao , Kaiqi Huang:
PanopticDepth: A Unified Framework for Depth-aware Panoptic Segmentation. 1622-1632
- Zimeng Zhao, Binghui Zuo, Wei Xie, Yangang Wang:
Stability-driven Contact Reconstruction From Monocular Color Images. 1633-1643
- Zhigang Jiang , Zhongzheng Xiang, Jinhua Xu, Ming Zhao:
LGT-Net: Indoor Panoramic Room Layout Estimation with Geometry-Aware Transformer Network. 1644-1653
- Tze Ho Elden Tse, Kwang In Kim, Ales Leonardis, Hyung Jin Chang :
Collaborative Learning for Hand and Object Reconstruction with Attention-guided Graph Convolution. 1654-1664
- Tak-Wai Hui:
RM-Depth: Unsupervised Learning of Recurrent Monocular Depth in Dynamic Scenes. 1665-1674
- Qing Lian, Botao Ye, Ruijia Xu, Weilong Yao, Tong Zhang:
Exploring Geometric Consistency for Monocular 3D Object Detection. 1675-1684
- Georgia Gkioxari, Nikhila Ravi, Justin Johnson:
Learning 3D Object Shape and Layout without 3D Supervision. 1685-1694
- Nikolay Patakin, Anna Vorontsova, Mikhail Artemyev, Anton Konushin:
Single-Stage 3D Geometry-Preserving Depth Estimation Model Training on Dataset Mixtures with Uncalibrated Stereo Data. 1695-1704
- Rawal Khirodkar, Shashank Tripathi, Kris Kitani:
Occluded Human Mesh Recovery. 1705-1715
- Junshu Tang, Zhijun Gong, Ran Yi, Yuan Xie, Lizhuang Ma:
LAKe-Net: Topology-Aware Point Cloud Completion by Localizing Aligned Keypoints. 1716-1725
- Wenbin Lin, Chengwei Zheng , Jun-Hai Yong, Feng Xu:
OcclusionFusion: Occlusion-aware Motion Estimation for Real-time Dynamic 3D Reconstruction. 1726-1735
- Yuhua Xu, Xiaoli Yang, Yushan Yu, Wei Jia, Zhaobi Chu, Yulan Guo:
Depth Estimation by Combining Binocular Stereo and Monocular Structured-Light. 1736-1745
- Mingtao Feng, Kendong Liu, Liang Zhang, Hongshan Yu, Yaonan Wang, Ajmal Mian :
Learning from Pixel-Level Noisy Label : A New Perspective for Light Field Saliency Detection. 1746-1756
- Wele Gedara Chaminda Bandara , Vishal M. Patel:
HyperTransformer: A Textural and Spectral Feature Fusion Transformer for Pan sharpening. 1757-1767
- Scott Workman, Muhammad Usman Rafique, Hunter Blanton, Nathan Jacobs :
Revisiting Near/Remote Sensing with Geospatial Attention. 1768-1777
- Gang Yang , Man Zhou, Keyu Yan, Aiping Liu, Xueyang Fu , Fan Wang:
Memory-augmented Deep Conditional Unfolding Network for Pan sharpening. 1778-1787
- Man Zhou, Keyu Yan, Jie Huang, Zihe Yang, Xueyang Fu , Feng Zhao:
Mutual Information-driven Pan-sharpening. 1788-1798
- Fengyu Yang, Chenyang Ma:
Sparse and Complete Latent Organization for Geospatial Semantic Segmentation. 1799-1808
- Dominik Muhle, Lukas Koestler, Nikolaus Demmel, Florian Bernard, Daniel Cremers :
The Probabilistic Normal Epipolar Constraint for Frame- To-Frame Rotation Optimization under Uncertain Feature Positions. 1809-1818
- Wentong Li, Yijie Chen, Kaixuan Hu, Jianke Zhu:
Oriented RepPoints for Aerial Object Detection. 1819-1828
- Christina Tsalicoglou, Thomas Rösgen:
Using 3D Topological Connectivity for Ghost Particle Reduction in Flow Reconstruction. 1829-1837
- Ngoc Long Nguyen, Jérémie Anger, Axel Davy, Pablo Arias, Gabriele Facciolo:
Self-Supervised Super-Resolution for Multi-Exposure Push-Frame Satellites. 1848-1858
- Xiaoguang Li, Qing Guo, Di Lin, Ping Li, Wei Feng, Song Wang:
MISF: Multi-level Interactive Siamese Filtering for High-Fidelity Image Inpainting. 1859-1868
- Si-Yuan Cao, Jianxin Hu, Ze-Hua Sheng, Hui-Liang Shen:
Iterative Deep Homography Estimation. 1869-1878
- Jingwen He, Wu Shi, Kai Chen, Lean Fu, Chao Dong:
GCFSR: a Generative and Controllable Face Super Resolution Method Without Facial and GAN Priors. 1879-1888
- Zhao Zhang, Huan Zheng, Richang Hong, Mingliang Xu, Shuicheng Yan, Meng Wang:
Deep Color Consistent Network for Low-Light Image Enhancement. 1889-1898

- Baisong Guo, Xiaoyun Zhang, Haoning Wu, Yu Wang, Ya Zhang , Yanfeng Wang:
LAR-SR: A Local Autoregressive Model for Image Super-Resolution. 1899-1908
- Bo Ji, Angela Yao:
Multi-Scale Memory-Based Video Deblurring. 1909-1918
- Jaewon Lee, Kyong Hwan Jin :
Local Texture Estimator for Implicit Representation Function. 1919-1928
- Qing Su, Shihao Ji:
ChiTransformer: Towards Reliable Stereo from Cues. 1929-1939
- Stefano Zorzi, Shabab Bazrafkan, Stefan Habenschuss, Friedrich Fraundorfer:
PolyWorld: Polygonal Building Extraction with Graph Neural Networks in Satellite Images. 1938-1947
- Jaihyun Koh, Jangho Lee, Sungroh Yoon:
BNUDC: A Two-Branched Deep Neural Network for Restoring Images from Under-Display Cameras. 1940-1949
- Metin Ersin Arican , Ozgur Kara, Gustav Bredell, Ender Konukoglu:
ISNAS-DIP: Image-Specific Neural Architecture Search for Deep Image Prior. 1950-1958
- Lingtong Kong, Boyuan Jiang, Donghao Luo, Wenqing Chu, Xiaoming Huang, Ying Tai, Chengjie Wang, Jie Yang:
IFRNet: Intermediate Feature Refine Network for Efficient Frame Interpolation. 1959-1968
- Riccardo de Lutio, Alexander Becker, Stefano D'Aronco, Stefania Russo, Jan D. Wegner, Konrad Schindler:
Learning Graph Regularisation for Guided Super-Resolution. 1969-1978
- Weixi Wang, Ji Li, Hui Ji:
Self-supervised Deep Image Restoration via Adaptive Stochastic Gradient Langevin Dynamics. 1979-1988
- Wenbo Zhao, Xianming Liu, Zhiwei Zhong, Junjun Jiang, Wei Gao , Ge Li, Xiangyang Ji:
Self-Supervised Arbitrary-Scale Point Clouds Upsampling via Implicit Neural Representation. 1989-1997
- Kwanyoung Kim, Taesung Kwon, Jong Chul Ye:
Noise Distribution Adaptive Self-Supervised Image Denoising using Tweedie Distribution and Score Matching. 1998-2006
- Xiang Chen, Jinshan Pan, Kui Jiang, Yufeng Li, Yufeng Huang, Caihua Kong, Longgang Dai, Zhentao Fan:
Unpaired Deep Image Deraining Using Dual Contrastive Learning. 2007-2016
- Zejin Wang, Jiazheng Liu, Guoqing Li, Hua Han:
Blind2Unblind: Self-Supervised Image Denoising with Visible Blind Spots. 2017-2026
- Yang Yang, Chaoyue Wang, Risheng Liu, Lin Zhang, Xiaojie Guo, Dacheng Tao:
Self-augmented Unpaired Image Dehazing via Density and Depth Decomposition. 2027-2036
- Zeyuan Chen, Yinbo Chen, Jingwen Liu, Xingqian Xu, Vedit Goel, Zhangyang Wang, Humphrey Shi , Xiaolong Wang:
VideoINR: Learning Video Implicit Neural Representation for Continuous Space-Time Super-Resolution. 2037-2047
- Ryuki Yamamoto, Hidekata Hontani, Akira Imakura, Tatsuya Yokota:
Fast Algorithm for Low-rank Tensor Completion in Delay-embedded Space. 2048-2056
- Cheng Zhang, Shaolin Su, Yu Zhu, Qingsen Yan, Jinqiu Sun, Yanning Zhang:
Exploring and Evaluating Image Restoration Potential in Dynamic Scenes. 2057-2066
- Pranjay Shyam, Kyung-Soo Kim, Kuk-Jin Yoon:
GIQE: Generic Image Quality Enhancement via Nth Order Iterative Degradation. 2067-2077
- Lai Jiang, Yifei Li, Shengxi Li, Mai Xu, Se Lei, Yichen Guo, Bo Huang:
Does text attract attention on e-commerce images: A novel saliency prediction dataset and method. 2078-2087
- Yi Zhang , Dasong Li, Ka Lung Law, Xiaogang Wang , Hongwei Qin, Hongsheng Li :
IDR: Self-Supervised Image Denoising via Iterative Data Refinement. 2088-2097
- Biwen Lei, Xiefan Guo, Hongyu Yang, Miaomiao Cui, Xuansong Xie, Di Huang:
ABPN: Adaptive Blend Pyramid Network for Real-Time Local Retouching of Ultra High-Resolution Photo. 2098-2107
- Salma Abdel Magid, Zudi Lin, Donglai Wei, Yulun Zhang , Jinjin Gu, Hanspeter Pfister :
Texture-based Error Analysis for Image Super-Resolution. 2108-2117
- Zongsheng Yue, Qian Zhao, Jianwen Xie, Lei Zhang, Deyu Meng, Kwan-Yee K. Wong:
Blind Image Super-resolution with Elaborate Degradation Modeling on Noise and Kernel. 2118-2128
- Hunsang Lee, Hyesong Choi, Kwanghoon Sohn, Dongbo Min:
KNN Local Attention for Image Restoration. 2129-2139
- Ruijun Gao, Qing Guo, Felix Juefei-Xu, Hongkai Yu, Huazhu Fu , Wei Feng, Yang Liu, Song Wang:
Can You Spot the Chameleon? Adversarially Camouflaging Images from Co-Salient Object Detection. 2140-2149
- Youwei Pang, Xiaoqi Zhao, Tian-Zhu Xiang, Lihe Zhang, Huchuan Lu:
Zoom In and Out: A Mixed-scale Triplet Network for Camouflaged Object Detection. 2150-2160
- Jennifer J. Sun, Serim Ryou, Roni H. Goldshmid, Brandon Weissbourd, John O. Dabiri, David J. Anderson, Ann Kennedy, Yisong Yue, Pietro Perona:
Self-Supervised Keypoint Discovery in Behavioral Videos. 2161-2170

- Weizhe Liu, Bugra Tekin, Huseyin Coskun, Vibhav Vineet, Pascal Fua, Marc Pollefeys [\(D\)](#): **Learning to Align Sequential Actions in the Wild.** 2171-2181
- Soma Nonaka, Shohei Nobuhara, Ko Nishino: **Dynamic 3D Gaze from Afar: Deep Gaze Estimation from Temporal Eye-Head-Body Coordination.** 2182-2191
- Danyang Tu, Xiongkuo Min, Huiyu Duan, Guodong Guo, Guangtao Zhai, Wei Shen: **End-to-End Human-Gaze-Target Detection with Transformers.** 2192-2200
- Albert Tseng, Jennifer J. Sun, Yisong Yue: **Automatic Synthesis of Diverse Weak Supervision Sources for Behavior Analysis.** 2201-2210
- Mihee Lee, Samuel S. Sohn, Seonghyeon Moon, Sejong Yoon, Mubbashir Kapadia, Vladimir Pavlovic: **MUSE-VAE: Multi-Scale VAE for Environment-Aware Long Term Trajectory Prediction.** 2211-2220
- Lihuan Li, Maurice Pagnucco, Yang Song: **Graph-based Spatial Transformer with Memory Replay for Multi-future Pedestrian Trajectory Prediction.** 2221-2231
- Ke Guo, Wenxi Liu, Jia Pan: **End-to-End Trajectory Distribution Prediction Based on Occupancy Grid Maps.** 2232-2241
- Hongchen Luo, Wei Zhai, Jing Zhang [\(D\)](#), Yang Cao, Dacheng Tao: **Learning Affordance Grounding from Exocentric Images.** 2242-2251
- Jaebong Jeong, Janghun Jo, Sunghyun Cho, Jaesik Park: **3D Scene Painting via Semantic Image Synthesis.** 2252-2262
- Jun Jia, Zhongpai Gao [\(D\)](#), Dandan Zhu, Xiongkuo Min, Guangtao Zhai, Xiaokang Yang: **Learning Invisible Markers for Hidden Codes in Offline-to-online Photography.** 2263-2272
- Lingteng Qiu, Zhangyang Xiong, Xuhao Wang, Kenkun Liu, Yihan Li, Guanying Chen [\(D\)](#), Xiaoguang Han, Shuguang Cui [\(D\)](#): **ETHSeg: An Amodel Instance Segmentation Network and a Real-world Dataset for X-Ray Waste Inspection.** 2273-2282
- Ayan Kumar Bhunia, Viswanatha Reddy Gajjala, Subhadeep Koley [\(D\)](#), Rohit Kundu, Aneeshan Sain [\(D\)](#), Tao Xiang, Yi-Zhe Song: **Doodle It Yourself: Class Incremental Learning by Drawing a Few Sketches.** 2283-2292
- Xiya Liu, Ziping Ma [\(D\)](#), Junxing Ma, Jian Zhang [\(D\)](#), Gerald Schaefer, Hui Fang: **Image Disentanglement Autoencoder for Steganography without Embedding.** 2293-2302
- Banghuai Li: **Adaptive Hierarchical Representation Learning for Long-Tailed Object Detection.** 2303-2312
- YuanFu Yang [\(D\)](#), Min Sun: **Semiconductor Defect Detection by Hybrid Classical-Quantum Deep Learning.** 2313-2322
- Yun He, Xinlin Ren, Danhang Tang, Yinda Zhang, Xiangyang Xue, Yanwei Fu [\(D\)](#): **Density-preserving Deep Point Cloud Compression.** 2323-2332
- Zheheng Jiang, Hossein Rahmani, Plamen Angelov, Sue Black, Bryan M. Williams: **Graph-context Attention Networks for Size-varied Deep Graph Matching.** 2333-2342
- Jeya Maria Jose Valanarasu, Rajeev Yasarla, Vishal M. Patel: **TransWeather: Transformer-based Restoration of Images Degraded by Adverse Weather Conditions.** 2343-2353
- Junke Wang, Zuxuan Wu, Jingjing Chen [\(D\)](#), Xintong Han, Abhinav Shrivastava, Ser-Nam Lim, Yu-Gang Jiang: **ObjectFormer for Image Manipulation Detection and Localization.** 2354-2363
- Qichen Fu, Xingyu Liu, Kris M. Kitani: **Sequential Voting with Relational Box Fields for Active Object Detection.** 2364-2373
- Fanjie Kong, Ricardo Henao: **Efficient Classification of Very Large Images with Tiny Objects.** 2374-2384
- Pinaki Nath Chowdhury, Ayan Kumar Bhunia, Viswanatha Reddy Gajjala, Aneeshan Sain [\(D\)](#), Tao Xiang, Yi-Zhe Song [\(D\)](#): **Partially Does It: Towards Scene-Level FG-SBIR with Partial Input.** 2385-2395
- Ming-Fang Chang, Yipu Zhao, Rajvi Shah, Jakob J. Engel, Michael Kaess [\(D\)](#), Simon Lucey [\(D\)](#): **Long-term Visual Map Sparsification with Heterogeneous GNN.** 2396-2405
- Ruize Han, Yiyang Gan, Jiacheng Li, Feifan Wang, Wei Feng, Song Wang: **Connecting the Complementary-view Videos: Joint Camera Identification and Subject Association.** 2406-2415
- Gwanghyun Kim, Taesung Kwon, Jong Chul Ye: **DiffusionCLIP: Text-Guided Diffusion Models for Robust Image Manipulation.** 2416-2425
- Yizhi Wang, Guo Pu, Wenhan Luo [\(D\)](#), Yixin Wang, Pengfei Xiong, Hongwen Kang, Zhouhui Lian: **Aesthetic Text Logo Synthesis via Content-aware Layout Inferring.** 2426-2435
- Gengyun Jia, Huabo Huang, Chaoyou Fu, Ran He: **Rethinking Image Cropping: Exploring Diverse Compositions from Global Views.** 2436-2445
- Jiakai Wang, Zixin Yin, Pengfei Hu, Aishan Liu, Renshuai Tao, Haotong Qin [\(D\)](#), Xianglong Liu, Dacheng Tao: **Defensive Patches for Robust Recognition in the Physical World.** 2446-2455

- Xun Jiang, Xing Xu, Jingran Zhang, Fumin Shen, Zuo Cao, Heng Tao Shen:
Semi-supervised Video Paragraph Grounding with Contrastive Encoder. 2456-2465
- Hao Ni, Jingkuan Song, Xiaopeng Luo, Feng Zheng, Wen Li, Heng Tao Shen:
Meta Distribution Alignment for Generalizable Person Re-Identification. 2477-2486
- Zhenpei Yang, Zhile Ren, Miguel Ángel Bautista, Zaiwei Zhang, Qi Shan, Qixing Huang:
FvOR: Robust Joint Shape and Pose Optimization for Few-view Object Reconstruction. 2487-2497
- Charig Yang, Weidi Xie, Andrew Zisserman:
It's About Time: Analog Clock Reading in the Wild. 2498-2507
- Samrudhdhi B. Rangrej, Chetan L. Srinidhi, James J. Clark:
Consistency driven Sequential Transformers Attention Model for Partially Observable Scenes. 2508-2517
- Ran Xu , Fangzhou Mu, Jayoung Lee, Preeti Mukherjee, Somali Chaterji, Saurabh Bagchi, Yin Li:
Smartadapt: Multi-branch Object Detection Framework for Videos on Mobiles. 2518-2528
- Han Joo Chae, Seunghwan Lee, Hyewon Son, Seungyeob Han, Taebin Lim:
Generating 3D Bio-Printable Patches Using Wound Segmentation and Reconstruction to Treat Diabetic Foot Ulcers. 2529-2539
- Hanjiang Hu, Zuxin Liu , Sharad Chitlangia, Akhil Agnihotri, Ding Zhao :
Investigating the Impact of Multi-LiDAR Placement on Object Detection for Autonomous Driving. 2540-2549
- Qihang Yu, Huiyu Wang, Dahun Kim, Siyuan Qiao, Maxwell D. Collins, Yukun Zhu, Hartwig Adam, Alan L. Yuille, Liang-Chieh Chen:
CMT-DeepLab: Clustering Mask Transformers for Panoptic Segmentation. 2550-2560
- Tsung-Wei Ke, Jyh-Jing Hwang, Yunhui Guo, Xudong Wang, Stella X. Yu:
Unsupervised Hierarchical Semantic Segmentation with Multiview Cosegmentation and Clustering Transformers. 2561-2571
- Tianfei Zhou, Wenguan Wang , Ender Konukoglu, Luc Van Gool:
Rethinking Semantic Segmentation: A Prototype View. 2572-2583
- Duo Peng, Yinjie Lei, Munawar Hayat, Yulan Guo, Wen Li:
Semantic-Aware Domain Generalized Segmentation. 2584-2595
- Sheng Liu, Kangning Liu, Weicheng Zhu, Yiqiu Shen , Carlos Fernandez-Granda:
Adaptive Early-Learning Correction for Segmentation from Noisy Annotations. 2596-2606
- Bowen Cheng, Omkar Parkhi, Alexander Kirillov:
Pointly-Supervised Instance Segmentation. 2607-2616
- Colin Gruber, Cyril Jazra, Wenjie Luo, Liangyan Gui, Alexander G. Schwing:
Joint Forecasting of Panoptic Segmentations with Difference Attention. 2617-2626
- Zheng Lin, Zheng-Peng Duan, Zhao Zhang , Chun-Le Guo, Ming-Ming Cheng :
FocusCut: Diving into a Focus View in Interactive Segmentation. 2627-2636
- Yanan Sun, Chi-Keung Tang, Yu-Wing Tai :
Human Instance Matting via Mutual Guidance and Multi-Instance Refinement. 2637-2646
- Vickie Ye, Zhengqi Li, Richard Tucker, Angjoo Kanazawa, Noah Snavely:
Deformable Sprites for Unsupervised Video Decomposition. 2647-2656
- Wonhui Park, Dongkwon Jin, Chang-Su Kim:
Eigencontours: Novel Contour Descriptors Based on Low-Rank Approximation. 2657-2665
- Weixiao Liu , Yuwei Wu, Sipu Ruan, Gregory S. Chirikjian:
Robust and Accurate Superquadric Recovery: a Probabilistic Approach. 2666-2675
- Morteza Rezanejad, Mohammad Khodadad, Hamidreza Mahyar, Herve Lombaert, Michael Gruninger, Dirk B. Walther, Kaleem Siddiqi:
Medial Spectral Coordinates for 3D Shape Analysis. 2676-2686
- Ozan Unal , Dengxin Dai, Luc Van Gool:
Scribble-Supervised LiDAR Semantic Segmentation. 2687-2697
- Thang Vu, Kookhoi Kim, Tung Minh Luu, Thanh Xuan Nguyen, Chang D. Yoo:
SoftGroup for 3D Instance Segmentation on Point Clouds. 2698-2707
- Vasileios Choutas, Lea Müller, Chun-Hao P. Huang, Siyu Tang , Dimitrios Tzionas, Michael J. Black:
Accurate 3D Body Shape Regression using Metric and Semantic Attributes. 2708-2718
- Yukang Cao, Guanying Chen , Kai Han, Wenqi Yang, Kwan-Yee K. Wong:
JIFF: Jointly-aligned Implicit Face Function for High Quality Single View Clothed Human Reconstruction. 2719-2729
- Jathushan Rajasegaran, Georgios Pavlakos, Angjoo Kanazawa, Jitendra Malik:
Tracking People by Predicting 3D Appearance, Location and Pose. 2730-2739
- Lixin Yang, Kailin Li , Xinyu Zhan, Jun Lv, Wenqiang Xu, Jiefeng Li, Cewu Lu:
ArtiBoost: Boosting Articulated 3D Hand-Object Pose Estimation via Online Exploration and Synthesis. 2740-2750
- Mengcheng Li, Liang An , Hongwen Zhang, Lianpeng Wu, Feng Chen, Tao Yu, Yebin Liu:
Interacting Attention Graph for Single Image Two-Hand Reconstruction. 2751-2760

- Stylianos Ploumpis, Stylianos Moschoglou, Vasileios Triantafyllou, Stefanos Zafeiriou: **3D human tongue reconstruction from single "in-the-wild" images.** 2761-2770
- Hansheng Chen, Pichao Wang, Fan Wang, Wei Tian, Lu Xiong, Hao Li: **EPro-PnP: Generalized End-to-End Probabilistic Perspective-n-Points for Monocular Object Pose Estimation.** 2771-2780
- Zhuoling Li, Zhan Qu, Yang Zhou, Jianzhuang Liu, Haoqian Wang, Lihui Jiang: **Diversity Matters: Fully Exploiting Depth Clues for Reliable Monocular 3D Object Detection.** 2781-2790
- Yuyan Li, Yuliang Guo, Zhixin Yan, Xinyu Huang, Ye Duan, Liu Ren: **OmniFusion: 360 Monocular Depth Estimation via Geometry-Aware Fusion.** 2791-2800
- Amanpreet Walia, Stefanie Walz, Mario Bijelic, Fahim Mannan, Frank D. Julca-Aguilar, Michael S. Langer, Werner Ritter, Felix Heide: **Gated2Gated: Self-Supervised Depth Estimation from Gated Images.** 2801-2811
- Rui Zhu, Zhengqin Li, Janarbek Matai, Fatih Porikli, Manmohan Chandraker: **IRISformer: Dense Vision Transformers for Single-Image Inverse Rendering in Indoor Scenes.** 2812-2821
- Tien Do, Khiem Vuong, Hyun Soo Park: **Egocentric Scene Understanding via Multimodal Spatial Rectifier.** 2822-2831
- Gwangbin Bae, Ignas Budvytis, Roberto Cipolla: **Multi-View Depth Estimation by Fusing Single-View Depth Probability with Multi-View Geometry.** 2832-2841
- Ilya Chugunov, Yuxuan Zhang, Zhihao Xia, Xuaner Zhang, Jiawen Chen, Felix Heide: **The Implicit Values of A Good Hand Shake: Handheld Multi-Frame Neural Depth Refinement.** 2842-2852
- Gengshan Yang, Minh Vo, Natalia Neverova, Deva Ramanan , Andrea Vedaldi, Hanbyul Joo: **BANMo: Building Animatable 3D Neural Models from Many Casual Videos.** 2853-2863
- Kanchana Ranasinghe, Muzammal Naseer, Salman Khan, Fahad Shahbaz Khan, Michael S. Ryoo: **Self-supervised Video Transformer.** 2864-2874
- Shusheng Yang, Xinggang Wang, Yu Li , Yuxin Fang, Jiemin Fang, Wenyu Liu, Xun Zhao, Ying Shan: **Temporally Efficient Vision Transformer for Video Instance Segmentation.** 2875-2885
- Su Ho Han, Sukjun Hwang, Seoung Wug Oh, Yeonchool Park, Hyunwoo Kim, Min-Jung Kim, Seon Joo Kim: **VISOLO: Grid-Based Space-Time Aggregation for Efficient Online Video Instance Segmentation.** 2886-2895
- Tengda Han, Weidi Xie, Andrew Zisserman: **Temporal Alignment Networks for Long-term Video.** 2896-2906
- Shyamal Buch, Cristóbal Eyzaguirre, Adrien Gaidon, Jiajun Wu, Li Fei-Fei, Juan Carlos Niebles: **Revisiting the "Video" in Video-Language Understanding.** 2907-2917
- Yicong Li, Xiang Wang, Junbin Xiao, Wei Ji, Tat-Seng Chua: **Invariant Grounding for Video Question Answering.** 2918-2927
- He Zhao, Isma Hadji, Nikita Dvornik, Konstantinos G. Derpanis, Richard P. Wildes, Allan D. Jepson: **P³IV: Probabilistic Procedure Planning from Instructional Videos with Weak Supervision.** 2928-2938
- Jinglin Xu, Yongming Rao, Xumin Yu, Guangyi Chen , Jie Zhou, Jiwen Lu: **FineDiving: A Fine-grained Dataset for Procedure-aware Action Quality Assessment.** 2939-2948
- Yinghao Xu, Fangyun Wei, Xiao Sun, Ceyuan Yang, Yujun Shen, Bo Dai, Bolei Zhou, Stephen Lin: **Cross-Model Pseudo-Labeling for Semi-Supervised Action Recognition.** 2949-2958
- Haodong Duan, Yue Zhao, Kai Chen, Dahua Lin, Bo Dai: **Revisiting Skeleton-based Action Recognition.** 2959-2968
- Wentao Bao , Qi Yu , Yu Kong: **OpenTAL: Towards Open Set Temporal Action Localization.** 2969-2979
- Mingfei Han, David Junhao Zhang, Yali Wang, Rui Yan, Lina Yao, Xiaojun Chang , Yu Qiao: **Dual-AI: Dual-path Actor Interaction Learning for Group Activity Recognition.** 2980-2989
- Haodong Duan, Nanxuan Zhao, Kai Chen, Dahua Lin: **TransRank: Self-supervised Video Representation Learning via Ranking-based Transformation Recognition.** 2990-3000
- Basile Van Hoorick, Purva Tendulkar, Dídac Surís, Dennis Park, Simon Stent, Carl Vondrick: **Revealing Occlusions with 4D Neural Fields.** 3001-3011
- Ali Athar, Jonathon Luiten, Alexander Hermans, Deva Ramanan , Bastian Leibe : **HODOR: High-level Object Descriptors for Object Re-segmentation in Video Learned from Static Images.** 3012-3021
- Juncheng Li, Junlin Xie, Long Qian, Linchao Zhu, Siliang Tang, Fei Wu, Yi Yang, Yueling Zhuang, Xin Eric Wang: **Compositional Temporal Grounding with Structured Variational Cross-Graph Correspondence Learning.** 3022-3031
- Ye Liu, Siyuan Li, Yang Wu, Chang Wen Chen, Ying Shan, Xiaohu Qie: **UMT: Unified Multi-modal Transformers for Joint Video Moment Retrieval and Highlight Detection.** 3032-3041
- Dayoung Gong, Joonseok Lee, Manjin Kim , Seong Jong Ha, Minsu Cho: **Future Transformer for Long-term Action Anticipation.** 3042-3051



- Zhaofan Qiu, Ting Yao, Chong-Wah Ngo, Tao Mei:
MLP-3D: A MLP-like 3D Architecture with Grouped Time Mixing. 3052-3062
- Fanyue Wei, Biao Wang, Tiezheng Ge, Yuning Jiang, Wen Li, Lixin Duan:
Learning Pixel-Level Distinctions for Video Highlight Detection. 3063-3072
- Tao Han, Lei Bai , Junyu Gao, Qi Wang, Wanli Ouyang:
DR.VIC: Decomposition and Reasoning for Video Individual Counting. 3073-3082
- Yi Zhou, Hui Zhang, Hana Lee, Shuyang Sun, Pingjun Li, Yangguang Zhu, ByungIn Yoo, Xiaojuan Qi, Jae-Joon Han:
Slot-VPS: Object-centric Representation Learning for Video Panoptic Segmentation. 3083-3093
- Kailai Zhou, Yibo Wang, Tao Lv, Yunqian Li, Linsen Chen, Qiu Shen, Xun Cao:
Explore Spatio-temporal Aggregation for Insubstantial Object Detection: Benchmark Dataset and Baseline. 3094-3105
- Xiao Lu , Yihong Cao, Sheng Liu, Chengjiang Long, Zipei Chen, Xuanyu Zhou, Yimin Yang, Chunxia Xiao:
Video Shadow Detection via Spatio-Temporal Interpolation Consistency Training. 3106-3115
- Guolei Sun, Yun Liu, Henghui Ding , Thomas Probst, Luc Van Gool:
Coarse-to-Fine Feature Mining for Video Semantic Segmentation. 3116-3127
- Zhaoyang Zeng, Yongsheng Luo, Zhenhua Liu, Fengyun Rao, Dian Li, Weidong Guo, Zhen Wen:
Tencent-MVSE: A Large-Scale Benchmark Dataset for Multi-Modal Video Similarity Evaluation. 3128-3137
- Roei Herzig, Elad Ben-Avraham, Karttikeya Mangalam, Amir Bar, Gal Chechik, Anna Rohrbach, Trevor Darrell, Amir Globerson:
Object-Region Video Transformers. 3138-3149
- Le Yang, Junwei Han, Dingwen Zhang:
Colar: Effective and Efficient Online Action Detection by Consulting Exemplars. 3150-3159
- Zhangyang Gao, Cheng Tan, Lirong Wu, Stan Z. Li:
SimVP: Simpler yet Better Video Prediction. 3160-3170
- Jisoo Jeong, Jamie Menjay Lin, Fatih Porikli, Nojun Kwak:
Imposing Consistency for Optical Flow Estimation. 3171-3181
- Fuchen Long, Zhaofan Qiu, Yingwei Pan, Ting Yao, Jiebo Luo , Tao Mei:
Stand-Alone Inter-Frame Attention in Video Models. 3182-3191
- Ze Liu, Jia Ning, Yue Cao, Yixuan Wei, Zheng Zhang, Stephen Lin, Han Hu:
Video Swin Transformer. 3192-3201
- Hitesh Sapkota, Qi Yu :
Bayesian Nonparametric Submodular Video Partition for Robust Anomaly Detection. 3202-3211
- Angchi Xu, Ling-An Zeng , Wei-Shi Zheng:
Likert Scoring with Grade Decoupling for Long-term Action Assessment. 3222-3231
- Yang Jin, Linchao Zhu, Yadong Mu:
Complex Video Action Reasoning via Learnable Markov Logic Network. 3232-3241
- Junfei Xiao, Longlong Jing, Lin Zhang, Ju He, Qi She, Zongwei Zhou, Alan L. Yuille, Yingwei Li:
Learning from Temporal Gradient for Semi-supervised Action Recognition. 3242-3252
- Jiafan Zhuang, Zilei Wang, Yuan Gao:
Semi-Supervised Video Semantic Segmentation with Inter-Frame Feature Reconstruction. 3253-3261
- Linjiang Huang, Liang Wang, Hongsheng Li :
Weakly Supervised Temporal Action Localization via Representative Snippet Knowledge Propagation. 3262-3271
- Shaowei Liu, Subarna Tripathi, Somdeb Majumdar, Xiaolong Wang:
Joint Hand Motion and Interaction Hotspots Prediction from Egocentric Videos. 3272-3282
- Mohit Goyal, Sahil Modi, Rishabh Goyal, Saurabh Gupta:
Human Hands as Probes for Interactive Object Understanding. 3283-3293
- Dan Liu, Libo Zhang, Yanjun Wu:
LD-ConGR: A Large RGB-D Video Dataset for Long-Distance Continuous Gesture Recognition. 3294-3302
- Alex Jinpeng Wang, Yixiao Ge, Guanyu Cai, Rui Yan, Xudong Lin, Ying Shan, Xiaohu Qie, Mike Zheng Shou:
Object-aware Video-language Pre-training for Retrieval. 3303-3312
- Zexing Du, Xue Wang , Guoqing Zhou, Qing Wang:
Fast and Unsupervised Action Boundary Detection for Action Segmentation. 3313-3322
- Shen Yan, Xuehan Xiong, Anurag Arnab, Zhichao Lu, Mi Zhang, Chen Sun, Cordelia Schmid:
Multiview Transformers for Video Recognition. 3323-3333
- Yuhang Shen , Ehsan Elhamifar:
Semi-Weakly-Supervised Learning of Complex Actions from Instructional Task Videos. 3334-3344
- Jiaqi Tang, Zhaoyang Liu, Chen Qian, Wayne Wu, Limin Wang:
Progressive Attention on Multi-Level Dense Difference Maps for Generic Event Boundary Detection. 3345-3354
- Daniel Geng, Max Hamilton , Andrew Owens:
Comparing Correspondences: Video Prediction with Correspondence-wise Losses. 3355-3366

- Seung Hyun Lee, Wonseok Roh, Wonmin Byeon, Sang Ho Yoon, Chanyoung Kim, Jinkyu Kim, Sangpil Kim [\(ID\)](#):
Sound-Guided Semantic Image Manipulation. 3367-3376
- Borong Liang [\(ID\)](#), Yan Pan, Zhizhi Guo, Hang Zhou, Zhibin Hong, Xiaoguang Han, Junyu Han, Jingtuo Liu, Errui Ding, Jingdong Wang:
Expressive Talking Head Generation with Granular Audio-Visual Control. 3377-3386
- Fa-Ting Hong, Longhao Zhang, Li Shen, Dan Xu:
Depth-Aware Generative Adversarial Network for Talking Head Video Generation. 3387-3396
- Jae Shin Yoon, Duygu Ceylan, Tuanfeng Y. Wang, Jingwan Lu, Jimei Yang, Zhixin Shu, Hyun Soo Park:
Learning Motion-Dependent Appearance for High-Fidelity Rendering of Dynamic Humans from a Single Camera. 3397-3407
- Yang Zhou, Jimei Yang, Dingzeyu Li, Jun Saito, Deepali Aneja, Evangelos Kalogerakis:
Audio-driven Neural Gesture Reenactment with Video Motion Graphs. 3408-3418
- Junfeng Lyu, Zhibo Wang, Feng Xu:
Portrait Eyeglasses and Shadow Removal by Leveraging 3D Synthetic Data. 3419-3429
- Ruili Feng, Cheng Ma, Chengji Shen, Xin Gao, Zhenjiang Liu, Xiaobo Li, Kairi Ou, Deli Zhao, Zheng-Jun Zha:
Weakly Supervised High-Fidelity Clothing Model Generation. 3430-3439
- You Xie, Huiqi Mao, Angela Yao, Nils Thuerey:
TemporalUV: Capturing Loose Clothing with Temporally Coherent UV Coordinates. 3440-3449
- Han Yang, Xinrui Yu, Ziwei Liu:
Full-Range Virtual Try-On with Recurrent Tri-Level Transform. 3450-3459
- Sen He, Yi-Zhe Song, Tao Xiang:
Style-Based Global Appearance Flow for Virtual Try-On. 3460-3469
- Xin Dong, Fuwei Zhao, Zhenyu Xie, Xijin Zhang, Daniel K. Du, Min Zheng, Xiang Long, Xiaodan Liang, Jianchao Yang:
Dressing in the Wild by Watching Dance Videos. 3470-3479
- Jinwoo Kim, Heeseok Oh, Seongjean Kim [\(ID\)](#), Hoseok Tong, Sanghoon Lee:
A Brand New Dance Partner: Music-Conditioned Pluralistic Dancing Controlled by Multiple Dance Genres. 3480-3490
- Yifang Men, Yuan Yao, Miaomiao Cui, Zhouhui Lian, Xuansong Xie, Xian-Sheng Hua:
Unpaired Cartoon Image Synthesis via Gated Cycle Mapping. 3491-3500
- Jingjing Ren, Qingqing Zheng, Yuanyuan Zhao, Xuemiao Xu, Chen Li:
DLFormer: Discrete Latent Transformer for Video Inpainting. 3501-3510
- Duolikun Danier, Fan Zhang, David Bull:
ST-MFNet: A Spatio-Temporal Multi-Flow Network for Frame Interpolation. 3511-3521
- Liying Lu, Ruizheng Wu, Huaijia Lin, Jiangbo Lu [\(ID\)](#), Jiaya Jia:
Video Frame Interpolation with Transformer. 3522-3532
- Dawit Mureja Argaw, In So Kweon:
Long-term Video Frame Interpolation via Feature Propagation. 3533-3542
- Ping Hu, Simon Niklaus, Stan Sclaroff, Kate Saenko:
Many-to-many Splatting for Efficient Video Frame Interpolation. 3543-3552
- Xuanchi Ren [\(ID\)](#), Xiaolong Wang:
Look Outside the Room: Synthesizing A Consistent Long-Term 3D Scene Video from A Single Image. 3553-3563
- Mengshun Hu, Kui Jiang, Liang Liao, Jing Xiao, Junjun Jiang, Zheng Wang:
Spatial-Temporal Space Hand-in-Hand: Spatial-Temporal Video Super-Resolution via Cycle-Projected Mutual Learning. 3564-3573
- Willi Menapace, Stéphane Lathuilière, Aliaksandr Siarohin, Christian Theobalt [\(ID\)](#), Sergey Tulyakov, Vladislav Golyanik, Elisa Ricci [\(ID\)](#):
Playable Environments: Video Manipulation in Space and Time. 3574-3583
- Lin Zhu [\(ID\)](#), Xiao Wang [\(ID\)](#), Yi Chang, Jianing Li, Tiejun Huang, Yonghong Tian:
Event-based Video Reconstruction via Potential-assisted Spiking Neural Network. 3584-3594
- Wei Yu, Wenxin Chen, Songheng Yin, Steve Easterbrook, Animesh Garg:
Modular Action Concept Grounding in Semantic Video Prediction. 3595-3604
- Ligong Han, Jian Ren, Hsin-Ying Lee, Francesco Barbieri, Kyle Olszewski, Shervin Minaee, Dimitris N. Metaxas, Sergey Tulyakov:
Show Me What and Tell Me How: Video Synthesis via Multimodal Conditioning. 3605-3615
- Ivan Skorokhodov [\(ID\)](#), Sergey Tulyakov, Mohamed Elhoseiny [\(ID\)](#):
StyleGAN-V: A Continuous Video Generator with the Price, Image Quality and Perks of StyleGAN2. 3616-3626
- Jiale Tao, Biao Wang, Borun Xu, Tiezheng Ge, Yuning Jiang, Wen Li, Lixin Duan:
Structure-Aware Motion Transfer with Deformable Anchor Model. 3627-3636
- Yoav Shalev, Lior Wolf:
Image Animation with Perturbed Masks. 3637-3646

- Jian Zhao, Hui Zhang:
Thin-Plate Spline Motion Model for Image Animation. 3647-3656
- Aniruddha Mahapatra, Kuldeep Kulkarni:
Controllable Animation of Fluid Elements in Still Images. 3657-3666
- Atsuhiro Noguchi, Umar Iqbal, Jonathan Tremblay, Tatsuya Harada, Orazio Gallo:
Watch It Move: Unsupervised Discovery of 3D Joints for Re-Posing of Articulated Objects. 3667-3677
- Peng Du, Jifeng Ning, Jiguang Cui, Shaoli Huang, Xinchao Wang , Jiaxin Wang:
Geometric Structure Preserving Warp for Natural Image Stitching. 3678-3686
- Pei Chen, Yangkang Zhang, Zejian Li , Lingyun Sun:
Few-Shot Incremental Learning for Label-to-Image Translation. 3687-3697
- Haiwei Chen , Jiayi Liu, Weikai Chen, Shichen Liu, Yajie Zhao:
Exemplar-based Pattern Synthesis with Implicit Periodic Field Network. 3698-3707
- Xianling Zhang, Nathan Tseng, Ameerah Syed, Rohan Bhasin, Nikita Jaipuria:
SIMBAR: Single Image-Based Scene Relighting For Effective Data Augmentation For Automated Driving Vision Tasks. 3708-3718
- Jiahao Yu, Li Chen, Mingrui Zhang, Mading Li:
SoftCollage: A Differentiable Probabilistic Tree Generator for Image Collage. 3719-3728
- Ning Kang, Shanzhao Qiu, Shifeng Zhang, Zhenguo Li, Shutao Xia:
PILC: Practical Image Lossless Compression with an End-to-end GPU Oriented Neural Framework. 3729-3738
- Klaus Greff, Francois Belletti, Lucas Beyer, Carl Doersch, Yilun Du, Daniel Duckworth, David J. Fleet, Dan Gnanapragasam, Florian Golemo, Charles Herrmann, Thomas Kipf, Abhijit Kundu, Dmitry Lagun, Issam H. Laradji, Hsueh-Ti Derek Liu, Henning Meyer, Yishu Miao, Derek Nowrouzezahrai, A. Cengiz Öztireli, Etienne Pot, Noha Radwan, Daniel Rebain, Sara Sabour, Mehdi S. M. Sajjadi, Matan Sela, Vincent Sitzmann, Austin Stone, Deqing Sun, Suhani Vora, Ziyu Wang, Tianhao Wu, Kwang Moo Yi, Fangcheng Zhong, Andrea Tagliasacchi:
Kubric: A scalable dataset generator. 3739-3751
- Manuel Rey-Area, Mingze Yuan, Christian Richardt:
360MonoDepth: High-Resolution 360° Monocular Depth Estimation. 3752-3762
- Kalyan Vasudev Alwala, Abhinav Gupta, Shubham Tulsiani:
Pretrain, Self-train, Distill: A simple recipe for Supersizing 3D Reconstruction. 3763-3772
- Tuo Cao, Fei Luo, Yanping Fu, Wenxiao Zhang, Shengjie Zheng, Chunxia Xiao:
DGCN: A Depth-Guided Edge Convolutional Network for End-to-End 6D Pose Estimation. 3773-3782
- Zequn Qin, Xi Li:
MonoGround: Detecting Monocular 3D Objects from the Ground. 3783-3792
- Xin Wen, Junsheng Zhou, Yu-Shen Liu, Hua Su, Zhen Dong, Zhizhong Han:
3D Shape Reconstruction from 2D Images with Disentangled Attribute Flow. 3793-3803
- Cho-Ying Wu, Jialiang Wang, Michael Hall, Ulrich Neumann, Shuochen Su:
Toward Practical Monocular Indoor Depth Estimation. 3804-3814
- Georgy Ponomatkin, Yann Labb  , Bryan C. Russell, Mathieu Aubry, Josef Sivic:
Focal Length and Object Pose Estimation via Render and Compare. 3815-3824
- Can Wang , Menglei Chai, Mingming He, Dongdong Chen, Jing Liao :
CLIP-NeRF: Text-and-Image Driven Manipulation of Neural Radiance Fields. 3825-3834
- Heming Zhu, Lingteng Qiu, Yuda Qiu, Xiaoguang Han:
Registering Explicit to Implicit: Towards High-Fidelity Garment mesh Reconstruction from Single Images. 3835-3844
- Soo Ye Kim, Jianming Zhang, Simon Niklaus, Yifei Fan, Simon Chen, Zhe Lin, Munchurl Kim:
Layered Depth Refinement with Mask Guidance. 3845-3855
- Jiacheng Chen, Yiming Qian, Yasutaka Furukawa:
HEAT: Holistic Edge Attention Transformer for Structured Reconstruction. 3856-3865
- Nadine R  egg, Silvia Zuffi, Konrad Schindler, Michael J. Black:
BARC: Learning to Regress 3D Dog Shape from Images by Exploiting Breed Information. 3866-3874
- Peixuan Li, Jieyu Jin:
Time3D: End-to-End Joint Monocular 3D Object Detection and Tracking for Autonomous Driving. 3875-3884
- Yufei Ye, Abhinav Gupta, Shubham Tulsiani:
What's in your hands? 3D Reconstruction of Generic Objects in Hands. 3885-3895
- Qianqian Wang, Zhengqi Li, David Salesin, Noah Snavely, Brian Curless, Janne Kontkanen:
3D Moments from Near-Duplicate Photos. 3896-3905
- Weihao Yuan, Xiaodong Gu, Zuozhuo Dai, Siyu Zhu, Ping Tan:
Neural Window Fully-connected CRFs for Monocular Depth Estimation. 3906-3915
- J  r  me Revaud, Vincent Leroy, Philippe Weinzaepfel, Boris Chidlovskii:
PUMP: Pyramidal and Uniqueness Matching Priors for Unsupervised Learning of Local Descriptors. 3916-3926

- Yannick Verdié, Jifei Song, Barnabé Mas, Benjamin Busam, Ales Leonardis, Steven McDonagh: **CroMo: Cross-Modal Learning for Monocular Depth Estimation.** 3927-3937
- Navami Kairanda, Edith Treitschk, Mohamed Elgharib, Christian Theobalt, Vladislav Golyanik: **\$\Phi\$-SfT: Shape-from-Template with a Physics-Based Deformation Model.** 3938-3948
- Hongwei Yi, Chun-Hao P. Huang, Dimitrios Tzionas, Muhammed Kocabas, Mohamed Hassan, Siyu Tang, Justus Thies, Michael J. Black: **Human-Aware Object Placement for Visual Environment Reconstruction.** 3949-3960
- Norman Müller, Andrea Simonelli, Lorenzo Porzi, Samuel Rota Bulò, Matthias Nießner, Peter Kontschieder: **AutoRF: Learning 3D Object Radiance Fields from Single View Observations.** 3961-3970
- Shengqu Cai, Anton Obukhov, Dengxin Dai, Luc Van Gool: **Pix2NeRF: Unsupervised Conditional \$\Phi\$-GAN for Single Image to Neural Radiance Fields Translation.** 3971-3980
- Anh-Quan Cao, Raoul de Charette: **MonoScene: Monocular 3D Semantic Scene Completion.** 3981-3991
- Felix Petersen, Bastian Goldluecke, Christian Borgelt, Oliver Deussen: **GenDR: A Generalized Differentiable Renderer.** 3992-4001
- Kuan-Chih Huang, Tsung-Han Wu, Hung-Ting Su, Winston H. Hsu: **MonoDTR: Monocular 3D Object Detection with Depth-Aware Transformer.** 4002-4011
- Can Gümeli, Angela Dai, Matthias Nießner: **ROCA: Robust CAD Model Retrieval and Alignment from a Single Image.** 4012-4021
- Chang Yu, Xiangyu Zhu, Xiaomei Zhang, Zidu Wang, Zhaoxiang Zhang, Zhen Lei: **HP-Capsule: Unsupervised Face Part Discovery by Hierarchical Parsing Capsule Network.** 4022-4031
- Xiang An, Jiankang Deng, Jia Guo, Ziyong Feng, Xuhan Zhu, Jing Yang, Tongliang Liu: **Killing Two Birds with One Stone: Efficient and Robust Training of Face Recognition CNNs by Partial FC.** 4032-4041
- Jiahao Xia, Weiwei Qu, Wenjian Huang, Jianguo Zhang, Xi Wang, Min Xu: **Sparse Local Patch Transformer for Robust Face Alignment and Landmarks Inherent Relation Learning.** 4042-4051
- Mingjie He, Jie Zhang, Shiguang Shan, Xilin Chen: **Enhancing Face Recognition with Self-Supervised 3D Reconstruction.** 4052-4061
- Chang Liu, Xiang Yu, Yi-Hsuan Tsai, Masoud Faraki, Ramin Moslemi, Manmohan Chandraker, Yun Fu: **Learning to Learn across Diverse Data Biases in Deep Face Recognition.** 4062-4072
- Kai Wang, Shuo Wang, Panpan Zhang, Zhipeng Zhou, Zheng Zhu, Xiaobo Wang, Xiaojiang Peng, Baigui Sun, Hao Li, Yang You: **An Efficient Training Approach for Very Large Scale Face Recognition.** 4073-4082
- Yang Liu, Fei Wang, Jiankang Deng, Zhipeng Zhou, Baigui Sun, Hao Li: **MogFace: Towards a Deeper Appreciation on Face Detection.** 4083-4092
- Shuai Jia, Chao Ma, Taiping Yao, Bangjie Yin, Shouhong Ding, Xiaokang Yang: **Exploring Frequency Adversarial Attacks for Face Forgery Detection.** 4093-4102
- Junyi Cao, Chao Ma, Taiping Yao, Shen Chen, Shouhong Ding, Xiaokang Yang: **End-to-End Reconstruction-Classification Learning for Face Forgery Detection.** 4103-4112
- Zhuo Wang, Zezheng Wang, Zitong Yu, Weihong Deng, Jiahong Li, Tingting Gao, Zhongyuan Wang: **Domain Generalization via Shuffled Style Assembly for Face Anti-Spoofing.** 4113-4123
- Chenqian Yan, Yuge Zhang, Quanlu Zhang, Yaming Yang, Xinyang Jiang, Yuqing Yang, Baoyuan Wang: **Privacy-preserving Online AutoML for Domain-Specific Face Detection.** 4124-4134
- Nataniel Ruiz, Adam Kortylewski, Weichao Qiu, Cihang Xie, Sarah Adel Bargal, Alan L. Yuille, Stan Sclaroff: **Simulated Adversarial Testing of Face Recognition Models.** 4135-4145
- Qingping Zheng, Jiankang Deng, Zheng Zhu, Ying Li, Stefanos Zafeiriou: **Decoupled Multi-task Learning with Cyclical Self-Regulation for Face Parsing.** 4146-4155
- Hangyu Li, Nannan Wang, Xi Yang, Xiaoyu Wang, Xinbo Gao: **Towards Semi-Supervised Deep Facial Expression Recognition with An Adaptive Confidence Margin.** 4156-4165
- Hui Li, Zidong Guo, Seon-Min Rhee, Seungju Han, Jae-Joon Han: **Towards Accurate Facial Landmark Detection via Cascaded Transformers.** 4166-4175
- Zitong Yu, Yuming Shen, Jingang Shi, Hengshuang Zhao, Philip H. S. Torr, Guoying Zhao: **PhysFormer: Facial Video-based Physiological Measurement with Temporal Difference Transformer.** 4176-4186
- Mingfang Zhang, Yunfei Liu, Feng Lu: **GazeOnce: Real-Time Multi-Person Gaze Estimation.** 4187-4196
- Yiwei Bao, Yunfei Liu, Haofei Wang, Feng Lu: **Generalizing Gaze Estimation with Rotation Consistency.** 4197-4206
- Andrew Z. Hou, Michel Sarkis, Ning Bi, Yiyi Tong, Xiaoming Liu: **Face Relighting with Geometrically Consistent Shadows.** 4207-4216
- Yiqian Wu, Yong-Liang Yang, Xiaogang Jin: **HairMapper: Removing Hair from Portraits Using GANs.** 4217-4226

- Zhenyu Zhang, Yanhao Ge, Ying Tai, Xiaoming Huang, Chengjie Wang, Hao Tang, Dongjin Huang, Zhifeng Xie:
Learning to Restore 3D Face from In-the-Wild Degraded Images. 4227-4237
- Yuchao Wang, Haochen Wang, Yujun Shen, Jingjing Fei, Wei Li, Guoqiang Jin, Liwei Wu, Rui Zhao, Xinyi Le:
Semi-Supervised Semantic Segmentation Using Unreliable Pseudo-Labels. 4238-4247
- Yuyuan Liu , Yu Tian, Yuanhong Chen, Fengbei Liu, Vasileios Belagiannis, Gustavo Carneiro :
Perturbed and Strict Mean Teachers for Semi-supervised Semantic Segmentation. 4248-4257
- Lihe Yang, Wei Zhuo, Lei Qi, Yinghuan Shi, Yang Gao:
ST++: Make Self-trainingWork Better for Semi-supervised Semantic Segmentation. 4258-4267
- Beomyoung Kim, Youngjoon Yoo, Chaeeun Rhee, Junmo Kim:
Beyond Semantic to Instance Segmentation: Weakly-Supervised Instance Segmentation via Semantic Knowledge Transfer and Self-Refinement. 4268-4277
- Qi Chen, Lingxiao Yang, Jianhuang Lai, Xiaohua Xie:
Self-supervised Image-specific Prototype Exploration for Weakly Supervised Semantic Segmentation. 4278-4288
- Tianfei Zhou, Meijie Zhang, Fang Zhao, Jianwu Li:
Regional Semantic Contrast and Aggregation for Weakly Supervised Semantic Segmentation. 4289-4299
- Lian Xu, Wanli Ouyang, Mohammed Bennamoun , Farid Boussaïd, Dan Xu:
Multi-class Token Transformer for Weakly Supervised Semantic Segmentation. 4300-4309
- Ye Du , Zehua Fu, Qingjie Liu, Yunhong Wang:
Weakly Supervised Semantic Segmentation by Pixel-to-Prototype Contrast. 4310-4319
- Minhyun Lee, Dongseob Kim, Hyunjung Shim:
Threshold Matters in WSSS: Manipulating the Activation for the Robust and Accurate Segmentation Model Against Thresholds. 4320-4329
- Yuyang Zhao, Zhun Zhong, Nicu Sebe , Gim Hee Lee:
Novel Class Discovery in Semantic Segmentation. 4330-4339
- Jin Kim, Jiyoung Lee , Jungin Park, Dongbo Min, Kwanghoon Sohn:
Pin the Memory: Learning to Generalize Semantic Segmentation. 4340-4350
- Shaohua Guo, Liang Liu, Zhenye Gan, Yabiao Wang , Wuhan Zhang, Chengjie Wang, Guannan Jiang, Wei Zhang, Ran Yi, Lizhuang Ma, Ke Xu :
ISDNet: Integrating Shallow and Deep Networks for Efficient Ultra-high Resolution Segmentation. 4351-4360
- Fabio Cermelli, Dario Fontanel, Antonio Tavera , Marco Ciccone, Barbara Caputo:
Incremental Learning in Semantic Segmentation from Image Labels. 4361-4371
- Justin Lazarow, Weijian Xu, Zhiwen Tu:
Instance Segmentation with Mask-supervised Polygonal Boundary Transformers. 4372-4381
- Chenming Zhu, Xuanye Zhang, Yanran Li, Liangdong Qiu, Kai Han, Xiaoguang Han:
SharpContour: A Contour-based Boundary Refinement Approach for Efficient and Accurate Instance Segmentation. 4382-4391
- Adrian Wolny, Qin Yu , Constantin Pape, Anna Kreshuk:
Sparse Object-level Supervision for Instance Segmentation with Pixel Embeddings. 4392-4401
- Lei Ke, Martin Danelljan, Xia Li, Yu-Wing Tai , Chi-Keung Tang, Fisher Yu:
Mask Transfiner for High-Quality Instance Segmentation. 4402-4411
- Weiyao Wang, Matt Feiszli, Heng Wang, Jitendra Malik, Du Tran:
Open-World Instance Segmentation: Exploiting Pseudo Ground Truth From Learned Pairwise Affinity. 4412-4422
- Tianheng Cheng, Xinggang Wang, Shaoyu Chen, Wenqiang Zhang, Qian Zhang, Chang Huang, Zhaoxiang Zhang, Wenyu Liu:
Sparse Instance Activation for Real-Time Instance Segmentation. 4423-4432
- Tao Zhang, Shiqing Wei, Shunping Ji:
E2EC: An End-to-End Contour-based Method for High-Quality High-Speed Instance Segmentation. 4433-4442
- Mina Ghadimi Atigh, Julian Schoep, Erman Acar , Nanne van Noord , Pascal Mettes:
Hyperbolic Image Segmentation. 4443-4452
- Dasol Han, Jaewook Yoo, Dokwan Oh:
SeeThroughNet: Resurrection of Auxiliary Loss by Preserving Class Probability Information. 4453-4462
- Kunliang Liu, Ouk Choi, Jianming Wang, Wonjun Hwang :
CDGNet: Class Distribution Guided Network for Human Parsing. 4463-4472
- Jinheng Xie, Xianxu Hou, Kai Ye, Linlin Shen:
CLIMS: Cross Language Image Matching for Weakly Supervised Semantic Segmentation. 4473-4482
- Olga Veksler, Yuri Boykov:
Sparse Non-local CRF. 4483-4493
- Yijie Zhong , Bo Li, Lv Tang, Senyun Kuang, Shuang Wu, Shouhong Ding:
Detecting Camouflaged Object in Frequency Domain. 4494-4503
- Wei Liao:
Progressive Minimal Path Method with Embedded CNN. 4504-4512

- Chang Liu , Chun Yang, Xu-Cheng Yin:
Open-Set Text Recognition via Character-Context Decoupling. 4513-4522
- Hao Liu, Xin Li, Bing Liu, Deqiang Jiang, Yinsong Liu, Bo Ren:
Neural Collaborative Graph Machines for Table Structure Recognition. 4523-4532
- Xiangwei Jiang, Rujiao Long, Nan Xue, Zhibo Yang, Cong Yao, Gui-Song Xia:
Revisiting Document Image Dewarping by Grid Regularization. 4533-4542
- Ye Yuan , Xiao Liu, Wondimu Dikubab, Hui Liu, Zhilong Ji, Zhongqin Wu, Xiang Bai:
Syntax-Aware Network for Handwritten Mathematical Expression Recognition. 4543-4552
- Jingqun Tang , Wenqing Zhang, Hongye Liu, Mingkun Yang, Bo Jiang, Guanglong Hu, Xiang Bai:
Few Could Be Better Than All: Feature Sampling and Grouping for Scene Text Detection. 4553-4562
- Chuhui Xue, Zichen Tian, Fangneng Zhan, Shijian Lu, Song Bai:
Fourier Document Restoration for Robust Document Dewarping and Recognition. 4563-4572
- Zhangxuan Gu, Changhua Meng, Ke Wang , Jun Lan, Weiqiang Wang, Ming Gu, Liqing Zhang:
XYLayoutLM: Towards Layout-Aware Multimodal Networks For Visually-Rich Document Understanding. 4573-4582
- Mingxin Huang , Yuliang Liu, Zhenghao Peng, Chongyu Liu, Dahua Lin, Shenggao Zhu, Nicholas Jing Yuan, Kai Ding, Lianwen Jin:
SwinTextSpotter: Scene Text Spotting via Better Synergy between Text Detection and Text Recognition. 4583-4593
- Yair Kitttenplon, Inbal Lavi, Sharon Fogel, Yarin Bar, R. Manmatha, Pietro Perona:
Towards Weakly-Supervised Text Spotting using a Multi-Task Transformer. 4594-4603
- Ahmed S. Nassar, Nikolaos Livathinos, Maksym Lysak, Peter W. J. Staar:
TableFormer: Table Structure Understanding with Transformers. 4604-4613
- Hao Wang, Junchao Liao , Tianheng Cheng, Zewen Gao, Hao Liu, Bo Ren, Xiang Bai, Wenyu Liu:
Knowledge Mining with Scene Text for Fine-Grained Recognition. 4614-4623
- Brandon Smock, Rohith Pesala, Robin Abraham:
PubTables-1M: Towards comprehensive table extraction from unstructured documents. 4624-4632
- Zhendong Yang, Zhe Li, Xiaohu Jiang, Yuan Gong, Zehuan Yuan, Danpei Zhao, Chun Yuan:
Focal and Global Knowledge Distillation for Detectors. 4633-4642
- Jiahao Fan, Huabin Liu, Wenjie Yang, John See , Aixin Zhang, Weiyao Lin:
Speed up Object Detection on Gigapixel-level Images with Patch Arrangement. 4643-4653
- Weixiang Hong, Jiangwei Lao, Wang Ren, Jian Wang, Jingdong Chen, Wei Chu:
Training Object Detectors from Scratch: An Empirical Study in the Era of Vision Transformer. 4652-4661
- Ahmet Iscen, Jack Valmadre , Anurag Arnab, Cordelia Schmid:
Learning with Neighbor Consistency for Noisy Labels. 4662-4671
- Chaoqun Wan, Xu Shen, Yonggang Zhang, Zhiheng Yin, Xinmei Tian, Feng Gao, Jianqiang Huang, Xian-Sheng Hua:
Meta Convolutional Neural Networks for Single Domain Generalization. 4672-4681
- Haowei Zhu, Wenjing Ke, Dong Li, Ji Liu, Lu Tian, Yi Shan:
Dual Cross-Attention Learning for Fine-Grained Visual Categorization and Object Re-Identification. 4682-4692
- Zhuangzhuang Chen, Jin Zhang, Zhuonan Lai, Jie Chen, Zun Liu, Jianqiang Li:
Geometry-Aware Guided Loss for Deep Crack Recognition. 4693-4702
- Qi Jia, Shuilian Yao, Yu Liu, Xin Fan, Risheng Liu, Zhongxuan Luo:
Segment, Magnify and Reiterate: Detecting Camouflaged Objects the Hard Way. 4703-4712
- Qinghang Hong, Fengming Liu, Dong Li, Ji Liu, Lu Tian, Yi Shan:
Dynamic Sparse R-CNN. 4713-4722
- Senqi Cao, Zhongfei Zhang:
Deep Hybrid Models for Out-of-Distribution Detection. 4723-4733
- Hongyang Gu, Jianmin Li, Guangyuan Fu, Chifong Wong, Xinghao Chen, Jun Zhu:
AutoLoss-GMS: Searching Generalized Margin-based Softmax Loss Function for Person Re-identification. 4734-4743
- Zhihang Wang, Feng Zhu, Shixiang Tang, Rui Zhao, Lihuo He, Jiangning Song:
Feature Erasing and Diffusion Network for Occluded Person Re-Identification. 4744-4753
- Emanuel Ben Baruch, Tal Ridnik, Itamar Friedman, Avi Ben-Cohen, Nadav Zamir, Asaf Noy, Lihi Zelnik-Manor:
Multi-label Classification with Partial Annotations using Class-aware Selective Loss. 4754-4762
- Duy-Kien Nguyen, Jihong Ju, Olaf Booij, Martin R. Oswald , Cees G. M. Snoek:
BoxeR: Box-Attention for 2D and 3D Transformers. 4763-4772
- Sai Rajeswar, Pau Rodríguez, Soumye Singhal, David Vázquez, Aaron C. Courville:
Multi-label Iterated Learning for Image Classification with Label Ambiguity. 4773-4783
- Zhuofan Xia , Xuran Pan, Shiji Song, Li Erran Li, Gao Huang:
Vision Transformer with Deformable Attention. 4784-4793

- Yanghao Li, Chao-Yuan Wu, Haoqi Fan, Karttikeya Mangalam, Bo Xiong, Jitendra Malik, Christoph Feichtenhofer: **MViTv2: Improved Multiscale Vision Transformers for Classification and Detection.** 4794-4804
- Binghui Chen, Pengyu Li, Xiang Chen, Biao Wang, Lei Zhang, Xian-Sheng Hua: **Dense Learning based Semi-Supervised Object Detection.** 4805-4814
- Yali Li, Shengjin Wang: **R(Det)²: Randomized Decision Routing for Object Detection.** 4815-4824
- Kareem Metwaly, Aerin Kim, Elliot Branson, Vishal Monga: **GlideNet: Global, Local and Intrinsic based Dense Embedding NETwork for Multi-category Attributes Prediction.** 4825-4836
- Jongmin Lee, Byungjin Kim, Minsu Cho: **Self-Supervised Equivariant Learning for Oriented Keypoint Detection.** 4837-4847
- Jingzhou Chen , Peng Wang, Jian Liu, Yuntao Qian: **Label Relation Graphs Enhanced Hierarchical Residual Network for Hierarchical Multi-Granularity Classification.** 4848-4857
- Xuehui Yu, Pengfei Chen , Di Wu, Najmul Hassan, Guorong Li, Junchi Yan, Humphrey Shi , Qixiang Ye, Zhenjun Han: **Object Localization under Single Coarse Point Supervision.** 4858-4867
- Gabriele Moreno Berton, Carlo Masone , Barbara Caputo: **Rethinking Visual Geo-localization for Large-Scale Applications.** 4868-4878
- Supreeth Narasimhaswamy, Thanh Nguyen, Mingzhen Huang, Minh Hoai: **Whose Hands are These? Hand Detection and Hand-Body Association in the Wild.** 4879-4889
- Yanan Wang, Xuezhi Liang, Shengcai Liao: **Cloning Outfits from Real-World Images to 3D Characters for Generalizable Person Re-Identification.** 4890-4899
- Xingxuan Zhang, Linjun Zhou, Renzhe Xu , Peng Cui, Zheyuan Shen, Haoxin Liu: **Towards Unsupervised Domain Generalization.** 4900-4910
- Haoqi Wang , Zhizhong Li, Litong Feng, Wayne Zhang : **ViM: Out-Of-Distribution with Virtual-logit Matching.** 4911-4920
- Arnav Chavan, Zhiqiang Shen, Zhuang Liu, Zechun Liu, Kwang-Ting Cheng, Eric P. Xing: **Vision Transformer Slimming: Multi-Dimension Searching in Continuous Optimization Space.** 4921-4931
- Zechun Liu, Kwang-Ting Cheng, Dong Huang, Eric P. Xing, Zhiqiang Shen: **Nonuniform-to-Uniform Quantization: Towards Accurate Quantization via Generalized Straight-Through Estimation.** 4932-4942
- Dongxu Li, Junnan Li, Hongdong Li, Juan Carlos Niebles, Steven C. H. Hoi: **Align and Prompt: Video-and-Language Pre-training with Entity Prompts.** 4943-4953
- Zihan Ding, Tianrui Hui, Junshi Huang, Xiaoming Wei, Jizhong Han, Si Liu: **Language-Bridged Spatial-Temporal Interaction for Referring Video Object Segmentation.** 4954-4963
- Jiannan Wu, Yi Jiang, Peize Sun, Zehuan Yuan, Ping Luo: **Language as Queries for Referring Video Object Segmentation.** 4964-4974
- Adam Botach, Evgenii Zheltonozhskii, Chaim Baskin : **End-to-End Referring Video Object Segmentation with Multimodal Transformers.** 4975-4985
- Dongming Wu , Xingping Dong, Ling Shao, Jianbing Shen: **Multi-Level Representation Learning with Semantic Alignment for Referring Video Object Segmentation.** 4986-4995
- Satya Krishna Gorti, Noël Vouitsis, Junwei Ma, Keyvan Golestan, Maksims Volkovs, Animesh Garg, Guangwei Yu: **X-Pool: Cross-Modal Language-Video Attention for Text-Video Retrieval.** 4996-5005
- Dohwan Ko, Joonmyung Choi, Juyeon Ko, Shinyeong Noh, Kyoung-Woon On, Eun-Sol Kim, Hyunwoo J. Kim: **Video-Text Representation Learning via Differentiable Weak Temporal Alignment.** 5006-5015
- Mattia Soldan , Alejandro Pardo, Juan León Alcázar, Fabian Caba Heilbron, Chen Zhao, Silvio Giancola, Bernard Ghanem : **MAD: A Scalable Dataset for Language Grounding in Videos from Movie Audio Descriptions.** 5016-5025
- Hongwei Xue, Tiankai Hang, Yanhong Zeng, Yuchong Sun, Bei Liu, Huan Yang, Jianlong Fu, Baining Guo: **Advancing High-Resolution Video-Language Representation with Large-Scale Video Transcriptions.** 5026-5035
- Mona Gandhi, Mustafa Omer Gul, Eva Prakash , Madeleine Grunde-McLaughlin, Ranjay Krishna, Maneesh Agrawala : **Measuring Compositional Consistency for Video Question Answering.** 5036-5045
- Paola Cascante-Bonilla, Hui Wu, Letao Wang, Rogério Feris, Vicente Ordonez: **Sim VQA: Exploring Simulated Environments for Visual Question Answering.** 5046-5056
- Feng Gao, Qing Ping, Govind Thattai, Aishwarya N. Reganti, Ying Nian Wu, Prem Natarajan: **Transform-Retrieve-Generate: Natural Language-Centric Outside-Knowledge Visual Question Answering.** 5057-5067
- Vipul Gupta, Zhuowan Li, Adam Kortylewski, Chenyu Zhang , Yingwei Li, Alan L. Yuille: **SwapMix: Diagnosing and Regularizing the Over-Reliance on Visual Context in Visual Question Answering.** 5068-5078

- Yang Ding, Jing Yu, Bang Liu, Yue Hu, Mingxin Cui, Qi Wu:
MuKEA: Multimodal Knowledge Extraction and Accumulation for Knowledge-based Visual Question Answering. 5079-5088
- Chenchen Jing, Yunde Jia, Yuwei Wu, Xinyu Liu, Qi Wu:
Maintaining Reasoning Consistency in Compositional Visual Question Answering. 5089-5098
- Aoxiong Yin, Zhou Zhao, Weike Jin, Meng Zhang, Xingshan Zeng, Xiaofei He:
MLSLT: Towards Multilingual Sign Language Translation. 5099-5109
- Yutong Chen, Fangyun Wei, Xiao Sun, Zhirong Wu, Stephen Lin:
A Simple Multi-Modality Transfer Learning Baseline for Sign Language Translation. 5110-5120
- Ronglai Zuo, Brian Mak :
C²SLR: Consistency-enhanced Continuous Sign Language Recognition. 5121-5130
- Ben Saunders , Necati Cihan Camgöz, Richard Bowden :
Signing at Scale: Learning to Co-Articulate Signs for Large-Scale Photo-Realistic Sign Language Production. 5131-5141
- Chuan Guo, Shihao Zou, Xinxin Zuo , Sen Wang , Wei Ji , Xingyu Li, Li Cheng:
Generating Diverse and Natural 3D Human Motions from Text. 5142-5151
- K. R. Prajwal, Triantafyllos Afouras, Andrew Zisserman:
Sub-word Level Lip Reading With Visual Attention. 5152-5162
- Ram Ramrakhya, Eric Undersander, Dhruv Batra, Abhishek Das:
Habitat-Web: Learning Embodied Object-Search Strategies from Human Demonstrations at Scale. 5163-5173
- Mengjun Cheng, Yipeng Sun, Longchao Wang, Xiongwei Zhu, Kun Yao, Jie Chen, Guoli Song, Junyu Han, Jingtuo Liu, Errui Ding, Jingdong Wang:
ViSTA: Vision and Scene Text Aggregation for Cross-Modal Retrieval. 5174-5183
- Simion-Vlad Bogolin, Ioana Croitoru, Hailin Jin, Yang Liu , Samuel Albanie:
Cross Modal Retrieval with Querybank Normalisation. 5184-5195
- Yuning Lu, Jianzhuang Liu, Yonggang Zhang, Yajing Liu, Xinmei Tian:
Prompt Distribution Learning. 5196-5205
- Yi Li, Rameswar Panda, Yoon Kim, Chun-Fu Richard Chen, Rogério Feris, David D. Cox, Nuno Vasconcelos :
VALHALLA: Visual Hallucination for Machine Translation. 5206-5216
- Yi-Lin Sung, Jaemin Cho, Mohit Bansal:
VL-ADAPTER: Parameter-Efficient Transfer Learning for Vision-and-Language Tasks. 5217-5227
- Tristan Thrush, Ryan Jiang, Max Bartolo, Amanpreet Singh, Adina Williams, Douwe Kiela, Candace Ross:
Winoground: Probing Vision and Language Models for Visio-Linguistic Compositionality. 5228-5238
- Qiang Chen, Qiman Wu, Jian Wang, Qinghao Hu, Tao Hu, Errui Ding, Jian Cheng, Jingdong Wang:
MixFormer: Mixing Features across Windows and Dimensions. 5239-5249
- Zhe Chen , Jing Zhang, Dacheng Tao:
Recurrent Glimpse-based Decoder for Detection with Transformer. 5250-5259
- Yinpeng Chen, Xiyang Dai, Dongdong Chen, Mengchen Liu, Xiaoyi Dong, Lu Yuan, Zicheng Liu:
Mobile-Former: Bridging MobileNet and Transformer. 5260-5269
- Sivan Harary, Eli Schwartz, Assaf Arbelle, Peter W. J. Staar, Shady Abu-Hussein, Elad Amrani, Roei Herzig, Amit Alfassy, Raja Giryes, Hilde Kuehne , Dina Katabi, Kate Saenko, Rogério Feris, Leonid Karlinsky:
Unsupervised Domain Generalization by Learning a Bridge Across Domains. 5270-5280
- Wuyang Li , Xinyu Liu , Yixuan Yuan :
SIGMA: Semantic-complete Graph Matching for Domain Adaptive Object Detection. 5281-5290
- Jiaxi Wu, Jiaxin Chen, Mengzhe He, Yiru Wang, Bo Li, Bingqi Ma, Weihao Gan, Wei Wu, Yali Wang, Di Huang:
Target-Relevant Knowledge Preservation for Multi-Source Domain Adaptive Object Detection. 5291-5300
- Zeren Sun, Fumin Shen, Dan Huang, Qiong Wang, Xiangbo Shu, Yazhou Yao, Jinhui Tang:
PNP: Robust Learning from Noisy Labels by Probabilistic Noise Prediction. 5301-5310
- Guangxing Han, Jiawei Ma, Shiyuan Huang, Long Chen , Shih-Fu Chang:
Few-Shot Object Detection with Fully Cross-Transformer. 5311-5320
- Su Been Lee, WonJun Moon, Jae-Pil Heo:
Task Discrepancy Maximization for Fine-grained Few-Shot Classification. 5321-5330
- Weizhe Liu, Nikita Durasov, Pascal Fua:
Leveraging Self-Supervision for Cross-Domain Crowd Counting. 5331-5342
- A. S. M. Iftekhar, Hao Chen, Kaustav Kundu, Xinyu Li, Joseph Tighe, Davide Modolo:
What to look at and where: Semantic and Spatial Refined Transformer for detecting human-object interactions. 5343-5353
- Ziteng Gao, Limin Wang, Bing Han, Sheng Guo:
AdaMixer: A Fast-Converging Query-Based Object Detector. 5354-5363

- Seongwon Lee Hongje Seong, Suhyeon Lee , Euntai Kim:
Correlation Verification for Image Retrieval. 5364-5374
- Jinrong Yang, Songtao Liu, Zeming Li, Xiaoping Li, Jian Sun:
Real-time Object Detection for Streaming Perception. 5375-5385
- Gabriele Moreno Berton, Riccardo Mereu, Gabriele Trivigno, Carlo Masone , Gabriela Csurka, Torsten Sattler, Barbara Caputo:
Deep Visual Geo-localization Benchmark. 5386-5397
- Ruoxi Shi, Xinyang Jiang, Caihua Shan, Yansen Wang, Dongsheng Li:
RendNet: Unified 2D/3D Recognizer with Latent Space Rendering. 5398-5407
- Xiaopei Wu, Liang Peng, Honghui Yang, Liang Xie, Chenxi Huang, Chengqi Deng, Haifeng Liu, Deng Cai:
Sparse Fuse Dense: Towards High Quality 3D Detection with Depth Completion. 5408-5417
- Yukang Chen, Yanwei Li, Xiangyu Zhang, Jian Sun, Jiaya Jia:
Focal Sparse Convolutional Networks for 3D Object Detection. 5418-5427
- Qiangeng Xu, Zexiang Xu, Julien Philip, Sai Bi, Zhixin Shu, Kalyan Sunkavalli, Ulrich Neumann:
Point-NeRF: Point-based Neural Radiance Fields. 5428-5438
- Xiaoshuai Zhang, Sai Bi, Kalyan Sunkavalli, Hao Su, Zexiang Xu:
NeRFusion: Fusing Radiance Fields for Large-Scale Scene Reconstruction. 5439-5448
- Cheng Sun, Min Sun, Hwann-Tzong Chen:
Direct Voxel Grid Optimization: Super-fast Convergence for Radiance Fields Reconstruction. 5449-5459
- Jonathan T. Barron, Ben Mildenhall, Dor Verbin, Pratul P. Srinivasan, Peter Hedman:
Mip-NeRF 360: Unbounded Anti-Aliased Neural Radiance Fields. 5460-5469
- Michael Niemeyer, Jonathan T. Barron, Ben Mildenhall, Mehdi S. M. Sajjadi, Andreas Geiger, Noha Radwan:
RegNeRF: Regularizing Neural Radiance Fields for View Synthesis from Sparse Inputs. 5470-5480
- Dor Verbin, Peter Hedman, Ben Mildenhall, Todd E. Zickler, Jonathan T. Barron, Pratul P. Srinivasan:
Ref-NeRF: Structured View-Dependent Appearance for Neural Radiance Fields. 5481-5490
- Sara Fridovich-Keil, Alex Yu, Matthew Tancik, Qinhong Chen, Benjamin Recht, Angjoo Kanazawa:
Plenoxels: Radiance Fields without Neural Networks. 5491-5500
- Haoyu Guo, Sida Peng, Haotong Lin, Qianqian Wang, Guofeng Zhang, Hujun Bao, Xiaowei Zhou:
Neural 3D Scene Reconstruction with the Manhattan-world Assumption. 5501-5510
- Tianye Li, Mira Slavcheva, Michael Zollhöfer, Simon Green, Christoph Lassner, Changil Kim, Tanner Schmidt, Steven Lovegrove, Michael Goesele, Richard A. Newcombe, Zhaoyang Lv:
Neural 3D Video Synthesis from Multi-view Video. 5511-5521
- Petr Hruby, Timothy Duff , Anton Leykin, Tomás Pajdla:
Learning to Solve Hard Minimal Problems. 5522-5532
- Yingjie Cai, Kwan-Yee Lin, Chao Zhang, Qiang Wang, Xiaogang Wang , Hongsheng Li :
Learning a Structured Latent Space for Unsupervised Point Cloud Completion. 5533-5543
- Yang Li, Tatsuya Harada:
Lepard: Learning partial point cloud matching in rigid and deformable scenes. 5544-5554
- Kai Zhang, Fujun Luan , Zhengqi Li, Noah Snavely:
IRON: Inverse Rendering by Optimizing Neural SDFs and Materials from Photometric Images. 5555-5564
- Damien Robert , Bruno Vallet, Loïc Landrieu:
Learning Multi-View Aggregation In the Wild for Large-Scale 3D Semantic Segmentation. 5565-5574
- Yu Zheng, Yueqi Duan, Jiwen Lu, Jie Zhou, Qi Tian:
HyperDet3D: Learning a Scene-conditioned 3D Object Detector. 5575-5584
- David Novotný, Ignacio Rocco, Samarth Sinha, Alexandre Carlier, Gael Kerchenbaum, Roman Shapovalov, Nikita Smetanin, Natalia Neverova, Benjamin Graham, Andrea Vedaldi:
KeyTr: Keypoint Transporter for 3D Reconstruction of Deformable Objects in Videos. 5585-5594
- Boyi Jiang, Yang Hong, Hujun Bao, Juyong Zhang:
SelfRecon: Self Reconstruction Your Digital Avatar from Monocular Video. 5595-5605
- Zhenyu Jiang, Cheng-Chun Hsu, Yuke Zhu:
Ditto: Building Digital Twins of Articulated Objects from Interaction. 5606-5616
- Yurui Zhu, Jie Huang, Xueyang Fu , Feng Zhao, Qibin Sun, Zheng-Jun Zha:
Bijective Mapping Network for Shadow Removal. 5617-5626
- Long Ma, Tengyu Ma, Risheng Liu, Xin Fan, Zhongxuan Luo:
Toward Fast, Flexible, and Robust Low-Light Image Enhancement. 5627-5636
- Dongdong Chen, Julián Tachella, Mike E. Davies:
Robust Equivariant Imaging: a fully unsupervised framework for learning to image from noisy and partial measurements. 5637-5646
- Jie Liang, Hui Zeng, Lei Zhang:
Details or Artifacts: A Locally Discriminative Learning Approach to Realistic Image Super-Resolution. 5647-5656

- Xiaoqian Xu, Pengxu Wei, Weikai Chen, Yang Liu, Mingzhi Mao, Liang Lin, Guanbin Li:
Dual Adversarial Adaptation for Cross-Device Real-World Image Super-Resolution. 5657-5666
- Youngho Yoon, Inchul Chung, Lin Wang , Kuk-Jin Yoon:
SphereSR: 360° Image Super-Resolution with Arbitrary Projection via Continuous Spherical Image Representation. 5667-5676
- Chengxu Liu, Huan Yang, Jianlong Fu, Xueming Qian:
Learning Trajectory-Aware Transformer for Video Super-Resolution. 5677-5686
- Zixiang Zhao , Jiangshe Zhang, Shuang Xu, Zudi Lin, Hanspeter Pfister :
Discrete Cosine Transform Network for Guided Depth Map Super-Resolution. 5687-5697
- Zhixuan Zhong, Liangyu Chai, Yang Zhou, Bailin Deng, Jia Pan, Shengfeng He :
Faithful Extreme Rescaling via Generative Prior Reciprocated Invertible Representations. 5698-5707
- Dailan He, Ziming Yang, Weikun Peng, Rui Ma, Hongwei Qin, Yan Wang:
ELIC: Efficient Learned Image Compression with Unevenly Grouped Space-Channel Contextual Adaptive Coding. 5708-5717
- Syed Waqas Zamir, Aditya Arora, Salman Khan, Munawar Hayat, Fahad Shahbaz Khan, Ming-Hsuan Yang:
Restormer: Efficient Transformer for High-Resolution Image Restoration. 5718-5729
- Lang Nie, Chunyu Lin, Kang Liao, Shuaicheng Liu , Yao Zhao:
Deep Rectangling for Image Stitching: A Learning Baseline. 5730-5738
- Shanel Gauthier, Benjamin Thérien, Laurent Alsène-Racicot, Muawiz Chaudhary, Irina Rish, Eugene Belilovsky, Michael Eickenberg, Guy Wolf:
Parametric Scattering Networks. 5739-5748
- Akshay Dudhane, Syed Waqas Zamir, Salman Khan, Fahad Shahbaz Khan, Ming-Hsuan Yang:
Burst Image Restoration and Enhancement. 5749-5758
- Zhengzhong Tu , Hossein Talebi, Han Zhang, Feng Yang, Peyman Milanfar, Alan C. Bovik, Yinxiao Li:
MAXIM: Multi-Axis MLP for Image Processing. 5759-5770
- Javier Hidalgo-Carrió, Guillermo Gallego, Davide Scaramuzza :
Event-aided Direct Sparse Odometry. 5771-5780
- Haisong Liu, Tao Lu, Yihui Xu, Jia Liu, Wenjie Li, Lijun Chen:
CamLiFlow: Bidirectional Camera-LiDAR Fusion for Joint Optical Flow and Scene Flow Estimation. 5781-5791
- Jinyuan Liu , Xin Fan, Zhanbo Huang, Guanyao Wu, Risheng Liu, Wei Zhong, Zhongxuan Luo:
Target-aware Dual Adversarial Learning and a Multi-scenario Multi-Modality Benchmark to Fuse Infrared and Visible for Object Detection. 5792-5801
- Chunle Guo, Qixin Yan, Saeed Anwar, Runmin Cong, Wenqi Ren, Chongyi Li:
Image Dehazing Transformer with Transmission-Aware 3D Position Embedding. 5802-5810
- Yuntong Ye, Changfeng Yu, Yi Chang, Lin Zhu , Xi-Le Zhao, Luxin Yan, Yonghong Tian:
Unsupervised Deraining: Where Contrastive Learning Meets Self-similarity. 5811-5820
- Huan Liu, Zijun Wu, Liangyan Li, Sadaf Salehkalaibar, Jun Chen, Keyan Wang:
Towards Multi-domain Single Image Dehazing via Test-time Training. 5821-5830
- Yi Li, Yi Chang, Yan Gao, Changfeng Yu, Luxin Yan:
Physically Disentangled Intra- and Inter-domain Adaptation for Varicolored Haze Removal. 5831-5840
- Yue Cao, Zhaolin Wan, Dongwei Ren, Zifei Yan, Wangmeng Zuo:
Incorporating Semi-Supervised and Positive-Unlabeled Learning for Boosting Full Reference Image Quality Assessment. 5841-5851
- Lina Guo, Xinjie Shi, Dailan He, Yuanyuan Wang, Rui Ma, Hongwei Qin, Yan Wang:
Practical Learned Lossless JPEG Recompression with Multi-Level Cross-Channel Entropy Model in the DCT Domain. 5852-5861
- Cong Huang , Jiahao Li, Bin Li, Dong Liu, Yan Lu:
Neural Compression-Based Feature Learning for Video Restoration. 5862-5871
- Xin Tian, Ke Xu , Xin Yang, Lin Du, Baocai Yin, Rynson W. H. Lau:
Bi-directional Object-Context Prioritization Learning for Saliency Ranking. 5872-5881
- Wenhui Wu, Jian Weng, Pingping Zhang, Xu Wang, Wenhan Yang, Jianmin Jiang:
URetinex-Net: Retinex-based Deep Unfolding Network for Low-light Image Enhancement. 5891-5900
- Jianqi Ma , Zhetong Liang, Lei Zhang:
A Text Attention Network for Spatial Deformation Robust Scene Text Image Super-resolution. 5901-5910
- Zhihao Hu, Guo Lu, Jinyang Guo, Shan Liu, Wei Jiang, Dong Xu:
Coarse-To-Fine Deep Video Coding with Hyperprior-Guided Mode Prediction. 5911-5920
- Yixuan Huang, Xiaoyun Zhang, Yu Fu, Siheng Chen, Ya Zhang , Yanfeng Wang, Dazhi He:
Task Decoupled Framework for Reference-based Super-Resolution. 5921-5930
- Huankang Guan , Jiaying Lin , Rynson W. H. Lau:
Learning Semantic Associations for Mirror Detection. 5931-5940

- Yu Zeng, Zhe Lin, Vishal M. Patel:
SketchEdit: Mask-Free Local Image Manipulation with Partial Sketches. 5941-5951
- Kelvin C. K. Chan, Shangchen Zhou, Xiangyu Xu, Chen Change Loy:
Investigating Tradeoffs in Real-World Video Super-Resolution. 5952-5961
- Kelvin C. K. Chan , Shangchen Zhou, Xiangyu Xu, Chen Change Loy:
BasicVSR++: Improving Video Super-Resolution with Enhanced Propagation and Alignment. 5962-5971
- Kaidong Zhang, Jingjing Fu, Dong Liu:
Inertia-Guided Flow Completion and Style Fusion for Video Inpainting. 5972-5981
- Jun-Hyuk Kim , Byeongho Heo, Jong-Seok Lee:
Joint Global and Local Hierarchical Priors for Learned Image Compression. 5982-5991
- Xiangtao Kong , Xina Liu, Jinjin Gu, Yu Qiao , Chao Dong:
Reflash Dropout in Image Super-Resolution. 5992-6002
- Yi Yu, Wenhan Yang, Yap-Peng Tan, Alex C. Kot:
Towards Robust Rain Removal Against Adversarial Attacks: A Comprehensive Benchmark Analysis and Beyond. 6003-6012
- Weiqi Zou, Yang Wang, Xueyang Fu , Yang Cao:
Dreaming to Prune Image Deraining Networks. 6013-6022
- Hochang Rhee, Yeong Il Jang, Seyun Kim, Nam Ik Cho:
LC-FDNet: Learned Lossless Image Compression with Frequency Decomposition Network. 6023-6032
- Jie Huang, Yajing Liu, Xueyang Fu , Man Zhou, Yang Wang, Feng Zhao, Zhiwei Xiong:
Exposure Normalization and Compensation for Multiple-Exposure Correction. 6033-6042
- Kun Zhou , Wenbo Li , Liying Lu, Xiaoguang Han, Jiangbo Lu :
Revisiting Temporal Alignment for Video Restoration. 6043-6052
- Zhenghao Chen, Guo Lu, Zhihao Hu, Shan Liu, Wei Jiang, Dong Xu:
LSVC: A Learning-based Stereo Video Compression Framework. 6063-6072
- Guo Lu, Tianxiong Zhong, Jing Geng, Qiang Hu, Dong Xu:
Learning based Multi-modality Image and Video Compression. 6073-6082
- Xin Tong , Xianghua Ying, Yongjie Shi, Ruibin Wang, Jinfa Yang:
Transformer Based Line Segment Classifier with Image Context for Real-Time Vanishing Point Detection in Manhattan World. 6083-6092
- Yancong Lin, Ruben Wiersma , Silvia L. Pintea, Klaus Hildebrandt, Elmar Eisemann, Jan C. van Gemert:
Deep vanishing point detection: Geometric priors make dataset variations vanish. 6093-6103
- Yeongwoo Nam, S. Mohammad Mostafavi I., Kuk-Jin Yoon, Jonghyun Choi :
Stereo Depth from Events Cameras: Concentrate and Focus on the Future. 6104-6113
- Ronald Clark:
Volumetric Bundle Adjustment for Online Photorealistic Scene Capture. 6114-6122
- Zhongzheng Ren, Aseem Agarwala, Bryan C. Russell, Alexander G. Schwing, Oliver Wang:
Neural Volumetric Object Selection. 6123-6132
- Ziyan Wang, Giljoo Nam, Tuur Stuyck, Stephen Lombardi, Michael Zollhöfer, Jessica K. Hodgins, Christoph Lassner:
HVH: Learning a Hybrid Neural Volumetric Representation for Dynamic Hair Performance Capture. 6133-6144
- Yuheng Jiang, Suyi Jiang, Guoxing Sun, Zhuo Su , Kaiwen Guo, Minye Wu, Jingyi Yu, Lan Xu:
NeuralHOFusion: Neural Volumetric Rendering under Human-object Interactions. 6145-6155
- Kejie Li, Yansong Tang, Victor Adrian Prisacariu, Philip H. S. Torr:
BNF-Fusion: Dense 3D Reconstruction using Bi-level Neural Volume Fusion. 6156-6165
- Wang Yifan , Carl Doersch, Relja Arandjelovic, João Carreira, Andrew Zisserman:
Input-level Inductive Biases for 3D Reconstruction. 6166-6176
- Markus Worchel, Rodrigo Diaz, Weiwen Hu, Oliver Schreer , Ingo Feldmann, Peter Eisert:
Multi-View Mesh Reconstruction with Neural Deferred Shading. 6177-6187
- Lukas Höller, Justin Johnson, Matthias Nießner:
StyleMesh: Style Transfer for Indoor 3D Scene Reconstructions. 6188-6198
- Haowen Wang, Mingyuan Wang, Zhengping Che, Zhiyuan Xu, Xiuquan Qiao, Mengshi Qi, Feifei Feng, Jian Tang:
RGB-Depth Fusion GAN for Indoor Depth Completion. 6199-6208
- Yiming Xie, Matheus Gadelha, Fengting Yang, Xiaowei Zhou, Huaizu Jiang:
PlanarRecon: Realtime 3D Plane Detection and Reconstruction from Posed Monocular Videos. 6209-6218
- Mehdi S. M. Sajjadi, Henning Meyer, Etienne Pot, Urs Bergmann, Klaus Greff, Noha Radwan, Suhani Vora, Mario Lucic, Daniel Duckworth, Alexey Dosovitskiy, Jakob Uszkoreit, Thomas A. Funkhouser, Andrea Tagliasacchi:
Scene Representation Transformer: Geometry-Free Novel View Synthesis Through Set-Latent Scene Representations. 6219-6228
- Xinguang Yan, Liqiang Lin, Niloy J. Mitra, Dani Lischinski, Daniel Cohen-Or, Hui Huang :
ShapeFormer: Transformer-based Shape Completion via Sparse Representation. 6229-6239



- Kyeongha Rho, Jinsung Ha, Youngjung Kim:
GuideFormer: Transformers for Image Guided Depth Completion. 6240-6249
- François Darmon, Bénédicte Bascle, Jean-Clément Devaux, Pascal Monasse, Mathieu Aubry:
Improving neural implicit surfaces geometry with patch warping. 6250-6259
- Jingyang Zhang, Yao Yao, Shiwei Li, Tian Fang, David McKinnon, Yanghai Tsin, Long Quan:
Critical Regularizations for Neural Surface Reconstruction in the Wild. 6260-6269
- Christiane Sommer, Lu Sang, David Schubert, Daniel Cremers :
Gradient-SDF: A Semi-Implicit Surface Representation for 3D Reconstruction. 6270-6279
- Dejan Azinovic, Ricardo Martin-Brualla, Dan B. Goldman, Matthias Nießner, Justus Thies:
Neural RGB-D Surface Reconstruction. 6280-6291
- Alexandre Boulch, Renaud Marlet:
POCO: Point Convolution for Surface Reconstruction. 6292-6304
- Baorui Ma, Yu-Shen Liu, Zhizhong Han:
Reconstructing Surfaces for Sparse Point Clouds with On-Surface Priors. 6305-6315
- Baorui Ma, Yu-Shen Liu, Matthias Zwicker, Zhizhong Han:
Surface Reconstruction from Point Clouds by Learning Predictive Context Priors. 6316-6327
- Yiming Zeng , Yue Qian, Qijian Zhang , Junhui Hou , Yixuan Yuan , Ying He:
IDEA-Net: Dynamic 3D Point Cloud Interpolation via Deep Embedding Alignment. 6328-6337
- Wen Chen, Haoang Li, Qiang Nie, Yun-Hui Liu:
Deterministic Point Cloud Registration via Novel Transformation Decomposition. 6338-6346
- Che Sun, Yunde Jia, Yi Guo, Yuwei Wu:
Global-Aware Registration of Less-Overlap RGB-D Scans. 6347-6356
- Mulin Yu, Florent Lafarge:
Finding Good Configurations of Planar Primitives in Unorganized Point Clouds. 6357-6366
- Hehe Fan, Xiaojun Chang , Wanyue Zhang, Yi Cheng, Ying Sun, Mohan S. Kankanhalli :
Self-Supervised Global-Local Structure Modeling for Point Cloud Domain Adaptation with Reliable Voted Pseudo Labels. 6367-6376
- Shaoyu Chen, Xinggang Wang, Tianheng Cheng, Wenqiang Zhang, Qian Zhang, Chang Huang, Wenyu Liu:
AzINorm: Exploiting the Radial Symmetry of Point Cloud for Azimuth-Normalized 3D Perception. 6377-6386
- Yingzhi Tang, Yue Qian, Qijian Zhang , Yiming Zeng , Junhui Hou , Xuefei Zhe:
WarpingGAN: Warping Multiple Uniform Priors for Adversarial 3D Point Cloud Generation. 6387-6395
- Mingzhen Huang, Supreeth Narasimhaswamy, Saif Vazir, Haibin Ling, Minh Hoai:
Forward Propagation, Backward Regression, and Pose Association for Hand Tracking in the Wild. 6396-6406
- Buzhen Huang , Liang Pan , Yuan Yang, Jingyi Ju, Yangang Wang:
Neural MoCon: Neural Motion Control for Physically Plausible Human Motion Capture. 6407-6416
- Takahiro Maeda, Norimichi Ukita:
MotionAug: Augmentation with Physical Correction for Human Motion Prediction. 6417-6426
- Tiezheng Ma, Yongwei Nie, Chengjiang Long, Qing Zhang, Guiqing Li:
Progressively Generating Better Initial Guesses Towards Next Stages for High-Quality Human Motion Prediction. 6427-6436
- Chongyang Zhong, Lei Hu, Zihao Zhang, Yongjing Ye, Shihong Xia:
Spatio-Temporal Gating-Adjacency GCN for Human Motion Prediction. 6437-6446
- Tim Salzmann , Marco Pavone , Markus Ryll:
Motron: Multimodal Probabilistic Human Motion Forecasting. 6447-6456
- Jianhua Sun, Yuxuan Li, Liang Chai, Haoshu Fang, Yong-Lu Li, Cewu Lu:
Human Trajectory Prediction with Momentary Observation. 6457-6466
- Inhwan Bae , Jin-Hwi Park, Hae-Gon Jeon :
Non-Probability Sampling Network for Stochastic Human Trajectory Prediction. 6467-6477
- Chenxin Xu, Weibo Mao, Wenjun Zhang, Siheng Chen:
Remember Intentions: Retrospective-Memory-based Trajectory Prediction. 6478-6487
- Chenxin Xu, Maosen Li, Zhenyang Ni, Ya Zhang , Siheng Chen:
GroupNet: Multiscale Hypergraph Neural Networks for Trajectory Prediction with Relational Reasoning. 6488-6497
- Zhangxing Bian , Allan Jabri, Alexei A. Efros, Andrew Owens:
Learning Pixel Trajectories with Multiscale Contrastive Random Walks. 6498-6509
- Yi Xu, Lichen Wang, Yizhou Wang, Yun Fu:
Adaptive Trajectory Prediction via Transferable GNN. 6510-6521
- Chaoyang Wang, Xueqian Li, Jhony Kaesemel Pontes, Simon Lucey :
Neural Prior for Trajectory Estimation. 6522-6532

- Qiao Sun , Xin Huang, Junru Gu, Brian C. Williams, Hang Zhao:
M2I: From Factored Marginal Trajectory Prediction to Interactive Prediction. 6533-6542
- Alessio Monti, Angelo Porrello, Simone Calderara, Pasquale Coscia, Lamberto Ballan , Rita Cucchiara:
How many Observations are Enough? Knowledge Distillation for Trajectory Forecasting. 6543-6552
- Chunnan Wang, Xiang Chen, Junzhe Wang, Hongzhi Wang:
ATPFL: Automatic Trajectory Prediction Model Design under Federated Learning Framework. 6553-6562
- Xinshuo Weng, Boris Ivanovic, Kris Kitani, Marco Pavone :
Whose Track Is It Anyway? Improving Robustness to Tracking Errors with Affinity-based Trajectory Prediction. 6563-6572
- Zhaoen Su, Chao Wang, David Bradley, Carlos Vallespi-Gonzalez, Carl Wellington, Nemanja Djuric:
Convolutions for Spatial Interaction Modeling. 6573-6582
- Tianxin Tao, Xiaohang Zhan, Zhongquan Chen, Michiel van de Panne:
Style-ERD: Responsive and Coherent Online Motion Style Transfer. 6583-6593
- Sachini Herath, David Caruso, Chen Liu, Yufan Chen, Yasutaka Furukawa:
Neural Inertial Localization. 6594-6603
- Xiya Cao, Caifa Zhou, Dandan Zeng, Yongliang Wang:
RIO: Rotation-equivariance supervised learning of robust inertial odometry. 6604-6613
- Jiahui Lei, Kostas Daniilidis:
CaDeX: Learning Canonical Deformation Coordinate Space for Dynamic Surface Representation via Neural Homeomorphism. 6614-6624
- Bastian Wandt, James J. Little, Helge Rhodin:
ElePose: Unsupervised 3D Human Pose Estimation by Predicting Camera Elevation and Learning Normalizing Flows on 2D Poses. 6625-6635
- Jiayi Chen , Yingda Yin, Tolga Birdal , Baoquan Chen, Leonidas J. Guibas, He Wang:
Projective Manifold Gradient Layer for Deep Rotation Regression. 6636-6645
- Noam Rotstein, Amit Bracha, Ron Kimmel:
Multimodal Colored Point Cloud to Image Alignment. 6646-6656
- Weixuan Tang, Danping Zou:
Multi-instance Point Cloud Registration by Efficient Correspondence Clustering. 6657-6666
- Zi Jian Yew, Gim Hee Lee:
REGTR: End-to-end Point Cloud Correspondences with Transformers. 6667-6676
- Manuel Kolmet, Qunjie Zhou, Aljosa Osep, Laura Leal-Taixé:
Text2Pos: Text-to-Point-Cloud Cross-Modal Localization. 6677-6686
- Jiachen Li , Bin Wang, Shiqiang Zhu, Xin Cao, Fan Zhong, Wenxuan Chen, Te Li, Jason Gu, Xueying Qin:
BCOT: A Markerless High-Precision 3D Object Tracking Benchmark. 6687-6696
- Haitao Lin, Zichang Liu, Chilam Cheang, Yanwei Fu , Guodong Guo, Xiangyang Xue:
SAR-Net: Shape Alignment and Recovery Network for Category-level 6D Object Pose and Size Estimation. 6697-6707
- Ningkai Mo, Wanshui Gan, Naoto Yokoya, Shifeng Chen:
ES6D: A Computation Efficient and Symmetry-Aware 6D Pose Regression Framework. 6708-6717
- Lahav Lipson, Zachary Teed, Ankit Goyal, Jia Deng:
Coupled Iterative Refinement for 6D Multi-Object Pose Estimation. 6718-6727
- Yongzhi Su , Mahdi Saleh, Torben Fetzer, Jason R. Rambach, Nassir Navab, Benjamin Busam , Didier Stricker, Federico Tombari:
ZebraPose: Coarse to Fine Surface Encoding for 6DoF Object Pose Estimation. 6728-6738
- Rasmus Laurvig Haugaard, Anders Glent Buch:
SurfEmb: Dense and Continuous Correspondence Distributions for Object Pose Estimation with Learnt Surface Embeddings. 6739-6748
- Ben Usman, Andrea Tagliasacchi, Kate Saenko, Avneesh Sud:
MetaPose: Fast 3D Pose from Multiple Views without 3D Supervision. 6749-6760
- Van Nguyen Nguyen , Yinlin Hu, Yang Xiao, Mathieu Salzmann, Vincent Lepetit:
Templates for 3D Object Pose Estimation Revisited: Generalization to New Objects and Robustness to Occlusions. 6761-6770
- Yan Di, Ruida Zhang, Zhiqiang Lou, Fabian Manhardt, Xiangyang Ji, Nassir Navab, Federico Tombari:
GPV-Pose: Category-level Object Pose Estimation via Geometry-guided Point-wise Voting. 6771-6781
- Yudi Dai , Yitai Lin, Chenglu Wen, Siqi Shen, Lan Xu, Jingyi Yu, Yuexin Ma, Cheng Wang:
HSC4D: Human-centered 4D Scene Capture in Large-scale Indoor-outdoor Space Using Wearable IMUs and LiDAR. 6782-6792
- Dingding Cai , Janne Heikkilä, Esa Rahtu :
OVE6D: Object Viewpoint Encoding for Depth-based 6D Object Pose Estimation. 6793-6803

- Yisheng He, Yao Wang, Haoqiang Fan, Jian Sun, Qifeng Chen:
FS6D: Few-Shot 6D Pose Estimation of Novel Objects. 6804-6814
- Jiaming Sun , Zihao Wang, Siyu Zhang, Xingyi He, Hongcheng Zhao, Guofeng Zhang, Xiaowei Zhou:
OnePose: One-Shot Object Pose Estimation without CAD Models. 6815-6824
- Ivan Shugurov, Fu Li, Benjamin Busam , Slobodan Ilic:
OSOP: A Multi-Stage One Shot Object Pose Estimation Framework. 6825-6834
- Chethan M. Parameshwara, Gokul Hari, Cornelia Fermüller, Nitin J. Sanket, Yiannis Aloimonos:
DiffPoseNet: Direct Differentiable Camera Pose Estimation. 6835-6844
- Manuel Stoiber, Martin Sundermeyer, Rudolph Triebel:
Iterative Corresponding Geometry: Fusing Region and Depth for Highly Efficient 3D Tracking of Textureless Objects. 6845-6855
- Yang You, Ruoxi Shi, Weiming Wang, Cewu Lu:
CPPF: Towards Robust Category-Level 9D Pose Estimation in the Wild. 6856-6865
- Mohamed Adel Musallam, Vincent Gaudillière, Miguel Ortiz del Castillo, Kassem Al Ismaeil, Djamila Aouada:
Leveraging Equivariant Features for Absolute Pose Regression. 6866-6876
- Seulki Park, Youngkyu Hong, Byeongho Heo, Sangdoo Yun, Jin Young Choi:
The Majority Can Help the Minority: Context-rich Minority Oversampling for Long-tailed Classification. 6877-6886
- Shaden Alshammari, Yu-Xiong Wang, Deva Ramanan , Shu Kong:
Long-Tailed Recognition via Weight Balancing. 6887-6897
- Jianggang Zhu, Zheng Wang , Jingjing Chen , Yi-Ping Phoebe Chen , Yu-Gang Jiang:
Balanced Contrastive Learning for Long-Tailed Visual Recognition. 6898-6907
- Tianhong Li, Peng Cao, Yuan Yuan, Lijie Fan, Yuzhe Yang , Rogério Feris, Piotr Indyk, Dina Katabi:
Targeted Supervised Contrastive Learning for Long-Tailed Recognition. 6908-6918
- Mengke Li, Yiu-Ming Cheung, Yang Lu:
Long-tailed Visual Recognition via Gaussian Clouded Logit Adjustment. 6919-6928
- Sarah Parisot, Pedro M. Esperança, Steven McDonagh , Tamas J. Madarasz, Yongxin Yang, Zhenguo Li:
Long-tail Recognition via Compositional Knowledge Transfer. 6929-6938
- Jun Li, Zichang Tan, Jun Wan, Zhen Lei, Guodong Guo:
Nested Collaborative Learning for Long-Tailed Visual Recognition. 6939-6948
- Alexander Long, Wei Yin, Thalaiyasingam Ajanthan, Vu Nguyen, Pulak Purkait, Ravi Garg , Alan Blair, Chunhua Shen, Anton van den Hengel:
Retrieval Augmented Classification for Long-Tail Visual Recognition. 6949-6959
- Bolian Li , Zongbo Han, Haining Li, Huazhu Fu , Changqing Zhang:
Trustworthy Long-Tailed Classification. 6960-6969
- Tong Wang, Yousong Zhu, Yingying Chen, Chaoyang Zhao, Bin Yu, Jinqiao Wang, Ming Tang:
C2AM Loss: Chasing a Better Decision Boundary for Long-Tail Object Detection. 6970-6979
- Bo Li , Yongqiang Yao, Jingru Tan, Gang Zhang, Fengwei Yu, Jianwei Lu, Ye Luo:
Equalized Focal Loss for Dense Long-Tailed Object Detection. 6980-6989
- Yin-Yin He, Peizhen Zhang, Xiu-Shen Wei, Xiangyu Zhang, Jian Sun:
Relieving Long-tailed Instance Segmentation via Pairwise Class Balance. 6990-6999
- Khoi Nguyen, Sinisa Todorovic:
iFS-RCNN: An Incremental Few-shot Instance Segmenter. 7000-7009
- Dat Huynh, Jason Kuen, Zhe Lin, Jiuxiang Gu, Ehsan Elhamifar:
Open-Vocabulary Instance Segmentation via Robust Cross-Modal Pseudo-Labeling. 7010-7021
- Xiaoqing Guo , Jie Liu , Tongliang Liu , Yixuan Yuan :
SimT: Handling Open-set Noise for Domain Adaptive Semantic Segmentation. 7022-7031
- Yahao Liu, Jinhong Deng, Jiale Tao, Tong Chu, Lixin Duan, Wen Li:
Undoing the Damage of Label Shift for Cross-domain Semantic Segmentation. 7032-7042
- Chang-Bin Zhang, Jia-Wen Xiao, Xialei Liu, Ying-Cong Chen, Ming-Ming Cheng :
Representation Compensation Networks for Continual Semantic Segmentation. 7043-7054
- Wenjian Wang, Lijuan Duan, Yuxi Wang, Qing En, Junsong Fan, Zhaoxiang Zhang:
Remember the Difference: Cross-Domain Few-Shot Semantic Segmentation via Meta-Memory Transfer. 7055-7064
- Xinyue Huo, Lingxi Xie, Hengtong Hu, Wengang Zhou, Houqiang Li, Qi Tian:
Domain-Agnostic Prior for Transfer Semantic Segmentation. 7065-7075
- Timo Lüdecke, Alexander S. Ecker:
Image Segmentation Using Text and Image Prompts. 7076-7086
- Xufeng Yao, Yang Bai, Xinyun Zhang, Yuechen Zhang, Qi Sun , Ran Chen, Ruiyu Li, Bei Yu:
PCL: Proxy-based Contrastive Learning for Domain Generalization. 7087-7097
- Wei Zhu, Le Lu, Jing Xiao, Mei Han, Jiebo Luo , Adam P. Harrison:
Localized Adversarial Domain Generalization. 7098-7108

- Chaoqi Chen, Jiongcheng Li, Xiaoguang Han, Xiaoqing Liu, Yizhou Yu:
Compound Domain Generalization via Meta-Knowledge Encoding. 7109-7119
- Juwon Kang, Sohyun Lee, Namyup Kim, Suha Kwak:
Style Neophile: Constantly Seeking Novel Styles for Domain Generalization. 7120-7130
- Rang Meng, Weijie Chen, Shicai Yang , Jie Song, Luojun Lin , Di Xie, Shiliang Pu, Xinchao Wang , Mingli Song, Yueling Zhuang:
Slimmable Domain Adaptation. 7131-7140
- Fan Wang, Zhongyi Han, Yongshun Gong, Yilong Yin:
Exploring Domain-Invariant Parameters for Source Free Domain Adaptation. 7141-7150
- Wei-Hong Li, Xialei Liu, Hakan Bilen :
Cross-domain Few-shot Learning with Task-specific Adapters. 7151-7160
- Shiyuan Huang, Jiawei Ma, Guangxing Han, Shih-Fu Chang:
Task-Adaptive Negative Envision for Few-Shot Open-Set Recognition. 7161-7170
- Lin Chen, Huaian Chen, Zhixiang Wei, Xin Jin, Xiao Tan, Yi Jin, Enhong Chen:
Reusing the Task-specific Classifier as a Discriminator: Discriminator-free Adversarial Domain Adaptation. 7171-7180
- Tao Sun, Cheng Lu, Tianshuo Zhang, Haibin Ling:
Safe Self-Refinement for Transformer-based Domain Adaptation. 7181-7190
- Qin Wang, Olga Fink , Luc Van Gool, Dengxin Dai:
Continual Test-Time Domain Adaptation. 7191-7201
- Ning Ding, Yixing Xu, Yehui Tang, Chao Xu, Yunhe Wang, Dacheng Tao:
Source-Free Domain Adaptation via Distribution Estimation. 7202-7212
- Yuefan Shen , Yanchao Yang, Mi Yan, He Wang, Youyi Zheng, Leonidas J. Guibas:
Domain Adaptation on Point Clouds via Geometry-Aware Implicit. 7213-7222
- Zhao Jin, Yinjie Lei, Naveed Akhtar , Haifeng Li, Munawar Hayat:
Deformation and Correspondence Aware Unsupervised Synthetic-to-Real Scene Flow Estimation for Point Clouds. 7223-7233
- Cheng Tan, Zhangyang Gao, Lirong Wu, Siyuan Li, Stan Z. Li:
Hyperspherical Consistency Regularization. 7234-7245
- Zhi Hou, Baosheng Yu, Dacheng Tao:
BatchFormer: Learning to Explore Sample Relationships for Robust Representation Learning. 7246-7256
- Rui Yu, Dawei Du, Rodney LaLonde, Daniel Davila, Christopher Funk , Anthony Hoogs, Brian Clipp:
Cascade Transformers for End-to-End Person Search. 7257-7266
- Chongzhi Zhang, Mingyuan Zhang, Shanghang Zhang, Daisheng Jin, Qiang Zhou, Zhongang Cai, Haiyu Zhao, Xianglong Liu, Ziwei Liu:
Delving Deep into the Generalization of Vision Transformers under Distribution Shifts. 7267-7276
- Youngwan Lee, Jonghee Kim, Jeffrey Willette, Sung Ju Hwang:
MPViT: Multi-Path Vision Transformer for Dense Prediction. 7277-7286
- Haochen Wang, Jiayi Shen, Yongtuo Liu, Yan Gao, Efstratios Gavves:
NFormer: Robust Person Re-identification with Neighbor Transformer. 7287-7297
- Yoonki Cho, Woo Jae Kim, Seunghoon Hong, Sung-Eui Yoon:
Part-based Pseudo Label Refinement for Unsupervised Person Re-identification. 7298-7308
- Wei Wu, Jiawei Liu, Kecheng Zheng, Qibin Sun, Zhengjun Zha:
Temporal Complementarity-Guided Reinforcement Learning for Image-to-Video Person Re-Identification. 7309-7318
- Yichen Lu, Mei Wang, Weihong Deng:
Augmented Geometric Distillation for Data-Free Incremental Person ReID. 7319-7328
- Shutao Bai, Bingpeng Ma, Hong Chang, Rui Huang, Xilin Chen:
Salient-to-Broad Transition for Video Person Re-identification. 7329-7338
- Qiang Zhang, Changzhou Lai, Jianan Liu, Nianchang Huang, Jungong Han:
FMCNet: Feature-Level Modality Compensation for Visible-Infrared Person Re-Identification. 7339-7348
- Shengcai Liao, Ling Shao:
Graph Sampling Based Deep Metric Learning for Generalizable Person Re-Identification. 7349-7358
- Xinyu Zhang , Dongdong Li, Zhigang Wang, Jian Wang, Errui Ding, Javen Qinfeng Shi , Zhaoxiang Zhang, Jingdong Wang:
Implicit Sample Extension for Unsupervised Person Re-Identification. 7359-7368
- Yibo Zhou:
Rethinking Reconstruction Autoencoder-Based Out-of-Distribution Detection. 7369-7377
- Choubo Ding, Guansong Pang , Chunhua Shen:
Catching Both Gray and Black Swans: Open-set Supervised Anomaly Detection. 7378-7388
- Xuhui Yang, Yaowei Wang, Ke Chen, Yong Xu, Yonghong Tian:
Fine-Grained Object Classification via Self-Supervised Pose Alignment. 7389-7398

- Aleksandr Ermolov, Leyla Mirvakhanova, Valentin Khrulkov, Nicu Sebe , Ivan V. Oseledets:
Hyperbolic Vision Transformers: Combining Improvements in Metric Learning. 7399-7409
- Karsten Roth, Oriol Vinyals, Zeynep Akata:
Non-isotropy Regularization for Proxy-based Deep Metric Learning. 7410-7420
- Sungyeon Kim, Dongwon Kim, Minsu Cho, Suha Kwak:
Self-Taught Metric Learning without Labels. 7421-7431
- Yanan Gu , Xu Yang, Kun Wei, Cheng Deng:
Not Just Selection, but Exploration: Online Class-Incremental Continual Learning via Dual View Consistency. 7432-7441
- K. J. Joseph, Salman Khan, Fahad Shahbaz Khan, Rao Muhammad Anwer, Vineeth N. Balasubramanian:
Energy-based Latent Aligner for Incremental Learning. 7442-7451
- Aneeshan Sain , Ayan Kumar Bhunia, Vaishnav Potlapalli, Pinaki Nath Chowdhury, Tao Xiang, Yi-Zhe Song :
Sketch3T: Test-Time Training for Zero-Shot SBIR. 7452-7461
- Ronghan Chen, Yang Cong:
The Devil is in the Pose: Ambiguity-free 3D Rotation-invariant Learning via Pose-aware Convolution. 7462-7471
- Lan Yang, Kaiyue Pang, Honggang Zhang, Yi-Zhe Song:
Finding Badly Drawn Bunnies. 7472-7481
- Sagar Vaze, Kai Han, Andrea Vedaldi, Andrew Zisserman:
Generalized Category Discovery. 7482-7491
- Yash Patel, Giorgos Tolias, Jirí Matas :
Recall@k Surrogate Loss with Large Batches and Similarity Mixup. 7492-7501
- Quan Zhang, Kaiheng Dang, Jian-Huang Lai, Zhan-Xiang Feng, Xiaohua Xie:
Modeling 3D Layout For Group Re-Identification. 7502-7510
- Chengzhi Mao, Kevin Xia, James Wang, Hao Wang, Junfeng Yang, Elias Bareinboim, Carl Vondrick:
Causal Transportability for Visual Recognition. 7511-7521
- Borui Zhang, Wenzhao Zheng, Jie Zhou, Jiwen Lu:
Attributable Visual Similarity Learning. 7522-7531
- Shenjian Gong, Shanshan Zhang, Jian Yang, Dengxin Dai, Bernt Schiele :
Bi-level Alignment for Cross-Domain Crowd Counting. 7532-7540
- Erkun Yang, Dongren Yao, Tongliang Liu , Cheng Deng:
Mutual Quantization for Cross-Modal Search with Noisy Labels. 7541-7550
- Matthew Wallingford, Hao Li, Alessandro Achille, Avinash Ravichandran, Charless C. Fowlkes, Rahul Bhotika, Stefano Soatto:
Task Adaptive Parameter Sharing for Multi-Task Learning. 7551-7560
- Xingyi Zhou, Vladlen Koltun, Philipp Krähenbühl:
Simple Multi-dataset Detection. 7561-7570
- Yu-Jhe Li, Xiaoliang Dai, Chih-Yao Ma, Yen-Cheng Liu, Kan Chen, Bichen Wu, Zijian He, Kris Kitani, Peter Vajda:
Cross-Domain Adaptive Teacher for Object Detection. 7571-7580
- Hanqing Yang, Sijia Cai, Hualian Sheng, Bing Deng, Jianqiang Huang, Xian-Sheng Hua, Yong Tang, Yu Zhang:
Balanced and Hierarchical Relation Learning for One-shot Object Detection. 7581-7590
- Yizhou Zhao, Xun Guo, Yan Lu:
Semantic-aligned Fusion Transformer for One-shot Object Detection. 7591-7601
- Shiming Chen, Ziming Hong, Guo-Sen Xie, Wenhan Yang, Qinmu Peng, Kai Wang, Jian Zhao, Xinge You:
MSDN: Mutually Semantic Distillation Network for Zero-Shot Learning. 7602-7611
- Peiliang Huang, Junwei Han, De Cheng, Dingwen Zhang:
Robust Region Feature Synthesizer for Zero-Shot Object Detection. 7612-7621
- Chao Xu, Jiangning Zhang, Miao Hua, Qian He, Zili Yi, Yong Liu:
Region-Aware Face Swapping. 7622-7631
- Yangyang Xu, Bailin Deng, Junle Wang, Yanqing Jing, Jia Pan, Shengfeng He :
High-resolution Face Swapping via Latent Semantics Disentanglement. 7632-7641
- Yang Zhao, Yu-Chuan Su, Chun-Te Chu, Yandong Li, Marius Renn, Yukun Zhu, Changyou Chen, Xuhui Jia:
Rethinking Deep Face Restoration. 7642-7651
- Feida Zhu, Junwei Zhu, Wenqing Chu, Xinyi Zhang, Xiaozhong Ji, Chengjie Wang, Ying Tai:
Blind Face Restoration via Integrating Face Shape and Generative Priors. 7652-7661
- Jingxiang Sun, Xuan Wang, Yong Zhang, Xiaoyu Li, Qi Zhang, Yebin Liu, Jue Wang:
FENeRF: Face Editing in Neural Radiance Fields. 7662-7672
- Yanbo Xu, Yueqin Yin, Liming Jiang, Qianyi Wu, Chengyao Zheng, Chen Change Loy, Bo Dai, Wayne Wu:
TransEditor: Transformer-Based Dual-Space GAN for Highly Controllable Facial Editing. 7673-7682

- Shuai Yang, Liming Jiang, Ziwei Liu, Chen Change Loy:
Pastiche Master: Exemplar-Based High-Resolution Portrait Style Transfer. 7683-7692
- Zijian Wang, Xingqun Qi, Kun Yuan, Muyi Sun:
Self-supervised Correlation Mining Network for Person Image Generation. 7693-7702
- Pengze Zhang, Lingxiao Yang, Jianhuang Lai, Xiaohua Xie:
Exploring Dual-task Correlation for Pose Guided Person Image Generation. 7703-7712
- Anna Frühstück , Krishna Kumar Singh, Eli Shechtman, Niloy J. Mitra, Peter Wonka, Jingwan Lu:
InsetGAN for Full-Body Image Generation. 7713-7722
- Chaojie Yang, Hanhui Li, Shengjie Wu, Shengkai Zhang, Haonan Yan, Nianhong Jiao, Jie Tang, Runnan Zhou, Xiaodan Liang, Tianxiang Zheng:
BodyGAN: General-purpose Controllable Neural Human Body Generation. 7723-7732
- Fuqiang Zhao, Wei Yang, Jiakai Zhang , Pei Lin, Yingliang Zhang , Jingyi Yu, Lan Xu:
HumanNeRF: Efficiently Generated Human Radiance Field from Sparse Inputs. 7733-7743
- Jianqiang Ren, Yuan Yao, Biwen Lei, Miaomiao Cui, Xuansong Xie:
Structure-Aware Flow Generation for Human Body Reshaping. 7744-7753
- Zuopeng Yang, Daqing Liu, Chaoyue Wang, Jie Yang, Dacheng Tao:
Modeling Image Composition for Complex Scene Generation. 7754-7763
- Sang-Heon Shim , Sangeek Hyun, Dae Hyun Bae, Jae-Pil Heo:
Local Attention Pyramid for Scene Image Generation. 7764-7772
- Bo Wang, Tao Wu, Minfeng Zhu , Peng Du:
Interactive Image Synthesis with Panoptic Layout Generation. 7773-7782
- Feixiang He, Yanlong Huang, He Wang:
iPLAN: Interactive and Procedural Layout Planning. 7783-7792
- Chen Song, Qixing Huang, Chandrajit Bajaj:
E-CIR: Event-Enhanced Continuous Intensity Recovery. 7793-7802
- Yuqi Sun, Shili Zhou, Ri Cheng, Weimin Tan, Bo Yan, Lang Fu:
Learning Robust Image-Based Rendering on Sparse Scene Geometry via Depth Completion. 7803-7813
- Yuan Liu, Sida Peng, Lingjie Liu, Qianqian Wang, Peng Wang, Christian Theobalt , Xiaowei Zhou, Wenping Wang:
Neural Rays for Occlusion-aware Image-based Rendering. 7814-7823
- Jinchao Yang, Fei Guo, Shuo Chen, Jun Li, Jian Yang:
Industrial Style Transfer with Large-scale Geometric Warping and Content Preservation. 7824-7833
- Tai-Yin Chiu, Danna Gurari:
PCA-Based Knowledge Distillation Towards Lightweight and Content-Style Balanced Photorealistic Style Transfer Models. 7834-7843
- Kyungjune Baek, Hyunjung Shim:
Commonality in Natural Images Rescues GANs: Pretraining GANs with Generic and Privacy-free Synthetic Data. 7844-7854
- Chengdong Dong, Ajay Kumar, Eryun Liu:
Think Twice Before Detecting GAN-generated Fake Images from their Spectral Domain Imprints. 7855-7864
- Youmin Xu, Chong Mou, Yujie Hu, Jingfen Xie, Jian Zhang:
Robust Invertible Image Steganography. 7865-7874
- Hongzu Su, Jingjing Li, Zhi Chen , Lei Zhu, Ke Lu:
Distinguishing Unseen from Seen for Generalized Zero-shot Learning. 7875-7884
- Licheng Tang, Yiyang Cai, Jiaming Liu, Zhibin Hong, Mingming Gong, Minhu Fan, Junyu Han, Jingtuo Liu, Errui Ding, Jingdong Wang:
Few-Shot Font Generation by Learning Fine-Grained Local Styles. 7885-7894
- Wei Liu, Fangyue Liu, Fei Ding, Qian He, Zili Yi:
XMP-Font: Self-Supervised Cross-Modality Pre-training for Few-Shot Font Generation. 7895-7904
- Caroline Chan, Frédo Durand, Phillip Isola:
Learning to generate line drawings that convey geometry and semantics. 7905-7915
- Jiawei Ren, Mingyuan Zhang, Cunjun Yu, Ziwei Liu:
Balanced MSE for Imbalanced Visual Regression. 7916-7925
- Andrea Agostinelli, Jasper R. R. Uijlings, Thomas Mensink, Vittorio Ferrari:
Transferability Metrics for Selecting Source Model Ensembles. 7926-7936
- Nanyang Ye, Kaican Li, Haoyue Bai, Runpeng Yu, Lanqing Hong, Fengwei Zhou, Zhenguo Li, Jun Zhu:
OoD-Bench: Quantifying and Understanding Two Dimensions of Out-of-Distribution Generalization. 7937-7948
- Mitchell Wortsman, Gabriel Ilharco, Jong Wook Kim, Mike Li, Simon Kornblith, Rebecca Roelofs, Raphael Gontijo Lopes, Hannaneh Hajishirzi, Ali Farhadi, Hongseok Namkoong, Ludwig Schmidt:
Robust fine-tuning of zero-shot models. 7949-7961
- Jiangtao Xie , Fei Long, Jiaming Lv, Qilong Wang, Peihua Li:
Joint Distribution Matters: Deep Brownian Distance Covariance for Few-Shot Classification. 7962-7971

- Zhenyi Wang, Li Shen, Tiehang Duan, Donglin Zhan, Le Fang, Mingchen Gao:
Learning to Learn and Remember Super Long Multi-Domain Task Sequence. 7972-7982
- Ming Xie, Yuxi Li, Yabiao Wang , Zekun Luo, Zhenye Gan, Zhongyi Sun, Mingmin Chi, Chengjie Wang, Pei Wang:
Learning Distinctive Margin toward Active Domain Adaptation. 7983-7992
- Jian Liang, Dapeng Hu, Jiashi Feng, Ran He:
DINE: Domain Adaptation from Single and Multiple Black-box Predictors. 7993-8003
- Shuaifeng Li , Mao Ye, Xiatian Zhu, Lihua Zhou, Lin Xiong:
Source-Free Object Detection by Learning to Overlook Domain Style. 8004-8013
- Hanlin Zhang, Yifan Zhang, Weiyang Liu, Adrian Weller, Bernhard Schölkopf, Eric P. Xing:
Towards Principled Disentanglement for Domain Generalization. 8014-8024
- Yabin Zhang, Minghan Li , Ruihuang Li, Kui Jia, Lei Zhang:
Exact Feature Distribution Matching for Arbitrary Style Transfer and Domain Generalization. 8025-8035
- Fangrui Lv, Jian Liang, Shuang Li, Bin Zang, Chi Harold Liu, Ziteng Wang, Di Liu:
Causality Inspired Representation Learning for Domain Generalization. 8036-8046
- Chunbo Lang , Gong Cheng , Binfei Tu, Junwei Han:
Learning What Not to Segment: A New Perspective on Few-Shot Segmentation. 8047-8057
- Binhui Xie, Longhui Yuan, Shuang Li, Chi Harold Liu, Xinjing Cheng:
Towards Fewer Annotations: Active Learning via Region Impurity and Prediction Uncertainty for Domain Adaptive Semantic Segmentation. 8058-8068
- Hanxiang Ren, Yanchao Yang, He Wang, Bokui Shen, Qingnan Fan, Youyi Zheng, C. Karen Liu, Leonidas J. Guibas:
ADeLA: Automatic Dense Labeling with Attention for Viewpoint Shift in Semantic Segmentation. 8069-8079
- Jiarui Cai, Mingze Xu, Wei Li, Yuanjun Xiong, Wei Xia, Zhuowen Tu, Stefano Soatto:
MeMOT: Multi-Object Tracking with Memory. 8080-8090
- QiuHong Shen, Lei Qiao, Jinyang Guo, Peixia Li, Xin Li, Bo Li, Weitao Feng, Weihao Gan, Wei Wu, Wanli Ouyang:
Unsupervised Learning of Accurate Siamese Tracking. 8091-8100
- Chaoda Zheng, Xu Yan, Haiming Zhang, Baoyuan Wang, Shenghui Cheng, Shuguang Cui , Zhen Li:
Beyond 3D Siamese Tracking: A Motion-Centric Paradigm for 3D Single Object Tracking in Point Clouds. 8101-8110
- Haofei Xu, Jing Zhang, Jianfei Cai, Hamid Rezatofighi, Dacheng Tao:
GMFlow: Learning Optical Flow via Global Matching. 8111-8120
- Abhishek Kumar , Oladayo S. Ajani , Swagatam Das, Rammohan Mallipeddi:
GridShift: A Faster Mode-seeking Algorithm for Image Segmentation and Object Tracking. 8121-8129
- Igor Santesteban , Miguel A. Otaduy, Dan Casas :
SNUG: Self-Supervised Neural Dynamic Garments. 8130-8140
- Wei Mao, Miaomiao Liu, Mathieu Salzmann:
Weakly-supervised Action Transition Learning for Stochastic Human Motion Prediction. 8141-8150
- Hengbo Ma, Jiachen Li, Ramtin Hosseini, Masayoshi Tomizuka, Chiho Choi:
Multi-Objective Diverse Human Motion Prediction with Knowledge Distillation. 8151-8161
- Taein Kwon , Bugra Tekin, Siyu Tang , Marc Pollefeys :
Context-Aware Sequence Alignment using 4D Skeletal Augmentation. 8162-8172
- Lachlan E. MacDonald, Sameera Ramasinghe, Simon Lucey :
Enabling Equivariance for Arbitrary Lie Groups. 8173-8182
- Ahmed Abbas , Paul Swoboda:
RAMA: A Rapid Multicut Algorithm on GPU. 8183-8192
- Peri Akiva, Matthew Purri, Matthew J. Leotta:
Self-Supervised Material and Texture Representation Learning for Remote Sensing Tasks. 8193-8205
- Xiaodong Gu, Chengzhou Tang, Weihao Yuan, Zuozhuo Dai, Siyu Zhu, Ping Tan:
RCP: Recurrent Closest Point for Point Cloud. 8206-8216
- Karren Yang, Dejan Markovic, Steven Krenn, Vasu Agrawal, Alexander Richard:
Audio-Visual Speech Codecs: Rethinking Audio-Visual Speech Enhancement by Re-Synthesis. 8217-8227
- Xiaokang Peng, Yake Wei, Andong Deng, Dong Wang, Di Hu:
Balanced Multimodal Learning via On-the-fly Gradient Modulation. 8228-8237
- Matthew Tancik, Vincent Casser, Xinchen Yan, Sabeek Pradhan, Ben P. Mildenhall, Pratul P. Srinivasan, Jonathan T. Barron, Henrik Kretzschmar:
Block-NeRF: Scalable Large Scene Neural View Synthesis. 8238-8248
- Luwei Yang, Rakesh Shrestha, Wenbo Li, Shuaicheng Liu , Guofeng Zhang, Zhaopeng Cui, Ping Tan:
SceneSqueezer: Learning to Compress Scene for Camera Relocalization. 8249-8258
- Mohammed Suhail, Carlos Esteves, Leonid Sigal, Ameesh Makadia:
Light Field Neural Rendering. 8259-8269

- Jacob Munkberg, Wenzheng Chen, Jon Hasselgren, Alex Evans, Tianchang Shen, Thomas Müller, Jun Gao, Sanja Fidler: **Extracting Triangular 3D Models, Materials, and Lighting From Images.** 8270-8280
- Marc Alexa: **Super-Fibonacci Spirals: Fast, Low-Discrepancy Sampling of SO(3).** 8281-8290
- Feng Cheng, Mingze Xu, Yuanjun Xiong, Hao Chen, Xinyu Li, Wei Li, Wei Xia: **Stochastic Backpropagation: A Memory Efficient Strategy for Training Video Models.** 8291-8300
- Kanghyun Choi, Hyeyoon Lee, Deokki Hong, Joonsang Yu, Noseong Park, Youngsok Kim, Jinho Lee: **It's All In the Teacher: Zero-Shot Quantization Brought Closer to the Teacher.** 8301-8311
- Fawaz Sammani, Tanmoy Mukherjee, Nikos Deligiannis: **NLX-GPT: A Model for Natural Language Explanations in Vision and Vision-Language Tasks.** 8312-8322
- Yu Yang, Seungbae Kim, Jungseock Joo: **Explaining Deep Convolutional Neural Networks via Latent Visual-Semantic Filter Attention.** 8323-8333
- Malik Boudiaf, Romain Müller, Ismail Ben Ayed, Luca Bertinetto: **Parameter-free Online Test-time Adaptation.** 8334-8343
- Sukmin Yun, Hankook Lee, Jaehyung Kim, Jinwoo Shin: **Patch-level Representation Learning for Self-supervised Vision Transformers.** 8344-8353
- Luke Melas-Kyriazi, Christian Rupprecht, Iro Laina, Andrea Vedaldi: **Deep Spectral Methods: A Surprisingly Strong Baseline for Unsupervised Semantic Segmentation and Localization.** 8354-8365
- Aditya Golatkar, Alessandro Achille, Yu-Xiang Wang, Aaron Roth, Michael Kearns, Stefano Soatto: **Mixed Differential Privacy in Computer Vision.** 8366-8376
- Jia-Wei Chen, Chia-Mu Yu, Ching-Chia Kao, Tzai-Wei Pang, Chun-Shien Lu: **DPGEN: Differentially Private Generative Energy-Guided Network for Natural Image Synthesis.** 8377-8386
- Matías Mendieta, Taojiannan Yang, Pu Wang, Minwoo Lee, Zhengming Ding, Chen Chen: **Local Learning Matters: Rethinking Data Heterogeneity in Federated Learning.** 8387-8396
- Nikhil Varma Keetha, Chen Wang, Yuheng Qiu, Kuan Xu, Sebastian A. Scherer: **AirObject: A Temporally Evolving Graph Embedding for Object Identification.** 8397-8406
- Chenhang He, Ruihuang Li, Shuai Li, Lei Zhang: **Voxel Set Transformer: A Set-to-Set Approach to 3D Object Detection from Point Clouds.** 8407-8417
- Chuandong Liu, Chenqiang Gao, Fangcen Liu, Jiang Liu, Deyu Meng, Xinbo Gao: **SS3D: Sparsely-Supervised 3D Object Detection from Point Cloud.** 8418-8427
- Xiuwei Xu, Yifan Wang, Yu Zheng, Yongming Rao, Jie Zhou, Jiwen Lu: **Back to Reality: Weakly-supervised 3D Object Detection with Shape-guided Label Enhancement.** 8428-8437
- Shengheng Deng, Zhihao Liang, Lin Sun, Kui Jia: **VISTA: Boosting 3D Object Detection via Dual Cross-View SpaTial Attention.** 8438-8447
- Lue Fan, Ziqi Pang, Tianyuan Zhang, Yu-Xiong Wang, Hang Zhao, Feng Wang, Naiyan Wang, Zhaoxiang Zhang: **Embracing Single Stride 3D Object Detector with Sparse Transformer.** 8448-8458
- Jordan S. K. Hu, Tianshu Kuai, Steven L. Waslander: **Point Density-Aware Voxels for LiDAR 3D Object Detection.** 8459-8468
- Yuenan Hou, Xinge Zhu, Yuexin Ma, Chen Change Loy, Yikang Li: **Point-to-Voxel Knowledge Distillation for LiDAR Semantic Segmentation.** 8469-8478
- Liyao Tang, Yibing Zhan, Zhe Chen, Baosheng Yu, Dacheng Tao: **Contrastive Boundary Learning for Point Cloud Segmentation.** 8479-8489
- Xin Lai, Jianhui Liu, Li Jiang, Liwei Wang, Hengshuang Zhao, Shu Liu, Xiaojuan Qi, Jiaya Jia: **Stratified Transformer for 3D Point Cloud Segmentation.** 8490-8499
- Jia-Xing Zhong, Kaichen Zhou, Qingyong Hu, Bing Wang, Niki Trigoni, Andrew Markham: **No Pain, Big Gain: Classify Dynamic Point Cloud Sequences with Static Models by Fitting Feature-level Space-time Surfaces.** 8500-8510
- Yujing Xue, Jiageng Mao, Minzhe Niu, Hang Xu, Michael Bi Mi, Wei Zhang, Xiaogang Wang, Xinchao Wang: **Point2Seq: Detecting 3D Objects as Sequences.** 8511-8520
- Changqing Zhou, Zhipeng Luo, Yueru Luo, Tianrui Liu, Liang Pan, Zhongang Cai, Haiyu Zhao, Shijian Lu: **PTTR: Relational 3D Point Cloud Object Tracking with Transformer.** 8521-8530
- Zetong Yang, Li Jiang, Yanan Sun, Bernt Schiele, Jiaya Jia: **A Unified Query-based Paradigm for Point Cloud Understanding.** 8531-8541
- Renrui Zhang, Ziyu Guo, Wei Zhang, Kunchang Li, Xupeng Miao, Bin Cui, Yu Qiao, Peng Gao, Hongsheng Li: **PointCLIP: Point Cloud Understanding by CLIP.** 8542-8552
- Zhihao Yuan, Xu Yan, Yinghong Liao, Yao Guo, Guanbin Li, Shuguang Cui, Zhen Li: **X -Trans2Cap: Cross-Modal Knowledge Transfer using Transformer for 3D Dense Captioning.** 8553-8563
- Zhenpei Yang, Zhile Ren, Qi Shan, Qixing Huang: **MVS2D: Efficient Multiview Stereo via Attention-Driven 2D Convolutions.** 8564-8574

- Yikang Ding, Wentao Yuan, Qingtian Zhu, Haotian Zhang, Xiangyue Liu, Yuanjiang Wang, Xiao Liu:
TransMVSNet: Global Context-aware Multi-view Stereo Network with Transformers. 8575-8584
- Junhua Xi, Yifei Shi, Yijie Wang, Yulan Guo, Kai Xu:
RayMVSNet: Learning Ray-based 1D Implicit Fields for Accurate Multi-View Stereo. 8585-8595
- Fangjinhua Wang, Silvano Galliani, Christoph Vogel, Marc Pollefeys :
IterMVS: Iterative Probability Estimation for Efficient Multi-View Stereo. 8596-8605
- Haiyan Wang, Will Hutchcroft, Yuguang Li, Zhiqiang Wan, Ivaylo Boyadzhiev, Yingli Tian , Sing Bing Kang:
PSMNet: Position-aware Stereo Merging Network for Room Layout Estimation. 8606-8615
- Jiayu Yang , José M. Álvarez, Miaomiao Liu:
Non-parametric Depth Distribution Modelling based Depth Inference for Multi-view Stereo. 8616-8624
- Shubham Goel, Georgia Gkioxari, Jitendra Malik:
Differentiable Stereopsis: Meshes from multiple views using differentiable rendering. 8625-8634
- Rui Peng, Rongjie Wang, Zhenyu Wang, Yawen Lai, Ronggang Wang:
Rethinking Depth Estimation for Multi-View Stereo: A Unified Representation. 8635-8644
- Shaoqian Wang, Bo Li , Yuchao Dai:
Efficient Multi-view Stereo by Iterative Dynamic Cost Volume. 8645-8654
- Jiachen Liu, Pan Ji, Nitin Bansal, Changjiang Cai , Qingan Yan, Xiaolei Huang, Yi Xu:
PlaneMVS: 3D Plane Reconstruction from Multi-View Stereo. 8655-8665
- Kaixuan Zhang, Kaiwei Che, Jianguo Zhang, Jie Cheng, Ziyang Zhang, Qinghai Guo, Liziwei Leng:
Discrete time convolution for fast event-based stereo. 8666-8676
- Taras Khakhulin, Denis Korzhenkov, Pavel Solovev, Gleb Sterkin, Andrei-Timotei Ardelean, Victor Lempitsky:
Stereo Magnification with Multi-Layer Images. 8677-8686
- Seungwook Kim, Juhong Min, Minsu Cho:
TransformMatcher: Match-to-Match Attention for Semantic Correspondence. 8687-8697
- Prune Truong, Martin Danelljan, Fisher Yu, Luc Van Gool:
Probabilistic Warp Consistency for Weakly-Supervised Semantic Correspondences. 8698-8708
- Liulei Li, Tianfei Zhou, Wenguan Wang , Lu Yang, Jianwu Li, Yi Yang:
Locality-Aware Inter-and Intra-Video Reconstruction for Self-Supervised Correspondence Learning. 8709-8720
- Christoph Mayer, Martin Danelljan, Goutam Bhat, Matthieu Paul, Danda Pani Paudel , Fisher Yu, Luc Van Gool:
Transforming Model Prediction for Tracking. 8721-8730
- Feng Tang, Qiang Ling:
Ranking-Based Siamese Visual Tracking. 8731-8740
- Fei Xie, Chunyu Wang, Guangting Wang, Yue Cao, Wankou Yang, Wenjun Zeng :
Correlation-Aware Deep Tracking. 8741-8750
- Zikun Zhou, Jianqiu Chen, Wenjie Pei, Kaige Mao, Hongpeng Wang, Zhenyu He:
Global Tracking via Ensemble of Local Trackers. 8751-8760
- Xingyi Zhou, Tianwei Yin, Vladlen Koltun, Philipp Krähenbühl:
Global Tracking Transformers. 8761-8770
- Fan Ma, Mike Zheng Shou, Linchao Zhu, Haoqi Fan, Yilei Xu, Yi Yang, Zhicheng Yan:
Unified Transformer Tracker for Object Tracking. 8771-8780
- Zikai Song, Junqing Yu, Yi-Ping Phoebe Chen , Wei Yang:
Transformer Tracking with Cyclic Shifting Window Attention. 8781-8790
- Jiqing Zhang, Bo Dong, Haiwei Zhang, Jianchuan Ding, Felix Heide, Baocai Yin, Xin Yang:
Spiking Transformers for Event-based Single Object Tracking. 8791-8800
- Jan-Nico Zaech, Alexander Liniger, Martin Danelljan, Dengxin Dai, Luc Van Gool:
Adiabatic Quantum Computing for Multi Object Tracking. 8801-8812
- Zikang Zhou , Luyao Ye , Jianping Wang, Kui Wu, Kejie Lu:
HiVT: Hierarchical Vector Transformer for Multi-Agent Motion Prediction. 8813-8823
- En Yu, Zhuoling Li, Shoudong Han:
Towards Discriminative Representation: Multi-view Trajectory Contrastive Learning for Online Multi-object Tracking. 8824-8833
- Tim Meinhardt, Alexander Kirillov, Laura Leal-Taixé, Christoph Feichtenhofer:
TrackFormer: Multi-Object Tracking with Transformers. 8834-8844
- Shuai Li, Yu Kong, Hamid Rezatofighi:
Learning of Global Objective for Network Flow in Multi-Object Tracking. 8845-8855
- Duy M. H. Nguyen, Roberto Henschel, Bodo Rosenhahn, Daniel Sonntag, Paul Swoboda:
LMGP: Lifted Multicut Meets Geometry Projections for Multi-Camera Multi-Object Tracking. 8856-8865
- Shuai Liu, Xin Li, Huchuan Lu, You He:
Multi-Object Tracking Meets Moving UAV. 8866-8875

- Pengyu Zhang, Jie Zhao, Dong Wang, Huchuan Lu, Xiang Ruan:
Visible-Thermal UAV Tracking: A Large-Scale Benchmark and New Baseline. 8876-8885
- Junjie Ye , Changhong Fu, Guangze Zheng , Danda Pani Paudel , Guang Chen:
Unsupervised Domain Adaptation for Nighttime Aerial Tracking. 8886-8895
- Ao Luo, Fan Yang, Xin Li, Shuaicheng Liu :
Learning Optical Flow with Kernel Patch Attention. 8896-8905
- Simon Schrodi, Tonmoy Saikia, Thomas Brox:
Towards Understanding Adversarial Robustness of Optical Flow Networks. 8906-8914
- Zihua Zheng, Ni Nie, Zhi Ling, Pengfei Xiong, Jiangyu Liu, Hao Wang, Jiankun Li:
DIP: Deep Inverse Patchmatch for High-Resolution Optical Flow. 8915-8924
- Hongyi Fan, Joe Kileel, Benjamin B. Kimia:
On the Instability of Relative Pose Estimation and RANSAC's Role. 8925-8933
- Haofei Zhang , Jiarui Duan, Mengqi Xue , Jie Song, Li Sun, Mingli Song:
Bootstrapping ViTs: Towards Liberating Vision Transformers from Pre-training. 8934-8943
- Zi-En Fan, Feng Lian, Jia-Ni Quan:
Global Sensing and Measurements Reuse for Image Compressed Sensing. 8944-8953
- Erchuan Zhang , David Suter , Ruwan B. Tennakoon , Tat-Jun Chin, Alireza Bab-Hadiashar, Giang Truong, Syed Zulqarnain Gilani :
Maximum Consensus by Weighted Influences of Monotone Boolean Functions. 8954-8962
- Luanyuan Dai, Yizhang Liu, Jiayi Ma, Lifang Wei, Taotao Lai, Changcai Yang, Riqing Chen:
MS2DG-Net: Progressive Correspondence Learning via Multiple Sparse Semantics Dynamic Graph. 8963-8972
- Jeeseung Park, Younggeun Kim:
Styleformer: Transformer based Generative Adversarial Networks with Style Vector. 8973-8982
- Fang Bai, Agniva Sengupta, Adrien Bartoli:
Scanline Homographies for Rolling-Shutter Plane Absolute Pose. 8983-8992
- Jingyi Xu, Hieu Le:
Generating Representative Samples for Few-Shot Classification. 8993-9003
- Arman Afrasiyabi, Hugo Larochelle, Jean-François Lalonde, Christian Gagné :
Matching Feature Sets for Few-Shot Image Classification. 9004-9014
- Junhao Dong , Yuan Wang, Jianhuang Lai, Xiaohua Xie:
Improving Adversarially Robust Few-shot Image Classification with Generalizable Representations. 9015-9024
- Li Yin, Juan M. Perez-Rua, Kevin J. Liang:
Sylph: A Hypernetwork Framework for Incremental Few-shot Object Detection. 9025-9035
- Da-Wei Zhou, Fu-Yun Wang, Han-Jia Ye, Liang Ma , Shiliang Pu, De-Chuan Zhan:
Forward Compatible Few-Shot Class-Incremental Learning. 9036-9046
- Michael Hersche, Geethan Karunaratne, Giovanni Cherubini, Luca Benini, Abu Sebastian, Abbas Rahimi :
Constrained Few-shot Class-incremental Learning. 9047-9057
- Shell Xu Hu, Da Li, Jan Stühmer, Minyoung Kim, Timothy M. Hospedales:
Pushing the Limits of Simple Pipelines for Few-Shot Learning: External Data and Fine-Tuning Make a Difference. 9058-9067
- Hao Zhu, Piotr Koniusz:
EASE: Unsupervised Discriminant Subspace Learning for Transductive Few-Shot Learning. 9068-9078
- Kevin J. Liang, Samrudhdhi B. Rangrej, Vladan Petrovic, Tal Hassner:
Few-shot Learning with Noisy Labels. 9079-9088
- Pan Li, Shaogang Gong, Chengjie Wang, Yanwei Fu :
Ranking Distance Calibration for Cross-Domain Few-Shot Learning. 9089-9098
- Moslem Yazdanpanah, Aamer Abdul Rahman, Muawiz Chaudhary, Christian Desrosiers, Mohammad Havaei, Eugene Belilovsky, Samira Ebrahimi Kahou:
Revisiting Learnable Affines for Batch Norm in Few-Shot Transfer Learning. 9099-9108
- Yangji He, Weihan Liang, Dongyang Zhao, Hong-Yu Zhou , Weifeng Ge, Yizhou Yu, Wenqiang Zhang:
Attribute Surrogates Learning and Spectral Tokens Pooling in Transformers for Few-shot Learning. 9109-9119
- Yu Xie, Yanwei Fu , Ying Tai, Yun Cao, Junwei Zhu, Chengjie Wang:
Learning to Memorize Feature Hallucination for One-Shot Image Generation. 9120-9129
- Yunqing Zhao, Henghui Ding , Houjing Huang, Ngai-Man Cheung:
A Closer Look at Few-shot Image Generation. 9130-9140
- Jiamin Wu, Tianzhu Zhang, Zhe Zhang, Feng Wu, Yongdong Zhang:
Motion-modulated Temporal Fragment Alignment Network For Few-Shot Action Recognition. 9141-9150
- Ruifei He, Shuyang Sun, Jihan Yang, Song Bai, Xiaojuan Qi:
Knowledge Distillation as Efficient Pre-training: Faster Convergence, Higher Data-efficiency, and Better Transferability. 9151-9161

- Michal Pndy, Andrea Agostinelli, Jasper R. R. Uijlings, Vittorio Ferrari, Thomas Mensink:
Transferability Estimation using Bhattacharyya Class Separability. 9162-9172
- Yizhou Wang, Shixiang Tang, Feng Zhu, Lei Bai , Rui Zhao, Donglian Qi, Wanli Ouyang:
Revisiting the Transferability of Supervised Pretraining: an MLP Perspective. 9173-9183
- Samarth Mishra, Rameswar Panda, Cheng Perng Phoo, Chun-Fu Richard Chen, Leonid Karlinsky, Kate Saenko, Venkatesh Saligrama , Rogrio Schmidt Feris:
Task2Sim: Towards Effective Pre-training and Transfer from Synthetic Data. 9184-9194
- Cdric Renggli, Andr Susano Pinto, Luka Rimanic, Joan Puigcerver, Carlos Riquelme, Ce Zhang, Mario Lucic:
Which Model to Transfer? Finding the Needle in the Growing Haystack. 9195-9204
- Yutaro Yamada, Mayu Otani:
Does Robustness on ImageNet Transfer to Downstream Tasks? 9205-9214
- Christos Matsoukas, Johan Fredin Haslum, Moein Sorkhei, Magnus Sderberg, Kevin Smith:
What Makes Transfer Learning Work for Medical Images: Feature Reuse & Other Factors. 9215-9224
- Akshita Gupta, Sanath Narayan, K. J. Joseph, Salman Khan, Fahad Shahbaz Khan, Mubarak Shah:
OW-DETR: Open-world Detection Transformer. 9225-9234
- Hari Chandana Kuchibhotla , Sumitra S. Malagi, Shivam Chandhok, Vineeth N. Balasubramanian:
Unseen Classes at a Later Time? No Problem. 9235-9244
- Binbin Yang, Xinch Deng, Han Shi, Changlin Li, Gengwei Zhang , Hang Xu, Shen Zhao, Liang Lin, Xiaodan Liang:
Continual Object Detection via Prototypical Task Correlation Guided Gating Mechanism. 9245-9254
- Christian Simon, Masoud Faraki, Yi-Hsuan Tsai, Xiang Yu, Samuel Schulter, Yumin Suh, Mehrtash Harandi, Manmohan Chandraker:
On Generalizing Beyond Domains in Cross-Domain Continual Learning. 9255-9264
- Jihwan Bang, Hyunseo Koh, Seulki Park, Hwanjun Song, Jung-Woo Ha , Jonghyun Choi :
Online Continual Learning on a Contaminated Data Stream with Blurry Task Boundaries. 9265-9274
- Arthur Douillard, Alexandre Ram, Guillaume Couairon, Matthieu Cord:
DyTox: Transformers for Continual Learning with DYnamic TOken eXpansion. 9275-9285
- Kai Zhu, Wei Zhai, Yang Cao, Jiebo Luo , Zhengjun Zha:
Self-Sustaining Representation Expansion for Non-Exemplar Class-Incremental Learning. 9286-9295
- Xia Kong, Zuodong Gao, Xiaofan Li, Ming Hong, Jun Liu, Chengjie Wang, Yuan Xie, Yanyun Qu:
En-Compactness: Self-Distillation Embedding & Contrastive Generation for Generalized Zero-Shot Learning. 9296-9305
- Wenjia Xu, Yongqin Xian, Jiuniu Wang , Bernt Schiele , Zeynep Akata:
VGSE: Visually-Grounded Semantic Embeddings for Zero-Shot Learning. 9306-9315
- Xiangyu Li, Xu Yang, Kun Wei, Cheng Deng, Muli Yang:
Siamese Contrastive Embedding Network for Compositional Zero-Shot Learning. 9316-9325
- Shyamgopal Karthik, Massimiliano Mancini , Zeynep Akata:
KG-SP: Knowledge Guided Simple Primitives for Open World Compositional Zero-Shot Learning. 9326-9335
- Yaogong Feng, Xiaowen Huang, Pengbo Yang, Jian Yu, Jitao Sang:
Non-generative Generalized Zero-shot Learning via Task-correlated Disentanglement and Controllable Samples Synthesis. 9336-9345
- N. Dinesh Reddy, Robert Tamburo, Srinivasa G. Narasimhan :
WALT: Watch And Learn 2D amodal representation from Time-lapse imagery. 9346-9356
- Pei Wang, Zhaowei Cai, Hao Yang, Gurumurthy Swaminathan, Nuno Vasconcelos , Bernt Schiele , Stefano Soatto:
Omni-DETR: Omni-Supervised Object Detection with Transformers. 9357-9366
- Liqiang He, Sinisa Todorovic:
DESTR: Object Detection with Split Transformer. 9367-9376
- Shuai Li, Chenchang He, Ruihuang Li, Lei Zhang:
A Dual Weighting Label Assignment Scheme for Object Detection. 9377-9386
- Jiaxi Wu, Jiaxin Chen, Di Huang:
Entropy-based Active Learning for Object Detection with Progressive Diversity Constraint. 9387-9396
- Zhaohui Zheng, Rongguang Ye, Ping Wang, Dongwei Ren, Wangmeng Zuo, Qibin Hou, Ming-Ming Cheng :
Localization Distillation for Dense Object Detection. 9397-9406
- Shilong Zhang, Zhuoran Yu, Liyang Liu, Xinjiang Wang, Aojun Zhou, Kai Chen:
Group R-CNN for Weakly Semi-supervised Object Detection with Points. 9407-9416
- Tao Feng, Mang Wang, Hangjie Yuan:
Overcoming Catastrophic Forgetting in Incremental Object Detection via Elastic Response Distillation. 9417-9426
- Jilan Xu, Junlin Hou, Yuejie Zhang, Rui Feng, Rui-Wei Zhao, Tao Zhang, Xuequan Lu , Shang Gao:
CREAM: Weakly Supervised Object Localization via Class RE-Activation Mapping. 9427-9436
- Khoa D. Doan , Peng Yang, Ping Li:
One Loss for Quantization: Deep Hashing with Discrete Wasserstein Distributional Matching. 9437-9447

- Jiale Cao, Yanwei Pang, Rao Muhammad Anwer, Hisham Cholakkal, Jin Xie, Mubarak Shah, Fahad Shahbaz Khan: **PSTR: End-to-End One-Step Person Search With Transformers.** 9448-9457
- Xiaoyi Dong, Jianmin Bao, Dongdong Chen, Ting Zhang, Weiming Zhang, Nenghai Yu, Dong Chen, Fang Wen, Baining Guo: **Protecting Celebrities from DeepFake with Identity Consistency Transformer.** 9458-9468
- Liwen Xu, Zhengtao Wang, Bin Wu, Simon Lui: **MDAN: Multi-level Dependent Attention Network for Visual Emotion Analysis.** 9469-9478
- Hui Wu, Min Wang, Wengang Zhou, Houqiang Li, Qi Tian: **Contextual Similarity Distillation for Asymmetric Image Retrieval.** 9479-9488
- Li Yang , Yan Xu, Chunfeng Yuan, Wei Liu, Bing Li, Weiming Hu: **Improving Visual Grounding with Visual-Linguistic Verification and Iterative Reasoning.** 9489-9498
- Junjie Liu, Junlong Liu, Shaotian Yan, Rongxin Jiang, Xiang Tian, Boxuan Gu, Yaowu Chen, Chen Shen, Jianqiang Huang: **MPC: Multi-view Probabilistic Clustering.** 9499-9508
- Xiang Zhang, Yongwen Su, Subarna Tripathi, Zhuowen Tu: **Text Spotting Transformers.** 9509-9518
- Min Shi, Hao Lu, Chen Feng , Chengxin Liu, Zhiguo Cao: **Represent, Compare, and Learn: A Similarity-Aware Framework for Class-Agnostic Counting.** 9519-9528
- Ahyun Seo, Byungjin Kim, Suha Kwak, Minsu Cho: **Reflection and Rotation Symmetry Detection via Equivariant Learning.** 9529-9538
- Yu-Ming Tang, Yi-Xing Peng, Wei-Shi Zheng: **Learning to Imagine: Diversify Memory for Incremental Learning using Unlabeled Data.** 9539-9548
- Yuanzhi Liang, Linchao Zhu, Xiaohan Wang, Yi Yang: **A Simple Episodic Linear Probe Improves Visual Recognition in the Wild.** 9549-9559
- Mengzhe He, Yali Wang, Jiaxi Wu, Yiru Wang, Hanqing Li, Bo Li, Weihao Gan, Wei Wu, Yu Qiao: **Cross Domain Object Detection by Target-Perceived Dual Branch Distillation.** 9560-9570
- Wenzhang Zhou, Dawei Du, Libo Zhang, Tiejian Luo, Yanjun Wu: **Multi-Granularity Alignment Domain Adaptation for Object Detection.** 9571-9580
- Jiaming Han, Yuqiang Ren, Jian Ding, Xingjia Pan, Ke Yan, Gui-Song Xia : **Expanding Low-Density Latent Regions for Open-Set Object Detection.** 9581-9590
- Tz-Ying Wu, Gurumurthy Swaminathan, Zhizhong Li, Avinash Ravichandran, Nuno Vasconcelos , Rahul Bhotika, Stefano Soatto: **Class-Incremental Learning with Strong Pre-trained Models.** 9591-9600
- Hengcan Shi, Munawar Hayat, Yicheng Wu , Jianfei Cai: **ProposalCLIP: Unsupervised Open-Category Object Proposal Generation via Exploiting CLIP Cues.** 9601-9610
- Enrico Fini, Victor G. Turrisi da Costa, Xavier Alameda-Pineda, Elisa Ricci , Karteek Alahari, Julien Mairal: **Self-Supervised Models are Continual Learners.** 9611-9620
- Zeyi Huang, Haohan Wang, Dong Huang, Yong Jae Lee, Eric P. Xing: **The Two Dimensions of Worst-case Training and Their Integrated Effect for Out-of-domain Generalization.** 9621-9631
- Matthew Gwilliam, Abhinav Shrivastava: **Beyond Supervised vs. Unsupervised: Representative Benchmarking and Analysis of Image Representation Learning.** 9632-9642
- Zhenda Xie, Zheng Zhang, Yue Cao, Yutong Lin, Jianmin Bao, Zhuliang Yao, Qi Dai, Han Hu: **SimMIM: a Simple Framework for Masked Image Modeling.** 9643-9653
- Guangrun Wang, Yansong Tang, Liang Lin, Philip H. S. Torr: **Semantic-Aware Auto-Encoders for Self-supervised Representation Learning.** 9654-9665
- Nazmul Karim, Mamshad Nayeem Rizve, Nazanin Rahnavard, Ajmal Mian, Mubarak Shah: **UNICON: Combating Label Noise Through Uniform Selection and Contrastive Learning.** 9666-9676
- Zesheng Ye , Lina Yao: **Contrastive Conditional Neural Processes.** 9677-9686
- Yuhang Zhang, Xiaopeng Zhang, Lingxi Xie, Jie Li, Robert C. Qiu, Hengtong Hu, Qi Tian: **One-bit Active Query with Contrastive Pairs.** 9687-9695
- Yuanfan Guo, Minghao Xu, Jiawen Li, Bingbing Ni, Xuanyu Zhu, Zhenbang Sun, Yi Xu: **HCSC: Hierarchical Contrastive Selective Coding.** 9696-9705
- Shuangrui Ding, Maomao Li, Tianyu Yang, Rui Qian, Haohang Xu, Qingyi Chen, Jue Wang, Hongkai Xiong: **Motion-aware Contrastive Video Representation Learning via Foreground-background Merging.** 9706-9716
- Fanyi Xiao, Kaustav Kundu, Joseph Tighe, Davide Modolo: **Hierarchical Self-supervised Representation Learning for Movie Understanding.** 9717-9726
- Hanqiu Deng, Xingyu Li: **Anomaly Detection via Reverse Distillation from One-Class Embedding.** 9727-9736

- Dahyun Kim, Jonghyun Choi :
Unsupervised Representation Learning for Binary Networks by Joint Classifier Learning. 9737-9746
- Zhen Zhao , Luping Zhou , Yue Duan , Lei Wang, Lei Qi, Yinghuan Shi:
DC-SSL: Addressing Mismatched Class Distribution in Semi-supervised Learning. 9747-9755
- Shuangtong Li, Tianyi Zhou, Xinmei Tian, Dacheng Tao:
Learning to Collaborate in Decentralized Learning of Personalized Models. 9756-9765
- Siwei Wang , Xinwang Liu, Li Liu, Wenxuan Tu, Xinzhong Zhu, Jiyuan Liu, Sihang Zhou, En Zhu:
Highly-efficient Incomplete Largescale Multiview Clustering with Consensus Bipartite Graph. 9766-9775
- Youngtaek Oh, Dong-Jin Kim, In So Kweon:
DASO: Distribution-Aware Semantics-Oriented Pseudo-label for Imbalanced Semi-Supervised Learning. 9776-9786
- Haoxiang Wang, Yite Wang, Ruoyu Sun, Bo Li:
Global Convergence of MAML and Theory-Inspired Neural Architecture Search for Few-Shot Learning. 9787-9798
- Aoxue Li, Peng Yuan, Zhenguo Li:
Semi-Supervised Object Detection via Multi-instance Alignment with Global Class Prototypes. 9799-9808
- Yen-Cheng Liu, Chih-Yao Ma, Zsolt Kira:
Unbiased Teacher v2: Semi-supervised Object Detection for Anchor-free and Anchor-based Detectors. 9809-9818
- Jingyi Zhang, Jiaxing Huang, Zichen Tian, Shijian Lu:
Spectral Unsupervised Domain Adaptation for Visual Recognition. 9819-9830
- Qing Chang, Junran Peng, Lingxi Xie, Jiajun Sun, Haoran Yin, Qi Tian, Zhaoxiang Zhang:
DATA: Domain-Aware and Task-Aware Self-supervised Learning. 9831-9840
- Arnav Chavan, Rishabh Tiwari, Udbhav Bamba, Deepak K. Gupta:
Dynamic Kernel Selection for Improved Generalization and Memory Efficiency in Meta-learning. 9841-9850
- Meitar Ronen, Shahaf E. Finder , Oren Freifeld:
DeepDPM: Deep Clustering With an Unknown Number of Clusters. 9851-9860
- R. Kenny Jones , Homer Walke, Daniel Ritchie :
PLAD: Learning to Infer Shape Programs with Pseudo-Labels and Approximate Distributions. 9861-9870
- Kushal Chauhan, Barath Mohan Umapathi , Pradeep Shenoy, Manish Gupta, Devarajan Sridharan:
Robust outlier detection by de-biasing VAE likelihoods. 9871-9880
- Corentin Sautier, Gilles Puy, Spyros Gidaris, Alexandre Boulch, Andrei Bursuc, Renaud Marlet:
Image-to-Lidar Self-Supervised Distillation for Autonomous Driving Data. 9881-9891
- Mohamed Afham, Isuru Dissanayake, Dinithi Dissanayake, Amaya Dharmasiri, Kanchana Thilakarathna, Ranga Rodrigo:
CrossPoint: Self-Supervised Cross-Modal Contrastive Learning for 3D Point Cloud Understanding. 9892-9902
- Huan Gao, Jichang Guo, Guoli Wang, Qian Zhang:
Cross-Domain Correlation Distillation for Unsupervised Domain Adaptation in Nighttime Semantic Segmentation. 9903-9913
- Lukas Hoyer, Dengxin Dai, Luc Van Gool:
DAFormer: Improving Network Architectures and Training Strategies for Domain-Adaptive Semantic Segmentation. 9914-9925
- Suhyeon Lee , Hongje Seong, Seongwon Lee , Euntai Kim:
WildNet: Learning Domain Generalized Semantic Segmentation from the Wild. 9926-9936
- Jiashuo Fan, Bin Gao, Huan Jin, Lihui Jiang:
UCC: Uncertainty guided Cross-head Cotraining for Semi-Supervised Semantic Segmentation. 9937-9946
- Donghyeon Kwon, Suha Kwak:
Semi-supervised Semantic Segmentation with Error Localization Network. 9947-9957
- Dayan Guan, Jiaxing Huang, Aoran Xiao, Shijian Lu:
Unbiased Subclass Regularization for Semi-Supervised Semantic Segmentation. 9958-9968
- Dahyun Kang, Minsu Cho:
Integrative Few-Shot Learning for Classification and Segmentation. 9969-9980
- Oindrila Saha, Zezhou Cheng, Subhransu Maji:
GANORCON: Are Generative Models Useful for Few-shot Segmentation? 9981-9990
- Tianyi Chen, Yunfei Zhang, Xiaoyang Huo, Si Wu, Yong Xu, Hau-San Wong:
SphericGAN: Semi-supervised Hyper-spherical Generative Adversarial Networks for Fine-grained Image Synthesis. 9991-10000
- Jiteng Mu, Shalini De Mello, Zhiding Yu, Nuno Vasconcelos , Xiaolong Wang, Jan Kautz, Sifei Liu:
CoordGAN: Self-Supervised Dense Correspondences Emerge from GANs. 10001-10010
- Ali Hatamizadeh, Hongxu Yin, Holger Roth, Wenqi Li, Jan Kautz, Daguang Xu, Pavlo Molchanov:
GradViT: Gradient Inversion of Vision Transformers. 10011-10020
- Innfarn Yoo, Huiwen Chang, Xiyang Luo, Ondrej Stava, Ce Liu, Peyman Milanfar, Feng Yang:
Deep 3D-to-2D Watermarking: Embedding Messages in 3D Meshes and Extracting Them from 2D Renderings. 10021-10030



- Yiqing Shen , Yuyin Zhou, Lequan Yu :
CD²-pFed: Cyclic Distillation-guided Channel Decoupling for Model Personalization in Federated Learning. 10031-10040
- Jiahao Lu, Xi Sheryl Zhang, Tianli Zhao, Xiangyu He, Jian Cheng:
APRIL: Finding the Achilles' Heel on Privacy for Vision Transformers. 10041-10050
- Liangqiong Qu, Yuyin Zhou, Paul Pu Liang, Yingda Xia, Feifei Wang, Ehsan Adeli , Li Fei-Fei, Daniel L. Rubin:
Rethinking Architecture Design for Tackling Data Heterogeneity in Federated Learning. 10051-10061
- Xiuwen Fang, Mang Ye :
Robust Federated Learning with Noisy and Heterogeneous Clients. 10062-10071
- Xin-Chun Li, Yichu Xu, Shaoming Song, Bingshuai Li, Yinchuan Li , Yunfeng Shao, De-Chuan Zhan:
Federated Learning with Position-Aware Neurons. 10072-10081
- Xiaosong Ma, Jie Zhang, Song Guo, Wenchao Xu:
Layer-wised Model Aggregation for Personalized Federated Learning. 10082-10091
- Minxue Tang, Xuefei Ning, Yitu Wang, Jingwei Sun, Yu Wang, Hai Helen Li, Yiran Chen:
FedCor: Correlation-Based Active Client Selection Strategy for Heterogeneous Federated Learning. 10092-10101
- Liang Gao, Huazhu Fu , Li Li, Yingwen Chen, Ming Xu, Cheng-Zhong Xu:
FedDC: Federated Learning with Non-IID Data via Local Drift Decoupling and Correction. 10102-10111
- Anda Cheng, Peisong Wang, Xi Sheryl Zhang, Jian Cheng:
Differentially Private Federated Learning with Local Regularization and Sparsification. 10112-10121
- Zhuohang Li, Jiaxin Zhang, Luyang Liu, Jian Liu:
Auditing Privacy Defenses in Federated Learning via Generative Gradient Leakage. 10122-10132
- Wenke Huang , Mang Ye , Bo Du:
Learn from Others and Be Yourself in Heterogeneous Federated Learning. 10133-10143
- Xiaoxiao Liang , Yiqun Lin, Huazhu Fu , Lei Zhu, Xiaomeng Li:
RSCFed: Random Sampling Consensus Federated Semi-supervised Learning. 10144-10153
- Jiahua Dong, Lixu Wang, Zhen Fang, Gan Sun, Shichao Xu, Xiao Wang, Qi Zhu:
Federated Class-Incremental Learning. 10154-10163
- Lin Zhang, Li Shen, Liang Ding, Dacheng Tao, Ling-Yu Duan:
Fine-tuning Global Model via Data-Free Knowledge Distillation for Non-IID Federated Learning. 10164-10173
- Jingyi Xu , Zihan Chen, Tony Q. S. Quek , Kai Fong Ernest Chong :
FedCorr: Multi-Stage Federated Learning for Label Noise Correction. 10174-10183
- Jingtao Li , Adnan Siraj Rakin, Xing Chen, Zhezhi He, Deliang Fan, Chaitali Chakrabarti:
ResSFL: A Resistance Transfer Framework for Defending Model Inversion Attack in Split Federated Learning. 10184-10192
- Saeed Khorram, Fuxin Li:
Cycle-Consistent Counterfactuals by Latent Transformations. 10193-10202
- Vipin Pillai, Soroush Abbasi Koohpayegani, Ashley Ouligan, Dennis Fong, Hamed Pirsiavash:
Consistent Explanations by Contrastive Learning. 10203-10212
- Sukrut Rao , Moritz Böhle, Bernt Schiele :
Towards Better Understanding Attribution Methods. 10213-10222
- Monish Keswani, Sriranjani Ramakrishnan, Nishant Reddy, Vineeth N. Balasubramanian:
Proto2Proto: Can you recognize the car, the way I do? 10223-10233
- Ashkan Khakzar, Pedram Khorsandi, Rozhin Nobahari, Nassir Navab:
Do Explanations Explain? Model Knows Best. 10234-10243
- Andong Wang, Wei-Ning Lee, Xiaojuan Qi:
HINT: Hierarchical Neuron Concept Explainer. 10244-10254
- Jon Donnelly , Alina Jade Barnett , Chaofan Chen:
Deformable ProtoPNet: An Interpretable Image Classifier Using Deformable Prototypes. 10255-10265
- Kshitij Dwivedi, Gemma Roig, Aniruddha Kembhavi, Roozbeh Mottaghi:
What do navigation agents learn about their environment? 10266-10275
- Anirban Sarkar, Deepak Vijaykeerthy, Anindya Sarkar, Vineeth N. Balasubramanian:
A Framework for Learning Ante-hoc Explainable Models via Concepts. 10276-10285
- Mahdi S. Hosseini, Mathieu Tuli, Konstantinos N. Plataniotis:
Exploiting Explainable Metrics for Augmented SGD. 10286-10296
- Yuxi Wu, Changhuai Chen, Jun Che, Shiliang Pu:
FAM: Visual Explanations for the Feature Representations from Deep Convolutional Networks. 10297-10306
- Wolfgang Stammer, Marius Memmel, Patrick Schramowski, Kristian Kersting:
Interactive Disentanglement: Learning Concepts by Interacting with their Prototype Representations. 10307-10318

- Moritz Böhle, Mario Fritz, Bernt Schiele :
B-cos Networks: Alignment is All We Need for Interpretability. 10319-10328
- Nathan Mankovich, Emily J. King , Chris Peterson, Michael Kirby:
The Flag Median and FlagIRLS. 10329-10337
- Sangwon Jung, Sanghyuk Chun, Taesup Moon:
Learning Fair Classifiers with Partially Annotated Group Labels. 10338-10347
- Shervin Ardesthir, Cristina Segalin, Nathan Kallus:
Estimating Structural Disparities for Face Models. 10348-10357
- Chirag Agarwal, Daniel D'souza, Sara Hooker:
Estimating Example Difficulty using Variance of Gradients. 10358-10368
- Zhibo Wang, Xiaowei Dong, Henry Xue, Zhifei Zhang, Weifeng Chiu, Tao Wei, Kui Ren:
Fairness-aware Adversarial Perturbation Towards Bias Mitigation for Deployed Deep Models. 10369-10378
- Sungho Park, Jework Lee, Pilhyeon Lee, Sunhee Hwang, Dohyung Kim, Hyeran Byun:
Fair Contrastive Learning for Facial Attribute Classification. 10379-10388
- Ganesh Del Grosso, Hamid Jalalzai, Georg Pichler, Catuscia Palamidessi, Pablo Piantanida:
Leveraging Adversarial Examples to Quantify Membership Information Leakage. 10389-10399
- Dominik Zietlow, Michael Lohaus, Guha Balakrishnan, Matthias Kleindessner, Francesco Locatello, Bernhard Schölkopf, Chris Russell :
Leveling Down in Computer Vision: Pareto Inefficiencies in Fair Deep Classifiers. 10400-10411
- Ronak Mehta, Sourav Pal, Vikas Singh, Sathya N. Ravi:
Deep Unlearning via Randomized Conditionally Independent Hessians. 10412-10421
- Vishnu Suresh Lokhande, Rudrasis Chakraborty, Sathya N. Ravi, Vikas Singh:
Equivariance Allows Handling Multiple Nuisance Variables When Analyzing Pooled Neuroimaging Datasets. 10422-10431
- Kirill Sirokin, Pablo Carballeira , Marcos Escudero-Viñolo :
A study on the distribution of social biases in self-supervised learning visual models. 10432-10441
- Cho-Ying Wu, Chin-Cheng Hsu, Ulrich Neumann:
Cross-Modal Perceptionist: Can Face Geometry be Gleaned from Voices? 10442-10451
- Xian Liu, Qianyi Wu, Hang Zhou, Yinghao Xu, Rui Qian, Xinyi Lin, Xiaowei Zhou, Wayne Wu, Bo Dai, Bolei Zhou:
Learning Hierarchical Cross-Modal Association for Co-Speech Gesture Generation. 10452-10462
- Yuanzhi Liang, Qianyu Feng, Linchao Zhu, Li Hu, Pan Pan, Yi Yang:
SEEG: Semantic Energized Co-speech Gesture Generation. 10463-10472
- Xixi Hu, Ziyang Chen, Andrew Owens:
Mix and Localize: Localizing Sound Sources in Mixtures. 10473-10482
- Akam Rahimi, Triantafyllos Afouras, Andrew Zisserman:
Reading to Listen at the Cocktail Party: Multi-Modal Speech Separation. 10483-10492
- Guande Wu, Jianzhe Lin, Cláudio T. Silva:
IntentVizor: Towards Generic Query Guided Interactive Video Summarization. 10493-10502
- Tsu-Jui Fu, Xin Eric Wang, Scott T. Grafton, Miguel P. Eckstein, William Yang Wang:
M³L: Language-based Video Editing via Multi-Modal Multi-Level Transformers. 10503-10512
- Chuang Gan, Yi Gu, Siyuan Zhou, Jeremy Schwartz, Seth Alter, James Traer, Dan Gutfreund, Joshua B. Tenenbaum, Josh H. McDermott, Antonio Torralba:
Finding Fallen Objects Via Asynchronous Audio-Visual Integration. 10513-10523
- Sangmin Lee , Hyung-Il Kim, Yong Man Ro:
Weakly Paired Associative Learning for Sound and Image Representations via Bimodal Associative Memory. 10524-10533
- Hao Jiang, Calvin Murdock, Vamsi Krishna Ithapu:
Egocentric Deep Multi-Channel Audio-Visual Active Speaker Localization. 10534-10542
- Otniel-Bogdan Mercea, Lukas Riesch, A. Sophia Koepke, Zeynep Akata:
Audiovisual Generalised Zero-shot Learning with Cross-modal Attention and Language. 10543-10553
- Dídac Surís, Carl Vondrick, Bryan Russell, Justin Salamon:
It's Time for Artistic Correspondence in Music and Video. 10554-10564
- Triantafyllos Afouras, Yuki M. Asano, Francois Fagan, Andrea Vedaldi, Florian Metze:
Self-supervised object detection from audio-visual correspondence. 10565-10576
- Michael Hassid, Michelle Tadmor Ramanovich, Brendan Shillingford, Miaosen Wang, Ye Jia, Tal Remez:
More than Words: In-the-Wild Visually-Driven Prosody for Text-to-Speech. 10577-10587
- Ruohan Gao, Zilin Si, Yen-Yu Chang, Samuel Clarke, Jeannette Bohg , Li Fei-Fei, Wenzhen Yuan, Jiajun Wu:
ObjectFolder 2.0: A Multisensory Object Dataset for Sim2Real Transfer. 10588-10598
- Dongran Yu, Bo Yang, Qianhao Wei, Anchen Li, Shirui Pan:
A Probabilistic Graphical Model Based on Neural-symbolic Reasoning for Visual Relationship Detection. 10599-

- 10608
 Konpat Preechakul, Nattanat Chatthee, Suttisak Wizadwongsa, Supasorn Suwajanakorn:
Diffusion Autoencoders: Toward a Meaningful and Decodable Representation. 10609-10619
- Seung Wook Kim, Karsten Kreis, Daiqing Li, Antonio Torralba, Sanja Fidler:
Polymorphic-GAN: Generating Aligned Samples across Multiple Domains with Learned Morph Maps. 10620-10630
- Ahmed Imtiaz Humayun, Randall Balestriero, Richard G. Baraniuk:
Polarity Sampling: Quality and Diversity Control of Pre-Trained Generative Networks via Singular Values. 10631-10640
- Nupur Kumari, Richard Zhang, Eli Shechtman, Jun-Yan Zhu:
Ensembling Off-the-shelf Models for GAN Training. 10641-10652
- Fangneng Zhan, Yingchen Yu, Rongliang Wu, Jiahui Zhang , Shijian Lu, Changgong Zhang:
Marginal Contrastive Correspondence for Guided Image Generation. 10653-10662
- Yu Deng, Jiaolong Yang , Jianfeng Xiang, Xin Tong :
GRAM: Generative Radiance Manifolds for 3D-Aware Image Generation. 10663-10673
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, Björn Ommer:
High-Resolution Image Synthesis with Latent Diffusion Models. 10674-10685
- Shuyang Gu, Dong Chen, Jianmin Bao, Fang Wen, Bo Zhang, Dongdong Chen, Lu Yuan, Baining Guo:
Vector Quantized Diffusion Model for Text-to-Image Synthesis. 10686-10696
- Jianan Wang, Guansong Lu, Hang Xu, Zhenguo Li, Chunjing Xu, Yanwei Fu :
ManiTrans: Entity-Level Text-Guided Image Manipulation via Token-wise Semantic Alignment and Generation. 10697-10707
- George Cazenavette, Tongzhou Wang, Antonio Torralba, Alexei A. Efros, Jun-Yan Zhu:
Dataset Distillation by Matching Training Trajectories. 10708-10717
- Geng Chen, Wendong Zhang , Han Lu, Siyu Gao, Yunbo Wang, Mingsheng Long , Xiaokang Yang:
Continual Predictive Learning from Videos. 10718-10727
- Long Mai, Feng Liu:
Motion-Adjustable Neural Implicit Video Representation. 10728-10737
- Narek Tumanyan, Omer Bar-Tal, Shai Bagon, Tali Dekel:
Splicing ViT Features for Semantic Appearance Transfer. 10738-10747
- Wenbo Li , Zhe Lin, Kun Zhou , Lu Qi, Yi Wang, Jiaya Jia:
MAT: Mask-Aware Transformer for Large Hole Image Inpainting. 10748-10758
- Abhijith Punnappurath, Abdullah Abuolaim, Abdelrahman Abdelhamed, Alex Levishtein, Michael S. Brown:
Day-to-Night Image Synthesis for Training Nighttime Neural ISPs. 10759-10768
- Jiseob Kim, Jihoon Lee, Byoung-Tak Zhang:
Smooth-Swap: A Simple Enhancement for Face-Swapping with Smoothness. 10769-10778
- Changyong Shu, Hemao Wu, Hang Zhou, Jiaming Liu, Zhibin Hong, Changxing Ding, Junyu Han, Jingtuo Liu, Errui Ding, Jingdong Wang:
Few-Shot Head Swapping in the Wild. 10779-10788
- Jianbin Jiang, Tan Wang, He Yan, Junhui Liu:
ClothFormer: Taming Video Virtual Try-on in All Module. 10789-10798
- Hongxu Yin, Arash Vahdat, José M. Álvarez, Arun Mallya, Jan Kautz, Pavlo Molchanov:
A-ViT: Adaptive Tokens for Efficient Vision Transformer. 10799-10808
- Weihao Yu , Mi Luo, Pan Zhou , Chenyang Si, Yichen Zhou, Xinchao Wang , Jiashi Feng, Shuicheng Yan:
MetaFormer is Actually What You Need for Vision. 10809-10819
- Karttikeya Mangalam, Haoqi Fan, Yanghao Li, Chao-Yuan Wu, Bo Xiong, Christoph Feichtenhofer, Jitendra Malik:
Reversible Vision Transformers. 10820-10830
- Moab Arar, Ariel Shamir, Amit H. Bermano :
Learned Queries for Efficient Local Attention. 10831-10842
- Sucheng Ren, Daquan Zhou , Shengfeng He , Jiashi Feng, Xinchao Wang :
Shunted Self-Attention via Multi-Scale Token Aggregation. 10843-10852
- Shaofei Cai, Liang Li, Xinzhe Han, Jiebo Luo , Zheng-Jun Zha, Qingming Huang:
Automatic Relation-aware Graph Network Proliferation. 10853-10863
- Peng Ye, Baopu Li, Yikang Li, Tao Chen, Jiayuan Fan, Wanli Ouyang:
β-DARTS: Beta-Decay Regularization for Differentiable Architecture Search. 10864-10873
- Junyi Pan, Chong Sun, Yizhou Zhou, Ying Zhang, Chen Li:
Distribution Consistent Neural Architecture Search. 10874-10883
- Qinjin Zhou, Kekai Sheng, Xiawu Zheng, Ke Li, Xing Sun, Yonghong Tian, Jie Chen, Rongrong Ji:
Training-free Transformer Architecture Search. 10884-10893
- Teppei Suzuki:
TeachAugment: Data Augmentation Optimization Using Teacher Knowledge. 10894-10904

- Sihao Lin, Hongwei Xie, Bing Wang, Kaicheng Yu, Xiaojun Chang , Xiaodan Liang, Gang Wang: **Knowledge Distillation via the Target-aware Transformer.** 10905-10914
- Lucas Beyer, Xiaohua Zhai, Amélie Royer, Larisa Markeeva, Rohan Anil, Alexander Kolesnikov: **Knowledge distillation: A good teacher is patient and consistent.** 10915-10924
- Yehui Tang, Kai Han, Jianyuan Guo , Chang Xu, Yanxi Li , Chao Xu, Yunhe Wang: **An Image Patch is a Wave: Phase-Aware Vision MLP.** 10925-10934
- Lingfeng Yang, Xiang Li, Renjie Song, Borui Zhao, Juntian Tao, Shihao Zhou, Jiajun Liang, Jian Yang: **Dynamic MLP for Fine-Grained Image Classification by Leveraging Geographical and Temporal Information.** 10935-10944
- Dripta S. Raychaudhuri, Yumin Suh, Samuel Schulter, Xiang Yu, Masoud Faraki, Amit K. Roy-Chowdhury, Manmohan Chandraker: **Controllable Dynamic Multi-Task Architectures.** 10945-10954
- Liunian Harold Li, Pengchuan Zhang, Haotian Zhang, Jianwei Yang, Chunyuan Li, Yiwu Zhong, Lijuan Wang, Lu Yuan, Lei Zhang, Jenq-Neng Hwang, Kai-Wei Chang, Jianfeng Gao: **Grounded Language-Image Pre-training.** 10955-10965
- Georg Bökmann , Fredrik Kahl, Axel Flinth: **ZZ-Net: A Universal Rotation-Equivariant Architecture for 2D Point Clouds.** 10966-10975
- Zhiwen Fan, Tianlong Chen, Peihao Wang, Zhangyang Wang: **CADTransformer: Panoptic Symbol Spotting Transformer for CAD Drawings.** 10976-10986
- Andrey Davydov, Anastasia Remizova, Victor Constantin, Sina Honari, Mathieu Salzmann, Pascal Fua: **Adversarial Parametric Pose Prior.** 10987-10995
- Zhengguang Liu, Runyang Feng, Haoming Chen , Shuang Wu, Yixing Gao, Yunjun Gao, Xiang Wang: **Temporal Feature Alignment and Mutual Information Maximization for Video-Based Human Pose Estimation.** 10996-11006
- Kehong Gong, Bingbing Li, Jianfeng Zhang, Tao Wang , Jing Huang, Michael Bi Mi, Jiashi Feng, Xinchao Wang : **PoseTriplet: Co-evolving 3D Human Pose Estimation, Imitation, and Hallucination under Self-supervision.** 11007-11017
- Kristijan Bartol, David Bojanic, Tomislav Petkovic: **Generalizable Human Pose Triangulation.** 11018-11027
- Ye Yuan, Umar Iqbal, Pavlo Molchanov, Kris Kitani, Jan Kautz: **GLAMR: Global Occlusion-Aware Human Mesh Recovery with Dynamic Cameras.** 11028-11039
- Li Siyao, Weijiang Yu, Tianpei Gu, Chunze Lin, Quan Wang, Chen Qian, Chen Change Loy, Ziwei Liu: **Bailando: 3D Dance Generation by Actor-Critic GPT with Choreographic Memory.** 11040-11049
- Dongkai Wang, Shiliang Zhang: **Contextual Instance Decoupling for Robust Multi-Person Pose Estimation.** 11050-11058
- Dahu Shi , Xing Wei, Liangqi Li, Ye Ren, Wenming Tan: **End-to-End Multi-Person Pose Estimation with Transformers.** 11059-11068
- Jia Gong, Zhipeng Fan , QiuHong Ke, Hossein Rahmani, Jun Liu: **Meta Agent Teaming Active Learning for Pose Estimation.** 11069-11079
- Shreyas Hampali, Sayan Deb Sarkar, Mahdi Rad, Vincent Lepetit: **Keypoint Transformer: Solving Joint Identification in Challenging Hands and Object Interactions for Accurate 3D Pose Estimation.** 11080-11090
- Wang Zeng, Sheng Jin, Wentao Liu, Chen Qian, Ping Luo, Wanli Ouyang, Xiaogang Wang : **Not All Tokens Are Equal: Human-centric Visual Analysis via Token Clustering Transformer.** 11091-11101
- Congcong Zhu , Xintong Wan, Shaorong Xie, Xiaoqiang Li, Yinzhen Gu: **Occlusion-robust Face Alignment using A Viewpoint-invariant Hierarchical Network Architecture.** 11102-11111
- Zhixiang Min, Naji Khosravan, Zachary Bessinger, Manjunath Narayana, Sing Bing Kang, Enrique Dunn , Ivaylo Boyadzhiev: **LASER: LATent SpacE Rendering for 2D Visual Localization.** 11112-11121
- Tien Do, Ondrej Miksik, Joseph DeGol, Hyun Soo Park, Sudipta N. Sinha: **Learning to Detect Scene Landmarks for Camera Localization.** 11122-11132
- Zheng Qin , Hao Yu, Changjian Wang, Yulan Guo, Yuxing Peng, Kai Xu: **Geometric Transformer for Fast and Robust Point Cloud Registration.** 11133-11142
- Liangzu Peng , Manolis C. Tsakiris, René Vidal: **ARCS: Accurate Rotation and Correspondence Search.** 11143-11153
- Yingda Yin, Yingcheng Cai, He Wang, Baoquan Chen: **FisherMatch: Semi-Supervised Rotation Regression via Entropy-based Filtering.** 11154-11163
- Xiaoke Jiang, Donghai Li, Hao Chen, Ye Zheng , Rui Zhao, Liwei Wu: **Uni6D: A Unified CNN Framework without Projection Breakdown for 6D Pose Estimation.** 11164-11174
- Kai Katsumata, Duc Minh Vo, Hideki Nakayama: **OSSGAN: Open-Set Semi-Supervised Image Generation.** 11175-11183

- Guanqi Ding, Xinzhe Han, Shuhui Wang, Shuzhe Wu, Xin Jin, Dandan Tu, Qingming Huang:
Attribute Group Editing for Reliable Few-shot Image Generation. 11184-11193
- Jiayu Xiao, Liang Li, Chaofei Wang, Zheng-Jun Zha, Qingming Huang:
Few Shot Generative Model Adaption via Relaxed Spatial Structural Alignment. 11194-11203
- Zhengyao Lv, Xiaoming Li, Zhenxing Niu, Bing Cao, Wangmeng Zuo:
Semantic-shape Adaptive Feature Modulation for Semantic Image Synthesis. 11204-11213
- Yupeng Shi, Xiao Liu, Yuxiang Wei, Zhongqin Wu, Wangmeng Zuo:
Retrieval-based Spatially Adaptive Normalization for Semantic Image Synthesis. 11214-11223
- Rhea Sanjay Sukthanker, Zhiwu Huang, Suryansh Kumar , Radu Timofte , Luc Van Gool:
Generative Flows with Invertible Attentions. 11224-11233
- Yuan-kui Li, Yun-Hsuan Lien, Yu-Shuen Wang:
Style-Structure Disentangled Features and Normalizing Flows for Diverse Icon Colorization. 11234-11243
- Yichun Shi, Xiao Yang, Yangyue Wan, Xiaohui Shen:
SemanticStyleGAN: Learning Compositional Generative Priors for Controllable Image Synthesis and Editing. 11244-11254
- Yao Ni, Piotr Koniusz, Richard I. Hartley, Richard Nock:
Manifold Learning Benefits GANs. 11255-11264
- Aye Phyu Phyu Aung, Xinrun Wang, Runsheng Yu, Bo An, Senthilnath Jayavelu , Xiaoli Li:
DO-GAN: A Double Oracle Framework for Generative Adversarial Networks. 11265-11274
- Jianyu Wang, Ceyuan Yang, Yinghao Xu, Yujun Shen, Hongdong Li, Bolei Zhou:
Improving GAN Equilibrium by Raising Spatial Awareness. 11275-11283
- Junho Kim, Yunjey Choi, Youngjung Uh :
Feature Statistics Mixing Regularization for Generative Adversarial Networks. 11284-11293
- Bowen Zhang, Shuyang Gu, Bo Zhang, Jianmin Bao, Dong Chen, Fang Wen, Yong Wang, Baining Guo:
StyleSwin: Transformer-based GAN for High-resolution Image Generation. 11294-11304
- Huiwen Chang, Han Zhang, Lu Jiang, Ce Liu, William T. Freeman:
MaskGIT: Masked Generative Image Transformer. 11305-11315
- Yingying Deng, Fan Tang, Weiming Dong, Chongyang Ma, Xingjia Pan, Lei Wang, Changsheng Xu:
StyTr²: Image Style Transfer with Transformers. 11316-11326
- Xueqi Hu, Qiusheng Huang, Zhengyi Shi, Siyuan Li, Changxin Gao, Li Sun, Qingli Li:
Style Transformer for Image Inversion and Editing. 11327-11336
- Qiankun Liu, Zhentao Tan, Dongdong Chen, Qi Chu, Xiyang Dai, Yinpeng Chen, Mengchen Liu, Lu Yuan, Nenghai Yu:
Reduce Information Loss in Transformers for Pluralistic Image Inpainting. 11337-11347
- Qiaole Dong , Chenjie Cao, Yanwei Fu :
Incremental Transformer Structure Enhanced Image Inpainting with Masking Positional Encoding. 11348-11358
- Jimeng Sun, Shuchen Weng, Zheng Chang, Si Li, Boxin Shi:
UniCoRN: A Unified Conditional Image Repainting Network. 11359-11368
- Tengfei Wang, Yong Zhang, Yanbo Fan, Jue Wang, Qifeng Chen:
High-Fidelity GAN Inversion for Image Attribute Editing. 11369-11378
- Tan M. Dinh, Anh Tuan Tran, Rang Nguyen, Binh-Son Hua:
HyperInverter: Improving StyleGAN Inversion via Hypernetwork. 11379-11388
- Gaurav Parmar, Yijun Li, Jingwan Lu, Richard Zhang, Jun-Yan Zhu, Krishna Kumar Singh:
Spatially-Adaptive Multilayer Selection for GAN Inversion and Editing. 11389-11399
- Gaurav Parmar, Richard Zhang, Jun-Yan Zhu:
On Aliased Resizing and Surprising Subtleties in GAN Evaluation. 11400-11410
- Wentao Wang, Li Niu, Jianfu Zhang, Xue Yang, Liqing Zhang:
Dual-path Image Inpainting with Auxiliary GAN Inversion. 11411-11420
- Yen-Chi Cheng , Chieh Hubert Lin, Hsin-Ying Lee, Jian Ren, Sergey Tulyakov, Ming-Hsuan Yang:
InOut: Diverse Image Outpainting via GAN Inversion. 11421-11430
- Naofumi Akimoto, Yushi Matsuo, Yoshimitsu Aoki:
Diverse Plausible 360-Degree Image Outpainting for Efficient 3DCG Background Creation. 11431-11440
- Jiacheng Li, Chang Chen, Zhiwei Xiong:
Contextual Outpainting with Object-Level Contrastive Learning. 11441-11450
- Andreas Lugmayr, Martin Danelljan, Andrés Romero, Fisher Yu, Radu Timofte , Luc Van Gool:
RePaint: Inpainting using Denoising Diffusion Probabilistic Models. 11451-11461
- Jooyoung Choi, Jungbeom Lee, Chaehun Shin, Sungwon Kim, Hyunwoo Kim, Sungroh Yoon:
Perception Prioritized Training of Diffusion Models. 11462-11471
- Yaniv Benny, Lior Wolf:
Dynamic Dual-Output Diffusion Models. 11472-11481

- Vikash Sehwag, Caner Hazirbas, Albert Gordo, Firat Ozgenel, Cristian Canton-Ferrer: **Generating High Fidelity Data from Low-density Regions using Diffusion Models.** 11482-11491
- Minghui Hu , Yujie Wang, Tat-Jen Cham , Jianfei Yang, Ponnuthurai N. Suganthan: **Global Context with Discrete Diffusion in Vector Quantised Modelling for Image Generation.** 11492-11501
- Chuanxia Zheng, Tat-Jen Cham , Jianfei Cai, Dinh Q. Phung: **Bridging Global Context Interactions for High-Fidelity Image Completion.** 11502-11512
- Doyup Lee, Chiheon Kim, Saehoon Kim, Minsu Cho, Wook-Shin Han: **Autoregressive Image Generation using Residual Quantization.** 11513-11522
- Evangelos Ntavelis, Mohamad Shahbazi, Iason Kastanis, Radu Timofte , Martin Danelljan, Luc Van Gool: **Arbitrary-Scale Image Synthesis.** 11523-11532
- Markos Georgopoulos, James Oldfield , Grigoris G. Chrysos, Yannis Panagakis: **Cluster-guided Image Synthesis with Unconditional Models.** 11533-11542
- Jie Liu, Yanqi Bao, Guo-Sen Xie, Huan Xiong, Jan-Jakob Sonke, Efstratios Gavves: **Dynamic Prototype Convolution Network for Few-Shot Semantic Segmentation.** 11543-11552
- Zhuotao Tian, Xin Lai, Li Jiang , Shu Liu, Michelle Shu, Hengshuang Zhao, Jiaya Jia: **Generalized Few-shot Semantic Segmentation.** 11553-11562
- Yuanwei Liu, Nian Liu , Qinglong Cao, Xiwen Yao, Junwei Han, Ling Shao: **Learning Non-target Knowledge for Few-shot Semantic Segmentation.** 11563-11572
- Jian Ding , Nan Xue, Gui-Song Xia, Dengxin Dai: **Decoupling Zero-Shot Semantic Segmentation.** 11573-11582
- Ruihuang Li, Shuai Li, Chenhang He, Yabin Zhang, Xu Jia, Lei Zhang: **Class-Balanced Pixel-Level Self-Labeling for Domain Adaptive Semantic Segmentation.** 11583-11593
- Xuehui Wang, Kai Zhao, Ruixin Zhang, Shouhong Ding, Yan Wang, Wei Shen: **ContrastMask: Contrastive Learning to Segment Every Thing.** 11594-11603
- R. Kenny Jones , Aalia Habib, Rana Hanocka, Daniel Ritchie : **The Neurally-Guided Shape Parser: Grammar-based Labeling of 3D Shape Regions with Approximate Inference.** 11604-11613
- Xueyi Liu, Xiaomeng Xu, Anyi Rao , Chuang Gan, Li Yi: **AutoGPart: Intermediate Supervision Search for Generalizable 3D Part Segmentation.** 11614-11624
- Zhan Xu, Matthew Fisher, Yang Zhou, Deepali Aneja, Rushikesh Dudhat, Li Yi, Evangelos Kalogerakis: **APES: Articulated Part Extraction from Sprite Sheets.** 11625-11634
- Alberto Bailoni, Constantin Pape, Nathan Hütsch, Steffen Wolf, Thorsten Beier, Anna Kreshuk, Fred A. Hamprecht: **GASP, a generalized framework for agglomerative clustering of signed graphs and its application to Instance Segmentation.** 11635-11645
- Ke Zhang, Xiahai Zhuang: **CycleMix: A Holistic Strategy for Medical Image Segmentation from Scribble Supervision.** 11646-11655
- Huisi Wu, Zhaoze Wang, Youyi Song, Lin Yang, Jing Qin: **Cross-patch Dense Contrastive Learning for Semi-supervised Segmentation of Cellular Nuclei in Histopathologic Images.** 11656-11665
- Zhang Chen , Zhiqiang Tian, Jihua Zhu, Ce Li, Shaoyi Du: **C-CAM: Causal CAM for Weakly Supervised Semantic Segmentation on Medical Image.** 11666-11675
- Zhaoqing Wang, Yu Lu, Qiang Li, Xunqiang Tao, Yandong Guo, Mingming Gong, Tongliang Liu : **CRIS: CLIP-Driven Referring Image Segmentation.** 11676-11685
- Gyutae Park, Sungjoon Son, Jaeyoung Yoo, Seho Kim, Nojun Kwak: **MatteFormer: Transformer-Based Image Matting via Prior-Tokens.** 11686-11696
- Yutong Dai, Brian L. Price, He Zhang, Chunhua Shen: **Boosting Robustness of Image Matting with Context Assembling and Strong Data Augmentation.** 11697-11706
- Chenxi Xie, Changqun Xia, Mingcan Ma, Zhirui Zhao, Xiaowu Chen, Jia Li: **Pyramid Grafting Network for One-Stage High Resolution Saliency Detection.** 11707-11716
- Yifan Wang, Wenbo Zhang, Lijun Wang, Ting Liu, Huchuan Lu: **Multi-Source Uncertainty Mining for Deep Unsupervised Saliency Detection.** 11717-11726
- Wangbo Zhao, Kai Wang, Xiangxiang Chu, Fuzhao Xue, Xinchao Wang , Yang You: **Modeling Motion with Multi-Modal Features for Text-Based Video Segmentation.** 11727-11736
- Zhaohua Zheng, Jianfang Li, Lingjie Zhu, Honghua Li, Frank Petzold , Ping Tan: **GAT-CADNet: Graph Attention Network for Panoptic Symbol Spotting in CAD Drawings.** 11737-11746
- Mahdi Saleh, Shun-Cheng Wu , Luca Cosmo, Nassir Navab, Benjamin Busam , Federico Tombari: **Bending Graphs: Hierarchical Shape Matching using Gated Optimal Transport.** 11747-11757
- Fenggen Yu, Zhiqin Chen, Manyi Li, Aditya Sanghi, Hooman Shayani, Ali Mahdavi-Amiri, Hao Zhang: **CAPRI-Net: Learning Compact CAD Shapes with Adaptive Primitive Assembly.** 11758-11768

- Chengjie Niu, Manyi Li, Kai Xu, Hao Zhang:
RIM-Net: Recursive Implicit Fields for Unsupervised Learning of Hierarchical Shape Structures. 11769-11778
- Zhipeng Bao, Pavel Tokmakov, Allan Jabri, Yu-Xiong Wang, Adrien Gaidon, Martial Hebert:
Discovering Objects that Can Move. 11779-11788
- Cheng Zhang, Haocheng Wan, Xinyi Shen, Zizhao Wu:
PatchFormer: An Efficient Point Transformer with Patch Attention. 11789-11798
- Jinke Li, Xiao He, Yang Wen, Yuan Gao, Xiaoqiang Cheng, Dan Zhang:
Panoptic-PHNet: Towards Real-Time and High-Precision LiDAR Panoptic Segmentation via Clustering Pseudo Heatmap. 11799-11808
- Ziyi Wang, Yongming Rao, Xumin Yu, Jie Zhou, Jiwen Lu:
SemAffiNet: Semantic-Affine Transformation for Point Cloud Segmentation. 11809-11819
- Cheng-Kun Yang, Ji-Jia Wu, Kai-Syun Chen, Yung-Yu Chuang, Yen-Yu Lin :
An MIL-Derived Transformer for Weakly Supervised Point Cloud Segmentation. 11820-11829
- Hanyu Shi, Jiacheng Wei, Ruibo Li, Fayao Liu, Guosheng Lin:
Weakly Supervised Segmentation on Outdoor 4D point clouds with Temporal Matching and Spatial Graph Propagation. 11830-11839
- Mikaela Angelina Uy, Yen-Yu Chang, Minhyuk Sung , Purvi Goel, Joseph Lambourne, Tolga Birdal , Leonidas J. Guibas:
Point2Cyl: Reverse Engineering 3D Objects from Point Clouds to Extrusion Cylinders. 11840-11850
- Jisoo Mok, Byunggoon Na, Ji-Hoon Kim, Dongyo Han, Sungroh Yoon:
Demystifying the Neural Tangent Kernel from a Practical Perspective: Can it be trusted for Neural Architecture Search without training? 11851-11860
- Miao Zhang, Shirui Pan, Xiaojun Chang , Steven Su, Jilin Hu , Gholamreza Haffari, Bin Yang :
BaLeNAS: Differentiable Architecture Search via the Bayesian Learning Rule. 11861-11870
- Minbin Huang, Zhijian Huang, Changlin Li, Xin Chen, Hang Xu, Zhenguo Li, Xiaodan Liang:
Arch-Graph: Acyclic Architecture Relation Predictor for Task-Transferable Neural Architecture Search. 11871-11881
- Han Xiao, Ziwei Wang , Zheng Zhu, Jie Zhou, Jiwen Lu:
Shapley-NAS: Discovering Operation Contribution for Neural Architecture Search. 11882-11891
- Tao Huang, Shan You, Fei Wang, Chen Qian, Changshui Zhang, Xiaogang Wang , Chang Xu:
GreedyNASv2: Greedier Search with a Greedy Path Filter. 11892-11901
- Xiaowu Zheng, Xiang Fei, Lei Zhang, Chenglin Wu, Fei Chao, Jianzhuang Liu, Wei Zeng , Yonghong Tian, Rongrong Ji:
Neural Architecture Search with Representation Mutual Information. 11902-11911
- Pengtao Xie, Xuefeng Du:
Performance-Aware Mutual Knowledge Distillation for Improving Neural Architecture Search. 11912-11922
- Defang Chen, Jian-Ping Mei, Hailin Zhang, Can Wang, Yan Feng, Chun Chen:
Knowledge Distillation with the Reused Teacher Classifier. 11923-11932
- Yiqing Shen , Liwu Xu, Yuzhe Yang, Yaqian Li, Yandong Guo:
Self-Distillation from the Last Mini-Batch for Consistency Regularization. 11933-11942
- Borui Zhao, Quan Cui, Renjie Song, Yiyu Qiu, Jiajun Liang:
Decoupled Knowledge Distillation. 11943-11952
- Xiaohan Ding, Xiangyu Zhang, Jungong Han, Guiguang Ding:
Scaling Up Your Kernels to 31×31: Revisiting Large Kernel Design in CNNs. 11953-11965
- Zhuang Liu, Hanzi Mao, Chao-Yuan Wu, Christoph Feichtenhofer, Trevor Darrell, Saining Xie:
A ConvNet for the 2020s. 11966-11976
- Pengzhen Ren , Changlin Li, Guangrun Wang, Yun Xiao, Qing Du, Xiaodan Liang, Xiaojun Chang :
Beyond Fixation: Dynamic Window Visual Transformer. 11977-11987
- Chenglin Yang, Yilin Wang, Jianming Zhang, He Zhang, Zijun Wei, Zhe Lin, Alan L. Yuille:
Lite Vision Transformer with Enhanced Self-Attention. 11988-11998
- Ze Liu, Han Hu, Yutong Lin, Zhuliang Yao, Zhenda Xie, Yixuan Wei, Jia Ning, Yue Cao, Zheng Zhang, Li Dong, Furu Wei, Baining Guo:
Swin Transformer V2: Scaling Up Capacity and Resolution. 11999-12009
- Tianlong Chen, Zhenyu Zhang, Yu Cheng, Ahmed Awadallah, Zhangyang Wang:
The Principle of Diversity: Training Stronger Vision Transformers Calls for Reducing All Levels of Redundancy. 12010-12020
- Deblina Bhattacharjee , Tong Zhang, Sabine Süsstrunk, Mathieu Salzmann:
MuLT: An End-to-End Multitask Learning Transformer. 12021-12031
- Xiaofeng Mao , Gege Qi, Yuefeng Chen, Xiaodan Li, Ranjie Duan, Shaokai Ye, Yuan He, Hui Xue:
Towards Robust Vision Transformer. 12032-12041
- Xianing Chen, Qiong Cao, Yujie Zhong, Jing Zhang, Shenghua Gao, Dacheng Tao:
DearKD: Data-Efficient Early Knowledge Distillation for Vision Transformers. 12042-12052
- Jiemin Fang, Lingxi Xie, Xinggang Wang, Xiaopeng Zhang, Wenyu Liu, Qi Tian:
MSG-Transformer: Exchanging Local Spatial Information by Manipulating Messenger Tokens. 12053-12062

- Hao Liu, Xinghua Jiang, Xin Li, Zhimin Bao, Deqiang Jiang, Bo Ren:
NomMer: Nominate Synergistic Context in Vision Transformer for Visual Recognition. 12063-12072
- Wenqiang Zhang, Zilong Huang, Guozhong Luo, Tao Chen, Xinggang Wang, Wenyu Liu, Gang Yu, Chunhua Shen:
TopFormer: Token Pyramid Transformer for Mobile Semantic Segmentation. 12073-12083
- Jiaqi Gu, Hyoukjun Kwon, Dilin Wang, Wei Ye, Meng Li, Yu-Hsin Chen, Liangzhen Lai, Vikas Chandra, David Z. Pan:
Multi-Scale High-Resolution Vision Transformer for Semantic Segmentation. 12084-12093
- Yikai Wang, TengQi Ye, Lele Cao , Wenbing Huang, Fuchun Sun, Fengxiang He, Dacheng Tao:
Bridged Transformer for Vision and Point Cloud 3D Object Detection. 12104-12113
- Xiaoyi Dong, Jianmin Bao, Dongdong Chen, Weiming Zhang, Nenghai Yu, Lu Yuan, Dong Chen, Baining Guo:
CSSWin Transformer: A General Vision Transformer Backbone with Cross-Shaped Windows. 12114-12124
- Jieneng Chen, Shuyang Sun, Ju He, Philip H. S. Torr, Alan L. Yuille, Song Bai:
TransMix: Attend to Mix for Vision Transformers. 12125-12134
- Jinnian Zhang, Houwen Peng, Kan Wu, Mengchen Liu, Bin Xiao, Jianlong Fu, Lu Yuan:
MiniViT: Compressing Vision Transformers with Weight Multiplexing. 12135-12144
- Mark Sandler, Andrey Zhmoginov, Max Vladimirov, Andrew Jackson:
Fine-tuning Image Transformers using Learnable Memory. 12145-12154
- Yehui Tang, Kai Han, Yunhe Wang, Chang Xu, Jianyuan Guo , Chao Xu, Dacheng Tao:
Patch Slimming for Efficient Vision Transformers. 12155-12164
- Jianyuan Guo , Kai Han, Han Wu , Yehui Tang, Xinghao Chen, Yunhe Wang, Chang Xu:
CMT: Convolutional Neural Networks Meet Vision Transformers. 12165-12175
- Yikai Wang, Xinghao Chen, Lele Cao , Wenbing Huang, Fuchun Sun, Yunhe Wang:
Multimodal Token Fusion for Vision Transformers. 12176-12185
- Kai Wang, Bo Zhao, Xiangyu Peng, Zheng Zhu, Shuo Yang, Shuo Wang, Guan Huang, Hakan Bilen , Xinchao Wang , Yang You:
CAFE: Learning to Condense Dataset by Aligning Features. 12186-12195
- Qian Lou, Yen-Chang Hsu, Burak Uzkent, Ting Hua, Yilin Shen, Hongxia Jin:
Lite-MDETR: A Lightweight Multi-Modal Detector. 12196-12205
- Zhengcong Fei, Xu Yan, Shuhui Wang, Qi Tian:
DeeCap: Dynamic Early Exiting for Efficient Image Captioning. 12206-12216
- Linnan Wang, Chenhan Yu, Satish Salian, Slawomir Kierat, Szymon Migacz, Alex Fit-Florea:
Searching the Deployable Convolution Neural Networks for GPUs. 12217-12226
- Amin Parvaneh, Ehsan Abbasnejad, Damien Teney, Reza Haffari, Anton van den Hengel, Javen Qinfeng Shi :
Active Learning by Feature Mixing. 12227-12236
- Maying Shen, Pavlo Molchanov, Hongxu Yin, José M. Álvarez:
When to Prune? A Policy towards Early Structural Pruning. 12237-12246
- Jian Meng, Li Yang, Jinwoo Shin, Deliang Fan, Jae-Sun Seo:
Contrastive Dual Gating: Learning Sparse Features With Contrastive Learning. 12247-12255
- Eugenia Iofinova, Alexandra Peste, Mark Kurtz, Dan Alistarh:
How Well Do Sparse ImageNet Models Transfer? 12256-12266
- Li Yang, Adnan Siraj Rakin, Deliang Fan:
RepNet: Efficient On-Device Learning via Feature Reprogramming. 12267-12276
- Zejiang Hou, Minghai Qin, Fei Sun, Xiaolong Ma, Kun Yuan, Yi Xu, Yen-Kuang Chen, Rong Jin, Yuan Xie, Sun-Yuan Kung:
CHEX: CHannel EXploration for CNN Model Compression. 12277-12288
- Miao Yin, Yang Sui , Wanzhao Yang, Xiao Zang, Yu Gong, Bo Yuan:
HODEC: Towards Efficient High-Order DEcomposed Convolutional Neural Networks. 12289-12298
- Lingchen Meng, Hengduo Li, Bor-Chun Chen, Shiyi Lan, Zuxuan Wu, Yu-Gang Jiang, Ser-Nam Lim:
AdaViT: Adaptive Vision Transformers for Efficient Image Recognition. 12299-12308
- Chuanguang Yang, Helong Zhou, Zhulin An , Xue Jiang, Yongjun Xu , Qian Zhang:
Cross-Image Relational Knowledge Distillation for Semantic Segmentation. 12309-12318
- Yongkeon Jeon, Chungman Lee, Eulrang Cho, Yeonju Ro:
Mr.BiQ: Post-Training Non-Uniform Quantization based on Minimizing the Reconstruction Error. 12319-12328
- Yunshan Zhong, Mingbao Lin, Gongrui Nan, Jianzhuang Liu, Baochang Zhang, Yonghong Tian, Rongrong Ji:
IntraQ: Learning Synthetic Images with Intra-Class Heterogeneity for Zero-Shot Network Quantization. 12329-12338
- Manoj Alwani, Yang Wang, Vashisht Madhavan:
DECORE: Deep Compression with Reinforcement Learning. 12339-12349
- Yong Liu, Siqi Mai, Xiangning Chen, Cho-Jui Hsieh, Yang You:
Towards Efficient and Scalable Sharpness-Aware Minimization. 12350-12360

- Simon Schaefer, Daniel Gehrig , Davide Scaramuzza :
AEGNN: Asynchronous Event-based Graph Neural Networks. 12361-12371
- Xinglong Sun, Ali Hassani, Zhangyang Wang, Gao Huang, Humphrey Shi :
DiSparse: Disentangled Sparsification for Multitask Model Compression. 12372-12382
- Anshul Mittal, Kunal Dahiya , Shreya Malani, Janani Ramaswamy, Seba Ann Kuruvilla, Jitendra Ajmera, Keng-hao Chang, Sumeet Agarwal, Purushottam Kar, Manik Varma:
Multi-modal Extreme Classification. 12383-12392
- Georgios Exarchakis , Omar Oubari, Gregor Lenz:
A sampling-based approach for efficient clustering in large datasets. 12393-12402
- Hyungjin Chung, Byeongsu Sim, Jong Chul Ye:
Come-Closer-Diffuse-Faster: Accelerating Conditional Diffusion Models for Inverse Problems through Stochastic Contraction. 12403-12412
- Longguang Wang, Xiaoyu Dong, Yingqian Wang , Li Liu, Wei An, Yulan Guo:
Learnable Lookup Table for Neural Network Quantization. 12413-12423
- Zhenhua Liu, Yunhe Wang, Kai Han, Siwei Ma, Wen Gao:
Instance-Aware Dynamic Neural Network Quantization. 12424-12433
- Qingyan Meng, Mingqing Xiao, Shen Yan, Yisen Wang, Zhouchen Lin, Zhi-Quan Luo:
Training High-Performance Low-Latency Spiking Neural Networks by Differentiation on Spike Representation. 12434-12443
- Sara Elkerdawy, Mostafa Elhoushi, Hong Zhang, Nilanjan Ray:
Fire Together Wire Together: A Dynamic Pruning Approach with Self-Supervised Mask Prediction. 12444-12453
- Linfeng Zhang, Xin Chen, Xiaobing Tu, Pengfei Wan, Ning Xu, Kaisheng Ma:
Wavelet Knowledge Distillation: Towards Efficient Image-to-Image Translation. 12454-12464
- Yichi Zhang , Zhiru Zhang, Lukasz Lew:
PokeBNN: A Binary Pursuit of Lightweight Accuracy. 12465-12475
- Changlin Li, Bohan Zhuang, Guangrun Wang, Xiaodan Liang, Xiaojun Chang , Yi Yang:
Automated Progressive Learning for Efficient Training of Vision Transformers. 12476-12486
- Mathias Parger, Chengcheng Tang, Christopher D. Twigg, Cem Keskin, Robert Wang, Markus Steinberger:
DeltaCNN: End-to-End CNN Inference of Sparse Frame Differences in Videos. 12487-12496
- Vladimir Chikin, Vladimir Kryzhanovskiy:
Channel Balancing for Accurate Quantization of Winograd Convolutions. 12497-12506
- Yan Shi, Junxiong Cai, Yoli Shavit, Tai-Jiang Mu , Wensen Feng, Kai Zhang:
ClusterGNN: Cluster-based Coarse-to-Fine Graph Neural Network for Efficient Feature Matching. 12507-12516
- Paul Wimmer, Jens Mehnert, Alexandru Condurache:
Interspace Pruning: Using Adaptive Filter Representations to Improve Training of Sparse CNNs. 12517-12527
- Ting-An Chen, De-Nian Yang, Ming-Syan Chen:
AlignQ: Alignment Quantization with ADMM-based Correlation Preservation. 12528-12537
- Jierun Chen, Tianlang He, Weipeng Zhuo, Li Ma, Sangtae Ha, S.-H. Gary Chan:
TVConv: Efficient Translation Variant Convolution for Layout-aware Visual Processing. 12538-12548
- Xin Dong, Barbara De Salvo, Meng Li, Chiao Liu, Zhongnan Qu, H. T. Kung, Ziyun Li:
SplitNets: Designing Neural Architectures for Efficient Distributed Computing on Head-Mounted Systems. 12549-12559
- Shian Du, Yihong Luo, Wei Chen, Jian Xu, Delu Zeng:
TO-FLOW: Efficient Continuous Normalizing Flows with Temporal Optimization adjoint with Moving Speed. 12560-12570
- Jieji Ren, Feishi Wang, Jiahao Zhang, Qian Zheng, Mingjun Ren, Boxin Shi:
DiLiGenT10²: A Photometric Stereo Benchmark Dataset with Controlled Shape and Material Variation. 12571-12580
- Satoshi Ikehata:
Universal Photometric Stereo Network using Global Lighting Contexts. 12581-12590
- Berk Kaya, Suryansh Kumar , Carlos Eduardo Porto de Oliveira, Vittorio Ferrari, Luc Van Gool:
Uncertainty-Aware Deep Multi-View Photometric Stereo. 12591-12601
- Daniel Lichy, Soumyadip Sengupta, David W. Jacobs:
Fast Light-Weight Near-Field Photometric Stereo. 12602-12611
- Haiyang Mei, Bo Dong, Wen Dong, Jiaxi Yang, Seung-Hwan Baek , Felix Heide, Pieter Peers, Xiaopeng Wei, Xin Yang:
Glass Segmentation using Intensity and Spectral Polarization Cues. 12612-12621
- Chenyang Lei, Chenyang Qi, Jiaxin Xie, Na Fan, Vladlen Koltun, Qifeng Chen:
Shape from Polarization for Complex Scenes in the Wild. 12622-12631
- Fengting Yang, Xiaolei Huang, Zihan Zhou:
Deep Depth from Focus with Differential Focus Volume. 12632-12641

- Lei Liu, Yuze Chen, Junchi Yan, Yinqiang Zheng :
Optimal LED Spectral Multiplexing for NIR2RGB Translation. 12642-12650
- Yasuto Nagase, Takahiro Kushida, Kenichiro Tanaka, Takuya Funatomi, Yasuhiro Mukaigawa:
Shape from Thermal Radiation: Passive Ranging Using Multi-spectral LWIR Measurements. 12651-12661
- Naama Pearl, Tali Treibitz, Simon Korman:
NAN: Noise-Aware NeRFs for Burst-Denoising. 12662-12671
- Yunhao Zou, Ying Fu:
Estimating Fine-Grained Noise Model via Contrastive Learning. 12672-12681
- Maksim Makarenko, Arturo Burguete-Lopez, Qizhou Wang, Fedor Getman, Silvio Giancola, Bernard Ghanem, Andrea Fratalocchi:
Real-time Hyperspectral Imaging in Hardware via Trained Metasurface Encoders. 12682-12692
- Wuyuan Xie, Tengcong Huang, Miaohui Wang:
MNSRNet: Multimodal Transformer Network for 3D Surface Super-Resolution. 12693-12702
- Zhen Li, Lingli Wang, Xiang Huang, Cihui Pan, Jiaqi Yang:
PhyIR: Physics-based Inverse Rendering for Panoramic Indoor Images. 12703-12713
- Yun-Chun Chen, Haoda Li, Dylan Turpin, Alec Jacobson, Animesh Garg:
Neural Shape Mating: Self-Supervised Object Assembly with Adversarial Shape Priors. 12714-12723
- Xinyu Xu, Yong-Lu Li, Cewu Lu:
Learning to Anticipate Future with Dynamic Context Removal. 12724-12734
- Siyuan Xiang, Anbang Yang, Yanfei Xue, Yaoqing Yang, Chen Feng :
Self-supervised Spatial Reasoning on Multi-View Line Drawings. 12735-12744
- Ruyang Liu, Hao Liu, Ge Li, Haodi Hou, Tinghao Yu, Tao Yang:
Contextual Debiasing for Visual Recognition with Causal Mechanisms. 12745-12755
- Yaqing Ding, Daniel Barath, Jian Yang, Zuzana Kukelova:
Relative Pose from a Calibrated and an Uncalibrated Smartphone Image. 12756-12765
- Guanting Dong, Yueyi Zhang, Hanlin Li, Xiaoyan Sun, Zhiwei Xiong:
Exploiting Rigidity Constraints for LiDAR Scene Flow Estimation. 12766-12775
- Zihan Zhu, Songyou Peng, Viktor Larsson, Weiwei Xu, Hujun Bao, Zhaopeng Cui, Martin R. Oswald , Marc Pollefeys :
NICE-SLAM: Neural Implicit Scalable Encoding for SLAM. 12776-12786
- Tony Ng, Hyo Jin Kim, Vincent T. Lee, Daniel DeTone, Tsun-Yi Yang, Tianwei Shen, Eddy Ilg, Vassileios Balntas, Krystian Mikolajczyk, Chris Sweeney:
NinjaDesc: Content-Concealing Visual Descriptors via Adversarial Learning. 12787-12797
- Axel Barroso Laguna, Yurun Tian, Krystian Mikolajczyk:
ScaleNet: A Shallow Architecture for Scale Estimation. 12798-12808
- Linfei Pan, Marc Pollefeys , Viktor Larsson:
Camera Pose Estimation using Implicit Distortion Models. 12809-12818
- Jianglong Ye , Yuntao Chen, Naiyan Wang, Xiaolong Wang:
GIFS: Neural Implicit Function for General Shape Representation. 12819-12829
- Tianyang Li, Xin Wen, Yu-Shen Liu, Hua Su, Zhizhong Han:
Learning Deep Implicit Functions for 3D Shapes with Dynamic Code Clouds. 12830-12840
- Pablo R. Palafox, Nikolaos Sarafianos, Tony Tung, Angela Dai:
SPAMs: Structured Implicit Parametric Models. 12841-12850
- Li Ma, Xiaoyu Li, Jing Liao , Qi Zhang, Xuan Wang, Jue Wang, Pedro V. Sander:
Deblur-NeRF: Neural Radiance Fields from Blurry Images. 12851-12860
- Abhijit Kundu , Kyle Genova, Xiaoqi Yin, Alireza Fathi, Caroline Pantofaru, Leonidas J. Guibas, Andrea Tagliasacchi, Frank Dellaert, Thomas A. Funkhouser:
Panoptic Neural Fields: A Semantic Object-Aware Neural Scene Representation. 12861-12871
- Kangle Deng , Andrew Liu, Jun-Yan Zhu, Deva Ramanan :
Depth-supervised NeRF: Fewer Views and Faster Training for Free. 12872-12881
- Barbara Roessle, Jonathan T. Barron, Ben Mildenhall, Pratul P. Srinivasan, Matthias Nießner:
Dense Depth Priors for Neural Radiance Fields from Sparse Input Views. 12882-12891
- Tao Hu, Shu Liu, Yilun Chen, Tiancheng Shen, Jiaya Jia:
EfficientNeRF - Efficient Neural Radiance Fields. 12892-12901
- Mijeong Kim, Seonguk Seo, Bohyung Han:
InfoNeRF: Ray Entropy Minimization for Few-Shot Neural Volume Rendering. 12902-12911
- Haithem Turki, Deva Ramanan , Mahadev Satyanarayanan:
Mega-NeRF: Scalable Construction of Large-Scale NeRFs for Virtual Fly-Throughs. 12912-12921
- Konstantinos Rematas, Andrew Liu, Pratul P. Srinivasan, Jonathan T. Barron, Andrea Tagliasacchi, Thomas A. Funkhouser, Vittorio Ferrari:
Urban Radiance Fields. 12922-12932

- Xingyu Chen, Qi Zhang, Xiaoyu Li, Yue Chen, Ying Feng, Xuan Wang, Jue Wang:
Hallucinated Neural Radiance Fields in the Wild. 12933-12942
- Titus Leistner, Radek Mackowiak, Lynton Ardizzone, Ullrich Köthe, Carsten Rother:
Towards Multimodal Depth Estimation from Light Fields. 12943-12951
- Xihao Chen, Zhiwei Xiong, Zhen Cheng, Jiayong Peng, Yueyi Zhang, Zheng-Jun Zha:
Degradation-agnostic Correspondence from Resolution-asymmetric Stereo. 12952-12961
- Donghun Kang, Hyeonjoong Jang, Jungeon Lee, Chong-Min Kyung, Min H. Kim:
Uniform Subdivision of Omnidirectional Camera Space for Efficient Spherical Stereo Matching. 12962-12970
- Gangwei Xu, Junda Cheng, Peng Guo, Xin Yang:
Attention Concatenation Volume for Accurate and Efficient Stereo Matching. 12971-12980
- Zhenxing Mi, Di Chang, Dan Xu:
Generalized Binary Search Network for Highly-Efficient Multi-View Stereo. 12981-12990
- Jiawei Zhang , Xiang Wang, Xiao Bai, Chen Wang, Lei Huang, Yimin Chen, Lin Gu, Jun Zhou, Tatsuya Harada, Edwin R. Hancock:
Revisiting Domain Generalized Stereo Matching Networks from a Feature Consistency Perspective. 12991-13001
- Biyang Liu, Huimin Yu, Guodong Qi :
GraftNet: Towards Domain Generalized Stereo Matching with a Broad-Spectrum and Task-Oriented Feature. 13002-13011
- Weiqin Chuah, Ruwan B. Tennakoon , Reza Hoseinnezhad, Alireza Bab-Hadiashar, David Suter :
ITSA: An Information-Theoretic Approach to Automatic Shortcut Avoidance and Domain Generalization in Stereo Matching Networks. 13012-13022
- Isabella Liu, Edward Yang, Jianyu Tao, Rui Chen, Xiaoshuai Zhang, Qing Ran, Zhu Liu, Hao Su:
ActiveZero: Mixed Domain Learning for Active Stereovision with Zero Annotation. 13023-13032
- Chengtang Yao, Lidong Yu:
FoggyStereo: Stereo Matching with Fog Volume Representation. 13033-13042
- Wen Guo, Xiaoyu Bie, Xavier Alameda-Pineda, Francesc Moreno-Noguer:
Multi-Person Extreme Motion Prediction. 13043-13054
- Nan Xue, Tianfu Wu , Gui-Song Xia, Liangpei Zhang:
Learning Local-Global Contextual Adaptation for Multi-Person Pose Estimation. 13055-13064
- Mohsen Gholami, Bastian Wandt, Helge Rhodin, Rabab Ward, Z. Jane Wang:
AdaptPose: Cross-Dataset Adaptation for 3D Human Pose Estimation by Learnable Motion Generation. 13065-13075
- Lei Jin, Chenyang Xu, Xiaojuan Wang, Yabo Xiao, Yandong Guo, Xuecheng Nie, Jian Zhao:
Single-Stage is Enough: Multi-Person Absolute 3D Pose Estimation. 13076-13085
- Zitian Wang, Xuecheng Nie, Xiaochao Qu, Yunpeng Chen, Si Liu:
Distribution-Aware Single-Stage Models for Multi-Person 3D Pose Estimation. 13086-13095
- Erik Gäßtner, Mykhaylo Andriluka, Hongyi Xu, Cristian Sminchisescu:
Trajectory Optimization for Physics-Based Reconstruction of 3d Human Pose from Monocular Video. 13096-13105
- Yu Zhan, Fenghai Li, Renliang Weng, Wongun Choi:
Ray3D: ray-based 3D human pose estimation for monocular absolute 3D localization. 13106-13115
- Yihan Wang, Muyang Li, Han Cai, Wei-Ming Chen, Song Han:
Lite Pose: Efficient Architecture Design for 2D Human Pose Estimation. 13116-13126
- Xia Xia Xu, Yingguo Gao, Ke Yan, Xue Lin, Qi Zou:
Location-Free Human Pose Estimation. 13127-13136
- Wenhao Li, Hong Liu, Hao Tang, Pichao Wang, Luc Van Gool:
MHFormer: Multi-Hypothesis Transformer for 3D Human Pose Estimation. 13137-13146
- Jian Wang, Lingjie Liu, Weipeng Xu, Kripasindhu Sarkar, Diogo C. Luvizon, Christian Theobalt :
Estimating Egocentric 3D Human Pose in the Wild with External Weak Supervision. 13147-13156
- Xinyu Yi, Yuxiao Zhou, Marc Habermann, Soshi Shimada, Vladislav Golyanik, Christian Theobalt , Feng Xu:
Physical Inertial Poser (PIP): Physics-aware Real-time Human Motion Tracking from Sparse Inertial Sensors. 13157-13168
- Zhijian Yang, Xiaoran Fan, Volkan Isler, Hyunsoo Park:
PoseKernelLifter: Metric Lifting of 3D Human Pose using Sound. 13169-13179
- Erik Gäßtner, Mykhaylo Andriluka, Erwin Coumans, Cristian Sminchisescu:
Differentiable Dynamics for Articulated 3d Human Motion Reconstruction. 13180-13190
- Marko Mihajlovic, Shunsuke Saito, Aayush Bansal, Michael Zollhöfer, Siyu Tang :
COAP: Compositional Articulated Occupancy of People. 13191-13200
- Wen-Li Wei, Jen-Chun Lin, Tyng-Luh Liu, Hong-Yuan Mark Liao:
Capturing Humans in Motion: Temporal-Attentive 3D Human Pose and Shape Estimation from Monocular Video. 13201-13210

- Zhi Chen, Kun Sun, Fan Yang, Wenbing Tao:
SC²-PCR: A Second Order Spatial Compatibility for Efficient and Robust Point Cloud Registration. 13211-13221
- Jinlu Zhang, Zhigang Tu, Jianyu Yang, Yujin Chen, Junsong Yuan :
MixSTE: Seq2seq Mixed Spatio-Temporal Encoder for 3D Human Pose Estimation in Video. 13222-13232
- Yu Sun, Wu Liu, Qian Bao, Yili Fu, Tao Mei, Michael J. Black:
Putting People in their Place: Monocular Regression of 3D People in Depth. 13233-13242
- Sadegh Aliakbarian, Pashmina Cameron, Federica Bogo, Andrew W. Fitzgibbon, Thomas J. Cashman:
FLAG: Flow-based 3D Avatar Generation from Sparse Observations. 13243-13252
- Omid Taheri, Vasileios Choutas, Michael J. Black, Dimitrios Tzionas:
GOAL: Generating 4D Whole-Body Motion for Hand-Object Grasping. 13253-13263
- Chun-Hao P. Huang, Hongwei Yi, Markus Höschle, Matvey Safroshkin, Tsvetelina Alexiadis, Senya Polikovsky, Daniel Scharstein, Michael J. Black:
Capturing and Inferring Dense Full-Body Human-Scene Contact. 13264-13275
- Anastasia Ianina , Nikolaos Sarafianos, Yuanlu Xu , Ignacio Rocco, Tony Tung:
BodyMap: Learning Full-Body Dense Correspondence Map. 13276-13285
- Yuliang Xiu, Jinlong Yang, Dimitrios Tzionas, Michael J. Black:
ICON: Implicit Clothed humans Obtained from Normals. 13286-13296
- Zhanhao Hu, Siyuan Huang, Xiaopei Zhu, Fuchun Sun, Bo Zhang, Xiaolin Hu:
Adversarial Texture for Fooling Person Detectors in the Physical World. 13297-13306
- Xiaopei Zhu, Zhanhao Hu, Siyuan Huang, Jianmin Li, Xiaolin Hu:
Infrared Invisible Clothing: Hiding from Infrared Detectors at Multiple Angles in Real World. 13307-13316
- Ziqi Wang, Marco Loog:
Enhancing Classifier Conservativeness and Robustness by Polynomality. 13317-13326
- Aniruddha Saha, Ajinkya Tejankar, Soroush Abbasi Koohpayegani, Hamed Pirsiavash:
Backdoor Attacks on Self-Supervised Learning. 13327-13336
- Xiangyu Qi, Tinghao Xie, Ruizhe Pan, Jifeng Zhu, Yong Yang, Kai Bu:
Towards Practical Deployment-Stage Backdoor Attack on Deep Neural Networks. 13337-13347
- Jiyang Guan, Zhuozhuo Tu, Ran He, Dacheng Tao:
Few-shot Backdoor Defense Using Shapley Estimation. 13348-13357
- Guanhong Tao, Guangyu Shen, Yingqi Liu, Shengwei An, Qiuling Xu, Shiqing Ma, Pan Li, Xiangyu Zhang:
Better Trigger Inversion Optimization in Backdoor Scanning. 13358-13368
- Binghui Wang, Youqi Li, Pan Zhou:
Bandits for Structure Perturbation-based Black-box Attacks to Graph Neural Networks with Theoretical Guarantees. 13369-13377
- Ozan Özdenizci, Robert Legenstein:
Improving Robustness Against Stealthy Weight Bit-Flip Attacks by Output Code Matching. 13378-13387
- Xiaojun Jia , Yong Zhang, Baoyuan Wu, Ke Ma, Jue Wang, Xiaochun Cao:
LAS-AT: Adversarial Training with Learnable Attack Strategy. 13388-13398
- Tao Li, Yingwen Wu, Sizhe Chen, Kun Fang , Xiaolin Huang:
Subspace Adversarial Training. 13399-13408
- Charles Herrmann, Kyle Sargent, Lu Jiang, Ramin Zabih, Huiwen Chang, Ce Liu, Dilip Krishnan, Deqing Sun:
Pyramid Adversarial Training Improves ViT Performance. 13409-13419
- Zirui Peng, Shaofeng Li, Guoxing Chen, Cheng Zhang, Haojin Zhu, Minhui Xue:
Fingerprinting Deep Neural Networks Globally via Universal Adversarial Perturbations. 13420-13429
- Haiwei Wu, Jiantao Zhou, Jinyu Tian, Jun Liu:
Robust Image Forgery Detection over Online Social Network Shared Images. 13430-13439
- Yusuke Hirota, Yuta Nakashima, Noa Garcia:
Quantifying Societal Bias Amplification in Image Captioning. 13440-13449
- Niv Granot, Ben Feinstein, Assaf Shocher, Shai Bagon, Michal Irani:
Drop the GAN: In Defense of Patches Nearest Neighbors as Single Image Generative Models. 13450-13459
- William S. Peebles, Jun-Yan Zhu, Richard Zhang, Antonio Torralba, Alexei A. Efros, Eli Shechtman:
GAN-Supervised Dense Visual Alignment. 13460-13471
- Yuxin Kong, Canjie Luo, Weihong Ma , Qiyuan Zhu, Shenggao Zhu, Nicholas Jing Yuan, Lianwen Jin:
Look Closer to Supervise Better: One-Shot Font Generation via Component-Based Discriminator. 13472-13481
- Oscar Michel, Roi Bar-On, Richard Liu, Sagie Benaim, Rana Hanocka:
Text2Mesh: Text-Driven Neural Stylization for Meshes. 13482-13492
- Roy Or-El, Xuan Luo, Mengyi Shan, Eli Shechtman, Jeong Joon Park , Ira Kemelmacher-Shlizerman:
StyleSDF: High-Resolution 3D-Consistent Image and Geometry Generation. 13493-13503
- Mariem Mezghanni, Théo Bodrito, Malika Boukenafed, Maks Ovsjanikov:
Physical Simulation Layer for Accurate 3D Modeling. 13504-13513

- Liao Wang, Jiakai Zhang , Xinhang Liu, Fuqiang Zhao, Yanshun Zhang, Yingliang Zhang , Minye Wu, Jingyi Yu, Lan Xu: **Fourier PlenOctrees for Dynamic Radiance Field Rendering in Real-time.** 13514-13524
- Yurui Ren, Xiaoqing Fan, Ge Li, Shan Liu, Thomas H. Li: **Neural Texture Extraction and Distribution for Controllable Person Image Synthesis.** 13525-13534
- Yufeng Zheng, Victoria Fernández Abrevaya, Marcel C. Bühler, Xu Chen, Michael J. Black, Otmar Hilliges: **I M Avatar: Implicit Morphable Head Avatars from Videos.** 13535-13545
- Qiang Wang, Yanhao Zhang, Yun Zheng, Pan Pan: **RCL: Recurrent Continuous Localization for Temporal Action Detection.** 13556-13565
- Nicolae-Catalin Ristea, Neelu Madan, Radu Tudor Ionescu, Kamal Nasrollahi, Fahad Shahbaz Khan, Thomas B. Moeslund , Mubarak Shah: **Self-Supervised Predictive Convolutional Attentive Block for Anomaly Detection.** 13566-13576
- Chao-Yuan Wu, Yanghao Li, Karttikeya Mangalam, Haoqi Fan, Bo Xiong, Jitendra Malik, Christoph Feichtenhofer: **MeMViT: Memory-Augmented Multiscale Vision Transformer for Efficient Long-Term Video Recognition.** 13577-13587
- Jiaojiao Zhao, Yanyi Zhang, Xinyu Li, Hao Chen, Bing Shuai, Mingze Xu, Chunhui Liu, Kaustav Kundu, Yuanjun Xiong, Davide Modolo, Ivan Marsic, Cees G. M. Snoek, Joseph Tighe: **TubeR: Tubelet Transformer for Video Action Detection.** 13588-13597
- Yutao Cui, Cheng Jiang, Limin Wang, Gangshan Wu: **MixFormer: End-to-End Tracking with Iterative Mixed Attention.** 13598-13608
- Feng Li, Hao Zhang , Shilong Liu, Jian Guo, Lionel M. Ni, Lei Zhang: **DN-DETR: Accelerate DETR Training by Introducing Query DeNoising.** 13609-13617
- Cristina Nader Vasconcelos, Vighnesh Birodkar, Vincent Dumoulin: **Proper Reuse of Image Classification Features Improves Object Detection.** 13618-13627
- Wu Zheng, Mingxuan Hong , Li Jiang , Chi-Wing Fu: **Boosting 3D Object Detection by Simulating Multimodality on Point Clouds.** 13628-13637
- Ruotong Wang, Yanqing Shen, Weiliang Zuo, Sanping Zhou, Nanning Zheng: **TransVPR: Transformer-Based Place Recognition with Multi-Level Attention Aggregation.** 13638-13647
- Nirat Saini, Khoi Pham , Abhinav Shrivastava: **Disentangling Visual Embeddings for Attributes and Objects.** 13648-13657
- Chenhongyi Yang, Zehao Huang , Naiyan Wang: **QueryDet: Cascaded Sparse Query for Accelerating High-Resolution Small Object Detection.** 13658-13667
- Xuefeng Du, Xin Wang, Gabriel Gozum, Yixuan Li: **Unknown-Aware Object Detection: Learning What You Don't Know from Videos in the Wild.** 13668-13678
- Nicola Garau , Niccoló Bisagno , Zeno Sambugaro, Nicola Conci: **Interpretable part-whole hierarchies and conceptual-semantic relationships in neural networks.** 13679-13688
- Gowthami Somepalli, Liam Fowl, Arpit Bansal, Ping-Yeh Chiang, Yehuda Dar, Richard G. Baraniuk, Micah Goldblum, Tom Goldstein: **Can Neural Nets Learn the Same Model Twice? Investigating Reproducibility and Double Descent from the Decision Boundary Perspective.** 13689-13698
- Jiacheng Cheng , Nuno Vasconcelos : **Calibrating Deep Neural Networks by Pairwise Constraints.** 13699-13708
- Chen Wang , Yuheng Qiu, Dasong Gao, Sebastian A. Scherer : **Lifelong Graph Learning.** 13709-13718
- Wanyu Lin, Hao Lan, Hao Wang, Baochun Li: **OrphicX: A Causality-Inspired Latent Variable Model for Interpreting Graph Neural Networks.** 13719-13728
- Stephen James, Kentaro Wada, Tristan Laidlow, Andrew J. Davison: **Coarse-to-Fine Q-attention: Efficient Learning for Visual Robotic Manipulation via Discretisation.** 13729-13738
- Jin-Man Park, Ue-Hwan Kim, Seon-Hoon Lee, Jong-Hwan Kim: **Dual Task Learning by Leveraging Both Dense Correspondence and Mis-Correspondence for Robust Change Detection With Imperfect Matches.** 13739-13749
- Brady Zhou, Philipp Krähenbühl: **Cross-view Transformers for real-time Map-view Semantic Segmentation.** 13750-13759
- Will Price, Carl Vondrick, Dima Damen: **UnweaveNet: Unweaving Activity Stories.** 13760-13769
- Reza Ghoddoosian, Isht Dwivedi , Nakul Agarwal, Chiho Choi, Behzad Dariush: **Weakly-Supervised Online Action Segmentation in Multi-View Instructional Videos.** 13770-13780
- Yunhua Zhang, Hazel Doughty , Ling Shao, Cees G. M. Snoek: **Audio-Adaptive Activity Recognition Across Video Domains.** 13781-13790
- Minghao Chen, Fangyun Wei, Chong Li, Deng Cai: **Frame-wise Action Representations for Long Videos via Sequence Contrastive Learning.** 13791-13800

- Florian Richter, Ryan K. Orosco, Michael C. Yip:
Image Based Reconstruction of Liquids from 2D Surface Detections. 13801-13810
- Zhiwu Qing, Shiwei Zhang, Ziyuan Huang , Yi Xu, Xiang Wang, Mingqian Tang, Changxin Gao, Rong Jin, Nong Sang:
Learning from Untrimmed Videos: Self-Supervised Video Representation Learning with Hierarchical Consistency. 13811-13821
- Hazel Doughty , Cees G. M. Snoek:
How Do You Do It? Fine-Grained Action Understanding with Pseudo-Adverbs. 13822-13832
- Sumith Kulal, Jiayuan Mao, Alex Aiken, Jiajun Wu:
Programmatic Concept Learning for Human Motion Description and Synthesis. 13833-13842
- Xudong Lin, Fabio Petroni, Gedas Bertasius, Marcus Rohrbach, Shih-Fu Chang, Lorenzo Torresani:
Learning To Recognize Procedural Activities with Distant Supervision. 13843-13853
- Xuelian Cheng, Huan Xiong, Deng-Ping Fan, Yiran Zhong, Mehrtash Harandi, Tom Drummond, Zongyuan Ge:
Implicit Motion Handling for Video Camouflaged Object Detection. 13854-13863
- Yiming Li, Xiaoshan Yang , Changsheng Xu:
Dynamic Scene Graph Generation via Anticipatory Pre-training. 13864-13873
- Kun Xia, Le Wang, Sanping Zhou, Nanning Zheng, Wei Tang:
Learning to Refactor Action and Co-occurrence Features for Temporal Action Localization. 13874-13883
- Jintao Lin, Haodong Duan, Kai Chen, Dahua Lin, Limin Wang:
OCSampler: Compressing Videos to One Clip with Single-step Sampling. 13884-13893
- Tianshan Liu , Kin-Man Lam:
A Hybrid Egocentric Activity Anticipation Framework via Memory-Augmented Recurrent and One-shot Representation Forecasting. 13894-13903
- Dahun Kim, Jun Xie, Huiyu Wang, Siyuan Qiao, Qihang Yu, Hong-Seok Kim, Hartwig Adam, In So Kweon, Liang-Chieh Chen:
TubeFormer-DeepLab: Video Mask Transformer. 13904-13914
- Bo He, Xitong Yang, Le Kang, Zhiyu Cheng, Xin Zhou, Abhinav Shrivastava:
ASM-Loc: Action-aware Segment Modeling for Weakly-Supervised Temporal Action Localization. 13915-13925
- Zheng Chang, Xinfeng Zhang, Shanshe Wang, Siwei Ma, Wen Gao:
STRPM: A Spatiotemporal Residual Predictive Model for High-Resolution Video Prediction. 13926-13935
- Tomás Soucek, Jean-Baptiste Alayrac, Antoine Miech, Ivan Laptev, Josef Sivic:
Look for the Change: Learning Object States and State-Modifying Actions from Untrimmed Web Videos. 13936-13946
- Congcong Li, Xinyao Wang, Longyin Wen, Dexiang Hong, Tiejian Luo, Libo Zhang:
End-to-End Compressed Video Representation Learning for Generic Event Boundary Detection. 13947-13956
- Liangzhe Yuan, Rui Qian, Yin Cui, Boqing Gong, Florian Schroff, Ming-Hsuan Yang, Hartwig Adam, Ting Liu:
Contextualized Spatio-Temporal Contrastive Learning with Self-Supervision. 13957-13966
- Guang Yu , Siqi Wang, Zhiping Cai, Xinwang Liu, Chuanfu Xu, Chengkun Wu:
Deep Anomaly Discovery from Unlabeled Videos via Normality Advantage and Self-Paced Refinement. 13967-13978
- Matthew Kowal, Mennatullah Siam, Md. Amirul Islam, Neil D. B. Bruce, Richard P. Wildes, Konstantinos G. Derpanis:
A Deeper Dive Into What Deep Spatiotemporal Networks Encode: Quantifying Static vs. Dynamic Information. 13979-13989
- Jue Wang, Gedas Bertasius, Du Tran, Lorenzo Torresani:
Long-Short Temporal Contrastive Learning of Video Transformers. 13990-14000
- Haoqian Wu, Keyu Chen, Yanan Luo, Ruizhi Qiao, Bo Ren, Haozhe Liu, Weicheng Xie, Linlin Shen:
Scene Consistency Representation Learning for Video Scene Segmentation. 14001-14010
- Can Zhang, Tianyu Yang, Junwu Weng, Meng Cao, Jue Wang, Yuexian Zou:
Unsupervised Pre-training for Temporal Action Localization Tasks. 14011-14021
- Taivanbat Badamdorj, Mrigank Rochan, Yang Wang, Li Cheng:
Contrastive Learning for Unsupervised Video Highlight Detection. 14022-14032
- Jue Wang, Lorenzo Torresani:
Deformable Video Transformer. 14033-14042
- Jiewen Yang , Xingbo Dong, Lijun Liu, Chao Zhang, Jiajun Shen , Dahai Yu:
Recurring the Transformer for Video Action Recognition. 14043-14053
- Zongyang Ma, Guan Luo , Jin Gao, Liang Li, Yuxin Chen, Shaoru Wang, Congxuan Zhang , Weiming Hu:
Open-Vocabulary One-Stage Detection with Hierarchical Visual-Language Knowledge Distillation. 14054-14063
- Yu Du, Fangyun Wei, Zihe Zhang, Miaojing Shi, Yue Gao, Guoqi Li:
Learning to Prompt for Open-Vocabulary Object Detection with Vision-Language Model. 14064-14073
- Amanda Cardoso Duarte , Samuel Albanie, Xavier Giró-i-Nieto, Gü̈l Varol:
Sign Language Video Retrieval with Free-Form Textual Queries. 14074-14084
- Sonam Goenka, Zhaocheng Zheng, Ayush Jaiswal, Rakesh Chada, Yue Wu, Varsha Hedau, Pradeep Natarajan:
FashionVLP: Vision Language Transformer for Fashion Retrieval with Feedback. 14085-14095

- Caiyuan Zheng, Hui Li, Seon-Min Rhee, Seungju Han , Jae-Joon Han, Peng Wang:
Pushing the Performance Limit of Scene Text Recognizer without Human Annotation. 14096-14105
- Jun Bao, Buyu Liu, Jun Yu:
ESCNet: Gaze Target Detection with the Understanding of 3D Scenes. 14106-14115
- Chunggi Lee, Seonwook Park , Heon Song, Jeongun Ryu, Sanghoon Kim, Haejoon Kim, Sérgio Pereira , Donggeun Yoo:
Interactive Multi-Class Tiny-Object Detection. 14116-14125
- Xiaoxu Feng, Xiwen Yao, Gong Cheng , Junwei Han:
Weakly Supervised Rotation-Invariant Aerial Object Detection Network. 14126-14135
- Youngwook Kim , Jae-Myung Kim, Zeynep Akata, Jungwoo Lee:
Large Loss Matters in Weakly Supervised Multi-Label Classification. 14136-14145
- Zhixiang Chi, Li Gu, Huan Liu, Yang Wang, Yuanhao Yu, Jin Tang:
MetaFSCIL: A Meta-Learning Approach for Few-Shot Class Incremental Learning. 14146-14155
- Xinlong Wang, Zhiding Yu, Shalini De Mello, Jan Kautz, Anima Anandkumar, Chunhua Shen, José M. Álvarez:
FreeSOLO: Learning to Segment Objects without Annotations. 14156-14166
- Dongli Xu , Jinhong Deng, Wen Li:
Revisiting AP Loss for Dense Object Detection: Adaptive Ranking Pair Selection. 14167-14176
- Hanjun Li, Xingjia Pan, Ke Yan, Fan Tang, Wei-Shi Zheng:
SIOD: Single Instance Annotated Per Category Per Image for Object Detection. 14177-14186
- Xinyu Liu , Wuyang Li , Qiushi Yang, Baopu Li, Yixuan Yuan :
Towards Robust Adaptive Object Detection under Noisy Annotations. 14187-14196
- Liang Zhao, Limin Wang:
Task-specific Inconsistency Alignment for Domain Adaptive Object Detection. 14197-14206
- Lin Sui, Chen-Lin Zhang, Jianxin Wu:
Salvage of Supervision in Weakly Supervised Object Detection. 14207-14216
- Prannay Kaul, Weidi Xie, Andrew Zisserman:
Label, Verify, Correct: A Simple Few Shot Object Detection Method. 14217-14227
- Pingyu Wu, Wei Zhai, Yang Cao:
Background Activation Suppression for Weakly Supervised Object Localization. 14228-14237
- Eunji Kim, Siwon Kim, Jungbeom Lee, Hyunwoo Kim, Sungroh Yoon:
Bridging the Gap between Classification and Localization for Weakly Supervised Object Localization. 14238-14247
- Muli Yang, Yuehua Zhu, Jiaping Yu, Aming Wu, Cheng Deng:
Divide and Conquer: Compositional Experts for Generalized Novel Class Discovery. 14248-14257
- Xin Jin, Tianyu He , Kecheng Zheng, Zhiheng Yin, Xu Shen, Zhen Huang, Ruoyu Feng, Jianqiang Huang, Zhibo Chen, Xian-Sheng Hua:
Cloth-Changing Person Re-identification from A Single Image with Gait Prediction and Regularization. 14258-14267
- Zhipeng Huang, Zhizheng Zhang, Cuiling Lan, Wenjun Zeng , Peng Chu, Quanzeng You, Jiang Wang, Zicheng Liu, Zheng-Jun Zha:
Lifelong Unsupervised Domain Adaptive Person Re-identification with Coordinated Anti-forgetting and Adaptation. 14268-14277
- Zizheng Yang, Xin Jin, Kecheng Zheng, Feng Zhao:
Unleashing Potential of Unsupervised Pre-Training with Intra-Identity Regularization for Person Re-Identification. 14278-14287
- Mouxing Yang, Zhenyu Huang, Peng Hu, Taihao Li, Jiancheng Lv, Xi Peng:
Learning with Twin Noisy Labels for Visible-Infrared Person Re-Identification. 14288-14297
- Karsten Roth, Latha Pemula, Joaquin Zepeda, Bernhard Schölkopf, Thomas Brox, Peter V. Gehler:
Towards Total Recall in Industrial Anomaly Detection. 14298-14308
- Yunqiu Xu, Yifan Sun, Zongxin Yang, Jiaxu Miao, Yi Yang:
H2FA R-CNN: Holistic and Hierarchical Feature Alignment for Cross-domain Weakly Supervised Object Detection. 14309-14319
- Xiao-Chang Liu , Yongliang Yang, Peter Hall :
Geometric and Textural Augmentation for Domain Gap Reduction. 14320-14330
- Jiangwei Xie, Shipeng Yan, Xuming He:
General Incremental Learning with Domain-aware Categorical Representations. 14331-14340
- Wenxuan Wang, Xuelin Qian, Yanwei Fu , Xiangyang Xue:
DST: Dynamic Substitute Training for Data-free Black-box Attack. 14341-14350
- Ruibin Wang, Yibo Yang, Dacheng Tao:
ART-Point: Improving Rotation Robustness of Point Cloud Classifiers via Adversarial Rotation. 14351-14360
- Binbin Chen, Weijie Chen, Shicai Yang , Yunyi Xuan, Jie Song, Di Xie, Shiliang Pu, Mingli Song, Yueting Zhuang:
Label Matching Semi-Supervised Object Detection. 14361-14370

- Deep Shankar Pandey, Qi Yu :
Multidimensional Belief Quantification for Label-Efficient Meta-Learning. 14371-14380
- Noo-Ri Kim, Jee-Hyong Lee:
Propagation Regularizer for Semi-supervised Learning with Extremely Scarce Labeled Samples. 14381-14390
- Yang Liu, Weifeng Zhang, Chao Xiang, Tu Zheng, Deng Cai, Xiaofei He:
Learning to Affiliate: Mutual Centralized Learning for Few-shot Classification. 14391-14400
- Fan Yang, Kai Wu, Shuyi Zhang, Guannan Jiang, Yong Liu, Feng Zheng, Wei Zhang, Chengjie Wang, Long Zeng:
Class-Aware Contrastive Semi-Supervised Learning. 14401-14410
- Chenxin Tao, Honghui Wang, Xizhou Zhu, Jiahua Dong, Shiji Song, Gao Huang, Jifeng Dai:
Exploring the Equivalence of Siamese Self-Supervised Learning via A Unified Gradient Framework. 14411-14420
- Chaoning Zhang , Kang Zhang, Trung X. Pham, Axi Niu, Zhinan Qiao, Chang D. Yoo, In So Kweon:
Dual Temperature Helps Contrastive Learning Without Many Negative Samples: Towards Understanding and Simplifying MoCo. 14421-14430
- Lang Huang, Shan You, Mingkai Zheng, Fei Wang, Chen Qian, Toshihiko Yamasaki:
Learning Where to Learn in Cross-View Self-Supervised Learning. 14431-14440
- Yunrui Zhao, Qianqian Xu, Yangbangyan Jiang, Peisong Wen, Qingming Huang:
Dist-PU: Positive-Unlabeled Learning from a Label Distribution Perspective. 14441-14450
- Mingkai Zheng, Shan You, Lang Huang, Fei Wang, Chen Qian, Chang Xu:
SimMatch: Semi-supervised Learning with Similarity Matching. 14451-14461
- Peng Mi, Jianghang Lin, Yiyi Zhou, Yunhang Shen , Gen Luo, Xiaoshuai Sun, Liujuan Cao, Rongrong Fu, Qiang Xu, Rongrong Ji:
Active Teacher for Semi-Supervised Object Detection. 14462-14471
- Ismail Elezi, Zhiding Yu, Anima Anandkumar, Laura Leal-Taixé, José M. Álvarez:
Not All Labels Are Equal: Rationalizing The Labeling Costs for Training Object Detection. 14472-14481
- Adrian Ziegler, Yuki M. Asano:
Self-Supervised Learning of Object Parts for Semantic Segmentation. 14482-14491
- Jongmok Kim, Jooyoung Jang, Seunghyeon Seo, Jisoo Jeong, Jongkeun Na, Nojun Kwak:
MUM: Mix Image Tiles and UnMix Feature Tiles for Semi-Supervised Object Detection. 14492-14501
- Qiushan Guo, Yao Mu, Jianyu Chen, Tianqi Wang, Yizhou Yu, Ping Luo:
Scale-Equivalent Distillation for Semi-Supervised Object Detection. 14502-14511
- Ed Pizzi, Sreya Dutta Roy, Sugosh Nagavara Ravindra, Priya Goyal, Matthijs Douze:
A Self-Supervised Descriptor for Image Copy Detection. 14512-14522
- Yangtao Wang, Xi Shen, Shell Xu Hu, Yuan Yuan, James L. Crowley, Dominique Vaufreydaz:
Self-Supervised Transformers for Unsupervised Object Discovery using Normalized Cut. 14523-14533
- Philip Chikontwe, Soopil Kim, Sang Hyun Park:
CAD: Co-Adapting Discriminative Features for Improved Few-Shot Classification. 14534-14543
- Jie Ling, Lei Liao, Meng Yang, Jia Shuai:
Semi-Supervised Few-shot Learning via Multi-Factor Clustering. 14544-14553
- Yue Fan, Dengxin Dai, Anna Kukleva, Bernt Schiele :
CoSSL: Co-Learning of Representation and Classifier for Imbalanced Semi-Supervised Learning. 14554-14564
- Rundong He, Zhongyi Han, Xiankai Lu, Yilong Yin:
Safe-Student for Safe Deep Semi-Supervised Learning with Unseen-Class Unlabeled Data. 14565-14574
- Sucheng Ren, Huiyu Wang, Zhengqi Gao, Shengfeng He , Alan L. Yuille, Yuyin Zhou, Cihang Xie :
A Simple Data Mixing Prior for Improving Self-Supervised Learning. 14575-14584
- Amir Bar, Xin Wang, Vadim Kantorov, Colorado J. Reed, Roei Herzig, Gal Chechik, Anna Rohrbach, Trevor Darrell, Amir Globerson:
DETReg: Unsupervised Pretraining with Region Priors for Object Detection. 14585-14595
- Arun Balajee Vasudevan, Dengxin Dai, Luc Van Gool:
Sound and Visual Representation Learning with Multiple Pretraining Tasks. 14596-14606
- Zhaowen Li, Yousong Zhu, Fan Yang, Wei Li, Chaoyang Zhao, Yingying Chen, Zhiyang Chen, Jiahao Xie, Liwei Wu, Rui Zhao, Ming Tang, Jinqiao Wang:
UniVIP: A Unified Framework for Self-Supervised Visual Pre-training. 14607-14616
- Lei Zhu, Qi She, Qian Chen, Yunfei You, Boyu Wang, Yanye Lu:
Weakly Supervised Object Localization as Domain Adaption. 14617-14626
- Xudong Wang, Zhirong Wu, Long Lian, Stella X. Yu:
Debiased Learning from Naturally Imbalanced Pseudo-Labels. 14627-14637
- Hui Tang, Kui Jia:
Towards Discovering the Effectiveness of Moderately Confident Samples for Semi-Supervised Learning. 14638-14647
- Chen Wei, Haoqi Fan, Saining Xie, Chao-Yuan Wu, Alan L. Yuille, Christoph Feichtenhofer:
Masked Feature Prediction for Self-Supervised Visual Pre-Training. 14648-14658

- Jeany Son:
Contrastive Learning for Space-time Correspondence via Self-cycle Consistency. 14659-14668
- Bing Shuai, Xinyu Li, Kaustav Kundu, Joseph Tighe:
Id-Free Person Similarity Learning. 14669-14679
- Akash Kumar, Yogesh Singh Rawat:
End-to-End Semi-Supervised Learning for Video Action Detection. 14680-14690
- Jungin Park, Jiyoung Lee , Ig-Jae Kim, Kwanghoon Sohn:
Probabilistic Representations for Video Contrastive Learning. 14691-14701
- Lijin Yang, Yifei Huang, Yusuke Sugano, Yoichi Sato:
Interact before Align: Leveraging Cross-Modal Knowledge for Domain Adaptive Action Recognition. 14702-14712
- Rui Wang, Dongdong Chen, Zuxuan Wu, Yinpeng Chen, Xiyang Dai, Mengchen Liu, Yu-Gang Jiang, Luwei Zhou, Lu Yuan:
BEVT: BERT Pretraining of Video Transformers. 14713-14723
- Muhammad Zaigham Zaheer, Arif Mahmood, Muhammad Haris Khan, Mattia Segù, Fisher Yu, Seung-Ik Lee:
Generative Cooperative Learning for Unsupervised Video Anomaly Detection. 14724-14734
- Muhammad Jehanzeb Mirza, Jakub Micorek, Horst Possegger , Horst Bischof:
The Norm Must Go On: Dynamic Unsupervised Domain Adaptation by Normalization. 14745-14755
- Ning Gao, Hanna Ziesche, Ngo Anh Vien, Michael Volpp, Gerhard Neumann:
What Matters For Meta-Learning Vision Regression Tasks? 14756-14766
- Ankit Goyal, Arsalan Mousavian, Chris Paxton, Yu-Wei Chao, Brian Okorn, Jia Deng, Dieter Fox:
IFOR: Iterative Flow Minimization for Robotic Object Rearrangement. 14767-14777
- Ziang Cao, Ziyuan Huang , Liang Pan , Shiwei Zhang, Ziwei Liu, Changhong Fu:
TCTrack: Temporal Contexts for Aerial Tracking. 14778-14788
- Liu Liu , Wenqiang Xu, Haoyuan Fu, Sucheng Qian, Qiaojun Yu, Yang Han, Cewu Lu:
AKB-48: A Real-World Articulated Object Knowledge Base. 14789-14798
- Guangchi Fang, Qingyong Hu, Hanyun Wang , Yiling Xu, Yulan Guo:
3DAC: Learning Attribute Compression for Point Clouds. 14799-14808
- Apoorv Khandelwal, Luca Weihs, Roozbeh Mottaghi, Aniruddha Kembhavi:
Simple but Effective: CLIP Embeddings for Embodied AI. 14809-14818
- Kai Ye, Siyan Dong, Qingnan Fan, He Wang, Li Yi, Fei Xia, Jue Wang, Baoquan Chen:
Multi-Robot Active Mapping via Neural Bipartite Graph Matching. 14819-14828
- Samir Yitzhak Gadre, Kiana Ehsani, Shuran Song , Roozbeh Mottaghi:
Continuous Scene Representations for Embodied AI. 14829-14839
- Klemen Kotar, Roozbeh Mottaghi:
Interactron: Embodied Adaptive Object Detection. 14840-14849
- Tommaso Campari, Leonardo Lamanna, Paolo Traverso, Luciano Serafini, Lamberto Ballan :
Online Learning of Reusable Abstract Models for Object Goal Navigation. 14850-14859
- Yan Xu, Kwan-Yee Lin, Guofeng Zhang, Xiaogang Wang , Hongsheng Li :
RNNPose: Recurrent 6-DoF Object Pose Refinement with Robust Correspondence Field Estimation and Pose Optimization. 14860-14870
- Taeyeop Lee, Byeong-Uk Lee, Inkyu Shin, Jaesung Choe, Ukcheol Shin , In So Kweon, Kuk-Jin Yoon:
UDA-COPE: Unsupervised Domain Adaptation for Category-level Object Pose Estimation. 14871-14880
- Nathaniel Merrill, Yuliang Guo, Xingxing Zuo, Xinyu Huang, Stefan Leutenegger, Xi Peng, Liu Ren, Guoquan Huang :
Symmetry and Uncertainty-Aware Object SLAM for 6DoF Object Pose Estimation. 14881-14890
- Xufang Pang, Feng Li, Ning Ding, Xiaopin Zhong:
Upright-Net: Learning Upright Orientation for 3D Point Cloud. 14891-14899
- Xueyu Wang, Jiajun Huang, Siqi Ma, Surya Nepal , Chang Xu:
DeepFake Disrupter: The Detector of DeepFake Is My Friend. 14900-14909
- Mengtian Li, Yuan Xie, Yunhang Shen , Bo Ke, Ruizhi Qiao, Bo Ren, Shaohui Lin, Lizhuang Ma:
HybridCR: Weakly-Supervised 3D Point Cloud Semantic Segmentation via Hybrid Contrastive Regularization. 14910-14919
- Sahar Abdelnabi, Rakibul Hasan, Mario Fritz:
Open-Domain, Content-based, Multi-modal Fact-checking of Out-of-Context Images via Online Resources. 14920-14929
- Alexandros Haliassos, Rodrigo Mira , Stavros Petridis, Maja Pantic:
Leveraging Real Talking Faces via Self-Supervision for Robust Forgery Detection. 14930-14942
- Jiang Liu, Alexander Levine, Chun Pong Lau , Rama Chellappa, Soheil Feizi:
Segment and Complete: Defending Object Detectors against Adversarial Patch Attacks with Robust Patch Detection. 14953-14962
- Yifeng Xiong , Jiadong Lin, Min Zhang, John E. Hopcroft, Kun He:
Stochastic Variance Reduced Ensemble Adversarial Attack for Boosting the Adversarial Transferability. 14963-



14972

- Jianping Zhang, Weibin Wu, Jen-tse Huang , Yizhan Huang, Wenxuan Wang, Yuxin Su, Michael R. Lyu: **Improving Adversarial Transferability via Neuron Attribution-based Attacks.** 14973-14982
- Yingqi Liu, Guangyu Shen, Guanhong Tao, Zhenting Wang, Shiqing Ma, Xiangyu Zhang: **Complex Backdoor Detection by Symmetric Feature Differencing.** 14983-14993
- Shengshan Hu, Xiaogeng Liu, Yechao Zhang, Minghui Li , Leo Yu Zhang, Hai Jin, Libing Wu: **Protecting Facial Privacy: Generating Adversarial Identity Masks via Style-robust Makeup Transfer.** 14994-15003
- Zikui Cai, Shantanu Rane, Alejandro E. Brito, Chengyu Song , Srikanth V. Krishnamurthy , Amit K. Roy-Chowdhury, M. Salman Asif: **Zero-Query Transfer Attacks on Context-Aware Object Detectors.** 15004-15014
- Yunjian Zhang, Yanwei Liu, Jinxia Liu, Jingbo Miao, Antonios Argyriou, Liming Wang, Zhen Xu: **360-Attack: Distortion-Aware Perturbations from Perspective-Views.** 15015-15024
- Mostafa Kahla, Si Chen, Hoang Anh Just, Ruoxi Jia: **Label-Only Model Inversion Attacks via Boundary Repulsion.** 15025-15033
- Daksh Thapar, Aditya Nigam, Chetan Arora: **Merry Go Round: Rotate a Frame and Fool a DNN.** 15034-15043
- Zhipeng Wei, Jingjing Chen , Zuxuan Wu, Yu-Gang Jiang: **Cross-Modal Transferable Adversarial Attacks from Images to Videos.** 15044-15053
- Zhenting Wang, Juan Zhai , Shiqing Ma: **BppAttack: Stealthy and Efficient Trojan Attacks against Deep Neural Networks via Image Quantization and Contrastive Adversarial Learning.** 15054-15063
- Chaoning Zhang , Philipp Benz, Adil Karjauv, Jae-Won Cho, Kang Zhang, In So Kweon: **Investigating Top-k White-Box and Transferable Black-box Attack.** 15064-15073
- Yan Feng, Baoyuan Wu, Yanbo Fan, Li Liu, Zhifeng Li, Shu-Tao Xia: **Boosting Black-Box Attack with Partially Transferred Conditional Adversarial Distribution.** 15074-15083
- Ye Liu, Yaya Cheng, Lianli Gao, Xianglong Liu, Qilong Zhang, Jingkuan Song: **Practical Evaluation of Adversarial Robustness via Adaptive Auto Attack.** 15084-15093
- Jie Zhang, Bo Li, Jianghe Xu, Shuang Wu, Shouhong Ding, Lei Zhang, Chao Wu: **Towards Efficient Data Free Blackbox Adversarial Attack.** 15094-15104
- Byung-Kwan Lee, Junho Kim, Yong Man Ro: **Masking Adversarial Damage: Finding Adversarial Saliency for Robust and Sparse Network.** 15105-15115
- Hadi Salman, Saachi Jain, Eric Wong, Aleksander Madry: **Certified Patch Robustness via Smoothed Vision Transformers.** 15116-15126
- Zhaoyu Chen , Bo Li, Jianghe Xu, Shuang Wu, Shouhong Ding, Wenqiang Zhang: **Towards Practical Certifiable Patch Defense with Vision Transformer.** 15127-15137
- Qingzhao Zhang, Shengtuo Hu, Jiachen Sun, Qi Alfred Chen, Z. Morley Mao: **On Adversarial Robustness of Trajectory Prediction for Autonomous Vehicles.** 15138-15147
- Gabriel Pérez S., Juan C. Pérez, Motasem Alfarra, Silvio Giancola, Bernard Ghanem : **3DeformRS: Certifying Spatial Deformations on Point Clouds.** 15148-15158
- Zachary Berger, Parth Agrawal, Tian Yu Liu, Stefano Soatto, Alex Wong : **Stereoscopic Universal Perturbations across Different Architectures and Datasets.** 15159-15169
- Tianlong Chen, Peihao Wang, Zhiwen Fan, Zhangyang Wang: **Aug-NeRF: Training Stronger Neural Radiance Fields with Triple-Level Physically-Grounded Augmentations.** 15170-15181
- Qiuling Xu, Guanhong Tao, Xiangyu Zhang: **Bounded Adversarial Attack on Deep Content Features.** 15182-15191
- Zhendong Zhao, Xiaojun Chen, Yuexin Xuan, Ye Dong, Dakui Wang, Kaitai Liang: **DEFEAT: Deep Hidden Feature Backdoor Attacks by Imperceptible Perturbation and Latent Representation Constraints.** 15192-15201
- Tianyu Pang , Huishuai Zhang, Di He, Yinpeng Dong, Hang Su, Wei Chen, Jun Zhu, Tie-Yan Liu: **Two Coupled Rejection Metrics Can Tell Adversarial Examples Apart.** 15202-15212
- Giulio Lovisotto, Nicole Finnie, Mauricio Munoz, Chaithanya Kumar Mummadli, Jan Hendrik Metzen: **Give Me Your Attention: Dot-Product Attention Considered Harmful for Adversarial Patch Robustness.** 15213-15222
- Junyoung Byun, Seungju Cho, Myung-Joon Kwon, Heeseon Kim, Changick Kim: **Improving the Transferability of Targeted Adversarial Examples through Object-Based Diverse Input.** 15223-15232
- Linjun Zhou, Peng Cui, Xingxuan Zhang, Yinan Jiang, Shiqiang Yang: **Adversarial Eigen Attack on BlackBox Models.** 15233-15241
- Qibing Ren, Qingquan Bao, Runzhong Wang, Junchi Yan: **Appearance and Structure Aware Robust Deep Visual Graph Matching: Attack, Defense and Beyond.** 15242-15251

- Gaojie Jin, Xinping Yi, Wei Huang , Sven Schewe , Xiaowei Huang:
Enhancing Adversarial Training with Second-Order Statistics of Weights. 15252-15262
- Sunandini Sanyal, Sravanti Addepalli, R. Venkatesh Babu :
Towards Data-Free Model Stealing in a Hard Label Setting. 15263-15272
- Kaidong Li, Ziming Zhang, Cuncong Zhong, Guanghui Wang:
Robust Structured Declarative Classifiers for 3D Point Clouds: Defending Adversarial Attacks with Implicit Gradients. 15273-15283
- Naufal Suryanto , Yongsu Kim, Hyoeun Kang, Harashta Tatimma Larasati, Youngyeo Yun, Thi-Thu-Huong Le , Hunmin Yang, Se-Yoon Oh, Howon Kim:
DTA: Physical Camouflage Attacks using Differentiable Transformation Network. 15284-15293
- Cheng Luo, Qinliang Lin, Weicheng Xie, Bizhu Wu, Jinheng Xie, Linlin Shen:
Frequency-driven Imperceptible Adversarial Attack on Semantic Similarity. 15294-15303
- Mo Zhou, Vishal M. Patel:
Enhancing Adversarial Robustness for Deep Metric Learning. 15304-15313
- Qidong Huang, Xiaoyi Dong, Dongdong Chen, Hang Zhou, Weiming Zhang, Nenghai Yu:
Shape-invariant 3D Adversarial Point Clouds. 15314-15323
- Yiqi Zhong, Xianming Liu, Deming Zhai, Junjun Jiang, Xiangyang Ji:
Shadows can be Dangerous: Stealthy and Effective Physical-world Adversarial Attack by Natural Phenomenon. 15324-15333
- Xuxiang Sun, Gong Cheng , Hongda Li, Lei Pei, Junwei Han:
Exploring Effective Data for Surrogate Training Towards Black-box Attack. 15334-15343
- Simin Chen, Zihe Song, Mirazul Haque, Cong Liu, Wei Yang:
NICGSlowDown: Evaluating the Efficiency Robustness of Neural Image Caption Generation Models. 15344-15353
- Matthew Walmer, Karan Sikka, Indranil Sur, Abhinav Shrivastava, Susmit Jha:
Dual-Key Multimodal Backdoors for Visual Question Answering. 15354-15364
- Vishal Asnani, Xi Yin, Tal Hassner, Sijia Liu, Xiaoming Liu:
Proactive Image Manipulation Detection. 15365-15374
- Bingqian Lin, Yi Zhu, Zicong Chen, Xiwen Liang, Jianzhuang Liu, Xiaodan Liang:
ADAPT: Vision-Language Navigation with Modality-Aligned Action Prompts. 15375-15385
- Jialu Li, Hao Tan, Mohit Bansal:
Envedit: Environment Editing for Vision-and-Language Navigation. 15386-15396
- Yanyuan Qiao, Yuankai Qi , Yicong Hong, Zheng Yu, Peng Wang, Qi Wu:
HOP: History-and-Order Aware Pretraining for Vision-and-Language Navigation. 15397-15406
- Su Wang, Ceslee Montgomery, Jordi Orbay, Vighnesh Birodkar, Aleksandra Faust , Izzeddin Gur, Natasha Jaques, Austin Waters, Jason Baldridge, Peter Anderson:
Less is More: Generating Grounded Navigation Instructions from Landmarks. 15407-15417
- Yicong Hong, Zun Wang, Qi Wu, Stephen Gould:
Bridging the Gap Between Learning in Discrete and Continuous Environments for Vision-and-Language Navigation. 15418-15428
- Jinyu Chen, Chen Gao, Erli Meng, Qiong Zhang, Si Liu:
Reinforced Structured State-Evolution for Vision-Language Navigation. 15429-15438
- Georgios Georgakis, Karl Schmeckpeper, Karan Wanchoo, Soham Dan, Eleni Miltsakaki, Dan Roth, Kostas Daniilidis:
Cross-modal Map Learning for Vision and Language Navigation. 15439-15449
- Hanqing Wang, Wei Liang, Jianbing Shen, Luc Van Gool, Wenguan Wang :
Counterfactual Cycle-Consistent Learning for Instruction Following and Generation in Vision-Language Navigation. 15450-15460
- Chan Hee Song, Jihyung Kil, Tai-Yu Pan, Brian M. Sadler, Wei-Lun Chao, Yu Su:
One Step at a Time: Long-Horizon Vision-and-Language Navigation with Milestones. 15461-15470
- Tao Liang, Guosheng Lin, Mingyang Wan, Tianrui Li , Guojun Ma, Fengmao Lv:
Expanding Large Pre-trained Unimodal Models with Multimodal Information Injection for Image-Text Multimodal Classification. 15471-15480
- Jiabo Ye, Junfeng Tian, Ming Yan, Xiaoshan Yang , Xuwu Wang, Ji Zhang, Liang He, Xin Lin:
Shifting More Attention to Visual Backbone: Query-modulated Refinement Networks for End-to-End Visual Grounding. 15481-15491
- Haojun Jiang, Yuanze Lin, Dongchen Han, Shiji Song, Gao Huang:
Pseudo-Q: Generating Pseudo Language Queries for Visual Grounding. 15492-15502
- Shijia Huang, Yilun Chen, Jiaya Jia, Liwei Wang:
Multi-View Transformer for 3D Visual Grounding. 15503-15512
- Sijia Chen, Baochun Li:
Multi-Modal Dynamic Graph Transformer for Visual Grounding. 15513-15522



- Effrosyni Mavroudi, René Vidal:
Weakly-Supervised Generation and Grounding of Visual Descriptions with Conditional Generative Models. 15523-15533
- Minghang Zheng, Yanjie Huang, Qingchao Chen , Yuxin Peng, Yang Liu:
Weakly Supervised Temporal Sentence Grounding with Gaussian-based Contrastive Proposal Learning. 15534-15543
- Chen Liang, Wenguan Wang , Tianfei Zhou, Yi Yang:
Visual Abductive Reasoning. 15544-15554
- Yifeng Zhang, Ming Jiang, Qi Zhao:
Query and Attention Augmentation for Knowledge-Based Explainable Reasoning. 15555-15564
- Shi Chen , Qi Zhao:
REX: Reasoning-aware and Grounded Explanation. 15565-15574
- Arushi Goel, Basura Fernando, Frank Keller, Hakan Bilen :
Not All Relations are Equal: Mining Informative Labels for Scene Graph Generation. 15575-15585
- Chao Lou, Wenjuan Han, Yuhuan Lin, Zilong Zheng:
Unsupervised Vision-Language Parsing: Seamlessly Bridging Visual Scene Graphs with Language Structures via Dependency Relationships. 15586-15595
- Chiao-An Yang, Cheng-Yo Tan, Wan-Cyuan Fan, Cheng-Fu Yang, Meng-Lin Wu, Yu-Chiang Frank Wang:
Scene Graph Expansion for Semantics-Guided Image Outpainting. 15596-15605
- Jinhui Yang, Xianyu Chen, Ming Jiang, Shi Chen , Louis Wang, Qi Zhao:
VisualHow: Multimodal Problem Solving. 15606-15616
- Amanpreet Singh, Ronghang Hu, Vedanuj Goswami, Guillaume Couairon, Wojciech Galuba, Marcus Rohrbach, Douwe Kiela:
FLAVA: A Foundational Language And Vision Alignment Model. 15617-15629
- Jiali Duan, Liqun Chen, Son Tran, Jinyu Yang, Yi Xu, Belinda Zeng, Trishul Chilimbi:
Multi-modal Alignment using Representation Codebook. 15630-15639
- Kun Zhang , Zhendong Mao, Quan Wang, Yongdong Zhang:
Negative-Aware Attention Framework for Image-Text Matching. 15640-15649
- Jinyu Yang, Jiali Duan, Son Tran, Yi Xu, Sampath Chanda, Liqun Chen, Belinda Zeng, Trishul Chilimbi, Junzhou Huang:
Vision-Language Pre-Training with Triple Contrastive Learning. 15650-15659
- Sibo Song, Jianqiang Wan, Zhibo Yang, Jun Tang, Wenqing Cheng, Xiang Bai, Cong Yao:
Vision-Language Pre-Training for Boosting Scene Text Detectors. 15660-15670
- Haoyu Lu, Nanyi Fei, Yuqi Huo, Yizhao Gao, Zhiwu Lu, Ji-Rong Wen:
COTS: Collaborative Two-Stream Vision-Language Pre-Training Model for Cross-Modal Retrieval. 15671-15680
- Zhi-Hao Lin, Wei-Chiu Ma, Hao-Yu Hsu, Yu-Chiang Frank Wang, Shenlong Wang:
NeurMiPs: Neural Mixture of Planar Experts for View Synthesis. 15681-15691
- Ang Cao, Chris Rockwell, Justin Johnson:
FWD: Real-time Novel View Synthesis with Forward Warping and Depth. 15692-15703
- Tewodros Habtegebrial, Christiano Couto Gava, Marcel Rogge, Didier Stricker , Varun Jampani:
SOMSI: Spherical Novel View Synthesis with Soft Occlusion Multi-Sphere Images. 15704-15713
- Shaohan Li, Yunpeng Shi , Gilad Lerman :
Fast, Accurate and Memory-Efficient Partial Permutation Synchronization. 15714-15722
- Daniel Barath, Luca Cavalli, Marc Pollefeys :
Learning to Find Good Models in RANSAC. 15723-15732
- Evgeniy Martyushev, Jana Vráblíková , Tomás Pajdla:
Optimizing Elimination Templates by Greedy Parameter Search. 15733-15743
- Chiang-Heng Chien, Hongyi Fan, Ahmad Abdelfattah, Elias P. Tsigaridas, Stanimire Tomov, Benjamin B. Kimia:
GPU-Based Homotopy Continuation for Minimal Problems in Computer Vision. 15744-15755
- Seong Hun Lee , Javier Civera:
HARA: A Hierarchical Approach for Robust Rotation Averaging. 15756-15765
- Heng Li, Zhaopeng Cui, Shuaicheng Liu, Ping Tan:
RAGO: Recurrent Graph Optimizer For Multiple Rotation Averaging. 15766-15775
- Pedro Miraldo, José Pedro Iglesias:
A Unified Model for Line Projections in Catadioptric Cameras with Rotationally Symmetric Mirrors. 15776-15785
- Dong Wei, Yi Wan, Yongjun Zhang, Xinyi Liu , Bin Zhang, Xiqi Wang:
ELSR: Efficient Line Segment Reconstruction with Planes and Points Guidance. 15786-15794
- Fangyin Wei, Rohan Chabra, Lingni Ma, Christoph Lassner, Michael Zollhöfer, Szymon Rusinkiewicz , Chris Sweeney, Richard A. Newcombe, Mira Slavcheva:
Self-supervised Neural Articulated Shape and Appearance Models. 15795-15805
- Hsiao-Yu Chen, Edith Tretschk, Tuur Stuyck, Petr Kadlecak, Ladislav Kavan, Etienne Vouga, Christoph Lassner:
Virtual Elastic Objects. 15806-15816

- Kunhong Li, Longguang Wang, Li Liu, Qing Ran, Kai Xu, Yulan Guo:
Decoupling Makes Weakly Supervised Local Feature Better. 15817-15827
- Karl D. D. Willis, Pradeep Kumar Jayaraman, Hang Chu, Yunsheng Tian, Yifei Li, Daniele Grandi, Aditya Sanghi, Linh Tran, Joseph G. Lambourne, Armando Solar-Lezama , Wojciech Matusik:
JoinABLE: Learning Bottom-up Assembly of Parametric CAD Joints. 15828-15839
- Jiancheng Yang , Udaranga Wickramasinghe, Bingbing Ni, Pascal Fua:
ImplicitAtlas: Learning Deformable Shape Templates in Medical Imaging. 15840-15850
- Ruizhi Shao, Hongwen Zhang , He Zhang, Mingjia Chen , Yanpei Cao, Tao Yu, Yebin Liu:
DoubleField: Bridging the Neural Surface and Radiance Fields for High-fidelity Human Reconstruction and Rendering. 15851-15861
- Tianhan Xu, Yasuhiro Fujita, Eiichi Matsumoto:
Surface-Aligned Neural Radiance Fields for Controllable 3D Human Synthesis. 15862-15871
- Zerong Zheng , Han Huang, Tao Yu, Hongwen Zhang , Yandong Guo, Yebin Liu:
Structured Local Radiance Fields for Human Avatar Modeling. 15872-15882
- Hao Zhao, Jinsong Zhang, Yu-Kun Lai, Zerong Zheng, Yingdi Xie, Yebin Liu, Kun Li:
High-Fidelity Human Avatars from a Single RGB Camera. 15883-15892
- Christian Diller, Thomas A. Funkhouser, Angela Dai:
Forecasting Characteristic 3D Poses of Human Actions. 15893-15902
- Wei-Chiu Ma, Anqi Joyce Yang, Shenlong Wang, Raquel Urtasun, Antonio Torralba:
Virtual Correspondence: Humans as a Cue for Extreme-View Geometry. 15903-15913
- Bharat Lal Bhatnagar, Xianghui Xie, Ilya A. Petrov, Cristian Sminchisescu, Christian Theobalt , Gerard Pons-Moll:
BEHAVE: Dataset and Method for Tracking Human Object Interactions. 15914-15925
- Xinkle Li, Henghui Ding , Zekun Tong, Yuwei Wu , Yeow Meng Chee :
Primitive3D: 3D Object Dataset Synthesis from Randomly Assembled Primitives. 15926-15936
- Fabio Tosi, Pierluigi Zama Ramirez , Matteo Poggi, Samuele Salti, Stefano Mattoccia , Luigi Di Stefano:
RGB-Multispectral Matching: Dataset, Learning Methodology, Evaluation. 15937-15947
- Ruslan Rakhimov , Andrei-Timotei Ardelean, Victor Lempitsky, Evgeny Burnaev :
NPBG++: Accelerating Neural Point-Based Graphics. 15948-15958
- Sheng Liu, Xiaohan Nie, Raffay Hamid:
Depth-Guided Sparse Structure-from-Motion for Movies and TV Shows. 15959-15968
- Denys Rozumnyi , Martin R. Oswald , Vittorio Ferrari, Marc Pollefeys :
Motion-from-Blur: 3D Shape and Motion Estimation of Motion-blurred Objects in Videos. 15969-15978
- Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, Ross B. Girshick:
Masked Autoencoders Are Scalable Vision Learners. 15979-15988
- Kieran A. Murphy, Varun Jampani, Srikanth Ramalingam, Ameesh Makadia:
Learning ABCs: Approximate Bijective Correspondence for isolating factors of variation with weak supervision. 15989-15999
- Yong Lin, Hanze Dong, Hao Wang, Tong Zhang:
Bayesian Invariant Risk Minimization. 16000-16009
- Xiangyu Peng, Kai Wang, Zheng Zhu, Mang Wang, Yang You:
Crafting Better Contrastive Views for Siamese Representation Learning. 16010-16019
- Haoqing Wang, Xun Guo, Zhi-Hong Deng, Yan Lu:
Rethinking Minimal Sufficient Representation in Contrastive Learning. 16020-16029
- Jie Xu, Huayi Tang, Yazhou Ren , Liang Peng, Xiaofeng Zhu, Lifang He:
Multi-level Feature Learning for Contrastive Multi-view Clustering. 16030-16039
- Yutong Bai, Xinlei Chen, Alexander Kirillov, Alan L. Yuille, Alexander C. Berg:
Point-Level Region Contrast for Object Detection Pre-Training. 16040-16049
- Minsoo Kang, Jaeyoo Park, Bohyung Han:
Class-Incremental Learning by Knowledge Distillation with Adaptive Feature Consolidation. 16050-16059
- Ramya Hebbalaguppe, Jatin Prakash, Neelabh Madan, Chetan Arora:
A Stitch in Time Saves Nine: A Train-Time Regularizing Loss for Improved Neural Network Calibration. 16060-16069
- Salar Hosseini Khorasgani, Yuxuan Chen, Florian Shkurti:
SLIC: Self-Supervised Learning with Iterative Clustering for Human Action Videos. 16070-16080
- Rohit Girdhar, Mannat Singh, Nikhila Ravi, Laurens van der Maaten, Armand Joulin, Ishan Misra:
Omnivore: A Single Model for Many Visual Modalities. 16081-16091
- Jae-Han Lee, Seungmin Jeon, Kwang Pyo Choi, Youngo Park, Chang-Su Kim:
DPICT: Deep Progressive Image Compression Using Trit-Planes. 16092-16101
- Eric R. Chan, Connor Z. Lin, Matthew A. Chan, Koki Nagano, Boxiao Pan, Shalini De Mello, Orazio Gallo, Leonidas J. Guibas, Jonathan Tremblay, Sameh Khamis, Tero Karras, Gordon Wetzstein:
Efficient Geometry-aware 3D Generative Adversarial Networks. 16102-16112



- Liang Chen, Yihang Lou, Jianzhong He, Tao Bai, Minghua Deng:
Geometric Anchor Correspondence Mining with Uncertainty Modeling for Universal Domain Adaptation. 16113-16122
- Richard J. Chen, Chengkuan Chen, Yicong Li, Tiffany Y. Chen, Andrew D. Trister, Rahul G. Krishnan, Faisal Mahmood:
Scaling Vision Transformers to Gigapixel Images via Hierarchical Self-Supervised Learning. 16123-16134
- Fangzhou Hong, Liang Pan , Zhongang Cai, Ziwei Liu:
Versatile Multi-Modal Pre-Training for Human-Centric Perception. 16135-16145
- Yuying Ge, Yixiao Ge, Xihui Liu, Dian Li, Ying Shan, Xiaohu Qie, Ping Luo:
Bridging Video-text Retrieval with Multiple Choice Questions. 16146-16155
- Karsten Roth, Oriol Vinyals, Zeynep Akata:
Integrating Language Guidance into Vision-based Deep Metric Learning. 16156-16168
- Ben Mildenhall, Peter Hedman, Ricardo Martin-Brualla, Pratul P. Srinivasan, Jonathan T. Barron:
NeRF in the Dark: High Dynamic Range View Synthesis from Noisy Raw Images. 16169-16178
- Liwen Wu, Jae Yong Lee, Anand Bhattad, Yu-Xiong Wang, David A. Forsyth:
DIVeR: Real-time and Accurate Neural Radiance Fields with Deterministic Integration for Volume Rendering. 16179-16188
- Chung-Yi Weng, Brian Curless, Pratul P. Srinivasan, Jonathan T. Barron, Ira Kemelmacher-Shlizerman:
HumanNeRF: Free-viewpoint Rendering of Moving People from Monocular Video. 16189-16199
- Junxuan Li, Hongdong Li:
Neural Reflectance for Shape Recovery with Shadow Handling. 16200-16209
- Berthy T. Feng , Alexander C. Ogren, Chiara Daraio, Katherine L. Bouman:
Visual Vibration Tomography: Estimating Interior Material Properties from Monocular Video. 16210-16219
- Kristina Monakhova, Stephan R. Richter, Laura Waller, Vladlen Koltun:
Dancing under the stars: video denoising in starlight. 16220-16230
- David B. Lindell, Dave Van Veen, Jeong Joon Park , Gordon Wetzstein:
Bacon: Band-limited Coordinate Networks for Multiscale Scene Representation. 16231-16241
- Jiankun Li, Peisen Wang, Pengfei Xiong, Tao Cai, Ziwei Yan, Lei Yang, Jiangyu Liu, Haoqiang Fan, Shuaicheng Liu :
Practical Stereo Matching via Cascaded Recurrent Network with Adaptive Correlation. 16242-16251
- Fangzhou Mu, Jian Wang, Yicheng Wu, Yin Li:
3D Photo Stylization: Learning to Generate Stylized Novel Views from a Single Image. 16252-16261
- Juewen Peng, Zhiguo Cao, Xianrui Luo, Hao Lu, Ke Xian, Jianming Zhang:
BokehMe: When Neural Rendering Meets Classical Rendering. 16262-16271
- Jay Whang, Mauricio Delbracio, Hossein Talebi, Chitwan Saharia, Alexandros G. Dimakis , Peyman Milanfar:
Deblurring via Stochastic Refinement. 16272-16282
- Lingyan Ruan , Bin Chen , Jizhou Li , Miu-Ling Lam:
Learning to Deblur using Light Field Generated and Real Defocus Images. 16283-16292
- Xu Ma, Yuqian Zhou, Xingqian Xu, Bin Sun, Valerii Filev, Nikita Orlov, Yun Fu, Humphrey Shi :
Towards Layer-wise Image Vectorization. 16293-16302
- Mark Sheinin, Dorian Chan, Matthew O'Toole, Srinivasa G. Narasimhan :
Dual-Shutter Optical Vibration Sensing. 16303-16312
- Jiaqu Li, Tao Yue, Sijie Zhao, Xuemei Hu:
Fisher Information Guidance for Learned Time-of-Flight Imaging. 16313-16322
- Shijie Lin, Yingiang Zhang , Lei Yu, Bin Zhou, Xiaowei Luo , Jia Pan:
Autofocus for Event Cameras. 16323-16332
- Ryan Po, Adithya Pediredla, Ioannis Gkioulekas :
Adaptive Gating for Single-Photon 3D Imaging. 16333-16342
- Martin Hahner, Christos Sakaridis , Mario Bijelic, Felix Heide, Fisher Yu, Dengxin Dai, Luc Van Gool:
LiDAR Snowfall Simulation for Robust 3D Object Detection. 16343-16353
- Rowan Zellers, Jiasen Lu, Ximing Lu, Youngjae Yu, Yanpeng Zhao, Mohammadreza Salehi, Aditya Kusupati, Jack Hessel, Ali Farhadi, Yejin Choi:
MERLOT RESERVE: Neural Script Knowledge through Vision and Language and Sound. 16354-16366
- Hao Jiang, Yadong Mu:
Joint Video Summarization and Moment Localization by Cross-Task Sample Transfer. 16367-16377
- Tanmay Gupta, Amita Kamath, Aniruddha Kembhavi, Derek Hoiem:
Towards General Purpose Vision Systems: An End-to-End Task-Agnostic Vision-Language Architecture. 16378-16388
- Joanna Materzynska, Antonio Torralba, David Bau :
Disentangling visual and written concepts in CLIP. 16389-16398
- Manling Li, Ruochen Xu, Shuohang Wang, Luowei Zhou, Xudong Lin, Chenguang Zhu, Michael Zeng, Heng Ji, Shih-Fu Chang:
CLIP-Event: Connecting Text and Images with Event Structures. 16399-16408



- Alex Andonian, Shixing Chen, Raffay Hamid:
Robust Cross-Modal Representation Learning with Progressive Self-Distillation. 16409-16420
- Antoine Yang, Antoine Miech, Josef Sivic, Ivan Laptev, Cordelia Schmid:
TubeDETR: Spatio-Temporal Video Grounding with Transformers. 16421-16432
- Junyu Luo , Jiahui Fu, Xianghao Kong , Chen Gao, Haibing Ren, Hao Shen, Huaxia Xia, Si Liu:
3D-SPS: Single-Stage 3D Visual Grounding via Referred Point Progressive Selection. 16433-16442
- Daigang Cai, Lichen Zhao, Jing Zhang , Lu Sheng , Dong Xu:
3DJCG: A Unified Framework for Joint Dense Captioning and Visual Grounding on 3D Point Clouds. 16443-16452
- Dídac Surís, Dave Epstein, Carl Vondrick:
Globetrotter: Connecting Languages by Connecting Images. 16453-16463
- Mingyang Zhou, Licheng Yu, Amanpreet Singh, Mengjiao Wang, Zhou Yu, Ning Zhang:
Unsupervised Vision-and-Language Pretraining via Retrieval-based Multi-Granular Alignment. 16464-16473
- Yingshan Chang, Guihong Cao, Mridu Narang, Jianfeng Gao, Hisami Suzuki, Yonatan Bisk:
WebQA: Multihop and Multimodal QA. 16474-16483
- Juil Koo, Ian Huang, Panos Achlioptas, Leonidas J. Guibas, Minhyuk Sung :
PartGlot: Learning Shape Part Segmentation from Language Reference Games. 16484-16493
- Ming Tao, Hao Tang, Fei Wu, Xiaoyuan Jing, Bing-Kun Bao, Changsheng Xu:
DF-GAN: A Simple and Effective Baseline for Text-to-Image Synthesis. 16494-16504
- Taehoon Kim, Gwangmo Song, Sihaeng Lee, Sangyun Kim, Yewon Seo, Soonyoung Lee, Seung Hwan Kim, Honglak Lee, Kyunghoon Bae:
L-Verse: Bidirectional Generation Between Image and Text. 16505-16515
- Shizhe Chen, Pierre-Louis Guhur, Makarand Tapaswi, Cordelia Schmid, Ivan Laptev:
Think Global, Act Local: Dual-scale Graph Transformer for Vision-and-Language Navigation. 16516-16526
- Ali Furkan Biten, Ron Litman, Yusheng Xie, Srikanth Appalaraju, R. Manmatha:
LaTr: Layout-Aware Transformer for Scene-Text VQA. 16527-16537
- Dim P. Papadopoulos , Enrique Mora, Nadiia Chepurko, Kuan Wei Huang, Ferda Ofli, Antonio Torralba:
Learning Program Representations for Food Images and Cooking Recipes. 16538-16548
- Xiao Wang, Haoqi Fan, Yuandong Tian, Daisuke Kihara, Xinlei Chen:
On the Importance of Asymmetry for Siamese Representation Learning. 16549-16558
- Tong Zhang,