[+] [–] ICCV 2017: Venice, Italy

IEEE International Conference on Computer Vision, ICCV 2017, Venice, Italy, October 22-29, 2017. IEEE Computer Society 2017, ISBN 978-1-5386-1032-9

Oral Session 1

Spotlight Session 1

Tz-Ying Wu, Ting-An Chien, Cheng-Sheng Chan, Chan-Wei Hu, Min Sun:
Anticipating Daily Intention Using On-wrist Motion Triggered Sensing. 48-56

Rui Zhu, Hamed Kiani Galoogahi, Chaoyang Wang, Simon Lucey:
Rethinking Reprojection: Closing the Loop for Pose-Aware Shape Reconstruction from a Single Image. 57-65

Alex Kendall, Hayk Martirosyan, Saumitro Dasgupta, Peter Henry:
End-to-End Learning of Geometry and Context for Deep Stereo Regression. 66-75

Janne Heikkilä:
Using Sparse Elimination for Solving Minimal Problems in Computer Vision. 76-84

Xiaoguang Han, Zhen Li, Haibin Huang, Evangelos Kalogerakis, Yizhou Yu:
High-Resolution Shape Completion Using Deep Neural Networks for Global Structure and Local Geometry Inference. 85-93

Dotan Kaufman, Gil Levi, Tal Hassner, Lior Wolf:
Temporal Tessellation: A Unified Approach for Video Analysis. 94-104

Chen Huang, Simon Lucey, Deva Ramanan:
Learning Policies for Adaptive Tracking with Deep Feature Cascades. 105-114

Yuki Shiba, Satoshi Ono, Ryo Furukawa, Shinsaku Hiura, Hiroshi Kawasaki:
Temporal Shape Super-Resolution by Intra-frame Motion Encoding Using High-fps Structured Light. 115-123

Poster 1

Henning Tjaden, Ulrich Schwanecke, Elmar Schömer:
Real-Time Monocular Pose Estimation of 3D Objects Using Temporally Consistent Local Color Histograms. 124-132

Tolga Birdal, Slobodan Ilic:
CAD Priors for Accurate and Flexible Instance Reconstruction. 133-142

Jaesik Park, Qian-Yi Zhou, Vladlen Koltun:
Colored Point Cloud Registration Revisited. 143-152

Marc Khoury, Qian-Yi Zhou, Vladlen Koltun:
Learning Compact Geometric Features. 153-161

Jeong-Kyun Lee, Jae-Won Yea, Min-Gyu Park, Kuk-Jin Yoon:
Joint Layout Estimation and Global Multi-view Registration for Indoor Reconstruction. 162-171

Rudrasis Chakraborty, Vikas Singh, Nagesh Adluru, Baba C. Vemuri:
A Geometric Framework for Statistical Analysis of Trajectories with Distinct Temporal Spans. 172-181

Liang Mi, Wen Zhang, Junwei Zhang, Yonghui Fan, Dhruman Goradia, Kewei Chen, Eric M. Reiman, Xianfeng Gu, Yalin Wang:

Framework. 341-349

Xihui Liu, Haiyu Zhao, Maoqing Tian, Lu Sheng, Jing Shao, Shuai Yi, Junjie Yan, Xiaogang Wang:
HydraPlus-Net: Attentive Deep Features for Pedestrian Analysis. 350-359

Yair Movshovitz-Attias, Alexander Toshev, Thomas K. Leung, Sergey Ioffe, Saurabh Singh:
No Fuss Distance Metric Learning Using Proxies. 360-368

Matteo Ruggero Ronchi, Pietro Perona:
Benchmarking and Error Diagnosis in Multi-instance Pose Estimation. 369-378

Zhongdao Wang, Luming Tang, Xihui Liu, Zhuliang Yao, Shuai Yi, Jing Shao, Junjie Yan, Shengjin Wang, Hongsheng Li, Xiaogang Wang:
Orientation Invariant Feature Embedding and Spatial Temporal Regularization for Vehicle Re-identification. 379-387

Ziad Al-Halah, Rainer Stiefelhagen, Kristen Grauman:
Fashion Forward: Forecasting Visual Style in Fashion. 388-397

Xingyi Zhou, Qixing Huang, Xiao Sun, Xiangyang Xue, Yichen Wei:
Towards 3D Human Pose Estimation in the Wild: A Weakly-Supervised Approach. 398-407

Xizhou Zhu, Yujie Wang, Jifeng Dai, Lu Yuan, Yichen Wei:
Flow-Guided Feature Aggregation for Video Object Detection. 408-417

Jong-Chyi Su, Chenyun Wu, Huaizu Jiang, Subhransu Maji:
Reasoning About Fine-Grained Attribute Phrases Using Reference Games. 418-427

Lachlan Tychsen-Smith, Lars Petersson:
DeNet: Scalable Real-Time Object Detection with Directed Sparse Sampling. 428-436

Fatih Çakir, Kun He, Sarah Adel Bargal, Stan Sclaroff:
MIHash: Online Hashing with Mutual Information. 437-445

Jiajun Lu, Theerasit Issaranon, David A. Forsyth:
SafetyNet: Detecting and Rejecting Adversarial Examples Robustly. 446-454

Arun Mallya, Svetlana Lazebnik:
Recurrent Models for Situation Recognition. 455-463

Zhouxia Wang, Tianshui Chen, Guanbin Li, Ruijia Xu, Liang Lin:
Multi-label Image Recognition by Recurrently Discovering Attentional Regions. 464-472

Pengtao Xie, Ruslan Salakhutdinov, Luntian Mou, Eric P. Xing:
Deep Determinantal Point Process for Large-Scale Multi-label Classification. 473-482

Yuke Zhu, Daniel Gordon, Eric Kolve, Dieter Fox, Li Fei-Fei, Abhinav Gupta, Roozbeh Mottaghi, Ali Farhadi:
Visual Semantic Planning Using Deep Successor Representations. 483-492

Florian Walch, Caner Hazirbas, Laura Leal-Taixé, Torsten Sattler, Sebastian Hilsenbeck, Daniel Cremers:
Image-Based Localization Using LSTMs for Structured Feature Correlation. 627-637

Jian Ren, Xiaohui Shen, Zhe L. Lin, Radomír Mech, David J. Foran:
Personalized Image Aesthetics. 638-647

Pauline Luc, Natalia Neverova, Camille Couprie, Jakob Verbeek, Yann LeCun:
Predicting Deeper into the Future of Semantic Segmentation. 648-657

Wei Wen, Cong Xu, Chunpeng Wu, Yandan Wang, Yiran Chen, Hai Li:
Coordinating Filters for Faster Deep Neural Networks. 658-666

Hsin-Ying Lee, Jia-Bin Huang, Maneesh Singh, Ming-Hsuan Yang:
Unsupervised Representation Learning by Sorting Sequences. 667-676

Seil Na, Sangho Lee, Jisung Kim, Gunhee Kim:
A Read-Write Memory Network for Movie Story Understanding. 677-685

Jingchun Cheng, Yi-Hsuan Tsai, Shengjin Wang, Ming-Hsuan Yang:
SegFlow: Joint Learning for Video Object Segmentation and Optical Flow. 686-695

Khurram Soomro, Mubarak Shah:
Unsupervised Action Discovery and Localization in Videos. 696-705

Ranjay Krishna, Kenji Hata, Frederic Ren, Li Fei-Fei, Juan Carlos Niebles:
Dense-Captioning Events in Videos. 706-715

Yemin Shi, Yonghong Tian, Yaowei Wang, Wei Zeng, Tiejun Huang:
Learning Long-Term Dependencies for Action Recognition with a Biologically-Inspired Deep Network. 716-725

Tan Yu, Zhenzhen Wang, Junsong Yuan:
Compressive Quantization for Fast Object Instance Search in Videos. 726-735

Hehe Fan, Xiaojun Chang, De Cheng, Yi Yang, Dong Xu, Alexander G. Hauptmann:
Complex Event Detection by Identifying Reliable Shots from Untrimmed Videos. 736-744

Wenhao He, Xu-Yao Zhang, Fei Yin, Cheng-Lin Liu:
Deep Direct Regression for Multi-oriented Scene Text Detection. 745-753

Oral Session 2

Pau Panareda Busto, Juergen Gall:
Open Set Domain Adaptation. 754-763

Jifeng Dai, Haozhi Qi, Yuwen Xiong, Yi Li, Guodong Zhang, Han Hu, Yichen Wei:
Deformable Convolutional Networks. 764-773

Song Bai, Zhichao Zhou, Jingdong Wang, Xiang Bai, Longin Jan Latecki, Qi Tian:

Ensemble Diffusion for Retrieval. 774-783

Xin Li, Zequn Jie, Wei Wang, Changsong Liu, Jimei Yang, Xiaohui Shen, Zhe Lin, Qiang Chen, Shuicheng Yan, Jiashi Feng:
FoveaNet: Perspective-Aware Urban Scene Parsing. 784-792

Christopher Funk, Yanxi Liu:
Beyond Planar Symmetry: Modeling Human Perception of Reflection and Rotation Symmetries in the Wild. 793-803

Spotlight Session 2

Ronghang Hu, Jacob Andreas, Marcus Rohrbach, Trevor Darrell, Kate Saenko:
Learning to Reason: End-to-End Module Networks for Visual Question Answering. 804-813

Yuhui Yuan, Kuiyuan Yang, Chao Zhang:
Hard-Aware Deeply Cascaded Embedding. 814-823

Kan Chen, Rama Kovvuri, Ram Nevatia:
Query-Guided Regression Network with Context Policy for Phrase Grounding. 824-832

Himalaya Jain, Joaquin Zepeda, Patrick Pérez, Rémi Gribonval:
SuBiC: A Supervised, Structured Binary Code for Image Search. 833-842

Chen Sun, Abhinav Shrivastava, Saurabh Singh, Abhinav Gupta:
Revisiting Unreasonable Effectiveness of Data in Deep Learning Era. 843-852

Christoph Lassner, Gerard Pons-Moll, Peter V. Gehler:
A Generative Model of People in Clothing. 853-862

Roman Klokov, Victor S. Lempitsky:
Escape from Cells: Deep Kd-Networks for the Recognition of 3D Point Cloud Models. 863-872

Siqi Liu, Zhenhai Zhu, Ning Ye, Sergio Guadarrama, Kevin Murphy:
Improved Image Captioning via Policy Gradient optimization of SPIDEr. 873-881

Poster Session 2

Pulak Purkait, Christopher Zach, Ales Leonardis:
Rolling Shutter Correction in Manhattan World. 882-890

David Avidar, David Malah, Meir Barzohar:
Local-to-Global Point Cloud Registration Using a Dictionary of Viewpoint Descriptors. 891-899

Chuhang Zou, Ersin Yumer, Jimei Yang, Duygu Ceylan, Derek Hoiem:
3D-PRNN: Generating Shape Primitives with Recurrent Neural Networks. 900-909

Tao Yu, Kaiwen Guo, Feng Xu, Yuan Dong, Zhaoqi Su, Jianhui Zhao, Jianguo Li, Qionghai Dai, Yebin Liu:
BodyFusion: Real-Time Capture of Human Motion and Surface Geometry Using a Single Depth Camera. 910-919

Qianggong Zhang, Tat-Jun Chin, David Suter:
Quasiconvex Plane Sweep for Triangulation with Outliers. 920-928

Pan Ji, Hongdong Li, Yuchao Dai, Ian D. Reid:
"Maximizing Rigidity" Revisited: A Convex Programming Approach for Generic 3D Shape Reconstruction from Multiple Perspective Views. 929-937

Xiaopeng Zheng, Chengfeng Wen, Na Lei, Ming Ma, Xianfeng Gu:
Surface Registration via Foliation. 938-947

Bingbing Zhuang, Loong-Fah Cheong, Gim Hee Lee:
Rolling-Shutter-Aware Differential SfM and Image Rectification. 948-956

Sotiris Nousias, François Chadebecq, Jonas Pichat, Pearse A. Keane, Sébastien Ourselin, Christos Bergeles:
Corner-Based Geometric Calibration of Multi-focus Plenoptic Cameras. 957-965

Qi Guo, Emma Alexander, Todd E. Zickler:
Focal Track: Depth and Accommodation with Oscillating Lens Deformation. 966-974

Mark Buckler, Suren Jayasuriya, Adrian Sampson:
Reconfiguring the Imaging Pipeline for Computer Vision. 975-984

Yujia Xue, Kang Zhu, Qiang Fu, Xilin Chen, Jingyi Yu:
Catadioptric HyperSpectral Light Field Imaging. 985-993

Hong-Xing Yu, Ancong Wu, Wei-Shi Zheng:
Cross-View Asymmetric Metric Learning for Unsupervised Person Re-Identification. 994-1002

Kang Wang, Qiang Ji:
Real Time Eye Gaze Tracking with 3D Deformable Eye-Face Model. 1003-1011

Inwoong Lee, Doyoung Kim, Seoungyoon Kang, Sanghoon Lee:
Ensemble Deep Learning for Skeleton-Based Action Recognition Using Temporal Sliding LSTM Networks. 1012-1020

Adrian Bulat, Georgios Tzimiropoulos:
How Far are We from Solving the 2D & 3D Face Alignment Problem? (and a Dataset of 230, 000 3D Facial Landmarks). 1021-1030

Aaron S. Jackson, Adrian Bulat, Vasileios Argyriou, Georgios Tzimiropoulos:
Large Pose 3D Face Reconstruction from a Single Image via Direct Volumetric CNN Regression. 1031-1039

Xialei Liu, Joost van de Weijer, Andrew D. Bagdanov:
RankIQA: Learning from Rankings for No-Reference Image Quality Assessment. 1040-1049

Xiaowu Chen, Anlin Zheng, Jia Li, Feng Lu:
Look, Perceive and Segment: Finding the Salient Objects in Images via Two-stream Fixation-Semantic CNNs. 1050-1058

Shengfeng He, Jianbo Jiao, Xiaodan Zhang, Guoqiang Han, Rynson W. H. Lau:
Delving into Salient Object Subitizing and Detection. 1059-1067

Ruichi Yu, Ang Li, Vlad I. Morariu, Larry S. Davis:
Visual Relationship Detection with Internal and External Linguistic Knowledge Distillation. 1068-1076

Jinshan Pan, Jiangxin Dong, Yu-Wing Tai, Zhixun Su, Ming-Hsuan Yang:
Learning Discriminative Data Fitting Functions for Blind Image Deblurring. 1077-1085

Wenqi Ren, Jinshan Pan, Xiaochun Cao, Ming-Hsuan Yang:
Video Deblurring via Semantic Segmentation and Pixel-Wise Non-linear Kernel. 1086-1094

Ruohan Gao, Kristen Grauman:
On-demand Learning for Deep Image Restoration. 1095-1104

Jun Xu, Lei Zhang, David Zhang, Xiangchu Feng:
Multi-channel Weighted Nuclear Norm Minimization for Real Color Image Denoising. 1105-1113

Dongdong Chen, Jing Liao, Lu Yuan, Nenghai Yu, Gang Hua:
Coherent Online Video Style Transfer. 1114-1123

Arko Barman, Shishir K. Shah:
SHaPE: A Novel Graph Theoretic Algorithm for Making Consensus-Based Decisions in Person Re-identification Systems. 1124-1133

Hamed Kiani Galoogahi, Ashton Fagg, Chen Huang, Deva Ramanan, Simon Lucey:
Need for Speed: A Benchmark for Higher Frame Rate Object Tracking. 1134-1143

Hamed Kiani Galoogahi, Ashton Fagg, Simon Lucey:
Learning Background-Aware Correlation Filters for Visual Tracking. 1144-1152

Zhu Teng, Junliang Xing, Qiang Wang, Congyan Lang, Songhe Feng, Yi Jin:
Robust Object Tracking Based on Temporal and Spatial Deep Networks. 1153-1162

Franziska Mueller, Dushyant Mehta, Oleksandr Sotnychenko, Srinath Sridhar, Dan Casas, Christian Theobalt:
Real-Time Hand Tracking under Occlusion from an Egocentric RGB-D Sensor. 1163-1172

Siyuan Qi, Siyuan Huang, Ping Wei, Song-Chun Zhu:
Predicting Human Activities Using Stochastic Grammar. 1173-1181

Anne S. Wannenwetsch, Margret Keuper, Stefan Roth:
ProbFlow: Joint Optical Flow and Uncertainty Estimation. 1182-1191

Thomas Möllenhoff, Daniel Cremers:
Sublabel-Accurate Discretization of Nonconvex Free-Discontinuity Problems. 1192-1200

Yinda Zhang, Mingru Bai, Pushmeet Kohli, Shahram Izadi, Jianxiong Xiao:
DeepContext: Context-Encoding Neural Pathways for 3D Holistic Scene Understanding. 1201-1210

Michael J. Wilber, Chen Fang, Hailin Jin, Aaron Hertzmann, John P. Collomosse, Serge J. Belongie:
BAM! The Behance Artistic Media Dataset for Recognition Beyond Photography. 1211-1220

Yu Chen, Chunhua Shen, Xiu-Shen Wei, Lingqiao Liu, Jian Yang:
Adversarial PoseNet: A Structure-Aware Convolutional Network for Human Pose Estimation. 1221-1230

Jiuxiang Gu, Gang Wang, Jianfei Cai, Tsuhan Chen:
An Empirical Study of Language CNN for Image Captioning. 1231-1240

Berkan Demirel, Ramazan Gokberk Cinbis, Nazli Ikizler-Cinbis:
Attributes2Classname: A Discriminative Model for Attribute-Based Unsupervised Zero-Shot Learning. 1241-1250

Marco Pedersoli, Thomas Lucas, Cordelia Schmid, Jakob Verbeek:
Areas of Attention for Image Captioning. 1251-1259

Zhoutong Zhang, Jiajun Wu, Qiujia Li, Zhengjia Huang, James Traer, Josh H. McDermott, Joshua B. Tenenbaum, William T. Freeman:
Generative Modeling of Audible Shapes for Object Perception. 1260-1269

Yikang Li, Wanli Ouyang, Bolei Zhou, Kun Wang, Xiaogang Wang:
Scene Graph Generation from Objects, Phrases and Region Captions. 1270-1279

Chenxi Liu, Zhe Lin, Xiaohui Shen, Jimei Yang, Xin Lu, Alan L. Yuille:
Recurrent Multimodal Interaction for Referring Image Segmentation. 1280-1289

Wei Yang, Shuang Li, Wanli Ouyang, Hongsheng Li, Xiaogang Wang:
Learning Feature Pyramids for Human Pose Estimation. 1290-1299

Chen Zhu, Yanpeng Zhao, Shuaiyi Huang, Kewei Tu, Yi Ma:
Structured Attentions for Visual Question Answering. 1300-1309

Debidatta Dwibedi, Ishan Misra, Martial Hebert:
Cut, Paste and Learn: Surprisingly Easy Synthesis for Instance Detection. 1310-1319

Di Lin, Guangyong Chen, Daniel Cohen-Or, Pheng-Ann Heng, Hui Huang:
Cascaded Feature Network for Semantic Segmentation of RGB-D Images. 1320-1328

Amal Rannen Triki, Rahaf Aljundi, Matthew B. Blaschko, Tinne Tuytelaars:
Encoder Based Lifelong Learning. 1329-1337

Xiaolong Wang, Kaiming He, Abhinav Gupta:
Transitive Invariance for Self-Supervised Visual Representation Learning. 1338-1347

Stepan Tulyakov, Anton Ivanov, François Fleuret:
Weakly Supervised Learning of Deep Metrics for Stereo Reconstruction. 1348-1357

Timnit Gebru, Judy Hoffman, Li Fei-Fei:
Fine-Grained Recognition in the Wild: A Multi-task Domain Adaptation Approach. 1358-1367

Yan Wang, Lingxi Xie, Chenxi Liu, Siyuan Qiao, Ya Zhang, Wenjun Zhang, Qi Tian, Alan L. Yuille:
SORT: Second-Order Response Transform for Visual Recognition. 1368-1377

Cihang Xie, Jianyu Wang, Zhishuai Zhang, Yuyin Zhou, Lingxi Xie, Alan L. Yuille:
Adversarial Examples for Semantic Segmentation and Object Detection. 1378-1387

Lingxi Xie, Alan L. Yuille:
Genetic CNN. 1388-1397

Yihui He, Xiangyu Zhang, Jian Sun:
Channel Pruning for Accelerating Very Deep Neural Networks. 1398-1406

Giorgio Roffo, Simone Melzi, Umberto Castellani, Alessandro Vinciarelli:
Infinite Latent Feature Selection: A Probabilistic Latent Graph-Based Ranking Approach. 1407-1415

Amir Mazaheri, Dong Zhang, Mubarak Shah:
Video Fill In the Blank Using LR/RL LSTMs with Spatial-Temporal Attentions. 1416-1425

Jia Li, Anlin Zheng, Xiaowu Chen, Bin Zhou:
Primary Video Object Segmentation via Complementary CNNs and Neighborhood Reversible Flow. 1426-1434

Tanya Marwah, Gaurav Mittal, Vineeth N. Balasubramanian:
Attentive Semantic Video Generation Using Captions. 1435-1443

Adrià Recasens, Carl Vondrick, Aditya Khosla, Antonio Torralba:
Following Gaze in Video. 1444-1452

Wenbo Li, Longyin Wen, Ming-Ching Chang, Ser Nam Lim, Siwei Lyu:
Adaptive RNN Tree for Large-Scale Human Action Recognition. 1453-1461

Masataka Yamaguchi, Kuniaki Saito, Yoshitaka Ushiku, Tatsuya Harada:
Spatio-Temporal Person Retrieval via Natural Language Queries. 1462-1471

Xintong Han, Zuxuan Wu, Phoenix X. Huang, Xiao Zhang, Menglong Zhu, Yuan Li, Yang Zhao, Larry S. Davis:
Automatic Spatially-Aware Fashion Concept Discovery. 1472-1480

Joseph DeGol, Timothy Bretl, Derek Hoiem:
ChromaTag: A Colored Marker and Fast Detection Algorithm. 1481-1490

Seong Joon Oh, Mario Fritz, Bernt Schiele:
Adversarial Image Perturbation for Privacy Protection A Game Theory Perspective. 1491-1500

Shangxuan Tian, Shijian Lu, Chongshou Li:
WeText: Scene Text Detection under Weak Supervision. 1501-1509

Vision for X Oral Session 3

Xun Huang, Serge J. Belongie:
Arbitrary Style Transfer in Real-Time with Adaptive Instance Normalization. 1510-1519

Qifeng Chen, Vladlen Koltun:
Photographic Image Synthesis with Cascaded Refinement Networks. 1520-1529

Wadim Kehl, Fabian Manhardt, Federico Tombari, Slobodan Ilic, Nassir Navab:
SSD-6D: Making RGB-Based 3D Detection and 6D Pose Estimation Great Again. 1530-1538

Lior Wolf, Yaniv Taigman, Adam Polyak:
Unsupervised Creation of Parameterized Avatars. 1539-1547

Karel Zimmermann, Tomás Petrícek, Vojtech Salanský, Tomás Svoboda:
Learning for Active 3D Mapping. 1548-1556

Poster Session 3

Jialiang Wang, Daniel Glasner, Todd E. Zickler:
Toward Perceptually-Consistent Stereo: A Scanline Study. 1557-1565

Weifeng Chen, Donglai Xiang, Jia Deng:
Surface Normals in the Wild. 1566-1575

Chao Zhou, Hong Zhang, Xiaoyong Shen, Jiaya Jia:
Unsupervised Learning of Stereo Matching. 1576-1584

Matan Sela, Elad Richardson, Ron Kimmel:
Unrestricted Facial Geometry Reconstruction Using Image-to-Image Translation. 1585-1594

Wilfried Hartmann, Silvano Galliani, Michal Havlena, Luc Van Gool, Konrad Schindler:
Learned Multi-patch Similarity. 1595-1603

Ryan Szeto, Jason J. Corso:
Click Here: Human-Localized Keypoints as Guidance for Viewpoint Estimation. 1604-1613

Alessio Tonioni, Matteo Poggi, Stefano Mattoccia, Luigi Di Stefano:
Unsupervised Adaptation for Deep Stereo. 1614-1622

Parikshit Sakurikar, P. J. Narayanan:
Composite Focus Measure for High Quality Depth Maps. 1623-1631

Xi Peng, Xiang Yu, Kihyuk Sohn, Dimitris N. Metaxas, Manmohan Chandraker:
Reconstruction-Based Disentanglement for Pose-Invariant Face Recognition. 1632-1641

Shengtao Xiao, Jiashi Feng, Luoqi Liu, Xuecheng Nie, Wei Wang, Shuicheng Yan, Ashraf A. Kassim:
Recurrent 3D-2D Dual Learning for Large-Pose Facial Landmark Detection. 1642-1651

Eirikur Agustsson, Radu Timofte, Luc Van Gool:
Anchored Regression Networks Applied to Age Estimation and Super Resolution. 1652-1661

Eryun Liu:
Infant Footprint Recognition. 1662-1669

Dong Gong, Mingkui Tan, Yanning Zhang, Anton van den Hengel, Qinfeng Shi:
Self-Paced Kernel Estimation for Robust Blind Image Deblurring. 1670-1679

Wenguan Wang, Jianbing Shen, Jianwen Xie, Fatih Porikli:
Super-Trajectory for Video Segmentation. 1680-1688

Shizhan Zhu, Sanja Fidler, Raquel Urtasun, Dahua Lin, Chen Change Loy:
Be Your Own Prada: Fashion Synthesis with Structural Coherence. 1689-1697

Huaibo Huang, Ran He, Zhenan Sun, Tieniu Tan:
Wavelet-SRNet: A Wavelet-Based CNN for Multi-scale Face Super Resolution. 1698-1706

George Leifman, Dmitry Rudoy, Tristan Swedish, Eduardo Bayro-Corrochano, Ramesh Raskar:
Learning Gaze Transitions from Depth to Improve Video Saliency Estimation. 1707-1716

Shuhang Gu, Deyu Meng, Wangmeng Zuo, Lei Zhang:
Joint Convolutional Analysis and Synthesis Sparse Representation for Single Image Layer Separation. 1717-1725

Seonghyeon Nam, Seon Joo Kim:
Modelling the Scene Dependent Imaging in Cameras with a Deep Neural Network. 1726-1734

Yi Chang, Luxin Yan, Sheng Zhong:
Transformed Low-Rank Model for Line Pattern Noise Removal. 1735-1743

Utkarsh Gaur, B. S. Manjunath:
Weakly Supervised Manifold Learning for Dense Semantic Object Correspondence. 1744-1752

Junfeng Yang, Xueyang Fu, Yuwen Hu, Yue Huang, Xinghao Ding, John W. Paisley:
PanNet: A Deep Network Architecture for Pan-Sharpening. 1753-1761

Shuang Li, Tong Xiao, Hongsheng Li, Wei Yang, Xiaogang Wang:
Identity-Aware Textual-Visual Matching with Latent Co-attention. 1908-1917

Yantao Shen, Tong Xiao, Hongsheng Li, Shuai Yi, Xiaogang Wang:
Learning Deep Neural Networks for Vehicle Re-ID with Visual-spatio-Temporal Path Proposals. 1918-1927

Yuncheng Li, Jianchao Yang, Yale Song, Liangliang Cao, Jiebo Luo, Li-Jia Li:
Learning from Noisy Labels with Distillation. 1928-1936

Zhiqiang Shen, Zhuang Liu, Jianguo Li, Yu-Gang Jiang, Yurong Chen, Xiangyang Xue:
DSOD: Learning Deeply Supervised Object Detectors from Scratch. 1937-1945

Bryan A. Plummer, Arun Mallya, Christopher M. Cervantes, Julia Hockenmaier, Svetlana Lazebnik:
Phrase Localization and Visual Relationship Detection with Comprehensive Image-Language Cues. 1946-1955

Wanli Ouyang, Kun Wang, Xin Zhu, Xiaogang Wang:
Chained Cascade Network for Object Detection. 1956-1964

Seokju Lee, Junsik Kim, Jae Shin Yoon, Seunghak Shin, Oleksandr Bailo, Namil Kim, Tae-Hee Lee, Hyun Seok Hong, Seung-Hoon Han, In So Kweon:
VPGNet: Vanishing Point Guided Network for Lane and Road Marking Detection and Recognition. 1965-1973

Gedas Bertasius, Hyun Soo Park, Stella X. Yu, Jianbo Shi:
Unsupervised Learning of Important Objects from First-Person Videos. 1974-1982

Kushal Kafle, Christopher Kanan:
An Analysis of Visual Question Answering Algorithms. 1983-1991

Dahjung Chung, Khalid Tahboub, Edward J. Delp:
A Two Stream Siamese Convolutional Neural Network for Person Re-identification. 1992-2000

Vicky Kalogeiton, Philippe Weinzaepfel, Vittorio Ferrari, Cordelia Schmid:
Joint Learning of Object and Action Detectors. 2001-2010

Yi-Hsin Chen, Wei-Yu Chen, Yu-Ting Chen, Bo-Cheng Tsai, Yu-Chiang Frank Wang, Min Sun:
No More Discrimination: Cross City Adaptation of Road Scene Segmenters. 2011-2020

Hang Zhao, Xavier Puig, Bolei Zhou, Sanja Fidler, Antonio Torralba:
Open Vocabulary Scene Parsing. 2021-2029

Steffen Wolf, Lukas Schott, Ullrich Köthe, Fred A. Hamprecht:
Learned Watershed: End-to-End Learning of Seeded Segmentation. 2030-2038

Yang Zhang, Philip David, Boqing Gong:
Curriculum Domain Adaptation for Semantic Segmentation of Urban

Scenes. 2039-2049

Rui Zhang, Sheng Tang, Yongdong Zhang, Jintao Li, Shuicheng Yan:
Scale-Adaptive Convolutions for Scene Parsing. 2050-2058

Ryo Yonetani, Vishnu Naresh Boddeti, Kris M. Kitani, Yoichi Sato:
Privacy-Preserving Visual Learning Using Doubly Permuted Homomorphic Encryption. 2059-2069

Carl Doersch, Andrew Zisserman:
Multi-task Self-Supervised Visual Learning. 2070-2079

Xiaojun Chen, Joshua Zhexue Huang, Feiping Nie, Renjie Chen, Qingyao Wu:
A Self-Balanced Min-Cut Algorithm for Image Clustering. 2080-2088

Peihua Li, Jiangtao Xie, Qilong Wang, Wangmeng Zuo:
Is Second-Order Information Helpful for Large-Scale Visual Recognition? 2089-2097

Yanghao Li, Naiyan Wang, Jiaying Liu, Xiaodi Hou:
Factorized Bilinear Models for Image Recognition. 2098-2106

Maxim Tatarchenko, Alexey Dosovitskiy, Thomas Brox:
Octree Generating Networks: Efficient Convolutional Architectures for High-resolution 3D Outputs. 2107-2115

Yan Zhang, Mete Ozay, Shuohao Li, Takayuki Okatani:
Truncating Wide Networks Using Binary Tree Architectures. 2116-2124

Fatemehsadat Saleh, Mohammad Sadegh Ali Akbarian, Mathieu Salzmann, Lars Petersson, Jose M. Alvarez:
Bringing Background into the Foreground: Making All Classes Equal in Weakly-Supervised Video Semantic Segmentation. 2125-2135

Pengfei Zhang, Cuiling Lan, Junliang Xing, Wenjun Zeng, Jianru Xue, Nanning Zheng:
View Adaptive Recurrent Neural Networks for High Performance Human Action Recognition from Skeleton Data. 2136-2145

Jean-Baptiste Alayrac, Josef Sivic, Ivan Laptev, Simon Lacoste-Julien:
Joint Discovery of Object States and Manipulation Actions. 2146-2155

Gunnar A. Sigurdsson, Olga Russakovsky, Abhinav Gupta:
What Actions are Needed for Understanding Human Actions in Videos? 2156-2165

Lin Sun, Kui Jia, Kevin Chen, Dit-Yan Yeung, Bertram E. Shi, Silvio Savarese:
Lattice Long Short-Term Memory for Human Action Recognition. 2166-2175

Jiong Yang, Junsong Yuan:
Common Action Discovery and Localization in Unconstrained Videos. 2176-2185

Jae Shin Yoon, François Rameau, Junsik Kim, Seokju Lee, Seunghak Shin, In So Kweon:
Pixel-Level Matching for Video Object Segmentation Using Convolutional Neural Networks. 2186-2195

Gedas Bertasius, Hyun Soo Park, Stella X. Yu, Jianbo Shi:
Am I a Baller? Basketball Performance Assessment from First-Person Videos. 2196-2204

Wenguan Wang, Jianbing Shen:
Deep Cropping via Attention Box Prediction and Aesthetics Assessment. 2205-2213

Chen Liu, Jiajun Wu, Pushmeet Kohli, Yasutaka Furukawa:
Raster-to-Vector: Revisiting Floorplan Transformation. 2214-2222

Michal Busta, Lukás Neumann, Jiri Matas:
Deep TextSpotter: An End-to-End Trainable Scene Text Localization and Recognition Framework. 2223-2231

Vision for X & Computational Photography Spotlight Session 3

Stephan R. Richter, Zeeshan Hayder, Vladlen Koltun:
Playing for Benchmarks. 2232-2241

Jun-Yan Zhu, Taesung Park, Phillip Isola, Alexei A. Efros:
Unpaired Image-to-Image Translation Using Cycle-Consistent Adversarial Networks. 2242-2251

Anton Osokin, Anatole Chessel, Rafael Edgardo Carazo-Salas, Federico Vaggi:
GANs for Biological Image Synthesis. 2252-2261

Pratul P. Srinivasan, Tongzhou Wang, Ashwin Sreelal, Ravi Ramamoorthi, Ren Ng:
Learning to Synthesize a 4D RGBD Light Field from a Single Image. 2262-2270

Stefan Heber, Wei Yu, Thomas Pock:
Neural EPI-Volume Networks for Shape from Light Field. 2271-2279

Guilin Liu, Duygu Ceylan, Ersin Yumer, Jimei Yang, Jyh-Ming Lien:
Material Editing Using a Physically Based Rendering Network. 2280-2288

Katherine L. Bouman, Vickie Ye, Adam B. Yedidia, Frédo Durand, Gregory W. Wornell, Antonio Torralba, William T. Freeman:
Turning Corners into Cameras: Principles and Methods. 2289-2297

Silvia Tozza, William A. P. Smith, Dizhong Zhu, Ravi Ramamoorthi, Edwin R. Hancock:
Linear Differential Constraints for Photo-Polarimetric Height Estimation. 2298-2306

Poster Session 4

Viktor Larsson, Kalle Åström, Magnus Oskarsson:
Polynomial Solvers for Saturated Ideals. 2307-2316

Weiyue Wang, Qiangui Huang, Suya You, Chao Yang, Ulrich Neumann:
Shape Inpainting Using 3D Generative Adversarial Network and Recurrent

Convolutional Networks. 2317-2325

Mengqi Ji, Juergen Gall, Haitian Zheng, Yebin Liu, Lu Fang:
SurfaceNet: An End-to-End 3D Neural Network for Multiview Stereopsis. 2326-2334

Viktor Larsson, Zuzana Kukelova, Yinqiang Zheng:
Making Minimal Solvers for Absolute Pose Estimation Compact and Robust. 2335-2343

Wuyuan Xie, Miaohui Wang, Xianbiao Qi, Lei Zhang:
3D Surface Detail Enhancement from a Single Normal Map. 2344-2352

Haoshu Fang, Shuqin Xie, Yu-Wing Tai, Cewu Lu:
RMPE: Regional Multi-person Pose Estimation. 2353-2362

Yongyi Lu, Cewu Lu, Chi-Keung Tang:
Online Video Object Detection Using Association LSTM. 2363-2371

Liangliang Nan, Peter Wonka:
PolyFit: Polygonal Surface Reconstruction from Point Clouds. 2372-2380

Lei Zhou, Siyu Zhu, Tianwei Shen, Jinglu Wang, Tian Fang, Long Quan:
Progressive Large Scale-Invariant Image Matching in Scale Space. 2381-2390

Liu Liu, Hongdong Li, Yuchao Dai:
Efficient Global 2D-3D Matching for Camera Localization in a Large-Scale 3D Map. 2391-2400

Sk. Mohammadul Haque, Venu Madhav Govindu:
Multi-view Non-rigid Refinement and Normal Selection for High Quality 3D Reconstruction. 2401-2409

Wei Shen, Bin Wang, Yuan Jiang, Yan Wang, Alan L. Yuille:
Multi-stage Multi-recursive-input Fully Convolutional Networks for Neuronal Boundary Detection. 2410-2419

Jiandong Tian, Zak Murez, Tong Cui, Zhen Zhang, David J. Kriegman, Ravi Ramamoorthi:
Depth and Image Restoration from Light Field in a Scattering Medium. 2420-2429

Ajay Nandoriya, Mohamed A. Elgharib, Changil Kim, Mohamed Hefeeda, Wojciech Matusik:
Video Reflection Removal Through Spatio-Temporal Optimization. 2430-2438

Jiahuan Zhou, Pei Yu, Wei Tang, Ying Wu:
Efficient Online Local Metric Adaptation via Negative Samples for Person Re-identification. 2439-2447

Zimo Liu, Dong Wang, Huchuan Lu:
Stepwise Metric Promotion for Unsupervised Video Person Re-identification. 2448-2457

Rui Huang, Shu Zhang, Tianyu Li, Ran He:
Beyond Face Rotation: Global and Local Perception GAN for Photorealistic

and Identity Preserving Frontal View Synthesis. 2458-2467

Giuseppe Lisanti, Niki Martinel, Alberto Del Bimbo, Gian Luca Foresti:
Group Re-identification via Unsupervised Transfer of Sparse Features Encoding. 2468-2477

Hamdi Dibeklioglu:
Visual Transformation Aided Contrastive Learning for Video-Based Kinship Verification. 2478-2487

Ming Lu, Hao Zhao, Anbang Yao, Feng Xu, Yurong Chen, Li Zhang:
Decoder Network over Lightweight Reconstructed Feature for Fast Semantic Style Transfer. 2488-2496

Jiangxin Dong, Jinshan Pan, Zhixun Su, Ming-Hsuan Yang:
Blind Image Deblurring with Outlier Handling. 2497-2505

Hamed R. Tavakoli, Rakshith Shetty, Ali Borji, Jorma Laaksonen:
Paying Attention to Descriptions Generated by Image Captioning Models. 2506-2515

Qifeng Chen, Jia Xu, Vladlen Koltun:
Fast Image Processing with Fully-Convolutional Networks. 2516-2525

Ding Liu, Zhaowen Wang, Yuchen Fan, Xianming Liu, Zhangyang Wang, Shiyu Chang, Thomas S. Huang:
Robust Video Super-Resolution with Learned Temporal Dynamics. 2526-2534

Wei Wei, Lixuan Yi, Qi Xie, Qian Zhao, Deyu Meng, Zongben Xu:
Should We Encode Rain Streaks in Video as Deterministic or Stochastic? 2535-2544

Lei Zhu, Chi-Wing Fu, Dani Lischinski, Pheng-Ann Heng:
Joint Bi-layer Optimization for Single-Image Rain Streak Removal. 2545-2553

Edoardo Remelli, Anastasia Tkach, Andrea Tagliasacchi, Mark Pauly:
Low-Dimensionality Calibration through Local Anisotropic Scaling for Robust Hand Model Personalization. 2554-2562

Andrii Maksai, Xinchao Wang, François Fleuret, Pascal Fua:
Non-Markovian Globally Consistent Multi-object Tracking. 2563-2573

Yibing Song, Chao Ma, Lijun Gong, Jiawei Zhang, Rynson W. H. Lau, Ming-Hsuan Yang:
CREST: Convolutional Residual Learning for Visual Tracking. 2574-2583

Katrin Lasinger, Christoph Vogel, Konrad Schindler:
Volumetric Flow Estimation for Incompressible Fluids Using the Stationary Stokes Equations. 2584-2592

Aseem Behl, Omid Hosseini Jafari, Siva Karthik Mustikovela, Hassan Abu Alhaija, Carsten Rother, Andreas Geiger:
Bounding Boxes, Segmentations and Object Coordinates: How Important is Recognition for 3D Scene Flow Estimation in Autonomous Driving Scenarios? 2593-2602

Zefan Li, Bingbing Ni, Wenjun Zhang, Xiaokang Yang, Wen Gao:
Performance Guaranteed Network Acceleration via High-Order Residual Quantization. 2603-2611

Jian Wang, Feng Zhou, Shilei Wen, Xiao Liu, Yuanqing Lin:
Deep Metric Learning with Angular Loss. 2612-2620

Xiao Sun, Jiaxiang Shang, Shuang Liang, Yichen Wei:
Compositional Human Pose Regression. 2621-2630

Hédi Ben-Younes, Rémi Cadène, Matthieu Cord, Nicolas Thome:
MUTAN: Multimodal Tucker Fusion for Visual Question Answering. 2631-2639

Nam N. Vo, Nathan Jacobs, James Hays:
Revisiting IM2GPS in the Deep Learning Era. 2640-2649

Wei-Chih Hung, Yi-Hsuan Tsai, Xiaohui Shen, Zhe L. Lin, Kalyan Sunkavalli, Xin Lu, Ming-Hsuan Yang:
Scene Parsing with Global Context Embedding. 2650-2658

Julieta Martinez, Rayat Hossain, Javier Romero, James J. Little:
A Simple Yet Effective Baseline for 3d Human Pose Estimation. 2659-2668

Junnan Li, Yongkang Wong, Qi Zhao, Mohan S. Kankanhalli:
Dual-Glance Model for Deciphering Social Relationships. 2669-2678

John P. Collomosse, Tu Bui, Michael J. Wilber, Chen Fang, Hailin Jin:
Sketching with Style: Visual Search with Sketches and Aesthetic Context. 2679-2687

Su Zhang, Yang Yang, Kun Yang, Yi Luo, Sim Heng Ong:
Point Set Registration with Global-Local Correspondence and Transformation Estimation. 2688-2696

John McCormac, Ankur Handa, Stefan Leutenegger, Andrew J. Davison:
SceneNet RGB-D: Can 5M Synthetic Images Beat Generic ImageNet Pre-training on Indoor Segmentation? 2697-2706

Scott Workman, Menghua Zhai, David J. Crandall, Nathan Jacobs:
A Unified Model for Near and Remote Sensing. 2707-2716

Haotian Xu, Ming Dong, Zichun Zhong:
Directionally Convolutional Networks for 3D Shape Segmentation. 2717-2726

Stavros Tsogkas, Sven J. Dickinson:
AMAT: Medial Axis Transform for Natural Images. 2727-2736

Ping Luo, Guangrun Wang, Liang Lin, Xiaogang Wang:
Deep Dual Learning for Semantic Image Segmentation. 2737-2745

Jun Hao Liew, Yunchao Wei, Wei Xiong, Sim Heng Ong, Jiashi Feng:
Regional Interactive Image Segmentation Networks. 2746-2754

Zhuang Liu, Jianguo Li, Zhiqiang Shen, Gao Huang, Shoumeng Yan, Changshui Zhang:

Learning Efficient Convolutional Networks through Network Slimming. 2755-2763

Jianmin Bao, Dong Chen, Fang Wen, Houqiang Li, Gang Hua:
CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training. 2764-2773

Jan Hendrik Metzen, Mummadi Chaithanya Kumar, Thomas Brox, Volker Fischer:
Universal Adversarial Perturbations Against Semantic Image Segmentation. 2774-2783

Philip Häusser, Thomas Frerix, Alexander Mordvintsev, Daniel Cremers:
Associative Domain Adaptation. 2784-2792

Justin Lazarow, Long Jin, Zhuowen Tu:
Introspective Neural Networks for Generative Modeling. 2793-2802

Wei Tang, Pei Yu, Jiahuan Zhou, Ying Wu:
Towards a Unified Compositional Model for Visual Pattern Modeling. 2803-2812

Xudong Mao, Qing Li, Haoran Xie, Raymond Y. K. Lau, Zhen Wang, Stephen Paul Smolley:
Least Squares Generative Adversarial Networks. 2813-2821

Lei Huang, Xianglong Liu, Yang Liu, Bo Lang, Dacheng Tao:
Centered Weight Normalization in Accelerating Training of Deep Neural Networks. 2822-2830

Guangcong Wang, Xiaohua Xie, Jianhuang Lai, Jiaxuan Zhuo:
Deep Growing Learning. 2831-2839

Ben Harwood, Vijay Kumar B. G, Gustavo Carneiro, Ian D. Reid, Tom Drummond:
Smart Mining for Deep Metric Learning. 2840-2848

Masaki Saito, Eiichi Matsumoto, Shunta Saito:
Temporal Generative Adversarial Nets with Singular Value Clipping. 2849-2858

R. Manmatha, Chao-Yuan Wu, Alexander J. Smola, Philipp Krähenbühl:
Sampling Matters in Deep Embedding Learning. 2859-2867

Zili Yi, Hao (Richard) Zhang, Ping Tan, Minglun Gong:
DualGAN: Unsupervised Dual Learning for Image-to-Image Translation. 2868-2876

Kang Zheng, Xiaochuan Fan, Yuewei Lin, Hao Guo, Hongkai Yu, Dazhou Guo, Song Wang:
Learning View-Invariant Features for Person Identification in Temporally Synchronized Videos Taken by Wearable Cameras. 2877-2885

Jonghwan Mun, Paul Hongsuck Seo, Ilchae Jung, Bohyung Han:
MarioQA: Answering Questions by Watching Gameplay Videos. 2886-2894

Xin Li, Mooi Choo Chuah:
SBGAR: Semantics Based Group Activity Recognition. 2895-2904

Davide Moltisanti, Michael Wray, Walterio W. Mayol-Cuevas, Dima Damen: **Trespassing the Boundaries: Labeling Temporal Bounds for Object Interactions in Egocentric Video.** 2905-2913

Radu Tudor Ionescu, Sorina Smeureanu, Bogdan Alexe, Marius Popescu: **Unmasking the Abnormal Events in Video.** 2914-2922

Mohammadreza Zolfaghari, Gabriel L. Oliveira, Nima Sedaghat, Thomas Brox: **Chained Multi-stream Networks Exploiting Pose, Motion, and Appearance for Action Classification and Detection.** 2923-2932

Yue Zhao, Yuanjun Xiong, Limin Wang, Zhirong Wu, Xiaoou Tang, Dahua Lin: **Temporal Action Detection with Structured Segment Networks.** 2933-2942

Yang Liu, Ping Wei, Song-Chun Zhu: **Jointly Recognizing Object Fluents and Tasks in Egocentric Videos.** 2943-2951

Hanqing Wang, Wei Liang, Lap-Fai Yu: **Transferring Objects: Joint Inference of Container and Human Pose.** 2952-2960

Jinkyu Kim, John F. Canny: **Interpretable Learning for Self-Driving Cars by Visualizing Causal Attention.** 2961-2969

Recognition 2 Oral Session 4

Abhishek Das, Satwik Kottur, José M. F. Moura, Stefan Lee, Dhruv Batra: **Learning Cooperative Visual Dialog Agents with Deep Reinforcement Learning.** 2970-2979

Kaiming He, Georgia Gkioxari, Piotr Dollár, Ross B. Girshick: **Mask R-CNN.** 2980-2988

Bo Dai, Sanja Fidler, Raquel Urtasun, Dahua Lin: **Towards Diverse and Natural Image Descriptions via a Conditional GAN.** 2989-2998

Tsung-Yi Lin, Priya Goyal, Ross B. Girshick, Kaiming He, Piotr Dollár: **Focal Loss for Dense Object Detection.** 2999-3007

Justin Johnson, Bharath Hariharan, Laurens van der Maaten, Judy Hoffman, Li Fei-Fei, C. Lawrence Zitnick, Ross B. Girshick: **Inferring and Executing Programs for Visual Reasoning.** 3008-3017

Spotlight Session 4

Kuo-Hao Zeng, William B. Shen, De-An Huang, Min Sun, Juan Carlos Niebles: **Visual Forecasting by Imitating Dynamics in Natural Sequences.** 3018-3027

Shenlong Wang, Min Bai, Gellért Máttyus, Hang Chu, Wenjie Luo, Bin Yang, Justin Liang, Joel Cheverie, Sanja Fidler, Raquel Urtasun:

TorontoCity: Seeing the World with a Million Eyes. 3028-3036

Bharath Hariharan, Ross B. Girshick:
Low-Shot Visual Recognition by Shrinking and Hallucinating Features. 3037-3046

Shaoli Huang, Mingming Gong, Dacheng Tao:
A Coarse-Fine Network for Keypoint Localization. 3047-3056

Christoph Feichtenhofer, Axel Pinz, Andrew Zisserman:
Detect to Track and Track to Detect. 3057-3065

Pan He, Weilin Huang, Tong He, Qile Zhu, Yu Qiao, Xiaolin Li:
Single Shot Text Detector with Regional Attention. 3066-3074

Necati Cihan Camgöz, Simon Hadfield, Oscar Koller, Richard Bowden:
SubUNets: End-to-End Hand Shape and Continuous Sign Language Recognition. 3075-3084

Isma Hadji, Richard P. Wildes:
A Spatiotemporal Oriented Energy Network for Dynamic Texture Recognition. 3085-3093

Poster Session 5

Paul Gay, Vaibhav Bansal, Cosimo Rubino, Alessio Del Bue:
Probabilistic Structure from Motion with Objects (PSfMO). 3094-3103

Hang Dai, Nick E. Pears, William A. P. Smith, Christian Duncan:
A 3D Morphable Model of Craniofacial Shape and Texture Variation. 3104-3112

Vincent Leroy, Jean-Sébastien Franco, Edmond Boyer:
Multi-view Dynamic Shape Refinement Using Local Temporal Integration. 3113-3122

Chiho Choi, Sangpil Kim, Karthik Ramani:
Learning Hand Articulations by Hallucinating Heat Distribution. 3123-3132

Robert Maier, Kihwan Kim, Daniel Cremers, Jan Kautz, Matthias Nießner:
Intrinsic3D: High-Quality 3D Reconstruction by Joint Appearance and Geometry Optimization with Spatially-Varying Lighting. 3133-3141

Chiho Choi, Sang Ho Yoon, Chin-Ning Chen, Karthik Ramani:
Robust Hand Pose Estimation during the Interaction with an Unknown Object. 3142-3151

Xinxin Zuo, Sen Wang, Jiangbin Zheng, Ruigang Yang:
Detailed Surface Geometry and Albedo Recovery from RGB-D Video under Natural Illumination. 3152-3161

Haoping Deng, Wangjiang Zhu:
Monocular Free-Head 3D Gaze Tracking with Deep Learning and Geometry Constraints. 3162-3171

Boaz Arad, Ohad Ben-Shahar:
Filter Selection for Hyperspectral Estimation. 3172-3180

Lixiong Chen, Yinqiang Zheng, Boxin Shi, Art Subpa-Asa, Imari Sato:
A Microfacet-Based Reflectance Model for Photometric Stereo with Highly Specular Surfaces. 3181-3189

Kaipeng Zhang, Zhanpeng Zhang, Hao Wang, Zhifeng Li, Yu Qiao, Wei Liu:
Detecting Faces Using Inside Cascaded Contextual CNN. 3190-3198

Anis Kacem, Mohamed Daoudi, Boulbaba Ben Amor, Juan Carlos Álvarez Paiva:
A Novel Space-Time Representation on the Positive Semidefinite Cone for Facial Expression Recognition. 3199-3208

Dieu Linh Tran, Robert Walecki, Ognjen Rudovic, Stefanos Eleftheriadis, Björn W. Schuller, Maja Pantic:
DeepCoder: Semi-Parametric Variational Autoencoders for Automatic Facial Action Coding. 3209-3218

Amin Jourabloo, Mao Ye, Xiaoming Liu, Liu Ren:
Pose-Invariant Face Alignment with a Single CNN. 3219-3228

James Thewlis, Hakan Bilen, Andrea Vedaldi:
Unsupervised Learning of Object Landmarks by Factorized Spatial Embeddings. 3229-3238

Liming Zhao, Xi Li, Yueting Zhuang, Jingdong Wang:
Deeply-Learned Part-Aligned Representations for Person Re-identification. 3239-3248

Jun-Tae Lee, Han-Ul Kim, Chul Lee, Chang-Su Kim:
Semantic Line Detection and Its Applications. 3249-3257

Qingnan Fan, Jiaolong Yang, Gang Hua, Baoquan Chen, David P. Wipf:
A Generic Deep Architecture for Single Image Reflection Removal and Image Smoothing. 3258-3267

Tiancheng Sun, Yifan Peng, Wolfgang Heidrich:
Revisiting Cross-Channel Information Transfer for Chromatic Aberration Correction. 3268-3276

Xiaoyong Shen, Hongyun Gao, Xin Tao, Chao Zhou, Jiaya Jia:
High-Quality Correspondence and Segmentation Estimation for Dual-Lens Smart-Phone Portraits. 3277-3286

Ming Jiang, Qi Zhao:
Learning Visual Attention to Identify People with Autism Spectrum Disorder. 3287-3296

Andrey Ignatov, Nikolay Kobyshev, Radu Timofte, Kenneth Vanhoey, Luc Van Gool:
DSLR-Quality Photos on Mobile Devices with Deep Convolutional Networks. 3297-3305

Yuval Bahat, Netalee Efrat, Michal Irani:
Non-uniform Blind Deblurring by Reblurring. 3306-3314

Takashi Shibata, Masayuki Tanaka, Masatoshi Okutomi:
Misalignment-Robust Joint Filter for Cross-Modal Image Pairs. 3315-3324

Wei Chen, Nan Song:
Low-Rank Tensor Completion: A Pseudo-Bayesian Learning Approach. 3325-3333

Tsun-Yi Yang, Jo-Han Hsu, Yen-Yu Lin, Yung-Yu Chuang:
DeepCD: Learning Deep Complementary Descriptors for Patch Representations. 3334-3342

Luka Cehovin Zajc, Alan Lukezic, Ales Leonardis, Matej Kristan:
Beyond Standard Benchmarks: Parameterizing Performance Evaluation in Visual Object Tracking. 3343-3351

Jacob Walker, Kenneth Marino, Abhinav Gupta, Martial Hebert:
The Pose Knows: Video Forecasting by Generating Pose Futures. 3352-3361

Panna Felsen, Pulkit Agrawal, Jitendra Malik:
What will Happen Next? Forecasting Player Moves in Sports Videos. 3362-3371

Mehdi Bahri, Yannis Panagakis, Stefanos Zafeiriou:
Robust Kronecker-Decomposable Component Analysis for Low-Rank Modeling. 3372-3381

Xiaodan Liang, Zhiting Hu, Hao Zhang, Chuang Gan, Eric P. Xing:
Recurrent Topic-Transition GAN for Visual Paragraph Generation. 3382-3391

Jun Li, Reinhard Klein, Angela Yao:
A Two-Streamed Network for Estimating Fine-Scaled Depth Maps from Single RGB Images. 3392-3400

Miaojing Shi, Holger Caesar, Vittorio Ferrari:
Weakly Supervised Object Localization Using Things and Stuff Transfer. 3401-3410

Zhichen Zhao, Huimin Ma, Shaodi You:
Single Image Action Recognition Using Semantic Body Part Actions. 3411-3419

Konstantin Shmelkov, Cordelia Schmid, Karteek Alahari:
Incremental Learning of Object Detectors without Catastrophic Forgetting. 3420-3429

Simone Palazzo, Concetto Spampinato, Isaak Kavasidis, Daniela Giordano, Mubarak Shah:
Generative Adversarial Networks Conditioned by Brain Signals. 3430-3438

Yining Li, Chen Huang, Xiaoou Tang, Chen Change Loy:
Learning to Disambiguate by Asking Discriminative Questions. 3439-3448

Ruth C. Fong, Andrea Vedaldi:
Interpretable Explanations of Black Boxes by Meaningful Perturbation. 3449-3457

Yeong Jun Koh, Chang-Su Kim:
CDTS: Collaborative Detection, Tracking, and Segmentation for Online Multiple Object Segmentation in Videos. 3621-3629

Se-Ho Lee, Won-Dong Jang, Chang-Su Kim:
Temporal Superpixels Based on Proximity-Weighted Patch Matching. 3630-3638

Ryota Hinami, Tao Mei, Shin'ichi Satoh:
Joint Detection and Recounting of Abnormal Events by Learning Deep Generic Knowledge. 3639-3647

Jiyang Gao, Zhenheng Yang, Chen Sun, Kan Chen, Ram Nevatia:
TURN TAP: Temporal Unit Regression Network for Temporal Action Proposals. 3648-3656

Gurkirt Singh, Suman Saha, Michael Sapienza, Philip H. S. Torr, Fabio Cuzzolin:
Online Real-Time Multiple Spatiotemporal Action Localisation and Prediction. 3657-3666

Heng Tao Shen, Chao Li, Jiewei Cao, Zi Huang, Lei Zhu:
Leveraging Weak Semantic Relevance for Complex Video Event Classification. 3667-3676

Rameswar Panda, Abir Das, Ziyan Wu, Jan Ernst, Amit K. Roy-Chowdhury:
Weakly Supervised Summarization of Web Videos. 3677-3686

Shanghang Zhang, Guanhang Wu, João Paulo Costeira, José M. F. Moura:
FCN-rLSTM: Deep Spatio-Temporal Neural Networks for Vehicle Counting in City Cameras. 3687-3696

Iryna Korshunova, Wenzhe Shi, Joni Dambre, Lucas Theis:
Fast Face-Swap Using Convolutional Neural Networks. 3697-3705

Tribhuvanesh Orekondy, Bernt Schiele, Mario Fritz:
Towards a Visual Privacy Advisor: Understanding and Predicting Privacy Risks in Images. 3706-3715

Face and Human Behaviour Analysis Oral Session 5

Nicholas Rhinehart, Kris M. Kitani:
First-Person Activity Forecasting with Online Inverse Reinforcement Learning. 3716-3725

Adrian Bulat, Georgios Tzimiropoulos:
Binarized Convolutional Landmark Localizers for Human Pose Estimation and Face Alignment with Limited Resources. 3726-3734

Ayush Tewari, Michael Zollhöfer, Hyeongwoo Kim, Pablo Garrido, Florian Bernard, Patrick Pérez, Christian Theobalt:
MoFA: Model-Based Deep Convolutional Face Autoencoder for Unsupervised Monocular Reconstruction. 3735-3744

Wenbin Du, Yali Wang, Yu Qiao:
RPAN: An End-to-End Recurrent Pose-Attention Network for Action

Recognition in Videos. 3745-3754

Chi Nhan Duong, Kha Gia Quach, Khoa Luu, T. Hoang Ngan Le, Marios Savvides:
Temporal Non-volume Preserving Approach to Facial Age-Progression and Age-Invariant Face Recognition. 3755-3763

Spotlight Session 5

Guosheng Hu, Yang Hua, Yang Yuan, Zhihong Zhang, Zheng Lu, Sankha S. Mukherjee, Timothy M. Hospedales, Neil Martin Robertson, Yongxin Yang:
Attribute-Enhanced Face Recognition with Neural Tensor Fusion Networks. 3764-3773

Zhedong Zheng, Liang Zheng, Yi Yang:
Unlabeled Samples Generated by GAN Improve the Person Re-identification Baseline in Vitro. 3774-3782

Congqi Cao, Yifan Zhang, Yi Wu, Hanqing Lu, Jian Cheng:
Egocentric Gesture Recognition Using Recurrent 3D Convolutional Neural Networks with Spatiotemporal Transformer Modules. 3783-3791

Wanglong Wu, Meina Kan, Xin Liu, Yi Yang, Shiguang Shan, Xilin Chen:
Recursive Spatial Transformer (ReST) for Alignment-Free Face Recognition. 3792-3800

Yongming Rao, Ji Lin, Jiwen Lu, Jie Zhou:
Learning Discriminative Aggregation Network for Video-Based Face Recognition. 3801-3810

Muhammad Haris Khan, John McDonagh, Georgios Tzimiropoulos:
Synergy between Face Alignment and Tracking via Discriminative Global Consensus Optimization. 3811-3819

Yifan Sun, Liang Zheng, Weijian Deng, Shengjin Wang:
SVDNet for Pedestrian Retrieval. 3820-3828

Zijing Zhao, Ajay Kumar:
Towards More Accurate Iris Recognition Using Deeply Learned Spatially Corresponding Features. 3829-3838

Poster Session 6

Maros Blaha, Mathias Rothermel, Martin R. Oswald, Torsten Sattler, Audrey Richard, Jan Dirk Wegner, Marc Pollefeys, Konrad Schindler:
Semantically Informed Multiview Surface Refinement. 3839-3847

Mahdi Rad, Vincent Lepetit:
BB8: A Scalable, Accurate, Robust to Partial Occlusion Method for Predicting the 3D Poses of Challenging Objects without Using Depth. 3848-3856

William Nguatem, Helmut Mayer:
Modeling Urban Scenes from Pointclouds. 3857-3866

Filippo Bergamasco, Luca Cosmo, Andrea Gasparetto, Andrea Albarelli, Andrea Torsello:
Parameter-Free Lens Distortion Calibration of Central Cameras. 3867-3875

Vassileios Balntas, Andreas Doumanoglou, Caner Sahin, Juil Sock, Rigas Kouskouridas, Tae-Kyun Kim:
Pose Guided RGBD Feature Learning for 3D Object Pose Estimation. 3876-3884

Andreas Schneider, Sandro Schönborn, Bernhard Egger, Lavrenti Frobeen, Thomas Vetter:
Efficient Global Illumination for Morphable Models. 3885-3893

Sean Ryan Fanello, Julien P. C. Valentin, Adarsh Kowdle, Christoph Rhemann, Vladimir Tankovich, Carlo Ciliberto, Philip Davidson, Shahram Izadi:
Low Compute and Fully Parallel Computer Vision with HashMatch. 3894-3903

Mathias Gallardo, Toby Collins, Adrien Bartoli:
Dense Non-rigid Structure-from-Motion and Shading with Unknown Albedos. 3904-3912

Lubor Ladicky, Olivier Saurer, SoHyeon Jeong, Fabio Maninchedda, Marc Pollefeys:
From Point Clouds to Mesh Using Regression. 3913-3922

Rui Wang, Martin Schwörer, Daniel Cremers:
Stereo DSO: Large-Scale Direct Sparse Visual Odometry with Stereo Cameras. 3923-3931

Minhaeng Lee, Charless C. Fowlkes:
Space-Time Localization and Mapping. 3932-3941

Renjie Wan, Boxin Shi, Ling-Yu Duan, Ah-Hwee Tan, Alex C. Kot:
Benchmarking Single-Image Reflection Removal Algorithms. 3942-3950

Yongming Rao, Jiwen Lu, Jie Zhou:
Attention-Aware Deep Reinforcement Learning for Video Face Recognition. 3951-3960

Bugra Tekin, Pablo Márquez-Neila, Mathieu Salzmann, Pascal Fua:
Learning to Fuse 2D and 3D Image Cues for Monocular Body Pose Estimation. 3961-3970

Shan Wu, Shangfei Wang, Bowen Pan, Qiang Ji:
Deep Facial Action Unit Recognition from Partially Labeled Data. 3971-3979

Chi Su, Jianing Li, Shiliang Zhang, Junliang Xing, Wen Gao, Qi Tian:
Pose-Driven Deep Convolutional Model for Person Re-identification. 3980-3989

Carlos Fabian Benitez-Quiroz, Yan Wang, Aleix M. Martínez:
Recognition of Action Units in the Wild with Deep Nets and a New Global-Local Loss. 3990-3999

Chandrasekhar Bhagavatula, Chenchen Zhu, Khoa Luu, Marios Savvides:
Faster than Real-Time Facial Alignment: A 3D Spatial Transformer Network Approach in Unconstrained Poses. 4000-4009

Xi Yin, Xiang Yu, Kihyuk Sohn, Xiaoming Liu, Manmohan Chandraker:
Towards Large-Pose Face Frontalization in the Wild. 4010-4019

Bolun Cai, Xianming Xu, Kailing Guo, Kui Jia, Bin Hu, Dacheng Tao:
A Joint Intrinsic-Extrinsic Prior Model for Retinex. 4020-4029

Mahesh Mohan M. R., A. N. Rajagopalan:
Going Unconstrained with Rolling Shutter Deblurring. 4030-4038

Tiantian Wang, Ali Borji, Lihe Zhang, Pingping Zhang, Huchuan Lu:
A Stagewise Refinement Model for Detecting Salient Objects in Images. 4039-4048

Shir Gur, Ohad Ben-Shahar:
From Square Pieces to Brick Walls: The Next Challenge in Solving Jigsaw Puzzles. 4049-4057

Tae Hyun Kim, Kyoung Mu Lee, Bernhard Schölkopf, Michael Hirsch:
Online Video Deblurring via Dynamic Temporal Blending Network. 4058-4067

Dingwen Zhang, Junwei Han, Yu Zhang:
Supervision by Fusion: Towards Unsupervised Learning of Deep Salient Object Detector. 4068-4076

Roberto Tron, Xiaowei Zhou, Carlos Esteves, Kostas Daniilidis:
Fast Multi-image Matching via Density-Based Clustering. 4077-4086

Agrim Gupta, Justin Johnson, Alexandre Alahi, Li Fei-Fei:
Characterizing and Improving Stability in Neural Style Transfer. 4087-4096

Venice Erin Liong, Jiwen Lu, Yap-Peng Tan, Jie Zhou:
Cross-Modal Deep Variational Hashing. 4097-4105

Xinlei Chen, Abhinav Gupta:
Spatial Memory for Context Reasoning in Object Detection. 4106-4116

Yuming Shen, Li Liu, Ling Shao, Jingkuan Song:
Deep Binaries: Encoding Semantic-Rich Cues for Efficient Textual-Visual Cross Retrieval. 4117-4126

Yu Liu, Yanming Guo, Erwin M. Bakker, Michael S. Lew:
Learning a Recurrent Residual Fusion Network for Multimodal Matching. 4127-4136

Anders Glent Buch, Lilita Kiforenko, Dirk Kraft:
Rotational Subgroup Voting and Pose Clustering for Robust 3D Object Recognition. 4137-4145

Yousong Zhu, Chaoyang Zhao, Jinqiao Wang, Xu Zhao, Yi Wu, Hanqing Lu:
CoupleNet: Coupling Global Structure with Local Parts for Object Detection. 4146-4154

Rakshith Shetty, Marcus Rohrbach, Lisa Anne Hendricks, Mario Fritz, Bernt Schiele:
Speaking the Same Language: Matching Machine to Human Captions by Adversarial Training. 4155-4164

Meng-Ru Hsieh, Yen-Liang Lin, Winston H. Hsu:
Drone-Based Object Counting by Spatially Regularized Regional Proposal Network. 4165-4173

Nikita Dvornik, Konstantin Shmelkov, Julien Mairal, Cordelia Schmid:
BlitzNet: A Real-Time Deep Network for Scene Understanding. 4174-4182

Ruiyu Li, Makarand Tapaswi, Renjie Liao, Jiaya Jia, Raquel Urtasun, Sanja Fidler:
Situation Recognition with Graph Neural Networks. 4183-4192

Ang Li, Allan Jabri, Armand Joulin, Laurens van der Maaten:
Learning Visual N-Grams from Web Data. 4193-4202

Chiori Hori, Takaaki Hori, Teng-Yok Lee, Ziming Zhang, Bret Harsham, John R. Hershey, Tim K. Marks, Kazuhiro Sumi:
Attention-Based Multimodal Fusion for Video Description. 4203-4212

Wei-Lin Hsiao, Kristen Grauman:
Learning the Latent "Look": Unsupervised Discovery of a Style-Coherent Embedding from Fashion Images. 4213-4222

Tanmay Gupta, Kevin J. Shih, Saurabh Singh, Derek Hoiem:
Aligned Image-Word Representations Improve Inductive Transfer Across Vision-Language Tasks. 4223-4232

Huajie Jiang, Ruiping Wang, Shiguang Shan, Yi Yang, Xilin Chen:
Learning Discriminative Latent Attributes for Zero-Shot Classification. 4233-4242

Hanwang Zhang, Zawlin Kyaw, Jinyang Yu, Shih-Fu Chang:
PPR-FCN: Weakly Supervised Visual Relation Detection via Parallel Pairwise R-FCN. 4243-4251

Margret Keuper:
Higher-Order Minimum Cost Lifted Multicuts for Motion Segmentation. 4252-4260

Haoyang Zhang, Xuming He:
Deep Free-Form Deformation Network for Object-Mask Registration. 4261-4269

Matteo Denitto, Simone Melzi, Manuele Bicego, Umberto Castellani, Alessandro Farinelli, Mário A. T. Figueiredo, Yanir Kleiman, Maks Ovsjanikov:
Region-Based Correspondence Between 3D Shapes via Spatially Smooth Biclustering. 4270-4279

Anoop Cherian, Panagiotis Stanitsas, Mehrtash Harandi, Vassilios Morellas, Nikos Papanikolopoulos:
Learning Discriminative αβ-Divergences for Positive Definite Matrices. 4280-4289

Sara Shaheen, Lama Affara, Bernard Ghanem:
Constrained Convolutional Sparse Coding for Parametric Based Reconstruction of Line Drawings. 4434-4442

Tomas Wilkinson, Jonas Lindström, Anders Brun:
Neural Ctrl-F: Segmentation-Free Query-by-String Word Spotting in Handwritten Manuscript Collections. 4443-4452

Video Analysis Oral Session 6

Pascal Mettes, Cees G. M. Snoek:
Spatial-Aware Object Embeddings for Zero-Shot Localization and Classification of Actions. 4453-4462

Raghudeep Gadde, Varun Jampani, Peter V. Gehler:
Semantic Video CNNs Through Representation Warping. 4463-4472

Ziwei Liu, Raymond A. Yeh, Xiaoou Tang, Yiming Liu, Aseem Agarwala:
Video Frame Synthesis Using Deep Voxel Flow. 4473-4481

Xin Tao, Hongyun Gao, Renjie Liao, Jue Wang, Jiaya Jia:
Detail-Revealing Deep Video Super-Resolution. 4482-4490

Pavel Tokmakov, Karteek Alahari, Cordelia Schmid:
Learning Video Object Segmentation with Visual Memory. 4491-4500

Low-Level Vision Oral Session 7

Mehdi S. M. Sajjadi, Bernhard Schölkopf, Michael Hirsch:
EnhanceNet: Single Image Super-Resolution Through Automated Texture Synthesis. 4501-4510

Ying-Cong Chen, Xiaoyong Shen, Jiaya Jia:
Makeup-Go: Blind Reversion of Portrait Edit. 4511-4519

Vu Nguyen, Tomas F. Yago Vicente, Maozheng Zhao, Minh Hoai, Dimitris Samaras:
Shadow Detection with Conditional Generative Adversarial Networks. 4520-4528

Jinsong Zhang, Jean-François Lalonde:
Learning High Dynamic Range from Outdoor Panoramas. 4529-4538

Seungryong Kim, Dongbo Min, Stephen Lin, Kwanghoon Sohn:
DCTM: Discrete-Continuous Transformation Matching for Semantic Flow. 4539-4548

Spotlight Session 6

Ying Tai, Jian Yang, Xiaoming Liu, Chunyan Xu:
MemNet: A Persistent Memory Network for Image Restoration. 4549-4557

Deng-Ping Fan, Ming-Ming Cheng, Yun Liu, Tao Li, Ali Borji:
Structure-Measure: A New Way to Evaluate Foreground Maps. 4558-4567

Donghyeon Cho, Jinsun Park, Tae-Hyun Oh, Yu-Wing Tai, In So Kweon:
Weakly- and Self-Supervised Learning for Content-Aware Deep Image Retargeting. 4568-4577

Eleonora Maset, Federica Arrigoni, Andrea Fusiello:
Practical and Efficient Multi-view Matching. 4578-4586

Yu-Sheng Lin, Wei-Chao Chen, Shao-Yi Chien:
Unrolled Memory Inner-Products: An Abstract GPU Operator for Efficient Vision-Related Computations. 4587-4595

Jakob Kruse, Carsten Rother, Uwe Schmidt:
Learning to Push the Limits of Efficient FFT-Based Image Deconvolution. 4596-4604

Xu Zhang, Felix X. Yu, Sanjiv Kumar, Shih-Fu Chang:
Learning Spread-Out Local Feature Descriptors. 4605-4613

Laurie Bose, Jianing Chen, Stephen J. Carey, Piotr Dudek, Walterio W. Mayol-Cuevas:
Visual Odometry for Pixel Processor Arrays. 4614-4622

Poster Session 7

Haesol Park, Kyoung Mu Lee:
Joint Estimation of Camera Pose, Depth, Deblurring, and Super-Resolution from a Blurred Image Sequence. 4623-4631

Jean Lahoud, Bernard Ghanem:
2D-Driven 3D Object Detection in RGB-D Images. 4632-4640

Yingliang Zhang, Peihong Yu, Wei Yang, Yuanxi Ma, Jingyi Yu:
Ray Space Features for Plenoptic Structure-from-Motion. 4641-4649

Ryo Furukawa, Ryusuke Sagawa, Hiroshi Kawasaki:
Depth Estimation Using Structured Light Flow - Analysis of Projected Pattern Flow on an Object's Surface. 4650-4658

Suryansh Kumar, Yuchao Dai, Hongdong Li:
Monocular Dense 3D Reconstruction of a Complex Dynamic Scene from Two Perspective Frames. 4659-4667

Luc Van Gool, Danda Pani Paudel, Adlane Habed:
Optimal Transformation Estimation with Semantic Cues. 4668-4677

Xikang Zhang, Bengisu Özbay, Mario Sznaier, Octavia I. Camps:
Dynamics Enhanced Multi-camera Motion Segmentation from Unsynchronized Videos. 4678-4686

Oscar Mendez Maldonado, Simon Hadfield, Nicolas Pugeault, Richard Bowden:

Taking the Scenic Route to 3D: Optimising Reconstruction from Moving Cameras. 4687-4695

W. Nicholas Greene, Nicholas Roy:
FLaME: Fast Lightweight Mesh Estimation Using Variational Smoothing on Delaunay Graphs. 4696-4704

Markus Rempfler, Jan-Hendrik Lange, Florian Jug, Corinna Blasse, Eugene W. Myers, Bjoern H. Menze, Bjoern Andres:
Efficient Algorithms for Moral Lineage Tracing. 4705-4714

Yan Jia, Yinqiang Zheng, Lin Gu, Art Subpa-Asa, Antony Lam, Yoichi Sato, Imari Sato:
From RGB to Spectrum for Natural Scenes via Manifold-Based Mapping. 4715-4723

K. Ram Prabhakar, V. Sai Srikar, R. Venkatesh Babu:
DeepFuse: A Deep Unsupervised Approach for Exposure Fusion with Extreme Exposure Image Pairs. 4724-4732

Ronald Yu, Shunsuke Saito, Haoxiang Li, Duygu Ceylan, Hao Li:
Learning Dense Facial Correspondences in Unconstrained Images. 4733-4742

Shuangjie Xu, Yu Cheng, Kang Gu, Yang Yang, Shiyu Chang, Pan Zhou:
Jointly Attentive Spatial-Temporal Pooling Networks for Video-Based Person Re-identification. 4743-4752

Yeong Won Kim, Chang-Ryeol Lee, Dae Yong Cho, Yong Hoon Kwon, Hyeok-Jae Choi, Kuk-Jin Yoon:
Automatic Content-Aware Projection for 360° Videos. 4753-4761

Thekke Madam Nimisha, Akash Kumar Singh, A. N. Rajagopalan:
Blur-Invariant Deep Learning for Blind-Deblurring. 4762-4770

Georgios Zoumpourlis, Alexandros Doumanoglou, Nicholas Vretos, Petros Daras:
Non-linear Convolution Filters for CNN-Based Learning. 4771-4779

Boyi Li, Xiulian Peng, Zhangyang Wang, Jizheng Xu, Dan Feng:
AOD-Net: All-in-One Dehazing Network. 4780-4788

Tushar Sandhan, Jin Young Choi:
Simultaneous Detection and Removal of High Altitude Clouds from an Image. 4789-4798

Matthias Kümmerer, Thomas S. A. Wallis, Leon A. Gatys, Matthias Bethge:
Understanding Low- and High-Level Contributions to Fixation Prediction. 4799-4808

Tong Tong, Gen Li, Xiejie Liu, Qinquan Gao:
Image Super-Resolution Using Dense Skip Connections. 4809-4817

Sunghyun Cho, Seungyong Lee:
Convergence Analysis of MAP Based Blur Kernel Estimation. 4818-4826

Gang Wang, Carlos Lopez-Molina, Bernard De Baets:
Blob Reconstruction Using Unilateral Second Order Gaussian Kernels with Application to High-ISO Long-Exposure Image Denoising. 4827-4835

Leonardo Galteri, Lorenzo Seidenari, Marco Bertini, Alberto Del Bimbo:
Deep Generative Adversarial Compression Artifact Removal. 4836-4845

Qi Chu, Wanli Ouyang, Hongsheng Li, Xiaogang Wang, Bin Liu, Nenghai Yu:
Online Multi-object Tracking Using CNN-Based Single Object Tracker with Spatial-Temporal Attention Mechanism. 4846-4855

Yuan Liao, Xiaoqing Lu, Chengcui Zhang, Yongtao Wang, Zhi Tang:
Mutual Enhancement for Detection of Multiple Logos in Sports Videos. 4856-4865

Jingyu Liu, Liang Wang, Ming-Hsuan Yang:
Referring Expression Generation and Comprehension via Attributes. 4866-4874

Chen-Yu Lee, Vijay Badrinarayanan, Tomasz Malisiewicz, Andrew Rabinovich:
RoomNet: End-to-End Room Layout Estimation. 4875-4884

Mahyar Najibi, Pouya Samangouei, Rama Chellappa, Larry S. Davis:
SSH: Single Stage Headless Face Detector. 4885-4894

Artem Babenko, Victor S. Lempitsky:
AnnArbor: Approximate Nearest Neighbors Using Arborescence Coding. 4895-4903

Ting Yao, Yingwei Pan, Yehao Li, Zhaofan Qiu, Tao Mei:
Boosting Image Captioning with Attributes. 4904-4912

Christian Zimmermann, Thomas Brox:
Learning to Estimate 3D Hand Pose from Single RGB Images. 4913-4921

Yang Song, Fan Zhang, Qing Li, Heng Huang, Lauren J. O'Donnell, Weidong Cai:
Locally-Transferred Fisher Vectors for Texture Classification. 4922-4930

Jianxiang Ma, Anlong Ming, Zilong Huang, Xinggang Wang, Yu Zhou:
Object-Level Proposals. 4931-4939

Dim P. Papadopoulos, Jasper R. R. Uijlings, Frank Keller, Vittorio Ferrari:
Extreme Clicking for Efficient Object Annotation. 4940-4949

Han Hu, Chengquan Zhang, Yuxuan Luo, Yuzhuo Wang, Junyu Han, Errui Ding:
WordSup: Exploiting Word Annotations for Character Based Text Detection. 4950-4959

Garrick Brazil, Xi Yin, Xiaoming Liu:
Illuminating Pedestrians via Simultaneous Detection and Segmentation. 4960-4969

Marcel Simon, Yang Gao, Trevor Darrell, Joachim Denzler, Erik Rodner:
Generalized Orderless Pooling Performs Implicit Salient Matching. 4970-4979

Jawadul H. Bappy, Amit K. Roy-Chowdhury, Jason Bunk, Lakshmanan Nataraj, B. S. Manjunath:
Exploiting Spatial Structure for Localizing Manipulated Image Regions. 4980-4989

Seungyong Lee, Seong-Jin Park, Ki-Sang Hong:
RDFNet: RGB-D Multi-level Residual Feature Fusion for Indoor Semantic Segmentation. 4990-4999

Gerhard Neuhold, Tobias Ollmann, Samuel Rota Bulò, Peter Kontschieder:
The Mapillary Vistas Dataset for Semantic Understanding of Street Scenes. 5000-5009

Yue Wu, Prem Natarajan:
Self-Organized Text Detection with Minimal Post-processing via Border Learning. 5010-5019

Monami Banerjee, Rudrasis Chakraborty, Baba C. Vemuri:
Sparse Exact PGA on Riemannian Manifolds. 5020-5028

Qiong Luo, Zhi Han, Xiai Chen, Yao Wang, Deyu Meng, Dong Liang, Yandong Tang:
Tensor RPCA by Bayesian CP Factorization with Complex Noise. 5029-5038

Guoli Song, Shuhui Wang, Qingming Huang, Qi Tian:
Multimodal Gaussian Process Latent Variable Models with Harmonization. 5039-5047

Adam W. Harley, Konstantinos G. Derpanis, Iasonas Kokkinos:
Segmentation-Aware Convolutional Networks Using Local Attention Masks. 5048-5057

Diego Marcos, Michele Volpi, Nikos Komodakis, Devis Tuia:
Rotation Equivariant Vector Field Networks. 5058-5067

Jian-Hao Luo, Jianxin Wu, Weiyao Lin:
ThiNet: A Filter Level Pruning Method for Deep Neural Network Compression. 5068-5076

Fabio Maria Carlucci, Lorenzo Porzi, Barbara Caputo, Elisa Ricci, Samuel Rota Bulò:
AutoDIAL: Automatic Domain Alignment Layers. 5077-5085

Zhanzhan Cheng, Fan Bai, Yunlu Xu, Gang Zheng, Shiliang Pu, Shuigeng Zhou:
Focusing Attention: Towards Accurate Text Recognition in Natural Images. 5086-5094

Emanuela Haller, Marius Leordeanu:
Unsupervised Object Segmentation in Video by Efficient Selection of Highly Probable Positive Features. 5095-5103

Prasoon Goyal, Zhiting Hu, Xiaodan Liang, Chenyu Wang, Eric P. Xing, Carnegie Mellon:
Nonparametric Variational Auto-Encoders for Hierarchical Representation Learning. 5104-5112

Siddhartha Chandra, Nicolas Usunier, Iasonas Kokkinos:
Dense and Low-Rank Gaussian CRFs Using Deep Embeddings. 5113-5122

Quan Gan, Shangfei Wang, Longfei Hao, Qiang Ji:
A Multimodal Deep Regression Bayesian Network for Affective Video Content Analyses. 5123-5132

Moein Shakeri, Hong Zhang:
Moving Object Detection in Time-Lapse or Motion Trigger Image Sequences Using Low-Rank and Invariant Sparse Decomposition. 5133-5141

Yizhe Zhu, Ahmed M. Elgammal:
A Multilayer-Based Framework for Online Background Subtraction with Freely Moving Cameras. 5142-5151

Mang Ye, Andy Jinhua Ma, Liang Zheng, Jiawei Li, Pong C. Yuen:
Dynamic Label Graph Matching for Unsupervised Video Re-identification. 5152-5160

Feng Xiong, Xingjian Shi, Dit-Yan Yeung:
Spatiotemporal Modeling for Crowd Counting in Videos. 5161-5169

Tae-Hyun Oh, Kyungdon Joo, Neel Joshi, Baoyuan Wang, In So Kweon, Sing Bing Kang:
Personalized Cinemagraphs Using Semantic Understanding and Collaborative Learning. 5170-5179

Stamatios Georgoulis, Konstantinos Rematas, Tobias Ritschel, Mario Fritz, Tinne Tuytelaars, Luc Van Gool:
What is Around the Camera? 5180-5188

Recognition 3 Oral Session 8

Julia Peyre, Ivan Laptev, Cordelia Schmid, Josef Sivic:
Weakly-Supervised Learning of Visual Relations. 5189-5198

Michael Opitz, Georg Waltner, Horst Possegger, Horst Bischof:
BIER - Boosting Independent Embeddings Robustly. 5199-5208

Xiaojuan Qi, Renjie Liao, Jiaya Jia, Sanja Fidler, Raquel Urtasun:
3D Graph Neural Networks for RGBD Semantic Segmentation. 5209-5218

Heliang Zheng, Jianlong Fu, Tao Mei, Jiebo Luo:
Learning Multi-attention Convolutional Neural Network for Fine-Grained Image Recognition. 5219-5227

David Novotný, Diane Larlus, Andrea Vedaldi:
Learning 3D Object Categories by Looking Around Them. 5228-5237

Spotlight Session 7

Matteo Poggi, Fabio Tosi, Stefano Mattoccia:
Quantitative Evaluation of Confidence Measures in a Machine Learning

World. 5238-5247

Hui Li, Peng Wang, Chunhua Shen:
Towards End-to-End Text Spotting with Convolutional Recurrent Neural Networks. 5248-5256

Seyed Hamid Rezatofighi, Vijay Kumar B. G, Anton Milan, Ehsan Abbasnejad, Anthony R. Dick, Ian D. Reid:
DeepSetNet: Predicting Sets with Deep Neural Networks. 5257-5266

Antoine Miech, Jean-Baptiste Alayrac, Piotr Bojanowski, Ivan Laptev, Josef Sivic:
Learning from Video and Text via Large-Scale Discriminative Clustering. 5267-5276

Jiyang Gao, Chen Sun, Zhenheng Yang, Ram Nevatia:
TALL: Temporal Activity Localization via Language Query. 5277-5285

Sou-Young Jin, Hang Su, Chris Stauffer, Erik G. Learned-Miller:
End-to-End Face Detection and Cast Grouping in Movies Using Erdös-Rényi Clustering. 5286-5295

Miriam W. Huijser, Jan C. van Gemert:
Active Decision Boundary Annotation with Deep Generative Models. 5296-5305

Vardan Papyan, Yaniv Romano, Michael Elad, Jeremias Sulam:
Convolutional Dictionary Learning via Local Processing. 5306-5314

Poster Session 8

Paul A. Beardsley, Gaurav Chaurasia:
Editable Parametric Dense Foliage from 3D Capture. 5315-5324

François Chadebecq, Francisco Vasconcelos, George Dwyer, Rene M. Lacher, Sébastien Ourselin, Tom Vercauteren, Danail Stoyanov:
Refractive Structure-from-Motion Through a Flat Refractive Interface. 5325-5333

Mike Roberts, Shital Shah, Debadeepta Dey, Anh Truong, Sudipta N. Sinha, Ashish Kapoor, Pat Hanrahan, Neel Joshi:
Submodular Trajectory Optimization for Aerial 3D Scanning. 5334-5343

Gil Ben-Artzi:
Camera Calibration by Global Constraints on the Motion of Silhouettes. 5344-5353

Hyowon Ha, Michal Perdoch, Hatem Alismail, In So Kweon, Yaser Sheikh:
Deltille Grids for Geometric Camera Calibration. 5354-5362

Wolfgang Stürzl:
A Lightweight Single-Camera Polarization Compass with Covariance Estimation. 5363-5371

Zhuo Hui, Kalyan Sunkavalli, Joon-Young Lee, Sunil Hadap, Jian Wang, Aswin C. Sankaranarayanan:

Reflectance Capture Using Univariate Sampling of BRDFs. 5372-5380

Guodong Xu, Yuhui Quan, Hui Ji:
Estimating Defocus Blur via Rank of Local Patches. 5381-5389

Ancong Wu, Wei-Shi Zheng, Hong-Xing Yu, Shaogang Gong, Jianhuang Lai:
RGB-Infrared Cross-Modality Person Re-identification. 5390-5399

Xiaokang Yu, Na Lei, Yalin Wang, Xianfeng Gu:
Intrinsic 3D Dynamic Surface Tracking based on Dynamic Ricci Flow and Teichmüller Map. 5400-5408

Xuelin Qian, Yanwei Fu, Yu-Gang Jiang, Tao Xiang, Xiangyang Xue:
Multi-scale Deep Learning Architectures for Person Re-identification. 5409-5418

Xiao Zhang, Zhiyuan Fang, Yandong Wen, Zhifeng Li, Yu Qiao:
Range Loss for Deep Face Recognition with Long-Tailed Training Data. 5419-5428

Shruti Nagpal, Maneet Singh, Richa Singh, Mayank Vatsa, Afzel Noore, Angshul Majumdar:
Face Sketch Matching via Coupled Deep Transform Learning. 5429-5438

Kyle Olszewski, Zimo Li, Chao Yang, Yi Zhou, Ronald Yu, Zeng Huang, Sitao Xiang, Shunsuke Saito, Pushmeet Kohli, Hao Li:
Realistic Dynamic Facial Textures from a Single Image Using GANs. 5439-5448

Ryan Dahl, Mohammad Norouzi, Jonathon Shlens:
Pixel Recursive Super Resolution. 5449-5458

Yanlin Qian, Ke Chen, Jarno Nikkanen, Joni-Kristian Kamarainen, Jiri Matas:
Recurrent Color Constancy. 5459-5467

Lei Zhu, Haibin Ling, Jin Wu, Huiping Deng, Jin Liu:
Saliency Pattern Detection by Ranking Structured Trees. 5468-5477

Yousef Atoum, Joseph Roth, Michael Bliss, Wende Zhang, Xiaoming Liu:
Monocular Video-Based Trailer Coupler Detection Using Multiplexer Convolutional Neural Network. 5478-5486

Heng Fan, Haibin Ling:
Parallel Tracking and Verifying: A Framework for Real-Time and High Accuracy Visual Tracking. 5487-5495

Xin Sun, Ngai-Man Cheung, Hongxun Yao, Yiluan Guo:
Non-rigid Object Tracking via Deformable Patches Using Shape-Preserved KCF and Level Sets. 5496-5504

Chen Wang, Charles Herrmann, Ramin Zabih:
A Discriminative View of MRF Pre-processing Algorithms. 5505-5514

Elias N. Zois, Ilias Theodorakopoulos, George Economou:
Offline Handwritten Signature Modeling and Verification Based on Archetypal Analysis. 5515-5524

Huseyin Coskun, Felix Achilles, Robert S. DiPietro, Nassir Navab, Federico Tombari:
Long Short-Term Memory Kalman Filters: Recurrent Neural Estimators for Pose Regularization. 5525-5533

Zhaofan Qiu, Ting Yao, Tao Mei:
Learning Spatio-Temporal Representation with Pseudo-3D Residual Networks. 5534-5542

Da Li, Yongxin Yang, Yi-Zhe Song, Timothy M. Hospedales:
Deeper, Broader and Artier Domain Generalization. 5543-5551

Jifei Song, Qian Yu, Yi-Zhe Song, Tao Xiang, Timothy M. Hospedales:
Deep Spatial-Semantic Attention for Fine-Grained Sketch-Based Image Retrieval. 5552-5561

Navaneeth Bodla, Bharat Singh, Rama Chellappa, Larry S. Davis:
Soft-NMS - Improving Object Detection with One Line of Code. 5562-5570

Aron Yu, Kristen Grauman:
Semantic Jitter: Dense Supervision for Visual Comparisons via Synthetic Images. 5571-5580

Xiaojie Jin, Xin Li, Huaxin Xiao, Xiaohui Shen, Zhe Lin, Jimei Yang, Yunpeng Chen, Jian Dong, Luoqi Liu, Zequn Jie, Jiashi Feng, Shuicheng Yan:
Video Scene Parsing with Predictive Feature Learning. 5581-5589

Scott Workman, Richard Souvenir, Nathan Jacobs:
Understanding and Mapping Natural Beauty. 5590-5599

Ke Sun, Cuiling Lan, Junliang Xing, Wenjun Zeng, Dong Liu, Jingdong Wang:
Human Pose Estimation Using Global and Local Normalization. 5600-5608

Zhangjie Cao, Mingsheng Long, Jianmin Wang, Philip S. Yu:
HashNet: Deep Learning to Hash by Continuation. 5609-5618

Edouard Oyallon, Eugene Belilovsky, Sergey Zagoruyko:
Scaling the Scattering Transform: Deep Hybrid Networks. 5619-5628

Takumi Kobayashi:
Flip-Invariant Motion Representation. 5629-5638

Salman H. Khan, Munawar Hayat, Fatih Porikli:
Scene Categorization with Spectral Features. 5639-5649

Xuelong Li, Di Hu, Xiaoqiang Lu:
Image2song: Song Retrieval via Bridging Image Content and Lyric Words. 5650-5659

Or Litany, Tal Remez, Emanuele Rodolà, Alexander M. Bronstein, Michael M. Bronstein:
Deep Functional Maps: Structured Prediction for Dense Shape Correspondence. 5660-5668

Nicholas I. Kolkin, Gregory Shakhnarovich, Eli Shechtman:
Training Deep Networks to be Spatially Sensitive. 5669-5678

Rui Hou, Chen Chen, Mubarak Shah:
Tube Convolutional Neural Network (T-CNN) for Action Detection in Videos. 5823-5832

Hossein Rahmani, Mohammed Bennamoun:
Learning Action Recognition Model from Depth and Skeleton Videos. 5833-5842

Raghav Goyal, Samira Ebrahimi Kahou, Vincent Michalski, Joanna Materzynska, Susanne Westphal, Heuna Kim, Valentin Haenel, Ingo Fründ, Peter Yianilos, Moritz Mueller-Freitag, Florian Hoppe, Christian Thurau, Ingo Bax, Roland Memisevic:
The "Something Something" Video Database for Learning and Evaluating Visual Common Sense. 5843-5851

Avi Singh, Larry Yang, Sergey Levine:
GPLAC: Generalizing Vision-Based Robotic Skills Using Weakly Labeled Images. 5852-5861

Wei Liu, Xiaogang Chen, Chunhua Shen, Zhi Liu, Jie Yang:
Semi-Global Weighted Least Squares in Image Filtering. 5862-5870

Xiaochuan Yin, Xiangwei Wang, Xiaoguo Du, Qijun Chen:
Scale Recovery for Monocular Visual Odometry Using Depth Estimated with Deep Convolutional Neural Fields. 5871-5879

Machine Learning Oral Session 9

Jianlong Chang, Lingfeng Wang, Gaofeng Meng, Shiming Xiang, Chunhong Pan:
Deep Adaptive Image Clustering. 5880-5888

Jen-Hao Rick Chang, Chun-Liang Li, Barnabás Póczos, B. V. K. Vijaya Kumar:
One Network to Solve Them All - Solving Linear Inverse Problems Using Deep Projection Models. 5889-5898

Mehdi Noroozi, Hamed Pirsiavash, Paolo Favaro:
Representation Learning by Learning to Count. 5899-5907

Han Zhang, Tao Xu, Hongsheng Li:
StackGAN: Text to Photo-Realistic Image Synthesis with Stacked Generative Adversarial Networks. 5908-5916

Kihyuk Sohn, Sifei Liu, Guangyu Zhong, Xiang Yu, Ming-Hsuan Yang, Manmohan Chandraker:
Unsupervised Domain Adaptation for Face Recognition in Unlabeled Videos. 5917-5925

dblp was originally created in 1993 at:

UNIVERSITÄT TRIER

since 2018, dblp has been operated and maintained by:

SCHLOSS DAGSTUHL
Leibniz-Zentrum für Informatik

the dblp computer science bibliography is funded and supported by:

Federal Ministry of Education and Research

RheinlandPfalz

SAARLAND

Leibniz Association

DFG Deutsche Forschungsgemeinschaft

nfdi
Nationale Forschungsdaten Infrastruktur