[+] [–] ICCV 2019: Seoul, South Korea

2019 IEEE/CVF International Conference on Computer Vision, ICCV 2019, Seoul, Korea (South), October 27 - November 2, 2019. IEEE 2019, ISBN 978-1-7281-4803-8

Poster 1.1

Deep Learning

Rizard Renanda Adhi Pramono, Yie-Tarng Chen, Wen-Hsien Fang:
Hierarchical Self-Attention Network for Action Localization in Videos. 61-70

Umar Riaz Muhammad, Yongxin Yang, Timothy M. Hospedales, Tao Xiang, Yi-Zhe Song:
Goal-Driven Sequential Data Abstraction. 71-80

Roberto Annunziata, Christos Sagonas, Jacques Calì:
Jointly Aligning Millions of Images With Deep Penalised Reconstruction Congealing. 81-90

Seungmin Lee, Dongwan Kim, Namil Kim, Seong-Gyun Jeong:
Drop to Adapt: Learning Discriminative Features for Unsupervised Domain Adaptation. 91-100

Youngdong Kim, Junho Yim, Juseung Yun, Junmo Kim:
NLNL: Negative Learning for Noisy Labels. 101-110

Shaokai Ye, Xue Lin, Kaidi Xu, Sijia Liu, Hao Cheng, Jan-Henrik Lambrechts, Huan Zhang, Aojun Zhou, Kaisheng Ma, Yanzhi Wang:
Adversarial Robustness vs. Model Compression, or Both? 111-120

Pu Zhao, Sijia Liu, Pin-Yu Chen, Nghia Hoang, Kaidi Xu, Bhavya Kailkhura, Xue Lin:
On the Design of Black-Box Adversarial Examples by Leveraging Gradient-Free Optimization and Operator Splitting Method. 121-130

Sagnik Das, Ke Ma, Zhixin Shu, Dimitris Samaras, Roy Shilkrot:
DewarpNet: Single-Image Document Unwarping With Stacked 3D and 2D Regression Networks. 131-140

Xu Zou, Sheng Zhong, Luxin Yan, Xiangyun Zhao, Jiahuan Zhou, Ying Wu:
Learning Robust Facial Landmark Detection via Hierarchical Structured Ensemble. 141-150

Zitong Yu, Wei Peng, Xiaobai Li, Xiaopeng Hong, Guoying Zhao:
Remote Heart Rate Measurement From Highly Compressed Facial Videos: An End-to-End Deep Learning Solution With Video Enhancement. 151-160

Tianyang Shi, Yi Yuan, Changjie Fan, Zhengxia Zou, Zhenwei Shi, Yong Liu:
Face-to-Parameter Translation for Game Character Auto-Creation. 161-170

Guha Balakrishnan, Adrian V. Dalca, Amy Zhao, John V. Guttag, Frédo Durand, William T. Freeman:
Visual Deprojection: Probabilistic Recovery of Collapsed Dimensions. 171-180

Yurui Ren, Xiaoming Yu, Ruonan Zhang, Thomas H. Li, Shan Liu, Ge Li:
StructureFlow: Image Inpainting via Structure-Aware Appearance Flow. 181-190

Md Mahfuzur Rahman Siddiquee, Zongwei Zhou, Nima Tajbakhsh, Ruibin Feng, Michael B. Gotway, Yoshua Bengio, Jianming Liang:
Learning Fixed Points in Generative Adversarial Networks: From Image-to-Image Translation to Disease Detection and Localization. 191-200

Zhengxia Zou, Wenyuan Li, Tianyang Shi, Zhenwei Shi, Jieping Ye:
Generative Adversarial Training for Weakly Supervised Cloud Matting. 201-210

Zheng Tang, Milind Naphade, Stan Birchfield, Jonathan Tremblay, William Hodge, Ratnesh Kumar, Shuo Wang, Xiaodong Yang:
PAMTRI: Pose-Aware Multi-Task Learning for Vehicle Re-Identification Using Highly Randomized Synthetic Data. 211-220

Eirikur Agustsson, Michael Tschannen, Fabian Mentzer, Radu Timofte, Luc Van Gool:
Generative Adversarial Networks for Extreme Learned Image Compression. 221-231

Yanbei Chen, Xiatian Zhu, Shaogang Gong:
Instance-Guided Context Rendering for Cross-Domain Person Re-Identification. 232-242

Mahmoud Afifi, Michael S. Brown:
What Else Can Fool Deep Learning? Addressing Color Constancy Errors on Deep Neural Network Performance. 243-252

Patrick Ebel, Eduard Trulls, Kwang Moo Yi, Pascal Fua, Anastasiia Mishchuk:
Beyond Cartesian Representations for Local Descriptors. 253-262

Muhamad Risqi Utama Saputra, Pedro Porto Buarque de Gusmão, Yasin Almalioglu, Andrew Markham, Niki Trigoni:
Distilling Knowledge From a Deep Pose Regressor Network. 263-272

Kyung-Rae Kim, Whan Choi, Yeong Jun Koh, Seong-Gyun Jeong, Chang-Su Kim:
Instance-Level Future Motion Estimation in a Single Image Based on Ordinal Regression. 273-282

Hang Zhou, Ziwei Liu, Xudong Xu, Ping Luo, Xiaogang Wang:
Vision-Infused Deep Audio Inpainting. 283-292

Zhen Dong, Zhewei Yao, Amir Gholami, Michael W. Mahoney, Kurt Keutzer:
HAWQ: Hessian AWare Quantization of Neural Networks With Mixed-Precision. 293-302

Jun-Ho Choi, Huan Zhang, Jun-Hyuk Kim, Cho-Jui Hsieh, Jong-Seok Lee:
Evaluating Robustness of Deep Image Super-Resolution Against Adversarial Attacks. 303-311

Kibok Lee, Kimin Lee, Jinwoo Shin, Honglak Lee:
Overcoming Catastrophic Forgetting With Unlabeled Data in the Wild. 312-321

Yisen Wang, Xingjun Ma, Zaiyi Chen, Yuan Luo, Jinfeng Yi, James Bailey:
Symmetric Cross Entropy for Robust Learning With Noisy Labels. 322-330

Avinash Ravichandran, Rahul Bhotika, Stefano Soatto:
Few-Shot Learning With Embedded Class Models and Shot-Free Meta Training. 331-339

Maneet Singh, Shruti Nagpal, Richa Singh, Mayank Vatsa:
Dual Directed Capsule Network for Very Low Resolution Image

Recognition. 340-349

Xiangyun Zhao, Yi Yang, Feng Zhou, Xiao Tan, Yuchen Yuan, Yingze Bao, Ying Wu:
Recognizing Part Attributes With Insufficient Data. 350-360

Jiaxin Li, Gim Hee Lee:
USIP: Unsupervised Stable Interest Point Detection From 3D Point Clouds. 361-370

Binghui Chen, Weihong Deng, Jiani Hu:
Mixed High-Order Attention Network for Person Re-Identification. 371-381

Rodrigo Ferreira Berriel, Stéphane Lathuilière, Moin Nabi, Tassilo Klein, Thiago Oliveira-Santos, Nicu Sebe, Elisa Ricci:
Budget-Aware Adapters for Multi-Domain Learning. 382-391

Tuong Do, Huy Tran, Thanh-Toan Do, Erman Tjiputra, Quang D. Tran:
Compact Trilinear Interaction for Visual Question Answering. 392-401

Ishan Nigam, Pavel Tokmakov, Deva Ramanan:
Towards Latent Attribute Discovery From Triplet Similarities. 402-410

Utkarsh Mall, Kevin Matzen, Bharath Hariharan, Noah Snavely, Kavita Bala:
GeoStyle: Discovering Fashion Trends and Events. 411-420

Haichao Zhang, Jianyu Wang:
Towards Adversarially Robust Object Detection. 421-430

Recognition

Junli Zhao, Xin Qi, Chengfeng Wen, Na Lei, Xianfeng Gu:
Automatic and Robust Skull Registration Based on Discrete Uniformization. 431-440

Zhimao Peng, Zechao Li, Junge Zhang, Yan Li, Guo-Jun Qi, Jinhui Tang:
Few-Shot Image Recognition With Knowledge Transfer. 441-449

Michael Wray, Gabriela Csurka, Diane Larlus, Dima Damen:
Fine-Grained Action Retrieval Through Multiple Parts-of-Speech Embeddings. 450-459

Peng Wang, Bingliang Jiao, Lu Yang, Yifei Yang, Shizhou Zhang, Wei Wei, Yanning Zhang:
Vehicle Re-Identification in Aerial Imagery: Dataset and Approach. 460-469

Krishna Regmi, Mubarak Shah:
Bridging the Domain Gap for Ground-to-Aerial Image Matching. 470-479

Mehran Khodabandeh, Arash Vahdat, Mani Ranjbar, William G. Macready:
A Robust Learning Approach to Domain Adaptive Object Detection. 480-490

Yin Bi, Aaron Chadha, Alhabib Abbas, Eirina Bourtsoulatze, Yiannis Andreopoulos:
Graph-Based Object Classification for Neuromorphic Vision Sensing. 491-501

Jiwoong Choi, Dayoung Chun, Hyun Kim, Hyuk-Jae Lee:
Gaussian YOLOv3: An Accurate and Fast Object Detector Using Localization

Uncertainty for Autonomous Driving. 502-511

Lezi Wang, Ziyan Wu, Srikrishna Karanam, Kuan-Chuan Peng, Rajat Vikram Singh, Bo Liu, Dimitris N. Metaxas:
Sharpen Focus: Learning With Attention Separability and Consistency. 512-521

Tianshui Chen, Muxin Xu, Xiaolu Hui, Hefeng Wu, Liang Lin:
Learning Semantic-Specific Graph Representation for Multi-Label Image Recognition. 522-531

Sergey Zakharov, Wadim Kehl, Slobodan Ilic:
DeceptionNet: Network-Driven Domain Randomization. 532-541

Jiaxu Miao, Yu Wu, Ping Liu, Yuhang Ding, Yi Yang:
Pose-Guided Feature Alignment for Occluded Person Re-Identification. 542-551

Tianyuan Yu, Da Li, Yongxin Yang, Timothy M. Hospedales, Tao Xiang:
Robust Person Re-Identification by Modelling Feature Uncertainty. 552-561

Arulkumar Subramaniam, Athira M. Nambiar, Anurag Mittal:
Co-Segmentation Inspired Attention Networks for Video-Based Person Re-Identification. 562-572

Huizi Mao, Xiaodong Yang, Bill Dally:
A Delay Metric for Video Object Detection: What Average Precision Fails to Tell. 573-582

Eden Belouadah, Adrian Popescu:
IL2M: Class Incremental Learning With Dual Memory. 583-592

Segmentation, Grouping, & Shape

Zhen Zhu, Mengdu Xu, Song Bai, Tengteng Huang, Xiang Bai:
Asymmetric Non-Local Neural Networks for Semantic Segmentation. 593-602

Zilong Huang, Xinggang Wang, Lichao Huang, Chang Huang, Yunchao Wei, Wenyu Liu:
CCNet: Criss-Cross Attention for Semantic Segmentation. 603-612

Shousheng Luo, Xue-Cheng Tai, Limei Huo, Yang Wang, Roland Glowinski:
Convex Shape Prior for Multi-Object Segmentation Using a Single Level Set Function. 613-621

Khoi Nguyen, Sinisa Todorovic:
Feature Weighting and Boosting for Few-Shot Segmentation. 622-631

Niv Haim, Nimrod Segol, Heli Ben-Hamu, Haggai Maron, Yaron Lipman:
Surface Networks via General Covers. 632-641

Naiyu Gao, Yanhu Shan, Yupei Wang, Xin Zhao, Yinan Yu, Ming Yang, Kaiqi Huang:
SSAP: Single-Shot Instance Segmentation With Affinity Pyramid. 642-651

Sifei Liu, Xueting Li, Varun Jampani, Shalini De Mello, Jan Kautz:
Learning Propagation for Arbitrarily-Structured Data. 652-661

Jun Hao Liew, Scott Cohen, Brian L. Price, Long Mai, Sim Heng Ong, Jiashi Feng:
MultiSeg: Semantically Meaningful, Scale-Diverse Segmentations From Minimal User Input. 662-670

Federica Arrigoni, Tomás Pajdla:
Robust Motion Segmentation From Pairwise Matches. 671-681

Haoshu Fang, Jianhua Sun, Runzhong Wang, Minghao Gou, Yong-Lu Li, Cewu Lu:
InstaBoost: Boosting Instance Segmentation via Probability Map Guided Copy-Pasting. 682-691

Face & Body

Mei Wang, Weihong Deng, Jiani Hu, Xunqiang Tao, Yaohai Huang:
Racial Faces in the Wild: Reducing Racial Bias by Information Maximization Adaptation Network. 692-702

Jingxiao Zheng, Ruichi Yu, Jun-Cheng Chen, Boyu Lu, Carlos Domingo Castillo, Rama Chellappa:
Uncertainty Modeling of Contextual-Connections Between Tracklets for Unconstrained Video-Based Face Recognition. 703-712

Xingxuan Zhang, Feng Cheng, Shilin Wang:
Spatio-Temporal Fusion Based Convolutional Sequence Learning for Lip Reading. 713-722

Yu Cheng, Bo Yang, Bo Wang, Yan Wending, Robby T. Tan:
Occlusion-Aware Networks for 3D Human Pose Estimation in Video. 723-732

Yong Zhang, Haiyong Jiang, Baoyuan Wu, Yanbo Fan, Qiang Ji:
Context-Aware Feature and Label Fusion for Facial Action Unit Intensity Estimation With Partially Labeled Data. 733-742

Chaoyang Wang, Chen Kong, Simon Lucey:
Distill Knowledge From NRSfM for Weakly Supervised 3D Pose Learning. 743-752

Yuan Yao, Yasamin Jafarian, Hyun Soo Park:
MONET: Multiview Semi-Supervised Keypoint Detection via Epipolar Divergence. 753-762

Gilwoo Lee, Zhiwei Deng, Shugao Ma, Takaaki Shiratori, Siddhartha S. Srinivasa, Yaser Sheikh:
Talking With Hands 16.2M: A Large-Scale Dataset of Synchronized Body-Finger Motion and Audio for Conversational Motion Analysis and Synthesis. 763-772

Lingxue Song, Dihong Gong, Zhifeng Li, Changsong Liu, Wei Liu:
Occlusion Robust Face Recognition Based on Mask Learning With Pairwise Differential Siamese Network. 773-782

Xuanyi Dong, Yi Yang:
Teacher Supervises Students How to Learn From Partially Labeled Images for Facial Landmark Detection. 783-792

Fu Xiong, Boshen Zhang, Yang Xiao, Zhiguo Cao, Taidong Yu, Joey Tianyi Zhou, Junsong Yuan:
A2J: Anchor-to-Joint Regression Network for 3D Articulated Pose Estimation From a Single Depth Image. 793-802

Georgios Pavlakos, Nikos Kolotouros, Kostas Daniilidis:
TexturePose: Supervising Human Mesh Estimation With Texture Consistency. 803-812

Christian Zimmermann, Duygu Ceylan, Jimei Yang, Bryan C. Russell, Max J. Argus, Thomas Brox:
FreiHAND: A Dataset for Markerless Capture of Hand Pose and Shape From Single RGB Images. 813-822

Nitin Saini, Eric Price, Rahul Tallamraju, Raffi Enficiaud, Roman Ludwig, Igor Martinovic, Aamir Ahmad, Michael J. Black:
Markerless Outdoor Human Motion Capture Using Multiple Autonomous Micro Aerial Vehicles. 823-832

Action & Video

Srijan Das, Rui Dai, Michal Koperski, Luca Minciullo, Lorenzo Garattoni, François Brémond, Gianpiero Francesca:
Toyota Smarthome: Real-World Activities of Daily Living. 833-842

Penghao Zhou, Mingmin Chi:
Relation Parsing Neural Network for Human-Object Interaction Detection. 843-851

Rohit Girdhar, Du Tran, Lorenzo Torresani, Deva Ramanan:
DistInit: Learning Video Representations Without a Single Labeled Video. 852-861

Fadime Sener, Angela Yao:
Zero-Shot Anticipation for Instructional Activities. 862-871

Tianhong Li, Lijie Fan, Mingmin Zhao, Yingcheng Liu, Dina Katabi:
Making the Invisible Visible: Action Recognition Through Walls and Occlusions. 872-881

Xudong Xu, Bo Dai, Dahua Lin:
Recursive Visual Sound Separation Using Minus-Plus Net. 882-891

Motion & Tracking

Fitsum A. Reda, Deqing Sun, Aysegul Dundar, Mohammad Shoeybi, Guilin Liu, Kevin J. Shih, Andrew Tao, Jan Kautz, Bryan Catanzaro:
Unsupervised Video Interpolation Using Cycle Consistency. 892-900

Tao Wang, Haibin Ling, Congyan Lang, Songhe Feng, Xiaohui Hou:
Deformable Surface Tracking by Graph Matching. 901-910

Janghoon Choi, Junseok Kwon, Kyoung Mu Lee:
Deep Meta Learning for Real-Time Target-Aware Visual Tracking. 911-920

Chiho Choi, Behzad Dariush:
Looking to Relations for Future Trajectory Forecast. 921-930

Zhao Yang, Qiang Wang, Luca Bertinetto, Song Bai, Weiming Hu, Philip H. S. Torr:
Anchor Diffusion for Unsupervised Video Object Segmentation. 931-940

Philipp Bergmann, Tim Meinhardt, Laura Leal-Taixé:
Tracking Without Bells and Whistles. 941-951

Scene Understanding

Zhaoyi Yan, Yuchen Yuan, Wangmeng Zuo, Xiao Tan, Yezhen Wang, Shilei Wen, Errui Ding:
Perspective-Guided Convolution Networks for Crowd Counting. 952-961

Yichao Zhou, Haozhi Qi, Yi Ma:
End-to-End Wireframe Parsing. 962-971

Yoshikatsu Nakajima, Byeongkeun Kang, Hideo Saito, Kris Kitani:
Incremental Class Discovery for Semantic Segmentation With RGBD Sensing. 972-981

Liang Du, Jingang Tan, Hongye Yang, Jianfeng Feng, Xiangyang Xue, Qibao Zheng, Xiaoqing Ye, Xiaolin Zhang:
SSF-DAN: Separated Semantic Feature Based Domain Adaptation Network for Semantic Segmentation. 982-991

Nicholas Weir, David Lindenbaum, Alexei Bastidas, Adam Van Etten, Varun Kumar Vijay, Sean McPherson, Jacob Shermeyer, Hanlin Tang:
SpaceNet MVOI: A Multi-View Overhead Imagery Dataset. 992-1001

Vishwanath Sindagi, Vishal M. Patel:
Multi-Level Bottom-Top and Top-Bottom Feature Fusion for Crowd Counting. 1002-1012

Yuenan Hou, Zheng Ma, Chunxiao Liu, Chen Change Loy:
Learning Lightweight Lane Detection CNNs by Self Attention Distillation. 1013-1021

Daniel Gordon, Abhishek Kadian, Devi Parikh, Judy Hoffman, Dhruv Batra:
SplitNet: Sim2Sim and Task2Task Transfer for Embodied Visual Navigation. 1022-1031

3D From Multiview & Sensors

Wentao Cheng, Weisi Lin, Kan Chen, Xinfeng Zhang:
Cascaded Parallel Filtering for Memory-Efficient Image-Based Localization. 1032-1041

Chao Wen, Yinda Zhang, Zhuwen Li, Yanwei Fu:
Pixel2Mesh++: Multi-View 3D Mesh Generation via Deformation. 1042-1051

Fotios Logothetis, Roberto Mecca, Roberto Cipolla:
A Differential Volumetric Approach to Multi-View Photometric Stereo. 1052-1061

Viktor Larsson, Torsten Sattler, Zuzana Kukelova, Marc Pollefeys:
Revisiting Radial Distortion Absolute Pose. 1062-1071

Tobias Würfl, André Aichert, Nicole Maass, Frank Dennerlein, Andreas K. Maier:
Estimating the Fundamental Matrix Without Point Correspondences With Application to Transmission Imaging. 1072-1081

Devesh Adlakha, Adlane Habed, Fabio Morbidi, Cédric Demonceaux, Michel de Mathelin:
QUARCH: A New Quasi-Affine Reconstruction Stratum From Vague Relative Camera Orientation Knowledge. 1082-1090

Dániel Baráth, Zuzana Kukelova:
Homography From Two Orientation- and Scale-Covariant Features. 1091-1099

Applications. Medical, & Robotics

Hyukryul Yang, Hao Ouyang, Vladlen Koltun, Qifeng Chen:
Hiding Video in Audio via Reversible Generative Models. 1100-1109

Yong Zhao, Shibiao Xu, Shuhui Bu, Hongkai Jiang, Pengcheng Han:
GSLAM: A General SLAM Framework and Benchmark. 1110-1120

Sang Jun Lee, Sung Soo Hwang:
Elaborate Monocular Point and Line SLAM With Robust Initialization. 1121-1129

Jia Wan, Antoni B. Chan:
Adaptive Density Map Generation for Crowd Counting. 1130-1139

Xingxu Yao, Dongyu She, Sicheng Zhao, Jie Liang, Yu-Kun Lai, Jufeng Yang:
Attention-Aware Polarity Sensitive Embedding for Affective Image Retrieval. 1140-1150

Chi Zhan, Dongyu She, Sicheng Zhao, Ming-Ming Cheng, Jufeng Yang:
Zero-Shot Emotion Recognition via Affective Structural Embedding. 1151-1160

Haoye Dong, Xiaodan Liang, Xiaohui Shen, Bowen Wu, Bing-Cheng Chen, Jian Yin:
FW-GAN: Flow-Navigated Warping GAN for Video Virtual Try-On. 1161-1170

Arnab Ghosh, Richard Zhang, Puneet K. Dokania, Oliver Wang, Alexei A. Efros, Philip H. S. Torr, Eli Shechtman:
Interactive Sketch & Fill: Multiclass Sketch-to-Image Translation. 1171-1180

Shi Chen, Qi Zhao:
Attention-Based Autism Spectrum Disorder Screening With Privileged Modality. 1181-1190

Jun-Tae Lee, Chang-Su Kim:
Image Aesthetic Assessment Based on Pairwise Comparison A Unified Approach to Score Regression, Binary Classification, and Personalization. 1191-1200

Zhenyu Wu, Karthik Suresh, Priya Narayanan, Hongyu Xu, Heesung Kwon, Zhangyang Wang:
Delving Into Robust Object Detection From Unmanned Aerial Vehicles: A Deep Nuisance Disentanglement Approach. 1201-1210

Adnan Siraj Rakin, Zhezhi He, Deliang Fan:
Bit-Flip Attack: Crushing Neural Network With Progressive Bit Search. 1211-1220

Vishwanath Sindagi, Rajeev Yasarla, Vishal M. Patel:
Pushing the Frontiers of Unconstrained Crowd Counting: New Dataset and Benchmark Method. 1221-1231

Yi Liu, Qiang Zhang, Dingwen Zhang, Jungong Han:
Employing Deep Part-Object Relationships for Salient Object Detection. 1232-1241

Vladimiros Sterzentsenko, Leonidas Saroglou, Anargyros Chatzitofis, Spiros Thermos, Nikolaos Zioulis, Alexandros Doumanoglou, Dimitrios Zarpalas, Petros Daras:
Self-Supervised Deep Depth Denoising. 1242-1251

Hanxiao Wang, Venkatesh Saligrama, Stan Sclaroff, Vitaly Ablavsky:
Cost-Aware Fine-Grained Recognition for IoTs Based on Sequential Fixations. 1252-1261

Ruichi Yu, Hongcheng Wang, Ang Li, Jingxiao Zheng, Vlad I. Morariu, Larry Davis:
Layout-Induced Video Representation for Recognizing Agent-in-Place Actions. 1262-1272

Trong-Nguyen Nguyen, Jean Meunier:
Anomaly Detection in Video Sequence With Appearance-Motion Correspondence. 1273-1283

Oral 1.2A

Architectures, Multi-Task Learning, Domain Adaptation

Saining Xie, Alexander Kirillov, Ross B. Girshick, Kaiming He:
Exploring Randomly Wired Neural Networks for Image Recognition. 1284-1293

Xin Chen, Lingxi Xie, Jun Wu, Qi Tian:
Progressive Differentiable Architecture Search: Bridging the Depth Gap Between Search and Evaluation. 1294-1303

Xiawu Zheng, Rongrong Ji, Lang Tang, Baochang Zhang, Jianzhuang Liu, Qi Tian:
Multinomial Distribution Learning for Effective Neural Architecture Search. 1304-1313

Andrew Howard, Ruoming Pang, Hartwig Adam, Quoc V. Le, Mark Sandler, Bo Chen, Weijun Wang, Liang-Chieh Chen, Mingxing Tan, Grace Chu, Vijay Vasudevan, Yukun Zhu:
Searching for MobileNetV3. 1314-1324

Markus Nagel, Mart van Baalen, Tijmen Blankevoort, Max Welling:
Data-Free Quantization Through Weight Equalization and Bias Correction. 1325-1334

Laurie Bose, Piotr Dudek, Jianing Chen, Stephen J. Carey, Walterio W. Mayol-Cuevas:
A Camera That CNNs: Towards Embedded Neural Networks on Pixel Processor Arrays. 1335-1344

Xiao Jin, Baoyun Peng, Yichao Wu, Yu Liu, Jiaheng Liu, Ding Liang, Junjie Yan, Xiaolin Hu:
Knowledge Distillation via Route Constrained Optimization. 1345-1354

Mary Phuong, Christoph Lampert:
Distillation-Based Training for Multi-Exit Architectures. 1355-1364

Frederick Tung, Greg Mori:
Similarity-Preserving Knowledge Distillation. 1365-1374

Gjorgji Strezoski, Nanne van Noord, Marcel Worring:
Many Task Learning With Task Routing. 1375-1384

Felix J. S. Bragman, Ryutaro Tanno, Sébastien Ourselin, Daniel C. Alexander, Manuel Jorge Cardoso:
Stochastic Filter Groups for Multi-Task CNNs: Learning Specialist and Generalist Convolution Kernels. 1385-1394

Anh Tuan Tran, Cuong V. Nguyen, Tal Hassner:
Transferability and Hardness of Supervised Classification Tasks. 1395-1405

Xingchao Peng, Qinxun Bai, Xide Xia, Zijun Huang, Kate Saenko, Bo Wang:
Moment Matching for Multi-Source Domain Adaptation. 1406-1415

Safa Cicek, Stefano Soatto:
Unsupervised Domain Adaptation via Regularized Conditional Alignment. 1416-1425

Ruijia Xu, Guanbin Li, Jihan Yang, Liang Lin:
Larger Norm More Transferable: An Adaptive Feature Norm Approach for Unsupervised Domain Adaptation. 1426-1435

Jogendra Nath Kundu, Nishank Lakkakula, Venkatesh Babu Radhakrishnan:
UM-Adapt: Unsupervised Multi-Task Adaptation Using Adversarial Cross-Task Distillation. 1436-1445

Da Li, Jianshu Zhang, Yongxin Yang, Cong Liu, Yi-Zhe Song, Timothy M. Hospedales:
Episodic Training for Domain Generalization. 1446-1455

Yi-Hsuan Tsai, Kihyuk Sohn, Samuel Schulter, Manmohan Chandraker:
Domain Adaptation for Structured Output via Discriminative Patch Representations. 1456-1465

Qin Wang, Wen Li, Luc Van Gool:
Semi-Supervised Learning by Augmented Distribution Alignment. 1466-1475

Lucas Beyer, Xiaohua Zhai, Avital Oliver, Alexander Kolesnikov:
S4L: Self-Supervised Semi-Supervised Learning. 1476-1485

Oral 1.2B

Multi-View Geometry, 3D Scene Understanding

Pablo Speciale, Johannes L. Schönberger, Sudipta N. Sinha, Marc Pollefeys:
Privacy Preserving Image Queries for Camera Localization. 1486-1496

Songyou Peng, Peter F. Sturm:
Calibration Wizard: A Guidance System for Camera Calibration Based on Modelling Geometric and Corner Uncertainty. 1497-1505

Tobias Gruber, Frank D. Julca-Aguilar, Mario Bijelic, Felix Heide:
Gated2Depth: Real-Time Dense Lidar From Gated Images. 1506-1516

Andrea Nicastro, Ronald Clark, Stefan Leutenegger:
X-Section: Cross-Section Prediction for Enhanced RGB-D Fusion. 1517-1526

Stepan Tulyakov, François Fleuret, Martin Kiefel, Peter V. Gehler, Michael Hirsch:
Learning an Event Sequence Embedding for Dense Event-Based Deep Stereo. 1527-1537

Rui Chen, Songfang Han, Jing Xu, Hao Su:
Point-Based Multi-View Stereo Network. 1538-1547

Xiangyu Xu, Enrique Dunn:
Discrete Laplace Operator Estimation for Dynamic 3D Reconstruction. 1548-1557

Chen Kong, Simon Lucey:
Deep Non-Rigid Structure From Motion. 1558-1567

Carlos Esteves, Yinshuang Xu, Christine Allen-Blanchette, Kostas Daniilidis:
Equivariant Multi-View Networks. 1568-1577

Jiageng Mao, Xiaogang Wang, Hongsheng Li:
Interpolated Convolutional Networks for 3D Point Cloud Understanding. 1578-1587

Mikaela Angelina Uy, Quang-Hieu Pham, Binh-Son Hua, Duc Thanh Nguyen, Sai-Kit Yeung:
Revisiting Point Cloud Classification: A New Benchmark Dataset and Classification Model on Real-World Data. 1588-1597

Tianhang Zheng, Changyou Chen, Junsong Yuan, Bo Li, Kui Ren:
PointCloud Saliency Maps. 1598-1606

Zhiyuan Zhang, Binh-Son Hua, Sai-Kit Yeung:
ShellNet: Efficient Point Cloud Convolutional Neural Networks Using Concentric Shells Statistics. 1607-1616

Jean-Michel Roufosse, Abhishek Sharma, Maks Ovsjanikov:
Unsupervised Deep Learning for Structured Shape Matching. 1617-1627

Nadav Dym, Shahar Z. Kovalsky:
Linearly Converging Quasi Branch and Bound Algorithms for Global Rigid Registration. 1628-1636

Zhipeng Cai, Tat-Jun Chin, Vladlen Koltun:
Consensus Maximization Tree Search Revisited. 1637-1645

Haoang Li, Ji Zhao, Jean-Charles Bazin, Wen Chen, Zhe Liu, Yunhui Liu:
Quasi-Globally Optimal and Efficient Vanishing Point Estimation in Manhattan World. 1646-1654

Yaqing Ding, Jian Yang, Jean Ponce, Hui Kong:
An Efficient Solution to the Homography-Based Relative Pose Problem With a Common Reference Direction. 1655-1664

Heng Yang, Luca Carlone:
A Quaternion-Based Certifiably Optimal Solution to the Wahba Problem With Outliers. 1665-1674

Timothy Duff, Kathlén Kohn, Anton Leykin, Tomás Pajdla:
PLMP - Point-Line Minimal Problems in Complete Multi-View Visibility. 1675-1684

Poster 1.2

Deep Learning

Jian Zhang, Chenglong Zhao, Bingbing Ni, Minghao Xu, Xiaokang Yang:
Variational Few-Shot Learning. 1685-1694

Sankha Subhra Mullick, Shounak Datta, Swagatam Das:
Generative Adversarial Minority Oversampling. 1695-1704

Dong Gong, Lingqiao Liu, Vuong Le, Budhaditya Saha, Moussa Reda Mansour, Svetha Venkatesh, Anton van den Hengel:
Memorizing Normality to Detect Anomaly: Memory-Augmented Deep Autoencoder for Unsupervised Anomaly Detection. 1705-1714

Zuoyue Li, Jan Dirk Wegner, Aurélien Lucchi:
Topological Map Extraction From Overhead Images. 1715-1724

Haokui Zhang, Ying Li, Yuanzhouhan Cao, Yu Liu, Chunhua Shen, Youliang Yan:
Exploiting Temporal Consistency for Real-Time Video Depth Estimation. 1725-1734

Hang Zhao, Chuang Gan, Wei-Chiu Ma, Antonio Torralba:
The Sound of Motions. 1735-1744

Youngjoo Jo, Jongyoul Park:
SC-FEGAN: Face Editing Generative Adversarial Network With User's Sketch and Color. 1745-1753

Hongwei Ge, Zehang Yan, Kai Zhang, Mingde Zhao, Liang Sun:
Exploring Overall Contextual Information for Image Captioning in Human-Like Cognitive Style. 1754-1763

Zhuoyuan Chen, Kavya Srinet, Charles R. Qi, Haoqi Fan, Jerry Ma, Larry Zitnick, Demi Guo, Tong Xiao, Saining Xie, Xinlei Chen, Arthur Szlam, Shubham Tulsiani, Haonan Yu, Jonathan Gray:
Order-Aware Generative Modeling Using the 3D-Craft Dataset. 1764-1773

Lingbo Liu, Zhilin Qiu, Guanbin Li, Shufan Liu, Wanli Ouyang, Liang Lin:
Crowd Counting With Deep Structured Scale Integration Network. 1774-1783

Tomer Cohen, Lior Wolf:
Bidirectional One-Shot Unsupervised Domain Mapping. 1784-1792

A. J. Piergiovanni, Anelia Angelova, Alexander Toshev, Michael S. Ryoo:
Evolving Space-Time Neural Architectures for Videos. 1793-1802

Jiahui Yu, Thomas S. Huang:
Universally Slimmable Networks and Improved Training Techniques. 1803-1811

Tonmoy Saikia, Yassine Marrakchi, Arber Zela, Frank Hutter, Thomas Brox:
AutoDispNet: Improving Disparity Estimation With AutoML. 1812-1823

Gidi Littwin, Lior Wolf:
Deep Meta Functionals for Shape Representation. 1824-1833

Yu Liu, Jihao Liu, Xiaogang Wang, Ailing Zeng:
Differentiable Kernel Evolution. 1834-1843

Mikolaj Binkowski, R. Devon Hjelm, Aaron C. Courville:
Batch Weight for Domain Adaptation With Mass Shift. 1844-1853

HyunJae Lee, Hyo-Eun Kim, Hyeonseob Nam:
SRM: A Style-Based Recalibration Module for Convolutional Neural Networks. 1854-1862

Xingang Pan, Xiaohang Zhan, Jianping Shi, Xiaoou Tang, Ping Luo:
Switchable Whitening for Deep Representation Learning. 1863-1871

Adria Ruiz, Jakob Verbeek:
Adaptative Inference Cost With Convolutional Neural Mixture Models. 1872-1881

Ilija Radosavovic, Justin Johnson, Saining Xie, Wan-Yen Lo, Piotr Dollár:
On Network Design Spaces for Visual Recognition. 1882-1890

Hao Li, Hong Zhang, Xiaojuan Qi, Ruigang Yang, Gao Huang:
Improved Techniques for Training Adaptive Deep Networks. 1891-1900

Yunyang Xiong, Ronak Mehta, Vikas Singh:
Resource Constrained Neural Network Architecture Search: Will a Submodularity Assumption Help? 1901-1910

Xiaohan Ding, Yuchen Guo, Guiguang Ding, Jungong Han:
ACNet: Strengthening the Kernel Skeletons for Powerful CNN via Asymmetric Convolution Blocks. 1911-1920

Byeongho Heo, Jeesoo Kim, Sangdoo Yun, Hyojin Park, Nojun Kwak, Jin Young Choi:

A Comprehensive Overhaul of Feature Distillation. 1921-1930

Recognition

Yew Siang Tang, Gim Hee Lee:
Transferable Semi-Supervised 3D Object Detection From RGB-D Data. 1931-1940

Sergey Zakharov, Ivan Shugurov, Slobodan Ilic:
DPOD: 6D Pose Object Detector and Refiner. 1941-1950

Zetong Yang, Yanan Sun, Shu Liu, Xiaoyong Shen, Jiaya Jia:
STD: Sparse-to-Dense 3D Object Detector for Point Cloud. 1951-1960

Hang Zhou, Kejiang Chen, Weiming Zhang, Han Fang, Wenbo Zhou, Nenghai Yu:
DUP-Net: Denoiser and Upsampler Network for 3D Adversarial Point Clouds Defense. 1961-1970

Tiancai Wang, Rao Muhammad Anwer, Hisham Cholakkal, Fahad Shahbaz Khan, Yanwei Pang, Ling Shao:
Learning Rich Features at High-Speed for Single-Shot Object Detection. 1971-1980

Julia Peyre, Josef Sivic, Ivan Laptev, Cordelia Schmid:
Detecting Unseen Visual Relations Using Analogies. 1981-1990

Andrea Simonelli, Samuel Rota Bulò, Lorenzo Porzi, Manuel Lopez-Antequera, Peter Kontschieder:
Disentangling Monocular 3D Object Detection. 1991-1999

Boyuan Jiang, Mengmeng Wang, Weihao Gan, Wei Wu, Junjie Yan:
STM: SpatioTemporal and Motion Encoding for Action Recognition. 2000-2009

Shuaiyi Huang, Qiuyue Wang, Songyang Zhang, Shipeng Yan, Xuming He:
Dynamic Context Correspondence Network for Semantic Alignment. 2010-2019

Akshayvarun Subramanya, Vipin Pillai, Hamed Pirsiavash:
Fooling Network Interpretation in Image Classification. 2020-2029

Yinan Zhao, Brian L. Price, Scott Cohen, Danna Gurari:
Unconstrained Foreground Object Search. 2030-2039

Jianwei Yang, Zhile Ren, Mingze Xu, Xinlei Chen, David J. Crandall, Devi Parikh, Dhruv Batra:
Embodied Amodal Recognition: Learning to Move to Perceive Objects. 2040-2050

Kaiyu Yang, Olga Russakovsky, Jia Deng:
SpatialSense: An Adversarially Crowdsourced Benchmark for Spatial Relation Recognition. 2051-2060

Xinlei Chen, Ross B. Girshick, Kaiming He, Piotr Dollár:
TensorMask: A Foundation for Dense Object Segmentation. 2061-2069

Peng-Tao Jiang, Qibin Hou, Yang Cao, Ming-Ming Cheng, Yunchao Wei, Hongkai Xiong:
Integral Object Mining via Online Attention Accumulation. 2070-2079

Segmentation, Grouping, & Shape

Vladislav Golyanik, Christian Theobalt, Didier Stricker:
Accelerated Gravitational Point Set Alignment With Altered Physical Laws. 2080-2089

Minghao Chen, Hongyang Xue, Deng Cai:
Domain Adaptation for Semantic Segmentation With Maximum Squares Loss. 2090-2099

Xiangyu Yue, Yang Zhang, Sicheng Zhao, Alberto L. Sangiovanni-Vincentelli, Kurt Keutzer, Boqing Gong:
Domain Randomization and Pyramid Consistency: Simulation-to-Real Generalization Without Accessing Target Domain Data. 2100-2110

Yi He, Jiayuan Shi, Chuan Wang, Haibin Huang, Jiaming Liu, Guanbin Li, Risheng Liu, Jue Wang:
Semi-Supervised Skin Detection by Network With Mutual Guidance. 2111-2120

Zuxuan Wu, Xin Wang, Joseph Gonzalez, Tom Goldstein, Larry Davis:
ACE: Adapting to Changing Environments for Semantic Segmentation. 2121-2130

Dmitrii Marin, Zijian He, Peter Vajda, Priyam Chatterjee, Sam S. Tsai, Fei Yang, Yuri Boykov:
Efficient Segmentation: Learning Downsampling Near Semantic Boundaries. 2131-2141

Wei Wang, Kaicheng Yu, Joachim Hugonot, Pascal Fua, Mathieu Salzmann:
Recurrent U-Net for Resource-Constrained Segmentation. 2142-2151

Krzysztof Lis, Krishna Kanth Nakka, Pascal Fua, Mathieu Salzmann:
Detecting the Unexpected via Image Resynthesis. 2152-2161

3D From Single View & RGBD

Jamie Watson, Michael Firman, Gabriel J. Brostow, Daniyar Turmukhambetov:
Self-Supervised Monocular Depth Hints. 2162-2171

Daeyun Shin, Zhile Ren, Erik B. Sudderth, Charless C. Fowlkes:
3D Scene Reconstruction With Multi-Layer Depth and Epipolar Transformers. 2172-2182

Tom van Dijk, Guido de Croon:
How Do Neural Networks See Depth in Single Images? 2183-2191

Zhi Li, Xuan Wang, Fei Wang, Peilin Jiang:
On Boosting Single-Frame 3D Human Pose Estimation via Monocular Videos. 2192-2201

Nilesh Kulkarni, Shubham Tulsiani, Abhinav Gupta:
Canonical Surface Mapping via Geometric Cycle Consistency. 2202-2211

Nilesh Kulkarni, Ishan Misra, Shubham Tulsiani, Abhinav Gupta:
3D-RelNet: Joint Object and Relational Network for 3D Prediction. 2212-2221

Alexander Grabner, Peter M. Roth, Vincent Lepetit:
GP2C: Geometric Projection Parameter Consensus for Joint 3D Pose and Focal Length Estimation in the Wild. 2222-2231

Face & Body

Valentin Gabeur, Jean-Sébastien Franco, Xavier Martin, Cordelia Schmid, Grégory Rogez:
Moulding Humans: Non-Parametric 3D Human Shape Estimation From Single Images. 2232-2241

Albert Pumarola, Jordi Sanchez, Gary P. T. Choi, Alberto Sanfeliu, Francesc Moreno:
3DPeople: Modeling the Geometry of Dressed Humans. 2242-2251

Nikos Kolotouros, Georgios Pavlakos, Michael J. Black, Kostas Daniilidis:
Learning to Reconstruct 3D Human Pose and Shape via Model-Fitting in the Loop. 2252-2261

Hai Ci, Chunyu Wang, Xiaoxuan Ma, Yizhou Wang:
Optimizing Network Structure for 3D Human Pose Estimation. 2262-2271

Yujun Cai, Liuhao Ge, Jun Liu, Jianfei Cai, Tat-Jen Cham, Junsong Yuan, Nadia Magnenat-Thalmann:
Exploiting Spatial-Temporal Relationships for 3D Pose Estimation via Graph Convolutional Networks. 2272-2281

Mohamed Hassan, Vasileios Choutas, Dimitrios Tzionas, Michael J. Black:
Resolving 3D Human Pose Ambiguities With 3D Scene Constraints. 2282-2292

Thiemo Alldieck, Gerard Pons-Moll, Christian Theobalt, Marcus A. Magnor:
Tex2Shape: Detailed Full Human Body Geometry From a Single Image. 2293-2303

Shunsuke Saito, Zeng Huang, Ryota Natsume, Shigeo Morishima, Hao Li, Angjoo Kanazawa:
PIFu: Pixel-Aligned Implicit Function for High-Resolution Clothed Human Digitization. 2304-2314

Xiaoxing Zeng, Xiaojiang Peng, Yu Qiao:
DF2Net: A Dense-Fine-Finer Network for Detailed 3D Face Reconstruction. 2315-2324

Saurabh Sharma, Pavan Teja Varigonda, Prashast Bindal, Abhishek Sharma, Arjun Jain:
Monocular 3D Human Pose Estimation by Generation and Ordinal Ranking. 2325-2334

Linlin Yang, Shile Li, Dongheui Lee, Angela Yao:
Aligning Latent Spaces for 3D Hand Pose Estimation. 2335-2343

Kun Zhou, Xiaoguang Han, Nianjuan Jiang, Kui Jia, Jiangbo Lu:
HEMlets Pose: Learning Part-Centric Heatmap Triplets for Accurate 3D Human Pose Estimation. 2344-2353

Xiong Zhang, Qiang Li, Hong Mo, Wenbo Zhang, Wen Zheng:
End-to-End Hand Mesh Recovery From a Monocular RGB Image. 2354-2364

Motion & Tracking

Wenwei Zhang, Hui Zhou, Shuyang Sun, Zhe Wang, Jianping Shi, Chen Change Loy:
Robust Multi-Modality Multi-Object Tracking. 2365-2374

Boris Ivanovic, Marco Pavone:
The Trajectron: Probabilistic Multi-Agent Trajectory Modeling With Dynamic Spatiotemporal Graphs. 2375-2384

Bin Yan, Haojie Zhao, Dong Wang, Huchuan Lu, Xiaoyun Yang:
'Skimming-Perusal' Tracking: A Framework for Real-Time and Robust Long-Term Tracking. 2385-2393

Kyle Min, Jason J. Corso:
TASED-Net: Temporally-Aggregating Spatial Encoder-Decoder Network for Video Saliency Detection. 2394-2403

Anurag Ranjan, Joel Janai, Andreas Geiger, Michael J. Black:
Attacking Optical Flow. 2404-2413

Computational Photography & Graphics

Chunyu Li, Yusuke Monno, Hironori Hidaka, Masatoshi Okutomi:
Pro-Cam SSfM: Projector-Camera System for Structure and Spectral Reflectance From Motion. 2414-2423

Bin He, Ce Wang, Boxin Shi, Lingyu Duan:
Mop Moiré Patterns Using MopNet. 2424-2432

Ruofan Zhou, Sabine Süsstrunk:
Kernel Modeling Super-Resolution on Real Low-Resolution Images. 2433-2443

Daiqian Ma, Renjie Wan, Boxin Shi, Alex C. Kot, Lingyu Duan:
Learning to Jointly Generate and Separate Reflections. 2444-2452

Zijun Deng, Lei Zhu, Xiaowei Hu, Chi-Wing Fu, Xuemiao Xu, Qing Zhang, Jing Qin, Pheng-Ann Heng:
Deep Multi-Model Fusion for Single-Image Dehazing. 2453-2462

Yuhui Quan, Shijie Deng, Yixin Chen, Hui Ji:
Deep Learning for Seeing Through Window With Raindrops. 2463-2471

Xiaowei Hu, Yitong Jiang, Chi-Wing Fu, Pheng-Ann Heng:
Mask-ShadowGAN: Learning to Remove Shadows From Unpaired Data. 2472-2481

Low-Level Vision & Optimization

Shangchen Zhou, Jiawei Zhang, Jinshan Pan, Wangmeng Zuo, Haozhe Xie, Jimmy S. J. Ren:
Spatio-Temporal Filter Adaptive Network for Video Deblurring. 2482-2491

Yang Liu, Jinshan Pan, Jimmy S. J. Ren, Zhixun Su:
Learning Deep Priors for Image Dehazing. 2492-2500

Xueyang Fu, Zheng-Jun Zha, Feng Wu, Xinghao Ding, John W. Paisley:
JPEG Artifacts Reduction via Deep Convolutional Sparse Coding. 2501-2510

Shuhang Gu, Yawei Li, Luc Van Gool, Radu Timofte:
Self-Guided Network for Fast Image Denoising. 2511-2520

Ziang Cheng, Yinqiang Zheng, Shaodi You, Imari Sato:
Non-Local Intrinsic Decomposition With Near-Infrared Priors. 2521-2530

Scene Understanding

Romain Cohendet, Claire-Hélène Demarty, Ngoc Q. K. Duong, Martin Engilberge:
VideoMem: Constructing, Analyzing, Predicting Short-Term and Long-Term Video Memorability. 2531-2540

Maciej Halber, Yifei Shi, Kai Xu, Thomas A. Funkhouser:
Rescan: Inductive Instance Segmentation for Indoor RGBD Scans. 2541-2550

Armen Avetisyan, Angela Dai, Matthias Nießner:
End-to-End CAD Model Retrieval and 9DoF Alignment in 3D Scans. 2551-2560

Tianhao Yang, Zheng-Jun Zha, Hanwang Zhang:
Making History Matter: History-Advantage Sequence Training for Visual Dialog. 2561-2569

Liu Liu, Hongdong Li, Yuchao Dai:
Stochastic Attraction-Repulsion Embedding for Large Scale Image Localization. 2570-2579

Ranjay Krishna, Vincent S. Chen, Paroma Varma, Michael S. Bernstein, Christopher Ré, Li Fei-Fei:
Scene Graph Prediction With Limited Labels. 2580-2590

Language & Reasoning

Ramprasaath Ramasamy Selvaraju, Stefan Lee, Yilin Shen, Hongxia Jin, Shalini Ghosh, Larry P. Heck, Dhruv Batra, Devi Parikh:
Taking a HINT: Leveraging Explanations to Make Vision and Language Models More Grounded. 2591-2600

Samyak Datta, Karan Sikka, Anirban Roy, Karuna Ahuja, Devi Parikh, Ajay Divakaran:
Align2Ground: Weakly Supervised Phrase Grounding Guided by Image-Caption Alignment. 2601-2610

Xuejing Liu, Liang Li, Shuhui Wang, Zheng-Jun Zha, Dechao Meng, Qingming Huang:
Adaptive Reconstruction Network for Weakly Supervised Referring Expression Grounding. 2611-2620

Ting Yao, Yingwei Pan, Yehao Li, Tao Mei:
Hierarchy Parsing for Image Captioning. 2621-2629

Antoine Miech, Dimitri Zhukov, Jean-Baptiste Alayrac, Makarand Tapaswi, Ivan Laptev, Josef Sivic:
HowTo100M: Learning a Text-Video Embedding by Watching Hundred Million Narrated Video Clips. 2630-2640

Bairui Wang, Lin Ma, Wei Zhang, Wenhao Jiang, Jingwen Wang, Wei Liu:
Controllable Video Captioning With POS Sequence Guidance Based on Gated Fusion Network. 2641-2650

3D From Multiview & Sensors

Yuxin Hou, Juho Kannala, Arno Solin:
Multi-View Stereo by Temporal Nonparametric Fusion. 2651-2660

Jiacheng Chen, Chen Liu, Jiaye Wu, Yasutaka Furukawa:
Floor-SP: Inverse CAD for Floorplans by Sequential Room-Wise Shortest Path. 2661-2670

Zhaopeng Cui, Viktor Larsson, Marc Pollefeys:
Polarimetric Relative Pose Estimation. 2671-2680

Seong Hun Lee, Javier Civera:
Closed-Form Optimal Two-View Triangulation Based on Angular Errors. 2681-2689

Haozhe Xie, Hongxun Yao, Xiaoshuai Sun, Shangchen Zhou, Shengping Zhang:
Pix2Vox: Context-Aware 3D Reconstruction From Single and Multi-View Images. 2690-2698

Image & Video Synthesis

Patrick Esser, Johannes Haux, Björn Ommer:
Unsupervised Robust Disentangling of Latent Characteristics for Image Synthesis. 2699-2709

Mohammad Saeed Rad, Behzad Bozorgtabar, Urs-Viktor Marti, Max Basler, Hazim Kemal Ekenel, Jean-Philippe Thiran:
SROBB: Targeted Perceptual Loss for Single Image Super-Resolution. 2710-2719

Haotian Zhang, Long Mai, Hailin Jin, Zhaowen Wang, Ning Xu, John P. Collomosse:
An Internal Learning Approach to Video Inpainting. 2720-2729

Sai Bi, Kalyan Sunkavalli, Federico Perazzi, Eli Shechtman, Vladimir G. Kim, Ravi Ramamoorthi:
Deep CG2Real: Synthetic-to-Real Translation via Image Disentanglement. 2730-2739

Yunseok Jang, Tianchen Zhao, Seunghoon Hong, Honglak Lee:
Adversarial Defense via Learning to Generate Diverse Attacks. 2740-2749

Atsuhiro Noguchi, Tatsuya Harada:
Image Generation From Small Datasets via Batch Statistics Adaptation.

2750-2758

Mengyao Zhai, Lei Chen, Frederick Tung, Jiawei He, Megha Nawhal, Greg Mori: **Lifelong GAN: Continual Learning for Conditional Image Generation.** 2759-2768

Applications. Medical, & Robotics

Yi Wu, Yuxin Wu, Aviv Tamar, Stuart Russell, Georgia Gkioxari, Yuandong Tian: **Bayesian Relational Memory for Semantic Visual Navigation.** 2769-2779

Fabian Brickwedde, Steffen Abraham, Rudolf Mester: **Mono-SF: Multi-View Geometry Meets Single-View Depth for Monocular Scene Flow Estimation of Dynamic Traffic Scenes.** 2780-2790

Zhaoyang Huang, Yan Xu, Jianping Shi, Xiaowei Zhou, Hujun Bao, Guofeng Zhang: **Prior Guided Dropout for Robust Visual Localization in Dynamic Environments.** 2791-2800

Manuel Martin, Alina Roitberg, Monica Haurilet, Matthias Horne, Simon Reiß, Michael Voit, Rainer Stiefelhagen: **Drive&Act: A Multi-Modal Dataset for Fine-Grained Driver Behavior Recognition in Autonomous Vehicles.** 2801-2810

Yan Xu, Xinge Zhu, Jianping Shi, Guofeng Zhang, Hujun Bao, Hongsheng Li: **Depth Completion From Sparse LiDAR Data With Depth-Normal Constraints.** 2811-2820

Nicholas Rhinehart, Rowan McAllister, Kris Kitani, Sergey Levine: **PRECOG: PREdiction Conditioned on Goals in Visual Multi-Agent Settings.** 2821-2830

Zhe Liu, Shunbo Zhou, Chuanzhe Suo, Peng Yin, Wen Chen, Hesheng Wang, Haoang Li, Yunhui Liu: **LPD-Net: 3D Point Cloud Learning for Large-Scale Place Recognition and Environment Analysis.** 2831-2840

Fei Xue, Xin Wang, Zike Yan, Qiuyuan Wang, Junqiu Wang, Hongbin Zha: **Local Supports Global: Deep Camera Relocalization With Sequence Enhancement.** 2841-2850

Shunkai Li, Fei Xue, Xin Wang, Zike Yan, Hongbin Zha: **Sequential Adversarial Learning for Self-Supervised Deep Visual Odometry.** 2851-2860

Ziyang Hong, Yvan R. Petillot, David Lane, Yishu Miao, Sen Wang: **TextPlace: Visual Place Recognition and Topological Localization Through Reading Scene Texts.** 2861-2870

Mingyu Ding, Zhe Wang, Jiankai Sun, Jianping Shi, Ping Luo: **CamNet: Coarse-to-Fine Retrieval for Camera Re-Localization.** 2871-2880

William B. Shen, Danfei Xu, Yuke Zhu, Li Fei-Fei, Leonidas J. Guibas, Silvio Savarese: **Situational Fusion of Visual Representation for Visual Navigation.** 2881-2890

Ziyuan Huang, Changhong Fu, Yiming Li, Fuling Lin, Peng Lu:
Learning Aberrance Repressed Correlation Filters for Real-Time UAV Tracking. 2891-2900

Arsalan Mousavian, Clemens Eppner, Dieter Fox:
6-DOF GraspNet: Variational Grasp Generation for Object Manipulation. 2901-2910

Namdar Homayounfar, Justin Liang, Wei-Chiu Ma, Jack Fan, Xinyu Wu, Raquel Urtasun:
DAGMapper: Learning to Map by Discovering Lane Topology. 2911-2920

Noa Garnett, Rafi Cohen, Tomer Pe'er, Roee Lahav, Dan Levi:
3D-LaneNet: End-to-End 3D Multiple Lane Detection. 2921-2930

Oral 2.1A

Feature Representations, Similarity Learning

Janis Postels, Francesco Ferroni, Huseyin Coskun, Nassir Navab, Federico Tombari:
Sampling-Free Epistemic Uncertainty Estimation Using Approximated Variance Propagation. 2931-2940

Hong Liu, Rongrong Ji, Jie Li, Baochang Zhang, Yue Gao, Yongjian Wu, Feiyue Huang:
Universal Adversarial Perturbation via Prior Driven Uncertainty Approximation. 2941-2949

Ruth Fong, Mandela Patrick, Andrea Vedaldi:
Understanding Deep Networks via Extremal Perturbations and Smooth Masks. 2950-2958

Mathilde Caron, Piotr Bojanowski, Julien Mairal, Armand Joulin:
Unsupervised Pre-Training of Image Features on Non-Curated Data. 2959-2968

Linguang Zhang, Szymon Rusinkiewicz:
Learning Local Descriptors With a CDF-Based Dynamic Soft Margin. 2969-2978

Minyoung Kim, Yuting Wang, Pritish Sahu, Vladimir Pavlovic:
Bayes-Factor-VAE: Hierarchical Bayesian Deep Auto-Encoder Models for Factor Disentanglement. 2979-2987

Wei Jiang, Weiwei Sun, Andrea Tagliasacchi, Eduard Trulls, Kwang Moo Yi:
Linearized Multi-Sampling for Differentiable Image Transformation. 2988-2997

Zhiqiang Tang, Xi Peng, Tingfeng Li, Yizhe Zhu, Dimitris N. Metaxas:
AdaTransform: Adaptive Data Transformation. 2998-3006

Jiaqi Wang, Kai Chen, Rui Xu, Ziwei Liu, Chen Change Loy, Dahua Lin:
CARAFE: Content-Aware ReAssembly of FEatures. 3007-3016

Dou Quan, Xuefeng Liang, Shuang Wang, Shaowei Wei, Yanfeng Li, Ning Huyan, Licheng Jiao:
AFD-Net: Aggregated Feature Difference Learning for Cross-Spectral Image Patch Matching. 3017-3026

Shupeng Su, Zhisheng Zhong, Chao Zhang:
Deep Joint-Semantics Reconstructing Hashing for Large-Scale Unsupervised Cross-Modal Retrieval. 3027-3035

Stanislav Morozov, Artem Babenko:
Unsupervised Neural Quantization for Compressed-Domain Similarity Search. 3036-3045

Soumava Kumar Roy, Mehrtash Harandi, Richard Nock, Richard I. Hartley:
Siamese Networks: The Tale of Two Manifolds. 3046-3055

Runzhong Wang, Junchi Yan, Xiaokang Yang:
Learning Combinatorial Embedding Networks for Deep Graph Matching. 3056-3065

Zhanghui Kuang, Yiming Gao, Guanbin Li, Ping Luo, Yimin Chen, Liang Lin, Wayne Zhang:
Fashion Retrieval via Graph Reasoning Networks on a Similarity Pyramid. 3066-3075

Oral 2.1B

Low Level Vision

Xin Deng, Ren Yang, Mai Xu, Pier Luigi Dragotti:
Wavelet Domain Style Transfer for an Effective Perception-Distortion Tradeoff in Single Image Super-Resolution. 3076-3085

Jianrui Cai, Hui Zeng, Hongwei Yong, Zisheng Cao, Lei Zhang:
Toward Real-World Single Image Super-Resolution: A New Benchmark and a New Model. 3086-3095

Wenlong Zhang, Yihao Liu, Chao Dong, Yu Qiao:
RankSRGAN: Generative Adversarial Networks With Ranker for Image Super-Resolution. 3096-3105

Peng Yi, Zhongyuan Wang, Kui Jiang, Junjun Jiang, Jiayi Ma:
Progressive Fusion Video Super-Resolution Network via Exploiting Non-Local Spatio-Temporal Correlations. 3106-3115

Soo Ye Kim, Jihyong Oh, Munchurl Kim:
Deep SR-ITM: Joint Learning of Super-Resolution and Inverse Tone-Mapping for 4K UHD HDR Applications. 3116-3125

Tatsuya Yokota, Kazuya Kawai, Muneyuki Sakata, Yuichi Kimura, Hidekata Hontani:
Dynamic PET Image Reconstruction Using Nonnegative Matrix Factorization Incorporated With Deep Image Prior. 3126-3135

Jerry Liu, Shenlong Wang, Raquel Urtasun:
DSIC: Deep Stereo Image Compression. 3136-3145

Yoojin Choi, Mostafa El-Khamy, Jungwon Lee:
Variable Rate Deep Image Compression With a Conditional Autoencoder. 3146-3154

Saeed Anwar, Nick Barnes:
Real Image Denoising With Feature Attention. 3155-3164

Abdelrahman Abdelhamed, Marcus A. Brubaker, Michael S. Brown:
Noise Flow: Noise Modeling With Conditional Normalizing Flows. 3165-3173

Ahmed Abbas, Paul Swoboda:
Bottleneck Potentials in Markov Random Fields. 3174-3183

Chen Chen, Qifeng Chen, Minh N. Do, Vladlen Koltun:
Seeing Motion in the Dark. 3184-3193

Huaizu Jiang, Deqing Sun, Varun Jampani, Zhaoyang Lv, Erik G. Learned-Miller, Jan Kautz:
SENSE: A Shared Encoder Network for Scene-Flow Estimation. 3194-3203

Poster 2.1

Deep Learning

Firas Shama, Roey Mechrez, Alon Shoshan, Lihi Zelnik-Manor:
Adversarial Feedback Loop. 3204-3213

Alon Shoshan, Roey Mechrez, Lihi Zelnik-Manor:
Dynamic-Net: Tuning the Objective Without Re-Training for Synthesis Tasks. 3214-3222

Xinyu Gong, Shiyu Chang, Yifan Jiang, Zhangyang Wang:
AutoGAN: Neural Architecture Search for Generative Adversarial Networks. 3223-3233

Han Shu, Yunhe Wang, Xu Jia, Kai Han, Hanting Chen, Chunjing Xu, Qi Tian, Chang Xu:
Co-Evolutionary Compression for Unpaired Image Translation. 3234-3243

Zeyu Feng, Chang Xu, Dacheng Tao:
Self-Supervised Representation Learning From Multi-Domain Data. 3244-3254

Michael Möller, Thomas Möllenhoff, Daniel Cremers:
Controlling Neural Networks via Energy Dissipation. 3255-3264

Hao Lu, Yutong Dai, Chunhua Shen, Songcen Xu:
Indices Matter: Learning to Index for Deep Image Matting. 3265-3274

Yunan Li, Qiguang Miao, Wanli Ouyang, Zhenxin Ma, Huijuan Fang, Chao Dong, Yi-Ning Quan:
LAP-Net: Level-Aware Progressive Network for Image Dehazing. 3275-3284

Irwan Bello, Barret Zoph, Quoc Le, Ashish Vaswani, Jonathon Shlens:
Attention Augmented Convolutional Networks. 3285-3294

Zechun Liu, Haoyuan Mu, Xiangyu Zhang, Zichao Guo, Xin Yang, Kwang-Ting Cheng, Jian Sun:
MetaPruning: Meta Learning for Automatic Neural Network Channel Pruning. 3295-3304

Yuefu Zhou, Ya Zhang, Yan-Feng Wang, Qi Tian:
Accelerate CNN via Recursive Bayesian Pruning. 3305-3314

Duo Li, Aojun Zhou, Anbang Yao:
HBONet: Harmonious Bottleneck on Two Orthogonal Dimensions. 3315-3324

Jinchi Huang, Lie Qu, Rongfei Jia, Binqiang Zhao:
O2U-Net: A Simple Noisy Label Detection Approach for Deep Neural Networks. 3325-3333

Dongmin Park, Seokil Hong, Bohyung Han, Kyoung Mu Lee:
Continual Learning by Asymmetric Loss Approximation With Single-Side Overestimation. 3334-3343

Weifeng Ge, Weilin Huang, Sheng Guo, Matthew R. Scott:
Label-PEnet: Sequential Label Propagation and Enhancement Networks for Weakly Supervised Instance Segmentation. 3344-3353

Ziteng Gao, Limin Wang, Gangshan Wu:
LIP: Local Importance-Based Pooling. 3354-3363

Takumi Kobayashi:
Global Feature Guided Local Pooling. 3364-3373

Jinghua Wang, Jianmin Jiang:
Conditional Coupled Generative Adversarial Networks for Zero-Shot Domain Adaptation. 3374-3383

Aamir Mustafa, Salman H. Khan, Munawar Hayat, Roland Goecke, Jianbing Shen, Ling Shao:
Adversarial Defense by Restricting the Hidden Space of Deep Neural Networks. 3384-3393

Juhong Min, Jongmin Lee, Jean Ponce, Minsu Cho:
Hyperpixel Flow: Semantic Correspondence With Multi-Layer Neural Features. 3394-3403

Weitao Wan, Jiansheng Chen, Tianpeng Li, Yiqing Huang, Jingqi Tian, Cheng Yu, Youze Xue:
Information Entropy Based Feature Pooling for Convolutional Neural Networks. 3404-3413

Yuning Chai:
Patchwork: A Patch-Wise Attention Network for Efficient Object Detection and Segmentation in Video Streams. 3414-3423

Siddhesh Khandelwal, Leonid Sigal:
AttentionRNN: A Structured Spatial Attention Mechanism. 3424-3433

Yunpeng Chen, Haoqi Fan, Bing Xu, Zhicheng Yan, Yannis Kalantidis, Marcus Rohrbach, Shuicheng Yan, Jiashi Feng:
Drop an Octave: Reducing Spatial Redundancy in Convolutional Neural Networks With Octave Convolution. 3434-3443

Sagie Benaim, Michael Khaitov, Tomer Galanti, Lior Wolf:
Domain Intersection and Domain Difference. 3444-3452

Oren Rippel, Sanjay Nair, Carissa Lew, Steve Branson, Alexander G. Anderson, Lubomir D. Bourdev:
Learned Video Compression. 3453-3462

Han Hu, Zheng Zhang, Zhenda Xie, Stephen Lin:
Local Relation Networks for Image Recognition. 3463-3472

Éloi Mehr, Ariane Jourdan, Nicolas Thome, Matthieu Cord, Vincent Guitteny:
DiscoNet: Shapes Learning on Disconnected Manifolds for 3D Editing. 3473-3482

Max Ehrlich, Larry Davis:
Deep Residual Learning in the JPEG Transform Domain. 3483-3492

Xinqi Zhu, Chang Xu, Langwen Hui, Cewu Lu, Dacheng Tao:
Approximated Bilinear Modules for Temporal Modeling. 3493-3502

Chengchao Shen, Mengqi Xue, Xinchao Wang, Jie Song, Li Sun, Mingli Song:
Customizing Student Networks From Heterogeneous Teachers via Adaptive Knowledge Amalgamation. 3503-3512

Hanting Chen, Yunhe Wang, Chang Xu, Zhaohui Yang, Chuanjian Liu, Boxin Shi, Chunjing Xu, Chao Xu, Qi Tian:
Data-Free Learning of Student Networks. 3513-3521

Yue Wang, Justin Solomon:
Deep Closest Point: Learning Representations for Point Cloud Registration. 3522-3531

Chao Zhang, Stephan Liwicki, William Smith, Roberto Cipolla:
Orientation-Aware Semantic Segmentation on Icosahedron Spheres. 3532-3540

Zhaoyang Zhang, Jingyu Li, Wenqi Shao, Zhanglin Peng, Ruimao Zhang, Xiaogang Wang, Ping Luo:
Differentiable Learning-to-Group Channels via Groupable Convolutional Neural Networks. 3541-3550

Ping Chao, Chao-Yang Kao, Yu-Shan Ruan, Chien-Hsiang Huang, Youn-Long Lin:
HarDNet: A Low Memory Traffic Network. 3551-3560

Junjun He, Zhongying Deng, Yu Qiao:
Dynamic Multi-Scale Filters for Semantic Segmentation. 3561-3571

Ravi Teja Mullapudi, Steven Chen, Keyi Zhang, Deva Ramanan, Kayvon Fatahalian:
Online Model Distillation for Efficient Video Inference. 3572-3581

Recognition

Kai Li, Martin Renqiang Min, Yun Fu:
Rethinking Zero-Shot Learning: A Conditional Visual Classification Perspective. 3582-3591

Senthil Purushwalkam, Maximilian Nickel, Abhinav Gupta, Marc'Aurelio Ranzato:
Task-Driven Modular Networks for Zero-Shot Compositional Learning. 3592-3601

Limeng Qiao, Yemin Shi, Jia Li, Yonghong Tian, Tiejun Huang, Yaowei Wang:
Transductive Episodic-Wise Adaptive Metric for Few-Shot Learning. 3602-3611

Wei Zhai, Yang Cao, Jing Zhang, Zheng-Jun Zha:
Deep Multiple-Attribute-Perceived Network for Real-World Texture Recognition. 3612-3621

Guan'an Wang, Tianzhu Zhang, Jian Cheng, Si Liu, Yang Yang, Zengguang Hou:
RGB-Infrared Cross-Modality Person Re-Identification via Joint Pixel and Feature Alignment. 3622-3631

Saurabh Singh, Abhinav Shrivastava:
EvalNorm: Estimating Batch Normalization Statistics for Evaluation. 3632-3640

Jianyuan Guo, Yuhui Yuan, Lang Huang, Chao Zhang, Jin-Ge Yao, Kai Han:
Beyond Human Parts: Dual Part-Aligned Representations for Person Re-Identification. 3641-3650

Qi Dong, Xiatian Zhu, Shaogang Gong:
Person Search by Text Attribute Query As Zero-Shot Learning. 3651-3660

Qing Liu, Lingxi Xie, Huiyu Wang, Alan L. Yuille:
Semantic-Aware Knowledge Preservation for Zero-Shot Sketch-Based Image Retrieval. 3661-3670

Hamed H. Aghdam, Abel Gonzalez-Garcia, Antonio M. López, Joost van de Weijer:
Active Learning for Deep Detection Neural Networks. 3671-3679

Xuanyi Dong, Yi Yang:
One-Shot Neural Architecture Search via Self-Evaluated Template Network. 3680-3689

Zuozhuo Dai, Mingqiang Chen, Xiaodong Gu, Siyu Zhu, Ping Tan:
Batch DropBlock Network for Person Re-Identification and Beyond. 3690-3700

Kaiyang Zhou, Yongxin Yang, Andrea Cavallaro, Tao Xiang:
Omni-Scale Feature Learning for Person Re-Identification. 3701-3711

Linfeng Zhang, Jiebo Song, Anni Gao, Jingwei Chen, Chenglong Bao, Kaisheng Ma:
Be Your Own Teacher: Improve the Performance of Convolutional Neural Networks via Self Distillation. 3712-3721

Nikita Dvornik, Julien Mairal, Cordelia Schmid:
Diversity With Cooperation: Ensemble Methods for Few-Shot Classification.

3722-3730

Cheng Xu, Zhaoqun Li, Qiang Qiu, Biao Leng, Jingfei Jiang:
Enhancing 2D Representation via Adjacent Views for 3D Shape Retrieval. 3731-3739

Kun Wei, Muli Yang, Hao Wang, Cheng Deng, Xianglong Liu:
Adversarial Fine-Grained Composition Learning for Unseen Attribute-Object Recognition. 3740-3748

Ruijie Quan, Xuanyi Dong, Yu Wu, Linchao Zhu, Yi Yang:
Auto-ReID: Searching for a Part-Aware ConvNet for Person Re-Identification. 3749-3758

Bryan Bryan, Yuan Gong, Yizhe Zhang, Christian Poellabauer:
Second-Order Non-Local Attention Networks for Person Re-Identification. 3759-3768

Segmentation, Grouping, & Shape

Zipeng Ye, Ran Yi, Minjing Yu, Yong-Jin Liu, Ying He:
Fast Computation of Content-Sensitive Superpixels and Supervoxels Using Q-Distances. 3769-3778

Dániel Baráth, Jiri Matas:
Progressive-X: Efficient, Anytime, Multi-Model Fitting Algorithm. 3779-3787

Yingyue Xu, Dan Xu, Xiaopeng Hong, Wanli Ouyang, Rongrong Ji, Min Xu, Guoying Zhao:
Structured Modeling of Joint Deep Feature and Prediction Refinement for Salient Object Detection. 3788-3797

Jinming Su, Jia Li, Yu Zhang, Changqun Xia, Yonghong Tian:
Selectivity or Invariance: Boundary-Aware Salient Object Detection. 3798-3807

Urbano Miguel Nunes, Yiannis Demiris:
Online Unsupervised Learning of the 3D Kinematic Structure of Arbitrary Rigid Bodies. 3808-3816

3D From Single View & RGBD

Bram Wallace, Bharath Hariharan:
Few-Shot Generalization for Single-Image 3D Reconstruction via Priors. 3817-3826

Clément Godard, Oisin Mac Aodha, Michael Firman, Gabriel J. Brostow:
Digging Into Self-Supervised Monocular Depth Estimation. 3827-3837

Jing Zhu, Yi Fang:
Learning Object-Specific Distance From a Monocular Image. 3838-3847

Geonho Cha, Minsik Lee, Songhwai Oh:
Unsupervised 3D Reconstruction Networks. 3848-3857

Dong Wook Shu, Sung Woo Park, Junseok Kwon:
3D Point Cloud Generative Adversarial Network Based on Tree Structured

Graph Convolutions. 3858-3867

Junjie Hu, Yan Zhang, Takayuki Okatani:
Visualization of Convolutional Neural Networks for Monocular Depth Estimation. 3868-3877

Action & Video

Ruohan Gao, Kristen Grauman:
Co-Separating Sounds of Visual Objects. 3878-3887

Tianwei Lin, Xiao Liu, Xin Li, Errui Ding, Shilei Wen:
BMN: Boundary-Matching Network for Temporal Action Proposal Generation. 3888-3897

Ziyi Liu, Le Wang, Qilin Zhang, Zhanning Gao, Zhenxing Niu, Nanning Zheng, Gang Hua:
Weakly Supervised Temporal Action Localization Through Contrast Based Evaluation Networks. 3898-3907

Chaoxu Guo, Bin Fan, Jie Gu, Qian Zhang, Shiming Xiang, Véronique Prinet, Chunhong Pan:
Progressive Sparse Local Attention for Video Object Detection. 3908-3917

Tete Xiao, Quanfu Fan, Danny Gutfreund, Mathew Monfort, Aude Oliva, Bolei Zhou:
Reasoning About Human-Object Interactions Through Dual Attention Networks. 3918-3927

Xiaohui Zeng, Renjie Liao, Li Gu, Yuwen Xiong, Sanja Fidler, Raquel Urtasun:
DMM-Net: Differentiable Mask-Matching Network for Video Object Segmentation. 3928-3937

Hao Wang, Cheng Deng, Junchi Yan, Dacheng Tao:
Asymmetric Cross-Guided Attention Network for Actor and Action Video Segmentation From Natural Language Query. 3938-3947

Huaijia Lin, Xiaojuan Qi, Jiaya Jia:
AGSS-VOS: Attention Guided Single-Shot Video Object Segmentation. 3948-3956

Jianing Li, Shiliang Zhang, Jingdong Wang, Wen Gao, Qi Tian:
Global-Local Temporal Representations for Video Person Re-Identification. 3957-3966

Chaowei Xiao, Ruizhi Deng, Bo Li, Taesung Lee, Benjamin Edwards, Jinfeng Yi, Dawn Song, Mingyan Liu, Ian M. Molloy:
AdvIT: Adversarial Frames Identifier Based on Temporal Consistency in Videos. 3967-3976

Motion & Tracking

Ziqin Wang, Jun Xu, Li Liu, Fan Zhu, Ling Shao:
RANet: Ranking Attention Network for Fast Video Object Segmentation. 3977-3986

Jiarui Xu, Yue Cao, Zheng Zhang, Han Hu:
Spatial-Temporal Relation Networks for Multi-Object Tracking. 3987-3997

Lianghua Huang, Xin Zhao, Kaiqi Huang:
Bridging the Gap Between Detection and Tracking: A Unified Approach. 3998-4008

Lichao Zhang, Abel Gonzalez-Garcia, Joost van de Weijer, Martin Danelljan, Fahad Shahbaz Khan:
Learning the Model Update for Siamese Trackers. 4009-4018

Linyu Zheng, Ming Tang, Yingying Chen, Jinqiao Wang, Hanqing Lu:
Fast-deepKCF Without Boundary Effect. 4019-4028

Computational Photography & Graphics

Xiuming Zhang, Jiayuan Mao, Yikai Li, William T. Freeman, Joshua B. Tenenbaum, Jiajun Wu:
Program-Guided Image Manipulators. 4029-4038

Pierre-André Brousseau, Sébastien Roy:
Calibration of Axial Fisheye Cameras Through Generic Virtual Central Models. 4039-4047

Vishwanath Saragadam, Raja Venkata, Jian Wang, Shree K. Nayar, Mohit Gupta:
Micro-Baseline Structured Light. 4048-4057

Xin Miao, Xin Yuan, Yunchen Pu, Vassilis Athitsos:
lambda-Net: Reconstruct Hyperspectral Images From a Snapshot Measurement. 4058-4068

Masako Kashiwagi, Nao Mishima, Tatsuo Kozakaya, Shinsaku Hiura:
Deep Depth From Aberration Map. 4069-4078

Lukas Murmann, Michaël Gharbi, Miika Aittala, Frédo Durand:
A Dataset of Multi-Illumination Images in the Wild. 4079-4088

Jie Song, Xu Chen, Otmar Hilliges:
Monocular Neural Image Based Rendering With Continuous View Control. 4089-4099

Marc Comino Trinidad, Ricardo Martin-Brualla, Florian Kainz, Janne Kontkanen:
Multi-View Image Fusion. 4100-4109

Low-Level & Optimization

Wei Wang, Xin Chen, Cheng Yang, Xiang Li, Xuemei Hu, Tao Yue:
Enhancing Low Light Videos by Exploring High Sensitivity Camera Noise. 4110-4118

Qifan Gao, Xiao Shu, Xiaolin Wu:
Deep Restoration of Vintage Photographs From Scanned Halftone Prints. 4119-4128

Qiqi Hou, Feng Liu:
Context-Aware Image Matting for Simultaneous Foreground and Alpha Estimation. 4129-4138

Wei Wang, Ruiming Guo, Yapeng Tian, Wenming Yang:
CFSNet: Toward a Controllable Feature Space for Image Restoration. 4139-4148

Wu Wang, Weihong Zeng, Yue Huang, Xinghao Ding, John W. Paisley:
Deep Blind Hyperspectral Image Fusion. 4149-4158

Sungmin Cha, Taesup Moon:
Fully Convolutional Pixel Adaptive Image Denoiser. 4159-4168

Hongyu Liu, Bin Jiang, Yi Xiao, Chao Yang:
Coherent Semantic Attention for Image Inpainting. 4169-4178

Yajun Qiu, Ruxin Wang, Dapeng Tao, Jun Cheng:
Embedded Block Residual Network: A Recursive Restoration Model for Single-Image Super-Resolution. 4179-4188

Shuhang Gu, Wen Li, Luc Van Gool, Radu Timofte:
Fast Image Restoration With Multi-Bin Trainable Linear Units. 4189-4198

Scene Understanding

Zenglin Shi, Pascal Mettes, Cees Snoek:
Counting With Focus for Free. 4199-4208

Behzad Bozorgtabar, Mohammad Saeed Rad, Dwarikanath Mahapatra, Jean-Philippe Thiran:
SynDeMo: Synergistic Deep Feature Alignment for Joint Learning of Depth and Ego-Motion. 4209-4218

Ke Li, Tianhao Zhang, Jitendra Malik:
Diverse Image Synthesis From Semantic Layouts via Conditional IMLE. 4219-4228

Yanwei Pang, Yazhao Li, Jianbing Shen, Ling Shao:
Towards Bridging Semantic Gap to Improve Semantic Segmentation. 4229-4238

Language & Reasoning

Lixin Liu, Jiajun Tang, Xiaojun Wan, Zongming Guo:
Generating Diverse and Descriptive Image Captions Using Visual Paraphrases. 4239-4248

Xu Yang, Hanwang Zhang, Jianfei Cai:
Learning to Collocate Neural Modules for Image Captioning. 4249-4259

Jyoti Aneja, Harsh Agrawal, Dhruv Batra, Alexander G. Schwing:
Sequential Latent Spaces for Modeling the Intention During Diverse Image Captioning. 4260-4269

Nilavra Bhattacharya, Qing Li, Danna Gurari:
Why Does a Visual Question Have Different Answers? 4270-4279

Mohit Bajaj, Lanjun Wang, Leonid Sigal:
G3raphGround: Graph-Based Language Grounding. 4280-4289

Ali Furkan Biten, Rubèn Tito, Andrés Mafla, Lluís Gómez i Bigorda, Marçal Rusiñol, C. V. Jawahar, Ernest Valveny, Dimosthenis Karatzas:
Scene Text Visual Question Answering. 4290-4300

3D From Multiview & Sensors

Lu Sheng, Dan Xu, Wanli Ouyang, Xiaogang Wang:
Unsupervised Collaborative Learning of Keyframe Detection and Visual Odometry Towards Monocular Deep SLAM. 4301-4310

Youze Xue, Jiansheng Chen, Weitao Wan, Yiqing Huang, Cheng Yu, Tianpeng Li, Jiayu Bao:
MVSCRF: Learning Multi-View Stereo With Conditional Random Fields. 4311-4320

Eric Brachmann, Carsten Rother:
Neural-Guided RANSAC: Learning Where to Sample Model Hypotheses. 4321-4330

Sergey Prokudin, Christoph Lassner, Javier Romero:
Efficient Learning on Point Clouds With Basis Point Sets. 4331-4340

Haibo Qiu, Chunyu Wang, Jingdong Wang, Naiyan Wang, Wenjun Zeng:
Cross View Fusion for 3D Human Pose Estimation. 4341-4350

Junbang Liang, Ming C. Lin:
Shape-Aware Human Pose and Shape Reconstruction Using Multi-View Images. 4351-4361

Di Yan, Henrique Morimitsu, Shan Gao, Xiangyang Ji:
Monocular Piecewise Depth Estimation in Dynamic Scenes by Exploiting Superpixel Relations. 4362-4371

Hajime Taira, Ignacio Rocco, Jirí Sedlár, Masatoshi Okutomi, Josef Sivic, Tomás Pajdla, Torsten Sattler, Akihiko Torii:
Is This the Right Place? Geometric-Semantic Pose Verification for Indoor Visual Localization. 4372-4382

Shivam Duggal, Shenlong Wang, Wei-Chiu Ma, Rui Hu, Raquel Urtasun:
DeepPruner: Learning Efficient Stereo Matching via Differentiable PatchMatch. 4383-4392

Image & Video Synthesis

Sijie Yan, Zhizhong Li, Yuanjun Xiong, Huahan Yan, Dahua Lin:
Convolutional Sequence Generation for Skeleton-Based Action Synthesis. 4393-4401

Seoung Wug Oh, Sungho Lee, Joon-Young Lee, Seon Joo Kim:
Onion-Peel Networks for Deep Video Completion. 4402-4411

Sungho Lee, Seoung Wug Oh, DaeYeun Won, Seon Joo Kim:
Copy-and-Paste Networks for Deep Video Inpainting. 4412-4420

Dmytro Kotovenko, Artsiom Sanakoyeu, Sabine Lang, Björn Ommer:
Content and Style Disentanglement for Artistic Style Transfer. 4421-4430

Oral 3.1A

Generative Modeling & Synthesis

Rameen Abdal, Yipeng Qin, Peter Wonka:
Image2StyleGAN: How to Embed Images Into the StyleGAN Latent Space? 4431-4440

Shuai Yang, Zhangyang Wang, Zhaowen Wang, Ning Xu, Jiaying Liu, Zongming Guo:
Controllable Artistic Text Style Transfer via Shape-Matching GAN. 4441-4450

Tai-Yin Chiu:
Understanding Generalized Whitening and Coloring Transform for Universal Style Transfer. 4451-4459

Cícero Nogueira dos Santos, Youssef Mroueh, Inkit Padhi, Pierre L. Dognin:
Learning Implicit Generative Models by Matching Perceptual Features. 4460-4469

Jiahui Yu, Zhe Lin, Jimei Yang, Xiaohui Shen, Xin Lu, Thomas S. Huang:
Free-Form Image Inpainting With Gated Convolution. 4470-4479

Xintong Han, Zuxuan Wu, Weilin Huang, Matthew R. Scott, Larry Davis:
FiNet: Compatible and Diverse Fashion Image Inpainting. 4480-4490

Assaf Shocher, Shai Bagon, Phillip Isola, Michal Irani:
InGAN: Capturing and Retargeting the "DNA" of a Natural Image. 4491-4500

David Bau, Jun-Yan Zhu, Jonas Wulff, William S. Peebles, Bolei Zhou, Hendrik Strobelt, Antonio Torralba:
Seeing What a GAN Cannot Generate. 4501-4510

Chieh Hubert Lin, Chia-Che Chang, Yu-Sheng Chen, Da-Cheng Juan, Wei Wei, Hwann-Tzong Chen:
COCO-GAN: Generation by Parts via Conditional Coordinating. 4511-4520

Hang Chu, Daiqing Li, David Acuna, Amlan Kar, Maria Shugrina, Xinkai Wei, Ming-Yu Liu, Antonio Torralba, Sanja Fidler:
Neural Turtle Graphics for Modeling City Road Layouts. 4521-4529

Michael Oechsle, Lars M. Mescheder, Michael Niemeyer, Thilo Strauss, Andreas Geiger:
Texture Fields: Learning Texture Representations in Function Space. 4530-4539

Guandao Yang, Xun Huang, Zekun Hao, Ming-Yu Liu, Serge J. Belongie, Bharath Hariharan:
PointFlow: 3D Point Cloud Generation With Continuous Normalizing Flows. 4540-4549

Amlan Kar, Aayush Prakash, Ming-Yu Liu, Eric Cameracci, Justin Yuan, Matt Rusiniak, David Acuna, Antonio Torralba, Sanja Fidler:
Meta-Sim: Learning to Generate Synthetic Datasets. 4550-4559

Oron Ashual, Lior Wolf:
Specifying Object Attributes and Relations in Interactive Scene Generation. 4560-4568

Tamar Rott Shaham, Tali Dekel, Tomer Michaeli:
SinGAN: Learning a Generative Model From a Single Natural Image. 4569-4579

Oral 3.1B

Vision, Language, & Text

Xin Wang, Jiawei Wu, Junkun Chen, Lei Li, Yuan-Fang Wang, William Yang Wang:
VaTeX: A Large-Scale, High-Quality Multilingual Dataset for Video-and-Language Research. 4580-4590

Yu Xiong, Qingqiu Huang, Lingfeng Guo, Hang Zhou, Bolei Zhou, Dahua Lin:
A Graph-Based Framework to Bridge Movies and Synopses. 4591-4600

Ajeet Kumar Singh, Anand Mishra, Shashank Shekhar, Anirban Chakraborty:
From Strings to Things: Knowledge-Enabled VQA Model That Can Read and Reason. 4601-4611

Long Chen, Hanwang Zhang, Jun Xiao, Xiangnan He, Shiliang Pu, Shih-Fu Chang:
Counterfactual Critic Multi-Agent Training for Scene Graph Generation. 4612-4622

Dong Huk Park, Trevor Darrell, Anna Rohrbach:
Robust Change Captioning. 4623-4632

Lun Huang, Wenmin Wang, Jie Chen, Xiaoyong Wei:
Attention on Attention for Image Captioning. 4633-4642

Sibei Yang, Guanbin Li, Yizhou Yu:
Dynamic Graph Attention for Referring Expression Comprehension. 4643-4652

Kunpeng Li, Yulun Zhang, Kai Li, Yuanyuan Li, Yun Fu:
Visual Semantic Reasoning for Image-Text Matching. 4653-4661

Josiah Wang, Lucia Specia:
Phrase Localization Without Paired Training Examples. 4662-4671

Daqing Liu, Hanwang Zhang, Feng Wu, Zheng-Jun Zha:
Learning to Assemble Neural Module Tree Networks for Visual Grounding. 4672-4681

Zhengyuan Yang, Boqing Gong, Liwei Wang, Wenbing Huang, Dong Yu, Jiebo Luo:

A Fast and Accurate One-Stage Approach to Visual Grounding. 4682-4692

Arka Sadhu, Kan Chen, Ram Nevatia:
Zero-Shot Grounding of Objects From Natural Language Queries. 4693-4702

Siyang Qin, Alessandro Bissacco, Michalis Raptis, Yasuhisa Fujii, Ying Xiao:
Towards Unconstrained End-to-End Text Spotting. 4703-4713

Jeonghun Baek, Geewook Kim, Junyeop Lee, Sungrae Park, Dongyoon Han, Sangdoo Yun, Seong Joon Oh, Hwalsuk Lee:
What Is Wrong With Scene Text Recognition Model Comparisons? Dataset and Model Analysis. 4714-4722

Poster 3.1

Deep Learning

Francesco Croce, Matthias Hein:
Sparse and Imperceivable Adversarial Attacks. 4723-4731

Qian Huang, Isay Katsman, Zeqi Gu, Horace He, Serge J. Belongie, Ser-Nam Lim:
Enhancing Adversarial Example Transferability With an Intermediate Level Attack. 4732-4741

Mateusz Michalkiewicz, Jhony Kaesemodel Pontes, Dominic Jack, Mahsa Baktashmotlagh, Anders P. Eriksson:
Implicit Surface Representations As Layers in Neural Networks. 4742-4751

Pablo Navarrete Michelini, Hanwen Liu, Yunhua Lu, Xingqun Jiang:
A Tour of Convolutional Networks Guided by Linear Interpreters. 4752-4761

João F. Henriques, Sébastien Ehrhardt, Samuel Albanie, Andrea Vedaldi:
Small Steps and Giant Leaps: Minimal Newton Solvers for Deep Learning. 4762-4771

Ameya Joshi, Amitangshu Mukherjee, Soumik Sarkar, Chinmay Hegde:
Semantic Adversarial Attacks: Parametric Transformations That Fool Deep Classifiers. 4772-4782

Yang Bai, Yan Feng, Yisen Wang, Tao Dai, Shutao Xia, Yong Jiang:
Hilbert-Based Generative Defense for Adversarial Examples. 4783-4792

Jang Hyun Cho, Bharath Hariharan:
On the Efficacy of Knowledge Distillation. 4793-4801

Simyung Chang, Seonguk Park, John Yang, Nojun Kwak:
Sym-Parameterized Dynamic Inference for Mixed-Domain Image Translation. 4802-4810

Shuang Wang, Yanfeng Li, Xuefeng Liang, Dou Quan, Bowu Yang, Shaowei Wei, Licheng Jiao:
Better and Faster: Exponential Loss for Image Patch Matching. 4811-4820

Rey Wiyatno, Anqi Xu:
Physical Adversarial Textures That Fool Visual Object Tracking. 4821-4830

Huidong Liu, Xianfeng Gu, Dimitris Samaras:
Wasserstein GAN With Quadratic Transport Cost. 4831-4840

Sven Gowal, Krishnamurthy Dvijotham, Robert Stanforth, Rudy Bunel, Chongli Qin, Jonathan Uesato, Relja Arandjelovic, Timothy Arthur Mann, Pushmeet Kohli:
Scalable Verified Training for Provably Robust Image Classification. 4841-4850

Ruihao Gong, Xianglong Liu, Shenghu Jiang, Tianxiang Li, Peng Hu, Jiazhen Lin, Fengwei Yu, Junjie Yan:
Differentiable Soft Quantization: Bridging Full-Precision and Low-Bit Neural Networks. 4851-4860

Chris Finlay, Aram-Alexandre Pooladian, Adam M. Oberman:
The LogBarrier Adversarial Attack: Making Effective Use of Decision Boundary Information. 4861-4869

Thalaiyasingam Ajanthan, Puneet K. Dokania, Richard Hartley, Philip H. S. Torr:
Proximal Mean-Field for Neural Network Quantization. 4870-4879

Hao-Yun Chen, Jhao-Hong Liang, Shih-Chieh Chang, Jia-Yu Pan, Yu-Ting Chen, Wei Wei, Da-Cheng Juan:
Improving Adversarial Robustness via Guided Complement Entropy. 4880-4888

Yujia Liu, Seyed-Mohsen Moosavi-Dezfooli, Pascal Frossard:
A Geometry-Inspired Decision-Based Attack. 4889-4897

Jie Li, Rongrong Ji, Hong Liu, Xiaopeng Hong, Yue Gao, Qi Tian:
Universal Perturbation Attack Against Image Retrieval. 4898-4907

Jiaxin Gu, Junhe Zhao, Xiaolong Jiang, Baochang Zhang, Jianzhuang Liu, Guodong Guo, Rongrong Ji:
Bayesian Optimized 1-Bit CNNs. 4908-4916

Kaiming He, Ross B. Girshick, Piotr Dollár:
Rethinking ImageNet Pre-Training. 4917-4926

Chaithanya Kumar Mummadi, Thomas Brox, Jan Hendrik Metzen:
Defending Against Universal Perturbations With Shared Adversarial Training. 4927-4936

Yiyou Sun, Sathya N. Ravi, Vikas Singh:
Adaptive Activation Thresholding: Dynamic Routing Type Behavior for Interpretability in Convolutional Neural Networks. 4937-4946

Andrei Kapishnikov, Tolga Bolukbasi, Fernanda B. Viégas, Michael Terry:
XRAI: Better Attributions Through Regions. 4947-4956

Thomas Brunner, Frederik Diehl, Michael Truong-Le, Alois C. Knoll:
Guessing Smart: Biased Sampling for Efficient Black-Box Adversarial Attacks. 4957-4965

Recognition

Yanwei Pang, Jin Xie, Muhammad Haris Khan, Rao Muhammad Anwer, Fahad Shahbaz Khan, Ling Shao:

Mask-Guided Attention Network for Occluded Pedestrian Detection. 4966-4974

Chuanchen Luo, Yuntao Chen, Naiyan Wang, Zhaoxiang Zhang:
Spectral Feature Transformation for Person Re-Identification. 4975-4984

Xiaofeng Liu, Zhenhua Guo, Site Li, Ping Jia, Lingsheng Kong, Jane You, B. V. K. Vijaya Kumar:
Permutation-Invariant Feature Restructuring for Correlation-Aware Image Set-Based Recognition. 4985-4995

Chufeng Tang, Lu Sheng, Zhaoxiang Zhang, Xiaolin Hu:
Improving Pedestrian Attribute Recognition With Weakly-Supervised Multi-Scale Attribute-Specific Localization. 4996-5005

Baoyun Peng, Xiao Jin, Dongsheng Li, Shunfeng Zhou, Yichao Wu, Jiaheng Liu, Zhaoning Zhang, Yu Liu:
Correlation Congruence for Knowledge Distillation. 5006-5015

Yiru Wang, Weihao Gan, Jie Yang, Wei Wu, Junjie Yan:
Dynamic Curriculum Learning for Imbalanced Data Classification. 5016-5025

Makarand Tapaswi, Marc T. Law, Sanja Fidler:
Video Face Clustering With Unknown Number of Clusters. 5026-5035

Giorgos Tolias, Filip Radenovic, Ondrej Chum:
Targeted Mismatch Adversarial Attack: Query With a Flower to Retrieve the Tower. 5036-5045

Wei-Lin Hsiao, Isay Katsman, Chao-Yuan Wu, Devi Parikh, Kristen Grauman:
Fashion++: Minimal Edits for Outfit Improvement. 5046-5055

Si Wu, Sihao Lin, Wenhao Wu, Mohamed Azzam, Hau-San Wong:
Semi-Supervised Pedestrian Instance Synthesis and Detection With Mutual Reinforcement. 5056-5065

Tao Hu, Pascal Mettes, Jia-Hong Huang, Cees Snoek:
SILCO: Show a Few Images, Localize the Common Object. 5066-5075

Jimmy Addison Lee, Peng Liu, Jun Cheng, Huazhu Fu:
A Deep Step Pattern Representation for Multimodal Retinal Image Registration. 5076-5085

Zhen Zhang, Wee Sun Lee:
Deep Graphical Feature Learning for the Feature Matching Problem. 5086-5095

Dong Lao, Ganesh Sundaramoorthi:
Minimum Delay Object Detection From Video. 5096-5105

Jérôme Revaud, Jon Almazán, Rafael S. Rezende, César Roberto de Souza:
Learning With Average Precision: Training Image Retrieval With a Listwise Loss. 5106-5115

Amirreza Shaban, Amir Rahimi, Shray Bansal, Stephen Gould, Byron Boots, Richard Hartley:

Learning to Find Common Objects Across Few Image Collections. 5116-5125

Lu Zhang, Xiangyu Zhu, Xiangyu Chen, Xu Yang, Zhen Lei, Zhiyong Liu:
Weakly Aligned Cross-Modal Learning for Multispectral Pedestrian Detection. 5126-5136

Jiangfan Han, Ping Luo, Xiaogang Wang:
Deep Self-Learning From Noisy Labels. 5137-5146

Marcelo Gennari Do Nascimento, Victor Prisacariu, Roger Fawcett:
DSConv: Efficient Convolution Operator. 5147-5156

Jiangfan Han, Xiaoyi Dong, Ruimao Zhang, Dongdong Chen, Weiming Zhang, Nenghai Yu, Ping Luo, Xiaogang Wang:
Once a MAN: Towards Multi-Target Attack via Learning Multi-Target Adversarial Network Once. 5157-5166

Segmentation, Grouping, & Shape

Wenqiang Xu, Haiyang Wang, Fubo Qi, Cewu Lu:
Explicit Shape Encoding for Real-Time Instance Segmentation. 5167-5176

Cheng-Yang Fu, Tamara L. Berg, Alexander C. Berg:
IMP: Instance Mask Projection for High Accuracy Semantic Segmentation of Things. 5177-5186

Linjie Yang, Yuchen Fan, Ning Xu:
Video Instance Segmentation. 5187-5196

Kunpeng Li, Yulun Zhang, Kai Li, Yuanyuan Li, Yun Fu:
Attention Bridging Network for Knowledge Transfer. 5197-5206

Wataru Shimoda, Keiji Yanai:
Self-Supervised Difference Detection for Weakly-Supervised Semantic Segmentation. 5207-5216

Bowen Cheng, Liang-Chieh Chen, Yunchao Wei, Yukun Zhu, Zilong Huang, Jinjun Xiong, Thomas S. Huang, Wen-Mei Hwu, Honghui Shi:
SPGNet: Semantic Prediction Guidance for Scene Parsing. 5217-5227

Towaki Takikawa, David Acuna, Varun Jampani, Sanja Fidler:
Gated-SCNN: Gated Shape CNNs for Semantic Segmentation. 5228-5237

Yongcheng Liu, Bin Fan, Gaofeng Meng, Jiwen Lu, Shiming Xiang, Chunhong Pan:
DensePoint: Learning Densely Contextual Representation for Efficient Point Cloud Processing. 5238-5247

Mennatullah Siam, Boris N. Oreshkin, Martin Jägersand:
AMP: Adaptive Masked Proxies for Few-Shot Segmentation. 5248-5257

Tarun Kalluri, Girish Varma, Manmohan Chandraker, C. V. Jawahar:
Universal Semi-Supervised Semantic Segmentation. 5258-5269

Statistics, Physics, Theory & Datasets

Long-Kai Huang, Jianda Chen, Sinno Jialin Pan:
Accelerate Learning of Deep Hashing With Gradient Attention. 5270-5279

Qing-Yuan Jiang, Yi He, Gen Li, Jian Lin, Lei Li, Wu-Jun Li:
SVD: A Large-Scale Short Video Dataset for Near-Duplicate Video Retrieval. 5280-5288

Hubert Lin, Paul Upchurch, Kavita Bala:
Block Annotation: Better Image Annotation With Sub-Image Decomposition. 5289-5299

Yanzhu Liu, Fan Wang, Adams Wai-Kin Kong:
Probabilistic Deep Ordinal Regression Based on Gaussian Processes. 5300-5308

Tianlu Wang, Jieyu Zhao, Mark Yatskar, Kai-Wei Chang, Vicente Ordonez:
Balanced Datasets Are Not Enough: Estimating and Mitigating Gender Bias in Deep Image Representations. 5309-5318

Pouya Bashivan, Mark Tensen, James J. DiCarlo:
Teacher Guided Architecture Search. 5319-5328

3D From Single View & RGBD

David Smith, Matthew Loper, Xiaochen Hu, Paris Mavroidis, Javier Romero:
FACSIMILE: Fast and Accurate Scans From an Image in Less Than a Second. 5329-5338

Yu Rong, Ziwei Liu, Cheng Li, Kaidi Cao, Chen Change Loy:
Delving Deep Into Hybrid Annotations for 3D Human Recovery in the Wild. 5339-5347

Yu Sun, Yun Ye, Wu Liu, Wenpeng Gao, Yili Fu, Tao Mei:
Human Mesh Recovery From Monocular Images via a Skeleton-Disentangled Representation. 5348-5357

Silvia Zuffi, Angjoo Kanazawa, Tanya Y. Berger-Wolf, Michael J. Black:
Three-D Safari: Learning to Estimate Zebra Pose, Shape, and Texture From Images "In the Wild". 5358-5367

Helisa Dhamo, Nassir Navab, Federico Tombari:
Object-Driven Multi-Layer Scene Decomposition From a Single Image. 5368-5377

Michael Niemeyer, Lars M. Mescheder, Michael Oechsle, Andreas Geiger:
Occupancy Flow: 4D Reconstruction by Learning Particle Dynamics. 5378-5388

Hou-Ning Hu, Qi-Zhi Cai, Dequan Wang, Ji Lin, Min Sun, Philipp Krähenbühl, Trevor Darrell, Fisher Yu:
Joint Monocular 3D Vehicle Detection and Tracking. 5389-5398

Face & Body

Bowen Shi, Aurora Martinez Del Rio, Jonathan Keane, Diane Brentari, Greg Shakhnarovich, Karen Livescu:
Fingerspelling Recognition in the Wild With Iterative Visual Attention. 5399-5408

Hang Dai, Ling Shao:
PointAE: Point Auto-Encoder for 3D Statistical Shape and Texture Modelling. 5409-5418

Bharat Lal Bhatnagar, Garvita Tiwari, Christian Theobalt, Gerard Pons-Moll:
Multi-Garment Net: Learning to Dress 3D People From Images. 5419-5429

Haiyong Jiang, Jianfei Cai, Jianmin Zheng:
Skeleton-Aware 3D Human Shape Reconstruction From Point Clouds. 5430-5440

Naureen Mahmood, Nima Ghorbani, Nikolaus F. Troje, Gerard Pons-Moll, Michael J. Black:
AMASS: Archive of Motion Capture As Surface Shapes. 5441-5450

Fei Wang, Sanping Zhou, Stanislav Panev, Jinsong Han, Dong Huang:
Person-in-WiFi: Fine-Grained Person Perception Using WiFi. 5451-5460

Keqiang Sun, Wayne Wu, Tinghao Liu, Shuo Yang, Quan Wang, Qiang Zhou, Zuochang Ye, Chen Qian:
FAB: A Robust Facial Landmark Detection Framework for Motion-Blurred Videos. 5461-5470

Bong-Nam Kang, Yonghyun Kim, Bongjin Jun, Daijin Kim:
Attentional Feature-Pair Relation Networks for Accurate Face Recognition. 5471-5480

Action & Video

Brais Martínez, Davide Modolo, Yuanjun Xiong, Joseph Tighe:
Action Recognition With Spatial-Temporal Discriminative Filter Banks. 5481-5490

Evangelos Kazakos, Arsha Nagrani, Andrew Zisserman, Dima Damen:
EPIC-Fusion: Audio-Visual Temporal Binding for Egocentric Action Recognition. 5491-5500

Phuc Xuan Nguyen, Deva Ramanan, Charless C. Fowlkes:
Weakly-Supervised Action Localization With Background Modeling. 5501-5510

Chenxu Luo, Alan L. Yuille:
Grouped Spatial-Temporal Aggregation for Efficient Action Recognition. 5511-5520

Tan Yu, Zhou Ren, Yuncheng Li, Enxu Yan, Ning Xu, Junsong Yuan:
Temporal Structure Mining for Weakly Supervised Action Detection. 5521-5530

Mingze Xu, Mingfei Gao, Yi-Ting Chen, Larry Davis, David J. Crandall:
Temporal Recurrent Networks for Online Action Detection. 5531-5540

Mingfei Gao, Mingze Xu, Larry Davis, Richard Socher, Caiming Xiong:
StartNet: Online Detection of Action Start in Untrimmed Videos. 5541-5550

Du Tran, Heng Wang, Matt Feiszli, Lorenzo Torresani:
Video Classification With Channel-Separated Convolutional Networks.

5551-5560

Harshala Gammulle, Simon Denman, Sridha Sridharan, Clinton Fookes:
Predicting the Future: A Jointly Learnt Model for Action Anticipation. 5561-5570

Low-Level & Optimization

Ziyi Shen, Wenguan Wang, Xiankai Lu, Jianbing Shen, Haibin Ling, Tingfa Xu, Ling Shao:
Human-Aware Motion Deblurring. 5571-5580

Lu Zhang, Zhe Lin, Jianming Zhang, Huchuan Lu, You He:
Fast Video Object Segmentation via Dynamic Targeting Network. 5581-5590

Sean I. Young, Aous Thabit Naman, Bernd Girod, David Taubman:
Solving Vision Problems via Filtering. 5591-5600

Ankit Raj, Yuqi Li, Yoram Bresler:
GAN-Based Projector for Faster Recovery With Convergence Guarantees in Linear Inverse Problems. 5601-5610

Deng-Ping Fan, Shengchuan Zhang, Yu-Huan Wu, Yun Liu, Ming-Ming Cheng, Bo Ren, Paul L. Rosin, Rongrong Ji:
Scoot: A Perceptual Metric for Facial Sketches. 5611-5621

Yawei Li, Shuhang Gu, Luc Van Gool, Radu Timofte:
Learning Filter Basis for Convolutional Neural Network Compression. 5622-5631

Daniel Gehrig, Antonio Loquercio, Konstantinos G. Derpanis, Davide Scaramuzza:
End-to-End Learning of Representations for Asynchronous Event-Based Data. 5632-5642

Guoqing Wang, Changming Sun, Arcot Sowmya:
ERL-Net: Entangled Representation Learning for Single Image De-Raining. 5643-5651

Oleg Voynov, Alexey Artemov, Vage Egiazarian, Alexandr Notchenko, Gleb Bobrovskikh, Evgeny Burnaev, Denis Zorin:
Perceptual Deep Depth Super-Resolution. 5652-5662

Scene Understanding

Iro Armeni, Zhi-Yang He, Amir Zamir, JunYoung Gwak, Jitendra Malik, Martin Fischer, Silvio Savarese:
3D Scene Graph: A Structure for Unified Semantics, 3D Space, and Camera. 5663-5672

Cheng Lin, Changjian Li, Wenping Wang:
Floorplan-Jigsaw: Jointly Estimating Scene Layout and Aligning Partial Scans. 5673-5682

Wei Yin, Yifan Liu, Chunhua Shen, Youliang Yan:
Enforcing Geometric Constraints of Virtual Normal for Depth Prediction. 5683-5692

Tiancai Wang, Rao Muhammad Anwer, Muhammad Haris Khan, Fahad Shahbaz Khan, Yanwei Pang, Ling Shao, Jorma Laaksonen:
Deep Contextual Attention for Human-Object Interaction Detection. 5693-5701

Wenguan Wang, Zhijie Zhang, Siyuan Qi, Jianbing Shen, Yanwei Pang, Ling Shao:
Learning Compositional Neural Information Fusion for Human Parsing. 5702-5712

Anran Zhang, Lei Yue, Jiayi Shen, Fan Zhu, Xiantong Zhen, Xianbin Cao, Ling Shao:
Attentional Neural Fields for Crowd Counting. 5713-5722

Lifeng Fan, Wenguan Wang, Song-Chun Zhu, Xinyu Tang, Siyuan Huang:
Understanding Human Gaze Communication by Spatio-Temporal Graph Reasoning. 5723-5732

Jean-Baptiste Alayrac, João Carreira, Relja Arandjelovic, Andrew Zisserman:
Controllable Attention for Structured Layered Video Decomposition. 5733-5742

Lore Goetschalckx, Alex Andonian, Aude Oliva, Phillip Isola:
GANalyze: Toward Visual Definitions of Cognitive Image Properties. 5743-5752

Language & Reasoning

Zhong Ji, Haoran Wang, Jungong Han, Yanwei Pang:
Saliency-Guided Attention Network for Image-Sentence Matching. 5753-5762

Zihao Wang, Xihui Liu, Hongsheng Li, Lu Sheng, Junjie Yan, Xiaogang Wang, Jing Shao:
CAMP: Cross-Modal Adaptive Message Passing for Text-Image Retrieval. 5763-5772

Yan Huang, Liang Wang:
ACMM: Aligned Cross-Modal Memory for Few-Shot Image and Sentence Matching. 5773-5782

Mohamed Elhoseiny, Mohamed Elfeki:
Creativity Inspired Zero-Shot Learning. 5783-5792

Mikihiro Tanaka, Takayuki Itamochi, Kenichi Narioka, Ikuro Sato, Yoshitaka Ushiku, Tatsuya Harada:
Generating Easy-to-Understand Referring Expressions for Target Identifications. 5793-5802

Jonatas Wehrmann, Maurício Armani Lopes, Douglas M. Souza, Rodrigo C. Barros:
Language-Agnostic Visual-Semantic Embeddings. 5803-5812

Nikolaos Sarafianos, Xiang Xu, Ioannis A. Kakadiaris:
Adversarial Representation Learning for Text-to-Image Matching. 5813-5823

Peng Gao, Haoxuan You, Zhanpeng Zhang, Xiaogang Wang, Hongsheng Li:
Multi-Modality Latent Interaction Network for Visual Question Answering. 5824-5834

3D From Multiview & Sensors

Axel Barroso Laguna, Edgar Riba, Daniel Ponsa, Krystian Mikolajczyk:
Key.Net: Keypoint Detection by Handcrafted and Learned CNN Filters. 5835-5843

Jiahui Zhang, Dawei Sun, Zixin Luo, Anbang Yao, Lei Zhou, Tianwei Shen, Yurong Chen, Hongen Liao, Long Quan:
Learning Two-View Correspondences and Geometry Using Order-Aware Network. 5844-5853

Michael Bloesch, Tristan Laidlow, Ronald Clark, Stefan Leutenegger, Andrew J. Davison:
Learning Meshes for Dense Visual SLAM. 5854-5863

Michael Strecke, Jörg Stückler:
EM-Fusion: Dynamic Object-Level SLAM With Probabilistic Data Association. 5864-5873

Jiahui Huang, Sheng Yang, Zishuo Zhao, Yu-Kun Lai, Shimin Hu:
ClusterSLAM: A SLAM Backend for Simultaneous Rigid Body Clustering and Motion Estimation. 5874-5883

Uttaran Bhattacharya, Venu Madhav Govindu:
Efficient and Robust Registration on the 3D Special Euclidean Group. 5884-5893

Yoni Kasten, Amnon Geifman, Meirav Galun, Ronen Basri:
Algebraic Characterization of Essential Matrices and Their Averaging in Multiview Settings. 5894-5902

Image & Video Synthesis

Wen Liu, Zhixin Piao, Jie Min, Wenhan Luo, Lin Ma, Shenghua Gao:
Liquid Warping GAN: A Unified Framework for Human Motion Imitation, Appearance Transfer and Novel View Synthesis. 5903-5912

Yu-Jing Lin, Po-Wei Wu, Che-Han Chang, Edward Y. Chang, Shih-Wei Liao:
RelGAN: Multi-Domain Image-to-Image Translation via Relative Attributes. 5913-5921

Ruizheng Wu, Xin Tao, Xiaodong Gu, Xiaoyong Shen, Jiaya Jia:
Attribute-Driven Spontaneous Motion in Unpaired Image Translation. 5922-5931

Caroline Chan, Shiry Ginosar, Tinghui Zhou, Alexei A. Efros:
Everybody Dance Now. 5932-5941

Yulun Zhang, Chen Fang, Yilin Wang, Zhaowen Wang, Zhe Lin, Yun Fu, Jimei Yang:
Multimodal Style Transfer via Graph Cuts. 5942-5950

Ming Lu, Hao Zhao, Anbang Yao, Yurong Chen, Feng Xu, Li Zhang:
A Closed-Form Solution to Universal Style Transfer. 5951-5960

Jingyuan Li, Fengxiang He, Lefei Zhang, Bo Du, Dacheng Tao:
Progressive Reconstruction of Visual Structure for Image Inpainting. 5961-5970

Oral 3.2A

Recognition, Detection, & Re-Identification

Samarth Sinha, Sayna Ebrahimi, Trevor Darrell:
Variational Adversarial Active Learning. 5971-5980

Yang Zou, Zhiding Yu, Xiaofeng Liu, B. V. K. Vijaya Kumar, Jinsong Wang:
Confidence Regularized Self-Training. 5981-5990

Serim Ryou, Seong-Gyun Jeong, Pietro Perona:
Anchor Loss: Modulating Loss Scale Based on Prediction Difficulty. 5991-6000

Chengxu Zhuang, Alex Lin Zhai, Daniel Yamins:
Local Aggregation for Unsupervised Learning of Visual Embeddings. 6001-6011

Zhennan Wang, Wenbin Zou, Chen Xu:
PR Product: A Substitute for Inner Product in Neural Networks. 6012-6021

Sangdoo Yun, Dongyoon Han, Sanghyuk Chun, Seong Joon Oh, Youngjoon Yoo, Junsuk Choe:
CutMix: Regularization Strategy to Train Strong Classifiers With Localizable Features. 6022-6031

Tianfu Wu, Xi Song:
Towards Interpretable Object Detection by Unfolding Latent Structures. 6032-6042

Jason Kuen, Federico Perazzi, Zhe L. Lin, Jianming Zhang, Yap-Peng Tan:
Scaling Object Detection by Transferring Classification Weights. 6043-6052

Yanghao Li, Yuntao Chen, Naiyan Wang, Zhaoxiang Zhang:
Scale-Aware Trident Networks for Object Detection. 6053-6062

Satoshi Kosugi, Toshihiko Yamasaki, Kiyoharu Aizawa:
Object-Aware Instance Labeling for Weakly Supervised Object Detection. 6063-6071

Lanlan Liu, Michael Muelly, Jia Deng, Tomas Pfister, Li-Jia Li:
Generative Modeling for Small-Data Object Detection. 6072-6080

Shafin Rahman, Salman H. Khan, Nick Barnes:
Transductive Learning for Zero-Shot Object Detection. 6081-6090

Seunghyeon Kim, Jaehoon Choi, Taekyung Kim, Changick Kim:
Self-Training and Adversarial Background Regularization for Unsupervised Domain Adaptive One-Stage Object Detection. 6091-6100

Suichan Li, Dapeng Chen, Bin Liu, Nenghai Yu, Rui Zhao:
Memory-Based Neighbourhood Embedding for Visual Recognition. 6101-6110

Yang Fu, Yunchao Wei, Guanshuo Wang, Yuqian Zhou, Honghui Shi, Thomas S. Huang:
Self-Similarity Grouping: A Simple Unsupervised Cross Domain Adaptation Approach for Person Re-Identification. 6111-6120

Zimo Liu, Jingya Wang, Shaogang Gong, Dacheng Tao, Huchuan Lu:
Deep Reinforcement Active Learning for Human-in-the-Loop Person Re-Identification. 6121-6130

Pirazh Khorramshahi, Amit Kumar, Neehar Peri, Sai Saketh Rambhatla, Jun-Cheng Chen, Rama Chellappa:
A Dual-Path Model With Adaptive Attention for Vehicle Re-Identification. 6131-6140

Zhiheng Ma, Xing Wei, Xiaopeng Hong, Yihong Gong:
Bayesian Loss for Crowd Count Estimation With Point Supervision. 6141-6150

Zhi-Qi Cheng, Jun-Xiu Li, Qi Dai, Xiao Wu, Alexander G. Hauptmann:
Learning Spatial Awareness to Improve Crowd Counting. 6151-6160

Oral 3.2B

Video & Action Understanding

Peixia Li, Boyu Chen, Wanli Ouyang, Dong Wang, Xiaoyun Yang, Huchuan Lu:
GradNet: Gradient-Guided Network for Visual Object Tracking. 6161-6170

Peng Chu, Haibin Ling:
FAMNet: Joint Learning of Feature, Affinity and Multi-Dimensional Assignment for Online Multiple Object Tracking. 6171-6180

Goutam Bhat, Martin Danelljan, Luc Van Gool, Radu Timofte:
Learning Discriminative Model Prediction for Tracking. 6181-6190

Ali Diba, Vivek Sharma, Luc Van Gool, Rainer Stiefelhagen:
DynamoNet: Dynamic Action and Motion Network. 6191-6200

Christoph Feichtenhofer, Haoqi Fan, Jitendra Malik, Kaiming He:
SlowFast Networks for Video Recognition. 6201-6210

Lichen Wang, Zhengming Ding, Zhiqiang Tao, Yunyu Liu, Yun Fu:
Generative Multi-View Human Action Recognition. 6211-6220

Wenhao Wu, Dongliang He, Xiao Tan, Shifeng Chen, Shilei Wen:
Multi-Agent Reinforcement Learning Based Frame Sampling for Effective Untrimmed Video Recognition. 6221-6230

Bruno Korbar, Du Tran, Lorenzo Torresani:
SCSampler: Sampling Salient Clips From Video for Efficient Action Recognition. 6231-6241

Jun Li, Peng Lei, Sinisa Todorovic:
Weakly Supervised Energy-Based Learning for Action Segmentation. 6242-6250

Antonino Furnari, Giovanni Maria Farinella:
What Would You Expect? Anticipating Egocentric Actions With Rolling-Unrolling LSTMs and Modality Attention. 6251-6260

Amir Rasouli, Iuliia Kotseruba, Toni Kunic, John K. Tsotsos:
PIE: A Large-Scale Dataset and Models for Pedestrian Intention Estimation and Trajectory Prediction. 6261-6270

Yingfan Huang, Huikun Bi, Zhaoxin Li, Tianlu Mao, Zhaoqi Wang:
STGAT: Modeling Spatial-Temporal Interactions for Human Trajectory Prediction. 6271-6280

Khoi-Nguyen C. Mac, Dhiraj Joshi, Raymond A. Yeh, Jinjun Xiong, Rogério Schmidt Feris, Minh N. Do:
Learning Motion in Feature Space: Locally-Consistent Deformable Convolution Networks for Fine-Grained Action Detection. 6281-6290

Yu Wu, Linchao Zhu, Yan Yan, Yi Yang:
Dual Attention Matching for Audio-Visual Event Localization. 6291-6299

Mahesh Subedar, Ranganath Krishnan, Paulo Lopez-Meyer, Omesh Tickoo, Jonathan Huang:
Uncertainty-Aware Audiovisual Activity Recognition Using Deep Bayesian Variational Inference. 6300-6309

Canmiao Fu, Wenjie Pei, Qiong Cao, Chaopeng Zhang, Yong Zhao, Xiaoyong Shen, Yu-Wing Tai:
Non-Local Recurrent Neural Memory for Supervised Sequence Modeling. 6310-6319

Min-Hung Chen, Zsolt Kira, Ghassan Alregib, Jaekwon Yoo, Ruxin Chen, Jian Zheng:
Temporal Attentive Alignment for Large-Scale Video Domain Adaptation. 6320-6329

Jiahui Pan, Jibin Gao, Wei-Shi Zheng:
Action Assessment by Joint Relation Graphs. 6330-6339

Ehsan Elhamifar, Zwe Naing:
Unsupervised Procedure Learning via Joint Dynamic Summarization. 6340-6349

Giorgos Kordopatis-Zilos, Symeon Papadopoulos, Ioannis Patras, Yiannis Kompatsiaris:
ViSiL: Fine-Grained Spatio-Temporal Video Similarity Learning. 6350-6359

Poster 3.2

Deep Learning

James Thewlis, Samuel Albanie, Hakan Bilen, Andrea Vedaldi:
Unsupervised Learning of Landmarks by Descriptor Vector Exchange. 6360-6370

Pavel Tokmakov, Yu-Xiong Wang, Martial Hebert:
Learning Compositional Representations for Few-Shot Recognition. 6371-6380

Kanglin Liu, Guoping Qiu, Wenming Tang, Fei Zhou:
Spectral Regularization for Combating Mode Collapse in GANs. 6381-6389

Priya Goyal, Dhruv Mahajan, Abhinav Gupta, Ishan Misra:
Scaling and Benchmarking Self-Supervised Visual Representation Learning. 6390-6399

Riccardo Spezialetti, Samuele Salti, Luigi Di Stefano:
Learning an Effective Equivariant 3D Descriptor Without Supervision. 6400-6409

Hugues Thomas, Charles R. Qi, Jean-Emmanuel Deschaud, Beatriz Marcotegui, François Goulette, Leonidas J. Guibas:
KPConv: Flexible and Deformable Convolution for Point Clouds. 6410-6419

Abdelaziz Djelouah, Joaquim Campos, Simone Schaub-Meyer, Christopher Schroers:
Neural Inter-Frame Compression for Video Coding. 6420-6428

Alessandro Achille, Michael Lam, Rahul Tewari, Avinash Ravichandran, Subhransu Maji, Charless C. Fowlkes, Stefano Soatto, Pietro Perona:
Task2Vec: Task Embedding for Meta-Learning. 6429-6438

Linxiao Yang, Ngai-Man Cheung, Jiaying Li, Jun Fang:
Deep Clustering by Gaussian Mixture Variational Autoencoders With Graph Embedding. 6439-6448

Qi Qian, Lei Shang, Baigui Sun, Juhua Hu, Tacoma Tacoma, Hao Li, Rong Jin:
SoftTriple Loss: Deep Metric Learning Without Triplet Sampling. 6449-6457

Fariborz Taherkhani, Hadi Kazemi, Ali Dabouei, Jeremy M. Dawson, Nasser M. Nasrabadi:
A Weakly Supervised Fine Label Classifier Enhanced by Coarse Supervision. 6458-6467

Munawar Hayat, Salman H. Khan, Syed Waqas Zamir, Jianbing Shen, Ling Shao:
Gaussian Affinity for Max-Margin Class Imbalanced Learning. 6468-6478

Jingjia Huang, Zhangheng Li, Nannan Li, Shan Liu, Ge Li:
AttPool: Towards Hierarchical Feature Representation in Graph Convolutional Networks via Attention Mechanism. 6479-6488

Baosheng Yu, Dacheng Tao:
Deep Metric Learning With Tuplet Margin Loss. 6489-6498

Yogesh Balaji, Rama Chellappa, Soheil Feizi:
Normalized Wasserstein for Mixture Distributions With Applications in Adversarial Learning and Domain Adaptation. 6499-6507

Jiequan Cui, Pengguang Chen, Ruiyu Li, Shu Liu, Xiaoyong Shen, Jiaya Jia:
Fast and Practical Neural Architecture Search. 6508-6517

Jiwoong Park, Minsik Lee, Hyung Jin Chang, Kyuewang Lee, Jin Young Choi:
Symmetric Graph Convolutional Autoencoder for Unsupervised Graph Representation Learning. 6518-6527

Chanho Ahn, Eunwoo Kim, Songhwai Oh:
Deep Elastic Networks With Model Selection for Multi-Task Learning. 6528-6537

Pierre Jacob, David Picard, Aymeric Histace, Edouard Klein:
Metric Learning With HORDE: High-Order Regularizer for Deep Embeddings. 6538-6547

Yaoyao Zhong, Weihong Deng:
Adversarial Learning With Margin-Based Triplet Embedding Regularization. 6548-6557

Recognition

Ahmed Samy Nassar, Sébastien Lefèvre, Jan Dirk Wegner:
Simultaneous Multi-View Instance Detection With Learned Geometric Soft-Constraints. 6558-6567

Kaiwen Duan, Song Bai, Lingxi Xie, Honggang Qi, Qingming Huang, Qi Tian:
CenterNet: Keypoint Triplets for Object Detection. 6568-6577

Chen Lin, Minghao Guo, Chuming Li, Xin Yuan, Wei Wu, Junjie Yan, Dahua Lin, Wanli Ouyang:
Online Hyper-Parameter Learning for Auto-Augmentation Strategy. 6578-6587

Haolan Xue, Chang Liu, Fang Wan, Jianbin Jiao, Xiangyang Ji, Qixiang Ye:
DANet: Divergent Activation for Weakly Supervised Object Localization. 6588-6597

Yao Ding, Yanzhao Zhou, Yi Zhu, Qixiang Ye, Jianbin Jiao:
Selective Sparse Sampling for Fine-Grained Image Recognition. 6598-6607

Shuai Li, Lingxiao Yang, Jianqiang Huang, Xian-Sheng Hua, Lei Zhang:
Dynamic Anchor Feature Selection for Single-Shot Object Detection. 6608-6617

Ye Xiang, Ying Fu, Pan Ji, Hua Huang:
Incremental Learning Using Conditional Adversarial Networks. 6618-6627

Jianyu Wang, Haichao Zhang:
Bilateral Adversarial Training: Towards Fast Training of More Robust Models Against Adversarial Attacks. 6628-6637

Fangyi Liu, Lei Zhang:
View Confusion Feature Learning for Person Re-Identification. 6638-6647

Hang Xu, Lewei Yao, Zhenguo Li, Xiaodan Liang, Wei Zhang:
Auto-FPN: Automatic Network Architecture Adaptation for Object Detection Beyond Classification. 6648-6657

Ziyang Wu, Yuwei Li, Lihua Guo, Kui Jia:
PARN: Position-Aware Relation Networks for Few-Shot Learning. 6658-6666

Zhenwei He, Lei Zhang:
Multi-Adversarial Faster-RCNN for Unrestricted Object Detection. 6667-6676

Hanming Deng, Yang Hua, Tao Song, Zongpu Zhang, Zhengui Xue, Ruhui Ma, Neil Martin Robertson, Haibing Guan:
Object Guided External Memory Network for Video Object Detection. 6677-6686

Xizhou Zhu, Dazhi Cheng, Zheng Zhang, Stephen Lin, Jifeng Dai:
An Empirical Study of Spatial Attention Mechanisms in Deep Networks. 6687-6696

Yang Liu, Jishun Guo, Deng Cai, Xiaofei He:
Attribute Attention for Semantic Disambiguation in Zero-Shot Learning. 6697-6706

Puneet Gupta, Esa Rahtu:
CIIDefence: Defeating Adversarial Attacks by Fusing Class-Specific Image Inpainting and Image Denoising. 6707-6716

Zheng Qin, Zeming Li, Zhaoning Zhang, Yiping Bao, Gang Yu, Yuxing Peng, Jian Sun:
ThunderNet: Towards Real-Time Generic Object Detection on Mobile Devices. 6717-6726

Zhanghan Ke, Daoye Wang, Qiong Yan, Jimmy S. J. Ren, Rynson W. H. Lau:
Dual Student: Breaking the Limits of the Teacher in Semi-Supervised Learning. 6727-6735

Han Sun, Zhiyuan Chen, Shiyang Yan, Lin Xu:
MVP Matching: A Maximum-Value Perfect Matching for Mining Hard Samples, With Application to Person Re-Identification. 6736-6746

Segmentation, Grouping, & Shape

Jun Fu, Jing Liu, Yuhang Wang, Yong Li, Yongjun Bao, Jinhui Tang, Hanqing Lu:
Adaptive Context Network for Scene Parsing. 6747-6756

Qing Lian, Lixin Duan, Fengmao Lv, Boqing Gong:
Constructing Self-Motivated Pyramid Curriculums for Cross-Domain Semantic Segmentation: A Non-Adversarial Approach. 6757-6766

Huikai Wu, Junge Zhang, Kaiqi Huang:
SparseMask: Differentiable Connectivity Learning for Dense Image Prediction. 6767-6776

Yawei Luo, Ping Liu, Tao Guan, Junqing Yu, Yi Yang:
Significance-Aware Information Bottleneck for Domain Adaptive Semantic Segmentation. 6777-6786

Anran Zhang, Jiayi Shen, Zehao Xiao, Fan Zhu, Xiantong Zhen, Xianbin Cao, Ling Shao:
Relational Attention Network for Crowd Counting. 6787-6796

Fan Zhang, Yanqin Chen, Zhihang Li, Zhibin Hong, Jingtuo Liu, Feifei Ma, Junyu Han, Errui Ding:
ACFNet: Attentional Class Feature Network for Semantic Segmentation. 6797-6806

Jungbeom Lee, Eunji Kim, Sungmin Lee, Jangho Lee, Sungroh Yoon:
Frame-to-Frame Aggregation of Active Regions in Web Videos for Weakly Supervised Semantic Segmentation. 6807-6817

Henghui Ding, Xudong Jiang, Ai Qun Liu, Nadia Magnenat-Thalmann, Gang Wang:
Boundary-Aware Feature Propagation for Scene Segmentation. 6818-6828

Jaehoon Choi, Taekyung Kim, Changick Kim:
Self-Ensembling With GAN-Based Data Augmentation for Domain Adaptation in Semantic Segmentation. 6829-6839

3D From Single View & RGBD

Fabian Manhardt, Diego Martín Arroyo, Christian Rupprecht, Benjamin Busam, Tolga Birdal, Nassir Navab, Federico Tombari:
Explaining the Ambiguity of Object Detection and 6D Pose From Visual Data. 6840-6849

Xinzhu Ma, Zhihui Wang, Haojie Li, Pengbo Zhang, Wanli Ouyang, Xin Fan:
Accurate Monocular 3D Object Detection via Color-Embedded 3D Reconstruction for Autonomous Driving. 6850-6859

Lorenzo Bertoni, Sven Kreiss, Alexandre Alahi:
MonoLoco: Monocular 3D Pedestrian Localization and Uncertainty Estimation. 6860-6870

Junsheng Zhou, Yuwang Wang, Kaihuai Qin, Wenjun Zeng:
Unsupervised High-Resolution Depth Learning From Videos With Dual Networks. 6871-6880

Face & Body

Rui Zhao, Kang Wang, Hui Su, Qiang Ji:
Bayesian Graph Convolution LSTM for Skeleton Based Action Recognition. 6881-6891

Arnaud Dapogny, Matthieu Cord, Kevin Bailly:
DeCaFA: Deep Convolutional Cascade for Face Alignment in the Wild. 6892-6900

Yichun Shi, Anil K. Jain:
Probabilistic Face Embeddings. 6901-6910

Petr Kellnhofer, Adrià Recasens, Simon Stent, Wojciech Matusik, Antonio Torralba:
Gaze360: Physically Unconstrained Gaze Estimation in the Wild. 6911-6920

Ancong Wu, Wei-Shi Zheng, Jian-Huang Lai:
Unsupervised Person Re-Identification by Camera-Aware Similarity Consistency Learning. 6921-6930

Zhe He, Adrian Spurr, Xucong Zhang, Otmar Hilliges:
Photo-Realistic Monocular Gaze Redirection Using Generative Adversarial Networks. 6931-6940

Xuecheng Nie, Yuncheng Li, Linjie Luo, Ning Zhang, Jiashi Feng:
Dynamic Kernel Distillation for Efficient Pose Estimation in Videos. 6941-6949

Xuecheng Nie, Jiashi Feng, Jianfeng Zhang, Shuicheng Yan:
Single-Stage Multi-Person Pose Machines. 6950-6959

Yujin Chen, Zhigang Tu, Liuhao Ge, Dejun Zhang, Ruizhi Chen, Junsong Yuan:
SO-HandNet: Self-Organizing Network for 3D Hand Pose Estimation With Semi-Supervised Learning. 6960-6969

Xinyao Wang, Liefeng Bo, Fuxin Li:
Adaptive Wing Loss for Robust Face Alignment via Heatmap Regression. 6970-6980

Gines Hidalgo Martinez, Yaadhav Raaj, Haroon Idrees, Donglai Xiang, Hanbyul Joo, Tomas Simon, Yaser Sheikh:
Single-Network Whole-Body Pose Estimation. 6981-6990

Lisha Chen, Hui Su, Qiang Ji:
Face Alignment With Kernel Density Deep Neural Network. 6991-7001

Action & Video

He Zhao, Rick Wildes:
Spatiotemporal Feature Residual Propagation for Action Prediction. 7002-7011

Fanyi Xiao, Haotian Liu, Yong Jae Lee:
Identity From Here, Pose From There: Self-Supervised Disentanglement and Generation of Objects Using Unlabeled Videos. 7012-7021

Jiajun Deng, Yingwei Pan, Ting Yao, Wengang Zhou, Houqiang Li, Tao Mei:
Relation Distillation Networks for Video Object Detection. 7022-7031

AmirHossein Habibian, Ties van Rozendaal, Jakub M. Tomczak, Taco Cohen:
Video Compression With Rate-Distortion Autoencoders. 7032-7041

Yi Xu, Longwen Gao, Kai Tian, Shuigeng Zhou, Huyang Sun:
Non-Local ConvLSTM for Video Compression Artifact Reduction. 7042-7051

Chuang Gan, Hang Zhao, Peihao Chen, David D. Cox, Antonio Torralba:
Self-Supervised Moving Vehicle Tracking With Stereo Sound. 7052-7061

Yuhua Chen, Cordelia Schmid, Cristian Sminchisescu:
Self-Supervised Learning With Geometric Constraints in Monocular Video: Connecting Flow, Depth, and Camera. 7062-7071

Jingwei Ji, Kaidi Cao, Juan Carlos Niebles:
Learning Temporal Action Proposals With Fewer Labels. 7072-7081

Ji Lin, Chuang Gan, Song Han:
TSM: Temporal Shift Module for Efficient Video Understanding. 7082-7092

Runhao Zeng, Wenbing Huang, Chuang Gan, Mingkui Tan, Yu Rong, Peilin Zhao, Junzhou Huang:
Graph Convolutional Networks for Temporal Action Localization. 7093-7102

Shiyao Wang, Hongchao Lu, Zhidong Deng:
Fast Object Detection in Compressed Video. 7103-7112

Motion & Tracking

Jason Y. Zhang, Panna Felsen, Angjoo Kanazawa, Jitendra Malik:
Predicting 3D Human Dynamics From Video. 7113-7122

Borui Wang, Ehsan Adeli, Hsu-Kuang Chiu, De-An Huang, Juan Carlos Niebles:
Imitation Learning for Human Pose Prediction. 7123-7132

Alejandro Hernandez Ruiz, Jürgen Gall, Francesc Moreno:
Human Motion Prediction via Spatio-Temporal Inpainting. 7133-7142

Emre Aksan, Manuel Kaufmann, Otmar Hilliges:
Structured Prediction Helps 3D Human Motion Modelling. 7143-7152

Computational Photography & Graphics

Kyle Genova, Forrester Cole, Daniel Vlasic, Aaron Sarna, William T. Freeman, Thomas A. Funkhouser:
Learning Shape Templates With Structured Implicit Functions. 7153-7163

Bingyao Huang, Haibin Ling:
CompenNet++: End-to-End Full Projector Compensation. 7164-7173

Marc-André Gardner, Yannick Hold-Geoffroy, Kalyan Sunkavalli, Christian Gagné, Jean-François Lalonde:
Deep Parametric Indoor Lighting Estimation. 7174-7182

Yuval Nirkin, Yosi Keller, Tal Hassner:
FSGAN: Subject Agnostic Face Swapping and Reenactment. 7183-7192

Hao Zhou, Sunil Hadap, Kalyan Sunkavalli, David Jacobs:
Deep Single-Image Portrait Relighting. 7193-7201

Ruihui Li, Xianzhi Li, Chi-Wing Fu, Daniel Cohen-Or, Pheng-Ann Heng:
PU-GAN: A Point Cloud Upsampling Adversarial Network. 7202-7211

Giorgos Bouritsas, Sergiy Bokhnyak, Stylianos Ploumpis, Stefanos Zafeiriou, Michael M. Bronstein:
Neural 3D Morphable Models: Spiral Convolutional Networks for 3D Shape Representation Learning and Generation. 7212-7221

Low-Level & Optimization

Yu Zeng, Yun-Zhi Zhuge, Huchuan Lu, Lihe Zhang:
Joint Learning of Saliency Detection and Weakly Supervised Semantic Segmentation. 7222-7232

Yi Zeng, Pingping Zhang, Zhe L. Lin, Jianming Zhang, Huchuan Lu:
Towards High-Resolution Salient Object Detection. 7233-7242

Timo Stoffregen, Guillermo Gallego, Tom Drummond, Lindsay Kleeman, Davide Scaramuzza:
Event-Based Motion Segmentation by Motion Compensation. 7243-7252

Yongri Piao, Wei Ji, Jingjing Li, Miao Zhang, Huchuan Lu:
Depth-Induced Multi-Scale Recurrent Attention Network for Saliency Detection. 7253-7262

Zhe Wu, Li Su, Qingming Huang:
Stacked Cross Refinement Network for Edge-Aware Salient Object Detection. 7263-7272

Haofeng Li, Guanqi Chen, Guanbin Li, Yizhou Yu:
Motion Guided Attention for Video Salient Object Detection. 7273-7282

Pengxiang Yan, Guanbin Li, Yuan Xie, Zhen Li, Chuan Wang, Tianshui Chen, Liang Lin:
Semi-Supervised Video Salient Object Detection Using Pseudo-Labels. 7283-7292

Sangryul Jeon, Dongbo Min, Seungryong Kim, Kwanghoon Sohn:
Joint Learning of Semantic Alignment and Object Landmark Detection. 7293-7302

Ruoteng Li, Robby T. Tan, Loong Fah Cheong, Angelica I. Avilés-Rivero, Qingnan Fan, Carola Schönlieb:
RainFlow: Optical Flow Under Rain Streaks and Rain Veiling Effect. 7303-7312

Xiaohong Liu, Yongrui Ma, Zhihao Shi, Jun Chen:
GridDehazeNet: Attention-Based Multi-Scale Network for Image Dehazing. 7313-7322

Haiyang Jiang, Yinqiang Zheng:
Learning to See Moving Objects in the Dark. 7323-7332

Scene Understanding

Jyh-Jing Hwang, Stella X. Yu, Jianbo Shi, Maxwell D. Collins, Tien-Ju Yang, Xiao Zhang, Liang-Chieh Chen:
SegSort: Segmentation by Discriminative Sorting of Segments. 7333-7343

Keng-Chi Liu, Yi-Ting Shen, Jan Klopp, Liang-Gee Chen:
What Synthesis Is Missing: Depth Adaptation Integrated With Weak Supervision for Indoor Scene Parsing. 7344-7353

Konstantin Sofiiuk, Olga Barinova, Anton Konushin:
AdaptIS: Adaptive Instance Selection Network. 7354-7362

Tuan-Hung Vu, Himalaya Jain, Maxime Bucher, Matthieu Cord, Patrick Pérez:
DADA: Depth-Aware Domain Adaptation in Semantic Segmentation. 7363-7372

Christos Sakaridis, Dengxin Dai, Luc Van Gool:
Guided Curriculum Model Adaptation and Uncertainty-Aware Evaluation for Semantic Nighttime Image Segmentation. 7373-7382

Yang Zhou, Zachary While, Evangelos Kalogerakis:
SceneGraphNet: Neural Message Passing for 3D Indoor Scene Augmentation. 7383-7391

Seyed Majid Azimi, Corentin Henry, Lars Sommer, Arne Schumann, Eleonora Vig:
SkyScapes - Fine-Grained Semantic Understanding of Aerial Scenes. 7392-7402

Language & Reasoning

Haoshuo Huang, Vihan Jain, Harsh Mehta, Alexander Ku, Gabriel Magalhães, Jason Baldridge, Eugene Ie:
Transferable Representation Learning in Vision-and-Language Navigation. 7403-7412

Iro Laina, Christian Rupprecht, Nassir Navab:
Towards Unsupervised Image Captioning With Shared Multimodal Embeddings. 7413-7423

Tanmay Gupta, Alexander G. Schwing, Derek Hoiem:
ViCo: Word Embeddings From Visual Co-Occurrences. 7424-7433

Boren Li, Boyu Zhuang, Mingyang Li, Jian Gu:
Seq-SG2SL: Inferring Semantic Layout From Scene Graph Through Sequence to Sequence Learning. 7434-7442

Badri N. Patro, Mayank Lunayach, Shivansh Patel, Vinay P. Namboodiri:
U-CAM: Visual Explanation Using Uncertainty Based Class Activation Maps. 7443-7452

Ding-Jie Chen, Songhao Jia, Yi-Chen Lo, Hwann-Tzong Chen, Tyng-Luh Liu:
See-Through-Text Grouping for Referring Image Segmentation. 7453-7462

Chen Sun, Austin Myers, Carl Vondrick, Kevin Murphy, Cordelia Schmid:
VideoBERT: A Joint Model for Video and Language Representation Learning. 7463-7472

Andrea Burns, Reuben Tan, Kate Saenko, Stan Sclaroff, Bryan A. Plummer:
Language Features Matter: Effective Language Representations for Vision-Language Tasks. 7473-7482

3D From Multiview & Sensors

Zhenyao Wu, Xinyi Wu, Xiaoping Zhang, Song Wang, Lili Ju:
Semantic Stereo Matching With Pyramid Cost Volumes. 7483-7492

Zhenyao Wu, Xinyi Wu, Xiaoping Zhang, Song Wang, Lili Ju:
Spatial Correspondence With Generative Adversarial Network: Learning Depth From Monocular Videos. 7493-7503

Ze Yang, Liwei Wang:
Learning Relationships for Multi-View 3D Object Recognition. 7504-7513

Xinwei He, Tengteng Huang, Song Bai, Xiang Bai:
View N-Gram Network for 3D Object Retrieval. 7514-7523

Eric Brachmann, Carsten Rother:
Expert Sample Consensus Applied to Camera Re-Localization. 7524-7533

Yutong Bai, Qing Liu, Lingxi Xie, Yan Zheng, Weichao Qiu, Alan L. Yuille:
Semantic Part Detection via Matching: Learning to Generalize to Novel Viewpoints From Limited Training Data. 7534-7544

Jinxian Liu, Bingbing Ni, Caiyuan Li, Jiancheng Yang, Qi Tian:
Dynamic Points Agglomeration for Hierarchical Point Sets Learning. 7545-7554

Image & Video Synthesis

Ning Yu, Larry Davis, Mario Fritz:
Attributing Fake Images to GANs: Learning and Analyzing GAN Fingerprints. 7555-7565

Qicheng Lao, Mohammad Havaei, Ahmad Pesaranghader, Francis Dutil, Lisa Di-Jorio, Thomas Fevens:
Dual Adversarial Inference for Text-to-Image Synthesis. 7566-7575

Mohamed Ilyes Lakhal, Oswald Lanz, Andrea Cavallaro:
View-LSTM: Novel-View Video Synthesis Through View Decomposition. 7576-7586

Thu Nguyen-Phuoc, Chuan Li, Lucas Theis, Christian Richardt, Yong-Liang Yang:
HoloGAN: Unsupervised Learning of 3D Representations From Natural Images. 7587-7596

Shuang Ma, Daniel McDuff, Yale Song:
Unpaired Image-to-Speech Synthesis With Multimodal Information Bottleneck. 7597-7606

Lluís Castrejón, Nicolas Ballas, Aaron C. Courville:
Improved Conditional VRNNs for Video Prediction. 7607-7616

Xiaosheng Yan, Yuanlong Yu, Feigege Wang, Wenxi Liu, Shengfeng He, Jia Pan:
Visualizing the Invisible: Occluded Vehicle Segmentation and Recovery. 7617-7626

Oral 4.1A

Single-View 3D Modeling, Pose Estimation

Rahul Garg, Neal Wadhwa, Sameer Ansari, Jonathan T. Barron:
Learning Single Camera Depth Estimation Using Dual-Pixels. 7627-7636

Pedro O. Pinheiro, Negar Rostamzadeh, Sungjin Ahn:
Domain-Adaptive Single-View 3D Reconstruction. 7637-7646

Kyle Olszewski, Sergey Tulyakov, Oliver J. Woodford, Hao Li, Linjie Luo:
Transformable Bottleneck Networks. 7647-7656

Johanna Wald, Armen Avetisyan, Nassir Navab, Federico Tombari, Matthias Nießner:

RIO: 3D Object Instance Re-Localization in Changing Indoor Environments. 7657-7666

Kiru Park, Timothy Patten, Markus Vincze:
Pix2Pose: Pixel-Wise Coordinate Regression of Objects for 6D Pose Estimation. 7667-7676

Zhigang Li, Gu Wang, Xiangyang Ji:
CDPN: Coordinates-Based Disentangled Pose Network for Real-Time RGB-Based 6-DoF Object Pose Estimation. 7677-7686

David Novotný, Nikhila Ravi, Benjamin Graham, Natalia Neverova, Andrea Vedaldi:
C3DPO: Canonical 3D Pose Networks for Non-Rigid Structure From Motion. 7687-7696

Yichao Zhou, Haozhi Qi, Yuexiang Zhai, Qi Sun, Zhili Chen, Li-Yi Wei, Yi Ma:
Learning to Reconstruct 3D Manhattan Wireframes From a Single Image. 7697-7706

Shichen Liu, Weikai Chen, Tianye Li, Hao Li:
Soft Rasterizer: A Differentiable Renderer for Image-Based 3D Reasoning. 7707-7716

Karim Iskakov, Egor Burkov, Victor S. Lempitsky, Yury Malkov:
Learnable Triangulation of Human Pose. 7717-7726

Denis Tomè, Patrick Peluse, Lourdes Agapito, Hernán Badino:
xR-EgoPose: Egocentric 3D Human Pose From an HMD Camera. 7727-7737

Zerong Zheng, Tao Yu, Yixuan Wei, Qionghai Dai, Yebin Liu:
DeepHuman: 3D Human Reconstruction From a Single Image. 7738-7748

Sicong Tang, Feitong Tan, Kelvin Cheng, Zhaoyang Li, Siyu Zhu, Ping Tan:
A Neural Network for Detailed Human Depth Estimation From a Single Image. 7749-7758

Yuanlu Xu, Song-Chun Zhu, Tony Tung:
DenseRaC: Joint 3D Pose and Shape Estimation by Dense Render-and-Compare. 7759-7769

Jue Wang, Shaoli Huang, Xinchao Wang, Dacheng Tao:
Not All Parts Are Created Equal: 3D Pose Estimation by Modeling Bi-Directional Dependencies of Body Parts. 7770-7779

Oral 4.1B

Computational Photography

Inchang Choi, Orazio Gallo, Alejandro J. Troccoli, Min H. Kim, Jan Kautz:
Extreme View Synthesis. 7780-7789

Xiaogang Xu, Ying-Cong Chen, Jiaya Jia:
View Independent Generative Adversarial Network for Novel View Synthesis. 7790-7799

Pingping Zhang, Wei Liu, Yinjie Lei, Huchuan Lu, Xiaoyun Yang:
Cascaded Context Pyramid for Full-Resolution 3D Semantic Scene Completion. 7800-7809

Numair Khan, Qian Zhang, Lucas Kasser, Henry Stone, Min H. Kim, James Tompkin:
View-Consistent 4D Light Field Superpixel Segmentation. 7810-7818

Hao Zhou, Xiang Yu, David Jacobs:
GLoSH: Global-Local Spherical Harmonics for Intrinsic Image Decomposition. 7819-7828

Satoshi Murai, Meng-Yu Kuo, Ryo Kawahara, Shohei Nobuhara, Ko Nishino:
Surface Normals and Shape From Water. 7829-7837

Jerin Geo James, Pranay Agrawal, Ajit Rajwade:
Restoration of Non-Rigidly Distorted Underwater Images Using a Combination of Compressive Sensing and Local Polynomial Image Representations. 7838-7847

Yajie Zhao, Zeng Huang, Tianye Li, Weikai Chen, Chloe LeGendre, Xinglei Ren, Ari Shapiro, Hao Li:
Learning Perspective Undistortion of Portraits. 7848-7858

Salman Siddique Khan, Adarsh V. R, Vivek Boominathan, Jasper Tan, Ashok Veeraraghavan, Kaushik Mitra:
Towards Photorealistic Reconstruction of Highly Multiplexed Lensless Images. 7859-7868

Mahesh Mohan M. R., Sharath Girish, Rajagopalan Ambasamudram:
Unconstrained Motion Deblurring for Dual-Lens Cameras. 7869-7878

Jongho Lee, Mohit Gupta:
Stochastic Exposure Coding for Handling Multi-ToF-Camera Interference. 7879-7887

Byeongjoo Ahn, Akshat Dave, Ashok Veeraraghavan, Ioannis Gkioulekas, Aswin C. Sankaranarayanan:
Convolutional Approximations to the General Non-Line-of-Sight Imaging Operator. 7888-7898

Joseph R. Bartels, Jian Wang, William Whittaker, Srinivasa G. Narasimhan:
Agile Depth Sensing Using Triangulation Light Curtains. 7899-7907

Anant Gupta, Atul Ingle, Mohit Gupta:
Asynchronous Single-Photon 3D Imaging. 7908-7917

Poster 4.1

Deep Learning

Yu-Jhe Li, Ci-Siang Lin, Yan-Bo Lin, Yu-Chiang Frank Wang:
Cross-Dataset Person Re-Identification via Unsupervised Pose Disentanglement and Adaptation. 7918-7928

Raphael Gontijo Lopes, David Ha, Douglas Eck, Jonathon Shlens:
A Learned Representation for Scalable Vector Graphics. 7929-7938

Assia Benbihi, Matthieu Geist, Cédric Pradalier:
ELF: Embedded Localisation of Features in Pre-Trained CNN. 7939-7948

Tianyang Xu, Zhen-Hua Feng, Xiao-Jun Wu, Josef Kittler:
Joint Group Feature Selection and Discriminative Filter Learning for Robust Visual Object Tracking. 7949-7959

Jing Lu, Chaofan Xu, Wei Zhang, Lingyu Duan, Tao Mei:
Sampling Wisely: Deep Image Embedding by Top-K Precision Optimization. 7960-7969

Bashir Sadeghi, Runyi Yu, Vishnu Boddeti:
On the Global Optima of Kernelized Adversarial Representation Learning. 7970-7978

Riccardo Volpi, Vittorio Murino:
Addressing Model Vulnerability to Distributional Shifts Over Image Transformation Sets. 7979-7988

Qianyu Feng, Guoliang Kang, Hehe Fan, Yi Yang:
Attract or Distract: Exploit the Margin of Open Set. 7989-7998

Biagio Brattoli, Karsten Roth, Björn Ommer:
MIC: Mining Interclass Characteristics for Improved Metric Learning. 7999-8008

Mohammad Sabokrou, Mohammad Khalooei, Ehsan Adeli:
Self-Supervised Representation Learning via Neighborhood-Relational Encoding. 8009-8018

Mohammad Tavakolian, Hamed Rezazadegan Tavakoli, Abdenour Hadid:
AWSD: Adaptive Weighted Spatiotemporal Distillation for Video Representation. 8019-8028

Pengfei Fang, Jieming Zhou, Soumava Kumar Roy, Lars Petersson, Mehrtash Harandi:
Bilinear Attention Networks for Person Retrieval. 8029-8038

Sanping Zhou, Fei Wang, Zeyi Huang, Jinjun Wang:
Discriminative Feature Learning With Consistent Attention Regularization for Person Re-Identification. 8039-8048

Kuniaki Saito, Donghyun Kim, Stan Sclaroff, Trevor Darrell, Kate Saenko:
Semi-Supervised Domain Adaptation via Minimax Entropy. 8049-8057

Spyros Gidaris, Andrei Bursuc, Nikos Komodakis, Patrick Pérez, Matthieu Cord:
Boosting Few-Shot Visual Learning With Self-Supervision. 8058-8067

Aditya Ganeshan, Vivek B. S., Venkatesh Babu Radhakrishnan:
FDA: Feature Disruptive Attack. 8068-8078

Lei Qi, Lei Wang, Jing Huo, Luping Zhou, Yinghuan Shi, Yang Gao:
A Novel Unsupervised Camera-Aware Domain Adaptation Framework for Person Re-Identification. 8079-8088

Yu-Jhe Li, Yun-Chun Chen, Yen-Yu Lin, Xiaofei Du, Yu-Chiang Frank Wang:
Recover and Identify: A Generative Dual Model for Cross-Resolution Person Re-Identification. 8089-8098

Ang Li, Huiyi Hu, Piotr Mirowski, Mehrdad Farajtabar:
Cross-View Policy Learning for Street Navigation. 8099-8108

Pierluigi Zama Ramirez, Alessio Tonioni, Samuele Salti, Luigi Di Stefano:
Learning Across Tasks and Domains. 8109-8118

Gil Avraham, Yan Zuo, Thanuja Dharmasiri, Tom Drummond:
EMPNet: Neural Localisation and Mapping Using Embedded Memory Points. 8119-8128

Guo-Jun Qi, Liheng Zhang, Chang Wen Chen, Qi Tian:
AVT: Unsupervised Learning of Transformation Equivariant Representations by Autoencoding Variational Transformations. 8129-8138

Anastasia Dubrovina, Fei Xia, Panos Achlioptas, Mira Shalah, Raphaël Groscot, Leonidas J. Guibas:
Composite Shape Modeling via Latent Space Factorization. 8139-8148

Jianlong Wu, Keyu Long, Fei Wang, Chen Qian, Cheng Li, Zhouchen Lin, Hongbin Zha:
Deep Comprehensive Correlation Mining for Image Clustering. 8149-8158

Kaveh Hassani, Mike Haley:
Unsupervised Multi-Task Feature Learning on Point Clouds. 8159-8170

Ruihuang Li, Changqing Zhang, Huazhu Fu, Xi Peng, Joey Tianyi Zhou, Qinghua Hu:
Reciprocal Multi-Layer Subspace Learning for Multi-View Clustering. 8171-8179

Tristan Aumentado-Armstrong, Stavros Tsogkas, Allan D. Jepson, Sven J. Dickinson:
Geometric Disentanglement for Generative Latent Shape Models. 8180-8189

Jogendra Nath Kundu, Maharshi Gor, Dakshit Agrawal, Venkatesh Babu Radhakrishnan:
GAN-Tree: An Incrementally Learned Hierarchical Generative Framework for Multi-Modal Data Distributions. 8190-8199

Jue Wang, Anoop Cherian:
GODS: Generalized One-Class Discriminative Subspaces for Anomaly Detection. 8200-8210

Shuyan Li, Zhixiang Chen, Jiwen Lu, Xiu Li, Jie Zhou:
Neighborhood Preserving Hashing for Scalable Video Retrieval. 8211-8220

Recognition

Xinyu Zhang, Jiewei Cao, Chunhua Shen, Mingyu You:
Self-Training With Progressive Augmentation for Unsupervised Cross-Domain Person Re-Identification. 8221-8230

Chenfeng Xu, Kai Qiu, Jianlong Fu, Song Bai, Yongchao Xu, Xiang Bai:
Learn to Scale: Generating Multipolar Normalized Density Maps for Crowd Counting. 8381-8389

Sudong Cai, Yulan Guo, Salman H. Khan, Jiwei Hu, Gongjian Wen:
Ground-to-Aerial Image Geo-Localization With a Hard Exemplar Reweighting Triplet Loss. 8390-8399

Kai Han, Andrea Vedaldi, Andrew Zisserman:
Learning to Discover Novel Visual Categories via Deep Transfer Clustering. 8400-8408

Chuming Li, Xin Yuan, Chen Lin, Minghao Guo, Wei Wu, Junjie Yan, Wanli Ouyang:
AM-LFS: AutoML for Loss Function Search. 8409-8418

Bingyi Kang, Zhuang Liu, Xin Wang, Fisher Yu, Jiashi Feng, Trevor Darrell:
Few-Shot Object Detection via Feature Reweighting. 8419-8428

Shuai Shao, Zeming Li, Tianyuan Zhang, Chao Peng, Gang Yu, Xiangyu Zhang, Jing Li, Jian Sun:
Objects365: A Large-Scale, High-Quality Dataset for Object Detection. 8429-8438

Wenhai Wang, Enze Xie, Xiaoge Song, Yuhang Zang, Wenjia Wang, Tong Lu, Gang Yu, Chunhua Shen:
Efficient and Accurate Arbitrary-Shaped Text Detection With Pixel Aggregation Network. 8439-8448

Lingxiao He, Yinggang Wang, Wu Liu, He Zhao, Zhenan Sun, Jiashi Feng:
Foreground-Aware Pyramid Reconstruction for Alignment-Free Occluded Person Re-Identification. 8449-8458

Fusheng Hao, Fengxiang He, Jun Cheng, Lei Wang, Jianzhong Cao, Dacheng Tao:
Collect and Select: Semantic Alignment Metric Learning for Few-Shot Learning. 8459-8468

Segmentation, Grouping, & Shape

Roy Uziel, Meitar Ronen, Oren Freifeld:
Bayesian Adaptive Superpixel Segmentation. 8469-8478

Kevin Duarte, Yogesh S. Rawat, Mubarak Shah:
CapsuleVOS: Semi-Supervised Video Object Segmentation Using Capsule Routing. 8479-8488

Zhiqin Chen, Kangxue Yin, Matthew Fisher, Siddhartha Chaudhuri, Hao Zhang:
BAE-NET: Branched Autoencoder for Shape Co-Segmentation. 8489-8498

Hsien-Yu Meng, Lin Gao, Yu-Kun Lai, Dinesh Manocha:
VV-Net: Voxel VAE Net With Group Convolutions for Point Cloud Segmentation. 8499-8507

Huan Wang, Luping Zhou, Lei Wang:
Miss Detection vs. False Alarm: Adversarial Learning for Small Object Segmentation in Infrared Images. 8508-8517

Bo Li, Zhengxing Sun, Qian Li, Yunjie Wu, Anqi Hu:
Group-Wise Deep Object Co-Segmentation With Co-Attention Recurrent Neural Network. 8518-8527

Statistics, Physics, Theory & Datasets

Sen He, Hamed Rezazadegan Tavakoli, Ali Borji, Nicolas Pugeault:
Human Attention in Image Captioning: Dataset and Analysis. 8528-8537

Zhenzhang Ye, Bjoern Haefner, Maolin Gao, Tao Wu, Yvain Quéau, Daniel Cremers:
Variational Uncalibrated Photometric Stereo Under General Lighting. 8538-8547

Qian Zheng, Yiming Jia, Boxin Shi, Xudong Jiang, Lingyu Duan, Alex C. Kot:
SPLINE-Net: Sparse Photometric Stereo Through Lighting Interpolation and Normal Estimation Networks. 8548-8557

Tao Zhang, Ying Fu, Lizhi Wang, Hua Huang:
Hyperspectral Image Reconstruction Using Deep External and Internal Learning. 8558-8567

Didier Bieler, Semih Günel, Pascal Fua, Helge Rhodin:
Gravity as a Reference for Estimating a Person's Height From Video. 8568-8576

Hieu M. Le, Dimitris Samaras:
Shadow Removal via Shadow Image Decomposition. 8577-8586

Ruqi Huang, Marie-Julie Rakotosaona, Panos Achlioptas, Leonidas J. Guibas, Maks Ovsjanikov:
OperatorNet: Recovering 3D Shapes From Difference Operators. 8587-8596

Soumyadip Sengupta, Jinwei Gu, Kihwan Kim, Guilin Liu, David W. Jacobs, Jan Kautz:
Neural Inverse Rendering of an Indoor Scene From a Single Image. 8597-8606

3D From Single View & RGBD

Yida Wang, David Joseph Tan, Nassir Navab, Federico Tombari:
ForkNet: Multi-Branch Volumetric Semantic Completion From a Single Depth Image. 8607-8616

Junsheng Zhou, Yuwang Wang, Kaihuai Qin, Wenjun Zeng:
Moving Indoor: Unsupervised Video Depth Learning in Challenging Environments. 8617-8626

Duc Nguyen, Seonghwa Choi, Woojae Kim, Sanghoon Lee:
GraphX-Convolution for Point Cloud Deformation in 2D-to-3D Conversion. 8627-8636

Jingwei Huang, Yichao Zhou, Thomas A. Funkhouser, Leonidas J. Guibas:
FrameNet: Learning Local Canonical Frames of 3D Surfaces From a Single RGB Image. 8637-8646

Yixin Chen, Siyuan Huang, Tao Yuan, Yixin Zhu, Siyuan Qi, Song-Chun Zhu:
Holistic++ Scene Understanding: Single-View 3D Holistic Scene Parsing and Human Pose Estimation With Human-Object Interaction and Physical Commonsense. 8647-8656

Action & Video

Quan Kong, Ziming Wu, Ziwei Deng, Martin Klinkigt, Bin Tong, Tomokazu Murakami:
MMAct: A Large-Scale Dataset for Cross Modal Human Action Understanding. 8657-8666

Hang Zhao, Antonio Torralba, Lorenzo Torresani, Zhicheng Yan:
HACS: Human Action Clips and Segments Dataset for Recognition and Temporal Localization. 8667-8677

Sanath Narayan, Hisham Cholakkal, Fahad Shahbaz Khan, Ling Shao:
3C-Net: Category Count and Center Loss for Weakly-Supervised Action Localization. 8678-8686

Tushar Nagarajan, Christoph Feichtenhofer, Kristen Grauman:
Grounded Human-Object Interaction Hotspots From Video. 8687-8696

Lei Wang, Piotr Koniusz, Du Huynh:
Hallucinating IDT Descriptors and I3D Optical Flow Features for Action Recognition With CNNs. 8697-8707

Computational Photography & Graphics

Zhewei Huang, Shuchang Zhou, Wen Heng:
Learning to Paint With Model-Based Deep Reinforcement Learning. 8708-8717

Carlo Innamorati, Bryan C. Russell, Danny M. Kaufman, Niloy J. Mitra:
Neural Re-Simulation for Generating Bounces in Single Images. 8718-8727

Maxim Maximov, Tobias Ritschel, Laura Leal-Taixé, Mario Fritz:
Deep Appearance Maps. 8728-8737

Erhan Gundogdu, Victor Constantin, Amrollah Seifoddini, Minh Dang, Mathieu Salzmann, Pascal Fua:
GarNet: A Two-Stream Network for Fast and Accurate 3D Cloth Draping. 8738-8747

Manuel Dahnert, Angela Dai, Leonidas J. Guibas, Matthias Nießner:
Joint Embedding of 3D Scan and CAD Objects. 8748-8757

Nadav Schor, Oren Katzir, Hao Zhang, Daniel Cohen-Or:
CompoNet: Learning to Generate the Unseen by Part Synthesis and Composition. 8758-8767

Chiyu Max Jiang, Dana Lynn Ona Lansigan, Philip Marcus, Matthias Nießner:
DDSL: Deep Differentiable Simplex Layer for Learning Geometric Signals. 8768-8777

Low-Level & Optimization

Jiaxing Zhao, Jiang-Jiang Liu, Deng-Ping Fan, Yang Cao, Jufeng Yang, Ming-Ming Cheng:
EGNet: Edge Guidance Network for Salient Object Detection. 8778-8787

David Berga, Xosé Ramón Fernández-Vidal, Xavier Otazu, Xosé M. Pardo:
SID4VAM: A Benchmark Dataset With Synthetic Images for Visual Attention Modeling. 8788-8797

Haochen Zhang, Dong Liu, Zhiwei Xiong:
Two-Stream Action Recognition-Oriented Video Super-Resolution. 8798-8807

Xin Yang, Haiyang Mei, Ke Xu, Xiaopeng Wei, Baocai Yin, Rynson W. H. Lau:
Where Is My Mirror? 8808-8817

Shaofan Cai, Xiaoshuai Zhang, Haoqiang Fan, Haibin Huang, Jiangyu Liu, Jiaming Liu, Jiaying Liu, Jue Wang, Jian Sun:
Disentangled Image Matting. 8818-8827

Riccardo de Lutio, Stefano D'Aronco, Jan Dirk Wegner, Konrad Schindler:
Guided Super-Resolution As Pixel-to-Pixel Transformation. 8828-8836

Tiantian Wang, Yongri Piao, Huchuan Lu, Xiao Li, Lihe Zhang:
Deep Learning for Light Field Saliency Detection. 8837-8847

Kai Zhao, Shanghua Gao, Wenguan Wang, Ming-Ming Cheng:
Optimizing the F-Measure for Threshold-Free Salient Object Detection. 8848-8856

Chaohao Xie, Shaohui Liu, Chao Li, Ming-Ming Cheng, Wangmeng Zuo, Xiao Liu, Shilei Wen, Errui Ding:
Image Inpainting With Learnable Bidirectional Attention Maps. 8857-8866

Thibaud Ehret, Axel Davy, Pablo Arias, Gabriele Facciolo:
Joint Demosaicking and Denoising by Fine-Tuning of Bursts of Raw Images. 8867-8876

Orest Kupyn, Tetiana Martyniuk, Junru Wu, Zhangyang Wang:
DeblurGAN-v2: Deblurring (Orders-of-Magnitude) Faster and Better. 8877-8886

Language & Reasoning

Lei Ke, Wenjie Pei, Ruiyu Li, Xiaoyong Shen, Yu-Wing Tai:
Reflective Decoding Network for Image Captioning. 8887-8896

Gilad Vered, Gal Oren, Yuval Atzmon, Gal Chechik:
Joint Optimization for Cooperative Image Captioning. 8897-8906

Tanzila Rahman, Bicheng Xu, Leonid Sigal:
Watch, Listen and Tell: Multi-Modal Weakly Supervised Dense Event Captioning. 8907-8916

Jingyi Hou, Xinxiao Wu, Wentian Zhao, Jiebo Luo, Yunde Jia:
Joint Syntax Representation Learning and Visual Cue Translation for Video Captioning. 8917-8926

Guang Li, Linchao Zhu, Ping Liu, Yi Yang:
Entangled Transformer for Image Captioning. 8927-8936

Panos Achlioptas, Leonidas J. Guibas, Noah D. Goodman, Judy Fan, Robert X. D. Hawkins:
Shapeglot: Learning Language for Shape Differentiation. 8937-8946

Harsh Agrawal, Peter Anderson, Karan Desai, Yufei Wang, Xinlei Chen, Rishabh Jain, Mark Johnson, Dhruv Batra, Devi Parikh, Stefan Lee:
nocaps: novel object captioning at scale. 8947-8956

3D From Multiview & Sensors

Christopher B. Choy, Jaesik Park, Vladlen Koltun:
Fully Convolutional Geometric Features. 8957-8965

Georgios Georgakis, Srikrishna Karanam, Ziyan Wu, Jana Kosecka:
Learning Local RGB-to-CAD Correspondences for Object Pose Estimation. 8966-8975

Ariel Gordon, Hanhan Li, Rico Jonschkowski, Anelia Angelova:
Depth From Videos in the Wild: Unsupervised Monocular Depth Learning From Unknown Cameras. 8976-8985

Changhee Won, Jongbin Ryu, Jongwoo Lim:
OmniMVS: End-to-End Learning for Omnidirectional Stereo Matching. 8986-8995

Chuangrong Chen, Xiaozhi Chen, Hui Cheng:
On the Over-Smoothing Problem of CNN Based Disparity Estimation. 8996-9004

Image & Video Synthesis

Hang Gao, Huazhe Xu, Qi-Zhi Cai, Ruth Wang, Fisher Yu, Trevor Darrell:
Disentangling Propagation and Generation for Video Prediction. 9005-9014

Badour Albahar, Jia-Bin Huang:
Guided Image-to-Image Translation With Bi-Directional Feature Transformation. 9015-9024

Haoye Dong, Xiaodan Liang, Xiaohui Shen, Bochao Wang, Hanjiang Lai, Jia Zhu, Zhiting Hu, Jian Yin:
Towards Multi-Pose Guided Virtual Try-On Network. 9025-9034

Jaejun Yoo, Youngjung Uh, Sanghyuk Chun, Byeongkyu Kang, Jung-Woo Ha:
Photorealistic Style Transfer via Wavelet Transforms. 9035-9044

Cong Yu, Yang Hu, Yan Chen, Bing Zeng:
Personalized Fashion Design. 9045-9054

Hyunsu Kim, Ho Young Jhoo, Eunhyeok Park, Sungjoo Yoo:
Tag2Pix: Line Art Colorization Using Text Tag With SECat and Changing Loss. 9055-9064

Ya-Liang Chang, Zhe Yu Liu, Kuan-Ying Lee, Winston H. Hsu:
Free-Form Video Inpainting With 3D Gated Convolution and Temporal

PatchGAN. 9065-9074

Applications, Medical & Robotics

Wei Feng, Wenhao He, Fei Yin, Xu-Yao Zhang, Cheng-Lin Liu:
TextDragon: An End-to-End Framework for Arbitrary Shaped Text Spotting. 9075-9084

Yipeng Sun, Jiaming Liu, Wei Liu, Junyu Han, Errui Ding, Jingtuo Liu:
Chinese Street View Text: Large-Scale Chinese Text Reading With Partially Supervised Learning. 9085-9094

Zhiliang Zeng, Xianzhi Li, Ying Kin Yu, Chi-Wing Fu:
Deep Floor Plan Recognition Using a Multi-Task Network With Room-Boundary-Guided Attention. 9095-9103

Fangneng Zhan, Chuhui Xue, Shijian Lu:
GA-DAN: Geometry-Aware Domain Adaptation Network for Scene Text Detection and Recognition. 9104-9114

Tianlang Chen, Zhaowen Wang, Ning Xu, Hailin Jin, Jiebo Luo:
Large-Scale Tag-Based Font Retrieval With Generative Feature Learning. 9115-9124

Linjie Xing, Zhi Tian, Weilin Huang, Matthew R. Scott:
Convolutional Character Networks. 9125-9135

Jiaqi Duan, Youjiang Xu, Zhanghui Kuang, Xiaoyu Yue, Hongbin Sun, Yue Guan, Wayne Zhang:
Geometry Normalization Networks for Accurate Scene Text Detection. 9136-9145

Mingkun Yang, Yushuo Guan, Minghui Liao, Xin He, Kaigui Bian, Song Bai, Cong Yao, Xiang Bai:
Symmetry-Constrained Rectification Network for Scene Text Recognition. 9146-9155

Oral 4.2A

Segmentation, Detection, 3D Scene Understanding

Daniel Bolya, Chong Zhou, Fanyi Xiao, Yong Jae Lee:
YOLACT: Real-Time Instance Segmentation. 9156-9165

Xia Li, Zhisheng Zhong, Jianlong Wu, Yibo Yang, Zhouchen Lin, Hong Liu:
Expectation-Maximization Attention Networks for Semantic Segmentation. 9166-9175

Yifan Zhao, Jia Li, Yu Zhang, Yonghong Tian:
Multi-Class Part Parsing With Joint Boundary-Semantic Awareness. 9176-9185

Runjin Chen, Hao Chen, Ge Huang, Jie Ren, Quanshi Zhang:
Explaining Neural Networks Semantically and Quantitatively. 9186-9195

Kaixin Wang, Jun Hao Liew, Yingtian Zou, Daquan Zhou, Jiashi Feng:
PANet: Few-Shot Image Semantic Segmentation With Prototype Alignment. 9196-9205

Weicheng Kuo, Anelia Angelova, Jitendra Malik, Tsung-Yi Lin:
ShapeMask: Learning to Segment Novel Objects by Refining Shape Priors. 9206-9215

Haiping Wu, Yuntao Chen, Naiyan Wang, Zhaoxiang Zhang:
Sequence Level Semantics Aggregation for Video Object Detection. 9216-9224

Seoung Wug Oh, Joon-Young Lee, Ning Xu, Seon Joo Kim:
Video Object Segmentation Using Space-Time Memory Networks. 9225-9234

Wenguan Wang, Xiankai Lu, Jianbing Shen, David J. Crandall, Ling Shao:
Zero-Shot Video Object Segmentation via Attentive Graph Neural Networks. 9235-9244

Xingyu Liu, Mengyuan Yan, Jeannette Bohg:
MeteorNet: Deep Learning on Dynamic 3D Point Cloud Sequences. 9245-9254

Jean Lahoud, Bernard Ghanem, Martin R. Oswald, Marc Pollefeys:
3D Instance Segmentation via Multi-Task Metric Learning. 9255-9265

Guohao Li, Matthias Müller, Ali K. Thabet, Bernard Ghanem:
DeepGCNs: Can GCNs Go As Deep As CNNs? 9266-9275

Charles R. Qi, Or Litany, Kaiming He, Leonidas J. Guibas:
Deep Hough Voting for 3D Object Detection in Point Clouds. 9276-9285

Garrick Brazil, Xiaoming Liu:
M3D-RPN: Monocular 3D Region Proposal Network for Object Detection. 9286-9295

Jens Behley, Martin Garbade, Andres Milioto, Jan Quenzel, Sven Behnke, Cyrill Stachniss, Jürgen Gall:
SemanticKITTI: A Dataset for Semantic Scene Understanding of LiDAR Sequences. 9296-9306

Senthil Kumar Yogamani, Christian Witt, Hazem Rashed, Sanjaya Nayak, Saquib Mansoor, Padraig Varley, Xavier Perrotton, Derek O'Dea, Patrick Pérez, Ciarán Hughes, Jonathan Horgan, Ganesh Sistu, Sumanth Chennupati, Michal Uricár, Stefan Milz, Martin Simon, Karl Amende:
WoodScape: A Multi-Task, Multi-Camera Fisheye Dataset for Autonomous Driving. 9307-9317

Dzung A. Doan, Yasir Latif, Tat-Jun Chin, Yu Liu, Thanh-Toan Do, Ian D. Reid:
Scalable Place Recognition Under Appearance Change for Autonomous Driving. 9318-9327

Felipe Codevilla, Eder Santana, Antonio M. López, Adrien Gaidon:
Exploring the Limitations of Behavior Cloning for Autonomous Driving. 9328-9337

Manolis Savva, Jitendra Malik, Devi Parikh, Dhruv Batra, Abhishek Kadian, Oleksandr Maksymets, Yili Zhao, Erik Wijmans, Bhavana Jain, Julian Straub, Jia Liu, Vladlen Koltun:
Habitat: A Platform for Embodied AI Research. 9338-9346

Oral 4.2B

Face & Body Modeling

Bangjie Yin, Luan Tran, Haoxiang Li, Xiaohui Shen, Xiaoming Liu:
Towards Interpretable Face Recognition. 9347-9356

Xiaobo Wang, Shuo Wang, Hailin Shi, Jun Wang, Tao Mei:
Co-Mining: Deep Face Recognition With Noisy Labels. 9357-9366

Seonwook Park, Shalini De Mello, Pavlo Molchanov, Umar Iqbal, Otmar Hilliges, Jan Kautz:
Few-Shot Adaptive Gaze Estimation. 9367-9376

Oran Gafni, Lior Wolf, Yaniv Taigman:
Live Face De-Identification in Video. 9377-9386

Wenqi Ren, Jiaolong Yang, Senyou Deng, David P. Wipf, Xiaochun Cao, Xin Tong:
Face Video Deblurring Using 3D Facial Priors. 9387-9396

Jingtan Piao, Chen Qian, Hongsheng Li:
Semi-Supervised Monocular 3D Face Reconstruction With End-to-End Shape-Preserved Domain Transfer. 9397-9406

Feng Liu, Luan Tran, Xiaoming Liu:
3D Face Modeling From Diverse Raw Scan Data. 9407-9417

Victoria Fernández Abrevaya, Adnane Boukhayma, Stefanie Wuhrer, Edmond Boyer:
A Decoupled 3D Facial Shape Model by Adversarial Training. 9418-9427

Anpei Chen, Zhang Chen, Guli Zhang, Kenny Mitchell, Jingyi Yu:
Photo-Realistic Facial Details Synthesis From Single Image. 9428-9438

Zhenliang He, Meina Kan, Shiguang Shan, Xilin Chen:
S2GAN: Share Aging Factors Across Ages and Share Aging Trends Among Individuals. 9439-9448

Ben Usman, Nick Dufour, Kate Saenko, Chris Bregler:
PuppetGAN: Cross-Domain Image Manipulation by Demonstration. 9449-9457

Egor Zakharov, Aliaksandra Shysheya, Egor Burkov, Victor S. Lempitsky:
Few-Shot Adversarial Learning of Realistic Neural Talking Head Models. 9458-9467

Bo Wan, Desen Zhou, Yongfei Liu, Rongjie Li, Xuming He:
Pose-Aware Multi-Level Feature Network for Human Object Interaction Detection. 9468-9477

Haodong Duan, Kwan-Yee Lin, Sheng Jin, Wentao Liu, Chen Qian, Wanli Ouyang:
TRB: A Novel Triplet Representation for Understanding 2D Human Body. 9478-9487

Wei Mao, Miaomiao Liu, Mathieu Salzmann, Hongdong Li:
Learning Trajectory Dependencies for Human Motion Prediction. 9488-9496

Jinkun Cao, Hongyang Tang, Haoshu Fang, Xiaoyong Shen, Yu-Wing Tai, Cewu Lu:
Cross-Domain Adaptation for Animal Pose Estimation. 9497-9506

Poster 4.2

Recognition

Jiyang Gao, Jiang Wang, Shengyang Dai, Li-Jia Li, Ram Nevatia:
NOTE-RCNN: NOise Tolerant Ensemble RCNN for Semi-Supervised Object Detection. 9507-9516

Qing Yu, Kiyoharu Aizawa:
Unsupervised Out-of-Distribution Detection by Maximum Classifier Discrepancy. 9517-9525

Yan Huang, Qiang Wu, Jingsong Xu, Yi Zhong:
SBSGAN: Suppression of Inter-Domain Background Shift for Person Re-Identification. 9526-9535

Jing Nie, Rao Muhammad Anwer, Hisham Cholakkal, Fahad Shahbaz Khan, Yanwei Pang, Ling Shao:
Enriched Feature Guided Refinement Network for Object Detection. 9536-9545

Guangyi Chen, Tianren Zhang, Jiwen Lu, Jie Zhou:
Deep Meta Metric Learning. 9546-9555

Chunluan Zhou, Ming Yang, Junsong Yuan:
Discriminative Feature Transformation for Occluded Pedestrian Detection. 9556-9565

Supreeth Narasimhaswamy, Zhengwei Wei, Yang Wang, Justin Zhang, Minh Hoai Nguyen:
Contextual Attention for Hand Detection in the Wild. 9566-9575

Xiaopeng Yan, Ziliang Chen, Anni Xu, Xiaoxi Wang, Xiaodan Liang, Liang Lin:
Meta R-CNN: Towards General Solver for Instance-Level Low-Shot Learning. 9576-9585

Chi Zhang, Guosheng Lin, Fayao Liu, Jiushuang Guo, Qingyao Wu, Rui Yao:
Pyramid Graph Networks With Connection Attentions for Region-Based One-Shot Semantic Segmentation. 9586-9594

Oisin Mac Aodha, Elijah Cole, Pietro Perona:
Presence-Only Geographical Priors for Fine-Grained Image Classification. 9595-9605

Junran Peng, Ming Sun, Zhaoxiang Zhang, Tieniu Tan, Junjie Yan:
POD: Practical Object Detection With Scale-Sensitive Network. 9606-9615

Joshua C. Peterson, Ruairidh M. Battleday, Thomas L. Griffiths, Olga Russakovsky:
Human Uncertainty Makes Classification More Robust. 9616-9625

Zhi Tian, Chunhua Shen, Hao Chen, Tong He:
FCOS: Fully Convolutional One-Stage Object Detection. 9626-9635

Guangyi Chen, Chunze Lin, Liangliang Ren, Jiwen Lu, Jie Zhou:
Self-Critical Attention Learning for Person Re-Identification. 9636-9645

Xinqian Gu, Bingpeng Ma, Hong Chang, Shiguang Shan, Xilin Chen:
Temporal Knowledge Propagation for Image-to-Video Person Re-Identification. 9646-9655

Ze Yang, Shaohui Liu, Han Hu, Liwei Wang, Stephen Lin:
RepPoints: Point Set Representation for Object Detection. 9656-9665

Haonan Luo, Guosheng Lin, Zichuan Liu, Fayao Liu, Zhenmin Tang, Yazhou Yao:
SegEQA: Video Segmentation Based Visual Attention for Embodied Question Answering. 9666-9675

Tanmay Gupta, Alexander G. Schwing, Derek Hoiem:
No-Frills Human-Object Interaction Detection: Factorization, Layout Encodings, and Training Techniques. 9676-9684

Keren Ye, Mingda Zhang, Adriana Kovashka, Wei Li, Danfeng Qin, Jesse Berent:
Cap2Det: Learning to Amplify Weak Caption Supervision for Object Detection. 9685-9694

Tomás Jenícek, Ondrej Chum:
No Fear of the Dark: Image Retrieval Under Varying Illumination Conditions. 9695-9703

Jiale Cao, Yanwei Pang, Jungong Han, Xuelong Li:
Hierarchical Shot Detector. 9704-9713

Aoxue Li, Tiange Luo, Tao Xiang, Weiran Huang, Liwei Wang:
Few-Shot Learning With Global Class Representations. 9714-9723

Junhyug Noh, Wonho Bae, Wonhee Lee, Jinhwan Seo, Gunhee Kim:
Better to Follow, Follow to Be Better: Towards Precise Supervision of Feature Super-Resolution for Small Object Detection. 9724-9733

Xiaoyan Li, Meina Kan, Shiguang Shan, Xilin Chen:
Weakly Supervised Object Detection With Segmentation Collaboration. 9734-9743

Mahyar Najibi, Bharat Singh, Larry Davis:
AutoFocus: Efficient Multi-Scale Inference. 9744-9754

Mykhailo Shvets, Wei Liu, Alexander C. Berg:
Leveraging Long-Range Temporal Relationships Between Proposals for Video Object Detection. 9755-9763

Huajie Jiang, Ruiping Wang, Shiguang Shan, Xilin Chen:
Transferable Contrastive Network for Generalized Zero-Shot Learning. 9764-9773

Yilun Chen, Shu Liu, Xiaoyong Shen, Jiaya Jia:
Fast Point R-CNN. 9774-9783

Georgia Gkioxari, Justin Johnson, Jitendra Malik:
Mesh R-CNN. 9784-9794

Yudong Chen, Zhihui Lai, Yujuan Ding, Kaiyi Lin, Wai Keung Wong:
Deep Supervised Hashing With Anchor Graph. 9795-9803

Hao Yang, Hao Wu, Hao Chen:
Detecting 11K Classes: Large Scale Object Detection Without Fine-Grained Bounding Boxes. 9804-9812

Chuchu Han, Jiacheng Ye, Yunshan Zhong, Xin Tan, Chi Zhang, Changxin Gao, Nong Sang:
Re-ID Driven Localization Refinement for Person Search. 9813-9822

Huu Le, Ming Xu, Tuan Hoang, Michael Milford:
Hierarchical Encoding of Sequential Data With Compact and Sub-Linear Storage Cost. 9823-9832

Gao Yan, Boxiao Liu, Nan Guo, Xiaochun Ye, Fang Wan, Haihang You, Dongrui Fan:
C-MIDN: Coupled Multiple Instance Detection Network With Segmentation Guidance for Weakly Supervised Object Detection. 9833-9842

Yizhe Zhu, Jianwen Xie, Bingchen Liu, Ahmed Elgammal:
Learning Feature-to-Feature Translator by Alternating Back-Propagation for Generative Zero-Shot Learning. 9843-9853

Leulseged Tesfaye Alemu, Mubarak Shah, Marcello Pelillo:
Deep Constrained Dominant Sets for Person Re-Identification. 9854-9863

Xu Ji, Andrea Vedaldi, João F. Henriques:
Invariant Information Clustering for Unsupervised Image Classification and Segmentation. 9864-9873

Statistics, Physics, Theory & Datasets

Masataka Yamaguchi, Go Irie, Takahito Kawanishi, Kunio Kashino:
Subspace Structure-Aware Spectral Clustering for Robust Subspace Clustering. 9874-9883

Bing Su, Jiahuan Zhou, Ying Wu:
Order-Preserving Wasserstein Discriminant Analysis. 9884-9893

Akash Abdu Jyothi, Thibaut Durand, Jiawei He, Leonid Sigal, Greg Mori:
LayoutVAE: Stochastic Scene Layout Generation From a Label Set. 9894-9903

Jie Zhou, Xinke Ma, Li Liang, Liu Yuhe, Shijin Xu, Sim Heng Ong, Yang Yang:
Robust Variational Bayesian Point Set Registration. 9904-9913

Chong You, Chun-Guang Li, Daniel P. Robinson, René Vidal:
Is an Affine Constraint Needed for Affine Subspace Clustering? 9914-9923

Yu-Xiong Wang, Deva Ramanan, Martial Hebert:
Meta-Learning to Detect Rare Objects. 9924-9933

Zhenhua Wang, Tong Liu, Qinfeng Shi, M. Pawan Kumar, Jianhua Zhang:
New Convex Relaxations for MRF Inference With Unknown Graphs. 9934-9942

Zhijie Deng, Yucen Luo, Jun Zhu:
Cluster Alignment With a Teacher for Unsupervised Domain Adaptation. 9943-9952

Luca Anthony Thiede, Pratik Prabhanjan Brahma:
Analyzing the Variety Loss in the Context of Probabilistic Trajectory Prediction. 9953-9962

3D From Single View & RGBD

Junyi Pan, Xiaoguang Han, Weikai Chen, Jiapeng Tang, Kui Jia:
Deep Mesh Reconstruction From Single RGB Images via Topology Modification Networks. 9963-9972

Wenqi Xian, Zhengqi Li, Noah Snavely, Matthew Fisher, Jonathan Eisenmann, Eli Shechtman:
UprightNet: Geometry-Aware Camera Orientation Estimation From Single Images. 9973-9982

Philipp Henzler, Niloy J. Mitra, Tobias Ritschel:
Escaping Plato's Cave: 3D Shape From Adversarial Rendering. 9983-9992

Di Qiu, Jiahao Pang, Wenxiu Sun, Chengxi Yang:
Deep End-to-End Alignment and Refinement for Time-of-Flight RGB-D Module. 9993-10002

Erickson Rangel do Nascimento, Guilherme A. Potje, Renato Martins, Felipe C. Chamone, Mario F. M. Campos, Ruzena Bajcsy:
GEOBIT: A Geodesic-Based Binary Descriptor Invariant to Non-Rigid Deformations for RGB-D Images. 10003-10011

Alan Lukezic, Ugur Kart, Jani Käpylä, Ahmed Durmush, Joni-Kristian Kamarainen, Jiri Matas, Matej Kristan:
CDTB: A Color and Depth Visual Object Tracking Dataset and Benchmark. 10012-10021

Yun Chen, Bin Yang, Ming Liang, Raquel Urtasun:
Learning Joint 2D-3D Representations for Depth Completion. 10022-10031

Face & Body

Shengju Qian, Kwan-Yee Lin, Wayne Wu, Yangxiaokang Liu, Quan Wang, Fumin Shen, Chen Qian, Ran He:
Make a Face: Towards Arbitrary High Fidelity Face Manipulation. 10032-10041

Peipei Li, Xiang Wu, Yibo Hu, Ran He, Zhenan Sun:
M2FPA: A Multi-Yaw Multi-Pitch High-Quality Dataset and Benchmark for Facial Pose Analysis. 10042-10050

Bingyu Liu, Weihong Deng, Yaoyao Zhong, Mei Wang, Jiani Hu, Xunqiang Tao, Yaohai Huang:
Fair Loss: Margin-Aware Reinforcement Learning for Deep Face Recognition. 10051-10060

Xiaowei Yuan, In Kyu Park:
Face De-Occlusion Using 3D Morphable Model and Generative Adversarial Network. 10061-10070

Sheng-Yu Wang, Oliver Wang, Richard Zhang, Andrew Owens, Alexei A. Efros:
Detecting Photoshopped Faces by Scripting Photoshop. 10071-10080

Ye Yuan, Kris Kitani:
Ego-Pose Estimation and Forecasting As Real-Time PD Control. 10081-10091

Jie Song, Bjoern Andres, Michael J. Black, Otmar Hilliges, Siyu Tang:
End-to-End Learning for Graph Decomposition. 10092-10101

Joseph P. Robinson, Yuncheng Li, Ning Zhang, Yun Fu, Sergey Tulyakov:
Laplace Landmark Localization. 10102-10111

Mingmin Zhao, Yingcheng Liu, Aniruddh Raghu, Hang Zhao, Tianhong Li, Antonio Torralba, Dina Katabi:
Through-Wall Human Mesh Recovery Using Radio Signals. 10112-10121

Hakan Cevikalp, Golara Ghorban Dordinejad:
Discriminatively Learned Convex Models for Set Based Face Recognition. 10122-10131

Gyeongsik Moon, Ju Yong Chang, Kyoung Mu Lee:
Camera Distance-Aware Top-Down Approach for 3D Multi-Person Pose Estimation From a Single RGB Image. 10132-10141

Jiyoung Lee, Seungryong Kim, Sunok Kim, Jungin Park, Kwanghoon Sohn:
Context-Aware Emotion Recognition Networks. 10142-10151

Shengju Qian, Keqiang Sun, Wayne Wu, Chen Qian, Jiaya Jia:
Aggregation via Separation: Boosting Facial Landmark Detector With Semi-Supervised Style Translation. 10152-10162

Felix Kuhnke, Jörn Ostermann:
Deep Head Pose Estimation Using Synthetic Images and Partial Adversarial Domain Adaption for Continuous Label Spaces. 10163-10172

Computational Photography & Graphics

Eden Sassoon, Tali Treibitz, Yoav Y. Schechner:
Flare in Interference-Based Hyperspectral Cameras. 10173-10181

Shipeng Zhang, Lizhi Wang, Ying Fu, Xiaoming Zhong, Hua Huang:
Computational Hyperspectral Imaging Based on Dimension-Discriminative Low-Rank Tensor Recovery. 10182-10191

Julie Chang, Gordon Wetzstein:
Deep Optics for Monocular Depth Estimation and 3D Object Detection. 10192-10201

Shirsendu Sukanta Halder, Jean-François Lalonde, Raoul de Charette:
Physics-Based Rendering for Improving Robustness to Rain. 10202-10211

Bin Ding, Chengjiang Long, Ling Zhang, Chunxia Xiao:
ARGAN: Attentive Recurrent Generative Adversarial Network for Shadow Detection and Removal. 10212-10221

Jiawei Ma, Xiao-Yang Liu, Zheng Shou, Xin Yuan:
Deep Tensor ADMM-Net for Snapshot Compressive Imaging. 10222-10231

Low-Level & Optimization

Thomas Probst, Danda Pani Paudel, Ajad Chhatkuli, Luc Van Gool:
Convex Relaxations for Consensus and Non-Minimal Problems in 3D Vision. 10232-10241

Christopher Zach, Guillaume Bourmaud:
Pareto Meets Huber: Efficiently Avoiding Poor Minima in Robust Estimation. 10242-10250

Yifan Sun, Jiacheng Zhuo, Arnav Mohan, Qixing Huang:
K-Best Transformation Synchronization. 10251-10260

Jonas Geiping, Michael Moeller:
Parametric Majorization for Data-Driven Energy Minimization Methods. 10261-10272

Xingchen Ma, Amal Rannen Triki, Maxim Berman, Christos Sagonas, Jacques Calì, Matthew B. Blaschko:
A Bayesian Optimization Framework for Neural Network Compression. 10273-10282

Florian Bernard, Johan Thunberg, Paul Swoboda, Christian Theobalt:
HiPPI: Higher-Order Projected Power Iterations for Scalable Multi-Matching. 10283-10292

Language & Reasoning

Ronghang Hu, Anna Rohrbach, Trevor Darrell, Kate Saenko:
Language-Conditioned Graph Networks for Relational Reasoning. 10293-10302

Alaaeldin El-Nouby, Shikhar Sharma, Hannes Schulz, R. Devon Hjelm, Layla El Asri, Samira Ebrahimi Kahou, Yoshua Bengio, Graham W. Taylor:
Tell, Draw, and Repeat: Generating and Modifying Images Based on Continual Linguistic Instruction. 10303-10311

Linjie Li, Zhe Gan, Yu Cheng, Jingjing Liu:
Relation-Aware Graph Attention Network for Visual Question Answering. 10312-10321

Jiuxiang Gu, Shafiq R. Joty, Jianfei Cai, Handong Zhao, Xu Yang, Gang Wang:
Unpaired Image Captioning via Scene Graph Alignments. 10322-10331

Yannick Le Cacheux, Hervé Le Borgne, Michel Crucianu:
Modeling Inter and Intra-Class Relations in the Triplet Loss for Zero-Shot Learning. 10332-10341

Rui Lu, Feng Xue, Menghan Zhou, Anlong Ming, Yu Zhou:
Occlusion-Shared and Feature-Separated Network for Occlusion Relationship Reasoning. 10342-10351

Yufei Ye, Maneesh Singh, Abhinav Gupta, Shubham Tulsiani:
Compositional Video Prediction. 10352-10361

Mohammed Suhail, Leonid Sigal:
Mixture-Kernel Graph Attention Network for Situation Recognition. 10362-10371

Reuben Tan, Mariya I. Vasileva, Kate Saenko, Bryan A. Plummer:
Learning Similarity Conditions Without Explicit Supervision. 10372-10381

Huikun Bi, Zhong Fang, Tianlu Mao, Zhaoqi Wang, Zhigang Deng:
Joint Prediction for Kinematic Trajectories in Vehicle-Pedestrian-Mixed Scenes. 10382-10391

Tingke Shen, Amlan Kar, Sanja Fidler:
Learning to Caption Images Through a Lifetime by Asking Questions. 10392-10401

Yuanzhi Liang, Yalong Bai, Wei Zhang, Xueming Qian, Li Zhu, Tao Mei:
VrR-VG: Refocusing Visually-Relevant Relationships. 10402-10411

3D From Multiview & Sensors

Andrea Romanoni, Matteo Matteucci:
TAPA-MVS: Textureless-Aware PAtchMatch Multi-View Stereo. 10412-10421

Armin Mustafa, Chris Russell, Adrian Hilton:
U4D: Unsupervised 4D Dynamic Scene Understanding. 10422-10431

Li Jiang, Hengshuang Zhao, Shu Liu, Xiaoyong Shen, Chi-Wing Fu, Jiaya Jia:
Hierarchical Point-Edge Interaction Network for Point Cloud Semantic Segmentation. 10432-10440

Zhizhong Han, Xiyang Wang, Yu-Shen Liu, Matthias Zwicker:
Multi-Angle Point Cloud-VAE: Unsupervised Feature Learning for 3D Point Clouds From Multiple Angles by Joint Self-Reconstruction and Half-to-Half Prediction. 10441-10450

Keyang Luo, Tao Guan, Lili Ju, Haipeng Huang, Yawei Luo:
P-MVSNet: Learning Patch-Wise Matching Confidence Aggregation for Multi-View Stereo. 10451-10460

Image & Video Synthesis

Yung-Han Ho, Chuan-Yuan Cho, Guo-Lun Jin, Wen-Hsiao Peng:
SME-Net: Sparse Motion Estimation for Parametric Video Prediction Through Reinforcement Learning. 10461-10469

Xintong Han, Weilin Huang, Xiaojun Hu, Matthew R. Scott:
ClothFlow: A Flow-Based Model for Clothed Person Generation. 10470-10479

Qiao Gu, Guanzhi Wang, Mang Tik Chiu, Yu-Wing Tai, Chi-Keung Tang:
LADN: Local Adversarial Disentangling Network for Facial Makeup and De-Makeup. 10480-10489

Tsun-Hsuan Wang, Yen-Chi Cheng, Chieh Hubert Lin, Hwann-Tzong Chen, Min Sun:
Point-to-Point Video Generation. 10490-10499

Hongchen Tan, Xiuping Liu, Xin Li, Yi Zhang, Baocai Yin:
Semantics-Enhanced Adversarial Nets for Text-to-Image Synthesis. 10500-10509

Ruiyun Yu, Xiaoqi Wang, Xiaohui Xie:
VTNFP: An Image-Based Virtual Try-On Network With Body and Clothing Feature Preservation. 10510-10519

Dilip Krishnan, Piotr Teterwak, Aaron Sarna, Aaron Maschinot, Ce Liu, David Belanger, William T. Freeman:
Boundless: Generative Adversarial Networks for Image Extension. 10520-10529

Wei Sun, Tianfu Wu:
Image Synthesis From Reconfigurable Layout and Style. 10530-10539

Kenan E. Ak, Ashraf A. Kassim, Joo-Hwee Lim, Jo Yew Tham:
Attribute Manipulation Generative Adversarial Networks for Fashion Images. 10540-10549

Ming-Yu Liu, Xun Huang, Arun Mallya, Tero Karras, Timo Aila, Jaakko Lehtinen, Jan Kautz:
Few-Shot Unsupervised Image-to-Image Translation. 10550-10559

Zongxin Yang, Jian Dong, Ping Liu, Yi Yang, Shuicheng Yan:
Very Long Natural Scenery Image Prediction by Outpainting. 10560-10569

Applications, Medical & Robotics

Ronak Mehta, Rudrasis Chakraborty, Vikas Singh, Yunyang Xiong:
Scaling Recurrent Models via Orthogonal Approximations in Tensor Trains. 10570-10578

Jinwoo Kim, Woojae Kim, Heeseok Oh, Seongmin Lee, Sanghoon Lee:
A Deep Cybersickness Predictor Based on Brain Signal Analysis for Virtual Reality Contents. 10579-10588

Botong Wu, Xinwei Sun, Lingjing Hu, Yizhou Wang:
Learning With Unsure Data for Medical Image Diagnosis. 10589-10598

Shengyu Zhao, Yue Dong, Eric I-Chao Chang, Yan Xu:
Recursive Cascaded Networks for Unsupervised Medical Image Registration. 10599-10609

Haoliang Sun, Ronak Mehta, Hao Henry Zhou, Zhichun Huang, Sterling C. Johnson, Vivek Prabhakaran, Vikas Singh:

DUAL-GLOW: Conditional Flow-Based Generative Model for Modality Transfer. 10610-10619

Rudrasis Chakraborty, Xingjian Zhen, Nicholas Vogt, Barbara B. Bendlin, Vikas Singh:
Dilated Convolutional Neural Networks for Sequential Manifold-Valued Data. 10620-10630

Jingyu Liu, Gangming Zhao, Yu Fei, Ming Zhang, Yizhou Wang, Yizhou Yu:
Align, Attend and Locate: Chest X-Ray Diagnosis via Contrast Induced Attention Network With Limited Supervision. 10631-10640

Xiaoping Wu, Ni Wen, Jie Liang, Yu-Kun Lai, Dongyu She, Ming-Ming Cheng, Jufeng Yang:
Joint Acne Image Grading and Counting via Label Distribution Learning. 10641-10650

Fengze Liu, Yingda Xia, Dong Yang, Alan L. Yuille, Daguang Xu:
An Alarm System for Segmentation Algorithm Based on Shape Model. 10651-10660

Lyndon Chan, Mahdi S. Hosseini, Corwyn Rowsell, Konstantinos N. Plataniotis, Savvas Damaskinos:
HistoSegNet: Semantic Segmentation of Histological Tissue Type in Whole Slide Images. 10661-10670

Yuyin Zhou, Zhe Li, Song Bai, Xinlei Chen, Mei Han, Chong Wang, Elliot K. Fishman, Alan L. Yuille:
Prior-Aware Neural Network for Partially-Supervised Multi-Organ Segmentation. 10671-10680

Gang Xu, Zhigang Song, Zhuo Sun, Calvin Ku, Zhe Yang, Cancheng Liu, Shuhao Wang, Jianpeng Ma, Wei Xu:
CAMEL: A Weakly Supervised Learning Framework for Histopathology Image Segmentation. 10681-10690

Seong Jae Hwang, Zirui Tao, Vikas Singh, Won Hwa Kim:
Conditional Recurrent Flow: Conditional Generation of Longitudinal Samples With Applications to Neuroimaging. 10691-10700

Shusuke Takahama, Yusuke Kurose, Yusuke Mukuta, Hiroyuki Abe, Masashi Fukayama, Akihiko Yoshizawa, Masanobu Kitagawa, Tatsuya Harada:
Multi-Stage Pathological Image Classification Using Semantic Segmentation. 10701-10710

Jiahua Dong, Yang Cong, Gan Sun, Dongdong Hou:
Semantic-Transferable Weakly-Supervised Endoscopic Lesions Segmentation. 10711-10720

Shir Gur, Lior Wolf, Lior Golgher, Pablo Blinder:
Unsupervised Microvascular Image Segmentation Using an Active Contours Mimicking Neural Network. 10721-10730

Prune Truong, Stefanos Apostolopoulos, Agata Mosinska, Samuel Stucky, Carlos Ciller, Sandro De Zanet:
GLAMpoints: Greedily Learned Accurate Match Points. 10731-10740