



We've just launched a new service: our brand new **dblp SPARQL query service**. Read more about it in our latest blog post or try out some of the SPARQL queries linked on the dblp web pages below.



search dblp

[+] 9th ICLR 2021: Virtual Event, Austria

> Home > Conferences and Workshops > ICLR

Dagstuhl

9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021. OpenReview.net 2021

[+] SPARQL queries

[–] Refine list

showing all 861 records

refine by search term

refine by author

Sergey Levine (16)
Yoshua Bengio (10)
Bernhard Schölkopf (8)
Zhangyang Wang (8)
J. Zico Kolter (8)
Aaron C. Courville (7)
Stefano Ermon (7)
Tie-Yan Liu (7)
Sung Ju Hwang (7)
Joshua B. Tenenbaum (7)
3,126 more options

refine by access
open (861)

Oral Presentations

- Marcin Andrychowicz, Anton Raichuk, Piotr Stanczyk, Manu Orsini, Sertan Girgin, Raphaël Marinier, Léonard Hussonot, Matthieu Geist, Olivier Pietquin, Marcin Michalski, Sylvain Gelly, Olivier Bachem:
What Matters for On-Policy Deep Actor-Critic Methods? A Large-Scale Study.
- Colin Wei, Kendrick Shen, Yining Chen, Tengyu Ma:
Theoretical Analysis of Self-Training with Deep Networks on Unlabeled Data.
- Dibya Ghosh, Abhishek Gupta, Ashwin Reddy, Justin Fu, Coline Manon Devin, Benjamin Eysenbach, Sergey Levine:
Learning to Reach Goals via Iterated Supervised Learning.
- Brenden K. Petersen, Mikel Landajuela, T. Nathan Mundhenk, Cláudio Prata Santiago, Sookyung Kim, Joanne Taery Kim:
Deep symbolic regression: Recovering mathematical expressions from data via risk-seeking policy gradients.
- Atsushi Nitanda, Taiji Suzuki:
Optimal Rates for Averaged Stochastic Gradient Descent under Neural Tangent Kernel Regime.
- Shuo Yang, Lu Liu, Min Xu:
Free Lunch for Few-shot Learning: Distribution Calibration.
- Mike Gartrell, Insu Han, Elvis Dohmatob, Jennifer Gillenwater, Victor-Emmanuel Brunel:
Scalable Learning and MAP Inference for Nonsymmetric Determinantal Point Processes.
- Deniz Oktay, Nick McGreivy, Joshua Aduol, Alex Beatson, Ryan P. Adams:
Randomized Automatic Differentiation.
- Luca Weihs, Aniruddha Kembhavi, Kiana Ehsani, Sarah M. Pratt, Winson Han, Alvaro Herrasti, Eric Kolve, Dustin Schwenk, Roozbeh Mottaghi, Ali Farhadi:
Learning Generalizable Visual Representations via Interactive Gameplay.
- Huy Tuan Pham, Phan-Minh Nguyen:
Global Convergence of Three-layer Neural Networks in the Mean Field Regime.
- Max B. Paulus, Chris J. Maddison, Andreas Krause:
Rao-Blackwellizing the Straight-Through Gumbel-Softmax Gradient Estimator.
- Krzysztof Marcin Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamás Sarlós, Peter Hawkins, Jared Quincy Davis, Afroz Mohiuddin, Lukasz Kaiser, David Benjamin Belanger, Lucy J. Colwell, Adrian Weller:
Rethinking Attention with Performers.
- Javier Antorán, Umang Bhatt, Tameem Adel, Adrian Weller, José Miguel Hernández-Lobato:
Getting a CLUE: A Method for Explaining Uncertainty Estimates.
- Xiaoxia Wu, Ethan Dyer, Behnam Neyshabur:
When Do Curricula Work?
- Durmus Alp Emre Acar, Yue Zhao , Ramon Matas Navarro, Matthew Mattina, Paul N. Whatmough, Venkatesh Saligrama:
Federated Learning Based on Dynamic Regularization.
- Jingfeng Zhang , Jianing Zhu, Gang Niu, Bo Han, Masashi Sugiyama, Mohan S. Kankanhalli:
Geometry-aware Instance-reweighted Adversarial Training.



- Jang-Hyun Kim, Wonho Choo, Hosan Jeong, Hyun Oh Song:
Co-Mixup: Saliency Guided Joint Mixup with Supermodular Diversity.
- Mikhail Yurochkin, Yuekai Sun:
SenSel: Sensitive Set Invariance for Enforcing Individual Fairness.
- Jeff Donahue, Sander Dieleman, Mikolaj Binkowski, Erich Elsen, Karen Simonyan:
End-to-end Adversarial Text-to-Speech.
- Bo Zhao, Konda Reddy Mopuri, Hakan Bilen :
Dataset Condensation with Gradient Matching.
- Ruochen Wang, Minhao Cheng , Xiangning Chen, Xiaocheng Tang, Cho-Jui Hsieh:
Rethinking Architecture Selection in Differentiable NAS.
- Muhammad Khalifa, Hady Elsahar, Marc Dymetman:
A Distributional Approach to Controlled Text Generation.
- Qiang Zhang, Tete Xiao, Alexei A. Efros, Lerrel Pinto, Xiaolong Wang:
Learning Cross-Domain Correspondence for Control with Dynamics Cycle-Consistency.
- Jonathan Gray, Adam Lerer, Anton Bakhtin, Noam Brown:
Human-Level Performance in No-Press Diplomacy via Equilibrium Search.
- Avi Singh, Huihan Liu, Gaoyue Zhou, Albert Yu, Nicholas Rhinehart, Sergey Levine:
Parrot: Data-Driven Behavioral Priors for Reinforcement Learning.
- Amy Zhang, Rowan Thomas McAllister, Roberto Calandra, Yarin Gal, Sergey Levine:
Learning Invariant Representations for Reinforcement Learning without Reconstruction.
- Xingang Pan, Bo Dai, Ziwei Liu, Chen Change Loy, Ping Luo:
Do 2D GANs Know 3D Shape? Unsupervised 3D Shape Reconstruction from 2D Image GANs.
- Lizhen Nie, Mao Ye, Qiang Liu, Dan Nicolae:
VCNet and Functional Targeted Regularization For Learning Causal Effects of Continuous Treatments.
- Suraj Srinivas, François Fleuret:
Rethinking the Role of Gradient-based Attribution Methods for Model Interpretability.
- Alexander Richard, Dejan Markovic, Israel D. Gebru, Steven Krenn, Gladstone Alexander Butler, Fernando De la Torre, Yaser Sheikh:
Neural Synthesis of Binaural Speech From Mono Audio.
- Zhifeng Kong, Wei Ping, Jiaji Huang, Kexin Zhao, Bryan Catanzaro:
DiffWave: A Versatile Diffusion Model for Audio Synthesis.
- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, Neil Houlsby:
An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale.
- Matthew Smart , Anton Zilman:
On the mapping between Hopfield networks and Restricted Boltzmann Machines.
- Glen Berseth , Daniel Geng, Coline Manon Devin, Nicholas Rhinehart, Chelsea Finn, Dinesh Jayaraman, Sergey Levine:
SMiRL: Surprise Minimizing Reinforcement Learning in Unstable Environments.
- John D. Co-Reyes, Yingjie Miao, Daiyi Peng, Esteban Real, Quoc V. Le, Sergey Levine, Honglak Lee, Aleksandra Faust:
Evolving Reinforcement Learning Algorithms.
- Xin Yuan, Pedro Henrique Pamplona Savarese, Michael Maire:
Growing Efficient Deep Networks by Structured Continuous Sparsification.
- Xizhou Zhu, Weijie Su, Lewei Lu, Bin Li, Xiaogang Wang, Jifeng Dai:
Deformable DETR: Deformable Transformers for End-to-End Object Detection.
- Ian Gemp, Brian McWilliams, Claire Vernade, Thore Graepel:
EigenGame: PCA as a Nash Equilibrium.
- Yuan Yin, Vincent Le Guen, Jérémie Donà, Emmanuel de Bézenac, Ibrahim Ayed, Nicolas Thome, Patrick Gallinari:
Augmenting Physical Models with Deep Networks for Complex Dynamics Forecasting.
- Erik Arakelyan , Daniel Daza, Pasquale Minervini, Michael Cochez:
Complex Query Answering with Neural Link Predictors.
- David A. Klindt, Lukas Schott, Yash Sharma, Ivan Ustyuzhaninov, Wieland Brendel, Matthias Bethge, Dylan M. Paiton:
Towards Nonlinear Disentanglement in Natural Data with Temporal Sparse Coding.
- Cheng Perng Phoo, Bharath Hariharan:
Self-training For Few-shot Transfer Across Extreme Task Differences.
- Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, Ben Poole:
Score-Based Generative Modeling through Stochastic Differential Equations.
- Biao Zhang, Ankur Bapna, Rico Sennrich , Orhan Firat:
Share or Not? Learning to Schedule Language-Specific Capacity for Multilingual Translation.
- Yuxuan Zhang, Wenzheng Chen, Huan Ling, Jun Gao, Yinan Zhang, Antonio Torralba, Sanja Fidler:
Image GANs meet Differentiable Rendering for Inverse Graphics and Interpretable 3D Neural Rendering.



Keyulu Xu, Mozhi Zhang, Jingling Li, Simon Shaolei Du, Ken-ichi Kawarabayashi, Stefanie Jegelka:
How Neural Networks Extrapolate: From Feedforward to Graph Neural Networks.

Zhengxian Lin, Kin-Ho Lam, Alan Fern:
Contrastive Explanations for Reinforcement Learning via Embedded Self Predictions.

Chenlin Meng, Jiaming Song, Yang Song, Shengjia Zhao, Stefano Ermon:
Improved Autoregressive Modeling with Distribution Smoothing.

Beidi Chen, Zichang Liu, Binghui Peng, Zhaozhuo Xu, Jonathan Lingjie Li, Tri Dao, Zhao Song, Anshumali Shrivastava, Christopher Ré:
MONGOOSE: A Learnable LSH Framework for Efficient Neural Network Training.

Gobinda Saha , Isha Garg, Kaushik Roy:
Gradient Projection Memory for Continual Learning.

Zhiyuan Li , Yi Zhang, Sanjeev Arora:
Why Are Convolutional Nets More Sample-Efficient than Fully-Connected Nets?

Ankit Vani, Max Schwarzer, Yuchen Lu, Eeshan Dhekane, Aaron C. Courville:
Iterated learning for emergent systematicity in VQA.

T. Konstantin Rusch, Siddhartha Mishra:
Coupled Oscillatory Recurrent Neural Network (coRNN): An accurate and (gradient) stable architecture for learning long time dependencies.

Spotlight Presentations

Zhenyu Liao , Romain Couillet, Michael W. Mahoney:
Sparse Quantized Spectral Clustering.

Binh Tang, David S. Matteson:
Graph-Based Continual Learning.

Marisa Kirisame, Steven Lyubomirsky, Altan Haan, Jennifer Brennan, Mike He, Jared Roesch, Tianqi Chen, Zachary Tatlock:
Dynamic Tensor Rematerialization.

Zirui Wang, Yulia Tsvetkov, Orhan Firat, Yuan Cao:
Gradient Vaccine: Investigating and Improving Multi-task Optimization in Massively Multilingual Models.

Yonggan Fu, Han Guo, Meng Li, Xin Yang, Yining Ding, Vikas Chandra, Yingyan Lin:
CPT: Efficient Deep Neural Network Training via Cyclic Precision.

Ainesh Bakshi, Chiranjib Bhattacharyya, Ravi Kannan, David P. Woodruff, Samson Zhou:
Learning a Latent Simplex in Input Sparsity Time.

Waïss Azizian, Marc Lelarge:
Expressive Power of Invariant and Equivariant Graph Neural Networks.

Tom Zahavy, André Barreto, Daniel J. Mankowitz, Shaobo Hou, Brendan O'Donoghue, Iurii Kemaev, Satinder Singh:
Discovering a set of policies for the worst case reward.

Stephen Tian, Suraj Nair, Frederik Ebert, Sudeep Dasari, Benjamin Eysenbach, Chelsea Finn, Sergey Levine:
Model-Based Visual Planning with Self-Supervised Functional Distances.

Pengfei Chen, Guangyong Chen, Junjie Ye, Jingwei Zhao, Pheng-Ann Heng:
Noise against noise: stochastic label noise helps combat inherent label noise.

Rishabh Agarwal, Marlos C. Machado , Pablo Samuel Castro, Marc G. Bellemare:
Contrastive Behavioral Similarity Embeddings for Generalization in Reinforcement Learning.

Zhisheng Xiao, Karsten Kreis, Jan Kautz, Arash Vahdat:
VAEBM: A Symbiosis between Variational Autoencoders and Energy-based Models.

Liam Li, Mikhail Khodak, Nina Balcan, Ameet Talwalkar:
Geometry-Aware Gradient Algorithms for Neural Architecture Search.

Talya Eden, Piotr Indyk, Shyam Narayanan, Ronitt Rubinfeld, Sandeep Silwal, Tal Wagner:
Learning-based Support Estimation in Sublinear Time.

Malik Tiomoko, Hafiz Tiomoko Ali, Romain Couillet:
Deciphering and Optimizing Multi-Task Learning: a Random Matrix Approach.

Nicola De Cao, Gautier Izacard, Sebastian Riedel, Fabio Petroni:
Autoregressive Entity Retrieval.

Faruk Ahmed, Yoshua Bengio, Harm van Seijen, Aaron C. Courville:
Systematic generalisation with group invariant predictions.

Max Olan Smith, Thomas Anthony, Michael P. Wellman:
Iterative Empirical Game Solving via Single Policy Best Response.

Da Xu, Yuting Ye, Chuanwei Ruan:
Understanding the role of importance weighting for deep learning.

-     Aditya Krishna Menon, Sadeep Jayasumana, Ankit Singh Rawat, Himanshu Jain, Andreas Veit, Sanjiv Kumar:
Long-tail learning via logit adjustment.
-     Guan-Horng Liu, Tianrong Chen, Evangelos A. Theodorou:
DDPNOpt: Differential Dynamic Programming Neural Optimizer.
-     Yikai Zhang, Songzhu Zheng, Pengxiang Wu, Mayank Goswami, Chao Chen:
Learning with Feature-Dependent Label Noise: A Progressive Approach.
-     Xinran Wang, Yu Xiang, Jun Gao, Jie Ding:
Information Laundering for Model Privacy.
-     Rui Zhao, Yang Gao, Pieter Abbeel, Volker Tresp, Wei Xu:
Mutual Information State Intrinsic Control.
-     Taiji Suzuki, Shunta Akiyama:
Benefit of deep learning with non-convex noisy gradient descent: Provable excess risk bound and superiority to kernel methods.
-     Linjun Zhang, Zhun Deng, Kenji Kawaguchi, Amirata Ghorbani, James Zou:
How Does Mixup Help With Robustness and Generalization?
-     Pratyush Maini, Mohammad Yaghini, Nicolas Papernot:
Dataset Inference: Ownership Resolution in Machine Learning.
-     Alexander Vargo, Fan Zhang, Mikhail Yurochkin, Yuekai Sun:
Individually Fair Gradient Boosting.
-     Shengyu Zhao, Jonathan Cui , Yilun Sheng, Yue Dong, Xiao Liang, Eric I-Chao Chang, Yan Xu:
Large Scale Image Completion via Co-Modulated Generative Adversarial Networks.
-     Nicklas Hansen, Rishabh Jangir, Yu Sun, Guillem Alenyà, Pieter Abbeel, Alexei A. Efros, Lerrel Pinto, Xiaolong Wang:
Self-Supervised Policy Adaptation during Deployment.
-     Pierre Foret, Ariel Kleiner, Hossein Mobahi, Behnam Neyshabur:
Sharpness-aware Minimization for Efficiently Improving Generalization.
-     Yoav Levine, Barak Lenz, Opher Lieber, Omri Abend, Kevin Leyton-Brown, Moshe Tennenholtz, Yoav Shoham:
PMI-Masking: Principled masking of correlated spans.
-     Rewon Child:
Very Deep VAEs Generalize Autoregressive Models and Can Outperform Them on Images.
-     Max Schwarzer, Ankesh Anand, Rishab Goel, R. Devon Hjelm, Aaron C. Courville, Philip Bachman:
Data-Efficient Reinforcement Learning with Self-Predictive Representations.
-     Xavier Puig, Tianmin Shu, Shuang Li, Zilin Wang, Yuan-Hong Liao, Joshua B. Tenenbaum, Sanja Fidler, Antonio Torralba:
Watch-And-Help: A Challenge for Social Perception and Human-AI Collaboration.
-     Yu Tian, Jian Ren, Menglei Chai, Kyle Olszewski, Xi Peng, Dimitris N. Metaxas, Sergey Tulyakov:
A Good Image Generator Is What You Need for High-Resolution Video Synthesis.
-     Siyi Hu, Fengda Zhu, Xiaojun Chang, Xiaodan Liang:
UPDeT: Universal Multi-agent RL via Policy Decoupling with Transformers.
-     Augustus Odena, Kensen Shi, David Bieber, Rishabh Singh, Charles Sutton, Hanjun Dai:
BUSTLE: Bottom-Up Program Synthesis Through Learning-Guided Exploration.
-     Yang Bai, Yuyuan Zeng, Yong Jiang, Shu-Tao Xia, Xingjun Ma, Yisen Wang:
Improving Adversarial Robustness via Channel-wise Activation Suppressing.
-     Ruosong Wang, Dean P. Foster, Sham M. Kakade:
What are the Statistical Limits of Offline RL with Linear Function Approximation?
-     Hanxun Huang, Xingjun Ma, Sarah Monazam Erfani, James Bailey, Yisen Wang:
Unlearnable Examples: Making Personal Data Unexploitable.
-     Tobias Pfaff, Meire Fortunato, Alvaro Sanchez-Gonzalez, Peter W. Battaglia:
Learning Mesh-Based Simulation with Graph Networks.
-     Xiu Su, Shan You, Tao Huang, Fei Wang, Chen Qian, Changshui Zhang, Chang Xu:
Locally Free Weight Sharing for Network Width Search.
-     Xiuyuan Cheng, Zichen Miao, Qiang Qiu:
Graph Convolution with Low-rank Learnable Local Filters.
-     Wonseok Jeon, Chen-Yang Su, Paul Barde, Thang Doan, Derek Nowrouzezahrai, Joelle Pineau:
Regularized Inverse Reinforcement Learning.
-     Michael Sejr Schlichtkrull, Nicola De Cao, Ivan Titov:
Interpreting Graph Neural Networks for NLP With Differentiable Edge Masking.
-     Nils Lukas, Yuxuan Zhang, Florian Kerschbaum:
Deep Neural Network Fingerprinting by Conferrable Adversarial Examples.
-     Dequan Wang, Evan Shelhamer, Shaoteng Liu, Bruno A. Olshausen, Trevor Darrell:
Tent: Fully Test-Time Adaptation by Entropy Minimization.



Nurit Spingarn, Ron Banner, Tomer Michaeli:
GAN "Steerability" without optimization.



Omer Yair, Tomer Michaeli:
Contrastive Divergence Learning is a Time Reversal Adversarial Game.



Xiaoling Hu, Yusu Wang, Fuxin Li, Dimitris Samaras, Chao Chen:
Topology-Aware Segmentation Using Discrete Morse Theory.



Zhen Qin, Le Yan, Honglei Zhuang, Yi Tay, Rama Kumar Pasumarthi, Xuanhui Wang, Michael Bendersky, Marc Najork:
Are Neural Rankers still Outperformed by Gradient Boosted Decision Trees?



Yoshua Bengio, Prateek Gupta, Tegan Maharaj, Nasim Rahaman, Martin Weiss, Tristan Deleu, Eilif Benjamin Müller, Meng Qu, Victor Schmidt, Pierre-Luc St-Charles, Hannah Alsdurf, Olexa Bilaniuk, David L. Buckeridge, Gaétan Marceau-Caron, Pierre Luc Carrier, Joumana Ghosn, Satya Ortiz-Gagne, Christopher J. Pal, Irina Rish, Bernhard Schölkopf, Abhinav Sharma, Jian Tang, Andrew Williams:
Predicting Infectiousness for Proactive Contact Tracing.



Zhuang Liu, Xuanlin Li, Bingyi Kang, Trevor Darrell:
Regularization Matters in Policy Optimization - An Empirical Study on Continuous Control.



Sejun Park, Chulhee Yun, Jaeho Lee, Jinwoo Shin:
Minimum Width for Universal Approximation.



Xinshuai Dong, Anh Tuan Luu, Rongrong Ji, Hong Liu:
Towards Robustness Against Natural Language Word Substitutions.



Kenji Kawaguchi:
On the Theory of Implicit Deep Learning: Global Convergence with Implicit Layers.



Giovanni Paolini, Ben Athiwaratkun, Jason Krone, Jie Ma, Alessandro Achille, Rishita Anubhai, Cícero Nogueira dos Santos, Bing Xiang, Stefano Soatto:
Structured Prediction as Translation between Augmented Natural Languages.



Amartya Sanyal, Puneet K. Dokania, Varun Kanade, Philip H. S. Torr:
How Benign is Benign Overfitting ?



Sanjeevan Ahilan, Peter Dayan:
Correcting experience replay for multi-agent communication.



Taylor Whittington Webb, Ishan Sinha, Jonathan D. Cohen:
Emergent Symbols through Binding in External Memory.



Naoyuki Terashita, Hiroki Ohashi, Yuichi Nonaka, Takashi Kanemaru:
Influence Estimation for Generative Adversarial Networks.



Zhiao Huang, Yuanming Hu, Tao Du, Siyuan Zhou, Hao Su, Joshua B. Tenenbaum, Chuang Gan:
PlasticineLab: A Soft-Body Manipulation Benchmark with Differentiable Physics.



Cheng Lu, Jianfei Chen, Chongxuan Li, Qiuhan Wang, Jun Zhu:
Implicit Normalizing Flows.



Mandela Patrick, Po-Yao Huang, Yuki Markus Asano, Florian Metze, Alexander G. Hauptmann, João F. Henriques, Andrea Vedaldi:
Support-set bottlenecks for video-text representation learning.



Deunsol Yoon, Sungsoon Hong, Byung-Jun Lee, Kee-Eung Kim:
Winning the L2RPN Challenge: Power Grid Management via Semi-Markov Afterstate Actor-Critic.



Nils Wandel, Michael Weinmann, Reinhard Klein:
Learning Incompressible Fluid Dynamics from Scratch - Towards Fast, Differentiable Fluid Models that Generalize.



Elliot Meyerson, Risto Miikkulainen:
The Traveling Observer Model: Multi-task Learning Through Spatial Variable Embeddings.



Felix Hill, Olivier Tieleman, Tamara von Glehn, Nathaniel Wong, Hamza Merzic, Stephen Clark:
Grounded Language Learning Fast and Slow.



Xudong Wang, Long Lian, Zhongqi Miao, Ziwei Liu, Stella X. Yu:
Long-tailed Recognition by Routing Diverse Distribution-Aware Experts.



Florian Tramèr, Dan Boneh:
Differentially Private Learning Needs Better Features (or Much More Data).



Anand Gopalakrishnan, Sjoerd van Steenkiste, Jürgen Schmidhuber:
Unsupervised Object Keypoint Learning using Local Spatial Predictability.



Sebastian Farquhar, Yarin Gal, Tom Rainforth:
On Statistical Bias In Active Learning: How and When to Fix It.



Tolga Ergen, Mert Pilancı:
Implicit Convex Regularizers of CNN Architectures: Convex Optimization of Two- and Three-Layer Networks in Polynomial Time.



Konstantin-Klemens Lurz, Mohammad Bashiri, Konstantin Willeke, Akshay Kumar Jagadish, Eric Wang, Edgar Y. Walker, Santiago A. Cadena, Taliah Muhammad, Erick Cobos, Andreas S. Tolias, Alexander S. Ecker, Fabian H. Sinz:
Generalization in data-driven models of primary visual cortex.



Markus Norman Rabe, Dennis Lee, Kshitij Bansal, Christian Szegedy:
Mathematical Reasoning via Self-supervised Skip-tree Training.



Aston Zhang, Yi Tay, Shuai Zhang, Alvin Chan, Anh Tuan Luu, Siu Cheung Hui, Jie Fu:
Beyond Fully-Connected Layers with Quaternions: Parameterization of Hypercomplex Multiplications with 1/n Parameters.



Khai Nguyen, Nhat Ho, Tung Pham, Hung Bui:
Distributional Sliced-Wasserstein and Applications to Generative Modeling.



Yu Sun, Jiaming Liu, Yiran Sun, Brendt Wohlberg, Ulugbek Kamilov:
Async-RED: A Provably Convergent Asynchronous Block Parallel Stochastic Method using Deep Denoising Priors.



Aayam Kumar Shrestha, Stefan Lee, Prasad Tadepalli, Alan Fern:
DeepAveragers: Offline Reinforcement Learning By Solving Derived Non-Parametric MDPs.



Bowen Jing, Stephan Eismann, Patricia Suriana, Raphael John Lamarre Townshend, Ron O. Dror:
Learning from Protein Structure with Geometric Vector Perceptrons.



Fumihiro Sasaki, Ryota Yamashina:
Behavioral Cloning from Noisy Demonstrations.



Haoyu Ma , Tianlong Chen, Ting-Kuei Hu, Chenyu You, Xiaohui Xie, Zhangyang Wang:
Undistillable: Making A Nasty Teacher That CANNOT teach students.



Kashif Rasul, Abdul-Saboor Sheikh, Ingmar Schuster, Urs M. Bergmann, Roland Vollgraf:
Multivariate Probabilistic Time Series Forecasting via Conditioned Normalizing Flows.



Denis Yarats, Ilya Kostrikov, Rob Fergus:
Image Augmentation Is All You Need: Regularizing Deep Reinforcement Learning from Pixels.



Chaojian Li, Zhongzhi Yu, Yonggan Fu, Yongan Zhang, Yang Zhao , Haoran You, Qixuan Yu, Yue Wang, Cong Hao, Yingyan Lin:
HW-NAS-Bench: Hardware-Aware Neural Architecture Search Benchmark.



Jacob Menick, Erich Elsen, Utku Evci, Simon Osindero, Karen Simonyan, Alex Graves:
Practical Real Time Recurrent Learning with a Sparse Approximation.



Hao Peng, Nikolaos Pappas, Dani Yogatama, Roy Schwartz, Noah A. Smith, Lingpeng Kong:
Random Feature Attention.



Ekdeep Singh Lubana, Robert P. Dick:
A Gradient Flow Framework For Analyzing Network Pruning.



Anirudh Goyal, Alex Lamb, Jordan Hoffmann, Shagun Sodhani, Sergey Levine, Yoshua Bengio, Bernhard Schölkopf:
Recurrent Independent Mechanisms.



Phillip Pope, Chen Zhu, Ahmed Abdelkader, Micah Goldblum, Tom Goldstein:
The Intrinsic Dimension of Images and Its Impact on Learning.



Anastasios Nikolas Angelopoulos, Stephen Bates, Michael I. Jordan, Jitendra Malik:
Uncertainty Sets for Image Classifiers using Conformal Prediction.



Akinori F. Ebihara, Taiki Miyagawa, Kazuyuki Sakurai, Hitoshi Imaoka:
Sequential Density Ratio Estimation for Simultaneous Optimization of Speed and Accuracy.



Jun Han, Martin Renqiang Min, Ligong Han, Li Erran Li, Xuan Zhang:
Disentangled Recurrent Wasserstein Autoencoder.



Daniel Hsu, Ziwei Ji, Matus Telgarsky, Lan Wang:
Generalization bounds via distillation.



Yanzhi Chen, Dinghuai Zhang, Michael U. Gutmann, Aaron C. Courville, Zhanxing Zhu:
Neural Approximate Sufficient Statistics for Implicit Models.



Sanghyun Hong, Yigitcan Kaya, Ionut-Vlad Modoranu, Tudor Dumitras:
A Panda? No, It's a Sloth: Slowdown Attacks on Adaptive Multi-Exit Neural Network Inference.



Asher Trockman, J. Zico Kolter:
Orthogonalizing Convolutional Layers with the Cayley Transform.



Irwan Bello:
LambdaNetworks: Modeling long-range Interactions without Attention.



Bilal Alsallakh, Narine Kokhlikyan, Vivek Miglani, Jun Yuan, Orion Reblitz-Richardson:
Mind the Pad - CNNs Can Develop Blind Spots.



Dong Bok Lee, Dongchan Min, Seanie Lee, Sung Ju Hwang:
Meta-GMVAE: Mixture of Gaussian VAE for Unsupervised Meta-Learning.



Aymeric Fromherz, Klas Leino, Matt Fredrikson, Bryan Parno, Corina S. Pasareanu:
Fast Geometric Projections for Local Robustness Certification.



Eli Ovits, Lior Wolf:
Fidelity-based Deep Adiabatic Scheduling.



Stanislav Morozov, Andrey Voynov, Artem Babenko:
On Self-Supervised Image Representations for GAN Evaluation.



Shangqing Liu, Yu Chen, Xiaofei Xie, Jing Kai Sio, Yang Liu:
Retrieval-Augmented Generation for Code Summarization via Hybrid GNN.



Andrii Zadaianchuk, Maximilian Seitzer, Georg Martius:
Self-supervised Visual Reinforcement Learning with Object-centric Representations.



Dominik Schmidt, Georgia Koppe, Zahra Monfared, Max Beutelspacher, Daniel Durstewitz:
Identifying nonlinear dynamical systems with multiple time scales and long-range dependencies.



He Zhao, Dinh Phung, Viet Huynh, Trung Le, Wray L. Buntine :
Neural Topic Model via Optimal Transport.



Aashaka Shah, Chao-Yuan Wu, Jayashree Mohan, Vijay Chidambaram , Philipp Krähenbühl:
Memory Optimization for Deep Networks.



Gege Qi, Lijun Gong, Yibing Song, Kai Ma, Yefeng Zheng:
Stabilized Medical Image Attacks.



Adam Gleave, Michael Dennis, Shane Legg, Stuart Russell, Jan Leike:
Quantifying Differences in Reward Functions.



Yutong Xie, Chence Shi, Hao Zhou, Yuwei Yang, Weinan Zhang, Yong Yu, Lei Li:
MARS: Markov Molecular Sampling for Multi-objective Drug Discovery.



Pim de Haan, Maurice Weiler, Taco Cohen, Max Welling:
Gauge Equivariant Mesh CNNs: Anisotropic convolutions on geometric graphs.



Naichen Shi, Dawei Li, Mingyi Hong, Ruoyu Sun:
RMSprop converges with proper hyper-parameter.

Poster Presentations



Yunsheng Li, Yinpeng Chen, Xiyang Dai, Mengchen Liu, Dongdong Chen, Ye Yu, Lu Yuan, Zicheng Liu, Mei Chen, Nuno Vasconcelos:
Revisiting Dynamic Convolution via Matrix Decomposition.



Philipp Liznerski, Lukas Ruff, Robert A. Vandermeulen , Billy Joe Franks, Marius Kloft, Klaus-Robert Müller:
Explainable Deep One-Class Classification.



Qiyu Wu, Chen Xing, Yatao Li, Guolin Ke, Di He, Tie-Yan Liu:
Taking Notes on the Fly Helps Language Pre-Training.



Alejandro Pimentel-Alarcón, Daniel L. Pimentel-Alarcón:
Mixed-Features Vectors and Subspace Splitting.



Huan Wang, Can Qin, Yulun Zhang, Yun Fu:
Neural Pruning via Growing Regularization.



Xiufeng Yang, Tanuj Kr Aasawat, Kazuki Yoshizoe:
Practical Massively Parallel Monte-Carlo Tree Search Applied to Molecular Design.



Yangming Li, Lemao Liu, Shuming Shi:
Empirical Analysis of Unlabeled Entity Problem in Named Entity Recognition.



Sam Buchanan, Dar Gilboa, John Wright:
Deep Networks and the Multiple Manifold Problem.



Jing Yang, Brais Martínez, Adrian Bulat, Georgios Tzimiropoulos:
Knowledge distillation via softmax regression representation learning.



Urvashi Khandelwal, Angela Fan, Dan Jurafsky, Luke Zettlemoyer, Mike Lewis:
Nearest Neighbor Machine Translation.



Renkun Ni, Hong-Min Chu, Oscar Castañeda, Ping-yeh Chiang, Christoph Studer, Tom Goldstein:
WrapNet: Neural Net Inference with Ultra-Low-Precision Arithmetic.



Mengye Ren, Michael Louis Iuzzolino, Michael Curtis Mozer, Richard S. Zemel:
Wandering within a world: Online contextualized few-shot learning.



Simon Shaolei Du, Wei Hu, Sham M. Kakade, Jason D. Lee, Qi Lei:
Few-Shot Learning via Learning the Representation, Provably.



Ke Sun, Zhanxing Zhu, Zhouchen Lin:
AdaGCN: Adaboosting Graph Convolutional Networks into Deep Models.



Alon Talmor, Ori Yoran, Amnon Catav, Dan Lahav, Yizhong Wang, Akari Asai, Gabriel Ilharco, Hannaneh Hajishirzi, Jonathan Berant:
MultiModalQA: complex question answering over text, tables and images.



Liran Katzir, Gal Elidan, Ran El-Yaniv:
Net-DNF: Effective Deep Modeling of Tabular Data.



Preetum Nakkiran, Prayaag Venkat, Sham M. Kakade, Tengyu Ma:
Optimal Regularization can Mitigate Double Descent.



- Hieu Pham, Xinyi Wang, Yiming Yang, Graham Neubig:
Meta Back-Translation.
- Jiayi Shen, Xiaohan Chen, Howard Heaton, Tianlong Chen, Jialin Liu, Wotao Yin, Zhangyang Wang:
Learning A Minimax Optimizer: A Pilot Study.
- Leon Lang , Maurice Weiler:
A Wigner-Eckart Theorem for Group Equivariant Convolution Kernels.
- Alex Tamkin, Mike Wu, Noah D. Goodman:
Viewmaker Networks: Learning Views for Unsupervised Representation Learning.
- Joan Puigcerver, Carlos Riquelme Ruiz, Basil Mustafa, Cédric Renggli, André Susano Pinto, Sylvain Gelly, Daniel Keysers, Neil Houlsby:
Scalable Transfer Learning with Expert Models.
- Abhishek Sinha, Kumar Ayush, Jiaming Song, Burak Uzkent, Hongxia Jin, Stefano Ermon:
Negative Data Augmentation.
- Sahil Singla, Soheil Feizi:
Fantastic Four: Differentiable and Efficient Bounds on Singular Values of Convolution Layers.
- Yanru Qu, Dinghan Shen, Yelong Shen, Sandra Sajeev, Weizhu Chen, Jiawei Han:
CoDA: Contrast-enhanced and Diversity-promoting Data Augmentation for Natural Language Understanding.
- Aniruddh Raghu, Maithra Raghu, Simon Kornblith, David Duvenaud, Geoffrey E. Hinton:
Teaching with Commentaries.
- Kevin J. Liang, Weituo Hao, Dinghan Shen, Yufan Zhou, Weizhu Chen, Changyou Chen, Lawrence Carin:
MixKD: Towards Efficient Distillation of Large-scale Language Models.
- Pengyu Cheng, Weituo Hao, Siyang Yuan, Shijing Si, Lawrence Carin:
FairFil: Contrastive Neural Debiasing Method for Pretrained Text Encoders.
- Marc Anton Finzi, Roberto Bondesan, Max Welling:
Probabilistic Numeric Convolutional Neural Networks.
- Eran Malach, Shai Shalev-Shwartz:
Computational Separation Between Convolutional and Fully-Connected Networks.
- Marius Mosbach, Maksym Andriushchenko, Dietrich Klakow:
On the Stability of Fine-tuning BERT: Misconceptions, Explanations, and Strong Baselines.
- Rabeeh Karimi Mahabadi, Yonatan Belinkov, James Henderson :
Variational Information Bottleneck for Effective Low-Resource Fine-Tuning.
- Jonas Geiping, Liam H. Fowl, W. Ronny Huang, Wojciech Czaja, Gavin Taylor, Michael Moeller, Tom Goldstein:
Witches' Brew: Industrial Scale Data Poisoning via Gradient Matching.
- Pengcheng He, Xiaodong Liu, Jianfeng Gao, Weizhu Chen:
Deberta: decoding-Enhanced Bert with Disentangled Attention.
- Yining Wang, Ruosong Wang, Simon Shaolei Du, Akshay Krishnamurthy:
Optimism in Reinforcement Learning with Generalized Linear Function Approximation.
- Elan Sopher Markowitz, Keshav Balasubramanian, Mehrnoosh Mirtaheri, Sami Abu-El-Haija, Bryan Perozzi, Greg Ver Steeg, Aram Galstyan:
Graph Traversal with Tensor Functionals: A Meta-Algorithm for Scalable Learning.
- Gaurav Shrivastava, Abhinav Shrivastava:
Diverse Video Generation using a Gaussian Process Trigger.
- Patrick Kidger, Terry J. Lyons:
Signatory: differentiable computations of the signature and logsignature transforms, on both CPU and GPU.
- Junnan Li, Caiming Xiong, Steven C. H. Hoi:
MoPro: Webly Supervised Learning with Momentum Prototypes.
- Lu Liu, William L. Hamilton, Guodong Long, Jing Jiang, Hugo Larochelle:
A Universal Representation Transformer Layer for Few-Shot Image Classification.
- Robert Dadashi, Léonard Hussenot, Matthieu Geist, Olivier Pietquin:
Primal Wasserstein Imitation Learning.
- Eric Wong , J. Zico Kolter:
Learning perturbation sets for robust machine learning.
- Jiaqi Ma, Bo Chang, Xuefei Zhang, Qiaozhu Mei:
CopulaGNN: Towards Integrating Representational and Correlational Roles of Graphs in Graph Neural Networks.
- Andy Shih, Arjun Sawhney, Jovana Kondic, Stefano Ermon, Dorsa Sadigh:
On the Critical Role of Conventions in Adaptive Human-AI Collaboration.
- Uri Alon, Eran Yahav:
On the Bottleneck of Graph Neural Networks and its Practical Implications.
- Kyle Aitken, Vinay Venkatesh Ramasesh, Ankush Garg, Yuan Cao, David Sussillo, Niru Maheswaranathan:
The geometry of integration in text classification RNNs.



- Jeremy Cohen, Simran Kaur, Yuanzhi Li, J. Zico Kolter, Ameet Talwalkar:
Gradient Descent on Neural Networks Typically Occurs at the Edge of Stability.
- Ossama Ahmed, Frederik Träuble, Anirudh Goyal, Alexander Neitz, Manuel Wuthrich, Yoshua Bengio, Bernhard Schölkopf, Stefan Bauer:
CausalWorld: A Robotic Manipulation Benchmark for Causal Structure and Transfer Learning.
- Kartik Ahuja, Jun Wang, Amit Dhurandhar, Karthikeyan Shanmugam, Kush R. Varshney:
Empirical or Invariant Risk Minimization? A Sample Complexity Perspective.
- Subham Sekhar Sahoo, Subhashini Venugopalan, Li Li, Rishabh Singh, Patrick Riley:
Scaling Symbolic Methods using Gradients for Neural Model Explanation.
- Brandon Cui, Yinlam Chow, Mohammad Ghavamzadeh:
Control-Aware Representations for Model-based Reinforcement Learning.
- Benjamin Eysenbach, Ruslan Salakhutdinov, Sergey Levine:
C-Learning: Learning to Achieve Goals via Recursive Classification.
- Preetum Nakkiran, Behnam Neyshabur, Hanie Sedghi:
The Deep Bootstrap Framework: Good Online Learners are Good Offline Generalizers.
- Matthew Willetts, Alexander Camuto, Tom Rainforth, Stephen J. Roberts, Christopher C. Holmes:
Improving VAEs' Robustness to Adversarial Attack.
- Kiana Ehsani, Daniel Gordon, Thomas Hai Dang Nguyen, Rozbeh Mottaghi, Ali Farhadi:
What Can You Learn From Your Muscles? Learning Visual Representation from Human Interactions.
- Ali Ayub, Alan R. Wagner:
EEC: Learning to Encode and Regenerate Images for Continual Learning.
- Jiaqi Yang, Wei Hu, Jason D. Lee, Simon Shaolei Du:
Impact of Representation Learning in Linear Bandits.
- Tsz-Him Cheung , Dit-Yan Yeung:
MODALS: Modality-agnostic Automated Data Augmentation in the Latent Space.
- Sina Alemohammad, Zichao Wang, Randall Balestrieri, Richard G. Baraniuk:
The Recurrent Neural Tangent Kernel.
- Chris Cannella, Mohammadreza Soltani, Vahid Tarokh:
Projected Latent Markov Chain Monte Carlo: Conditional Sampling of Normalizing Flows.
- Aviv Navon, Aviv Shamsian, Ethan Fetaya, Gal Chechik:
Learning the Pareto Front with Hypernetworks.
- Ali Harakeh, Steven L. Waslander:
Estimating and Evaluating Regression Predictive Uncertainty in Deep Object Detectors.
- Yuli Slavutsky, Yuval Benjamini:
Predicting Classification Accuracy When Adding New Unobserved Classes.
- Yuhang Li, Ruihao Gong, Xu Tan, Yang Yang, Peng Hu, Qi Zhang, Fengwei Yu, Wei Wang, Shi Gu:
BRECQ: Pushing the Limit of Post-Training Quantization by Block Reconstruction.
- Will Sussman Grathwohl, Jacob Jin Kelly, Milad Hashemi, Mohammad Norouzi, Kevin Swersky, David Duvenaud:
No MCMC for me: Amortized sampling for fast and stable training of energy-based models.
- Daya Guo, Shuo Ren, Shuai Lu, Zhangyin Feng, Duyu Tang, Shujie Liu, Long Zhou, Nan Duan, Alexey Svyatkovskiy, Shengyu Fu, Michele Tufano, Shao Kun Deng, Colin B. Clement, Dawn Drain, Neel Sundaresan, Jian Yin, Dixin Jiang, Ming Zhou:
GraphCodeBERT: Pre-training Code Representations with Data Flow.
- Homanga Bharadhwaj, Aviral Kumar, Nicholas Rhinehart, Sergey Levine, Florian Shkurti, Animesh Garg:
Conservative Safety Critics for Exploration.
- Linfeng Zhang, Kaisheng Ma :
Improve Object Detection with Feature-based Knowledge Distillation: Towards Accurate and Efficient Detectors.
- Da Xu, Chuanwei Ruan, Evren Körpeoglu, Sushant Kumar, Kannan Achan:
A Temporal Kernel Approach for Deep Learning with Continuous-time Information.
- Yamini Bansal, Gal Kaplun, Boaz Barak:
For self-supervised learning, Rationality implies generalization, provably.
- Dongkwan Kim, Alice Oh:
How to Find Your Friendly Neighborhood: Graph Attention Design with Self-Supervision.
- Ricards Marcinkevics , Julia E. Vogt:
Interpretable Models for Granger Causality Using Self-explaining Neural Networks.
- Allan Zhou, Tom Knowles, Chelsea Finn:
Meta-learning Symmetries by Reparameterization.
- Saehyung Lee, Changhwa Park, Hyungyu Lee, Jihun Yi, Jonghyun Lee, Sungroh Yoon:
Removing Undesirable Feature Contributions Using Out-of-Distribution Data.
- Justin Bayer, Maximilian Soelch, Atanas Mirchev, Baris Kayalibay, Patrick van der Smagt :
Mind the Gap when Conditioning Amortised Inference in Sequential Latent-Variable Models.



- David Holzmüller:
On the Universality of the Double Descent Peak in Ridgeless Regression.
- Ching-Yao Chuang, Youssef Mroueh:
Fair Mixup: Fairness via Interpolation.
- Yao-Hung Hubert Tsai, Yue Wu, Ruslan Salakhutdinov, Louis-Philippe Morency:
Self-supervised Learning from a Multi-view Perspective.
- Samuel Lavoie-Marchildon, Faruk Ahmed, Aaron C. Courville:
Integrating Categorical Semantics into Unsupervised Domain Translation.
- Louis Thiry, Michael Arbel, Eugene Belilovsky, Edouard Oyallon:
The Unreasonable Effectiveness of Patches in Deep Convolutional Kernels Methods.
- Wenhui Chen, Ming-Wei Chang, Eva Schlinger, William Yang Wang, William W. Cohen:
Open Question Answering over Tables and Text.
- Kazuaki Hanawa, Sho Yokoi, Satoshi Hara, Kentaro Inui:
Evaluation of Similarity-based Explanations.
- Zeke Xie, Issei Sato, Masashi Sugiyama:
A Diffusion Theory For Deep Learning Dynamics: Stochastic Gradient Descent Exponentially Favors Flat Minima.
- Zixiang Chen, Yuan Cao, Difan Zou, Quanquan Gu:
How Much Over-parameterization Is Sufficient to Learn Deep ReLU Networks?
- Jad Rahme, Samy Jelassi, S. Matthew Weinberg :
Auction Learning as a Two-Player Game.
- Huan Zhang, Hongge Chen, Duane S. Boning, Cho-Jui Hsieh:
Robust Reinforcement Learning on State Observations with Learned Optimal Adversary.
- Shauharda Khadka, Estelle Aflalo, Mattias Marder, Avrech Ben-David, Santiago Miret, Shie Mannor, Tamir Hazan, Hanlin Tang, Somdeb Majumdar:
Optimizing Memory Placement using Evolutionary Graph Reinforcement Learning.
- Ruihan Yang, Yibo Yang, Joseph Marino, Stephan Mandt:
Hierarchical Autoregressive Modeling for Neural Video Compression.
- Amanda Bower, Hamid Eftekhari, Mikhail Yurochkin, Yuekai Sun:
Individually Fair Rankings.
- Minkai Xu, Shitong Luo, Yoshua Bengio, Jian Peng, Jian Tang:
Learning Neural Generative Dynamics for Molecular Conformation Generation.
- Jan Hendrik Metzen, Maksym Yatsura:
Efficient Certified Defenses Against Patch Attacks on Image Classifiers.
- Arda Sahiner, Morteza Mardani, Batu Ozturkler, Mert Pilanci, John M. Pauly:
Convex Regularization behind Neural Reconstruction.
- Jiawang Bai, Baoyuan Wu, Yong Zhang, Yiming Li, Zhifeng Li, Shu-Tao Xia:
Targeted Attack against Deep Neural Networks via Flipping Limited Weight Bits.
- Thomas M. Sutter , Imant Daunhauer, Julia E. Vogt:
Generalized Multimodal ELBO.
- Daniele Bracale, Stefano Favaro, Sandra Fortini, Stefano Peluchetti:
Large-width functional asymptotics for deep Gaussian neural networks.
- El Mahdi El Mhamdi, Rachid Guerraoui , Sébastien Rouault:
Distributed Momentum for Byzantine-resilient Stochastic Gradient Descent.
- Mangal Prakash, Alexander Krull, Florian Jug:
Fully Unsupervised Diversity Denoising with Convolutional Variational Autoencoders.
- Aviv Navon, Idan Achituv, Haggai Maron, Gal Chechik, Ethan Fetaya:
Auxiliary Learning by Implicit Differentiation.
- Dan A. Calian, Daniel J. Mankowitz, Tom Zahavy, Zhongwen Xu, Junhyuk Oh, Nir Levine, Timothy A. Mann:
Balancing Constraints and Rewards with Meta-Gradient D4PG.
- Yannis Flet-Berliac, Johan Ferret, Olivier Pietquin, Philippe Preux, Matthieu Geist:
Adversarially Guided Actor-Critic.
- Xiangxiang Chu, Xiaoxing Wang, Bo Zhang, Shun Lu, Xiaolin Wei, Junchi Yan:
DARTS-: Robustly Stepping out of Performance Collapse Without Indicators.
- Anna Golubeva, Guy Gur-Ari, Behnam Neyshabur:
Are wider nets better given the same number of parameters?
- Shikuang Deng , Shi Gu:
Optimal Conversion of Conventional Artificial Neural Networks to Spiking Neural Networks.
- Alberto Bietti, Francis R. Bach:
Deep Equals Shallow for ReLU Networks in Kernel Regimes.



- Chen Cai, Dingkang Wang, Yusu Wang:
Graph Coarsening with Neural Networks.
- Reinhard Heckel, Fatih Furkan Yilmaz:
Early Stopping in Deep Networks: Double Descent and How to Eliminate it.
- Feng Zhou, Yixuan Zhang, Jun Zhu:
Efficient Inference of Flexible Interaction in Spiking-neuron Networks.
- Alexandre Ramé, Matthieu Cord:
DICE: Diversity in Deep Ensembles via Conditional Redundancy Adversarial Estimation.
- Yanbang Wang, Yen-Yu Chang, Yunyu Liu, Jure Leskovec, Pan Li:
Inductive Representation Learning in Temporal Networks via Causal Anonymous Walks.
- Yuji Roh, Kangwook Lee, Steven Euijong Whang, Changho Suh:
FairBatch: Batch Selection for Model Fairness.
- Byung-Jun Lee, Jongmin Lee, Kee-Eung Kim:
Representation Balancing Offline Model-based Reinforcement Learning.
- Wei Deng, Qi Feng, Georgios Karagiannis, Guang Lin, Faming Liang:
Accelerating Convergence of Replica Exchange Stochastic Gradient MCMC via Variance Reduction.
- Jacob Buckman, Carles Gelada, Marc G. Bellemare:
The Importance of Pessimism in Fixed-Dataset Policy Optimization.
- Carl Allen, Ivana Balazevic, Timothy M. Hospedales:
Interpreting Knowledge Graph Relation Representation from Word Embeddings.
- Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Lukas Gruber, Markus Holzleitner, Thomas Adler, David P. Kreil, Michael K. Kopp, Günter Klambauer, Johannes Brandstetter, Sepp Hochreiter:
Hopfield Networks is All You Need.
- Cheng Wang, Carolin Lawrence, Mathias Niepert:
Uncertainty Estimation and Calibration with Finite-State Probabilistic RNNs.
- Vaishnav Nagarajan, Anders Andreassen, Behnam Neyshabur:
Understanding the failure modes of out-of-distribution generalization.
- Shuhei Kurita, Kyunghyun Cho:
Generative Language-Grounded Policy in Vision-and-Language Navigation with Bayes' Rule.
- Avik Pal, Jonah Philion, Yuan-Hong Liao, Sanja Fidler:
Emergent Road Rules In Multi-Agent Driving Environments.
- Alexander Korotin, Vage Egiazarian, Arip Asadulaev, Alexander Safin, Evgeny Burnaev:
Wasserstein-2 Generative Networks.
- Yanchao Sun, Da Huo, Furong Huang:
Vulnerability-Aware Poisoning Mechanism for Online RL with Unknown Dynamics.
- Francesco Tonolini, Pablo Garcia Moreno, Andreas C. Damianou, Roderick Murray-Smith:
Tomographic Auto-Encoder: Unsupervised Bayesian Recovery of Corrupted Data.
- William Taylor Bakst, Nobuyuki Morioka, Erez Louidor:
Monotonic Kronecker-Factored Lattice.
- Neil Zeghidour, Olivier Teboul, Félix de Chaumont Quirky, Marco Tagliasacchi:
LEAF: A Learnable Frontend for Audio Classification.
- Maruan Al-Shedivat, Jennifer Gillenwater, Eric P. Xing, Afshin Rostamizadeh:
Federated Learning via Posterior Averaging: A New Perspective and Practical Algorithms.
- Daochen Zha, Wenye Ma, Lei Yuan, Xia Hu, Ji Liu:
Rank the Episodes: A Simple Approach for Exploration in Procedurally-Generated Environments.
- Kapil Vaidya, Eric Knorr, Michael Mitzenmacher, Tim Kraska:
Partitioned Learned Bloom Filters.
- Lee Xiong, Chenyan Xiong, Ye Li, Kwok-Fung Tang, Jialin Liu, Paul N. Bennett, Junaid Ahmed, Arnold Overwijk:
Approximate Nearest Neighbor Negative Contrastive Learning for Dense Text Retrieval.
- Lucio M. Dery, Yann N. Dauphin, David Grangier:
Auxiliary Task Update Decomposition: the Good, the Bad and the neutral.
- Vikash Sehwag, Mung Chiang, Prateek Mittal:
SSD: A Unified Framework for Self-Supervised Outlier Detection.
- Valerie Chen, Abhinav Gupta, Kenneth Marino:
Ask Your Humans: Using Human Instructions to Improve Generalization in Reinforcement Learning.
- Tianyi Zhang, Felix Wu, Arzoo Katiyar, Kilian Q. Weinberger, Yoav Artzi:
Revisiting Few-sample BERT Fine-tuning.
- Tian Li, Ahmad Beirami, Maziar Sanjabi, Virginia Smith:
Tilted Empirical Risk Minimization.



Lin Chen, Sheng Xu:

Deep Neural Tangent Kernel and Laplace Kernel Have the Same RKHS.

Andrea Dittadi, Frederik Träuble, Francesco Locatello, Manuel Wuthrich, Vaibhav Agrawal, Ole Winther, Stefan Bauer, Bernhard Schölkopf:

On the Transfer of Disentangled Representations in Realistic Settings.

David Widmann, Fredrik Lindsten, Dave Zachariah:

Calibration tests beyond classification.

Aditya Krishna Menon, Ankit Singh Rawat, Sanjiv Kumar:

Overparameterisation and worst-case generalisation: friend or foe?

Edgar Schönfeld, Vadim Sushko, Dan Zhang, Juergen Gall, Bernt Schiele , Anna Khoreva:

You Only Need Adversarial Supervision for Semantic Image Synthesis.

Ekin Akyürek, Afra Feyza Akyürek, Jacob Andreas:

Learning to Recombine and Resample Data For Compositional Generalization.

Benjamin David Haeffele, Chong You, René Vidal:

A Critique of Self-Expressive Deep Subspace Clustering.

Yuhuai Wu, Albert Q. Jiang, Jimmy Ba, Roger Baker Grosse:

INT: An Inequality Benchmark for Evaluating Generalization in Theorem Proving.

Jack Prescott, Xiao Zhang, David E. Evans:

Improved Estimation of Concentration Under ℓ_p -Norm Distance Metrics Using Half Spaces.

Sashank J. Reddi, Zachary Charles, Manzil Zaheer, Zachary Garrett, Keith Rush, Jakub Konecný, Sanjiv Kumar, Hugh Brendan McMahan:

Adaptive Federated Optimization.

Haoye Lu, Yongyi Mao, Amiya Nayak:

On the Dynamics of Training Attention Models.

Xiaorui Liu, Yao Li , Rongrong Wang, Jiliang Tang, Ming Yan:

Linear Convergent Decentralized Optimization with Compression.

Emilio Parisotto, Ruslan Salakhutdinov:

Efficient Transformers in Reinforcement Learning using Actor-Learner Distillation.

Dmitry Krotov, John J. Hopfield:

Large Associative Memory Problem in Neurobiology and Machine Learning.

Sanjay Kariyappa, Atul Prakash , Moinuddin K. Qureshi:

Protecting DNNs from Theft using an Ensemble of Diverse Models.

Ziyi Chen, Yi Zhou, Tengyu Xu, Yingbin Liang:

Proximal Gradient Descent-Ascent: Variable Convergence under $K\ddot{L}$ Geometry.

Xinjie Fan, Shujian Zhang, Korawat Tanwisuth, Xiaoning Qian, Mingyuan Zhou :

Contextual Dropout: An Efficient Sample-Dependent Dropout Module.

Sourya Basu, Govardana Sachitanandam Ramachandran, Nitish Shirish Keskar, Lav R. Varshney:

Mirostat: a Neural Text decoding Algorithm that directly controls perplexity.

Rishabh Joshi, Vidhisha Balachandran, Shikhar Vashishth, Alan W. Black, Yulia Tsvetkov:

DialoGraph: Incorporating Interpretable Strategy-Graph Networks into Negotiation Dialogues.

Satya Narayan Shukla, Benjamin M. Marlin:

Multi-Time Attention Networks for Irregularly Sampled Time Series.

Yang Zhao, Jianwen Xie, Ping Li:

Learning Energy-Based Generative Models via Coarse-to-Fine Expanding and Sampling.

Kangle Deng, Aayush Bansal, Deva Ramanan:

Unsupervised Audiovisual Synthesis via Exemplar Autoencoders.

Jeffrey Li, Vaishnavh Nagarajan, Gregory Plumb, Ameet Talwalkar:

A Learning Theoretic Perspective on Local Explainability.

Zhiyuan Fang, Jianfeng Wang, Lijuan Wang, Lei Zhang, Yezhou Yang, Zicheng Liu:

SEED: Self-supervised Distillation For Visual Representation.

Lu Liu, Tianyi Zhou, Guodong Long, Jing Jiang, Xuanyi Dong, Chengqi Zhang:

Isometric Propagation Network for Generalized Zero-shot Learning.

Jindong Gu, Baoyuan Wu, Volker Tresp:

Effective and Efficient Vote Attack on Capsule Networks.

Kaidi Cao, Yining Chen, Junwei Lu, Nikos Aréchiga, Adrien Gaidon, Tengyu Ma:

Heteroskedastic and Imbalanced Deep Learning with Adaptive Regularization.

Alexander Korotin, Lingxiao Li, Justin Solomon, Evgeny Burnaev:

Continuous Wasserstein-2 Barycenter Estimation without Minimax Optimization.

Weitong Zhang, Dongruo Zhou, Lihong Li, Quanquan Gu:

Neural Thompson Sampling.



- Daniel Kunin, Javier Sagastuy-Breña, Surya Ganguli, Daniel L. K. Yamins, Hidenori Tanaka:
Neural Mechanics: Symmetry and Broken Conservation Laws in Deep Learning Dynamics.
- Brian Chmiel, Liad Ben-Uri, Moran Shkolnik, Elad Hoffer, Ron Banner, Daniel Soudry:
Neural gradients are near-lognormal: improved quantized and sparse training.
- Tonghan Wang, Tarun Gupta, Anuj Mahajan, Bei Peng, Shimon Whiteson, Chongjie Zhang:
RODE: Learning Roles to Decompose Multi-Agent Tasks.
- Wonkwang Lee, Whie Jung, Han Zhang, Ting Chen, Jing Yu Koh, Thomas E. Huang, Hyungsuk Yoon, Honglak Lee, Seunghoon Hong:
Revisiting Hierarchical Approach for Persistent Long-Term Video Prediction.
- Sebastian Kaltenbach, Phaedon-Stelios Koutsourelakis:
Physics-aware, probabilistic model order reduction with guaranteed stability.
- Jianhong Wang, Yuan Zhang, Tae-Kyun Kim , Yunjie Gu:
Modelling Hierarchical Structure between Dialogue Policy and Natural Language Generator with Option Framework for Task-oriented Dialogue System.
- Giambattista Parascandolo, Alexander Neitz, Antonio Orvieto, Luigi Gresele, Bernhard Schölkopf:
Learning explanations that are hard to vary.
- Oliver J. Cobb, Christopher G. R. Wallis, Augustine N. Mavor-Parker, Augustin Marignier, Matthew A. Price, Mayeul d'Avezac, Jason D. McEwen:
Efficient Generalized Spherical CNNs.
- Jan Schuchardt, Aleksandar Bojchevski, Johannes Klicpera, Stephan Günnemann:
Collective Robustness Certificates: Exploiting Interdependence in Graph Neural Networks.
- Fabrizio Pittorino, Carlo Lucibello, Christoph Feinauer, Gabriele Perugini, Carlo Baldassi, Elizaveta Demyanenko, Riccardo Zecchina:
Entropic gradient descent algorithms and wide flat minima.
- Haibo Yang, Minghong Fang, Jia Liu:
Achieving Linear Speedup with Partial Worker Participation in Non-IID Federated Learning.
- Phillip Lippe , Efstratios Gavves:
Categorical Normalizing Flows via Continuous Transformations.
- Arash Tavakoli, Mehdi Fatemi, Petar Kormushev:
Learning to Represent Action Values as a Hypergraph on the Action Vertices.
- Mohammad Taha Bahadori, David Heckerman:
Debiasing Concept-based Explanations with Causal Analysis.
- Jorge A. Mendez, Eric Eaton:
Lifelong Learning of Compositional Structures.
- Hyung Won Chung, Thibault Févry, Henry Tsai, Melvin Johnson, Sebastian Ruder:
Rethinking Embedding Coupling in Pre-trained Language Models.
- Songwei Ge, Vedanuj Goswami, Larry Zitnick, Devi Parikh:
Creative Sketch Generation.
- Kaidi Cao, Maria Brbic, Jure Leskovec:
Concept Learners for Few-Shot Learning.
- Kaiyang Zhou, Yongxin Yang, Yu Qiao, Tao Xiang:
Domain Generalization with MixStyle.
- Sachin Mehta, Marjan Ghazvininejad, Srinivasan Iyer, Luke Zettlemoyer, Hannaneh Hajishirzi:
DeLighT: Deep and Light-weight Transformer.
- Zuyue Fu, Zhuoran Yang, Zhaoran Wang:
Single-Timescale Actor-Critic Provably Finds Globally Optimal Policy.
- Danijar Hafner, Timothy P. Lillicrap, Mohammad Norouzi, Jimmy Ba:
Mastering Atari with Discrete World Models.
- Ricky T. Q. Chen, Brandon Amos, Maximilian Nickel:
Learning Neural Event Functions for Ordinary Differential Equations.
- Ali Borji:
Contemplating Real-World Object Classification.
- Ricky T. Q. Chen, Brandon Amos, Maximilian Nickel:
Neural Spatio-Temporal Point Processes.
- Ahmed M. Alaa, Alex James Chan, Mihaela van der Schaar:
Generative Time-series Modeling with Fourier Flows.
- Yihan Wang, Beining Han, Tonghan Wang, Heng Dong, Chongjie Zhang:
DOP: Off-Policy Multi-Agent Decomposed Policy Gradients.
- Elan Rosenfeld, Pradeep Kumar Ravikumar, Andrej Risteski:
The Risks of Invariant Risk Minimization.



- Minjia Zhang, Menghao Li, Chi Wang, Mingqin Li:
DynaTune: Dynamic Tensor Program Optimization in Deep Neural Network Compilation.
- Tianyu Pang, Xiao Yang, Yinpeng Dong, Hang Su, Jun Zhu:
Bag of Tricks for Adversarial Training.
- Hao Cheng, Zhaowei Zhu, Xingyu Li, Yifei Gong, Xing Sun, Yang Liu:
Learning with Instance-Independent Label Noise: A Sample Sieve Approach.
- Xiaoyu Chen, Jiachen Hu, Lihong Li, Liwei Wang:
Efficient Reinforcement Learning in Factored MDPs with Application to Constrained RL.
- Yen-Cheng Liu, Chih-Yao Ma, Zijian He, Chia-Wen Kuo, Kan Chen, Peizhao Zhang, Bichen Wu, Zsolt Kira, Peter Vajda:
Unbiased Teacher for Semi-Supervised Object Detection.
- Yige Li, Xixiang Lyu, Nodens Koren, Lingjuan Lyu, Bo Li, Xingjun Ma:
Neural Attention Distillation: Erasing Backdoor Triggers from Deep Neural Networks.
- Seanie Lee, Dong Bok Lee, Sung Ju Hwang:
Contrastive Learning with Adversarial Perturbations for Conditional Text Generation.
- Jiaheng Wei, Yang Liu:
When Optimizing f-Divergence is Robust with Label Noise.
- Sameera Ramasinghe, Kanchana Nisal Ranasinghe, Salman H. Khan, Nick Barnes, Stephen Gould:
Conditional Generative Modeling via Learning the Latent Space.
- Richard Yuanzhe Pang, He He:
Text Generation by Learning from Demonstrations.
- Haozhi Qi, Xiaolong Wang, Deepak Pathak, Yi Ma, Jitendra Malik:
Learning Long-term Visual Dynamics with Region Proposal Interaction Networks.
- Rishabh Tiwari, Udbhav Bamba, Arnav Chavan, Deepak K. Gupta:
ChipNet: Budget-Aware Pruning with Heaviside Continuous Approximations.
- Mrigank Raman, Aaron Chan, Siddhant Agarwal, Peifeng Wang, Hansen Wang, Sungchul Kim, Ryan A. Rossi, Handong Zhao, Nedim Lipka, Xiang Ren:
Learning to Deceive Knowledge Graph Augmented Models via Targeted Perturbation.
- Manli Zhang, Jianhong Zhang, Zhiwu Lu, Tao Xiang, Mingyu Ding, Songfang Huang:
IEPT: Instance-Level and Episode-Level Pretext Tasks for Few-Shot Learning.
- Wei Tao, Sheng Long, Gaowei Wu, Qing Tao:
The Role of Momentum Parameters in the Optimal Convergence of Adaptive Polyak's Heavy-ball Methods.
- Pierre Stock, Angela Fan, Benjamin Graham, Edouard Grave, Rémi Gribonval, Hervé Jégou, Armand Joulin:
Training with Quantization Noise for Extreme Model Compression.
- Kimon Antonakopoulos, Elena Veronica Belmega, Panayotis Mertikopoulos:
Adaptive Extra-Gradient Methods for Min-Max Optimization and Games.
- Gautier Izacard, Edouard Grave:
Distilling Knowledge from Reader to Retriever for Question Answering.
- Zhenggang Tang, Chao Yu, Boyuan Chen, Huazhe Xu, Xiaolong Wang, Fei Fang, Simon Shaolei Du, Yu Wang, Yi Wu:
Discovering Diverse Multi-Agent Strategic Behavior via Reward Randomization.
- Niels Bruun Ipsen, Pierre-Alexandre Mattei, Jes Frellsen:
not-MIWAE: Deep Generative Modelling with Missing not at Random Data.
- Rianne van den Berg, Alexey A. Gritsenko, Mostafa Dehghani, Casper Kaae Sønderby, Tim Salimans:
IDF++: Analyzing and Improving Integer Discrete Flows for Lossless Compression.
- Alihan Hüyük, Daniel Jarrett, Cem Tekin, Mihaela van der Schaar:
Explaining by Imitating: Understanding Decisions by Interpretable Policy Learning.
- Andres Campero, Roberta Raileanu, Heinrich Küttler, Joshua B. Tenenbaum, Tim Rocktäschel, Edward Grefenstette:
Learning with AMIGo: Adversarially Motivated Intrinsic Goals.
- Rui Wang, Robin Walters, Rose Yu:
Incorporating Symmetry into Deep Dynamics Models for Improved Generalization.
- Christopher A. Choquette-Choo, Natalie Dullerud, Adam Dziedzic, Yunxiang Zhang, Somesh Jha, Nicolas Papernot, Xiao Wang:
CaPC Learning: Confidential and Private Collaborative Learning.
- Bogdan Georgiev, Lukas Franken, Mayukh Mukherjee:
Heating up decision boundaries: isocapacitory saturation, adversarial scenarios and generalization bounds.
- Renjie Liao, Raquel Urtasun, Richard S. Zemel:
A PAC-Bayesian Approach to Generalization Bounds for Graph Neural Networks.
- Daniel Jarrett, Jinsung Yoon, Ioana Bica, Zhaozhi Qian, Ari Ercole, Mihaela van der Schaar:
Clairvoyance: A Pipeline Toolkit for Medical Time Series.
- Yao-Hung Hubert Tsai, Martin Q. Ma, Muqiao Yang, Han Zhao, Louis-Philippe Morency, Ruslan Salakhutdinov:
Self-supervised Representation Learning with Relative Predictive Coding.



- Justin Fu, Sergey Levine:
Offline Model-Based Optimization via Normalized Maximum Likelihood Estimation.
- Alistair Letcher:
On the Impossibility of Global Convergence in Multi-Loss Optimization.
- Sean Fox, Seyedramin Rasoulinezhad, Julian Faraone, David Boland, Philip H. W. Leong:
A Block Minifloat Representation for Training Deep Neural Networks.
- Matthew L. Leavitt, Ari S. Morcos:
Selectivity considered harmful: evaluating the causal impact of class selectivity in DNNs.
- Chao Shang, Jie Chen, Jinbo Bi:
Discrete Graph Structure Learning for Forecasting Multiple Time Series.
- Joshua David Robinson, Ching-Yao Chuang, Suvrit Sra, Stefanie Jegelka:
Contrastive Learning with Hard Negative Samples.
- Simon Carbonnelle, Christophe De Vleeschouwer:
Intraclass clustering: an implicit learning ability that regularizes DNNs.
- Wenbo Gong, Yingzhen Li, José Miguel Hernández-Lobato:
Sliced Kernelized Stein Discrepancy.
- Jiaming Song, Chenlin Meng, Stefano Ermon:
Denoising Diffusion Implicit Models.
- Jesse Zhang, Haonan Yu, Wei Xu:
Hierarchical Reinforcement Learning by Discovering Intrinsic Options.
- Wenhan Xiong, Xiang Lorraine Li, Srini Iyer, Jingfei Du, Patrick S. H. Lewis, William Yang Wang, Yashar Mehdad, Scott Yih, Sebastian Riedel, Douwe Kiela, Barlas Oguz:
Answering Complex Open-Domain Questions with Multi-Hop Dense Retrieval.
- Helong Zhou, Liangchen Song, Jiajie Chen, Ye Zhou, Guoli Wang, Junsong Yuan, Qian Zhang:
Rethinking Soft Labels for Knowledge Distillation: A Bias-Variance Tradeoff Perspective.
- Tianjian Meng, Xiaohan Chen, Yifan Jiang, Zhangyang Wang:
A Design Space Study for LISTA and Beyond.
- Tete Xiao, Xiaolong Wang, Alexei A. Efros, Trevor Darrell:
What Should Not Be Contrastive in Contrastive Learning.
- Thao Nguyen, Maithra Raghu, Simon Kornblith:
Do Wide and Deep Networks Learn the Same Things? Uncovering How Neural Network Representations Vary with Width and Depth.
- Changan Chen, Sagnik Majumder, Ziad Al-Halah, Ruohan Gao, Santhosh Kumar Ramakrishnan, Kristen Grauman:
Learning to Set Waypoints for Audio-Visual Navigation.
- Olga Moskvyak, Frédéric Maire, Feras Dayoub , Mahsa Baktashmotlagh:
Semi-supervised Keypoint Localization.
- Wuyang Chen, Xinyu Gong, Zhangyang Wang:
Neural Architecture Search on ImageNet in Four GPU Hours: A Theoretically Inspired Perspective.
- Benjamin Eysenbach, Shreyas Chaudhari, Swapnil Asawa, Sergey Levine, Ruslan Salakhutdinov:
Off-Dynamics Reinforcement Learning: Training for Transfer with Domain Classifiers.
- Wonyong Jeong, Jaehong Yoon, Eunho Yang, Sung Ju Hwang:
Federated Semi-Supervised Learning with Inter-Client Consistency & Disjoint Learning.
- Zengyi Qin, Kaiqing Zhang, Yuxiao Chen, Jingkai Chen, Chuchu Fan:
Learning Safe Multi-agent Control with Decentralized Neural Barrier Certificates.
- Jingfeng Wu, Difan Zou, Vladimir Braverman, Quanquan Gu:
Direction Matters: On the Implicit Bias of Stochastic Gradient Descent with Moderate Learning Rate.
- Kanika Madan, Nan Rosemary Ke, Anirudh Goyal, Bernhard Schölkopf, Yoshua Bengio:
Fast And Slow Learning Of Recurrent Independent Mechanisms.
- Ziang Yan, Yiwen Guo, Jian Liang, Changshui Zhang:
Policy-Driven Attack: Learning to Query for Hard-label Black-box Adversarial Examples.
- Nikunj Saunshi, Sadhika Malladi, Sanjeev Arora:
A Mathematical Exploration of Why Language Models Help Solve Downstream Tasks.
- Junwen Bai, Weiran Wang, Yingbo Zhou, Caiming Xiong:
Representation Learning for Sequence Data with Deep Autoencoding Predictive Components.
- Chulhee Yun, Shankar Krishnan, Hossein Mobahi:
A unifying view on implicit bias in training linear neural networks.
- Nanxuan Zhao, Zhirong Wu, Rynson W. H. Lau, Stephen Lin:
What Makes Instance Discrimination Good for Transfer Learning?
- Yichen Qian, Zhiyu Tan, Xiuyu Sun, Ming Lin, Dongyang Li, Zhenhong Sun, Hao Li, Rong Jin:
Learning Accurate Entropy Model with Global Reference for Image Compression.



Peidong Liu, Gengwei Zhang, Bochao Wang, Hang Xu, Xiaodan Liang, Yong Jiang, Zhenguo Li:
Loss Function Discovery for Object Detection via Convergence-Simulation Driven Search.

Tao Zhuo, Mohan S. Kankanhalli:
Effective Abstract Reasoning with Dual-Contrast Network.

Da Yu, Huishuai Zhang, Wei Chen, Tie-Yan Liu:
Do not Let Privacy Overbill Utility: Gradient Embedding Perturbation for Private Learning.

David W. Zhang, Gertjan J. Burghouts, Cees G. M. Snoek:
Set Prediction without Imposing Structure as Conditional Density Estimation.

Yaling Tao, Kentaro Takagi, Kouta Nakata:
Clustering-friendly Representation Learning via Instance Discrimination and Feature Decorrelation.

Daniel Zügner, Tobias Kirschstein, Michele Catasta, Jure Leskovec, Stephan Günnemann:
Language-Agnostic Representation Learning of Source Code from Structure and Context.

Jongheon Jeong, Jinwoo Shin:
Training GANs with Stronger Augmentations via Contrastive Discriminator.

Samyadeep Basu, Phillip Pope, Soheil Feizi:
Influence Functions in Deep Learning Are Fragile.

John Zarka, Florentin Guth, Stéphane Mallat:
Separation and Concentration in Deep Networks.

Manoj Kumar, Dirk Weissenborn, Nal Kalchbrenner:
Colorization Transformer.

Michael R. Zhang, Thomas Paine, Ofir Nachum, Cosmin Paduraru, George Tucker, Ziyu Wang, Mohammad Norouzi:
Autoregressive Dynamics Models for Offline Policy Evaluation and Optimization.

Xiaoxiao Li, Meirui Jiang, Xiaofei Zhang, Michael Kamp, Qi Dou:
FedBN: Federated Learning on Non-IID Features via Local Batch Normalization.

Amy Zhang, Shagun Sodhani, Khimya Khetarpal, Joelle Pineau:
Learning Robust State Abstractions for Hidden-Parameter Block MDPs.

Yufan Zhou, Zhenyi Wang, Jiayi Xian, Changyou Chen, Jinhui Xu:
Meta-Learning with Neural Tangent Kernels.

Benjamin Ehret, Christian Henning, Maria R. Cervera, Alexander Meulemans, Johannes von Oswald, Benjamin F. Grewe:
Continual learning in recurrent neural networks.

Grégoire Mialon, Dexiong Chen, Alexandre d'Aspremont, Julien Mairal:
A Trainable Optimal Transport Embedding for Feature Aggregation and its Relationship to Attention.

Ioana Bica, Daniel Jarrett, Alihan Huyuk, Mihaela van der Schaar:
Learning "What-if" Explanations for Sequential Decision-Making.

Adam Foster, Rattana Pukdee, Tom Rainforth:
Improving Transformation Invariance in Contrastive Representation Learning.

Christopher Frye, Damien de Mijolla, Tom Begley, Laurence Cowton, Megan Stanley, Ilya Feige:
Shapley explainability on the data manifold.

Kai Yuanqing Xiao, Logan Engstrom, Andrew Ilyas, Aleksander Madry:
Noise or Signal: The Role of Image Backgrounds in Object Recognition.

Peiye Zhuang, Oluwasanmi Koyejo, Alexander G. Schwing:
Enjoy Your Editing: Controllable GANs for Image Editing via Latent Space Navigation.

Cassidy Laidlaw, Sahil Singla, Soheil Feizi:
Perceptual Adversarial Robustness: Defense Against Unseen Threat Models.

Mohamed S. Abdelfattah, Abhinav Mehrotra, Lukasz Dudziak, Nicholas Donald Lane:
Zero-Cost Proxies for Lightweight NAS.

Michael Kleinman, Alessandro Achille, Daksh Idnani, Jonathan C. Kao:
Usable Information and Evolution of Optimal Representations During Training.

Ben Adlam, Jaehoon Lee, Lechao Xiao, Jeffrey Pennington, Jasper Snoek:
Exploring the Uncertainty Properties of Neural Networks' Implicit Priors in the Infinite-Width Limit.

Cory Stephenson, Suchismita Padhy, Abhinav Ganesh, Yue Hui, Hanlin Tang, SueYeon Chung:
On the geometry of generalization and memorization in deep neural networks.

Alexander Levine, Soheil Feizi:
Deep Partition Aggregation: Provable Defenses against General Poisoning Attacks.

Priya L. Donti, David Rolnick, J. Zico Kolter:
DC3: A learning method for optimization with hard constraints.

Zhiqiang Shen, Zechun Liu, Dejia Xu, Zitian Chen, Kwang-Ting Cheng, Marios Savvides:
Is Label Smoothing Truly Incompatible with Knowledge Distillation: An Empirical Study.

-     Yingwei Li, Qihang Yu, Mingxing Tan, Jieru Mei, Peng Tang, Wei Shen, Alan L. Yuille, Cihang Xie:
Shape-Texture Debiased Neural Network Training.
-     Lucy Chai, Jonas Wulff, Phillip Isola:
Using latent space regression to analyze and leverage compositionality in GANs.
-     Mohak Bhardwaj, Sanjiban Choudhury, Byron Boots:
Blending MPC & Value Function Approximation for Efficient Reinforcement Learning.
-     Karan Goel, Albert Gu, Yixuan Li, Christopher Ré:
Model Patching: Closing the Subgroup Performance Gap with Data Augmentation.
-     Yihao Feng, Ziyang Tang, Na Zhang, Qiang Liu:
Non-asymptotic Confidence Intervals of Off-policy Evaluation: Primal and Dual Bounds.
-     Seyed-Iman Mirzadeh, Mehrdad Farajtabar, Dilan Görür, Razvan Pascanu, Hassan Ghasemzadeh:
Linear Mode Connectivity in Multitask and Continual Learning.
-     Zhenlin Xu, Deyi Liu, Junlin Yang, Colin Raffel, Marc Niethammer:
Robust and Generalizable Visual Representation Learning via Random Convolutions.
-     Pedro Hermosilla, Marco Schäfer , Matej Lang, Gloria Fackelmann, Pere-Pau Vázquez, Barbora Kozlíková, Michael Krone, Tobias Ritschel, Timo Ropinski:
Intrinsic-Extrinsic Convolution and Pooling for Learning on 3D Protein Structures.
-     Atanas Mirchev, Baris Kayalibay, Patrick van der Smagt , Justin Bayer:
Variational State-Space Models for Localisation and Dense 3D Mapping in 6 DoF.
-     Byeongho Heo, Sanghyuk Chun, Seong Joon Oh, Dongyo Han, Sangdoo Yun, Gyawan Kim, Youngjung Uh, Jung-Woo Ha:
AdamP: Slowing Down the Slowdown for Momentum Optimizers on Scale-invariant Weights.
-     Tsung Wei Tsai, Chongxuan Li, Jun Zhu:
MiCE: Mixture of Contrastive Experts for Unsupervised Image Clustering.
-     Deyao Zhu, Mohamed Zahran, Li Erran Li, Mohamed Elhoseiny:
HalentNet: Multimodal Trajectory Forecasting with Hallucinative Intents.
-     Balázs Kégl, Gabriel Hurtado, Albert Thomas:
Model-based micro-data reinforcement learning: what are the crucial model properties and which model to choose?
-     Yuanyuan Yuan, Shuai Wang, Junping Zhang:
Private Image Reconstruction from System Side Channels Using Generative Models.
-     Quang Pham, Chenghao Liu, Doyen Sahoo, Steven C. H. Hoi:
Contextual Transformation Networks for Online Continual Learning.
-     Xin Wang, Jie Ren , Shuyun Lin, Xiangming Zhu, Yisen Wang, Quanshi Zhang:
A Unified Approach to Interpreting and Boosting Adversarial Transferability.
-     Mary Phuong, Christoph H. Lampert:
The inductive bias of ReLU networks on orthogonally separable data.
-     Laurence Aitchison:
A statistical theory of cold posteriors in deep neural networks.
-     Jinhua Zhu, Lijun Wu, Yingce Xia, Shufang Xie, Tao Qin, Wengang Zhou, Houqiang Li, Tie-Yan Liu:
IOT: Instance-wise Layer Reordering for Transformer Structures.
-     Axel Sauer, Andreas Geiger:
Counterfactual Generative Networks.
-     Jonathan Pilault, Amine Elhattami, Christopher J. Pal:
Conditionally Adaptive Multi-Task Learning: Improving Transfer Learning in NLP Using Fewer Parameters & Less Data.
-     Liyang Liu, Yi Li, Zhanghui Kuang, Jing-Hao Xue, Yimin Chen, Wenming Yang, Qingmin Liao, Wayne Zhang:
Towards Impartial Multi-task Learning.
-     James Lucas, Mengye Ren, Irene Raissa Kameni, Toniann Pitassi, Richard S. Zemel:
Theoretical bounds on estimation error for meta-learning.
-     Edoardo Cetin, Oya Çeliktutan:
Domain-Robust Visual Imitation Learning with Mutual Information Constraints.
-     Sana Tonekaboni, Danny Eytan, Anna Goldenberg:
Unsupervised Representation Learning for Time Series with Temporal Neighborhood Coding.
-     Priya L. Donti, Melrose Roderick, Mahyar Fazlyab, J. Zico Kolter:
Enforcing robust control guarantees within neural network policies.
-     Shuang Ma, Zhaoyang Zeng, Daniel McDuff, Yale Song:
Active Contrastive Learning of Audio-Visual Video Representations.
-     Sangho Lee, Youngjae Yu, Gunhee Kim, Thomas M. Breuel, Jan Kautz, Yale Song:
Parameter Efficient Multimodal Transformers for Video Representation Learning.



- Soufiane Hayou, Jean-Francois Ton, Arnaud Doucet, Yee Whye Teh:
Robust Pruning at Initialization.
- Ted Moskovitz, Michael Arbel, Ferenc Huszar, Arthur Gretton:
Efficient Wasserstein Natural Gradients for Reinforcement Learning.
- Boli Chen , Yao Fu, Guangwei Xu, Pengjun Xie, Chuanqi Tan, Moshia Chen, Liping Jing:
Probing BERT in Hyperbolic Spaces.
- Ren Wang, Kaidi Xu, Sijia Liu, Pin-Yu Chen, Tsui-Wei Weng, Chuang Gan, Meng Wang:
On Fast Adversarial Robustness Adaptation in Model-Agnostic Meta-Learning.
- Vinay Venkatesh Ramasesh, Ethan Dyer, Maithra Raghu:
Anatomy of Catastrophic Forgetting: Hidden Representations and Task Semantics.
- Zongbo Han, Changqing Zhang, Huazhu Fu, Joey Tianyi Zhou:
Trusted Multi-View Classification.
- Kibok Lee, Yian Zhu, Kihyuk Sohn, Chun-Liang Li, Jinwoo Shin, Honglak Lee:
i-Mix: A Domain-Agnostic Strategy for Contrastive Representation Learning.
- Mikhail Khodak, Neil A. Tenenholz, Lester Mackey, Nicolò Fusi:
Initialization and Regularization of Factorized Neural Layers.
- Dongsu Zhang, Changwoon Choi, Jeonghwan Kim, Young Min Kim:
Learning to Generate 3D Shapes with Generative Cellular Automata.
- Youngjae Yu, Sangho Lee, Gunhee Kim, Yale Song:
Self-Supervised Learning of Compressed Video Representations.
- Anirudh Goyal, Alex Lamb, Phanideep Gampa, Philippe Beaudoin, Charles Blundell, Sergey Levine, Yoshua Bengio, Michael Curtis Mozer:
Factorizing Declarative and Procedural Knowledge in Structured, Dynamical Environments.
- Sarah M. Hooper, Michael Wornow, Ying Hang Seah, Peter Kellman, Hui Xue, Frederic Sala, Curtis P. Langlotz, Christopher Ré:
Cut out the annotator, keep the cutout: better segmentation with weak supervision.
- Yi Ren, Chenxu Hu, Xu Tan, Tao Qin, Sheng Zhao, Zhou Zhao, Tie-Yan Liu:
FastSpeech 2: Fast and High-Quality End-to-End Text to Speech.
- Xiangpeng Wei, Rongxiang Weng, Yue Hu, Luxi Xing, Heng Yu, Weihua Luo:
On Learning Universal Representations Across Languages.
- Yong Liu, Jiankun Liu, Shuqiang Wang:
Effective Distributed Learning with Random Features: Improved Bounds and Algorithms.
- Kanil Patel, William H. Beluch, Bin Yang, Michael Pfeiffer, Dan Zhang:
Multi-Class Uncertainty Calibration via Mutual Information Maximization-based Binning.
- Alexander Nordcliff, Cristian Bodnar, Ben Day, Jacob Moss, Pietro Liò:
Neural ODE Processes.
- Ziying Li , Shuwen Yang, Guojie Song, Lingsheng Cai:
Conformation-Guided Molecular Representation with Hamiltonian Neural Networks.
- Dihan Zheng, Sia Huat Tan, Xiaowen Zhang, Zuoqiang Shi, Kaisheng Ma , Chenglong Bao:
An Unsupervised Deep Learning Approach for Real-World Image Denoising.
- Andrey Malinin, Liudmila Prokhorenko, Aleksei Ustimenko:
Uncertainty in Gradient Boosting via Ensembles.
- Gustav Sourek, Filip Zelezný, Ondrej Kuzelka:
Lossless Compression of Structured Convolutional Models via Lifting.
- Johannes von Oswald, Seijin Kobayashi, João Sacramento, Alexander Meulemans, Christian Henning, Benjamin F. Grewe:
Neural networks with late-phase weights.
- Dennis Müller, Cezary Kaliszyk :
Disambiguating Symbolic Expressions in Informal Documents.
- George Dasoulas , Johannes F. Lutzeyer, Michalis Vazirgiannis:
Learning Parametrised Graph Shift Operators.
- Adam Fisch, Tal Schuster, Tommi S. Jaakkola, Regina Barzilay:
Efficient Conformal Prediction via Cascaded Inference with Expanded Admission.
- Xuxi Chen, Zhenyu Zhang, Yongduo Sui, Tianlong Chen:
GANs Can Play Lottery Tickets Too.
- Katharina Ott, Prateek Katiyar, Philipp Hennig, Michael Tiemann:
ResNet After All: Neural ODEs and Their Numerical Solution.
- Fredrik Carlsson, Amaru Cuba Gyllensten, Evangelia Gogoulou, Erik Ylipää Hellqvist, Magnus Sahlgren:
Semantic Re-tuning with Contrastive Tension.
- Xiaojie Guo, Yuanqi Du, Liang Zhao:
Property Controllable Variational Autoencoder via Invertible Mutual Dependence.



- Edward De Brouwer, Adam Arany, Jaak Simm, Yves Moreau:
Latent Convergent Cross Mapping.
- Eli Chien , Jianhao Peng, Pan Li, Olgica Milenkovic:
Adaptive Universal Generalized PageRank Graph Neural Network.
- Yatin Nandwani, Deepanshu Jindal, Mausam, Parag Singla:
Neural Learning of One-of-Many Solutions for Combinatorial Problems in Structured Output Spaces.
- Vitaly Kurin, Maximilian Igli, Tim Rocktäschel, Wendelin Boehmer, Shimon Whiteson:
My Body is a Cage: the Role of Morphology in Graph-Based Incompatible Control.
- Hong-You Chen, Wei-Lun Chao:
FedBE: Making Bayesian Model Ensemble Applicable to Federated Learning.
- Juntang Zhuang, Nicha C. Dvornek, Sekhar Tatikonda, James S. Duncan:
MALI: A memory efficient and reverse accurate integrator for Neural ODEs.
- Thomas Bird, Friso H. Kingma, David Barber:
Reducing the Computational Cost of Deep Generative Models with Binary Neural Networks.
- Sang Michael Xie, Ananya Kumar, Robbie Jones, Fereshte Khani, Tengyu Ma, Percy Liang:
In-N-Out: Pre-Training and Self-Training using Auxiliary Information for Out-of-Distribution Robustness.
- Kuilin Chen, Chi-Guhn Lee:
Incremental few-shot learning via vector quantization in deep embedded space.
- Wuyang Chen, Zhiding Yu, Shalini De Mello, Sifei Liu, José M. Álvarez, Zhangyang Wang, Anima Anandkumar:
Contrastive Syn-to-Real Generalization.
- Sayna Ebrahimi, Suzanne Petryk, Akash Gokul, William Gan, Joseph E. Gonzalez, Marcus Rohrbach, Trevor Darrell:
Remembering for the Right Reasons: Explanations Reduce Catastrophic Forgetting.
- Yangchen Pan , Kirby Banman, Martha White:
Fuzzy Tiling Activations: A Simple Approach to Learning Sparse Representations Online.
- Adrian Bulat, Brais Martínez, Georgios Tzimiropoulos:
High-Capacity Expert Binary Networks.
- David Lindner , Rohin Shah, Pieter Abbeel, Anca D. Dragan:
Learning What To Do by Simulating the Past.
- Pau de Jorge, Amartya Sanyal, Harkirat S. Behl, Philip H. S. Torr, Grégory Rogez, Puneet K. Dokania:
Progressive Skeletonization: Trimming more fat from a network at initialization.
- Vin Sachidananda, Ziyi Yang, Chenguang Zhu:
Filtered Inner Product Projection for Crosslingual Embedding Alignment.
- Dmitriy Smirnov, Mikhail Bessmeltsev, Justin Solomon:
Learning Manifold Patch-Based Representations of Man-Made Shapes.
- Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, Jacob Steinhardt:
Aligning AI With Shared Human Values.
- Jason Ramapuram, Yan Wu, Alexandros Kalousis :
Kanerva++: Extending the Kanerva Machine With Differentiable, Locally Block Allocated Latent Memory.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, Jacob Steinhardt:
Measuring Massive Multitask Language Understanding.
- Zhuotong Chen, Qianxiao Li, Zheng Zhang:
Towards Robust Neural Networks via Close-loop Control.
- Subha Maity, Songkai Xue, Mikhail Yurochkin, Yuekai Sun:
Statistical inference for individual fairness.
- Yi Tay, Zhe Zhao, Dara Bahri, Donald Metzler, Da-Cheng Juan:
HyperGrid Transformers: Towards A Single Model for Multiple Tasks.
- Shaocong Ma, Ziyi Chen, Yi Zhou, Shaofeng Zou:
Greedy-GQ with Variance Reduction: Finite-time Analysis and Improved Complexity.
- Sitan Chen, Xiaoxiao Li, Zhao Song, Danyang Zhuo:
On InstaHide, Phase Retrieval, and Sparse Matrix Factorization.
- Bowen Pan, Rameswar Panda, Camilo Luciano Fosco, Chung-Ching Lin, Alex J. Andonian, Yue Meng, Kate Saenko, Aude Oliva, Rogério Feris:
VA-RED2: Video Adaptive Redundancy Reduction.
- Myeongjang Pyeon, Jihwan Moon, Taeyoung Hahn, Gunhee Kim:
SEDONA: Search for Decoupled Neural Networks toward Greedy Block-wise Learning.
- Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, Matthew J. Hausknecht:
ALFWorld: Aligning Text and Embodied Environments for Interactive Learning.
- Yuchen Lu, Yikang Shen, Siyuan Zhou, Aaron C. Courville, Joshua B. Tenenbaum, Chuang Gan:
Learning Task Decomposition with Ordered Memory Policy Network.



Francisco Utrera, Evan Kravitz, N. Benjamin Erichson, Rajiv Khanna, Michael W. Mahoney:
Adversarially-Trained Deep Nets Transfer Better: Illustration on Image Classification.

Jiayi Shen, Haotao Wang, Shupeng Gui, Jianchao Tan, Zhangyang Wang, Ji Liu:
UMEC: Unified model and embedding compression for efficient recommendation systems.

Bingyi Kang, Yu Li, Sa Xie, Zehuan Yuan, Jiashi Feng:
Exploring Balanced Feature Spaces for Representation Learning.

Kartik Gupta, Amir Rahimi, Thalaiyasingam Ajanthan, Thomas Mensink, Cristian Sminchisescu, Richard Hartley:
Calibration of Neural Networks using Splines.

Khai Nguyen, Son Nguyen, Nhat Ho, Tung Pham, Hung Bui:
Improving Relational Regularized Autoencoders with Spherical Sliced Fused Gromov Wasserstein.

Guolin Ke, Di He, Tie-Yan Liu:
Rethinking Positional Encoding in Language Pre-training.

Xuanlin Li, Brandon Trabucco, Dong Huk Park, Michael Luo, Sheng Shen, Trevor Darrell, Yang Gao:
Discovering Non-monotonic Autoregressive Orderings with Variational Inference.

Fabian Otto, Philipp Becker, Ngo Anh Vien, Hanna Carolin Maria Ziesche, Gerhard Neumann:
Differentiable Trust Region Layers for Deep Reinforcement Learning.

A. F. M. Shahab Uddin, Mst. Sirazam Monira, Wheemyung Shin, TaeChoong Chung, Sung-Ho Bae:
SaliencyMix: A Saliency Guided Data Augmentation Strategy for Better Regularization.

Donald Joseph Hejna III, Pieter Abbeel, Lerrel Pinto:
Task-Agnostic Morphology Evolution.

Imanol Schlag, Tsendsuren Munkhdalai, Jürgen Schmidhuber:
Learning Associative Inference Using Fast Weight Memory.

Sergei Ivanov, Liudmila Prokhorenkova:
Boost then Convolve: Gradient Boosting Meets Graph Neural Networks.

Shyam Anil Tailor, Javier Fernández-Marqués, Nicholas Donald Lane:
Degree-Quant: Quantization-Aware Training for Graph Neural Networks.

Duong H. Le, Binh-Son Hua:
Network Pruning That Matters: A Case Study on Retraining Variants.

Hao Li, Chenxin Tao, Xizhou Zhu, Xiaogang Wang, Gao Huang, Jifeng Dai:
Auto Seg-Loss: Searching Metric Surrogates for Semantic Segmentation.

Erik Scharwächter, Jonathan Lennartz, Emmanuel Müller:
Differentiable Segmentation of Sequences.

Zhenfang Chen, Jiayuan Mao, Jiajun Wu, Kwan-Yee Kenneth Wong, Joshua B. Tenenbaum, Chuang Gan:
Grounding Physical Concepts of Objects and Events Through Dynamic Visual Reasoning.

Liyuan Xu, Yutian Chen, Siddarth Srinivasan, Nando de Freitas, Arnaud Doucet, Arthur Gretton:
Learning Deep Features in Instrumental Variable Regression.

Changhao Shi, Chester Holtz, Gal Mishne:
Online Adversarial Purification based on Self-supervised Learning.

Junchi Yu, Tingyang Xu, Yu Rong, Yatao Bian, Junzhou Huang, Ran He:
Graph Information Bottleneck for Subgraph Recognition.

Ishaan Gulrajani, David Lopez-Paz:
In Search of Lost Domain Generalization.

Tianyi Zhou, Shengjie Wang, Jeff A. Bilmes:
Robust Curriculum Learning: from clean label detection to noisy label self-correction.

Jun-Tae Lee, Mihir Jain, Hyoungwoo Park, Sungrack Yun:
Cross-Attentional Audio-Visual Fusion for Weakly-Supervised Action Localization.

Alvin Chan, Yew-Soon Ong, Bill Pung, Aston Zhang, Jie Fu:
CoCon: A Self-Supervised Approach for Controlled Text Generation.

Neel Dey, Antong Chen, Soheil Ghafurian:
Group Equivariant Generative Adversarial Networks.

Eugene Kharitonov, Rahma Chaabouni:
What they do when in doubt: a study of inductive biases in seq2seq learners.

Alexander Neitz, Giambattista Parascandolo, Bernhard Schölkopf:
A teacher-student framework to distill future trajectories.

André Hottung, Bhanu Bhandari, Kevin Tierney:
Learning a Latent Search Space for Routing Problems using Variational Autoencoders.

Paulo Tabuada, Bahman Gharesifard:
Universal approximation power of deep residual neural networks via nonlinear control theory.



- Nadav Dym, Haggai Maron:
On the Universality of Rotation Equivariant Point Cloud Networks.
- Kunchang Li , Xianhang Li, Yali Wang, Jun Wang, Yu Qiao:
CT-Net: Channel Tensorization Network for Video Classification.
- Jonathan Cornford, Damjan Kalajdzievski, Marco Leite, Amélie Lamarquette, Dimitri Michael Kullmann, Blake Aaron Richards:
Learning to live with Dale's principle: ANNs with separate excitatory and inhibitory units.
- Andrey Malinin, Mark J. F. Gales:
Uncertainty Estimation in Autoregressive Structured Prediction.
- Roshan Rao, Joshua Meier, Tom Sercu, Sergey Ovchinnikov, Alexander Rives:
Transformer protein language models are unsupervised structure learners.
- Hengrui Cai, Rui Song, Wenbin Lu:
ANOCE: Analysis of Causal Effects with Multiple Mediators via Constrained Structural Learning.
- Ingmar Schubert, Ozgur S. Oguz, Marc Toussaint:
Plan-Based Relaxed Reward Shaping for Goal-Directed Tasks.
- Xin Ding , Yongwei Wang, Zuheng Xu, William J. Welch, Z. Jane Wang:
CcGAN: Continuous Conditional Generative Adversarial Networks for Image Generation.
- Thomas Fischbacher, Luciano Sbaiz:
Single-Photon Image Classification.
- Sven Gowal, Po-Sen Huang, Aäron van den Oord, Timothy A. Mann, Pushmeet Kohli:
Self-supervised Adversarial Robustness for the Low-label, High-data Regime.
- Guang Zhao, Edward R. Dougherty, Byung-Jun Yoon, Francis J. Alexander, Xiaoning Qian:
Uncertainty-aware Active Learning for Optimal Bayesian Classifier.
- Kevin Xie, Homanga Bharadhwaj, Danijar Hafner, Animesh Garg, Florian Shkurti:
Latent Skill Planning for Exploration and Transfer.
- Valerii Iakovlev, Markus Heinonen, Harri Lähdesmäki:
Learning continuous-time PDEs from sparse data with graph neural networks.
- Andrew Brock, Soham De, Samuel L. Smith:
Characterizing signal propagation to close the performance gap in unnormalized ResNets.
- Tianlong Chen, Zhenyu Zhang, Sijia Liu, Shiyu Chang, Zhangyang Wang:
Robust Overfitting may be mitigated by properly learned smoothening.
- Tianlong Chen, Zhenyu Zhang, Sijia Liu, Shiyu Chang, Zhangyang Wang:
Long Live the Lottery: The Existence of Winning Tickets in Lifelong Learning.
- Gregor N. C. Simm, Robert Pinsler, Gábor Csányi, José Miguel Hernández-Lobato:
Symmetry-Aware Actor-Critic for 3D Molecular Design.
- Yuliang Zou, Zizhao Zhang, Han Zhang, Chun-Liang Li, Xiao Bian, Jia-Bin Huang, Tomas Pfister:
PseudoSeg: Designing Pseudo Labels for Semantic Segmentation.
- Abhinav Mehrotra, Alberto Gil C. P. Ramos, Sourav Bhattacharya, Lukasz Dudziak, Ravichander Vipperla, Thomas Chau, Mohamed S. Abdelfattah, Samin Ishtiaq, Nicholas Donald Lane:
NAS-Bench-ASR: Reproducible Neural Architecture Search for Speech Recognition.
- Alessandro De Palma, Harkirat S. Behl, Rudy Bunel, Philip H. S. Torr, M. Pawan Kumar:
Scaling the Convex Barrier with Active Sets.
- Tanner Fiez, Lillian J. Ratliff:
Local Convergence Analysis of Gradient Descent Ascent with Finite Timescale Separation.
- Pablo Morales-Alvarez, Daniel Hernández-Lobato, Rafael Molina, José Miguel Hernández-Lobato:
Activation-level uncertainty in deep neural networks.
- Tom Veniat, Ludovic Denoyer, Marc'Aurelio Ranzato:
Efficient Continual Learning with Modular Networks and Task-Driven Priors.
- Shyamgopal Karthik, Ameya Prabhu, Puneet K. Dokania, Vineet Gandhi:
No Cost Likelihood Manipulation at Test Time for Making Better Mistakes in Deep Networks.
- Christian H. X. Ali Mehmeti-Göpel, David Hartmann, Michael Wand:
Ringing ReLUs: Harmonic Distortion Analysis of Nonlinear Feedforward Networks.
- Henry Gouk, Timothy M. Hospedales, Massimiliano Pontil:
Distance-Based Regularisation of Deep Networks for Fine-Tuning.
- Enrico Marchesini, Davide Corsi, Alessandro Farinelli:
Genetic Soft Updates for Policy Evolution in Deep Reinforcement Learning.
- Bingchen Liu, Yizhe Zhu, Kunpeng Song, Ahmed Elgammal:
Towards Faster and Stabilized GAN Training for High-fidelity Few-shot Image Synthesis.
- Wenda Li, Lei Yu, Yuhuai Wu, Lawrence C. Paulson:
IsarStep: a Benchmark for High-level Mathematical Reasoning.

-     James Diffenderfer, Bhavya Kailkhura:
Multi-Prize Lottery Ticket Hypothesis: Finding Accurate Binary Neural Networks by Pruning A Randomly Weighted Network.
-     Carles Domingo-Enrich, Fabian Pedregosa, Damien Scieur:
Average-case Acceleration for Bilinear Games and Normal Matrices.
-     Chi Wang, Qingyun Wu, Silu Huang, Amin Saied:
Economic Hyperparameter Optimization with Blended Search Strategy.
-     Huanrui Yang, Lin Duan, Yiran Chen, Hai Li:
BSQ: Exploring Bit-Level Sparsity for Mixed-Precision Neural Network Quantization.
-     Yuchen Jin, Tianyi Zhou, Liangyu Zhao, Yibo Zhu, Chuanxiong Guo, Marco Canini, Arvind Krishnamurthy:
AutoLRS: Automatic Learning-Rate Schedule by Bayesian Optimization on the Fly.
-     Jesse Vig, Ali Madani, Lav R. Varshney, Caiming Xiong, Richard Socher, Nazneen Fatema Rajani:
BERTology Meets Biology: Interpreting Attention in Protein Language Models.
-     Reuben Feinman, Brenden M. Lake:
Learning Task-General Representations with Generative Neuro-Symbolic Modeling.
-     Yunhao Ge, Sami Abu-El-Haija, Gan Xin, Laurent Itti:
Zero-shot Synthesis with Group-Supervised Learning.
-     Erik Jones, Shiori Sagawa, Pang Wei Koh, Ananya Kumar, Percy Liang:
Selective Classification Can Magnify Disparities Across Groups.
-     Armen Aghajanyan, Akshat Shrivastava, Anchit Gupta, Naman Goyal, Luke Zettlemoyer, Sonal Gupta:
Better Fine-Tuning by Reducing Representational Collapse.
-     Marton Havasi, Rodolphe Jenatton, Stanislav Fort, Jeremiah Zhe Liu, Jasper Snoek, Balaji Lakshminarayanan, Andrew Mingbo Dai, Dustin Tran:
Training independent subnetworks for robust prediction.
-     Sreejan Kumar, Ishita Dasgupta, Jonathan D. Cohen, Nathaniel D. Daw, Thomas L. Griffiths:
Meta-Learning of Structured Task Distributions in Humans and Machines.
-     Haotong Qin, Zhongang Cai, Mingyuan Zhang, Yifu Ding, Haiyu Zhao, Shuai Yi, Xianglong Liu, Hao Su:
BiPointNet: Binary Neural Network for Point Clouds.
-     Justin Fu, Mohammad Norouzi, Ofir Nachum, George Tucker, Ziyu Wang, Alexander Novikov, Mengjiao Yang, Michael R. Zhang, Yutian Chen, Aviral Kumar, Cosmin Paduraru, Sergey Levine, Tom Le Paine:
Benchmarks for Deep Off-Policy Evaluation.
-     Keiran Paster, Sheila A. McIlraith, Jimmy Ba:
Planning from Pixels using Inverse Dynamics Models.
-     Namhoon Lee, Thalaiyasingam Ajanthan, Philip H. S. Torr, Martin Jaggi:
Understanding the effects of data parallelism and sparsity on neural network training.
-     Ioannis Exarchos, Marcus Aloysius Pereira, Ziyi Wang, Evangelos A. Theodorou:
NOVAS: Non-convex Optimization via Adaptive Stochastic Search for End-to-end Learning and Control.
-     Duy-Kien Nguyen, Vedanuj Goswami, Xinlei Chen:
MoVie: Revisiting Modulated Convolutions for Visual Counting and Beyond.
-     Angtian Wang, Adam Kortylewski, Alan L. Yuille:
NeMo: Neural Mesh Models of Contrastive Features for Robust 3D Pose Estimation.
-     Lei Chen, Zhengdao Chen, Joan Bruna:
On Graph Neural Networks versus Graph-Augmented MLPs.
-     Jiahui Yu, Wei Han, Anmol Gulati, Chung-Cheng Chiu, Bo Li, Tara N. Sainath, Yonghui Wu, Ruoming Pang:
Dual-mode ASR: Unify and Improve Streaming ASR with Full-context Modeling.
-     Alaa Maalouf, Harry Lang, Daniela Rus, Dan Feldman:
Deep Learning meets Projective Clustering.
-     Yann Bouteiller, Simon Ramstedt, Giovanni Beltrame, Christopher J. Pal, Jonathan Binas:
Reinforcement Learning with Random Delays.
-     Xingyu Cai, Jiaji Huang, Yuchen Bian, Kenneth Church:
Isotropy in the Contextual Embedding Space: Clusters and Manifolds.
-     Chao Pan, Siheng Chen, Antonio Ortega:
Spatio-Temporal Graph Scattering Transform.
-     Tatsuya Matsushima, Hiroki Furuta, Yutaka Matsuo, Ofir Nachum, Shixiang Gu:
Deployment-Efficient Reinforcement Learning via Model-Based Offline Optimization.
-     J. Krishna Murthy, Miles Macklin, Florian Golemo, Vikram Voleti, Linda Petrini, Martin Weiss, Breandan Considine, Jérôme Parent-Lévesque, Kevin Xie, Kenny Erleben, Liam Paull, Florian Shkurti, Derek Nowrouzezahrai, Sanja Fidler:
gradSim: Differentiable simulation for system identification and visuomotor control.
-     Cheng-Yu Hsieh, Chih-Kuan Yeh, Xuanqing Liu, Pradeep Kumar Ravikumar, Seungyeon Kim, Sanjiv Kumar, Cho-Jui Hsieh:
Evaluations and Methods for Explanation through Robustness Analysis.



Meng Qu, Junkun Chen, Louis-Pascal A. C. Xhonneux, Yoshua Bengio, Jian Tang:
RNNLogic: Learning Logic Rules for Reasoning on Knowledge Graphs.



Yuchen Liang, Chaitanya K. Ryali, Benjamin Hoover , Leopold Grinberg, Saket Navlakha, Mohammed J. Zaki, Dmitry Krotov:

Can a Fruit Fly Learn Word Embeddings?



Zichao Yan, William L. Hamilton, Mathieu Blanchette:
Neural representation and generation for RNA secondary structures.



Tuan Anh Nguyen, Anh Tuan Tran:
WaNet - Imperceptible Warping-based Backdoor Attack.



Valeriia Cherepanova, Micah Goldblum, Harrison Foley, Shiyuan Duan, John P. Dickerson, Gavin Taylor, Tom Goldstein:
LowKey: Leveraging Adversarial Attacks to Protect Social Media Users from Facial Recognition.



Victor Sanh, Thomas Wolf, Yonatan Belinkov, Alexander M. Rush :
Learning from others' mistakes: Avoiding dataset biases without modeling them.



Junnan Li, Pan Zhou, Caiming Xiong, Steven C. H. Hoi:
Prototypical Contrastive Learning of Unsupervised Representations.



Harsh Mehta, Ashok Cutkosky , Behnam Neyshabur:
Extreme Memorization via Scale of Initialization.



Benedikt Boecking, Willie Neiswanger, Eric P. Xing, Artur Dubrawski:
Interactive Weak Supervision: Learning Useful Heuristics for Data Labeling.



Kuan Fang, Yuke Zhu, Silvio Savarese, Li Fei-Fei:
Adaptive Procedural Task Generation for Hard-Exploration Problems.



Shivangi Mahto, Vy Ai Vo, Javier S. Turek, Alexander Huth:
Multi-timescale Representation Learning in LSTM Language Models.



Jungo Kasai, Nikolaos Pappas, Hao Peng, James Cross, Noah A. Smith:
Deep Encoder, Shallow Decoder: Reevaluating Non-autoregressive Machine Translation.



Asish Ghoshal, Xilun Chen, Sonal Gupta, Luke Zettlemoyer, Yashar Mehdad:
Learning Better Structured Representations Using Low-rank Adaptive Label Smoothing.



Charles Lovering, Rohan Jha, Tal Linzen, Ellie Pavlick:
Predicting Inductive Biases of Pre-Trained Models.



Binghong Chen, Tianzhe Wang, Chengtao Li, Hanjun Dai, Le Song:
Molecule Optimization by Explainable Evolution.



Paul Pu Liang, Manzil Zaheer, Yuan Wang, Amr Ahmed:
Anchor & Transform: Learning Sparse Embeddings for Large Vocabularies.



Hehe Fan, Xin Yu, Yuhang Ding, Yi Yang, Mohan S. Kankanhalli:
PSTNet: Point Spatio-Temporal Convolution on Point Cloud Sequences.



Makoto Kawano, Wataru Kumagai, Akiyoshi Sannai, Yusuke Iwasawa, Yutaka Matsuo:
Group Equivariant Conditional Neural Processes.



Shun-ichi Amari, Jimmy Ba, Roger Baker Grosse, Xuechen Li, Atsushi Nitanda, Taiji Suzuki, Denny Wu, Ji Xu:
When does preconditioning help or hurt generalization?



Hung Le, Nancy F. Chen, Steven C. H. Hoi:
Learning Reasoning Paths over Semantic Graphs for Video-grounded Dialogues.



Ning Ding, Xiaobin Wang, Yao Fu, Guangwei Xu, Rui Wang, Pengjun Xie, Ying Shen, Fei Huang, Haitao Zheng, Rui Zhang:
Prototypical Representation Learning for Relation Extraction.



Jaeho Lee, Sejun Park, Sangwoo Mo, Sungsoo Ahn, Jinwoo Shin:
Layer-adaptive Sparsity for the Magnitude-based Pruning.



Abdul Fatir Ansari, Ming Liang Ang, Harold Soh:
Refining Deep Generative Models via Discriminator Gradient Flow.



Divyansh Kaushik, Amrit Sethuraman, Eduard H. Hovy, Zachary Chase Lipton:
Explaining the Efficacy of Counterfactually Augmented Data.



N. Benjamin Erichson, Omri Azencot, Alejandro F. Queiruga, Liam Hodgkinson, Michael W. Mahoney:
Lipschitz Recurrent Neural Networks.



Panagiotis Kyriakis, Iordanis Fostiroopoulos, Paul Bogdan:
Learning Hyperbolic Representations of Topological Features.



Yoonhyung Lee, Joongbo Shin, Kyomin Jung:
Bidirectional Variational Inference for Non-Autoregressive Text-to-Speech.



Núria Armengol Urpí, Sebastian Curi, Andreas Krause :
Risk-Averse Offline Reinforcement Learning.



David W. Romero, Jean-Baptiste Cordonnier:
Group Equivariant Stand-Alone Self-Attention For Vision.

-     Samuel Horváth, Peter Richtárik:
A Better Alternative to Error Feedback for Communication-Efficient Distributed Learning.
-     Qunxi Zhu, Yao Guo, Wei Lin:
Neural Delay Differential Equations.
-     Tom Joy, Sebastian M. Schmon, Philip H. S. Torr, Siddharth Narayanaswamy, Tom Rainforth:
Capturing Label Characteristics in VAEs.
-     Benjamin Paassen , Daniele Grattarola, Daniele Zambon, Cesare Alippi, Barbara Hammer:
Graph Edit Networks.
-     Boxin Wang, Shuohang Wang, Yu Cheng, Zhe Gan, Ruoxi Jia, Bo Li, Jingjing Liu:
InfoBERT: Improving Robustness of Language Models from An Information Theoretic Perspective.
-     Xiangning Chen, Ruochen Wang, Minhao Cheng , Xiaocheng Tang, Cho-Jui Hsieh:
DrNAS: Dirichlet Neural Architecture Search.
-     Jaekyeom Kim , Minjung Kim, Dongyeon Woo, Gunhee Kim:
Drop-Bottleneck: Learning Discrete Compressed Representation for Noise-Robust Exploration.
-     Youngsoo Jang, Seokin Seo, Jongmin Lee, Kee-Eung Kim:
Monte-Carlo Planning and Learning with Language Action Value Estimates.
-     Xiaobo Xia, Tongliang Liu, Bo Han, Chen Gong, Nannan Wang, Zongyuan Ge, Yi Chang:
Robust early-learning: Hindering the memorization of noisy labels.
-     Seungjun Lee, Haesang Yang, Woojae Seong:
Identifying Physical Law of Hamiltonian Systems via Meta-Learning.
-     Mingyang Yi, Lu Hou, Lifeng Shang, Xin Jiang, Qun Liu, Zhi-Ming Ma:
Reweighting Augmented Samples by Minimizing the Maximal Expected Loss.
-     Xuebo Liu, Longyue Wang, Derek F. Wong, Liang Ding, Lidia S. Chao, Zhaopeng Tu:
Understanding and Improving Encoder Layer Fusion in Sequence-to-Sequence Learning.
-     Dorđe Miladinovic, Aleksandar Stanic, Stefan Bauer, Jürgen Schmidhuber, Joachim M. Buhmann:
Spatial Dependency Networks: Neural Layers for Improved Generative Image Modeling.
-     Seon-Ho Lee, Chang-Su Kim:
Deep Repulsive Clustering of Ordered Data Based on Order-Identity Decomposition.
-     Yulin Wang , Zanlin Ni, Shiji Song, Le Yang, Gao Huang:
Revisiting Locally Supervised Learning: an Alternative to End-to-end Training.
-     Shiying Xiong, Yunjin Tong, Xingzhe He, Shuqi Yang, Cheng Yang, Bo Zhu:
Nonseparable Symplectic Neural Networks.
-     Sam Bond-Taylor, Chris G. Willcocks:
Gradient Origin Networks.
-     Youngmin Oh, Kimin Lee, Jinwoo Shin, Eunho Yang, Sung Ju Hwang:
Learning to Sample with Local and Global Contexts in Experience Replay Buffer.
-     Dipendra Misra, Qinghua Liu, Chi Jin, John Langford:
Provable Rich Observation Reinforcement Learning with Combinatorial Latent States.
-     Yunwen Lei, Yiming Ying:
Sharper Generalization Bounds for Learning with Gradient-dominated Objective Functions.
-     Hayeon Lee, Eunyoung Hyung, Sung Ju Hwang:
Rapid Neural Architecture Search by Learning to Generate Graphs from Datasets.
-     Yuge Shi, Brooks Paige, Philip H. S. Torr, N. Siddharth:
Relating by Contrasting: A Data-efficient Framework for Multimodal Generative Models.
-     Tehrim Yoon, Sumin Shin, Sung Ju Hwang, Eunho Yang:
FedMix: Approximation of Mixup under Mean Augmented Federated Learning.
-     Noel Loo, Siddharth Swaroop, Richard E. Turner:
Generalized Variational Continual Learning.
-     Liang Ding, Longyue Wang, Xuebo Liu, Derek F. Wong, Dacheng Tao, Zhaopeng Tu:
Understanding and Improving Lexical Choice in Non-Autoregressive Translation.
-     Michael Volpp, Fabian Flürenbrock, Lukas Großberger , Christian Daniel, Gerhard Neumann:
Bayesian Context Aggregation for Neural Processes.
-     Taehwan Kwon:
Variational Intrinsic Control Revisited.
-     David G. T. Barrett, Benoit Dherin:
Implicit Gradient Regularization.
-     Guoqing Liu, Chuheng Zhang, Li Zhao, Tao Qin, Jinhua Zhu, Jian Li, Nenghai Yu, Tie-Yan Liu:
Return-Based Contrastive Representation Learning for Reinforcement Learning.

-     Alex James Chan, Mihaela van der Schaar:
Scalable Bayesian Inverse Reinforcement Learning.
-     Judy Borowski, Roland Simon Zimmermann, Judith Schepers, Robert Geirhos, Thomas S. A. Wallis, Matthias Bethge, Wieland Brendel:
Exemplary Natural Images Explain CNN Activations Better than State-of-the-Art Feature Visualization.
-     Jiaojiao Zhao, Cees G. M. Snoek:
LiftPool: Bidirectional ConvNet Pooling.
-     Alexia Jolicoeur-Martineau, Rémi Piché-Taillefer, Ioannis Mitliagkas, Remi Tachet des Combes:
Adversarial score matching and improved sampling for image generation.
-     Maximilian Igl, Gregory Farquhar, Jelena Luketina, Wendelin Boehmer, Shimon Whiteson:
Transient Non-stationarity and Generalisation in Deep Reinforcement Learning.
-     Samuel L. Smith, Benoit Dherin, David G. T. Barrett, Soham De:
On the Origin of Implicit Regularization in Stochastic Gradient Descent.
-     Muhammet Balciłar, Guillaume Renton, Pierre Héroux, Benoit Gaüzère, Sébastien Adam, Paul Honeine:
Analyzing the Expressive Power of Graph Neural Networks in a Spectral Perspective.
-     Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, Michael Carbin:
Pruning Neural Networks at Initialization: Why Are We Missing the Mark?
-     Tsiry Mayet, Anne Lambert, Pascal Leguyadec, Françoise Le Bolzer, François Schnitzler:
SkipW: Resource Adaptable RNN with Strict Upper Computational Limit.
-     Efthymios Tzinis, Scott Wisdom, Aren Jansen, Shawn Hershey, Tal Remez, Dan Ellis, John R. Hershey:
Into the Wild with AudioScope: Unsupervised Audio-Visual Separation of On-Screen Sounds.
-     Lin Ning, Guoyang Chen, Weifeng Zhang, Xipeng Shen:
Simple Augmentation Goes a Long Way: ADRL for DNN Quantization.
-     Martin Wistuba, Josif Grabocka:
Few-Shot Bayesian Optimization with Deep Kernel Surrogates.
-     Mingjian Chen, Xu Tan, Bohan Li, Yanqing Liu, Tao Qin, Sheng Zhao, Tie-Yan Liu:
AdaSpeech: Adaptive Text to Speech for Custom Voice.
-     Enmao Diao, Jie Ding, Vahid Tarokh:
HeteroFL: Computation and Communication Efficient Federated Learning for Heterogeneous Clients.
-     Konstantinos Vougioukas, Stavros Petridis, Maja Pantic:
DINO: A Conditional Energy-Based GAN for Domain Translation.
-     Tsung-Wei Ke, Jyh-Jing Hwang, Stella X. Yu:
Universal Weakly Supervised Segmentation by Pixel-to-Segment Contrastive Learning.
-     Yujia Liu, Stefano D'Aronco, Konrad Schindler, Jan Dirk Wegner:
PC2WF: 3D Wireframe Reconstruction from Raw Point Clouds.
-     Adam Uri Yaari, Maxwell Sherman, Oliver Clarke Priebe, Po-Ru Loh, Boris Katz, Andrei Barbu, Bonnie Berger:
Multi-resolution modeling of a discrete stochastic process identifies causes of cancer.
-     Panteha Naderian, Gabriel Loaiza-Ganem, Harry J. Braviner, Anthony L. Caterini, Jesse C. Cresswell, Tong Li, Animesh Garg:
C-Learning: Horizon-Aware Cumulative Accessibility Estimation.
-     Rui Wang, Xiaoqian Wang, David I. Inouye :
Shapley Explanation Networks.
-     Milton Llera Montero, Casimir J. H. Ludwig, Rui Ponte Costa, Gaurav Malhotra, Jeffrey S. Bowers:
The role of Disentanglement in Generalisation.
-     Aojun Zhou, Yukun Ma, Junnan Zhu, Jianbo Liu, Zhijie Zhang, Kun Yuan, Wenxiu Sun, Hongsheng Li:
Learning N: M Fine-grained Structured Sparse Neural Networks From Scratch.
-     Siavash Khodadadeh, Sharare Zehtabian, Saeed Vahidian, Weijia Wang, Bill Lin, Ladislau Bölöni:
Unsupervised Meta-Learning through Latent-Space Interpolation in Generative Models.
-     Atin Ghosh, Alexandre H. Thiéry:
On Data-Augmentation and Consistency-Based Semi-Supervised Learning.
-     Yordan Hristov, Subramanian Ramamoorthy:
Learning from Demonstration with Weakly Supervised Disentanglement.
-     Freya Behrens, Jonathan Sauder, Peter Jung:
Neurally Augmented ALISTA.
-     Md. Amirul Islam, Matthew Kowal, Patrick Esser, Sen Jia, Björn Ommer, Konstantinos G. Derpanis, Neil D. B. Bruce:
Shape or Texture: Understanding Discriminative Features in CNNs.
-     Chin-Wei Huang, Ricky T. Q. Chen, Christos Tsirigotis, Aaron C. Courville:
Convex Potential Flows: Universal Probability Distributions with Optimal Transport and Convex Optimization.
-     Soheil Kolouri, Navid Naderializadeh, Gustavo K. Rohde, Heiko Hoffmann :
Wasserstein Embedding for Graph Learning.



Alberto Bernacchia:

Meta-learning with negative learning rates.Maxwell I. Nye, Yewen Pu, Matthew Bowers, Jacob Andreas, Joshua B. Tenenbaum, Armando Solar-Lezama :
Representing Partial Programs with Blended Abstract Semantics.

Huang Fang, Zhenan Fan, Michael P. Friedlander:

Fast convergence of stochastic subgradient method under interpolation.Yujia Xie, Yixiu Mao, Simiao Zuo, Hongteng Xu, Xiaojing Ye, Tuo Zhao, Hongyuan Zha:
A Hypergradient Approach to Robust Regression without Correspondence.Jessica B. Hamrick, Abram L. Friesen, Feryal M. P. Behbahani, Arthur Guez, Fabio Viola, Sims Witherspoon, Thomas Anthony, Lars Holger Buesing, Petar Velickovic, Theophane Weber:
On the role of planning in model-based deep reinforcement learning.

Robin Walters, Jinxi Li, Rose Yu:

Trajectory Prediction using Equivariant Continuous Convolution.Ahmed Akakzia, Cédric Colas, Pierre-Yves Oudeyer, Mohamed Chetouani, Olivier Sigaud:
Grounding Language to Autonomously-Acquired Skills via Goal Generation.

Yun Kuen Cheung, Yixin Tao:

Chaos of Learning Beyond Zero-sum and Coordination via Game Decompositions.Masanobu Horie, Naoki Morita, Toshiaki Hishinuma, Yu Ihara, Naoto Mitsume:
Isometric Transformation Invariant and Equivariant Graph Convolutional Networks.

Junyi Zhu , Matthew B. Blaschko:

R-GAP: Recursive Gradient Attack on Privacy.

Timothy Castiglia, Anirban Das , Stacy Patterson:

Multi-Level Local SGD: Distributed SGD for Heterogeneous Hierarchical Networks.

Dmitry Lepikhin, Hyoukjoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, Zhifeng Chen:

GShard: Scaling Giant Models with Conditional Computation and Automatic Sharding.

Garrett Honke, Irina Higgins, Nina Thigpen, Vladimir Miskovic, Katie Link, Sunny Duan, Pramod Gupta, Julia Klawohn, Greg Hajcak:

Representation learning for improved interpretability and classification accuracy of clinical factors from EEG.

Rizal Fathony, Anit Kumar Sahu, Devin Willmott, J. Zico Kolter:

Multiplicative Filter Networks.

Róbert Csordás, Sjoerd van Steenkiste, Jürgen Schmidhuber:

Are Neural Nets Modular? Inspecting Functional Modularity Through Differentiable Weight Masks.

Paul Michel, Tatsunori Hashimoto, Graham Neubig:

Modeling the Second Player in Distributionally Robust Optimization.

Marcel Neunhoeffer, Steven Wu , Cynthia Dwork:

Private Post-GAN Boosting.

Rafael Valle, Kevin J. Shih, Ryan Prenger, Bryan Catanzaro:

Flowtron: an Autoregressive Flow-based Generative Network for Text-to-Speech Synthesis.

Ziyu Yao , Frank F. Xu, Pengcheng Yin, Huan Sun, Graham Neubig:

Learning Structural Edits via Incremental Tree Transformations.

Jörg K. H. Franke, Gregor Köhler, André Biedenkapp, Frank Hutter:

Sample-Efficient Automated Deep Reinforcement Learning.

Yilun Du, Kevin A. Smith, Tomer D. Ullman, Joshua B. Tenenbaum, Jiajun Wu:

Unsupervised Discovery of 3D Physical Objects from Video.

Andrea Agazzi, Jianfeng Lu:

Global optimality of softmax policy gradient with single hidden layer neural networks in the mean-field regime.

Cristina Pinneri, Shambhuraj Sawant, Sebastian Blaes, Georg Martius:

Extracting Strong Policies for Robotics Tasks from Zero-Order Trajectory Optimizers.

Will Dabney, Georg Ostrovski, André Barreto:

Temporally-Extended ϵ -Greedy Exploration.

Samuel Ritter, Ryan Faulkner, Laurent Sartran, Adam Santoro, Matthew M. Botvinick, David Raposo:

Rapid Task-Solving in Novel Environments.

Raphael Gontijo Lopes, Sylvia J. Smullin, Ekin Dogus Cubuk, Ethan Dyer:

Tradeoffs in Data Augmentation: An Empirical Study.

Ahsan Mahmood, Junier Oliva, Martin Andreas Styner:

Multiscale Score Matching for Out-of-Distribution Detection.

Yogesh Balaji, Mohammadmahdi Sajedi, Neha Mukund Kalibhat, Mucong Ding, Dominik Stöger, Mahdi Soltanolkotabi , Soheil Feizi:

Understanding Over-parameterization in Generative Adversarial Networks.



- Mathieu Chalvidal, Matthew Ricci, Rufin VanRullen, Thomas Serre:
Go with the flow: Adaptive control for Neural ODEs.
- Chen-Yu Wei, Chung-Wei Lee, Mengxiao Zhang, Haipeng Luo:
Linear Last-iterate Convergence in Constrained Saddle-point Optimization.
- François Charton, Amaury Hayat, Guillaume Lample:
Learning advanced mathematical computations from examples.
- Nanxin Chen, Yu Zhang, Heiga Zen, Ron J. Weiss, Mohammad Norouzi, William Chan:
WaveGrad: Estimating Gradients for Waveform Generation.
- Matan Atzmon, Yaron Lipman:
 SALD: Sign Agnostic Learning with Derivatives.
- Michael Arbel, Liang Zhou, Arthur Gretton:
Generalized Energy Based Models.
- Yi Tay, Mostafa Dehghani, Samira Abnar, Yikang Shen, Dara Bahri, Philip Pham, Jinfeng Rao, Liu Yang, Sebastian Ruder, Donald Metzler:
Long Range Arena : A Benchmark for Efficient Transformers.
- Boyuan Chen, Yu Li, Sunand Raghupathi, Hod Lipson:
Beyond Categorical Label Representations for Image Classification.
- Shiyang Li, Semih Yavuz, Kazuma Hashimoto, Jia Li, Tong Niu, Nazneen Fatema Rajani, Xifeng Yan, Yingbo Zhou, Caiming Xiong:
CoCo: Controllable Counterfactuals for Evaluating Dialogue State Trackers.
- Mitch Hill, Jonathan Craig Mitchell, Song-Chun Zhu:
Stochastic Security: Adversarial Defense Using Long-Run Dynamics of Energy-Based Models.
- Jensen Gao, Siddharth Reddy, Glen Berseth, Nicholas Hardy, Nikhilesh Natraj, Karunesh Ganguly, Anca D. Dragan, Sergey Levine:
X2T: Training an X-to-Text Typing Interface with Online Learning from User Feedback.
- Hsiang-Yun Sherry Chien, Jinhan Zhang, Christopher J. Honey:
Mapping the Timescale Organization of Neural Language Models.
- Jérémie Donà, Jean-Yves Franceschi, Sylvain Lamprier, Patrick Gallinari:
PDE-Driven Spatiotemporal Disentanglement.
- Anurag Ajay, Aviral Kumar, Pulkit Agrawal, Sergey Levine, Ofir Nachum:
OPAL: Offline Primitive Discovery for Accelerating Offline Reinforcement Learning.
- Chaithanya Kumar Mummadgi, Ranjitha Subramaniam, Robin Hutmacher, Julien Vitay, Volker Fischer, Jan Hendrik Metzen:
Does enhanced shape bias improve neural network robustness to common corruptions?
- Veronika Thost, Jie Chen:
Directed Acyclic Graph Neural Networks.
- Jianhao Wang, Zhizhou Ren, Terry Liu, Yang Yu, Chongjie Zhang:
QPLEX: Duplex Dueling Multi-Agent Q-Learning.
- Ruiqi Gao, Yang Song, Ben Poole, Ying Nian Wu, Diederik P. Kingma:
Learning Energy-Based Models by Diffusion Recovery Likelihood.
- S. Chandra Mouli, Bruno Ribeiro:
Neural Networks for Learning Counterfactual G-Invariances from Single Environments.
- Arthur Argenson, Gabriel Dulac-Arnold:
Model-Based Offline Planning.
- Peizhao Li, Yifei Wang, Han Zhao, Pengyu Hong, Hongfu Liu:
On Dyadic Fairness: Exploring and Mitigating Bias in Graph Connections.
- Jingzhao Zhang, Aditya Krishna Menon, Andreas Veit, Srinadh Bhojanapalli, Sanjiv Kumar, Suvrit Sra:
Coping with Label Shift via Distributionally Robust Optimisation.
- Jinjie Zhang, Rayan Saab:
Faster Binary Embeddings for Preserving Euclidean Distances.
- Kihyuk Sohn, Chun-Liang Li, Jinsung Yoon, Minho Jin, Tomas Pfister:
Learning and Evaluating Representations for Deep One-Class Classification.
- Mike Wu, Milan Mosse, Chengxu Zhuang, Daniel Yamins, Noah D. Goodman:
Conditional Negative Sampling for Contrastive Learning of Visual Representations.
- Benyou Wang, Lifeng Shang, Christina Lioma, Xin Jiang, Hao Yang, Qun Liu, Jakob Grue Simonsen:
On Position Embeddings in BERT.
- Namyeong Kwon, Hwidong Na, Gabriel Huang, Simon Lacoste-Julien:
Repurposing Pretrained Models for Robust Out-of-domain Few-Shot Learning.
- Timothy Nguyen, Zhourong Chen, Jaehoon Lee:
Dataset Meta-Learning from Kernel Ridge-Regression.



- Yue Meng, Rameswar Panda, Chung-Ching Lin, Prasanna Sattigeri, Leonid Karlinsky, Kate Saenko, Aude Oliva, Rogério Feris:
AdaFuse: Adaptive Temporal Fusion Network for Efficient Action Recognition.
- Atish Agarwala, Abhimanyu Das, Brendan Juba, Rina Panigrahy, Vatsal Sharan, Xin Wang, Qiuyi Zhang:
One Network Fits All? Modular versus Monolithic Task Formulations in Neural Networks.
- Arda Sahiner, Tolga Ergen, John M. Pauly, Mert Pilancı:
Vector-output ReLU Neural Network Problems are Copositive Programs: Convex Analysis of Two Layer Networks and Polynomial-time Algorithms.
- Mamshad Nayeem Rizve, Kevin Duarte, Yogesh S. Rawat, Mubarak Shah:
In Defense of Pseudo-Labeling: An Uncertainty-Aware Pseudo-label Selection Framework for Semi-Supervised Learning.
- Nanyi Fei, Zhiwu Lu, Tao Xiang, Songfang Huang:
MELR: Meta-Learning via Modeling Episode-Level Relationships for Few-Shot Learning.
- Jing An, Lexing Ying, Yuhua Zhu:
Why resampling outperforms reweighting for correcting sampling bias with stochastic gradients.
- Changmin Yu, Timothy Behrens, Neil Burgess:
Prediction and generalisation over directed actions by grid cells.
- Honglu Zhou, Asim Kadav, Farley Lai, Alexandru Niculescu-Mizil, Martin Renqiang Min, Mubbashir Kapadia, Hans Peter Graf:
Hopper: Multi-hop Transformer for Spatiotemporal Reasoning.
- Joshua C. Chang, Patrick Fletcher, Jungmin Han, Ted L. Chang, Shashaank Vattikuti, Bart Desmet, Ayah Zirkly, Carson C. Chow:
Sparse encoding for more-interpretable feature-selecting representations in probabilistic matrix factorization.
- Ruozi Huang, Huang Hu, Wei Wu, Kei Sawada, Mi Zhang, Dixin Jiang:
Dance Revolution: Long-Term Dance Generation with Music via Curriculum Learning.
- Sangdon Park, Shuo Li, Insup Lee, Osbert Bastani:
PAC Confidence Predictions for Deep Neural Network Classifiers.
- Shibani Santurkar, Dimitris Tsipras, Aleksander Madry:
BREEDS: Benchmarks for Subpopulation Shift.
- Yingxue Zhou, Steven Wu , Arindam Banerjee:
Bypassing the Ambient Dimension: Private SGD with Gradient Subspace Identification.
- Daniel James Lenton, Stephen James, Ronald Clark, Andrew J. Davison:
End-to-End Egospheric Spatial Memory.
- Sharon Zhou, Eric Zelikman, Fred Lu, Andrew Y. Ng, Gunnar E. Carlsson, Stefano Ermon:
Evaluating the Disentanglement of Deep Generative Models through Manifold Topology.
- Tao Yu, Rui Zhang, Alex Polozov, Christopher Meek, Ahmed Hassan Awadallah:
SCoRe: Pre-Training for Context Representation in Conversational Semantic Parsing.
- Xuezhe Ma, Xiang Kong, Shanghang Zhang, Eduard H. Hovy:
Decoupling Global and Local Representations via Invertible Generative Flows.
- Wangchunshu Zhou, Dong-Ho Lee, Ravi Kiran Selvam, Seyeon Lee, Xiang Ren:
Pre-training Text-to-Text Transformers for Concept-centric Common Sense.
- Kyriakos Axiotis, Maxim Sviridenko:
Local Search Algorithms for Rank-Constrained Convex Optimization.
- Qian Huang, Horace He, Abhay Singh, Ser-Nam Lim, Austin R. Benson:
Combining Label Propagation and Simple Models out-performs Graph Neural Networks.
- Beliz Gunel, Jingfei Du, Alexis Conneau, Veselin Stoyanov:
Supervised Contrastive Learning for Pre-trained Language Model Fine-tuning.
- Qian Lou , Yilin Shen, Hongxia Jin, Lei Jiang:
SAFENet: A Secure, Accurate and Fast Neural Network Inference.
- Fatemeh Sheikholeslami, Ali Lotfi, J. Zico Kolter:
Provably robust classification of adversarial examples with detection.
- Joseph D. Viviano, Becks Simpson, Francis Dutil, Yoshua Bengio, Joseph Paul Cohen:
Saliency is a Possible Red Herring When Diagnosing Poor Generalization.
- Zongyi Li, Nikola Borislavov Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew M. Stuart, Anima Anandkumar:
Fourier Neural Operator for Parametric Partial Differential Equations.
- Yeming Wen, Ghassen Jerfel, Rafael Muller, Michael W. Dusenberry, Jasper Snoek, Balaji Lakshminarayanan, Dustin Tran:
Combining Ensembles and Data Augmentation Can Harm Your Calibration.
- Tharun Medini, Beidi Chen, Anshumali Shrivastava:
SOLAR: Sparse Orthogonal Learned and Random Embeddings.



- Ruihan Zhao, Kevin Lu, Pieter Abbeel, Stas Tiomkin:
Efficient Empowerment Estimation for Unsupervised Stabilization.
- Abdul Wasay, Stratos Idreos:
More or Less: When and How to Build Convolutional Neural Network Ensembles.
- Heming Du, Xin Yu, Liang Zheng :
VTNet: Visual Transformer Network for Object Goal Navigation.
- Ivan Skorokhodov, Mohamed Elhoseiny:
Class Normalization for (Continual)? Generalized Zero-Shot Learning.
- Yijie Guo, Shengyu Feng, Nicolas Le Roux, Ed H. Chi, Honglak Lee, Minmin Chen:
Batch Reinforcement Learning Through Continuation Method.
- Zhiyuan Li , Yuping Luo, Kaifeng Lyu:
Towards Resolving the Implicit Bias of Gradient Descent for Matrix Factorization: Greedy Low-Rank Learning.
- Jonathan Frankle, David J. Schwab, Ari S. Morcos:
Training BatchNorm and Only BatchNorm: On the Expressive Power of Random Features in CNNs.
- Manas Sahni, Shreya Varshini, Alind Khare, Alexey Tumanov:
CompOFA - Compound Once-For-All Networks for Faster Multi-Platform Deployment.
- Yu-Ying Chou, Hsuan-Tien Lin, Tyng-Luh Liu:
Adaptive and Generative Zero-Shot Learning.
- Ayya Alieva, Aiden Aceves, Jialin Song, Stephen Mayo, Yisong Yue, Yuxin Chen:
Learning to Make Decisions via Submodular Regularization.
- Utkarsh Ojha, Krishna Kumar Singh, Yong Jae Lee:
Generating Furry Cars: Disentangling Object Shape and Appearance across Multiple Domains.
- Aviral Kumar, Rishabh Agarwal, Dibya Ghosh, Sergey Levine:
Implicit Under-Parameterization Inhibits Data-Efficient Deep Reinforcement Learning.
- Mihir Prabhudesai, Shamit Lal, Darshan Patil, Hsiao-Yu Tung, Adam W. Harley, Katerina Fragkiadaki:
Disentangling 3D Prototypical Networks for Few-Shot Concept Learning.
- Yilun Xu, Yang Song, Sahaj Garg, Linyuan Gong, Rui Shu, Aditya Grover, Stefano Ermon:
Anytime Sampling for Autoregressive Models via Ordered Autoencoding.
- Zhou Xian, Shamit Lal, Hsiao-Yu Tung, Emmanouil Antonios Platanios, Katerina Fragkiadaki:
HyperDynamics: Meta-Learning Object and Agent Dynamics with Hypernetworks.
- Siyang Yuan, Pengyu Cheng, Ruiyi Zhang, Weituo Hao, Zhe Gan, Lawrence Carin:
Improving Zero-Shot Voice Style Transfer via Disentangled Representation Learning.
- Tao Yu, Chien-Sheng Wu, Xi Victoria Lin, Bailin Wang, Yi Chern Tan, Xinyi Yang, Dragomir R. Radev, Richard Socher, Caiming Xiong:
GraPPA: Grammar-Augmented Pre-Training for Table Semantic Parsing.
- Chirag Pabbaraju, Ezra Winston, J. Zico Kolter:
Estimating Lipschitz constants of monotone deep equilibrium models.
- Hrayr Harutyunyan, Alessandro Achille, Giovanni Paolini, Orchid Majumder, Avinash Ravichandran, Rahul Bhotika, Stefano Soatto:
Estimating informativeness of samples with Smooth Unique Information.
- Alvin Wan, Lisa Dunlap, Daniel Ho, Jihan Yin, Scott Lee, Suzanne Petryk, Sarah Adel Bargal, Joseph E. Gonzalez:
NBDT: Neural-Backed Decision Tree.
- Jinheon Baek, Minki Kang, Sung Ju Hwang:
Accurate Learning of Graph Representations with Graph Multiset Pooling.
- Zeyuan Allen-Zhu, Faeze Ebrahimiaghazani, Jerry Li, Dan Alistarh:
Byzantine-Resilient Non-Convex Stochastic Gradient Descent.
- Ying-Jun Du, Xiantong Zhen, Ling Shao, Cees G. M. Snoek:
MetaNorm: Learning to Normalize Few-Shot Batches Across Domains.
- Brennan Shacklett, Erik Wijmans, Aleksei Petrenko, Manolis Savva, Dhruv Batra, Vladlen Koltun, Kayvon Fatahalian:
Large Batch Simulation for Deep Reinforcement Learning.
- Michael Zhang, Karan Sapra, Sanja Fidler, Serena Yeung, José M. Álvarez:
Personalized Federated Learning with First Order Model Optimization.
- Arlei Lopes da Silva, Furkan Kocayusufoglu, Saber Jafarpour, Francesco Bullo, Ananthram Swami, Ambuj K. Singh:
Combining Physics and Machine Learning for Network Flow Estimation.
- Tri Dao, Govinda M. Kamath, Vasilis Syrgkanis, Lester Mackey:
Knowledge Distillation as Semiparametric Inference.
- Yannis Flet-Berliac, Reda Ouhamma, Odalric-Ambrym Maillard, Philippe Preux:
Learning Value Functions in Deep Policy Gradients using Residual Variance.
- Xinyue Chen, Che Wang, Zijian Zhou, Keith W. Ross:
Randomized Ensembled Double Q-Learning: Learning Fast Without a Model.



- Changhoon Kim, Yi Ren, Yezhou Yang:
Decentralized Attribution of Generative Models.
- Weijian Xu, Yifan Xu, Huaijin Wang, Zhuowen Tu:
Attentional Constellation Nets for Few-Shot Learning.
- Michael Dann , John Thangarajah:
Adapting to Reward Progressivity via Spectral Reinforcement Learning.
- Martin Trimmel, Henning Petzka, Cristian Sminchisescu:
TropEx: An Algorithm for Extracting Linear Terms in Deep Neural Networks.
- Kevin Lu, Aditya Grover, Pieter Abbeel, Igor Mordatch:
Reset-Free Lifelong Learning with Skill-Space Planning.
- Yu Cheng, Honghao Lin:
Robust Learning of Fixed-Structure Bayesian Networks in Nearly-Linear Time.
- Christopher Hahn, Frederik Schmitt, Jens U. Kreber, Markus Norman Rabe, Bernd Finkbeiner:
Teaching Temporal Logics to Neural Networks.
- Nasim Rahaman, Anirudh Goyal, Muhammad Waleed Gondal, Manuel Wuthrich, Stefan Bauer, Yash Sharma, Yoshua Bengio, Bernhard Schölkopf:
Spatially Structured Recurrent Modules.
- Jake Snell, Richard S. Zemel:
Bayesian Few-Shot Classification with One-vs-Each Pólya-Gamma Augmented Gaussian Processes.
- Francesco Faccio, Louis Kirsch, Jürgen Schmidhuber:
Parameter-Based Value Functions.
- Ryohei Shimizu, Yusuke Mukuta, Tatsuya Harada:
Hyperbolic Neural Networks++.
- Calypso Herrera, Florian Krach , Josef Teichmann:
Neural Jump Ordinary Differential Equations: Consistent Continuous-Time Prediction and Filtering.
- Csaba Tóth, Patric Bonnier, Harald Oberhauser:
Seq2Tens: An Efficient Representation of Sequences by Low-Rank Tensor Projections.
- Lanqing Li, Rui Yang, Dijun Luo:
FOCAL: Efficient Fully-Offline Meta-Reinforcement Learning via Distance Metric Learning and Behavior Regularization.
- Zhong Li, Jiequn Han, Weinan E, Qianxiao Li:
On the Curse of Memory in Recurrent Neural Networks: Approximation and Optimization Analysis.
- Shashank Srikant, Sijia Liu, Tamara Mitrovska, Shiyu Chang, Quanfu Fan, Gaoyuan Zhang, Una-May O'Reilly:
Generating Adversarial Computer Programs using Optimized Obfuscations.
- Jaehoon Oh, Hyungjun Yoo, ChangHwan Kim, Se-Young Yun:
BOIL: Towards Representation Change for Few-shot Learning.
- Hao Zhang, Sen Li, Yinchao Ma, Mingjie Li, Yichen Xie, Quanshi Zhang:
Interpreting and Boosting Dropout from a Game-Theoretic View.
- Jovana Mitrovic, Brian McWilliams, Jacob C. Walker, Lars Holger Buesing, Charles Blundell:
Representation Learning via Invariant Causal Mechanisms.
- Dániel Zombori, Balázs Bánhelyi, Tibor Csendes, István Megyeri, Márk Jelasity:
Fooling a Complete Neural Network Verifier.
- Sungmin Cha, Hsiang Hsu, Taebaek Hwang, Flávio P. Calmon, Taesup Moon:
CPR: Classifier-Projection Regularization for Continual Learning.
- Chen Wei, Huiyu Wang, Wei Shen, Alan L. Yuille:
CO2: Consistent Contrast for Unsupervised Visual Representation Learning.
- Sungmin Cha, Taeeon Park, Byeongjoon Kim, Jongduk Baek, Taesup Moon:
GAN2GAN: Generative Noise Learning for Blind Denoising with Single Noisy Images.
- Siyuan Li, Lulu Zheng, Jianhao Wang, Chongjie Zhang:
Learning Subgoal Representations with Slow Dynamics.
- Zhipeng Bao, Yu-Xiong Wang, Martial Hebert:
Bowtie Networks: Generative Modeling for Joint Few-Shot Recognition and Novel-View Synthesis.
- Tatjana Chavdarova , Matteo Pagliardini, Sebastian U. Stich, François Fleuret, Martin Jaggi:
Taming GANs with Lookahead-Minmax.
- Mark Niklas Müller, Mislav Balunovic, Martin T. Vechev:
Certify or Predict: Boosting Certified Robustness with Compositional Architectures.
- Peter Davies, Vijaykrishna Gurunanthan, Niusha Moshrefi, Saleh Ashkboos, Dan Alistarh:
New Bounds For Distributed Mean Estimation and Variance Reduction.
- Bin Xin Ru, Xingchen Wan, Xiaowen Dong, Michael A. Osborne:
Interpretable Neural Architecture Search via Bayesian Optimisation with Weisfeiler-Lehman Kernels.



- Hideaki Hayashi, Seiichi Uchida:
A Discriminative Gaussian Mixture Model with Sparsity.
- Woojun Kim, Jongeui Park, Youngchul Sung:
Communication in Multi-Agent Reinforcement Learning: Intention Sharing.
- Zhengyang Geng, Meng-Hao Guo, Hongxu Chen, Xia Li, Ke Wei, Zhouchen Lin:
Is Attention Better Than Matrix Decomposition?
- Kaidi Xu, Huan Zhang, Shiqi Wang, Yihan Wang, Suman Jana, Xue Lin, Cho-Jui Hsieh:
Fast and Complete: Enabling Complete Neural Network Verification with Rapid and Massively Parallel Incomplete Verifiers.
- Bin Xu, Huan Zhang, Shiqi Wang, Yihan Wang, Suman Jana, Xue Lin, Cho-Jui Hsieh:
A Geometric Analysis of Deep Generative Image Models and Its Applications.
- Yunfei Li, Yilin Wu, Huazhe Xu, Xiaolong Wang, Yi Wu:
Solving Compositional Reinforcement Learning Problems via Task Reduction.
- Kangkang Lu, Cuong Manh Nguyen, Xun Xu, Kiran Chari, Yu Jing Goh, Chuan-Sheng Foo:
ARMoured: Adversarially Robust MOdels using Unlabeled data by REgularizing Diversity.
- Esther Derman, Gal Dalal, Shie Mannor:
Acting in Delayed Environments with Non-Stationary Markov Policies.
- Ties van Rozendaal, Iris A. M. Huijben, Taco Cohen:
Overfitting for Fun and Profit: Instance-Adaptive Data Compression.
- Siyi Liu, Chen Gao, Yihong Chen, Depeng Jin, Yong Li:
Learnable Embedding sizes for Recommender Systems.
- Fei Deng, Zhuo Zhi, Donghun Lee, Sungjin Ahn:
Generative Scene Graph Networks.
- Yann N. Dauphin, Ekin Dogus Cubuk:
Deconstructing the Regularization of BatchNorm.
- Xiongwei Wu, Doyen Sahoo, Steven C. H. Hoi:
PolarNet: Learning to Optimize Polar Keypoints for Keypoint Based Object Detection.
- Hao Zhu, Piotr Koniusz:
Simple Spectral Graph Convolution.
- Zhen Han, Peng Chen, Yunpu Ma, Volker Tresp:
Explainable Subgraph Reasoning for Forecasting on Temporal Knowledge Graphs.
- Like Hui, Mikhail Belkin:
Evaluation of Neural Architectures trained with square Loss vs Cross-Entropy in Classification Tasks.

last updated on 2024-12-11 20:32 CET by the dblp team

all metadata released as open data under CC0 1.0 license

see also: [Terms of Use](#) | [Privacy Policy](#) | [Imprint](#)

dblp was originally created in 1993 at:



since 2018, dblp has been operated and maintained by:



the dblp computer science bibliography is funded and supported by:

