



We've just launched a new service: our brand new **dblp SPARQL query service**. Read more about it in our latest blog post or try out some of the SPARQL queries linked on the dblp web pages below.



[+] 12th ICLR 2024: Vienna, Austria []

> Home > Conferences and Workshops > ICLR

Dagstuhl

The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024. OpenReview.net 2024

[+] SPARQL queries

[–] Refine list

showing all 2,261 records

refine by search term

refine by author

Tongliang Liu (14)
Junchi Yan (13)
Yu Qiao 0001 (13)
Bo Han 0003 (13)
Yoshua Bengio (11)
Mihaela van der Schaar (10)
Zhangyang Wang (10)
Furong Huang (10)
Chelsea Finn (10)
Joshua M. Susskind (9)
9,201 more options

refine by access

open (2,261)

Accept (oral)

- Yonatan Oren, Nicole Meister, Niladri S. Chatterji, Faisal Ladhak, Tatsunori Hashimoto:
Proving Test Set Contamination in Black-Box Language Models.
- Yapei Chang, Kyle Lo, Tanya Goyal, Mohit Iyyer:
BoooookScore: A systematic exploration of book-length summarization in the era of LLMs.
- Zahra Kadkhodaie, Florentin Guth, Eero P. Simoncelli, Stéphane Mallat:
Generalization in diffusion models arises from geometry-adaptive harmonic representations.
- Satwik Bhattacharya, Arkil Patel, Phil Blunsom, Varun Kanade:
Understanding In-Context Learning in Transformers and LLMs by Learning to Learn Discrete Functions.
- Gautam Reddy:
The mechanistic basis of data dependence and abrupt learning in an in-context classification task.
- Yang Song, Prafulla Dhariwal:
Improved Techniques for Training Consistency Models.
- Thaddäus Wiedemer, Jack Brady, Alexander Panfilov, Attila Juhos, Matthias Bethge, Wieland Brendel:
Provable Compositional Generalization for Object-Centric Learning.
- Ching Fang, Kim Stachenfeld:
Predictive auxiliary objectives in deep RL mimic learning in the brain.
- Haoqi Yuan, Zhancun Mu, Feiyang Xie, Zongqing Lu:
Pre-Training Goal-based Models for Sample-Efficient Reinforcement Learning.
- Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, Peter Henderson:
Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!
- Ido Amos, Jonathan Berant, Ankit Gupta:
Never Train from Scratch: Fair Comparison of Long-Sequence Models Requires Data-Driven Priors.
- Yixiao Li, Yifan Yu, Chen Liang, Nikos Karampatziakis, Pengcheng He, Weizhu Chen, Tuo Zhao:
LoftQ: LoRA-Fine-Tuning-aware Quantization for Large Language Models.
- Miltiadis Kofinas, Boris Knyazev, Yan Zhang, Yunlu Chen, Gertjan J. Burghouts, Efstratios Gavves, Cees G. M. Snoek, David W. Zhang:
Graph Neural Networks for Learning Equivariant Representations of Neural Networks.
- Zaishuo Xia, Han Yang, Binghui Wang, Jinyuan Jia:
GNNCert: Deterministic Certification of Graph Neural Networks against Adversarial Perturbations.
- Hyungho Na, Yunkyeong Seo, Il-Chul Moon:
Efficient Episodic Memory Utilization of Cooperative Multi-Agent Reinforcement Learning.
- Yogesh Verma, Markus Heinonen, Vikas Garg:
ClimODE: Climate and Weather Forecasting with Physics-informed Neural ODEs.
- Hengrui Zhang, Jian Zhang, Zhengyuan Shen, Balasubramaniam Srinivasan, Xiao Qin, Christos Faloutsos, Huzeifa Rangwala, George Karypis:
Mixed-Type Tabular Data Synthesis with Score-based Diffusion in Latent Space.



- Linlu Qiu, Liwei Jiang, Ximing Lu, Melanie Sclar, Valentina Pyatkin, Chandra Bhagavatula, Bailin Wang, Yoon Kim, Yejin Choi, Nouha Dziri, Xiang Ren:
Phenomenal Yet Puzzling: Testing Inductive Reasoning Capabilities of Language Models with Hypothesis Refinement.
- Bohang Zhang, Jingchu Gai, Yiheng Du, Qiwei Ye, Di He, Liwei Wang:
Beyond Weisfeiler-Lehman: A Quantitative Framework for GNN Expressiveness.
- Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, Jianfeng Gao:
MathVista: Evaluating Mathematical Reasoning of Foundation Models in Visual Contexts.
- Nathan C. Frey, Daniel Berenberg, Karina Zadorozhny, Joseph Kleinhenz, Julien Lafrance-Vanassee, Isidro Hötzl, Yan Wu, Stephen Ra, Richard Bonneau, Kyunghyun Cho, Andreas Loukas, Vladimir Glorijevic, Saeed Saremi:
Protein Discovery with Discrete Walk-Jump Sampling.
- Kensen Shi, Joey Hong, Yinlin Deng, Pengcheng Yin, Manzil Zaheer, Charles Sutton:
ExeDec: Execution Decomposition for Compositional Generalization in Neural Program Synthesis.
- Yeming Wen, Swarat Chaudhuri:
Batched Low-Rank Adaptation of Foundation Models.
- Atsushi Shimizu, Xiaoou Cheng, Christopher Musco, Jonathan Weare:
Improved Active Learning via Dependent Leverage Score Sampling.
- Suyu Ge, Yunan Zhang, Liyuan Liu, Minjia Zhang, Jiawei Han, Jianfeng Gao:
Model Tells You What to Discard: Adaptive KV Cache Compression for LLMs.
- Galen Andrew, Peter Kairouz, Sewoong Oh, Alina Oprea, Hugh Brendan McMahan, Vinith Menon Suriyakumar:
One-shot Empirical Privacy Estimation for Federated Learning.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, Karthik R. Narasimhan:
SWE-bench: Can Language Models Resolve Real-world Github Issues?
- Seyed-Iman Mirzadeh, Keivan Alizadeh-Vahid, Sachin Mehta, Carlo C. del Mundo, Oncel Tuzel, Golnoosh Samei, Mohammad Rastegari, Mehrdad Farajtabar:
ReLU Strikes Back: Exploiting Activation Sparsity in Large Language Models.
- Yiding Jiang, Christina Baek, J. Zico Kolter:
On the Joint Interaction of Models, Data, and Features.
- Ismail Yunus Akhalwaya, Shashanka Ubaru, Kenneth L. Clarkson, Mark S. Squillante, Vishnu Jejjala, Yang-Hui He, Kugendran Naidoo, Vasileios Kalantzis, Lior Horesh:
Topological data analysis on noisy quantum computers.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, Hannaneh Hajishirzi:
Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection.
- Anshuman Chhabra, Peizhao Li, Prasant Mohapatra, Hongfu Liu:
"What Data Benefits My Classifier?" Enhancing Model Performance and Interpretability through Influence-Based Data Selection.
- Tianrong Chen, Jiatao Gu, Laurent Dinh, Evangelos A. Theodorou, Joshua M. Susskind, Shuangfei Zhai:
Generative Modeling with Phase Stochastic Bridge.
- Zengwei Yao, Liyong Guo, Xiaoyu Yang, Wei Kang, Fangjun Kuang, Yifan Yang, Zengrui Jin, Long Lin, Daniel Povey:
Zipformer: A faster and better encoder for automatic speech recognition.
- Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Jinlin Wang, Ceyao Zhang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, Lingfeng Xiao, Chenglin Wu, Jürgen Schmidhuber:
MetaGPT: Meta Programming for A Multi-Agent Collaborative Framework.
- Kim-Celine Kahl, Carsten T. Lüth, Maximilian Zenk, Klaus H. Maier-Hein, Paul F. Jaeger:
ValUES: A Framework for Systematic Validation of Uncertainty Estimation in Semantic Segmentation.
- Xudong Shen, Chao Du, Tianyu Pang, Min Lin, Yongkang Wong, Mohan S. Kankanhalli:
Finetuning Text-to-Image Diffusion Models for Fairness.
- Shangbin Feng, Weijia Shi, Yuyang Bai, Vidhisha Balachandran, Tianxing He, Yulia Tsvetkov:
Knowledge Card: Filling LLMs' Knowledge Gaps with Plug-in Specialized Language Models.
- Seohong Park, Oleh Rybkin, Sergey Levine:
METRA: Scalable Unsupervised RL with Metric-Aware Abstraction.
- Yichen Wu, Long-Kai Huang, Renzhen Wang, Deyu Meng, Ying Wei:
Meta Continual Learning Revisited: Implicitly Enhancing Online Hessian Approximation via Variance Reduction.
- Bo Zhao, Robert M. Gower, Robin Walters, Rose Yu:
Improving Convergence and Generalization Using Parameter Symmetries.
- Ricky T. Q. Chen, Yaron Lipman:
Flow Matching on General Geometries.
- Zhen Liu, Yao Feng, Yuliang Xiu, Weiyang Liu, Liam Paull, Michael J. Black, Bernhard Schölkopf:
Ghost on the Shell: An Expressive Representation of General 3D Shapes.



Pablo Pernias, Dominic Rampas, Mats L. Richter, Christopher Pal, Marc Aubreville:
Würstchen: An Efficient Architecture for Large-Scale Text-to-Image Diffusion Models.



Yuxuan Song, Jingjing Gong, Hao Zhou, Mingyue Zheng, Jingjing Liu, Wei-Ying Ma:
Unified Generative Modeling of 3D Molecules with Bayesian Flow Networks.



Mitchell Wortsman, Peter J. Liu, Lechao Xiao, Katie E. Everett, Alexander A. Alemi, Ben Adlam, John D. Co-Reyes, Izzeddin Gur, Abhishek Kumar, Roman Novak, Jeffrey Pennington, Jascha Sohl-Dickstein, Kelvin Xu, Jaehoon Lee, Justin Gilmer, Simon Kornblith:
Small-scale proxies for large-scale Transformer training instabilities.



Pascal Chang, Jingwei Tang, Markus Gross, Vinicius C. Azevedo:
How I Warped Your Noise: a Temporally-Correlated Noise Prior for Diffusion Models.



Timothée Darcet, Maxime Oquab, Julien Mairal, Piotr Bojanowski:
Vision Transformers Need Registers.



Sergei Solonets, Daniil Sintysyn, Lukas von Stumberg, Nikita Araslanov, Daniel Cremers:
An Analytical Solution to Gauss-Newton Loss for Direct Image Alignment.



Hyosoon Jang, Minsu Kim, Sungsoo Ahn:
Learning Energy Decompositions for Partial Inference in GFlowNets.



Ian Gemp, Luke Marris, Georgios Piliouras:
Approximating Nash Equilibria in Normal-Form Games via Stochastic Optimization.



Giorgio Mariani, Irene Tallini, Emilian Postolache, Michele Mancusi, Luca Cosmo, Emanuele Rodolà:
Multi-Source Diffusion Models for Simultaneous Music Generation and Separation.



Haiming Wang, Huajian Xin, Chuanyang Zheng, Zhengying Liu, Qingxing Cao, Yinya Huang, Jing Xiong, Han Shi, Enze Xie, Jian Yin, Zhenguo Li, Xiaodan Liang:
LEGO-Prover: Neural Theorem Proving with Growing Libraries.



Marius Memmel, Andrew Wagenmaker, Chunling Zhu, Dieter Fox, Abhishek Gupta:
ASID: Active Exploration for System Identification in Robotic Manipulation.



Germain Kolossov, Andrea Montanari, Pulkit Tandon:
Towards a statistical theory of data selection under weak supervision.



Mohammad Reza Samsami, Artem Zholus, Janarthanan Rajendran, Sarath Chandar:
Mastering Memory Tasks with World Models.



Gabriel Cardoso, Yazid Janati El Idrissi, Sylvain Le Corff, Eric Moulines:
Monte Carlo guided Denoising Diffusion models for Bayesian linear inverse problems.



Xian Li, Ping Yu, Chunting Zhou, Timo Schick, Omer Levy, Luke Zettlemoyer, Jason Weston, Mike Lewis:
Self-Alignment with Instruction Backtranslation.



Sherry Yang, Yilun Du, Seyed Kamyar Seyed Ghasemipour, Jonathan Tompson, Leslie Pack Kaelbling, Dale Schuurmans, Pieter Abbeel:
Learning Interactive Real-World Simulators.



Shuo He, Chaojie Wang, Guowu Yang, Lei Feng:
Candidate Label Set Pruning: A Data-centric Perspective for Deep Partial-label Learning.



Jonathan Richens, Tom Everitt:
Robust agents learn causal world models.



Jen-tse Huang, Wenxuan Wang, Eric John Li, Man Ho Lam, Shujie Ren, Youliang Yuan, Wenxiang Jiao, Zhaopeng Tu, Michael R. Lyu:
On the Humanity of Conversational AI: Evaluating the Psychological Portrayal of LLMs.



Jisu Nam, Gyuseong Lee, Sunwoo Kim, Hyeonsu Kim, Hyoungwon Cho, Seyeon Kim, Seungryong Kim:
Diffusion Model for Dense Matching.



Shashanka Venkataramanan, Mamshad Nayeem Rizve, João Carreira, Yuki M. Asano, Yannis Avrithis:
Is ImageNet worth 1 video? Learning strong image encoders from 1 long unlabelled video.



Panagiotis Eustratiadis, Lukasz Dudziak, Da Li, Timothy M. Hospedales:
Neural Fine-Tuning Search for Few-Shot Learning.



Qiuhan Zeng, Changjian Shui, Long-Kai Huang, Peng Liu, Xi Chen, Charles Ling, Boyu Wang:
Latent Trajectory Learning for Limited Timestamps under Distribution Shift over Time.



Ruoyu Chen, Hua Zhang, Siyuan Liang, Jingzhi Li, Xiaochun Cao:
Less is More: Fewer Interpretable Region via Submodular Subset Selection.



Jason Y. Zhang, Amy Lin, Moneish Kumar, Tzu-Hsuan Yang, Deva Ramanan, Shubham Tulsiani:
Cameras as Rays: Pose Estimation via Ray Diffusion.



Jie Hu, Vishwaraj Doshi, Do Young Eun:
Accelerating Distributed Stochastic Optimization via Self-Repellent Random Walks.



Yuxin Wen, Yuchen Liu, Chen Chen, Lingjuan Lyu:
Detecting, Explaining, and Mitigating Memorization in Diffusion Models.



Sebastian Pineda-Arango, Fabio Ferreira, Arlind Kadra, Frank Hutter, Josif Grabocka:
Quick-Tune: Quickly Learning Which Pretrained Model to Finetune and How.



- Yukang Chen, Shengju Qian, Haotian Tang, Xin Lai, Zhijian Liu, Song Han, Jiaya Jia: **LongLoRA: Efficient Fine-tuning of Long-Context Large Language Models.**
- Edward J. Hu, Moksh Jain, Eric Elmoznino, Younesse Kaddar, Guillaume Lajoie, Yoshua Bengio, Nikolay Malkin: **Amortizing intractable inference in large language models.**
- Ahmad Faiz, Sotaro Kaneda, Ruhan Wang, Rita Chukwunyere Osi, Prateek Sharma, Fan Chen, Lei Jiang: **LLMCarbon: Modeling the End-to-End Carbon Footprint of Large Language Models.**
- Izzeddin Gur, Hiroki Furuta, Austin V. Huang, Mustafa Safdari, Yutaka Matsuo, Douglas Eck, Aleksandra Faust: **A Real-World WebAgent with Planning, Long Context Understanding, and Program Synthesis.**
- Zhantao Yang, Ruili Feng, Han Zhang, Yujun Shen, Kai Zhu, Lianghua Huang, Yifei Zhang, Yu Liu, Deli Zhao, Jingren Zhou, Fan Cheng: **Lipschitz Singularities in Diffusion Models.**
- Yossi Gandelsman, Alexei A. Efros, Jacob Steinhardt: **Interpreting CLIP's Image Representation via Text-Based Decomposition.**
- Yang He, Lingao Xiao, Joey Tianyi Zhou, Ivor W. Tsang: **Multisize Dataset Condensation.**
- Jiaxiang Tang, Jiawei Ren, Hang Zhou, Ziwei Liu, Gang Zeng: **DreamGaussian: Generative Gaussian Splattering for Efficient 3D Content Creation.**
- Yicong Hong, Kai Zhang, Jiuxiang Gu, Sai Bi, Yang Zhou, Difan Liu, Feng Liu, Kalyan Sunkavalli, Trung Bui, Hao Tan: **LRM: Large Reconstruction Model for Single Image to 3D.**
- Wenxuan Li, Alan L. Yuille, Zongwei Zhou: **How Well Do Supervised 3D Models Transfer to Medical Imaging Tasks?**
- Haoyue Dai, Ignavier Ng, Gongxu Luo, Peter Spirtes, Petar Stojanov, Kun Zhang: **Gene Regulatory Network Inference in the Presence of Dropouts: a Causal View.**
- Yubo Zhuang, Xiaohui Chen, Yun Yang, Richard Y. Zhang: **Statistically Optimal K-means Clustering via Nonnegative Low-rank Semidefinite Programming.**
- André F. Cruz , Moritz Hardt: **Unprocessing Seven Years of Algorithmic Fairness.**
- Ziheng Qin, Kai Wang, Zangwei Zheng, Jianyang Gu, Xiangyu Peng, Zhaopan Xu, Daquan Zhou, Lei Shang, Baigui Sun, Xuansong Xie, Yang You: **InfoBatch: Lossless Training Speed Up by Unbiased Dynamic Data Pruning.**
- Yijie Lin, Jie Zhang, Zhenyu Huang, Jia Liu, Zujie Wen, Xi Peng: **Multi-granularity Correspondence Learning from Long-term Noisy Videos.**

Accept (spotlight)

- Sravanthi Gurugubelli, Sundeep Prabhakar Chepuri: **SaNN: Simple Yet Powerful Simplicial-aware Neural Networks.**
- Robin Staab, Mark Vero, Mislav Balunovic, Martin T. Vechev: **Beyond Memorization: Violating Privacy via Inference with Large Language Models.**
- Jasper Dekoninck, Marc Fischer, Luca Beurer-Kellner, Martin T. Vechev: **Controlled Text Generation via Language Model Arithmetic.**
- Nan Chen, Zemin Liu, Bryan Hooi, Bingsheng He, Rizal Fathony, Jun Hu, Jia Chen: **Consistency Training with Learnable Data Augmentation for Graph Anomaly Detection with Limited Supervision.**
- Juno Kim, Kakei Yamamoto, Kazusato Oko, Zhuoran Yang, Taiji Suzuki: **Symmetric Mean-field Langevin Dynamics for Distributional Minimax Problems.**
- Suhan Shetty, Teng Xue, Sylvain Calinon: **Generalized Policy Iteration using Tensor Approximation for Hybrid Control.**
- Maksim Velikanov, Maxim Panov, Dmitry Yarotsky: **Generalization error of spectral algorithms.**
- Haoxuan Li, Chunyuan Zheng, Yanghao Xiao, Peng Wu, Zhi Geng, Xu Chen, Peng Cui: **Debiased Collaborative Filtering with Kernel-Based Causal Balancing.**
- Cassidy Laidlaw, Banghua Zhu, Stuart Russell, Anca D. Dragan: **The Effective Horizon Explains Deep RL Performance in Stochastic Environments.**
- Ainaz Eftekhari, Kuo-Hao Zeng, Jiafei Duan, Ali Farhadi, Aniruddha Kembhavi, Ranjay Krishna: **Selective Visual Representations Improve Convergence and Generalization for Embodied AI.**
- Rui Zheng, Wei Shen, Yuan Hua, Wenbin Lai, Shihan Dou, Yuhao Zhou, Zhiheng Xi, Xiao Wang, Haoran Huang, Tao Gui, Qi Zhang, Xuanjing Huang: **Improving Generalization of Alignment with Human Preferences through Group Invariant Learning.**
- Gregory Kang Ruey Lau, Apivich Hemachandra, See-Kiong Ng, Bryan Kian Hsiang Low: **PINNACLE: PINN Adaptive ColLocation and Experimental points selection.**



- Mikhail Khodak, Edmond Chow, Maria-Florina Balcan, Ameet Talwalkar:
Learning to Relax: Setting Solver Parameters Across a Sequence of Linear System Instances.
- Guodong Wang, Yunhong Wang, Xiuguo Bao, Di Huang:
Rotation Has Two Sides: Evaluating Data Augmentation for Deep One-class Classification.
- Lin-Han Jia, Lan-Zhe Guo, Zhi Zhou, Yufeng Li:
Realistic Evaluation of Semi-supervised Learning Algorithms in Open Environments.
- Denizalp Goktas, Amy Greenwald, Sadie Zhao, Alec Koppel, Sumitra Ganesh:
Efficient Inverse Multiagent Learning.
- Tianqi Du, Yifei Wang, Yisen Wang:
On the Role of Discrete Tokenization in Visual Representation Learning.
- Athul Paul Jacob, Yikang Shen, Gabriele Farina, Jacob Andreas:
The Consensus Game: Language Model Generation via Equilibrium Search.
- Jake Grigsby, Linxi Fan, Yuke Zhu:
AMAGO: Scalable In-Context Reinforcement Learning for Adaptive Agents.
- Zhuqing Liu, Xin Zhang, Jia Liu, Zhengyuan Zhu, Songtao Lu:
PILOT: An $\mathcal{O}(1/K)$ -Convergent Approach for Policy Evaluation with Nonlinear Function Approximation.
- Ted Moskovitz, Aaditya K. Singh, DJ Strouse, Tuomas Sandholm, Ruslan Salakhutdinov, Anca D. Dragan, Stephen Marcus McAleer:
Confronting Reward Model Overoptimization with Constrained RLHF.
- Vimal Thilak, Chen Huang, Omid Saremi, Laurent Dinh, Hanlin Goh, Preetum Nakkiran, Joshua M. Susskind, Eta Littwin:
LiDAR: Sensing Linear Probing Performance in Joint Embedding SSL Architectures.
- Jiayang Liu, Yiming Bu, Daniel Tso, Qinru Qiu:
Improved Efficiency Based on Learned Saccade and Continuous Scene Reconstruction From Foveated Visual Sampling.
- Danny Halawi, Jean-Stanislas Denain, Jacob Steinhardt:
Overthinking the Truth: Understanding how Language Models Process False Demonstrations.
- Ibraheem Muhammad Moosa, Rui Zhang, Wenpeng Yin:
MT-Ranker: Reference-free machine translation evaluation by inter-system ranking.
- Zayne Sprague, Xi Ye, Kaj Bostrom, Swarat Chaudhuri, Greg Durrett:
MuSR: Testing the Limits of Chain-of-thought with Multistep Soft Reasoning.
- Philip Amortila, Dylan J. Foster, Nan Jiang, Ayush Sekhari, Tengyang Xie:
Harnessing Density Ratios for Online Reinforcement Learning.
- Hanqi Zhou, Robert Bamler, Charley M. Wu, Álvaro Tejero-Cantero:
Predictive, scalable and interpretable knowledge tracing on structured domains.
- Irene Cannistraci, Luca Moschella, Marco Fumero, Valentino Maiorca, Emanuele Rodolà:
From Bricks to Bridges: Product of Invariances to Enhance Latent Space Communication.
- Sumeet Batra, Bryon Tjanaka, Matthew Christopher Fontaine, Aleksei Petrenko, Stefanos Nikolaidis, Gaurav S. Sukhatme:
Proximal Policy Gradient Arborescence for Quality Diversity Reinforcement Learning.
- Sadegh Mahdavi, Renjie Liao, Christos Thrampoulidis:
Memorization Capacity of Multi-Head Attention in Transformers.
- Jack Merullo, Carsten Eickhoff, Ellie Pavlick:
Circuit Component Reuse Across Tasks in Transformer Language Models.
- Henry Li, Ronen Basri, Yuval Kluger:
Likelihood Training of Cascaded Diffusion Models via Hierarchical Volume-preserving Maps.
- Divyat Mahajan, Ioannis Mitliagkas, Brady Neal, Vasilis Syrgkanis:
Empirical Analysis of Model Selection for Heterogeneous Causal Effect Estimation.
- Ali Shahin Shamsabadi, Gefei Tan, Tudor Cebere, Aurélien Bellet, Hamed Haddadi, Nicolas Papernot, Xiao Wang, Adrian Weller:
Confidential-DPproof: Confidential Proof of Differentially Private Training.
- Weijia Shi, Sewon Min, Maria Lomeli, Chunting Zhou, Margaret Li, Xi Victoria Lin, Noah A. Smith, Luke Zettlemoyer, Wentai Yih, Mike Lewis:
In-Context Pretraining: Language Modeling Beyond Document Boundaries.
- Yanai Elazar, Akshita Bhagia, Ian Magnusson, Abhilasha Ravichander, Dustin Schwenk, Alane Suhr, Evan Pete Walsh, Dirk Groeneveld, Luca Soldaini, Sameer Singh, Hannaneh Hajishirzi, Noah A. Smith, Jesse Dodge:
What's In My Big Data?
- Victor Livernoche, Vineet Jain, Yashar Hezaveh, Siyamak Ravanbakhsh:
On Diffusion Modeling for Anomaly Detection.
- Arman Isajanyan, Artur Shatveryan, David Kocharian, Zhangyang Wang, Humphrey Shi:
Social Reward: Evaluating and Enhancing Generative AI through Million-User Feedback from an Online Creative Community.



- Aakash Lahoti, Stefani Karp, Ezra Winston, Aarti Singh, Yuanzhi Li:
Role of Locality and Weight Sharing in Image-Based Tasks: A Sample Complexity Separation between CNNs, LCNs, and FCNs.
- Lizhang Chen, Bo Liu, Kaizhao Liang, Qiang Liu:
Lion Secretly Solves a Constrained Optimization: As Lyapunov Predicts.
- Ronak Mehta, Vincent Roulet, Krishna Pillutla, Zaïd Harchaoui:
Distributionally Robust Optimization with Bias and Variance Reduction.
- Garrett Tanzer, Mirac Suzgun, Eline Visser, Dan Jurafsky, Luke Melas-Kyriazi:
A Benchmark for Learning to Translate a New Language from One Grammar Book.
- Sinong Geng, Aldo Pacchiano, Andrey Kolobov, Ching-An Cheng:
Improving Offline RL by Blending Heuristics.
- Jeonghye Kim, Suyoung Lee, Woojun Kim, Youngchul Sung:
Decision ConvFormer: Local Filtering in MetaFormer is Sufficient for Decision Making.
- Jingfeng Wu, Difan Zou, Zixiang Chen, Vladimir Braverman, Quanquan Gu, Peter L. Bartlett:
How Many Pretraining Tasks Are Needed for In-Context Learning of Linear Regression?
- Lei Li, Yekun Chai, Shuhuan Wang, Yu Sun, Hao Tian, Ningyu Zhang, Hua Wu:
Tool-Augmented Reward Modeling.
- Fan-Ming Luo, Tian Xu, Xingchen Cao, Yang Yu:
Reward-Consistent Dynamics Models are Strongly Generalizable for Offline Reinforcement Learning.
- Haolin Liu, Chen-Yu Wei, Julian Zimmert:
Towards Optimal Regret in Adversarial Linear MDPs with Bandit Feedback.
- Harshit Sikchi, Qingqing Zheng, Amy Zhang, Scott Niekum:
Dual RL: Unification and New Methods for Reinforcement and Imitation Learning.
- Seyed Amir Hossein Saberi, Amir Najafi, Alireza Heidari, Mohammad Hosein Movasaghinia, Abolfazl S. Motahari, Babak H. Khalaj:
Out-Of-Domain Unlabeled Data Improves Generalization.
- Yangsibo Huang, Samyak Gupta, Mengzhou Xia, Kai Li, Danqi Chen:
Catastrophic Jailbreak of Open-source LLMs via Exploiting Generation.
- Jingyu Chen, Runlin Lei, Zhewei Wei:
PolyGCL: GRAPH CONTRASTIVE LEARNING via Learnable Spectral Polynomial Filters.
- Shanqi Liu, Dong Xing, Pengjie Gu, Xinrun Wang, Bo An, Yong Liu:
Solving Homogeneous and Heterogeneous Cooperative Tasks with Greedy Sequential Execution.
- Chongyi Zheng, Benjamin Eysenbach, Homer Rich Walke, Patrick Yin, Kuan Fang, Ruslan Salakhutdinov, Sergey Levine:
Stabilizing Contrastive RL: Techniques for Robotic Goal Reaching from Offline Data.
- Dingling Yao, Danru Xu, Sébastien Lachapelle, Sara Magliacane, Perouz Taslakian, Georg Martius, Julius von Kügelgen, Francesco Locatello:
Multi-View Causal Representation Learning with Partial Observability.
- Sohan Patnaik, Heril Changwal, Milan Aggarwal, Sumit Bhatia, Yaman Kumar, Balaji Krishnamurthy:
CABINET: Content Relevance-based Noise Reduction for Table Question Answering.
- Josef Dai, Xuehai Pan, Ruiyang Sun, Jiaming Ji, Xinbo Xu, Mickel Liu, Yizhou Wang, Yaodong Yang:
Safe RLHF: Safe Reinforcement Learning from Human Feedback.
- Ruqi Bai, Saurabh Bagchi, David I. Inouye:
Benchmarking Algorithms for Federated Domain Generalization.
- Aditya Bhatt, Daniel Palenicek, Boris Belousov, Max Argus, Artemij Amiranashvili, Thomas Brox, Jan Peters:
CrossQ: Batch Normalization in Deep Reinforcement Learning for Greater Sample Efficiency and Simplicity.
- Xuefeng Liu, Takuma Yoneda, Rick Stevens, Matthew R. Walter, Yuxin Chen:
Blending Imitation and Reinforcement Learning for Robust Policy Improvement.
- Zhengyao Jiang, Yingchen Xu, Nolan Wagener, Yicheng Luo, Michael Janner, Edward Grefenstette, Tim Rocktäschel, Yuandong Tian:
H-GAP: Humanoid Control with a Generalist Planner.
- Alaa Saade, Steven Kapturowski, Daniele Calandriello, Charles Blundell, Pablo Sprechmann, Leopoldo Sarra, Oliver Groth, Michal Valko, Bilal Piot:
Unlocking the Power of Representations in Long-term Novelty-based Exploration.
- Hong Wang, Zhongkai Hao, Jie Wang, Zijie Geng, Zhen Wang, Bin Li, Feng Wu:
Accelerating Data Generation for Neural Operators via Krylov Subspace Recycling.
- Yunhe Zhang, Yan Sun, Jinyu Cai, Jicong Fan:
Deep Orthogonal Hypersphere Compression for Anomaly Detection.
- Chenjie Mao, Qiaosheng Zhang, Zhen Wang, Xuelong Li:
On the Role of General Function Approximation in Offline Reinforcement Learning.
- Goro Kobayashi, Tatsuki Kurabayashi, Sho Yokoi, Kentaro Inui:
Analyzing Feed-Forward Blocks in Transformers through the Lens of Attention Maps.

-     Pratik Patil, Daniel LeJeune:
Asymptotically Free Sketched Ridge Ensembles: Risks, Cross-Validation, and Tuning.
-     Shuai Fu, Xiequn Wang, Qiushi Huang, Yu Zhang:
Nemesis: Normalizing the Soft-prompt Vectors of Vision-Language Models.
-    Xuming Hu, Junzhe Chen, Xiaochuan Li, Yufei Guo, Lijie Wen, Philip S. Yu, Zhijiang Guo:
Towards Understanding Factual Knowledge of Large Language Models.
-    Chunsan Hong, Byunghee Cha, Tae-Hyun Oh:
CAS: A Probability-Based Approach for Universal Condition Alignment Score.
-    Hu Xu, Saining Xie, Xiaoqing Ellen Tan, Po-Yao Huang, Russell Howes, Vasu Sharma, Shang-Wen Li, Gargi Ghosh, Luke Zettlemoyer, Christoph Feichtenhofer:
Demystifying CLIP Data.
-    Huafeng Qin, Xin Jin, Yun Jiang, Mounîm A. El-Yacoubi, Xinbo Gao:
Adversarial AutoMixup.
-    Junmo Cho, Jaesik Yoon, Sungjin Ahn:
Spatially-Aware Transformers for Embodied Agents.
-    Yanwei Wang, Tsun-Hsuan Wang, Jiayuan Mao, Michael Hagenow, Julie Shah:
Grounding Language Plans in Demonstrations Through Counterfactual Perturbations.
-    Xiangyu Liu, Chenghao Deng, Yanchao Sun, Yongyuan Liang, Furong Huang:
Beyond Worst-case Attacks: Robust RL with Adaptive Defense via Non-dominated Policies.
-    Pingzhi Li, Zhenyu Zhang, Prateek Yadav, Yi-Lin Sung, Yu Cheng, Mohit Bansal, Tianlong Chen:
Merge, Then Compress: Demystify Efficient SMoE with Hints from Its Routing Policy.
-    Lin Chen, Michal Lukasik, Wittawat Jitkrittum, Chong You, Sanjiv Kumar:
On Bias-Variance Alignment in Deep Models.
-    Yang Liu, Jiashun Cheng, Haihong Zhao, Tingyang Xu, Peilin Zhao, Fugee Tsung, Jia Li, Yu Rong:
SEGNO: Generalizing Equivariant Graph Neural Networks with Physical Inductive Biases.
-    Runtian Zhai, Rattana Pukdee, Roger Jin, Maria-Florina Balcan, Pradeep Kumar Ravikumar:
Spectrally Transformed Kernel Regression.
-    Xin Zheng, Dongjin Song, Qingsong Wen, Bo Du, Shirui Pan:
Online GNN Evaluation Under Test-time Graph Distribution Shifts.
-    Hao Chen, Jindong Wang, Ankit Shah, Ran Tao, Hongxin Wei, Xing Xie, Masashi Sugiyama, Bhiksha Raj:
Understanding and Mitigating the Label Noise in Pre-training on Downstream Tasks.
-    Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, Yuntian Deng:
WildChat: 1M ChatGPT Interaction Logs in the Wild.
-    Tsung-Wei Ke, Sangwoo Mo, Stella X. Yu:
Learning Hierarchical Image Segmentation For Recognition and By Recognition.
-    Erfan Shayegani, Yue Dong, Nael B. Abu-Ghazaleh:
Jailbreak in pieces: Compositional Adversarial Attacks on Multi-Modal Language Models.
-    Kyungmin Lee, Kihyuk Sohn, Jinwoo Shin:
DreamFlow: High-quality text-to-3D generation by Approximating Probability Flow.
-    Brian DuSell, David Chiang:
Stack Attention: Improving the Ability of Transformers to Model Hierarchical Patterns.
-    Xuhui Zhou, Hao Zhu, Leena Mathur, Ruohong Zhang, Haofei Yu, Zhengyang Qi, Louis-Philippe Morency, Yonatan Bisk, Daniel Fried, Graham Neubig, Maarten Sap:
SOTOPIA: Interactive Evaluation for Social Intelligence in Language Agents.
-    Edward S. Hu, James Springer, Oleh Rybkin, Dinesh Jayaraman:
Privileged Sensing Scaffolds Reinforcement Learning.
-    Dominik Schmidt , Minqi Jiang:
Learning to Act without Actions.
-    Tianjian Li, Haoran Xu, Philipp Koehn, Daniel Khashabi, Kenton Murray:
Error Norm Truncation: Robust Training in the Presence of Data Noise for Text Generation Models.
-    Matt Barnes, Matthew Abueg, Oliver F. Lange, Matt Deeds, Jason Trader, Denali Molitor, Markus Wulfmeier, Shawn O'Banion:
Massively Scalable Inverse Reinforcement Learning in Google Maps.
-    Yian Wang, Juntian Zheng, Zhehuan Chen, Zhou Xian, Gu Zhang, Chao Liu, Chuang Gan:
Thin-Shell Object Manipulations With Differentiable Physics Simulations.
-    Vaidehi Patil, Peter Hase, Mohit Bansal:
Can Sensitive Information Be Deleted From LLMs? Objectives for Defending Against Extraction Attacks.
-    Harikrishna Narasimhan, Aditya Krishna Menon, Wittawat Jitkrittum, Neha Gupta, Sanjiv Kumar:
Learning to Reject Meets Long-tail Learning.



- Katherine L. Hermann, Hossein Mobahi, Thomas Fel, Michael Curtis Mozer:
On the Foundations of Shortcut Learning.
- Roman Pogodin, Jonathan Cornford, Arna Ghosh, Gauthier Gidel, Guillaume Lajoie, Blake Aaron Richards:
Synaptic Weight Distributions Depend on the Geometry of Plasticity.
- Derek Lim, Haggai Maron, Marc T. Law, Jonathan Lorraine , James Lucas:
Graph Metanetworks for Processing Diverse Neural Architectures.
- Peiran Yu, Junyi Li, Heng Huang:
Dropout Enhanced Bilevel Training.
- Christopher A. Choquette-Choo, Arun Ganesh, Thomas Steinke, Abhradeep Guha Thakurta:
Privacy Amplification for Matrix Mechanisms.
- Thomas Kleine Buening, Aadhirupa Saha, Christos Dimitrakakis, Haifeng Xu:
Bandits Meet Mechanism Design to Combat Clickbait in Online Recommendation.
- Dipendra Misra, Akanksha Saran, Tengyang Xie, Alex Lamb, John Langford:
Towards Principled Representation Learning from Videos for Reinforcement Learning.
- Noga Alon, Dmitrii Avdiukhin, Dor Elboim, Orr Fischer, Grigory Yaroslavtsev:
Optimal Sample Complexity of Contrastive Learning.
- Wenlong Chen, Yegor Klochkov, Yang Liu:
Post-hoc bias scoring is optimal for fair classification.
- Pengfei He, Han Xu, Jie Ren, Yingqian Cui, Shenglai Zeng, Hui Liu, Charu C. Aggarwal, Jiliang Tang:
Sharpness-Aware Data Poisoning Attack.
- Maryam Haghhighat, Peyman Moghadam, Shaheer Mohamed, Piotr Koniusz:
Pre-training with Random Orthogonal Projection Image Modeling.
- Fabricio Arend Torres, Marcello Massimo Negri, Marco Inversi, Jonathan Aellen, Volker Roth:
Lagrangian Flow Networks for Conservation Laws.
- Evan Hernandez, Arnab Sen Sharma, Tal Haklay, Kevin Meng, Martin Wattenberg, Jacob Andreas, Yonatan Belinkov, David Bau:
Linearity of Relation Decoding in Transformer Language Models.
- Lorenzo Loconte, Aleksanteri M. Sladek, Stefan Mengel, Martin Trapp, Arno Solin, Nicolas Gillis, Antonio Vergari:
Subtractive Mixture Models via Squaring: Representation and Learning.
- Jiawei Ge, Shange Tang, Jianqing Fan, Chi Jin:
On the Provable Advantage of Unsupervised Pretraining.
- Albert Bou, Matteo Bettini, Sebastian Dittert, Vikash Kumar, Shagun Sodhani, Xiaomeng Yang, Gianni De Fabritiis, Vincent Moens:
TorchRL: A data-driven decision-making library for PyTorch.
- Rui Yang, Han Zhong, Jiawei Xu, Amy Zhang, Chongjie Zhang, Lei Han, Tong Zhang:
Towards Robust Offline Reinforcement Learning under Diverse Data Corruption.
- James Harrison, John Willes, Jasper Snoek:
Variational Bayesian Last Layers.
- Md Mofijul Islam, Alexi Gladstone, Riashat Islam, Tariq Iqbal:
EQA-MX: Embodied Question Answering using Multimodal Expression.
- Jiawei Zhou, Xiaoguang Li, Lifeng Shang, Xin Jiang, Qun Liu, Lei Chen:
Retrieval-based Disentangled Representation Learning with Natural Language Supervision.
- Montgomery Bohde, Meng Liu, Alexandra Saxton, Shuiwang Ji:
On the Markov Property of Neural Algorithmic Reasoning: Analyses and Methods.
- Tom Sherborne, Naomi Saphra , Pradeep Dasigi, Hao Peng:
TRAM: Bridging Trust Regions and Sharpness Aware Minimization.
- Olga Fourkioti, Mat De Vries, Chris Bakal:
CAMIL: Context-Aware Multiple Instance Learning for Cancer Detection and Subtyping in Whole Slide Images.
- Maximilian Seitzer, Sjoerd van Steenkiste, Thomas Kipf, Klaus Greff, Mehdi S. M. Sajjadi:
DyST: Towards Dynamic Neural Scene Representations on Real-World Videos.
- Jie Hao, Xiaochuan Gong, Mingrui Liu:
Bilevel Optimization under Unbounded Smoothness: A New Algorithm and Convergence Analysis.
- Valentyn Melnychuk, Dennis Frauen, Stefan Feuerriegel:
Bounds on Representation-Induced Confounding Bias for Treatment Effect Estimation.
- Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, Christopher Potts:
DSPy: Compiling Declarative Language Model Calls into State-of-the-Art Pipelines.
- Wenhan Cao, Wei Pan:
Impact of Computation in Integral Reinforcement Learning for Continuous-Time Control.



- Advait Harshal Gadhikar, Rebekka Burkholz:
Masks, Signs, And Learning Rate Rewinding.
- Zhan Zhuang, Yu Zhang, Ying Wei:
Gradual Domain Adaptation via Gradient Flow.
- Jiarong Liu, Yifan Zhong, Siyi Hu, Haobo Fu, Qiang Fu, Xiaojun Chang, Yaodong Yang:
Maximum Entropy Heterogeneous-Agent Reinforcement Learning.
- Junyi An, Chao Qu, Zhipeng Zhou, Fenglei Cao, Yinghui Xu, Yuan Qi, Furao Shen:
Hybrid Directional Graph Neural Network for Molecules.
- Zhengmian Hu, Lichang Chen, Xidong Wu, Yihan Wu, Hongyang Zhang, Heng Huang:
Unbiased Watermark for Large Language Models.
- Neehal Tumma, Mathias Lechner, Noel Loo, Ramin M. Hasani, Daniela Rus:
Leveraging Low-Rank and Sparse Recurrent Connectivity for Robust Closed-Loop Control.
- Size Wu, Wenwei Zhang, Lumin Xu, Sheng Jin, Xiangtai Li, Wentao Liu, Chen Change Loy:
CLIPSelf: Vision Transformer Distills Itself for Open-Vocabulary Dense Prediction.
- Wei-Bang Jiang, Li-Ming Zhao, Bao-Liang Lu:
Large Brain Model for Learning Generic Representations with Tremendous EEG Data in BCI.
- Yehui Tang, Hao Xiong, Nianzu Yang, Tailong Xiao, Junchi Yan:
Towards LLM4QPE: Unsupervised Pretraining of Quantum Property Estimation and A Benchmark.
- Guikun Xu, Yongquan Jiang, PengChuan Lei, Yan Yang, Jim Chen:
GTMGC: Using Graph Transformer to Predict Molecule's Ground-State Conformation.
- Yihang Chen, Fanghui Liu, Yiping Lu, Grigorios Chrysos, Volkan Cevher:
Generalization of Scaled Deep ResNets in the Mean-Field Regime.
- Krzysztof Kacprzyk, Samuel Holt, Jeroen Berrevoets , Zhaozhi Qian, Mihaela van der Schaar:
ODE Discovery for Longitudinal Heterogeneous Treatment Effects Inference.
- Christian Gumbusch, Noor Sajid, Georg Martius, Martin V. Butz:
Learning Hierarchical World Models with Adaptive Temporal Abstractions from Discrete Latent Dynamics.
- Avni Kothari, Bogdan Kulynych, Tsui-Wei Weng, Berk Ustun :
Prediction without Preclusion: Recourse Verification with Reachable Sets.
- Liyuan Mao, Haoran Xu, Weinan Zhang, Xianyuan Zhan:
ODICE: Revealing the Mystery of Distribution Correction Estimation via Orthogonal-gradient Update.
- Ziming Hong, Zhenyi Wang, Li Shen, Yu Yao, Zhuo Huang, Shiming Chen, Chuanwu Yang, Mingming Gong, Tongliang Liu:
Improving Non-Transferable Representation Learning by Harnessing Content and Style.
- Donghao Luo, Xue Wang:
ModernTCN: A Modern Pure Convolution Structure for General Time Series Analysis.
- Yingtian Zou, Kenji Kawaguchi, Yingnan Liu, Jiashuo Liu, Mong-Li Lee, Wynne Hsu:
Towards Robust Out-of-Distribution Generalization Bounds via Sharpness.
- Lirong Wu, Yijun Tian, Yufei Huang, Siyuan Li, Haitao Lin, Nitesh V. Chawla, Stan Z. Li:
MAPE-PPI: Towards Effective and Efficient Protein-Protein Interaction Prediction via Microenvironment-Aware Protein Embedding.
- Xue Jiang, Feng Liu, Zhen Fang, Hong Chen, Tongliang Liu, Feng Zheng, Bo Han:
Negative Label Guided OOD Detection with Pretrained Vision-Language Models.
- Lijia Yu, Xiao-Shan Gao, Lijun Zhang:
Optimal robust Memorization with ReLU Neural Networks.
- Hadi Beik-Mohammadi, Søren Hauberg, Georgios Arvanitidis, Nadia Figueroa, Gerhard Neumann, Leonel Rozo:
Neural Contractive Dynamical Systems.
- Vivien Cabannes, Elvis Dohmatob, Alberto Bietti:
Scaling Laws for Associative Memories.
- Tianbao Xie, Siheng Zhao, Chen Henry Wu, Yitao Liu, Qian Luo, Victor Zhong, Yanchao Yang, Tao Yu:
Text2Reward: Reward Shaping with Language Models for Reinforcement Learning.
- Alexander Theus, Olin Geimer, Friedrich Wicke, Thomas Hofmann, Sotiris Anagnostidis, Sidak Pal Singh:
Towards Meta-Pruning via Optimal Transport.
- Yi Wang, Yinan He, Yizhuo Li, Kunchang Li, Jiashuo Yu, Xin Ma, Xinhao Li, Guo Chen, Xinyuan Chen, Yaohui Wang, Ping Luo, Ziwei Liu, Yali Wang, Limin Wang, Yu Qiao:
InternVid: A Large-scale Video-Text Dataset for Multimodal Understanding and Generation.
- Suhwan Choi, Myeongho Jeon, Yeonjung Hwang, Jeonglyul Oh, Sungjun Lim, Joonseok Lee, Myungjoo Kang:
Dictionary Contrastive Learning for Efficient Local Supervision without Auxiliary Networks.
- Yang Yang, Wenhui Wang, Zhe Chen, Jifeng Dai, Liang Zheng:
Bounding Box Stability against Feature Dropout Reflects Detector Generalization across Environments.
- Ce Ju, Reinmar J. Kobler, Liyao Tang, Cuntai Guan, Motoaki Kawanabe:
Deep Geodesic Canonical Correlation Analysis for Covariance-Based Neuroimaging Data.



Yameng Peng, Andy Song, Haytham M. Fayek, Vic Ciesielski, Xiaojun Chang:
SWAP-NAS: Sample-Wise Activation Patterns for Ultra-fast NAS.



Jiayuan Gu, Sean Kirmani, Paul Wohlhart, Yao Lu, Montserrat Gonzalez Arenas, Kanishka Rao, Wenhao Yu, Chuyuan Fu, Keerthana Gopalakrishnan, Zhuo Xu, Priya Sundaresan, Peng Xu, Hao Su, Karol Hausman, Chelsea Finn, Quan Vuong, Ted Xiao:

RT-Trajectory: Robotic Task Generalization via Hindsight Trajectory Sketches.



Kai Shen, Zeqian Ju, Xu Tan, Eric Liu, Yichong Leng, Lei He, Tao Qin, Sheng Zhao, Jiang Bian:
NaturalSpeech 2: Latent Diffusion Models are Natural and Zero-Shot Speech and Singing Synthesizers.



Manish Prajapat, Mojmir Mutny, Melanie N. Zeilinger, Andreas Krause:
Submodular Reinforcement Learning.



Jiahuan Yan, Bo Zheng, Hongxia Xu, Yiheng Zhu, Danny Z. Chen, Jimeng Sun, Jian Wu, Jintai Chen:
Making Pre-trained Language Models Great on Tabular Prediction.



Bowen Song, Soo Min Kwon, Zecheng Zhang, Xinyu Hu, Qing Qu, Liyue Shen:
Solving Inverse Problems with Latent Diffusion Models via Hard Data Consistency.



Arnav Gudibande, Eric Wallace, Charlie Snell, Xinyang Geng, Hao Liu, Pieter Abbeel, Sergey Levine, Dawn Song:
The False Promise of Imitating Proprietary Language Models.



Thomas T. C. K. Zhang, Leonardo Felipe Toso, James Anderson, Nikolai Matni:
Sample-Efficient Linear Representation Learning from Non-IID Non-Isotropic Data.



Zhipeng Xie, Yahe Li:
Information Retention via Learning Supplemental Features.



Yuan Feng, Yukun Cao, Hairu Wang, Xike Xie, S. Kevin Zhou:
Mayfly: a Neural Data Structure for Graph Stream Summarization.



Hanyu Zhou, Yi Chang, Haoyue Liu, Wending Yan, Yuxing Duan, Zhiwei Shi, Luxin Yan:
Exploring the Common Appearance-Boundary Adaptation for Nighttime Optical Flow.



Yijue Dai, Wenzhong Yan, Feng Yin:
Graphical Multioutput Gaussian Process with Attention.



Seunghan Lee, Taeyoung Park, Kibok Lee:
Soft Contrastive Learning for Time Series.



Somnath Basu Roy Chowdhury, Nicholas Monath, Ahmad Beirami, Rahul Kidambi, Kumar Avinava Dubey, Amr Ahmed, Snigdha Chaturvedi:
Enhancing Group Fairness in Online Settings Using Oblique Decision Forests.



Hongbin Huang, Minghua Chen, Xiao Qiao:
Generative Learning for Financial Time Series with Irregular and Scale-Invariant Patterns.



Yuchuan Tian, Hanting Chen, Xutao Wang, Zheyuan Bai, Qinghua Zhang, Ruifeng Li, Chao Xu, Yunhe Wang:
Multiscale Positive-Unlabeled Detection of AI-Generated Texts.



Shiqiang Wang, Mingyue Ji:
A Lightweight Method for Tackling Unknown Participation Statistics in Federated Averaging.



Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J. Maddison, Tatsunori Hashimoto:
Identifying the Risks of LM Agents with an LM-Emulated Sandbox.



Jiayi Wei, Greg Durrett, Isil Dillig:
Coeditor: Leveraging Repo-level Diffs for Code Auto-editing.



Zhijian Xu, Ailing Zeng, Qiang Xu:
FITS: Modeling Time Series with 10k Parameters.



Longhui Yu, Weisen Jiang, Han Shi, Jincheng Yu, Zhengying Liu, Yu Zhang, James T. Kwok, Zhenguo Li, Adrian Weller, Weiyang Liu:
MetaMath: Bootstrap Your Own Mathematical Questions for Large Language Models.



Xiao Hu, Jianxiong Li, Xianyuan Zhan, Qing-Shan Jia, Ya-Qin Zhang:
Query-Policy Misalignment in Preference-Based Reinforcement Learning.



Joonhun Lee, Myeongho Jeon, Myungjoo Kang, Kyunghyun Park:
Feature-aligned N-BEATS with Sinkhorn divergence.



Taehyeon Kim, Joonkee Kim, Gihun Lee, Se-Young Yun:
Instructive Decoding: Instruction-Tuned Large Language Models are Self-Refiner from Noisy Instructions.



Zixi Wei, Senlin Shu, Yuzhou Cao, Hongxin Wei, Bo An, Lei Feng:
Consistent Multi-Class Classification from Multiple Unlabeled Datasets.



Hongwei Ren, Yue Zhou, Xiaopeng Lin, Yulong Huang, Haotian Fu, Jie Song, Bojun Cheng:
SpikePoint: An Efficient Point-based Spiking Neural Network for Event Cameras Action Recognition.



Shida Wang, Zhong Li, Qianxiao Li:
Inverse Approximation Theory for Nonlinear Recurrent Neural Networks.



Haanvid Lee, Tri Wahyu Guntara, Jongmin Lee, Yung-Kyun Noh, Kee-Eung Kim:
Kernel Metric Learning for In-Sample Off-Policy Evaluation of Deterministic RL Policies.

-     Yuchen Hu, Chen Chen, Chao-Han Huck Yang, Ruizhe Li, Chao Zhang, Pin-Yu Chen, Engsiong Chng: **Large Language Models are Efficient Learners of Noise-Robust Speech Recognition.**
-     Minyoung Park, Mirae Do, YeonJae Shin, Jaeseok Yoo, Jongkwang Hong, Joongrock Kim, Chul Lee: **H2O-SDF: Two-phase Learning for 3D Indoor Reconstruction using Object Surface Fields.**
-     Ke Xue, Ren-Jian Wang, Pengyi Li, Dong Li, Jianye Hao, Chao Qian: **Sample-Efficient Quality-Diversity by Cooperative Coevolution.**
-     Sewon Min, Suchin Gururangan, Eric Wallace, Weijia Shi, Hannaneh Hajishirzi, Noah A. Smith, Luke Zettlemoyer: **SILO Language Models: Isolating Legal Risk In a Nonparametric Datastore.**
-     Hang Xu, Kai Li, Haobo Fu, Qiang Fu, Junliang Xing, Jian Cheng: **Dynamic Discounted Counterfactual Regret Minimization.**
-     Dante Everaert, Christopher Potts: **GIO: Gradient Information Optimization for Training Dataset Selection.**
-     Kazem Meidani, Parshin Shojaee, Chandan K. Reddy, Amir Barati Farimani: **SNIP: Bridging Mathematical Symbolic and Numeric Realms with Unified Pre-training.**
-     Karsten Roth, Lukas Thede, A. Sophia Koepke, Oriol Vinyals, Olivier J. Hénaff, Zeynep Akata: **Fantastic Gains and Where to Find Them: On the Existence and Prospect of General Knowledge Transfer between Any Pretrained Model.**
-     Annan Yu, Arnur Nigmatov, Dmitriy Morozov, Michael W. Mahoney, N. Benjamin Erichson: **Robustifying State-space Models for Long Sequences via Approximate Diagonalization.**
-     Wenhao Zhan, Masatoshi Uehara, Nathan Kallus, Jason D. Lee, Wen Sun: **Provable Offline Preference-Based Reinforcement Learning.**
-     Niloofar Miresghallah, Hyunwoo Kim , Xuhui Zhou, Yulia Tsvetkov, Maarten Sap, Reza Shokri, Yejin Choi: **Can LLMs Keep a Secret? Testing Privacy Implications of Language Models via Contextual Integrity Theory.**
-     Wenhao Zhan, Masatoshi Uehara, Wen Sun, Jason D. Lee: **Provable Reward-Agnostic Preference-Based Reinforcement Learning.**
-     Qiyu Kang, Kai Zhao, Qin Xu Ding, Feng Ji, Xuhao Li, Wenfei Liang, Yang Song, Wee Peng Tay: **Unleashing the Potential of Fractional Calculus in Graph Neural Networks with FROND.**
-     S. Chandra Mouli, Muhammad Ashraful Alam, Bruno Ribeiro: **MetaPhysiCa: Improving OOD Robustness in Physics-informed Machine Learning.**
-     Kimia Hamidieh, Haoran Zhang, Swami Sankaranarayanan, Marzyeh Ghassemi: **Views Can Be Deceiving: Improved SSL Through Feature Space Augmentation.**
-     Annie S. Chen, Yoonho Lee, Amirth Setlur, Sergey Levine, Chelsea Finn: **Project and Probe: Sample-Efficient Adaptation by Interpolating Orthogonal Features.**
-     Zihan Wang, Arthur Jacot: **Implicit bias of SGD in L2-regularized linear DNNs: One-way jumps from high to low rank.**
-     Ziyu Wang, Lejun Min, Gus Xia: **Whole-Song Hierarchical Generation of Symbolic Music Using Cascaded Diffusion Models.**
-     Jiuding Sun, Chantal Shaib, Byron C. Wallace: **Evaluating the Zero-shot Robustness of Instruction-tuned Language Models.**
-     Michael Kleinman, Alessandro Achille, Stefano Soatto: **Critical Learning Periods Emerge Even in Deep Linear Networks.**
-     Ethan Steinberg, Jason Alan Fries, Yizhe Xu, Nigam Shah: **MOTOR: A Time-to-Event Foundation Model For Structured Medical Records.**
-     Lirui Wang, Yiyang Ling, Zhecheng Yuan, Mohit Shridhar, Chen Bao, Yuzhe Qin, Bailin Wang, Huazhe Xu, Xiaolong Wang: **GenSim: Generating Robotic Simulation Tasks via Large Language Models.**
-     Runtian Zhai, Bingbin Liu, Andrej Risteski, J. Zico Kolter, Pradeep Kumar Ravikumar: **Understanding Augmentation-based Self-Supervised Representation Learning via RKHS Approximation and Regression.**
-     Angelica Chen, Ravid Schwartz-Ziv, Kyunghyun Cho, Matthew L. Leavitt, Naomi Saphra : **Sudden Drops in the Loss: Syntax Acquisition, Phase Transitions, and Simplicity Bias in MLMs.**
-     Avishek Joey Bose, Tara Akhound-Sadegh, Guillaume Huguet, Kilian Fatras, Jarrid Rector-Brooks, Cheng-Hao Liu, Andrei Cristian Nica, Maksym Korablyov, Michael M. Bronstein, Alexander Tong: **SE(3)-Stochastic Flow Matching for Protein Backbone Generation.**
-     Junyuan Hong, Jiachen T. Wang, Chenhui Zhang, Zhangheng Li, Bo Li, Zhangyang Wang: **DP-OPT: Make Large Language Model Your Privacy-Preserving Prompt Engineer.**
-     Marc Rußwurm, Konstantin Klemmer, Esther Rolf, Robin Zbinden, Devis Tuia: **Geographic Location Encoding with Spherical Harmonics and Sinusoidal Representation Networks.**
-     Carl Hvarfner, Frank Hutter, Luigi Nardi: **A General Framework for User-Guided Bayesian Optimization.**



- Yiheng Xu, Hongjin Su, Chen Xing, Boyu Mi, Qian Liu, Weijia Shi, Binyuan Hui, Fan Zhou, Yitao Liu, Tianbao Xie, Zhoujun Cheng, Siheng Zhao, Lingpeng Kong, Bailin Wang, Caiming Xiong, Tao Yu:
Lemur: Harmonizing Natural Language and Code for Language Agents.
- Antoine Gonon, Nicolas Brisebarre, Elisa Riccietti, Rémi Gribonval:
A path-norm toolkit for modern networks: consequences, promises and challenges.
- Jinyi Hu, Yuan Yao, Chongyi Wang, Shan Wang, Yinxu Pan, Qianyu Chen, Tianyu Yu, Hanghao Wu, Yue Zhao, Haoye Zhang, Xu Han, Yankai Lin, Jiao Xue, Dahai Li, Zhiyuan Liu, Maosong Sun:
Large Multilingual Models Pivot Zero-Shot Multimodal Learning across Languages.
- Joan Puigcerver, Carlos Riquelme Ruiz, Basil Mustafa, Neil Houlsby:
From Sparse to Soft Mixtures of Experts.
- Shrinivas Ramasubramanian, Harsh Rangwani, Sho Takemori, Kunal Samanta, Yuhei Umeda, Venkatesh Babu Radhakrishnan:
Selective Mixup Fine-Tuning for Optimizing Non-Decomposable Objectives.
- Pengfei Zheng , Yonggang Zhang, Zhen Fang, Tongliang Liu, Defu Lian, Bo Han:
NoiseDiffusion: Correcting Noise for Image Interpolation with Diffusion Models beyond Spherical Linear Interpolation.
- Dustin Podell, Zion English, Kyle Lacey, Andreas Blattmann, Tim Dockhorn, Jonas Müller, Joe Penna, Robin Rombach:
SDXL: Improving Latent Diffusion Models for High-Resolution Image Synthesis.
- Dan Haramati, Tal Daniel, Aviv Tamar:
Entity-Centric Reinforcement Learning for Object Manipulation from Pixels.
- Wei Yao, Chengming Yu, Shangzhi Zeng, Jin Zhang:
Constrained Bi-Level Optimization: Proximal Lagrangian Value Function Approach and Hessian-free Algorithm.
- Joseph Early, Gavin K. C. Cheung, Kurt Cutajar, Hanting Xie, Jas Kandola, Niall Twomey:
Inherently Interpretable Time Series Classification via Multiple Instance Learning.
- Christos Louizos, Matthias Reisser, Denis Korzhakov:
A Mutual Information Perspective on Federated Contrastive Learning.
- Yan Sun, Jicong Fan:
MMD Graph Kernel: Effective Metric Learning for Graphs via Maximum Mean Discrepancy.
- Margalit Glasgow:
SGD Finds then Tunes Features in Two-Layer Neural Networks with near-Optimal Sample Complexity: A Case Study in the XOR problem.
- Kaijie Zhu, Jiaao Chen, Jindong Wang, Neil Zhenqiang Gong, Diyi Yang, Xing Xie:
DyVal: Dynamic Evaluation of Large Language Models for Reasoning Tasks.
- Tim Franzmeyer, Stephen Marcus McAleer, João F. Henriques, Jakob Nicolaus Foerster, Philip Torr, Adel Bibi, Christian Schröder de Witt:
Illusory Attacks: Information-theoretic detectability matters in adversarial attacks.
- William Wei Wang, Dongqi Han, Xufang Luo, Dongsheng Li:
Addressing Signal Delay in Deep Reinforcement Learning.
- Jiayan Teng, Wendi Zheng, Ming Ding, Wenyi Hong, Jianqiao Wangni, Zhuoyi Yang, Jie Tang:
Relay Diffusion: Unifying diffusion process across resolutions for image synthesis.
- Yingqing He, Shaoshu Yang, Haoxin Chen, Xiaodong Cun, Menghan Xia, Yong Zhang, Xintao Wang, Ran He, Qifeng Chen, Ying Shan:
ScaleCrafter: Tuning-free Higher-Resolution Visual Generation with Diffusion Models.
- Guowei Xu, Ruijie Zheng, Yongyuan Liang, Xiyao Wang, Zhecheng Yuan, Tianying Ji, Yu Luo, Xiaoyu Liu, Jiaxin Yuan, Pu Hua, Shuzhen Li, Yanjie Ze, Hal Daumé III, Furong Huang, Huazhe Xu:
DrM: Mastering Visual Reinforcement Learning through Dormant Ratio Minimization.
- Nuoya Xiong, Lijun Ding, Simon Shaolei Du:
How Over-Parameterization Slows Down Gradient Descent in Matrix Sensing: The Curses of Symmetry and Initialization.
- Yuxiang Tuo, Wangmeng Xiang, Jun-Yan He, Yifeng Geng, Xuansong Xie:
AnyText: Multilingual Visual Text Generation and Editing.
- Yingwei Ma, Yue Liu, Yue Yu, Yuanliang Zhang, Yu Jiang, Changjian Wang, Shanshan Li:
At Which Training Stage Does Code Data Help LLMs Reasoning?
- Joo Chan Lee, Daniel Rho, Seungtae Nam, Jong Hwan Ko, Eunbyung Park:
Coordinate-Aware Modulation for Neural Fields.
- Kaichao You, Guo Qin, Anchang Bao, Meng Cao, Ping Huang, Jiulong Shan, Mingsheng Long:
Efficient ConvBN Blocks for Transfer Learning and Beyond.
- Ashmit Khandelwal, Aditya Agrawal, Aanisha Bhattacharyya, Yaman Kumar, Somesh Singh, Uttaran Bhattacharya, Ishita Dasgupta, Stefano Petrangeli, Rajiv Ratn Shah, Changyou Chen, Balaji Krishnamurthy:
Large Content And Behavior Models To Understand, Simulate, And Optimize Content And Behavior.
- Chaoqi Wang, Yibo Jiang, Chenghao Yang, Han Liu, Yuxin Chen:
Beyond Reverse KL: Generalizing Direct Preference Optimization with Diverse Divergence Constraints.



- Seonghyeon Ye, Doyoung Kim, Sungdong Kim, Hyeonbin Hwang, Seungone Kim, Yongrae Jo, James Thorne, Juho Kim, Minjoon Seo:
FLASK: Fine-grained Language Model Evaluation based on Alignment Skill Sets.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Tianle Li, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zhuohan Li, Zi Lin, Eric P. Xing, Joseph E. Gonzalez, Ion Stoica, Hao Zhang:
LMSYS-Chat-1M: A Large-Scale Real-World LLM Conversation Dataset.
- Yefei He, Jing Liu, Weijia Wu, Hong Zhou, Bohan Zhuang:
EfficientDM: Efficient Quantization-Aware Fine-Tuning of Low-Bit Diffusion Models.
- Qingqing Cao, Sewon Min, Yizhong Wang, Hannaneh Hajishirzi:
BTR: Binary Token Representations for Efficient Retrieval Augmented Language Models.
- Ziqi Pang, Ziyang Xie, Yunze Man, Yu-Xiong Wang:
Frozen Transformers in Language Models Are Effective Visual Encoder Layers.
- Junyan Cheng, Peter Chin:
SocioDojo: Building Lifelong Analytical Agents with Real-world Text and Time Series.
- Alexander Shypula, Aman Madaan, Yimeng Zeng, Uri Alon, Jacob R. Gardner, Yiming Yang, Milad Hashemi, Graham Neubig, Parthasarathy Ranganathan, Osbert Bastani, Amir Yazdanbakhsh:
Learning Performance-Improving Code Edits.
- Khai Nguyen, Nicola Barilett, Nhat Ho:
Quasi-Monte Carlo for 3D Sliced Wasserstein.
- Thien Le, Luana Ruiz, Stefanie Jegelka:
A Poincaré Inequality and Consistency Results for Signal Sampling on Large Graphs.
- Yihan Du, R. Srikant, Wei Chen:
Cascading Reinforcement Learning.
- Benjamin Lyo, Cristina Savin:
Complex priors and flexible inference in recurrent circuits with dendritic nonlinearities.
- Bobak T. Kiani, Thien Le, Hannah Lawrence, Stefanie Jegelka, Melanie Weber:
On the hardness of learning under symmetries.
- Haochen Luo, Jindong Gu, Fengyuan Liu, Philip Torr:
An Image Is Worth 1000 Lies: Transferability of Adversarial Images across Prompts on Vision-Language Models.
- Hao Liu, Jiarui Feng, Lecheng Kong, Ningyue Liang, Dacheng Tao, Yixin Chen, Muhan Zhang:
One For All: Towards Training One Graph Model For All Classification Tasks.
- Yining Jiao, Carlton J. Zdanski, Julia S. Kimbell, Andrew Prince, Cameron Worden, Samuel Kirse, Christopher Rutter, Benjamin Shields, William Dunn, Jisan Mahmud, Marc Niethammer:
NAISR: A 3D Neural Additive Model for Interpretable Shape Representation.
- Depen Morwani, Benjamin L. Edelman, Costin-Andrei Oncescu, Rosie Zhao, Sham M. Kakade:
Feature emergence via margin maximization: case studies in algebraic tasks.
- Quentin Bertrand, Avishek Joey Bose, Alexandre Duplessis, Marco Jiralerspong, Gauthier Gidel:
On the Stability of Iterative Retraining of Generative Models on their own Data.
- Priyank Jaini, Kevin Clark, Robert Geirhos:
Intriguing Properties of Generative Classifiers.
- Matteo Pirotta, Andrea Tirinzoni, Ahmed Touati, Alessandro Lazaric, Yann Ollivier:
Fast Imitation via Behavior Foundation Models.
- Yucheng Yang, Tianyi Zhou, Qiang He, Lei Han, Mykola Pechenizkiy, Meng Fang:
Task Adaptation from Skills: Information Geometry, Disentanglement, and New Objectives for Unsupervised Reinforcement Learning.
- Kun Wang, Hao Wu, Yifan Duan, Guibin Zhang, Kai Wang, Xiaojiang Peng, Yu Zheng, Yuxuan Liang, Yang Wang:
NuwaDynamics: Discovering and Updating in Causal Spatio-Temporal Modeling.
- Ling Pan, Moksh Jain, Kanika Madan, Yoshua Bengio:
Pre-Training and Fine-Tuning Generative Flow Networks.
- Weigao Sun, Zhen Qin, Weixuan Sun, Shidi Li, Dong Li, Xuyang Shen, Yu Qiao, Yiran Zhong:
CO2: Efficient Distributed Training with Full Communication-Computation Overlap.
- Seyedmorteza Sadat, Jakob Buhmann, Derek Bradley, Otmar Hilliges, Romann M. Weber:
CADS: Unleashing the Diversity of Diffusion Models through Condition-Annealed Sampling.
- Wenbo Li, Xin Yu, Kun Zhou, Yibing Song, Zhe Lin:
Image Inpainting via Iteratively Decoupled Probabilistic Modeling.
- Yongchao Du, Min Wang, Wengang Zhou, Shuping Hui, Houqiang Li:
Image2Sentence based Asymmetrical Zero-shot Composed Image Retrieval.
- Neta Shaul, Juan C. Pérez, Ricky T. Q. Chen, Ali K. Thabet, Albert Pumarola, Yaron Lipman:
Bespoke Solvers for Generative Flow Models.

-     Antoine Bambade, Fabian Schramm, Adrien B. Taylor, Justin Carpentier:
Leveraging augmented-Lagrangian techniques for differentiating over infeasible quadratic programs in machine learning.
-     Stéphane d'Ascoli, Sören Becker, Philippe Schwaller, Alexander Mathis, Niki Kilbertus:
ODEFormer: Symbolic Regression of Dynamical Systems with Transformers.
-     Shi Fu, Fengxiang He, Xinmei Tian, Dacheng Tao:
Convergence of Bayesian Bilevel Optimization.
-     Kaizhi Yang, Xiaoshuai Zhang, Zhiao Huang, Xuejin Chen, Zexiang Xu, Hao Su:
MovingParts: Motion-based 3D Part Discovery in Dynamic Radiance Field.
-     Illyes Batatia, Lars L. Schaaf, Gábor Csányi, Christoph Ortner, Felix A. Faber:
Equivariant Matrix Function Neural Networks.
-     Jiatong Shi, Hirofumi Inaguma, Xutai Ma, Illia Kulikov, Anna Y. Sun:
Multi-resolution HuBERT: Multi-resolution Speech Self-Supervised Learning with Masked Unit Prediction.
-     Trung Q. Trinh, Markus Heinonen, Luigi Acerbi, Samuel Kaski:
Input-gradient space particle inference for neural network ensembles.
-     Yilan Zhang, Yingxue Xu, Jianqi Chen, Fengying Xie, Hao Chen:
Prototypical Information Bottlenecking and Disentangling for Multimodal Cancer Survival Prediction.
-     Yinya Huang, Xiaohan Lin, Zhengying Liu, Qingxing Cao, Huajian Xin, Haiming Wang, Zhenguo Li, Linqi Song, Xiaodan Liang:
MUSTARD: Mastering Uniform Synthesis of Theorem and Proof Data.
-     Chenhao Li, Elijah Stanger-Jones, Steve Heim, Sangbae Kim:
FLD: Fourier Latent Dynamics for Structured Motion Representation and Learning.
-     Xiong Xu, Kunzhe Huang, Yiming Li, Zhan Qin, Kui Ren:
Towards Reliable and Efficient Backdoor Trigger Inversion via Decoupling Benign Features.
-     Young-Jae Park, Minseok Seo, Doyi Kim, Hyeri Kim, Sanghoon Choi, Beomkyu Choi, Jeongwon Ryu, Sohee Son, Hae-Gon Jeon, Yeji Choi:
Long-Term Typhoon Trajectory Prediction: A Physics-Conditioned Approach Without Reanalysis Data.
-     Jiawei Liang, Siyuan Liang, Aishan Liu, Xiaojun Jia, Junhao Kuang, Xiaochun Cao:
Poisoned Forgery Face: Towards Backdoor Attacks on Face Forgery Detection.
-     Zeqi Xiao, Tai Wang, Jingbo Wang, Jinkun Cao, Wenwei Zhang, Bo Dai, Dahua Lin, Jiangmiao Pang:
Unified Human-Scene Interaction via Prompted Chain-of-Contacts.
-     Hangting Ye, Wei Fan, Xiaozhuang Song, Shun Zheng, He Zhao, Dandan Guo, Yi Chang:
PTaRL: Prototype-based Tabular Representation Learning via Space Calibration.
-     YongKyung Oh, Dongyoung Lim, Sungil Kim:
Stable Neural Stochastic Differential Equations in Analyzing Irregular Time Series Data.
-     Jingcheng Niu, Andrew Liu, Zining Zhu, Gerald Penn:
What does the Knowledge Neuron Thesis Have to do with Knowledge?
-     Jadi Adams, Shireen Y. Elhabian:
Point2SSM: Learning Morphological Variations of Anatomies from Point Clouds.
-     Huaxiu Yao, Xinyu Yang, Xinyi Pan, Shengchao Liu, Pang Wei Koh, Chelsea Finn:
Improving Domain Generalization with Domain Relations.
-     Wufei Ma, Qihao Liu, Jiahao Wang, Angtian Wang, Xiaoding Yuan, Yi Zhang, Zihao Xiao, Guofeng Zhang, Beijia Lu, Ruxiao Duan, Yongrui Qi, Adam Kortylewski, Yaoyao Liu, Alan L. Yuille:
Generating Images with 3D Annotations Using Diffusion Models.
-     Gérard Ben Arous, Reza Gheissari, Jiaoyang Huang, Aukosh Jagannath:
High-dimensional SGD aligns with emerging outlier eigenspaces.
-     Zishun Yu, Yunzhe Tao, Liyu Chen, Tao Sun, Hongxia Yang:
B-Coder: Value-Based Deep Reinforcement Learning for Program Synthesis.
-     Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, Yu Su:
Adaptive Chameleon or Stubborn Sloth: Revealing the Behavior of Large Language Models in Knowledge Conflicts.
-     Minyoung Kim, Timothy M. Hospedales:
A Hierarchical Bayesian Model for Few-Shot Meta Learning.
-     Andrew Engel, Zhichao Wang, Natalie Frank, Ioana Dumitriu, Sutanay Choudhury, Anand D. Sarwate, Tony Chiang:
Faithful and Efficient Explanations for Neural Networks via Neural Tangent Kernel Surrogate Models.
-     Anastasios Nikolas Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, Tal Schuster:
Conformal Risk Control.
-     Ilia Igashov, Arne Schneuing, Marwin H. S. Segler, Michael M. Bronstein, Bruno E. Correia:
RetroBridge: Modeling Retrosynthesis with Markov Bridges.
-     Chenguo Lin, Yadong Mu:
InstructScene: Instruction-Driven 3D Indoor Scene Synthesis with Semantic Graph Prior.



Sigal Raab, Inbal Leibovitch, Guy Tevet, Moab Arar, Amit Haim Bermano, Daniel Cohen-Or:
Single Motion Diffusion.



Hongpeng Cao, Yanbing Mao, Lui Sha, Marco Caccamo:
Physics-Regulated Deep Reinforcement Learning: Invariant Embeddings.



Han Zhang, Xiaofan Gui, Shun Zheng, Ziheng Lu, Yuqi Li, Jiang Bian:
BatteryML: An Open-source Platform for Machine Learning on Battery Degradation.



Jin Su, Chenchen Han, Yuyang Zhou, Junjie Shan, Xibin Zhou, Fajie Yuan:
SaProt: Protein Language Modeling with Structure-aware Vocabulary.



Junsong Chen, Jincheng Yu, Chongjian Ge, Lewei Yao, Enze Xie, Zhongdao Wang, James T. Kwok, Ping Luo, Huchuan Lu, Zhenguo Li:
PixArt- α : Fast Training of Diffusion Transformer for Photorealistic Text-to-Image Synthesis.



Yang Bai, Xinxing Xu, Yong Liu, Salman Khan, Fahad Shahbaz Khan, Wangmeng Zuo, Rick Siow Mong Goh, Chun-Mei Feng:
Sentence-level Prompts Benefit Composed Image Retrieval.



Tailin Wu, Takashi Maruyama, Long Wei, Tao Zhang, Yilun Du, Gianluca Iaccarino, Jure Leskovec:
Compositional Generative Inverse Design.



Sejun Park, Sanghyuk Chun, Wonyeol Lee:
What does automatic differentiation compute for neural networks?



Wenqi Shao, Mengzhao Chen, Zhao Yang Zhang, Peng Xu, Lirui Zhao, Zhiqian Li, Kaipeng Zhang, Peng Gao, Yu Qiao, Ping Luo:
OmniQuant: Omnidirectionally Calibrated Quantization for Large Language Models.



Haoxuan You, Haotian Zhang, Zhe Gan, Xianzhi Du, Bowen Zhang, Zirui Wang, Liangliang Cao, Shih-Fu Chang, Yinfei Yang:
Ferret: Refer and Ground Anything Anywhere at Any Granularity.



Chongyu Fan, Jiancheng Liu, Yihua Zhang, Eric Wong, Dennis Wei, Sijia Liu:
SalUn: Empowering Machine Unlearning via Gradient-based Weight Saliency in Both Image Classification and Generation.



Weiran Yao, Shelby Heinecke, Juan Carlos Niebles, Zhiwei Liu, Yihao Feng, Le Xue, Rithesh R. N., Zeyuan Chen, Jianguo Zhang, Devansh Arpit, Ran Xu, Phil Mui, Huan Wang, Caiming Xiong, Silvio Savarese:
Retroformer: Retrospective Large Language Agents with Policy Gradient Optimization.



Stylianos Poulakakis-Daktylidis, Hadi Jamali Rad:
BECLR: Batch Enhanced Contrastive Few-Shot Learning.



Josh Alman, Zhao Song:
How to Capture Higher-order Correlations? Generalizing Matrix Softmax Attention to Kronecker Computation.



Runpei Dong, Chunrui Han, Yuang Peng, Zekun Qi, Zheng Ge, Jinrong Yang, Liang Zhao, Jianjian Sun, Hongyu Zhou, Haoran Wei, Xiangwen Kong, Xiangyu Zhang, Kaisheng Ma, Li Yi:
DreamLLM: Synergistic Multimodal Comprehension and Creation.



Po-Chen Ko, Jiayuan Mao, Yilun Du, Shao-Hua Sun, Joshua B. Tenenbaum:
Learning to Act from Actionless Videos through Dense Correspondences.



Jeongyeol Kwon, Dohyun Kwon, Stephen Wright, Robert D. Nowak:
On Penalty Methods for Nonconvex Bilevel Optimization and First-Order Stochastic Approximation.



Elias Frantar, Carlos Riquelme Ruiz, Neil Houlsby, Dan Alistarh, Utku Evci:
Scaling Laws for Sparsely-Connected Foundation Models.



Joe Benton, Valentin De Bortoli, Arnaud Doucet, George Deligiannidis:
Nearly d-Linear Convergence Bounds for Diffusion Models via Stochastic Localization.



Chong Mou, Xintao Wang, Jiechong Song, Ying Shan, Jian Zhang:
DragonDiffusion: Enabling Drag-style Manipulation on Diffusion Models.



Junsheng Zhou, Jinsheng Wang, Baorui Ma, Yu-Shen Liu, Tiejun Huang, Xinlong Wang:
Uni3D: Exploring Unified 3D Representation at Scale.



Siyuan Qi, Shuo Chen, Yixin Li, Xiangyu Kong, Junqi Wang, Bangcheng Yang, Pring Wong, Yifan Zhong, Xiaoyuan Zhang, Zhaowei Zhang, Nian Liu, Yaodong Yang, Song-Chun Zhu:
CivRealm: A Learning and Reasoning Odyssey in Civilization for Decision-Making Agents.



Simon Ging, María Alejandra Bravo, Thomas Brox:
Open-ended VQA benchmarking of Vision-Language models by exploiting Classification datasets and their semantic hierarchy.



Xuelun Shen, Zhipeng Cai, Wei Yin, Matthias Müller, Zijun Li, Kaixuan Wang, Xiaozi Chen, Cheng Wang:
GIM: Learning Generalizable Image Matcher From Internet Videos.



Yuan Liu, Cheng Lin, Zijiao Zeng, Xiaoxiao Long, Lingjie Liu, Taku Komura, Wenping Wang:
SyncDreamer: Generating Multiview-consistent Images from a Single-view Image.



Yufeng Zhang, Hang Yu, Jianguo Li, Weiyao Lin:
Finite-State Autoregressive Entropy Coding for Efficient Learned Lossless Compression.



Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, Maosong Sun:

**ToolLM: Facilitating Large Language Models to Master 16000+ Real-world APIs.**

Yuanqing Huang, Yingui Wang, Le Yang, Lei Wang:
Enhanced Face Recognition using Intra-class Incoherence Constraint.

Jonghyun Lee, Dahuin Jung, Saehyung Lee, Junsung Park, Juhyeon Shin, Uiwon Hwang, Sungroh Yoon:
Entropy is not Enough for Test-Time Adaptation: From the Perspective of Disentangled Factors.

Wenlong Zhang, Xiaohui Li, Xiangyu Chen, Xiaoyun Zhang, Yu Qiao, Xiao-Ming Wu, Chao Dong:
SEAL: A Framework for Systematic Evaluation of Real-World Super-Resolution.

Yaxuan Zhu, Jianwen Xie, Ying Nian Wu, Ruiqi Gao:
Learning Energy-Based Models by Cooperative Diffusion Recovery Likelihood.

Xiang Yue, Xingwei Qu, Ge Zhang, Yao Fu, Wenhao Huang, Huan Sun, Yu Su, Wenhui Chen:
MAmmoTH: Building Math Generalist Models through Hybrid Instruction Tuning.

Shahriar Golchin, Mihai Surdeanu:
Time Travel in LLMs: Tracing Data Contamination in Large Language Models.

Rembert Daems, Manfred Opper, Guillaume Crevecoeur, Tolga Birdal:
Variational Inference for SDEs Driven by Fractional Noise.

Pierre Marion, Yu-Han Wu, Michael Eli Sander, Gérard Biau:
Implicit regularization of deep residual networks towards neural ODEs.

Meng-Chieh Lee, Haiyang Yu, Jian Zhang, Vassilis N. Ioannidis, Xiang Song, Soji Adeshina, Da Zheng, Christos Faloutsos:
NetInfoF Framework: Measuring and Exploiting Network Usable Information.

Yaoming Wang, Jin Li, Xiaopeng Zhang, Bowen Shi, Chenglin Li, Wenrui Dai, Hongkai Xiong, Qi Tian:
BarLeRla: An Efficient Tuning Framework for Referring Image Segmentation.

Minsu Kim, Taeyoung Yun, Emmanuel Bengio, Dinghuai Zhang, Yoshua Bengio, Sungsoo Ahn, Jinkyoo Park:
Local Search GFlowNets.

Shengjie Luo, Tianlang Chen, Aditi S. Krishnapriyan:
Enabling Efficient Equivariant Operations in the Fourier Basis via Gaunt Tensor Products.

Tongda Xu, Ziran Zhu, Dailan He, Yanghao Li, Lina Guo, Yuanyuan Wang, Zhe Wang, Hongwei Qin, Yan Wang, Jingjing Liu, Ya-Qin Zhang:
Idempotence and Perceptual Image Compression.

Chengrui Li, Yule Wang, Weihan Li, Anqi Wu:
Forward χ^2 Divergence Based Variational Importance Sampling.

Nirmit Joshi, Gal Vardi, Nathan Srebro:
Noisy Interpolation Learning with Shallow Univariate ReLU Networks.

Zhiqiu Xu, Yanjie Chen, Kirill Vishniakov, Yida Yin, Zhiqiang Shen, Trevor Darrell, Lingjie Liu, Zhuang Liu:
Initializing Models with Larger Ones.

Yinghao Xu, Hao Tan, Fujun Luan, Sai Bi, Peng Wang, Jiahao Li, Zifan Shi, Kalyan Sunkavalli, Gordon Wetzstein, Zexiang Xu, Kai Zhang:
DMV3D: Denoising Multi-view Diffusion Using 3D Large Reconstruction Model.

Haoheng Lan, Jindong Gu, Philip Torr, Hengshuang Zhao:
Influencer Backdoor Attack on Semantic Segmentation.

Peng Wang, Hao Tan, Sai Bi, Yinghao Xu, Fujun Luan, Kalyan Sunkavalli, Wenping Wang, Zexiang Xu, Kai Zhang:
PF-LRM: Pose-Free Large Reconstruction Model for Joint Pose and Shape Prediction.

Zeyu Tang, Jialu Wang, Yang Liu, Peter Spirtes, Kun Zhang:
Procedural Fairness Through Decoupling Objectionable Data Generating Components.

Xinghang Li, Minghuan Liu, Hanbo Zhang, Cunjun Yu, Jie Xu, Hongtao Wu, Chilam Cheang, Ya Jing, Weinan Zhang, Huaping Liu, Hang Li, Tao Kong:
Vision-Language Foundation Models as Effective Robot Imitators.

Niklas Muennighoff, Qian Liu, Armel Randy Zebaze, Qinkai Zheng, Binyuan Hui, Terry Yue Zhuo, Swayam Singh, Xiangru Tang, Leandro von Werra, Shayne Longpre:
OctoPack: Instruction Tuning Code Large Language Models.

Haoning Wu, Zicheng Zhang, Erli Zhang, Chaofeng Chen, Liang Liao, Annan Wang, Chunyi Li, Wenxiu Sun, Qiong Yan, Guangtao Zhai, Weisi Lin:
Q-Bench: A Benchmark for General-Purpose Foundation Models on Low-level Vision.

Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, Mingsheng Long:
iTransformer: Inverted Transformers Are Effective for Time Series Forecasting.

Weian Mao, Muzhi Zhu, Zheng Sun, Shuaike Shen, Lin Yuanbo Wu, Hao Chen, Chunhua Shen:
De novo Protein Design Using Geometric Vector Field Networks.

Jingyang Qiao, Zhizhong Zhang, Xin Tan, Chengwei Chen, Yanyun Qu, Yong Peng, Yuan Xie:
Prompt Gradient Projection for Continual Learning.

Mengyuan Chen, Junyu Gao, Changsheng Xu:
R-EDL: Relaxing Nonessential Settings of Evidential Deep Learning.



- Yibing Liu, Chris Xing Tian, Haoliang Li, Lei Ma, Shiqi Wang:
Neuron Activation Coverage: Rethinking Out-of-distribution Detection and Generalization.
- Marko Mihajlovic, Sergey Prokudin, Marc Pollefeys, Siyu Tang :
ResFields: Residual Neural Fields for Spatiotemporal Signals.
- Nicklas Hansen, Hao Su, Xiaolong Wang:
TD-MPC2: Scalable, Robust World Models for Continuous Control.
- Xinmeng Huang, Ping Li, Xiaoyun Li:
Stochastic Controlled Averaging for Federated Learning with Communication Compression.
- Yuwei Guo, Ceyuan Yang, Anyi Rao, Zhengyang Liang, Yaohui Wang, Yu Qiao, Maneesh Agrawala, Dahua Lin, Bo Dai:
AnimateDiff: Animate Your Personalized Text-to-Image Diffusion Models without Specific Tuning.
- Tsu-Jui Fu, Wenze Hu, Xianzhi Du, William Yang Wang, Yinfei Yang, Zhe Gan:
Guiding Instruction-based Image Editing via Multimodal Large Language Models.
- Bingchen Zhao, Haoqin Tu, Chen Wei, Jieru Mei, Cihang Xie:
Tuning LayerNorm in Attention: Towards Efficient Multi-Modal LLM Finetuning.
- Zhengyi Luo, Jinkun Cao, Josh Merel, Alexander Winkler, Jing Huang, Kris M. Kitani, Weipeng Xu:
Universal Humanoid Motion Representations for Physics-Based Control.
- Quentin Delfosse, Patrick Schramowski, Martin Mundt, Alejandro Molina, Kristian Kersting:
Adaptive Rational Activations to Boost Deep Reinforcement Learning.
- Diyang Li, Charles Ling, Zhiqiang Xu, Huan Xiong, Bin Gu:
Learning No-Regret Sparse Generalized Linear Models with Varying Observation(s).
- Sam Toyer, Olivia Watkins, Ethan Adrian Mendes, Justin Svegliato, Luke Bailey, Tiffany Wang, Isaac Ong, Karim Elmaaroufi, Pieter Abbeel, Trevor Darrell, Alan Ritter, Stuart Russell:
Tensor Trust: Interpretable Prompt Injection Attacks from an Online Game.
- Tianhong Li, Sangnie Bhardwaj, Yonglong Tian, Han Zhang, Jarred Barber, Dina Katabi, Guillaume Lajoie, Huiwen Chang, Dilip Krishnan:
Leveraging Unpaired Data for Vision-Language Generative Models via Cycle Consistency.
- François Charton:
Learning the greatest common divisor: explaining transformer predictions.
- Steeven Janny, Madiha Nadri, Julie Digne, Christian Wolf:
Space and time continuous physics simulation from partial observations.
- Shaofei Cai, Bowei Zhang, Zihao Wang, Xiaojian Ma, Anji Liu, Yitao Liang:
GROOT: Learning to Follow Instructions by Watching Gameplay Videos.
- Ganlin Yang, Guoqiang Wei, Zhizheng Zhang, Yan Lu, Dong Liu:
Mask-Based Modeling for Neural Radiance Fields.
- Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, Minlie Huang:
Large Language Models Are Not Robust Multiple Choice Selectors.
- Aryaman Reddi, Maximilian Tölle, Jan Peters, Georgia Chalvatzaki, Carlo D'Eramo:
Robust Adversarial Reinforcement Learning via Bounded Rationality Curricula.
- Juncheng Li, Kaihang Pan, Zhiqi Ge, Minghe Gao, Wei Ji, Wenqiao Zhang, Tat-Seng Chua, Siliang Tang, Hanwang Zhang, Yuetong Zhuang:
Fine-tuning Multimodal LLMs to Follow Zero-shot Demonstrative Instructions.
- Sen Cui, Abudukelimu Wuerkaixi, Weishen Pan, Jian Liang, Lei Fang, Changshui Zhang, Fei Wang:
CLAP: Collaborative Adaptation for Patchwork Learning.
- Xinyu Shi, Jianhao Ding, Zecheng Hao, Zhaofei Yu:
Towards Energy Efficient Spiking Neural Networks: An Unstructured Pruning Framework.
- Yaoyu Zhu, Jianhao Ding, Tiejun Huang, Xiaodong Xie, Zhaofei Yu:
Online Stabilization of Spiking Neural Networks.
- Tim Lebailly, Thomas Stegmüller, Behzad Bozorgtabar, Jean-Philippe Thiran, Tinne Tuytelaars:
CrlBo: Self-Supervised Learning via Cross-Image Object-Level Bootstrapping.
- Zhenhui Ye, Tianyun Zhong, Yi Ren, Jiaqi Yang, Weichuang Li, Jiawei Huang, Ziyue Jiang, Jinzheng He, Rongjie Huang, Jinglin Liu, Chen Zhang, Xiang Yin, Zejun Ma, Zhou Zhao:
Real3D-Portrait: One-shot Realistic 3D Talking Portrait Synthesis.

Accept (poster)

- Yury Gorishniy, Ivan Rubachev, Nikolay Kartashev, Daniil Shlenskii, Akim Kotelnikov, Artem Babenko:
TabR: Tabular Deep Learning Meets Nearest Neighbors.
- Sijia Chen, Baochun Li, Di Niu:
Boosting of Thoughts: Trial-and-Error Problem Solving with Large Language Models.
- Zichen Liu, Chao Du, Wee Sun Lee, Min Lin:
Locality Sensitive Sparse Encoding for Learning World Models Online.

-     Binghui Xie, Yatao Bian, Kaiwen Zhou, Yongqiang Chen, Peilin Zhao, Bo Han, Wei Meng, James Cheng: **Enhancing Neural Subset Selection: Integrating Background Information into Set Representations.**
-     Jungin Park, Jiyoung Lee, Kwanghoon Sohn: **Bridging Vision and Language Spaces with Assignment Prediction.**
-     Junlong Li, Shichao Sun, Weizhe Yuan, Run-Ze Fan, Hai Zhao, Pengfei Liu: **Generative Judge for Evaluating Alignment.**
-     Tuo Xu, Lei Zou: **Rethinking and Extending the Probabilistic Inference Capacity of GNNs.**
-     Nishant Jain, Karthikeyan Shanmugam, Pradeep Shenoy: **Learning model uncertainty as variance-minimizing instance weights.**
-     Ravi Francesco Srinivasan, Francesca Mignacco , Martino Sorbaro, Maria Refinetti, Avi Cooper, Gabriel Kreiman, Giorgia Dellaferreira: **Forward Learning with Top-Down Feedback: Empirical and Analytical Characterization.**
-     Yeongwoo Song, Hawoong Jeong: **Towards Cross Domain Generalization of Hamiltonian Representation via Meta Learning.**
-     Wei Liu , Weihao Zeng, Keqing He, Yong Jiang, Junxian He: **What Makes Good Data for Alignment? A Comprehensive Study of Automatic Data Selection in Instruction Tuning.**
-     Kesen Zhao, Liang Zhang: **Causality-Inspired Spatial-Temporal Explanations for Dynamic Graph Neural Networks.**
-     Xiao Zhang, Ji Wu: **Dissecting learning and forgetting in language model finetuning.**
-     Mouxing Yang, Yunfan Li, Changqing Zhang, Peng Hu, Xi Peng: **Test-time Adaptation against Multi-modal Reliability Bias.**
-     Mridul Gupta, Sahil Manchanda, Hariprasad Kodamana, Sayan Ranu: **Mirage: Model-agnostic Graph Distillation for Graph Classification.**
-     Chenchen Gu, Xiang Lisa Li, Percy Liang, Tatsunori Hashimoto: **On the Learnability of Watermarks for Language Models.**
-     Bao Nguyen, Binh Nguyen, Viet Anh Nguyen: **Bellman Optimal Stepsize Straightening of Flow-Matching Models.**
-     Anirudh Buvanesh, Rahul Chand, Jatin Prakash, Bhawna Paliwal, Mudit Dhawan, Neelabh Madan, Deepesh Hada, Vedit Jain, Sonu Mehta, Yashoteja Prabhu, Manish Gupta, Ramachandran Ramjee, Manik Varma: **Enhancing Tail Performance in Extreme Classifiers by Label Variance Reduction.**
-     Aleksandar Makelov, Georg Lange, Atticus Geiger, Neel Nanda: **Is This the Subspace You Are Looking for? An Interpretability Illusion for Subspace Activation Patching.**
-     Daniil Tiapkin, Denis Belomestny, Daniele Calandriello, Eric Moulines, Alexey Naumov, Pierre Perrault, Michal Valko, Pierre Ménard: **Demonstration-Regularized RL.**
-     Yue Deng, Wenzuan Zhang, Sinno Jialin Pan, Lidong Bing: **Multilingual Jailbreak Challenges in Large Language Models.**
-     Juno Kim, Jaehyuk Kwon, Mincheol Cho, Hyunjong Lee, Joong-Ho Won: **\$t^3\$-Variational Autoencoder: Learning Heavy-tailed Data with Student's t and Power Divergence.**
-     Zehao Dong, Muhan Zhang, Philip R. O. Payne, Michael A. Province, Carlos Cruchaga, Tianyu Zhao, Fuhai Li, Yixin Chen: **Rethinking the Power of Graph Canonization in Graph Representation Learning with Stability.**
-     Artem Tsypin, Leonid Ugadiarov, Kuzma Khrabrov, Alexander Telepov, Egor Rumiantsev, Alexey Skrynnik, Aleksandr Panov, Dmitry P. Petrov, Elena Tutubalina, Artur Kadurin: **Gradual Optimization Learning for Conformational Energy Minimization.**
-     Qihang Zhou, Guansong Pang, Yu Tian, Shibo He, Jiming Chen: **AnomalyCLIP: Object-agnostic Prompt Learning for Zero-shot Anomaly Detection.**
-     Linan Yue, Qi Liu, Yichao Du, Li Wang, Weibo Gao, Yanqing An: **Towards Faithful Explanations: Boosting Rationalization with Shortcuts Discovery.**
-     Yoonyoung Cho, Junhyek Han, Yoonjae Cho, Beomjoon Kim: **CORN: Contact-based Object Representation for Nonprehensile Manipulation of General Unseen Objects.**
-     Dongwon Son, Jaehyung Kim, Sanghyeon Son, Beomjoon Kim: **An Intuitive Multi-Frequency Feature Representation for SO(3)-Equivariant Networks.**
-     Jifan Yu, Xiaozhi Wang, Shangqing Tu, Shulin Cao, Daniel Zhang-Li, Xin Lv, Hao Peng, Zijun Yao, Xiaohan Zhang, Hanming Li, Chunyang Li, Zheyuan Zhang, Yushi Bai, Yantao Liu, Amy Xin, Kaifeng Yun, Linlu Gong, Nianyi Lin, Jianhui Chen, Zhili Wu, Yunjia Qi, Weikai Li, Yong Guan, Kaisheng Zeng, Ji Qi, Hailong Jin, Jinxin Liu, Yu Gu, Yuan Yao, Ning Ding, Lei Hou, Zhiyuan Liu, Bin Xu, Jie Tang, Juanzi Li: **KoLA: Carefully Benchmarking World Knowledge of Large Language Models.**
-     Yunchong Song, Siyuan Huang, Xinbing Wang, Chenghu Zhou, Zhouhan Lin: **Graph Parsing Networks.**



- Hyunwook Lee, Sungahn Ko:
TESTAM: A Time-Enhanced Spatio-Temporal Attention Model with Mixture of Experts.
- Florian Frantzen, Michael T. Schaub:
Learning From Simplicial Data Based on Random Walks and 1D Convolutions.
- Lingfeng Liu, Dong Ni, Hangjie Yuan:
LUM-ViT: Learnable Under-sampling Mask Vision Transformer for Bandwidth Limited Optical Signal Acquisition.
- Saeed Saadatnejad, Yang Gao, Kaouther Messaoud, Alexandre Alahi:
Social-Transmotion: Promptable Human Trajectory Prediction.
- Erik Englesson, Hossein Azizpour:
Robust Classification via Regression for Learning with Noisy Labels.
- Christopher Mohri, Daniel Andor, Eunsol Choi, Michael Collins, Anqi Mao, Yutao Zhong:
Learning to Reject with a Fixed Predictor: Application to Decontextualization.
- Urszula Julia Komorowska, Simon V. Mathis, Kieran Didi, Francisco Vargas, Pietro Lio, Mateja Jamnik:
Dynamics-Informed Protein Design with Structure Conditioning.
- Wei Zhuo, Zemin Liu, Bryan Hooi, Bingsheng He, Guang Tan, Rizal Fathony, Jia Chen:
Partitioning Message Passing for Graph Fraud Detection.
- Niels Mündler, Jingxuan He, Slobodan Jenko, Martin T. Vechev:
Self-contradictory Hallucinations of Large Language Models: Evaluation, Detection and Mitigation.
- Ganesh Ramachandra Kini, Vala Vakilian, Tina Behnia, Jaidev Gill, Christos Thrampoulidis:
Symmetric Neural-Collapse Representations with Supervised Contrastive Loss: The Impact of ReLU and Batching.
- Jacob S. Prince, Gabriel Fajardo, George A. Alvarez, Talia Konkle:
Manipulating dropout reveals an optimal balance of efficiency and robustness in biological and machine visual systems.
- Dogyun Park, Siyeon Kim, Sojin Lee, Hyunwoo J. Kim:
DDMI: Domain-agnostic Latent Diffusion Models for Synthesizing High-Quality Implicit Neural Representations.
- Prateek Chanda, Shrey Modi, Ganesh Ramakrishnan:
Bayesian Coreset Optimization for Personalized Federated Learning.
- Tao Ge, Jing Hu, Lei Wang, Xun Wang, Si-Qing Chen, Furu Wei:
In-context Autoencoder for Context Compression in a Large Language Model.
- Fuxiao Liu, Kevin Lin, Linjie Li, Jianfeng Wang, Yaser Yacoob, Lijuan Wang:
Mitigating Hallucination in Large Multi-Modal Models via Robust Instruction Tuning.
- Michael S. Albergo, Nicholas Matthew Boffi, Michael Lindsey, Eric Vanden-Eijnden:
Multimarginal Generative Modeling with Stochastic Interpolants.
- Lifan Zhao, Yanyan Shen:
Rethinking Channel Dependence for Multivariate Time Series Forecasting: Learning from Leading Indicators.
- Parth Sarthi, Salman Abdullah, Aditi Tuli, Shubh Khanna, Anna Goldie, Christopher D. Manning:
RAPTOR: Recursive Abstractive Processing for Tree-Organized Retrieval.
- Zhenan Fan, Huang Fang, Xinglu Wang, Zirui Zhou, Jian Pei, Michael P. Friedlander, Yong Zhang:
Fair and Efficient Contribution Valuation for Vertical Federated Learning.
- Madhur Panwar, Kabir Ahuja, Navin Goyal:
In-Context Learning through the Bayesian Prism.
- Hao Liu, Matei Zaharia, Pieter Abbeel:
RingAttention with Blockwise Transformers for Near-Infinite Context.
- Hao Liu, Carmelo Sferrazza, Pieter Abbeel:
Chain of Hindsight aligns Language Models with Feedback.
- Peter Müller, Lukas Faber, Karolis Martinkus, Roger Wattenhofer:
GraphChef: Decision-Tree Recipes to Explain Graph Neural Networks.
- Riku Togashi, Tatsushi Oka, Naoto Ohsaka, Tetsuro Morimura:
Safe Collaborative Filtering.
- Hanlin Zhu, Baihe Huang, Stuart Russell:
On Representation Complexity of Model-based and Model-free Reinforcement Learning.
- Mohamed Elsayed , A. Rupam Mahmood:
Addressing Loss of Plasticity and Catastrophic Forgetting in Continual Learning.
- Ayan Sengupta, Shantanu Dixit, Md. Shad Akhtar, Tanmoy Chakraborty:
A Good Learner can Teach Better: Teacher-Student Collaborative Knowledge Distillation.
- Aliyah R. Hsu, Yeshwanth Cherapanamjeri, Briton Park, Tristan Naumann, Anobel Y. Odisho, Bin Yu:
Diagnosing Transformers: Illuminating Feature Spaces for Clinical Decision-Making.
- Samyak Jain, Robert Kirk, Ekdeep Singh Lubana, Robert P. Dick, Hidenori Tanaka, Tim Rocktäschel, Edward Grefenstette, David Scott Krueger:
Mechanistically analyzing the effects of fine-tuning on procedurally defined tasks.



- Tianyang Liu, Canwen Xu, Julian J. McAuley:
RepoBench: Benchmarking Repository-Level Code Auto-Completion Systems.
- Ming Zhong, Chenxin An, Weizhu Chen, Jiawei Han, Pengcheng He:
Seeking Neural Nuggets: Knowledge Transfer in Large Language Models from a Parametric Perspective.
- Xiaoxin He, Xavier Bresson, Thomas Laurent, Adam Perold, Yann LeCun, Bryan Hooi:
Harnessing Explanations: LLM-to-LM Interpreter for Enhanced Text-Attributed Graph Representation Learning.
- Jaehyung Kim, Jaehyun Nam, Sangwoo Mo, Jongjin Park, Sang-Woo Lee, Minjoon Seo, Jung-Woo Ha, Jinwoo Shin:
SuRe: Summarizing Retrievals using Answer Candidates for Open-domain QA of LLMs.
- Peng Xu, Wei Ping, Xianchao Wu, Lawrence McAfee, Chen Zhu, Zihan Liu, Sandeep Subramanian, Evelina Bakhturina, Mohammad Shoeybi, Bryan Catanzaro:
Retrieval meets Long Context Large Language Models.
- Yiheng Du, Nithin Chalapathi, Aditi S. Krishnapriyan:
Neural Spectral Methods: Self-supervised learning in the spectral domain.
- Xichen Pan, Li Dong, Shaohan Huang, Zhiliang Peng, Wenhui Chen, Furu Wei:
Kosmos-G: Generating Images in Context with Multimodal Large Language Models.
- Jean-Rémy Conti, Stéphan Cléménçon:
Assessing Uncertainty in Similarity Scoring: Performance & Fairness in Face Recognition.
- Xin Liu, Muhammad Khalifa, Lu Wang:
LitCab: Lightweight Language Model Calibration over Short- and Long-form Responses.
- Xingxuan Li, Ruochen Zhao, Yew Ken Chia, Bosheng Ding, Shafiq Joty, Soujanya Poria, Lidong Bing:
Chain-of-Knowledge: Grounding Large Language Models via Dynamic Knowledge Adapting over Heterogeneous Sources.
- Chengxing Jia, Chenxiao Gao, Hao Yin, Fuxiang Zhang, Xiong-Hui Chen, Tian Xu, Lei Yuan, Zongzhang Zhang, Zhi-Hua Zhou, Yang Yu:
Policy Rehearsal: Training Generalizable Policies for Reinforcement Learning.
- Ru Peng, Heming Zou, Haobo Wang, Yawen Zeng, Zenan Huang, Junbo Zhao:
Energy-based Automated Model Evaluation.
- Jian Kang, Yinglong Xia, Ross Maciejewski, Jiebo Luo, Hanghang Tong:
Deceptive Fairness Attacks on Graphs via Meta Learning.
- Thomas Tian, Chenfeng Xu, Masayoshi Tomizuka, Jitendra Malik, Andrea Bajcsy:
What Matters to You? Towards Visual Representation Alignment for Robot Learning.
- Junyi Li, Feihu Huang, Heng Huang:
FedDA: Faster Adaptive Gradient Methods for Federated Constrained Optimization.
- Chunshu Wu, Ruibing Song , Chuan Liu, Yunan Yang , Ang Li, Michael C. Huang, Tong Geng:
Extending Power of Nature from Binary to Real-Valued Graph Learning in Real World.
- Quoc Phong Nguyen, Bryan Kian Hsiang Low, Patrick Jaillet:
Meta-VBO: Utilizing Prior Tasks in Optimizing Risk Measures with Gaussian Processes.
- Bojan Karlas, David Dao, Matteo Interlandi, Sebastian Schelter, Wentao Wu, Ce Zhang:
Data Debugging with Shapley Importance over Machine Learning Pipelines.
- Qiwei Di, Heyang Zhao, Jiafan He, Quanquan Gu:
Pessimistic Nonlinear Least-Squares Value Iteration for Offline Reinforcement Learning.
- Dingli Yu, Simran Kaur, Arushi Gupta, Jonah Brown-Cohen, Anirudh Goyal, Sanjeev Arora:
SKILL-MIX: a Flexible and Expandable Family of Evaluations for AI Models.
- Xinran Gu, Kaifeng Lyu, Sanjeev Arora, Jingzhao Zhang, Longbo Huang:
A Quadratic Synchronization Rule for Distributed Deep Learning.
- Tong Zhou, Shaolei Ren, Xiaolin Xu:
ArchLock: Locking DNN Transferability at the Architecture Level with a Zero-Cost Binary Predictor.
- Ashutosh Baheti, Ximing Lu, Faeze Brahman, Ronan Le Bras, Maarten Sap, Mark O. Riedl:
Leftover Lunch: Advantage-based Offline Reinforcement Learning for Language Models.
- Fangyuan Xu, Weijia Shi, Eunsol Choi:
RECOMP: Improving Retrieval-Augmented LMs with Context Compression and Selective Augmentation.
- Sachin Kumar, Chan Young Park, Yulia Tsvetkov:
Gen-Z: Generative Zero-Shot Text Classification with Contextualized Label Descriptions.
- Eric J. Bigelow, Ekdeep Singh Lubana, Robert P. Dick, Hidenori Tanaka, Tomer D. Ullman:
In-Context Learning Dynamics with Random Binary Sequences.
- Kaifeng Lyu, Jikai Jin, Zhiyuan Li, Simon Shaolei Du, Jason D. Lee, Wei Hu:
Dichotomy of Early and Late Phase Implicit Biases Can Provably Induce Grokking.
- Haoxuan Li, Chunyuan Zheng, Sihaoy Ding, Peng Wu, Zhi Geng, Fuli Feng, Xiangnan He:
Be Aware of the Neighborhood Effect: Modeling Selection Bias under Interference.



Xinyuan Wang, Chenxi Li, Zhen Wang, Fan Bai, Haotian Luo, Jiayou Zhang, Nebojsa Jojic, Eric P. Xing, Zhiting Hu:
PromptAgent: Strategic Planning with Language Models Enables Expert-level Prompt Optimization.



Patrik Okanovic, Roger Waleffe, Vasilis Mageirakos, Konstantinos E. Nikolakakis, Amin Karbasi, Dionysios S. Kalogerias, Nezihe Merve Gürel, Theodoros Rekatsinas:
Repeated Random Sampling for Minimizing the Time-to-Accuracy of Learning.



Shashank Gupta, Vaishnavi Shrivastava, Ameet Deshpande, Ashwin Kalyan, Peter Clark, Ashish Sabharwal, Tushar Khot:
Bias Runs Deep: Implicit Reasoning Biases in Persona-Assigned LLMs.



Renyu Zhang, Aly A. Khan, Yuxin Chen, Robert L. Grossman:
Enhancing Instance-Level Image Classification with Set-Level Labels.



Quinn LeBlanc Fisher, Haoming Meng, Vardan Papyan:
Pushing Boundaries: Mixup's Influence on Neural Collapse.



Dongjin Kim, Dongwoo Jung, Sungyong Baik, Tae Hyun Kim:
sRGB Real Noise Modeling via Noise-Aware Sampling with Normalizing Flows.



Linlin Yu, Yifei Lou, Feng Chen:
Uncertainty-aware Graph-based Hyperspectral Image Classification.



Denizalp Goktas, David C. Parkes, Ian Gemp, Luke Marrs, Georgios Piliouras, Romuald Elie, Guy Lever, Andrea Tacchetti:
Generative Adversarial Equilibrium Solvers.



Raphael Poulain, Rahmatollah Beheshti:
Graph Transformers on EHRs: Better Representation Improves Downstream Performance.



Zi Wang, Bin Hu, Aaron J. Havens, Alexandre Araujo, Yang Zheng, Yudong Chen, Somesh Jha:
On the Scalability and Memory Efficiency of Semidefinite Programs for Lipschitz Constant Estimation of Neural Networks.



Yuanfeng Ji, Chongjian Ge, Weikai Kong, Enze Xie, Zhengying Liu, Zhenguo Li, Ping Luo:
Large Language Models as Automated Aligners for benchmarking Vision-Language Models.



Jaehyeon Kim, Keon Lee, Seungjun Chung, Jaewoong Cho:
CLaM-TTS: Improving Neural Codec Language Model for Zero-Shot Text-to-Speech.



Zahra Babaiee, Peyman M. Kiasari, Daniela Rus, Radu Grosu:
Unveiling the Unseen: Identifiable Clusters in Trained Depthwise Convolutional Kernels.



Hyunju Kang, Geonhee Han, Hogun Park:
UNR-Explainer: Counterfactual Explanations for Unsupervised Node Representation Learning Models.



Yisheng Xiao, Juntao Li, Zechen Sun, Zechang Li, Qingrong Xia, Xinyu Duan, Zhefeng Wang, Min Zhang:
Are Bert Family Good Instruction Followers? A Study on Their Potential And Limitations.



Kazuki Irie, Anand Gopalakrishnan, Jürgen Schmidhuber:
Exploring the Promise and Limits of Real-Time Recurrent Learning.



Defu Cao, Furong Jia, Sercan Ö. Arik, Tomas Pfister, Yixiang Zheng, Wen Ye, Yan Liu:
TEMPO: Prompt-based Generative Pre-trained Transformer for Time Series Forecasting.



Nithin Chalapathi, Yiheng Du, Aditi S. Krishnapriyan:
Scaling physics-informed hard constraints with mixture-of-experts.



Haoyu Han, Xiaorui Liu, Li Ma, MohamadAli Torkamani, Hui Liu, Jiliang Tang, Makoto Yamada:
Structural Fairness-aware Active Learning for Graph Neural Networks.



Qing Li, Yixin Zhu, Yitao Liang, Ying Nian Wu, Song-Chun Zhu, Siyuan Huang:
Neural-Symbolic Recursive Machine for Systematic Generalization.



Jaemin Cho, Yushi Hu, Jason M. Baldridge, Roopal Garg, Peter Anderson, Ranjay Krishna, Mohit Bansal, Jordi Pont-Tuset, Su Wang:
Davidsonian Scene Graph: Improving Reliability in Fine-grained Evaluation for Text-to-Image Generation.



Zhiyuan Liu, Hong Liu, Denny Zhou, Tengyu Ma:
Chain of Thought Empowers Transformers to Solve Inherently Serial Problems.



Yingyu Lin, Yian Ma, Yu-Xiang Wang, Rachel Redberg, Zhiqi Bu:
Tractable MCMC for Private Learning with Pure and Gaussian Differential Privacy.



Yi Li, Honghao Lin, David P. Woodruff:
Optimal Sketching for Residual Error Estimation for Matrix and Vector Norms.



Xunpeng Huang, Hanze Dong, Yifan Hao, Yian Ma, Tong Zhang:
Reverse Diffusion Monte Carlo.



Charilaos I. Kanatsoulis, Alejandro Ribeiro:
Counting Graph Substructures with Graph Neural Networks.



Catarina G. Belém, Preethi Seshadri, Yasaman Razeghi, Sameer Singh:
Are Models Biased on Text without Gender-related Language?



Faeze Brahman, Chandra Bhagavatula, Valentina Pyatkin, Jena D. Hwang, Xiang Lorraine Li, Hirona Jacqueline Arai, Soumya Sanyal, Keisuke Sakaguchi, Xiang Ren, Yejin Choi:
PlaSma: Procedural Knowledge Models for Language-based Planning and Re-Planning.



Nima Shoghi, Adeesh Kolluru, John R. Kitchin, Zachary W. Ulissi, C. Lawrence Zitnick, Brandon M. Wood:
From Molecules to Materials: Pre-training Large Generalizable Models for Atomic Property Prediction.



Dominique Beaini, Shenyang Huang, Joao Alex Cunha, Zhiyi Li, Gabriela Moisescu-Pareja, Oleksandr Dymov, Samuel Maddrell-Mander, Callum McLean, Frederik Wenkel, Luis Müller, Jama Hussein Mohamud, Ali Parviz, Michael Craig, Michal Koziarski, Jiarui Lu, Zhaocheng Zhu, Cristian Gabellini, Kerstin Klaser, Josef Dean, Cas Wognum, Maciej Sypetkowski, Guillaume Rabusseau, Reihaneh Rabbany, Jian Tang, Christopher Morris, Mirco Ravanelli, Guy Wolf, Prudencio Tossou, Hadrien Mary, Therence Bois, Andrew W. Fitzgibbon, Blazej Banaszewski, Chad Martin, Dominic Masters:
Towards Foundational Models for Molecular Learning on Large-Scale Multi-Task Datasets.



Chenchen Cai, Xu Zhang, Edoardo M. Airoldi:
Independent-Set Design of Experiments for Estimating Treatment and Spillover Effects under Network Interference.



Daniel Y. Fu, Hermann Kumbong, Eric Nguyen, Christopher Ré:
FlashFFTConv: Efficient Convolutions for Long Sequences with Tensor Cores.



Lucas D. Lingle:
Transformer-VQ: Linear-Time Transformers via Vector Quantization.



Michael Zhang, Kush Bhatia, Hermann Kumbong, Christopher Ré:
The Hedgehog & the Porcupine: Expressive Linear Attentions with Softmax Mimicry.



Awni Altabaa, Taylor Whittington Webb, Jonathan D. Cohen, John Lafferty:
Abstractors and relational cross-attention: An inductive bias for explicit relational reasoning in Transformers.



Daouda Sow, Sen Lin, Zhangyang Wang, Yingbin Liang:
Doubly Robust Instance-Reweighted Adversarial Training.



Kevin Black, Michael Janner, Yilun Du, Ilya Kostrikov, Sergey Levine:
Training Diffusion Models with Reinforcement Learning.



Chenyu Zhang, Han Wang, Aritra Mitra, James Anderson:
Finite-Time Analysis of On-Policy Heterogeneous Federated Reinforcement Learning.



Zhong Zheng, Fengyu Gao, Lingzhou Xue, Jing Yang:
Federated Q-Learning: Linear Regret Speedup with Low Communication Cost.



Lingfeng Shen, Sihao Chen, Linfeng Song, Lifeng Jin, Baolin Peng, Haitao Mi, Daniel Khashabi, Dong Yu:
The Trickle-down Impact of Reward Inconsistency on RLHF.



Xu Ma, Xiyang Dai, Jianwei Yang, Bin Xiao, Yinpeng Chen, Yun Fu, Lu Yuan:
Efficient Modulation for Vision Networks.



Tai-Yu Pan, Chenyang Ma, Tianle Chen, Cheng Perng Phoo, Katie Z. Luo, Yurong You, Mark Campbell, Kilian Q. Weinberger, Bharath Hariharan, Wei-Lun Chao:
Pre-training LiDAR-based 3D Object Detectors through Colorization.



Eric Mitchell, Rafael Rafailov, Archit Sharma, Chelsea Finn, Christopher D. Manning:
An Emulator for Fine-tuning Large Language Models using Small Language Models.



Chengyu Dong, Liyuan Liu, Jingbo Shang:
Toward Student-oriented Teacher Network Training for Knowledge Distillation.



Wes Gurnee, Max Tegmark:
Language Models Represent Space and Time.



Ahmed Abdulaal, Adamos Hadjivasiliou, Nina Montaña Brown, Tiantian He, Ayodeji Ijishakin, Ivana Drobnjak, Daniel C. Castro, Daniel C. Alexander:
Causal Modelling Agents: Causal Graph Discovery through Synergising Metadata- and Data-driven Reasoning.



Chengyu Dong, Liyuan Liu, Hao Cheng, Jingbo Shang, Jianfeng Gao, Xiaodong Liu:
Fast-ELECTRA for Efficient Pre-training.



Amin Rakhsha, Mete Kemertas, Mohammad Ghavamzadeh, Amir-massoud Farahmand:
Maximum Entropy Model Correction in Reinforcement Learning.



Mauricio Tec, Ana Trisovic, Michelle Audirac, Sophie Woodward, Jie Kate Hu, Naeem Khoshnevis, Francesca Dominici:
SpaCE: The Spatial Confounding Environment.



Charlotte Nicks, Eric Mitchell, Rafael Rafailov, Archit Sharma, Christopher D. Manning, Chelsea Finn, Stefano Ermon:
Language Model Detectors Are Easily Optimized Against.



Kevin Black, Mitsuhiro Nakamoto, Pranav Atreya, Homer Rich Walke, Chelsea Finn, Aviral Kumar, Sergey Levine:
Zero-Shot Robotic Manipulation with Pre-Trained Image-Editing Diffusion Models.



Chang Chen, Fei Deng, Kenji Kawaguchi, Caglar Gulcehre, Sungjin Ahn:
Simple Hierarchical Planning with Diffusion.



Jihao Andreas Lin, Shreyas Padhy, Javier Antorán, Austin Tripp, Alexander Terenin, Csaba Szepesvári, José Miguel Hernández-Lobato, David Janz:
Stochastic Gradient Descent for Gaussian Processes Done Right.



Jingyun Xiao, Ran Liu, Eva L. Dyer:
GAFormer: Enhancing Timeseries Transformers Through Group-Aware Embeddings.



- Christina Baek, J. Zico Kolter, Aditi Raghunathan:
Why is SAM Robust to Label Noise?
- Zijian Liu, Zhengyuan Zhou:
Revisiting the Last-Iterate Convergence of Stochastic Gradient Methods.
- Eric Qu, Yansen Wang, Xufang Luo, Wenqiang He, Kan Ren, Dongsheng Li:
CNN Kernels Can Be the Best Shapelets.
- Katherine Tian, Eric Mitchell, Huaxiu Yao, Christopher D. Manning, Chelsea Finn:
Fine-Tuning Language Models for Factuality.
- Runyu Zhang, Yang Hu, Na Li:
Soft Robust MDPs and Risk-Sensitive MDPs: Equivalence, Policy Gradient, and Sample Complexity.
- Greg Yang, Dingli Yu, Chen Zhu, Soufiane Hayou:
Tensor Programs VI: Feature Learning in Infinite Depth Neural Networks.
- Ganghua Wang, Xun Xian, Ashish Kundu, Jayanth Srinivasa, Xuan Bi, Mingyi Hong, Jie Ding:
Demystifying Poisoning Backdoor Attacks from a Statistical Perspective.
- Guanting Chen, Xiaocheng Li, Chunlin Sun, Hanzhao Wang:
Learning to Make Adherence-aware Advice.
- Anson Bastos, Kuldeep Singh, Abhishek Nadgeri, Manish Singh, Toyotaro Suzumura:
Beyond Spatio-Temporal Representations: Evolving Fourier Transform for Temporal Graphs.
- Aniket Rajiv Didolkar, Anirudh Goyal, Yoshua Bengio:
Cycle Consistency Driven Object Discovery.
- Nanda H. Krishna , Colin Bredenberg, Daniel Levenstein, Blake Aaron Richards, Guillaume Lajoie:
Sufficient conditions for offline reactivation in recurrent neural networks.
- Namyoung Park, Xing Wang, Antoine Simoulin, Shuai Yang, Grey Yang, Ryan A. Rossi, Puja Trivedi, Nesreen K. Ahmed:
Forward Learning of Graph Neural Networks.
- Yash J. Patel, Akash Kundu, Mateusz Ostaszewski, Xavier Bonet-Monroig, Vedran Dunjko, Onur Danaci:
Curriculum reinforcement learning for quantum architecture search under hardware errors.
- Prasanna Mayilvahanan, Thaddäus Wiedemer, Evgenia Rusak, Matthias Bethge, Wieland Brendel:
Does CLIP's generalization performance mainly stem from high train-test similarity?
- Mohammad Pedramfar, Yididiya Y. Nadew, Christopher John Quinn, Vaneet Aggarwal:
Unified Projection-Free Algorithms for Adversarial DR-Submodular Optimization.
- Biao Zhang, Zhongtao Liu, Colin Cherry, Orhan Firat:
When Scaling Meets LLM Finetuning: The Effect of Data, Model and Finetuning Method.
- Anton Bushuiev, Roman Bushuiev, Petr Kouba, Anatolii Filkin, Marketa Gabrielova, Michal Gabriel, Jirí Sedlár, Tomás Pluskal, Jirí Damborský, Stanislav Mazurenko, Josef Sivic:
Learning to design protein-protein interactions with enhanced generalization.
- Samuel Holt, Max Ruiz Luyten, Mihaela van der Schaar:
L2MAC: Large Language Model Automatic Computer for Extensive Code Generation.
- Zhen Xiang, Fengqing Jiang, Zidi Xiong, Bhaskar Ramasubramanian, Radha Poovendran, Bo Li:
BadChain: Backdoor Chain-of-Thought Prompting for Large Language Models.
- Wenxi Wang, Yang Hu, Mohit Tiwari, Sarfraz Khurshid, Kenneth L. McMillan, Risto Miikkulainen:
NeuroBack: Improving CDCL SAT Solving using Graph Neural Networks.
- Siyan Zhao, John Dang, Aditya Grover:
Group Preference Optimization: Few-Shot Alignment of Large Language Models.
- Yizhi Li, Ruibin Yuan, Ge Zhang, Yinghao Ma, Xingran Chen, Hanzi Yin, Chenghao Xiao, Chenghua Lin, Anton Ragni, Emmanouil Benetos, Norbert Gyenge, Roger B. Dannenberg, Ruibo Liu, Wenhui Chen, Gus Xia, Yemin Shi, Wenhao Huang, Zili Wang, Yike Guo, Jie Fu:
MERT: Acoustic Music Understanding Model with Large-Scale Self-supervised Training.
- Qiwei Di, Tao Jin, Yue Wu, Heyang Zhao, Farzad Farnoud, Quanquan Gu:
Variance-aware Regret Bounds for Stochastic Contextual Dueling Bandits.
- Rares Cristian, Georgia Perakis:
A Discretization Framework for Robust Contextual Stochastic Optimization.
- Xuheng Li, Yihe Deng, Jingfeng Wu, Dongruo Zhou, Quanquan Gu:
Risk Bounds of Accelerated SGD for Overparameterized Linear Regression.
- Matteo Alleman, Jack W. Lindsey, Stefano Fusi:
Task structure and nonlinearity jointly determine learned representational geometry.
- Zhangir Azerbayev, Hailey Schoelkopf, Keiran Paster, Marco Dos Santos, Stephen Marcus McAleer, Albert Q. Jiang, Jia Deng, Stella Biderman, Sean Welleck:
Llemma: An Open Language Model for Mathematics.
- Kevin Clark, Paul Vicol, Kevin Swersky, David J. Fleet:
Directly Fine-Tuning Diffusion Models on Differentiable Rewards.



- Jiawei Ge, Shange Tang, Jianqing Fan, Cong Ma, Chi Jin:
Maximum Likelihood Estimation is All You Need for Well-Specified Covariate Shift.
- Changwoo Lee, Hun-Seok Kim:
Differentiable Learning of Generalized Structured Matrices for Efficient Deep Neural Networks.
- Huan He, William Hao, Yuanzhe Xi, Yong Chen, Bradley A. Malin, Joyce C. Ho:
A Flexible Generative Model for Heterogeneous Tabular EHR with Missing Modality.
- Karim Hamade, Reid McIlroy-Young, Siddhartha Sen, Jon M. Kleinberg, Ashton Anderson:
Designing Skill-Compatible AI: Methodologies and Frameworks in Chess.
- David Valensi, Esther Derman, Shie Mannor, Gal Dalal:
Tree Search-Based Policy Optimization under Stochastic Execution Delay.
- Christopher Fifty, Dennis Duan, Ronald G. Jenkins, Ehsan Amid, Jure Leskovec, Christopher Ré, Sebastian Thrun:
Context-Aware Meta-Learning.
- Marcus J. Min, Yangruibo Ding, Luca Buratti, Saurabh Pujar, Gail E. Kaiser, Suman Jana, Baishakhi Ray:
Beyond Accuracy: Evaluating Self-Consistency of Code Large Language Models with IdentityChain.
- Aleksei Ustimenko, Aleksandr Beznosikov:
Ito Diffusion Approximation of Universal Ito Chains for Sampling, Optimization and Boosting.
- Athul Paul Jacob, Abhishek Gupta, Jacob Andreas:
Modeling Boundedly Rational Agents with Latent Inference Budgets.
- Vijaya Raghavan T. Ramkumar, Bahram Zonooz, Elahe Arani:
The Effectiveness of Random Forgetting for Robust Generalization.
- Thanh-Tung Le, Khai Nguyen, Shanlin Sun, Kun Han, Nhat Ho, Xiaohui Xie:
Diffeomorphic Mesh Deformation via Efficient Optimal Transport for Cortical Surface Reconstruction.
- Mircea Mironenco, Patrick Forré:
Lie Group Decompositions for Equivariant Neural Networks.
- Minh Hoang, Carl Kingsford:
Efficient Heterogeneous Meta-Learning via Channel Shuffling Modulation.
- Darshil Doshi, Aritra Das, Tianyu He, Andrey Gromov:
To Grok or not to Grok: Disentangling Generalization and Memorization on Corrupted Algorithmic Datasets.
- Dongqi Fu, Zhigang Hua, Yan Xie, Jin Fang, Si Zhang, Kaan Sancak, Hao Wu, Andrey Malevich, Jingrui He, Bo Long:
VCR-Graphomer: A Mini-batch Graph Transformer via Virtual Connections.
- Quoc Phong Nguyen, Wan Theng Ruth Chew, Le Song, Bryan Kian Hsiang Low, Patrick Jaillet:
Optimistic Bayesian Optimization with Unknown Constraints.
- Shoumik Saha, Wenxiao Wang, Yigitcan Kaya, Soheil Feizi, Tudor Dumitras:
DRSM: De-Randomized Smoothing on Malware Classifier Providing Certified Robustness.
- Keller Jordan:
On the Variance of Neural Network Training with respect to Test Sets and Distributions.
- Tennison Liu, Nicolás Astorga, Nabeel Seedat, Mihaela van der Schaar:
Large Language Models to Enhance Bayesian Optimization.
- Sagar Shrestha, Xiao Fu:
Towards Identifiable Unsupervised Domain Translation: A Diversified Distribution Matching Approach.
- Xuan Zhang, Jacob Helwig, Yuchao Lin, Yaochen Xie, Cong Fu, Stephan Wojtowytsh, Shuiwang Ji:
SineNet: Learning Temporal Dynamics in Time-Dependent Partial Differential Equations.
- Xiaoqi Wang, Han-Wei Shen:
GNNBoundary: Towards Explaining Graph Neural Networks through the Lens of Decision Boundaries.
- Biswadeep Chakraborty, Beomseok Kang, Harshit Kumar, Saibal Mukhopadhyay:
Sparse Spiking Neural Network: Exploiting Heterogeneity in Timescales for Pruning Recurrent SNN.
- Yihao Xue, Eric Gan, Jiayi Ni, Siddharth Joshi, Baharan Mirzasoleiman:
Investigating the Benefits of Projection Head for Representation Learning.
- Morteza Mardani, Jiaming Song, Jan Kautz, Arash Vahdat:
A Variational Perspective on Solving Inverse Problems with Diffusion Models.
- Zhijing Jin, Jiarui Liu, Zhiheng Lyu, Spencer Poff, Mrinmaya Sachan, Rada Mihalcea, Mona T. Diab, Bernhard Schölkopf:
Can Large Language Models Infer Causation from Correlation?
- Atsushi Nitanda, Kazusato Oko, Taiji Suzuki, Denny Wu:
Improved statistical and computational complexity of the mean-field Langevin dynamics under structured data.
- Florence Regol, Joud Chataoui, Mark Coates:
Jointly-Learned Exit and Inference for a Dynamic Neural Network.
- Yihao Xue, Siddharth Joshi, Dang Nguyen, Baharan Mirzasoleiman:
Understanding the Robustness of Multi-modal Contrastive Learning to Distribution Shift.



- Chris Cundy, Stefano Ermon:
SequenceMatch: Imitation Learning for Autoregressive Sequence Modelling with Backtracking.
- Linara Adilova, Maksym Andriushchenko, Michael Kamp, Asja Fischer , Martin Jaggi:
Layer-wise linear mode connectivity.
- Yuhao Mao, Mark Niklas Müller, Marc Fischer, Martin T. Vechev:
Understanding Certified Training with Interval Bound Propagation.
- Joey Hong, Anca D. Dragan, Sergey Levine:
Offline RL with Observation Histories: Analyzing and Improving Sample Complexity.
- Giovanni de Felice, Andrea Cini, Daniele Zambon, Vladimir V. Gusev, Cesare Alippi:
Graph-based Virtual Sensing from Sparse and Partial Multivariate Observations.
- Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli, Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum, Jonas Geiping, Tom Goldstein:
NEFTune: Noisy Embeddings Improve Instruction Finetuning.
- Tim De Ryck, Florent Bonnet , Siddhartha Mishra, Emmanuel de Bézenac:
An operator preconditioning perspective on training in physics-informed machine learning.
- Yihan Wang, Si Si, Daliang Li, Michal Lukasik, Felix Yu, Cho-Jui Hsieh, Inderjit S. Dhillon, Sanjiv Kumar:
Two-stage LLM Fine-tuning with Less Specialization and More Generalization.
- Alessandro De Palma, Rudy Bunel, Krishnamurthy (Dj) Dvijotham, M. Pawan Kumar, Robert Stanforth, Alessio Lomuscio:
Expressive Losses for Verified Robustness via Convex Combinations.
- Tianze Luo, Zhanfeng Mo, Sinno Jialin Pan:
Learning Adaptive Multiresolution Transforms via Meta-Framelet-based Graph Convolutional Network.
- Jin Peng Zhou, Yuhuai Wu, Qiyang Li, Roger Baker Grosse:
REFACTOR: Learning to Extract Theorems from Proofs.
- Noa Moriel, Matthew Ricci, Mor Nitzan:
Let's do the time-warp-attend: Learning topological invariants of dynamical systems.
- Xinyu Zhao, Xuxi Chen, Yu Cheng, Tianlong Chen:
Sparse MoE with Language Guided Routing for Multilingual Machine Translation.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, Luke Zettlemoyer:
Detecting Pretraining Data from Large Language Models.
- Jin Peng Zhou, Charles Staats, Wenda Li, Christian Szegedy, Kilian Q. Weinberger, Yuhuai Wu:
Don't Trust: Verify - Grounding LLM Quantitative Reasoning with Autoformalization.
- Chawin Sitawarin, Jaewon Chang, David Huang, Wesson Altoyan, David A. Wagner:
PubDef: Defending Against Transfer Attacks From Public Models.
- Jonas Belouadi , Anne Lauscher, Steffen Eger:
AutomaTikZ: Text-Guided Synthesis of Scientific Vector Graphics with TikZ.
- Jacob Mitchell Springer, Vaishnavh Nagarajan, Aditi Raghunathan:
Sharpness-Aware Minimization Enhances Feature Quality via Balanced Learning.
- Canyu Chen, Kai Shu:
Can LLM-Generated Misinformation Be Detected?
- Dipanjyoti Paul, Arpita Chowdhury, Xinqi Xiong, Feng-Ju Chang, David Edward Carlyn, Samuel Stevens, Kaiya Provost, Anuj Karpatne, Bryan Carstens, Daniel I. Rubenstein, Charles V. Stewart, Tanya Y. Berger-Wolf, Yu Su, Wei-Lun Chao:
A Simple Interpretable Transformer for Fine-Grained Image Classification and Analysis.
- Sheng-Jun Huang, Yi Li, Yiming Sun, Ying-Peng Tang:
One-shot Active Learning Based on Lewis Weight Sampling for Multiple Deep Models.
- Khalid Oublal, Saïd Ladjal, David Benhaim, Emmanuel Le-borgne, François Roueff:
Disentangling Time Series Representations via Contrastive Independence-of-Support on I-Variational Inference.
- Jiashun Jin, Zheng Tracy Ke, Gabriel Moryoussef, Jiajun Tang, Jingming Wang:
Improved algorithm and bounds for successive projection.
- Anand Siththaranjan, Cassidy Laidlaw, Dylan Hadfield-Menell:
Distributional Preference Learning: Understanding and Accounting for Hidden Context in RLHF.
- Dean A. Pospisil, Brett W. Larsen, Sarah E. Harvey, Alex H. Williams:
Estimating Shape Distances on Neural Representations with Limited Samples.
- Ashutosh Singh, Ricardo Augusto Borsoi, Deniz Erdogmus, Tales Imbiriba:
Learning semilinear neural operators: A unified recursive framework for prediction and data assimilation.
- Yecheng Jason Ma, William Liang, Guanzhi Wang, De-An Huang, Osbert Bastani, Dinesh Jayaraman, Yuke Zhu, Linxi Fan, Anima Anandkumar:
Eureka: Human-Level Reward Design via Coding Large Language Models.
- Sina Baharlouei, Shivam Patel, Meisam Razaviyayn:
f-FERM: A Scalable Framework for Robust Fair Empirical Risk Minimization.



Prakhar Kaushik, Aayush Mishra, Adam Kortylewski, Alan L. Yuille:
Source-Free and Image-Only Unsupervised Domain Adaptation for Category Level Object Pose Estimation.

Matthew Finlayson, John Hewitt, Alexander Koller, Swabha Swamyamdipta, Ashish Sabharwal:
Closing the Curious Case of Neural Text Degeneration.

Brian Hu Zhang, Gabriele Farina, Tuomas Sandholm:
Mediator Interpretation and Faster Learning Algorithms for Linear Correlated Equilibria in General Sequential Games.

Siming Yan, Yuqi Yang, Yu-Xiao Guo, Hao Pan, Peng-Shuai Wang, Xin Tong, Yang Liu, Qixing Huang:
3D Feature Prediction for Masked-AutoEncoder-Based Point Cloud Pretraining.

Suhas Kotha, Jacob Mitchell Springer, Aditi Raghunathan:
Understanding Catastrophic Forgetting in Language Models via Implicit Inference.

Beatrice Bevilacqua, Moshe Eliasof, Eli A. Meirom, Bruno Ribeiro, Haggai Maron:
Efficient Subgraph GNNs by Learning Effective Selection Policies.

Victor Geadah, International Brain Laboratory, Jonathan W. Pillow:
Parsing neural dynamics with infinite recurrent switching linear dynamical systems.

Luotian Yuan, Yemin Yu, Ying Wei, Yongwei Wang, Zhihua Wang, Fei Wu:
Active Retrosynthetic Planning Aware of Route Quality.

Jiahai Feng, Jacob Steinhardt:
How do Language Models Bind Entities in Context?

Patricia Pauli, Aaron J. Havens, Alexandre Araujo, Siddharth Garg, Farshad Khorrami, Frank Allgöwer, Bin Hu:
Novel Quadratic Constraints for Extending LipSDP beyond Slope-Restricted Activations.

Luca Eyring, Dominik Klein, Théo Uscidda, Giovanni Palla, Niki Kilbertus, Zeynep Akata, Fabian J. Theis:
Unbalancedness in Neural Monge Maps Improves Unpaired Domain Translation.

T. Mitchell Roddenberry, Vishwanath Saragadam, Maarten V. de Hoop, Richard G. Baraniuk:
Implicit Neural Representations and the Algebra of Complex Wavelets.

Nick Richardson, Deniz Oktay, Yaniv Ovadia, James C. Bowden, Ryan P. Adams:
Fiber Monte Carlo.

Shreyas Havaldar, Navodita Sharma, Shubhi Sareen, Karthikeyan Shanmugam, Aravindan Raghuvir:
Learning from Label Proportions: Bootstrapping Supervised Learners via Belief Propagation.

Orren Karniol-Tambour, David M. Zoltowski, E. Mika Diamanti, Lucas Pinto, Carlos D. Brody, David W. Tank, Jonathan W. Pillow:
Modeling state-dependent communication between brain regions with switching nonlinear dynamical systems.

Yi-Fu Wu, Minseung Lee, Sungjin Ahn:
Neural Language of Thought Models.

Animesh Basak Chowdhury, Marco Romanelli, Benjamin Tan, Ramesh Karri, Siddharth Garg:
Retrieval-Guided Reinforcement Learning for Boolean Circuit Minimization.

Xinlu Zhang, Shiyang Li, Xianjun Yang, Chenxin Tian, Yao Qin, Linda Ruth Petzold:
Enhancing Small Medical Learners with Privacy-preserving Contextual Prompting.

Gabryel Mason-Williams, Fredrik Dahlqvist:
What Makes a Good Prune? Maximal Unstructured Pruning for Maximal Cosine Similarity.

Yiming Gao, Feiyu Liu, Liang Wang, Dehua Zheng, Zhenjie Lian, Weixuan Wang, Wenjin Yang, Siqin Li, Xianliang Wang, Wenhui Chen, Jing Dai, Qiang Fu, Yang Wei, Lanxiao Huang, Wei Liu:
Enhancing Human Experience in Human-Agent Collaboration: A Human-Centered Modeling Approach Based on Positive Human Gain.

Hubert Siuzdak:
Vocos: Closing the gap between time-domain and Fourier-based neural vocoders for high-quality audio synthesis.

Lijia Zhou, James B. Simon, Gal Vardi, Nathan Srebro:
An Agnostic View on the Cost of Overfitting in (Kernel) Ridge Regression.

Maximilian Fleissner, Leena Chennuru Vankadara, Debarghya Ghoshdastidar:
Explaining Kernel Clustering via Decision Trees.

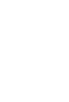
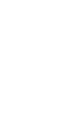
Ruiquan Huang, Yuan Cheng, Jing Yang, Vincent Tan, Yingbin Liang:
Provable Benefits of Multi-task RL under Non-Markovian Decision Making Processes.

Ilyass Hammouamri, Ismail Khalfaoui Hassani, Timothée Masquelier:
Learning Delays in Spiking Neural Networks using Dilated Convolutions with Learnable Spacings.

Lukas Berglund, Meg Tong, Maximilian Kaufmann, Mikita Balesni, Asa Cooper Stickland, Tomasz Korbak, Owain Evans:
The Reversal Curse: LLMs trained on "A is B" fail to learn "B is A".

Xiaogeng Liu, Nan Xu, Muhan Chen, Chaowei Xiao:
AutoDAN: Generating Stealthy Jailbreak Prompts on Aligned Large Language Models.

Xinyan Chen, Yang Li, Runzhong Wang, Junchi Yan:
MixSATGEN: Learning Graph Mixing for SAT Instance Generation.

-      Tommaso Salvatori, Yuhang Song, Yordan Yordanov, Beren Millidge, Lei Sha, Cornelius Emde, Zhenghua Xu, Rafal Bogacz, Thomas Lukasiewicz:
A Stable, Fast, and Fully Automatic Learning Algorithm for Predictive Coding Networks.
-      Hyunho Kim, Jong-Seok Lee:
Scalable Monotonic Neural Networks.
-      Seungcheol Park, Hojun Choi, U Kang:
Accurate Retraining-free Pruning for Pretrained Encoder-based Language Models.
-      Haopeng Sun, Lumin Xu, Sheng Jin, Ping Luo, Chen Qian, Wentao Liu:
PROGRAM: PROotype GRAph Model based Pseudo-Label Learning for Test-Time Adaptation.
-      Xinyu Yang, Weixin Liang, James Zou:
Navigating Dataset Documentations in AI: A Large-Scale Analysis of Dataset Cards on HuggingFace.
-      Weiqiang He, Hendrik Fichtenberger, Pan Peng:
A Differentially Private Clustering Algorithm for Well-Clustered Graphs.
-      Yuto Nishimura, Taiji Suzuki:
Minimax optimality of convolutional neural networks for infinite dimensional input-output problems and separation from kernel methods.
-      Erik Schultheis, Wojciech Kotlowski, Marek Wydmuch, Rohit Babbar, Strom Borman, Krzysztof Dembczynski:
Consistent algorithms for multi-label classification with macro-at-k metrics.
-      Tamir David Hay, Lior Wolf:
Dynamic Layer Tying for Parameter-Efficient Transformers.
-      Dongjun Kim, Chieh-Hsin Lai, Wei-Hsiang Liao, Naoki Murata, Yuhta Takida, Toshimitsu Uesaka, Yutong He, Yuki Mitsufuji, Stefano Ermon:
Consistency Trajectory Models: Learning Probability Flow ODE Trajectory of Diffusion.
-      Nabeel Seedat, Fergus Imrie, Mihaela van der Schaar:
Dissecting Sample Hardness: A Fine-Grained Analysis of Hardness Characterization Methods for Data-Centric AI.
-      Soumyadeep Pal, Yuguang Yao, Ren Wang, Bingquan Shen, Sijia Liu:
Backdoor Secrets Unveiled: Identifying Backdoor Data with Optimized Scaled Prediction Consistency.
-      Peter Richtárik, Elnur Gasanov, Konstantin Burlachenko:
Error Feedback Reloaded: From Quadratic to Arithmetic Mean of Smoothness Constants.
-      Ben Eisner, Yi Yang, Todor Davchev, Mel Vecerík, Jonathan Scholz, David Held:
Deep SE(3)-Equivariant Geometric Reasoning for Precise Placement Tasks.
-      Haoze Wu, Clark W. Barrett, Nina Narodytska:
Lemur: Integrating Large Language Models in Automated Program Verification.
-      Ilker Kesenci, Andrea Pedrotti, Mustafa Dogan, Michele Cafagna, Emre Can Acikgoz, Letitia Parcalabescu, Iacer Calixto, Anette Frank, Albert Gatt, Aykut Erdem, Erkut Erdem:
ViLMA: A Zero-Shot Benchmark for Linguistic and Temporal Grounding in Video-Language Models.
-      Gregory Dexter, Borja Ocejo, S. Sathiya Keerthi, Aman Gupta, Ayan Acharya, Rajiv Khanna:
A Precise Characterization of SGD Stability Using Loss Surface Geometry.
-      Sreyan Ghosh, Ashish Seth, Sonal Kumar, Utkarsh Tyagi, Chandra Kiran Reddy Evuru, Ramaneshwaran S., Sakshi Singh, Oriol Nieto, Ramani Duraiswami, Dinesh Manocha:
CompA: Addressing the Gap in Compositional Reasoning in Audio-Language Models.
-      Woomin Song, Seunghyun Oh, Sangwoo Mo, Jaehyung Kim, Sukmin Yun, Jung-Woo Ha, Jinwoo Shin:
Hierarchical Context Merging: Better Long Context Understanding for Pre-trained LLMs.
-      Ruoyu Wang, Yongqi Yang, Zhihao Qian, Ye Zhu, Yu Wu:
Diffusion in Diffusion: Cyclic One-Way Diffusion for Text-Vision-Conditioned Generation.
-      Yassine Abbahaddou, Sofiane Ennadif, Johannes F. Lutzeyer, Michalis Vazirgiannis, Henrik Boström:
Bounding the Expected Robustness of Graph Neural Networks Subject to Node Feature Attacks.
-      Harshit Sikchi, Rohan Chitnis, Ahmed Touati, Alborz Geramifard, Amy Zhang, Scott Niekum:
Score Models for Offline Goal-Conditioned Reinforcement Learning.
-      Kostadin Garov , Dimitar Iliev Dimitrov , Nikola Jovanovic, Martin T. Vechev:
Hiding in Plain Sight: Disguising Data Stealing Attacks in Federated Learning.
-      Moyang Li, Peng Wang, Lingzhe Zhao, Bangyan Liao, Peidong Liu:
USB-NeRF: Unrolling Shutter Bundle Adjusted Neural Radiance Fields.
-      Linyi Yang, Shuibai Zhang, Zhuohao Yu, Guangsheng Bao, Yidong Wang, Jindong Wang, Ruochen Xu, Wei Ye, Xing Xie, Weizhu Chen, Yue Zhang:
Supervised Knowledge Makes Large Language Models Better In-context Learners.
- Mintong Kang, Nezihe Merve Gürel, Linyi Li, Bo Li:
COLEP: Certifiably Robust Learning-Reasoning Conformal Prediction via Probabilistic Circuits.
- Chongyi Zheng, Ruslan Salakhutdinov, Benjamin Eysenbach:
Contrastive Difference Predictive Coding.



- Geyang Guo, Ranchi Zhao, Tianyi Tang, Xin Zhao, Ji-Rong Wen:
Beyond Imitation: Leveraging Fine-grained Quality Signals for Alignment.
- Brandon Trabucco, Kyle Doherty, Max Gurinas, Ruslan Salakhutdinov:
Effective Data Augmentation With Diffusion Models.
- Krzysztof Kacprzyk, Tennison Liu, Mihaela van der Schaar:
Towards Transparent Time Series Forecasting.
- Jian-Feng Cai, Yu Long, Ruixue Wen, Jiaxi Ying:
A Fast and Provable Algorithm for Sparse Phase Retrieval.
- Jiaxin Yin , Yuanyuan Qiao, Zitang Zhou, Xiangchao Wang, Jie Yang:
MCM: Masked Cell Modeling for Anomaly Detection in Tabular Data.
- Mahdi Karami:
HiGen: Hierarchical Graph Generative Networks.
- Mao Hong, Zhengling Qi, Yanxun Xu:
A Policy Gradient Method for Confounded POMDPs.
- Peizhong Ju, Arnob Ghosh, Ness B. Shroff:
Achieving Fairness in Multi-Agent MDP Using Reinforcement Learning.
- Yiwei Li, Peiwen Yuan, Shaoxiong Feng, Boyuan Pan, Xinglin Wang, Bin Sun, Heda Wang, Kan Li:
Escape Sky-high Cost: Early-stopping Self-Consistency for Multi-step Reasoning.
- Maxwell A. Xu, Alexander Moreno, Hui Wei, Benjamin M. Marlin, James Matthew Rehg:
REBAR: Retrieval-Based Reconstruction for Time-series Contrastive Learning.
- Seyed Saman Saboktsayr, Gonzalo Mateos, Mariano Tepper:
CoLiDE: Concomitant Linear DAG Estimation.
- Arda Sahiner, Tolga Ergen, Batu Ozturkler, John M. Pauly, Morteza Mardani, Mert Pilanci:
Scaling Convex Neural Networks with Burer-Monteiro Factorization.
- Yazheng Yang, Yuqi Wang, Guang Liu, Ledell Wu, Qi Liu:
UniTabE: A Universal Pretraining Protocol for Tabular Foundation Model in Data Science.
- Qianqian Dong, Zhiying Huang, Qi Tian, Chen Xu, Tom Ko, Yunlong Zhao, Siyuan Feng, Tang Li, Kexin Wang, Xuxin Cheng, Fengpeng Yue, Ye Bai, Xi Chen, Lu Lu, Zejun Ma, Yuping Wang, Mingxuan Wang, Yuxuan Wang:
PolyVoice: Language Models for Speech to Speech Translation.
- Dong Huang, Qingwen Bu:
Adversarial Feature Map Pruning for Backdoor.
- Maximilian Baader, Mark Niklas Müller, Yuhao Mao, Martin T. Vechev:
Expressivity of ReLU-Networks under Convex Relaxations.
- Koichi Namekata, Amirmojtaba Sabour, Sanja Fidler, Seung Wook Kim:
EmerDiff: Emerging Pixel-level Semantic Knowledge in Diffusion Models.
- Guanzheng Chen, Xin Li, Zaiqiao Meng, Shangsong Liang, Lidong Bing:
CLEX: Continuous Length Extrapolation for Large Language Models.
- Robert L. Peach, Matteo Vinao-Carl, Nir Grossman, Michael David, Emma Mallas, David J. Sharp, Paresh A. Malhotra, Pierre Vanderghenst, Adam Gosztolai:
Implicit Gaussian process representation of vector fields over arbitrary latent manifolds.
- Pablo Barceló, Alexander Kozachinskiy, Anthony Widjaja Lin, Vladimir V. Podolskii:
Logical Languages Accepted by Transformer Encoders with Hard Attention.
- Yu Tian, Min Shi, Yan Luo, Ava Kouhana, Tobias Elze, Mengyu Wang:
FairSeg: A Large-Scale Medical Image Segmentation Dataset for Fairness Learning Using Segment Anything Model with Fair Error-Bound Scaling.
- Haowen Wang, Tao Sun, Congyun Jin, Yingbo Wang, Yibo Fan, Yunqi Xu, Yuliang Du, Cong Fan:
Customizable Combination of Parameter-Efficient Modules for Multi-Task Learning.
- Yulu Gan, Sungwoo Park, Alexander Schubert, Anthony Philippakis, Ahmed M. Alaa:
InstructCV: Instruction-Tuned Text-to-Image Diffusion Models as Vision Generalists.
- Xiangchi Yuan, Chunhui Zhang, Yijun Tian, Yanfang Ye, Chuxu Zhang:
Mitigating Emergent Robustness Degradation while Scaling Graph Learning.
- T. Anderson Keller, Lyle Muller, Terrence J. Sejnowski, Max Welling:
Traveling Waves Encode The Recent Past and Enhance Sequence Learning.
- Shruthi Gowda, Bahram Zonooz, Elahe Arani:
Conserve-Update-Revise to Cure Generalization and Robustness Trade-off in Adversarial Training.
- Mudit Verma, Katherine Metcalf:
Hindsight PRIORs for Reward Learning from Human Preferences.
- Tuan Le, Julian Cremer, Frank Noé, Djork-Arné Clevert, Kristof T. Schütt:
Navigating the Design Space of Equivariant Diffusion-Based Generative Models for De Novo 3D Molecule Generation.



- Suhyeon Lee, Won Jun Kim, Jinho Chang, Jong Chul Ye:
LLM-CXR: Instruction-Finetuned LLM for CXR Image Understanding and Generation.
- Jianfei Yang, Hanjie Qian, Yuecong Xu, Kai Wang, Lihua Xie:
Can We Evaluate Domain Adaptation Models Without Target-Domain Labels?
- Samuel Pinilla, Jeyan Thiyyagalingam:
Global Optimality for Non-linear Constrained Restoration Problems via Invexity.
- Wenyu Jiang, Hao Cheng, Mingcai Chen, Chongjun Wang, Hongxin Wei:
DOS: Diverse Outlier Sampling for Out-of-Distribution Detection.
- Byeongjun Park, Sangmin Woo, Hyojun Go, Jin-Young Kim, Changick Kim:
Denoising Task Routing for Diffusion Models.
- Thomas Coste, Usman Anwar, Robert Kirk, David Krueger:
Reward Model Ensembles Help Mitigate Overoptimization.
- Han Li, Shaohui Li, Wenrui Dai, Chenglin Li, Junni Zou, Hongkai Xiong:
Frequency-Aware Transformer for Learned Image Compression.
- Xun Jiang, Zhuomin Chai, Yuxiang Zhao, Yibo Lin, Runsheng Wang, Ru Huang:
CircuitNet 2.0: An Advanced Dataset for Promoting Machine Learning Innovations in Realistic Chip Design Environment.
- Heng Dong, Junyu Zhang, Chongjie Zhang:
Leveraging Hyperbolic Embeddings for Coarse-to-Fine Robot Design.
- Oscar Sainz, Iker García-Ferrero, Rodrigo Agerri, Oier Lopez de Lacalle, German Rigau, Eneko Agirre:
GoLLIE: Annotation Guidelines improve Zero-Shot Information-Extraction.
- Nguyen Hung-Quang, Yingjie Lao, Tung Pham, Kok-Seng Wong, Khoa D. Doan:
Understanding the Robustness of Randomized Feature Defense Against Query-Based Adversarial Attacks.
- Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks V. S. Lakshmanan, Ahmed Hassan Awadallah:
Hybrid LLM: Cost-Efficient and Quality-Aware Query Routing.
- Sorawit Saengkyongam, Elan Rosenfeld, Pradeep Kumar Ravikumar, Niklas Pfister, Jonas Peters:
Identifying Representations for Intervention Extrapolation.
- Yinan Zheng, Jianxiong Li, Dongjie Yu, Yujie Yang, Shengbo Eben Li, Xianyuan Zhan, Jingjing Liu:
Safe Offline Reinforcement Learning with Feasibility-Guided Diffusion Model.
- Yifei Wang, Jizhe Zhang, Yisen Wang:
Do Generated Data Always Help Contrastive Learning?
- Jason Chun Lok Li, Steven Tin Sui Luo, Le Xu, Ngai Wong:
ASMR: Activation-Sharing Multi-Resolution Coordinate Networks for Efficient Inference.
- Naoya Hasegawa, Issei Sato:
Exploring Weight Balancing on Long-Tailed Recognition Problem.
- Penghui Qi, Xinyi Wan, Guangxing Huang, Min Lin:
Zero Bubble (Almost) Pipeline Parallelism.
- Florian Grötschla, Joël Mathys, Robert Veres, Roger Wattenhofer:
CoRe-GD: A Hierarchical Framework for Scalable Graph Visualization with GNNs.
- Yeda Song, Dongwook Lee, Gunhee Kim:
Compositional Conservatism: A Transductive Approach in Offline Reinforcement Learning.
- Zhilong Zhang, Yihao Sun, Junyin Ye, Tian-Shuo Liu, Jiaji Zhang, Yang Yu:
Flow to Better: Offline Preference-based Reinforcement Learning via Preferred Trajectory Generation.
- Eshant English, Matthias Kirchler, Christoph Lippert:
Kernelised Normalising Flows.
- Thong Thanh Nguyen, Xiaobao Wu, Xinshuai Dong, Cong-Duy T. Nguyen, See-Kiong Ng, Anh Tuan Luu:
Topic Modeling as Multi-Objective Contrastive Optimization.
- Jangho Park, Gihyun Kwon, Jong Chul Ye:
ED-NeRF: Efficient Text-Guided Editing of 3D Scene With Latent Space NeRF.
- Xueyang Tang, Song Guo, Jie Zhang, Jingcai Guo:
Learning Personalized Causally Invariant Representations for Heterogeneous Federated Clients.
- Zhe Wu, Haofei Lu, Junliang Xing, You Wu, Renye Yan, Yaozhong Gan, Yuanchun Shi:
PAE: Reinforcement Learning from External Knowledge for Efficient Exploration.
- Fei Shen, Hu Ye, Jun Zhang, Cong Wang, Xiao Han, Yang Wei:
Advancing Pose-Guided Image Synthesis with Progressive Conditional Diffusion Models.
- Arvind V. Mahankali, Tatsunori Hashimoto, Tengyu Ma:
One Step of Gradient Descent is Provably the Optimal In-Context Learner with One Layer of Linear Self-Attention.
- Seungone Kim, Jamin Shin, Yejin Choi, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Seongjin Shin, Sungdong Kim, James Thorne, Minjoon Seo:

**Prometheus: Inducing Fine-Grained Evaluation Capability in Language Models.**

- Seong Jin Cho, Gwangsu Kim, Junghyun Lee, Jinwoo Shin, Chang D. Yoo:
Querying Easily Flip-flopped Samples for Deep Active Learning.
- Taewon Park, Inchul Choi, Minho Lee:
Attention-based Iterative Decomposition for Tensor Product Representation.
- Ioannis Mavrothalassitis, Stratis Skoulakis, Leollo Tadesse Dadi, Volkan Cevher:
Efficient Continual Finite-Sum Minimization.
- Suresh Bishnoi, Jayadeva, Sayan Ranu, N. M. Anoop Krishnan:
BroGNet: Momentum-Conserving Graph Neural Stochastic Differential Equation for Learning Brownian Dynamics.
- Grigory Khromov, Sidak Pal Singh:
Some Fundamental Aspects about Lipschitz Continuity of Neural Networks.
- Tim R. Davidson, Veniamin Veselovsky, Michal Kosinski, Robert West:
Evaluating Language Model Agency Through Negotiations.
- Ori Yoran, Tomer Wolfson, Ori Ram, Jonathan Berant:
Making Retrieval-Augmented Language Models Robust to Irrelevant Context.
- Jinxi Xiang, Ricong Huang, Jun Zhang, Guanbin Li, Xiao Han, Yang Wei:
VersVideo: Leveraging Enhanced Temporal Diffusion Models for Versatile Video Generation.
- Shaofei Shen, Chenhao Zhang, Yawen Zhao, Alina Bialkowski, Weitong Chen, Miao Xu:
Label-Agnostic Forgetting: A Supervision-Free Unlearning in Deep Models.
- Ziwei Luo, Fredrik K. Gustafsson, Zheng Zhao, Jens Sjölund, Thomas B. Schön:
Controlling Vision-Language Models for Multi-Task Image Restoration.
- Robin Van De Water, Hendrik Schmidt, Paul W. G. Elbers, Patrick Thoral, Bert Arnrich, Patrick Rockenschaub:
Yet Another ICU Benchmark: A Flexible Multi-Center Framework for Clinical ML.
- Jan Schneider, Pierre Schumacher, Simon Guist, Le Chen, Daniel F. B. Haeufle, Bernhard Schölkopf, Dieter Büchler:
Identifying Policy Gradient Subspaces.
- Yisong Huang, Jin Li, Xinlong Chen, Yang-Geng Fu:
Training Graph Transformers via Curriculum-Enhanced Attention Distillation.
- Feiyang Ye, Yueming Lyu, Xuehao Wang, Yu Zhang, Ivor W. Tsang:
Adaptive Stochastic Gradient Algorithm for Black-box Multi-Objective Learning.
- Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, Jie Tang:
AgentBench: Evaluating LLMs as Agents.
- Sen Pei:
Image Background Serves as Good Proxy for Out-of-distribution Data.
- Zinan Lin, Sivakanth Gopi, Janardhan Kulkarni, Harsha Nori, Sergey Yekhanin:
Differentially Private Synthetic Data via Foundation Model APIs 1: Images.
- Wu Ran, Peirong Ma, Zhiqian He, Hao Ren, Hong Lu:
Harnessing Joint Rain-Detail-aware Representations to Eliminate Intricate Rains.
- Thiziri Nait Saada, Alireza Naderi, Jared Tanner:
Beyond IID weights: sparse and low-rank deep Neural Networks are also Gaussian Processes.
- Philip Quirke, Fazl Barez:
Understanding Addition in Transformers.
- Iosif Sakos, Stefanos Leonardos, Stelios Andrew Stavroulakis, Will Overman, Ioannis Panageas, Georgios Piliouras:
Beating Price of Anarchy and Gradient Descent without Regret in Potential Games.
- Aditya Desai, Anshumali Shrivastava:
In defense of parameter sharing for model-compression.
- Binwu Wang, Pengkun Wang, Wei Xu, Xu Wang, Yudong Zhang, Kun Wang, Yang Wang:
Kill Two Birds with One Stone: Rethinking Data Augmentation for Deep Long-tailed Learning.
- Yang Deng, Wenxuan Zhang, Wai Lam, See-Kiong Ng, Tat-Seng Chua:
Plug-and-Play Policy Planner for Large Language Model Powered Dialogue Agents.
- Aditya Chattopadhyay, Kwan Ho Ryan Chan, René Vidal:
Bootstrapping Variational Information Pursuit with Large Language and Vision Models for Interpretable Image Classification.
- Rohan Deb, Yikun Ban, Shiliang Zuo, Jingrui He, Arindam Banerjee:
Contextual Bandits with Online Neural Regression.
- Zhiyuan Zeng, Jiatong Yu, Tianyu Gao, Yu Meng, Tanya Goyal, Danqi Chen:
Evaluating Large Language Models at Evaluating Instruction Following.
- Ananya Kumar, Ruqi Shen, Sébastien Bubeck, Suriya Gunasekar:
How to Fine-Tune Vision Models with SGD.

-     Weiyu Sun, Xinyu Zhang, Hao Lu, Ying-Cong Chen, Ting Wang, Jinghui Chen, Lu Lin: **Backdoor Contrastive Learning via Bi-level Trigger Optimization.**
-     Souradip Chakraborty, Amrit S. Bedi, Alec Koppel, Huazheng Wang, Dinesh Manocha, Mengdi Wang, Furong Huang: **PARL: A Unified Framework for Policy Alignment in Reinforcement Learning from Human Feedback.**
-     Weidong Huang, Jiaming Ji, Chunhe Xia, Borong Zhang, Yaodong Yang: **SafeDreamer: Safe Reinforcement Learning with World Models.**
-     Min Zhang, Haoxuan Li, Fei Wu, Kun Kuang: **MetaCoCo: A New Few-Shot Classification Benchmark with Spurious Correlation.**
-     Chuheng Zhang, Xiangsen Wang, Wei Jiang, Xianliang Yang, Siwei Wang, Lei Song, Jiang Bian: **Whittle Index with Multiple Actions and State Constraint for Inventory Management.**
-     Liu Yang, Kangwook Lee, Robert D. Nowak, Dimitris Papailiopoulos: **Looped Transformers are Better at Learning Learning Algorithms.**
-     Milan Papez, Martin Rektoris, Václav Smídl, Tomás Pevný: **Sum-Product-Set Networks: Deep Tractable Models for Tree-Structured Graphs.**
-     Jiahao Zhang, Tao Lin, Weiqiang Zheng, Zhe Feng, Yifeng Teng, Xiaotie Deng: **Learning Thresholds with Latent Values and Censored Feedback.**
-     Ivan Lee, Nan Jiang, Taylor Berg-Kirkpatrick: **Is attention required for ICL? Exploring the Relationship Between Model Architecture and In-Context Learning Ability.**
-     Tokio Kajitsuka, Issei Sato: **Are Transformers with One Layer Self-Attention Using Low-Rank Weight Matrices Universal Approximators?**
-     June Yong Yang, Geondo Park, Joowon Kim, Hyeongwon Jang, Eunho Yang: **Language-Interfaced Tabular Oversampling via Progressive Imputation and Self-Authentication.**
-     Bhaskar Mukhoty, Hilal AlQuabeh, Giulia De Masi, Huan Xiong, Bin Gu: **Certified Adversarial Robustness for Rate Encoded Spiking Neural Networks.**
-     Rui Ye, Yixin Du, Zhenyang Ni, Yanfeng Wang, Siheng Chen: **Fake It Till Make It: Federated Learning with Consensus-Oriented Generation.**
-     Xinyuan Chen, Yaohui Wang, Lingjun Zhang, Shaobin Zhuang, Xin Ma, Jiashuo Yu, Yali Wang, Dahua Lin, Yu Qiao, Ziwei Liu: **SEINE: Short-to-Long Video Diffusion Model for Generative Transition and Prediction.**
-     Zichuan Liu, Yingying Zhang, Tianchun Wang, Zefan Wang, Dongsheng Luo, Mengnan Du, Min Wu, Yi Wang, Chunlin Chen, Lunting Fan, Qingsong Wen: **Explaining Time Series via Contrastive and Locally Sparse Perturbations.**
-     Mert Kosan, Samidha Verma, Burouj Armgaan, Khushbu Pahwa, Ambuj K. Singh, Sourav Medya, Sayan Ranu: **GNNX-BENCH: Unravelling the Utility of Perturbation-based GNN Explainers through In-depth Benchmarking.**
-     Avinash Kori, Francesco Locatello, Fabio De Sousa Ribeiro, Francesca Toni, Ben Glocker: **Grounded Object-Centric Learning.**
-     Yinan Huang, William Lu, Joshua Robinson, Yu Yang, Muhan Zhang, Stefanie Jegelka, Pan Li: **On the Stability of Expressive Positional Encodings for Graphs.**
-     Tian Qiu, Wenxiang Xu, Lin Chen, Linyun Zhou, Zunlei Feng, Mingli Song: **Dynamic Neural Response Tuning.**
-     Kyuyoung Kim, Jongheon Jeong, Minyong An, Mohammad Ghavamzadeh, Krishnamurthy Dj Dvijotham, Jinwoo Shin, Kimin Lee: **Confidence-aware Reward Optimization for Fine-tuning Text-to-Image Models.**
-     Huy Nguyen, Pedram Akbarian, Fanqi Yan, Nhat Ho: **Statistical Perspective of Top-K Sparse Softmax Gating Mixture of Experts.**
-     Yan Liu, Yu Liu, Xiaokang Chen, Pin-Yu Chen, Daoguang Zan, Min-Yen Kan, Tsung-Yi Ho: **The Devil is in the Neurons: Interpreting and Mitigating Social Biases in Language Models.**
-     Yu Chen, Yihan Du, Pihe Hu, Siwei Wang, Desheng Wu, Longbo Huang: **Provably Efficient Iterated CVaR Reinforcement Learning with Function Approximation and Human Feedback.**
-     Hyungi Lee, Giung Nam, Edwin Fong, Juho Lee: **Enhancing Transfer Learning with Flexible Nonparametric Posterior Sampling.**
-     Yuchen Zhuang, Xiang Chen, Tong Yu, Saayan Mitra, Victor S. Bursztyn, Ryan A. Rossi, Somdeb Sarkhel, Chao Zhang: **ToolChain*: Efficient Action Space Navigation in Large Language Models with A* Search.**
-     Behzad Shayegh, Yanshuai Cao, Xiaodan Zhu, Jackie C. K. Cheung, Lili Mou: **Ensemble Distillation for Unsupervised Constituency Parsing.**
-     Hattie Zhou, Arwen Bradley, Etai Littwin, Noam Razin, Omid Saremi, Joshua M. Susskind, Samy Bengio, Preetum Nakkiran: **What Algorithms can Transformers Learn? A Study in Length Generalization.**
-     Hien Dang, Tho Tran Huu, Tan Minh Nguyen, Nhat Ho: **Beyond Vanilla Variational Autoencoders: Detecting Posterior Collapse in Conditional and Hierarchical Variational Autoencoders.**



- Xu Han, Caihua Shan, Yifei Shen, Can Xu, Han Yang, Xiang Li, Dongsheng Li:
Training-free Multi-objective Diffusion Model for 3D Molecule Generation.
- Xiongye Xiao, Gengshuo Liu, Gaurav Gupta, Defu Cao, Shixuan Li, Yaxing Li, Tianqing Fang, Mingxi Cheng, Paul Bogdan:
Neuro-Inspired Information-Theoretic Hierarchical Perception for Multimodal Learning.
- Hamidreza Almasi, Harsh Mishra, Balajee Vamanan, Sathya N. Ravi:
Flag Aggregator: Scalable Distributed Training under Failures and Augmented Losses using Convex Optimization.
- Yongchan Kwon, Eric Wu, Kevin Wu, James Zou:
DataInf: Efficiently Estimating Data Influence in LoRA-tuned LLMs and Diffusion Models.
- Rui Qiao, Bryan Kian Hsiang Low:
Understanding Domain Generalization: A Noise Robustness Perspective.
- Yifei Wang, Qi Zhang, Yaoyu Guo, Yisen Wang:
Non-negative Contrastive Learning.
- Sehyun Kwon, Jaeseung Park, Minkyu Kim, Jaewoong Cho, Ernest K. Ryu, Kangwook Lee:
Image Clustering Conditioned on Text Criteria.
- Christopher A. Choquette-Choo, Krishnamurthy Dj Dvijotham, Krishna Pillutla, Arun Ganesh, Thomas Steinke, Abhradeep Guha Thakurta:
Correlated Noise Provably Beats Independent Noise for Differentially Private Learning.
- Zhaowei Zhu, Jialu Wang, Hao Cheng, Yang Liu:
Unmasking and Improving Data Credibility: A Study with Datasets for Training Harmless Language Models.
- Haiquan Qiu, Yongqi Zhang, Yong Li, Quanming Yao:
Understanding Expressivity of GNN in Rule Learning.
- Shunyu Yao, Howard Chen, Austin W. Hanjie, Runzhe Yang, Karthik R. Narasimhan:
COLLIE: Systematic Construction of Constrained Text Generation Tasks.
- Zhenfang Chen, Rui Sun, Wenjun Liu, Yining Hong, Chuang Gan:
GENOME: Generative Neuro-Symbolic Visual Reasoning by Growing and Reusing Modules.
- Noam Razin, Hattie Zhou, Omid Saremi, Vimal Thilak, Arwen Bradley, Preetum Nakkiran, Joshua M. Susskind, Eta Littwin:
Vanishing Gradients in Reinforcement Finetuning of Language Models.
- Changbin Li, Kangshuo Li, Yuzhe Ou, Lance M. Kaplan, Audun Jøsang, Jin-Hee Cho, Dong Hyun Jeong, Feng Chen:
Hyper Evidential Deep Learning to Quantify Composite Classification Uncertainty.
- Jacek Karwowski, Oliver Hayman, Xingjian Bai, Klaus Kiendlhofer, Charlie Griffin, Joar Max Viktor Skalse:
Goodhart's Law in Reinforcement Learning.
- Huayu Chen, Cheng Lu, Zhengyi Wang, Hang Su, Jun Zhu:
Score Regularized Policy Optimization through Diffusion Behavior.
- Zhenfeng He, Yao Shu, Zhongxiang Dai, Bryan Kian Hsiang Low:
Robustifying and Boosting Training-Free Neural Architecture Search.
- Aya Abdelsalam Ismail, Julius Adebayo, Héctor Corrada Bravo, Stephen Ra, Kyunghyun Cho:
Concept Bottleneck Generative Models.
- Renze Lou, Kai Zhang, Jian Xie, Yuxuan Sun, Janice Ahn, Hanzi Xu, Yu Su, Wenpeng Yin:
MUFFIN: Curating Multi-Faceted Instructions for Improving Instruction Following.
- Giung Nam, Byeongho Heo, Juho Lee:
Lipsum-FT: Robust Fine-Tuning of Zero-Shot Models Using Random Text Guidance.
- Ilan Naiman, N. Benjamin Erichson, Pu Ren, Michael W. Mahoney, Omri Azencot:
Generative Modeling of Regular and Irregular Time Series Data via Koopman VAEs.
- Saujas Vaduguru, Daniel Fried, Yewen Pu:
Generating Pragmatic Examples to Train Neural Program Synthesizers.
- Runzhe Wu, Wen Sun:
Making RL with Preference-based Feedback Efficient via Randomization.
- Zhou Lu, Qiuyi Zhang, Xinyi Chen, Fred Zhang, David P. Woodruff, Elad Hazan:
Adaptive Regret for Bandits Made Possible: Two Queries Suffice.
- Ziqi Wang, Chengpeng Hu, Jialin Liu, Xin Yao:
Negatively Correlated Ensemble Reinforcement Learning for Online Diverse Game Level Generation.
- Yining Li, Peizhong Ju, Ness B. Shroff:
Achieving Sample and Computational Efficient Reinforcement Learning by Action Space Reduction via Grouping.
- Zihao Wang, Eshaan Nichani, Jason D. Lee:
Learning Hierarchical Polynomials with Three-Layer Neural Networks.
- Lionel Wong, Jiayuan Mao, Pratyusha Sharma, Zachary S. Siegel, Jiahai Feng, Noa Korneev, Joshua B. Tenenbaum, Jacob Andreas:
Learning Grounded Action Abstractions from Language.
- Youhan Lee, Hasun Yu, Jaemyung Lee, Jaehoon Kim:
Pre-training Sequence, Structure, and Surface Features for Comprehensive Protein Representation Learning.



- Jinxuan Wang, Shiting Xu, Tong Zhang:
A unique M-pattern for micro-expression spotting in long videos.
- Siqi Kou, Lei Gan, Dequan Wang, Chongxuan Li, Zhijie Deng:
BayesDiff: Estimating Pixel-wise Uncertainty in Diffusion via Bayesian Inference.
- Adyasha Maharana, Prateek Yadav, Mohit Bansal:
D2 Pruning: Message Passing for Balancing Diversity & Difficulty in Data Pruning.
- Pengcheng Jiang, Cao Xiao, Adam Cross, Jimeng Sun:
GraphCare: Enhancing Healthcare Predictions with Personalized Knowledge Graphs.
- Tom Yan, Chicheng Zhang:
The Human-AI Substitution game: active learning from a strategic labeler.
- Gabriele Corso, Arthur Deng, Nicholas Polizzi, Regina Barzilay, Tommi S. Jaakkola:
Deep Confident Steps to New Pockets: Strategies for Docking Generalization.
- Xuefei Ning, Zinan Lin, Zixuan Zhou, Zifu Wang, Huazhong Yang, Yu Wang:
Skeleton-of-Thought: Prompting LLMs for Efficient Parallel Generation.
- Lijun Yu, José Lezama, Nitesh Bharadwaj Gundavarapu, Luca Versari, Kihyuk Sohn, David Minnen, Yong Cheng, Agrim Gupta, Xiuye Gu, Alexander G. Hauptmann, Boqing Gong, Ming-Hsuan Yang, Irfan Essa, David A. Ross, Lu Jiang:
Language Model Beats Diffusion - Tokenizer is key to visual generation.
- Yeqi Gao, Lianke Qin, Zhao Song, Yitan Wang:
A Sublinear Adversarial Training Algorithm.
- Diego Gomez, Michael Bowling, Marlos C. Machado:
Proper Laplacian Representation Learning.
- Han Guo, Philip Greengard, Eric P. Xing, Yoon Kim:
LQ-LoRA: Low-rank plus Quantized Matrix Decomposition for Efficient Language Model Finetuning.
- Meng Liu, Yue Liu, Ke Liang, Wenzuan Tu, Siwei Wang, Sihang Zhou, Xinwang Liu:
Deep Temporal Graph Clustering.
- Junyan Li, Delin Chen, Yining Hong, Zhenfang Chen, Peihao Chen, Yikang Shen, Chuang Gan:
CoVLM: Composing Visual Entities and Relationships in Large Language Models Via Communicative Decoding.
- Ruqi Yu, Shulei Wang:
Treatment Effects Estimation By Uniform Transformer.
- Joongkyu Lee, Min-hwan Oh:
Demystifying Linear MDPs and Novel Dynamics Aggregation Framework.
- Jian Chen, Ruiyi Zhang, Yufan Zhou, Changyou Chen:
Towards Aligned Layout Generation via Diffusion Model with Aesthetic Constraints.
- Nan Yin, Mengzhu Wang, Zhenghan Chen, Li Shen, Huan Xiong, Bin Gu, Xiao Luo:
DREAM: Dual Structured Exploration with Mixup for Open-set Graph Domain Adaption.
- Chuanhao Li, Chong Liu, Yu-Xiang Wang:
Communication-Efficient Federated Non-Linear Bandit Optimization.
- Yidong Wang, Zhuohao Yu, Wenjin Yao, Zhengran Zeng, Linyi Yang, Cunxiang Wang, Hao Chen, Chaoya Jiang, Rui Xie, Jindong Wang, Xing Xie, Wei Ye, Shikun Zhang, Yue Zhang:
PandaLM: An Automatic Evaluation Benchmark for LLM Instruction Tuning Optimization.
- Sipeng Zheng, Jiazheng Liu, Yicheng Feng, Zongqing Lu:
Steve-Eye: Equipping LLM-based Embodied Agents with Visual Perception in Open Worlds.
- Yujia Wang, Yuanpu Cao, Jingcheng Wu, Ruoyu Chen, Jinghui Chen:
Tackling the Data Heterogeneity in Asynchronous Federated Learning with Cached Update Calibration.
- Sheng Li, Chao Wu, Ao Li, Yanzhi Wang, Xulong Tang, Geng Yuan:
Waxing-and-Waning: a Generic Similarity-based Framework for Efficient Self-Supervised Learning.
- Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, Qingwei Lin, Dixin Jiang:
WizardLM: Empowering Large Pre-Trained Language Models to Follow Complex Instructions.
- Sheng Jin, Xueying Jiang, Jiaxing Huang, Lewei Lu, Shijian Lu:
LLMs Meet VLMs: Boost Open Vocabulary Object Detection with Fine-grained Descriptors.
- Qiongyi Zhou, Changde Du, Shengpei Wang, Huiguang He:
CLIP-MUSED: CLIP-Guided Multi-Subject Visual Neural Information Semantic Decoding.
- Zihan Ding, Chi Jin:
Consistency Models as a Rich and Efficient Policy Class for Reinforcement Learning.
- Joey Hejna, Rafael Rafailov, Harshit Sikchi, Chelsea Finn, Scott Niekum, W. Bradley Knox, Dorsa Sadigh:
Contrastive Preference Learning: Learning from Human Feedback without Reinforcement Learning.
- Yinbin Han, Meisam Razaviyayn, Renyuan Xu:
Neural Network-Based Score Estimation in Diffusion Models: Optimization and Generalization.
- Rasool Fakoor, Jonas Mueller, Zachary Chase Lipton, Pratik Chaudhari, Alex Smola:
Time-Varying Propensity Score to Bridge the Gap between the Past and Present.



- Shenyu Lu, Yipei Wang, Xiaoqian Wang:
Debiasing Attention Mechanism in Transformer without Demographics.
- Xi Victoria Lin, Xilun Chen, Mingda Chen, Weijia Shi, Maria Lomeli, Richard James, Pedro Rodriguez, Jacob Kahn, Gergely Szilvassy, Mike Lewis, Luke Zettlemoyer, Wen-tau Yih:
RA-DIT: Retrieval-Augmented Dual Instruction Tuning.
- Michael-Andrei Panaiteescu-Liess, Yigitcan Kaya, Sicheng Zhu, Furong Huang, Tudor Dumitras:
Like Oil and Water: Group Robustness Methods and Poisoning Defenses May Be at Odds.
- Vishaal Udandarao, Max F. Burg, Samuel Albanie, Matthias Bethge:
Visual Data-Type Understanding does not emerge from scaling Vision-Language Models.
- Eslam Mohamed Bakr, Mohamed Ayman, Mahmoud Ahmed, Habib Slim, Mohamed Elhoseiny:
CoT3DRef: Chain-of-Thoughts Data-Efficient 3D Visual Grounding.
- Aravind Gollakota, Adam R. Klivans, Konstantinos Stavropoulos, Arsen Vasilyan:
An Efficient Tester-Learner for Halfspaces.
- Tales Henrique Carvalho, Kenneth Tjhia, Levi Lelis:
Reclaiming the Source of Programmatic Policies: Programmatic versus Latent Spaces.
- Utkarsh Mall, Cheng Perng Phoo, Meilin Kelsey Liu, Carl Vondrick, Bharath Hariharan, Kavita Bala:
Remote Sensing Vision-Language Foundation Models without Annotations via Ground Remote Alignment.
- Blaise Delattre, Alexandre Araujo, Quentin Barthélémy, Alexandre Allauzen:
The Lipschitz-Variance-Margin Tradeoff for Enhanced Randomized Smoothing.
- Vincent Grari, Thibault Laugel, Tatsunori Hashimoto, Sylvain Lamprier, Marcin Detyniecki:
On the Fairness ROAD: Robust Optimization for Adversarial Debiasing.
- Weiyu Liu, Geng Chen, Joy Hsu, Jiayuan Mao, Jiajun Wu:
Learning Planning Abstractions from Language.
- Sahana Ramnath, Brihi Joshi, Skyler Hallinan, Ximing Lu, Liunian Harold Li, Aaron Chan, Jack Hessel, Yejin Choi, Xiang Ren:
Tailoring Self-Rationalizers with Multi-Reward Distillation.
- Hongxin Zhang, Weihua Du, Jiaming Shan, Qinhong Zhou, Yilun Du, Joshua B. Tenenbaum, Tianmin Shu, Chuang Gan:
Building Cooperative Embodied Agents Modularly with Large Language Models.
- Philippe Chlenski, Ethan Turok, Antonio Khalil Moretti, Itsik Pe'er:
Fast Hyperboloid Decision Tree Algorithms.
- Qinzi Zhang, Hoang Tran, Ashok Cutkosky:
Private Zeroth-Order Nonsmooth Nonconvex Optimization.
- Hong Liu, Zhiyuan Li, David Leo Wright Hall, Percy Liang, Tengyu Ma:
Sophia: A Scalable Stochastic Second-order Optimizer for Language Model Pre-training.
- Shengchao Liu, Jiongxiao Wang, Yijin Yang, Chengpeng Wang, Ling Liu, Hongyu Guo, Chaowei Xiao:
Conversational Drug Editing Using Retrieval and Domain Feedback.
- Amitis Shidani, R. Devon Hjelm, Jason Ramapuram, Russell Webb, Eeshan Gunesh Dhekane, Dan Busbridge:
Poly-View Contrastive Learning.
- Lazar Valkov, Akash Srivastava, Swarat Chaudhuri, Charles Sutton:
A Probabilistic Framework for Modular Continual Learning.
- Mengmeng Xu, Mattia Soldan, Jialin Gao, Shuming Liu, Juan-Manuel Pérez-Rúa, Bernard Ghanem:
Boundary Denoising for Video Activity Localization.
- Rafael A. Rivera Soto, Kailin Koch, Aleem Khan, Barry Y. Chen, Marcus Bishop, Nicholas Andrews:
Few-Shot Detection of Machine-Generated Text using Style Representations.
- Shuyang Yu, Junyuan Hong, Haobo Zhang, Haotao Wang, Zhangyang Wang, Jiayu Zhou:
Safe and Robust Watermark Injection with a Single OoD Image.
- Chenmien Tan, Ge Zhang, Jie Fu:
Massive Editing for Large Language Models via Meta Learning.
- Josue Ortega Caro, Antonio Henrique de Oliveira Fonseca, Syed Asad Rizvi, Matteo Rosati, Christopher L. Averill, James Cross, Prateek Mittal, Emanuele Zappala, Rahul Madhav Dhodapkar, Chadi Abdallah, David van Dijk:
BrainLM: A foundation model for brain activity recordings.
- Xuxi Chen, Yu Yang, Zhangyang Wang, Baharan Mirzasoleiman:
Data Distillation Can Be Like Vodka: Distilling More Times For Better Quality.
- Amin Mansouri, Jason S. Hartford, Yan Zhang, Yoshua Bengio:
Object centric architectures enable efficient causal representation learning.
- Yuhao Huang, Qingsong Wang, Akwum Onwunta, Bao Wang:
Efficient Score Matching with Deep Equilibrium Layers.
- Nikita Srivatsan, Sofía Samaniego, Omar Florez, Taylor Berg-Kirkpatrick:
Alt-Text with Context: Improving Accessibility for Images on Twitter.
- Alihan Hüyük, Qiyo Wei, Alicia Curth, Mihaela van der Schaar:
Defining Expertise: Applications to Treatment Effect Estimation.



- Mingxiao Li, Tingyu Qu, Ruicong Yao, Wei Sun, Marie-Francine Moens:
Alleviating Exposure Bias in Diffusion Models through Sampling with Shifted Time Steps.
- Shih-Hsin Wang, Yung-Chang Hsu, Justin M. Baker, Andrea L. Bertozzi, Jack Xin, Bao Wang:
Rethinking the Benefits of Steerable Features in 3D Equivariant Graph Neural Networks.
- Wei Mao, Richard Hartley, Mathieu Salzmann, Miaomiao Liu:
Neural SDF Flow for 3D Reconstruction of Dynamic Scenes.
- Hongwei Wen, Annika Betken, Hanyuan Hang:
Class Probability Matching with Calibrated Networks for Label Shift Adaption.
- Zilin Si, Gu Zhang, Qingwei Ben, Branden Romero, Zhou Xian, Chao Liu, Chuang Gan:
DIFFTACTILE: A Physics-based Differentiable Tactile Simulator for Contact-rich Robotic Manipulation.
- Tian Yu Liu, Aditya Golatkar, Stefano Soatto:
Tangent Transformers for Composition, Privacy and Removal.
- Linfeng Ye, Shayan Mohajer Hamidi, Renhao Tan, En-Hui Yang:
Bayes Conditional Distribution Estimation for Knowledge Distillation Based on Conditional Mutual Information.
- Melanie Sclar, Yejin Choi, Yulia Tsvetkov, Alane Suhr:
Quantifying Language Models' Sensitivity to Spurious Features in Prompt Design or: How I learned to start worrying about prompt formatting.
- Arpit Bansal, Hong-Min Chu, Avi Schwarzschild, Soumyadip Sengupta, Micah Goldblum, Jonas Geiping, Tom Goldstein:
Universal Guidance for Diffusion Models.
- Archiki Prasad, Elias Stengel-Eskin, Mohit Bansal:
Rephrase, Augment, Reason: Visual Grounding of Questions for Vision-Language Models.
- Haozhe Jiang, Qiwen Cui, Zhihan Xiong, Maryam Fazel, Simon Shaolei Du:
A Black-box Approach for Non-stationary Multi-agent Reinforcement Learning.
- Thomas P. Zollo, Todd Morrill, Zhun Deng, Jake Snell, Toniann Pitassi, Richard S. Zemel:
Prompt Risk Control: A Rigorous Framework for Responsible Deployment of Large Language Models.
- Ivan Grega, Ilyes Batatia, Gábor Csányi, Sri Karlapati, Vikram S. Deshpande:
Energy-conserving equivariant GNN for elasticity of lattice architected metamaterials.
- Junhyung Lyle Kim, Mohammad Taha Toghani, César A. Uribe, Anastasios Kyrillidis:
Adaptive Federated Learning with Auto-Tuned Clients.
- Tianyu Li, Hyunyoung Jung, Matthew C. Gombolay, Yong Kwon Cho, Sehoon Ha:
CrossLoco: Human Motion Driven Control of Legged Robots via Guided Unsupervised Reinforcement Learning.
- Simran Arora, Sabri Eyuboglu, Aman Timalsina, Isys Johnson, Michael Poli, James Zou, Atri Rudra, Christopher Ré:
Zoology: Measuring and Improving Recall in Efficient Language Models.
- Chau Pham, Boyi Liu, Yingxiang Yang, Zhengyu Chen, Tianyi Liu, Jianbo Yuan, Bryan A. Plummer, Zhaoran Wang, Hongxia Yang:
Let Models Speak Ciphers: Multiagent Debate through Embeddings.
- Alexandru Meterez, Amir Joudaki, Francesco Orabona, Alexander Immer, Gunnar Rätsch, Hadi Daneshmand:
Towards Training Without Depth Limits: Batch Normalization Without Gradient Explosion.
- Harikrishna Narasimhan, Aditya Krishna Menon, Wittawat Jitkrittum, Sanjiv Kumar:
Plugin estimators for selective classification with out-of-distribution detection.
- Mike Lasby, Anna Golubeva, Utku Evci, Mihai Nica, Yani Ioannou:
Dynamic Sparse Training with Structured Sparsity.
- Arnab Kumar Mondal, Siba Smarak Panigrahi, Sai Rajeswar, Kaleem Siddiqi, Siamak Ravanbakhsh:
Efficient Dynamics Modeling in Interactive Environments with Koopman Theory.
- Thomas Soares Mullen, Marine Schimel, Guillaume Hennequin, Christian K. Machens, Michael B. Orger, Adrien Jouary :
Learning interpretable control inputs and dynamics underlying animal locomotion.
- Mohammadreza M. Kalan, Samory Kpotufe:
Tight Rates in Supervised Outlier Transfer Learning.
- Frederic Koehler, Thuy-Duong Vuong:
Sampling Multimodal Distributions with the Vanilla Score: Benefits of Data-Based Initialization.
- Zhang-Wei Hong, Idan Shenfeld, Tsun-Hsuan Wang, Yung-Sung Chuang, Aldo Pareja, James R. Glass, Akash Srivastava, Pulkit Agrawal:
Curiosity-driven Red-teaming for Large Language Models.
- Victor Akinwande, Yiding Jiang, Dylan Sam, J. Zico Kolter:
Understanding prompt engineering may not require rethinking generalization.
- Jonathan D. Chang, Dhruv Sreenivas, Yingbing Huang, Kianté Brantley, Wen Sun:
Adversarial Imitation Learning via Boosting.
- Adrián Bazaga, Pietro Lio, Gos Micklem:
Unsupervised Pretraining for Fact Verification by Language Model Distillation.

-     Sara Ghazanfari, Alexandre Araujo, Prashanth Krishnamurthy, Farshad Khorrami, Siddharth Garg:
LipSim: A Provably Robust Perceptual Similarity Metric.
-     Megan Richards, Polina Kirichenko, Diane Bouchacourt, Mark Ibrahim:
Does Progress On Object Recognition Benchmarks Improve Generalization on Crowdsourced, Global Data?
-     Arjun Ashok, Étienne Marcotte, Valentina Zantedeschi, Nicolas Chapados, Alexandre Drouin:
TACTiS-2: Better, Faster, Simpler Attentional Copulas for Multivariate Time Series.
-     Panagiotis Theodoropoulos, Guan-Horng Liu, Tianrong Chen, Augustinos D. Saravacos, Evangelos A. Theodorou:
A robust differential Neural ODE Optimizer.
-     Tatjana Chavdarova , Tong Yang, Matteo Pagliardini, Michael I. Jordan:
A Primal-Dual Approach to Solving Variational Inequalities with General Constraints.
-     Saurabh Garg, Mehrdad Farajtabar, Hadi Pouransari, Raviteja Vemulapalli, Sachin Mehta, Oncel Tuzel, Vaishaal Shankar, Fartash Faghri:
TiC-CLIP: Continual Training of CLIP Models.
-     Mehrdad Saberi, Vinu Sankar Sadasivan, Keivan Rezaei, Aounon Kumar, Atoosa Malemir Chegini, Wenxiao Wang, Soheil Feizi:
Robustness of AI-Image Detectors: Fundamental Limits and Practical Attacks.
-     Duong Minh Le, Yang Chen, Alan Ritter, Wei Xu:
Constrained Decoding for Cross-lingual Label Projection.
-     Yujia Bao, Srinivasan Sivanandan, Theofanis Karaletsos:
Channel Vision Transformers: An Image Is Worth 1 x 16 x 16 Words.
-     Noel Loo, Ramin M. Hasani, Mathias Lechner, Alexander Amini, Daniela Rus:
Understanding Reconstruction Attacks with the Neural Tangent Kernel and Dataset Distillation.
-     Yi-Lin Sung, Jaehong Yoon, Mohit Bansal:
ECoFLaP: Efficient Coarse-to-Fine Layer-Wise Pruning for Vision-Language Models.
-     Andrei Lupu, Chris Lu, Jarek Liesen, Robert Tjarko Lange, Jakob Nicolaus Foerster:
Behaviour Distillation.
-     Youbang Sun, Zitao Li, Yaliang Li, Bolin Ding:
Improving LoRA in Privacy-preserving Federated Learning.
-     Ming-Yang Zhou, Zichao Yan, Elliot Layne, Nikolay Malkin, Dinghuai Zhang, Moksh Jain, Mathieu Blanchette, Yoshua Bengio:
PhyloGFN: Phylogenetic inference with generative flow networks.
-     Junbo Li, Zichen Miao, Qiang Qiu, Ruqi Zhang:
Training Bayesian Neural Networks with Sparse Subspace Variational Inference.
-     Federico Barbero, Ameya Velingker, Amin Saberi, Michael M. Bronstein, Francesco Di Giovanni:
Locality-Aware Graph Rewiring in GNNs.
-     Zhixiang Chi, Li Gu, Tao Zhong, Huan Liu, Yuanhao Yu, Konstantinos N. Plataniotis, Yang Wang:
Adapting to Distribution Shift by Visual Domain Prompt Generation.
-     Zhenbang Wu, Anant Dadu, Nicholas J. Tustison, Brian B. Avants, Mike A. Nalls, Jimeng Sun, Faraz Faghri:
Multimodal Patient Representation Learning with Missing Modalities and Labels.
-     Lorenz Richter, Julius Berner:
Improved sampling via learned diffusions.
-     Deyao Zhu, Jun Chen, Xiaoqian Shen, Xiang Li, Mohamed Elhoseiny:
MiniGPT-4: Enhancing Vision-Language Understanding with Advanced Large Language Models.
-     Sepanta Zeighami, Cyrus Shahabi:
Towards Establishing Guaranteed Error for Learned Database Operations.
-     Gabriele Sarti, Grzegorz Chrupala, Malvina Nissim, Arianna Bisazza:
Quantifying the Plausibility of Context Reliance in Neural Machine Translation.
-     Yongsheng Mei, Mahdi Imani, Tian Lan:
Bayesian Optimization through Gaussian Cox Process Models for Spatio-temporal Data.
-     Clément Bonnet, Daniel Luo, Donal Byrne, Shikha Surana, Sasha Abramowitz, Paul Duckworth, Vincent Coyette, Laurence Illing Midgley, Elshadai Tegegn, Tristan Kalloniatis, Omayma Mahjoub, Matthew Macfarlane, Andries P. Smit, Nathan Grinsztajn, Raphaël Boige, Cemlyn N. Waters, Mohamed A. Mimouni, Ulrich A. Mbou Sob, Ruan de Kock, Siddarth Singh, Daniel Furelos-Blanco, Victor Le, Arnu Pretorius, Alexandre Laterre:
Jumanji: a Diverse Suite of Scalable Reinforcement Learning Environments in JAX.
-     Raj Ghugare, Santiago Miret, Adriana Hugessen, Mariano Phielipp, Glen Berseth:
Searching for High-Value Molecules Using Reinforcement Learning and Transformers.
-     Ernst Röell, Bastian Rieck:
Differentiable Euler Characteristic Transforms for Shape Classification.
-     Mingxuan Li, Junzhe Zhang, Elias Bareinboim:
Causally Aligned Curriculum Learning.



- Oren Mangoubi, Nisheeth K. Vishnoi:
Faster Sampling from Log-Concave Densities over Polytopes via Efficient Linear Solvers.
- Zeyu Zhou, Ruqi Bai, Sean Kulinski, Murat Kocaoglu, David I. Inouye:
Towards Characterizing Domain Counterfactuals for Invertible Latent Causal Models.
- Tanishq Kumar, Blake Bordelon, Samuel J. Gershman, Cengiz Pehlevan:
Grokking as the transition from lazy to rich training dynamics.
- Mirco Mutti, Riccardo De Santi, Marcello Restelli, Alexander Marx, Giorgia Ramponi:
Exploiting Causal Graph Priors with Posterior Sampling for Reinforcement Learning.
- Neha Gupta, Harikrishna Narasimhan, Wittawat Jitkrittum, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar:
Language Model Cascades: Token-Level Uncertainty And Beyond.
- Runzhe Wang, Sadhika Malladi, Tianhao Wang, Kaifeng Lyu, Zhiyuan Li:
The Marginal Value of Momentum for Small Learning Rate SGD.
- Hannah Lawrence, Mitchell Tong Harris:
Learning Polynomial Problems with $SL(2, \mathbb{R})$ -Equivariance.
- Hanqing Zeng, Hanjia Lyu, Diyi Hu, Yinglong Xia, Jiebo Luo:
Mixture of Weak and Strong Experts on Graphs.
- Reese Pathak, Rajat Sen, Weihao Kong, Abhimanyu Das:
Transformers can optimally learn regression mixture models.
- Ming-Yu Chung, Sheng-Yen Chou , Chia-Mu Yu, Pin-Yu Chen, Sy-Yen Kuo, Tsung-Yi Ho:
Rethinking Backdoor Attacks on Dataset Distillation: A Kernel Method Perspective.
- Xiaogang Jia, Denis Blessing, Xinkai Jiang, Moritz Reuss, Atalay Donat, Rudolf Lioutikov, Gerhard Neumann:
Towards Diverse Behaviors: A Benchmark for Imitation Learning with Human Demonstrations.
- Mehdi Zadem, Sergio Mover, Sao Mai Nguyen:
Reconciling Spatial and Temporal Abstractions for Goal Representation.
- Soham Gadgil, Ian Connick Covert, Su-In Lee:
Estimating Conditional Mutual Information for Dynamic Feature Selection.
- Rachit Bansal, Bidisha Samanta, Siddharth Dalmia, Nitish Gupta, Sriram Ganapathy, Abhishek Bapna, Prateek Jain, Partha Talukdar:
LLM Augmented LLMs: Expanding Capabilities through Composition.
- Libin Zhu, Chaoyue Liu, Adityanarayanan Radhakrishnan, Mikhail Belkin:
Quadratic models for understanding catapult dynamics of neural networks.
- Arian Rokkum Jamasb, Alex Morehead, Chaitanya K. Joshi, Zuobai Zhang, Kieran Didi, Simon V. Mathis, Charles Harris, Jian Tang, Jianlin Cheng, Pietro Lio, Tom L. Blundell:
Evaluating Representation Learning on the Protein Structure Universe.
- Pratyush Maini, Sachin Goyal, Zachary Chase Lipton, J. Zico Kolter, Aditi Raghunathan:
T-MARS: Improving Visual Representations by Circumventing Text Feature Learning.
- Lukas Blecher, Guillem Cucurull, Thomas Scialom, Robert Stojnic:
Nougat: Neural Optical Understanding for Academic Documents.
- Enric Boix-Adserà, Omid Saremi, Emmanuel Abbe, Samy Bengio, Eta Littwin, Joshua M. Susskind:
When can transformers reason with abstract symbols?
- Fan Shi, Bin Li, Xiangyang Xue:
Towards Generative Abstract Reasoning: Completing Raven's Progressive Matrix via Rule Abstraction and Selection.
- Marco Pacini, Xiaowen Dong, Bruno Lepri, Gabriele Santin:
A Characterization Theorem for Equivariant Networks with Point-wise Activations.
- Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, Vaishnav Nagarajan:
Think before you speak: Training Language Models With Pause Tokens.
- Bahare Fatemi, Jonathan Halcrow, Bryan Perozzi:
Talk like a Graph: Encoding Graphs for Large Language Models.
- Fan Wu, Huseyin A. Inan, Arturs Backurs, Varun Chandrasekaran, Janardhan Kulkarni, Robert Sim:
Privately Aligning Language Models with Reinforcement Learning.
- Bowen Peng, Jeffrey Quesnelle, Honglu Fan, Enrico Shippole:
YaRN: Efficient Context Window Extension of Large Language Models.
- Xun Tang, Michael Shavlovsky, Holakou Rahamanian, Elisa Tardini, Kiran Koshy Thekumpampil, Tesi Xiao, Lexing Ying:
Accelerating Sinkhorn algorithm with sparse Newton iterations.
- Shanda Li, Chong You, Guru Guruganesh, Joshua Ainslie, Santiago Ontañón, Manzil Zaheer, Sumit Sanghai, Yiming Yang, Sanjiv Kumar, Srinadh Bhojanapalli:
Functional Interpolation for Relative Positions improves Long Context Transformers.
- Stephanie Fu, Mark Hamilton, Laura E. Brandt , Axel Feldmann, Zhoutong Zhang, William T. Freeman:
FeatUp: A Model-Agnostic Framework for Features at Any Resolution.

-    Sarthak Yadav, Sergios Theodoridis, Lars Kai Hansen, Zheng-Hua Tan:
Masked Autoencoders with Multi-Window Local-Global Attention Are Better Audio Learners.

   Stephen Marcus McAleer, JB Lanier, Kevin A. Wang, Pierre Baldi, Tuomas Sandholm, Roy Fox:
Toward Optimal Policy Population Growth in Two-Player Zero-Sum Games.

   Robert Huben, Hoagy Cunningham, Logan Riggs, Aidan Ewart, Lee Sharkey:
Sparse Autoencoders Find Highly Interpretable Features in Language Models.

   Wei-Cheng Huang, Chun-Fu Richard Chen, Hsiang Hsu:
OVOR: OnePrompt with Virtual Outlier Regularization for Rehearsal-Free Class-Incremental Learning.

   Saeed Saremi, Ji Won Park, Francis R. Bach:
Chain of Log-Concave Markov Chains.

   Zhiwei Xu, Yutong Wang, Spencer Frei, Gal Vardi, Wei Hu:
Benign Overfitting and Grokking in ReLU Networks for XOR Cluster Data.

   Chenyu Liu, Xinliang Zhou, Zhengri Zhu, Liming Zhai, Ziyu Jia, Yang Liu:
VBH-GNN: Variational Bayesian Heterogeneous Graph Neural Networks for Cross-subject Emotion Recognition.

   Prithvijit Chattopadhyay, Bharat Goyal, Boglarka Ecsedi, Viraj Prabhu, Judy Hoffman:
AUGCAL: Improving Sim2Real Adaptation by Uncertainty Calibration on Augmented Synthetic Images.

   Herbie Bradley, Andrew Dai, Hannah Benita Teufel, Jenny Zhang, Koen Oostermeijer, Marco Bellagente, Jeff Clune, Kenneth O. Stanley, Gr  gory Schott, Joel Lehman:
Quality-Diversity through AI Feedback.

   Zhepeng Cen, Zuxin Liu, Zitong Wang, Yihang Yao, Henry Lam, Ding Zhao:
Learning from Sparse Offline Datasets via Conservative Density Estimation.

   Rohan Subramani, Marcus Williams, Max Heitmann, Halfdan Holm, Charlie Griffin, Joar Max Viktor Skalse:
On the Expressivity of Objective-Specification Formalisms in Reinforcement Learning.

   Gianluca M. Bencomo, Jake Snell, Thomas L. Griffiths:
Implicit Maximum a Posteriori Filtering via Adaptive Optimization.

   Junfeng Long, Zirui Wang, Quanyi Li, Liu Cao, Jiawei Gao, Jiangmiao Pang:
Hybrid Internal Model: Learning Agile Legged Locomotion with Simulated Robot Response.

   Safa Messaoud, Billel Mokeddem, Zhenghai Xue, Linsey Pang, Bo An, Haipeng Chen, Sanjay Chawla:
S2AC: Energy-Based Reinforcement Learning with Stein Soft Actor Critic.

   Manju Garimella, Denizhan Pak, Justin N. Wood, Samantha Marie Waters Wood:
A Newborn Embodied Turing Test for Comparing Object Segmentation Across Animals and Machines.

   Saba Ghaffari, Ehsan Saleh, Alexander G. Schwing, Yu-Xiong Wang, Martin D. Burke, Saurabh Sinha:
Robust Model-Based Optimization for Challenging Fitness Landscapes.

   Shikai Fang, Madison Cooley, Da Long, Shibo Li, Mike Kirby, Shandian Zhe:
Solving High Frequency and Multi-Scale PDEs with Gaussian Processes.

   Keiran Paster, Marco Dos Santos, Zhangir Azerbayev, Jimmy Ba:
OpenWebMath: An Open Dataset of High-Quality Mathematical Web Text.

   Dhruva Tirumala, Thomas Lampe, Jos   Enrique Chen, Tuomas Haarnoja, Sandy H. Huang, Guy Lever, Ben Moran, Tim Hertweck, Leonard Hasenclever, Martin A. Riedmiller, Nicolas Heess, Markus Wulfmeier:
Replay across Experiments: A Natural Extension of Off-Policy RL.

   Zixiang Chen, Yihe Deng, Yuanzhi Li, Quanquan Gu:
Understanding Transferable Representation Learning and Zero-shot Transfer in CLIP.

   Gabriel della Maggiora, Luis Alberto Croquevielle, Nikita Deshpande, Harry Horsley, Thomas Heinis, Artur Yakimovich:
Conditional Variational Diffusion Models.

   Peiyan Hu, Yue Wang, Zhi-Ming Ma:
Better Neural PDE Solvers Through Data-Free Mesh Movers.

   Claudio Battiloro, Indro Spinelli, Lev Telyatnikov, Michael M. Bronstein, Simone Scardapane, Paolo Di Lorenzo:
From Latent Graph to Latent Topology Inference: Differentiable Cell Complex Module.

   Frederikke Isa Marin, Felix Teufel, Marc Horlacher, Dennis Madsen, Dennis Pultz, Ole Winther, Wouter Boomsma:
BEND: Benchmarking DNA Language Models on Biologically Meaningful Tasks.

   Xu Zheng, Farhad Shirani, Tianchun Wang, Wei Cheng, Zhuomin Chen, Haifeng Chen, Hua Wei, Dongsheng Luo:
Towards Robust Fidelity for Evaluating Explainability of Graph Neural Networks.

   Maxime Wabbartha, Joelle Pineau:
Piecewise Linear Parametrization of Policies: Towards Interpretable Deep Reinforcement Learning.

   Keita Suzuki, Taiji Suzuki:
Optimal criterion for feature learning of two-layer linear neural network in high dimensional interpolation regime.

   Arip Asadulaev, Alexander Korotin, Vage Egiazarian, Petr Mokrov, Evgeny Burnaev:
Neural Optimal Transport with General Cost Functionals.

-      Yu Wang, Tong Zhao, Yuying Zhao, Yunchao Liu, Xueqi Cheng, Neil Shah, Tyler Derr:
A Topological Perspective on Demystifying GNN-Based Link Prediction Performance.
-      Firas Al-Hafez, Guoping Zhao, Jan Peters, Davide Tateo:
Time-Efficient Reinforcement Learning with Stochastic Stateful Policies.
-      Ge Li, Hongyi Zhou, Dominik Roth, Serge Thilges, Fabian Otto, Rudolf Lioutikov, Gerhard Neumann:
Open the Black Box: Step-based Policy Updates for Temporally-Correlated Episodic Reinforcement Learning.
-      Mitja Nikolaus:
Emergent Communication with Conversational Repair.
-      Gourav Datta, Zeyu Liu, Peter Anthony Beerel:
Can we get the best of both Binary Neural Networks and Spiking Neural Networks for Efficient Computer Vision?
-      Hao Xiong, Yehui Tang, Yunlin He, Wei Tan, Junchi Yan:
Node2ket: Efficient High-Dimensional Network Embedding in Quantum Hilbert Space.
-      Marco Federici, Patrick Forré, Ryota Tomioka, Bastiaan S. Veeling:
Latent Representation and Simulation of Markov Processes via Time-Lagged Information Bottleneck.
-      António Farinhas, Chrysoula Zerva, Dennis Ulmer, André F. T. Martins:
Non-Exchangeable Conformal Risk Control.
-      Ruiquan Huang, Yingbin Liang, Jing Yang:
Provably Efficient UCB-type Algorithms For Learning Predictive State Representations.
-      Peter Sorrenson, Felix Draxler, Armand Rousselot, Sander Hummerich, Lea Zimmermann, Ullrich Köthe:
Lifting Architectural Constraints of Injective Flows.
-      Tianzhe Chu, Shengbang Tong, Tianjiao Ding, Xili Dai, Benjamin David Haeffele, René Vidal, Yi Ma:
Image Clustering via the Principle of Rate Reduction in the Age of Pretrained Models.
-      Tianxin Wei, Bowen Jin, Ruirui Li, Hansi Zeng, Zhengyang Wang, Jianhui Sun, Qingyu Yin, Hanqing Lu, Suhang Wang, Jingrui He, Xianfeng Tang:
Towards Unified Multi-Modal Personalization: Large Vision-Language Models for Generative Recommendation and Beyond.
-      Matthew Morris, Bernardo Cuenca Grau, Ian Horrocks:
Orbit-Equivariant Graph Neural Networks.
-      Daniil E. Kirilenko, Vitaliy Vorobyov, Alexey K. Kovalev, Aleksandr Panov:
Object-Centric Learning with Slot Mixture Module.
-      Yu-Lin Tsai, Chia-Yi Hsu, Chulin Xie, Chih-Hsun Lin, Jia-You Chen, Bo Li, Pin-Yu Chen, Chia-Mu Yu, Chun-Ying Huang:
Ring-A-Bell! How Reliable are Concept Removal Methods For Diffusion Models?
-      Bang An, Sicheng Zhu, Michael-Andrei Panaiteescu-Liess, Chaithanya Kumar Mummadipalli, Furong Huang:
PerceptionCLIP: Visual Classification by Inferring and Conditioning on Contexts.
-      Dawid Jan Kopiczko, Tijmen Blankevoort, Yuki M. Asano:
VeRA: Vector-based Random Matrix Adaptation.
-      Qi Zhao, Shijie Wang, Ce Zhang, Changcheng Fu, Minh Quan Do, Nakul Agarwal, Kwonjoon Lee, Chen Sun:
AntGPT: Can Large Language Models Help Long-term Action Anticipation from Videos?
-      Junhao Hu, Weijie Gan, Zhixin Sun, Hongyu An, Ulugbek Kamilov:
A Plug-and-Play Image Registration Network.
-      Haixin Wang, Jiaxin Li, Anubhav Dwivedi, Kentaro Hara, Tailin Wu:
BENO: Boundary-embedded Neural Operators for Elliptic PDEs.
-      Konstantin Hess, Valentyn Melnychuk, Dennis Frauen, Stefan Feuerriegel:
Bayesian Neural Controlled Differential Equations for Treatment Effect Estimation.
-      Ruibo Liu, Ruixin Yang, Chenyan Jia, Ge Zhang, Diyi Yang, Soroush Vosoughi:
Training Socially Aligned Language Models on Simulated Social Interactions.
-      Aaron Spieler, Nasim Rahaman, Georg Martius, Bernhard Schölkopf, Anna Levina:
The Expressive Leaky Memory Neuron: an Efficient and Expressive Phenomenological Neuron Model Can Solve Long-Horizon Tasks.
-      Zhepei Wei, Chuanhao Li, Tianze Ren, Haifeng Xu, Hongning Wang:
Incentivized Truthful Communication for Federated Bandits.
-      Salar Abbaspourazad, Oussama Elachqar, Andrew C. Miller, Saba Emrani, Udhayakumar Nallasamy, Ian Shapiro:
Large-scale Training of Foundation Models for Wearable Biosignals.
-      Dinghuai Zhang, Ricky T. Q. Chen, Cheng-Hao Liu, Aaron C. Courville, Yoshua Bengio:
Diffusion Generative Flow Samplers: Improving learning signals through partial trajectory optimization.
-      Ge Gao, Qitong Gao, Xi Yang, Song Ju, Miroslav Pajic, Min Chi:
On Trajectory Augmentations for Off-Policy Evaluation.
-      Alain Rakotomamonjy, Kimia Nadjahi, Liva Ralaivola:
Federated Wasserstein Distance.



- Haojie Huang, Owen Howell, Dian Wang, Xupeng Zhu, Robert Platt, Robin Walters:
Fourier Transporter: Bi-Equivariant Robotic Manipulation in 3D.
- Cong Liu, David Ruhe, Floor Eijkelboom, Patrick Forré:
Clifford Group Equivariant Simplicial Message Passing Networks.
- Hongtao Wu, Ya Jing, Chilam Cheang, Guangzeng Chen, Jiafeng Xu, Xinghang Li, Minghuan Liu, Hang Li, Tao Kong:
Unleashing Large-Scale Video Generative Pre-training for Visual Robot Manipulation.
- Shyamgopal Karthik, Karsten Roth, Massimiliano Mancini, Zeynep Akata:
Vision-by-Language for Training-Free Compositional Image Retrieval.
- Miao Lu, Beining Wu, Xiaodong Yang, Difan Zou:
Benign Oscillation of Stochastic Gradient Descent with Large Learning Rate.
- Hao Cheng, Qingsong Wen, Yang Liu, Liang Sun:
RobustTSF: Towards Theory and Design of Robust Time Series Forecasting with Anomalies.
- Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, Jelena Luketina, Eric Hambro, Edward Grefenstette, Roberta Raileanu:
Understanding the Effects of RLHF on LLM Generalisation and Diversity.
- Tianyu He, Junliang Guo, Runyi Yu, Yuchi Wang, Jialiang Zhu, Kaikai An, Leyi Li, Xu Tan, Chunyu Wang, Han Hu, HsiangTao Wu, Sheng Zhao, Jiang Bian:
GAIA: Zero-shot Talking Avatar Generation.
- Amirhossein Vahidi, Simon Schoßer, Lisa Wimmer, Yawei Li, Bernd Bischl, Eyke Hüllermeier, Mina Rezaei:
Probabilistic Self-supervised Representation Learning via Scoring Rules Minimization.
- Allan Jabri, Sjoerd van Steenkiste, Emiel Hoogeboom, Mehdi S. M. Sajjadi, Thomas Kipf:
DORSal: Diffusion for Object-centric Representations of Scenes et al.
- Archibald Fraikin, Adrien Bennetot, Stéphanie Allassonnière:
T-Rep: Representation Learning for Time Series using Time-Embeddings.
- Sara Klein, Simon Weissmann, Leif Döring:
Beyond Stationarity: Convergence Analysis of Stochastic Softmax Policy Gradient Methods.
- Saleh Ashkboos, Maximilian L. Croci, Marcelo Gennari Do Nascimento, Torsten Hoefer, James Hensman:
SliceGPT: Compress Large Language Models by Deleting Rows and Columns.
- Stefano B. Blumberg, Paddy J. Slator, Daniel C. Alexander:
Experimental Design for Multi-Channel Imaging via Task-Driven Feature Selection.
- Edward Milsom, Ben Anson, Laurence Aitchison:
Convolutional Deep Kernel Machines.
- Hannah Kriesel, Leon Sick, Tristan Payer, Tim Bergner, Kavitha Shaga Devan, Clarissa Read, Paul Walther, Timo Ropinski, Pedro Hermosilla:
Weakly Supervised Virus Capsid Detection with Image-Level Annotations in Electron Microscopy Images.
- Ziyang Xiao, Dongxiang Zhang, Yangjun Wu, Lilin Xu, Yuan Jessica Wang, Xiongwei Han, Xiaojin Fu, Tao Zhong, Jia Zeng, Mingli Song, Gang Chen:
Chain-of-Experts: When LLMs Meet Complex Operations Research Problems.
- John Kirchenbauer, Jonas Geiping, Yuxin Wen, Manli Shu, Khalid Saifullah, Kezhi Kong, Kasun Fernando, Aniruddha Saha, Micah Goldblum, Tom Goldstein:
On the Reliability of Watermarks for Large Language Models.
- Juan Elenter, Luiz F. O. Chamon, Alejandro Ribeiro:
Near-Optimal Solutions of Constrained Learning Problems.
- Alizée Pace, Hugo Yèche, Bernhard Schölkopf, Gunnar Rätsch, Guy Tennenholz:
Delphic Offline Reinforcement Learning under Nonidentifiable Hidden Confounding.
- Jiayuan Ye, Anastasia Borovykh, Soufiane Hayou, Reza Shokri:
Leave-one-out Distinguishability in Machine Learning.
- Haobo Song, Hao Zhao, Soumajit Majumder, Tao Lin:
Increasing Model Capacity for Free: A Simple Strategy for Parameter Efficient Fine-tuning.
- Yohann Benchetrit, Hubert J. Banville, Jean-Remi King:
Brain decoding: toward real-time reconstruction of visual perception.
- Yury Nahshan, Joseph Kampeas, Emir Haleva:
Linear Log-Normal Attention with Unbiased Concentration.
- Finn Rietz, Erik Schaffernicht, Stefan Heinrich, Johannes A. Stork:
Prioritized Soft Q-Decomposition for Lexicographic Reinforcement Learning.
- Petr Mokrov, Alexander Korotin, Alexander Kolesov, Nikita Gushchin, Evgeny Burnaev:
Energy-guided Entropic Neural Optimal Transport.
- Li Ren, Chen Chen, Liqiang Wang, Kien A. Hua:
Learning Semantic Proxies from Visual Prompts for Parameter-Efficient Fine-Tuning in Deep Metric Learning.
- Nate Gruver, Anuroop Sriram, Andrea Madotto, Andrew Gordon Wilson, C. Lawrence Zitnick, Zachary W. Ulissi:
Fine-Tuned Language Models Generate Stable Inorganic Materials as Text.

-     Tianyu Huang, Yihan Zeng, Bowen Dong, Hang Xu, Songcen Xu, Rynson W. H. Lau, Wangmeng Zuo: **TextField3D: Towards Enhancing Open-Vocabulary 3D Generation with Noisy Text Fields.**
-     Renat Sergazinov, Elizabeth Chun, Valeriya Rogovchenko, Nathaniel J. Fernandes, Nicholas Kasman, Irina Gaynanova: **GlucoBench: Curated List of Continuous Glucose Monitoring Datasets with Prediction Benchmarks.**
-     Zequn Yang, Yake Wei, Ce Liang, Di Hu: **Quantifying and Enhancing Multi-modal Robustness with Modality Preference.**
-     Artem Agafonov, Dmitry Kamzolov, Alexander V. Gasnikov, Ali Kavis, Kimon Antonakopoulos, Volkan Cevher, Martin Takáć: **Advancing the Lower Bounds: an Accelerated, Stochastic, Second-order Method with Optimal Adaptation to Inexactness.**
-     Eran Rosenbluth, Jan Tönshoff, Martin Ritzert, Berke Kisin, Martin Grohe: **Distinguished In Uniform: Self-Attention Vs. Virtual Nodes.**
-     Weida Li, Yaoliang Yu: **Faster Approximation of Probabilistic and Distributional Values via Least Squares.**
-     Ziyang Yu, Wenbing Huang, Yang Liu: **Rigid Protein-Protein Docking via Equivariant Elliptic-Paraboloid Interface Prediction.**
-     Zhaoxuan Wu , Mohammad Mohammadi Amiri, Ramesh Raskar, Bryan Kian Hsiang Low: **Incentive-Aware Federated Learning with Training-Time Model Rewards.**
-     Chenyu Wang, Sharut Gupta, Caroline Uhler, Tommi S. Jaakkola: **Removing Biases from Molecular Representations via Information Maximization.**
-     Xiaohuan Pei, Yanxi Li, Minjing Dong , Chang Xu: **Neural Architecture Retrieval.**
-     Hancheng Min, Enrique Mallada, René Vidal: **Early Neuron Alignment in Two-layer ReLU Networks with Small Initialization.**
-     Xianghong Fang, Jian Li, Qiang Sun, Benyou Wang: **Rethinking the Uniformity Metric in Self-Supervised Learning.**
-     Kai Cui, Sascha Hauck, Christian Fabian, Heinz Koeppl: **Learning Decentralized Partially Observable Mean Field Control for Artificial Collective Behavior.**
-     Haiyan Jiang, Vincent Zoonekynd, Giulia De Masi, Bin Gu, Huan Xiong: **TAB: Temporal Accumulated Batch Normalization in Spiking Neural Networks.**
-     Shengzhong Zhang, Wenjie Yang, Xinyuan Cao, Hongwei Zhang, Zengfeng Huang: **StructComp: Substituting propagation with Structural Compression in Training Graph Contrastive Learning.**
-     Gabriele Corso, Yilun Xu, Valentin De Bortoli, Regina Barzilay, Tommi S. Jaakkola: **Particle Guidance: non-I.I.D. Diverse Sampling with Diffusion Models.**
-     Max Maria Losch, Mohamed Omran, David Stutz, Mario Fritz, Bernt Schiele: **On Adversarial Training without Perturbing all Examples.**
-     Chen Gan, Zihao Yin, Kelei He, Yang Gao, Junfeng Zhang: **Diving Segmentation Model into Pixels.**
-     Xinyue Liu, Hualin Zhang, Bin Gu, Hong Chen: **General Stability Analysis for Zeroth-Order Optimization Algorithms.**
-     Zihao Yin, Chen Gan, Kelei He, Yang Gao, Junfeng Zhang: **Hybrid Sharing for Multi-Label Image Classification.**
-     Aaron Zweig, Joan Bruna: **Symmetric Single Index Learning.**
-     Bo Li, Xiaowen Jiang, Mikkel N. Schmidt, Tommy Sonne Alstrøm, Sebastian U. Stich: **An improved analysis of per-sample and per-update clipping in federated learning.**
-     Emmeran Johnson, Ciara Pike-Burke, Patrick Rebeschini: **Sample-Efficiency in Multi-Batch Reinforcement Learning: The Need for Dimension-Dependent Adaptivity.**
-     Matthias Lanzinger, Pablo Barceló: **On the Power of the Weisfeiler-Leman Test for Graph Motif Parameters.**
-     Wenwen Si, Sangdon Park, Insup Lee, Edgar Dobriban, Osbert Bastani: **PAC Prediction Sets Under Label Shift.**
-     Wenhao Wang, Muhammad Ahmad Kaleem, Adam Dziedzic, Michael Backes, Nicolas Papernot, Franziska Boenisch: **Memorization in Self-Supervised Learning Improves Downstream Generalization.**
-     Zhixuan Lin, Pierluca D'Oro, Evgenii Nikishin, Aaron C. Courville: **The Curse of Diversity in Ensemble-Based Exploration.**
-     Yixin Cheng, Grigorios Chrysos, Markos Georgopoulos, Volkan Cevher: **Multilinear Operator Networks.**
-     Ayesha Vermani, Il Memming Park , Josue Nassar: **Leveraging Generative Models for Unsupervised Alignment of Neural Time Series Data.**



- Wenxuan Zhou, Sheng Zhang, Yu Gu, Muhan Chen, Hoifung Poon:
UniversalNER: Targeted Distillation from Large Language Models for Open Named Entity Recognition.
- Antonis Antoniades, Yiyi Yu, Joseph Canzano, William Yang Wang, Spencer L. Smith:
Neuroformer: Multimodal and Multitask Generative Pretraining for Brain Data.
- Zifan Wu, Bo Tang, Qian Lin, Chao Yu, Shangqin Mao, Qianlong Xie, Xingxing Wang, Dong Wang:
Off-Policy Primal-Dual Safe Reinforcement Learning.
- Yifan Lu, Yue Hu, Yiqi Zhong, Dequan Wang, Yanfeng Wang, Siheng Chen:
An Extensible Framework for Open Heterogeneous Collaborative Perception.
- Benjie Wang, Joel Jennings, Wenbo Gong:
Neural structure learning with stochastic differential equations.
- Joar Max Viktor Skalse, Lucy Farnik, Sumeet Ramesh Motwani, Erik Jenner, Adam Gleave, Alessandro Abate:
STARC: A General Framework For Quantifying Differences Between Reward Functions.
- Grégoire Mialon, Clémentine Fourrier, Thomas Wolf, Yann LeCun, Thomas Scialom:
GAIA: a benchmark for General AI Assistants.
- Chaoming Wang, Tianqiu Zhang, Sichao He , Hongyaoxing Gu, Shangyang Li, Si Wu:
A differentiable brain simulator bridging brain simulation and brain-inspired computing.
- Hadar Sivan, Moshe Gabel, Assaf Schuster:
FOSI: Hybrid First and Second Order Optimization.
- Christian Fabian, Kai Cui, Heinz Koepll:
Learning Mean Field Games on Sparse Graphs: A Hybrid Graphex Approach.
- Harold Benoit, Liangze Jiang, Andrei Atanov, Oguzhan Fatih Kar, Mattia Rigotti, Amir Zamir:
Unraveling the Key Components of OOD Generalization via Diversification.
- Emanuele Palumbo, Laura Manduchi, Sonia Laguna, Daphné Chopard, Julia E. Vogt:
Deep Generative Clustering with Multimodal Diffusion Variational Autoencoders.
- Ruizhe Shi, Yuyao Liu, Yanjie Ze, Simon Shaolei Du, Huazhe Xu:
Unleashing the Power of Pre-trained Language Models for Offline Reinforcement Learning.
- Bowen Jing, Tommi S. Jaakkola, Bonnie Berger:
Equivariant Scalar Fields for Molecular Docking with Fast Fourier Transforms.
- Richard Ngo, Lawrence Chan, Sören Mindermann:
The Alignment Problem from a Deep Learning Perspective.
- Matthew Thomas Jackson, Chris Lu, Louis Kirsch, Robert Tjarko Lange, Shimon Whiteson, Jakob Nicolaus Foerster:
Discovering Temporally-Aware Reinforcement Learning Algorithms.
- Yeongyeon Na, Minje Park, Yunwon Tae, Sunghoon Joo:
Guiding Masked Representation Learning to Capture Spatio-Temporal Relationship of Electrocardiogram.
- Mihaela C. Stoian, Salijona Dyrmishi, Maxime Cordy, Thomas Lukasiewicz, Eleonora Giunchiglia:
How Realistic Is Your Synthetic Data? Constraining Deep Generative Models for Tabular Data.
- Yanbo Wang, Jian Liang, Ran He:
Towards Eliminating Hard Label Constraints in Gradient Inversion Attacks.
- Marten Lienen, David Lüdke, Jan Hansen-Palmus, Stephan Günnemann:
From Zero to Turbulence: Generative Modeling for 3D Flow Simulation.
- Chao Chen, Kai Liu, Ze Chen, Yi Gu, Yue Wu, Mingyuan Tao, Zhihang Fu, Jieping Ye:
INSIDE: LLMs' Internal States Retain the Power of Hallucination Detection.
- Hong Chen, Yipeng Zhang, Simin Wu, Xin Wang, Xuguang Duan, Yuwei Zhou, Wenwu Zhu:
DisenBooth: Identity-Preserving Disentangled Tuning for Subject-Driven Text-to-Image Generation.
- Axel Laborieux, Friedemann Zenke:
Improving equilibrium propagation without weight symmetry through Jacobian homeostasis.
- Jianshu Hu , Yunpeng Jiang, Paul Weng:
Revisiting Data Augmentation in Deep Reinforcement Learning.
- Aoran Wang, Jun Pang:
Structural Inference with Dynamics Encoding and Partial Correlation Coefficients.
- Kelly Maggs, Celia Hacker, Bastian Rieck:
Simplicial Representation Learning with Neural k-Forms.
- Liyang Zhu, Meng Ding, Vaneet Aggarwal, Jinhui Xu, Di Wang:
Improved Analysis of Sparse Linear Regression in Local Differential Privacy Model.
- Haotian Xue, Chumeng Liang, Xiaoyu Wu, Yongxin Chen:
Toward effective protection against diffusion-based mimicry through score distillation.
- Ryan Wong, Necati Cihan Camgöz, Richard Bowden:
Sign2GPT: Leveraging Large Language Models for Gloss-Free Sign Language Translation.

-     Shujian Yu, Xi Yu, Sigurd Løkse, Robert Jenssen, José C. Príncipe:
Cauchy-Schwarz Divergence Information Bottleneck for Regression.
-     Changli Tang, Wenyi Yu, Guangzhi Sun, Xianzhao Chen, Tian Tan, Wei Li, Lu Lu, Zejun Ma, Chao Zhang:
SALMONN: Towards Generic Hearing Abilities for Large Language Models.
-     Michalis K. Titsias, Alexandre Galashov, Amal Rannen-Triki, Razvan Pascanu, Yee Whye Teh, Jörg Bornschein:
Kalman Filter for Online Classification of Non-Stationary Data.
-     Maciej Mikula, Szymon Tworkowski, Szymon Antoniak, Bartosz Piotrowski, Albert Q. Jiang, Jin Peng Zhou, Christian Szegedy, Lukasz Kucinski, Piotr Milos, Yuhuai Wu:
Magnushammer: A Transformer-Based Approach to Premise Selection.
-     Whie Jung, Jaehoon Yoo, Sungjin Ahn, Seunghoon Hong:
Learning to Compose: Improving Object Centric Learning by Injecting Compositionality.
-     Alexander Korotin, Nikita Gushchin, Evgeny Burnaev:
Light Schrödinger Bridge.
-     Marc Rigter, Minqi Jiang, Ingmar Posner:
Reward-Free Curricula for Training Robust World Models.
-     Jeff Guo, Philippe Schwaller:
Beam Enumeration: Probabilistic Explainability For Sample Efficient Self-conditioned Molecular Design.
-     Gundeep Arora, Srujana Merugu, Anoop Saladi, Rajeev Rastogi:
Leveraging Uncertainty Estimates To Improve Classifier Performance.
-     Hee Suk Yoon, Eunseop Yoon, Joshua Tian Jin Tee, Mark A. Hasegawa-Johnson, Yingzhen Li, Chang D. Yoo:
C-TPT: Calibrated Test-Time Prompt Tuning for Vision-Language Models via Text Feature Dispersion.
-     Ahmet Iscen, Mathilde Caron, Alireza Fathi, Cordelia Schmid:
Retrieval-Enhanced Contrastive Vision-Text Models.
-     Ilan Price, Nicholas Daultry Ball, Adam C. Jones, Samuel C. H. Lam, Jared Tanner:
Deep Neural Network Initialization with Sparsity Inducing activations.
-     Jonathan Scott, Hossein Zakerinia, Christoph H. Lampert:
PeFLL: Personalized Federated Learning by Learning to Learn.
-     Max Zimmer, Christoph Spiegel, Sebastian Pokutta:
Sparse Model Soups: A Recipe for Improved Pruning via Model Averaging.
-     Louis Béthune, Thomas Massena, Thibaut Boissin, Aurélien Bellet, Franck Mamalet, Yannick Prudent, Corentin Friedrich, Mathieu Serrurier, David Vigouroux:
DP-SGD Without Clipping: The Lipschitz Neural Network Way.
-     Haque Ishfaq, Qingfeng Lan, Pan Xu, A. Rupam Mahmood, Doina Precup, Anima Anandkumar, Kamyar Azizzadenesheli:
Provable and Practical: Efficient Exploration in Reinforcement Learning via Langevin Monte Carlo.
-     Tom Hosking, Phil Blunsom, Max Bartolo:
Human Feedback is not Gold Standard.
-     Han Zhou , Xingchen Wan, Lev Proleev, Diana Mincu, Jilin Chen, Katherine A. Heller, Subhrajit Roy:
Batch Calibration: Rethinking Calibration for In-Context Learning and Prompt Engineering.
-     Paul Hagemann, Johannes Hertrich, Fabian Altekrüger, Robert Beinert, Jannis Chemeddine, Gabriele Steidl:
Posterior Sampling Based on Gradient Flows of the MMD with Negative Distance Kernel.
-     Oguz Kaan Yüksel, Etienne Boursier, Nicolas Flammarion:
First-order ANIL provably learns representations despite overparametrisation.
-     Linhao Luo, Yuan-Fang Li, Gholamreza Haffari, Shirui Pan:
Reasoning on Graphs: Faithful and Interpretable Large Language Model Reasoning.
-     Lorenzo Pacchiardi, Alex James Chan, Sören Mindermann, Ilan Moscovitz, Alexa Y. Pan, Yarin Gal, Owain Evans, Jan Markus Brauner:
How to Catch an AI Liar: Lie Detection in Black-Box LLMs by Asking Unrelated Questions.
-     Han Zhang, Yu Lei, Lin Gui, Min Yang, Yulan He, Hui Wang, Ruifeng Xu:
CPPO: Continual Learning for Reinforcement Learning with Human Feedback.
-     Francesco Bacchicocchi, Matteo Castiglioni, Alberto Marchesi, Nicola Gatti:
Learning Optimal Contracts: How to Exploit Small Action Spaces.
-     Yuanwen Yue, Sabarinath Mahadevan, Jonas Schult, Francis Engelmann, Bastian Leibe, Konrad Schindler, Theodora Kontogianni:
AGILE3D: Attention Guided Interactive Multi-object 3D Segmentation.
-     Longwei Zou, Han Zhang, Yangdong Deng:
A Multi-Level Framework for Accelerating Training Transformer Models.
-     Federico Cacciamani, Matteo Castiglioni, Nicola Gatti:
Online Information Acquisition: Hiring Multiple Agents.
-     Nhu-Thuat Tran, Hady W. Lauw:
Learning Multi-Faceted Prototypical User Interests.



Yongyuan Liang, Yanchao Sun, Ruijie Zheng, Xiangyu Liu, Benjamin Eysenbach, Tuomas Sandholm, Furong Huang, Stephen Marcus McAleer:
Game-Theoretic Robust Reinforcement Learning Handles Temporally-Coupled Perturbations.

Martino Bernasconi, Matteo Castiglioni, Andrea Celli, Federico Fusco:
Bandits with Replenishable Knapsacks: the Best of both Worlds.

Siyuan Guo, Jonas Bernhard Wildberger, Bernhard Schölkopf:
Out-of-Variable Generalisation for Discriminative Models.

Yeongmin Kim, Byeonghu Na, Minsang Park, JoonHo Jang, Dongjun Kim, Wanmo Kang, Il-Chul Moon:
Training Unbiased Diffusion Models From Biased Dataset.

Jianfa Lai, Zhifan Li, Dongming Huang, Qian Lin:
The optimality of kernel classifiers in Sobolev space.

Masanori Koyama, Kenji Fukumizu, Kohei Hayashi, Takeru Miyato:
Neural Fourier Transform: A General Approach to Equivariant Representation Learning.

Feng Hong, Jiangchao Yao, Yueming Lyu, Zhihan Zhou, Ivor W. Tsang, Ya Zhang, Yanfeng Wang:
On Harmonizing Implicit Subpopulations.

Yaniv Blumenfeld, Itay Hubara, Daniel Soudry:
Towards Cheaper Inference in Deep Networks with Lower Bit-Width Accumulators.

Donggyu Lee, Sangwon Jung, Taesup Moon:
Continual Learning in the Presence of Spurious Correlations: Analyses and a Simple Baseline.

Yufei Kuang, Jie Wang, Haoyang Liu, Fangzhou Zhu, Xijun Li, Jia Zeng, Jianye Hao, Bin Li, Feng Wu:
Rethinking Branching on Exact Combinatorial Optimization Solver: The First Deep Symbolic Discovery Framework.

Guocheng Qian, Jinjie Mai, Abdullah Hamdi, Jian Ren, Aliaksandr Siarohin, Bing Li, Hsin-Ying Lee, Ivan Skorokhodov, Peter Wonka, Sergey Tulyakov, Bernard Ghanem:
Magic123: One Image to High-Quality 3D Object Generation Using Both 2D and 3D Diffusion Priors.

Kejun Tang, Jiayu Zhai, Xiaoliang Wan, Chao Yang:
Adversarial Adaptive Sampling: Unify PINN and Optimal Transport for the Approximation of PDEs.

Chendi Qian, Andrei Manolache, Kareem Ahmed, Zhe Zeng, Guy Van den Broeck, Mathias Niepert, Christopher Morris:
Probabilistically Rewired Message-Passing Neural Networks.

Yanzhou Li, Tianlin Li, Kangjie Chen, Jian Zhang, Shangqing Liu, Wenhan Wang, Tianwei Zhang, Yang Liu:
BadEdit: Backdooring Large Language Models by Model Editing.

Yonggang Zhang, Zhiqin Yang, Xinmei Tian, Nannan Wang, Tongliang Liu, Bo Han:
Robust Training of Federated Models with Extremely Label Deficiency.

Guoqiang Zhang, Kenta Niwa, W. Bastiaan Kleijn:
On Accelerating Diffusion-Based Sampling Processes via Improved Integration Approximation.

James Chapman, Lennie Wells, Ana Lawry Aguila:
Unconstrained Stochastic CCA: Unifying Multiview and Self-Supervised Learning.

Ishita Mediratta, Qingfei You, Minqi Jiang, Roberta Raileanu:
The Generalization Gap in Offline Reinforcement Learning.

Sina Khajehabdollahi, Roxana Zeraati, Emmanouil Giannakakis, Tim Jakob Schäfer, Georg Martius, Anna Levina:
Emergent mechanisms for long timescales depend on training curriculum and affect performance in memory tasks.

Francisco Andrade, Gabriel Peyré, Clarice Poon:
Sparsistency for inverse optimal transport.

Minyang Hu, Hong Chang, Bingpeng Ma, Shiguang Shan, Xilin Chen:
Scalable Modular Network: A Framework for Adaptive Learning via Agreement Routing.

Oren Katzir, Or Patashnik, Daniel Cohen-Or, Dani Lischinski:
Noise-free Score Distillation.

Hao Wang, Chenyi Zhang, Tongyang Li:
Near-Optimal Quantum Algorithm for Minimizing the Maximal Loss.

Yanqin Jiang, Li Zhang, Jin Gao, Weiming Hu, Yao Yao:
Consistent4D: Consistent 360° Dynamic Object Generation from Monocular Video.

Moritz Akiya Zanger, Wendelin Boehmer, Matthijs T. J. Spaan:
Diverse Projection Ensembles for Distributional Reinforcement Learning.

Hanlei Zhang, Xin Wang, Hua Xu, Qianrui Zhou , Kai Gao, Jianhua Su, Jinyue Zhao, Wenrui Li, Yanting Chen:
MIIntRec2.0: A Large-scale Benchmark Dataset for Multimodal Intent Recognition and Out-of-scope Detection in Conversations.

Xiaosen Zheng, Tianyu Pang, Chao Du, Jing Jiang, Min Lin:
Intriguing Properties of Data Attribution on Diffusion Models.

Ahmad Bdeir, Kristian Schwethelm , Niels Landwehr:
Fully Hyperbolic Convolutional Neural Networks for Computer Vision.



- Zhiyu Zhu, Xinyi Wang, Zhibo Jin, Jiayu Zhang, Huaming Chen:
Enhancing Transferable Adversarial Attacks on Vision Transformers through Gradient Normalization Scaling and High-Frequency Adaptation.
- Min Xue, Artur Andrzejak, Marla Leuther:
An interpretable error correction method for enhancing code-to-code translation.
- Jianlan Luo, Perry Dong, Yuexiang Zhai, Yi Ma, Sergey Levine:
RLIF: Interactive Imitation Learning as Reinforcement Learning.
- Samir Khaki , Konstantinos N. Plataniotis:
The Need for Speed: Pruning Transformers with One Recipe.
- Ruijie Wu, Xiaojian Ma, Zhenliang Zhang, Wei Wang, Qing Li, Song-Chun Zhu, Yizhou Wang:
Bongard-OpenWorld: Few-Shot Reasoning for Free-form Visual Concepts in the Real World.
- Sihang Li, Zhiyuan Liu , Yanchen Luo, Xiang Wang, Xiangnan He, Kenji Kawaguchi, Tat-Seng Chua, Qi Tian:
Towards 3D Molecule-Text Interpretation in Language Models.
- Amro Abbas, Evgenia Rusak, Kushal Tirumala, Wieland Brendel, Kamalika Chaudhuri, Ari S. Morcos:
Effective pruning of web-scale datasets based on complexity of concept clusters.
- Zhiyu Zhu, Huaming Chen, Jiayu Zhang, Xinyi Wang, Zhibo Jin, Jason Xue, Flora D. Salim:
AttEXplore: Attribution for Explanation with model parameters eXploration.
- Viet Quoc Vo, Ehsan Abbasnejad, Damith Ranasinghe:
Brusleattack: a Query-Efficient Score-based Black-Box Sparse Adversarial Attack.
- Vincent Leroy, Jérôme Revaud, Thomas Lucas, Philippe Weinzaepfel:
Win-Win: Training High-Resolution Vision Transformers from Two Windows.
- Sihan Chen, Xingjian He, Handong Li, Xiaojie Jin, Jiashi Feng, Jing Liu:
COSA: Concatenated Sample Pretrained Vision-Language Foundation Model.
- Seyyede Fatemeh Seyyedsalehi, Mahdieh Soleymani Baghshah, Hamid R. Rabiee:
SOInter: A Novel Deep Energy-Based Interpretation Method for Explaining Structured Output Models.
- Aiwei Liu, Leyi Pan, Xuming Hu, Shuang Li, Lijie Wen, Irwin King, Philip S. Yu:
An Unforgeable Publicly Verifiable Watermark for Large Language Models.
- Duanyi Yao, Songze Li, Ye Xue, Jin Liu:
Constructing Adversarial Examples for Vertical Federated Learning: Optimal Client Corruption through Multi-Armed Bandit.
- Enming Liang, Minghua Chen:
Generative Learning for Solving Non-Convex Problem with Multi-Valued Input-Solution Mapping.
- Xianghao Kong, Ollie Liu, Han Li, Dani Yogatama, Greg Ver Steeg:
Interpretable Diffusion via Information Decomposition.
- Karim Abdel Sadek, Marek Eliás:
Algorithms for Caching and MTS with reduced number of predictions.
- Jiseok Chae, Kyuwon Kim, Donghwan Kim:
Two-timescale Extragradient for Finding Local Minimax Points.
- Sheng Xu, Guiliang Liu:
Uncertainty-aware Constraint Inference in Inverse Constrained Reinforcement Learning.
- Xuanlei Zhao, Shenggan Cheng, Guangyang Lu, Haotian Zhou, Bin Jia, Yang You:
AutoChunk: Automated Activation Chunk for Memory-Efficient Deep Learning Inference.
- Juncheng Liu, Bryan Hooi, Kenji Kawaguchi, Yiwei Wang, Chaosheng Dong, Xiaokui Xiao:
Scalable and Effective Implicit Graph Neural Networks on Large Graphs.
- Bowen Cao, Deng Cai, Leyang Cui, Xuxin Cheng, Wei Bi, Yuexian Zou, Shuming Shi:
Retrieval is Accurate Generation.
- Changwen Zhang, Wenli Ouyang, Hao Yuan, Liming Gong, Yong Sun, Ziao Guo, Zhichen Dong, Junchi Yan:
Towards Imitation Learning to Branch for MIP: A Hybrid Reinforcement Learning based Sample Augmentation Approach.
- Tianyu Fan, Lirong Wu, Yufei Huang, Haitao Lin, Cheng Tan, Zhangyang Gao, Stan Z. Li:
Decoupling Weighing and Selecting for Integrating Multiple Graph Pre-training Tasks.
- Zikai Xiao, Zihan Chen, Liyinglan Liu, Yang Feng, Joey Tianyi Zhou, Jian Wu, Wanlu Liu, Howard Hao Yang, Zuozhu Liu:
FedLoGe: Joint Local and Generic Federated Learning under Long-tailed Data.
- Robert Jenssen:
MAP IT to Visualize Representations.
- Jianhao Shen, Ye Yuan, Srbuhi Mirzoyan, Ming Zhang, Chenguang Wang:
Measuring Vision-Language STEM Skills of Neural Models.
- David Brellmann, Eloïse Berthier, David Filliat, Goran Frehse:
On Double Descent in Reinforcement Learning with LSTD and Random Features.

-     Raphaël Avalos, Florent Delgrange, Ann Nowé, Guillermo A. Pérez, Diederik M. Roijers:
The Wasserstein Believer: Learning Belief Updates for Partially Observable Environments through Reliable Latent Space Models.
-     Aming Wu, Cheng Deng:
Modulated Phase Diffusor: Content-Oriented Feature Synthesis for Detecting Unknown Objects.
-     Haruka Kiyohara, Ren Kishimoto, Kosuke Kawakami, Ken Kobayashi, Kazuhide Nakata, Yuta Saito:
Towards Assessing and Benchmarking Risk-Return Tradeoff of Off-Policy Evaluation.
-     Na Li, Yuchen Jiao, Hangguan Shan, Shefeng Yan:
Provable Memory Efficient Self-Play Algorithm for Model-free Reinforcement Learning.
-     Moritz Imfeld, Jacopo Graldi, Marco Giordano, Thomas Hofmann, Sotiris Anagnostidis, Sidak Pal Singh:
Transformer Fusion with Optimal Transport.
-     Xun Wu, Shaohan Huang, Furu Wei:
Mixture of LoRA Experts.
-     Hila Manor, Tomer Michaeli:
On the Posterior Distribution in Denoising: Application to Uncertainty Quantification.
-     Weidi Xu, Jingwei Wang, Lele Xie, Jianshan He, Hongting Zhou, Taifeng Wang, Xiaopei Wan, Jingdong Chen, Chao Qu, Wei Chu:
LogicMP: A Neuro-symbolic Approach for Encoding First-order Logic Constraints.
-     Marcel Binz, Eric Schulz:
Turning large language models into cognitive models.
-     Shashanka Venkataramanan, Amir Ghodrati, Yuki M. Asano, Fatih Porikli, Amirhossein Habibian:
Skip-Attention: Improving Vision Transformers by Paying Less Attention.
-     Xiang Lisa Li, Vaishnavi Shrivastava, Siyan Li, Tatsunori Hashimoto, Percy Liang:
Benchmarking and Improving Generator-Validator Consistency of Language Models.
-     Guang Lin, Chao Li, Jianhai Zhang, Toshihisa Tanaka, Qibin Zhao:
Adversarial Training on Purification (AToP): Advancing Both Robustness and Generalization.
-     Kun Lei, Zhengmao He, Chenhao Lu, Kaizhe Hu, Yang Gao, Huazhe Xu:
Uni-O4: Unifying Online and Offline Deep Reinforcement Learning with Multi-Step On-Policy Optimization.
-     Hanan Gani, Shariq Farooq Bhat, Muzammal Naseer, Salman Khan, Peter Wonka:
LLM Blueprint: Enabling Text-to-Image Generation with Complex and Detailed Prompts.
-     Xiaotian Han, Jianfeng Chi, Yu Chen, Qifan Wang, Han Zhao, Na Zou, Xia Hu:
FFB: A Fair Fairness Benchmark for In-Processing Group Fairness Methods.
-     Kai Lagemann, Christian Lagemann, Sach Mukherjee:
Invariance-based Learning of Latent Dynamics.
-     Youn-Yeol Yu, Jeongwhan Choi, Woojin Cho, Kookjin Lee, Nayong Kim, Kiseok Chang, ChangSeung Woo, Ilho Kim, SeokWoo Lee, Joon-Young Yang, Sooyoung Yoon, Noseong Park:
Learning Flexible Body Collision Dynamics with Hierarchical Contact Mesh Transformer.
-     Yiqun Yao, Zheng Zhang, Jing Li, Yequan Wang:
Masked Structural Growth for 2x Faster Language Model Pre-training.
-     Haozhe Zhao, Zefan Cai, Shuzheng Si, Xiaojian Ma, Kaikai An, Liang Chen, Zixuan Liu, Sheng Wang, Wenjuan Han, Baobao Chang:
MMICL: Empowering Vision-language Model with Multi-Modal In-Context Learning.
-     Xinyu Hu, Pengfei Tang, Simiao Zuo, Zihan Wang, Bowen Song, Qiang Lou, Jian Jiao, Denis Charles:
Evoke: Evoking Critical Thinking Abilities in LLMs via Reviewer-Author Prompt Editing.
-     Byeonghu Na, Yeongmin Kim, HeeSun Bae, Jung Hyun Lee, Se Jung Kwon, Wanmo Kang, Il-Chul Moon:
Label-Noise Robust Diffusion Models.
-     Siqiao Xue, Xiaoming Shi, Zhixuan Chu, Yan Wang, Hongyan Hao, Fan Zhou, Caigao Jiang, Chen Pan, James Y. Zhang, Qingsong Wen, Jun Zhou, Hongyuan Mei:
EasyTPP: Towards Open Benchmarking Temporal Point Processes.
-     Bo Peng, Yadan Luo, Yonggang Zhang, Yixuan Li, Zhen Fang:
ConjNorm: Tractable Density Estimation for Out-of-Distribution Detection.
-     Haoying Li, Jixin Zhao, Shangchen Zhou, Huajun Feng, Chongyi Li, Chen Change Loy:
Adaptive Window Pruning for Efficient Local Motion Deblurring.
-     Enshu Liu, Xuefei Ning, Huazhong Yang, Yu Wang:
A Unified Sampling Framework for Solver Searching of Diffusion Probabilistic Models.
-     Neria Uzan, Nir Weinberger:
A representation-learning game for classes of prediction tasks.
-     Jicong Fan, Rui Chen, Zhao Zhang, Chris Ding:
Neuron-Enhanced AutoEncoder Matrix Completion and Collaborative Filtering: Theory and Practice.
-     Sam Bond-Taylor, Chris G. Willcocks:
∞-Diff: Infinite Resolution Diffusion with Subsampled Mollified States.

-     Zhongwang Zhang, Yuqing Li, Tao Luo, Zhi-Qin John Xu:
Stochastic Modified Equations and Dynamics of Dropout Algorithm.
-     Peiyan Zhang, Haoyang Liu, Chaozhuo Li, Xing Xie, Sunghun Kim, Haohan Wang:
Foundation Model-oriented Robustness: Robust Image Model Evaluation with Pretrained Models.
-     Daiki Chijiwa:
Transferring Learning Trajectories of Neural Networks.
-     Kashif Rasul, Andrew Bennett, Pablo Vicente, Umang Gupta, Hena Ghonia, Anderson Schneider, Yuriy Nevmyvaka:
VQ-TR: Vector Quantized Attention for Time Series Forecasting.
-     Yujee Song, Donghyun Lee, Rui Meng, Won Hwa Kim:
Decoupled Marked Temporal Point Process using Neural Ordinary Differential Equations.
-     Seon-Ho Lee, Nyeong-Ho Shin, Chang-Su Kim:
Unsupervised Order Learning.
-     Hoyong Kim, Kangil Kim:
Fixed Non-negative Orthogonal Classifier: Inducing Zero-mean Neural Collapse with Feature Dimension Separation.
-     Hao Wang, Yongsheng Yu, Tiejian Luo, Heng Fan, Libo Zhang:
MaGIC: Multi-modality Guided Image Completion.
-     Hang Yu, Cong Liao, Ruolan Liu, Jianguo Li, Yun Hu, Xinze Wang:
AmortizedPeriod: Attention-based Amortized Inference for Periodicity Identification.
-     Tianyu Guo, Wei Hu, Song Mei, Huan Wang, Caiming Xiong, Silvio Savarese, Yu Bai:
How Do Transformers Learn In-Context Beyond Simple Functions? A Case Study on Learning with Representations.
-     Jiawei Yang, Boris Ivanovic, Or Litany, Xinshuo Weng, Seung Wook Kim, Boyi Li, Tong Che, Danfei Xu, Sanja Fidler, Marco Pavone, Yue Wang:
EmerNeRF: Emergent Spatial-Temporal Scene Decomposition via Self-Supervision.
-     Yiliu Wang, Wei Chen, Milan Vojnovic:
Combinatorial Bandits for Maximum Value Reward Function under Value-Index Feedback.
-     Xiaoxiao Sun, Yue Yao, Shengjin Wang, Hongdong Li, Liang Zheng:
Alice Benchmarks: Connecting Real World Re-Identification with the Synthetic.
-     Andrew Szot, Max Schwarzer, Harsh Agrawal, Bogdan Mazoure, Rin Metcalf, Walter Talbott, Natalie Mackraz, R. Devon Hjelm, Alexander T. Toshev:
Large Language Models as Generalizable Policies for Embodied Tasks.
-     Yilun Du, Sherry Yang, Pete Florence, Fei Xia, Ayzaan Wahid, Brian Ichter, Pierre Sermanet, Tianhe Yu, Pieter Abbeel, Joshua B. Tenenbaum, Leslie Pack Kaelbling, Andy Zeng, Jonathan Tompson:
Video Language Planning.
-     Bin Zhu, Bin Lin, Munan Ning, Yang Yan, Jiaxi Cui, Hongfa Wang, Yatian Pang, Wenhao Jiang, Junwu Zhang, Zongwei Li, Caiwan Zhang, Zhifeng Li, Wei Liu, Li Yuan:
LanguageBind: Extending Video-Language Pretraining to N-modality by Language-based Semantic Alignment.
-     Fei Kong, Jinhao Duan, Ruipeng Ma, Heng Tao Shen, Xiaoshuang Shi, Xiaofeng Zhu, Kaidi Xu:
An Efficient Membership Inference Attack for the Diffusion Model by Proximal Initialization.
-     Yin Fang, Xiaozhuan Liang, Ningyu Zhang, Kangwei Liu, Rui Huang, Zhuo Chen, Xiaohui Fan, Huajun Chen:
Mol-Instructions: A Large-Scale Biomolecular Instruction Dataset for Large Language Models.
-     Yiyang Ma, Huan Yang, Wenhan Yang, Jianlong Fu, Jiaying Liu:
Solving Diffusion ODEs with Optimal Boundary Conditions for Better Image Super-Resolution.
-     Soobin Um, Suhyeon Lee, Jong Chul Ye:
Don't Play Favorites: Minority Guidance for Diffusion Models.
-     Adel Javanmard, Lin Chen, Vahab Mirrokni, Ashwinkumar Badanidiyuru, Gang Fu:
Learning from Aggregate responses: Instance Level versus Bag Level Loss Functions.
-     Qinzhong Zhou, Sunli Chen, Yisong Wang, Haozhe Xu, Weihua Du, Hongxin Zhang, Yilun Du, Joshua B. Tenenbaum, Chuang Gan:
HAZARD Challenge: Embodied Decision Making in Dynamically Changing Environments.
-     Bin Lu, Tingyan Ma, Xiaoying Gan, Xinbing Wang, Yunqiang Zhu, Chenghu Zhou, Shiyu Liang:
Temporal Generalization Estimation in Evolving Graphs.
-     Licong Lin, Yu Bai, Song Mei:
Transformers as Decision Makers: Provable In-Context Reinforcement Learning via Supervised Pretraining.
-     Ziwei Guan, Yi Zhou, Yingbin Liang:
On the Hardness of Online Nonconvex Optimization with Single Oracle Feedback.
-     Shengjie Zhou, Lue Tao, Yuzhou Cao, Tao Xiang, Bo An, Lei Feng:
On the Vulnerability of Adversarially Trained Models Against Two-faced Attacks.
-     Wentao Wu, Aleksei Timofeev, Chen Chen, Bowen Zhang, Kun Duan, Shuangning Liu, Yantao Zheng, Jonathon Shlens, Xianzhi Du, Yinfai Yang:

**MOFI: Learning Image Representations from Noisy Entity Annotated Images.**

- Kushagra Pandey, Maja Rudolph, Stephan Mandt:
Efficient Integrators for Diffusion Generative Models.
- Yanqiao Zhu, Jeehyun Hwang, Keir Adams, Zhen Liu, Bozhao Nan, Brock Stenfors, Yuanqi Du, Jatin Chauhan, Olaf Wiest, Olexandr Isayev, Connor W. Coley, Yizhou Sun, Wei Wang:
Learning Over Molecular Conformer Ensembles: Datasets and Benchmarks.
- Ge Yan, Yaniv Romano, Tsui-Wei Weng:
Provably Robust Conformal Prediction with Improved Efficiency.
- Xianjun Yang, Wei Cheng, Yue Wu, Linda Ruth Petzold, William Yang Wang, Haifeng Chen:
DNA-GPT: Divergent N-Gram Analysis for Training-Free Detection of GPT-Generated Text.
- Zhenheng Tang, Yonggang Zhang, Shaohuai Shi, Xinmei Tian, Tongliang Liu, Bo Han, Xiaowen Chu:
FedImpro: Measuring and Improving Client Update in Federated Learning.
- Jonghyun Lee, Hansam Cho, Young Joon Yoo, Seoung Bum Kim, Yonghyun Jeong:
Compose and Conquer: Diffusion-Based 3D Depth Aware Composable Image Synthesis.
- Sebastian Shenghong Tay, Chuan-Sheng Foo, Daisuke Urano, Richalynn Leong, Bryan Kian Hsiang Low:
A Unified Framework for Bayesian Optimization under Contextual Uncertainty.
- Zhaoyi Zhou, Chunling Zhu, Runlong Zhou, Qiwen Cui, Abhishek Gupta, Simon Shaolei Du:
Free from Bellman Completeness: Trajectory Stitching via Model-based Return-conditioned Supervised Learning.
- Changyou Chen, Han Ding, Bunyamin Sisman, Yi Xu, Ouye Xie, Benjamin Z. Yao, Son Dinh Tran, Belinda Zeng:
Diffusion Models for Multi-Task Generative Modeling.
- Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, Graham Neubig:
WebArena: A Realistic Web Environment for Building Autonomous Agents.
- Qingyan Guo, Rui Wang, Junliang Guo, Bei Li, Kaitao Song, Xu Tan, Guoqing Liu, Jiang Bian, Yujiu Yang:
Connecting Large Language Models with Evolutionary Algorithms Yields Powerful Prompt Optimizers.
- Moonseok Choi, Hyungi Lee, Giung Nam, Juho Lee:
Sparse Weight Averaging with Multiple Particles for Iterative Magnitude Pruning.
- Katja Schwarz, Seung Wook Kim, Jun Gao, Sanja Fidler, Andreas Geiger, Karsten Kreis:
WildFusion: Learning 3D-Aware Latent Diffusion Models in View Space.
- Ziqi Xu, Debo Cheng, Jiuyong Li, Jixue Liu, Lin Liu, Kui Yu:
Causal Inference with Conditional Front-Door Adjustment and Identifiable Variational Autoencoder.
- Suttisak Wizadwongsu, Worameth Chinchuthakun, Pramook Khungurn, Amit Raj, Supasorn Suwajanakorn:
Diffusion Sampling with Momentum for Mitigating Divergence Artifacts.
- Shurui Gui, Xiner Li, Shuiwang Ji:
Active Test-Time Adaptation: Theoretical Analyses and An Algorithm.
- Zibin Dong, Yifu Yuan, Jianye Hao, Fei Ni, Yao Mu, Yan Zheng, Yujing Hu, Tangjie Lv, Changjie Fan, Zhipeng Hu:
AlignDiff: Aligning Diverse Human Preferences via Behavior-Customisable Diffusion Model.
- Yong Wu, Yanwei Fu, Shouyan Wang, Xinwei Sun:
Doubly Robust Proximal Causal Learning for Continuous Treatments.
- Chengrui Li, Soon Ho Kim, Chris Rodgers, Hannah Choi, Anqi Wu:
One-hot Generalized Linear Model for Switching Brain State Discovery.
- Cong Lei, Yuxuan Du, Peng Mi, Jun Yu, Tongliang Liu:
Neural Auto-designer for Enhanced Quantum Kernels.
- Satoki Ishikawa, Ryo Karakida:
On the Parameterization of Second-Order Optimization Effective towards the Infinite Width.
- Shen Nie, Hanzhong Allan Guo, Cheng Lu, Yuhao Zhou, Chenyu Zheng, Chongxuan Li:
The Blessing of Randomness: SDE Beats ODE in General Diffusion-based Image Editing.
- Hung Le, Hailin Chen, Amrita Saha, Akash Gokul, Doyen Sahoo, Shafiq Joty:
CodeChain: Towards Modular Code Generation Through Chain of Self-revisions with Representative Sub-modules.
- Xiangyan Liu, Rongxue Li, Wei Ji, Tao Lin:
Towards Robust Multi-Modal Reasoning via Model Selection.
- Yongchao Zhou, Kaifeng Lyu, Ankit Singh Rawat, Aditya Krishna Menon, Afshin Rostamizadeh, Sanjiv Kumar, Jean-François Kagy, Rishabh Agarwal:
DistillSpec: Improving Speculative Decoding via Knowledge Distillation.
- Kuan Li, Yiwen Chen, Yang Liu, Jin Wang, Qing He, Minhao Cheng, Xiang Ao:
Boosting the Adversarial Robustness of Graph Neural Networks: An OOD Perspective.
- Jing Xiong, Zixuan Li, Chuanyang Zheng, Zhijiang Guo, Yichun Yin, Enze Xie, Zhicheng Yang, Qingxing Cao, Haiming Wang, Xiongwei Han, Jing Tang, Chengming Li, Xiaodan Liang:
DQ-LoRe: Dual Queries with Low Rank Approximation Re-ranking for In-Context Learning.

-     Jianlang Chen, Xuhong Ren, Qing Guo, Felix Juefei-Xu, Di Lin, Wei Feng, Lei Ma, Jianjun Zhao:
LRR: Language-Driven Resamplable Continuous Representation against Adversarial Tracking Attacks.
-     Hongjun Wang, Sagar Vaze, Kai Han:
SPTNet: An Efficient Alternative Framework for Generalized Category Discovery with Spatial Prompt Tuning.
-     Mengzhou Xia, Tianyu Gao, Zhiyuan Zeng, Danqi Chen:
Sheared LLaMA: Accelerating Language Model Pre-training via Structured Pruning.
-     Hang Yin, Zihao Wang, Yangqiu Song:
Rethinking Complex Queries on Knowledge Graphs with Neural Link Predictors.
-     Tao Yu, Toni J. B. Liu, Albert Tseng, Christopher De Sa:
Shadow Cones: A Generalized Framework for Partial Order Embeddings.
-     Yikun Ban, Ishika Agarwal, Ziwei Wu, Yada Zhu, Kommy Woldemariam, Hanghang Tong, Jingrui He:
Neural Active Learning Beyond Bandits.
-     Yanbang Wang, Jon M. Kleinberg:
From Graphs to Hypergraphs: Hypergraph Projection and its Reconstruction.
-     Yuyao Zhang, Lan Wei, Nikolaos M. Freris:
Synergistic Patch Pruning for Vision Transformer: Unifying Intra- & Inter-Layer Patch Importance.
-     Yi-Rui Yang, Chang-Wei Shi, Wu-Jun Li:
On the Effect of Batch Size in Byzantine-Robust Distributed Learning.
-     Murong Yue, Jie Zhao, Min Zhang, Liang Du, Ziyu Yao:
Large Language Model Cascades with Mixture of Thought Representations for Cost-Efficient Reasoning.
-     Beomsu Kim, Gihyun Kwon, Kwanyoung Kim, Jong Chul Ye:
Unpaired Image-to-Image Translation via Neural Schrödinger Bridge.
-     Guozheng Ma, Lu Li, Sen Zhang, Zixuan Liu, Zhen Wang, Yixin Chen, Li Shen, Xueqian Wang, Dacheng Tao:
Revisiting Plasticity in Visual Reinforcement Learning: Data, Modules and Training Stages.
-     Haoran Xu, Young Jin Kim, Amr Sharaf, Hany Hassan Awadalla:
A Paradigm Shift in Machine Translation: Boosting Translation Performance of Large Language Models.
-     Renjie Pi, Lewei Yao, Jianhua Han, Xiaodan Liang, Wei Zhang, Hang Xu:
Ins-DetCLIP: Aligning Detection Model to Follow Human-Language Instruction.
-     Chenxiang Ma, Jibin Wu, Chenyang Si, Kay Chen Tan:
Scaling Supervised Local Learning with Augmented Auxiliary Networks.
-     Jiajun Ma, Tianyang Hu, Wenjia Wang, Jiacheng Sun:
Elucidating the design space of classifier-guided diffusion generation.
-     Ashwinee Panda, Christopher A. Choquette-Choo, Zhengming Zhang, Yaoqing Yang, Prateek Mittal:
Teach LLMs to Phish: Stealing Private Information from Language Models.
-     Lirui Wang, Kaiqing Zhang, Allan Zhou, Max Simchowitz, Russ Tedrake:
Robot Fleet Learning via Policy Merging.
-     Ruizhe Liu, Qian Luo, Yanchao Yang:
InfoCon: Concept Discovery with Generative and Discriminative Informativeness.
-     Jianliang He, Han Zhong, Zhuoran Yang:
Sample-efficient Learning of Infinite-horizon Average-reward MDPs with General Function Approximation.
-     Lingbing Guo, Zhuo Chen, Jiaoyan Chen, Yin Fang, Wen Zhang, Huajun Chen:
Revisit and Outstrip Entity Alignment: A Perspective of Generative Models.
-     Ilya Shenbin, Sergey I. Nikolenko:
ImplicitSLIM and How it Improves Embedding-based Collaborative Filtering.
-     Boya Shi, Zhengqin Xu, Shuai Jia, Chao Ma:
Prompt Learning with Quaternion Networks.
-     Yilang Zhang , Georgios B. Giannakis:
Meta-Learning Priors Using Unrolled Proximal Networks.
-     Tserendorj Adiya, Jae Shin Yoon, Jungeun Lee, Sanghun Kim, Hwasup Lim:
Bidirectional Temporal Diffusion Model for Temporally Consistent Human Animation.
-     Felix Petersen, Aashwin Ananda Mishra, Hilde Kuehne, Christian Borgelt, Oliver Deussen, Mikhail Yurochkin:
Uncertainty Quantification via Stable Distribution Propagation.
-     Xiang Fu, Tian Xie, Andrew S. Rosen, Tommi S. Jaakkola, Jake Smith:
MOFDiff: Coarse-grained Diffusion for Metal-Organic Framework Design.
-     Weiming Zhuang, Lingjuan Lyu:
FedWon: Triumphing Multi-domain Federated Learning Without Normalization.
-     Zuxin Liu, Jesse Zhang, Kavosh Asadi, Yao Liu, Ding Zhao, Shoham Sabach, Rasool Fakoor:
TAIL: Task-specific Adapters for Imitation Learning with Large Pretrained Models.

-     Chang Liu, Zhichen Dong, Haobo Ma, Weilin Luo, Xijun Li, Bowen Pang, Jia Zeng, Junchi Yan:
L2P-MIP: Learning to Presolve for Mixed Integer Programming.
-     Sichao Li, Rong Wang, Quanling Deng, Amanda S. Barnard:
Exploring the cloud of feature interaction scores in a Rashomon set.
-     Yunyang Li, Yusong Wang, Lin Huang, Han Yang, Xinran Wei, Jia Zhang, Tong Wang, Zun Wang, Bin Shao, Tie-Yan Liu:
Long-Short-Range Message-Passing: A Physics-Informed Framework to Capture Non-Local Interaction for Scalable Molecular Dynamics Simulation.
-     Yucen Lily Li, Tim G. J. Rudner, Andrew Gordon Wilson:
A Study of Bayesian Neural Network Surrogates for Bayesian Optimization.
-     Michihiro Yasunaga, Xinyun Chen, Yujia Li, Panupong Pasupat, Jure Leskovec, Percy Liang, Ed H. Chi, Denny Zhou:
Large Language Models as Analogical Reasoners.
-     Rabia Gondur, Usama Bin Sikandar, Evan Schaffer, Mikio Christian Aoi , Stephen L. Keeley:
Multi-modal Gaussian Process Variational Autoencoders for Neural and Behavioral Data.
-     Haoyu Lu, Yuqi Huo, Guoxing Yang, Zhiwu Lu, Wei Zhan, Masayoshi Tomizuka, Mingyu Ding:
UniAdapter: Unified Parameter-Efficient Transfer Learning for Cross-modal Modeling.
-     Yu-Yu Wu, Hung-Jui Wang, Shang-Tse Chen:
Annealing Self-Distillation Rectification Improves Adversarial Training.
-     Yitian Zhang, Yue Bai, Huan Wang, Yizhou Wang, Yun Fu:
Don't Judge by the Look: Towards Motion Coherent Video Representation.
-     Ming Jin, Shiyu Wang, Lintao Ma, Zhixuan Chu, James Y. Zhang, Xiaoming Shi, Pin-Yu Chen, Yuxuan Liang, Yuan-Fang Li, Shirui Pan, Qingsong Wen:
Time-LLM: Time Series Forecasting by Reprogramming Large Language Models.
-     Jiangmeng Li, Fei Song, Yifan Jin, Wenwen Qiang, Changwen Zheng, Fuchun Sun, Hui Xiong:
BayesPrompt: Prompting Large-Scale Pre-Trained Language Models on Few-shot Inference via Debiased Domain Abstraction.
-     Yifei Zhou, Ayush Sekhari, Yuda Song, Wen Sun:
Offline Data Enhanced On-Policy Policy Gradient with Provable Guarantees.
-     Yunzhen Feng, Shanmukha Ramakrishna Vedantam, Julia Kempe:
Embarrassingly Simple Dataset Distillation.
-     Cheng Tan, Yijie Zhang , Zhangyang Gao, Bozhen Hu, Siyuan Li, Zicheng Liu, Stan Z. Li:
RDesign: Hierarchical Data-efficient Representation Learning for Tertiary Structure-based RNA Design.
-     Yuxiang Lai, Yi Zhou, Xinghong Liu, Tao Zhou:
Memory-Assisted Sub-Prototype Mining for Universal Domain Adaptation.
-     Tianle Cai, Xuezhi Wang, Tengyu Ma, Xinyun Chen, Denny Zhou:
Large Language Models as Tool Makers.
-     Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Nan Duan, Weizhu Chen:
CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing.
-     Sung Moon Ko, Sumin Lee, Dae-Woong Jeong, Woohyun Lim, Sehui Han:
Geometrically Aligned Transfer Encoder for Inductive Transfer in Regression Tasks.
-     Elias Stengel-Eskin, Kyle Rawlins, Benjamin Van Durme:
Zero and Few-shot Semantic Parsing with Ambiguous Inputs.
-     Yunhui Jang, Dongwoo Kim, Sungsoo Ahn:
Graph Generation with K2-trees.
-     Yutong He, Naoki Murata, Chieh-Hsin Lai, Yuhta Takida, Toshimitsu Uesaka, Dongjun Kim, Wei-Hsiang Liao, Yuki Mitsufuji, J. Zico Kolter, Ruslan Salakhutdinov, Stefano Ermon:
Manifold Preserving Guided Diffusion.
-     Xiyuan Wang, Haotong Yang, Muhan Zhang:
Neural Common Neighbor with Completion for Link Prediction.
-     Xiangming Zhu, Huayu Deng, Haochen Yuan, Yunbo Wang, Xiaokang Yang:
Latent Intuitive Physics: Learning to Transfer Hidden Physics from A 3D Video.
-     Tong Wu, Ashwinee Panda, Jiachen T. Wang, Prateek Mittal:
Privacy-Preserving In-Context Learning for Large Language Models.
-     Caixia Yan, Xiaojun Chang, Zhihui Li, Lina Yao, Minnan Luo, Qinghua Zheng:
Masked Distillation Advances Self-Supervised Transformer Architecture Search.
-     Kibum Kim, Kanghoon Yoon, Yeonjun In, Jinyoung Moon, Donghyun Kim, Chanyoung Park:
Adaptive Self-training Framework for Fine-grained Scene Graph Generation.
-     Chuanqing Wang, Di Wu, Chaoming Fang, Jie Yang, Mohamad Sawan:
Exploring Effective Stimulus Encoding via Vision System Modeling for Visual Prostheses.
-     Mikhail Galkin, Xinyu Yuan, Hesham Mostafa, Jian Tang, Zhaocheng Zhu:
Towards Foundation Models for Knowledge Graph Reasoning.

-     Anne Harrington, Vasha DuTell, Mark Hamilton, Ayush Tewari, Simon Stent, William T. Freeman, Ruth Rosenholtz: **COCO-Periph: Bridging the Gap Between Human and Machine Perception in the Periphery.**
-     Gen Li, Yuting Wei, Yuxin Chen, Yuejie Chi: **Towards Non-Asymptotic Convergence for Diffusion-Based Generative Models.**
-     James B. Simon, Dhruva Karkada, Nikhil Ghosh, Mikhail Belkin: **More is Better: when Infinite Overparameterization is Optimal and Overfitting is Obligatory.**
-     Zhiqing Sun, Yikang Shen, Hongxin Zhang, Qinhong Zhou, Zhenfang Chen, David Daniel Cox, Yiming Yang, Chuang Gan: **SALMON: Self-Alignment with Instructable Reward Models.**
-     Yabo Zhang, Yuxiang Wei, Dongsheng Jiang, Xiaopeng Zhang, Wangmeng Zuo, Qi Tian: **ControlVideo: Training-free Controllable Text-to-video Generation.**
-     Marina Zhang, Owen S. Vallis, Aysegul Bumin, Tanay Vakharia, Elie Burszttein: **RETSim: Resilient and Efficient Text Similarity.**
-     Zhangyang Gao, Cheng Tan, Xingran Chen, Yijie Zhang , Jun Xia, Siyuan Li, Stan Z. Li: **KW-Design: Pushing the Limit of Protein Design via Knowledge Refinement.**
-     Hongyi Wang, Felipe Maia Polo, Yuekai Sun, Souvik Kundu, Eric P. Xing, Mikhail Yurochkin: **Fusing Models with Complementary Expertise.**
-     Jose Javier Gonzalez Ortiz, John V. Guttag, Adrian V. Dalca: **Magnitude Invariant Parametrizations Improve Hypernetwork Learning.**
-     Xinshuai Dong, Biwei Huang, Ignavier Ng, Xiangchen Song, Yujia Zheng, Songyao Jin, Roberto Legaspi, Peter Spirtes, Kun Zhang: **A Versatile Causal Discovery Framework to Allow Causally-Related Hidden Variables.**
-     Xinyun Chen, Maxwell Lin, Nathanael Schärli, Denny Zhou: **Teaching Large Language Models to Self-Debug.**
-     Asaf Yehudai, Boaz Carmeli, Yosi Mass, Ofir Ariv, Nathaniel Mills, Eyal Shnarch, Leshem Choshen: **Achieving Human Parity in Content-Grounded Datasets Generation.**
-     Zhiwei Li, Guodong Long, Tianyi Zhou: **Federated Recommendation with Additive Personalization.**
-     Yuhang Liu, Zhen Zhang, Dong Gong, Mingming Gong, Biwei Huang, Anton van den Hengel, Kun Zhang, Javen Qinfeng Shi: **Identifiable Latent Polynomial Causal Models through the Lens of Change.**
-     Shuai Zhao , Xiaohan Wang, Linchao Zhu, Yi Yang: **Test-Time Adaptation with CLIP Reward for Zero-Shot Generalization in Vision-Language Models.**
-     Changmao Li, Jeffrey Flanigan: **Future Language Modeling from Temporal Document History.**
-     Siyuan Li, Weiyang Jin, Zedong Wang, Fang Wu, Zicheng Liu, Cheng Tan, Stan Z. Li: **SemiReward: A General Reward Model for Semi-supervised Learning.**
-     Maryam Toloubidokhti, Yubo Ye, Ryan Missel, Xiajun Jiang, Nilesh Kumar, Ruby Shrestha, Linwei Wang: **DATS: Difficulty-Aware Task Sampler for Meta-Learning Physics-Informed Neural Networks.**
-     Xi Weng, Yunhao Ni, Tengwei Song, Jie Luo, Rao Muhammad Anwer, Salman Khan, Fahad Shahbaz Khan, Lei Huang: **Modulate Your Spectrum in Self-Supervised Learning.**
-     Roger Creus Castanyer, Joshua Romoff, Glen Berseth: **Improving Intrinsic Exploration by Creating Stationary Objectives.**
-     Tri Dao: **FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning.**
-     Hanxun Huang, Ricardo J. G. B. Campello, Sarah Monazam Erfani, Xingjun Ma, Michael E. Houle, James Bailey: **LDReg: Local Dimensionality Regularized Self-Supervised Learning.**
-     Pangpang Liu, Yichuan Zhao: **Empirical Likelihood for Fair Classification.**
-     Jaroslaw Blasiok, Preetum Nakkiran: **Smooth ECE: Principled Reliability Diagrams via Kernel Smoothing.**
-     Sherry Yang, Yilun Du, Bo Dai, Dale Schuurmans, Joshua B. Tenenbaum, Pieter Abbeel: **Probabilistic Adaptation of Black-Box Text-to-Video Models.**
-     Zihan Zhang, Jason D. Lee, Yuxin Chen, Simon Shaolei Du: **Horizon-Free Regret for Linear Markov Decision Processes.**
-     Xiaxia Wang, David Jaime Tena Cucala, Bernardo Cuenca Grau, Ian Horrocks: **Faithful Rule Extraction for Differentiable Rule Learning Models.**
-     Frank Cole, Yulong Lu: **Score-based generative models break the curse of dimensionality in learning a family of sub-Gaussian distributions.**
-     Sherry Yang, KwangHwan Cho, Amil Merchant, Pieter Abbeel, Dale Schuurmans, Igor Mordatch, Ekin Dogus Cubuk: **Scalable Diffusion for Materials Generation.**



Qinyu Zhao, Ming Xu, Kartik Gupta, Akshay Asthana, Liang Zheng, Stephen Gould:
Towards Optimal Feature-Shaping Methods for Out-of-Distribution Detection.



Celine Lee, Abdulrahman Mahmoud, Michal Kurek, Simone Campanoni, David Brooks, Stephen Chong, Gu-Yeon Wei, Alexander M. Rush:
Guess & Sketch: Language Model Guided Transpilation.



Martin Klissarov, Pierluca D'Oro, Shagun Sodhani, Roberta Raileanu, Pierre-Luc Bacon, Pascal Vincent, Amy Zhang, Mikael Henaff:
Motif: Intrinsic Motivation from Artificial Intelligence Feedback.



Xingyu Zhou, Sayak Ray Chowdhury:
On Differentially Private Federated Linear Contextual Bandits.



Chuan Guo, Yuxuan Mu, Xinxin Zuo, Peng Dai, Youliang Yan, Juwei Lu, Li Cheng:
Generative Human Motion Stylization in Latent Space.



Zhongxia Yan, Cathy Wu:
Neural Neighborhood Search for Multi-agent Path Finding.



Erik Jones, Hamid Palangi, Cláisse Simões, Varun Chandrasekaran, Subhabrata Mukherjee, Arindam Mitra, Ahmed Hassan Awadallah, Ece Kamar:
Teaching Language Models to Hallucinate Less with Synthetic Tasks.



Ziqi Wang, Le Hou, Tianjian Lu, Yuexin Wu, Yunxuan Li, Hongkun Yu, Heng Ji:
Enabling Lanuguage Models to Implicitly Learn Self-Improvement.



Vladislav Lalin, Sherin Muckatira, Namrata Shivagunde, Anna Rumshisky:
ReLoRA: High-Rank Training Through Low-Rank Updates.



Andi Peng, Ilia Sucholutsky, Belinda Z. Li, Theodore R. Sumers, Thomas L. Griffiths, Jacob Andreas, Julie Shah:
Learning with Language-Guided State Abstractions.



Qiying Yu, Yudi Zhang, Yuyan Ni, Shikun Feng, Yanyan Lan, Hao Zhou, Jingjing Liu:
Multimodal Molecular Pretraining via Modality Blending.



Nicholas Corrado, Josiah P. Hanna:
Understanding when Dynamics-Invariant Data Augmentations Benefit Model-free Reinforcement Learning Updates.



Yuandong Tian, Yiping Wang, Zhenyu Zhang, Beidi Chen, Simon Shaolei Du:
JoMA: Demystifying Multilayer Transformers via Joint Dynamics of MLP and Attention.



Ziping Xu, Zifan Xu, Runxuan Jiang, Peter Stone, Ambuj Tewari:
Sample Efficient Myopic Exploration Through Multitask Reinforcement Learning with Diverse Tasks.



Paul Pu Liang, Chun Kai Ling, Yun Cheng, Alexander Obolenskiy, Yudong Liu, Rohan Pandey, Alex Wilf, Louis-Philippe Morency, Russ Salakhutdinov:
Multimodal Learning Without Labeled Multimodal Data: Guarantees and Applications.



Nishant Yadav, Nicholas Monath, Manzil Zaheer, Rob Fergus, Andrew McCallum:
Adaptive Retrieval and Scalable Indexing for k-NN Search with Cross-Encoders.



Pratyusha Sharma, Jordan T. Ash, Dipendra Misra:
The Truth is in There: Improving Reasoning in Language Models with Layer-Selective Rank Reduction.



Aochuan Chen, Yimeng Zhang, Jinghan Jia, James Diffenderfer, Konstantinos Parasyris, Jiancheng Liu, Yihua Zhang, Zheng Zhang, Bhavya Kailkhura, Sijia Liu:
DeepZero: Scaling Up Zeroth-Order Optimization for Deep Model Training.



Lifeng Shen, Weiyu Chen, James T. Kwok:
Multi-Resolution Diffusion Models for Time Series Forecasting.



Ibrahim Alabdulmohsin, Xiao Wang, Andreas Peter Steiner, Priya Goyal, Alexander D'Amour, Xiaohua Zhai:
CLIP the Bias: How Useful is Balancing Data in Multimodal Learning?



Yangming Li, Mihaela van der Schaar:
On Error Propagation of Diffusion Models.



Samuel Sokota, Gabriele Farina, David J. Wu, Hengyuan Hu, Kevin A. Wang, J. Zico Kolter, Noam Brown:
The Update-Equivalence Framework for Decision-Time Planning.



Yangming Li, Boris van Breugel, Mihaela van der Schaar:
Soft Mixture Denoising: Beyond the Expressive Bottleneck of Diffusion Models.



Badih Ghazi, Yangsibo Huang, Pritish Kamath, Ravi Kumar, Pasin Manurangsi, Chiyuan Zhang:
LabelDP-Pro: Learning with Label Differential Privacy via Projections.



Rohan Sharma, Kaiyi Ji, Zhiqiang Xu, Changyou Chen:
AUC-CL: A Batchsize-Robust Framework for Self-Supervised Contrastive Representation Learning.



Nayoung Lee, Kartik Sreenivasan, Jason D. Lee, Kangwook Lee, Dimitris Papailiopoulos:
Teaching Arithmetic to Small Transformers.



Tianyu Du, Luca Melis, Ting Wang:
ReMasker: Imputing Tabular Data with Masked Autoencoding.



- Xufeng Cai, Ahmet Alacaoglu, Jelena Diakonikolas:
Variance Reduced Halpern Iteration for Finite-Sum Monotone Inclusions.
- Mehdi Fatemi, Sindhu C. M. Gowda:
A Dynamical View of the Question of Why.
- Tongzhou Mu, Minghua Liu, Hao Su:
DrS: Learning Reusable Dense Rewards for Multi-Stage Tasks.
- Tongxin Yin, Jean-Francois Ton, Ruocheng Guo, Yuanshun Yao, Mingyan Liu, Yang Liu:
Fair Classifiers that Abstain without Harm.
- Yaning Jia, Chunhui Zhang, Soroush Vosoughi:
Aligning Relational Learning with Lipschitz Fairness.
- Yuan Gong , Hongyin Luo, Alexander H. Liu, Leonid Karlinsky, James R. Glass:
Listen, Think, and Understand.
- Shengyao Lu, Keith G. Mills, Jiao He, Bang Liu, Di Niu:
GOAt: Explaining Graph Neural Networks via Graph Output Attribution.
- Adam Lechowicz, Rik Sengupta, Bo Sun, Shahin Kamali, Mohammad Hajiesmaili:
Time Fairness in Online Knapsack Problems.
- Ziheng Chen, Yue Song, Yunmei Liu, Nicu Sebe:
A Lie Group Approach to Riemannian Batch Normalization.
- Chen Qiu, Xingyu Li, Chaithanya Kumar Mummadi, Madan Ravi Ganesh, Zhenzhen Li, Lu Peng, Wan-Yi Lin:
Federated Text-driven Prompt Generation for Vision-Language Models.
- Ziyi Chen, Yi Zhou, Heng Huang:
On the Hardness of Constrained Cooperative Multi-Agent Reinforcement Learning.
- Hanmin Li, Avetik G. Karagulyan, Peter Richtárik:
Det-CGD: Compressed Gradient Descent with Matrix Stepsizes for Non-Convex Optimization.
- Yixin Fang, Faming Liang:
Causal-StoNet: Causal Inference for High-Dimensional Complex Data.
- MinGyu Choi, Changhee Lee:
Conditional Information Bottleneck Approach for Time Series Imputation.
- Katie Kang, Amrit Sethuraman, Claire J. Tomlin, Sergey Levine:
Deep Neural Networks Tend To Extrapolate Predictably.
- Kevin Yang, Dan Klein, Asli Celikyilmaz, Nanyun Peng, Yuandong Tian:
RLCD: Reinforcement Learning from Contrastive Distillation for LM Alignment.
- Zhiwei Tang, Dmitry Rybin, Tsung-Hui Chang:
Zeroth-Order Optimization Meets Human Feedback: Provable Learning via Ranking Oracles.
- Julius Kunze, Daniel Severo, Giulio Zani, Jan-Willem van de Meent, James Townsend:
Entropy Coding of Unordered Data Structures.
- Atharva Sehgal, Arya Grayeli, Jennifer J. Sun, Swarat Chaudhuri:
Neurosymbolic Grounding for Compositional World Models.
- Kezhi Kong, Jian Zhang, Zhengyuan Shen, Balasubramaniam Srinivasan, Chuan Lei, Christos Faloutsos, Huzeifa Rangwala, George Karypis:
OpenTab: Advancing Large Language Models as Open-domain Table Reasoners.
- Anna Bair, Hongxu Yin, Maying Shen, Pavlo Molchanov, José M. Álvarez:
Adaptive Sharpness-Aware Pruning for Robust Sparse Networks.
- Mara Finkelstein, Markus Freitag:
MBR and QE Finetuning: Training-time Distillation of the Best and Most Expensive Decoding Methods.
- Yavuz Faruk Bakman, Duygu Nur Yaldiz, Yahya H. Ezzeldin, Salman Avestimehr:
Federated Orthogonal Training: Mitigating Global Catastrophic Forgetting in Continual Federated Learning.
- Sina Alemdamghan, Josue Casco-Rodriguez, Lorenzo Luzi, Ahmed Imtiaz Humayun, Hossein Babaei, Daniel LeJeune, Ali Siahkoohi, Richard G. Baraniuk:
Self-Consuming Generative Models Go MAD.
- Hila Chefer, Oran Lang, Mor Geva, Volodymyr Polosukhin, Assaf Shocher, Michal Irani, Inbar Mosseri, Lior Wolf:
The Hidden Language of Diffusion Models.
- Siddarth Venkatraman, Shivesh Khaitan, Ravi Tej Akella, John Dolan, Jeff Schneider, Glen Berseth:
Reasoning with Latent Diffusion in Offline Reinforcement Learning.
- Yuzhou Gu, Zhao Song, Junze Yin, Lichen Zhang:
Low Rank Matrix Completion via Robust Alternating Minimization in Nearly Linear Time.
- Sohyun An, Hayeon Lee, Jaehyeong Jo, Seanie Lee, Sung Ju Hwang:
DiffusionNAG: Predictor-guided Neural Architecture Generation with Diffusion Models.
- Longkang Li, Ignavier Ng, Gongxu Luo, Biwei Huang, Guangyi Chen, Tongliang Liu, Bin Gu, Kun Zhang:
Federated Causal Discovery from Heterogeneous Data.



- Hongzhi Wen, Wenzhuo Tang, Xinnan Dai, Jiayuan Ding, Wei Jin , Yuying Xie, Jiliang Tang:
CellPLM: Pre-training of Cell Language Model Beyond Single Cells.
- Edwin Zhang, Yujie Lu, Shinda Huang, William Yang Wang, Amy Zhang:
Language Control Diffusion: Efficiently Scaling through Space, Time, and Tasks.
- Jiachen Sun, Haizhong Zheng, Qingzhao Zhang, Atul Prakash , Zuoqing Mao, Chaowei Xiao:
CALICO: Self-Supervised Camera-LiDAR Contrastive Pre-training for BEV Perception.
- Wuyang Chen, Junru Wu, Zhangyang Wang, Boris Hanin:
Principled Architecture-aware Scaling of Hyperparameters.
- Xuandong Zhao, Prabhanjan Vijendra Ananth, Lei Li, Yu-Xiang Wang:
Provable Robust Watermarking for AI-Generated Text.
- Jiarui Lu, Boztas Zhong, Zuobai Zhang, Jian Tang:
Str2Str: A Score-based Framework for Zero-shot Protein Conformation Sampling.
- Liyiming Ke, Yunchu Zhang, Abhay Deshpande, Siddhartha S. Srinivasa, Abhishek Gupta:
CCIL: Continuity-Based Data Augmentation for Corrective Imitation Learning.
- Kaustubh Sridhar, Souradeep Dutta, Dinesh Jayaraman, James Weimer, Insup Lee:
Memory-Consistent Neural Networks for Imitation Learning.
- Luciano Dyballa, Samuel Lang, Alexandra Haslund-Gourley, Eviatar Yemini, Steven W. Zucker:
Learning dynamic representations of the functional connectome in neurobiological networks.
- Hsiao-Ru Pan, Bernhard Schölkopf:
Skill or Luck? Return Decomposition via Advantage Functions.
- Xiaolin Sun, Zizhan Zheng:
Belief-Enriched Pessimistic Q-Learning against Adversarial State Perturbations.
- Yue Wu, Xuan Tang, Tom M. Mitchell, Yuanzhi Li:
SmartPlay : A Benchmark for LLMs as Intelligent Agents.
- Debo Cheng, Ziqi Xu, Jiuyong Li, Lin Liu, Jixue Liu, Thuc Duy Le:
Conditional Instrumental Variable Regression with Representation Learning for Causal Inference.
- Apratim Bhattacharyya, Sunny Panchal, Reza Pourreza, Mingu Lee, Pulkit Madan, Roland Memisevic:
Look, Remember and Reason: Grounded Reasoning in Videos with Language Models.
- Shengcao Cao, Jiuxiang Gu, Jason Kuen, Hao Tan, Ruiyi Zhang, Handong Zhao, Ani Nenkova, Liangyan Gui, Tong Sun, Yu-Xiong Wang:
SOHES: Self-supervised Open-world Hierarchical Entity Segmentation.
- Yi-Lun Liao, Brandon M. Wood, Abhishek Das, Tess E. Smidt:
EquiformerV2: Improved Equivariant Transformer for Scaling to Higher-Degree Representations.
- Harsh Chaudhari, Giorgio Severi, Alina Oprea, Jonathan R. Ullman:
Chameleon: Increasing Label-Only Membership Leakage with Adaptive Poisoning.
- Yongtao Wu, Fanghui Liu, Carl-Johann Simon-Gabriel, Grigoris Chrysos, Volkan Cevher:
Robust NAS under adversarial training: benchmark, theory, and beyond.
- Victor Quach, Adam Fisch, Tal Schuster, Adam Yala, Jae Ho Sohn, Tommi S. Jaakkola, Regina Barzilay:
Conformal Language Modeling.
- Yu Meng , Jitin Krishnan, Sinong Wang, Qifan Wang, Yunling Mao, Han Fang, Marjan Ghazvininejad, Jiawei Han, Luke Zettlemoyer:
Representation Deficiency in Masked Language Modeling.
- Peihao Wang, Shenghao Yang, Shu Li, Zhangyang Wang, Pan Li:
Polynomial Width is Sufficient for Set Representation with High-dimensional Features.
- Xu Zheng, Tianchun Wang, Wei Cheng, Aitian Ma, Haifeng Chen, Mo Sha, Dongsheng Luo:
Parametric Augmentation for Time Series Contrastive Learning.
- Michal Zajac, Tinne Tuytelaars, Gido M. van de Ven:
Prediction Error-based Classification for Class-Incremental Learning.
- Moritz Hardt, Yu Sun:
Test-Time Training on Nearest Neighbors for Large Language Models.
- Sanghyuk Chun:
Improved Probabilistic Image-Text Representations.
- Yulei Niu, Wenliang Guo, Long Chen, Xudong Lin, Shih-Fu Chang:
SCHEMA: State CHangEs MATter for Procedure Planning in Instructional Videos.
- Tinghao Xie, Xiangyu Qi, Ping He, Yiming Li, Jiachen T. Wang, Prateek Mittal:
BaDExpert: Extracting Backdoor Functionality for Accurate Backdoor Input Detection.
- Theo X. Olausson, Jeevana Priya Inala, Chenglong Wang, Jianfeng Gao, Armando Solar-Lezama:
Is Self-Repair a Silver Bullet for Code Generation?
- Juncai He, Xinliang Liu, Jinchao Xu:
MgNO: Efficient Parameterization of Linear Operators via Multigrid.



- Debangshu Banerjee, Avaljot Singh, Gagandeep Singh:
Interpreting Robustness Proofs of Deep Neural Networks.
- Xinwei Zhang, Zhiqi Bu, Steven Wu, Mingyi Hong:
Differentially Private SGD Without Clipping Bias: An Error-Feedback Approach.
- Jong-Hoon Ahn, Akshay Vashist:
A Linear Algebraic Framework for Counterfactual Generation.
- Zhihan Zhou, Yanrong Ji, Weijian Li, Pratik Dutta, Ramana V. Davuluri, Han Liu:
DNABERT-2: Efficient Foundation Model and Benchmark For Multi-Species Genomes.
- Huangjie Zheng, Zhendong Wang, Jianbo Yuan, Guanghan Ning, Pengcheng He, Quanzeng You, Hongxia Yang, Mingyuan Zhou:
Learning Stackable and Skippable LEGO Bricks for Efficient, Reconfigurable, and Variable-Resolution Diffusion Modeling.
- Yuning You, Ruida Zhou, Jiwoong Park, Haotian Xu, Chao Tian, Zhangyang Wang, Yang Shen:
Latent 3D Graph Diffusion.
- Alekса Sukovic, Goran Radanovic:
Reward Design for Justifiable Sequential Decision-Making.
- Mucong Ding, Bang An, Yuancheng Xu, Anirudh Satheesh, Furong Huang:
SAFLEX: Self-Adaptive Augmentation via Feature Label Extrapolation.
- Gabriel Grand, Lionel Wong, Matthew Bowers, Theo X. Olausson, Muxin Liu, Joshua B. Tenenbaum, Jacob Andreas:
LILO: Learning Interpretable Libraries by Compressing and Documenting Code.
- Jitao Gu, Shuangfei Zhai, Yizhe Zhang, Joshua M. Susskind, Navdeep Jaitly:
Matryoshka Diffusion Models.
- Arjit Sehanobish, Krzysztof Marcin Choromanski, Yunfan Zhao, Kumar Avinava Dubey, Valerii Likhoshevstov:
Scalable Neural Network Kernels.
- Nico Daheim, Thomas Möllenhoff, Edoardo M. Ponti, Iryna Gurevych, Mohammad Emtiyaz Khan:
Model Merging by Uncertainty-Based Gradient Matching.
- Cristian Meo, Louis Mahon, Anirudh Goyal, Justin Dauwels:
αTC-VAE: On the relationship between Disentanglement and Diversity.
- Yuyang Hu, Mauricio Delbracio, Peyman Milanfar, Ulugbek Kamilov:
A Restoration Network as an Implicit Prior.
- Tingchen Fu, Lemao Liu, Deng Cai, Guoping Huang, Shuming Shi, Rui Yan:
The Reasonableness Behind Unreasonable Translation Capability of Large Language Model.
- Frank Shih, Faming Liang:
Fast Value Tracking for Deep Reinforcement Learning.
- Krishna Acharya, Eshwar Ram Arunachaleswaran, Sampath Kannan, Aaron Roth, Juba Ziani:
Oracle Efficient Algorithms for Groupwise Regret.
- Zijun Wu, Yongkang Wu, Lili Mou:
Zero-Shot Continuous Prompt Transfer: Generalizing Task Semantics Across Language Models.
- Sepehr Dehdashtian, Lan Wang, Vishnu Boddeti:
FairerCLIP: Debiasing CLIP's Zero-Shot Predictions using Functions in RKHSs.
- Guan-Horng Liu, Yaron Lipman, Maximilian Nickel, Brian Karrer, Evangelos A. Theodorou, Ricky T. Q. Chen:
Generalized Schrödinger Bridge Matching.
- Sihyun Yu, Weili Nie, De-An Huang, Boyi Li, Jinwoo Shin, Anima Anandkumar:
Efficient Video Diffusion Models via Content-Frame Motion-Latent Decomposition.
- Zohar Rimon, Tom Jurgenson, Orr Krupnik, Gilad Adler, Aviv Tamar:
MAMBA: an Effective World Model Approach for Meta-Reinforcement Learning.
- Blake Bordelon, Lorenzo Noci, Mufan Bill Li, Boris Hanin, Cengiz Pehlevan:
Depthwise Hyperparameter Transfer in Residual Networks: Dynamics and Scaling Limit.
- Yifan Jiang, Hao Tang, Jen-Hao Rick Chang, Liangchen Song, Zhangyang Wang, Liangliang Cao:
Efficient-3Dim: Learning a Generalizable Single-image Novel-view Synthesizer in One Day.
- Dejiao Zhang, Wasi Uddin Ahmad, Ming Tan, Hantian Ding, Ramesh Nallapati, Dan Roth, Xiaofei Ma, Bing Xiang:
Code Representation Learning at Scale.
- Anthony Bardou, Patrick Thiran, Thomas Begin:
Relaxing the Additivity Constraints in Decentralized No-Regret High-Dimensional Bayesian Optimization.
- Ajay Kumar Jaiswal, Zhe Gan, Xianzhi Du, Bowen Zhang, Zhangyang Wang, Yinfei Yang:
Compressing LLMs: The Truth is Rarely Pure and Never Simple.
- Subha Maity, Mayank Agarwal, Mikhail Yurochkin, Yuekai Sun:
An Investigation of Representation and Allocation Harms in Contrastive Learning.
- Zhiqian Lan, Yuxuan Jiang, Yao Mu, Chen Chen, Shengbo Eben Li:
SEPT: Towards Efficient Scene Representation Learning for Motion Prediction.



- Andreas Bergmeister, Karolis Martinkus, Nathanaël Perraudin, Roger Wattenhofer:
Efficient and Scalable Graph Generation through Iterative Local Expansion.
- Simon Schug, Seijin Kobayashi, Yassir Akram, Maciej Wolczyk, Alexandra Proca, Johannes von Oswald, Razvan Pascanu, João Sacramento, Angelika Steger:
Discovering modular solutions that generalize compositionally.
- Xuan Li, Zhanke Zhou, Jiangchao Yao, Yu Rong, Lu Zhang, Bo Han:
Neural Atoms: Propagating Long-range Interaction in Molecular Graphs through Efficient Communication Channel.
- Raman Dutt, Ondrej Bohdal, Sotirios A. Tsaftaris, Timothy M. Hospedales:
FairTune: Optimizing Parameter Efficient Fine Tuning for Fairness in Medical Image Analysis.
- Mengkang Hu, Yao Mu, Xinmiao Yu, Mingyu Ding, Shiguang Wu, Wenqi Shao, Qiguang Chen, Bin Wang, Yu Qiao, Ping Luo:
Tree-Planner: Efficient Close-loop Task Planning with Large Language Models.
- Chen Chen, Ruizhe Li, Yuchen Hu, Sabato Marco Siniscalchi, Pin-Yu Chen, Engsiong Chng, Chao-Han Huck Yang:
It's Never Too Late: Fusing Acoustic Information into Large Language Models for Automatic Speech Recognition.
- Milad Aghajohari, Juan Agustin Duque, Tim Cooijmans, Aaron C. Courville:
LOQA: Learning with Opponent Q-Learning Awareness.
- Ethan Baron, Itamar Zimerman, Lior Wolf:
A 2-Dimensional State Space Layer for Spatial Inductive Bias.
- William F. Whitney, Tatiana Lopez-Guevara, Tobias Pfaff, Yulia Rubanova, Thomas Kipf, Kim Stachenfeld, Kelsey R. Allen:
Learning 3D Particle-based Simulators from RGB-D Videos.
- Jiacheng Lin, Meng Xu, Zhihua Xiong, Huangang Wang:
CAMBranch: Contrastive Learning with Augmented MILPs for Branching.
- Sunghyeon Woo, Sunwoo Lee, Dongsuk Jeon:
ALAM: Averaged Low-Precision Activation for Memory-Efficient Training of Transformer Models.
- Maresa Schröder, Dennis Frauen, Stefan Feuerriegel:
Causal Fairness under Unobserved Confounding: A Neural Sensitivity Framework.
- Jinyang Jiang, Zeliang Zhang, Chenliang Xu, Zhaofei Yu, Yijie Peng:
One Forward is Enough for Neural Network Training via Likelihood Ratio Method.
- Xiangyu Zeng, Jie Lin, Piao Hu, Ruizheng Huang, Zhicheng Zhang:
A Framework for Inference Inspired by Human Memory Mechanisms.
- Diego Martinez-Taboada, Edward Kennedy:
Counterfactual Density Estimation using Kernel Stein Discrepancies.
- Vishakh Padmakumar, He He:
Does Writing with Language Models Reduce Content Diversity?
- Soroush H. Zargarbashi, Aleksandar Bojchevski:
Conformal Inductive Graph Neural Networks.
- Zhihang Yuan, Yuzhang Shang, Zhen Dong:
PB-LLM: Partially Binarized Large Language Models.
- David A. R. Robin, Kevin Scaman, Marc Lelarge:
Random Sparse Lifts: Construction, Analysis and Convergence of finite sparse networks.
- Mahan Fathi, Clement Gehring, Jonathan Pilault, David Kanaa, Pierre-Luc Bacon, Ross Goroshin:
Course Correcting Koopman Representations.
- Fran Jelenic, Josip Jukic, Martin Tutek, Mate Puljiz, Jan Snajder:
Out-of-Distribution Detection by Leveraging Between-Layer Transformation Smoothness.
- Darshan Patil, Janarthanan Rajendran, Glen Berseth, Sarath Chandar:
Intelligent Switching for Reset-Free RL.
- Yumeng Li, Margret Keuper, Dan Zhang, Anna Khoreva:
Adversarial Supervision Makes Layout-to-Image Diffusion Models Thrive.
- Ted Zadouri, Ahmet Üstün, Arash Ahmadian, Beyza Ermis, Acyr Locatelli, Sara Hooker:
Pushing Mixture of Experts to the Limit: Extremely Parameter Efficient MoE for Instruction Tuning.
- Zhanke Zhou, Yongqi Zhang, Jiangchao Yao, Quanming Yao, Bo Han:
Less is More: One-shot Subgraph Reasoning on Large-scale Knowledge Graphs.
- Zhirui Chen, P. N. Karthik, Yeow Meng Chee, Vincent Y. F. Tan:
Fixed-Budget Differentially Private Best Arm Identification.
- Shahar Lutati, Eliya Nachmani, Lior Wolf:
Separate and Diffuse: Using a Pretrained Diffusion Model for Better Source Separation.
- Muthu Chidambaram, Rong Ge:
On the Limitations of Temperature Scaling for Distributions with Overlaps.
- Seungjae Shin, HeeSun Bae, Byeonghu Na, Yoon-Yeong Kim, Il-Chul Moon:
Unknown Domain Inconsistency Minimization for Domain Generalization.



- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, Bryan Hooi:
Can LLMs Express Their Uncertainty? An Empirical Evaluation of Confidence Elicitation in LLMs.
- Jonathan Brokman, Roy Betser, Rotem Turjeman, Tom Berkov, Ido Cohen, Guy Gilboa:
Enhancing Neural Training via a Correlated Dynamics Model.
- Simon N. Segert:
Flat Minima in Linear Estimation and an Extended Gauss Markov Theorem.
- Yunhui Jang, Seul Lee, Sungsoo Ahn:
A Simple and Scalable Representation for Graph Generation.
- Weihao Tan, Wentao Zhang, Shanqi Liu, Longtao Zheng, Xinrun Wang, Bo An:
True Knowledge Comes from Practice: Aligning Large Language Models with Embodied Environments via Reinforcement Learning.
- Zihao Tang, Zheqi Lv, Shengyu Zhang, Yifan Zhou, Xinyu Duan, Fei Wu, Kun Kuang:
AuG-KD: Anchor-Based Mixup Generation for Out-of-Domain Knowledge Distillation.
- Tian Jin, Nolan Clement, Xin Dong, Vaishnav Nagarajan, Michael Carbin, Jonathan Ragan-Kelley, Gintare Karolina Dziugaite:
The Cost of Scaling Down Large Language Models: Reducing Model Size Affects Memory before In-context Learning.
- Di Wu, Jun Bai, Yiliao Song, Junjun Chen, Wei Zhou, Yong Xiang, Atul Sajjanhar:
FedInverse: Evaluating Privacy Leakage in Federated Learning.
- Yukai Shi, Jianan Wang, He Cao, Boshi Tang, Xianbiao Qi, Tianyu Yang, Yukun Huang, Shilong Liu, Lei Zhang, Heung-Yeung Shum:
TOSS: High-quality Text-guided Novel View Synthesis from a Single Image.
- Xinyu Tang, Richard Shin, Huseyin A. Inan, Andre Manoel, Fatemehsadat Mireshghallah, Zinan Lin, Sivakanth Gopi, Janardhan Kulkarni, Robert Sim:
Privacy-Preserving In-Context Learning with Differentially Private Few-Shot Generation.
- Noa Rubin, Inbar Seroussi, Zohar Ringel:
Grokking as a First Order Phase Transition in Two Layer Networks.
- Mang Ning, Mingxiao Li, Jianlin Su, Albert Ali Salah, Itir Önal Ertugrul:
Elucidating the Exposure Bias in Diffusion Models.
- Rui Pan, Yuxing Liu, Xiaoyu Wang, Tong Zhang:
Accelerated Convergence of Stochastic Heavy Ball Method under Anisotropic Gradient Noise.
- Erik J. Bekkers, Sharvaree P. Vadgama, Rob Hesselink, Putri A. van der Linden, David W. Romero:
Fast, Expressive SE(n) Equivariant Networks through Weight-Sharing in Position-Orientation Space.
- Jonas Seng, Matej Zecevic, Devendra Singh Dhami, Kristian Kersting:
Learning Large DAGs is Harder than you Think: Many Losses are Minimal for the Wrong DAG.
- Jingyang Zhang, Shiwei Li, Yuanxun Lu, Tian Fang, David McKinnon, Yanghai Tsin, Long Quan, Yao Yao:
JointNet: Extending Text-to-Image Diffusion for Dense Distribution Modeling.
- Rhys Gould, Euan Ong, George Ogden, Arthur Conmy:
Successor Heads: Recurring, Interpretable Attention Heads In The Wild.
- Kai-Po Chang, Chi-Pin Huang, Wei-Yuan Cheng, Fu-En Yang, Chien-Yi Wang, Yung-Hsuan Lai, Yu-Chiang Frank Wang:
RAPPER: Reinforced Rationale-Prompted Paradigm for Natural Language Explanation in Visual Question Answering.
- Ivan Butakov, Aleksander Tolmachev, Sofia Malanchuk, Anna Neopryatnaya, Alexey A. Frolov, Kirill V. Andreev:
Information Bottleneck Analysis of Deep Neural Networks via Lossy Compression.
- Fanqi Wan, Xinting Huang, Deng Cai, Xiaojun Quan, Wei Bi, Shuming Shi:
Knowledge Fusion of Large Language Models.
- Mingkun Yang, Ran Zhu, Qing Wang, Jie Yang:
FedTrans: Client-Transparent Utility Estimation for Robust Federated Learning.
- Noa Cohen, Hila Manor, Yuval Bahat, Tomer Michaeli:
From Posterior Sampling to Meaningful Diversity in Image Restoration.
- Dennis Frauen, Fergus Imrie, Alicia Curth, Valentyn Melnychuk, Stefan Feuerriegel, Mihaela van der Schaar:
A Neural Framework for Generalized Causal Sensitivity Analysis.
- Geraud Nangue Tasse, Devon Jarvis, Steven James, Benjamin Rosman:
Skill Machines: Temporal Logic Skill Composition in Reinforcement Learning.
- Zijian Feng, Han Zhang Zhou, Zixiao Zhu, Junlang Qian, Kezhi Mao:
Unveiling and Manipulating Prompt Influence in Large Language Models.
- Hanni Cheng, Ya Cong, Weihao Jiang, Shiliang Pu:
Learning to solve Class-Constrained Bin Packing Problems via Encoder-Decoder Model.
- Assaf Shocher, Amil David, Yossi Gandelsman, Inbar Mosseri, Michael Rubinstein, Alexei A. Efros:
Idempotent Generative Network.



- Driton Salihu, Adam Misik, Yuankai Wu, Constantin Patsch, Fabián Seguel, Eckehard G. Steinbach:**
DeepSPF: Spherical SO(3)-Equivariant Patches for Scan-to-CAD Estimation.
- Qianxu Wang, Haotong Zhang, Congyue Deng, Yang You, Hao Dong, Yixin Zhu, Leonidas J. Guibas:**
SparseDFF: Sparse-View Feature Distillation for One-Shot Dexterous Manipulation.
- Shangyu Wu, Ying Xiong, Yufei Cui, Xue Liu, Buzhou Tang, Tei-Wei Kuo, Chun Jason Xue:**
ReFusion: Improving Natural Language Understanding with Computation-Efficient Retrieval Representation Fusion.
- Haozhao Wang, Haoran Xu, Yichen Li, Yuan Xu** **Ruixuan Li, Tianwei Zhang:**
FedCDA: Federated Learning with Cross-rounds Divergence-aware Aggregation.
- Tim Franzmeyer, Edith Elkind, Philip Torr, Jakob Nicolaus Foerster, João F. Henriques:**
Select to Perfect: Imitating desired behavior from large multi-agent data.
- Sangyu Han, Yearim Kim, Nojun Kwak:**
Respect the model: Fine-grained and Robust Explanation with Sharing Ratio Decomposition.
- Zhongqi Yue, Jiankun Wang, Qianru Sun, Lei Ji, Eric I-Chao Chang, Hanwang Zhang:**
Exploring Diffusion Time-steps for Unsupervised Representation Learning.
- Weiyu Li, Rui Chen, Xuelin Chen, Ping Tan:**
SweetDreamer: Aligning Geometric Priors in 2D diffusion for Consistent Text-to-3D.
- Jielong Yan, Yifan Feng, Shihui Ying, Yue Gao:**
Hypergraph Dynamic System.
- Mahdi Kallel, Debabrota Basu, Riad Akroud, Carlo D'Eramo:**
Augmented Bayesian Policy Search.
- Quan Sun, Qiying Yu, Yufeng Cui, Fan Zhang, Xiaosong Zhang, Yueze Wang, Hongcheng Gao, Jingjing Liu, Tiejun Huang, Xinlong Wang:**
Emu: Generative Pretraining in Multimodality.
- HeeSun Bae, Seungjae Shin, Byeonghu Na, Il-Chul Moon:**
Dirichlet-based Per-Sample Weighting by Transition Matrix for Noisy Label Learning.
- Christian Koke, Daniel Cremers:**
HoloNets: Spectral Convolutions do extend to Directed Graphs.
- Yifu Yuan, Jianye Hao, Yi Ma, Zibin Dong, Hebin Liang, Jinyi Liu, Zhixin Feng, Kai Zhao, Yan Zheng:**
Uni-RLHF: Universal Platform and Benchmark Suite for Reinforcement Learning with Diverse Human Feedback.
- Sameera Ramasinghe, Violetta Shevchenko, Gil Avraham, Hisham Husain, Anton van den Hengel:**
Improving the Convergence of Dynamic NeRFs via Optimal Transport.
- Haoran Deng, Yang Yang, Jiahe Li, Cheng Chen, Weihao Jiang, Shiliang Pu:**
Fast Updating Truncated SVD for Representation Learning with Sparse Matrices.
- Juan Rocamonde, Victoriano Montesinos, Elvis Nava, Ethan Perez, David Lindner:**
Vision-Language Models are Zero-Shot Reward Models for Reinforcement Learning.
- Michal Geyer, Omer Bar-Tal, Shai Bagon, Tali Dekel:**
TokenFlow: Consistent Diffusion Features for Consistent Video Editing.
- Yukun Huang, Jianan Wang, Yukai Shi, Boshi Tang, Xianbiao Qi, Lei Zhang:**
DreamTime: An Improved Optimization Strategy for Diffusion-Guided 3D Generation.
- Xue Wang, Tian Zhou, Qingsong Wen, Jinyang Gao, Bolin Ding, Rong Jin:**
CARD: Channel Aligned Robust Blend Transformer for Time Series Forecasting.
- Yilong Xu, Yang Liu, Hao Sun:**
Reinforcement Symbolic Regression Machine.
- Mengxi Ya, Yiming Li, Tao Dai, Bin Wang, Yong Jiang, Shu-Tao Xia:**
Towards Faithful XAI Evaluation via Generalization-Limited Backdoor Watermark.
- Rui Jiao, Wenbing Huang, Yu Liu, Deli Zhao, Yang Liu:**
Space Group Constrained Crystal Generation.
- Shengchao Hu, Li Shen, Ya Zhang, Dacheng Tao:**
Learning Multi-Agent Communication from Graph Modeling Perspective.
- Qihan Liu, Jianing Ye, Xiaoteng Ma, Jun Yang, Bin Liang, Chongjie Zhang:**
Efficient Multi-agent Reinforcement Learning by Planning.
- Mingyuan Sun, Donghao Zhang, Zongyuan Ge, Jiaxu Wang, Jia Li, Zheng Fang, Renjing Xu:**
EventRPG: Event Data Augmentation with Relevance Propagation Guidance.
- Dong Wei, Huaijiang Sun, Bin Li, Xiaoning Sun, Shengxiang Hu, Weiqing Li, Jianfeng Lu:**
NeRM: Learning Neural Representations for High-Framerate Human Motion Synthesis.
- Yang Jin, Kun Xu, Kun Xu, Liwei Chen, Chao Liao, Jianchao Tan, Quzhe Huang, Bin Chen, Chengru Song, Dai Meng, Di Zhang, Wenwu Ou, Kun Gai, Yadong Mu:**
Unified Language-Vision Pretraining in LLM with Dynamic Discrete Visual Tokenization.



- Ziteng Wang, Jianfei Chen, Jun Zhu:
Efficient Backpropagation with Variance Controlled Adaptive Sampling.
- Fabian Akkerman, Julius Luy, Wouter van Heeswijk, Maximilian Schiffer:
Dynamic Neighborhood Construction for Structured Large Discrete Action Spaces.
- Sheng Shen, Le Hou, Yanqi Zhou, Nan Du, Shayne Longpre, Jason Wei, Hyung Won Chung, Barret Zoph, William Fedus, Xinyun Chen, Tu Vu, Yuexin Wu, Wuyang Chen, Albert Webson, Yunxuan Li, Vincent Y. Zhao, Hongkun Yu, Kurt Keutzer, Trevor Darrell, Denny Zhou:
Mixture-of-Experts Meets Instruction Tuning: A Winning Combination for Large Language Models.
- Licheng Wen, Daocheng Fu, Xin Li, Xinyu Cai, Tao Ma, Pinlong Cai, Min Dou, Botian Shi, Liang He, Yu Qiao:
DiLu: A Knowledge-Driven Approach to Autonomous Driving with Large Language Models.
- Jinsung Jeon, Hyundong Jin, Jonghyun Choi, Sanghyun Hong, Dongeun Lee, Kookjin Lee, Noseong Park:
PAC-FNO: Parallel-Structured All-Component Fourier Neural Operators for Recognizing Low-Quality Images.
- Lifan Yuan, Yangyi Chen, Xingyao Wang, Yi Fung, Hao Peng, Heng Ji:
CRAFT: Customizing LLMs by Creating and Retrieving from Specialized Toolsets.
- Erdun Gao, Howard D. Bondell, Wei Huang, Mingming Gong:
A Variational Framework for Estimating Continuous Treatment Effects with Measurement Error.
- Xingchao Liu, Xiwen Zhang, Jianzhu Ma, Jian Peng, Qiang Liu:
InstaFlow: One Step is Enough for High-Quality Diffusion-Based Text-to-Image Generation.
- Zhen Yang, Ganggui Ding, Wen Wang, Hao Chen, Bohan Zhuang, Chunhua Shen:
Object-Aware Inversion and Reassembly for Image Editing.
- Yuhta Takida, Masaaki Imaizumi, Takashi Shibuya, Chieh-Hsin Lai, Toshimitsu Uesaka, Naoki Murata, Yuki Mitsufuji:
SAN: Inducing Metrizability of GAN with Discriminative Normalized Linear Layer.
- Hritik Bansal, John Dang, Aditya Grover:
Peering Through Preferences: Unraveling Feedback Acquisition for Aligning Large Language Models.
- Nuoya Xiong, Zhihan Liu, Zhaoran Wang, Zhuoran Yang:
Sample-Efficient Multi-Agent RL: An Optimization Perspective.
- Namjun Kim, Chanho Min, Sejun Park:
Minimum width for universal approximation using ReLU networks on compact domain.
- Dong Bok Lee, Seanie Lee, Joonho Ko, Kenji Kawaguchi, Juho Lee, Sung Ju Hwang:
Self-Supervised Dataset Distillation for Transfer Learning.
- Ge Yan, Hongxu Chen, Kaisen Pan, Junchi Yan:
Rethinking the symmetry-preserving circuits for constrained variational quantum algorithms.
- Lean Wang, Wenkai Yang, Deli Chen, Hao Zhou, Yankai Lin, Fandong Meng, Jie Zhou, Xu Sun:
Towards Codable Watermarking for Injecting Multi-Bits Information to LLMs.
- Ruocheng Wang, Eric Zelikman, Gabriel Poesia, Yewen Pu, Nick Haber, Noah D. Goodman:
Hypothesis Search: Inductive Reasoning with Language Models.
- Haozhe Ji, Pei Ke, Hongning Wang, Minlie Huang:
Language Model Decoding as Direct Metrics Optimization.
- Bowen Gao, Yinjun Jia, Yuanle Mo, Yuyan Ni, Wei-Ying Ma, Zhi-Ming Ma, Yanyan Lan:
Self-supervised Pocket Pretraining via Protein Fragment-Surroundings Alignment.
- Ke Wang, Houxing Ren, Aojun Zhou, Zimu Lu, Sichun Luo, Weikang Shi, Renrui Zhang, Linqi Song, Mingjie Zhan, Hongsheng Li:
MathCoder: Seamless Code Integration in LLMs for Enhanced Mathematical Reasoning.
- Aojun Zhou, Ke Wang, Zimu Lu, Weikang Shi, Sichun Luo, Zipeng Qin, Shaoqing Lu, Anya Jia, Linqi Song, Mingjie Zhan, Hongsheng Li:
Solving Challenging Math Word Problems Using GPT-4 Code Interpreter with Code-based Self-Verification.
- Yu-Ju Tsai, Yu-Lun Liu, Lu Qi, Kelvin C. K. Chan, Ming-Hsuan Yang:
Dual Associated Encoder for Face Restoration.
- Samar Khanna, Patrick Liu, Linqi Zhou, Chenlin Meng, Robin Rombach, Marshall Burke, David B. Lobell, Stefano Ermon:
DiffusionSat: A Generative Foundation Model for Satellite Imagery.
- Jingxiang Sun, Bo Zhang, Ruizhi Shao, Lizhen Wang, Wen Liu, Zhenda Xie, Yebin Liu:
DreamCraft3D: Hierarchical 3D Generation with Bootstrapped Diffusion Prior.
- Xiaoming Zhao, Alex Colburn, Fangchang Ma, Miguel Ángel Bautista, Joshua M. Susskind, Alexander G. Schwing:
Pseudo-Generalized Dynamic View Synthesis from a Video.
- Reza Esfandiarpoor, Stephen H. Bach:
Follow-Up Differential Descriptions: Language Models Resolve Ambiguities for Image Classification.
- Xiaoyi Liu, Duxin Chen, Wenjia Wei, Xia Zhu, Wenwu Yu:
Interpretable Sparse System Identification: Beyond Recent Deep Learning Techniques on Time-Series Prediction.
- Shikai Fang, Xin Yu, Zheng Wang, Shibo Li, Mike Kirby, Shandian Zhe:
Functional Bayesian Tucker Decomposition for Continuous-indexed Tensor Data.



- Alexander H. Liu, Matthew Le, Apoorv Vyas, Bowen Shi, Andros Tjandra, Wei-Ning Hsu:
Generative Pre-training for Speech with Flow Matching.
- Huaijin Wu, Wei Liu, Yatao Bian, Jiaxiang Wu, Nianzu Yang, Junchi Yan:
EBMDock: Neural Probabilistic Protein-Protein Docking via a Differentiable Energy Model.
- Yuxiao Cheng, Ziqian Wang, Tingxiong Xiao, Qin Zhong, Jinli Suo, Kunlun He:
CausalTime: Realistically Generated Time-series for Benchmarking of Causal Discovery.
- Shikun Feng, Minghao Li, Yinjun Jia, Wei-Ying Ma, Yanyan Lan:
Protein-ligand binding representation learning from fine-grained interactions.
- Aliasghar Khani, Saeid Asgari Taghanaki, Aditya Sanghi, Ali Mahdavi-Amiri, Ghassan Hamarneh:
SLiMe: Segment Like Me.
- Binchi Zhang , Yushun Dong, Chen Chen, Yada Zhu, Minnan Luo, Jundong Li:
Adversarial Attacks on Fairness of Graph Neural Networks.
- Songning Lai, Lijie Hu, Junxiao Wang, Laure Berti-Équille, Di Wang:
Faithful Vision-Language Interpretation via Concept Bottleneck Models.
- Arturs Backurs, Zinan Lin, Sepideh Mahabadi, Sandeep Silwal, Jakub Tarnawski:
Efficiently Computing Similarities to Private Datasets.
- Yuyan Ni, Shikun Feng, Wei-Ying Ma, Zhi-Ming Ma, Yanyan Lan:
Sliced Denoising: A Physics-Informed Molecular Pre-Training Method.
- Xiangyu Dong, Xingyi Zhang, Sibo Wang:
Rayleigh Quotient Graph Neural Networks for Graph-level Anomaly Detection.
- Chuyu Zhang, Hui Ren, Xuming He:
P2OT: Progressive Partial Optimal Transport for Deep Imbalanced Clustering.
- Zhiyu Mei, Wei Fu, Jiaxuan Gao, Guangju Wang, Huachen Zhang, Yi Wu:
SRL: Scaling Distributed Reinforcement Learning to Over Ten Thousand Cores.
- Zhenyi Wang, Yan Li, Li Shen, Heng Huang:
A Unified and General Framework for Continual Learning.
- Xinyao Fan, Yueying Wu, Chang Xu, Yuhao Huang, Weiqing Liu, Jiang Bian:
MG-TSD: Multi-Granularity Time Series Diffusion Models with Guided Learning Process.
- Meirui Jiang, Anjie Le, Xiaoxiao Li, Qi Dou:
Heterogeneous Personalized Federated Learning by Local-Global Updates Mixing via Convergence Rate.
- Kang Liu:
SetCSE: Set Operations using Contrastive Learning of Sentence Embeddings.
- Chenhui Deng, Zichao Yue, Zhiru Zhang:
Polynormer: Polynomial-Expressive Graph Transformer in Linear Time.
- Qihan Ren, Jiayang Gao, Wen Shen, Quanshi Zhang:
Where We Have Arrived in Proving the Emergence of Sparse Interaction Primitives in DNNs.
- Yueru Luo, Shuguang Cui, Zhen Li:
DV-3DLane: End-to-end Multi-modal 3D Lane Detection with Dual-view Representation.
- Xihai Luo, Wei Xu, Balu Nadiga, Yihui Ren, Shinjae Yoo:
Continuous Field Reconstruction from Sparse Observations with Implicit Neural Networks.
- Yifeng Fan, Yongqiang Li, Bo Chen:
Weaker MVI Condition: Extragradient Methods with Multi-Step Exploration.
- Guy Tennenholtz, Yinlam Chow, Chih-Wei Hsu, Jihwan Jeong, Lior Shani, Azamat Tulepbergenov, Deepak Ramachandran, Martin Mladenov, Craig Boutilier:
Demystifying Embedding Spaces using Large Language Models.
- Yuchen Zeng, Kangwook Lee:
The Expressive Power of Low-Rank Adaptation.
- Jiahao Nie, Zhiwei He, Xudong Lv, Xueyi Zhou, Dong-Kyu Chae, Fei Xie:
Towards Category Unification of 3D Single Object Tracking on Point Clouds.
- Yutong Wu, Han Qiu, Shangwei Guo, Jiwei Li, Tianwei Zhang:
You Only Query Once: An Efficient Label-Only Membership Inference Attack.
- Albert Xu, Jhih-Yi Hsieh, Bhaskar Vundurthy, Nithya Kemp, Eliana Cohen, Lu Li, Howie Choset:
Mathematical Justification of Hard Negative Mining via Isometric Approximation Theorem.
- Junyan Cheng, Peter Chin:
Bridging Neural and Symbolic Representations with Transitional Dictionary Learning.
- Kethmi Hirushini Hettige, Jiahao Ji, Shili Xiang, Cheng Long, Gao Cong, Jingyuan Wang:
AirPhyNet: Harnessing Physics-Guided Neural Networks for Air Quality Prediction.
- Suning Huang, Boyuan Chen, Huazhe Xu, Vincent Sitzmann:
DittoGym: Learning to Control Soft Shape-Shifting Robots.



- Junwei Su, Difan Zou, Chuan Wu:
PRES: Toward Scalable Memory-Based Dynamic Graph Neural Networks.
- Hyunjin Seo, Jihun Yun, Eunho Yang:
TEDDY: Trimming Edges with Degree-based Discrimination Strategy.
- Yichen Li, Yilun Du, Chao Liu, Chao Liu, Francis Williams, Michael Foshey, Benjamin Eckart, Jan Kautz, Joshua B. Tenenbaum, Antonio Torralba, Wojciech Matusik:
Learning to Jointly Understand Visual and Tactile Signals.
- Gen Li, Lu Yin, Jie Ji, Wei Niu, Minghai Qin, Bin Ren, Linke Guo, Shiwei Liu, Xiaolong Ma:
NeurRev: Train Better Sparse Neural Network Practically via Neuron Revitalization.
- Kathan Shah, Chawin Sitawarin:
SPDER: Semiperiodic Damping-Enabled Object Representation.
- Daniel Severo, Lucas Theis, Johannes Ballé:
The Unreasonable Effectiveness of Linear Prediction as a Perceptual Metric.
- Tatsunori Taniai, Ryo Igarashi, Yuta Suzuki, Naoya Chiba, Kotaro Saito, Yoshitaka Ushiku, Kanta Ono:
Crystalformer: Infinitely Connected Attention for Periodic Structure Encoding.
- Hyunsu Kim, Jongmin Yoon, Juho Lee:
Fast Ensembling with Diffusion Schrödinger Bridge.
- Sobhan Mohammadpour, Emma Frejinger, Pierre-Luc Bacon:
Decoupling regularization from the action space.
- Shuo Chen, Gang Niu, Chen Gong, Okan Koc, Jian Yang, Masashi Sugiyama:
Robust Similarity Learning with Difference Alignment Regularization.
- Yuhang Zang, Hanlin Goh, Joshua M. Susskind, Chen Huang:
Overcoming the Pitfalls of Vision-Language Model Finetuning for OOD Generalization.
- Mahsa Keramati, Lili Meng, R. David Evans:
ConR: Contrastive Regularizer for Deep Imbalanced Regression.
- Jintang Li, Huizhe Zhang, Ruofan Wu, Zulun Zhu, Baokun Wang, Changhua Meng, Zibin Zheng, Liang Chen:
A Graph is Worth 1-bit Spikes: When Graph Contrastive Learning Meets Spiking Neural Networks.
- Alexander Robey, Fabian Latorre, George J. Pappas, Hamed Hassani, Volkan Cevher:
Adversarial Training Should Be Cast as a Non-Zero-Sum Game.
- Murtaza Dalal, Tarun Chiruvolu, Devendra Singh Chaplot, Ruslan Salakhutdinov:
Plan-Seq-Learn: Language Model Guided RL for Solving Long Horizon Robotics Tasks.
- Xinyu Yuan, Yan Qiao:
Diffusion-TS: Interpretable Diffusion for General Time Series Generation.
- Jiecheng Lu, Xu Han, Shihao Yang:
ARM: Refining Multivariate Forecasting with Adaptive Temporal-Contextual Learning.
- Sai Surya Duvvuri, Devvrit, Rohan Anil, Cho-Jui Hsieh, Inderjit S. Dhillon:
Combining Axes Preconditioners through Kronecker Approximation for Deep Learning.
- Hsiang Hsu, Guihong Li, Shaohan Hu, Chun-Fu Chen:
Dropout-Based Rashomon Set Exploration for Efficient Predictive Multiplicity Estimation.
- Sophia Huiwen Sun, Rose Yu:
Copula Conformal prediction for multi-step time series prediction.
- Fredrik Carlsson, Johan Broberg, Erik Hillbom, Magnus Sahlgren, Joakim Nivre:
Branch-GAN: Improving Text Generation with (not so) Large Language Models.
- Lukas Fesser, Melanie Weber:
Effective Structural Encodings via Local Curvature Profiles.
- Gabriele Tiboni, Pascal Klink, Jan Peters, Tatiana Tommasi, Carlo D'Eramo, Georgia Chalvatzaki:
Domain Randomization via Entropy Maximization.
- Jenny Zhang, Joel Lehman, Kenneth O. Stanley, Jeff Clune:
OMNI: Open-endedness via Models of human Notions of Interestingness.
- Guihong Li, Hsiang Hsu, Chun-Fu Chen, Radu Marculescu:
Machine Unlearning for Image-to-Image Generative Models.
- Hailey Joren, Charles T. Marx, Berk Ustun :
Classification with Conceptual Safeguards.
- Kwangjun Ahn, Xiang Cheng, Minhak Song, Chulhee Yun, Ali Jadbabaie, Suvrit Sra:
Linear attention is (maybe) all you need (to understand Transformer optimization).
- Yichun Shi, Peng Wang, Jianglong Ye, Long Mai, Kejie Li, Xiao Yang:
MVDream: Multi-view Diffusion for 3D Generation.
- Minjun Sung, Sambhu H. Karumanchi, Aditya Gahlawat, Naira Hovakimyan:
Robust Model Based Reinforcement Learning Using L1 Adaptive Control.

-     Nathan Godey, Éric Villemonte de la Clergerie, Benoît Sagot:
Headless Language Models: Learning without Predicting with Contrastive Weight Tying.
-     Nils Lukas, Abdulrahman Diaa, Lucas Fenaux, Florian Kerschbaum:
Leveraging Optimization for Adaptive Attacks on Image Watermarks.
-     Guanhua Wang, Heyang Qin, Sam Ade Jacobs, Xiaoxia Wu, Connor Holmes, Zhewei Yao, Samyam Rajbhandari, Olatunji Ruwase, Feng Yan, Lei Yang, Yuxiong He:
ZeRO++: Extremely Efficient Collective Communication for Large Model Training.
-     Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, Denny Zhou:
Large Language Models Cannot Self-Correct Reasoning Yet.
-     Lunjun Zhang, Yuwen Xiong, Ze Yang, Sergio Casas, Rui Hu, Raquel Urtasun:
Copilot4D: Learning Unsupervised World Models for Autonomous Driving via Discrete Diffusion.
-     Benjamin Schneider, Nils Lukas, Florian Kerschbaum:
Universal Backdoor Attacks.
-     Daniel Goldfarb, Itay Evron, Nir Weinberger, Daniel Soudry, Paul Hand:
The Joint Effect of Task Similarity and Overparameterization on Catastrophic Forgetting - An Analytical Model.
-     Haitao Yang, Xiangru Huang, Bo Sun, Chandrajit L. Bajaj, Qixing Huang:
GenCorres: Consistent Shape Matching via Coupled Implicit-Explicit Shape Generative Models.
-     Xingyao Wang, Zihan Wang, Jiateng Liu, Yangyi Chen, Lifan Yuan, Hao Peng, Heng Ji:
MINT: Evaluating LLMs in Multi-turn Interaction with Tools and Language Feedback.
-     Alon Ziv, Itai Gat, Gaël Le Lan, Tal Remez, Felix Kreuk, Jade Copet, Alexandre Défossez, Gabriel Synnaeve, Yossi Adi:
Masked Audio Generation using a Single Non-Autoregressive Transformer.
-     Federico Bianchi, Mirac Suzgun, Giuseppe Attanasio, Paul Röttger, Dan Jurafsky, Tatsunori Hashimoto, James Zou:
Safety-Tuned LLaMAs: Lessons From Improving the Safety of Large Language Models that Follow Instructions.
-     Bo Zhou, Ruiwei Jiang, Siqian Shen:
Learning to Solve Bilevel Programs with Binary Tender.
-     Ahmed A. A. Elhag, Yuyang Wang, Joshua M. Susskind, Miguel Ángel Bautista:
Manifold Diffusion Fields.
-     Justin Dumouchelle, Esther Julien, Jannis Kurtz, Elias Boutros Khalil:
Neur2RO: Neural Two-Stage Robust Optimization.
-     Elías Abad-Rocamora, Fanghui Liu, Grigoris Chrysos, Pablo M. Olmos, Volkan Cevher:
Efficient local linearity regularization to overcome catastrophic overfitting.
-     Qiuyi Chen, Mark D. Fuge:
Compressing Latent Space via Least Volume.
-     Anke Tang, Li Shen, Yong Luo, Yibing Zhan, Han Hu, Bo Du, Yixin Chen, Dacheng Tao:
Parameter-Efficient Multi-Task Model Fusion with Partial Linearization.
-     Xavier Puig, Eric Undersander, Andrew Szot, Mikael Dallaire Cote, Tsung-Yen Yang, Ruslan Partsey, Ruta Desai, Alexander Clegg, Michal Hlavac, So Yeon Min, Vladimir Vondrus, Théophile Gervet, Vincent-Pierre Berges, John M. Turner, Oleksandr Maksymets, Zsolt Kira, Mrinal Kalakrishnan, Jitendra Malik, Devendra Singh Chaplot, Unnat Jain, Dhruv Batra, Akshara Rai, Roozbeh Mottaghi:
Habitat 3.0: A Co-Habitat for Humans, Avatars, and Robots.
-     Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, Karl Cobbe:
Let's Verify Step by Step.
-     Druv Pai, Sam Buchanan, Ziyang Wu, Yaodong Yu, Yi Ma:
Masked Completion via Structured Diffusion with White-Box Transformers.
-     Kai Hu, Klas Leino, Zifan Wang, Matt Fredrikson:
A Recipe for Improved Certifiable Robustness.
-     Yichuan Li, Xiyao Ma, Sixing Lu, Kyumin Lee, Xiaohu Liu, Chenlei Guo:
MEND: Meta Demonstration Distillation for Efficient and Effective In-Context Learning.
-     Tycho F. A. van der Ouderaa, Markus Nagel, Mart van Baalen, Tijmen Blankevoort:
The LLM Surgeon.
-     Jonathan Vacher, Pascal Mamassian:
Perceptual Scales Predicted by Fisher Information Metrics.
-     John X. Morris, Wenting Zhao, Justin T. Chiu, Vitaly Shmatikov, Alexander M. Rush:
Language Model Inversion.
-     Hyungjin Chung, Suhyeon Lee, Jong Chul Ye:
Decomposed Diffusion Sampler for Accelerating Large-Scale Inverse Problems.
-     Mingxuan Liu, Subhankar Roy, Wenjing Li, Zhun Zhong, Nicu Sebe, Elisa Ricci:
Democratizing Fine-grained Visual Recognition with Large Language Models.
-     Lichang Chen, Shiyang Li, Jun Yan, Hai Wang, Kalpa Gunaratna, Vikas Yadav, Zheng Tang, Vijay Srinivasan, Tianyi Zhou, Heng Huang, Hongxia Jin:

**AlpaGasus: Training a Better Alpaca with Fewer Data.**

Isaac Reid, Krzysztof Marcin Choromanski, Eli Berger, Adrian Weller:
General Graph Random Features.

Insu Han, Rajesh Jayaram, Amin Karbasi, Vahab Mirrokni, David P. Woodruff, Amir Zandieh :
HyperAttention: Long-context Attention in Near-Linear Time.

Isaac Reid, Eli Berger, Krzysztof Marcin Choromanski, Adrian Weller:
Repelling Random Walks.

Patrick Schnell, Nils Thuerey:
Stabilizing Backpropagation Through Time to Learn Complex Physics.

Seamus Somerstep, Yuekai Sun, Yaakov Ritov:
Learning in reverse causal strategic environments with ramifications on two sided markets.

Kaixuan Ji, Qingyue Zhao, Jiafan He, Weitong Zhang, Quanquan Gu:
Horizon-free Reinforcement Learning in Adversarial Linear Mixture MDPs.

Haomin Zhuang, Mingxian Yu, Hao Wang, Yang Hua, Jian Li, Xu Yuan:
Backdoor Federated Learning by Poisoning Backdoor-Critical Layers.

Jiajun He, Gergely Flamich, Zongyu Guo, José Miguel Hernández-Lobato:
RECOMBINER: Robust and Enhanced Compression with Bayesian Implicit Neural Representations.

Huigen Ye, Hua Xu, Hongyan Wang:
Light-MILPopt: Solving Large-scale Mixed Integer Linear Programs with Lightweight Optimizer and Small-scale Training Dataset.

Tim Dettmers, Ruslan Svirchevski, Vage Egiazarian, Denis Kuznedelev, Elias Frantar, Saleh Ashkboos, Alexander Borzunov, Torsten Hoefer, Dan Alistarh:
SpQR: A Sparse-Quantized Representation for Near-Lossless LLM Weight Compression.

Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schieber, Da Yan, Miranda Zhang, Ethan Perez:
Towards Understanding Sycophancy in Language Models.

Edouard Yvinec, Arnaud Dapogny, Kevin Bailly:
Network Memory Footprint Compression Through Jointly Learnable Codebooks and Mappings.

Marlène Careil, Matthew J. Muckley, Jakob Verbeek, Stéphane Lathuilière:
Towards image compression with perfect realism at ultra-low bitrates.

Xiaoran Liu, Hang Yan, Chenxin An, Xipeng Qiu, Dahua Lin:
Scaling Laws of RoPE-based Extrapolation.

Yichao Shen, Zigang Geng, Yuhui Yuan, Yutong Lin, Ze Liu, Chunyu Wang, Han Hu, Nanning Zheng, Baining Guo:
V-DETR: DETR with Vertex Relative Position Encoding for 3D Object Detection.

Zixian Huang, Gengyang Xiao, Yu Gu, Gong Cheng:
A Branching Decoder for Set Generation.

Alex Fang, Albin Madappally Jose, Amit Jain, Ludwig Schmidt, Alexander T. Toshev, Vaishaal Shankar:
Data Filtering Networks.

Wei-Hong Li, Steven McDonagh, Ales Leonardis, Hakan Bilen:
Multi-task Learning with 3D-Aware Regularization.

Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, Mike Lewis:
Efficient Streaming Language Models with Attention Sinks.

Hongcheng Guo, Jian Yang, Jiaheng Liu, Liqun Yang, Linzheng Chai, Jiaqi Bai, Junran Peng, Xiaorong Hu, Chao Chen, Dongfeng Zhang, Xu Shi, Tieqiao Zheng, Liangfan Zheng, Bo Zhang, Ke Xu, Zhoujun Li:
OWL: A Large Language Model for IT Operations.

Zhihe Yang, Yunjian Xu:
DMBP: Diffusion model-based predictor for robust offline reinforcement learning against state observation perturbations.

Johannes Hertrich, Christian Wald, Fabian Altekrüger, Paul Hagemann:
Generative Sliced MMD Flows with Riesz Kernels.

Maxim Khanov, Jirayu Buraphacheep, Yixuan Li:
ARGS: Alignment as Reward-Guided Search.

Dongyoung Go, Tomasz Korbak, Germán Kruszewski, Jos Rozen, Marc Dymetman:
Compositional Preference Models for Aligning LMs.

Zhikai Chen, Haitao Mao, Hongzhi Wen, Haoyu Han, Wei Jin, Haiyang Zhang, Hui Liu, Jiliang Tang:
Label-free Node Classification on Graphs with Large Language Models (LLMs).

Khai Nguyen, Nhat Ho:
Sliced Wasserstein Estimation with Control Variates.

Yiwei Zhang, Guo Lu, Yunuo Chen, Shen Wang, Yibo Shi, Jing Wang, Li Song:
Neural Rate Control for Learned Video Compression.

-     Ahmed Hendawy, Jan Peters, Carlo D'Eramo:
Multi-Task Reinforcement Learning with Mixture of Orthogonal Experts.
-     Site Bai, Brian Bullins:
Local Composite Saddle Point Optimization.
-     Tianci Liu, Haoyu Wang, Feijie Wu, Hengtong Zhang, Pan Li, Lu Su, Jing Gao:
Towards Poisoning Fair Representations.
-     Ling Yang, Zhilong Zhang, Zhaochen Yu, Jingwei Liu, Minkai Xu, Stefano Ermon, Bin Cui:
Cross-Modal Contextualized Diffusion Models for Text-Guided Visual Generation and Editing.
-     Shaopeng Fu, Di Wang:
Theoretical Analysis of Robust Overfitting for Wide DNNs: An NTK Approach.
-     Zhenghan Fang, Sam Buchanan, Jeremias Sulam:
What's in a Prior? Learned Proximal Networks for Inverse Problems.
-     Michael Gastpar, Ido Nachum, Jonathan Shafer, Thomas Weinberger:
Fantastic Generalization Measures are Nowhere to be Found.
-     Xiang Lan, Hanshu Yan, Shenda Hong, Mengling Feng:
Towards Enhancing Time Series Contrastive Learning: A Dynamic Bad Pair Mining Approach.
-     Eliya Segev, Maya Alroy, Ronen Katsir, Noam Wies, Ayana Shenhav, Yael Ben-Oren, David Zar, Oren Tadmor, Jacob Bitterman, Amnon Shashua, Tal Rosenwein:
Align With Purpose: Optimize Desired Properties in CTC Models with a General Plug-and-Play Framework.
-     Yonghao Song, Bingchuan Liu, Xiang Li, Nanlin Shi, Yijun Wang, Xiaorong Gao:
Decoding Natural Images from EEG for Object Recognition.
-     Rocio Diaz Martin, Ivan Vladimir Medri, Yikun Bai, Xinran Liu, Kangbai Yan, Gustavo K. Rohde, Soheil Kolouri:
LCOT: Linear Circular Optimal Transport.
-     Zhoubo Li, Ningyu Zhang, Yunzhi Yao, Mengru Wang, Xi Chen, Huajun Chen:
Unveiling the Pitfalls of Knowledge Editing for Large Language Models.
-     Jianhui Li, Shilong Liu, Zidong Liu, Yikai Wang, Kaiwen Zheng, Jinghui Xu, Jianmin Li, Jun Zhu:
InstructPix2NeRF: Instructed 3D Portrait Editing from a Single Image.
-     Omer Nahum, Gali Noti, David C. Parkes, Nir Rosenfeld:
Decongestion by Representation: Learning to Improve Economic Welfare in Marketplaces.
-     Beatrix Miranda Ginn Nielsen, Anders Christensen, Andrea Dittadi, Ole Winther:
DiffEnc: Variational Diffusion with a Learned Encoder.
-     Zitao Song, Wendi Ren, Shuang Li:
Amortized Network Intervention to Steer the Excitatory Point Processes.
-     Lorenz Vaitl, Ludwig Winkler, Lorenz Richter, Pan Kessel:
Fast and unified path gradient estimators for normalizing flows.
-     Peng Xu, Wenqi Shao, Mengzhao Chen, Shitao Tang, Kaipeng Zhang, Peng Gao, Fengwei An, Yu Qiao, Ping Luo:
BESA: Pruning Large Language Models with Blockwise Parameter-Efficient Sparsity Allocation.
-     Vladimir R. Kostic, Pietro Novelli, Riccardo Grazzi, Karim Lounici, Massimiliano Pontil:
Learning invariant representations of time-homogeneous stochastic dynamical systems.
-     Nikolas Patris, Ioannis Panageas:
Learning Nash Equilibria in Rank-1 Games.
-     Yuan Gao, Rustem Islamov, Sebastian U. Stich:
EControl: Fast Distributed Optimization with Compression and Error Control.
-     Li Meng, Morten Goodwin, Anis Yazidi, Paal E. Engelstad:
State Representation Learning Using an Unbalanced Atlas.
-     Austin Tripp, Krzysztof Maziarz, Sarah Lewis, Marwin H. S. Segler, José Miguel Hernández-Lobato:
Retro-fallback: retrosynthetic planning in an uncertain world.
-     Guangchi Fang, Qingyong Hu, Longguang Wang, Yulan Guo:
ACRF: Compressing Explicit Neural Radiance Fields via Attribute Compression.
-     Sifan Zhou, Liang Li, Xinyu Zhang, Bo Zhang, Shipeng Bai, Miao Sun, Ziyu Zhao, Xiaobo Lu, Xiangxiang Chu:
LiDAR-PTQ: Post-Training Quantization for Point Cloud 3D Object Detection.
-     Pol Lababarrie, Adrien Chan-Hon-Tong, Stéphane Herbin, Milad Leyli-Abadi:
Optimal transport based adversarial patch to leverage large scale attack transferability.
-     Li Siyao, Tianpei Gu, Zhitao Yang, Zhengyu Lin, Ziwei Liu, Henghui Ding, Lei Yang, Chen Change Loy:
Duolando: Follower GPT with Off-Policy Reinforcement Learning for Dance Accompaniment.
-     Pum Jun Kim, Seojun Kim, Jaejun Yoo:
STREAM: Spatio-TempoRal Evaluation and Analysis Metric for Video Generative Models.
-     Jorge Fernandez-de-Cossío-Díaz, Clément Roussel, Simona Cocco, Rémi Monasson:
Accelerated Sampling with Stacked Restricted Boltzmann Machines.



- Anji Liu, Mathias Niepert, Guy Van den Broeck:
Image Inpainting via Tractable Steering of Diffusion Models.
- Yuhui Xu, Lingxi Xie, Xiaotao Gu, Xin Chen, Heng Chang, Hengheng Zhang, Zhengsu Chen, Xiaopeng Zhang, Qi Tian:
QA-LoRA: Quantization-Aware Low-Rank Adaptation of Large Language Models.
- Duy-Kien Nguyen, Yanghao Li, Vaibhav Aggarwal, Martin R. Oswald, Alexander Kirillov, Cees G. M. Snoek, Xinlei Chen:
R-MAE: Regions Meet Masked Autoencoders.
- Yuxin Li, Wenchao Chen, Xinyue Hu, Bo Chen, Baolin Sun, Mingyuan Zhou:
Transformer-Modulated Diffusion Models for Probabilistic Multivariate Time Series Forecasting.
- Guan Wang, Sijie Cheng, Xianyuan Zhan, Xiangang Li, Sen Song, Yang Liu:
OpenChat: Advancing Open-source Language Models with Mixed-Quality Data.
- Jannik Kossen, Yarin Gal, Tom Rainforth:
In-Context Learning Learns Label Relationships but Is Not Conventional Learning.
- Yoni Shafir, Guy Tevet, Roy Kapon, Amit Haim Bermano:
Human Motion Diffusion as a Generative Prior.
- Xinzhe Yuan, William de Vazelhes, Bin Gu, Huan Xiong:
New Insight of Variance reduce in Zero-Order Hard-Thresholding: Mitigating Gradient Error and Expansivity Contradictions.
- Enneng Yang, Zhenyi Wang, Li Shen, Shiwei Liu, Guibing Guo, Xingwei Wang, Dacheng Tao:
AdaMerging: Adaptive Model Merging for Multi-Task Learning.
- Zhilu Zhang, Haoyu Wang, Shuai Liu, Xiaotao Wang, Lei Lei, Wangmeng Zuo:
Self-Supervised High Dynamic Range Imaging with Multi-Exposure Images in Dynamic Scenes.
- Yue Huang, Jiawen Shi, Yuan Li, Chenrui Fan, Siyuan Wu, Qihui Zhang, Yixin Liu, Pan Zhou, Yao Wan, Neil Zhenqiang Gong, Lichao Sun :
MetaTool Benchmark for Large Language Models: Deciding Whether to Use Tools and Which to Use.
- Yiyang Chen, Zhedong Zheng, Wei Ji, Leigang Qu, Tat-Seng Chua:
Composed Image Retrieval with Text Feedback via Multi-grained Uncertainty Regularization.
- Takuya Furusawa:
Mean Field Theory in Deep Metric Learning.
- Panagiotis Dimitrakopoulos, Giorgos Sfikas, Christophoros Nikou:
Implicit Neural Representation Inference for Low-Dimensional Bayesian Deep Learning.
- Yiting Chen, Zhanpeng Zhou, Junchi Yan:
Going Beyond Neural Network Feature Similarity: The Network Feature Complexity and Its Interpretation Using Category Theory.
- Jiaxin Lu, Zetian Jiang, Tianzhe Wang, Junchi Yan:
M3C: A Framework towards Convergent, Flexible, and Unsupervised Learning of Mixture Graph Matching and Clustering.
- Jiachun Pan, Jun Hao Liew, Vincent Y. F. Tan, Jiashi Feng, Hanshu Yan:
AdjointDPM: Adjoint Sensitivity Method for Gradient Backpropagation of Diffusion Probabilistic Models.
- Renrui Zhang, Jiaming Han, Chris Liu, Aojun Zhou, Pan Lu, Yu Qiao, Hongsheng Li, Peng Gao:
LLaMA-Adapter: Efficient Fine-tuning of Large Language Models with Zero-initialized Attention.
- Junyoung Seo, Wooseok Jang, Minseop Kwak, Inès Hyeonsu Kim, Jaehoon Ko, Junho Kim, Jin-Hwa Kim, Jiyong Lee, Seungryong Kim:
Let 2D Diffusion Model Know 3D-Consistency for Robust Text-to-3D Generation.
- Ziteng Gao, Zhan Tong, Limin Wang, Mike Zheng Shou:
SparseFormer: Sparse Visual Recognition via Limited Latent Tokens.
- Faisal Hamman, Sanghamitra Dutta:
Demystifying Local & Global Fairness Trade-offs in Federated Learning Using Partial Information Decomposition.
- Zhilin Huang, Ling Yang, Xiangxin Zhou, Zhilong Zhang, Wentao Zhang, Xiawu Zheng, Jie Chen, Yu Wang, Bin Cui, Wenming Yang:
Protein-Ligand Interaction Prior for Binding-aware 3D Molecule Diffusion Models.
- Jiashuo Sun, Chengjin Xu, Lumingyan Tang, Saizhuo Wang, Chen Lin, Yeyun Gong, Lionel M. Ni, Heung-Yeung Shum, Jian Guo:
Think-on-Graph: Deep and Responsible Reasoning of Large Language Model on Knowledge Graph.
- Yujie Mo, Feiping Nie, Ping Hu, Heng Tao Shen, Zheng Zhang, Xinchao Wang, Xiaofeng Zhu:
Self-Supervised Heterogeneous Graph Learning: a Homophily and Heterogeneity View.
- Haoxuan You, Mandy Guo, Zhecan Wang, Kai-Wei Chang, Jason M. Baldridge, Jiahui Yu:
CoBIT: A Contrastive Bi-directional Image-Text Generation Model.
- Ziqi Gao, Xiangguo Sun, Zijing Liu, Yu Li, Hong Cheng, Jia Li:
Protein Multimer Structure Prediction via Prompt Learning.
- Yin Fang, Ningyu Zhang, Zhuo Chen, Lingbing Guo, Xiaohui Fan, Huajun Chen:
Domain-Agnostic Molecular Generation with Chemical Feedback.



- Long Lian, Baifeng Shi, Adam Yala, Trevor Darrell, Boyi Li:
LLM-grounded Video Diffusion Models.
- Tao Dai, Beiliang Wu, Peiyuan Liu, Naiqi Li, Jigang Bao, Yong Jiang, Shu-Tao Xia:
Periodicity Decoupling Framework for Long-term Series Forecasting.
- Yuyang Liu, Weijun Dong, Yingdong Hu, Chuan Wen, Zhao-Heng Yin, Chongjie Zhang, Yang Gao:
Imitation Learning from Observation with Automatic Discount Scheduling.
- Jongwon Jeong, Hoyeop Lee, Hyui Geon Yoon, Beomyoung Lee, Junhee Heo, Geonsoo Kim, Kim Jin Seon:
iGraphMix: Input Graph Mixup Method for Node Classification.
- Hansam Cho, Jonghyun Lee, Seoung Bum Kim, Tae-Hyun Oh, Yonghyun Jeong:
Noise Map Guidance: Inversion with Spatial Context for Real Image Editing.
- Ilmin Kang, HyounYoung Bae, Kangil Kim:
Label-Focused Inductive Bias over Latent Object Features in Visual Classification.
- Tomoya Murata, Kenta Niwa, Takumi Fukami, lifan Tyu:
Simple Minimax Optimal Byzantine Robust Algorithm for Nonconvex Objectives with Uniform Gradient Heterogeneity.
- Dongming Wu, Jiahao Chang, Fan Jia, Yingfei Liu, Tiancai Wang, Jianbing Shen:
TopoMLP: A Simple yet Strong Pipeline for Driving Topology Reasoning.
- Renrui Zhang, Zhengkai Jiang, Ziyu Guo, Shilin Yan, Junting Pan, Hao Dong, Yu Qiao, Peng Gao, Hongsheng Li:
Personalize Segment Anything Model with One Shot.
- Jung-Chun Liu, Chi-Hsien Chang, Shao-Hua Sun, Tian-Li Yu:
Integrating Planning and Deep Reinforcement Learning via Automatic Induction of Task Substructures.
- Dawei Zhu, Nan Yang, Liang Wang, Yifan Song, Wenhao Wu, Furu Wei, Sujian Li:
PoSE: Efficient Context Window Extension of LLMs via Positional Skip-wise Training.
- Jaewoo Choi, Youngwoo Yoon, Hyobin Ong, Jaehong Kim, Minsu Jang:
LoTa-Bench: Benchmarking Language-oriented Task Planners for Embodied Agents.
- Xinhua Cheng, Tianyu Yang, Jianan Wang, Yu Li, Lei Zhang, Jian Zhang, Li Yuan:
Progressive3D: Progressively Local Editing for Text-to-3D Content Creation with Complex Semantic Prompts.
- Bill Yuchen Lin, Abhilasha Ravichander, Ximing Lu, Nouha Dziri, Melanie Sclar, Khyathi Raghavi Chandu, Chandra Bhagavatula, Yejin Choi:
The Unlocking Spell on Base LLMs: Rethinking Alignment via In-Context Learning.
- Fred Zhang, Neel Nanda:
Towards Best Practices of Activation Patching in Language Models: Metrics and Methods.
- Chaohua Shi, Kexin Huang, Lu Gan, Hongqing Liu, Mingrui Zhu, Nannan Wang, Xinbo Gao:
On the Analysis of GAN-based Image-to-Image Translation with Gaussian Noise Injection.
- Ziyao Wang, Jianyu Wang, Ang Li:
FedHyper: A Universal and Robust Learning Rate Scheduler for Federated Learning with Hypergradient Descent.
- Haiping Wang, Yuan Liu, Bing Wang, Yujing Sun, Zhen Dong, Wenping Wang, Bisheng Yang:
FreeReg: Image-to-Point Cloud Registration Leveraging Pretrained Diffusion Models and Monocular Depth Estimators.
- Huaixiu Steven Zheng, Swaroop Mishra, Xinyun Chen, Heng-Tze Cheng, Ed H. Chi, Quoc V. Le, Denny Zhou:
Take a Step Back: Evoking Reasoning via Abstraction in Large Language Models.
- Max Ku, Tianle Li, Kai Zhang, Yujie Lu, Xingyu Fu, Wenwen Zhuang, Wenhui Chen:
ImagenHub: Standardizing the evaluation of conditional image generation models.
- Kai Cheng, Xiaoxiao Long, Wei Yin, Jin Wang, Zhiqiang Wu, Yuexin Ma, Kaixuan Wang, Xiaozhi Chen, Xuejin Chen:
UC-NERF: Neural Radiance Field for Under-Calibrated Multi-View Cameras in Autonomous Driving.
- Daixuan Cheng, Shaohan Huang, Furu Wei:
Adapting Large Language Models via Reading Comprehension.
- Zhenting Wang, Chen Chen, Lingjuan Lyu, Dimitris N. Metaxas, Shiqing Ma:
DIAGNOSIS: Detecting Unauthorized Data Usages in Text-to-image Diffusion Models.
- Yite Wang, Jiahao Su, Hanlin Lu, Cong Xie, Tianyi Liu, Jianbo Yuan, Haibin Lin, Ruoyu Sun, Hongxia Yang:
LEMON: Lossless model expansion.
- Zhengbo Wang, Jian Liang, Lijun Sheng, Ran He, Zilei Wang, Tieniu Tan:
A Hard-to-Beat Baseline for Training-free CLIP-based Adaptation.
- Yuxian Gu, Li Dong, Furu Wei, Minlie Huang:
MiniLLM: Knowledge Distillation of Large Language Models.
- Kai Huang, Hanyun Yin, Heng Huang, Wei Gao:
Towards Green AI in Fine-tuning Large Language Models via Adaptive Backpropagation.
- Ziteng Sun, Ananda Theertha Suresh, Aditya Krishna Menon:
The importance of feature preprocessing for differentially private linear optimization.



Peng Chen, Yingying Zhang, Yunyao Cheng, Yang Shu, Yihang Wang, Qingsong Wen, Bin Yang, Chenjuan Guo:
Pathformer: Multi-scale Transformers with Adaptive Pathways for Time Series Forecasting.



Leo Feng, Frederick Tung, Hossein Hajimirsadeghi, Yoshua Bengio, Mohamed Osama Ahmed:
Tree Cross Attention.



Gunho Park, Baeseong Park, Minsub Kim, Sungjae Lee, Jeonghoon Kim, Beomseok Kwon, Se Jung Kwon, Byeongwook Kim, Youngjoo Lee, Dongsoo Lee:
LUT-GEMM: Quantized Matrix Multiplication based on LUTs for Efficient Inference in Large-Scale Generative Language Models.



Elan Rosenfeld, Andrej Risteski:
Outliers with Opposing Signals Have an Outsized Effect on Neural Network Optimization.



William Rudman, Carsten Eickhoff:
Stable Anisotropic Regularization.



Qin Zhang, Linghan Xu, Jun Fang, Qingming Tang, Ying Nian Wu, Joseph Tighe, Yifan Xing:
Threshold-Consistent Margin Loss for Open-World Deep Metric Learning.



Emanuele Aiello, Lili Yu, Yixin Nie, Armen Aghajanyan, Barlas Oguz:
Jointly Training Large Autoregressive Multimodal Models.



Xiangyu Liu, Souradip Chakraborty, Yanchao Sun, Furong Huang:
Rethinking Adversarial Policies: A Generalized Attack Formulation and Provable Defense in RL.



Runqi Lin, Chaojian Yu, Bo Han, Tongliang Liu:
On the Over-Memorization During Natural, Robust and Catastrophic Overfitting.



Peter West, Ximing Lu, Nouha Dziri, Faeze Brahman, Linjie Li, Jena D. Hwang, Liwei Jiang, Jillian Fisher, Abhilasha Ravichander, Khyathi Raghavi Chandu, Benjamin Newman, Pang Wei Koh, Allyson Ettinger, Yejin Choi:
The Generative AI Paradox: "What It Can Create, It May Not Understand".



Fengrui Tian, Yueqi Duan, Angtian Wang, Jianfei Guo, Shaoyi Du:
Semantic Flow: Learning Semantic Fields of Dynamic Scenes from Monocular Videos.



Haitao Mao, Juanhui Li, Harry Shomer, Bingheng Li, Wenqi Fan, Yao Ma, Tong Zhao, Neil Shah, Jiliang Tang:
Revisiting Link Prediction: a data perspective.



Zilong Wang, Hao Zhang, Chun-Liang Li, Julian Martin Eisenschlos, Vincent Perot, Zifeng Wang, Lesly Miculicich, Yasuhisa Fujii, Jingbo Shang, Chen-Yu Lee, Tomas Pfister:
Chain-of-Table: Evolving Tables in the Reasoning Chain for Table Understanding.



Linqi Zhou, Aaron Lou, Samar Khanna, Stefano Ermon:
Denoising Diffusion Bridge Models.



Shubham Ugare, Tarun Suresh, Debangshu Banerjee, Gagandeep Singh, Sasa Misailovic:
Incremental Randomized Smoothing Certification.



Artur Back de Luca, Kimon Fountoulakis, Shenghao Yang:
Local Graph Clustering with Noisy Labels.



Enyi Jiang, Yibo Jacky Zhang, Sanmi Koyejo:
Principled Federated Domain Adaptation: Gradient Projection and Auto-Weighting.



Kiarash Shamsi, Farimah Poursafaei, Shenyang Huang, Tran Gia Bao Ngo, Baris Coskunuzer, Cuneyt Gurcan Akcora:
GraphPulse: Topological representations for temporal graph property prediction.



Qingru Zhang, Chandan Singh, Liyuan Liu, Xiaodong Liu, Bin Yu, Jianfeng Gao, Tuo Zhao:
Tell Your Model Where to Attend: Post-hoc Attention Steering for LLMs.



Keivan Rezaei, Mehrdad Saberi, Mazda Moayeri, Soheil Feizi:
PRIME: Prioritizing Interpretability in Failure Mode Extraction.



Christian Horvat, Jean-Pascal Pfister:
On gauge freedom, conservativity and intrinsic dimensionality estimation in diffusion models.



Sidhika Balachandar, Nikhil Garg, Emma Pierson:
Domain constraints improve risk prediction when outcome data is missing.



Yat Long Lo, Biswa Sengupta, Jakob Nicolaus Foerster, Michael Noukhovitch:
Learning Multi-Agent Communication with Contrastive Learning.



Raj Ghugare, Matthieu Geist, Glen Berseth, Benjamin Eysenbach:
Closing the Gap between TD Learning and Supervised Learning - A Generalisation Point of View.



Shangmin Guo, Yi Ren, Stefano V. Albrecht, Kenny Smith:
IpNTK: Better Generalisation with Less Data via Sample Interaction During Learning.



Wenxuan Zhang, Youssef Mohamed, Bernard Ghanem, Philip Torr, Adel Bibi, Mohamed Elhoseiny:
Continual Learning on a Diet: Learning from Sparsely Labeled Streams Under Constrained Computation.



Gaurav Shrivastava, Ser-Nam Lim, Abhinav Shrivastava:
Video Decomposition Prior: Editing Videos Layer by Layer.



Mustafa Shukor, Alexandre Ramé, Corentin Dancette, Matthieu Cord:
Beyond task performance: evaluating and reducing the flaws of large multimodal models with in-context learning.

-     Adam Block, Dylan J. Foster, Akshay Krishnamurthy, Max Simchowitz, Cyril Zhang:
Butterfly Effects of SGD Noise: Error Amplification in Behavior Cloning and Autoregression.
-    Ruinan Jin, Shuai Li, Baoxiang Wang:
On Stationary Point Convergence of PPO-Clip.
-    Min Lin:
Automatic Functional Differentiation in JAX.
-    Zilinghan Li, Pranshu Chaturvedi, Shilan He, Han Chen, Gagandeep Singh, Volodymyr V. Kindratenko, Eliu A. Huerta, Kibaek Kim, Ravi K. Madduri:
FedCompass: Efficient Cross-Silo Federated Learning on Heterogeneous Client Devices Using a Computing Power-Aware Scheduler.
-    Jiuxiang Gu, Xiangxi Shi, Jason Kuen, Lu Qi, Ruiyi Zhang, Anqi Liu, Ani Nenkova, Tong Sun:
ADOPD: A Large-Scale Document Page Decomposition Dataset.
-    Yulai Zhao, Wenhao Zhan, Xiaoyan Hu, Ho-fung Leung, Farzan Farnia, Wen Sun, Jason D. Lee:
Provably Efficient CVaR RL in Low-rank MDPs.
-    Xiyao Wang, Ruijie Zheng, Yanchao Sun, Ruonan Jia, Wichayaporn Wongkamjan, Huazhe Xu, Furong Huang:
COPlanner: Plan to Roll Out Conservatively but to Explore Optimistically for Model-Based RL.
-    Chuan Wen, Dinesh Jayaraman, Yang Gao:
Can Transformers Capture Spatial Relations between Objects?
-    Marien Renaud, Jiaming Liu, Valentin De Bortoli, Andrés Almansa, Ulugbek Kamilov:
Plug-and-Play Posterior Sampling under Mismatched Measurement and Prior Models.
-    Jiaxu Wang, Ziyi Zhang, Renjing Xu:
Learning Robust Generalizable Radiance Field with Visibility and Feature Augmented Point Representation.
-    Shih-Ying Yeh , Yu-Guan Hsieh, Zhidong Gao, Bernard B. W. Yang, Giyeong Oh, Yanmin Gong:
Navigating Text-To-Image Customization: From LyCORIS Fine-Tuning to Model Evaluation.
-    Fabian Mentzer, David Minnen, Eirikur Agustsson, Michael Tschannen:
Finite Scalar Quantization: VQ-VAE Made Simple.
-    Matthieu Blanke, Marc Lelarge:
Interpretable Meta-Learning of Physical Systems.
-    Noam Levi, Alon Beck, Yohai Bar-Sinai:
Grokking in Linear Estimators - A Solvable Model that Groks without Understanding.
-    Qihao Liu, Adam Kortylewski, Yutong Bai, Song Bai, Alan L. Yuille:
Discovering Failure Modes of Text-guided Diffusion Models via Adversarial Search.
-    Roi Benita, Michael Elad, Joseph Keshet:
DiffAR: Denoising Diffusion Autoregressive Model for Raw Speech Waveform Generation.
-    Tianqi Liu, Yao Zhao, Rishabh Joshi, Misha Khalman, Mohammad Saleh, Peter J. Liu, Jialu Liu:
Statistical Rejection Sampling Improves Preference Optimization.
-    Takuya Ito, Soham Dan, Mattia Rigotti, James R. Kozloski, Murray Campbell:
On the generalization capacity of neural networks during generic multimodal reasoning.
-    Cheng Shi, Sibei Yang:
The Devil is in the Object Boundary: Towards Annotation-free Instance Segmentation using Foundation Models.
-    Wanru Zhao, Yihong Chen, Royson Lee, Xinchi Qiu, Yan Gao, Hongxiang Fan, Nicholas Donald Lane:
Breaking Physical and Linguistic Borders: Multilingual Federated Prompt Tuning for Low-Resource Languages.
-    Eduardo Dadalto Câmara Gomes, Marco Romanelli, Georg Pichler, Pablo Piantanida:
A Data-Driven Measure of Relative Uncertainty for Misclassification Detection.
-    Max F. Burg, Thomas Zenkel, Michaela Vystrcilová, Jonathan Oesterle, Larissa Höfling, Konstantin F. Willeke, Jan Lause, Sarah Müller, Paul G. Fahey, Zhiwei Ding, Kelli Restivo, Shashwat Sridhar, Tim Gollisch, Philipp Berens, Andreas S. Tolias, Thomas Euler, Matthias Bethge, Alexander S. Ecker:
Most discriminative stimuli for functional cell type clustering.
-    Xiaodan Chen, Xiucheng Li, Bo Liu, Zhipun Li:
Biased Temporal Convolution Graph Network for Time Series Forecasting with Missing Values.
-    Jung Hwan Heo, Jeonghoon Kim, Beomseok Kwon, Byeongwook Kim, Se Jung Kwon, Dongsoo Lee:
Rethinking Channel Dimensions to Isolate Outliers for Low-bit Weight Quantization of Large Language Models.
-    Xiangxin Zhou, Xiwei Cheng, Yuwei Yang, Yu Bao, Liang Wang, Quanquan Gu:
DecompOpt: Controllable and Decomposed Diffusion Models for Structure-based Molecular Optimization.
-    Yair Ori Gat, Nitay Calderon, Amir Feder, Alexander Chapanin, Amit Sharma, Roi Reichart:
Faithful Explanations of Black-box NLP Models Using LLM-generated Counterfactuals.
-    Robin Louiset, Edouard Duchesnay, Antoine Grigis, Pietro Gori:
Separating common from salient patterns with Contrastive Representation Learning.
-    Junwoo Park, Daehoon Gwak, Jaegul Choo, Edward Choi:
Self-Supervised Contrastive Learning for Long-term Forecasting.



- Aiwei Liu, Leyi Pan, Xuming Hu, Shiao Meng, Lijie Wen:
A Semantic Invariant Robust Watermark for Large Language Models.
- Zhiyuan Liu, Yi Wang, Zhiren Wang:
Fast Equilibrium of SGD in Generic Situations.
- Eliya Nachmani, Alon Levkovich, Roy Hirsch, Julian Salazar, Chulayuth Asawaroengchai, Soroosh Mariooryad, Ehud Rivlin, R. J. Skerry-Ryan, Michelle Tadmor Ramanovich:
Spoken Question Answering and Speech Continuation Using Spectrogram-Powered LLM.
- Francisco Vargas, Shreyas Padhy, Denis Blessing, Nikolas Nüsken:
Transport meets Variational Inference: Controlled Monte Carlo Diffusions.
- Hyungyu Lee, Saehyung Lee, Hyemi Jang, Junsung Park, Ho Bae, Sungroh Yoon:
DAFA: Distance-Aware Fair Adversarial Training.
- Yuexiao Ma, Huixia Li, Xiawu Zheng, Feng Ling, Xuefeng Xiao, Rui Wang, Shilei Wen, Fei Chao, Rongrong Ji:
AffineQuant: Affine Transformation Quantization for Large Language Models.
- Hansheng Xue, Vijini Mallawaarachchi, Lexing Xie, Vaibhav Rajan:
Encoding Unitig-level Assembly Graphs with Heterophilous Constraints for Metagenomic Contigs Binning.
- Uiwon Hwang, Jonghyun Lee, Juhyeon Shin, Sungroh Yoon:
SF(DA)2: Source-free Domain Adaptation Through the Lens of Data Augmentation.
- Song Xia, Yi Yu, Xudong Jiang, Henghui Ding:
Mitigating the Curse of Dimensionality for Certified Robustness via Dual Randomized Smoothing.
- Zhenwen Dai, Federico Tomasi, Sina Ghiassian:
In-context Exploration-Exploitation for Reinforcement Learning.
- Jun Nie, Yonggang Zhang, Zhen Fang, Tongliang Liu, Bo Han, Xinmei Tian:
Out-of-Distribution Detection with Negative Prompts.
- Gianluca Scarpellini, Ksenia Konyushkova, Claudio Fantacci, Thomas Paine, Yutian Chen, Misha Denil:
π2vec: Policy Representation with Successor Features.
- Feng Lu, Lijun Zhang, Xiangyuan Lan, Shuting Dong, Yaowei Wang, Chun Yuan:
Towards Seamless Adaptation of Pre-trained Models for Visual Place Recognition.
- Xiaohu Huang, Hao Zhou, Kun Yao, Kai Han:
FROSTER: Frozen CLIP is A Strong Teacher for Open-Vocabulary Action Recognition.
- Aoqi Zuo, Yiqing Li, Susan Wei, Mingming Gong:
Interventional Fairness on Partially Known Causal Graphs: A Constrained Optimization Approach.
- Weiyun Wang, Min Shi, Qingyun Li, Wenhui Wang, Zhenhang Huang, Linjie Xing, Zhe Chen, Hao Li, Xizhou Zhu, Zhiguo Cao, Yushi Chen, Tong Lu, Jifeng Dai, Yu Qiao:
The All-Seeing Project: Towards Panoptic Visual Recognition and Understanding of the Open World.
- Xiaoxiao Sun, Xingjian Leng, Zijian Wang, Yang Yang, Zi Huang, Liang Zheng:
CIFAR-10-Warehouse: Broad and More Realistic Testbeds in Model Generalization Analysis.
- Karan Mirakh, Sourav Ghosh, Dipanjan Das, Brojeshwar Bhowmick:
Task Planning for Visual Room Rearrangement under Partial Observability.
- Yi Heng Lim, Qi Zhu, Joshua Selfridge, Muhammad Firmansyah Kasim:
Parallelizing non-linear sequential models over the sequence length.
- Tianjiao Zhang, Huangjie Zheng, Jiangchao Yao, Xiangfeng Wang, Mingyuan Zhou, Ya Zhang, Yanfeng Wang:
Long-tailed Diffusion Models with Oriented Calibration.
- Dongyang Liu, Meina Kan, Shiguang Shan, Xilin Chen:
A Simple Romance Between Multi-Exit Vision Transformer and Token Reduction.
- Shengbo Wang, José H. Blanchet, Peter W. Glynn:
Optimal Sample Complexity for Average Reward Markov Decision Processes.
- Yizhou Jiang, Kunlin Hu, Tianren Zhang, Haichuan Gao, Yuqian Liu, Ying Fang, Feng Chen:
Spatio-Temporal Approximation: A Training-Free SNN Conversion for Transformers.
- Siqi Liu, Luke Marrs, Georgios Piliouras, Ian Gemp, Nicolas Heess:
NfgTransformer: Equivariant Representation Learning for Normal-form Games.
- Keming Lu, Hongyi Yuan, Zheng Yuan, Runji Lin, Junyang Lin, Chuanqi Tan, Chang Zhou, Jingren Zhou:
#InsTag: Instruction Tagging for Analyzing Supervised Fine-tuning of Large Language Models.
- Aleksandar Petrov, Philip Torr, Adel Bibi:
When Do Prompting and Prefix-Tuning Work? A Theory of Capabilities and Limitations.
- Jianhao Yan, Jin Xu, Chiyu Song, Chenming Wu, Yafu Li, Yue Zhang:
Understanding In-Context Learning from Repetitions.
- Hugo Cui, Florent Krzakala, Eric Vanden-Eijnden, Lenka Zdeborová:
Analysis of Learning a Flow-based Generative Model from Limited Sample Complexity.
- Hengjia Li, Yang Liu, Linxuan Xia, Yuqi Lin, Wenxiao Wang, Tu Zheng, Zheng Yang, Xiaohui Zhong, Xiaobo Ren, Xiaofei He:
Few-shot Hybrid Domain Adaptation of Image Generator.



- Yuxin Dong, Tieliang Gong, Hong Chen, Shujian Yu, Chen Li:
Rethinking Information-theoretic Generalization: Loss Entropy Induced PAC Bounds.
- Yaofo Chen, Shuaicheng Niu, Yaowei Wang, Shoukai Xu, Hengjie Song, Mingkui Tan:
Towards Robust and Efficient Cloud-Edge Elastic Model Adaptation via Selective Entropy Distillation.
- Marah I Abdin, Suriya Gunasekar, Varun Chandrasekaran, Jerry Li, Mert Yüksekgönül, Rahee Ghosh Peshawaria, Ranjita Naik, Besmira Nushi:
KITAB: Evaluating LLMs on Constraint Satisfaction for Information Retrieval.
- Jingyan Chen, Guanghui Zhu, Chunfeng Yuan, Yihua Huang:
Boosting Graph Anomaly Detection with Adaptive Message Passing.
- Giulio Franzese, Mustapha Bounoua, Pietro Michiardi:
MINDE: Mutual Information Neural Diffusion Estimation.
- Jae-Hong Lee, Joon-Hyuk Chang:
Continual Momentum Filtering on Parameter Space for Online Test-time Adaptation.
- Ziqi Gao, Tao Feng, Jiaxuan You, Chenyi Zi, Yan Zhou, Chen Zhang, Jia Li:
Deep Reinforcement Learning for Modelling Protein Complexes.
- Maarten Buyl, MaryBeth Defrance, Tijl De Bie:
fairret: a Framework for Differentiable Fairness Regularization Terms.
- Tomasz Limisiewicz, David Marecek, Tomás Musil:
Debiasing Algorithm through Model Adaptation.
- Yoni Choukroun, Lior Wolf:
A Foundation Model for Error Correction Codes.
- Xianfan Gu, Chuan Wen, Weirui Ye, Jiaming Song, Yang Gao:
Seer: Language Instructed Video Prediction with Latent Diffusion Models.
- Xuan Son Nguyen, Shuo Yang, Aymeric Histace:
Matrix Manifold Neural Networks++.
- Siyu Ren, Zhiyong Wu, Kenny Q. Zhu:
Emo: Earth Mover Distance Optimization for Auto-Regressive Language Modeling.
- Rui Li, Guoyin Wang, Jiwei Li:
Are Human-generated Demonstrations Necessary for In-context Learning?
- Naman Jain, Tianjun Zhang, Wei-Lin Chiang, Joseph E. Gonzalez, Koushik Sen, Ion Stoica:
LLM-Assisted Code Cleaning For Training Accurate Code Generators.
- Haoyue Bai, Yifei Ming, Julian Katz-Samuels, Yixuan Li:
HYPO: Hyperspherical Out-Of-Distribution Generalization.
- Jaemoo Choi, Jaewoong Choi, Myungjoo Kang:
Analyzing and Improving Optimal-Transport-based Adversarial Networks.
- Jiafei Lyu, Xiaoteng Ma, Le Wan, Runze Liu, Li Xiu, Zongqing Lu:
SEABO: A Simple Search-Based Method for Offline Imitation Learning.
- Yuxin Zhang, Lirui Zhao, Mingbao Lin, Yunyun Sun, Yiwu Yao, Xingjia Han, Jared Tanner, Shiwei Liu, Rongrong Ji:
Dynamic Sparse No Training: Training-Free Fine-tuning for Sparse LLMs.
- Bobby He, Thomas Hofmann:
Simplifying Transformer Blocks.
- Yuan Gao, Weizhong Zhang, Wenhan Luo, Lin Ma, Jin-Gang Yu, Gui-Song Xia, Jiayi Ma:
Aux-NAS: Exploiting Auxiliary Labels with Negligibly Extra Inference Cost.
- Qian Wang, Zhen Zhang, Zemin Liu, Shengliang Lu, Bingqiao Luo, Bingsheng He:
EX-Graph: A Pioneering Dataset Bridging Ethereum and X.
- Rong Dai, Yonggang Zhang, Ang Li, Tongliang Liu, Xun Yang, Bo Han:
Enhancing One-Shot Federated Learning Through Data and Ensemble Co-Boosting.
- Wenlong Liu, Tianyu Yang, Yuhang Wang, Qizhi Yu, Lei Zhang:
Symbol as Points: Panoptic Symbol Spotting via Point-based Representation.
- Sunwoo Kim, Shinhwan Kang, Fanchen Bu, Soo Yong Lee, Jaemin Yoo, Kijung Shin:
HypeBoy: Generative Self-Supervised Representation Learning on Hypergraphs.
- Dyah Adila, Changho Shin, Linrong Cai, Frederic Sala:
Zero-Shot Robustification of Zero-Shot Models.
- Junchi Yu, Ran He, Zhitao Ying:
Thought Propagation: an Analogical Approach to Complex Reasoning with Large Language Models.
- Yuxing Tian, Yiyuan Qi, Fan Guo:
FreeDyG: Frequency Enhanced Continuous-Time Dynamic Graph Model for Link Prediction.
- Vint Lee, Pieter Abbeel, Youngwoon Lee:
DreamSmooth: Improving Model-based Reinforcement Learning via Reward Smoothing.



Haonan Yu, Wei Xu:
VONet: Unsupervised Video Object Learning With Parallel U-Net Attention and Object-wise Sequential VAE.



Bo Zhang, Xinyu Cai, Jiakang Yuan, Donglin Yang, Jianfei Guo, Xiangchao Yan, Renqiu Xia, Botian Shi, Min Dou, Tao Chen, Si Liu, Junchi Yan, Yu Qiao:
ReSimAD: Zero-Shot 3D Domain Transfer for Autonomous Driving with Source Reconstruction and Target Simulation.



Zeren Chen, Ziqin Wang, Zhen Wang, Huayang Liu, Zhenfei Yin, Si Liu, Lu Sheng, Wanli Ouyang, Jing Shao:
Octavius: Mitigating Task Interference in MLLMs via LoRA-MoE.



Daniel Geng, Andrew Owens:
Motion Guidance: Diffusion-Based Image Editing with Differentiable Motion Estimators.



Meraj Hashemizadeh, Juan Ramirez, Rohan Sukumaran, Golnoosh Farnadi, Simon Lacoste-Julien, Jose Gallego-Posada:
Balancing Act: Constraining Disparate Impact in Sparse Models.



Mouïn Ben Ammar, Nacim Belkhir, Sebastian Popescu, Antoine Manzanera, Gianni Franchi:
NECO: NEural Collapse Based Out-of-distribution detection.



Yifan Feng, Yihe Luo, Shihui Ying, Yue Gao:
LightHGNN: Distilling Hypergraph Neural Networks into MLPs for 100x Faster Inference.



Ryo Ueda, Tadahiro Taniguchi:
Lewis's Signaling Game as beta-VAE For Natural Word Lengths and Segments.



Ruipeng Zhang, Ziqing Fan, Jiangchao Yao, Ya Zhang, Yanfeng Wang:
Domain-Inspired Sharpness-Aware Minimization Under Domain Shifts.



Yuying Ge, Sijie Zhao, Ziyun Zeng, Yixiao Ge, Chen Li, Xintao Wang, Ying Shan:
Making LLaMA SEE and Draw with SEED Tokenizer.



Haruo Hosoya:
A Cognitive Model for Learning Abstract Relational Structures from Memory-based Decision-Making Tasks.



Yanpeng Zhao, Siyu Gao, Yunbo Wang, Xiaokang Yang:
DynaVol: Unsupervised Learning for Dynamic Scenes through Object-Centric Voxelization.



Chenxi Sun, Hongyan Li, Yaliang Li, Shenda Hong:
TEST: Text Prototype Aligned Embedding to Activate LLM's Ability for Time Series.



Rundi Wu, Ruoshi Liu, Carl Vondrick, Changxi Zheng:
Sin3DM: Learning a Diffusion Model from a Single 3D Textured Shape.



Sotirios Panagiotis Chytas, Vishnu Suresh Lokhande, Vikas Singh:
Pooling Image Datasets with Multiple Covariate Shift and Imbalance.



Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos Garea, Matthieu Geist, Olivier Bachem:
On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes.



Qiang He, Tianyi Zhou, Meng Fang, Setareh Maghsudi:
Adaptive Regularization of Representation Rank as an Implicit Constraint of Bellman Equation.



Puja Trivedi, Mark Heimann, Rushil Anirudh, Danai Koutra, Jayaraman J. Thiagarajan:
Accurate and Scalable Estimation of Epistemic Uncertainty for Graph Neural Networks.



Ameya Daigavane, Song Kim, Mario Geiger, Tess E. Smidt:
Symphony: Symmetry-Equivariant Point-Centered Spherical Harmonics for 3D Molecule Generation.



Lukas Muttenthaler, Robert A. Vandermeulen, Qiuyi Zhang, Thomas Unterthiner, Klaus-Robert Müller:
Set Learning for Accurate and Calibrated Models.



Haozhe Chen, Junfeng Yang, Carl Vondrick, Chengzhi Mao:
INViTÉ: INterpret and Control Vision-Language Models with Text Explanations.



Jonah Philion, Xue Bin Peng, Sanja Fidler:
Trajeglish: Traffic Modeling as Next-Token Prediction.



Tian Yu Liu, Matthew Trager, Alessandro Achille, Pramuditha Perera, Luca Zancato, Stefano Soatto:
Meaning Representations from Trajectories in Autoregressive Models.



Lei You, Hei Victor Cheng:
SWAP: Sparse Entropic Wasserstein Regression for Robust Network Pruning.



Minh Pham, Kelly O. Marshall, Niv Cohen, Govind Mittal, Chinmay Hegde:
Circumventing Concept Erasure Methods For Text-To-Image Generative Models.



Chunjin Song, Bastian Wandt, Helge Rhodin:
Pose Modulated Avatars from Video.



Nicholas Konz, Maciej A. Mazurowski:
The Effect of Intrinsic Dataset Properties on Generalization: Unraveling Learning Differences Between Natural and Medical Images.



Keqiang Yan, Cong Fu, Xiaofeng Qian, Xiaoning Qian, Shuiwang Ji:
Complete and Efficient Graph Transformers for Crystal Material Property Prediction.



- Zheng Ding, Mengqi Zhang, Jiajun Wu, Zhuowen Tu:
Patched Denoising Diffusion Models For High-Resolution Image Synthesis.
- Soroush Abbasi Koohpayegani, Navaneet K. L., Parsa Nooralinejad, Soheil Kolouri, Hamed Pirsiavash:
NOLA: Compressing LoRA using Linear Combination of Random Basis.
- Mahdi Alikhasi, Levi Lelis:
Unveiling Options with Neural Network Decomposition.
- Junzhe Zhu, Peiye Zhuang, Sanmi Koyejo:
HIFA: High-fidelity Text-to-3D Generation with Advanced Diffusion Guidance.
- Yuren Cong, Mengmeng Xu, Christian Simon, Shoufa Chen, Jiawei Ren, Yanping Xie, Juan-Manuel Pérez-Rúa, Bodo Rosenhahn, Tao Xiang, Sen He:
FLATTEN: optical FLow-guided ATTENTION for consistent text-to-video editing.
- Xuefeng Du, Zhen Fang, Ilias Diakonikolas, Yixuan Li:
How Does Unlabeled Data Provably Help Out-of-Distribution Detection?
- Jean-Pierre R. Falet, Hae Beom Lee, Nikolay Malkin, Chen Sun, Dragos Secrieru, Dinghuai Zhang, Guillaume Lajoie, Yoshua Bengio:
Delta-AI: Local objectives for amortized inference in sparse graphical models.
- Hao Zhang, Fang Li, Samyak Rawlekar, Narendra Ahuja:
Learning Implicit Representation for Reconstructing Articulated Objects.
- Andrew Kirjner, Jason Yim, Raman Samusevich, Shahar Bracha, Tommi S. Jaakkola, Regina Barzilay, Ila R. Fiete:
Improving protein optimization with smoothed fitness landscapes.
- Vijay Lingam, Mohammad Sadegh Akhondzadeh, Aleksandar Bojchevski:
Rethinking Label Poisoning for GNNs: Pitfalls and Attacks.
- Songyao Jin, Feng Xie, Guangyi Chen, Biwei Huang, Zhengming Chen, Xinshuai Dong, Kun Zhang:
Structural Estimation of Partially Observed Linear Non-Gaussian Acyclic Model: A Practical Approach with Identifiability.
- Shicheng Liu, Minghui Zhu:
Meta Inverse Constrained Reinforcement Learning: Convergence Guarantee and Generalization Analysis.
- Shengyi Huang, Jiayi Weng, Rujikorn Charakorn, Min Lin, Zhongwen Xu, Santiago Ontañón:
Cleanba: A Reproducible and Efficient Distributed Reinforcement Learning Platform.
- Harry Zhao, Safa Alver, Harm van Seijen, Romain Laroche, Doina Precup, Yoshua Bengio:
Consciousness-Inspired Spatio-Temporal Abstractions for Better Generalization in Reinforcement Learning.
- Adam X. Yang, Maxime Robeyns, Xi Wang, Laurence Aitchison:
Bayesian Low-rank Adaptation for Large Language Models.
- Aidan Scannell, Riccardo Mereu, Paul Edmund Chang, Ella Tamir, Joni Pajarinen, Arno Solin:
Function-space Parameterization of Neural Networks for Sequential Learning.
- Shitong Duan, Xiaoyuan Yi, Peng Zhang, Tun Lu, Xing Xie, Ning Gu:
Denevil: towards Deciphering and Navigating the Ethical Values of Large Language Models via Instruction Learning.
- William Merrill, Ashish Sabharwal:
The Expressive Power of Transformers with Chain of Thought.
- Prajjwal Bhargava, Rohan Chitnis, Alborz Geramifard, Shagun Sodhani, Amy Zhang:
When should we prefer Decision Transformers for Offline Reinforcement Learning?
- Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu, Wei Xue, Shanghang Zhang, Jie Fu, Zhiyuan Liu:
ChatEval: Towards Better LLM-based Evaluators through Multi-Agent Debate.
- Tianwei Ni, Benjamin Eysenbach, Erfan Seyedsalehi, Michel Ma, Clement Gehring, Aditya Mahajan, Pierre-Luc Bacon:
Bridging State and History Representations: Understanding Self-Predictive RL.
- Jiaxu Zhang, Shaoli Huang, Zhigang Tu, Xin Chen, Xiaohang Zhan, Gang Yu, Ying Shan:
TapMo: Shape-aware Motion Generation of Skeleton-free Characters.
- Ronghao Dang, Jiangyan Feng, Haodong Zhang, Chongjian Ge, Lin Song, Lijun Gong, Chengju Liu, Qijun Chen, Feng Zhu, Rui Zhao, Yibing Song:
InstructDET: Diversifying Referring Object Detection with Generalized Instructions.
- Yuhui Li, Fangyun Wei, Jinjing Zhao, Chao Zhang, Hongyang Zhang:
RAIN: Your Language Models Can Align Themselves without Finetuning.
- Zhongpai Gao, Huayi Zhou, Abhishek Sharma, Meng Zheng, Benjamin Planche, Terrence Chen, Ziyan Wu:
PBADet: A One-Stage Anchor-Free Approach for Part-Body Association.
- Hugo Lebeau, Mohamed El Amine Seddik, José Henrique de Morais Goulart:
Performance Gaps in Multi-view Clustering under the Nested Matrix-Tensor Model.
- Cong Zhang, Zhiguang Cao, Wen Song, Yaoxin Wu, Jie Zhang:
Deep Reinforcement Learning Guided Improvement Heuristic for Job Shop Scheduling.
- Siqi Zhang, Sayantan Choudhury, Sebastian U. Stich, Nicolas Loizou:
Communication-Efficient Gradient Descent-Accent Methods for Distributed Variational Inequalities: Unified



Analysis and Local Updates.

- Rebekka Burkholz:
Batch normalization is sufficient for universal function approximation in CNNs.
- Shengding Hu, Xin Liu, Xu Han, Xinrong Zhang, Chaoqun He, Weilin Zhao, Yankai Lin, Ning Ding, Zebin Ou, Guoyang Zeng, Zhiyuan Liu, Maosong Sun:
Predicting Emergent Abilities with Infinite Resolution Evaluation.
- Dingyuan Shi, Yongxin Tong, Zimu Zhou, Ke Xu, Zheng Wang, Jieping Ye:
Graph-constrained diffusion for End-to-End Path Planning.
- Longtao Zheng, Rundong Wang, Xinrun Wang, Bo An:
Synapse: Trajectory-as-Exemplar Prompting with Memory for Computer Control.
- Simone Magistri, Tomaso Trinci, Albin Soutif-Cormerais, Joost van de Weijer, Andrew D. Bagdanov:
Elastic Feature Consolidation For Cold Start Exemplar-Free Incremental Learning.
- Scott Sussex, Pier Giuseppe Sessa, Anastasia Makarova, Andreas Krause:
Adversarial Causal Bayesian Optimization.
- Xin Yu, Yuan-Chen Guo, Yangguang Li, Ding Liang, Song-Hai Zhang, Xiaojuan Qi:
Text-to-3D with Classifier Score Distillation.
- Abudukelimu Wuerkaixi, Sen Cui, Jingfeng Zhang, Kunda Yan, Bo Han, Gang Niu, Lei Fang, Changshui Zhang, Masashi Sugiyama:
Accurate Forgetting for Heterogeneous Federated Continual Learning.
- Li Yang, Ruizheng Wu, Jiyong Li, Ying-Cong Chen:
GNeRP: Gaussian-guided Neural Reconstruction of Reflective Objects with Noisy Polarization Priors.
- Jia-Wang Bian, Wenjing Bian, Victor Adrian Prisacariu, Philip Torr:
Porf: Pose residual field for accurate Neural surface Reconstruction.
- Alexander Ashcroft, Ayan Das, Yulia Gryaditskaya, Zhiyu Qu, Yi-Zhe Song:
Modelling complex vector drawings with stroke-clouds.
- Yong Lin, Lu Tan, Yifan Hao, Honam Wong, Hanze Dong, Weizhong Zhang, Yujiu Yang, Tong Zhang:
Spurious Feature Diversification Improves Out-of-distribution Generalization.
- Hyeonho Jeong, Jong Chul Ye:
Ground-A-Video: Zero-shot Grounded Video Editing using Text-to-image Diffusion Models.
- Bohao Peng, Zhuotao Tian, Shu Liu, Ming-Chang Yang, Jiaya Jia:
Scalable Language Model with Generalized Continual Learning.
- Ziqiang Li, Hong Sun, Pengfei Xia, Heng Li, Beihao Xia, Yi Wu, Bin Li:
Efficient Backdoor Attacks for Deep Neural Networks in Real-world Scenarios.
- Rene Winchenbach, Nils Thuerey:
Symmetric Basis Convolutions for Learning Lagrangian Fluid Mechanics.
- Jason Piquenot, Aldo Moscatelli, Maxime Berar, Pierre Héroux, Romain Raveaux, Jean-Yves Ramel, Sébastien Adam:
G2N2 : Weisfeiler and Lehman go grammatical.
- Zhaomin Wu, Junyi Hou , Bingsheng He:
VertiBench: Advancing Feature Distribution Diversity in Vertical Federated Learning Benchmarks.
- Hiroki Furuta, Kuang-Huei Lee, Ofir Nachum, Yutaka Matsuo, Aleksandra Faust, Shixiang Shane Gu, Izzeddin Gur:
Multimodal Web Navigation with Instruction-Finetuned Foundation Models.
- Jianhao Yuan, Jie Zhang, Shuyang Sun, Philip Torr, Bo Zhao:
Real-Fake: Effective Training Data Synthesis Through Distribution Matching.
- Abhra Chaudhuri, Serban Georgescu, Anjan Dutta:
Learning Conditional Invariances through Non-Commutativity.
- Youliang Yuan, Wenxiang Jiao, Wenxuan Wang, Jen-tse Huang, Pinjia He, Shuming Shi, Zhaopeng Tu:
GPT-4 Is Too Smart To Be Safe: Stealthy Chat with LLMs via Cipher.
- Qingyue Zhao, Banghua Zhu:
Towards the Fundamental Limits of Knowledge Transfer over Finite Domains.
- Szu-Wei Fu, Kuo-Hsuan Hung, Yu Tsao, Yu-Chiang Frank Wang:
Self-Supervised Speech Quality Estimation and Enhancement Using Only Clean Speech.
- Yun-Hin Chan, Rui Zhou, Running Zhao, Zhihan Jiang, Edith C. H. Ngai:
Internal Cross-layer Gradients for Extending Homogeneity to Heterogeneity in Federated Learning.
- Zhiqian Tan, Yifan Zhang, Jingqin Yang, Yang Yuan:
Contrastive Learning is Spectral Clustering on Similarity Graph.
- Linus Bleistein, Agathe Guilloux:
On the Generalization and Approximation Capacities of Neural Controlled Differential Equations.
- Yingtao Zhang, Jialin Zhao, Wenjing Wu, Alessandro Muscoloni, Carlo Vittorio Cannistraci:
Epitopological learning and Cannistraci-Hebb network shape intelligence brain-inspired theory for ultra-sparse advantage in deep learning.

-      Yingtao Zhang, Haoli Bai, Haokun Lin, Jialin Zhao, Lu Hou, Carlo Vittorio Cannistraci: **Plug-and-Play: An Efficient Post-training Pruning Method for Large Language Models.**
-      Javier Rando, Florian Tramèr: **Universal Jailbreak Backdoors from Poisoned Human Feedback.**
-      Xindi Yang, Zeke Xie, Xiong Zhou, Boyu Liu, Buhua Liu, Yi Liu, Haoran Wang, Yunfeng Cai, Mingming Sun: **Neural Field Classifiers via Target Encoding and Classification Loss.**
-      Zhiyuan Cheng, Hongjun Choi, Shiwei Feng, James Chenhao Liang, Guanhong Tao, Dongfang Liu, Michael Zuzak, Xiangyu Zhang: **Fusion Is Not Enough: Single Modal Attacks on Fusion Models for 3D Object Detection.**
-      Yochai Yemini, Aviv Shamsian, Lior Bracha, Sharon Gannot, Ethan Fetaya: **LipVoicer: Generating Speech from Silent Videos Guided by Lip Reading.**
-      Daniel Bolya, Chaitanya Ryali, Judy Hoffman, Christoph Feichtenhofer: **Window Attention is Bugged: How not to Interpolate Position Embeddings.**
-      Lingxuan Wu, Xiao Yang, Yinpeng Dong, Liuwei Xie, Hang Su, Jun Zhu: **Embodied Active Defense: Leveraging Recurrent Feedback to Counter Adversarial Patches.**
-      Lukas Struppek, Dominik Hintersdorf, Kristian Kersting: **Be Careful What You Smooth For: Label Smoothing Can Be a Privacy Shield but Also a Catalyst for Model Inversion Attacks.**
-      Lin Yong, Fan Zhou, Lu Tan, Lintao Ma, Jianmeng Liu, Yansu He, Yuan Yuan, Yu Liu, James Y. Zhang, Yujiu Yang, Hao Wang: **Continuous Invariance Learning.**
-      George Stoica, Daniel Bolya, Jakob Bjorner, Pratik Ramesh, Taylor Hearn, Judy Hoffman: **Ziplt! Merging Models from Different Tasks without Training.**
-      Bowen Shi, Xiaopeng Zhang, Yaoming Wang, Jin Li, Wenrui Dai, Junni Zou, Hongkai Xiong, Qi Tian: **Hybrid Distillation: Connecting Masked Autoencoders with Contrastive Learners.**
-      Cheng Han, Qifan Wang, Yiming Cui, Wenguan Wang, Lifu Huang, Siyuan Qi, Dongfang Liu: **Facing the Elephant in the Room: Visual Prompt Tuning or Full finetuning?**
-      Zhiliang Peng, Wenhui Wang, Li Dong, Yaru Hao, Shaohan Huang, Shuming Ma, Qixiang Ye, Furu Wei: **Grounding Multimodal Large Language Models to the World.**
-      Tianyuan Zou, Zixuan Gu, Yu He, Hideaki Takahashi, Yang Liu, Ya-Qin Zhang: **VFLAIR: A Research Library and Benchmark for Vertical Federated Learning.**
-      Yue Cao, Tianlin Li, Xiaofeng Cao, Ivor W. Tsang, Yang Liu, Qing Guo: **IRAD: Implicit Representation-driven Image Resampling against Adversarial Attacks.**
-      Yihuan Mao, Chengjie Wu, Xi Chen, Hao Hu, Ji Jiang, Tianze Zhou, Tangjie Lv, Changjie Fan, Zhipeng Hu, Yi Wu, Yujing Hu, Chongjie Zhang: **Stylized Offline Reinforcement Learning: Extracting Diverse High-Quality Behaviors from Heterogeneous Datasets.**
-      Jiacheng Guo, Minshuo Chen, Huan Wang, Caiming Xiong, Mengdi Wang, Yu Bai: **Sample-Efficient Learning of POMDPs with Multiple Observations In Hindsight.**
-      Zecheng Wang, Che Wang, Zixuan Dong, Keith W. Ross: **Pre-training with Synthetic Data Helps Offline Reinforcement Learning.**
-      Weize Chen, Yusheng Su, Jingwei Zuo, Cheng Yang, Chenfei Yuan, Chi-Min Chan, Heyang Yu, Yaxi Lu, Yi-Hsin Hung, Chen Qian, Yujia Qin, Xin Cong, Ruobing Xie, Zhiyuan Liu, Maosong Sun, Jie Zhou: **AgentVerse: Facilitating Multi-Agent Collaboration and Exploring Emergent Behaviors.**
-      Yuzhen Mao, Martin Ester, Ke Li: **IceFormer: Accelerated Inference with Long-Sequence Transformers on CPUs.**
-      Wenhao Li: **Efficient Planning with Latent Diffusion.**
-      Etash Kumar Guha, Shlok Natarajan, Thomas Möllenhoff, Mohammad Emtiyaz Khan, Eugène Ndiaye: **Conformal Prediction via Regression-as-Classification.**
-      Zihan Zhong, Zhiqiang Tang, Tong He, Haoyang Fang, Chun Yuan: **Convolution Meets LoRA: Parameter Efficient Finetuning for Segment Anything Model.**
-      Xinyu Huang, Youcai Zhang, Jinyu Ma, Weiwei Tian, Rui Feng, Yuejie Zhang, Yaqian Li, Yandong Guo, Lei Zhang: **Tag2Text: Guiding Vision-Language Model via Image Tagging.**
-      Haowei Lin, Yijia Shao, Weinan Qian, Ningxin Pan, Yiduo Guo, Bing Liu: **Class Incremental Learning via Likelihood Ratio Based Task Prediction.**
-      Jun Chen, Haishan Ye, Mengmeng Wang, Tianxin Huang, Guang Dai, Ivor W. Tsang, Yong Liu: **Decentralized Riemannian Conjugate Gradient Method on the Stiefel Manifold.**
-      Jiechao Guan, Hui Xiong: **Improved Regret Bounds for Non-Convex Online-Within-Online Meta Learning.**
-      Ziheng Cheng, Xinmeng Huang, Pengfei Wu, Kun Yuan: **Momentum Benefits Non-iid Federated Learning Simply and Provably.**



- Shuhai Zhang, Yiliao Song, Jiahao Yang, Yuanqing Li, Bo Han, Mingkui Tan:
Detecting Machine-Generated Texts by Multi-Population Aware Optimization for Maximum Mean Discrepancy.
- Aadirupa Saha, Branislav Kveton:
Only Pay for What Is Uncertain: Variance-Adaptive Thompson Sampling.
- Xuan Ju, Ailing Zeng, Yuxuan Bian, Shaoteng Liu, Qiang Xu:
PnP Inversion: Boosting Diffusion-based Editing with 3 Lines of Code.
- Linwei Tao, Younan Zhu, Haolan Guo, Minjing Dong , Chang Xu:
A Benchmark Study on Calibration.
- Chengzhi Cao, Yinghao Fu, Sheng Xu, Ruimao Zhang, Shuang Li:
Enhancing Human-AI Collaboration Through Logic-Guided Reasoning.
- Haodong Lu, Dong Gong, Shuo Wang, Jason Xue, Lina Yao, Kristen Moore:
Learning with Mixture of Prototypes for Out-of-Distribution Detection.
- Leo Zhiyuan Zhao, Xueying Ding, B. Aditya Prakash:
PINNsFormer: A Transformer-Based Framework For Physics-Informed Neural Networks.
- Zican Hu, Zongzhang Zhang, Huaxiong Li, Chunlin Chen, Hongyu Ding, Zhi Wang:
Attention-Guided Contrastive Role Representations for Multi-agent Reinforcement Learning.
- Saptarshi Chakraborty, Peter L. Bartlett:
A Statistical Analysis of Wasserstein Autoencoders for Intrinsically Low-dimensional Data.
- Mert Yüksekgönül, Varun Chandrasekaran, Erik Jones, Suriya Gunasekar, Ranjita Naik, Hamid Palangi, Ece Kamar, Besmira Nushi:
Attention Satisfies: A Constraint-Satisfaction Lens on Factual Errors of Language Models.
- Sharut Gupta, Joshua Robinson, Derek Lim, Soledad Villar, Stefanie Jegelka:
Structuring Representation Geometry with Rotationally Equivariant Contrastive Learning.
- Dehao Yuan, Furong Huang, Cornelia Fermüller, Yiannis Aloimonos:
Decodable and Sample Invariant Continuous Object Encoder.
- Haitz Sáez de Ocáriz Borde, Anastasis Kratsios:
Neural Snowflakes: Universal Latent Graph Inference via Trainable Latent Geometries.
- Sharut Gupta, Stefanie Jegelka, David Lopez-Paz, Kartik Ahuja:
Context is Environment.
- Jungtaek Kim, Jeongbeen Yoon, Minsu Cho:
Generalized Neural Sorting Networks with Error-Free Differentiable Swap Functions.
- Zhaoyuan Yang, Zhengyang Yu, Zhiwei Xu, Jaskirat Singh, Jing Zhang, Dylan Campbell, Peter H. Tu, Richard Hartley:
IMPUS: Image Morphing with Perceptually-Uniform Sampling Using Diffusion Models.
- Yash Chandak, Shiv Shankar, Vasilis Syrgkanis, Emma Brunskill:
Adaptive Instrument Design for Indirect Experiments.
- Johnathan Xie, Yoonho Lee, Annie S. Chen, Chelsea Finn:
Self-Guided Masked Autoencoders for Domain-Agnostic Self-Supervised Learning.
- Hao Sun, Alihan Hüyük, Mihaela van der Schaar:
Query-Dependent Prompt Evaluation and Optimization with Offline Inverse RL.
- Zifeng Wang, Zichen Wang, Balasubramaniam Srinivasan, Vassilis N. Ioannidis, Huzeifa Rangwala, Rishita Anubhai:
BioBridge: Bridging Biomedical Foundation Models via Knowledge Graphs.
- Zhuoyan Xu, Zhenmei Shi, Junyi Wei, Fangzhou Mu, Yin Li, Yingyu Liang:
Towards Few-Shot Adaptation of Foundation Models via Multitask Finetuning.
- Hwanwoo Kim, Xin Zhang, Jiwei Zhao, Qinglong Tian:
ReTaSA: A Nonparametric Functional Estimation Approach for Addressing Continuous Target Shift.
- Xuangeng Chu, Yu Li, Ailing Zeng, Tianyu Yang, Lijian Lin, Yunfei Liu, Tatsuya Harada:
GPAvatar: Generalizable and Precise Head Avatar from Image(s).
- Congpei Qiu, Tong Zhang, Yanhao Wu, Wei Ke, Mathieu Salzmann, Sabine Süsstrunk:
Mind Your Augmentation: The Key to Decoupling Dense Self-Supervised Learning.
- Bolian Li, Ruqi Zhang:
Entropy-MCMC: Sampling from Flat Basins with Ease.
- Jiale Zhang, Yulun Zhang, Jinjin Gu, Jiahua Dong, Linghe Kong, Xiaokang Yang:
Xformer: Hybrid X-Shaped Transformer for Image Denoising.
- Seunghan Lee, Taeyoung Park, Kibok Lee:
Learning to Embed Time Series Patches Independently.
- Wei Huang, Ye Shi, Zhongyi Cai, Taiji Suzuki:
Understanding Convergence and Generalization in Federated Learning through Feature Learning Theory.
- Joar Max Viktor Skalse, Alessandro Abate:
Quantifying the Sensitivity of Inverse Reinforcement Learning to Misspecification.



- Quanrui Rao, Lin Wang, Wuying Liu:
Rethinking CNN's Generalization to Backdoor Attack from Frequency Domain.
- Guangyi Chen, Yuke Li, Xiao Liu, Zijian Li, Eman Al Suradi, Donglai Wei, Kun Zhang:
LLCP: Learning Latent Causal Processes for Reasoning-based Video Question Answer.
- Yichong Leng, Zhifang Guo, Kai Shen, Zeqian Ju, Xu Tan, Eric Liu, Yufei Liu, Dongchao Yang, Leying Zhang, Kaitao Song, Lei He, Xiangyang Li, Sheng Zhao, Tao Qin, Jiang Bian:
PromptTTS 2: Describing and Generating Voices with Text Prompt.
- Samuel Pegg, Kai Li, Xiaolin Hu:
RTFS-Net: Recurrent Time-Frequency Modelling for Efficient Audio-Visual Speech Separation.
- Jiaxin Cheng, Tianjun Xiao, Tong He:
Consistent Video-to-Video Transfer Using Synthetic Dataset.
- Simin Li, Jun Guo, Jingqiao Xiu, Ruixiao Xu, Xin Yu, Jiakai Wang, Aishan Liu, Yaodong Yang, Xianglong Liu:
Byzantine Robust Cooperative Multi-Agent Reinforcement Learning as a Bayesian Game.
- Jijin Hu, Ke Li, Yonggang Qi, Yi-Zhe Song:
Scale-Adaptive Diffusion Model for Complex Sketch Synthesis.
- Guangsheng Bao, Yanbin Zhao, Zhiyang Teng, Linyi Yang, Yue Zhang:
Fast-DetectGPT: Efficient Zero-Shot Detection of Machine-Generated Text via Conditional Probability Curvature.
- Lu Chen, Siyu Lou, Benhao Huang, Quanshi Zhang:
Defining and extracting generalizable interaction primitives from DNNs.
- Manh Luong, Khai Nguyen, Nhat Ho, Gholamreza Haffari, Dinh Phung, Lizhen Qu:
Revisiting Deep Audio-Text Retrieval Through the Lens of Transportation.
- Yang Liu, Muzhi Zhu, Hengtao Li, Hao Chen, Xinlong Wang, Chunhua Shen:
Matcher: Segment Anything with One Shot Using All-Purpose Feature Matching.
- Zecheng Tang, Chenfei Wu, Juntao Li, Nan Duan:
LayoutNUWA: Revealing the Hidden Layout Expertise of Large Language Models.
- Yiyang Zhou, Chenhang Cui, Jaehong Yoon, Linjun Zhang, Zhun Deng, Chelsea Finn, Mohit Bansal, Huaxiu Yao:
Analyzing and Mitigating Object Hallucination in Large Vision-Language Models.
- Mingqing Xiao, Qingyan Meng, Zongpeng Zhang, Di He, Zhouchen Lin:
Hebbian Learning based Orthogonal Projection for Continual Learning of Spiking Neural Networks.
- Yung-Sung Chuang, Yujia Xie, Hongyin Luo, Yoon Kim, James R. Glass, Pengcheng He:
DoLa: Decoding by Contrasting Layers Improves Factuality in Large Language Models.
- Tianhao Wu, Chuanxia Zheng, Tat-Jen Cham:
PanoDiffusion: 360-degree Panorama Outpainting via Diffusion.
- Jiaming Liu, Senqiao Yang, Peidong Jia, Renrui Zhang, Ming Lu, Yandong Guo, Wei Xue, Shanghang Zhang:
ViDA: Homeostatic Visual Domain Adapter for Continual Test Time Adaptation.
- Shaokun Zhang, Xiaobo Xia, Zhaoqing Wang, Ling-Hao Chen, Jiale Liu, Qingyun Wu, Tongliang Liu:
IDEAL: Influence-Driven Selective Annotations Empower In-Context Learners in Large Language Models.
- Yuhui Zhang, Elaine Sui, Serena Yeung:
Connect, Collapse, Corrupt: Learning Cross-Modal Tasks with Uni-Modal Data.
- Sunghwan Hong, Seokju Cho, Seungryong Kim, Stephen Lin:
Unifying Feature and Cost Aggregation with Transformers for Semantic and Visual Correspondence.
- Yang He, Joey Tianyi Zhou:
Data-independent Module-aware Pruning for Hierarchical Vision Transformers.
- Kai Xu, Rongyu Chen, Gianni Franchi, Angela Yao:
Scaling for Training Time and Post-hoc Out-of-distribution Detection Enhancement.
- Mingjie Sun, Zhuang Liu, Anna Bair, J. Zico Kolter:
A Simple and Effective Pruning Approach for Large Language Models.
- Rohin Manvi, Samar Khanna, Gengchen Mai, Marshall Burke, David B. Lobell, Stefano Ermon:
GeoLLM: Extracting Geospatial Knowledge from Large Language Models.
- Jiahao Li, Hao Tan, Kai Zhang, Zexiang Xu, Fujun Luan, Yinghao Xu, Yicong Hong, Kalyan Sunkavalli, Greg Shakhnarovich, Sai Bi:
Instant3D: Fast Text-to-3D with Sparse-view Generation and Large Reconstruction Model.
- Qinbin Li, Chulin Xie, Xiaojun Xu, Xiaoyuan Liu, Ce Zhang, Bo Li, Bingsheng He, Dawn Song:
Effective and Efficient Federated Tree Learning on Hybrid Data.
- Kaixiang Zheng, En-Hui Yang:
Knowledge Distillation Based on Transformed Teacher Matching.
- Cheng Han, James Chenhao Liang, Qifan Wang, Majid Rabbani, Sohail A. Dianat, Raghuvir Rao, Ying Nian Wu, Dongfang Liu:
Image Translation as Diffusion Visual Programmers.



- Chengzhi Mao, Carl Vondrick, Hao Wang, Junfeng Yang:
Radar: geneRative AI Detection via Rewriting.
- Zhengxiang Shi, Aldo Lipani:
DePT: Decomposed Prompt Tuning for Parameter-Efficient Fine-tuning.
- Siming Yan, Chen Song, Youkang Kong, Qixing Huang:
Multi-View Representation is What You Need for Point-Cloud Pre-Training.
- Haoyu Lu, Guoxing Yang, Nanyi Fei, Yuqi Huo, Zhiwu Lu, Ping Luo, Mingyu Ding:
VDT: General-purpose Video Diffusion Transformers via Mask Modeling.
- Yanqi Bao, Tianyu Ding, Jing Huo, Wenbin Li, Yuxin Li, Yang Gao:
InsertNeRF: Instilling Generalizability into NeRF with HyperNet Modules.
- Xiang Hu, Qingyang Zhu, Kewei Tu, Wei Wu:
Augmenting Transformers with Recursively Composed Multi-grained Representations.
- Zipeng Wang, Xuehui Yu, Xumeng Han, Wenwen Yu, Zhixun Huang, Jianbin Jiao, Zhenjun Han:
P2Seg: Pointly-supervised Segmentation via Mutual Distillation.
- Shiyu Wang, Haixu Wu, Xiaoming Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y. Zhang, Jun Zhou:
TimeMixer: Decomposable Multiscale Mixing for Time Series Forecasting.
- Shaofeng Zhang, Jinfang Huang, Qiang Zhou, Zhibin Wang, Fan Wang, Jiebo Luo, Junchi Yan:
Continuous-Multiple Image Outpainting in One-Step via Positional Query and A Diffusion-based Approach.
- Linwei Chen, Lin Gu, Ying Fu:
When Semantic Segmentation Meets Frequency Aliasing.
- Yili Wang, Kaixiong Zhou, Ninghao Liu, Ying Wang, Xin Wang:
Efficient Sharpness-Aware Minimization for Molecular Graph Transformer Models.
- Chenjie Cao, Xinlin Ren, Yanwei Fu:
MVSFormer++: Revealing the Devil in Transformer's Details for Multi-View Stereo.
- Yuan Yuan, Chenyang Shao, Jingtao Ding, Depeng Jin, Yong Li:
Spatio-Temporal Few-Shot Learning via Diffusive Neural Network Generation.
- Soichiro Kumano, Hiroshi Kera, Toshihiko Yamasaki:
Theoretical Understanding of Learning from Adversarial Perturbations.
- Guibin Zhang, Kun Wang, Wei Huang, Yanwei Yue, Yang Wang, Roger Zimmermann, Aojun Zhou, Dawei Cheng, Jin Zeng, Yuxuan Liang:
Graph Lottery Ticket Automated.
- Haonan Qiu, Menghan Xia, Yong Zhang, Yingqing He, Xintao Wang, Ying Shan, Ziwei Liu:
FreeNoise: Tuning-Free Longer Video Diffusion via Noise Rescheduling.
- Shikun Sun, Longhui Wei, Zhicai Wang, Zixuan Wang, Junliang Xing, Jia Jia, Qi Tian:
Inner Classifier-Free Guidance and Its Taylor Expansion for Diffusion Models.
- Changyao Tian, Chenxin Tao, Jifeng Dai, Hao Li, Ziheng Li, Lewei Lu, Xiaogang Wang, Hongsheng Li, Gao Huang, Xizhou Zhu:
ADDP: Learning General Representations for Image Recognition and Generation with Alternating Denoising Diffusion Process.
- Yixuan Weng, Minjun Zhu, Fei Xia, Bin Li, Shizhu He, Kang Liu, Jun Zhao:
Mastering Symbolic Operations: Augmenting Language Models with Compiled Neural Networks.
- Zeyu Liu, Gourav Datta, Anni Li, Peter Anthony Beerel:
LMUFormer: Low Complexity Yet Powerful Spiking Model With Legendre Memory Units.
- Jiawei Sun, Kailai Li, Ruoxin Chen, Jie Li, Chentao Wu, Yue Ding, Junchi Yan:
InterpGNN: Understand and Improve Generalization Ability of Transductive GNNs through the Lens of Interplay between Train and Test Nodes.
- Dennis Wu, Jerry Yao-Chieh Hu, Weijian Li, Bo-Yu Chen, Han Liu:
STanHop: Sparse Tandem Hopfield Model for Memory-Enhanced Time Series Prediction.
- Kuofeng Gao, Yang Bai, Jindong Gu, Shu-Tao Xia, Philip Torr, Zhifeng Li, Wei Liu:
Inducing High Energy-Latency of Large Vision-Language Models with Verbose Images.
- Haeyong Kang, Jaehong Yoon, Dahyun Kim, Sung Ju Hwang, Chang D. Yoo:
Progressive Fourier Neural Representation for Sequential Video Compilation.
- Tingting Jiang, Qi Xu, Xuming Ran, Jiangrong Shen, Pan Lv, Qiang Zhang, Gang Pan:
Adaptive deep spiking neural network with global-local learning via balanced excitatory and inhibitory mechanism.
- Zehao Dou, Yang Song:
Diffusion Posterior Sampling for Linear Inverse Problem Solving: A Filtering Perspective.
- Yuhan Helena Liu, Aristide Baratin, Jonathan Cornford, Stefan Mihalas, Eric Shea-Brown, Guillaume Lajoie:
How connectivity structure shapes rich and lazy learning in neural circuits.
- Xilie Xu, Keyi Kong, Ning Liu, Lizhen Cui, Di Wang, Jingfeng Zhang, Mohan S. Kankanhalli:
An LLM can Fool Itself: A Prompt-Based Adversarial Attack.



- Xilie Xu, Jingfeng Zhang, Mohan S. Kankanhalli: **AutoLoRa: An Automated Robust Fine-Tuning Framework.**
- Shuai Yang, Yukang Chen, Luozhou Wang, Shu Liu, Ying-Cong Chen: **Denoising Diffusion Step-aware Models.**
- Feiyang Kang, Hoang Anh Just, Yifan Sun, Himanshu Jahagirdar, Yuanzhi Zhang, Rongxing Du, Anit Kumar Sahu, Ruoxi Jia: **Get more for less: Principled Data Selection for Warming Up Fine-Tuning in LLMs.**
- Yang Fu, Shalini De Mello, Xuetong Li, Amey Kulkarni, Jan Kautz, Xiaolong Wang, Sifei Liu: **3D Reconstruction with Generalizable Neural Fields using Scene Priors.**
- Xiu-Chuan Li, Kun Zhang, Tongliang Liu: **Causal Structure Recovery with Latent Variables under Milder Distributional and Graphical Assumptions.**
- Hsi-Ai Tsao, Lei Hsiung, Pin-Yu Chen, Si Liu, Tsung-Yi Ho: **AutoVP: An Automated Visual Prompting Framework and Benchmark.**
- Yuanhao Xiong, Long Zhao, Boqing Gong, Ming-Hsuan Yang, Florian Schroff, Ting Liu, Cho-Jui Hsieh, Liangzhe Yuan: **Structured Video-Language Modeling with Temporal Grouping and Spatial Grounding.**
- Ziyang Luo, Can Xu, Pu Zhao, Qingfeng Sun, Xiubo Geng, Wenxiang Hu, Chongyang Tao, Jing Ma, Qingwei Lin, Dixin Jiang: **WizardCoder: Empowering Code Large Language Models with Evol-Instruct.**
- Yihang Chen, Lukas Mauch: **Order-Preserving GFlowNets.**
- Ling Yang, Ye Tian, Minkai Xu, Zhongyi Liu, Shenda Hong, Wei Qu, Wentao Zhang, Bin Cui, Muhan Zhang, Jure Leskovec: **VQGraph: Rethinking Graph Representation Space for Bridging GNNs and MLPs.**
- Nilesh Gupta, Devvrit, Ankit Singh Rawat, Srinadh Bhojanapalli, Prateek Jain, Inderjit S. Dhillon: **Dual-Encoders for Extreme Multi-label Classification.**
- Ali Hatamizadeh, Greg Heinrich, Hongxu Yin, Andrew Tao, José M. Álvarez, Jan Kautz, Pavlo Molchanov: **FasterViT: Fast Vision Transformers with Hierarchical Attention.**
- Qi Yan, Raihan Seraj, Jiawei He, Lili Meng, Tristan Sylvain: **AutoCast++: Enhancing World Event Prediction with Zero-shot Ranking-based Context Retrieval.**
- Thomas Laurent, James von Brecht, Xavier Bresson: **Feature Collapse.**
- Eric Todd, Millicent L. Li, Arnab Sen Sharma, Aaron Mueller, Byron C. Wallace, David Bau: **Function Vectors in Large Language Models.**
- Nikhil Prakash, Tamar Rott Shaham, Tal Haklay, Yonatan Belinkov, David Bau: **Fine-Tuning Enhances Existing Mechanisms: A Case Study on Entity Tracking.**
- Zhiwei Deng, Ting Chen, Yang Li: **Perceptual Group Tokenizer: Building Perception with Iterative Grouping.**
- William Yang, Byron Zhang, Olga Russakovsky: **ImageNet-OOD: Deciphering Modern Out-of-Distribution Detection Algorithms.**
- Yi Sui, Tongzi Wu, Jesse C. Cresswell, Ga Wu, George Stein, Xiao Shi Huang, Xiaochen Zhang, Maksims Volkovs: **Self-supervised Representation Learning from Random Data Projectors.**
- Matan Atzmon, Jiahui Huang, Francis Williams, Or Litany: **Approximately Piecewise E(3) Equivariant Point Networks.**
- Luke Nicholas Darlow, Qiwen Deng, Ahmed Hassan, Martin Asenov, Rajkarn Singh, Artjom Joosen, Adam Barker, Amos J. Storkey: **DAM: Towards a Foundation Model for Forecasting.**
- Tanvir Mahmud, Saeed Amizadeh, Kazuhito Koishida, Diana Marculescu: **Weakly-supervised Audio Separation via Bi-modal Semantic Similarity.**
- Marco Jiralerspong, Bilun Sun, Danilo Vucetic, Tianyu Zhang, Yoshua Bengio, Gauthier Gidel, Nikolay Malkin: **Expected flow networks in stochastic environments and two-player zero-sum games.**
- Chen Geng, Hong-Xing Yu, Sida Peng, Xiaowei Zhou, Jiajun Wu: **Neural Polynomial Gabor Fields for Macro Motion Analysis.**
- Titas Anciukevicius, Fabian Manhardt, Federico Tombari, Paul Henderson: **Denoising Diffusion via Image-Based Rendering.**
- Hanwen Jiang, Zhenyu Jiang, Yue Zhao, Qixing Huang: **LEAP: Liberate Sparse-View 3D Modeling from Camera Poses.**
- Grégoire Delétang, Anian Ruoss, Paul-Ambroise Duquenne, Elliot Catt, Tim Genewein, Christopher Mattern, Jordi Grau-Moya, Li Kevin Wenliang, Matthew Aitchison, Laurent Orseau, Marcus Hutter, Joel Veness: **Language Modeling Is Compression.**
- Francis Engelmann, Fabian Manhardt, Michael Niemeyer, Keisuke Tateno, Federico Tombari: **OpenNeRF: Open Set 3D Neural Scene Segmentation with Pixel-Wise Features and Rendered Novel Views.**



Guillaume Bono, Leonid Antsfeld, Assem Sadek, Gianluca Monaci, Christian Wolf:
Learning with a Mole: Transferable latent spatial representations for navigation without reconstruction.

Stone Tao, Arth Shukla, Tse-kai Chan, Hao Su:
Reverse Forward Curriculum Learning for Extreme Sample and Demo Efficiency.

Jianzhe Lin, Maurice Diesendruck, Liang Du, Robin Abraham:
BatchPrompt: Accomplish more with less.

Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V. Le, Denny Zhou, Xinyun Chen:
Large Language Models as Optimizers.

Elisa Kreiss, Eric Zelikman, Christopher Potts, Nick Haber:
ContextRef: Evaluating Referenceless Metrics for Image Description Generation.

Federico Errica, Mathias Niepert:
Tractable Probabilistic Graph Representation Learning with Graph-Induced Sum-Product Networks.

Xian Liu, Jian Ren, Aliaksandr Siarohin, Ivan Skorokhodov, Yanyu Li, Dahua Lin, Xihui Liu, Ziwei Liu, Sergey Tulyakov:
HyperHuman: Hyper-Realistic Human Generation with Latent Structural Diffusion.

Kyle Vedder, Neehar Peri, Nathaniel Chodosh, Ishan Khatri, Eric Eaton, Dinesh Jayaraman, Yang Liu, Deva Ramanan, James Hays:
ZeroFlow: Scalable Scene Flow via Distillation.

Jiayu Xiao, Henglei Lv, Liang Li, Shuhui Wang, Qingming Huang:
R&B: Region and Boundary Aware Zero-shot Grounded Text-to-image Generation.

Xinyue Xu, Yi Qin, Lu Mi, Hao Wang, Xiaomeng Li:
Energy-Based Concept Bottleneck Models: Unifying Prediction, Concept Intervention, and Probabilistic Interpretations.

Xin Zhang, Dong Zhang, Shimin Li, Yaqian Zhou, Xipeng Qiu:
SpeechTokenizer: Unified Speech Tokenizer for Speech Language Models.

Zijie Pan, Jiachen Lu, Xiatian Zhu, Li Zhang:
Enhancing High-Resolution 3D Generation through Pixel-wise Gradient Clipping.

Yuan-Hong Liao, David Acuna, Rafid Mahmood, James Lucas, Viraj Prabhu, Sanja Fidler:
Transferring Labels to Solve Annotation Mismatches Across Object Detection Datasets.

Yannis Kalantidis, Mert Bülent Sarıyıldız, Rafael S. Rezende, Philippe Weinzaepfel, Diane Larlus, Gabriela Csurka:
Weatherproofing Retrieval for Localization with Generative AI and Geometric Consistency.

Jie Xiao, Ruili Feng, Han Zhang, Zhiheng Liu, Zhantao Yang, Yurui Zhu, Xueyang Fu, Kai Zhu, Yu Liu, Zheng-Jun Zha:
DreamClean: Restoring Clean Image Using Deep Diffusion Prior.

Nan Ding, Tomer Levinboim, Jialin Wu, Sebastian Goodman, Radu Soricut:
CausalLM is not optimal for in-context learning.

Guillaume Bono, Leonid Antsfeld, Boris Chidlovskii, Philippe Weinzaepfel, Christian Wolf:
End-to-End (Instance)-Image Goal Navigation through Correspondence as an Emergent Phenomenon.

Chunming He, Kai Li, Yachao Zhang, Yulun Zhang, Chenyu You, Zhenhua Guo, Xiu Li, Martin Danelljan, Fisher Yu:
Strategic Preys Make Acute Predators: Enhancing Camouflaged Object Detectors by Generating Camouflaged Objects.

Shuvendu Roy, Ali Etemad:
Consistency-guided Prompt Learning for Vision-Language Models.

Zeyu Yang, Hongye Yang, Zijie Pan, Li Zhang:
Real-time Photorealistic Dynamic Scene Representation and Rendering with 4D Gaussian Splatting.

Sharon Lee, Yunzhi Zhang, Shangzhe Wu, Jiajun Wu:
Language-Informed Visual Concept Learning.

Byeonghwi Kim, Minhyuk Seo, Jonghyun Choi:
Online Continual Learning for Interactive Instruction Following Agents.

Samyadeep Basu, Nanxuan Zhao, Vlad I. Morariu, Soheil Feizi, Varun Manjunatha:
Localizing and Editing Knowledge In Text-to-Image Generative Models.

Weiyang Liu, Zeju Qiu, Yao Feng, Yuliang Xiu, Yuxuan Xue, Longhui Yu, Haiwen Feng, Zhen Liu, Juyeon Heo, Songyou Peng, Yandong Wen, Michael J. Black, Adrian Weller, Bernhard Schölkopf:
Parameter-Efficient Orthogonal Finetuning via Butterfly Factorization.

Marin Scalbert, Maria Vakalopoulou, Florent Couzinie-Devy:
Towards domain-invariant Self-Supervised Learning with Batch Styles Standardization.

Tianyu Li, Peijin Jia, Bangjun Wang, Li Chen, Kun Jiang, Junchi Yan, Hongyang Li:
LaneSegNet: Map Learning with Lane Segment Perception for Autonomous Driving.

Man Yao , Jiakui Hu, Tianxiang Hu, Yifan Xu, Zhaokun Zhou, Yonghong Tian, Bo Xu, Guoqi Li:
Spike-driven Transformer V2: Meta Spiking Neural Network Architecture Inspiring the Design of Next-generation Neuromorphic Chips.

Andrew F. Luo, Margaret M. Henderson, Michael J. Tarr, Leila Wehbe:
BrainSCUBA: Fine-Grained Natural Language Captions of Visual Cortex Selectivity.



- Xueyi Liu, Li Yi:
GeneOH Diffusion: Towards Generalizable Hand-Object Interaction Denoising via Denoising Diffusion.
- Senmao Li, Joost van de Weijer, Taihang Hu, Fahad Shahbaz Khan, Qibin Hou, Yaxing Wang, Jian Yang:
Get What You Want, Not What You Don't: Image Content Suppression for Text-to-Image Diffusion Models.
- Chengming Hu, Haolun Wu, Xuan Li, Chen Ma, Xi Chen, Boyu Wang, Jun Yan, Xue Liu:
Less or More From Teacher: Exploiting Trilateral Geometry For Knowledge Distillation.
- Yufei Gu, Xiaoqing Zheng, Tomaso Aste:
Unraveling the Enigma of Double Descent: An In-depth Analysis through the Lens of Learned Feature Space.
- Xingyu Liu, Deepak Pathak, Ding Zhao:
Meta-Evolve: Continuous Robot Evolution for One-to-many Policy Transfer.
- Bowen Yin, Xuying Zhang, Zhong-Yu Li, Li Liu, Ming-Ming Cheng, Qibin Hou:
DFormer: Rethinking RGBD Representation Learning for Semantic Segmentation.
- Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Minlie Huang, Nan Duan, Weizhu Chen:
ToRA: A Tool-Integrated Reasoning Agent for Mathematical Problem Solving.
- Ganchao Wei:
Bayesian Bi-clustering of Neural Spiking Activity with Latent Structures.
- Sascha Marton, Stefan Lüdtke, Christian Bartelt, Heiner Stuckenschmidt:
GRANDE: Gradient-Based Decision Tree Ensembles for Tabular Data.
- Takeru Miyato, Bernhard Jaeger, Max Welling, Andreas Geiger:
GTA: A Geometry-Aware Attention Mechanism for Multi-View Transformers.
- Zihao Zhu, Mingda Zhang, Shaokui Wei, Bingzhe Wu, Baoyuan Wu:
VDC: Versatile Data Cleanser based on Visual-Linguistic Inconsistency by Multimodal Large Language Models.
- Ziying Guo, Kai Wang, George Cazenavette, Hui Li, Kaipeng Zhang, Yang You:
Towards Lossless Dataset Distillation via Difficulty-Aligned Trajectory Matching.
- Jiacheng Chen, Zeyuan Ma, Hongshu Guo, Yining Ma, Jie Zhang, Yue-Jiao Gong:
SYMBOL: Generating Flexible Black-Box Optimizers through Symbolic Equation Learning.
- Heejun Lee, Jina Kim, Jeffrey Willette, Sung Ju Hwang:
SEA: Sparse Linear Attention with Estimated Attention Mask.
- Xiong Zhou, Xianming Liu, Feilong Zhang, Gang Wu, Deming Zhai, Junjun Jiang, Xiangyang Ji:
Zero-Mean Regularized Spectral Contrastive Learning: Implicitly Mitigating Wrong Connections in Positive-Pair Graphs.
- Xiong Zhou, Xianming Liu, Hao Yu, Jialiang Wang, Zeke Xie, Junjun Jiang, Xiangyang Ji:
Variance-enlarged Poisson Learning for Graph-based Semi-Supervised Learning with Extremely Sparse Labeled Data.
- Jiyang Zheng, Yu Yao, Bo Han, Dadong Wang, Tongliang Liu:
Enhancing Contrastive Learning for Ordinal Regression via Ordinal Content Preserved Data Augmentation.
- Ning Miao, Yee Whye Teh, Tom Rainforth:
SelfCheck: Using LLMs to Zero-Shot Check Their Own Step-by-Step Reasoning.
- Yiming Xie, Varun Jampani, Lei Zhong, Deqing Sun, Huaizu Jiang:
OmniControl: Control Any Joint at Any Time for Human Motion Generation.
- Renbo Tu, Colin White, Jean Kossaifi, Boris Bonev, Gennady Pekhimenko, Kamyar Azizzadenesheli, Anima Anandkumar:
Guaranteed Approximation Bounds for Mixed-Precision Neural Operators.
- Junoh Lee, Hyunjun Jung, Jin-Hwi Park, Inhwan Bae, Hae-Gon Jeon:
Geometry-Aware Projective Mapping for Unbounded Neural Radiance Fields.
- David Ireland, Giovanni Montana:
REValueD: Regularised Ensemble Value-Decomposition for Factorisable Markov Decision Processes.
- Borui Zhang, Wenzhao Zheng, Jie Zhou, Jiwen Lu:
Path Choice Matters for Clear Attributions in Path Methods.
- Xingbin Liu, Jinghao Zhou, Tao Kong, Xianming Lin, Rongrong Ji:
Exploring Target Representations for Masked Autoencoders.
- Yuka Hashimoto, Sho Sonoda, Isao Ishikawa, Atsushi Nitanda, Taiji Suzuki:
Koopman-based generalization bound: New aspect for full-rank weights.
- Kai Chen, Chunwei Wang, Kuo Yang, Jianhua Han, Lanqing Hong, Fei Mi, Hang Xu, Zhengying Liu, Wenyong Huang, Zhenguo Li, Dit-Yan Yeung, Lifeng Shang:
Gaining Wisdom from Setbacks: Aligning Large Language Models via Mistake Analysis.
- Ruiyuan Gao, Kai Chen, Enze Xie, Lanqing Hong, Zhenguo Li, Dit-Yan Yeung, Qiang Xu:
MagicDrive: Street View Generation with Diverse 3D Geometry Control.
- Siyuan Li, Zedong Wang, Zicheng Liu, Cheng Tan, Haitao Lin, Di Wu, Zhiyuan Chen, Jiangbin Zheng, Stan Z. Li:
MogaNet: Multi-order Gated Aggregation Network.

-     Kai Chen, Enze Xie, Zhe Chen, Yibo Wang, Lanqing Hong, Zhenguo Li, Dit-Yan Yeung:
GeoDiffusion: Text-Prompted Geometric Control for Object Detection Data Generation.
-     Devavrat Tomar, Guillaume Vray, Jean-Philippe Thiran, Behzad Bozorgtabar:
Un-Mixing Test-Time Normalization Statistics: Combatting Label Temporal Correlation.
-     Riccardo Massidda, Francesco Landolfi, Martina Cinquini, Davide Bacciu:
Constraint-Free Structure Learning with Smooth Acyclic Orientations.
-     Zhipeng Zhou, Liu Liu, Peilin Zhao, Wei Gong:
Pareto Deep Long-Tailed Recognition: A Conflict-Averse Solution.
-     Yuxue Yang, Lue Fan, Zhaoxiang Zhang:
MixSup: Mixed-grained Supervision for Label-efficient LiDAR-based 3D Object Detection.
-     Zhentao Tan, Xiaodan Li, Yue Wu, Qi Chu, Le Lu, Nenghai Yu, Jieping Ye:
Boosting Vanilla Lightweight Vision Transformers via Re-parameterization.
-     Yixuan He, Gesine Reinert, David Wipf, Mihai Cucuringu:
Robust Angular Synchronization via Directed Graph Neural Networks.
-     Haotian Yan, Ming Wu, Chuang Zhang:
Multi-Scale Representations by Varying Window Attention for Semantic Segmentation.
-     Kai Yi, Nidham Gazagnadou, Peter Richtárik, Lingjuan Lyu:
FedP3: Federated Personalized and Privacy-friendly Network Pruning under Model Heterogeneity.
-     Jang-Hyun Kim, Junyoung Yeom, Sangdoo Yun, Hyun Oh Song:
Compressed Context Memory for Online Language Model Interaction.
-     An-Chieh Cheng, Xuetong Li, Sifei Liu, Xiaolong Wang:
TUVF: Learning Generalizable Texture UV Radiance Fields.
-     Adriano Cardace, Pierluigi Zama Ramirez, Francesco Ballerini, Allan Zhou, Samuele Salti, Luigi Di Stefano:
Neural Processing of Tri-Plane Hybrid Neural Fields.
-     Ziang Cao, Fangzhou Hong, Tong Wu, Liang Pan, Ziwei Liu:
Large-Vocabulary 3D Diffusion Model with Transformer.
-     Yusuke Sekikawa, Shingo Yashima:
SAS: Structured Activation Sparsification.
-     Zecheng Hao, Xinyu Shi, Zihan Huang, Tong Bu, Zhao Fei Yu, Tiejun Huang:
A Progressive Training Framework for Spiking Neural Networks with Learnable Multi-hierarchical Model.
-     Ziyue Jiang, Jinglin Liu, Yi Ren, Jinzheng He, Zhenhui Ye, Shengpeng Ji, Qian Yang, Chen Zhang, Pengfei Wei, Chunfeng Wang, Xiang Yin, Zejun Ma, Zhou Zhao:
Mega-TTS 2: Boosting Prompting Mechanisms for Zero-Shot Speech Synthesis.
-     Olivier Laurent, Emanuel Aldea, Gianni Franchi:
A Symmetry-Aware Exploration of Bayesian Neural Network Posteriors.
-     Zecheng Hao, Tong Bu, Xinyu Shi, Zihan Huang, Zhao Fei Yu, Tiejun Huang:
Threaten Spiking Neural Networks through Combining Rate and Temporal Information.
-     Jing Liu, Ruihao Gong, Xiuying Wei, Zhiwei Dong, Jianfei Cai, Bohan Zhuang:
QLLM: Accurate and Efficient Low-Bitwidth Quantization for Large Language Models.
-     Chen Zhao, Tong Zhang, Mathieu Salzmann:
3D-Aware Hypothesis & Verification for Generalizable Relative Object Pose Estimation.
-     Jing-Cheng Pang, Pengyuan Wang, Kaiyuan Li, Xiong-Hui Chen, Jiacheng Xu, Zongzhang Zhang, Yang Yu:
Language Model Self-improvement by Reinforcement Learning Contemplation.
-     Grzegorz Rypesc, Sebastian Cygert, Valeriya Khan, Tomasz Trzcinski, Bartosz Zielinski, Bartłomiej Twardowski:
Divide and not forget: Ensemble of selectively trained experts in Continual Learning.
-     Yuheng Jing, Kai Li, Bingyun Liu, Yifan Zang, Haobo Fu, Qiang Fu, Junliang Xing, Jian Cheng:
Towards Offline Opponent Modeling with In-context Learning.
-     Suqin Yuan, Lei Feng, Tongliang Liu:
Early Stopping Against Label Noise Without Validation Data.
-     Zheng Chen, Yulun Zhang, Jinjin Gu, Linghe Kong, Xiaokang Yang:
Recursive Generalization Transformer for Image Super-Resolution.
-     Huanran Chen, Yichi Zhang, Yinpeng Dong, Xiao Yang, Hang Su, Jun Zhu:
Rethinking Model Ensemble in Transfer-based Adversarial Attacks.
-     Lu Yu, Avetik G. Karagulyan, Arnak S. Dalalyan:
Langevin Monte Carlo for strongly log-concave distributions: Randomized midpoint revisited.
-     Xurui Li, Ziming Huang, Feng Xue, Yu Zhou:
MuSc: Zero-Shot Industrial Anomaly Classification and Segmentation with Mutual Scoring of the Unlabeled Images.
-     Manley Roberts, Himanshu Thakur, Christine Herlihy, Colin White, Samuel Dooley:
To the Cutoff... and Beyond? A Longitudinal Perspective on LLM Data Contamination.



Shiqian Li, Kewen Wu, Chi Zhang, Yixin Zhu:
I-PHYRE: Interactive Physical Reasoning.



Mingzhen Huang, Shan Jia, Zhou Zhou, Yan Ju, Jialing Cai, Siwei Lyu:
Exposing Text-Image Inconsistency Using Diffusion Models.

VALID HTML last updated on 2024-12-11 20:32 CET by the dblp team

all metadata released as open data under CC0 1.0 license

see also: Terms of Use | Privacy Policy | Imprint

dblp was originally created in 1993 at:



since 2018, dblp has been operated and maintained by:



the dblp computer science bibliography is funded and supported by:



SAARLAND

