



We've just launched a new service: our brand new **dblp SPARQL query service**. Read more about it in our latest blog post or try out some of the SPARQL queries linked on the dblp web pages below.

X




## [+] 38th ICML 2021: Virtual Event

> Home > Conferences and Workshops > ICML

Dagstuhl

- Marina Meila, Tong Zhang:  
**Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event.** Proceedings of Machine Learning Research 139, PMLR 2021
- Majid Abdolshah, Hung Le, Thommen George Karimpanal, Sunil Gupta, Santu Rana, Svetha Venkatesh:  
**A New Representation of Successor Features for Transfer across Dissimilar Environments.** 1-9
- Kuruge Darshana Abeyrathna, Bimal Bhattacharai, Morten Goodwin, Saeed Rahimi Gorji, Ole-Christoffer Granmo, Lei Jiao, Rupsa Saha, Rohan Kumar Yadav:  
**Massively Parallel and Asynchronous Tsetlin Machine Architecture Supporting Almost Constant-Time Scaling.** 10-20
- Durmus Alp Emre Acar, Yue Zhao, Ruizhao Zhu, Ramon Matas Navarro, Matthew Mattina, Paul N. Whatmough, Venkatesh Saligrama:  
**Debiasing Model Updates for Improving Personalized Federated Training.** 21-31
- Durmus Alp Emre Acar, Ruizhao Zhu, Venkatesh Saligrama:  
**Memory Efficient Online Meta Learning.** 32-42
- Jayadev Acharya, Ziteng Sun, Huanyu Zhang:  
**Robust Testing and Estimation under Manipulation Attacks.** 43-53
- Idan Achituve, Aviv Navon, Yochai Yemini, Gal Chechik, Ethan Fetaya:  
**GP-Tree: A Gaussian Process Classifier for Few-Shot Incremental Learning.** 54-65
- David Acuna, Guojun Zhang, Marc T. Law, Sanja Fidler:  
**f-Domain Adversarial Learning: Theory and Algorithms.** 66-75
- Darius Afchar, Vincent Guigue, Romain Hennequin:  
**Towards Rigorous Interpretations: a Formalisation of Feature Attribution.** 76-86
- Naman Agarwal, Surbhi Goel, Cyril Zhang:  
**Acceleration via Fractal Learning Rate Schedules.** 87-99
- Naman Agarwal, Elad Hazan, Anirudha Majumdar, Karan Singh:  
**A Regret Minimization Approach to Iterative Learning Control.** 100-109
- Sushant Agarwal, Shahin Jabbari, Chirag Agarwal, Sohini Upadhyay, Steven Wu, Himabindu Lakkaraju:  
**Towards the Unification and Robustness of Perturbation and Gradient Based Explanations.** 110-119
- Abhinav Aggarwal, Shiva Prasad Kasiviswanathan, Zekun Xu, Oluwaseyi Feyisetan, Nathanael Teissier:  
**Label Inference Attacks from Log-loss Scores.** 120-129
- Laurence Aitchison, Adam X. Yang, Sebastian W. Ober:  
**Deep Kernel Processes.** 130-140
- Ali Akbari, Muhammad Awais, Manijeh Bashar, Josef Kittler:  
**How Does Loss Function Affect Generalization Performance of Deep Learning? Application to Human Age Estimation.** 141-151
- Shunta Akiyama, Taiji Suzuki:  
**On Learnability via Gradient Method for Two-Layer ReLU Neural Networks in Teacher-Student Setting.** 152-162
- Maxwell Mbilla Aladago, Lorenzo Torresani:  
**Slot Machines: Discovering Winning Combinations of Random Weights in Neural Networks.** 163-174
- Ferran Alet, Javier Lopez-Contreras, James Koppel, Maxwell I. Nye, Armando Solar-Lezama , Tomás Lozano-Pérez, Leslie Pack Kaelbling, Joshua B. Tenenbaum:  
**A large-scale benchmark for few-shot program induction and synthesis.** 175-186

[+] **SPARQL queries**

[–] **Refine list**

showing all 1,184 records

refine by search term

refine by author

Masashi Sugiyama (14)  
Sergey Levine (13)  
Gang Niu 0001 (11)  
Zhaoran Wang 0001 (11)  
Zhuoran Yang (11)  
Quanquan Gu (8)  
Percy Liang (7)  
Chelsea Finn (7)  
Andreas Krause 0001 (7)  
Csaba Szepesvári (7)  
3,881 more options

refine by access

open (1,184)



- Ayya Alieva, Ashok Cutkosky , Abhimanyu Das:  
**Robust Pure Exploration in Linear Bandits with Limited Budget.** 187-195
- Foivos Alimisis, Peter Davies, Dan Alistarh:  
**Communication-Efficient Distributed Optimization with Quantized Preconditioners.** 196-206
- Pierre Alquier:  
**Non-Exponentially Weighted Aggregation: Regret Bounds for Unbounded Loss Functions.** 207-218
- David Alvarez-Melis, Nicolò Fusi:  
**Dataset Dynamics via Gradient Flows in Probability Space.** 219-230
- Georgios Amanatidis, Federico Fusco, Philip Lazos, Stefano Leonardi, Alberto Marchetti-Spaccamela, Rebecca Reiffenhäuser :  
**Submodular Maximization subject to a Knapsack Constraint: Combinatorial Algorithms with Near-optimal Adaptive Complexity.** 231-242
- Sanae Amani, Christos Thrampoulidis, Lin Yang :  
**Safe Reinforcement Learning with Linear Function Approximation.** 243-253
- Luca Ambrogioni, Gianluigi Silvestri, Marcel van Gerven:  
**Automatic variational inference with cascading flows.** 254-263
- Sebastian E. Ament, Carla P. Gomes:  
**Sparse Bayesian Learning via Stepwise Regression.** 264-274
- Susan Amin, Maziar Gomrokchi, Hossein Aboutalebi, Harsh Satija, Doina Precup:  
**Locally Persistent Exploration in Continuous Control Tasks with Sparse Rewards.** 275-285
- Nishanth V. Anand, Doina Precup:  
**Preferential Temporal Difference Learning.** 286-296
- Fidel Ernesto Diaz Andino, Maria Kokkou, Mateus de Oliveira Oliveira, Farhad Vadiee:  
**Unitary Branching Programs: Learnability and Lower Bounds.** 297-306
- Brandon Araki, Xiao Li, Kiran Vodrahalli, Jonathan A. DeCastro, Micah J. Fry, Daniela Rus:  
**The Logical Options Framework.** 307-317
- Michael Arbel, Alexander G. de G. Matthews, Arnaud Doucet:  
**Annealed Flow Transport Monte Carlo.** 318-330
- David Arbour, Drew Dimmery, Arjun Sondhi:  
**Permutation Weighting.** 331-341
- Ludovic Arnould, Claire Boyer, Erwan Scornet:  
**Analyzing the tree-layer structure of Deep Forests.** 342-350
- Raman Arora, Peter L. Bartlett, Poorya Mianjy, Nathan Srebro:  
**Dropout: Explicit Forms and Capacity Control.** 351-361
- Artem Artemev, David R. Burt, Mark van der Wilk:  
**Tighter Bounds on the Log Marginal Likelihood of Gaussian Process Regression Using Conjugate Gradients.** 362-372
- Dilip Arumugam, Benjamin Van Roy:  
**Deciding What to Learn: A Rate-Distortion Approach.** 373-382
- Hilal Asi, John C. Duchi, Alireza Fallah, Omid Javidbakht, Kunal Talwar:  
**Private Adaptive Gradient Methods for Convex Optimization.** 383-392
- Hilal Asi, Vitaly Feldman, Tomer Koren, Kunal Talwar:  
**Private Stochastic Convex Optimization: Optimal Rates in L1 Geometry.** 393-403
- Alexia Atsidakou, Orestis Papadigenopoulos, Soumya Basu, Constantine Caramanis, Sanjay Shakkottai:  
**Combinatorial Blocking Bandits with Stochastic Delays.** 404-413
- Julien Audiffren:  
**Dichotomous Optimistic Search to Quantify Human Perception.** 414-424
- Dmitrii Avdiukhin, Shiva Prasad Kasiviswanathan:  
**Federated Learning under Arbitrary Communication Patterns.** 425-435
- Rotem Zamir Aviv, Ido Hakimi, Assaf Schuster, Kfir Yehuda Levy:  
**Asynchronous Distributed Learning : Adapting to Gradient Delays without Prior Knowledge.** 436-445
- Kyriakos Axiotis, Adam Karczmarz , Anish Mukherjee, Piotr Sankowski, Adrian Vladu:  
**Decomposable Submodular Function Minimization via Maximum Flow.** 446-456
- Sergül Aydöre, William Brown, Michael Kearns, Krishnaram Kenthapadi, Luca Melis, Aaron Roth, Amaresh Ankit Siva:  
**Differentially Private Query Release Through Adaptive Projection.** 457-467
- Shahar Azulay, Edward Moroshko, Mor Shpigel Nacson, Blake E. Woodworth, Nathan Srebro, Amir Globerson, Daniel Soudry:  
**On the Implicit Bias of Initialization Shape: Beyond Infinitesimal Mirror Descent.** 468-477
- Zahra Babaiee, Ramin M. Hasani, Mathias Lechner, Daniela Rus, Radu Grosu:  
**On-Off Center-Surround Receptive Fields for Accurate and Robust Image Classification.** 478-489



- Gregor Bachmann, Seyed-Mohsen Moosavi-Dezfooli, Thomas Hofmann:  
**Uniform Convergence, Adversarial Spheres and a Simple Remedy.** 490-499
- Arturs Backurs, Piotr Indyk, Cameron Musco, Tal Wagner:  
**Faster Kernel Matrix Algebra via Density Estimation.** 500-510
- Kishan Panaganti Badrinath, Dileep Kalathil:  
**Robust Reinforcement Learning using Least Squares Policy Iteration with Provable Performance Guarantees.** 511-520
- Akhil Bagaria, Jason K. Senthil, George Konidaris:  
**Skill Discovery for Exploration and Planning using Deep Skill Graphs.** 521-531
- Dara Bahri, Heinrich Jiang:  
**Locally Adaptive Label Smoothing Improves Predictive Churn.** 532-542
- Yu Bai, Minshuo Chen, Pan Zhou, Tuo Zhao, Jason D. Lee, Sham M. Kakade, Huan Wang, Caiming Xiong:  
**How Important is the Train-Validation Split in Meta-Learning?** 543-553
- Shaojie Bai, Vladlen Koltun, J. Zico Kolter:  
**Stabilizing Equilibrium Models by Jacobian Regularization.** 554-565
- Yu Bai, Song Mei, Huan Wang, Caiming Xiong:  
**Don't Just Blame Over-parametrization for Over-confidence: Theoretical Analysis of Calibration in Binary Classification.** 566-576
- Chenjia Bai, Lingxiao Wang, Lei Han, Jianye Hao, Animesh Garg, Peng Liu, Zhaoran Wang:  
**Principled Exploration via Optimistic Bootstrapping and Backward Induction.** 577-587
- Yunsheng Bai, Derek Xu, Yizhou Sun, Wei Wang:  
**GLSearch: Maximum Common Subgraph Detection via Learning to Search.** 588-598
- Muhammet Balcīlar, Pierre Héroux, Benoit Gaüzère, Pascal Vasseur, Sébastien Adam, Paul Honeine:  
**Breaking the Limits of Message Passing Graph Neural Networks.** 599-608
- Eric Balkanski, Sharon Qian, Yaron Singer:  
**Instance Specific Approximations for Submodular Maximization.** 609-618
- Philip J. Ball, Cong Lu, Jack Parker-Holder, Stephen J. Roberts:  
**Augmented World Models Facilitate Zero-Shot Dynamics Generalization From a Single Offline Environment.** 619-629
- Santiago R. Balseiro, Haihao Lu, Vahab S. Mirrokni:  
**Regularized Online Allocation Problems: Fairness and Beyond.** 630-639
- Yujia Bao, Shiyu Chang, Regina Barzilay:  
**Predict then Interpolate: A Simple Algorithm to Learn Stable Classifiers.** 640-650
- Fan Bao, Kun Xu, Chongxuan Li, Lanqing Hong, Jun Zhu, Bo Zhang:  
**Variational (Gradient) Estimate of the Score Function in Energy-based Latent Variable Models.** 651-661
- Amir Bar, Roei Herzig, Xiaolong Wang, Anna Rohrbach, Gal Chechik, Trevor Darrell, Amir Globerson:  
**Compositional Video Synthesis with Action Graphs.** 662-673
- Nadav Barak, Sivan Sabato:  
**Approximating a Distribution Using Weight Queries.** 674-683
- Aseem Baranwal, Kimon Fountoulakis, Aukosh Jagannath:  
**Graph Convolution for Semi-Supervised Classification: Improved Linear Separability and Out-of-Distribution Generalization.** 684-693
- Burak Bartan, Mert Pilancı:  
**Training Quantized Neural Networks to Global Optimality via Semidefinite Programming.** 694-704
- Soumya Basu, Karthik Abinav Sankararaman, Abishek Sankararaman:  
**Beyond  $\log^2(T)$  regret for decentralized bandits in matching markets.** 705-715
- Dorian Baudry, Romain Gautron, Emilie Kaufmann, Odalric Maillard:  
**Optimal Thompson Sampling strategies for support-aware CVaR bandits.** 716-726
- Dorian Baudry, Yoan Russac, Olivier Cappé:  
**On Limited-Memory Subsampling Strategies for Bandits.** 727-737
- Matthias Bauer, Andriy Mnih:  
**Generalized Doubly Reparameterized Gradient Estimators.** 738-747
- Dominique Beaini, Saro Passaro, Vincent Létourneau, William L. Hamilton, Gabriele Corso, Pietro Lió:  
**Directional Graph Networks.** 748-758
- Alexis Bellot, Mihaela van der Schaar:  
**Policy Analysis using Synthetic Controls in Continuous-Time.** 759-768
- Gregory W. Benton, Wesley J. Maddox, Sanae Lotfi, Andrew Gordon Wilson:  
**Loss Surface Simplexes for Mode Connecting Volumes and Fast Ensembling.** 769-779
- Berkay Berabi, Jingxuan He, Veselin Raychev, Martin T. Vechev:  
**TFix: Learning to Fix Coding Errors with a Text-to-Text Transformer.** 780-791



- Jeroen Berrevoets, Ahmed M. Alaa, Zhaozhi Qian, James Jordon, Alexander E. S. Gimson, Mihaela van der Schaar:  
**Learning Queueing Policies for Organ Transplantation Allocation using Interpretable Counterfactual Survival Analysis.** 792-802
- Patrice Bertail, Stéphan Cléménçon, Yannick Guyonvarch, Nathan Noiry:  
**Learning from Biased Data: A Semi-Parametric Approach.** 803-812
- Gedas Bertasius, Heng Wang, Lorenzo Torresani:  
**Is Space-Time Attention All You Need for Video Understanding?** 813-824
- Antonin Berthon, Bo Han, Gang Niu, Tongliang Liu, Masashi Sugiyama:  
**Confidence Scores Make Instance-dependent Label-noise Learning Possible.** 825-836
- Beatrice Bevilacqua, Yangze Zhou, Bruno Ribeiro:  
**Size-Invariant Graph Representations for Graph Classification Extrapolations.** 837-851
- Sourabh Bhadane, Aaron B. Wagner, Jayadev Acharya:  
**Principal Bit Analysis: Autoencoding with Schur-Concave Loss.** 852-862
- Arjun Nitin Bhagoji, Daniel Cullina, Vikash Sehwag, Prateek Mittal:  
**Lower Bounds on Cross-Entropy Loss in the Presence of Test-time Adversaries.** 863-873
- Aditya Bhaskara, Aravinda Kanchana Ruwanpathirana, Maheshakya Wijewardena:  
**Additive Error Guarantees for Weighted Low Rank Approximation.** 874-883
- Robi Bhattacharjee, Somesh Jha, Kamalika Chaudhuri:  
**Sample Complexity of Robust Linear Classification on Separated Data.** 884-893
- Chiranjib Bhattacharyya, Ravindran Kannan, Amit Kumar:  
**Finding k in Latent k- polytope.** 894-903
- Hangrui Bi, Hengyi Wang, Chence Shi, Connor W. Coley, Jian Tang, Hongyu Guo:  
**Non-Autoregressive Electron Redistribution Modeling for Reaction Prediction.** 904-913
- André Biedenkapp, Raghu Rajan, Frank Hutter, Marius Lindauer:  
**TempoRL: Learning When to Act.** 914-924
- Jakub Bielawski, Thiparat Chotibut, Fryderyk Falniowski, Grzegorz Kosiorowski, Michal Misiurewicz, Georgios Piliouras:  
**Follow-the-Regularized-Leader Routes to Chaos in Routing Games.** 925-935
- Luca Biggio, Tommaso Bendinelli, Alexander Neitz, Aurélien Lucchi, Giambattista Parascandolo:  
**Neural Symbolic Regression that scales.** 936-945
- Max Biggs, Wei Sun, Markus Ettl:  
**Model Distillation for Revenue Optimization: Interpretable Personalized Pricing.** 946-956
- Marin Bilos, Stephan Günnemann:  
**Scalable Normalizing Flows for Permutation Invariant Densities.** 957-967
- Ilai Bistritz, Nicholas Bambos:  
**Online Learning for Load Balancing of Unknown Monotone Resource Allocation Games.** 968-979
- Johan Björck, Xiangyu Chen, Christopher De Sa, Carla P. Gomes, Kilian Q. Weinberger:  
**Low-Precision Reinforcement Learning: Running Soft Actor-Critic in Half Precision.** 980-991
- Davis W. Blalock, John V. Guttag:  
**Multiplying Matrices Without Multiplying.** 992-1004
- Avrim Blum, Nika Haghtalab, Richard Lanas Phillips, Han Shao:  
**One for One, or All for All: Equilibria and Optimality of Collaboration in Federated Learning.** 1005-1014
- Erik Bodin, Zhenwen Dai, Neill W. Campbell, Carl Henrik Ek:  
**Black-box density function estimation using recursive partitioning.** 1015-1025
- Cristian Bodnar, Fabrizio Frasca, Yuguang Wang, Nina Otter, Guido F. Montúfar, Pietro Lió, Michael M. Bronstein:  
**Weisfeiler and Lehman Go Topological: Message Passing Simplicial Networks.** 1026-1037
- Roberto Bondesan, Max Welling:  
**The Hintons in your Neural Network: a Quantum Field Theory View of Deep Learning.** 1038-1048
- David Brandfonbrener, William F. Whitney, Rajesh Ranganath, Joan Bruna:  
**Offline Contextual Bandits with Overparameterized Models.** 1049-1058
- Andy Brock, Soham De, Samuel L. Smith, Karen Simonyan:  
**High-Performance Large-Scale Image Recognition Without Normalization.** 1059-1071
- James A. Brofos, Roy R. Lederman:  
**Evaluating the Implicit Midpoint Integrator for Riemannian Hamiltonian Monte Carlo.** 1072-1081
- Ethan A. Brooks, Janarthanan Rajendran, Richard L. Lewis, Satinder Singh:  
**Reinforcement Learning of Implicit and Explicit Control Flow Instructions.** 1082-1091
- Jonathan Brophy, Daniel Lowd:  
**Machine Unlearning for Random Forests.** 1092-1104
- Daniel S. Brown, Jordan Schneider , Anca D. Dragan, Scott Niekum:  
**Value Alignment Verification.** 1105-1115



- David Bruns-Smith:  
**Model-Free and Model-Based Policy Evaluation when Causality is Uncertain.** 1116-1126
- Francois Buet-Golfouse:  
**Narrow Margins: Classification, Margins and Fat Tails.** 1127-1135
- Mark Bun, Marek Eliás, Janardhan Kulkarni:  
**Differentially Private Correlation Clustering.** 1136-1146
- Vivien A. Cabannes, Francis R. Bach, Alessandro Rudi:  
**Disambiguation of Weak Supervision leading to Exponential Convergence rates.** 1147-1157
- Diana Cai, Trevor Campbell, Tamara Broderick:  
**Finite mixture models do not reliably learn the number of components.** 1158-1169
- Tianle Cai, Ruiqi Gao, Jason D. Lee, Qi Lei:  
**A Theory of Label Propagation for Subpopulation Shift.** 1170-1182
- Xu Cai, Selwyn Gomes, Jonathan Scarlett:  
**Lenient Regret and Good-Action Identification in Gaussian Process Bandits.** 1183-1192
- HanQin Cai, Yuchen Lou, Daniel McKenzie, Wotao Yin:  
**A Zeroth-Order Block Coordinate Descent Algorithm for Huge-Scale Black-Box Optimization.** 1193-1203
- Tianle Cai, Shengjie Luo, Keyulu Xu, Di He, Tie-Yan Liu, Liwei Wang:  
**GraphNorm: A Principled Approach to Accelerating Graph Neural Network Training.** 1204-1215
- Xu Cai, Jonathan Scarlett:  
**On Lower Bounds for Standard and Robust Gaussian Process Bandit Optimization.** 1216-1226
- Romain Camilleri, Kevin Jamieson, Julian Katz-Samuels:  
**High-dimensional Experimental Design and Kernel Bandits.** 1227-1237
- Andrew Campbell, Wenlong Chen, Vincent Stimper, José Miguel Hernández-Lobato, Yichuan Zhang:  
**A Gradient Based Strategy for Hamiltonian Monte Carlo Hyperparameter Optimization.** 1238-1248
- Alexander Camuto, Xiaoyu Wang, Lingjiong Zhu, Chris C. Holmes, Mert Gürbüzbalaban, Umut Simsekli:  
**Asymmetric Heavy Tails and Implicit Bias in Gaussian Noise Injections.** 1249-1260
- Yue Cao, Payel Das, Vijil Chenthamarakshan, Pin-Yu Chen, Igor Melnyk, Yang Shen:  
**Fold2Seq: A Joint Sequence(1D)-Fold(3D) Embedding-based Generative Model for Protein Design.** 1261-1271
- Yuzhou Cao, Lei Feng, Yitian Xu, Bo An, Gang Niu, Masashi Sugiyama:  
**Learning from Similarity-Confidence Data.** 1272-1282
- Alejandro Carderera, Jelena Diakonikolas, Cheuk Yin Lin, Sebastian Pokutta:  
**Parameter-free Locally Accelerated Conditional Gradients.** 1283-1293
- Mathieu Carrière, Frédéric Chazal, Marc Glisse, Yuichi Ike, Hariprasad Kannan, Yuhei Umeda:  
**Optimizing persistent homology based functions.** 1294-1303
- Asaf B. Cassel, Tomer Koren:  
**Online Policy Gradient for Model Free Learning of Linear Quadratic Regulators with  $\sqrt{T}$  Regret.** 1304-1313
- Matteo Castiglioni, Alberto Marchesi, Andrea Celli, Nicola Gatti:  
**Multi-Receiver Online Bayesian Persuasion.** 1314-1323
- Amnon Catav, Boyang Fu, Yazeed Zoabi, Ahuva Weiss-Meilik, Noam Shomron, Jason Ernst, Sriram Sankararaman, Ran Gilad-Bachrach:  
**Marginal Contribution Feature Importance - an Axiomatic Approach for Explaining Data.** 1324-1335
- Charlotte Caucheteux, Alexandre Gramfort, Jean-Remi King:  
**Disentangling syntax and semantics in the brain with deep networks.** 1336-1348
- L. Elisa Celis, Lingxiao Huang, Vijay Keswani, Nisheeth K. Vishnoi:  
**Fair Classification with Noisy Protected Attributes: A Framework with Provable Guarantees.** 1349-1361
- Leonardo Cellia, Massimiliano Pontil, Claudio Gentile:  
**Best Model Identification: A Rested Bandit Formulation.** 1362-1372
- Johan Samir Obando-Ceron, Pablo Samuel Castro:  
**Revisiting Rainbow: Promoting more insightful and inclusive deep reinforcement learning research.** 1373-1383
- Edoardo Cetin, Oya Çeliktutan:  
**Learning Routines for Effective Off-Policy Reinforcement Learning.** 1384-1394
- Ciwan Ceylan , Salla Franzén, Florian T. Pokorný:  
**Learning Node Representations Using Stationary Flow Prediction on Large Payment and Cash Transaction Networks.** 1395-1406
- Ben Chamberlain, James Rowbottom, Maria I. Gorinova, Michael M. Bronstein, Stefan Webb, Emanuele Rossi:  
**GRAND: Graph Neural Diffusion.** 1407-1418
- Ines Chami, Albert Gu, Dat Nguyen, Christopher Ré:  
**HoroPCA: Hyperbolic Dimensionality Reduction via Horospherical Projections.** 1419-1429
- Elliot Chané-Sane, Cordelia Schmid, Ivan Laptev:  
**Goal-Conditioned Reinforcement Learning with Imagined Subgoals.** 1430-1440



- Alisa Chang, Badih Ghazi, Ravi Kumar, Pasin Manurangsi:  
**Locally Private k-Means in One Round.** 1441-1451
- Michael Chang, Sidhant Kaushik, Sergey Levine, Tom Griffiths:  
**Modularity in Reinforcement Learning via Algorithmic Independence in Credit Assignment.** 1452-1462
- Nadine Chang, Zhiding Yu, Yu-Xiong Wang, Animashree Anandkumar, Sanja Fidler, José M. Álvarez:  
**Image-Level or Object-Level? A Tale of Two Resampling Strategies for Long-Tailed Detection.** 1463-1472
- Sudhanshu Chanpuriya, Cameron Musco, Konstantinos Sotiropoulos, Charalampos E. Tsourakakis:  
**DeepWalking Backwards: From Embeddings Back to Graphs.** 1473-1483
- Devendra Singh Chaplot, Deepak Pathak, Jitendra Malik:  
**Differentiable Spatial Planning using Transformers.** 1484-1495
- Henry Charlesworth, Giovanni Montana:  
**Solving Challenging Dexterous Manipulation Tasks With Trajectory Optimisation and Reinforcement Learning.** 1496-1506
- Nontawat Charoenphakdee, Zhenghang Cui, Yivan Zhang, Masashi Sugiyama:  
**Classification with Rejection Based on Cost-sensitive Classification.** 1507-1517
- Yevgen Chebotar, Karol Hausman, Yao Lu, Ted Xiao, Dmitry Kalashnikov, Jacob Varley, Alex Irpan, Benjamin Eysenbach, Ryan Julian, Chelsea Finn, Sergey Levine:  
**Actionable Models: Unsupervised Offline Reinforcement Learning of Robotic Skills.** 1518-1528
- Xu Chen:  
**Unified Robust Semi-Supervised Variational Autoencoder.** 1529-1538
- Boyuan Chen, Pieter Abbeel, Deepak Pathak:  
**Unsupervised Learning of Visual 3D Keypoints for Control.** 1539-1549
- Rui Chen, Sanjeeb Dash, Tian Gao:  
**Integer Programming for Causal Structure Learning in the Presence of Latent Variables.** 1550-1560
- Yifang Chen, Simon S. Du, Kevin Jamieson:  
**Improved Corruption Robust Algorithms for Episodic Reinforcement Learning.** 1561-1570
- Yongxin Chen, Jiaoqiao Fan, Amirhossein Taghvaei:  
**Scalable Computations of Wasserstein Barycenter via Input Convex Neural Networks.** 1571-1581
- Yunlu Chen, Basura Fernando, Hakan Bilen , Thomas Mensink, Efstratios Gavves:  
**Neural Feature Matching in Implicit 3D Representations.** 1582-1593
- Shixiang Chen, Alfredo García, Mingyi Hong, Shahin Shahrampour:  
**Decentralized Riemannian Gradient Descent on the Stiefel Manifold.** 1594-1605
- Chao Chen, Haoyu Geng, Nianzu Yang, Junchi Yan, Daiyue Xue, Jianping Yu, Xiaokang Yang:  
**Learning Self-Modulating Attention in Continuous Time Space with Applications to Sequential Recommendation.** 1606-1616
- Mayee F. Chen, Karan Goel, Nimit Sharad Sohoni, Fait Poms, Kayvon Fatahalian, Christopher Ré:  
**Mandoline: Model Evaluation under Distribution Shift.** 1617-1629
- Xiaohui Chen, Xu Han, Jiajing Hu, Francisco J. R. Ruiz, Li-Ping Liu:  
**Order Matters: Probabilistic Modeling of Node Sequence for Graph Generation.** 1630-1639
- Dian Chen, Hongxin Hu, Qian Wang, Yinli Li, Cong Wang, Chao Shen, Qi Li:  
**CARTL: Cooperative Adversarially-Robust Transfer Learning.** 1640-1650
- Liyu Chen, Haipeng Luo:  
**Finding the Stochastic Shortest Path with Low Regret: the Adversarial Cost and Unknown Transition Case.** 1651-1660
- Xinyun Chen, Petros Maniatis, Rishabh Singh, Charles Sutton, Hanjun Dai, Max Lin, Denny Zhou:  
**SpreadsheetCoder: Formula Prediction from Semi-structured Context.** 1661-1672
- Shuo Chen, Gang Niu, Chen Gong, Jun Li, Jian Yang, Masashi Sugiyama:  
**Large-Margin Contrastive Learning with Distance Polarization Regularizer.** 1673-1683
- Yuzhou Chen, Ignacio Segovia-Dominguez, Yulia R. Gel:  
**Z-GCNets: Time Zigzags at Graph Convolutional Networks for Time Series Forecasting.** 1684-1694
- Tianlong Chen, Yongduo Sui, Xuxi Chen, Aston Zhang, Zhangyang Wang:  
**A Unified Lottery Ticket Hypothesis for Graph Neural Networks.** 1695-1706
- Wei Chen, Xiaoming Sun, Jialin Zhang, Zhijie Zhang:  
**Network Inference and Influence Maximization from Samples.** 1707-1716
- Renyi Chen, Molei Tao:  
**Data-driven Prediction of General Hamiltonian Dynamics via Learning Exactly-Symplectic Maps.** 1717-1727
- Tyler Chen, Thomas Trogdon, Shashanka Ubaru:  
**Analysis of stochastic Lanczos quadrature for spectrum approximation.** 1728-1739
- Tianrong Chen, Ziyi Wang, Ioannis Exarchos, Evangelos A. Theodorou:  
**Large-Scale Multi-Agent Deep FBSDEs.** 1740-1748



- Xinyang Chen, Sinan Wang, Jianmin Wang , Mingsheng Long :  
**Representation Subspace Distance for Domain Adaptation Regression.** 1749-1759
- Pei-Hung Chen, Wei Wei, Cho-Jui Hsieh, Bo Dai:  
**Overcoming Catastrophic Forgetting by Bayesian Generative Regularization.** 1760-1770
- Xiangyu Chen, Min Ye:  
**Cyclically Equivariant Neural Decoders for Cyclic Codes.** 1771-1780
- Jintai Chen, Hongyun Yu, Chengde Qian, Danny Z. Chen, Jian Wu:  
**A Receptor Skeleton for Capsule Neural Networks.** 1781-1790
- Yiming Chen, Kun Yuan, Yingya Zhang, Pan Pan, Yinghui Xu, Wotao Yin:  
**Accelerating Gossip SGD with Periodic Global Averaging.** 1791-1802
- Jianfei Chen, Lianmin Zheng, Zhewei Yao, Dequan Wang, Ion Stoica, Michael W. Mahoney, Joseph Gonzalez:  
**ActNN: Reducing Training Memory Footprint via 2-Bit Activation Compressed Training.** 1803-1813
- Wuxinlin Cheng, Chenhui Deng, Zhiqiang Zhao, Yaohui Cai, Zhiru Zhang, Zhuo Feng:  
**SPADE: A Spectral Method for Black-Box Adversarial Robustness Evaluation.** 1814-1824
- Yong Cheng, Wei Wang, Lu Jiang, Wolfgang Macherey:  
**Self-supervised and Supervised Joint Training for Resource-rich Machine Translation.** 1825-1835
- Giovanni Cherubin, Konstantinos Chatzikokolakis, Martin Jaggi:  
**Exact Optimization of Conformal Predictors via Incremental and Decremental Learning.** 1836-1845
- James Cheshire, Pierre Ménard, Alexandra Carpentier:  
**Problem Dependent View on Structured Thresholding Bandit Problems.** 1846-1854
- Yun Kuen Cheung, Georgios Piliouras:  
**Online Optimization in Games via Control Theory: Connecting Regret, Passivity and Poincaré Recurrence.** 1855-1865
- Jianfeng Chi, Yuan Tian, Geoffrey J. Gordon, Han Zhao:  
**Understanding and Mitigating Accuracy Disparity in Regression.** 1866-1876
- Steve Chien, Prateek Jain, Walid Krichene, Steffen Rendle, Shuang Song, Abhradeep Thakurta, Li Zhang:  
**Private Alternating Least Squares: Practical Private Matrix Completion with Tighter Rates.** 1877-1887
- Flavio Chierichetti, Ravi Kumar, Andrew Tomkins:  
**Light RUMs.** 1888-1897
- Narsimha Reddy Chilkuri, Chris Eliasmith:  
**Parallelizing Legendre Memory Unit Training.** 1898-1907
- Uthsav Chitra, Kimberly Ding, Jasper C. H. Lee, Benjamin J. Raphael:  
**Quantifying and Reducing Bias in Maximum Likelihood Estimation of Structured Anomalies.** 1908-1919
- Jakub Chłedowski, Adam Polak, Bartosz Szabucki, Konrad Tomasz Zolna:  
**Robust Learning-Augmented Caching: An Experimental Study.** 1920-1930
- Jaemin Cho, Jie Lei, Hao Tan, Mohit Bansal:  
**Unifying Vision-and-Language Tasks via Text Generation.** 1931-1942
- Youngwon Choi, Sungdong Lee, Joong-Ho Won:  
**Learning from Nested Data with Ornstein Auto-Encoders.** 1943-1952
- Jongwook Choi, Archit Sharma, Honglak Lee, Sergey Levine, Shixiang Shane Gu:  
**Variational Empowerment as Representation Learning for Goal-Conditioned Reinforcement Learning.** 1953-1963
- Christopher A. Choquette-Choo, Florian Tramèr , Nicholas Carlini, Nicolas Papernot:  
**Label-Only Membership Inference Attacks.** 1964-1974
- Jishnu Ray Chowdhury, Cornelia Caragea:  
**Modeling Hierarchical Structures with Continuous Recursive Neural Networks.** 1975-1988
- Filippos Christianos, Georgios Papoudakis, Arrasy Rahman, Stefano V. Albrecht:  
**Scaling Multi-Agent Reinforcement Learning with Selective Parameter Sharing.** 1989-1998
- Wesley Chung, Valentin Thomas, Marlos C. Machado , Nicolas Le Roux:  
**Beyond Variance Reduction: Understanding the True Impact of Baselines on Policy Optimization.** 1999-2009
- Julien Grand-Clément, Christian Kroer:  
**First-Order Methods for Wasserstein Distributionally Robust MDP.** 2010-2019
- Karl Cobbe, Jacob Hilton, Oleg Klimov, John Schulman:  
**Phasic Policy Gradient.** 2020-2027
- Samuel Cohen, Brandon Amos, Yaron Lipman:  
**Riemannian Convex Potential Maps.** 2028-2038
- Alain-Sam Cohen, Rama Cont, Alain Rossier, Renyuan Xu:  
**Scaling Properties of Deep Residual Networks.** 2039-2048
- Edith Cohen, Haim Kaplan, Yishay Mansour, Uri Stemmer, Eliad Tsfadia:  
**Differentially-Private Clustering of Easy Instances.** 2049-2059



- Vincent Cohen-Addad, Rémi de Joannis de Verclos, Guillaume Lagarde:  
**Improving Ultrametrics Embeddings Through Coresets.** 2060-2068
- Vincent Cohen-Addad, Silvio Lattanzi, Slobodan Mitrovic, Ashkan Norouzi-Fard, Nikos Parotsidis, Jakub Tarnawski:  
**Correlation Clustering in Constant Many Parallel Rounds.** 2069-2078
- Fabien Collas, Ekhine Irurozki:  
**Concentric mixtures of Mallows models for top-k rankings: sampling and identifiability.** 2079-2088
- Liam Collins, Hamed Hassani, Aryan Mokhtari, Sanjay Shakkottai:  
**Exploiting Shared Representations for Personalized Federated Learning.** 2089-2099
- Adrien Corenflos, James Thornton, George Deligiannidis, Arnaud Doucet:  
**Differentiable Particle Filtering via Entropy-Regularized Optimal Transport.** 2100-2111
- José Correa, Andrés Cristi, Paul Duetting, Ashkan Norouzi-Fard:  
**Fairness and Bias in Online Selection.** 2112-2121
- Corinna Cortes, Mehryar Mohri, Ananda Theertha Suresh:  
**Relative Deviation Margin Bounds.** 2122-2131
- Corinna Cortes, Mehryar Mohri, Ananda Theertha Suresh, Ningshan Zhang:  
**A Discriminative Technique for Multiple-Source Adaptation.** 2132-2143
- Amanda Coston, Ashesh Rambachan, Alexandra Chouldechova:  
**Characterizing Fairness Over the Set of Good Models Under Selective Labels.** 2144-2155
- Romain Couillet, Florent Chatelain, Nicolas Le Bihan:  
**Two-way kernel matrix puncturing: towards resource-efficient PCA and spectral clustering.** 2156-2165
- Jonathan Crabbé, Mihaela van der Schaar:  
**Explaining Time Series Predictions with Dynamic Masks.** 2166-2177
- Zac Cranko, Zhan Shi, Xinhua Zhang, Richard Nock, Simon Kornblith:  
**Generalised Lipschitz Regularisation Equals Distributional Robustness.** 2178-2188
- Elliot Creager, Jörn-Henrik Jacobsen, Richard S. Zemel:  
**Environment Inference for Invariant Learning.** 2189-2200
- Francesco Croce, Matthias Hein:  
**Mind the Box:  $\ell_1$ -APGD for Sparse Adversarial Attacks on Image Classifiers.** 2201-2211
- Wentao Cui, Yuhong Guo:  
**Parameterless Transductive Feature Re-representation for Few-Shot Learning.** 2212-2221
- Shuang Cui, Kai Han, Tianshuai Zhu, Jing Tang, Benwei Wu, He Huang:  
**Randomized Algorithms for Submodular Function Maximization with a k-System Constraint.** 2222-2232
- Jingyi Cui, Hanyuan Hang, Yisen Wang, Zhouchen Lin:  
**GBHT: Gradient Boosting Histogram Transform for Density Estimation.** 2233-2243
- Chris Cummins, Zacharias V. Fisches, Tal Ben-Nun, Torsten Hoefer, Michael F. P. O'Boyle, Hugh Leather:  
**PrograML: A Graph-based Program Representation for Data Flow Analysis and Compiler Optimizations.** 2244-2253
- Sebastian Curi, Ilija Bogunovic, Andreas Krause :  
**Combining Pessimism with Optimism for Robust and Efficient Model-Based Deep Reinforcement Learning.** 2254-2264
- Mihaela Curmei, Sarah Dean, Benjamin Recht:  
**Quantifying Availability and Discovery in Recommender Systems via Stochastic Reachability.** 2265-2275
- Ashok Cutkosky , Christoph Dann, Abhimanyu Das, Claudio Gentile, Aldo Pacchiano, Manish Purohit:  
**Dynamic Balancing for Model Selection in Bandits and RL.** 2276-2285
- Stéphane d'Ascoli, Hugo Touvron, Matthew L. Leavitt, Ari S. Morcos, Giulio Biroli, Levent Sagun:  
**ConViT: Improving Vision Transformers with Soft Convolutional Inductive Biases.** 2286-2296
- Tommaso d'Orsi, Gleb Novikov, David Steurer :  
**Consistent regression when oblivious outliers overwhelm.** 2297-2306
- Robert Dadashi, Shideh Rezaifar, Nino Vieillard, Léonard Hussenot, Olivier Pietquin, Matthieu Geist:  
**Offline Reinforcement Learning with Pseudometric Learning.** 2307-2318
- Shabnam Daghanghi, Tharun Medini, Nicholas Meisburger, Beidi Chen, Mengnan Zhao, Anshumali Shrivastava:  
**A Tale of Two Efficient and Informative Negative Sampling Distributions.** 2319-2329
- Kunal Dahiya, Ananya Agarwal, Deepak Saini, Gururaj K, Jian Jiao, Amit Singh, Sumeet Agarwal, Purushottam Kar, Manik Varma:  
**SiameseXML: Siamese Networks meet Extreme Classifiers with 100M Labels.** 2330-2340
- Yogesh Dahiya, Fedor V. Fomin, Fahad Panolan, Kirill Simonov:  
**Fixed-Parameter and Approximation Algorithms for PCA with Outliers.** 2341-2351
- Biwei Dai , Uros Seljak:  
**Sliced Iterative Normalizing Flows.** 2352-2364
- Tuan Dam, Carlo D'Eramo, Jan Peters, Joni Pajarinen:  
**Convex Regularization in Monte-Carlo Tree Search.** 2365-2375



- Christopher R. Dance, Julien Perez, Théo Cachet:  
**Demonstration-Conditioned Reinforcement Learning for Few-Shot Imitation.** 2376-2387
- Mohamad H. Danesh, Anurag Koul, Alan Fern, Saeed Khorram:  
**Re-understanding Finite-State Representations of Recurrent Policy Networks.** 2388-2397
- Amir Daneshmand, Gesualdo Scutari, Pavel E. Dvurechensky, Alexander V. Gasnikov:  
**Newton Method over Networks is Fast up to the Statistical Precision.** 2398-2409
- Dominic Danks, Christopher Yau:  
**BasisDeVAE: Interpretable Simultaneous Dimensionality Reduction and Feature-Level Clustering with Derivative-Based Variational Autoencoders.** 2410-2420
- Giannis Daras, Joseph Dean, Ajil Jalal, Alex Dimakis:  
**Intermediate Layer Optimization for Inverse Problems using Deep Generative Models.** 2421-2432
- Mohammad Zalbagi Darestani, Akshay S. Chaudhari, Reinhard Heckel:  
**Measuring Robustness in Deep Learning Based Compressive Sensing.** 2433-2444
- Lokesh Chandra Das, Myounggyu Won:  
**SAINT-ACC: Safety-Aware Intelligent Adaptive Cruise Control for Autonomous Vehicles Using Deep Reinforcement Learning.** 2445-2455
- George Dasoulas, Kevin Scaman, Aladin Virmaux:  
**Lipschitz normalization for self-attention layers with application to graph neural networks.** 2456-2466
- Jyotikrishna Dass, Rabi N. Mahapatra:  
**Householder Sketch for Accurate and Accelerated Least-Mean-Squares Solvers.** 2467-2477
- Deepesh Data, Suhas N. Diggavi:  
**Byzantine-Resilient High-Dimensional SGD with Local Iterations on Heterogeneous Data.** 2478-2488
- Jared Quincy Davis, Albert Gu, Krzysztof Choromanski, Tri Dao, Christopher Ré, Chelsea Finn, Percy Liang:  
**Catformer: Designing Stable Transformers via Sensitivity Analysis.** 2489-2499
- Quinlan Dawkins, Tianxi Li, Haifeng Xu:  
**Diffusion Source Identification on Networks with Statistical Confidence.** 2500-2509
- Erik A. Daxberger, Eric T. Nalisnick, James Urquhart Allingham, Javier Antorán, José Miguel Hernández-Lobato:  
**Bayesian Deep Learning via Subnetwork Inference.** 2510-2521
- Giacomo De Palma, Bobak Toussi Kiani, Seth Lloyd:  
**Adversarial Robustness Guarantees for Random Deep Neural Networks.** 2522-2534
- Filip de Roos, Alexandra Gessner, Philipp Hennig:  
**High-Dimensional Gaussian Process Inference with Derivatives.** 2535-2545
- Lucas Deecke, Lukas Ruff, Robert A. Vandermeulen, Hakan Bilen (D):  
**Transfer-Based Semantic Anomaly Detection.** 2546-2558
- Javier Dehesa, Andrew Vidler, Julian A. Padget, Christof Lutteroth:  
**Grid-Functioned Neural Networks.** 2559-2567
- Erik D. Demaine, Adam Hesterberg, Frederic Koehler, Jayson Lynch, John C. Urschel:  
**Multidimensional Scaling: Approximation and Complexity.** 2568-2578
- Weijian Deng, Stephen Gould, Liang Zheng (D):  
**What Does Rotation Prediction Tell Us about Classifier Accuracy under Varying Testing Environments?** 2579-2589
- Zhun Deng, Hangfeng He, Weijie J. Su:  
**Toward Better Generalization Bounds with Locally Elastic Stability.** 2590-2600
- Yuan Deng, Sébastien Lahaie, Vahab S. Mirrokni, Song Zuo:  
**Revenue-Incentive Tradeoffs in Dynamic Reserve Pricing.** 2601-2610
- Don Kurian Dennis, Tian Li, Virginia Smith:  
**Heterogeneity for the Win: One-Shot Federated Clustering.** 2611-2620
- Mohammad Mahdi Derakhshani, Xiantong Zhen, Ling Shao, Cees Snoek:  
**Kernel Continual Learning.** 2621-2631
- Aryan Deshwal, Syrine Belakaria, Janardhan Rao Doppa:  
**Bayesian Optimization over Hybrid Spaces.** 2632-2643
- Sam Devlin, Raluca Georgescu, Ida Momennejad, Jaroslaw Rzepecki, Evelyn Zuniga, Gavin Costello, Guy Leroy, Ali Shaw, Katja Hofmann:  
**Navigation Turing Test (NTT): Learning to Evaluate Human-Like Navigation.** 2644-2653
- Laurens Devos, Wannes Meert, Jesse Davis:  
**Versatile Verification of Tree Ensembles.** 2654-2664
- Oussama Dhifallah, Yue M. Lu:  
**On the Inherent Regularization Effects of Noise Injection During Training.** 2665-2675
- Laxman Dhulipala, David Eisenstat, Jakub Lacki, Vahab S. Mirrokni, Jessica Shi:  
**Hierarchical Agglomerative Graph Clustering in Nearly-Linear Time.** 2676-2686



- Ilias Diakonikolas, Vasilis Kontonis, Christos Tzamos, Ali Vakilian , Nikos Zarifis:  
**Learning Online Algorithms with Distributional Advice.** 2687-2696
- Theo Diamandis, Yonina C. Eldar, Alireza Fallah, Farzan Farnia, Asuman E. Ozdaglar:  
**A Wasserstein Minimax Framework for Mixed Linear Regression.** 2697-2706
- Charles Dickens, Connor Pryor, Eriq Augustine, Alexander Miller, Lise Getoor:  
**Context-Aware Online Collective Inference for Templatized Graphical Models.** 2707-2716
- Aleksandar Dimitriev, Mingyuan Zhou:  
**ARMS: Antithetic-REINFORCE-Multi-Sample Gradient for Binary Variables.** 2717-2727
- Fan Ding, Jianzhu Ma, Jinbo Xu, Yexiang Xue:  
**XOR-CD: Linearly Convergent Constrained Structure Generation.** 2728-2738
- Tianyu Ding, Zhihui Zhu, René Vidal, Daniel P. Robinson:  
**Dual Principal Component Pursuit for Robust Subspace Learning: Theory and Algorithms for a Holistic Approach.** 2739-2748
- Tuan Dinh, Kangwook Lee:  
**Coded-InvNet for Resilient Prediction Serving Systems.** 2749-2759
- Vincent Divol, Théo Lacombe:  
**Estimation and Quantization of Expected Persistence Diagrams.** 2760-2770
- Carles Domingo-Enrich, Alberto Bietti, Eric Vanden-Eijnden, Joan Bruna:  
**On Energy-Based Models with Overparametrized Shallow Neural Networks.** 2771-2782
- Omar Darwiche Domingues, Pierre Ménard, Matteo Pirotta, Emilie Kaufmann, Michal Valko:  
**Kernel-Based Reinforcement Learning: A Finite-Time Analysis.** 2783-2792
- Yihe Dong, Jean-Baptiste Cordonnier, Andreas Loukas:  
**Attention is not all you need: pure attention loses rank doubly exponentially with depth.** 2793-2803
- Konstantin Donhauser, Mingqi Wu, Fanny Yang:  
**How rotational invariance of common kernels prevents generalization in high dimensions.** 2804-2814
- Radu-Alexandru Dragomir, Mathieu Even, Hadrien Hendrikx:  
**Fast Stochastic Bregman Gradient Methods: Sharp Analysis and Variance Reduction.** 2815-2825
- Simon S. Du, Sham M. Kakade, Jason D. Lee, Shachar Lovett, Gaurav Mahajan, Wen Sun, Ruosong Wang:  
**Bilinear Classes: A Structural Framework for Provable Generalization in RL.** 2826-2836
- Yilun Du, Shuang Li, Joshua B. Tenenbaum, Igor Mordatch:  
**Improved Contrastive Divergence Training of Energy-Based Models.** 2837-2848
- Cunxiao Du, Zhaopeng Tu, Jing Jiang:  
**Order-Agnostic Cross Entropy for Non-Autoregressive Machine Translation.** 2849-2859
- Elbert Du, Franklyn Wang, Michael Mitzenmacher:  
**Putting the "Learning" into Learning-Augmented Algorithms for Frequency Estimation.** 2860-2869
- Yali Du, Xue Yan, Xu Chen, Jun Wang, Haifeng Zhang:  
**Estimating  $\alpha$ -Rank from A Few Entries with Low Rank Matrix Completion.** 2870-2879
- Xuefeng Du, Jingfeng Zhang, Bo Han, Tongliang Liu, Yu Rong, Gang Niu, Junzhou Huang, Masashi Sugiyama:  
**Learning Diverse-Structured Networks for Adversarial Robustness.** 2880-2891
- Yaqi Duan, Chi Jin, Zhiyuan Li:  
**Risk Bounds and Rademacher Complexity in Batch Reinforcement Learning.** 2892-2902
- Zhibin Duan, Dongsheng Wang, Bo Chen, Chaojie Wang, Wenchao Chen, Yewen Li, Jie Ren, Mingyuan Zhou:  
**Sawtooth Factorial Topic Embeddings Guided Gamma Belief Network.** 2903-2913
- Arkopal Dutt, Andrey Y. Lokhov, Marc Vuffray, Sidhant Misra:  
**Exponential Reduction in Sample Complexity with Learning of Ising Model Dynamics.** 2914-2925
- Adrien Ecoffet, Joel Lehman:  
**Reinforcement Learning Under Moral Uncertainty.** 2926-2936
- Yonathan Efroni, Nadav Merlis, Aadhirupa Saha, Shie Mannor:  
**Confidence-Budget Matching for Sequential Budgeted Learning.** 2937-2947
- Theresa Eimer, André Biedenkapp, Frank Hutter, Marius Lindauer:  
**Self-Paced Context Evaluation for Contextual Reinforcement Learning.** 2948-2958
- Bryn Elesedy, Sheheryar Zaidi:  
**Provably Strict Generalisation Benefit for Equivariant Models.** 2959-2969
- Patrick Emami, Pan He, Sanjay Ranka, Anand Rangarajan :  
**Efficient Iterative Amortized Inference for Learning Symmetric and Disentangled Multi-Object Representations.** 2970-2981
- Melikasadat Emami, Mojtaba Sahraee-Ardakan, Parthe Pandit, Sundeep Rangan, Alyson K. Fletcher:  
**Implicit Bias of Linear RNNs.** 2982-2992
- Tolga Ergen, Mert Pilanci:  
**Global Optimality Beyond Two Layers: Training Deep ReLU Networks via Convex Programs.** 2993-3003



- Tolga Ergen, Mert Pilancı:  
**Revealing the Structure of Deep Neural Networks via Convex Duality.** 3004-3014
- Aleksandr Ermolov, Aliaksandr Siarohin, Enver Sangineto, Nicu Sebe:  
**Whitening for Self-Supervised Representation Learning.** 3015-3024
- Federico Errica, Davide Bacci, Alessio Micheli:  
**Graph Mixture Density Networks.** 3025-3035
- Yasaman Esfandiari, Sin Yong Tan, Zhanhong Jiang, Aditya Balu, Ethan Herron, Chinmay Hegde, Soumik Sarkar:  
**Cross-Gradient Aggregation for Decentralized Learning from Non-IID Data.** 3036-3046
- Panagiotis Eustratiadis, Henry Gouk, Da Li, Timothy M. Hospedales:  
**Weight-covariance alignment for adversarially robust neural networks.** 3047-3056
- Zalan Fabian, Reinhard Heckel, Mahdi Soltanolkotabi :  
**Data augmentation for deep learning based accelerated MRI reconstruction with limited data.** 3057-3067
- Xuhui Fan, Bin Li, Yaqiong Li, Scott A. Sisson:  
**Poisson-Randomised DirBN: Large Mutation is Needed in Dirichlet Belief Networks.** 3068-3077
- Ying Fan, Yifei Ming:  
**Model-based Reinforcement Learning for Continuous Control with Posterior Sampling.** 3078-3087
- Linxi Fan, Guanzhi Wang, De-An Huang, Zhiding Yu, Li Fei-Fei, Yuke Zhu, Animashree Anandkumar:  
**SECANT: Self-Expert Cloning for Zero-Shot Generalization of Visual Policies.** 3088-3099
- Guanhua Fang, Ping Li:  
**On Estimation in Latent Variable Models.** 3100-3110
- Guanhua Fang, Ping Li:  
**On Variational Inference in Biclustering Models.** 3111-3121
- Zhen Fang, Jie Lu, Anjin Liu, Feng Liu, Guangquan Zhang:  
**Learning Bounds for Open-Set Learning.** 3122-3132
- Shikai Fang, Zheng Wang, Zhimeng Pan, Ji Liu, Shandian Zhe:  
**Streaming Bayesian Deep Tensor Factorization.** 3133-3142
- Amir Massoud Farahmand, Mohammad Ghavamzadeh:  
**PID Accelerated Value Iteration Algorithm.** 3143-3153
- Vivek F. Farias, Andrew A. Li, Tianyi Peng:  
**Near-Optimal Entrywise Anomaly Detection for Low-Rank Matrices with Sub-Exponential Noise.** 3154-3163
- Gabriele Farina, Andrea Celli, Nicola Gatti, Tuomas Sandholm:  
**Connecting Optimal Ex-Ante Collusion in Teams to Extensive-Form Correlation: Faster Algorithms and Positive Complexity Results.** 3164-3173
- Farzan Farnia, Asuman E. Ozdaglar:  
**Train simultaneously, generalize better: Stability of gradient-based minimax learners.** 3174-3185
- Kilian Fatras, Thibault Séjourné, Rémi Flamary, Nicolas Courty:  
**Unbalanced minibatch Optimal Transport; applications to Domain Adaptation.** 3186-3197
- Yingjie Fei, Zhuoran Yang, Zhaoran Wang:  
**Risk-Sensitive Reinforcement Learning with Function Approximation: A Debiasing Approach.** 3198-3207
- Vitaly Feldman, Kunal Talwar:  
**Lossless Compression of Efficient Private Local Randomizers.** 3208-3219
- Zhili Feng, Praneeth Kacham, David P. Woodruff:  
**Dimensionality Reduction for the Sum-of-Distances Metric.** 3220-3229
- Zhe Feng, Sébastien Lahaie, Jon Schneider, Jinchao Ye:  
**Reserve Price Optimization for First Price Auctions in Display Advertising.** 3230-3239
- Ruili Feng, Zhouchen Lin, Jiapeng Zhu, Deli Zhao, Jingren Zhou, Zheng-Jun Zha:  
**Uncertainty Principles of Encoding GANs.** 3240-3251
- Lei Feng, Senlin Shu, Nan Lu, Bo Han, Miao Xu , Gang Niu, Bo An, Masashi Sugiyama:  
**Pointwise Binary Classification with Pairwise Confidence Comparisons.** 3252-3262
- Fei Feng, Wotao Yin, Alekh Agarwal, Lin Yang :  
**Provably Correct Optimization and Exploration with Non-linear Policies.** 3263-3273
- Haozhe Feng, Zhaoyang You, Minghao Chen, Tianye Zhang, Minfeng Zhu, Fei Wu, Chao Wu, Wei Chen:  
**KD3A: Unsupervised Multi-Source Decentralized Domain Adaptation via Knowledge Distillation.** 3274-3283
- Ruili Feng, Deli Zhao, Zheng-Jun Zha:  
**Understanding Noise Injection in GANs.** 3284-3293
- Matthias Fey, Jan Eric Lenssen, Frank Weichert, Jure Leskovec:  
**GNNAutoScale: Scalable and Expressive Graph Neural Networks via Historical Embeddings.** 3294-3304
- Angelos Filos, Clare Lyle, Yarin Gal, Sergey Levine, Natasha Jaques, Gregory Farquhar:  
**PsiPhi-Learning: Reinforcement Learning with Demonstrations using Successor Features and Inverse Temporal Difference Learning.** 3305-3317



- Marc Finzi, Max Welling, Andrew Gordon Wilson:  
**A Practical Method for Constructing Equivariant Multilayer Perceptrons for Arbitrary Matrix Groups.** 3318-3328
- Adam Fisch, Tal Schuster, Tommi S. Jaakkola, Regina Barzilay:  
**Few-Shot Conformal Prediction with Auxiliary Tasks.** 3329-3339
- Marc Fischer, Maximilian Baader, Martin T. Vechev:  
**Scalable Certified Segmentation via Randomized Smoothing.** 3340-3351
- Jonas Fischer, Anna Oláh, Jilles Vreeken:  
**What's in the Box? Exploring the Inner Life of Neural Networks with Robust Rules.** 3352-3362
- Genevieve Flaspoehler, Francesco Orabona, Judah Cohen, Soukaina Mouatadid, Miruna Oprescu, Paulo Orenstein, Lester Mackey:  
**Online Learning with Optimism and Delay.** 3363-3373
- Xavier Fontaine, Pierre Perrault, Michal Valko, Vianney Perchet:  
**Online A-Optimal Design and Active Linear Regression.** 3374-3383
- Adam Foster, Desi R. Ivanova, Ilyas Malik, Tom Rainforth:  
**Deep Adaptive Design: Amortizing Sequential Bayesian Experimental Design.** 3384-3395
- Dimitris Fotakis, Georgios Piliouras, Stratis Skoulakis:  
**Efficient Online Learning for Dynamic k-Clustering.** 3396-3406
- Yann Fraboni, Richard Vidal, Laetitia Kameni, Marco Lorenzi:  
**Clustered Sampling: Low-Variance and Improved Representativity for Clients Selection in Federated Learning.** 3407-3416
- Spencer Frei, Yuan Cao, Quanquan Gu:  
**Agnostic Learning of Halfspaces with Gradient Descent via Soft Margins.** 3417-3426
- Spencer Frei, Yuan Cao, Quanquan Gu:  
**Provable Generalization of SGD-trained Neural Networks of Any Width in the Presence of Adversarial Label Noise.** 3427-3438
- Tobias Freidling, Benjamin Poignard, Héctor Climente-González, Makoto Yamada:  
**Post-selection inference with HSIC-Lasso.** 3439-3448
- Thomas Frerix, Dmitrii Kochkov, Jamie A. Smith, Daniel Cremers, Michael P. Brenner, Stephan Hoyer:  
**Variational Data Assimilation with a Learned Inverse Observation Operator.** 3449-3458
- Christian Fröhlich, Alexandra Gessner, Philipp Hennig, Bernhard Schölkopf, Georgios Arvanitidis:  
**Bayesian Quadrature on Riemannian Data Manifolds.** 3459-3468
- Cheng Fu, Hanxian Huang, Xinyun Chen, Yuandong Tian, Jishen Zhao:  
**Learn-to-Share: A Hardware-friendly Transfer Learning Framework Exploiting Computation and Parameter Sharing.** 3469-3479
- Xiang Fu, Ge Yang, Pulkit Agrawal, Tommi S. Jaakkola:  
**Learning Task Informed Abstractions.** 3480-3491
- Yonggan Fu, Qixuan Yu, Meng Li, Vikas Chandra, Yingyan Lin:  
**Double-Win Quant: Aggressively Winning Robustness of Quantized Deep Neural Networks via Random Precision Training and Inference.** 3492-3504
- Yonggan Fu, Yongan Zhang, Yang Zhang, David D. Cox, Yingyan Lin:  
**Auto-NBA: Efficient and Effective Search Over the Joint Space of Networks, Bitwidths, and Accelerators.** 3505-3517
- Scott Fujimoto, David Meger, Doina Precup:  
**A Deep Reinforcement Learning Approach to Marginalized Importance Sampling with the Successor Representation.** 3518-3529
- Marco Fumero, Luca Cosmo, Simone Melzi, Emanuele Rodolà:  
**Learning disentangled representations via product manifold projection.** 3530-3540
- Hiroki Furuta, Tatsuya Matsushima, Tadashi Kozuno, Yutaka Matsuo, Sergey Levine, Ofir Nachum, Shixiang Shane Gu:  
**Policy Information Capacity: Information-Theoretic Measure for Task Complexity in Deep Reinforcement Learning.** 3541-3552
- Yansong Gao, Pratik Chaudhari:  
**An Information-Geometric Distance on the Space of Tasks.** 3553-3563
- Ruize Gao, Feng Liu, Jingfeng Zhang, Bo Han, Tongliang Liu, Gang Niu, Masashi Sugiyama:  
**Maximum Mean Discrepancy Test is Aware of Adversarial Attacks.** 3564-3575
- Qingzhe Gao, Bin Wang, Libin Liu, Baoquan Chen:  
**Unsupervised Co-part Segmentation through Assembly.** 3576-3586
- Yi Gao, Min-Ling Zhang:  
**Discriminative Complementary-Label Learning with Weighted Loss.** 3587-3597
- Saurabh Garg, Sivaraman Balakrishnan, J. Zico Kolter, Zachary C. Lipton:  
**RATT: Leveraging Unlabeled Data to Guarantee Generalization.** 3598-3609
- Saurabh Garg, Joshua Zhanson, Emilio Parisotto, Adarsh Prasad, J. Zico Kolter, Zachary C. Lipton, Sivaraman Balakrishnan, Ruslan Salakhutdinov, Pradeep Ravikumar:

**On Proximal Policy Optimization's Heavy-tailed Gradients.** 3610-3619

Damien Garreau, Dina Mardaoui:  
**What does LIME really see in images?** 3620-3629



Camille-Sovanney Gauthier, Romaric Gaudel, Élisa Fromont, Boammani Aser Lombo:  
**Parametric Graph for Unimodal Ranking Bandit.** 3630-3639



Floris Geerts, Filip Mazowiecki, Guillermo A. Pérez:  
**Let's Agree to Degree: Comparing Graph Convolutional Networks in the Message-Passing Framework.** 3640-3649



Tomas Geffner, Justin Domke:  
**On the difficulty of unbiased alpha divergence minimization.** 3650-3659



Amanda Gentzel, Purva Pruthi, David D. Jensen:  
**How and Why to Use Experimental Data to Evaluate Methods for Observational Causal Inference.** 3660-3671



Ganesh Ghalme, Vineet Nair, Itay Eilat, Inbal Talgam-Cohen, Nir Rosenfeld:  
**Strategic Classification in the Dark.** 3672-3681



Seyed Kamyar Seyed Ghasemipour, Dale Schuurmans, Shixiang Shane Gu:  
**EMaQ: Expected-Max Q-Learning Operator for Simple Yet Effective Offline and Online RL.** 3682-3691



Badih Ghazi, Ravi Kumar, Pasin Manurangsi, Rasmus Pagh, Amer Sinha:  
**Differentially Private Aggregation in the Shuffle Model: Almost Central Accuracy in Almost a Single Message.** 3692-3701



Rohan Ghuge, Anupam Gupta, Viswanath Nagarajan:  
**The Power of Adaptivity for Stochastic Submodular Cover.** 3702-3712



Jennifer Gillenwater, Matthew Joseph, Alex Kulesza:  
**Differentially Private Quantiles.** 3713-3722



Grzegorz Gluch, Rüdiger L. Urbanke:  
**Query Complexity of Adversarial Attacks.** 3723-3733



Florin Gogianu, Tudor Berariu, Mihaela Rosca, Claudia Clopath, Lucian Busoniu, Razvan Pascanu:  
**Spectral Normalisation for Deep Reinforcement Learning: An Optimisation Perspective.** 3734-3744



Tomer Golany, Kira Radinsky, Daniel Freedman, Saar Minha:  
**12-Lead ECG Reconstruction via Koopman Operators.** 3745-3754



Muhammad Waleed Gondal, Shruti Joshi, Nasim Rahaman, Stefan Bauer, Manuel Wuthrich, Bernhard Schölkopf:  
**Function Contrastive Learning of Transferable Meta-Representations.** 3755-3765



Wenbo Gong, Kaibo Zhang, Yingzhen Li, José Miguel Hernández-Lobato:  
**Active Slices for Sliced Stein Discrepancy.** 3766-3776



Sruthi Gorantla, Amit Deshpande, Anand Louis:  
**On the Problem of Underranking in Group-Fair Ranking.** 3777-3787



Eduard Gorbunov, Konstantin Burlachenko, Zhize Li, Peter Richtárik:  
**MARINA: Faster Non-Convex Distributed Learning with Compression.** 3788-3798



Martijn Gösgens, Alexey Tikhonov, Liudmila Prokhorenkova:  
**Systematic Analysis of Cluster Similarity Indices: How to Validate Validation Measures.** 3799-3808



Ankit Goyal, Hei Law, Bowei Liu, Alejandro Newell, Jia Deng:  
**Revisiting Point Cloud Shape Classification with a Simple and Effective Baseline.** 3809-3820



Florian Graf, Christoph D. Hofer, Marc Niethammer, Roland Kwitt:  
**Dissecting Supervised Contrastive Learning.** 3821-3830



Will Grathwohl, Kevin Swersky, Milad Hashemi, David Duvenaud, Chris J. Maddison:  
**Oops I Took A Gradient: Scalable Sampling for Discrete Distributions.** 3831-3841



Ido Greenberg, Shie Mannor:  
**Detecting Rewards Deterioration in Episodic Reinforcement Learning.** 3842-3853



Jiaqi Gu, Guosheng Yin:  
**Crystallization Learning with the Delaunay Triangulation.** 3854-3863



Chaoyu Guan, Xin Wang, Wenwu Zhu:  
**AutoAttend: Automated Attention Representation Search.** 3864-3874



Limor Gultchin, David S. Watson, Matt J. Kusner, Ricardo Silva:  
**Operationalizing Complex Causes: A Pragmatic View of Mediation.** 3875-3885



Sergey Guminov, Pavel E. Dvurechensky, Nazarii Tupitsa, Alexander V. Gasnikov:  
**On a Combination of Alternating Minimization and Nesterov's Momentum.** 3886-3898



Hongyi Guo, Zuyue Fu, Zhuoran Yang, Zhaoran Wang:  
**Decentralized Single-Timescale Actor-Critic on Zero-Sum Two-Player Stochastic Games.** 3899-3909



Wenbo Guo, Xian Wu, Sui Huang, Xinyu Xing:  
**Adversarial Policy Learning in Two-player Competitive Games.** 3910-3919



Zongyu Guo, Zhizheng Zhang, Runsen Feng, Zhibo Chen:  
**Soft then Hard: Rethinking the Quantization in Neural Image Compression.** 3920-3929



- Tarun Gupta, Anuj Mahajan, Bei Peng, Wendelin Boehmer, Shimon Whiteson:  
**UneVEN: Universal Value Exploration for Multi-Agent Reinforcement Learning.** 3930-3941
- Chirag Gupta, Aaditya Ramdas:  
**Distribution-Free Calibration Guarantees for Histogram Binning without Sample Splitting.** 3942-3952
- Shantanu Gupta, Hao Wang, Zachary C. Lipton, Yuyang Wang:  
**Correcting Exposure Bias for Link Recommendation.** 3953-3963
- Mert Gürbüzbalaban, Umut Simsekli, Lingjiong Zhu:  
**The Heavy-Tail Phenomenon in SGD.** 3964-3975
- Nezihe Merve Gürel, Xiangyu Qi, Luka Rimanic, Ce Zhang, Bo Li:  
**Knowledge Enhanced Machine Learning Pipeline against Diverse Adversarial Attacks.** 3976-3987
- András György, Pooria Joulani:  
**Adapting to Delays and Data in Adversarial Multi-Armed Bandits.** 3988-3997
- Hassan Hafez-Kolahi, Behrad Moniri, Shohreh Kasaei, Mahdieh Soleymani Baghshah:  
**Rate-Distortion Analysis of Minimum Excess Risk in Bayesian Learning.** 3998-4007
- Nadav Hallak, Panayotis Mertikopoulos, Volkan Cevher :  
**Regret Minimization in Stochastic Non-Convex Learning via a Proximal-Gradient Approach.** 4008-4017
- Seungyul Han, Youngchul Sung:  
**Diversity Actor-Critic: Sample-Aware Entropy Regularization for Sample-Efficient Exploration.** 4018-4029
- Yanjun Han, Yining Wang, Xi Chen:  
**Adversarial Combinatorial Bandits with General Non-linear Reward Functions.** 4030-4039
- Mengyue Hang, Jennifer Neville, Bruno Ribeiro:  
**A Collective Learning Framework to Boost GNN Expressiveness for Node Classification.** 4040-4050
- Austin W. Hanjie, Victor Zhong, Karthik Narasimhan:  
**Grounding Language to Entities and Dynamics for Generalization in Reinforcement Learning.** 4051-4062
- Botao Hao, Yaqi Duan, Tor Lattimore, Csaba Szepesvári, Mengdi Wang:  
**Sparse Feature Selection Makes Batch Reinforcement Learning More Sample Efficient.** 4063-4073
- Botao Hao, Xiang Ji, Yaqi Duan, Hao Lu, Csaba Szepesvári, Mengdi Wang:  
**Bootstrapping Fitted Q-Evaluation for Off-Policy Inference.** 4074-4084
- Yi Hao, Alon Orlitsky:  
**Compressed Maximum Likelihood.** 4085-4095
- Jason S. Hartford, Victor Veitch, Dhanya Sridhar, Kevin Leyton-Brown:  
**Valid Causal Inference with (Some) Invalid Instruments.** 4096-4106
- Tatsunori Hashimoto:  
**Model Performance Scaling with Multiple Data Sources.** 4107-4116
- Jakob Drachmann Havtorn, Jes Frellsen, Søren Hauberg, Lars Maaløe:  
**Hierarchical VAEs Know What They Don't Know.** 4117-4128
- Jonathan Hayase, Weihao Kong, Raghav Somani, Sewoong Oh:  
**Defense against backdoor attacks via robust covariance estimation.** 4129-4139
- Elad Hazan, Karan Singh:  
**Boosting for Online Convex Optimization.** 4140-4149
- Chaoyang He, Shen Li, Mahdi Soltanolkotabi , Salman Avestimehr:  
**PipeTransformer: Automated Elastic Pipelining for Distributed Training of Large-scale Models.** 4150-4159
- Yuhang He, Niki Trigoni, Andrew Markham:  
**SoundDet: Polyphonic Moving Sound Event Detection and Localization from Raw Waveform.** 4160-4170
- Jiafan He, Dongruo Zhou, Quanquan Gu:  
**Logarithmic Regret for Reinforcement Learning with Linear Function Approximation.** 4171-4180
- Mohsen Heidari, Jithin K. Sreedharan, Gil I. Shamir, Wojciech Szpankowski:  
**Finding Relevant Information via a Discrete Fourier Expansion.** 4181-4191
- Amélie Héliou, Matthieu Martin, Panayotis Mertikopoulos, Thibaud Rahier:  
**Zeroth-Order Non-Convex Learning via Hierarchical Dual Averaging.** 4192-4202
- Ryan Henderson, Djork-Arné Clevert, Floriane Montanari:  
**Improving Molecular Graph Neural Network Explainability with Orthonormalization and Induced Sparsity.** 4203-4213
- Matteo Hessel, Ivo Danihelka, Fabio Viola, Arthur Guez, Simon Schmitt, Laurent Sifre, Theophane Weber, David Silver, Hado van Hasselt:  
**Muesli: Combining Improvements in Policy Optimization.** 4214-4226
- Sophie Hilgard, Nir Rosenfeld, Mahzarin R. Banaji, Jack Cao, David C. Parkes:  
**Learning Representations by Humans, for Humans.** 4227-4238
- Gaurush Hiranandani, Jatin Mathur, Harikrishna Narasimhan, Mahdi Milani Fard, Sanmi Koyejo:  
**Optimizing Black-box Metrics with Iterative Example Weighting.** 4239-4249



- Roy Hirsch, Ran Gilad-Bachrach:  
**Trees with Attention for Set Prediction Tasks.** 4250-4261
- Liam Hodgkinson, Michael W. Mahoney:  
**Multiplicative Noise and Heavy Tails in Stochastic Optimization.** 4262-4274
- Pieter-Jan Hoedt, Frederik Kratzert, Daniel Klotz, Christina Halmich, Markus Holzleitner, Grey Nearing, Sepp Hochreiter, Günter Klambauer:  
**MC-LSTM: Mass-Conserving LSTM.** 4275-4286
- Derek Hoiem, Tanmay Gupta, Zhizhong Li, Michal Shlapentokh-Rothman:  
**Learning Curves for Analysis of Deep Networks.** 4287-4296
- Peter Holderrieth, Michael J. Hutchinson, Yee Whye Teh:  
**Equivariant Learning of Stochastic Fields: Gaussian Processes and Steerable Conditional Neural Processes.** 4297-4307
- Joey Hong, David Dohan, Rishabh Singh, Charles Sutton, Manzil Zaheer:  
**Latent Programmer: Discrete Latent Codes for Program Synthesis.** 4308-4318
- Sangwoo Hong, Heecheol Yang, Youngseok Yoon, Taehyun Cho, Jungwoo Lee:  
**Chebyshev Polynomial Codes: Task Entanglement-based Coding for Distributed Matrix Multiplication.** 4319-4327
- Hossein Hosseini, Hyunsin Park, Sungrock Yun, Christos Louizos, Joseph Soriaga, Max Welling:  
**Federated Learning of User Verification Models Without Sharing Embeddings.** 4328-4336
- Ya-Ping Hsieh, Panayotis Mertikopoulos, Volkan Cevher :  
**The Limits of Min-Max Optimization Algorithms: Convergence to Spurious Non-Critical Sets.** 4337-4348
- Jiachen Hu, Xiaoyu Chen, Chi Jin, Lihong Li, Liwei Wang:  
**Near-Optimal Representation Learning for Linear Bandits and Linear RL.** 4349-4358
- Zhengmian Hu, Heng Huang:  
**On the Random Conjugate Kernel and Neural Tangent Kernel.** 4359-4368
- Hengyuan Hu, Adam Lerer, Brandon Cui, Luis Pineda, Noam Brown, Jakob N. Foerster:  
**Off-Belief Learning.** 4369-4379
- Hao Hu, Jianing Ye, Guangxiang Zhu, Zhizhou Ren, Chongjie Zhang:  
**Generalizable Episodic Memory for Deep Reinforcement Learning.** 4380-4390
- Kaixun Hua, Mingfei Shi, Yankai Cao:  
**A Scalable Deterministic Global Optimization Algorithm for Clustering Problems.** 4391-4401
- Haiying Huang, Adnan Darwiche:  
**On Recovering from Modeling Errors Using Testing Bayesian Networks.** 4402-4411
- Jiawei Huang, Ruomin Huang, Wenjie Liu, Nikolaos M. Freris, Hu Ding:  
**A Novel Sequential Coreset Method for Gradient Descent Algorithms.** 4412-4422
- Baihe Huang, Xiaoxiao Li, Zhao Song, Xin Yang:  
**FL-NTK: A Neural Tangent Kernel-based Framework for Federated Learning Analysis.** 4423-4434
- Hengguan Huang, Hongfu Liu, Hao Wang, Chang Xiao, Ye Wang:  
**STRODE: Stochastic Boundary Ordinary Differential Equation.** 4435-4445
- Minhui Huang, Shiqian Ma, Lifeng Lai:  
**A Riemannian Block Coordinate Descent Method for Computing the Projection Robust Wasserstein Distance.** 4446-4455
- Minhui Huang, Shiqian Ma, Lifeng Lai:  
**Projection Robust Wasserstein Barycenters.** 4456-4465
- Itay Hubara, Yury Nahshan, Yair Hanani, Ron Banner, Daniel Soudry:  
**Accurate Post Training Quantization With Small Calibration Sets.** 4466-4475
- Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Mohammadamin Barekatain, Simon Schmitt, David Silver:  
**Learning and Planning in Complex Action Spaces.** 4476-4486
- Drew A. Hudson, Larry Zitnick:  
**Generative Adversarial Transformers.** 4487-4499
- Zeshan M. Hussain, Rahul G. Krishnan, David A. Sontag:  
**Neural Pharmacodynamic State Space Modeling.** 4500-4510
- Léonard Hussenot, Marcin Andrychowicz, Damien Vincent, Robert Dadashi, Anton Raichuk, Sabela Ramos, Nikola Momchev, Sertan Girgin, Raphaël Marinier, Lukasz Stafiniak, Manu Orsini, Olivier Bachem, Matthieu Geist, Olivier Pietquin:  
**Hyperparameter Selection for Imitation Learning.** 4511-4522
- Todd Huster, Jeremy E. J. Cohen, Zinan Lin, Kevin Chan, Charles A. Kamhoua, Nandi O. Leslie, Cho-Yu Jason Chiang, Vyas Sekar:  
**Pareto GAN: Extending the Representational Power of GANs to Heavy-Tailed Distributions.** 4523-4532
- Michael J. Hutchinson, Charline Le Lan, Sheheryar Zaidi, Emilien Dupont, Yee Whye Teh, Hyunjik Kim:  
**LieTransformer: Equivariant Self-Attention for Lie Groups.** 4533-4543



- Shahana Ibrahim, Xiao Fu:  
**Crowdsourcing via Annotator Co-occurrence Imputation and Provable Symmetric Nonnegative Matrix Factorization.** 4544-4554
- Maximilian Ilse, Jakub M. Tomczak, Patrick Forré:  
**Selecting Data Augmentation for Simulating Interventions.** 4555-4562
- Alexander Immer, Matthias Bauer, Vincent Fortuin, Gunnar Rätsch, Mohammad Emtiyaz Khan:  
**Scalable Marginal Likelihood Estimation for Model Selection in Deep Learning.** 4563-4573
- Yu Inatsu, Shogo Iwazaki, Ichiro Takeuchi:  
**Active Learning for Distributionally Robust Level-Set Estimation.** 4574-4584
- Hedda Cohen Indelman, Tamir Hazan:  
**Learning Randomly Perturbed Structured Predictors for Direct Loss Minimization.** 4585-4595
- Shariq Iqbal, Christian A. Schröder de Witt, Bei Peng, Wendelin Boehmer, Shimon Whiteson, Fei Sha:  
**Randomized Entity-wise Factorization for Multi-Agent Reinforcement Learning.** 4596-4606
- Haque Ishfaq, Qiwen Cui, Viet Nguyen, Alex Ayoub, Zhuoran Yang, Zhaoran Wang, Doina Precup, Lin Yang:  
**Randomized Exploration in Reinforcement Learning with General Value Function Approximation.** 4607-4616
- Rustem Islamov, Xun Qian, Peter Richtárik:  
**Distributed Second Order Methods with Fast Rates and Compressed Communication.** 4617-4628
- Pavel Izmailov, Sharad Vikram, Matthew D. Hoffman, Andrew Gordon Wilson:  
**What Are Bayesian Neural Network Posteriors Really Like?** 4629-4640
- Zachary Izzo, Lexing Ying, James Zou:  
**How to Learn when Data Reacts to Your Model: Performative Gradient Descent.** 4641-4650
- Andrew Jaegle, Felix Gimeno, Andy Brock, Oriol Vinyals, Andrew Zisserman, João Carreira:  
**Perceiver: General Perception with Iterative Attention.** 4651-4664
- Andrew Jaegle, Yury Sulsky, Arun Ahuja, Jake Bruce, Rob Fergus, Greg Wayne:  
**Imitation by Predicting Observations.** 4665-4676
- Jafar Jafarov, Sanchit Kalhan, Konstantin Makarychev, Yury Makarychev:  
**Local Correlation Clustering with Asymmetric Classification Errors.** 4677-4686
- Meena Jagadeesan, Celestine Mendler-Dünner, Moritz Hardt:  
**Alternative Microfoundations for Strategic Classification.** 4687-4697
- Ayush Jain, Alon Orlitsky:  
**Robust Density Estimation from Batches: The Best Things in Life are (Nearly) Free.** 4698-4708
- Ajil Jalal, Sushrut Karmalkar, Alex Dimakis, Eric Price:  
**Instance-Optimal Compressed Sensing via Posterior Sampling.** 4709-4720
- Ajil Jalal, Sushrut Karmalkar, Jessica Hoffmann, Alex Dimakis, Eric Price:  
**Fairness for Image Generation with Uncertain Sensitive Attributes.** 4721-4732
- Hamid Jalalzai, Rémi Leluc:  
**Feature Clustering for Support Identification in Extreme Regions.** 4733-4743
- Kyoungseok Jang, Kwang-Sung Jun, Se-Young Yun, Wanmo Kang:  
**Improved Regret Bounds of Bilinear Bandits using Action Space Analysis.** 4744-4754
- Daniel Jarrett, Alihan Hüyük, Mihaela van der Schaar:  
**Inverse Decision Modeling: Learning Interpretable Representations of Behavior.** 4755-4771
- Stanislaw Jastrzebski, Devansh Arpit, Oliver Åstrand, Giancarlo Kerg, Huan Wang, Caiming Xiong, Richard Socher, Kyunghyun Cho, Krzysztof J. Geras:  
**Catastrophic Fisher Explosion: Early Phase Fisher Matrix Impacts Generalization.** 4772-4784
- Zaynah Javed, Daniel S. Brown, Satvik Sharma, Jerry Zhu, Ashwin Balakrishna, Marek Petrik, Anca D. Dragan, Ken Goldberg:  
**Policy Gradient Bayesian Robust Optimization for Imitation Learning.** 4785-4796
- Rajesh Jayaram, Alireza Samadian, David P. Woodruff, Peng Ye:  
**In-Database Regression in Input Sparsity Time.** 4797-4806
- Vivek Jayaram, John Thickstun:  
**Parallel and Flexible Sampling from Autoregressive Models via Langevin Dynamics.** 4807-4818
- Taewon Jeong, Heeyoung Kim:  
**Objective Bound Conditional Gaussian Process for Bayesian Optimization.** 4819-4828
- Andrew Jesson, Sören Mindermann, Yarin Gal, Uri Shalit:  
**Quantifying Ignorance in Individual-Level Causal-Effect Estimates under Hidden Confounding.** 4829-4838
- Nandan Kumar Jha, Zahra Ghodsi, Siddharth Garg, Brandon Reagen:  
**DeepReDuce: ReLU Reduction for Fast Private Inference.** 4839-4849
- Aditi Jha, Michael J. Morais, Jonathan W. Pillow :  
**Factor-analytic inverse regression for high-dimension, small-sample dimensionality reduction.** 4850-4859



- Ziwei Ji, Nathan Srebro, Matus Telgarsky:  
**Fast margin maximization via dual acceleration.** 4860-4869
- Geng Ji, Debora Sujono, Erik B. Sudderth:  
**Marginalized Stochastic Natural Gradients for Black-Box Variational Inference.** 4870-4881
- Kaiyi Ji, Junjie Yang, Yingbin Liang:  
**Bilevel Optimization: Convergence Analysis and Enhanced Design.** 4882-4892
- Sheng Jia, Ehsan Nezhadarya, Yuhuai Wu, Jimmy Ba:  
**Efficient Statistical Tests: A Neural Tangent Kernel Approach.** 4893-4903
- Chao Jia, Yinfai Yang, Ye Xia, Yi-Ting Chen, Zarana Parekh, Hieu Pham, Quoc V. Le, Yun-Hsuan Sung, Zhen Li, Tom Duerig:  
**Scaling Up Visual and Vision-Language Representation Learning With Noisy Text Supervision.** 4904-4916
- Bin-Bin Jia, Min-Ling Zhang:  
**Multi-Dimensional Classification via Sparse Label Encoding.** 4917-4926
- Ziyu Jiang, Tianlong Chen, Bobak J. Mortazavi, Zhangyang Wang:  
**Self-Damaging Contrastive Learning.** 4927-4939
- Minqi Jiang, Edward Grefenstette, Tim Rockäschel:  
**Prioritized Level Replay.** 4940-4950
- Yuankun Jiang, Chenglin Li, Wenrui Dai, Junni Zou, Hongkai Xiong:  
**Monotonic Robust Policy Optimization with Model Discrepancy.** 4951-4960
- Haotian Jiang, Zhong Li, Qianxiao Li:  
**Approximation Theory of Convolutional Architectures for Time Series Modelling.** 4961-4970
- Shuli Jiang, Dennis Li, Irene Mengze Li, Arvind V. Mahankali, David P. Woodruff:  
**Streaming and Distributed Algorithms for Robust Column Subset Selection.** 4971-4981
- Yifei Jiang, Yi Li, Yiming Sun, Jiaxin Wang, David P. Woodruff:  
**Single Pass Entrywise-Transformed Low Rank Approximation.** 4982-4991
- Jiechuan Jiang, Zongqing Lu:  
**The Emergence of Individuality.** 4992-5001
- Zhihao Jiang, Pinyan Lu, Zhihao Gavin Tang, Yuhao Zhang:  
**Online Selection Problems against Constrained Adversary.** 5002-5012
- Heinrich Jiang, Afshin Rostamizadeh:  
**Active Covering.** 5013-5022
- Ray Jiang, Tom Zahavy, Zhongwen Xu, Adam White, Matteo Hessel, Charles Blundell, Hado van Hasselt:  
**Emphatic Algorithms for Deep Reinforcement Learning.** 5023-5033
- Ziheng Jiang, Chiyuan Zhang, Kunal Talwar, Michael C. Mozer:  
**Characterizing Structural Regularities of Labeled Data in Overparameterized Models.** 5034-5044
- Tianyuan Jin, Keke Huang, Jing Tang, Xiaokui Xiao:  
**Optimal Streaming Algorithms for Multi-Armed Bandits.** 5045-5054
- Yujia Jin, Aaron Sidford:  
**Towards Tight Bounds on the Sample Complexity of Average-reward MDPs.** 5055-5064
- Tianyuan Jin, Jing Tang, Pan Xu, Keke Huang, Xiaokui Xiao, Quanquan Gu:  
**Almost Optimal Anytime Algorithm for Batched Multi-Armed Bandits.** 5065-5073
- Tianyuan Jin, Pan Xu, Jieming Shi , Xiaokui Xiao, Quanquan Gu:  
**MOTS: Minimax Optimal Thompson Sampling.** 5074-5083
- Ying Jin, Zhuoran Yang, Zhaoran Wang:  
**Is Pessimism Provably Efficient for Offline RL?** 5084-5096
- Mingxuan Jing, Wenbing Huang, Fuchun Sun, Xiaojian Ma, Tao Kong, Chuang Gan, Lei Li:  
**Adversarial Option-Aware Hierarchical Imitation Learning.** 5097-5106
- Changhun Jo, Kangwook Lee:  
**Discrete-Valued Latent Preference Matrix Estimation with Graph Side Information.** 5107-5117
- Matt Jordan, Alex Dimakis:  
**Provable Lipschitz Certification for Generative Models.** 5118-5126
- Martin Jørgensen, Søren Hauberg:  
**Isometric Gaussian Process Latent Variable Model for Dissimilarity Data.** 5127-5136
- Peizhong Ju, Xiaojun Lin, Ness B. Shroff:  
**On the Generalization Power of Overfitted Two-Layer Neural Tangent Kernel Models.** 5137-5147
- Kwang-Sung Jun, Lalit Jain, Houssam Nassif, Blake Mason:  
**Improved Confidence Bounds for the Linear Logistic Model and Applications to Bandits.** 5148-5157
- Ji Hyung Jung, Hye Won Chung, Ji Oon Lee:  
**Detection of Signal in the Spiked Rectangular Models.** 5158-5167



- Yonghan Jung, Jin Tian, Elias Bareinboim:  
**Estimating Identifiable Causal Effects on Markov Equivalence Class through Double Machine Learning.** 5168-5179
- Mustafa Devrim Kaba, Chong You, Daniel P. Robinson, Enrique Mallada, René Vidal:  
**A Nullspace Property for Subspace-Preserving Recovery.** 5180-5188
- Anil Kag, Venkatesh Saligrama:  
**Training Recurrent Neural Networks via Forward Propagation Through Time.** 5189-5200
- Peter Kairouz, Ziyu Liu, Thomas Steinke:  
**The Distributed Discrete Gaussian Mechanism for Federated Learning with Secure Aggregation.** 5201-5212
- Peter Kairouz, Brendan McMahan, Shuang Song, Om Thakkar, Abhradeep Thakurta, Zheng Xu:  
**Practical and Private (Deep) Learning Without Sampling or Shuffling.** 5213-5225
- Hiroshi Kajino:  
**A Differentiable Point Process with Its Application to Spiking Neural Networks.** 5226-5235
- Vasileios Kalantzis, Georgios Kollias, Shashanka Ubaru, Athanasios N. Nikolakopoulos, Lior Horesh, Kenneth L. Clarkson:  
**Projection techniques to update the truncated SVD of evolving matrices with applications.** 5236-5246
- Nathan Kallus, Yuta Saito, Masatoshi Uehara:  
**Optimal Off-Policy Evaluation from Multiple Logging Policies.** 5247-5256
- Angeliki Kamoutsi, Goran Banjac, John Lygeros:  
**Efficient Performance Bounds for Primal-Dual Reinforcement Learning from Demonstrations.** 5257-5268
- Anthimos Vardis Kandiros, Yuval Dagan, Nishanth Dikkala, Surbhi Goel, Constantinos Daskalakis:  
**Statistical Estimation from Dependent Data.** 5269-5278
- Sanyam Kapoor, Marc Finzi, Ke Alexander Wang, Andrew Gordon Wilson:  
**SKling on Simplices: Kernel Interpolation on the Permutohedral Lattice for Scalable Gaussian Processes.** 5279-5289
- Sanyam Kapoor, Theofanis Karaletsos, Thang D. Bui:  
**Variational Auto-Regressive Gaussian Processes for Continual Learning.** 5290-5300
- Nikos Karampatziakis, Paul Mineiro, Aaditya Ramdas:  
**Off-Policy Confidence Sequences.** 5301-5310
- Sai Praneeth Karimireddy, Lie He, Martin Jaggi:  
**Learning from History for Byzantine Robust Optimization.** 5311-5319
- Masahiro Kato, Takeshi Teshima:  
**Non-Negative Bregman Divergence Minimization for Deep Direct Density Ratio Estimation.** 5320-5333
- Julian Katz-Samuels, Jifan Zhang, Lalit Jain, Kevin Jamieson:  
**Improved Algorithms for Agnostic Pool-based Active Classification.** 5334-5344
- Yigitcan Kaya, Tudor Dumitras:  
**When Does Data Augmentation Help With Membership Inference Attacks?** 5345-5355
- Ehsan Kazemi, Shervin Minaee, Moran Feldman, Amin Karbasi:  
**Regularized Submodular Maximization at Scale.** 5356-5366
- Varun A. Kelkar, Mark A. Anastasio:  
**Prior Image-Constrained Reconstruction using Style-Based Generative Models.** 5367-5377
- T. Anderson Keller, Jorn W. T. Peters, Priyank Jaini, Emiel Hoogeboom, Patrick Forré, Max Welling:  
**Self Normalizing Flows.** 5378-5387
- Henry Kenlay, Dorina Thanou, Xiaowen Dong:  
**Interpretable Stability Bounds for Spectral Graph Filters.** 5388-5397
- Thomas Kerdreux, Lewis Liu, Simon Lacoste-Julien, Damien Scieur:  
**Affine Invariant Analysis of Frank-Wolfe on Strongly Convex Sets.** 5398-5408
- David Khachaturov, Ilia Shumailov, Yiren Zhao, Nicolas Papernot, Ross J. Anderson:  
**Markpainting: Adversarial Machine Learning meets Inpainting.** 5409-5419
- Sajad Khodadadian, Zaiwei Chen, Siva Theja Maguluri:  
**Finite-Sample Analysis of Off-Policy Natural Actor-Critic Algorithm.** 5420-5431
- Valentin Khrulkov, Artem Babenko, Ivan V. Oseledets:  
**Functional Space Analysis of Local GAN Convergence.** 5432-5442
- Patrick Kidger, Ricky T. Q. Chen, Terry J. Lyons:  
**"Hey, that's not an ODE": Faster ODE Adjoint via Seminorms.** 5443-5452
- Patrick Kidger, James Foster, Xuechen Li, Terry J. Lyons:  
**Neural SDEs as Infinite-Dimensional GANs.** 5453-5463
- KrishnaTeja Killamsetty, Durga Sivasubramanian, Ganesh Ramakrishnan, Abir De, Rishabh K. Iyer:  
**GRAD-MATCH: Gradient Matching based Data Subset Selection for Efficient Deep Model Training.** 5464-5474
- Kwang In Kim:  
**Improving Predictors via Combination Across Diverse Task Categories.** 5475-5485



- Junsu Kim, Sungsoo Ahn, Hankook Lee, Jinwoo Shin:  
**Self-Improved Retrosynthetic Planning.** 5486-5495
- Kuno Kim, Shivam Garg, Kirankumar Shiragur, Stefano Ermon:  
**Reward Identification in Inverse Reinforcement Learning.** 5496-5505
- Sehoon Kim, Amir Gholami, Zhewei Yao, Michael W. Mahoney, Kurt Keutzer:  
**I-BERT: Integer-only BERT Quantization.** 5506-5518
- Taehyeong Kim, Injune Hwang, Hyundo Lee, Hyunseo Kim, Won-Seok Choi, Joseph J. Lim, Byoung-Tak Zhang:  
**Message Passing Adaptive Resonance Theory for Online Active Semi-supervised Learning.** 5519-5529
- Jaehyeon Kim, Jungil Kong, Juhee Son:  
**Conditional Variational Autoencoder with Adversarial Learning for End-to-End Text-to-Speech.** 5530-5540
- Dong-Ki Kim, Miao Liu, Matthew Riemer, Chuangchuang Sun, Marwa Abdulhai, Golnaz Habibi, Sebastian Lopez-Cot, Gerald Tesauro, Jonathan P. How:  
**A Policy Gradient Algorithm for Learning to Learn in Multiagent Reinforcement Learning.** 5541-5550
- Timothy D. Kim, Thomas Z. Luo, Jonathan W. Pillow , Carlos D. Brody:  
**Inferring Latent Dynamics Underlying Neural Population Activity via Neural Differential Equations.** 5551-5561
- Hyunjik Kim, George Papamakarios, Andriy Mnih:  
**The Lipschitz Constant of Self-Attention.** 5562-5571
- Jaekyeom Kim , Seohong Park, Gunhee Kim:  
**Unsupervised Skill Discovery with Bottleneck Option Learning.** 5572-5582
- Wonjae Kim, Bokyung Son, Ildoo Kim:  
**ViLT: Vision-and-Language Transformer Without Convolution or Region Supervision.** 5583-5594
- Johannes Kirschner, Andreas Krause :  
**Bias-Robust Bayesian Optimization via Dueling Bandits.** 5595-5605
- Dani Kiyasseh, Tingting Zhu, David A. Clifton:  
**CLOCS: Contrastive Learning of Cardiac Signals Across Space, Time, and Patients.** 5606-5615
- Johannes Klicpera, Marten Lienen, Stephan Günnemann:  
**Scalable Optimal Transport in High Dimensions for Graph Distances, Embedding Alignment, and More.** 5616-5627
- Frederic Koehler, Viraj Mehta, Andrej Risteski:  
**Representational aspects of depth and conditioning in normalizing flows.** 5628-5636
- Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanas Phillips, Irena Gao, Tony Lee, Etienne David, Ian Stavness, Wei Guo, Berton Earnshaw, Imran S. Haque, Sara M. Beery, Jure Leskovec, Anshul Kundaje, Emma Pierson, Sergey Levine, Chelsea Finn, Percy Liang:  
**WILDS: A Benchmark of in-the-Wild Distribution Shifts.** 5637-5664
- Vladimir Kolmogorov, Thomas Pock:  
**One-sided Frank-Wolfe algorithms for saddle problems.** 5665-5675
- Abi Komanduru, Jean Honorio :  
**A Lower Bound for the Sample Complexity of Inverse Reinforcement Learning.** 5676-5685
- Lingjing Kong, Tao Lin, Anastasia Koloskova, Martin Jaggi, Sebastian U. Stich:  
**Consensus Control for Decentralized Deep Learning.** 5686-5696
- Mikhail Konobeev, Ilja Kuzborskij, Csaba Szepesvári:  
**A Distribution-dependent Analysis of Meta Learning.** 5697-5706
- Anna-Kathrin Kopetzki, Bertrand Charpentier, Daniel Zügner, Sandhya Giri, Stephan Günnemann:  
**Evaluating Robustness of Predictive Uncertainty Estimation: Are Dirichlet-based Models Reliable?** 5707-5718
- Anna Korba, Pierre-Cyril Aubin-Frankowski, Szymon Majewski, Pierre Albin:  
**Kernel Stein Discrepancy Descent.** 5719-5730
- Jack Kosaian, Amar Phanishayee, Matthai Philipose, Debadatta Dey, Rashmi Vinayak:  
**Boosting the Throughput and Accelerator Utilization of Specialized CNN Inference Beyond Increasing Batch Size.** 5731-5741
- Adam R. Kosiorek, Heiko Strathmann, Daniel Zoran, Pol Moreno, Rosalia Schneider, Sona Mokrá, Danilo Jimenez Rezende:  
**NeRF-VAE: A Geometry Aware 3D Scene Generative Model.** 5742-5752
- Jannik Kossen, Sebastian Farquhar, Yarin Gal, Tom Rainforth:  
**Active Testing: Sample-Efficient Model Evaluation.** 5753-5763
- James E. Kostas, Yash Chandak, Scott M. Jordan, Georgios Theocharous, Philip S. Thomas:  
**High Confidence Generalization for Reinforcement Learning.** 5764-5773
- Ilya Kostrikov, Rob Fergus, Jonathan Tompson, Ofir Nachum:  
**Offline Reinforcement Learning with Fisher Divergence Critic Regularization.** 5774-5783
- Dmitry Kovalev, Egor Shulgin, Peter Richtárik, Alexander Rogozin, Alexander V. Gasnikov:  
**ADOM: Accelerated Decentralized Optimization Method for Time-Varying Networks.** 5784-5793
- Tadashi Kozuno, Yunhao Tang, Mark Rowland, Rémi Munos, Steven Kapturowski, Will Dabney, Michal Valko, David Abel:  
**Revisiting Peng's Q( $\lambda$ ) for Modern Reinforcement Learning.** 5794-5804



- Sanath Kumar Krishnamurthy, Vitor Hadad, Susan Athey:  
**Adapting to misspecification in contextual bandits with offline regression oracles.** 5805-5814
- David Krueger, Ethan Caballero, Jörn-Henrik Jacobsen, Amy Zhang, Jonathan Binas, Dinghuai Zhang, Rémi Le Priol, Aaron C. Courville:  
**Out-of-Distribution Generalization via Risk Extrapolation (REx).** 5815-5826
- Arun K. Kuchibhotla, Qinqing Zheng:  
**Near-Optimal Confidence Sequences for Bounded Random Variables.** 5827-5837
- Tejas Kulkarni, Joonas Jälkö, Antti Koskela, Samuel Kaski, Antti Honkela:  
**Differentially Private Bayesian Inference for Generalized Linear Models.** 5838-5849
- Abhishek Kumar, Sunabha Chatterjee, Piyush Rai:  
**Bayesian Structural Adaptation for Continual Learning.** 5850-5860
- Abhishek Kumar, Harikrishna Narasimhan, Andrew Cotter:  
**Implicit rate-constrained optimization of non-decomposable objectives.** 5861-5871
- Christian Kümmerle, Claudio Mayrink Verdun:  
**A Scalable Second Order Method for Ill-Conditioned Matrix Completion from Few Samples.** 5872-5883
- Branislav Kveton, Mikhail Konobeev, Manzil Zaheer, Chih-Wei Hsu, Martin Mladenov, Craig Boutilier, Csaba Szepesvári:  
**Meta-Thompson Sampling.** 5884-5893
- Minae Kwon, Siddharth Karamcheti, Mariano-Florentino Cuellar, Dorsa Sadigh:  
**Targeted Data Acquisition for Evolving Negotiation Agents.** 5894-5904
- Jungmin Kwon, Jeongseop Kim, Hyunseo Park, In Kwon Choi:  
**ASAM: Adaptive Sharpness-Aware Minimization for Scale-Invariant Learning of Deep Neural Networks.** 5905-5914
- Eduardo Sany Laber, Lucas Murtinho:  
**On the price of explainability for some clustering problems.** 5915-5925
- Jonathan Lacotte, Yifei Wang, Mert Pilancı:  
**Adaptive Newton Sketch: Linear-time Optimization with Quadratic Convergence and Effective Hessian Dimensionality.** 5926-5936
- Pierre Laforgue, Guillaume Staerman, Stéphan Cléménçon:  
**Generalization Bounds in the Presence of Outliers: a Median-of-Means Study.** 5937-5947
- Thanh Chi Lam, Trong Nghia Hoang, Bryan Kian Hsiang Low, Patrick Jaillet:  
**Model Fusion for Personalized Learning.** 5948-5958
- Maximilian Lam, Gu-Yeon Wei, David Brooks, Vijay Janapa Reddi, Michael Mitzenmacher:  
**Gradient Disaggregation: Breaking Privacy in Federated Learning by Reconstructing the User Participant Matrix.** 5959-5968
- Tal Lancewicki, Shahar Segal, Tomer Koren, Yishay Mansour:  
**Stochastic Multi-Armed Bandits with Unrestricted Delay Distributions.** 5969-5978
- Mikel Landajuela, Brenden K. Petersen, Sookyung Kim, Cláudio P. Santiago, Ruben Glatt, T. Nathan Mundhenk, Jacob F. Pettit, Daniel M. Faissol:  
**Discovering symbolic policies with deep reinforcement learning.** 5979-5989
- Hunter Lang, David A. Sontag, Aravindan Vijayaraghavan:  
**Graph Cuts Always Find a Global Optimum for Potts Models (With a Catch).** 5990-5999
- Jan-Hendrik Lange, Paul Swoboda:  
**Efficient Message Passing for 0-1 ILPs with Binary Decision Diagrams.** 6000-6010
- Kasper Green Larsen, Rasmus Pagh, Jakub Tetek:  
**CountSketches, Feature Hashing and the Median of Three.** 6011-6020
- Sophie Laternus, Philipp Berens:  
**MorphVAE: Generating Neural Morphologies from 3D-Walks using a Variational Autoencoder with Spherical Latent Space.** 6021-6031
- Nevena Lazic, Dong Yin, Yasin Abbasi-Yadkori, Csaba Szepesvári:  
**Improved Regret Bound and Experience Replay in Regularized Policy Iteration.** 6032-6042
- Trung Le, Tuan Nguyen, Nhat Ho, Hung Bui, Dinh Phung:  
**LAMDA: Label Matching Deep Domain Adaptation.** 6043-6054
- Armin Lederer, Alejandro Jose Ordóñez Conejo, Korbinian Maier, Wenxin Xiao, Jonas Umlauft, Sandra Hirche:  
**Gaussian Process-Based Real-Time Learning for Safety Critical Applications.** 6055-6064
- Seungwon Lee, Sima Behpour, Eric Eaton:  
**Sharing Less is More: Lifelong Learning in Deep Networks with Selective Layer Transfer.** 6065-6075
- Joshua K. Lee, Yuheng Bu, Deepsa Rajan, Prasanna Sattigeri, Rameswar Panda, Subhro Das, Gregory W. Wornell:  
**Fair Selective Classification Via Sufficiency.** 6076-6086
- Moontae Lee, Sungjun Cho, Kun Dong, David Mimno, David Bindel:  
**On-the-fly Rectification for Robust Large-Vocabulary Topic Inference.** 6087-6097
- Dong Hoon Lee, Sae-Young Chung:  
**Unsupervised Embedding Adaptation via Early-Stage Feature Reconstruction for Few-Shot Classification.** 6098-



6108

- Sebastian Lee, Sebastian Goldt, Andrew M. Saxe:  
**Continual Learning in the Teacher-Student Setup: Impact of Task Similarity.** 6109-6119
- Jongmin Lee, Wonseok Jeon, Byung-Jun Lee, Joelle Pineau, Kee-Eung Kim:  
**OptiDICE: Offline Policy Optimization via Stationary Distribution Correction Estimation.** 6120-6130
- Kimin Lee, Michael Laskin, Aravind Srinivas, Pieter Abbeel:  
**SUNRISE: A Simple Unified Framework for Ensemble Learning in Deep Reinforcement Learning.** 6131-6141
- Chung-Wei Lee, Haipeng Luo, Chen-Yu Wei, Mengxiao Zhang, Xiaojin Zhang:  
**Achieving Near Instance-Optimality and Minimax-Optimality in Stochastic and Adversarial Linear Bandits Simultaneously.** 6142-6151
- Kimin Lee, Laura M. Smith, Pieter Abbeel:  
**PEBBLE: Feedback-Efficient Interactive Reinforcement Learning via Relabeling Experience and Unsupervised Pre-training.** 6152-6163
- Qi Lei, Wei Hu, Jason D. Lee:  
**Near-Optimal Linear Regression under Distribution Shift.** 6164-6174
- Yunwen Lei, Zhenhuan Yang, Tianbao Yang, Yiming Ying:  
**Stability and Generalization of Stochastic Gradient Methods for Minimax Problems.** 6175-6186
- Joel Z. Leibo, Edgar A. Duéñez-Guzmán, Alexander Vezhnevets, John P. Agapiou, Peter Sunehag, Raphael Koster, Jayd Matyas, Charlie Beattie, Igor Mordatch, Thore Graepel:  
**Scalable Evaluation of Multi-Agent Reinforcement Learning with Melting Pot.** 6187-6199
- Benedict J. Leimkuhler, Tiffany J. Vlaar, Timothée Pouchon, Amos J. Storkey:  
**Better Training using Weight-Constrained Stochastic Dynamics.** 6200-6211
- Klas Leino, Zifan Wang, Matt Fredrikson:  
**Globally-Robust Neural Networks.** 6212-6222
- Renato Paes Leme, Balasubramanian Sivan, Yifeng Teng, Pratik Worah:  
**Learning to Price Against a Moving Target.** 6223-6232
- Maud Lemercier, Christopher Salvi, Thomas Cass, Edwin V. Bonilla, Theodoros Damoulas, Terry J. Lyons:  
**SigGPDE: Scaling Sparse Gaussian Processes on Sequential Data.** 6233-6242
- Sagi Levanon, Nir Rosenfeld:  
**Strategic Classification Made Practical.** 6243-6253
- Alexander Levine, Soheil Feizi:  
**Improved, Deterministic Smoothing for  $L_1$  Certified Robustness.** 6254-6264
- Mike Lewis, Shruti Bhosale, Tim Dettmers, Naman Goyal, Luke Zettlemoyer:  
**BASE Layers: Simplifying Training of Large, Sparse Models.** 6265-6274
- José Lezama, Wei Chen, Qiang Qiu:  
**Run-Sort-ReRun: Escaping Batch Size Limitations in Sliced Wasserstein Generative Models.** 6275-6285
- Zhize Li, Hongyan Bao, Xiangliang Zhang, Peter Richtárik:  
**PAGE: A Simple and Optimal Probabilistic Gradient Estimator for Nonconvex Optimization.** 6286-6295
- Gen Li, Changxiao Cai, Yuxin Chen, Yuantao Gu, Yuting Wei, Yuejie Chi:  
**Tightening the Dependence on Horizon in the Sample Complexity of Q-Learning.** 6296-6306
- Wenshuo Li, Hanting Chen, Mingqiang Huang, Xinghao Chen, Chunjing Xu, Yunhe Wang:  
**Winograd Algorithm for AdderNet.** 6307-6315
- Yuhang Li, Shikuang Deng, Xin Dong, Ruihao Gong, Shi Gu:  
**A Free Lunch From ANN: Towards Efficient, Accurate Spiking Neural Networks Calibration.** 6316-6325
- Xiling Li, Rafael Dowsley, Martine De Cock:  
**Privacy-Preserving Feature Selection with Secure Multiparty Computation.** 6326-6336
- Gen Li, Yuantao Gu:  
**Theory of Spectral Method for Union of Subspaces-Based Random Geometry Graph.** 6337-6345
- Kevin Li, Abhishek Gupta, Ashwin Reddy, Vitchyr H. Pong, Aurick Zhou, Justin Yu, Sergey Levine:  
**MURAL: Meta-Learning Uncertainty-Aware Rewards for Outcome-Driven Reinforcement Learning.** 6346-6356
- Tian Li, Shengyuan Hu, Ahmad Beirami, Virginia Smith:  
**Ditto: Fair and Robust Federated Learning Through Personalization.** 6357-6368
- Xiaoyun Li, Ping Li:  
**Quantization Algorithms for Random Fourier Features.** 6369-6380
- Bo Li, Lijun Li, Ankang Sun, Chenhao Wang, Yingfan Wang:  
**Approximate Group Fairness for Clustering.** 6381-6391
- Shaojie Li, Yong Liu:  
**Sharper Generalization Bounds for Clustering.** 6392-6402
- Xuefeng Li, Tongliang Liu, Bo Han, Gang Niu, Masashi Sugiyama:  
**Provably End-to-end Label-noise Learning without Anchor Points.** 6403-6413



- Duanshun Li, Jing Liu, Dongeun Lee, Ali Seyedmazloom, Giridhar Kaushik, Kookjin Lee, Noseong Park:  
**A Novel Method to Solve Neural Knapsack Problems.** 6414-6424
- Haoran Li, Wei Lu:  
**Mixed Cross Entropy Loss for Neural Machine Translation.** 6425-6436
- Guohao Li, Matthias Müller, Bernard Ghanem, Vladlen Koltun:  
**Training Graph Neural Networks with 1000 Layers.** 6437-6449
- Yang Li, Junier Oliva:  
**Active Feature Acquisition with Generative Surrogate Models.** 6450-6459
- Yang Li, Junier Oliva:  
**Partially Observed Exchangeable Modeling.** 6460-6470
- Zhong Li, Minxue Pan, Tian Zhang, Xuandong Li:  
**Testing DNN-based Autonomous Driving Systems under Critical Environmental Conditions.** 6471-6482
- Xiaocheng Li, Chunlin Sun, Yinyu Ye:  
**The Symmetry between Arms and Knapsacks: A Primal-Dual Approach for Bandits with Knapsacks.** 6483-6492
- Mengmeng Li, Tobias Sutter, Daniel Kuhn:  
**Distributionally Robust Optimization with Markovian Data.** 6493-6503
- Xiang Li, Shusen Wang, Kun Chen, Zhihua Zhang:  
**Communication-Efficient Distributed SVD via Local Power Iterations.** 6504-6514
- Bo Li, Qili Wang, Gim Hee Lee:  
**FILTRA: Rethinking Steerable CNN by Filter Transform.** 6515-6522
- Shi Li, Jiayi Xian:  
**Online Unrelated Machine Load Balancing with Predictions Revisited.** 6523-6532
- Zeng Li, Chuanlong Xie, Qinwen Wang:  
**Asymptotic Normality and Confidence Intervals for Prediction Risk of the Min-Norm Least Squares Estimator.** 6533-6542
- Zhuohan Li, Siyuan Zhuang, Shiyuan Guo, Danyang Zhuo , Hao Zhang, Dawn Song, Ion Stoica:  
**TeraPipe: Token-Level Pipeline Parallelism for Training Large-Scale Language Models.** 6543-6552
- Xiaoyu Li, Zhenxun Zhuang, Francesco Orabona:  
**A Second look at Exponential and Cosine Step Sizes: Simplicity, Adaptivity, and Performance.** 6553-6564
- Paul Pu Liang, Chiyu Wu, Louis-Philippe Morency, Ruslan Salakhutdinov:  
**Towards Understanding and Mitigating Social Biases in Language Models.** 6565-6576
- Kaizhao Liang, Jacky Y. Zhang, Boxin Wang, Zhuolin Yang, Sanmi Koyejo, Bo Li:  
**Uncovering the Connections Between Adversarial Transferability and Knowledge Transferability.** 6577-6587
- Tung-Che Liang, Jin Zhou, Yun-Sheng Chan, Tsung-Yi Ho , Krishnendu Chakrabarty, Cy Lee:  
**Parallel Droplet Control in MEDA Biochips using Multi-Agent Reinforcement Learning.** 6588-6599
- Peiyuan Liao, Han Zhao, Keyulu Xu, Tommi S. Jaakkola, Geoffrey J. Gordon, Stefanie Jegelka, Ruslan Salakhutdinov:  
**Information Obfuscation of Graph Neural Networks.** 6600-6610
- Gabriele Libardi, Gianni De Fabritiis, Sebastian Dittert:  
**Guided Exploration with Proximal Policy Optimization using a Single Demonstration.** 6611-6620
- Valerii Likhoshesterov, Xingyou Song, Krzysztof Choromanski, Jared Quincy Davis, Adrian Weller:  
**Debiasing a First-order Heuristic for Approximate Bi-level Optimization.** 6621-6630
- Chi-Heng Lin, Mehdi Azabou, Eva L. Dyer:  
**Making transport more robust and interpretable by moving data through a small number of anchor points.** 6631-6641
- Xiang Lin, Simeng Han, Shafiq R. Joty:  
**Straight to the Gradient: Learning to Use Novel Tokens for Neural Text Generation.** 6642-6653
- Tao Lin, Sai Praneeth Karimireddy, Sebastian U. Stich, Martin Jaggi:  
**Quasi-global Momentum: Accelerating Decentralized Deep Learning on Heterogeneous Data.** 6654-6665
- Wanyu Lin, Hao Lan, Baochun Li:  
**Generative Causal Explanations for Graph Neural Networks.** 6666-6679
- Wu Lin, Frank Nielsen, Mohammad Emtiyaz Khan, Mark Schmidt:  
**Tractable structured natural-gradient descent using local parameterizations.** 6680-6691
- Dominik Linzner, Heinz Koepll:  
**Active Learning of Continuous-time Bayesian Networks through Interventions.** 6692-6701
- Yaron Lipman:  
**Phase Transitions, Distance Functions, and Implicit Neural Representations.** 6702-6712
- Florian List:  
**The Earth Mover's Pinball Loss: Quantiles for Histogram-Valued Regression.** 6713-6724
- Yang Liu:  
**Understanding Instance-Level Label Noise: Disparate Impacts and Treatments.** 6725-6735



- Hao Liu, Pieter Abbeel:  
**APS: Active Pretraining with Successor Features.** 6736-6747
- Yang Liu, Jeremy Bernstein, Markus Meister, Yisong Yue:  
**Learning by Turning: Neural Architecture Aware Optimisation.** 6748-6758
- Guan-Horng Liu, Tianrong Chen, Evangelos A. Theodorou:  
**Dynamic Game Theoretic Neural Optimizer.** 6759-6769
- Hao Liu, Minshuo Chen, Tuo Zhao, Wenjing Liao:  
**Besov Function Approximation and Binary Classification on Low-Dimensional Manifolds Using Convolutional Residual Networks.** 6770-6780
- Evan Zheran Liu, Behzad Haghgoo, Annie S. Chen, Aditi Raghunathan, Pang Wei Koh, Shiori Sagawa, Percy Liang, Chelsea Finn:  
**Just Train Twice: Improving Group Robustness without Training Group Information.** 6781-6792
- Siqi Liu, Milos Hauskrecht:  
**Event Outlier Detection in Continuous Time.** 6793-6803
- Jiashuo Liu, Zheyuan Hu, Peng Cui, Bo Li, Zheyuan Shen:  
**Heterogeneous Risk Minimization.** 6804-6814
- Linfeng Liu, Michael C. Hughes, Soha Hassoun, Liping Liu:  
**Stochastic Iterative Graph Matching.** 6815-6825
- Iou-Jen Liu, Unnat Jain, Raymond A. Yeh, Alexander G. Schwing:  
**Cooperative Exploration for Multi-Agent Deep Reinforcement Learning.** 6826-6836
- Xiaorui Liu, Wei Jin, Yao Ma, Yaxin Li, Hua Liu, Yiqi Wang, Ming Yan, Jiliang Tang:  
**Elastic Graph Neural Networks.** 6837-6849
- Xinwang Liu, Li Liu, Qing Liao, Siwei Wang, Yi Zhang, Wenxuan Tu, Chang Tang, Jiyuan Liu, En Zhu:  
**One Pass Late Fusion Multi-view Clustering.** 6850-6859
- Bo Liu, Qiang Liu, Peter Stone, Animesh Garg, Yuke Zhu, Anima Anandkumar:  
**Coach-Player Multi-agent Reinforcement Learning for Dynamic Team Composition.** 6860-6870
- Yiwei Liu, Jiamou Liu, Kaibin Wan, Zhan Qin, Zijian Zhang, Bakhadyr Khoussainov, Liehuang Zhu:  
**From Local to Global Norm Emergence: Dissolving Self-reinforcing Substructures with Incremental Social Instruments.** 6871-6881
- Risheng Liu, Xuan Liu, Xiaoming Yuan, Shangzhi Zeng, Jin Zhang:  
**A Value-Function-based Interior-point Method for Non-convex Bi-level Optimization.** 6882-6892
- Shiwei Liu, Decebal Constantin Mocanu, Yulong Pei, Mykola Pechenizkiy:  
**Selfish Sparse RNN Training.** 6893-6904
- Rui Liu, Alex Olshevsky:  
**Temporal Difference Learning as Gradient Splitting.** 6905-6913
- Zifan Liu, Jongho Park, Theodoros Rekatsinas, Christos Tzamos:  
**On Robust Mean Estimation under Coordinate-level Corruption.** 6914-6924
- Evan Zheran Liu, Aditi Raghunathan, Percy Liang, Chelsea Finn:  
**Decoupling Exploration and Exploitation for Meta-Reinforcement Learning without Sacrifices.** 6925-6935
- Zechun Liu, Zhiqiang Shen, Shichao Li, Koen Helwegen, Dong Huang, Kwang-Ting Cheng:  
**How Do Adam and Training Strategies Help BNNs Optimization.** 6936-6946
- Yuhang Liu, Shiliang Sun:  
**SagaNet: A Small Sample Gated Network for Pediatric Cancer Diagnosis.** 6947-6956
- Boyang Liu, Mengying Sun, Ding Wang, Pang-Ning Tan, Jiayu Zhou:  
**Learning Deep Neural Networks under Agnostic Corrupted Supervision.** 6957-6967
- Terrance Liu, Giuseppe Vietri, Thomas Steinke, Jonathan R. Ullman, Zhiwei Steven Wu :  
**Leveraging Public Data for Practical Private Query Release.** 6968-6977
- Hanwen Liu, Zhenyu Weng, Yuesheng Zhu:  
**Watermarking Deep Neural Networks with Greedy Residuals.** 6978-6988
- Shiwei Liu, Lu Yin, Decebal Constantin Mocanu, Mykola Pechenizkiy:  
**Do We Actually Need Dense Over-Parameterization? In-Time Over-Parameterization in Sparse Training.** 6989-7000
- Qinghua Liu, Tiancheng Yu, Yu Bai, Chi Jin :  
**A Sharp Analysis of Model-based Reinforcement Learning with Self-Play.** 7001-7010
- Ning Liu, Geng Yuan, Zhengping Che, Xuan Shen, Xiaolong Ma, Qing Jin, Jian Ren, Jian Tang, Sijia Liu, Yanzhi Wang:  
**Lottery Ticket Preserves Weight Correlation: Is It Desirable or Not?** 7011-7020
- Liyang Liu, Shilong Zhang, Zhanghui Kuang, Aojun Zhou, Jing-Hao Xue, Xinjiang Wang, Yimin Chen, Wenming Yang, Qingmin Liao, Wayne Zhang:  
**Group Fisher Pruning for Practical Network Compression.** 7021-7032
- Lewis Liu, Yufeng Zhang, Zhuoran Yang, Reza Babanezhad, Zhaoran Wang:  
**Infinite-Dimensional Optimization for Zero-Sum Games via Variational Transport.** 7033-7044



- Kangqiao Liu, Ziyin Liu, Masahito Ueda:  
**Noise and Fluctuation of Finite Learning Rate Stochastic Gradient Descent.** 7045-7056
- Xutong Liu, Jinhang Zuo, Xiaowei Chen, Wei Chen, John C. S. Lui:  
**Multi-layered Network Exploration via Random Walks: From Offline Optimization to Online Learning.** 7057-7066
- Antoine Liutkus, Ondrej Cífká, Shih-Lun Wu, Umut Simsekli, Yi-Hsuan Yang, Gaél Richard:  
**Relative Positional Encoding for Transformers with Linear Complexity.** 7067-7079
- Alfonso Lobos, Paul Grigas, Zheng Wen:  
**Joint Online Learning and Decision-making via Dual Mirror Descent.** 7080-7089
- Federico López, Beatrice Pozzetti, Steve Trettel, Michael Strube, Anna Wienhard:  
**Symmetric Spaces for Graph Embeddings: A Finsler-Riemannian Approach.** 7090-7101
- Qian Lou , Lei Jiang:  
**HEMET: A Homomorphic-Encryption-Friendly Privacy-Preserving Mobile Neural Network Architecture.** 7102-7110
- Yucheng Lu, Christopher De Sa:  
**Optimal Complexity in Decentralized Training.** 7111-7123
- Yang Young Lu, Wenbo Guo, Xinyu Xing, William Stafford Noble:  
**DANCE: Enhancing saliency maps using decoys.** 7124-7133
- Nan Lu, Shida Lei, Gang Niu, Issei Sato, Masashi Sugiyama:  
**Binary Classification from Multiple Unlabeled Datasets via Surrogate Set Classification.** 7134-7144
- Yucheng Lu, Youngsuk Park, Lifan Chen, Yuyang Wang, Christopher De Sa, Dean P. Foster:  
**Variance Reduced Training with Stratified Sampling for Forecasting Models.** 7145-7155
- Yang Young Lu, Timothy C. Yu, Giancarlo Bonora, William Stafford Noble:  
**ACE: Explaining cluster from an adversarial perspective.** 7156-7167
- James Lucas, Juhan Bae, Michael R. Zhang, Stanislav Fort, Richard S. Zemel, Roger B. Grosse:  
**On Monotonic Linear Interpolation of Neural Network Parameters.** 7168-7179
- Denis Lukovnikov, Asja Fischer:  
**Improving Breadth-Wise Backpropagation in Graph Neural Networks Helps Learning Long-Range Dependencies.** 7180-7191
- Youzhi Luo, Keqiang Yan, Shuiwang Ji:  
**GraphDF: A Discrete Flow Model for Molecular Graph Generation.** 7192-7203
- Andrei Lupu, Brandon Cui, Hengyuan Hu, Jakob N. Foerster:  
**Trajectory Diversity for Zero-Shot Coordination.** 7204-7213
- Shahar Lutati, Lior Wolf:  
**HyperHyperNetwork for the Design of Antenna Arrays.** 7214-7223
- Michael Lutter, Shie Mannor, Jan Peters, Dieter Fox, Animesh Garg:  
**Value Iteration in Continuous Actions, States and Time.** 7224-7234
- Xingchen Ma , Matthew B. Blaschko:  
**Meta-Cal: Well-controlled Post-hoc Calibration by Ranking.** 7235-7245
- Baorui Ma, Zhizhong Han, Yu-Shen Liu, Matthias Zwicker:  
**Neural-Pull: Learning Signed Distance Function from Point clouds by Learning to Pull Space onto Surface.** 7246-7257
- Shaojun Ma, Shu Liu, Hongyuan Zha, Haomin Zhou:  
**Learning Stochastic Behaviour from Aggregate Data.** 7258-7267
- Peter Macgregor, He Sun:  
**Local Algorithms for Finding Densely Connected Clusters.** 7268-7278
- Divyam Madaan, Jinwoo Shin, Sung Ju Hwang:  
**Learning to Generate Noise for Multi-Attack Robustness.** 7279-7289
- Mauro Maggioni, Jason Miller, Hongda Qiu, Ming Zhong:  
**Learning Interaction Kernels for Agent Systems on Riemannian Manifolds.** 7290-7300
- Anuj Mahajan, Mikayel Samvelyan, Lei Mao , Viktor Makoviychuk, Animesh Garg, Jean Kossaifi, Shimon Whiteson, Yuke Zhu, Animashree Anandkumar:  
**Tesseract: Tensorised Actors for Multi-Agent Reinforcement Learning.** 7301-7312
- Divyat Mahajan, Shruti Tople, Amit Sharma:  
**Domain Generalization using Causal Matching.** 7313-7324
- Vien V. Mai, Mikael Johansson :  
**Stability and Convergence of Stochastic Gradient Clipping: Beyond Lipschitz Continuity and Smoothness.** 7325-7335
- Carol Mak, Fabian Zaiser , Luke Ong:  
**Nonparametric Hamiltonian Monte Carlo.** 7336-7347
- Maggie Makar, Lauren West, David Hooper, Eric Horvitz, Erica Shenoy, John V. Guttag :  
**Exploiting structured data for learning contagious diseases under incomplete testing.** 7348-7357



- Konstantin Makarychev, Liren Shan:  
**Near-Optimal Algorithms for Explainable k-Medians and k-Means.** 7358-7367
- Ashok Vardhan Makkula, Xiyang Liu, Mohammad Vahid Jamali, Hessam Mahdavifar, Sewoong Oh, Pramod Viswanath:  
**KO codes: inventing nonlinear encoding and decoding for reliable wireless communication via deep-learning.** 7368-7378
- Eran Malach, Pritish Kamath, Emmanuel Abbe, Nathan Srebro:  
**Quantifying the Benefit of Using Differentiable Learning over Tangent Kernels.** 7379-7389
- Shehryar Malik, Usman Anwar, Alireza Aghasi, Ali Ahmed:  
**Inverse Constrained Reinforcement Learning.** 7390-7399
- Osman Asif Malik, Stephen Becker:  
**A Sampling-Based Method for Tensor Ring Decomposition.** 7400-7411
- Dhruv Malik, Aldo Pacchiano, Vishwak Srinivasan, Yuanzhi Li:  
**Sample Efficient Reinforcement Learning In Continuous State Spaces: A Perspective Beyond Linearity.** 7412-7422
- Gustavo Malkomes, Bolong Cheng, Eric Hans Lee, Mike McCourt:  
**Beyond the Pareto Efficient Frontier: Constraint Active Search for Multiobjective Experimental Design.** 7423-7434
- Xueyu Mao , Deepayan Chakrabarti, Purnamrita Sarkar:  
**Consistent Nonparametric Methods for Network Assisted Covariate Estimation.** 7435-7446
- Weichao Mao, Kaiqing Zhang, Ruihao Zhu, David Simchi-Levi, Tamer Basar:  
**Near-Optimal Model-Free Reinforcement Learning in Non-Stationary Episodic MDPs.** 7447-7458
- Aymen Al Marjani, Alexandre Proutière:  
**Adaptive Sampling for Best Policy Identification in Markov Decision Processes.** 7459-7468
- João Marques-Silva, Thomas Gerspacher, Martin C. Cooper, Alexey Ignatiev, Nina Narodytska:  
**Explanations for Monotonic Classifiers.** 7469-7479
- Luke Marris, Paul Muller, Marc Lanctot, Karl Tuyls, Thore Graepel:  
**Multi-Agent Training beyond Zero-Sum with Correlated Equilibrium Meta-Solvers.** 7480-7491
- Natalia Martínez, Martín Bertrán, Afroditi Papadaki, Miguel R. D. Rodrigues, Guillermo Sapiro:  
**Blind Pareto Fairness and Subgroup Robustness.** 7492-7501
- Atalanti-Anastasia Mastakouri, Bernhard Schölkopf, Dominik Janzing:  
**Necessary and sufficient conditions for causal feature selection in time series with latent common causes.** 7502-7511
- Afsaneh Mastouri, Yuchen Zhu, Limor Gultchin, Anna Korba, Ricardo Silva, Matt J. Kusner, Arthur Gretton, Krikamol Muandet:  
**Proximal Causal Learning with Kernels: Two-Stage Estimation and Moment Restriction.** 7512-7523
- Andreas Maurer, Daniela Angela Parletta, Andrea Paudice, Massimiliano Pontil:  
**Robust Unsupervised Learning via L-statistic Minimization.** 7524-7533
- Alessio Mazzetto, Cyrus Cousins, Dylan Sam, Stephen H. Bach, Eli Upfal:  
**Adversarial Multi Class Learning under Weak Supervision with Performance Guarantees.** 7534-7543
- Mohammad Mehrabi, Adel Javanmard, Ryan A. Rossi, Anup B. Rao, Tung Mai:  
**Fundamental Tradeoffs in Distributionally Adversarial Training.** 7544-7554
- Jincheng Mei, Yue Gao, Bo Dai, Csaba Szepesvári, Dale Schuurmans:  
**Leveraging Non-uniformity in First-order Non-convex Optimization.** 7555-7564
- Eli A. Meirom, Haggai Maron, Shie Mannor, Gal Chechik:  
**Controlling Graph Dynamics with Reinforcement Learning and Graph Neural Networks.** 7565-7577
- Gabriel Mel, Surya Ganguli:  
**A theory of high dimensional regression with arbitrary correlations between input features and target functions: sample complexity, multiple descent curves and a hierarchy of phase transitions.** 7578-7587
- Joe Mellor, Jack Turner, Amos J. Storkey, Elliot J. Crowley:  
**Neural Architecture Search without Training.** 7588-7598
- Pierre Ménard, Omar Darwiche Domingues, Anders Jonsson, Emilie Kaufmann, Edouard Leurent, Michal Valko:  
**Fast active learning for pure exploration in reinforcement learning.** 7599-7608
- Pierre Ménard, Omar Darwiche Domingues, Xuedong Shang, Michal Valko:  
**UCB Momentum Q-learning: Correcting the bias without forgetting.** 7609-7618
- Tao Meng, Kai-Wei Chang:  
**An Integer Linear Programming Framework for Mining Constraints from Data.** 7619-7631
- Aditya Krishna Menon, Ankit Singh Rawat, Sashank J. Reddi, Seunghyeon Kim, Sanjiv Kumar:  
**A statistical perspective on distillation.** 7632-7642
- Amil Merchant, Luke Metz, Samuel S. Schoenholz, Ekin D. Cubuk:  
**Learn2Hop: Learned Optimization on Rough Landscapes.** 7643-7653
- Thomas Mesnard, Theophane Weber, Fabio Viola, Shantanu Thakoor, Alaa Saade, Anna Harutyunyan, Will Dabney, Thomas S. Stepleton, Nicolas Heess, Arthur Guez, Eric Moulines, Marcus Hutter, Lars Buesing, Rémi Munos:  
**Counterfactual Credit Assignment in Model-Free Reinforcement Learning.** 7654-7664



- Alberto Maria Metelli, Giorgia Ramponi, Alessandro Concetti, Marcello Restelli:  
**Provably Efficient Learning of Transferable Rewards.** 7665-7676
- Laurent Meunier, Meyer Scetbon, Rafael Pinot, Jamal Atif, Yann Chevaleyre:  
**Mixed Nash Equilibria in the Adversarial Examples Game.** 7677-7687
- David Henry Mguni, Yutong Wu, Yali Du, Yaodong Yang, Ziyi Wang, Minne Li, Ying Wen, Joel Jennings, Jun Wang:  
**Learning in Nonzero-Sum Stochastic Games with Potentials.** 7688-7699
- Chenfeng Miao, Shuang Liang, Zhengchen Liu, Minchuan Chen, Jun Ma, Shaojun Wang, Jing Xiao:  
**EfficientTTS: An Efficient and High-Quality Text-to-Speech Architecture.** 7700-7709
- John Miller, Juan C. Perdomo, Tijana Zrnic:  
**Outside the Echo Chamber: Optimizing the Performative Risk.** 7710-7720
- John Miller, Rohan Taori, Aditi Raghunathan, Shiori Sagawa, Pang Wei Koh, Vaishaal Shankar, Percy Liang, Yair Carmon, Ludwig Schmidt:  
**Accuracy on the Line: on the Strong Correlation Between Out-of-Distribution and In-Distribution Generalization.** 7721-7735
- Ming Min, Ruimeng Hu:  
**Signatured Deep Fictitious Play for Mean Field Games with Common Noise.** 7736-7747
- Dongchan Min, Dong Bok Lee, Eunho Yang, Sung Ju Hwang:  
**Meta-StyleSpeech : Multi-Speaker Adaptive Text-to-Speech Generation.** 7748-7759
- Hancheng Min, Salma Tarmouni, René Vidal, Enrique Mallada:  
**On the Explicit Role of Initialization on the Convergence and Implicit Bias of Overparametrized Linear Networks.** 7760-7768
- Graziano Mita, Maurizio Filippone, Pietro Michiardi:  
**An Identifiable Double VAE For Disentangled Representations.** 7769-7779
- Eric Mitchell, Rafael Rafailov, Xue Bin Peng, Sergey Levine, Chelsea Finn:  
**Offline Meta-Reinforcement Learning with Advantage Weighting.** 7780-7791
- Taiki Miyagawa, Akinori F. Ebihara:  
**The Power of Log-Sum-Exp: Sequential Density Ratio Matrix Estimation for Speed-Accuracy Optimization.** 7792-7804
- Miguel Zamora, Momchil Peychev, Sehoon Ha, Martin T. Vechev, Stelian Coros:  
**PODS: Policy Optimization via Differentiable Simulation.** 7805-7817
- Dustin Morrill, Ryan D'Orazio, Marc Lanctot, James R. Wright, Michael Bowling, Amy R. Greenwald:  
**Efficient Deviation Types and Learning for Hindsight Rationality in Extensive-Form Games.** 7818-7828
- James Morrill, Christopher Salvi, Patrick Kidger, James Foster:  
**Neural Rough Differential Equations for Long Time Series.** 7829-7838
- Michal Moshkovitz, Yao-Yuan Yang, Kamalika Chaudhuri:  
**Connecting Interpretability and Robustness in Decision Trees through Separation.** 7839-7849
- Debarghya Mukherjee, Aritra Guha, Justin M. Solomon, Yuekai Sun, Mikhail Yurochkin:  
**Outlier-Robust Optimal Transport.** 7850-7860
- Alexander Munteanu, Simon Omlor, David P. Woodruff:  
**Oblivious Sketching for Logistic Regression.** 7861-7871
- Tomoya Murata, Taiji Suzuki:  
**Bias-Variance Reduced Local SGD for Less Heterogeneous Federated Learning.** 7872-7881
- Kieran A. Murphy, Carlos Esteves, Varun Jampani, Srikumar Ramalingam, Ameesh Makadia:  
**Implicit-PDF: Non-Parametric Representation of Probability Distributions on the Rotation Manifold.** 7882-7893
- Mojmir Mutny, Andreas Krause:  
**No-regret Algorithms for Capturing Events in Poisson Point Processes.** 7894-7904
- Ofir Nabati, Tom Zahavy, Shie Mannor:  
**Online Limited Memory Neural-Linear Bandits with Likelihood Matching.** 7905-7915
- Akira Nakagawa, Keizo Kato, Taiji Suzuki:  
**Quantitative Understanding of VAE as a Non-linearly Scaled Isometric Embedding.** 7916-7926
- Daniel Wontae Nam, Younghoon Kim, Chan Y. Park:  
**GMAC: A Distributional Perspective on Actor-Critic Framework.** 7927-7936
- Deepak Narayanan, Amar Phanishayee, Kaiyu Shi, Xie Chen, Matei Zaharia:  
**Memory-Efficient Pipeline-Parallel DNN Training.** 7937-7947
- Shyam Narayanan, Sandeep Silwal, Piotr Indyk, Or Zamir:  
**Randomized Dimensionality Reduction for Facility Location and Single-Linkage Clustering.** 7948-7957
- Charlie Nash, Jacob Menick, Sander Dieleman, Peter W. Battaglia:  
**Generating images with sparse representations.** 7958-7968
- Viacheslav Natarovskii, Daniel Rudolf, Björn Sprungk:  
**Geometric convergence of elliptical slice sampling.** 7969-7978



- Niv Nayman, Yonathan Aflalo, Asaf Noy, Lih Zelnik:  
**HardCoRe-NAS: Hard Constrained differentiable Neural Architecture Search.** 7979-7990
- Kamal Ndousse, Douglas Eck, Sergey Levine, Natasha Jaques:  
**Emergent Social Learning via Multi-agent Reinforcement Learning.** 7991-8004
- Willie Neiswanger, Ke Alexander Wang, Stefano Ermon:  
**Bayesian Algorithm Execution: Estimating Computable Properties of Black-box Functions Using Mutual Information.** 8005-8015
- Hadi Nekoei, Akilesh Badrinarayanan, Aaron C. Courville, Sarath Chandar:  
**Continuous Coordination As a Realistic Scenario for Lifelong Learning.** 8016-8024
- Mark W. Nemecek, Ron Parr:  
**Policy Caches with Successor Features.** 8025-8033
- Elias Chaibub Neto:  
**Causality-aware counterfactual confounding adjustment as an alternative to linear residualization in anticausal prediction tasks based on linear learners.** 8034-8044
- Dung Daniel T. Ngo, Logan Stapleton, Vasilis Syrgkanis, Steven Wu :  
**Incentivizing Compliance with Algorithmic Instruments.** 8045-8055
- Quynh Nguyen:  
**On the Proof of Global Convergence of Gradient Descent for Deep ReLU Networks with Linear Widths.** 8056-8062
- Quoc Phong Nguyen, Zhongxiang Dai, Bryan Kian Hsiang Low, Patrick Jaillet:  
**Value-at-Risk Optimization with Gaussian Processes.** 8063-8072
- Xuan-Phi Nguyen, Shafiq R. Joty, Thanh-Tung Nguyen, Kui Wu, Ai Ti Aw:  
**Cross-model Back-translated Distillation for Unsupervised Machine Translation.** 8073-8083
- Vu Nguyen, Tam Le, Makoto Yamada, Michael A. Osborne:  
**Optimal Transport Kernels for Sequential and Parallel Neural Architecture Search.** 8084-8095
- Khanh Nguyen, Dipendra Misra, Robert E. Schapire, Miroslav Dudík, Patrick Shafto:  
**Interactive Learning from Activity Description.** 8096-8108
- Quan Nguyen, Arghavan Modiri, Roman Garnett:  
**Nonmyopic Multifidelity Active Search.** 8109-8118
- Quynh Nguyen, Marco Mondelli, Guido F. Montúfar:  
**Tight Bounds on the Smallest Eigenvalue of the Neural Tangent Kernel for Deep ReLU Networks.** 8119-8129
- Tung D. Nguyen, Rui Shu, Tuan Pham, Hung Bui, Stefano Ermon:  
**Temporal Predictive Coding For Model-Based Planning In Latent Space.** 8130-8139
- Dung Nguyen, Anil Vullikanti :  
**Differentially Private Densest Subgraph Detection.** 8140-8151
- Renkun Ni, Micah Goldblum, Amr Sharaf, Kezhi Kong, Tom Goldstein:  
**Data Augmentation for Meta-Learning.** 8152-8161
- Alexander Quinn Nichol, Prafulla Dhariwal:  
**Improved Denoising Diffusion Probabilistic Models.** 8162-8171
- Sloan Nietert, Ziv Goldfeld, Kengo Kato:  
**Smooth p-Wasserstein Distance: Structure, Empirical Approximation, and Statistical Applications.** 8172-8183
- Shuaicheng Niu, Jiaxiang Wu, Guanghui Xu, Yifan Zhang, Yong Guo, Peilin Zhao, Peng Wang, Mingkui Tan:  
**AdaXpert: Adapting Neural Architecture for Growing Data.** 8184-8194
- Kenta Niwa, Guoqiang Zhang, W. Bastiaan Kleijn, Noboru Harada, Hiroshi Sawada, Akinori Fujino:  
**Asynchronous Decentralized Optimization With Implicit Stochastic Variance Reduction.** 8195-8204
- Albert No, Taeho Yoon, Sehyun Kwon, Ernest K. Ryu:  
**WGAN with an Infinitely Wide Generator Has No Spurious Stationary Points.** 8205-8215
- Richard Nock, Stephen Hardy, Wilko Henecka, Hamish Ivey-Law, Jakub Nabaglo, Giorgio Patrini, Guillaume Smith, Brian Thorne:  
**The Impact of Record Linkage on Learning from Feature Partitioned Data.** 8216-8226
- Harsha Nori, Rich Caruana, Zhiqi Bu, Judy Hanwen Shen, Janardhan Kulkarni:  
**Accuracy, Interpretability, and Differential Privacy via Explainable Boosting.** 8227-8237
- Chris Nota, Philip S. Thomas, Bruno C. da Silva:  
**Posterior Value Functions: Hindsight Baselines for Policy Gradient Methods.** 8238-8247
- Sebastian W. Ober, Laurence Aitchison:  
**Global inducing point variational posteriors for Bayesian neural networks and deep Gaussian processes.** 8248-8259
- Michael Oberst, Nikolaj Thams, Jonas Peters, David A. Sontag:  
**Regularizing towards Causal Invariance: Linear Models with Proxies.** 8260-8270
- Min-hwan Oh, Garud Iyengar, Assaf Zeevi:  
**Sparsity-Agnostic Lasso Bandit.** 8271-8280



-     Alon Oring, Zohar Yakhini, Yacov Hel-Or:  
**Autoencoder Image Interpolation by Shaping the Latent Space.** 8281-8290
  -     Samet Oymak, Mingchen Li, Mahdi Soltanolkotabi :  
**Generalization Guarantees for Neural Architecture Search with Train-Validation Split.** 8291-8301
  -     Sherjil Ozair, Yazhe Li, Ali Razavi, Ioannis Antonoglou, Aäron van den Oord, Oriol Vinyals:  
**Vector Quantized Models for Planning.** 8302-8313
  -     Ozan Özdenizci, Robert Legenstein:  
**Training Adversarially Robust Sparse Networks via Bayesian Connectivity Sampling.** 8314-8324
  -     Avik Pal, Yingbo Ma, Viral B. Shah, Christopher Vincent Rackauckas:  
**Opening the Blackbox: Accelerating Neural Differential Equations by Regularizing Internal Solver Heuristics.** 8325-8335
  -     Soumyasundar Pal, Liheng Ma, Yingxue Zhang, Mark Coates:  
**RNN with Particle Flow for Probabilistic Spatio-temporal Forecasting.** 8336-8348
  -     Mengjie Pan, Tyler H. McCormick, Bailey K. Fosdick:  
**Inference for Network Regression Models with Community Structure.** 8349-8358
  -     Bo Pang, Ying Nian Wu:  
**Latent Space Energy-Based Model of Symbol-Vector Coupling for Text Generation and Classification.** 8359-8370
  -     Matteo Papini, Andrea Tirinzoni, Marcello Restelli, Alessandro Lazaric, Matteo Pirotta:  
**Leveraging Good Representations in Linear Contextual Bandits.** 8371-8380
  -     Sung Woo Park, Junseok Kwon:  
**Wasserstein Distributional Normalization For Robust Distributional Certification of Noisy Labeled Data.** 8381-8390
  -     Yookoon Park, Sangho Lee, Gunhee Kim, David M. Blei:  
**Unsupervised Representation Learning via Neural Activation Coding.** 8391-8400
  -     Junhyung Park, Uri Shalit, Bernhard Schölkopf, Krikamol Muandet:  
**Conditional Distributional Treatment Effect with Kernel Conditional Mean Embeddings and U-Statistic Regression.** 8401-8412
  -     Sung Woo Park, Dong Wook Shu, Junseok Kwon:  
**Generative Adversarial Networks for Markovian Temporal Dynamics: Stochastic Continuous Data Generation.** 8413-8421
  -     Axel Parmentier, Thibaut Vidal:  
**Optimal Counterfactual Explanations in Tree Ensembles.** 8422-8431
  -     Shreyas Malakarjun Patil, Constantine Dovrolis:  
**PHEW : Constructing Sparse Networks that Learn Fast and Generalize Well without Training Data.** 8432-8442
  -     Anselm Paulus, Michal Rolínek, Vít Musil, Brandon Amos, Georg Martius:  
**CombOptNet: Fit the Right NP-Hard Problem by Learning Integer Programming Constraints.** 8443-8453
  -     Oren Peer, Chen Tessler, Nadav Merlis, Ron Meir:  
**Ensemble Bootstrapping for Q-Learning.** 8454-8463
  -     Liangzu Peng, Boshi Wang, Manolis C. Tsakiris:  
**Homomorphic Sensing: Sparsity and Noise.** 8464-8475
  -     Dinglan Peng, Shuxin Zheng, Yatao Li, Guolin Ke, Di He, Tie-Yan Liu:  
**How could Neural Networks understand Programs?** 8476-8486
  -     Sikha Pentyala, Rafael Dowsley, Martine De Cock:  
**Privacy-Preserving Video Classification with Convolutional Neural Networks.** 8487-8499
  -     Ethan Perez, Douwe Kiela, Kyunghyun Cho:  
**Rissanen Data Analysis: Examining Dataset Characteristics via Description Length.** 8500-8513
  -     Nicolas Perez Nieves, Yaodong Yang, Oliver Slumbers, David Henry Mguni, Ying Wen, Jun Wang:  
**Modelling Behavioural Diversity for Learning in Open-Ended Games.** 8514-8524
  -     Julien Pérolat, Rémi Munos, Jean-Baptiste Lespiau, Shayegan Omidshafiei, Mark Rowland, Pedro A. Ortega, Neil Burch, Thomas W. Anthony, David Balduzzi, Bart De Vylder, Georgios Piliouras, Marc Lanctot, Karl Tuyls:  
**From Poincaré Recurrence to Convergence in Imperfect Information Games: Finding Equilibrium via Regularization.** 8525-8535
  -     Adeel Pervez, Efstratios Gavves:  
**Spectral Smoothing Unveils Phase Transitions in Hierarchical Variational Autoencoders.** 8536-8545
  -     Felix Petersen , Christian Borgelt, Hilde Kuehne, Oliver Deussen:  
**Differentiable Sorting Networks for Scalable Sorting and Ranking Supervision.** 8546-8555
  -     Aleksei Petrenko, Erik Wijmans, Brennan Shacklett, Vladlen Koltun:  
**Megaverse: Simulating Embodied Agents at One Million Experiences per Second.** 8556-8566
  -     My Phan, Philip S. Thomas, Erik G. Learned-Miller:  
**Towards Practical Mean Bounds for Small Samples.** 8567-8576



- Vincent Plassier, Maxime Vono, Alain Durmus, Eric Moulines:  
**DG-LMC: A Turn-key and Scalable Synchronous Distributed MCMC Algorithm via Langevin Monte Carlo within Gibbs.** 8577-8587
- Petra Poklukar, Anastasiia Varava, Danica Kragic:  
**GeomCA: Geometric Evaluation of Data Representations.** 8588-8598
- Vadim Popov, Ivan Vovk, Vladimir Gogoryan, Tasnima Sadekova, Mikhail A. Kudinov:  
**Grad-TTS: A Diffusion Probabilistic Model for Text-to-Speech.** 8599-8608
- Andres Potapczynski, Luhuan Wu, Dan Biderman, Geoff Pleiss, John P. Cunningham:  
**Bias-Free Scalable Gaussian Processes via Randomized Truncations.** 8609-8619
- Ilan Price, Jared Tanner:  
**Dense for the Price of Sparse: Improved Performance of Sparsely Initialized Networks via a Subspace Offset.** 8620-8629
- Weizhen Qi, Yeyun Gong, Jian Jiao, Yu Yan, Weizhu Chen, Dayiheng Liu, Kewen Tang, Houqiang Li, Jiusheng Chen, Ruofei Zhang, Ming Zhou, Nan Duan:  
**BANG: Bridging Autoregressive and Non-autoregressive Generation with Large Scale Pretraining.** 8630-8639
- Xin Qian, Diego Klabjan:  
**A Probabilistic Approach to Neural Network Pruning.** 8640-8649
- Kaizhi Qian, Yang Zhang, Shiyu Chang, Jinjun Xiong, Chuang Gan, David D. Cox, Mark Hasegawa-Johnson:  
**Global Prosody Style Transfer Without Text Transcriptions.** 8650-8660
- Yi-Ling Qiao, Junbang Liang, Vladlen Koltun, Ming C. Lin:  
**Efficient Differentiable Simulation of Articulated Bodies.** 8661-8671
- Gang Qiao, Weijie J. Su, Li Zhang:  
**Oneshot Differentially Private Top-k Selection.** 8672-8681
- Zengyi Qin, Yuxiao Chen, Chuchu Fan:  
**Density Constrained Reinforcement Learning.** 8682-8692
- Tian Qin, Tian-Zuo Wang, Zhi-Hua Zhou:  
**Budgeted Heterogeneous Treatment Effect Estimation.** 8693-8702
- Chen Qiu, Timo Pfrommer, Marius Kloft, Stephan Mandt, Maja Rudolph:  
**Neural Transformation Learning for Deep Anomaly Detection Beyond Images.** 8703-8714
- Shuang Qiu, Xiaohan Wei, Jieping Ye, Zhaoran Wang, Zhuoran Yang:  
**Provably Efficient Fictitious Play Policy Optimization for Zero-Sum Markov Games with Structured Transitions.** 8715-8725
- Zhaofan Qiu, Ting Yao, Chong-Wah Ngo, Tao Mei:  
**Optimization Planning for 3D ConvNets.** 8726-8736
- Shuang Qiu, Jieping Ye, Zhaoran Wang, Zhuoran Yang:  
**On Reward-Free RL with Kernel and Neural Function Approximations: Single-Agent MDP and Markov Game.** 8737-8747
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, Ilya Sutskever:  
**Learning Transferable Visual Models From Natural Language Supervision.** 8748-8763
- Jayaram Raghuram, Varun Chandrasekaran, Somesh Jha, Suman Banerjee:  
**A General Framework For Detecting Anomalous Inputs to DNN Classifiers.** 8764-8775
- Arrasy Rahman, Niklas Höpner, Filippos Christianos, Stefano V. Albrecht:  
**Towards Open Ad Hoc Teamwork Using Graph-based Policy Learning.** 8776-8786
- Roberta Raileanu, Rob Fergus:  
**Decoupling Value and Policy for Generalization in Reinforcement Learning.** 8787-8798
- Anand Rajagopalan, Fabio Vitale, Danny Vainstein, Gui Citovsky, Cecilia M. Procopiuc, Claudio Gentile:  
**Hierarchical Clustering of Data Streams: Scalable Algorithms and Approximation Guarantees.** 8799-8809
- Alain Rakotomamonjy, Liva Ralaivola:  
**Differentially Private Sliced Wasserstein Distance.** 8810-8820
- Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, Ilya Sutskever:  
**Zero-Shot Text-to-Image Generation.** 8821-8831
- Syama Sundar Rangapuram, Lucien D. Werner, Konstantinos Benidis, Pedro Mercado, Jan Gasthaus, Tim Januschowski:  
**End-to-End Learning of Coherent Probabilistic Forecasts for Hierarchical Time Series.** 8832-8843
- Roshan Rao, Jason Liu, Robert Verkuil, Joshua Meier, John F. Canny, Pieter Abbeel, Tom Sercu, Alexander Rives:  
**MSA Transformer.** 8844-8856
- Kashif Rasul, Calvin Seward, Ingmar Schuster, Roland Vollgraf:  
**Autoregressive Denoising Diffusion Models for Multivariate Probabilistic Time Series Forecasting.** 8857-8868
- Neale Ratzlaff, Qinxun Bai, Fuxin Li, Wei Xu:  
**Generative Particle Variational Inference via Estimation of Functional Gradients.** 8869-8879



- Netanel Raviv, Aidan Kelley, Minzhe Guo, Yevgeniy Vorobeychik:  
**Enhancing Robustness of Neural Networks through Fourier Stabilization.** 8880-8889
- Ankit Singh Rawat, Aditya Krishna Menon, Wittawat Jitkrittum, Sadeep Jayasumana, Felix X. Yu, Sashank J. Reddi, Sanjiv Kumar:  
**Disentangling Sampling and Labeling Bias for Learning in Large-output Spaces.** 8890-8901
- Dripta S. Raychaudhuri, Sujoy Paul, Jeroen van Baar, Amit K. Roy-Chowdhury:  
**Cross-domain Imitation from Observations.** 8902-8912
- Noam Razin, Asaf Maman, Nadav Cohen:  
**Implicit Regularization in Tensor Factorization.** 8913-8924
- Maria Refinetti, Stéphane d'Ascoli, Ruben Ohana, Sebastian Goldt :  
**Align, then memorise: the dynamics of learning with feedback alignment.** 8925-8935
- Maria Refinetti, Sebastian Goldt , Florent Krzakala , Lenka Zdeborová:  
**Classifying high-dimensional Gaussian mixtures: Where kernel methods fail and neural networks succeed.** 8936-8947
- Konstantinos Rematas, Ricardo Martin-Brualla, Vittorio Ferrari:  
**Sharf: Shape-conditioned Radiance Fields from a Single View.** 8948-8958
- Hongyu Ren, Hanjun Dai, Bo Dai, Xinyun Chen, Michihiro Yasunaga, Haitian Sun, Dale Schuurmans, Jure Leskovec, Denny Zhou:  
**LEGO: Latent Execution-Guided Reasoning for Multi-Hop Question Answering on Knowledge Graphs.** 8959-8970
- Jie Ren , Mingjie Li, Zexu Liu, Quanshi Zhang:  
**Interpreting and Disentangling Feature Components of Various Complexity from DNNs.** 8971-8981
- Jiaxiang Ren, Zijie Zhang , Jiayin Jin, Xin Zhao, Sixing Wu, Yang Zhou, Yelong Shen, Tianshi Che, Ruoming Jin, Dejing Dou:  
**Integrated Defense for Resilient Graph Matching.** 8982-8997
- Lorenz Richter, Leon Sallandt, Nikolas Nüsken:  
**Solving high-dimensional parabolic PDEs using the tensor train format.** 8998-9009
- Geovani Rizk, Albert Thomas, Igor Colin, Rida Laraki, Yann Chevaleyre:  
**Best Arm Identification in Graphical Bilinear Bandits.** 9010-9019
- T. Mitchell Roddenberry, Nicholas Glaze, Santiago Segarra:  
**Principled Simplicial Neural Networks for Trajectory Prediction.** 9020-9029
- Geoffrey Roeder, Luke Metz, Durk Kingma:  
**On Linear Identifiability of Learned Representations.** 9030-9039
- Esther Rolf, Theodora T. Worledge, Benjamin Recht, Michael I. Jordan:  
**Representation Matters: Assessing the Importance of Subgroup Allocations in Training Data.** 9040-9051
- Clément Romac, Rémy Portelas, Katja Hofmann, Pierre-Yves Oudeyer:  
**TeachMyAgent: a Benchmark for Automatic Curriculum Learning in Deep RL.** 9052-9063
- Mihaela Rosca, Yan Wu, Benoit Dherin, David Barrett:  
**Discretization Drift in Two-Player Games.** 9064-9074
- Jonathan S. Rosenfeld, Jonathan Frankle, Michael Carbin, Nir Shavit:  
**On the Predictability of Pruning Across Scales.** 9075-9083
- Andrew Slavin Ross, Finale Doshi-Velez:  
**Benchmarks, Algorithms, and Metrics for Hierarchical Disentanglement.** 9084-9094
- Karsten Roth, Timo Milbich, Björn Ommer, Joseph Paul Cohen, Marzyeh Ghassemi:  
**Simultaneous Similarity-based Self-Distillation for Deep Metric Learning.** 9095-9106
- Guy N. Rothblum, Gal Yona:  
**Multi-group Agnostic PAC Learnability.** 9107-9115
- Jonas Rothfuss, Vincent Fortuin, Martin Josifoski, Andreas Krause:  
**PACOH: Bayes-Optimal Meta-Learning with PAC-Guarantees.** 9116-9126
- Chloé Rouyer, Yevgeny Seldin, Nicolò Cesa-Bianchi:  
**An Algorithm for Stochastic and Adversarial Bandits with Switching Costs.** 9127-9135
- Yangjun Ruan, Karen Ullrich, Daniel Severo, James Townsend, Ashish Khisti, Arnaud Doucet, Alireza Makhzani, Chris J. Maddison:  
**Improving Lossless Compression Rates via Monte Carlo Bits-Back Coding.** 9136-9147
- Tim G. J. Rudner, Oscar Key, Yarin Gal, Tom Rainforth:  
**On Signal-to-Noise Ratio Issues in Variational Inference for Deep Gaussian Processes.** 9148-9156
- Miguel Ruiz-Garcia, Ge Zhang, Samuel S. Schoenholz, Andrea J. Liu:  
**Tilting the playing field: Dynamical loss functions for machine learning.** 9157-9167
- T. Konstantin Rusch, Siddhartha Mishra:  
**UnICORNN: A recurrent model for learning very long time dependencies.** 9168-9178
- Oleh Rybkin, Kostas Daniilidis, Sergey Levine:  
**Simple and Effective VAE Training with Calibrated Decoders.** 9179-9189



- Oleh Rybkin, Chuning Zhu, Anusha Nagabandi, Kostas Daniilidis, Igor Mordatch, Sergey Levine:  
**Model-Based Reinforcement Learning via Latent-Space Collocation.** 9190-9201
- Durga Sivasubramanian, Rishabh K. Iyer, Ganesh Ramakrishnan, Abir De:  
**Training Data Subset Selection for Regression with Controlled Generalization Error.** 9202-9212
- Sara Sabour, Andrea Tagliasacchi, Soroosh Yazdani, Geoffrey E. Hinton, David J. Fleet:  
**Unsupervised Part Representation by Flow Capsules.** 9213-9223
- Mher Safaryan, Peter Richtárik:  
**Stochastic Sign Descent Methods: New Algorithms and Better Theory.** 9224-9234
- Aadirupa Saha, Tomer Koren, Yishay Mansour:  
**Adversarial Dueling Bandits.** 9235-9244
- Aadirupa Saha, Tomer Koren, Yishay Mansour:  
**Dueling Convex Optimization.** 9245-9254
- Aadirupa Saha, Nagarajan Natarajan, Praneeth Netrapalli, Prateek Jain:  
**Optimal regret algorithm for Pseudo-1d Bandit Convex Optimization.** 9255-9264
- Mojtaba Sahraee-Ardakan, Tung Mai, Anup B. Rao, Ryan A. Rossi, Sundeep Rangan, Alyson K. Fletcher:  
**Asymptotics of Ridge Regression in Convolutional Models.** 9265-9275
- Michael E. Sander, Pierre Albin, Mathieu Blondel, Gabriel Peyré:  
**Momentum Residual Neural Networks.** 9276-9287
- Mark Sandler, Max Vladmyrov, Andrey Zhmoginov, Nolan Miller, Tom Madams, Andrew Jackson, Blaise Agüera y Arcas:  
**Meta-Learning Bidirectional Update Rules.** 9288-9300
- Elad Sarafian, Shai Keynan, Sarit Kraus:  
**Recomposing the Reinforcement Learning Building Blocks with Hypernetworks.** 9301-9312
- Roei Sarussi, Alon Brutzkus, Amir Globerson:  
**Towards Understanding Learning in Neural Networks with Linear Teachers.** 9313-9322
- Victor Garcia Satorras, Emiel Hoogeboom, Max Welling:  
**E(n) Equivariant Graph Neural Networks.** 9323-9332
- Nikunj Saunshi, Arushi Gupta, Wei Hu:  
**A Representation Learning Perspective on the Importance of Train-Validation Splitting in Meta-Learning.** 9333-9343
- Meyer Scetbon, Marco Cuturi, Gabriel Peyré:  
**Low-Rank Sinkhorn Factorization.** 9344-9354
- Imanol Schlag, Kazuki Irie, Jürgen Schmidhuber:  
**Linear Transformers Are Secretly Fast Weight Programmers.** 9355-9366
- Robin M. Schmidt, Frank Schneider, Philipp Hennig:  
**Descending through a Crowded Valley - Benchmarking Deep Learning Optimizers.** 9367-9376
- Kristof Schütt, Oliver T. Unke, Michael Gastegger:  
**Equivariant message passing for the prediction of tensorial properties and molecular spectra.** 9377-9388
- Avi Schwarzschild, Micah Goldblum, Arjun Gupta, John P. Dickerson, Tom Goldstein:  
**Just How Toxic is Data Poisoning? A Unified Benchmark for Backdoor and Data Poisoning Attacks.** 9389-9398
- Damien Scieur, Youngsung Kim:  
**Connecting Sphere Manifolds Hierarchically for Regularization.** 9399-9409
- Jenny Denise Seidenschwarz, Ismail Elezi, Laura Leal-Taixé:  
**Learning Intra-Batch Connections for Deep Metric Learning.** 9410-9421
- Rajat Sen, Alexander Rakhlin, Lexing Ying, Rahul Kidambi, Dean P. Foster, Daniel N. Hill, Inderjit S. Dhillon:  
**Top-k eXtreme Contextual Bandits with Arm Hierarchy.** 9422-9433
- Flore Sentenac, Jialin Yi, Clément Calauzènes, Vianney Perchet, Milan Vojnovic:  
**Pure Exploration and Regret Minimization in Matching Bandits.** 9434-9442
- Younggyo Seo, Lili Chen, Jinwoo Shin, Honglak Lee, Pieter Abbeel, Kimin Lee:  
**State Entropy Maximization with Random Encoders for Efficient Exploration.** 9443-9454
- Pier Giuseppe Sessa, Ilija Bogunovic, Andreas Krause, Maryam Kamgarpour:  
**Online Submodular Resource Allocation with Applications to Rebalancing Shared Mobility Systems.** 9455-9464
- Rutav M. Shah, Vikash Kumar:  
**RRL: Resnet as representation for Reinforcement Learning.** 9465-9476
- Mehran Shakerinava, Siamak Ravanbakhsh:  
**Equivariant Networks for Pixelized Spheres.** 9477-9488
- Aviv Shamsian, Aviv Navon, Ethan Fetaya, Gal Chechik:  
**Personalized Federated Learning using Hypernetworks.** 9489-9502
- Jie Shen:  
**On the Power of Localized Perceptron for Label-Optimal Learning of Halfspaces with Adversarial Noise.** 9503-9514



- Jie Shen:  
**Sample-Optimal PAC Learning of Halfspaces with Malicious Noise.** 9515-9524
- Guangyu Shen, Yingqi Liu, Guanhong Tao, Shengwei An, Qiuling Xu, Siyuan Cheng, Shiqing Ma, Xiangyu Zhang:  
**Backdoor Scanning for Deep Neural Networks through K-Arm Optimization.** 9525-9536
- Simon P. Shen, Yecheng Jason Ma, Omer Gottesman, Finale Doshi-Velez:  
**State Relevance for Off-Policy Evaluation.** 9537-9546
- Han Shi, Jiahui Gao, Xiaozhe Ren, Hang Xu, Xiaodan Liang, Zhenguo Li, James Tin-Yau Kwok:  
**SparseBERT: Rethinking the Importance Analysis in Self-attention.** 9547-9557
- Chence Shi, Shitong Luo, Minkai Xu, Jian Tang:  
**Learning Gradient Fields for Molecular Conformation Generation.** 9558-9568
- Ruian Shi, Quaid Morris:  
**Segmenting Hybrid Trajectories using Latent ODEs.** 9569-9579
- Chengchun Shi, Runzhe Wan, Victor Chernozhukov, Rui Song:  
**Deeply-Debiased Off-Policy Interval Estimation.** 9580-9591
- Sheng-Min Shih, Pin-Ju Tien, Zohar Karnin:  
**GANMEX: One-vs-One Attributions using GAN-based Model Explainability.** 9592-9602
- Jaewoong Shin, Haebeom Lee, Boqing Gong, Sung Ju Hwang:  
**Large-Scale Meta-Learning with Continual Trajectory Shifting.** 9603-9613
- Tianmin Shu, Abhishek Bhandwaldar, Chuang Gan, Kevin A. Smith, Shari Liu, Dan Gutfreund, Elizabeth S. Spelke, Joshua B. Tenenbaum, Tomer D. Ullman:  
**AGENT: A Benchmark for Core Psychological Reasoning.** 9614-9625
- Yang Shu, Zhi Kou, Zhangjie Cao, Jianmin Wang , Mingsheng Long :  
**Zoo-Tuning: Adaptive Transfer from A Zoo of Models.** 9626-9637
- Changjian Shui, Zijian Li, Jiaqi Li , Christian Gagné, Charles X. Ling, Boyu Wang:  
**Aggregating From Multiple Target-Shifted Sources.** 9638-9648
- Nian Si, Karthyek Murthy, Jose H. Blanchet, Viet Anh Nguyen:  
**Testing Group Fairness via Optimal Transport Projections.** 9649-9659
- Sahil Sidhekh, Aroof Aimen, Narayanan C. Krishnan:  
**On Characterizing GAN Convergence Through Proximal Duality Gap.** 9660-9670
- Houssem Sifaou, Abla Kammoun, Mohamed-Slim Alouini:  
**A Precise Performance Analysis of Support Vector Regression.** 9671-9680
- Aaron Sim, Maciej Wiatrak, Angus Brayne, Páidí Creed, Saee Paliwal:  
**Directed Graph Embeddings in Pseudo-Riemannian Manifolds.** 9681-9690
- Rachael Hwee Ling Sim, Yehong Zhang, Bryan Kian Hsiang Low, Patrick Jaillet:  
**Collaborative Bayesian Optimization with Fair Regret.** 9691-9701
- David Simchi-Levi, Zeyu Zheng, Feng Zhu:  
**Dynamic Planning and Learning under Recovering Rewards.** 9702-9711
- Carl-Johann Simon-Gabriel, Noman Ahmed Sheikh, Andreas Krause :  
**PopSkipJump: Decision-Based Attack for Probabilistic Classifiers.** 9712-9721
- Berfin Simsek, François Ged, Arthur Jacot, Francesco Spadaro, Clément Hongler, Wulfram Gerstner, Johann Brea:  
**Geometry of the Loss Landscape in Overparameterized Neural Networks: Symmetries and Invariances.** 9722-9732
- Raghav Singal, George Michailidis, Hoiyi Ng:  
**Flow-based Attribution in Graphical Models: A Recursive Shapley Approach.** 9733-9743
- Gautam Singh, Skand Vishwanath Peri, Junghyun Kim, Hyunseok Kim, Sungjin Ahn:  
**Structured World Belief for Reinforcement Learning in POMDP.** 9744-9755
- Sahil Singla, Soheil Feizi:  
**Skew Orthogonal Convolutions.** 9756-9766
- Shagun Sodhani, Amy Zhang, Joelle Pineau:  
**Multi-Task Reinforcement Learning with Context-based Representations.** 9767-9779
- Sungryull Sohn, Sungtae Lee, Jongwook Choi, Harm van Seijen, Mehdi Fatemi, Honglak Lee:  
**Shortest-Path Constrained Reinforcement Learning for Sparse Reward Tasks.** 9780-9790
- Yang Song, Chenlin Meng, Renjie Liao, Stefano Ermon:  
**Accelerating Feedforward Computation via Parallel Nonlinear Equation Solving.** 9791-9800
- Yuda Song, Wen Sun:  
**PC-MLP: Model-based Reinforcement Learning with Policy Cover Guided Exploration.** 9801-9811
- Zhao Song, David P. Woodruff, Zheng Yu, Lichen Zhang:  
**Fast Sketching of Polynomial Kernels of Polynomial Degree.** 9812-9823
- Chaobing Song, Stephen J. Wright, Jelena Diakonikolas:  
**Variance Reduction via Primal-Dual Accelerated Dual Averaging for Nonsmooth Convex Finite-Sums.** 9824-9834



- Zhao Song, Zheng Yu:  
**Oblivious Sketching-based Central Path Method for Linear Programming.** 9835-9847
- Sumedh A. Sontakke, Arash Mehrjou, Laurent Itti, Bernhard Schölkopf:  
**Causal Curiosity: RL Agents Discovering Self-supervised Experiments for Causal Representation Learning.** 9848-9858
- Alessandro Sordoni, Nouha Dziri, Hannes Schulz, Geoffrey J. Gordon, Philip Bachman, Remi Tachet des Combes:  
**Decomposed Mutual Information Estimation for Contrastive Representation Learning.** 9859-9869
- Adam Stooke, Kimin Lee, Pieter Abbeel, Michael Laskin:  
**Decoupling Representation Learning from Reinforcement Learning.** 9870-9879
- Xiu Su, Shan You, Mingkai Zheng, Fei Wang, Chen Qian, Changshui Zhang, Chang Xu:  
**K-shot NAS: Learnable Weight-Sharing for NAS with K-shot Supernets.** 9880-9890
- Kazuya Sugiyama, Vo Nguyen Le Duy, Ichiro Takeuchi:  
**More Powerful and General Selective Inference for Stepwise Feature Selection using Homotopy Method.** 9891-9901
- Sainbayar Sukhbaatar, Da Ju, Spencer Poff, Stephen Roller, Arthur Szlam, Jason Weston, Angela Fan:  
**Not All Memories are Created Equal: Learning to Forget by Expiring.** 9902-9912
- Cecilia Summers, Michael J. Dinneen:  
**Nondeterminism and Instability in Neural Network Optimization.** 9913-9922
- Ming Sun, Haoxuan Dou, Baopu Li, Junjie Yan, Wanli Ouyang, Lei Cui:  
**AutoSampling: Search for Effective Data Sampling Schedules.** 9923-9933
- Peize Sun, Yi Jiang, Enze Xie, Wenqi Shao, Zehuan Yuan, Changhu Wang, Ping Luo:  
**What Makes for End-to-End Object Detection?** 9934-9944
- Wei-Fang Sun, Cheng-Kuang Lee, Chun-Yi Lee:  
**DFAC Framework: Factorizing the Value Function via Quantile Mixture for Multi-Agent Distributional Q-Learning.** 9945-9954
- Shengyang Sun, Jiaxin Shi, Andrew Gordon Wilson, Roger B. Grosse:  
**Scalable Variational Gaussian Processes via Harmonic Kernel Decomposition.** 9955-9965
- Haitian Sun, Patrick Verga, Bhuwan Dhingra, Ruslan Salakhutdinov, William W. Cohen:  
**Reasoning Over Virtual Knowledge Bases With Open Predicate Relations.** 9966-9977
- Ravi Sundaram, Anil Vullikanti, Haifeng Xu, Fan Yao:  
**PAC-Learning for Strategic Classification.** 9978-9988
- Wesley Suttle, Kaiqing Zhang, Zhuoran Yang, Ji Liu, David N. Kraemer:  
**Reinforcement Learning for Cost-Aware Markov Decision Processes.** 9989-9999
- Fnu Suya, Saeed Mahloujifar, Anshuman Suri, David Evans, Yuan Tian:  
**Model-Targeted Poisoning Attacks with Provable Convergence.** 10000-10010
- Atsushi Suzuki, Atsushi Nitanda, Jing Wang, Linchuan Xu, Kenji Yamanishi, Marc Cavazza:  
**Generalization Error Bound for Hyperbolic Ordinal Embedding.** 10011-10021
- Gokul Swamy, Sanjiban Choudhury, J. Andrew Bagnell, Steven Wu:  
**Of Moments and Matching: A Game-Theoretic Framework for Closing the Imitation Gap.** 10022-10032
- Saifuddin Syed, Vittorio Romaniello, Trevor Campbell, Alexandre Bouchard-Côté:  
**Parallel tempering on optimized paths.** 10033-10042
- Saeid Asgari Taghanaki, Kristy Choi, Amir Hosein Khasahmadi, Anirudh Goyal:  
**Robust Representation Learning via Perceptual Similarity Metrics.** 10043-10053
- Ashraf Tahmasbi, Ellango Jothimurugesan, Srikantha Tirthapura, Phillip B. Gibbons:  
**DriftSurf: Stable-State / Reactive-State Learning under Concept Drift.** 10054-10064
- Kai Sheng Tai, Peter Bailis, Gregory Valiant:  
**Sinkhorn Label Allocation: Semi-Supervised Classification via Annealed Self-Training.** 10065-10075
- Sho Takemori, Masahiro Sato:  
**Approximation Theory Based Methods for RKHS Bandits.** 10076-10085
- Yuki Takezawa, Ryoma Sato, Makoto Yamada:  
**Supervised Tree-Wasserstein Distance.** 10086-10095
- Mingxing Tan, Quoc V. Le:  
**EfficientNetV2: Smaller Models and Faster Training.** 10096-10106
- Anshoo Tandon, Aldric H. J. Han, Vincent Y. F. Tan:  
**SGA: A Robust Algorithm for Partial Recovery of Tree-Structured Graphical Models with Noisy Samples.** 10107-10117
- Hanlin Tang, Shaoduo Gan, Ammar Ahmad Awan, Samyam Rajbhandari, Conglong Li, Xiangru Lian, Ji Liu, Ce Zhang, Yuxiong He:  
**1-bit Adam: Communication Efficient Large-Scale Training with Adam's Convergence Speed.** 10118-10129
- Yunhao Tang, Mark Rowland, Rémi Munos, Michal Valko:  
**Taylor Expansion of Discount Factors.** 10130-10140



- Yunzhe Tao, Sahika Genc, Jonathan Chung, Tao Sun, Sunil Mallya:  
**REPAINT: Knowledge Transfer in Deep Reinforcement Learning.** 10141-10152
- Salma Tarmoun, Guilherme Fran  a, Benjamin D. Haeffele, Ren   Vidal:  
**Understanding the Dynamics of Gradient Flow in Overparameterized Linear models.** 10153-10161
- Bahar Taskesen, Man-Chung Yue, Jose H. Blanchet, Daniel Kuhn, Viet Anh Nguyen:  
**Sequential Domain Adaptation by Synthesizing Distributionally Robust Experts.** 10162-10172
- Zenna Tavares, James Koppel, Xin Zhang, Ria Das, Armando Solar-Lezama :  
**A Language for Counterfactual Generative Models.** 10173-10182
- Yi Tay, Dara Bahri, Donald Metzler, Da-Cheng Juan, Zhe Zhao, Che Zheng:  
**Synthesizer: Rethinking Self-Attention for Transformer Models.** 10183-10192
- Yi Tay, Mostafa Dehghani, Vamsi Aribandi, Jai Prakash Gupta, Philip Pham, Zhen Qin, Dara Bahri, Da-Cheng Juan, Donald Metzler:  
**OmniNet: Omnidirectional Representations from Transformers.** 10193-10202
- Jiaye Teng, Zeren Tan, Yang Yuan:  
**T-SCI: A Two-Stage Conformal Inference Algorithm with Guaranteed Coverage for Cox-MLP.** 10203-10213
- D  vid Terj  k:  
**Moreau-Yosida f-divergences.** 10214-10224
- Piotr Teterwak, Chiyuan Zhang, Dilip Krishnan, Michael C. Mozer:  
**Understanding Invariance via Feedforward Inversion of Discriminatively Trained Classifiers.** 10225-10235
- Brijen Thananjeyan, Kirthevasan Kandasamy, Ion Stoica, Michael I. Jordan, Ken Goldberg, Joseph Gonzalez:  
**Resource Allocation in Multi-armed Bandit Exploration: Overcoming Sublinear Scaling with Adaptive Parallelism.** 10236-10246
- Achille Thin, Nikita Kotelevskii, Arnaud Doucet, Alain Durmus, Eric Moulines, Maxim Panov:  
**Monte Carlo Variational Auto-Encoders.** 10247-10257
- Christian B. Thygesen, Christian Skj  dt Steenmans, Ahmad Salim Al-Sibahi, Lys Sanz Moreta, Anders Bundg  rd S  rensen, Thomas Hamelryck:  
**Efficient Generative Modelling of Protein Structure Fragments using a Deep Markov Model.** 10258-10267
- Yuandong Tian, Xinlei Chen, Surya Ganguli:  
**Understanding self-supervised learning dynamics without contrastive pairs.** 10268-10278
- Yi Tian, Yuanhao Wang, Tiancheng Yu, Suvrit Sra:  
**Online Learning in Unknown Markov Games.** 10279-10288
- Louis C. Tiao, Aaron Klein, Matthias W. Seeger, Edwin V. Bonilla, C  dric Archambeau, Fabio Ramos:  
**BORE: Bayesian Optimization by Density-Ratio Estimation.** 10289-10300
- Conor Tillinghast, Shandian Zhe:  
**Nonparametric Decomposition of Sparse Tensors.** 10301-10311
- David Tolpin, Yuan Zhou, Tom Rainforth, Hongseok Yang:  
**Probabilistic Programs with Stochastic Conditioning.** 10312-10323
- Nergis Tomen, Silvia-Laura Pintea, Jan van Gemert:  
**Deep Continuous Networks.** 10324-10335
- Alexander Tong, Guillaume Huguet, Amine Natik, Kincaid MacDonald, Manik Kuchroo, Ronald R. Coifman, Guy Wolf, Smita Krishnaswamy:  
**Diffusion Earth Mover's Distance and Distribution Embeddings.** 10336-10346
- Hugo Touvron, Matthieu Cord, Matthijs Douze, Francisco Massa, Alexandre Sablayrolles, Herv   J  gou:  
**Training data-efficient image transformers & distillation through attention.** 10347-10357
- Brandon Trabucco, Aviral Kumar, Xinyang Geng, Sergey Levine:  
**Conservative Objective Models for Effective Offline Model-Based Optimization.** 10358-10368
- Gia-Lac Tran, Dimitrios Milios, Pietro Michiardi, Maurizio Filippone:  
**Sparse within Sparse Gaussian Processes using Neighbor Information.** 10369-10378
- Trang H. Tran, Lam M. Nguyen, Quoc Tran-Dinh:  
**SMG: A Shuffling Gradient-Based Method with Momentum.** 10379-10389
- Hung Tran-The, Sunil Gupta, Santu Rana, Svetha Venkatesh :  
**Bayesian Optimistic Optimisation with Exponentially Decaying Regret.** 10390-10400
- Frederik Tr  uble, Elliot Creager, Niki Kilbertus, Francesco Locatello, Andrea Dittadi, Anirudh Goyal, Bernhard Sch  lkopf, Stefan Bauer:  
**On Disentangled Representations Learned from Correlated Data.** 10401-10412
- Johannes Treutlein, Michael Dennis, Caspar Oesterheld, Jakob N. Foerster:  
**A New Formalism, Method and Open Issues for Zero-Shot Coordination.** 10413-10423
- Eleni Triantafillou, Hugo Larochelle, Richard S. Zemel, Vincent Dumoulin:  
**Learning a Universal Template for Few-shot Dataset Generalization.** 10424-10433



- Nilesh Tripuraneni, Chi Jin, Michael I. Jordan:  
**Provable Meta-Learning of Linear Representations.** 10434-10443
- William Trouleau, Jalal Etesami, Matthias Grossglauser, Negar Kiyavash, Patrick Thiran:  
**Cumulants of Hawkes Processes are Robust to Observation Noise.** 10444-10454
- Shubham Tulsiani, Abhinav Gupta:  
**PixelTransformer: Sample Conditioned Signal Generation.** 10455-10464
- Jalaj Upadhyay, Sarvagya Upadhyay:  
**A Framework for Private Matrix Analysis in Sliding Window Model.** 10465-10475
- Ilnura Usmanova, Maryam Kamgarpour, Andreas Krause, Kfir Y. Levy:  
**Fast Projection Onto Convex Smooth Constraints.** 10476-10486
- Aleksei Ustimenko, Liudmila Prokhorenko:  
**SGLB: Stochastic Gradient Langevin Boosting.** 10487-10496
- Pashootan Vaezipoor, Andrew C. Li, Rodrigo Toro Icarte, Sheila A. McIlraith:  
**LTL2Action: Generalizing LTL Instructions for Multi-Task RL.** 10497-10508
- Hans Van Gorp, Iris A. M. Huijben, Bastiaan S. Veeling, Nicola Pezzotti, Ruud J. G. van Sloun:  
**Active Deep Probabilistic Subsampling.** 10509-10518
- Ramakrishna Vedantam, Arthur Szlam, Maximilian Nickel, Ari Morcos, Brenden M. Lake:  
**CURI: A Benchmark for Productive Concept Learning Under Uncertainty.** 10519-10529
- Vikas Verma, Thang Luong, Kenji Kawaguchi, Hieu Pham, Quoc V. Le:  
**Towards Domain-Agnostic Contrastive Learning.** 10530-10541
- Sagar Verma, Jean-Christophe Pesquet:  
**Sparsifying Networks via Subdifferential Inclusion.** 10542-10552
- Paul Vicol, Luke Metz, Jascha Sohl-Dickstein:  
**Unbiased Gradient Estimation in Unrolled Computation Graphs with Persistent Evolution Strategies.** 10553-10563
- Cédric Vincent-Cuaz, Titouan Vayer, Rémi Flamary, Marco Corneli, Nicolas Courty:  
**Online Graph Dictionary Learning.** 10564-10574
- Marin Vlastelica P., Michal Rolínek, Georg Martius:  
**Neuro-algorithmic Policies Enable Fast Combinatorial Generalization.** 10575-10585
- Daniël Vos, Sicco Verwer:  
**Efficient Training of Robust Decision Trees Against Adversarial Examples.** 10586-10595
- Andrey Voynov, Stanislav Morozov, Artem Babenko:  
**Object Segmentation Without Labels with Large-Scale Generative Models.** 10596-10606
- Robbie Vreugdenhil, Viet Anh Nguyen, Armin Eftekhari, Peyman Mohajerin Esfahani:  
**Principal Component Hierarchy for Sparse Quadratic Programs.** 10607-10616
- Neha S. Wadia, Daniel Duckworth, Samuel S. Schoenholz, Ethan Dyer, Jascha Sohl-Dickstein:  
**Whitening and Second Order Optimization Both Make Information in the Dataset Unusable During Training, and Can Reduce or Prevent Generalization.** 10617-10629
- Nolan Wagener, Byron Boots, Ching-An Cheng:  
**Safe Reinforcement Learning Using Advantage-Based Intervention.** 10630-10640
- Andrew J. Wagenmaker, Max Simchowitz, Kevin G. Jamieson:  
**Task-Optimal Exploration in Linear Dynamical Systems.** 10641-10652
- Yi Wan, Abhishek Naik, Richard S. Sutton:  
**Learning and Planning in Average-Reward Markov Decision Processes.** 10653-10662
- Xingchen Wan, Vu Nguyen, Huong Ha, Bin Xin Ru, Cong Lu, Michael A. Osborne:  
**Think Global and Act Local: Bayesian Optimisation over High-Dimensional Categorical and Mixed Search Spaces.** 10663-10674
- Zi Wang:  
**Zero-Shot Knowledge Distillation from a Decision-Based Black-Box Model.** 10675-10685
- Lequn Wang, Yiwei Bai, Wen Sun, Thorsten Joachims:  
**Fairness of Exposure in Stochastic Bandits.** 10686-10696
- Yixin Wang, David M. Blei:  
**A Proxy Variable View of Shared Confounding.** 10697-10707
- Jiali Wang, He Chen, Rujun Jiang, Xudong Li, Zihao Li:  
**Fast Algorithms for Stackelberg Prediction Game with Least Squares Loss.** 10708-10716
- Wenxiao Wang, Minghao Chen, Shuai Zhao, Long Chen, Jinming Hu, Haifeng Liu, Deng Cai, Xiaofei He, Wei Liu:  
**Accelerate CNNs from Three Dimensions: A Comprehensive Pruning Framework.** 10717-10726
- Xin Wang, Shuyi Fan, Kun Kuang, Wenwu Zhu:  
**Explainable Automated Graph Representation Learning with Hyperparameter Importance.** 10727-10737
- Ximei Wang, Jinghan Gao, Mingsheng Long, Jianmin Wang:  
**Self-Tuning for Data-Efficient Deep Learning.** 10738-10748



- Jing Wang, Xin Geng:  
**Label Distribution Learning Machine.** 10749-10759
- Dilin Wang, Chengyue Gong, Meng Li, Qiang Liu, Vikas Chandra:  
**AlphaNet: Improved Training of Supernet with Alpha-Divergence.** 10760-10771
- Weichen Wang, Jiequn Han, Zhuoran Yang, Zhaoran Wang:  
**Global Convergence of Policy Gradient for Linear-Quadratic Mean-Field Control/Game in Continuous Time.** 10772-10782
- Yu Wang, Alfred Olivier Hero:  
**SG-PALM: a Fast Physically Interpretable Tensor Graphical Model.** 10783-10793
- Gefei Wang, Yuling Jiao, Qian Xu, Yang Wang, Can Yang:  
**Deep Generative Learning via Schrödinger Bridge.** 10794-10804
- Y. Samuel Wang, Si Kai Lee, Panos Toulis, Mladen Kolar :  
**Robust Inference for High-Dimensional Linear Models via Residual Randomization.** 10805-10815
- Jun-Kun Wang, Chi-Heng Lin, Jacob D. Abernethy:  
**A Modular Analysis of Provable Acceleration via Polyak's Momentum: Training a Wide ReLU Network and a Deep Linear Network.** 10816-10827
- Peng Wang, Huikang Liu, Zirui Zhou, Anthony Man-Cho So:  
**Optimal Non-Convex Exact Recovery in Stochastic Block Model via Projected Power Method.** 10828-10838
- Mengfan Wang, Boyu Lyu, Guoqiang Yu:  
**ConvexVST: A Convex Optimization Approach to Variance-stabilizing Transformation.** 10839-10848
- Bohan Wang, Qi Meng, Wei Chen, Tie-Yan Liu:  
**The Implicit Bias for Adaptive Optimization Algorithms on Homogeneous Neural Networks.** 10849-10858
- Yunjuan Wang, Poorya Mianjy, Raman Arora:  
**Robust Learning for Data Poisoning Attacks.** 10859-10869
- Alexander Wang, Mengye Ren, Richard S. Zemel:  
**SketchEmbedNet: Learning Novel Concepts by Imitating Drawings.** 10870-10881
- Angelina Wang, Olga Russakovsky:  
**Directional Bias Amplification.** 10882-10893
- Yutong Wang, Clayton Scott:  
**An exact solver for the Weston-Watkins SVM subproblem.** 10894-10904
- Xiangjun Wang, Junxiao Song, Penghui Qi, Peng Peng, Zhenkun Tang, Wei Zhang, Weimin Li, Xiongjun Pi, Jujie He, Chao Gao, Haitao Long, Quan Yuan:  
**SCC: an efficient deep reinforcement learning agent mastering the game of StarCraft II.** 10905-10915
- Daochen Wang, Aarthi Sundaram, Robin Kothari, Ashish Kapoor, Martin Roetteler:  
**Quantum algorithms for reinforcement learning with a generative model.** 10916-10926
- Jiayi Wang, Raymond K. W. Wong, Xiaojun Mao, Kwun Chuen Gary Chan:  
**Matrix Completion with Model-free Weighting.** 10927-10936
- Chengyi Wang, Yu Wu, Yao Qian, Ken'ichi Kumatani, Shujie Liu, Furu Wei, Michael Zeng, Xuedong Huang:  
**UniSpeech: Unified Speech Representation Learning with Labeled and Unlabeled Data.** 10937-10947
- Ruosong Wang, Yifan Wu, Ruslan Salakhutdinov, Sham M. Kakade:  
**Instabilities of Offline RL with Pre-Trained Neural Representation.** 10948-10960
- Yunke Wang, Chang Xu, Bo Du, Honglak Lee:  
**Learning to Weight Imperfect Demonstrations.** 10961-10970
- Yujing Wang, Yaming Yang, Jiangang Bai, Mingliang Zhang, Jing Bai, Jing Yu, Ce Zhang, Gao Huang, Yunhai Tong:  
**Evolving Attention with Residual Convolutions.** 10971-10980
- Xiang Wang, Shuai Yuan, Chenwei Wu, Rong Ge:  
**Guarantees for Tuning the Step Size using a Learning-to-Learn Approach.** 10981-10990
- Haoxiang Wang, Han Zhao, Bo Li:  
**Bridging Multi-Task Learning and Meta-Learning: Towards Efficient Training and Effective Adaptation.** 10991-11002
- Kaixin Wang, Kuangqi Zhou, Qixin Zhang , Jie Shao, Bryan Hooi, Jiashi Feng:  
**Towards Better Laplacian Representation in Reinforcement Learning with Generalized Graph Drawing.** 11003-11012
- Andrew Warrington, Jonathan Wilder Lavington, Adam Scibior, Mark Schmidt, Frank Wood:  
**Robust Asymmetric Learning in POMDPs.** 11013-11023
- Tomoki Watanabe, Paolo Favaro:  
**A Unified Generative Adversarial Network Training via Self-Labeling and Self-Attention.** 11024-11034
- Dennis Wei:  
**Decision-Making Under Selective Labels: Optimal Finite-Domain Policies and Beyond.** 11035-11046
- Song Wei, Yao Xie, Dobromir Rahnev:  
**Inferring serial correlation with dynamic backgrounds.** 11047-11057



- Ying Wei, Peilin Zhao, Junzhou Huang:  
**Meta-learning Hyperparameter Performance Prediction with Neural Processes.** 11058-11067
- Eli N. Weinstein, Debora S. Marks:  
**A Structured Observation Distribution for Generative Biological Sequence Prediction and Forecasting.** 11068-11079
- Gail Weiss , Yoav Goldberg, Eran Yahav:  
**Thinking Like Transformers.** 11080-11090
- Hongwei Wen, Jingyi Cui, Hanyuan Hang, Jiabin Liu, Yisen Wang, Zhouchen Lin:  
**Leveraged Weighted Loss for Partial Label Learning.** 11091-11100
- Junfeng Wen, Saurabh Kumar, Ramki Gummadi, Dale Schuurmans:  
**Characterizing the Gap Between Actor-Critic and Policy Gradient.** 11101-11111
- Zixin Wen, Yuanzhi Li:  
**Toward Understanding the Feature Learning Process of Self-supervised Contrastive Learning.** 11112-11122
- Chuan Wen, Jierui Lin, Jianing Qian, Yang Gao, Dinesh Jayaraman:  
**Keyframe-Focused Visual Imitation Learning.** 11123-11133
- Jack Weston, Raphael Lenain, Udeepa Meepegama, Emil Fristed:  
**Learning de-identified representations of prosody from raw audio.** 11134-11145
- Jay Whang, Qi Lei, Alex Dimakis:  
**Solving Inverse Problems with a Flow-based Noise Model.** 11146-11157
- Jay Whang, Erik M. Lindgren, Alex Dimakis:  
**Composing Normalizing Flows for Inverse Problems.** 11158-11169
- Noam Wies, Yoav Levine, Daniel Jannai, Amnon Shashua:  
**Which transformer architecture fits my data? A vocabulary bottleneck in self-attention.** 11170-11181
- Mateusz Wilinski, Andrey Y. Lokhov:  
**Prediction-Centric Learning of Independent Cascade Dynamics from Partial Observations.** 11182-11192
- Catherine Wong, Kevin Ellis, Joshua B. Tenenbaum, Jacob Andreas:  
**Leveraging Language to Learn Program Abstractions and Search Heuristics.** 11193-11204
- Eric Wong, Shibani Santurkar, Aleksander Madry:  
**Leveraging Sparse Linear Layers for Debuggable Deep Networks.** 11205-11216
- Mitchell Wortsman, Maxwell Horton, Carlos Guestrin, Ali Farhadi, Mohammad Rastegari:  
**Learning Neural Network Subspaces.** 11217-11227
- Hao Wu, Babak Esmaeili , Michael L. Wick, Jean-Baptiste Tristan, Jan-Willem van de Meent:  
**Conjugate Energy-Based Models.** 11228-11239
- Ruihan Wu, Chuan Guo, Felix Wu, Rahul Kidambi, Laurens van der Maaten, Kilian Q. Weinberger:  
**Making Paper Reviewing Robust to Bid Manipulation Attacks.** 11240-11250
- Yuhuai Wu, Markus N. Rabe, Wenda Li, Jimmy Ba, Roger B. Grosse, Christian Szegedy:  
**LIME: Learning Inductive Bias for Primitives of Mathematical Reasoning.** 11251-11262
- Qingyun Wu, Chi Wang, John Langford, Paul Mineiro, Marco Rossi:  
**ChaCha for Online AutoML.** 11263-11273
- Xueqing Wu, Lewen Wang, Yingce Xia, Weiqing Liu, Lijun Wu, Shufang Xie, Tao Qin, Tie-Yan Liu:  
**Temporally Correlated Task Scheduling for Sequence Learning.** 11274-11284
- Songhua Wu, Xiaobo Xia, Tongliang Liu, Bo Han, Mingming Gong, Nannan Wang, Haifeng Liu, Gang Niu:  
**Class2Simi: A Noise Reduction Perspective on Learning with Noisy Labels.** 11285-11295
- Tianhao Wu, Yunchang Yang, Simon S. Du, Liwei Wang:  
**On Reinforcement Learning with Adversarial Corruption and Its Application to Block MDP.** 11296-11306
- Yi-Fu Wu, Jaesik Yoon, Sungjin Ahn:  
**Generative Video Transformer: Can Objects be the Words?** 11307-11318
- Yue Wu, Shuangfei Zhai, Nitish Srivastava, Joshua M. Susskind, Jian Zhang, Ruslan Salakhutdinov, Hanlin Goh:  
**Uncertainty Weighted Actor-Critic for Offline Reinforcement Learning.** 11319-11328
- Qitian Wu, Hengrui Zhang, Xiaofeng Gao, Junchi Yan, Hongyuan Zha:  
**Towards Open-World Recommendation: An Inductive Model-based Collaborative Filtering Approach.** 11329-11339
- Markus Wulfmeier, Dushyant Rao, Roland Hafner, Thomas Lampe, Abbas Abdolmaleki, Tim Hertweck, Michael Neunert, Dhruva Tirumala, Noah Y. Siegel, Nicolas Heess, Martin A. Riedmiller:  
**Data-efficient Hindsight Off-policy Option Learning.** 11340-11350
- Zehao Xiao, Jiayi Shen, Xiantong Zhen, Ling Shao, Cees Snoek:  
**A Bit More Bayesian: Domain-Invariant Learning with Uncertainty.** 11351-11361
- Chenjun Xiao, Yifan Wu, Jincheng Mei, Bo Dai, Tor Lattimore, Lihong Li, Csaba Szepesvári, Dale Schuurmans:  
**On the Optimality of Batch Policy Optimization Algorithms.** 11362-11371
- Chulin Xie, Minghao Chen, Pin-Yu Chen, Bo Li:  
**CRFL: Certifiably Robust Federated Learning against Backdoor Attacks.** 11372-11382



- Xiaofei Xie, Wenbo Guo, Lei Ma, Wei Le, Jian Wang, Lingjun Zhou, Yang Liu, Xinyu Xing:  
**RNNRepair: Automatic RNN Repair via Model-based Analysis.** 11383-11392
- Annie Xie, James Harrison, Chelsea Finn:  
**Deep Reinforcement Learning amidst Continual Structured Non-Stationarity.** 11393-11403
- Tengyang Xie, Nan Jiang:  
**Batch Value-function Approximation with Only Realizability.** 11404-11413
- Tengyang Xie, John Langford, Paul Mineiro, Ida Momennejad:  
**Interaction-Grounded Learning.** 11414-11423
- Sang Michael Xie, Tengyu Ma, Percy Liang:  
**Composed Fine-Tuning: Freezing Pre-Trained Denoising Autoencoders for Improved Generalization.** 11424-11435
- Qiaomin Xie, Zhuoran Yang, Zhaoran Wang, Andreea Minca:  
**Learning While Playing in Mean-Field Games: Convergence and Optimality.** 11436-11447
- Zeke Xie, Li Yuan, Zhanxing Zhu, Masashi Sugiyama:  
**Positive-Negative Momentum: Manipulating Stochastic Gradient Noise to Improve Generalization.** 11448-11458
- Ran Xin, Usman A. Khan, Soumya Kar:  
**A Hybrid Variance-Reduced Method for Decentralized Stochastic Non-Convex Optimization.** 11459-11469
- Shengzhou Xiong, Yihua Tan, Guoyou Wang:  
**Explore Visual Concept Formation for Image Classification.** 11470-11479
- Tengyu Xu, Yingbin Liang, Guanghui Lan:  
**CRPO: A New Approach for Safe Reinforcement Learning with Convergence Guarantee.** 11480-11491
- Han Xu, Xiaorui Liu, Yixin Li, Anil K. Jain, Jiliang Tang:  
**To be Robust or to be Fair: Towards Fairness in Adversarial Training.** 11492-11501
- Wenkai Xu, Takeru Matsuda:  
**Interpretable Stein Goodness-of-fit Tests on Riemannian Manifold.** 11502-11513
- Da Xu, Chuanwei Ruan, Evren Körpeoglu, Sushant Kumar, Kannan Achan:  
**Rethinking Neural vs. Matrix-Factorization Collaborative Filtering: the Theoretical Perspectives.** 11514-11524
- Yi Xu, Lei Shang, Jinxing Ye, Qi Qian, Yufeng Li, Baigui Sun, Hao Li, Rong Jin:  
**Dash: Semi-Supervised Learning with Dynamic Thresholding.** 11525-11536
- Minkai Xu, Wujie Wang, Shitong Luo, Chence Shi, Yoshua Bengio, Rafael Gómez-Bombarelli, Jian Tang:  
**An End-to-End Framework for Molecular Conformation Generation via Bilevel Programming.** 11537-11547
- Minghao Xu, Hang Wang, Bingbing Ni, Hongyu Guo, Jian Tang:  
**Self-supervised Graph-level Representation Learning with Local and Global Structure.** 11548-11558
- Chen Xu, Yao Xie:  
**Conformal prediction interval for dynamic time-series.** 11559-11569
- Jiaming Xu, Kuang Xu, Dana Yang:  
**Learner-Private Convex Optimization.** 11570-11580
- Tengyu Xu, Zhuoran Yang, Zhaoran Wang, Yingbin Liang:  
**Doubly Robust Off-Policy Actor-Critic: Convergence and Optimality.** 11581-11591
- Keyulu Xu, Mozhi Zhang, Stefanie Jegelka, Kenji Kawaguchi:  
**Optimization of Graph Neural Networks: Implicit Acceleration by Skip Connections and More Depth.** 11592-11602
- Kan Xu, Xuanyi Zhao, Hamsa Bastani, Osbert Bastani:  
**Group-Sparse Matrix Factorization for Transfer Learning of Word Embeddings.** 11603-11612
- Jingjing Xu, Liang Zhao, Junyang Lin, Rundong Gao, Xu Sun, Hongxia Yang:  
**KNAS: Green Neural Architecture Search.** 11613-11625
- Yassine Yaakoubi, François Fleuret, Simon Lacoste-Julien:  
**Structured Convolutional Kernel Networks for Airline Crew Scheduling.** 11626-11636
- Ikko Yamane, Junya Honda, Florian Yger, Masashi Sugiyama:  
**Mediated Uncoupled Learning: Learning Functions without Direct Input-output Correspondences.** 11637-11647
- Yu Yan, Jiusheng Chen, Weizhen Qi, Nikhil Bhendawade, Yeyun Gong, Nan Duan, Ruofei Zhang:  
**EL-Attention: Memory Efficient Lossless Attention for Generation.** 11648-11658
- Zuoyu Yan, Tengfei Ma, Liangcai Gao, Zhi Tang, Chao Chen:  
**Link Prediction with Persistent Homology: An Interactive View.** 11659-11669
- Shen Yan, Kaiqiang Song, Fei Liu, Mi Zhang:  
**CATE: Computation-aware Neural Architecture Encoding with Transformers.** 11670-11681
- Zeyu Yan, Fei Wen, Rendong Ying, Chao Ma, Peilin Liu:  
**On Perceptual Lossy Compression: The Cost of Perceptual Reconstruction and An Optimal Training Framework.** 11682-11692
- Hanshu Yan, Jingfeng Zhang, Gang Niu, Jiashi Feng, Vincent Y. F. Tan, Masashi Sugiyama:  
**CIFS: Improving Adversarial Robustness of CNNs via Channel-wise Importance-based Feature Selection.** 11693-11703



Zitong Yang, Yu Bai, Song Mei:  
**Exact Gap between Generalization Error and Uniform Convergence in Random Feature Models.** 11704-11715

Chunxue Yang, Xiaohui Bei:  
**Learning Optimal Auctions with Correlated Valuations from Samples.** 11716-11726

Greg Yang, Edward J. Hu:  
**Tensor Programs IV: Feature Learning in Infinite-Width Neural Networks.** 11727-11737

Xiaolong Yang, Xiaohong Jia, Dihong Gong, Dong-Ming Yan, Zhifeng Li, Wei Liu:  
**LARNet: Lie Algebra Residual Network for Face Recognition.** 11738-11750

Yi-Rui Yang, Wu-Jun Li:  
**BASGD: Buffered Asynchronous SGD for Byzantine Learning.** 11751-11761

Greg Yang, Eta Littwin:  
**Tensor Programs IIb: Architectural Universality Of Neural Tangent Kernel Training Dynamics.** 11762-11772

Yongyi Yang, Tang Liu, Yangkun Wang, Jinjing Zhou, Quan Gan, Zhewei Wei, Zheng Zhang, Zengfeng Huang , David Wipf:  
**Graph Neural Networks Inspired by Classical Iterative Algorithms.** 11773-11783

Mengjiao Yang, Ofir Nachum:  
**Representation Matters: Offline Pretraining for Sequential Decision Making.** 11784-11794

Tsung-Yen Yang, Justinian Rosca, Karthik Narasimhan, Peter J. Ramadge:  
**Accelerating Safe Reinforcement Learning with Constraint-mismatched Baseline Policies.** 11795-11807

Chao-Han Huck Yang, Yun-Yun Tsai, Pin-Yu Chen:  
**Voice2Series: Reprogramming Acoustic Models for Time Series Classification.** 11808-11819

Zhiyong Yang, Qianqian Xu, Shilong Bao, Yuan He, Xiaochun Cao, Qingming Huang:  
**When All We Need is a Piece of the Pie: A Generic Framework for Optimizing Two-way Partial AUC.** 11820-11829

Xue Yang, Junchi Yan, Qi Ming, Wentao Wang, Xiaopeng Zhang, Qi Tian:  
**Rethinking Rotated Object Detection with Gaussian Wasserstein Distance Loss.** 11830-11841

Yuzhe Yang, Kaiwen Zha, Ying-Cong Chen, Hao Wang, Dina Katabi:  
**Delving into Deep Imbalanced Regression.** 11842-11851

Yukun Yang, Wenrui Zhang, Peng Li:  
**Backpropagated Neighborhood Aggregation for Accurate Training of Spiking Neural Networks.** 11852-11862

Lingxiao Yang, Ru-Yuan Zhang, Lida Li, Xiaohua Xie:  
**SimAM: A Simple, Parameter-Free Attention Module for Convolutional Neural Networks.** 11863-11874

Zhewei Yao, Zhen Dong, Zhangcheng Zheng, Amir Gholami, Jiali Yu, Eric Tan, Leyuan Wang, Qijing Huang, Yida Wang, Michael W. Mahoney, Kurt Keutzer:  
**HAWQ-V3: Dyadic Neural Network Quantization.** 11875-11886

Huaxiu Yao, Long-Kai Huang, Linjun Zhang, Ying Wei, Li Tian, James Zou, Junzhou Huang, Zhenhui Li:  
**Improving Generalization in Meta-learning via Task Augmentation.** 11887-11897

Junwen Yao, Jonas Mueller, Jane-Ling Wang:  
**Deep Learning for Functional Data Analysis with Adaptive Basis Layers.** 11898-11908

Pau Ching Yap, Hippolyt Ritter, David Barber:  
**Addressing Catastrophic Forgetting in Few-Shot Problems.** 11909-11919

Denis Yarats, Rob Fergus, Alessandro Lazaric, Lerrel Pinto:  
**Reinforcement Learning with Prototypical Representations.** 11920-11931

Dmitry Yarotsky:  
**Elementary superexpressive activations.** 11932-11940

Michihiro Yasunaga, Percy Liang:  
**Break-It-Fix-It: Unsupervised Learning for Program Repair.** 11941-11952

Eric C. Yeats, Yiran Chen, Hai Li:  
**Improving Gradient Regularization using Complex-Valued Neural Networks.** 11953-11963

Hugo Yèche, Gideon Dresdner, Francesco Locatello, Matthias Hüser, Gunnar Rätsch:  
**Neighborhood Contrastive Learning Applied to Online Patient Monitoring.** 11964-11974

Gilad Yehudai, Ethan Fetaya, Eli A. Meirom, Gal Chechik, Haggai Maron:  
**From Local Structures to Size Generalization in Graph Neural Networks.** 11975-11986

Mingyang Yi, Lu Hou, Jiacheng Sun, Lifeng Shang, Xin Jiang, Qun Liu, Zhiming Ma:  
**Improved OOD Generalization via Adversarial Training and Pretraining.** 11987-11997

Xinlei Yi, Xiuxian Li, Tao Yang, Lihua Xie, Tianyou Chai, Karl Henrik Johansson:  
**Regret and Cumulative Constraint Violation Analysis for Online Convex Optimization with Long Term Constraints.** 11998-12008

Çağatay Yıldız, Markus Heinonen, Harri Lähdesmäki:  
**Continuous-time Model-based Reinforcement Learning.** 12009-12018

Rong Yin, Yong Liu, Weiping Wang, Dan Meng:  
**Distributed Nyström Kernel Learning with Communications.** 12019-12028



- Ryo Yonetani, Tatsunori Taniai, Mohammadamin Barekatain, Mai Nishimura, Asako Kanezaki:  
**Path Planning using Neural A\* Search.** 12029-12039
- Jihyeong Yoo, Qifeng Chen:  
**SinIR: Efficient General Image Manipulation with Single Image Reconstruction.** 12040-12050
- Boseon Yoo, Jiwoo Lee, Janghoon Ju, Seijun Chung , Soyeon Kim, Jaesik Choi:  
**Conditional Temporal Neural Processes with Covariance Loss.** 12051-12061
- Jongmin Yoon, Sung Ju Hwang, Juho Lee:  
**Adversarial Purification with Score-based Generative Models.** 12062-12072
- Jaehong Yoon, Wonyong Jeong, Giwoong Lee, Eunho Yang, Sung Ju Hwang:  
**Federated Continual Learning with Weighted Inter-client Transfer.** 12073-12086
- Sangwoong Yoon, Yung-Kyun Noh, Frank Chongwoo Park:  
**Autoencoding Under Normalization Constraints.** 12087-12097
- Taeho Yoon, Ernest K. Ryu:  
**Accelerated Algorithms for Smooth Convex-Concave Minimax Problems with O(1/k^2) Rate on Squared Gradient Norm.** 12098-12109
- Shuhei M. Yoshida, Takashi Takenouchi, Masashi Sugiyama:  
**Lower-Bounded Proper Losses for Weakly Supervised Classification.** 12110-12120
- Yuning You, Tianlong Chen, Yang Shen, Zhangyang Wang:  
**Graph Contrastive Learning Automated.** 12121-12132
- Kaichao You, Yong Liu, Jianmin Wang , Mingsheng Long :  
**LogME: Practical Assessment of Pre-trained Models for Transfer Learning.** 12133-12143
- Xuchen You, Xiaodi Wu:  
**Exponentially Many Local Minima in Quantum Neural Networks.** 12144-12155
- Yue Yu, Tian Gao, Naiyu Yin, Qiang Ji:  
**DAGs with No Curl: An Efficient DAG Structure Learning Approach.** 12156-12166
- Tiancheng Yu, Yi Tian, Jingzhao Zhang, Suvrit Sra:  
**Provably Efficient Algorithms for Multi-Objective Competitive RL.** 12167-12176
- Zhongjie Yu, Fabrizio Ventola, Kristian Kersting:  
**Whittle Networks: A Deep Likelihood Model for Time Series.** 12177-12186
- Tianshu Yu, Runzhong Wang, Junchi Yan, Baoxin Li:  
**Deep Latent Graph Matching.** 12187-12197
- Jiaqian Yu, Jingtao Xu, Yiwei Chen, Weiming Li, Qiang Wang, ByungIn Yoo, Jae-Joon Han:  
**Learning Generalized Intersection Over Union for Dense Pixelwise Prediction.** 12198-12207
- Da Yu, Huishuai Zhang, Wei Chen, Jian Yin, Tie-Yan Liu:  
**Large Scale Private Learning via Low-rank Reparametrization.** 12208-12218
- Zhuoning Yuan, Zhishuai Guo, Yi Xu, Yiming Ying, Tianbao Yang:  
**Federated Deep AUC Maximization for Heterogeneous Data with a Constant Communication Complexity.** 12219-12229
- Chia-Hung Yuan, Shan-Hung Wu:  
**Neural Tangent Generalization Attacks.** 12230-12240
- Hao Yuan, Haiyang Yu, Jie Wang, Kang Li, Shuiwang Ji:  
**On Explainability of Graph Neural Networks via Subgraph Explorations.** 12241-12252
- Honglin Yuan, Manzil Zaheer, Sashank J. Reddi:  
**Federated Composite Optimization.** 12253-12266
- Alp Yurtsever, Varun Mangalick, Suvrit Sra:  
**Three Operator Splitting with a Nonconvex Loss Function.** 12267-12277
- Santiago Zanella Béguelin, Shruti Tople, Andrew Paverd, Boris Köpf:  
**Grey-box Extraction of Natural Language Models.** 12278-12286
- Andrea Zanette:  
**Exponential Lower Bounds for Batch Reinforcement Learning: Batch RL can be Exponentially Harder than Online RL.** 12287-12297
- Valentina Zantedeschi, Matt J. Kusner, Vlad Niculae:  
**Learning Binary Decision Trees by Argmin Differentiation.** 12298-12309
- Jure Zbontar, Li Jing, Ishan Misra, Yann LeCun, Stéphane Deny:  
**Barlow Twins: Self-Supervised Learning via Redundancy Reduction.** 12310-12320
- Zhanpeng Zeng, Yunyang Xiong, Sathya N. Ravi, Shailesh Acharya, Glenn Moo Fung, Vikas Singh:  
**You Only Sample (Almost) Once: Linear Cost Self-Attention Via Bernoulli Sampling.** 12321-12332
- Daochen Zha, Jingru Xie, Wenye Ma, Sheng Zhang, Xiangru Lian, Xia Hu, Ji Liu:  
**DouZero: Mastering DouDizhu with Self-Play Deep Reinforcement Learning.** 12333-12344



- Runtian Zhai, Chen Dan, J. Zico Kolter, Pradeep Ravikumar:  
**DORO: Distributional and Outlier Robust Optimization.** 12345-12355
- Dinghuai Zhang, Kartik Ahuja, Yilun Xu, Yisen Wang, Aaron C. Courville:  
**Can Subnetwork Structure Be the Key to Out-of-Distribution Generalization?** 12356-12367
- Bohang Zhang, Tianle Cai, Zhou Lu, Di He, Liwei Wang:  
**Towards Certifying L-infinity Robustness using Neural Networks with L-inf-dist Neurons.** 12368-12379
- Zhenyu Zhang, Xuxi Chen, Tianlong Chen, Zhangyang Wang:  
**Efficient Lottery Ticket Finding: Less Data is More.** 12380-12390
- Xuezhou Zhang, Yiding Chen, Xiaojin Zhu, Wen Sun:  
**Robust Policy Gradient against Strong Data Corruption.** 12391-12401
- Zihan Zhang, Simon S. Du, Xiangyang Ji:  
**Near Optimal Reward-Free Reinforcement Learning.** 12402-12412
- Shujian Zhang, Xinjie Fan, Bo Chen, Mingyuan Zhou :  
**Bayesian Attention Belief Networks.** 12413-12426
- Lily H. Zhang, Mark Goldstein, Rajesh Ranganath:  
**Understanding Failures in Out-of-Distribution Detection with Deep Generative Models.** 12427-12436
- Hang Zhang, Yeyun Gong, Yelong Shen, Weisheng Li, Jiancheng Lv, Nan Duan, Weizhu Chen:  
**Poolingformer: Long Document Modeling with Pooling Attention.** 12437-12446
- Honghua Zhang, Brendan Juba, Guy Van den Broeck:  
**Probabilistic Generating Circuits.** 12447-12457
- Wanrong Zhang, Gautam Kamath, Rachel Cummings:  
**PAPRIKA: Private Online False Discovery Rate Control.** 12458-12467
- Mingyuan Zhang, Jane H. Lee, Shivani Agarwal:  
**Learning from Noisy Labels with No Change to the Training Process.** 12468-12478
- Jiawei Zhang, Linyi Li, Huichen Li, Xiaolu Zhang, Shuang Yang, Bo Li:  
**Progressive-Scale Boundary Blackbox Attack via Projective Gradient Estimation.** 12479-12490
- Tianhao Zhang , Yueheng Li, Chen Wang, Guangming Xie, Zongqing Lu:  
**FOP: Factorizing Optimal Joint Policy of Maximum-Entropy Multi-Agent Reinforcement Learning.** 12491-12500
- Yivan Zhang, Gang Niu, Masashi Sugiyama:  
**Learning Noise Transition Matrix from Only Noisy Labels via Total Variation Regularization.** 12501-12512
- Mengyan Zhang, Cheng Soon Ong:  
**Quantile Bandits for Best Arms Identification.** 12513-12523
- Shufei Zhang, Zhuang Qian, Kaizhu Huang, Qiufeng Wang, Rui Zhang, Xinping Yi:  
**Towards Better Robust Generalization with Shift Consistency Regularization.** 12524-12534
- Yiming Zhang, Keith W. Ross:  
**On-Policy Deep Reinforcement Learning for the Average-Reward Criterion.** 12535-12545
- Zhaoyang Zhang, Wenqi Shao, Jinwei Gu, Xiaogang Wang, Ping Luo:  
**Differentiable Dynamic Quantization with Mixed Precision and Adaptive Resolution.** 12546-12556
- Miao Zhang, Steven W. Su, Shirui Pan, Xiaojun Chang, M. Ehsan Abbasnejad, Reza Haffari:  
**iDARTS: Differentiable Architecture Search with Stochastic Implicit Gradients.** 12557-12566
- Yijie Zhang, Herke van Hoof :  
**Deep Coherent Exploration for Continuous Control.** 12567-12577
- Shantong Zhang, Yi Wan, Richard S. Sutton, Shimon Whiteson:  
**Average-Reward Off-Policy Policy Evaluation with Function Approximation.** 12578-12588
- Mengjiao Zhang, Shusen Wang:  
**Matrix Sketching for Secure Collaborative Machine Learning.** 12589-12599
- Jin Zhang, Jianhao Wang, Hao Hu, Tong Chen, Yingfeng Chen, Changjie Fan, Chongjie Zhang:  
**MetaCURE: Meta Reinforcement Learning with Empowerment-Driven Exploration.** 12600-12610
- Lunjun Zhang, Ge Yang, Bradly C. Stadie:  
**World Model as a Graph: Learning Latent Landmarks for Planning.** 12611-12620
- Shantong Zhang, Hengshuai Yao, Shimon Whiteson:  
**Breaking the Deadly Triad with a Target Network.** 12621-12631
- Shumao Zhang, Pengchuan Zhang, Thomas Y. Hou:  
**Multiscale Invertible Generative Networks for High-Dimensional Bayesian Inference.** 12632-12641
- Qian Zhang, Yilin Zheng, Jean Honorio :  
**Meta Learning for Support Recovery in High-dimensional Precision Matrix Estimation.** 12642-12652
- Zihan Zhang, Yuan Zhou, Xiangyang Ji:  
**Model-Free Reinforcement Learning: from Clipped Pseudo-Regret to Sample Complexity.** 12653-12662



- Zhu Zhang, Chang Zhou, Jianxin Ma, Zhijie Lin, Jingren Zhou, Hongxia Yang, Zhou Zhao:  
**Learning to Rehearse in Long Sequence Memorization.** 12663-12673
- Bo Zhao, Hakan Bilen :  
**Dataset Condensation with Differentiable Siamese Augmentation.** 12674-12685
- Junbo Zhao, Mingfeng Ou, Linji Xue, Yunkai Cui, Sai Wu, Gang Chen:  
**Joining datasets via data augmentation in the label space for neural networks.** 12686-12696
- Zihao Zhao, Eric Wallace, Shi Feng, Dan Klein, Sameer Singh:  
**Calibrate Before Use: Improving Few-shot Performance of Language Models.** 12697-12706
- Yiyang Zhao, Linnan Wang, Yuandong Tian, Rodrigo Fonseca, Tian Guo:  
**Few-Shot Neural Architecture Search.** 12707-12718
- Xin Zhao, Zeru Zhang, Zijie Zhang , Lingfei Wu, Jiayin Jin, Yang Zhou, Ruoming Jin, Dejing Dou, Da Yan:  
**Expressive 1-Lipschitz Neural Networks for Robust Multiple Graph Learning against Adversarial Attacks.** 12719-12735
- Renjie Zheng, Junkun Chen, Mingbo Ma, Liang Huang:  
**Fused Acoustic and Text Encoding for Multimodal Bilingual Pretraining and Speech Translation.** 12736-12746
- Wenbo Zheng, Lan Yan, Chao Gou, Fei-Yue Wang:  
**Two Heads are Better Than One: Hypergraph-Enhanced Graph Reasoning for Visual Event Ratiocination.** 12747-12760
- Xuebin Zheng, Bingxin Zhou, Junbin Gao, Yuguang Wang, Pietro Lió, Ming Li, Guido Montúfar:  
**How Framelets Enhance Graph Neural Networks.** 12761-12771
- Zixin Zhong, Wang Chi Cheung, Vincent Y. F. Tan:  
**Probabilistic Sequential Shrinking: A Best Arm Identification Algorithm for Stochastic Bandits with Corruptions.** 12772-12781
- Fangwei Zhong, Peng Sun, Wenhan Luo, Tingyun Yan, Yizhou Wang:  
**Towards Distraction-Robust Active Visual Tracking.** 12782-12792
- Dongruo Zhou, Jiafan He, Quanquan Gu:  
**Provably Efficient Reinforcement Learning for Discounted MDPs with Feature Mapping.** 12793-12802
- Aurick Zhou, Sergey Levine:  
**Amortized Conditional Normalized Maximum Likelihood: Reliable Out of Distribution Uncertainty Estimation.** 12803-12812
- Fan Zhou, Ping Li:  
**Optimal Estimation of High Dimensional Smooth Additive Function Based on Noisy Observations.** 12813-12823
- Tianchen Zhou, Jia Liu, Chaosheng Dong, Jingyuan Deng:  
**Incentivized Bandit Learning with Self-Reinforcing User Preferences.** 12824-12834
- Dawei Zhou , Tongliang Liu, Bo Han, Nannan Wang, Chunlei Peng, Xinbo Gao:  
**Towards Defending against Adversarial Examples via Attack-Invariant Features.** 12835-12845
- Xiong Zhou, Xianming Liu, Junjun Jiang, Xin Gao, Xiangyang Ji:  
**Asymmetric Loss Functions for Learning with Noisy Labels.** 12846-12856
- Chunting Zhou, Xuezhe Ma , Paul Michel, Graham Neubig:  
**Examining and Combating Spurious Features under Distribution Shift.** 12857-12867
- Mingkang Zhu, Tianlong Chen, Zhangyang Wang:  
**Sparse and Imperceptible Adversarial Attack via a Homotopy Algorithm.** 12868-12877
- Zhuangdi Zhu, Junyuan Hong, Jiayu Zhou:  
**Data-Free Knowledge Distillation for Heterogeneous Federated Learning.** 12878-12889
- Chunjiang Zhu, Qingqing Liu, Jinbo Bi:  
**Spectral vertex sparsifiers and pair-wise spanners over distributed graphs.** 12890-12900
- Hao Zhu, Graham Neubig, Yonatan Bisk:  
**Few-shot Language Coordination by Modeling Theory of Mind.** 12901-12911
- Zhaowei Zhu, Yiwen Song, Yang Liu:  
**Clusterability as an Alternative to Anchor Points When Learning with Noisy Labels.** 12912-12923
- Xinqi Zhu, Chang Xu, Dacheng Tao:  
**Commutative Lie Group VAE for Disentanglement Learning.** 12924-12934
- Huiping Zhuang, Zhenyu Weng, Fulin Luo, Kar-Ann Toj, Haizhou Li, Zhiping Lin:  
**Accumulated Decoupled Learning with Gradient Staleness Mitigation for Convolutional Neural Networks.** 12935-12944
- Dominik Zietlow, Michal Rolínek, Georg Martius:  
**Demystifying Inductive Biases for (Beta)-VAE Based Architectures.** 12945-12954
- Itamar Zimerman, Eliya Nachmani, Lior Wolf:  
**Recovering AES Keys with a Deep Cold Boot Attack.** 12955-12966
- Matthieu Zimmer, Claire Glanois, Umer Siddique, Paul Weng:  
**Learning Fair Policies in Decentralized Cooperative Multi-Agent Reinforcement Learning.** 12967-12978



- Roland S. Zimmermann, Yash Sharma, Steffen Schneider, Matthias Bethge, Wieland Brendel:  
**Contrastive Learning Inverts the Data Generating Process.** 12979-12990
- Luisa M. Zintgraf, Leo Feng, Cong Lu, Maximilian Igl, Kristian Hartikainen, Katja Hofmann, Shimon Whiteson:  
**Exploration in Approximate Hyper-State Space for Meta Reinforcement Learning.** 12991-13001
- Difan Zou, Spencer Frei, Quanquan Gu:  
**Provable Robustness of Adversarial Training for Learning Halfspaces with Noise.** 13002-13011
- Difan Zou, Quanquan Gu:  
**On the Convergence of Hamiltonian Monte Carlo with Stochastic Gradients.** 13012-13022
- Aaron Zweig, Joan Bruna:  
**A Functional Perspective on Learning Symmetric Functions with Neural Networks.** 13023-13032

---

**VALID HTML** last updated on 2024-12-11 20:31 CET by the dblp team

all metadata released as open data under CC0 1.0 license

see also: [Terms of Use](#) | [Privacy Policy](#) | [Imprint](#)

dblp was originally created in 1993 at:



since 2018, dblp has been operated and maintained by:



the dblp computer science bibliography is funded and supported by:

