



We've just launched a new service: our brand new **dblp SPARQL query service**. Read more about it in our latest blog post or try out some of the SPARQL queries linked on the dblp web pages below.

X




## [+] 40th ICML 2023: Honolulu, HI, USA

> Home > Conferences and Workshops > ICML

Dagstuhl

- Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, Jonathan Scarlett:  
**International Conference on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA.** Proceedings of Machine Learning Research 202, PMLR 2023
- Anders Aamand, Alexandr Andoni, Justin Y. Chen, Piotr Indyk, Shyam Narayanan, Sandeep Silwal:  
**Data Structures for Density Estimation.** 1-18
- Ahmed Abbas, Paul Swoboda:  
**ClusterFuG: Clustering Fully connected Graphs by Multicut.** 19-30
- Emmanuel Abbe, Samy Bengio, Aryo Lotfi, Kevin Rizk:  
**Generalization on the Unseen, Logic Reasoning and Degree Curriculum.** 31-60
- Amirhesam Abedsoltan, Mikhail Belkin, Parthe Pandit:  
**Toward Large Kernel Models.** 61-78
- Axel Abels, Tom Lenaerts, Vito Trianni, Ann Nowé:  
**Expertise Trees Resolve Knowledge Limitations in Collective Decision-Making.** 79-90
- Naoufal Acharki, Ramiro Lugo, Antoine Bertoncello, Josselin Garnier:  
**Comparison of meta-learners for estimating multi-valued treatment heterogeneous effects.** 91-132
- Steven Adams, Andrea Patane, Morteza Lahijanian, Luca Laurenti:  
**BNN-DP: Robustness Certification of Bayesian Neural Networks via Dynamic Programming.** 133-151
- Atish Agarwala, Yann N. Dauphin:  
**SAM operates far from home: eigenvalue regularization as a dynamical phenomenon.** 152-168
- Atish Agarwala, Fabian Pedregosa, Jeffrey Pennington:  
**Second-order regression models exhibit progressive sharpening to the edge of stability.** 169-195
- Andrea Agazzi, Jianfeng Lu, Sayan Mukherjee:  
**Global optimality of Elman-type RNNs in the mean-field regime.** 196-227
- Pranjal Aggarwal, Ameet Deshpande, Karthik R. Narasimhan:  
**SemSup-XC: Semantic Supervision for Zero and Few-shot Extreme Classification.** 228-247
- Mehran Aghabozorgi, Shichong Peng, Ke Li:  
**Adaptive IMLE for Few-shot Pretraining-free Generative Modelling.** 248-264
- Armen Aghajanyan, Lili Yu, Alexis Conneau, Wei-Ning Hsu, Karen Hambardzumyan, Susan Zhang, Stephen Roller, Naman Goyal, Omer Levy, Luke Zettlemoyer:  
**Scaling Laws for Generative Mixed-Modal Language Models.** 265-279
- Anass Aghbalou, Guillaume Staerman:  
**Hypothesis Transfer Learning with Surrogate Classification Losses: Generalization Bounds through Algorithmic Stability.** 280-303
- Virginia Aglietti, Alan Malek, Ira Ktena, Silvia Chiappa:  
**Constrained Causal Bayesian Optimization.** 304-321
- Elisabeth Agoritsas, Giovanni Catania, Aurélien Decelle, Beatriz Seoane:  
**Explaining the effects of non-convergent MCMC in the training of Energy-Based Models.** 322-336
- Gati V. Aher, Rosa I. Arriaga, Adam Tauman Kalai:  
**Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies.** 337-371
- Kartik Ahuja, Divyat Mahajan, Yixin Wang, Yoshua Bengio:  
**Interventional Causal Representation Learning.** 372-407

[+] **SPARQL queries**

[–] **Refine list**

showing all 1,829 records

refine by search term

refine by author

Quanquan Gu (14)  
Yoshua Bengio (12)  
Rémi Munos (11)  
Yunhao Tang (10)  
Bernhard Schölkopf (10)  
Michal Valko (9)  
Bo Han 0003 (9)  
Zhangyang Wang (9)  
Tongliang Liu (9)  
Pieter Abbeel (9)  
6,504 more options

refine by access

open (1,829)



- Elisabeth Ailer, Jason S. Hartford, Niki Kilbertus:  
**Sequential Underspecified Instrument Selection for Cause-Effect Estimation.** 408-420
- Matthew Aitchison, Penny Sweetser, Marcus Hutter:  
**Atari-5: Distilling the Arcade Learning Environment down to Five Games.** 421-438
- Naveed Akhtar, Mohammad A. A. K. Jalwana:  
**Towards credible visual model interpretation with path attribution.** 439-457
- Ahmet Alacaoglu, Hanbaek Lyu:  
**Convergence of First-Order Methods for Constrained Nonconvex Optimization with Dependent Data.** 458-489
- Mohammad Mahmudul Alam, Edward Raff, Stella Biderman, Tim Oates, James Holt:  
**Recasting Self-Attention with Holographic Reduced Representations.** 490-507
- Wael Alghamdi, Juan Felipe Gómez, Shahab Asoodeh, Flávio P. Calmon, Oliver Kosut, Lalitha Sankar:  
**The Saddle-Point Method in Differential Privacy.** 508-528
- Christian H. X. Ali Mehmeti-Göpel, Jan Disselhoff:  
**Nonlinear Advantage: Trained Networks Might Not Be As Complex as You Think.** 529-546
- James Urquhart Allingham, Jie Ren, Michael W. Dusenberry, Xiuye Gu, Yin Cui, Dustin Tran, Jeremiah Zhe Liu, Balaji Lakshminarayanan:  
**A Simple Zero-shot Prompt Weighting Technique to Improve Prompt Ensembling in Text-Image Models.** 547-568
- Youssef Allouah , Rachid Guerraoui, Nirupam Gupta, Rafael Pinot, John Stephan:  
**On the Privacy-Robustness-Utility Trilemma in Distributed Learning.** 569-626
- Baris Alparslan, Sinan Yildirim , S. İlker Birbil:  
**Differentially Private Distributed Bayesian Linear Regression with MCMC.** 627-641
- Matías Altamirano, François-Xavier Briol, Jeremias Knoblauch:  
**Robust and Scalable Bayesian Online Changepoint Detection.** 642-663
- Fabian Altekrüger, Johannes Hertrich, Gabriele Steidl:  
**Neural Wasserstein Gradient Flows for Discrepancies with Riesz Kernels.** 664-690
- Sanae Amani, Tor Lattimore, András György, Lin Yang:  
**Distributed Contextual Linear Bandits with Minimax Optimal Communication Cost.** 691-717
- Alan Nawzad Amin, Eli N. Weinstein, Debora Susan Marks:  
**A Kernelized Stein Discrepancy for Biological Sequences.** 718-767
- Philip Amortila, Nan Jiang, Csaba Szepesvári:  
**The Optimal Approximation Factors in Misspecified Off-Policy Value Function Estimation.** 768-790
- Brandon Amos, Giulia Luise, Samuel Cohen, Ievgen Redko:  
**Meta Optimal Transport.** 791-813
- Ioannis Anagnostides, Gabriele Farina, Tuomas Sandholm:  
**Near-Optimal  $\Phi$ -Regret Learning in Extensive-Form Games.** 814-839
- Maksym Andriushchenko, Francesco Croce, Maximilian Müller, Matthias Hein, Nicolas Flammarion:  
**A Modern Look at the Relationship between Sharpness and Generalization.** 840-902
- Maksym Andriushchenko, Aditya Vardhan Varre, Loucas Pillaud-Vivien, Nicolas Flammarion:  
**SGD with Large Step Sizes Learns Sparse Features.** 903-925
- Abdul Fatir Ansari, Alvin Heng, Andre Lim, Harold Soh:  
**Neural Continuous-Discrete State Space Models for Irregularly-Sampled Time Series.** 926-951
- Antonios Antoniadis, Joan Boyar, Marek Eliás, Lene Monrad Favrholdt, Ruben Hoeksma, Kim S. Larsen, Adam Polak, Bertrand Simon:  
**Paging with Succinct Predictions.** 952-968
- Antonios Antoniadis, Christian Coester, Marek Eliás, Adam Polak, Bertrand Simon:  
**Mixing Predictions for Online Metric Algorithms.** 969-983
- Imad Aouali, Victor-Emmanuel Brunel, David Rohde, Anna Korba:  
**Exponential Smoothing for Off-Policy Learning.** 984-1017
- Jamil Arbas, Hassan Ashtiani, Christopher Liaw:  
**Polynomial Time and Private Learning of Unbounded Gaussian Mixture Models.** 1018-1040
- Sohei Arisaka, Qianxiao Li:  
**Principled Acceleration of Iterative Numerical Methods Using Machine Learning.** 1041-1059
- Raman Arora, Raef Bassily, Tomás González, Cristóbal Guzmán, Michael Menart, Enayat Ullah:  
**Faster Rates of Convergence to Stationary Points in Differentially Private Optimization.** 1060-1092
- Nader Asadi, MohammadReza Davari, Sudhir Mudur, Rahaf Aljundi, Eugene Belilovsky:  
**Prototype-Sample Relation Distillation: Towards Replay-Free Continual Learning.** 1093-1106
- Hilal Asi, Vitaly Feldman, Tomer Koren, Kunal Talwar:  
**Near-Optimal Algorithms for Private Online Optimization in the Realizable Regime.** 1107-1120
- Hilal Asi, Jonathan R. Ullman, Lydia Zakynthinou :  
**From Robustness to Privacy and Back.** 1121-1146



- Amit Attia, Tomer Koren:  
**SGD with AdaGrad Stepsizes: Full Adaptivity with High Probability to Unknown Parameters, Unbounded Gradients and Affine Variance.** 1147-1171
- Idan Attias, Steve Hanneke:  
**Adversarially Robust PAC Learnability of Real-Valued Functions.** 1172-1199
- Mattia Atzeni, Mrinmaya Sachan, Andreas Loukas:  
**Infusing Lattice Symmetry Priors in Attention Mechanisms for Sample-Efficient Abstract Geometric Reasoning.** 1200-1217
- Yuval Atzmon, Eli A. Meirom, Shie Mannor, Gal Chechik:  
**Learning to Initiate and Reason in Event-Driven Cascading Processes.** 1218-1243
- Julien Aubert, Luc Lehéritier, Patricia Reynaud-Bouret:  
**On the convergence of the MLE as an estimator of the learning rate in the Exp3 algorithm.** 1244-1275
- Pavel Avdeyev, Chenlai Shi, Yuhao Tan, Ksenia Dudnyk, Jian Zhou:  
**Dirichlet Diffusion Score Model for Biological Sequence Generation.** 1276-1301
- Kyriakos Axiotis, Maxim Sviridenko:  
**Gradient Descent Converges Linearly for Logistic Regression on Separable Data.** 1302-1319
- Alexis Ayme, Claire Boyer, Aymeric Dieuleveut, Erwan Scornet:  
**Naive imputation implicitly regularizes high-dimensional linear models.** 1320-1340
- Mehdi Azabou, Venkataramana Ganesh, Shantanu Thakoor, Chi-Heng Lin, Lakshmi Sathidevi, Ran Liu, Michal Valko, Petar Velickovic, Eva L. Dyer:  
**Half-Hop: A graph upsampling approach for slowing down message passing.** 1341-1360
- Abdus Salam Azad, Izzeddin Gur, Jasper Emhoff, Nathaniel Alexis, Aleksandra Faust, Pieter Abbeel, Ion Stoica:  
**CLUTR: Curriculum Learning via Unsupervised Task Representation Learning.** 1361-1395
- Jinheon Baek, Wonyong Jeong, Jiongdao Jin, Jaehong Yoon, Sung Ju Hwang:  
**Personalized Subgraph Federated Learning.** 1396-1415
- Alexei Baevski, Arun Babu, Wei-Ning Hsu, Michael Auli:  
**Efficient Self-supervised Learning with Contextualized Target Representations for Vision, Speech and Language.** 1416-1429
- Charlotte Baey, Maud Delattre, Estelle Kuhn, Jean-Benoist Leger, Sarah Lemler:  
**Efficient preconditioned stochastic gradient descent for estimation in latent variable models.** 1430-1453
- Haoyue Bai, Gregory Canal, Xuefeng Du, Jeongyeol Kwon, Robert D. Nowak, Yixuan Li:  
**Feed Two Birds with One Scone: Exploiting Wild Data for Both Out-of-Distribution Generalization and Detection.** 1454-1471
- Yushi Bai, Xin Lv, Juanzi Li, Lei Hou:  
**Answering Complex Logical Queries on Knowledge Graphs via Query Computation Tree Optimization.** 1472-1491
- Yikun Bai, Ivan Vladimir Medri, Rocio Diaz Martin, Rana Muhammad Shahroz Khan, Soheil Kolouri:  
**Linear optimal partial transport embedding.** 1492-1520
- Justin M. Baker, Qingsong Wang, Cory D. Hauck, Bao Wang:  
**Implicit Graph Neural Networks: A Monotone Operator Viewpoint.** 1521-1548
- Ainesh Bakshi, Allen Liu, Ankur Moitra, Morris Yau:  
**Tensor Decompositions Meet Control Theory: Learning General Mixtures of Linear Dynamical Systems.** 1549-1563
- Oleg Balabanov, Matthias Beaupère, Laura Grigori, Victor Lederer:  
**Block Subsampled Randomized Hadamard Transform for Nyström Approximation on Distributed Architectures.** 1564-1576
- Philip J. Ball, Laura M. Smith, Ilya Kostrikov, Sergey Levine:  
**Efficient Online Reinforcement Learning with Offline Data.** 1577-1594
- Marin Balliu, Quentin Berthet:  
**Mirror Sinkhorn: Fast Online Optimization on Transport Polytopes.** 1595-1613
- András Balogh, Márk Jelasity:  
**On the Functional Similarity of Robust and Non-Robust Neural Representations.** 1614-1635
- Santiago R. Balseiro, Rachitesh Kumar, Vahab Mirrokni, Balasubramanian Sivan, Di Wang:  
**Robust Budget Pacing with a Single Sample.** 1636-1659
- Kiarash Banihashem, Leyla Biabani, Samira Goudarzi, MohammadTaghi Hajiaghayi, Peyman Jabbarzade, Morteza Monemizadeh:  
**Dynamic Constrained Submodular Optimization with Polylogarithmic Update Time.** 1660-1691
- Fan Bao, Shen Nie, Kaiwen Xue, Chongxuan Li, Shi Pu, Yaole Wang, Gang Yue, Yue Cao, Hang Su, Jun Zhu:  
**One Transformer Fits All Distributions in Multi-Modal Diffusion at Scale.** 1692-1717
- Wenxuan Bao, Haohan Wang, Jun Wu, Jingrui He:  
**Optimizing the Collaboration Structure in Cross-Silo Federated Learning.** 1718-1736
- Omer Bar-Tal, Lior Yariv, Yaron Lipman, Tali Dekel:  
**MultiDiffusion: Fusing Diffusion Paths for Controlled Image Generation.** 1737-1752



- Anas Barakat, Ilyas Fatkhullin, Niao He:  
**Reinforcement Learning with General Utilities: Simpler Variance Reduction and Large State-Action Space.** 1753-1800
- Pietro Barbiero, Gabriele Ciravegna, Francesco Giannini, Mateo Espinosa Zarlenga, Lucie Charlotte Magister, Alberto Tonda, Pietro Lio, Frédéric Precioso, Mateja Jamnik, Giuseppe Marra:  
**Interpretable Neural-Symbolic Concept Reasoning.** 1801-1825
- Burak Bartan, Haoming Li, Harris Teague, Christopher Lott, Bistra Dilkina:  
**Moccasin: Efficient Tensor Rematerialization for Neural Networks.** 1826-1837
- Raef Bassily, Ziteng Sun:  
**User-level Private Stochastic Convex Optimization with Optimal Rates.** 1838-1851
- Soumya Basu, Ankit Singh Rawat, Manzil Zaheer:  
**A Statistical Perspective on Retrieval-Based Models.** 1852-1886
- Jakob Bauer, Kate Baumli, Feryal M. P. Behbahani, Avishkar Bhoopchand, Nathalie Bradley-Schmieg, Michael Chang, Natalie Clay, Adrian Collister, Vibhavari Dasagi, Lucy Gonzalez, Karol Gregor, Edward Hughes, Sheleem Kashem, Maria Loks-Thompson, Hannah Openshaw, Jack Parker-Holder, Shreya Pathak, Nicolas Perez Nieves, Nemanja Rakicevic, Tim Rocktäschel, Yannick Schroecker, Satinder Singh, Jakub Sygnowski, Karl Tuyls, Sarah York, Alexander Zacherl, Lei M. Zhang:  
**Human-Timescale Adaptation in an Open-Ended Task Space.** 1887-1935
- Jerome Baum, Heishiro Kanagawa, Arthur Gretton:  
**A Kernel Stein Test of Goodness of Fit for Sequential Models.** 1936-1953
- Yahav Bechavod, Aaron Roth:  
**Individually Fair Learning with One-Sided Feedback.** 1954-1977
- Sören Becker, Michal Klein, Alexander Neitz, Giambattista Parascandolo, Niki Kilbertus:  
**Predicting Ordinary Differential Equations with Transformers.** 1978-2002
- Daniel Beechey , Thomas M. S. Smith , Özgür Simsek:  
**Explaining Reinforcement Learning with Shapley Values.** 2003-2014
- Maysam Behmanesh, Maximilian Krahn, Maks Ovsjanikov:  
**TIDE: Time Derivative Diffusion for Deep Learning on Graphs.** 2015-2030
- Riade Benbaki, Wenyu Chen, Xiang Meng, Hussein Hazimeh, Natalia Ponomareva, Zhe Zhao, Rahul Mazumder:  
**Fast as CHITA: Neural Network Pruning with Combinatorial Optimization.** 2031-2049
- Christopher M. Bender, Yifeng Shi, Marc Niethammer, Junier Oliva:  
**Continuously Parameterized Mixture Models.** 2050-2062
- Tommaso Bendinelli, Luca Biggio, Pierre-Alexandre Kamienny:  
**Controllable Neural Symbolic Regression.** 2063-2077
- Viktor Bengs, Eyke Hüllermeier, Willem Waegeman:  
**On Second-Order Scoring Rules for Epistemic Uncertainty Quantification.** 2078-2091
- M. Amine Bennouna, Ryan Lucas, Bart P. G. Van Parys :  
**Certified Robust Neural Networks: Generalization and Corruption Resistance.** 2092-2112
- Renato Berlinghieri, Brian L. Trippe, David R. Burt, Ryan James Giordano, Kaushik Srinivasan , Tamay M. Özgökmen, Junfei Xia, Tamara Broderick:  
**Gaussian processes at the Helm(holtz): A more fluid model for ocean currents.** 2113-2163
- Martino Bernasconi, Matteo Castiglioni, Andrea Celli, Alberto Marchesi, Francesco Trovò, Nicola Gatti:  
**Optimal Rates and Efficient Algorithms for Online Bayesian Persuasion.** 2164-2183
- Martino Bernasconi, Matteo Castiglioni, Alberto Marchesi, Francesco Trovò, Nicola Gatti:  
**Constrained Phi-Equilibria.** 2184-2205
- Jeroen Berrevoets , Nabeel Seedat, Fergus Imrie, Mihaela van der Schaar:  
**Differentiable and Transportable Structure Learning.** 2206-2233
- Arturs Berzins:  
**Polyhedral Complex Extraction from ReLU Networks using Edge Subdivision.** 2234-2244
- Louis Béthune, Paul Novello, Guillaume Coiffier, Thibaut Boissin, Mathieu Serrurier, Quentin Vincenot, Andres Troya-Galvis:  
**Robust One-Class Classification with Signed Distance Function using 1-Lipschitz Neural Networks.** 2245-2271
- Beatrice Bevilacqua, Kyriacos Nikiforou, Borja Ibarz, Ioana Bica, Michela Paganini, Charles Blundell, Jovana Mitrovic, Petar Velickovic:  
**Neural Algorithmic Reasoning with Causal Regularisation.** 2272-2288
- Ayush Bharti, Masha Naslidnyk, Oscar Key, Samuel Kaski, François-Xavier Briol:  
**Optimally-weighted Estimators of the Maximum Mean Discrepancy for Likelihood-Free Inference.** 2289-2312
- Aditya Bhaskara, Ashok Cutkosky, Ravi Kumar, Manish Purohit:  
**Bandit Online Linear Optimization with Hints and Queries.** 2313-2336
- Aadyot Bhatnagar, Huan Wang, Caiming Xiong, Yu Bai:  
**Improved Online Conformal Prediction via Strongly Adaptive Online Learning.** 2337-2363



- Robi Bhattacharjee, Sanjoy Dasgupta, Kamalika Chaudhuri:  
**Data-Copying in Generative Models: A Formal Framework.** 2364-2396
- Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O'Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, Oskar van der Wal:  
**Pythia: A Suite for Analyzing Large Language Models Across Training and Scaling.** 2397-2430
- Vaibhav Bihani, Sahil Manchanda, Srikanth Sastry, Sayan Ranu, N. M. Anoop Krishnan:  
**StriderNet: A Graph Reinforcement Learning Approach to Optimize Atomic Structures on Rough Energy Landscapes.** 2431-2451
- Marin Bilos, Kashif Rasul, Anderson Schneider, Yuriy Nevmyvaka, Stephan Günnemann:  
**Modeling Temporal Data as Continuous Functions with Stochastic Process Diffusion.** 2452-2470
- Julian Bitterwolf, Maximilian Müller, Matthias Hein:  
**In or Out? Fixing ImageNet Out-of-Distribution Detection Evaluation.** 2471-2506
- Ondrej Biza, Sjoerd van Steenkiste, Mehdi S. M. Sajjadi, Gamaleldin Fathy Elsayed, Aravindh Mahendran, Thomas Kipf:  
**Invariant Slot Attention: Object Discovery with Slot-Centric Reference Frames.** 2507-2527
- Mitchell Black, Zhengchao Wan, Amir Nayyeri, Yusu Wang:  
**Understanding Oversquashing in GNNs through the Lens of Effective Resistance.** 2528-2547
- Charlie Blake, Douglas Orr, Carlo Luschi:  
**Unit Scaling: Out-of-the-Box Low-Precision Training.** 2548-2576
- Matthieu Blanke, Marc Lelarge:  
**FLEX: an Adaptive Exploration Algorithm for Nonlinear Systems.** 2577-2591
- Markus Bläser:  
**Not all Strongly Rayleigh Distributions Have Small Probabilistic Generating Circuits.** 2592-2602
- Linus Bleistein, Adeline Fermanian, Anne-Sophie Jannot, Agathe Guilloux:  
**Learning the Dynamics of Sparsely Observed Interacting Systems.** 2603-2640
- Niclas Boehmer, L. Elisa Celis, Lingxiao Huang, Anay Mehrotra, Nisheeth K. Vishnoi:  
**Subset Selection Based On Multiple Rankings in the Presence of Bias: Effectiveness of Fairness Constraints for Multiwinner Voting Score Functions.** 2641-2688
- Niclas Boehmer, Piotr Faliszewski, Sonja Kraiczy:  
**Properties of the Mallows Model Depending on the Number of Alternatives: A Warning for an Experimentalist.** 2689-2711
- David Boetius, Stefan Leue, Tobias Sutter:  
**A Robust Optimisation Perspective on Counterexample-Guided Repair of Neural Networks.** 2712-2737
- Simone Bombari, Shayan Kiyani, Marco Mondelli:  
**Beyond the Universal Law of Robustness: Sharper Laws for Random Features and Neural Tangent Kernels.** 2738-2776
- Clément Bonet, Benoît Malézieux, Alain Rakotomamonjy, Lucas Drumetz, Thomas Moreau, Matthieu Kowalski, Nicolas Courty:  
**Sliced-Wasserstein on Symmetric Positive Definite Matrices for M/EEG Signals.** 2777-2805
- Boris Bonev, Thorsten Kurth, Christian Hundt, Jaideep Pathak, Maximilian Baust, Karthik Kashinath, Anima Anandkumar:  
**Spherical Fourier Neural Operators: Learning Stable Dynamics on the Sphere.** 2806-2823
- Victor Boone, Bruno Gaujal:  
**The Regret of Exploration and the Control of Bad Episodes in Reinforcement Learning.** 2824-2856
- Akhilan Boopathy, Kevin Liu, Jaedong Hwang, Shu Ge, Asaad Mohammedsaleh, Ila Fiete:  
**Model-agnostic Measure of Generalization Difficulty.** 2857-2884
- Shahine Bouabid, Jake Fawkes, Dino Sejdinovic:  
**Returning The Favour: When Regression Benefits From Probabilistic Causal Knowledge.** 2885-2913
- Malik Boudiaf, Tom Denton, Bart van Merriënboer, Vincent Dumoulin, Eleni Triantafillou:  
**In Search for a Generalizable Method for Source Free Domain Adaptation.** 2914-2931
- Adam Bouland, Yosheb M. Getachew, Yujia Jin, Aaron Sidford, Kevin Tian:  
**Quantum Speedups for Zero-Sum Games via Improved Dynamic Gibbs Sampling.** 2932-2952
- Victor Boutin, Thomas Fel, Lakshya Singhal, Rishav Mukherji, Akash Nagaraj, Julien Colin , Thomas Serre:  
**Diffusion Models as Artists: Are we Closing the Gap between Humans and Machines?** 2953-3002
- Michael Bowling, John D. Martin, David Abel, Will Dabney:  
**Settling the Reward Hypothesis.** 3003-3020
- Manuel Brack, Patrick Schramowski, Björn Deiseroth, Kristian Kersting:  
**ILLUME: Rationalizing Vision-Language Models through Human Interactions.** 3021-3037
- Jack Brady, Roland S. Zimmermann, Yash Sharma, Bernhard Schölkopf, Julius von Kügelgen, Wieland Brendel:  
**Provably Learning Object-Centric Representations.** 3038-3062
- Gecia Bravo Hermsdorff:  
**Quantifying Human Priors over Social and Navigation Networks.** 3063-3105



- Pierre Bréchet, Katerina Papagiannouli, Jing An, Guido Montúfar:  
**Critical Points and Convergence Analysis of Generative Deep Linear Networks Trained with Bures-Wasserstein Loss.** 3106-3147
- Trenton Bricken, Rylan Schaeffer, Bruno A. Olshausen, Gabriel Kreiman:  
**Emergence of Sparse Representations from Noise.** 3148-3191
- Zhiqi Bu, Yu-Xiang Wang, Sheng Zha, George Karypis:  
**Differentially Private Optimization on Large Model at Small Cost.** 3192-3218
- Alexander Bukharin, Tianyi Liu, Shengjie Wang, Simiao Zuo, Weihao Gao, Wen Yan, Tuo Zhao:  
**Machine Learning Force Fields with Data Cost Aware Training.** 3219-3232
- Róbert Istvan Busa-Fekete, Andrés Muñoz Medina, Umar Syed, Sergei Vassilvitskii:  
**Label differential privacy and private training data release.** 3233-3251
- Vivien Cabannes, Bobak Toussi Kiani, Randall Balestrieri, Yann LeCun, Alberto Bietti:  
**The SSL Interplay: Augmentations, Inductive Bias, and Generalization.** 3252-3298
- Federico Cacciamani, Matteo Castiglioni, Nicola Gatti:  
**Online Mechanism Design for Information Acquisition.** 3299-3326
- Vittorio Caggiano, Sudeep Dasari, Vikash Kumar:  
**MyoDex: A Generalizable Prior for Dexterous Manipulation.** 3327-3346
- Francesco Cagnetta, Alessandro Favero, Matthieu Wyart:  
**What Can Be Learnt With Wide Convolutional Neural Networks?** 3347-3379
- Ruichu Cai, Zhiyi Huang , Wei Chen, Zhifeng Hao, Kun Zhang:  
**Causal Discovery with Latent Confounders Based on Higher-Order Cumulants.** 3380-3407
- Chen Cai, Truong Son Hy, Rose Yu, Yusu Wang:  
**On the Connection Between MPNN and Graph Transformer.** 3408-3430
- Dongqi Cai, Yangyuxuan Kang, Anbang Yao, Yurong Chen:  
**Ske2Grid: Skeleton-to-Grid Representation Learning for Action Recognition.** 3431-3441
- Yuchao Cai, Yuheng Ma, Yiwei Dong, Hanfang Yang:  
**Extrapolated Random Tree for Regression.** 3442-3468
- Xufeng Cai, Chaobing Song, Stephen J. Wright, Jelena Diakonikolas:  
**Cyclic Block Coordinate Descent With Variance Reduction for Composite Nonconvex Optimization.** 3469-3494
- Ruiqi Cai, Zhenyu Zhang, Zhangyang Wang:  
**Robust Weight Signatures: Gaining Robustness as Easy as Patching Weights?** 3495-3506
- Yang Cai, Weiqiang Zheng:  
**Doubly Optimal No-Regret Learning in Monotone Games.** 3507-3524
- Mine Melodi Caliskan, Francesco Chini, Setareh Maghsudi:  
**Multi-Agent Learning from Learners.** 3525-3540
- Yadi Cao, Menglei Chai, Minchen Li, Chenfanfu Jiang:  
**Efficient Learning of Mesh-Based Physical Simulation with Bi-Stride Multi-Scale Graph Neural Network.** 3541-3558
- Jian Cao, Myeongjong Kang, Felix Jimenez, Huiyan Sang, Florian Tobias Schaefer, Matthias Katzfuss:  
**Variational Sparse Inverse Cholesky Approximation for Latent Gaussian Processes via Double Kullback-Leibler Minimization.** 3559-3576
- Shengcao Cao, Mengtian Li, James Hays, Deva Ramanan, Yu-Xiong Wang, Liangyan Gui:  
**Learning Lightweight Object Detectors via Multi-Teacher Progressive Distillation.** 3577-3598
- Steven Cao, Percy Liang, Gregory Valiant:  
**One-sided Matrix Completion from Two Observations Per Row.** 3599-3624
- Gabriel Cardoso, Yazid Janati El Idrissi, Sylvain Le Corff , Eric Moulines, Jimmy Olsson:  
**State and parameter learning with PARIS particle Gibbs.** 3625-3675
- Thomas Carta, Clément Romac, Thomas Wolf, Sylvain Lamprier, Olivier Sigaud, Pierre-Yves Oudeyer:  
**Grounding Large Language Models in Interactive Environments with Online Reinforcement Learning.** 3676-3713
- Nicolas Castanet, Olivier Sigaud, Sylvain Lamprier:  
**Stein Variational Goal Generation for adaptive Exploration in Multi-Goal Reinforcement Learning.** 3714-3731
- Alberto Castellini, Federico Bianchi, Edoardo Zorzi, Thiago D. Simão, Alessandro Farinelli, Matthijs T. J. Spaan:  
**Scalable Safe Policy Improvement via Monte Carlo Tree Search.** 3732-3756
- Timothy Castiglia, Yi Zhou, Shiqiang Wang, Swanand Kadhe, Nathalie Baracaldo, Stacy Patterson:  
**LESS-VFL: Communication-Efficient Feature Selection for Vertical Federated Learning.** 3757-3781
- Rémi Catellier, Samuel Vaïter, Damien Garreau:  
**On the Robustness of Text Vectorizers.** 3782-3814
- Juan Cerviño, Luiz F. O. Chamon, Benjamin David Haeffele, René Vidal, Alejandro Ribeiro:  
**Learning Globally Smooth Functions on Manifolds.** 3815-3854
- Jaeyoung Cha, Jaewook Lee, Chulhee Yun:  
**Tighter Lower Bounds for Shuffling SGD: Random Permutations and Beyond.** 3855-3912



- Jaehoon Cha, Jeyan Thiyagalingam:  
**Orthogonality-Enforced Latent Space in Autoencoders: An Approach to Learning Disentangled Representations.** 3913-3948
- Souradip Chakraborty, Amrit S. Bedi, Alec Koppel, Mengdi Wang, Furong Huang, Dinesh Manocha:  
**STEERING : Stein Information Directed Exploration for Model-Based Reinforcement Learning.** 3949-3978
- Sunrit Chakraborty, Saptarshi Roy, Ambuj Tewari:  
**Thompson Sampling for High-Dimensional Sparse Linear Contextual Bandits.** 3979-4008
- Yash Chandak, Shantanu Thakoor, Zhaohan Daniel Guo, Yunhao Tang, Rémi Munos, Will Dabney, Diana L. Borsa:  
**Representations and Exploration for Deep Reinforcement Learning using Singular Value Decomposition.** 4009-4034
- Paul Edmund Chang, Prakhar Verma , S. T. John, Arno Solin, Mohammad Emtiyaz Khan:  
**Memory-Based Dual Gaussian Processes for Sequential Learning.** 4035-4054
- Huiwen Chang, Han Zhang, Jarred Barber, Aaron Maschinot, José Lezama, Lu Jiang, Ming-Hsuan Yang, Kevin Patrick Murphy, William T. Freeman, Michael Rubinstein, Yuanzhen Li, Dilip Krishnan:  
**Muse: Text-To-Image Generation via Masked Generative Transformers.** 4055-4075
- Chen-Hao Chao, Wei-Fang Sun, Bo-Wun Cheng, Chun-Yi Lee:  
**On Investigating the Conservative Property of Score-Based Generative Models.** 4076-4095
- Vasileios Charisopoulos, Hossein Esfandiari, Vahab Mirrokni:  
**Robust and private stochastic linear bandits.** 4096-4115
- Anamay Chaturvedi, Huy L. Nguyen, Thy Dinh Nguyen:  
**Streaming Submodular Maximization with Differential Privacy.** 4116-4143
- Kamalika Chaudhuri, Kartik Ahuja, Martín Arjovsky, David Lopez-Paz:  
**Why does Throwing Away Data Improve Worst-Group Error?** 4144-4188
- Ronshee Chawla, Daniel Vial, Sanjay Shakkottai, R. Srikant:  
**Collaborative Multi-Agent Heterogeneous Multi-Armed Bandits.** 4189-4217
- Fengdi Che, Gautham Vasan, A. Rupam Mahmood:  
**Correcting discount-factor mismatch in on-policy policy gradient methods.** 4218-4240
- Tianshi Che, Yang Zhou, Zijie Zhang, Lingjuan Lyu, Ji Liu, Da Yan, Dejing Dou, Jun Huan:  
**Fast Federated Machine Unlearning with Nonlinear Functional Theory.** 4241-4268
- David Cheikhi, Daniel Russo:  
**On the Statistical Benefits of Temporal Difference Learning.** 4269-4293
- Zhengdao Chen:  
**Multi-Layer Neural Networks as Trainable Ladders of Hilbert Spaces.** 4294-4329
- Lei Chen, Joan Bruna:  
**Beyond the Edge of Stability via Two-step Gradient Updates.** 4330-4391
- Kuan-Yu Chen, Ping-Han Chiang, Hsin-Rung Chou, Ting-Wei Chen, Tien-Hao Chang:  
**Trompt: Towards a Better Deep Neural Network for Tabular Data.** 4392-4434
- Du Chen, Geoffrey A. Chua:  
**Differentially Private Stochastic Convex Optimization under a Quantile Loss Function.** 4435-4461
- Sitan Chen, Giannis Daras, Alex Dimakis:  
**Restoration-Degradation Beyond Linear Diffusions: A Non-Asymptotic Analysis For DDIM-type Samplers.** 4462-4484
- Yu Chen, Wei Deng, Shikai Fang, Fengpei Li, Nicole Tianjiao Yang, Yikai Zhang, Kashif Rasul, Shandian Zhe, Anderson Schneider, Yuriy Nevmyvaka:  
**Provably Convergent Schrödinger Bridge with Applications to Probabilistic Time Series Imputation.** 4485-4513
- Siyuan Chen, Pratik Pramod Fegade, Tianqi Chen, Phillip B. Gibbons, Todd C. Mowry:  
**ED-Batch: Efficient Automatic Batching of Dynamic Neural Networks via Learned Finite State Machines.** 4514-4528
- Yanzhi Chen, Michael U. Gutmann, Adrian Weller:  
**Is Learning Summary Statistics Necessary for Likelihood-free Inference?** 4529-4544
- Runfa Chen, Jiaqi Han, Fuchun Sun, Wenbing Huang:  
**Subequivariant Graph Reinforcement Learning in 3D Environments.** 4545-4565
- Hanxiao Chen, Meng Hao, Hongwei Li, Kangjie Chen, Guowen Xu, Tianwei Zhang, Xilin Zhang:  
**GuardHFL: Privacy Guardian for Heterogeneous Federated Learning.** 4566-4584
- Xiaohui Chen, Jiaxing He, Xu Han, Liping Liu:  
**Efficient and Degree-Guided Graph Generation via Discrete Diffusion Modeling.** 4585-4610
- Shiming Chen, Wenjin Hou, Ziming Hong, Xiaohan Ding, Yibing Song, Xinge You, Tongliang Liu, Kun Zhang:  
**Evolving Semantic Prototype Improves Generative Zero-Shot Learning.** 4611-4622
- Yimeng Chen, Tianyang Hu, Fengwei Zhou, Zhenguo Li, Zhi-Ming Ma:  
**Explore and Exploit the Diverse Knowledge in Model Zoo for Domain Generalization.** 4623-4640



- Xuxing Chen, Minhui Huang, Shiqian Ma, Krishna Balasubramanian:  
**Decentralized Stochastic Bilevel Optimization with Improved per-Iteration Complexity.** 4641-4671
- Minshuo Chen, Kaixuan Huang, Tuo Zhao, Mengdi Wang:  
**Score Approximation, Estimation and Distribution Recovery of Diffusion Models on Low-Dimensional Data.** 4672-4712
- Ziyu Chen, Markos A. Katsoulakis, Luc Rey-Bellet, Wei Zhu:  
**Sample Complexity of Probability Divergences under Group Symmetry.** 4713-4734
- Hongrui Chen, Holden Lee, Jianfeng Lu:  
**Improved Analysis of Score-based Generative Modeling: User-Friendly Bounds under Minimal Smoothness Assumptions.** 4735-4763
- Yineng Chen, Zuchao Li, Lefei Zhang, Bo Du, Hai Zhao:  
**Bidirectional Looking with A Novel Double Exponential Moving Average to Adaptive and Non-adaptive Momentum Optimizers.** 4764-4803
- Lu Chen, Siyu Lou , Keyan Zhang, Jin Huang , Quanshi Zhang:  
**HarsanyiNet: Computing Accurate Shapley Values in a Single Forward Propagation.** 4804-4825
- Keyi Chen, Francesco Orabona:  
**Generalized Implicit Follow-The-Regularized-Leader.** 4826-4838
- Dexiong Chen, Paolo Pellizzoni, Karsten M. Borgwardt:  
**Fisher Information Embedding for Node and Graph Learning.** 4839-4855
- Jiaxuan Chen, Yu Qi, Gang Pan:  
**Rethinking Visual Reconstruction: Experience-Based Content Completion Guided by Visual Cues.** 4856-4866
- Jiefeng Chen, Jayaram Raghuram, Jihye Choi, Xi Wu, Yingyu Liang, Somesh Jha:  
**Stratified Adversarial Robustness with Rejection.** 4867-4894
- Jiayu Chen, Dipesh Tamboli, Tian Lan, Vaneet Aggarwal:  
**Multi-task Hierarchical Adversarial Inverse Reinforcement Learning.** 4895-4920
- Yatong Chen, Zeyu Tang, Kun Zhang, Yang Liu:  
**Model Transferability with Responsive Decision Subjects.** 4921-4952
- Liyu Chen, Andrea Tirinzoni, Alessandro Lazaric, Matteo Pirotta:  
**Layered State Discovery for Incremental Autonomous Exploration.** 4953-5001
- Sijia Chen, Wei-Wei Tu, Peng Zhao, Lijun Zhang:  
**Optimistic Online Mirror Descent for Bridging Stochastic and Adversarial Online Convex Optimization.** 5002-5035
- Xuxi Chen, Nelson Vadovi, Tianlong Chen, Zhangyang Wang:  
**Learning to Optimize Differentiable Games.** 5036-5051
- Yurong Chen, Qian Wang , Zhijian Duan, Haoran Sun, Zhaohua Chen, Xiang Yan, Xiaotie Deng:  
**Coordinated Dynamic Bidding in Repeated Second-Price Auctions with Budgets.** 5052-5086
- Changyu Chen, Xiting Wang, Yiqiao Jin, Victor Ye Dong, Li Dong, Jie Cao, Yi Liu, Rui Yan:  
**Semi-Offline Reinforcement Learning for Optimized Text Generation.** 5087-5103
- Fan Chen, Huan Wang, Caiming Xiong, Song Mei, Yu Bai:  
**Lower Bounds for Learning in Revealing POMDPs.** 5104-5161
- Honglin Chen, Rundi Wu , Eitan Grinspun, Changxi Zheng, Peter Yichen Chen:  
**Implicit Neural Spatial Representations for Time-dependent PDEs.** 5162-5177
- Sanyuan Chen, Yu Wu, Chengyi Wang, Shujie Liu, Daniel Tompkins, Zhuo Chen, Wanxiang Che, Xiangzhan Yu, Furu Wei:  
**BEATs: Audio Pre-Training with Acoustic Tokenizers.** 5178-5193
- Siyu Chen, Jibang Wu, Yifan Wu, Zhuoran Yang:  
**Learning to Incentivize Information Acquisition: Proper Scoring Rules Meet Principal-Agent Model.** 5194-5218
- Lesi Chen, Jing Xu, Luo Luo:  
**Faster Gradient-Free Algorithms for Nonsmooth Nonconvex Stochastic Optimization.** 5219-5233
- Daoyuan Chen, Liuyi Yao, Dawei Gao, Bolin Ding, Yaliang Li:  
**Efficient Personalized Federated Learning via Sparse Model-Adaptation.** 5234-5256
- Yifan Chen, Rentian Yao, Yun Yang, Jie Chen:  
**A Gromov-Wasserstein Geometric View of Spectrum-Preserving Graph Coarsening.** 5257-5281
- Dangxing Chen, Weicheng Ye:  
**How to address monotonicity for model risk management?** 5282-5295
- Xin Chen, Yicheng Zeng, Siyue Yang, Qiang Sun:  
**Sketched Ridgeless Linear Regression: The Role of Downsampling.** 5296-5326
- Dingyang Chen, Qi Zhang:  
**Context-Aware Bayesian Network Actor-Critic Methods for Cooperative Multi-Agent Reinforcement Learning.** 5327-5350
- Can Chen, Yingxue Zhang, Xue Liu, Mark Coates:  
**Bidirectional Learning for Offline Model-based Biological Sequence Design.** 5351-5366



- Tianqi Chen, Mingyuan Zhou:  
**Learning to Jump: Thinning and Thickening Latent Counts for Generative Modeling.** 5367-5382
- Wuyang Chen, Yanqi Zhou, Nan Du, Yanping Huang, James Laudon, Zhifeng Chen, Claire Cui:  
**Lifelong Language Pretraining with Distribution-Specialized Experts.** 5383-5395
- Ziyi Chen, Yi Zhou, Yingbin Liang, Zhaosong Lu:  
**Generalized-Smooth Nonconvex Optimization is As Efficient As Smooth Nonconvex Optimization.** 5396-5427
- Xin Cheng , Yuzhou Cao, Ximing Li, Bo An, Lei Feng:  
**Weakly Supervised Regression with Interval Targets.** 5428-5448
- Chin-Yi Cheng, Forrest Huang, Gang Li, Yang Li:  
**PLay: Parametrically Conditioned Layout Generation using Latent Diffusion.** 5449-5471
- Minhao Cheng , Rui Min, Haochen Sun, Pin-Yu Chen:  
**Identification of the Adversary from a Single Adversarial Example.** 5472-5484
- Xiaotong Cheng, Cheng Pan, Setareh Maghsudi:  
**Parallel Online Clustering of Bandits via Hedonic Game.** 5485-5503
- Yong Cheng, Yu Zhang, Melvin Johnson, Wolfgang Macherey, Ankur Bapna:  
**Mu<sup>2</sup>SLAM: Multitask, Multilingual Speech and Language Models.** 5504-5520
- Duo Cheng, Xingyu Zhou, Bo Ji:  
**Understanding the Role of Feedback in Online Learning with Switching Costs.** 5521-5543
- David Chiang, Peter Cholak, Anand Pillay:  
**Tighter Bounds on the Expressivity of Transformer Encoders.** 5544-5562
- Muthu Chidambaram, Xiang Wang, Chenwei Wu, Rong Ge:  
**Provably Learning Diverse Features in Multi-View Data with Midpoint Mixup.** 5563-5599
- Muthu Chidambaram, Chenwei Wu, Yu Cheng, Rong Ge:  
**Hiding Data Helps: On the Benefits of Masking for Sparse Coding.** 5600-5615
- Eli Chien, Jiong Zhang, Cho-Jui Hsieh, Jyun-Yu Jiang, Wei-Cheng Chang, Olgica Milenkovic, Hsiang-Fu Yu:  
**PINA: Leveraging Side Information in eXtreme Multi-label Classification via Predicted Instance Neighborhood Aggregation.** 5616-5630
- Hong-Ming Chiu, Richard Y. Zhang:  
**Tight Certification of Adversarially Trained Neural Networks via Nonconvex Low-Rank Semidefinite Relaxations.** 5631-5660
- Cheol Jun Cho, Edward F. Chang, Gopala Krishna Anumanchipalli:  
**Neural Latent Aligner: Cross-trial Alignment for Learning Representations of Complex, Naturalistic Neural Data.** 5661-5676
- Yae Jee Cho, Pranay Sharma , Gauri Joshi, Zheng Xu, Satyen Kale, Tong Zhang:  
**On the Convergence of Federated Averaging with Cyclic Client Participation.** 5677-5721
- Jeongwhan Choi, Seoyoung Hong, Noseong Park, Sung-Bae Cho:  
**GREAD: Graph Neural Reaction-Diffusion Networks.** 5722-5747
- Hee Min Choi, Hyo Kang, Dokwan Oh:  
**Is Overfitting Necessary for Implicit Video Representation?** 5748-5770
- Young-Geun Choi, Gi-Soo Kim, Yunseo Choi, Wooseong Cho, Myunghee Cho Paik, Min-hwan Oh:  
**Semi-Parametric Contextual Pricing Algorithm using Cox Proportional Hazards Model.** 5771-5786
- Jaemoo Choi, Yesom Park, Myungjoo Kang:  
**Restoration based Generative Models.** 5787-5816
- Jihye Choi, Jayaram Raghuram, Ryan Feng, Jiefeng Chen, Somesh Jha, Atul Prakash :  
**Concept-based Explanations for Out-of-Distribution Detectors.** 5817-5837
- Davin Choo, Themistoklis Gouleakis, Arnab Bhattacharyya:  
**Active causal structure learning with advice.** 5838-5867
- Davin Choo, Kirankumar Shiragur:  
**New metrics and search algorithms for weighted causal DAGs.** 5868-5903
- Nicolas Chopin, Andras Fulop, Jeremy Heng, Alexandre H. Thiery:  
**Computational Doob h-transforms for Online Filtering of Discretely Observed Diffusions.** 5904-5923
- Christopher A. Choquette-Choo, Hugh Brendan McMahan, J. Keith Rush, Abhradeep Guha Thakurta:  
**Multi-Epoch Matrix Factorization Mechanisms for Private Machine Learning.** 5924-5963
- Krzysztof Marcin Choromanski:  
**Taming graph kernels with random features.** 5964-5977
- Krzysztof Marcin Choromanski, Arijit Sehanobish, Han Lin, Yunfan Zhao, Eli Berger, Tetiana Parshakova, Alvin Pan, David Watkins, Tianyi Zhang, Valerii Likhoshesterov, Somnath Basu Roy Chowdhury, Kumar Avinava Dubey, Deepali Jain, Tamás Sarlós, Snigdha Chaturvedi, Adrian Weller:  
**Efficient Graph Field Integrators Meet Point Clouds.** 5978-6004



- Era Choshen, Aviv Tamar:  
**ContraBAR: Contrastive Bayes-Adaptive Deep RL.** 6005-6027
- Rishav Chourasia, Neil Shah:  
**Forget Unlearning: Towards True Data-Deletion in Machine Learning.** 6028-6073
- Mohammed Nowaz Rabbani Chowdhury, Shuai Zhang, Meng Wang, Sijia Liu, Pin-Yu Chen:  
**Patch-level Routing in Mixture-of-Experts is Provably Sample-efficient for Convolutional Neural Networks.** 6074-6114
- Rhea Chowers , Yair Weiss:  
**What do CNNs Learn in the First Layer and Why? A Linear Systems Perspective.** 6115-6139
- Dimitrios Christofidellis, Giorgio Giannone, Jannis Born, Ole Winther, Teodoro Laino, Matteo Manica:  
**Unifying Molecular and Textual Representations via Multi-task Language Modelling.** 6140-6157
- Xu Chu, Yujie Jin, Xin Wang, Shanghang Zhang, Yasha Wang, Wenwu Zhu, Hong Mei:  
**Wasserstein Barycenter Matching for Graph Size Generalization of Message Passing Neural Networks.** 6158-6184
- Yu-Min Chu, Chieh Liu, Ting-I Hsieh, Hwann-Tzong Chen, Tyng-Luh Liu:  
**Shape-Guided Dual-Memory Learning for 3D Anomaly Detection.** 6185-6194
- Jianing Chu, Shu Yang, Wenbin Lu:  
**Multiply Robust Off-policy Evaluation and Learning under Truncation by Death.** 6195-6227
- Ching-Yao Chuang, Stefanie Jegelka, David Alvarez-Melis:  
**InfoOT: Information Maximizing Optimal Transport.** 6228-6242
- Bilal Chughtai, Lawrence Chan, Neel Nanda:  
**A Toy Model of Universality: Reverse Engineering how Networks Learn Group Operations.** 6243-6267
- Jase Clarkson:  
**Distribution Free Prediction Sets for Node Classification.** 6268-6278
- Lee Cohen, Saeed Sharifi-Malvajerdi, Kevin Stangl, Ali Vakilian, Juba Ziani:  
**Sequential Strategic Screening.** 6279-6295
- David Cohen, Tal Shnitzer, Yuval Kluger, Ronen Talmon:  
**Few-Sample Feature Selection via Feature Manifold Learning.** 6296-6319
- Elijah Cole, Grant Van Horn, Christian Lange, Alexander Shepard, Patrick Leary, Pietro Perona, Scott Loarie, Oisin Mac Aodha:  
**Spatial Implicit Neural Representations for Global-Scale Species Mapping.** 6320-6342
- Andrea Coletta, Svitlana Vytenko, Tucker Balch:  
**K-SHAP: Policy Clustering Algorithm for Anonymous Multi-Agent State-Action Pairs.** 6343-6363
- Armand Comas Massague, Yilun Du, Christian Fernandez Lopez, Sandesh Ghimire, Mario Sznaier, Joshua B. Tenenbaum, Octavia I. Camps:  
**Inferring Relational Potentials in Interacting Systems.** 6364-6383
- Bethany Connolly, Kim Moore, Tobias Schwedes, Alexander Adam, Gary Willis, Ilya Feige, Christopher Frye:  
**Task-specific experimental design for treatment effect estimation.** 6384-6401
- Elisabetta Cornacchia, Elchanan Mossel:  
**A Mathematical Model for Curriculum Learning for Parities.** 6402-6423
- Ian Connick Covert, Wei Qiu, Mingyu Lu, Nayoon Kim, Nathan J. White, Su-In Lee:  
**Learning to Maximize Mutual Information for Dynamic Feature Selection.** 6424-6447
- Jingyi Cui, Weiran Huang, Yifei Wang, Yisen Wang:  
**Rethinking Weak Supervision in Helping Contrastive Learning.** 6448-6467
- Hugo Cui, Florent Krzakala , Lenka Zdeborová:  
**Bayes-optimal Learning of Deep Random Networks of Extensive-width.** 6468-6521
- Junbiao Cui, Jianqing Liang, Qin Yue, Jiye Liang:  
**A General Representation Learning Framework with Generalization Performance Guarantees.** 6522-6544
- Yuning Cui, Wenqi Ren, Sining Yang, Xiaochun Cao, Alois Knoll:  
**IRNeXt: Rethinking Convolutional Network Design for Image Restoration.** 6545-6564
- Justin Cui, Ruochen Wang, Si Si, Cho-Jui Hsieh:  
**Scaling Up Dataset Distillation to ImageNet-1K with Constant Memory.** 6565-6590
- Yiming Cui, Linjie Yang, Haichao Yu:  
**Learning Dynamic Query Combinations for Transformer-based Object Detection and Segmentation.** 6591-6602
- Alicia Curth, Alihan Hüyük, Mihaela van der Schaar:  
**Adaptive Identification of Populations with Treatment Benefit in Clinical Trials: Machine Learning Challenges and Solutions.** 6603-6622
- Alicia Curth, Mihaela van der Schaar:  
**In Search of Insights, Not Magic Bullets: Towards Demystification of the Model Selection Dilemma in Heterogeneous Treatment Effect Estimation.** 6623-6642
- Ashok Cutkosky, Harsh Mehta, Francesco Orabona:  
**Optimal Stochastic Non-smooth Non-convex Optimization through Online-to-Non-convex Conversion.** 6643-6670



- Marco Cuturi, Michal Klein, Pierre Ablin:  
**Monge, Bregman and Occam: Interpretable Optimal Transport in High-Dimensions with Feature-Sparse Maps.** 6671-6682
- Edwige Cyffers, Aurélien Bellet, Debabrota Basu:  
**From Noisy Fixed-Point Iterations to Private ADMM for Centralized and Federated Learning.** 6683-6711
- Yanbo Dai, Songze Li:  
**Chameleon: Adapting to Peer Images for Planting Durable Backdoors in Federated Learning.** 6712-6725
- Yan Dai, Haipeng Luo, Chen-Yu Wei, Julian Zimmert:  
**Refined Regret for Adversarial MDPs with Linear Function Approximation.** 6726-6759
- Sihui Dai, Saeed Mahloujifar, Chong Xiang, Vikash Sehwag, Pin-Yu Chen, Prateek Mittal:  
**MultiRobustBench: Benchmarking Robustness Against Multiple Attacks.** 6760-6785
- Rui Dai, Yonggang Zhang, Zhen Fang, Bo Han, Xinmei Tian:  
**Moderately Distributional Exploration for Domain Generalization.** 6786-6817
- Brett Daley, Martha White, Christopher Amato, Marlos C. Machado:  
**Trajectory-Aware Eligibility Traces for Off-Policy Reinforcement Learning.** 6818-6835
- Hadi Daneshmand, Jason D. Lee, Chi Jin:  
**Efficient displacement convex optimization with particle gradient descent.** 6836-6854
- Ronghao Dang, Lu Chen, Liuyi Wang, Zongtao He, Chengju Liu, Qijun Chen:  
**Multiple Thinking Achieving Meta-Ability Decoupling for Object Navigation.** 6855-6872
- Hien Dang, Tho Tran Huu, Stanley J. Osher, Hung Tran-The, Nhat Ho, Tan Minh Nguyen:  
**Neural Collapse in Deep Linear Networks: From Balanced to Imbalanced Data.** 6873-6947
- Christoph Dann, Yishay Mansour, Mehryar Mohri:  
**Reinforcement Learning Can Be More Efficient with Multiple Rewards.** 6948-6967
- Christoph Dann, Chen-Yu Wei, Julian Zimmert:  
**Best of Both Worlds Policy Optimization.** 6968-7008
- Ayan Das, Stathi Fotiadis, Anil Batra, Farhang Nabieei, Fengting Liao, Sattar Vakili, Da-Shan Shiu, Alberto Bernacchia:  
**Image generation with shortest path diffusion.** 7009-7024
- Abhimanyu Das, Ayush Jain, Weihao Kong, Rajat Sen:  
**Efficient List-Decodable Regression using Batches.** 7025-7065
- Rudrajit Das, Satyen Kale, Zheng Xu, Tong Zhang, Sujay Sanghavi:  
**Beyond Uniform Lipschitz Condition in Differentially Private Optimization.** 7066-7101
- Rudrajit Das, Sujay Sanghavi:  
**Understanding Self-Distillation in the Presence of Label Noise.** 7102-7140
- Shounak Datta, Sankha Subhra Mullick, Anish Chakrabarty, Swagatam Das:  
**Interval Bound Interpolation for Few-shot Learning with Few Tasks.** 7141-7166
- Samuel Daulton, Maximilian Balandat, Eytan Bakshy:  
**Hypervolume Knowledge Gradient: A Lookahead Approach for Multi-Objective Bayesian Optimization with Partial Information.** 7167-7204
- Sami Davies, Benjamin Moseley, Heather Newman:  
**Fast Combinatorial Algorithms for Min Max Correlation Clustering.** 7205-7230
- Sami Davies, Benjamin Moseley, Sergei Vassilvitskii, Yuyan Wang:  
**Predictive Flows for Faster Ford-Fulkerson.** 7231-7248
- Thomas Davies, Zhengchao Wan, Rubén J. Sánchez-García:  
**The Persistent Laplacian for Data Science: Evaluating Higher-Order Persistent Spectral Representations of Data.** 7249-7263
- Arka Daw, Jie Bu, Sifan Wang, Paris Perdikaris, Anuj Karpatne:  
**Mitigating Propagation Failures in Physics-informed Neural Networks using Retain-Resample-Release (R3) Sampling.** 7264-7302
- Hassan Dbouk, Naresh R. Shanbhag:  
**On the Robustness of Randomized Ensembles to Adversarial Perturbations.** 7303-7328
- Michiel de Jong, Yury Zemlyanskiy, Nicholas FitzGerald, Joshua Ainslie, Sumit Sanghai, Fei Sha, William W. Cohen:  
**Pre-computed memory or on-the-fly encoding? A hybrid approach to retrieval augmentation makes the most of your compute.** 7329-7342
- Antonio Henrique de Oliveira Fonseca, Emanuele Zappala, Josue Ortega Caro, David van Dijk :  
**Continuous Spatiotemporal Transformer.** 7343-7365
- Ashwin De Silva, Rahul Ramesh, Carey E. Priebe, Pratik Chaudhari, Joshua T. Vogelstein:  
**The Value of Out-of-Distribution Data.** 7366-7389
- Fabio De Sousa Ribeiro, Tian Xia, Miguel Monteiro, Nick Pawlowski, Ben Glocker:  
**High Fidelity Image Counterfactuals with Probabilistic Causal Models.** 7390-7425



- Antoine Dedieu, Guangyao Zhou, Dileep George, Miguel Lázaro-Gredilla:  
**Learning Noisy OR Bayesian Networks with Max-Product Belief Propagation.** 7426-7448
- Aaron Defazio, Konstantin Mishchenko:  
**Learning-Rate-Free Learning by D-Adaptation.** 7449-7479
- Mostafa Dehghani, Josip Djolonga, Basil Mustafa, Piotr Padlewski, Jonathan Heek, Justin Gilmer, Andreas Peter Steiner, Mathilde Caron, Robert Geirhos, Ibrahim Alabdulmohsin, Rodolphe Jenatton, Lucas Beyer, Michael Tschannen, Anurag Arnab, Xiao Wang, Carlos Riquelme Ruiz, Matthias Minderer, Joan Puigcerver, Utku Evci, Manoj Kumar, Sjoerd van Steenkiste, Gamaleldin Fathy Elsayed, Aravindh Mahendran, Fisher Yu, Avital Oliver, Fantine Huot, Jasmijn Bastings, Mark Collier, Alexey A. Gritsenko, Vighnesh Birodkar, Cristina Nader Vasconcelos, Yi Tay, Thomas Mensink, Alexander Kolesnikov, Filip Pavetic, Dustin Tran, Thomas Kipf, Mario Lucic, Xiaohua Zhai, Daniel Keysers, Jeremiah J. Harmsen, Neil Houlsby:  
**Scaling Vision Transformers to 22 Billion Parameters.** 7480-7512
- Blaise Delattre, Quentin Barthélemy, Alexandre Araujo, Alexandre Allauzen:  
**Efficient Bound of Lipschitz Constant for Convolutional Layers by Gram Iteration.** 7513-7532
- Emir Demirovic, Emmanuel Hebrard, Louis Jean:  
**Blossom: an Anytime Algorithm for Computing Optimal Decision Trees.** 7533-7562
- Chang Deng, Kevin Bello, Bryon Aragam, Pradeep Kumar Ravikumar:  
**Optimizing NOTEARS Objectives via Topological Swaps.** 7563-7595
- Danruo Deng, Guangyong Chen, Yang Yu, Furui Liu, Pheng-Ann Heng:  
**Uncertainty Estimation by Fisher Information-based Evidential Deep Learning.** 7596-7616
- Yuan Deng, Negin Golrezaei, Patrick Jaillet, Jason Cheuk Nam Liang, Vahab Mirrokni:  
**Multi-channel Autobidding with Budget and ROI Constraints.** 7617-7644
- Shikuang Deng, Hao Lin, Yuhang Li, Shi Gu:  
**Surrogate Module Learning: Reduce the Gradient Error Accumulation in Training Spiking Neural Networks.** 7645-7657
- Weijian Deng, Yumin Suh, Stephen Gould, Liang Zheng:  
**Confidence and Dispersity Speak: Characterizing Prediction Matrix for Unsupervised Accuracy Estimation.** 7658-7674
- Ailin Deng, Miao Xiong, Bryan Hooi:  
**Great Models Think Alike: Improving Model Reliability via Inter-Model Latent Agreement.** 7675-7693
- Karan Desai, Maximilian Nickel, Tanmay Rajpurohit, Justin Johnson, Shanmukha Ramakrishna Vedantam:  
**Hyperbolic Image-text Representations.** 7694-7731
- Aditya Desai, Keren Zhou , Anshumali Shrivastava:  
**Hardware-Aware Compression with Random Operation Access Specific Tile (ROAST) Hashing.** 7732-7749
- Tim Dettmers, Luke Zettlemoyer:  
**The case for 4-bit precision: k-bit Inference Scaling Laws.** 7750-7774
- Siddartha Devic, David Kempe, Vatsal Sharan, Aleksandra Korolova:  
**Fairness in Matching under Uncertainty.** 7775-7794
- Nikita Dhawan, Sicong Huang, Juhan Bae, Roger Baker Grosse:  
**Efficient Parametric Approximations of Neural Network Function Space Distance.** 7795-7812
- Victor Dheur, Souhaib Ben Taieb:  
**A Large-Scale Study of Probabilistic Calibration in Neural Network Regression.** 7813-7836
- Qiwei Di, Jiafan He, Dongruo Zhou, Quanquan Gu:  
**Nearly Minimax Optimal Regret for Learning Linear Mixture Stochastic Shortest Path.** 7837-7864
- Francesco Di Giovanni, Lorenzo Giusti, Federico Barbero, Giulia Luise, Pietro Lio, Michael M. Bronstein:  
**On Over-Squashing in Message Passing Neural Networks: The Impact of Width, Depth, and Topology.** 7865-7885
- Ilias Diakonikolas, Daniel Kane, Ankit Pensia, Thanasis Pittas:  
**Nearly-Linear Time and Streaming Algorithms for Outlier-Robust PCA.** 7886-7921
- Ilias Diakonikolas, Daniel Kane, Lisheng Ren:  
**Near-Optimal Cryptographic Hardness of Agnostically Learning Halfspaces and ReLU Regression under Gaussian Marginals.** 7922-7938
- Nathaniel Lee Diamant, Alex M. Tseng, Kangway V. Chuang, Tommaso Biancalani, Gabriele Scalia:  
**Improving Graph Generation by Restricting Graph Bandwidth.** 7939-7959
- Michael Ziyang Diao, Krishna Balasubramanian, Sinho Chewi, Adil Salim:  
**Forward-Backward Gaussian Variational Inference via JKO in the Bures-Wasserstein Space.** 7960-7991
- Travis Dick, Alex Kulesza, Ziteng Sun, Ananda Theertha Suresh:  
**Subset-Based Instance Optimality in Private Estimation.** 7992-8014
- Nikolaos Dimitriadis, Pascal Frossard, François Fleuret:  
**Pareto Manifold Learning: Tackling multiple tasks via ensembles of single-task models.** 8015-8052
- Wenhao Ding, Tong Che, Ding Zhao, Marco Pavone:  
**Bayesian Reparameterization of Reward-Conditioned Reinforcement Learning with Energy-based Models.** 8053-8066



- Lisang Ding, Kexin Jin, Bicheng Ying, Kun Yuan, Wotao Yin:  
**DSGD-CECA: Decentralized SGD with Communication-Optimal Exact Consensus Algorithm.** 8067-8089
- Zheng Ding, Jieke Wang, Zhuowen Tu:  
**Open-Vocabulary Universal Image Segmentation with MaskCLIP.** 8090-8102
- Ziluo Ding, Wanpeng Zhang, Junpeng Yue, Xiangjun Wang, Tiejun Huang, Zongqing Lu:  
**Entity Divider with Language Grounding in Multi-Agent Reinforcement Learning.** 8103-8119
- AnhDung Dinh, Daochang Liu, Chang Xu:  
**PixelAsParam: A Gradient View on Diffusion Sampling with Guidance.** 8120-8137
- Nikita Doikov, El Mahdi Chayti, Martin Jaggi:  
**Second-Order Optimization with Lazy Hessians.** 8138-8161
- Nikita Doikov, Anton Rodomanov:  
**Polynomial Preconditioning for Gradient Methods.** 8162-8187
- Ricardo Dominguez-Olmedo, Amir-Hossein Karimi, Georgios Arvanitidis, Bernhard Schölkopf:  
**On Data Manifolds Entailed by Structural Causal Models.** 8188-8201
- Mingze Dong, Yuval Kluger:  
**Towards Understanding and Reducing Graph Structural Noise for GNNs.** 8202-8226
- Peiyan Dong, Zhenglun Kong, Xin Meng, Peng Zhang, Hao Tang, Yanzhi Wang, Chih-Hsien Chou:  
**SpeedDETR: Speed-aware Transformers for End-to-end Object Detection.** 8227-8243
- Chengyu Dong, Liyuan Liu, Hao Cheng, Jingbo Shang, Jianfeng Gao, Xiaodong Liu:  
**Understand and Modularize Generator Optimization in ELECTRA-style Pretraining.** 8244-8259
- Ruijiang Dong, Feng Liu, Haoang Chi, Tongliang Liu, Mingming Gong, Gang Niu, Masashi Sugiyama, Bo Han:  
**Diversity-enhancing Generative Network for Few-shot Hypothesis Adaptation.** 8260-8275
- Juncheng Dong, Weibin Mo, Zhengling Qi, Cong Shi, Ethan X. Fang, Vahid Tarokh:  
**PASTA: Pessimistic Assortment Optimization.** 8276-8295
- Yijun Dong, Yuege Xie, Rachel A. Ward:  
**Adaptively Weighted Data Augmentation Consistency Regularization for Robust Optimization under Concept Shift.** 8296-8316
- Jialin Dong, Lin Yang:  
**Does Sparsity Help in Learning Misspecified Linear Bandits?** 8317-8333
- Heng Dong, Junyu Zhang, Tonghan Wang, Chongjie Zhang:  
**Symmetry-Aware Robot Design with Structured Subgroups.** 8334-8355
- Ron Dorfman, Shay Vargaftik, Yaniv Ben-Itzhak, Kfir Yehuda Levy:  
**DoCoFL: Downlink Compression for Cross-Device Federated Learning.** 8356-8388
- Will Dorrell, Maria Yuffa, Peter E. Latham:  
**Meta-Learning the Inductive Bias of Simple Neural Circuits.** 8389-8402
- Vishwaraj Doshi, Jie Hu, Do Young Eun:  
**Self-Repellent Random Walks on General Graphs - Achieving Minimal Sampling Variance via Nonlinear Markov Chains.** 8403-8423
- Matthew Dowling, Yuan Zhao, Il Memming Park:  
**Linear Time GPs for Inferring Latent Trajectories from Neural Spike Trains.** 8424-8448
- Felix Draxler, Lars Kühmichel, Armand Rousselot, Jens Müller, Christoph Schnörr, Ullrich Köthe:  
**On the Convergence Rate of Gaussianization with Random Rotations.** 8449-8468
- Danny Driess, Fei Xia, Mehdi S. M. Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, Wenlong Huang, Yevgen Chebotar, Pierre Serbanet, Daniel Duckworth, Sergey Levine, Vincent Vanhoucke, Karol Hausman, Marc Toussaint, Klaus Greff, Andy Zeng, Igor Mordatch, Pete Florence:  
**PaLM-E: An Embodied Multimodal Language Model.** 8469-8488
- Yilun Du, Conor Durkan, Robin Strudel, Joshua B. Tenenbaum, Sander Dieleman, Rob Fergus, Jascha Sohl-Dickstein, Arnaud Doucet, Will Sussman Grathwohl:  
**Reduce, Reuse, Recycle: Compositional Generation with Energy-Based Diffusion Models and MCMC.** 8489-8510
- Yihan Du, Longbo Huang, Wen Sun:  
**Multi-task Representation Learning for Pure Exploration in Linear Bandits.** 8511-8564
- Chao Du, Tianbo Li, Tianyu Pang, Shuicheng Yan, Min Lin:  
**Nonparametric Generative Modeling with Conditional Sliced-Wasserstein Flows.** 8565-8584
- Jin-Hong Du, Pratik Patil, Arun K. Kuchibhotla:  
**Subsample Ridge Ensembles: Equivalences and Generalized Cross-Validation.** 8585-8631
- Chenzhuang Du, Jiaye Teng, Tingle Li, Yichen Liu, Tianyuan Yuan, Yue Wang, Yang Yuan, Hang Zhao:  
**On Uni-Modal Feature Learning in Supervised Multi-Modal Learning.** 8632-8656
- Yuqing Du, Olivia Watkins, Zihan Wang, Cédric Colas, Trevor Darrell, Pieter Abbeel, Abhishek Gupta, Jacob Andreas:  
**Guiding Pretraining in Reinforcement Learning with Large Language Models.** 8657-8677



- Weitao Du, He Zhang, Tao Yang, Yuanqi Du:  
**A Flexible Diffusion Model.** 8678-8696
- Chenguang Duan, Yuling Jiao, Lican Kang, Xiliang Lu, Jerry Zhijian Yang:  
**Fast Excess Risk Rates via Offset Rademacher Complexity.** 8697-8716
- Jinhao Duan, Fei Kong, Shiqi Wang, Xiaoshuang Shi, Kaidi Xu:  
**Are Diffusion Models Vulnerable to Membership Inference Attacks?** 8717-8730
- Zhibin Duan, Xinyang Liu, Yudi Su, Yishi Xu, Bo Chen, Mingyuan Zhou:  
**Bayesian Progressive Deep Topic Model with Knowledge Informed Textual Data Coarsening Process.** 8731-8746
- Zhijian Duan, Yunxuan Ma, Xiaotie Deng:  
**Are Equivariant Equilibrium Approximators Beneficial?** 8747-8778
- Yann Dubois, Tatsunori Hashimoto, Percy Liang:  
**Evaluating Self-Supervised Learning via Risk Decomposition.** 8779-8820
- Paul Duetting, Federico Fusco, Silvio Lattanzi, Ashkan Norouzi-Fard, Morteza Zadimoghaddam:  
**Fully Dynamic Submodular Maximization over Matroids.** 8821-8835
- Paul Duetting, Guru Guruganesh, Jon Schneider, Joshua Ruizhi Wang:  
**Optimal No-Regret Learning for One-Sided Lipschitz Functions.** 8836-8850
- Benoit Dufumier, Carlo Alberto Barbano, Robin Louiset, Edouard Duchesnay, Pietro Gori:  
**Integrating Prior Knowledge in Contrastive Learning with Kernel.** 8851-8878
- Owen M. Dugan, Peter Y. Lu, Rumen Dangovski, Di Luo, Marin Soljacic:  
**Q-Flow: Generative Modeling for Differential Equations of Open Quantum Dynamics with Normalizing Flows.** 8879-8901
- Lyndon R. Duong , David Lipshutz, David J. Heeger, Dmitri B. Chklovskii, Eero P. Simoncelli:  
**Adaptive Whitening in Neural Populations with Gain-modulating Interneurons.** 8902-8921
- Benjamin Dupuis, George Deligiannidis, Umut Simsekli:  
**Generalization Bounds using Data-Dependent Fractal Dimensions.** 8922-8968
- Arkadiy Dushatskiy, Alexander Chebykin, Tanja Alderliesten, Peter A. N. Bosman:  
**Multi-Objective Population Based Training.** 8969-8989
- Vincent Dutordoir, Alan Saul, Zoubin Ghahramani, Fergus Simpson:  
**Neural Diffusion Processes.** 8990-9012
- Alexandre Duval, Victor Schmidt, Alex Hernández-García, Santiago Miret, Fragkiskos D. Malliaros, Yoshua Bengio, David Rolnick:  
**FAENet: Frame Averaging Equivariant GNN for Materials Modeling.** 9013-9033
- Javier E. Santos, Zachary R. Fox, Nicholas Lubbers, Yen Ting Lin:  
**Blackout Diffusion: Generative Diffusion Models in Discrete-State Spaces.** 9034-9059
- Eduard Eiben, Robert Ganian, Iyad A. Kanj, Sebastian Ordyniak, Stefan Szeider:  
**The Computational Complexity of Concise Hypersphere Classification.** 9060-9070
- Floor Eijkelboom, Rob Hesselink, Erik J. Bekkers:  
**E(n) Equivariant Message Passing Simplicial Networks.** 9071-9081
- Itay Eilat, Nir Rosenfeld:  
**Performative Recommendation: Diversifying Content via Strategic Incentives.** 9082-9103
- Theresa Eimer, Marius Lindauer , Roberta Raileanu:  
**Hyperparameters in Reinforcement Learning and How To Tune Them.** 9104-9149
- Marwa El Halabi, Federico Fusco, Ashkan Norouzi-Fard, Jakab Tardos, Jakub Tarnawski:  
**Fairness in Streaming Submodular Maximization over a Matroid Constraint.** 9150-9171
- Marwa El Halabi, George Orfanides, Tim Hoheisel:  
**Difference of submodular minimization via DC programming.** 9172-9201
- Moshe Eliasof, Fabrizio Frasca, Beatrice Bevilacqua, Eran Treister, Gal Chechik, Haggai Maron:  
**Graph Positional Encoding via Random Feature Propagation.** 9202-9223
- Moshe Eliasof, Lars Ruthotto , Eran Treister:  
**Improving Graph Neural Networks with Learnable Propagation Operators.** 9224-9245
- Dor Elimelech, Wasim Huleihel:  
**Phase Transitions in the Detection of Correlated Databases.** 9246-9266
- Yury Elkin, Vitaliy Kurlin:  
**A new near-linear time algorithm for k-nearest neighbor search using a compressed cover tree.** 9267-9311
- Mark Endo, Joy Hsu, Jiaman Li, Jiajun Wu:  
**Motion Question Answering via Modular Motion Programs.** 9312-9328
- Joseph Enguehard:  
**Learning Perturbations to Explain Time Series Predictions.** 9329-9342
- Liad Erez, Tal Lancewicki, Uri Sherman, Tomer Koren, Yishay Mansour:  
**Regret Minimization and Convergence to Equilibria in General-sum Markov Games.** 9343-9373



- Emmanuel Esposito, Saeed Masoudian, Hao Qiu, Dirk van der Hoeven, Nicolò Cesa-Bianchi, Yevgeny Seldin:  
**Delayed Bandits: When Do Intermediate Observations Help?** 9374-9395
- Carlos Esteves, Jean-Jacques E. Slotine, Ameesh Makadia:  
**Scaling Spherical CNNs.** 9396-9411
- Mathieu Even:  
**Stochastic Gradient Descent under Markovian Sampling Schemes.** 9412-9439
- Itay Evron, Edward Moroshko, Gon Buzaglo, Maroun Khriesh, Badea Marjeh, Nathan Srebro, Daniel Soudry:  
**Continual Learning in Linear Classification on Separable Data.** 9440-9484
- Benjamin Eysenbach, Matthieu Geist, Sergey Levine, Ruslan Salakhutdinov:  
**A Connection between One-Step RL and Critic Regularization in Reinforcement Learning.** 9485-9507
- Lukas Faber, Roger Wattenhofer:  
**Neural Status Registers.** 9508-9522
- Matthew Fahrbach, Adel Javanmard, Vahab Mirrokni, Pratik Worah:  
**Learning Rate Schedules in the Presence of Distribution Shift.** 9523-9546
- Gregory Faleto, Jacob Bien:  
**Predicting Rare Events by Shrinking Towards Proportional Odds.** 9547-9602
- Xuhui Fan, Edwin V. Bonilla, Terence J. O'Kane, Scott A. Sisson:  
**Free-Form Variational Inference for Gaussian Process State-Space Models.** 9603-9622
- Ying Fan, Kangwook Lee:  
**Optimizing DDPM Sampling with Shortcut Fine-Tuning.** 9623-9639
- Chenglin Fan, Ping Li, Xiaoyun Li:  
**LSDS++ : Dual Sampling for Accelerated k-means++.** 9640-9649
- Zhenan Fan, Xinglu Wang, Oleksandr Yakovenko, Abdullah Ali Sivas, Owen Ren, Yong Zhang, Zirui Zhou:  
**Smart Initial Basis Selection for Linear Programs.** 9650-9664
- Vladimir Fanaskov, Tianchi Yu, Alexander Rudikov, Ivan V. Oseledets:  
**General Covariance Data Augmentation for Neural PDE Solvers.** 9665-9688
- Ora Nova Fandina, Mikael Møller Høgsgaard, Kasper Green Larsen:  
**The Fast Johnson-Lindenstrauss Transform Is Even Faster.** 9689-9715
- Guanhua Fang, Ping Li:  
**Regression with Label Permutation in Generalized Linear Model.** 9716-9760
- Sadegh Farhadkhani, Rachid Guerraoui, Nirupam Gupta, Lê-Nguyễn Hoang, Rafael Pinot, John Stephan:  
**Robust Collaborative Learning with Linear Gradient Overhead.** 9761-9813
- Oluwadamilola Fasina, Guillaume Huguet, Alexander Tong, Yanlei Zhang, Guy Wolf, Maximilian Nickel, Ian Adelstein, Smita Krishnaswamy:  
**Neural FIM for learning Fisher information metrics from point cloud data.** 9814-9826
- Ilyas Fatkhullin, Anas Barakat, Anastasia Kireeva, Niao He:  
**Stochastic Policy Gradient Methods: Improved Sample Complexity for Fisher-non-degenerate Policies.** 9827-9869
- Jonathan Feldstein, Modestas Jurcius, Efthymia Tsamoura:  
**Parallel Neurosymbolic Integration with Concordia.** 9870-9885
- Mattie Fellows, Matthew J. A. Smith, Shimon Whiteson:  
**Why Target Networks Stabilise Temporal Difference Methods.** 9886-9909
- Dieqiao Feng, Yuanqi Du, Carla P. Gomes, Bart Selman:  
**Weighted Sampling without Replacement for Deep Top-k Classification.** 9910-9920
- Zhe Feng, Christopher Liaw, Zixin Zhou:  
**Improved Online Learning Algorithms for CTR Prediction in Ad Auctions.** 9921-9937
- Shikun Feng, Yuyan Ni, Yanyan Lan, Zhi-Ming Ma, Wei-Ying Ma:  
**Fractional Denoising for 3D Molecular Pre-training.** 9938-9961
- Ying Feng, David P. Woodruff:  
**Improved Algorithms for White-Box Adversarial Streams.** 9962-9975
- Songtao Feng, Ming Yin, Ruiquan Huang, Yu-Xiang Wang, Jing Yang, Yingbin Liang:  
**Non-stationary Reinforcement Learning under General Function Approximation.** 9976-10007
- Vasilii Feofanov, Malik Tiomoko, Aladin Virmaux:  
**Random Matrix Analysis to Balance between Supervised and Unsupervised Learning under the Low Density Separation Assumption.** 10008-10033
- Aaron M. Ferber, Taoan Huang, Daochen Zha, Martin Schubert, Benoit Steiner, Bistra Dilkina, Yuandong Tian:  
**SurCo: Learning Linear SURrogates for COmbinatorial Nonlinear Optimization Problems.** 10034-10052
- Patrick Fernandes, Behrooz Ghorbani, Xavier Garcia, Markus Freitag, Orhan Firat:  
**Scaling Laws for Multilingual Neural Machine Translation.** 10053-10071
- Hendrik Fichtenberger, Monika Henzinger, Jalaj Upadhyay:  
**Constant Matters: Fine-grained Error Bound on Differentially Private Continual Observation.** 10072-10092



- Côme Fiegel, Pierre Ménard, Tadashi Kozuno, Rémi Munos, Vianney Perchet, Michal Valko:  
**Adapting to game trees in zero-sum imperfect information games.** 10093-10135
- Marc Anton Finzi, Anudhyan Boral, Andrew Gordon Wilson, Fei Sha, Leonardo Zepeda-Núñez:  
**User-defined Event Sampling and Uncertainty Quantification in Diffusion Models for Physical Dynamical Systems.** 10136-10152
- Alessandro Fontanella, Antreas Antoniou, Wenwen Li, Joanna M. Wardlaw, Grant Mair, Emanuele Trucco, Amos J. Storkey:  
**ACAT: Adversarial Counterfactual Attention for Classification and Detection in Medical Imaging.** 10153-10169
- Alexandre Forel, Axel Parmentier, Thibaut Vidal:  
**Explainable Data-Driven Optimization: From Context to Decision and Back Again.** 10170-10187
- Dylan J. Foster, Noah Golowich, Sham M. Kakade:  
**Hardness of Independent Learning and Sparse Equilibrium Computation in Markov Games.** 10188-10221
- Stathi Fotiadis, Mario Lino Valencia, Shunlong Hu, Stef Garasto, Chris D. Cantwell, Anil Anthony Bharath:  
**Disentangled Generative Models for Robust Prediction of System Dynamics.** 10222-10248
- Louis Fournier, Stéphane Rivaud, Eugene Belilovsky, Michael Eickenberg, Edouard Oyallon:  
**Can Forward Gradient Match Backpropagation?** 10249-10264
- Ayoub Foussoul, Vineet Goyal, Orestis Papadigenopoulos, Assaf Zeevi:  
**Last Switch Dependent Bandits with Monotone Payoff Functions.** 10265-10284
- Emanuele Francazi, Marco Baity-Jesi, Aurélien Lucchi:  
**A Theoretical Analysis of the Learning Dynamics under Class Imbalance.** 10285-10322
- Elias Frantar, Dan Alistarh:  
**SparseGPT: Massive Language Models Can be Accurately Pruned in One-Shot.** 10323-10337
- Benjamin Freed, Siddarth Venkatraman, Guillaume Adrien Sartoretti, Jeff Schneider, Howie Choset:  
**Learning Temporally AbstractWorld Models without Online Experimentation.** 10338-10356
- Gideon Joseph Freund, Elad Sarafian, Sarit Kraus:  
**A Coupled Flow Approach to Imitation Learning.** 10357-10372
- Daniel Y. Fu, Elliot L. Epstein, Eric Nguyen, Armin W. Thomas, Michael Zhang, Tri Dao, Atri Rudra, Christopher Ré:  
**Simple Hardware-Efficient Long Convolutions for Sequence Modeling.** 10373-10391
- Yang Fu , Ishan Misra, Xiaolong Wang:  
**MonoNeRF: Learning Generalizable NeRFs from Monocular Videos without Camera Poses.** 10392-10404
- Yao Fu, Run Peng, Honglak Lee:  
**Go Beyond Imagination: Maximizing Episodic Reachability with World Models.** 10405-10420
- Yao Fu, Hao Peng, Litu Ou, Ashish Sabharwal, Tushar Khot:  
**Specializing Smaller Language Models towards Multi-Step Reasoning.** 10421-10430
- Qiang Fu, Dongchu Xu, Ashia Camage Wilson:  
**Accelerated Stochastic Optimization Methods under Quasar-convexity.** 10431-10460
- Haotian Fu, Shangqun Yu, Saket Tiwari, Michael Littman, George Konidaris:  
**Meta-learning Parameterized Skills.** 10461-10481
- Yonggan Fu, Ye Yuan, Souvik Kundu, Shang Wu, Shunyao Zhang, Yingyan Celine Lin:  
**NeRFool: Uncovering the Vulnerability of Generalizable Neural Radiance Fields against Adversarial Perturbations.** 10482-10493
- Daniel Furelos-Blanco, Mark Law, Anders Jonsson, Krysia Broda, Alessandra Russo:  
**Hierarchies of Reward Machines.** 10494-10541
- Advait Harshal Gadhikar, Sohom Mukherjee, Rebekka Burkholz:  
**Why Random Pruning Is All We Need to Start Sparse.** 10542-10570
- Quentin Gallouédec, Emmanuel Dellandréa:  
**Cell-Free Latent Go-Explore.** 10571-10586
- Daniele Gammelli, James Harrison, Kaidi Yang, Marco Pavone, Filipe Rodrigues, Francisco C. Pereira:  
**Graph Reinforcement Learning for Network Control via Bi-Level Optimization.** 10587-10610
- Arun Ganesh, Mahdi Haghifam, Milad Nasr, Sewoong Oh, Thomas Steinke, Om Thakkar, Abhradeep Guha Thakurta, Lun Wang:  
**Why Is Public Pretraining Necessary for Private Model Training?** 10611-10627
- Roy Ganz, Bahjat Kawar, Michael Elad:  
**Do Perceptually Aligned Gradients Imply Robustness?** 10628-10648
- Wenzhi Gao, Dongdong Ge , Chunlin Sun, Yinyu Ye:  
**Solving Linear Programs with Fast Online Learning Algorithms.** 10649-10675
- Yihang Gao , Yiqi Gu, Michael Ng:  
**Gradient Descent Finds the Global Optima of Two-Layer Physics-Informed Neural Networks.** 10676-10707
- Nicholas Gao, Stephan Günnemann:  
**Generalizing Neural Wave Functions.** 10708-10726



- Ruijiang Gao, Himabindu Lakkaraju:  
**On the Impact of Algorithmic Recourse on Social Segregation.** 10727-10743
- Rui Gao , Weiwei Liu:  
**DDGR: Continual Learning with Deep Diffusion-based Generative Replay.** 10744-10763
- Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, Graham Neubig:  
**PAL: Program-aided Language Models.** 10764-10799
- Irena Gao, Shiori Sagawa, Pang Wei Koh, Tatsunori Hashimoto, Percy Liang:  
**Out-of-Domain Robustness via Targeted Augmentations.** 10800-10834
- Leo Gao, John Schulman, Jacob Hilton:  
**Scaling Laws for Reward Model Overoptimization.** 10835-10866
- Xavier Garcia, Yamini Bansal, Colin Cherry, George F. Foster, Maxim Krikun, Melvin Johnson, Orhan Firat:  
**The Unreasonable Effectiveness of Few-shot Learning for Machine Translation.** 10867-10878
- Saurabh Garg, Nick Erickson, James Sharpnack, Alex Smola, Sivaraman Balakrishnan, Zachary Chase Lipton:  
**RLSbench: Domain Adaptation Under Relaxed Label Shift.** 10879-10928
- Quentin Garrido, Randall Balestrieri, Laurent Najman, Yann LeCun:  
**RankMe: Assessing the Downstream Performance of Pretrained Self-Supervised Representations by Their Rank.** 10929-10974
- Quentin Garrido, Laurent Najman, Yann LeCun:  
**Self-supervised learning of Split Invariant Equivariant representations.** 10975-10996
- Adrià Gascón, Peter Kairouz, Ziteng Sun, Ananda Theertha Suresh:  
**Federated Heavy Hitter Recovery under Linear Sketching.** 10997-11012
- Mudit Gaur, Vaneet Aggarwal, Mridul Agarwal:  
**On the Global Convergence of Fitted Q-Iteration with Two-layer Neural Network Parametrization.** 11013-11049
- Lin Ge, Jitao Wang, Chengchun Shi, Zhenke Wu, Rui Song:  
**A Reinforcement Learning Framework for Dynamic Mediation Analysis.** 11050-11097
- Tomas Geffner, George Papamakarios, Andriy Mnih:  
**Compositional Score Modeling for Simulation-Based Inference.** 11098-11116
- Jonas Geiping, Tom Goldstein:  
**Cramming: Training a Language Model on a single GPU in one day.** 11117-11143
- Simon Geisler, Yujia Li, Daniel J. Mankowitz, Ali Taylan Cemgil, Stephan Günnemann, Cosmin Paduraru:  
**Transformers Meet Directed Graphs.** 11144-11172
- Tim Genewein, Grégoire Delétang, Anian Ruoss, Li Kevin Wenliang, Elliot Catt, Vincent Dutordoir, Jordi Grau-Moya, Laurent Orseau, Marcus Hutter, Joel Veness:  
**Memory-Based Meta-Learning on Non-Stationary Distributions.** 11173-11195
- Chuqin Geng, Nham Le, Xiaojie Xu, Zhaoxue Wang, Arie Gurfinkel, Xujie Si:  
**Towards Reliable Neural Specifications.** 11196-11212
- Matthias Gerstgrasser, David C. Parkes:  
**Oracles & Followers: Stackelberg Equilibria in Deep Multi-Agent Reinforcement Learning.** 11213-11236
- Mehrdad Ghadiri, Matthew Fahrbach, Gang Fu, Vahab Mirrokni:  
**Approximately Optimal Core Shapes for Tensor Decompositions.** 11237-11254
- Salah Ghamizi, Jingfeng Zhang, Maxime Cordy, Mike Papadakis, Masashi Sugiyama, Yves Le Traon:  
**GAT: Guided Adversarial Training with Pareto-optimal Auxiliary Tasks.** 11255-11282
- Badih Ghazi, Pritish Kamath, Ravi Kumar, Pasin Manurangsi, Raghu Meka, Chiyuan Zhang:  
**On User-Level Private Convex Optimization.** 11283-11299
- Gaurav Rohit Ghosal, Amrith Setlur, Daniel S. Brown, Anca D. Dragan, Aditi Raghunathan:  
**Contextual Reliability: When Different Features Matter in Different Contexts.** 11300-11320
- Dibya Ghosh, Chethan Anand Bhateja, Sergey Levine:  
**Reinforcement Learning from Passive Data via Latent Intentions.** 11321-11339
- Atiyo Ghosh, Antonio Andrea Gentile, Mario Dagrada, Chul Lee, Seong-Hyok Sean Kim, Hyukgeun Cha, Yunjun Choi, Dongho Kim, Jeong-II Kye, Vincent Emanuel Elfving:  
**Harmonic Neural Networks.** 11340-11359
- Shantanu Ghosh, Ke Yu, Forough Arabshahi, Kayhan Batmanghelich:  
**Dividing and Conquering a BlackBox to a Mixture of Interpretable Models: Route, Interpret, Repeat.** 11360-11397
- Angeliki Giannou, Shashank Rajput, Jy-yong Sohn, Kangwook Lee, Jason D. Lee, Dimitris Papailiopoulos:  
**Looped Transformers as Programmable Computers.** 11398-11442
- Luca Giuliani, Eleonora Misino, Michele Lombardi:  
**Generalized Disparate Impact for Configurable Fairness Solutions in ML.** 11443-11458
- Ira Globus-Harris, Declan Harrison, Michael Kearns, Aaron Roth, Jessica Sorrell:  
**Multicalibration as Boosting for Regression.** 11459-11492



- Manuel Glöckler, Michael Deistler, Jakob H. Macke:  
**Adversarial robustness of amortized Bayesian inference.** 11493-11524
- Kevin Gmelin, Shikhar Bahl, Russell Mendonca, Deepak Pathak:  
**Efficient RL via Disentangled Environment and Agent Representations.** 11525-11545
- Dongyoung Go, Tomasz Korbak, Germán Kruszewski, Jos Rozen, Nahyeon Ryu, Marc Dymetman:  
**Aligning Language Models with Preferences through f-divergence Minimization.** 11546-11583
- Morgane Goibert, Clément Calauzènes, Ekhine Irurozki, Stéphan Cléménçon:  
**Robust Consensus in Ranking Data Analysis: Definitions, Properties and Computational Issues.** 11584-11597
- Weiyuan Gong, Scott Aaronson:  
**Learning Distributions over Quantum Measurement Outcomes.** 11598-11613
- Eduard Gorbunov, Adrien B. Taylor, Samuel Horváth, Gauthier Gidel:  
**Convergence of Proximal Point and Extragradient-Based Methods Beyond Monotonicity: the Case of Negative Comonotonicity.** 11614-11641
- Shirin Goshtasbpour, Victor Cohen, Fernando Pérez-Cruz:  
**Adaptive Annealed Importance Sampling with Constant Rate Progress.** 11642-11658
- Devon R. Graham, Kevin Leyton-Brown, Tim Roughgarden:  
**Formalizing Preferences Over Runtime Distributions.** 11659-11682
- Vincent Peter Grande, Michael T. Schaub:  
**Topological Point Cloud Clustering.** 11683-11697
- Louis Grenioux, Alain Oliviero Durmus, Eric Moulines, Marylou Gabrié:  
**On Sampling with Approximate Transport Maps.** 11698-11733
- J. Elisenda Grigsby, Kathryn Lindsey , David Rolnick:  
**Hidden Symmetries of ReLU Networks.** 11734-11760
- Kaja Gruntkowska, Alexander Tyurin, Peter Richtárik:  
**EF21-P and Friends: Improved Theoretical Communication Complexity for Distributed Optimization with Bidirectional Compression.** 11761-11807
- Jiatao Gu, Alex Trevithick, Kai-En Lin, Joshua M. Susskind, Christian Theobalt, Lingjie Liu, Ravi Ramamoorthi:  
**NerfDiff: Single-image View Synthesis with NeRF-guided Distillation from 3D-aware Diffusion.** 11808-11826
- Jiaqi Guan, Xiangxin Zhou, Yuwei Yang, Yu Bao, Jian Peng, Jianzhu Ma, Qiang Liu, Liang Wang, Quanquan Gu:  
**DecompDiff: Diffusion Models with Decomposed Priors for Structure-Based Drug Design.** 11827-11846
- Aritra Guha, Nhat Ho, XuanLong Nguyen:  
**On Excess Mass Behavior in Gaussian Mixture Models with Orlicz-Wasserstein Distances.** 11847-11870
- Etash Kumar Guha, Eugène Ndiaye, Xiaoming Huo:  
**Conformalization of Sparse Generalized Linear Models.** 11871-11887
- Chuan Guo, Kamalika Chaudhuri, Pierre Stock, Michael G. Rabbat:  
**Privacy-Aware Compression for Federated Learning Through Numerical Mechanism Design.** 11888-11904
- Yaming Guo, Kai Guo, Xiaofeng Cao, Tieru Wu, Yi Chang:  
**Out-of-Distribution Generalization of Federated Learning via Implicit Invariant Relationships.** 11905-11933
- Zhishuai Guo, Rong Jin, Jiebo Luo, Tianbao Yang:  
**FeDXL: Provable Federated Learning for Deep X-Risk Optimization.** 11934-11966
- Jiacheng Guo, Zihao Li, Huazheng Wang, Mengdi Wang, Zhuoran Yang, Xuezhou Zhang:  
**Provably Efficient Representation Learning with Tractable Planning in Low-Rank POMDP.** 11967-11997
- Chuan Guo, Alexandre Sablayrolles, Maziar Sanjabi:  
**Analyzing Privacy Leakage in Machine Learning via Multiple Hypothesis Testing: A Lesson From Fano.** 11998-12011
- Zhichun Guo, William Shiao, Shichang Zhang, Yozen Liu, Nitesh V. Chawla, Neil Shah, Tong Zhao:  
**Linkless Link Prediction via Relational Distillation.** 12012-12033
- Yongxin Guo, Xiaoying Tang, Tao Lin:  
**FedBR: Improving Federated Learning on Heterogeneous Data via Local Learning Bias Reduction.** 12034-12054
- Minghao Guo, Veronika Thost, Samuel W. Song, Adithya Balachandran, Payel Das, Jie Chen, Wojciech Matusik:  
**Hierarchical Grammar-Induced Geometry for Data-Efficient Molecular Property Prediction.** 12055-12076
- Yuhe Guo, Zhewei Wei:  
**Graph Neural Networks with Learnable and Optimal Polynomial Bases.** 12077-12097
- Daya Guo, Canwen Xu, Nan Duan, Jian Yin, Julian J. McAuley:  
**LongCoder: A Long-Range Pre-trained Language Model for Code Completion.** 12098-12107
- Xingzhuo Guo, Yuchen Zhang, Jianmin Wang, Mingsheng Long:  
**Estimating Heterogeneous Treatment Effects: Mutual Information Bounds and Learning Algorithms.** 12108-12121
- Lan-Zhe Guo, Zhi Zhou , Yufeng Li, Zhi-Hua Zhou:  
**Identifying Useful Learnwares for Heterogeneous Label Spaces.** 12122-12131
- Shivam Gupta, Jasper C. H. Lee, Eric Price:  
**High-dimensional Location Estimation via Norm Concentration for Subgamma Vectors.** 12132-12164



- Shubham Gupta, Sahil Manchanda, Sayan Ranu, Srikanta J. Bedathur:  
**GRAFENNE: Learning on Graphs with Heterogeneous and Dynamic Feature Sets.** 12165-12181
- Chirag Gupta, Aaditya Ramdas:  
**Online Platt Scaling with Calibeating.** 12182-12204
- Naresh Kumar Gurulingan, Bahram Zonooz, Elahe Arani:  
**Multi-Task Structural Learning using Local Task Similarity induced Neuron Creation and Removal.** 12205-12223
- Florentin Guth, Etienne Lempereur, Joan Bruna, Stéphane Mallat:  
**Conditionally Strongly Log-Concave Generative Models.** 12224-12251
- Benjamin Gutteridge, Xiaowen Dong, Michael M. Bronstein, Francesco Di Giovanni:  
**DRew: Dynamically Rewired Message Passing with Delay.** 12252-12267
- Marie Guyomard, Susana Barbosa, Lionel Fillatre:  
**Kernel Logistic Regression Approximation of an Understandable ReLU Neural Network.** 12268-12291
- Soroush H. Zargarbashi, Simone Antonelli, Aleksandar Bojchevski:  
**Conformal Prediction Sets for Graph Neural Networks.** 12292-12318
- Seungwoong Ha, Hawoong Jeong:  
**Social learning spontaneously emerges by searching optimal heuristics with deep reinforcement learning.** 12319-12338
- Daniel Haider, Martin Ehler, Péter Balázs:  
**Convex Geometry of ReLU-layers, Injectivity on the Ball and Local Reconstruction.** 12339-12350
- Faisal Hamman, Erfan Noorani, Saumitra Mishra, Daniele Magazzeni, Sanghamitra Dutta:  
**Robust Counterfactual Explanations for Neural Networks With Probabilistic Guarantees.** 12351-12367
- Boran Han:  
**Wrapped Cauchy Distributed Angular Softmax for Long-Tailed Visual Recognition.** 12368-12388
- Hyeongrok Han, Siwon Kim, Hyun-Soo Choi, Sungroh Yoon:  
**On the Impact of Knowledge Distillation for Model Interpretability.** 12389-12410
- Haoyu Han, Xiaorui Liu, Haitao Mao, MohamadAli Torkamani, Feng Shi, Victor Lee, Jiliang Tang:  
**Alternately Optimized Graph Neural Networks.** 12411-12429
- Yena Han, Tomaso A. Poggio, Brian Cheung:  
**System Identification of Neural Systems: If We Got It Right, Would We Know?** 12430-12444
- Jonas Berg Hansen, Filippo Maria Bianchi:  
**Total Variation Graph Neural Networks.** 12445-12468
- Derek Hansen, Danielle C. Maddix, Shima Alizadeh, Gaurav Gupta, Michael W. Mahoney:  
**Learning Physical Models that Can Respect Conservation Laws.** 12469-12510
- Nicklas Hansen, Zhecheng Yuan, Yanjie Ze, Tongzhou Mu, Aravind Rajeswaran, Hao Su, Huazhe Xu, Xiaolong Wang:  
**On Pre-Training for Visuo-Motor Control: Revisiting a Learning-from-Scratch Baseline.** 12511-12526
- Botao Hao, Rahul Jain, Tor Lattimore, Benjamin Van Roy, Zheng Wen:  
**Leveraging Demonstrations to Improve Online Learning: Quality Matters.** 12527-12545
- Xiaoran Hao, Patrick Shafto:  
**Coupled Variational Autoencoder.** 12546-12555
- Zhongkai Hao, Zhengyi Wang, Hang Su, Chengyang Ying, Yinpeng Dong, Songming Liu, Ze Cheng, Jian Song, Jun Zhu:  
**GNOT: A General Neural Operator Transformer for Operator Learning.** 12556-12569
- Moritz Hardt, Eric Mazumdar, Celestine Mendler-Dünner, Tijana Zrnic:  
**Algorithmic Collective Action in Machine Learning.** 12570-12586
- Marc Häkkinen, Markus Lange-Hegermann, Bogdan Raita:  
**Gaussian Process Priors for Systems of Linear Partial Differential Equations with Constant Coefficients.** 12587-12615
- Hilaf Hasson, Danielle C. Maddix, Bernie Wang, Gaurav Gupta, Youngsuk Park:  
**Theoretical Guarantees of Learning Ensembling Strategies with Applications to Time Series Forecasting.** 12616-12632
- Ali Hatamizadeh, Hongxu Yin, Greg Heinrich, Jan Kautz, Pavlo Molchanov:  
**Global Context Vision Transformers.** 12633-12646
- Martin B. Haugh, Raghav Singal:  
**Counterfactual Analysis in Dynamic Latent State Models.** 12647-12677
- Satoshi Hayakawa, Harald Oberhauser, Terry J. Lyons:  
**Sampling-based Nyström Approximation and Kernel Quadrature.** 12678-12699
- Soufiane Hayou, Greg Yang:  
**Width and Depth Limits Commute in Residual Networks.** 12700-12723
- Xiaoxin He, Bryan Hooi, Thomas Laurent, Adam Perold, Yann LeCun, Xavier Bresson:  
**A Generalization of ViT/MLP-Mixer to Graphs.** 12724-12745



- Huan He, Owen Queen, Teddy Koker, Consuelo Cuevas, Theodoros Tsiligkaridis, Marinka Zitnik: **Domain Adaptation for Time Series Under Feature and Label Shifts.** 12746-12774
- Dongxiao He, Jitao Zhao, Rui Guo, Zhiyong Feng, Di Jin, Yuxiao Huang, Zhen Wang, Weixiong Zhang: **Contrastive Learning Meets Homophily: Two Birds with One Stone.** 12775-12789
- Jiafan He, Heyang Zhao, Dongruo Zhou, Quanquan Gu: **Nearly Minimax Optimal Reinforcement Learning for Linear Markov Decision Processes.** 12790-12822
- S. Ashwin Hebbar, Viraj Vivek Nadkarni, Ashok Vardhan Makkula, Suma Bhat, Sewoong Oh, Pramod Viswanath: **CRISP: Curriculum based Sequential neural decoders for Polar code family.** 12823-12845
- Jonathan Hehir, Daniel Ting, Graham Cormode: **Sketch-Flip-Merge: Mergeable Sketches for Private Distinct Counting.** 12846-12865
- Florian Heinrichs, Marvin Heim, Corinna Weber: **Functional Neural Networks: Shift invariant models for functional data with applications to EEG classification.** 12866-12881
- Joey Hejna, Jensen Gao, Dorsa Sadigh: **Distance Weighted Supervised Learning for Offline Interaction Data.** 12882-12906
- Jacob Helwig, Xuan Zhang, Cong Fu, Jerry Kurtin, Stephan Wojtowytzsch, Shuiwang Ji: **Group Equivariant Fourier Neural Operators for Partial Differential Equations.** 12907-12930
- Apivich Hemachandra, Zhongxiang Dai, Jasraj Singh, See-Kiong Ng, Bryan Kian Hsiang Low: **Training-Free Neural Active Learning with Initialization-Robustness Guarantees.** 12931-12971
- Mikael Henaff, Minqi Jiang, Roberta Raileanu: **A Study of Global and Episodic Bonuses for Exploration in Contextual MDPs.** 12972-12999
- Hwan Heo, Taekyung Kim, Jiyoung Lee, Jaewon Lee, Soohyun Kim, Hyunwoo J. Kim, Jin-Hwa Kim: **Robust Camera Pose Refinement for Multi-Resolution Hash Encoding.** 13000-13016
- Florian Hess, Zahra Monfared, Manuel Brenner, Daniel Durstewitz: **Generalized Teacher Forcing for Learning Chaotic Dynamics.** 13017-13049
- Caglar Hizli, S. T. John, Anne Tuulikki Juuti, Tuure Tapani Saarinen, Kirsi Hannele Pietiläinen, Pekka Marttinen: **Causal Modeling of Policy Interventions From Treatment-Outcome Sequences.** 13050-13084
- Liam Hodgkinson, Christopher van der Heide, Fred Roosta, Michael W. Mahoney: **Monotonicity and Double Descent in Uncertainty Estimation with Gaussian Processes.** 13085-13117
- Mikael Møller Høgsgaard, Kasper Green Larsen, Martin Ritzert: **AdaBoost is not an Optimal Weak to Strong Learner.** 13118-13140
- Rasmus Kjær Høier, D. Staudt, Christopher Zach: **Dual Propagation: Accelerating Contrastive Hebbian Learning with Dyadic Neurons.** 13141-13156
- Joey Hong, Branislav Kveton, Manzil Zaheer, Sumeet Katariya, Mohammad Ghavamzadeh: **Multi-Task Off-Policy Learning from Bandit Feedback.** 13157-13173
- Ilgee Hong, Sen Na, Michael W. Mahoney, Mladen Kolar: **Constrained Optimization via Exact Augmented Lagrangian and Randomized Iterative Sketching.** 13174-13198
- Junyuan Hong, Yi Zeng, Shuyang Yu, Lingjuan Lyu, Ruoxi Jia, Jiayu Zhou: **Revisiting Data-Free Knowledge Distillation with Poisoned Teachers.** 13199-13212
- Emiel Hoogeboom, Jonathan Heek, Tim Salimans: **simple diffusion: End-to-end diffusion for high resolution images.** 13213-13232
- Guy Horowitz, Nir Rosenfeld: **Causal Strategic Classification: A Tale of Two Shifts.** 13233-13253
- Ramtin Hosseini, Li Zhang, Bhanu Garg, Pengtao Xie: **Fair and Accurate Decision Making through Group-Aware Learning.** 13254-13269
- Sèdja Salomon Hotegni, Sepideh Mahabadi, Ali Vakilian: **Approximation Algorithms for Fair Range Clustering.** 13270-13284
- Elizabeth Mary Hou, Gregory David Castañón: **Decoding Layer Saliency in Language Transformers.** 13285-13308
- Bairu Hou, Joe O'Connor, Jacob Andreas, Shiyu Chang, Yang Zhang: **PromptBoosting: Black-Box Text Classification with Ten Forward Passes.** 13309-13324
- Boya Hou, Sina Sanjari, Nathan Dahlin, Subhonmesh Bose, Umesh Vaidya: **Sparse Learning of Dynamical Systems in RKHS: An Operator-Theoretic Approach.** 13325-13352
- Yunlong Hou, Vincent Y. F. Tan, Zixin Zhong: **Probably Anytime-Safe Stochastic Combinatorial Semi-Bandits.** 13353-13409
- Ignacio Hounie, Luiz F. O. Chamon, Alejandro Ribeiro: **Automatic Data Augmentation via Invariance-Constrained Learning.** 13410-13433
- Yu-Guan Hsieh, Shiva Prasad Kasiviswanathan, Branislav Kveton, Patrick Blöbaum: **Thompson Sampling with Diffusion Generative Prior.** 13434-13468



- Zhengmian Hu, Heng Huang:  
**Tighter Analysis for ProxSkip.** 13469-13496
- Lunjia Hu , Inbal Rachel Livni Navon, Omer Reingold, Chutong Yang:  
**Omnipredictors for Constrained Optimization.** 13497-13527
- Edward J. Hu, Nikolay Malkin, Moksh Jain, Katie E. Everett, Alexandros Graikos, Yoshua Bengio:  
**GFlowNet-EM for Learning Compositional Latent Variable Models.** 13528-13549
- Quanqi Hu, Zi-Hao Qiu, Zhishuai Guo, Lijun Zhang, Tianbao Yang:  
**Blockwise Stochastic Variance-Reduced Methods with Parallel Speedup for Multi-Block Bilevel Optimization.** 13550-13583
- Hengyuan Hu, Dorsa Sadigh:  
**Language Instructed Reinforcement Learning for Human-AI Coordination.** 13584-13598
- Yuan-Ting Hu, Alexander G. Schwing, Raymond A. Yeh:  
**Surface Snapping Optimization Layer for Single Image Object Shape Reconstruction.** 13599-13609
- Zixuan Hu, Li Shen, Zhenyi Wang, Baoyuan Wu, Chun Yuan, Dacheng Tao:  
**Learning to Learn from APIs: Black-Box Data-Free Meta-Learning.** 13610-13627
- Yingdong Hu, Renhao Wang, Li Erran Li, Yang Gao:  
**For Pre-Trained Vision Models in Motor Control, Not All Policy Learning Methods are Created Equal.** 13628-13651
- Zhengmian Hu, Xidong Wu, Heng Huang:  
**Beyond Lipschitz Smoothness: A Tighter Analysis for Nonconvex Optimization.** 13652-13678
- Yuzheng Hu, Fan Wu, Hongyang Zhang, Han Zhao:  
**Understanding the Impact of Adversarial Robustness on Accuracy Disparity.** 13679-13709
- Audrey Huang, Jinglin Chen, Nan Jiang:  
**Reinforcement Learning in Low-rank MDPs with Density Features.** 13710-13752
- Lianghua Huang, Di Chen, Yu Liu, Yujun Shen, Deli Zhao, Jingren Zhou:  
**Composer: Creative and Controllable Image Synthesis with Composable Conditions.** 13753-13773
- Zizheng Huang, Haoxing Chen, Ziqi Wen, Chao Zhang, Huaxiong Li, Bo Wang, Chunlin Chen:  
**Model-Aware Contrastive Learning: Towards Escaping the Dilemmas.** 13774-13790
- Tianyi Huang, Shenghui Cheng, Stan Z. Li, Zhengjun Zhang:  
**High-dimensional Clustering onto Hamiltonian Cycle.** 13791-13813
- Jiatai Huang, Yan Dai, Longbo Huang:  
**Banker Online Mirror Descent: A Universal Approach for Delayed Online Bandit Learning.** 13814-13844
- Junyu Huang, Qilong Feng, Ziyun Huang , Jinhui Xu, Jianxin Wang:  
**Fast Algorithms for Distributed k-Clustering with Outliers.** 13845-13868
- Taoan Huang, Aaron M. Ferber, Yuandong Tian, Bistra Dilkina, Benoit Steiner:  
**Searching Large Neighborhoods for Integer Linear Programs with Contrastive Learning.** 13869-13890
- Lingxiao Huang, Ruiyuan Huang, Zengfeng Huang, Xuan Wu:  
**On Coresets for Clustering in Small Dimensional Euclidean spaces.** 13891-13915
- Rongjie Huang, Jiawei Huang, Dongchao Yang, Yi Ren, Luping Liu, Mingze Li, Zhenhui Ye, Jinglin Liu, Xiang Yin, Zhou Zhao:  
**Make-An-Audio: Text-To-Audio Generation with Prompt-Enhanced Diffusion Models.** 13916-13932
- Lingxiao Huang, Shaofeng H.-C. Jiang, Jianing Lou:  
**The Power of Uniform Sampling for k-Median.** 13933-13956
- Zhaohuang, Litian Liang, Zhan Ling, Xuanlin Li, Chuang Gan, Hao Su:  
**Reparameterized Policy Learning for Multimodal Trajectory Optimization.** 13957-13975
- Yufan Huang, C. Seshadhri, David F. Gleich:  
**Theoretical Bounds on the Network Community Profile from Low-rank Semi-definite Programming.** 13976-13992
- Xinquan Huang , Wenlei Shi, Qi Meng, Yue Wang, Xiaotian Gao, Jia Zhang, Tie-Yan Liu:  
**NeuralStagger: Accelerating Physics-constrained Neural PDE Solver with Spatial-temporal Decomposition.** 13993-14006
- Jialei Huang, Zhao-Heng Yin, Yingdong Hu, Yang Gao:  
**Policy Contrastive Imitation Learning.** 14007-14022
- Tianjin Huang, Lu Yin, Zhenyu Zhang, Li Shen, Meng Fang, Mykola Pechenizkiy, Zhangyang Wang, Shiwei Liu:  
**Are Large Kernels Better Teachers than Transformers for ConvNets?** 14023-14038
- Minhui Huang, Dewei Zhang, Kaiyi Ji:  
**Achieving Linear Speedup in Non-IID Federated Bilevel Learning.** 14039-14059
- Ruiquan Huang, Huanyu Zhang, Luca Melis, Milan Shen, Meisam Hejazinia, Jing Yang:  
**Federated Linear Contextual Bandits with User-level Differential Privacy.** 14060-14095
- Minyoung Huh, Brian Cheung, Pulkit Agrawal, Phillip Isola:  
**Straightening Out the Straight-Through Estimator: Overcoming Optimization Challenges in Vector Quantized Networks.** 14096-14113



- Like Hui, Mikhail Belkin, Stephen Wright:  
**Cut your Losses with Squentropy.** 14114-14131
- Iris A. M. Huijben, Arthur Andreas Nijdam, Sebastiaan Overeem, Merel M. van Gilst, Ruud van Sloun:  
**SOM-CPC: Unsupervised Contrastive Learning with Self-Organizing Maps for Structured Representations of High-Rate Time Series.** 14132-14152
- Pierre Humbert, Batiste Le Bars, Aurélien Bellet, Sylvain Arlot:  
**One-Shot Federated Conformal Prediction.** 14153-14177
- Aamal Abbas Hussain, Francesco Belardinelli, Dario Paccagnan:  
**The Impact of Exploration on Convergence and Performance of Multi-Agent Q-Learning Dynamics.** 14178-14202
- Taehyun Hwang, Kyuwook Chai, Min-hwan Oh:  
**Combinatorial Neural Bandits.** 14203-14236
- Geonho Hwang, Jaewoong Choi, Hyunsoo Cho, Myungjoo Kang:  
**MAGANet: Achieving Combinatorial Generalization by Modeling a Group Action.** 14237-14248
- HyeongJoo Hwang, Seokin Seo, Youngsoo Jang, Sungyoon Kim, Geon-Hyeong Kim, Seunghoon Hong, Kee-Eung Kim:  
**Information-Theoretic State Space Model for Multi-View Reinforcement Learning.** 14249-14282
- Shahana Ibrahim, Xiao Fu, Rebecca A. Hutchinson, Eugene Seo:  
**Under-Counted Tensor Completion with Neural Incorporation of Attributes.** 14283-14315
- Alexander Immer, Christoph Schultheiss, Julia E. Vogt, Bernhard Schölkopf, Peter Bühlmann, Alexander Marx:  
**On the Identifiability and Estimation of Causal Location-Scale Noise Models.** 14316-14332
- Alexander Immer, Tycho F. A. van der Ouderaa, Mark van der Wilk, Gunnar Rätsch, Bernhard Schölkopf:  
**Stochastic Marginal Likelihood Gradients using Neural Tangent Kernels.** 14333-14352
- Jacob Imola, Alessandro Epasto, Mohammad Mahdian, Vincent Cohen-Addad, Vahab Mirrokni:  
**Differentially Private Hierarchical Clustering with Provable Approximation Guarantees.** 14353-14375
- Brian Irwin , Eldad Haber, Raviv Gal, Avi Ziv:  
**Neural Network Accelerated Implicit Filtering: Integrating Neural Network Surrogates With Provably Convergent Derivative Free Optimization Methods.** 14376-14389
- Riashat Islam, Manan Tomar, Alex Lamb, Yonathan Efroni, Hongyu Zang, Aniket Rajiv Didolkar, Dipendra Misra, Xin Li, Harm van Seijen, Remi Tachet des Combes, John Langford:  
**Principled Offline RL in the Presence of Rich Exogenous Information.** 14390-14421
- Thibaut Issenhuth, Ugo Tanielian, Jérémie Mary, David Picard:  
**Unveiling the Latent Space Geometry of Push-Forward Generative Models.** 14422-14444
- Desi R. Ivanova, Joel Jennings, Tom Rainforth, Cheng Zhang, Adam Foster:  
**CO-BED: Information-Theoretic Contextual Optimization via Bayesian Experimental Design.** 14445-14464
- Maor Ivgi, Oliver Hinder, Yair Carmon:  
**DoG is SGD's Best Friend: A Parameter-Free Dynamic Step Size Schedule.** 14465-14499
- Gaurav Iyer, Boris Hanin, David Rolnick:  
**Maximal Initial Learning Rates in Deep ReLU Networks.** 14500-14530
- Zachary Izzo, Ruishan Liu, James Zou:  
**Data-Driven Subgroup Identification for Linear Regression.** 14531-14552
- Sashank J. Reddi, Sobhan Miryoosefi, Stefani Karp, Shankar Krishnan, Satyen Kale, Seungyeon Kim, Sanjiv Kumar:  
**Efficient Training of Language Models using Few-Shot Learning.** 14553-14568
- Allan Jabri, David J. Fleet, Ting Chen:  
**Scalable Adaptive Computation for Iterative Generation.** 14569-14589
- Andrew Jacobsen, Ashok Cutkosky:  
**Unconstrained Online Learning with Unbounded Losses.** 14590-14630
- Moksh Jain, Sharath Chandra Raparthi, Alex Hernández-García, Jarrid Rector-Brooks, Yoshua Bengio, Santiago Miret, Emmanuel Bengio:  
**Multi-Objective GFlowNets.** 14631-14653
- Palak Jain, Sofya Raskhadnikova, Satchit Sivakumar, Adam D. Smith:  
**The Price of Differential Privacy under Continual Observation.** 14654-14678
- Ajay Kumar Jaiswal, Shiwei Liu, Tianlong Chen, Ying Ding, Zhangyang Wang:  
**Graph Ladling: Shockingly Simple Parallel GNN Training without Intermediate Communication.** 14679-14690
- Ajay Kumar Jaiswal, Shiwei Liu, Tianlong Chen, Ying Ding, Zhangyang Wang:  
**Instant Soup: Cheap Pruning Ensembles in A Single Pass Can Draw Lottery Tickets from Large Models.** 14691-14701
- Joel Jang, Seungone Kim, Seonghyeon Ye, Doyoung Kim, Lajanugen Logeswaran, Moontae Lee, Kyungjae Lee, Minjoon Seo:  
**Exploring the Benefits of Training Expert Language Models over Instruction Tuning.** 14702-14729
- Jinhyeok Jang, Woo-han Yun, Won Hwa Kim, Youngwoo Yoon, Jaehong Kim, Jaeyeon Lee, ByungOk Han:  
**Learning to Boost Training by Periodic Nowcasting Near Future Weights.** 14730-14757
- Faris Janjos, Lars Rosenbaum, Maxim Dolgov, J. Marius Zoellner:  
**Unscented Autoencoder.** 14758-14779



- Daniel Jarrett, Corentin Tallec, Florent Altché, Thomas Mesnard, Rémi Munos, Michal Valko:  
**Curiosity in Hindsight: Intrinsic Exploration in Stochastic Environments.** 14780-14816
- Kishaan Jeeveswaran, Prashant Shivaram Bhat, Bahram Zonooz, Elahe Arani:  
**BiRT: Bio-inspired Replay in Vision Transformers for Continual Learning.** 14817-14835
- Hyeonsu Jeong, Hye Won Chung:  
**Recovering Top-Two Answers and Confusion Probability in Multi-Choice Crowdsourcing.** 14836-14868
- Feng Ji, See Hian Lee, Hanyang Meng, Kai Zhao, Jielong Yang, Wee Peng Tay:  
**Leveraging Label Non-Uniformity for Node Classification in Graph Neural Networks.** 14869-14885
- Lin-Han Jia, Lan-Zhe Guo, Zhi Zhou , Jie-Jing Shao, Yuke Xiang, Yufeng Li:  
**Bidirectional Adaptation for Robust Semi-Supervised Learning with Inconsistent Data Distributions.** 14886-14901
- Su Jia, Nishant Oli, Ian Anderson, Paul Duff, Andrew A. Li, R. Ravi:  
**Short-lived High-volume Bandits.** 14902-14929
- Su Jia, Qian Xie, Nathan Kallus, Peter I. Frazier:  
**Smooth Non-stationary Bandits.** 14930-14944
- Haiyan Jiang, Srinivas Anumasa, Giulia De Masi, Huan Xiong, Bin Gu:  
**A Unified Optimization Framework of ANN-SNN Conversion: Towards Optimal Mapping from Activation Values to Firing Rates.** 14945-14974
- Yunfan Jiang, Agrim Gupta, Zichen Zhang, Guanzhi Wang, Yongqiang Dou, Yanjun Chen, Li Fei-Fei, Anima Anandkumar, Yuke Zhu, Linxi Fan:  
**VIMA: Robot Manipulation with Multimodal Prompts.** 14975-15022
- Ziyang Jiang, Zhuoran Hou, Yiling Liu, Yiman Ren, Keyu Li, David E. Carlson:  
**Estimating Causal Effects using a Multi-task Deep Ensemble.** 15023-15040
- Bowen Jiang, Bo Jiang, Jian Li, Tao Lin, Xinbing Wang, Chenghu Zhou:  
**Online Restless Bandits with Unobserved States.** 15041-15066
- Xue Jiang, Feng Liu, Zhen Fang, Hong Chen, Tongliang Liu, Feng Zheng, Bo Han:  
**Detecting Out-of-distribution Data through In-distribution Class Prior.** 15067-15088
- Yulun Jiang, Chen Liu, Zhichao Huang, Mathieu Salzmann, Sabine Süsstrunk:  
**Towards Stable and Efficient Adversarial Training against  $\ell_1$  Bounded Adversarial Attacks.** 15089-15104
- Wei Jiang, Jiayu Qin, Lingyu Wu, Changyou Chen, Tianbao Yang, Lijun Zhang:  
**Learning Unnormalized Statistical Models via Compositional Optimization.** 15105-15124
- Ziwei Jiang, Lai Wei, Murat Kocaoglu:  
**Approximate Causal Effect Identification under Weak Confounding.** 15125-15143
- Guangyuan Jiang, Manjie Xu, Shiji Xin, Wei Liang, Yujia Peng, Chi Zhang, Yixin Zhu:  
**MEWL: Few-shot multimodal word learning with referential uncertainty.** 15144-15169
- Chenbo Jiang, Jie Yang, Shuai He, Yu-Kun Lai, Lin Gao:  
**NeuralSlice: Neural 3D Triangle Mesh Reconstruction via Slicing 4D Tetrahedral Meshes.** 15170-15185
- Weisen Jiang, Yu Zhang, James T. Kwok:  
**Effective Structured Prompting by Meta-Learning and Representative Verbalizer.** 15186-15199
- Jikai Jin, Zhiyuan Li , Kaifeng Lyu, Simon Shaolei Du, Jason D. Lee:  
**Understanding Incremental Learning of Gradient Descent: A Fine-grained Analysis of Matrix Sensing.** 15200-15238
- Tianyuan Jin, Xianglin Yang, Xiaokui Xiao, Pan Xu:  
**Thompson Sampling with Less Exploration is Fast and Optimal.** 15239-15261
- Daniel D. Johnson, Daniel Tarlow, Christian Walder:  
**R-U-SURE? Uncertainty-Aware Code Suggestions By Maximizing Utility Across Random User Intents.** 15262-15306
- Erik Jones, Anca D. Dragan, Aditi Raghunathan, Jacob Steinhardt:  
**Automatically Auditing Large Language Models via Discrete Optimization.** 15307-15329
- Chaitanya K. Joshi, Cristian Bodnar, Simon V. Mathis, Taco Cohen, Pietro Lio:  
**On the Expressive Power of Geometric Graph Neural Networks.** 15330-15355
- Siddharth Joshi, Baharan Mirzasoleiman:  
**Data-Efficient Contrastive Self-supervised Learning: Most Beneficial Examples for Supervised Learning Contribute the Least.** 15356-15370
- Kishor Jothimurugan, Steve Hsu, Osbert Bastani, Rajeev Alur:  
**Robust Subtask Learning for Compositional Generalization.** 15371-15387
- Amir Joudaki, Hadi Daneshmand, Francis R. Bach:  
**On Bridging the Gap between Mean Field and Finite Width Deep Random Multilayer Perceptron with Batch Normalization.** 15388-15400
- Nikola Jovanovic, Mislav Balunovic, Dimitar Iliev Dimitrov, Martin T. Vechev:  
**FARE: Provably Fair Representation Learning with Practical Certificates.** 15401-15420
- Seungjin Jung, Seungmo Seo, Yonghyun Jeong, Jongwon Choi:  
**Scaling of Class-wise Training Losses for Post-hoc Calibration.** 15421-15434



- Yeonsung Jung, Hajin Shim, June Yong Yang, Eunho Yang:  
**Fighting Fire with Fire: Contrastive Debiasing without Bias-free Data via Generative Bias-transformation.** 15435-15450
- Yonghan Jung, Jin Tian, Elias Bareinboim:  
**Estimating Joint Treatment Effects by Combining Multiple Experiments.** 15451-15527
- Mateusz Maria Jurewicz, Graham W. Taylor, Leon Derczynski:  
**The Catalog Problem: Clustering and Ordering Variable-Sized Sets.** 15528-15545
- Sékou-Oumar Kaba, Arnab Kumar Mondal, Yan Zhang, Yoshua Bengio, Siamak Ravanbakhsh:  
**Equivariance with Learned Canonicalization Functions.** 15546-15566
- Hiroshi Kajino, Kohei Miyaguchi, Takayuki Osogami:  
**Biases in Evaluation of Molecular Optimization Methods and Bias Reduction Strategies.** 15567-15585
- Alkis Kalavasis, Amin Karbasi, Shay Moran, Grigoris Velekas:  
**Statistical Indistinguishability of Learning Algorithms.** 15586-15622
- Neha Mukund Kalibhat, Shweta Bhardwaj, C. Bayan Bruss, Hamed Firooz, Maziar Sanjabi, Soheil Feizi:  
**Identifying Interpretable Subspaces in Image Representations.** 15623-15638
- David Kaltenpoth, Jilles Vreeken:  
**Nonlinear Causal Discovery with Latent Confounders.** 15639-15654
- Pierre-Alexandre Kamienny, Guillaume Lample, Sylvain Lamprier, Marco Virgolin:  
**Deep Generative Symbolic Regression with Monte-Carlo-Tree-Search.** 15655-15668
- Sekitoshi Kanai, Shin'ya Yamaguchi, Masanori Yamada, Hiroshi Takahashi, Kentaro Ohno, Yasutoshi Ida:  
**One-vs-the-Rest Loss to Focus on Important Samples in Adversarial Training.** 15669-15695
- Nikhil Kandpal, Haikang Deng, Adam Roberts, Eric Wallace, Colin Raffel:  
**Large Language Models Struggle to Learn Long-Tail Knowledge.** 15696-15707
- Nikhil Kandpal, Brian Lester, Mohammed Muqeeth, Anisha Mascarenhas, Monty Evans, Vishal Baskaran, Tenghao Huang, Haokun Liu, Colin Raffel:  
**Git-Theta: A Git Extension for Collaborative Development of Machine Learning Models.** 15708-15719
- Ayano Kaneda, Osman Akar, Jingyu Chen, Victoria Alicia Trevino Kala, David Hyde, Joseph Teran:  
**A Deep Conjugate Direction Method for Iteratively Solving Linear Systems.** 15720-15736
- Juwon Kang, Nayeong Kim, Donghyeon Kwon, Jungseul Ok, Suha Kwak:  
**Leveraging Proxy of Training Data for Test-Time Adaptation.** 15737-15752
- Yachen Kang, Diyuan Shi, Jinxin Liu, Li He, Donglin Wang:  
**Beyond Reward: Offline Preference-guided Policy Optimization.** 15753-15768
- Siteng Kang, Zhan Shi, Xinhua Zhang:  
**Poisoning Generative Replay in Continual Learning to Promote Forgetting.** 15769-15785
- Qiyu Kang, Kai Zhao, Yang Song, Sijie Wang, Wee Peng Tay:  
**Node Embedding from Neural Hamiltonian Orbits in Graph Neural Networks.** 15786-15808
- Ryo Karakida, Tomoumi Takase, Tomohiro Hayase, Kazuki Osawa:  
**Understanding Gradient Regularization in Deep Learning: Efficient Finite-Difference Computation and Implicit Bias.** 15809-15827
- Amin Karbasi, Nikki Lijing Kuang, Yi-An Ma, Siddharth Mitra:  
**Langevin Thompson Sampling with Logarithmic Communication: Bandits and Reinforcement Learning.** 15828-15860
- Amir-Hossein Karimi, Krikamol Muandet, Simon Kornblith, Bernhard Schölkopf, Been Kim:  
**On the Relationship Between Explanation and Prediction: A Causal View.** 15861-15883
- Sanjay Kariyappa, Chuan Guo, Kiwan Maeng, Wenjie Xiong , G. Edward Suh, Moinuddin K. Qureshi, Hsien-Hsin S. Lee:  
**Cocktail Party Attack: Breaking Aggregation-Based Privacy in Federated Learning Using Independent Component Analysis.** 15884-15899
- Arjun Karuvally , Terrence J. Sejnowski, Hava T. Siegelmann:  
**General Sequential Episodic Memory Model.** 15900-15910
- Takayuki Katsuki, Takayuki Osogami:  
**Regression with Sensor Data Containing Incomplete Observations.** 15911-15927
- Ilya Kaufman, Omri Azencot:  
**Data Representations' Study of Latent Image Manifolds.** 15928-15945
- Prannay Kaul, Weidi Xie, Andrew Zisserman:  
**Multi-Modal Classifiers for Open-Vocabulary Object Detection.** 15946-15969
- Chinmaya Kausik, Kevin Tan , Ambuj Tewari:  
**Learning Mixtures of Markov Chains and MDPs.** 15970-16017
- Isaac Kauvar, Chris Doyle, Linqi Zhou, Nick Haber:  
**Curious Replay for Model-based Adaptation.** 16018-16048
- Kenji Kawaguchi, Zhun Deng, Xu Ji, Jiaoyang Huang:  
**How Does Information Bottleneck Help Deep Learning?** 16049-16096



- Yuta Kawakami, Manabu Kuroki, Jin Tian:  
**Instrumental Variable Estimation of Average Partial Causal Effects.** 16097-16130
- Zeki Kazan, Kaiyan Shi, Adam Groce, Andrew P. Bray:  
**The Test of Tests: A Framework for Differentially Private Hypothesis Testing.** 16131-16151
- Chuyang Ke, Jean Honorio:  
**Exact Inference in High-order Structured Prediction.** 16152-16167
- T. Anderson Keller, Max Welling:  
**Neural Wave Machines: Learning Spatiotemporally Structured Representations with Locally Coupled Oscillatory Recurrent Neural Networks.** 16168-16189
- Hamza Keurti, Hsiao-Ru Pan, Michel Besserve, Benjamin F. Grewe, Bernhard Schölkopf:  
**Homomorphism AutoEncoder - Learning Group Structured Representations from Observed Transitions.** 16190-16215
- Alaa Khaddaj, Guillaume Leclerc, Aleksandar Makelov, Kristian Georgiev, Hadi Salman, Andrew Ilyas, Aleksander Madry:  
**Rethinking Backdoor Attacks.** 16216-16236
- Adam Khakhar, Stephen Mell, Osbert Bastani:  
**PAC Prediction Sets for Large Language Models of Code.** 16237-16249
- Mohammad Khalafi, Digvijay Boob:  
**Accelerated Primal-Dual Methods for Convex-Strongly-Concave Saddle Point Problems.** 16250-16270
- Mohammad Mahdi Khalili, Xueru Zhang, Mahed Abroshan:  
**Loss Balancing for Fair Supervised Learning.** 16271-16290
- Prashant Khanduri, Ioannis C. Tsaknakis, Yihua Zhang, Jia Liu, Sijia Liu, Jiawei Zhang, Mingyi Hong:  
**Linearly Constrained Bilevel Optimization: A Smoothed Implicit Gradient Approach.** 16291-16325
- Mahyar Khayatkhoei, Wael Abd-Almageed:  
**Emergent Asymmetry of Precision and Recall for Measuring Fidelity and Diversity of Generative Models in High Dimensions.** 16326-16343
- Mikhail Khodak, Kareem Amin, Travis Dick, Sergei Vassilvitskii:  
**Learning-augmented private algorithms for multiple quantile release.** 16344-16376
- Jihye Kim, Aristide Baratin, Yan Zhang, Simon Lacoste-Julien:  
**CrossSplit: Mitigating Label Noise Memorization through Data Splitting.** 16377-16392
- Timothy Doyeon Kim, Tankut Can, Kamesh Krishnamurthy:  
**Trainability, Expressivity and Interpretability in Gated Neural ODEs.** 16393-16423
- Yoon-Yeong Kim, Youngjae Cho, JoonHo Jang, Byeongghu Na, Yeongmin Kim, Kyungwoo Song, Wanmo Kang, Il-Chul Moon:  
**SAAL: Sharpness-Aware Active Learning.** 16424-16440
- Jigang Kim, Daesol Cho, H. Jin Kim:  
**Demonstration-free Autonomous Reinforcement Learning via Implicit and Bidirectional Curriculum.** 16441-16457
- Wonyoung Kim, Garud Iyengar, Assaf Zeevi:  
**Improved Algorithms for Multi-period Multi-class Packing Problems with Bandit Feedback.** 16458-16501
- Jinuk Kim, Yeonwoo Jeong, Deokjae Lee, Hyun Oh Song:  
**Efficient Latency-Aware CNN Depth Compression via Two-Stage Dynamic Programming.** 16502-16520
- Eunji Kim, Dahuin Jung, Sangha Park, Siwon Kim, Sungroh Yoon:  
**Probabilistic Concept Bottleneck Models.** 16521-16540
- Haeyeon Kim, Minsu Kim, Federico Berto, Joung Ho Kim, Jinkyoo Park:  
**DevFormer: A Symmetric Transformer for Context-Aware Device Placement.** 16541-16566
- Dongjun Kim, Yeongmin Kim, Se Jung Kwon, Wanmo Kang, Il-Chul Moon:  
**Refining Generative Process with Discriminator Guidance in Score-based Diffusion Models.** 16567-16598
- Junghoon Kim, Taejoon Kim, David J. Love, Christopher G. Brinton:  
**Robust Non-Linear Feedback Coding via Power-Constrained Deep Learning.** 16599-16618
- Woojun Kim, Jeonghye Kim, Youngchul Sung:  
**LESSON: Learning to Integrate Exploration Strategies for Reinforcement Learning via an Option Framework.** 16619-16638
- Taebum Kim, Hyoungjoo Kim, Gyeong-In Yu, Byung-Gon Chun:  
**BPipe: Memory-Balanced Pipeline Parallelism for Training Large Language Models.** 16639-16653
- Seunghyun Kim, Hyunsu Kim, Eunggu Yun, Hwangrae Lee, Jaehun Lee, Juho Lee:  
**Probabilistic Imputation for Time-series Classification with Missing Data.** 16654-16667
- Seongun Kim, Kyowoon Lee, Jaesik Choi:  
**Variational Curriculum Reinforcement Learning for Unsupervised Discovery of Skills.** 16668-16695
- Byungjoo Kim, Suyoung Lee, Seanie Lee, Sooel Son, Sung Ju Hwang:  
**Margin-based Neural Network Watermarking.** 16696-16711
- Hyunsu Kim, Hyungi Lee, Hongseok Yang, Juho Lee:  
**Regularizing Towards Soft Equivariance Under Mixed Symmetries.** 16712-16727



- Byeongchan Kim, Min-hwan Oh:  
**Model-based Offline Reinforcement Learning with Count-based Conservatism.** 16728-16746
- Chanyeong Kim, Jongwoong Park, Hyunglip Bae, Woo Chang Kim:  
**Transformer-based Stagewise Decomposition for Large-Scale Multistage Stochastic Optimization.** 16747-16770
- Junetae Kim, Kyoungsuk Park, Hanseok Jeong, Youngwook Kim, Jeongseon Kim, Sun-Young Kim:  
**SurProGenes: Survival Risk-Ordered Representation of Cancer Patients and Genes for the Identification of Prognostic Genes.** 16771-16786
- Seungwook Kim, Chunghyun Park, Yoonwoo Jeong, Jaesik Park, Minsu Cho:  
**Stable and Consistent Prediction of 3D Characteristic Orientation via Invariant Residual Learning.** 16787-16806
- Jaehyung Kim, Jinwoo Shin, Dongyeop Kang:  
**Prefer to Classify: Improving Text Classifiers via Auxiliary Preference Learning.** 16807-16828
- Woojun Kim, Youngchul Sung:  
**An Adaptive Entropy-Regularization Framework for Multi-Agent Reinforcement Learning.** 16829-16852
- Kyurae Kim , Kaiwen Wu, Jisu Oh, Jacob R. Gardner:  
**Practical and Matching Gradient Variance Bounds for Black-Box Variational Bayesian Inference.** 16853-16876
- Gyuhak Kim , Changnan Xiao, Tatsuya Konishi, Bing Liu:  
**Learnability and Algorithm for Continual Learning.** 16877-16896
- Jungbin Kim, Insoon Yang:  
**Unifying Nesterov's Accelerated Gradient Methods for Convex and Strongly Convex Objective Functions.** 16897-16954
- Beomsu Kim, Jong Chul Ye:  
**Denoising MCMC for Accelerating Diffusion-Based Generative Models.** 16955-16977
- Dale Kim, Qing Zhou:  
**Structure Learning of Latent Factors via Clique Search on Correlation Thresholded Graphs.** 16978-16996
- Kwangho Kim, José R. Zubizarreta:  
**Fair and Robust Estimation of Heterogeneous Treatment Effects for Policy Learning.** 16997-17014
- Dhamma Kimpara, Rafael M. Frongillo, Bo Waggoner:  
**Proper Losses for Discrete Generative Models.** 17015-17040
- Yuri Kinoshita, Kenta Oono, Kenji Fukumizu, Yuichi Yoshida, Shin-ichi Maeda:  
**Controlling Posterior Collapse by an Inverse Lipschitz Constraint on the Decoder Network.** 17041-17060
- John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, Tom Goldstein:  
**A Watermark for Large Language Models.** 17061-17084
- Michael Kirchhof, Enkelejda Kasneci, Seong Joon Oh:  
**Probabilistic Contrastive Learning Recovers the Correct Aleatoric Uncertainty of Ambiguous Inputs.** 17085-17104
- Matthias Kirchler, Christoph Lippert, Marius Kloft:  
**Training Normalizing Flows from Dependent Data.** 17105-17121
- Varsha Kishore, Chao Wan, Justin Lovelace, Yoav Artzi, Kilian Q. Weinberger:  
**IncDSI: Incrementally Updatable Document Retrieval.** 17122-17134
- Toshinori Kitamura, Tadashi Kozuno, Yunhao Tang, Nino Vieillard, Michal Valko, Wenhao Yang, Jincheng Mei, Pierre Ménard, Mohammad Gheshlaghi Azar, Rémi Munos, Olivier Pietquin, Matthieu Geist, Csaba Szepesvári, Wataru Kumagai, Yutaka Matsuo:  
**Regularization and Variance-Weighted Regression Achieves Minimax Optimality in Linear MDPs: Theory and Practice.** 17135-17175
- Leo Klerner, Tim G. J. Rudner, Michael Reutlinger, Torsten Schindler, Garrett M. Morris, Charlotte M. Deane, Yee Whye Teh:  
**Drug Discovery under Covariate Shift with Domain-Informed Prior Distributions over Functions.** 17176-17197
- Martin Klissarov, Marlos C. Machado:  
**Deep Laplacian-based Options for Temporally-Extended Exploration.** 17198-17217
- Marina Knittel, Max Springer, John P. Dickerson, MohammadTaghi Hajiaghayi:  
**Generalized Reductions: Making any Hierarchical Clustering Fair and Balanced with Low Cost.** 17218-17242
- Boris Knyazev, Doha Hwang, Simon Lacoste-Julien:  
**Can We Scale Transformers to Predict Parameters of Diverse ImageNet Models?** 17243-17259
- Tomás Kocák, Alexandra Carpentier:  
**Online Learning with Feedback Graphs: The True Shape of Regret.** 17260-17282
- Jing Yu Koh, Ruslan Salakhutdinov, Daniel Fried:  
**Grounding Language Models to Images for Multimodal Inputs and Outputs.** 17283-17300
- Jonas Köhler, Michele Invernizzi, Pim de Haan, Frank Noé:  
**Rigid Body Flows for Sampling Molecular Crystal Structures.** 17301-17326
- Nils Kohring, Fabian Raoul Pieroth, Martin Bichler:  
**Enabling First-Order Gradient-Based Learning for Equilibrium Computation in Markets.** 17327-17342
- Anastasia Koloskova, Hadrien Hendrikx, Sebastian U. Stich:  
**Revisiting Gradient Clipping: Stochastic bias and tight convergence guarantees.** 17343-17363



- Christian Komusiewicz, Pascal Kunz, Frank Sommer, Manuel Sorge:  
**On Computing Optimal Tree Ensembles.** 17364-17374
- Kezhi Kong, Juhai Chen, John Kirchenbauer, Renkun Ni, C. Bayan Bruss, Tom Goldstein:  
**GOAT: A Global Transformer on Large-scale Graphs.** 17375-17390
- Lingkai Kong, Jiaying Cui, Haotian Sun, Yuchen Zhuang, B. Aditya Prakash, Chao Zhang:  
**Autoregressive Diffusion Model for Graph Generation.** 17391-17408
- Xiangzhe Kong, Wenbing Huang, Yang Liu:  
**End-to-End Full-Atom Antibody Design.** 17409-17429
- Insung Kong, Yuha Park, Joonhyuk Jung, Kwonsang Lee, Yongdai Kim:  
**Covariate balancing using the integral probability metric for causal inference.** 17430-17461
- Insung Kong, Dongyoon Yang, Jongjin Lee, Ilsang Ohn, Gyuseung Baek, Yongdai Kim:  
**Masked Bayesian Neural Networks : Theoretical Guarantee and its Posterior Inference.** 17462-17491
- Tatsuya Konishi, Mori Kurokawa, Chihiro Ono, Zixuan Ke, Gyuhak Kim , Bing Liu:  
**Parameter-Level Soft-Masking for Continual Learning.** 17492-17505
- Tomasz Korbak, Kejian Shi, Angelica Chen, Rasika Vinayak Bhalerao, Christopher L. Buckley, Jason Phang, Samuel R. Bowman, Ethan Perez:  
**Pretraining Language Models with Human Preferences.** 17506-17533
- Ezgi Korkmaz, Jonah Brown-Cohen:  
**Detecting Adversarial Directions in Deep Reinforcement Learning to Make Robust Decisions.** 17534-17543
- Arthur Kosmala, Johannes Gasteiger, Nicholas Gao, Stephan Günnemann:  
**Ewald-based Long-Range Message Passing for Molecular Graphs.** 17544-17563
- Akim Kotelnikov, Dmitry Baranchuk, Ivan Rubachev, Artem Babenko:  
**TabDDPM: Modelling Tabular Data with Diffusion Models.** 17564-17579
- Vignesh Kothapalli:  
**Randomized Schur Complement Views for Graph Contrastive Learning.** 17580-17614
- Yiwen Kou, Zixiang Chen, Yuanzhou Chen, Quanquan Gu:  
**Benign Overfitting in Two-layer ReLU Convolutional Neural Networks.** 17615-17659
- Batuhan Koyuncu, Pablo Sánchez-Martín, Ignacio Peis, Pablo M. Olmos, Isabel Valera:  
**Variational Mixture of HyperGenerators for Learning Distributions over Functions.** 17660-17683
- Itai Kreisler, Mor Shpigel Nacson, Daniel Soudry, Yair Carmon:  
**Gradient Descent Monotonically Decreases the Sharpness of Gradient Flow Solutions in Scalar Networks and Beyond.** 17684-17744
- Heiner Kremer, Yassine Nemmour, Bernhard Schölkopf, Jia-Jie Zhu:  
**Estimation Beyond Data Reweighting: Kernel Method of Moments.** 17745-17783
- Walid Krichene, Prateek Jain, Shuang Song, Mukund Sundararajan, Abhradeep Guha Thakurta, Li Zhang:  
**Multi-Task Differential Privacy Under Distribution Skew.** 17784-17807
- Satyapriya Krishna, Jiaqi Ma, Himabindu Lakkaraju:  
**Towards Bridging the Gaps between the Right to Explanation and the Right to be Forgotten.** 17808-17826
- Sanjukta Krishnagopal, Luana Ruiz:  
**Graph Neural Tangent Kernel: Convergence on Large Graphs.** 17827-17841
- Siddarth Krishnamoorthy, Satvik Mehul Mashkaria, Aditya Grover:  
**Diffusion Models for Black-Box Optimization.** 17842-17857
- Dmitrii Krylov, Pooya Khajeh, Junhan Ouyang, Thomas Reeves, Tongkai Liu, Hiba Ajmal, Hamidreza Aghasi, Roy Fox:  
**Learning to Design Analog Circuits to Meet Threshold Specifications.** 17858-17873
- Qi Kuang, Zhoufan Zhu, Liwen Zhang, Fan Zhou:  
**Variance Control for Distributional Reinforcement Learning.** 17874-17895
- Kalle Kujanpää, Joni Pajarinen, Alexander Ilin:  
**Hierarchical Imitation Learning with Vector Quantized Models.** 17896-17919
- Vladimir Kulikov, Shahar Yadin, Matan Kleiner, Tomer Michaeli:  
**SinDDM: A Single Image Denoising Diffusion Model.** 17920-17930
- Sean Kulinski, David I. Inouye:  
**Towards Explaining Distribution Shifts.** 17931-17952
- Manoj Kumar, Anurag Sharma, Shashwat Saxena, Sandeep Kumar:  
**Featured Graph Coarsening with Similarity Guarantees.** 17953-17975
- Andrey Kurenkov, Michael Lingelbach, Tanmay Agarwal, Emily Jin, Chengshu Li, Ruohan Zhang, Li Fei-Fei, Jiajun Wu, Silvio Savarese, Roberto Martín-Martín:  
**Modeling Dynamic Environments with Scene Graph Memory.** 17976-17993
- Emirhan Kurtulus, Zichao Li, Yann N. Dauphin, Ekin Dogus Cubuk:  
**Tied-Augment: Controlling Representation Similarity Improves Data Augmentation.** 17994-18007



- Oskar Kviman, Ricky Molén, Alexandra Hotti, Semih Kurt, Víctor Elvira, Jens Lagergren:  
**Cooperation in the Latent Space: The Benefits of Adding Mixture Components in Variational Autoencoders.** 18008-18022
- Minseop Kwak, Jiuhan Song, Seungryong Kim:  
**GeCoNeRF: Few-shot Neural Radiance Fields via Geometric Consistency.** 18023-18036
- Sehyun Kwon, Joo Young Choi, Ernest K. Ryu:  
**Rotation and Translation Invariant Representation Learning with Implicit Neural Representations.** 18037-18056
- Jeongyeol Kwon, Yonathan Efroni, Constantine Caramanis, Shie Mannor:  
**Reward-Mixing MDPs with Few Latent Contexts are Learnable.** 18057-18082
- Jeongyeol Kwon, Dohyun Kwon, Stephen Wright, Robert D. Nowak:  
**A Fully First-Order Method for Stochastic Bilevel Optimization.** 18083-18113
- Dohyun Kwon, Hanbaek Lyu:  
**Complexity of Block Coordinate Descent with Proximal Regularization and Applications to Wasserstein CP-dictionary Learning.** 18114-18134
- Yongchan Kwon, James Zou:  
**Data-OOB: Out-of-bag Estimate as a Simple and Efficient Data Value.** 18135-18152
- Aqeel Labash, Florian Stelzer, Daniel Majoral, Raul Vicente Zafra:  
**Emergence of Adaptive Circadian Rhythms in Deep Reinforcement Learning.** 18153-18170
- Sébastien Lachapelle, Tristan Deleu, Divyat Mahajan, Ioannis Mitliagkas, Yoshua Bengio, Simon Lacoste-Julien, Quentin Bertrand:  
**Synergies between Disentanglement and Sparsity: Generalization and Identifiability in Multi-Task Learning.** 18171-18206
- Steinar Laenen, Bogdan-Adrian Manghiuc, He Sun:  
**Nearly-Optimal Hierarchical Clustering for Well-Clustered Graphs.** 18207-18249
- Marc Lafon, Elias Ramzi, Clément Rambour, Nicolas Thome:  
**Hybrid Energy Based Model in the Feature Space for Out-of-Distribution Detection.** 18250-18268
- Salem Lahlou, Tristan Deleu, Pablo Lemos, Dinghuai Zhang, Alexandra Volokhova, Alex Hernández-García, Léna Néhale Ezzine, Yoshua Bengio, Nikolay Malkin:  
**A theory of continuous generative flow networks.** 18269-18300
- Jinlin Lai, Javier Burroni, Hui Guan, Daniel Sheldon:  
**Automatically marginalized MCMC in probabilistic programming.** 18301-18318
- Yuhang Lai, Chengxi Li, Yiming Wang, Tianyi Zhang, Ruiqi Zhong, Luke Zettlemoyer, Wen-Tau Yih, Daniel Fried, Sida I. Wang, Tao Yu:  
**DS-1000: A Natural and Reliable Benchmark for Data Science Code Generation.** 18319-18345
- Yao Lai, Jinxin Liu, Zhentao Tang, Bin Wang, Jianye Hao, Ping Luo:  
**ChiPFormer: Transferable Chip Placement via Offline Decision Transformer.** 18346-18364
- Chieh-Hsin Lai, Yuhta Takida, Naoki Murata, Toshimitsu Uesaka, Yuki Mitsufuji, Stefano Ermon:  
**FP-Diffusion: Improving Score-based Diffusion Models by Enforcing the Underlying Score Fokker-Planck Equation.** 18365-18398
- Clément Lalanne, Aurélien Garivier, Rémi Gribonval:  
**Private Statistical Estimation of Many Quantiles.** 18399-18418
- Henry Lam, Zhenyuan Liu:  
**Bootstrap in High Dimension with Low Computation.** 18419-18453
- Kevin H. Lam, Christian J. Walder, Spiridon I. Penev, Richard Nock:  
**LegendreTron: Uprising Proper Multiclass Loss Learning.** 18454-18470
- Andre Lamurias, Alessandro Tibo, Katja Hose, Mads Albertsen, Thomas Dyhre Nielsen:  
**Metagenomic Binning using Connectivity-constrained Variational Autoencoders.** 18471-18481
- Tal Lancewicki, Aviv Rosenberg, Dmitry Sotnikov:  
**Delay-Adapted Policy Optimization and Improved Regret for Adversarial MDP with Delayed Bandit Feedback.** 18482-18534
- Robert Tjarko Lange, Henning Sprekeler:  
**Lottery Tickets in Evolutionary Optimization: On Sparse Backpropagation-Free Trainability.** 18535-18547
- Romain Laroche, Remi Tachet des Combes:  
**On the Occupancy Measure of Non-Markovian Policies in Continuous MDPs.** 18548-18562
- Alexandra Anna Lassota, Alexander Lindermayr, Nicole Megow, Jens Schlöter:  
**Minimalistic Predictions to Schedule Jobs with Online Precedence Constraints.** 18563-18583
- Silvio Lattanzi, Ola Svensson, Sergei Vassilvitskii:  
**Speeding Up Bellman Ford via Minimum Violation Permutations.** 18584-18598
- Niklas Lauffer, Ameesh Shah, Micah Carroll, Michael D. Dennis, Stuart Russell:  
**Who Needs to Know? Minimal Knowledge for Optimal Coordination.** 18599-18613



- Jonathan Wilder Lavington, Sharan Vaswani, Reza Babanezhad Harikandeh, Mark Schmidt, Nicolas Le Roux:  
**Target-based Surrogates for Stochastic Optimization.** 18614-18651
- Connor Lawless, Oktay Günlük:  
**Cluster Explanation via Polyhedral Descriptions.** 18652-18666
- Phuong-Hang Le, Hongyu Gong, Changhan Wang, Juan Pino, Benjamin Lecouteux, Didier Schwab:  
**Pre-training for Speech Translation: CTC Meets Optimal Transport.** 18667-18685
- Charline Le Lan, Stephen Tu, Mark Rowland, Anna Harutyunyan, Rishabh Agarwal, Marc G. Bellemare, Will Dabney:  
**Bootstrapped Representations in Reinforcement Learning.** 18686-18713
- Tosca Lechner, Ruth Urner, Shai Ben-David:  
**Strategic Classification with Unknown User Manipulations.** 18714-18732
- Jonathan Lee, Alekh Agarwal, Christoph Dann, Tong Zhang:  
**Learning in POMDPs is Sample-Efficient with Hindsight Observability.** 18733-18773
- Soo Yong Lee, Fanchen Bu, Jaemin Yoo, Kijung Shin:  
**Towards Deep Attention in Graph Neural Networks: Problems and Remedies.** 18774-18795
- Jaejun Lee , Chanyoung Chung, Joyce Jiyoung Whang:  
**InGram: Inductive Knowledge Graph Embedding via Relation Graphs.** 18796-18809
- Jongyeong Lee , Junya Honda, Chao-Kai Chiang, Masashi Sugiyama:  
**Optimality of Thompson Sampling with Noninformative Priors for Pareto Bandits.** 18810-18851
- Namkyeong Lee, Dongmin Hyun, Gyoung S. Na, Sungwon Kim, Junseok Lee, Chanyoung Park:  
**Conditional Graph Information Bottleneck for Molecular Relational Learning.** 18852-18871
- Seul Lee, Jaehyeong Jo, Sung Ju Hwang:  
**Exploring Chemical Space with Score-based Out-of-distribution Generation.** 18872-18892
- Kenton Lee, Mandar Joshi, Iulia Raluca Turc, Hexiang Hu, Fangyu Liu, Julian Martin Eisenschlos, Urvashi Khandelwal, Peter Shaw, Ming-Wei Chang, Kristina Toutanova:  
**Pix2Struct: Screenshot Parsing as Pretraining for Visual Language Understanding.** 18893-18912
- Jung Hyun Lee, Jeonghoon Kim, Se Jung Kwon, Dongsoo Lee:  
**FlexRound: Learnable Rounding based on Element-wise Division for Post-Training Quantization.** 18913-18939
- Chaejeong Lee, Jayoung Kim, Noseong Park:  
**CoDi: Co-evolving Contrastive Diffusion Models for Mixed-type Tabular Synthesis.** 18940-18956
- Sangyun Lee, Beomsu Kim, Jong Chul Ye:  
**Minimizing Trajectory Curvature of ODE-based Generative Models.** 18957-18973
- Hangbin Lee, Youngjo Lee:  
**H-Likelihood Approach to Deep Neural Networks with Temporal-Spatial Random Effects for High-Cardinality Categorical Features.** 18974-18987
- Hojoon Lee, Koanho Lee, Dongyoon Hwang, Hyunho Lee, Byungkun Lee, Jaegul Choo:  
**On the Importance of Feature Decorrelation for Unsupervised Representation Learning in Reinforcement Learning.** 18988-19009
- Seewoo Lee, Garam Lee, Jung Woo Kim, Junbum Shin, Mun-Kyu Lee:  
**HETAL: Efficient Privacy-preserving Transfer Learning with Homomorphic Encryption.** 19010-19035
- Yoonjoo Lee, Kyungjae Lee, Sunghyun Park, Dasol Hwang, Jaehyeon Kim, Hong-In Lee, Moontae Lee:  
**QASA: Advanced Question Answering on Scientific Articles.** 19036-19052
- Donghwan Lee, Behrad Moniri, Xinpeng Huang, Edgar Dobriban, Hamed Hassani:  
**Demystifying Disagreement-on-the-Line in High Dimensions.** 19053-19093
- Wonyeol Lee, Sejun Park, Alex Aiken:  
**On the Correctness of Automatic Differentiation for Neural Networks with Machine-Representable Parameters.** 19094-19140
- Sungyoon Lee, Jinseong Park , Jaewook Lee:  
**Implicit Jacobian regularization weighted with impurity of probability output.** 19141-19184
- Sang-Hyun Lee , Seung-Woo Seo:  
**Unsupervised Skill Discovery for Learning Shared Structures across Changing Environments.** 19185-19199
- Yunwen Lei, Tianbao Yang, Yiming Ying, Ding-Xuan Zhou:  
**Generalization Analysis for Contrastive Representation Learning.** 19200-19227
- Gal Leibovich, Guy Jacob, Or Avner, Gal Novik, Aviv Tamar:  
**Learning Control by Iterative Inversion.** 19228-19255
- Pablo Lemos, Adam Coogan, Yashar Hezaveh, Laurence Perreault Levasseur:  
**Sampling-Based Accuracy Testing of Posterior Estimators for General Inference.** 19256-19273
- Yaniv Leviathan, Matan Kalman, Yossi Matias:  
**Fast Inference from Transformers via Speculative Decoding.** 19274-19286
- Orin Levy, Alon Cohen, Asaf B. Cassel, Yishay Mansour:  
**Efficient Rate Optimal Regret for Adversarial Contextual MDPs Using Online Function Approximation.** 19287-19314



- Dan Ley, Saumitra Mishra, Daniele Magazzeni:  
**GLOBE-CE: A Translation Based Approach for Global Counterfactual Explanations.** 19315-19342
- Guihong Li, Kartikeya Bhardwaj, Yuedong Yang, Radu Marculescu:  
**TIPS: Topologically Important Path Sampling for Anytime Neural Networks.** 19343-19359
- Anqi Li, Byron Boots, Ching-An Cheng:  
**MAHALO: Unifying Offline Reinforcement Learning and Imitation Learning from Observations.** 19360-19384
- Alexander Cong Li, Ellis Langham Brown, Alexei A. Efros, Deepak Pathak:  
**Internet Explorer: Targeted Representation Learning on the Open Web.** 19385-19406
- Yuxin Li, Wenchao Chen, Bo Chen, Dongsheng Wang, Long Tian, Mingyuan Zhou:  
**Prototype-oriented unsupervised anomaly detection for multivariate time series.** 19407-19424
- Yichen Li, Peter Yichen Chen, Tao Du, Wojciech Matusik:  
**Learning Preconditioners for Conjugate Gradient PDE Solvers.** 19425-19439
- Zechu Li, Tao Chen, Zhang-Wei Hong, Anurag Ajay, Pulkit Agrawal:  
**Parallel Q-Learning: Scaling Off-policy Reinforcement Learning under Massively Parallel Simulation.** 19440-19459
- Li'ang Li, Yifei Duan, Guanghua Ji, Yongqiang Cai:  
**Minimum Width of Leaky-ReLU Neural Networks for Uniform Universal Approximation.** 19460-19470
- Tianlin Li, Qing Guo, Aishan Liu, Mengnan Du, Zhiming Li, Yang Liu:  
**FAIRER: Fairness as Decision Rationale Alignment.** 19471-19489
- Pengyi Li, Jianye Hao, Hongyao Tang, Yan Zheng, Xian Fu:  
**RACE: Improve Multi-Agent Reinforcement Learning with Representation Asymmetry and Collaborative Evolution.** 19490-19503
- Qinbin Li, Bingsheng He, Dawn Song:  
**Adversarial Collaborative Learning on Non-IID Features.** 19504-19526
- Donghao Li, Ruiquan Huang, Cong Shen, Jing Yang:  
**Near-optimal Conservative Exploration in Reinforcement Learning under Episode-wise Constraints.** 19527-19564
- Yingcong Li, Muhammed Emrullah Ildiz, Dimitris Papailiopoulos, Samet Oymak:  
**Transformers as Algorithms: Generalization and Stability in In-context Learning.** 19565-19594
- Rui Li, S. T. John, Arno Solin:  
**Improving Hyperparameter Learning under Approximate Inference in Gaussian Process Models.** 19595-19615
- Shouheng Li, Dongwoo Kim, Qing Wang:  
**Local Vertex Colouring Graph Neural Networks.** 19616-19637
- Xiaoyun Li, Ping Li:  
**Analysis of Error Feedback in Federated Non-Convex Optimization with Biased Compression: Fast Convergence and Partial Participation.** 19638-19688
- Yuchen Li, Yuanzhi Li, Andrej Risteski:  
**How Do Transformers Learn Topic Structure: Towards a Mechanistic Understanding.** 19689-19729
- Junnan Li, Dongxu Li, Silvio Savarese, Steven C. H. Hoi:  
**BLIP-2: Bootstrapping Language-Image Pre-training with Frozen Image Encoders and Large Language Models.** 19730-19742
- Junfan Li , Shizhong Liao:  
**Nearly Optimal Algorithms with Sublinear Computational Complexity for Online Kernel Regression.** 19743-19766
- Zexi Li, Tao Lin, Xinyi Shang, Chao Wu:  
**Revisiting Weighted Aggregation in Federated Learning with Neural Networks.** 19767-19788
- Shaojie Li, Yong Liu:  
**Distribution-dependent McDiarmid-type Inequalities for Functions of Unbounded Interaction.** 19789-19810
- Jian Li, Yong Liu, Weiping Wang:  
**Optimal Convergence Rates for Agnostic Nyström Kernel Learning.** 19811-19836
- Yige Li, Xixiang Lyu, Xingjun Ma, Nodens Koren, Lingjuan Lyu, Bo Li, Yu-Gang Jiang:  
**Reconstructive Neuron Pruning for Backdoor Defense.** 19837-19854
- Shibo Li, Michael Penwarden, Yiming Xu, Conor Tillinghast, Akil Narayan, Mike Kirby, Shandian Zhe:  
**Meta Learning of Interface Conditions for Multi-Domain Physics-Informed Neural Networks.** 19855-19881
- Aodong Li, Chen Qiu, Marius Kloft, Padhraic Smyth, Stephan Mandt, Maja Rudolph:  
**Deep Anomaly Detection under Labeling Budget Constraints.** 19882-19910
- Jiahang Li, Yakun Song, Xiang Song, David Wipf:  
**On the Initialization of Graph Neural Networks.** 19911-19931
- Xiaoxiao Li, Zhao Song, Jiaming Yang:  
**Federated Adversarial Learning: A Framework with Convergence Analysis.** 19932-19959
- Ming Li, Sho Sonoda, Feilong Cao, Yu Guang Wang, Jiye Liang:  
**How Powerful are Shallow Neural Networks with Bandlimited Random Weights?** 19960-19981



- Xiantao Li, Chunhao Wang:  
**Efficient Quantum Algorithms for Quantum Optimal Control.** 19982-19994
- Yunfan Li, Yiran Wang, Yu Cheng, Lin Yang:  
**Low-Switching Policy Gradient with Exploration via Online Sensitivity Sampling.** 19995-20034
- Wenhao Li, Xiangfeng Wang, Bo Jin, Hongyuan Zha:  
**Hierarchical Diffusion for Offline Decision Making.** 20035-20064
- Wanshan Li, Daren Wang, Alessandro Rinaldo:  
**Divide and Conquer Dynamic Programming: An Almost Linear Time Change Point Detection Methodology in High Dimensions.** 20065-20148
- Siyuan Li, Di Wu, Fang Wu, Zelin Zang, Stan Z. Li:  
**Architecture-Agnostic Masked Image Modeling - From ViT back to CNN.** 20149-20167
- Peizhao Li, Ethan Xia, Hongfu Liu:  
**Learning Antidote Data to Individual Unfairness.** 20168-20181
- Haoxuan Li, Yanghao Xiao, Chunyuan Zheng, Peng Wu, Peng Cui:  
**Propensity Matters: Measuring and Enhancing Balancing for Recommendation.** 20182-20194
- Yuwen Li, Miao Xiong, Bryan Hooi:  
**GraphCleaner: Detecting Mislabelled Samples in Popular Graph Learning Benchmarks.** 20195-20209
- Zichong Li, Yanbo Xu, Simiao Zuo, Haoming Jiang, Chao Zhang, Tuo Zhao, Hongyuan Zha:  
**SMURF-THP: Score Matching-based Uncertainty quantification for Transformer Hawkes Process.** 20210-20220
- Shengshi Li, Lin Yang:  
**Horizon-free Learning for Markov Decision Processes and Games: Stochastically Bounded Rewards and Improved Bounds.** 20221-20252
- Weixin Li, Xiaodong Yang:  
**Transcendental Idealism of Planner: Evaluating Perception from Planning Perspective for Autonomous Driving.** 20253-20275
- Pengfei Li, Jianyi Yang, Shaolei Ren:  
**Learning for Edge-Weighted Online Bipartite Matching with Robustness Guarantees.** 20276-20295
- Songze Li, Duanyi Yao, Jin Liu:  
**FedVS: Straggler-Resilient and Privacy-Preserving Vertical Federated Learning for Split Models.** 20296-20311
- Jiayang Li, Jing Yu , Boyi Liu, Yu Marco Nie, Zhaoran Wang:  
**Achieving Hierarchy-Free Approximation for Bilevel Programs with Equilibrium Constraints.** 20312-20335
- Yixiao Li, Yifan Yu, Qingru Zhang, Chen Liang, Pengcheng He, Weizhu Chen, Tuo Zhao:  
**LoSparse: Structured Compression of Large Language Models based on Low-Rank and Sparse Approximation.** 20336-20350
- Chris Junchi Li, Huizhuo Yuan, Gauthier Gidel, Quanquan Gu, Michael I. Jordan:  
**Nesterov Meets Optimism: Rate-Optimal Separable Minimax Optimization.** 20351-20383
- Chao Li, Junhua Zeng, Chunmei Li, Cesar F. Caiafa, Qibin Zhao:  
**Alternating Local Enumeration (TnALE): Solving Tensor Network Structure Search with Fewer Evaluations.** 20384-20411
- Qiyang Li, Yuexiang Zhai, Yi Ma, Sergey Levine:  
**Understanding the Complexity Gains of Single-Task RL with a Curriculum.** 20412-20451
- Mingjie Li, Quanshi Zhang:  
**Does a Neural Network Really Encode Symbolic Concepts?** 20452-20469
- Yang Li, Shao Zhang , Jichen Sun, Yali Du, Ying Wen, Xinbing Wang, Wei Pan:  
**Cooperative Open-ended Learning Framework for Zero-Shot Coordination.** 20470-20484
- Jiachen Li, Edwin Zhang, Ming Yin, Qinrun Bai, Yu-Xiang Wang, William Yang Wang:  
**Offline Reinforcement Learning with Closed-Form Policy Improvement Operators.** 20485-20528
- Shaoang Li, Lan Zhang, Yingqi Yu, Xiangyang Li:  
**Optimal Arms Identification with Knapsacks.** 20529-20555
- Mengdi Li , Xufeng Zhao , Jae Hee Lee, Cornelius Weber, Stefan Wermter:  
**Internally Rewarded Reinforcement Learning.** 20556-20574
- Haoxuan Li, Chunyuan Zheng, Yixiao Cao, Zhi Geng, Yue Liu, Peng Wu:  
**Trustworthy Policy Learning under the Counterfactual No-Harm Criterion.** 20575-20598
- Shuangtong Li, Tianyi Zhou, Xinmei Tian, Dacheng Tao:  
**Structured Cooperative Learning with Graphical Model Priors.** 20599-20622
- Enming Liang, Minghua Chen, Steven H. Low:  
**Low Complexity Homeomorphic Projection to Ensure Neural-Network Solution Feasibility for Optimization over (Non-)Convex Set.** 20623-20649
- Weixuan Liang, Xinwang Liu, Yong Liu, Chuan Ma, Yunping Zhao, Zhe Liu, En Zhu:  
**Consistency of Multiple Kernel Clustering.** 20650-20676



- Hao Liang, Zhi-Quan Luo:  
**A Distribution Optimization Framework for Confidence Bounds of Risk Measures.** 20677-20705
- Weixin Liang, Yining Mao, Yongchan Kwon, Xinyu Yang, James Zou:  
**Accuracy on the Curve: On the Nonlinear Correlation of ML Performance Between Data Subpopulations.** 20706-20724
- Zhixuan Liang , Yao Mu, Mingyu Ding, Fei Ni, Masayoshi Tomizuka, Ping Luo:  
**AdaptDiffuser: Diffusion Models as Adaptive Self-evolving Planners.** 20725-20745
- Youwei Liang, Kevin Stone, Ali Shameli, Chris Cummins, Mostafa Elhoushi, Jiadong Guo, Benoit Steiner, Xiaomeng Yang, Pengtao Xie, Hugh James Leather, Yuandong Tian:  
**Learning Compiler Pass Orders using Coreset and Normalized Value Prediction.** 20746-20762
- Chumeng Liang, Xiaoyu Wu, Yang Hua, Jiaru Zhang, Yiming Xue, Tao Song, Zhengui Xue, Ruhui Ma, Haibing Guan:  
**Adversarial Example Does Good: Preventing Painting Imitation from Diffusion Models via Adversarial Examples.** 20763-20786
- James Chenhao Liang, Tianfei Zhou, Dongfang Liu, Wenguan Wang:  
**CLUSTSEG: Clustering for Universal Segmentation.** 20787-20809
- Ziyi Liang, Yanfei Zhou, Matteo Sesia:  
**Conformal Inference is (almost) Free for Neural Networks Trained with Early Stopping.** 20810-20851
- Chen Liang, Simiao Zuo, Qingru Zhang, Pengcheng He, Weizhu Chen, Tuo Zhao:  
**Less is More: Task-aware Layer-wise Distillation for Language Model Compression.** 20852-20867
- Luofeng Liao, Christian Kroer:  
**Statistical Inference and A/B Testing for First-Price Pacing Equilibria.** 20868-20905
- Christopher Liao, Theodoros Tsiligkaridis, Brian Kulis:  
**Supervised Metric Learning to Rank for Retrieval via Contextual Similarity Optimization.** 20906-20938
- Yun-Hsuan Lien, Ping-Chun Hsieh, Yu-Shuen Wang:  
**Revisiting Domain Randomization via Relaxed State-Adversarial Policy Optimization.** 20939-20949
- Valentin Liévin, Andreas Geert Motzfeldt, Ida Riis Jensen, Ole Winther:  
**Variational Open-Domain Question Answering.** 20950-20977
- Yeqing Lin, Mohammed AlQuraishi:  
**Generating Novel, Designable, and Diverse Protein Structures by Equivariantly Diffusing Oriented Residue Clouds.** 20978-21002
- Ya-Wei Eileen Lin, Ronald R. Coifman, Gal Mishne , Ronen Talmon:  
**Hyperbolic Diffusion Embedding and Distance for Hierarchical Representation Learning.** 21003-21025
- Wu Lin, Valentin Duruisseaux, Melvin Leok, Frank Nielsen, Mohammad Emtiyaz Khan, Mark Schmidt:  
**Simplifying Momentum-based Positive-definite Submanifold Optimization with Applications to Deep Learning.** 21026-21050
- Zhenghao Lin, Yeyun Gong, Yelong Shen, Tong Wu, Zhihao Fan, Chen Lin, Nan Duan, Weizhu Chen:  
**Text Generation with Diffusion Language Models: A Pre-training Approach with Continuous Paragraph Denoise.** 21051-21064
- Weiwei Lin, Chenhang He, Man-Wai Mak, Youzhi Tu:  
**Self-supervised Neural Factor Analysis for Disentangling Utterance-level Speech Representations.** 21065-21077
- Sen Lin, Peizhong Ju, Yingbin Liang, Ness B. Shroff:  
**Theory on Forgetting and Generalization of Continual Learning.** 21078-21100
- Cheuk Yin Lin, Chaobing Song, Jelena Diakonikolas:  
**Accelerated Cyclic Coordinate Dual Averaging with Extrapolation for Composite Convex Optimization.** 21101-21126
- Qian Lin, Bo Tang, Zifan Wu , Chao Yu, Shangqin Mao, Qianlong Xie, Xingxing Wang, Dong Wang:  
**Safe Offline Reinforcement Learning with Real-Time Budget Constraints.** 21127-21152
- Alexander Lin, Bahareh Tolooshams, Yves F. Atchadé, Demba E. Ba:  
**Probabilistic Unrolling: Scalable, Inverse-Free Maximum Likelihood Estimation for Latent Gaussian Models.** 21153-21181
- Zhen Lin, Shubhendu Trivedi, Cao Xiao, Jimeng Sun:  
**Fast Online Value-Maximizing Prediction Sets with Conformal Cost Control.** 21182-21203
- Chieh Hubert Lin, Hung-Yu Tseng, Hsin-Ying Lee, Maneesh Kumar Singh, Ming-Hsuan Yang:  
**Unveiling The Mask of Position-Information Pattern Through the Mist of Image Features.** 21204-21222
- Xiaoqiang Lin, Xinyi Xu, See-Kiong Ng, Chuan-Sheng Foo, Bryan Kian Hsiang Low:  
**Fair yet Asymptotically Equal Collaborative Learning.** 21223-21259
- Yuchao Lin, Keqiang Yan, Youzhi Luo, Yi Liu, Xiaoning Qian, Shuiwang Ji:  
**Efficient Approximations of Complete Interatomic Potentials for Crystal Property Prediction.** 21260-21287
- Xi Lin, Zhiyuan Yang, Xiaoyuan Zhang, Qingfu Zhang:  
**Continuation Path Learning for Homotopy Optimization.** 21288-21311



- Alexander Lindermayr, Nicole Megow, Martin Rapp:  
**Speed-Oblivious Online Scheduling: Knowing (Precise) Speeds is not Necessary.** 21312-21334
- Hongyi Ling, Zhimeng Jiang, Meng Liu, Shuiwang Ji, Na Zou:  
**Graph Mixup with Soft Alignments.** 21335-21349
- Chen Ling, Junji Jiang, Junxiang Wang, My T. Thai, Renhao Xue, James Song, Meikang Qiu, Liang Zhao:  
**Deep Graph Representation Learning and Optimization for Influence Maximization.** 21350-21361
- Hao Liu, Pieter Abbeel:  
**Emergent Agentic Transformer from Chain of Hindsight Experience.** 21362-21374
- Tommy Liu, Amanda S. Barnard:  
**Shapley Based Residual Decomposition for Instance Analysis.** 21375-21387
- Tennison Liu, Jeroen Berrevoets , Zhaozhi Qian, Mihaela van der Schaar:  
**Learning Representations without Compositional Assumptions.** 21388-21403
- Yuchen Liu, Chen Chen, Lingjuan Lyu, Fangzhao Wu, Sai Wu, Gang Chen:  
**Byzantine-Robust Learning on Heterogeneous Data via Gradient Splitting.** 21404-21425
- Jialin Liu, Xiaohan Chen, Zhangyang Wang, Wotao Yin, HanQin Cai:  
**Towards Constituting Mathematical Structures for Learning to Optimize.** 21426-21449
- Haohe Liu, Zehua Chen, Yi Yuan, Xinhao Mei, Xubo Liu, Danilo P. Mandic, Wenwu Wang, Mark D. Plumbley:  
**AudioLDM: Text-to-Audio Generation with Latent Diffusion Models.** 21450-21474
- Yang Liu, Hao Cheng, Kun Zhang:  
**Identifiability of Label Noise Transition Matrix.** 21475-21496
- Shengchao Liu, Weitao Du, Zhi-Ming Ma, Hongyu Guo, Jian Tang:  
**A Group Symmetric Stochastic Differential Equation Model for Molecule Multi-modal Pretraining.** 21497-21526
- Xing Liu, Andrew B. Duncan, Axel Gandy:  
**Using Perturbation to Improve Goodness-of-Fit Tests based on Kernelized Stein Discrepancy.** 21527-21547
- Zhiheng Liu, Ruili Feng, Kai Zhu, Yifei Zhang, Kecheng Zheng, Yu Liu, Deli Zhao, Jingren Zhou, Yang Cao:  
**Cones: Concept Neurons in Diffusion Models for Customized Generation.** 21548-21566
- Weiming Liu, Haobo Fu, Qiang Fu, Wei Yang:  
**Opponent-Limited Online Search for Imperfect Information Games.** 21567-21585
- Zuxin Liu, Zijian Guo, Zhepeng Cen, Huan Zhang, Yihang Yao, Hanjiang Hu, Ding Zhao:  
**Towards Robust and Safe Reinforcement Learning with Benign Off-policy Data.** 21586-21610
- Zuxin Liu, Zijian Guo, Yihang Yao, Zhepeng Cen, Wenhao Yu, Tingnan Zhang, Ding Zhao:  
**Constrained Decision Transformer for Offline Safe Reinforcement Learning.** 21611-21630
- Liang Liu, Yanan Guo, Youtao Zhang, Jun Yang:  
**Understanding and Defending Patched-based Adversarial Attacks for Vision Transformer.** 21631-21657
- Songming Liu, Zhongkai Hao, Chengyang Ying, Hang Su, Ze Cheng, Jun Zhu:  
**NUNO: A General Framework for Learning Parametric PDEs with Non-Uniform Data.** 21658-21671
- Guan-Ting Liu, En-Pei Hu, Pu-Jen Cheng, Hung-Yi Lee, Shao-Hua Sun:  
**Hierarchical Programmatic Reinforcement Learning via Learning to Compose Programs.** 21672-21697
- Yi Liu, Qirui Hu , Lei Ding, Linglong Kong:  
**Online Local Differential Private Quantile Inference via Self-normalization.** 21698-21714
- Dianbo Liu, Moksh Jain, Bonaventure F. P. Dossou, Qianli Shen, Salem Lahlou, Anirudh Goyal, Nikolay Malkin, Chris Chinene Emezue, Dinghuai Zhang, Nadhir Hassen, Xu Ji, Kenji Kawaguchi, Yoshua Bengio:  
**GFlowOut: Dropout with Generative Flow Networks.** 21715-21729
- Zhihong Liu, Hoang Anh Just, Xiangyu Chang, Xi Chen, Ruoxi Jia:  
**2D-Shapley: A Framework for Fragmented Data Valuation.** 21730-21755
- Chenxi Liu, Kun Kuang:  
**Causal Structure Learning for Latent Intervened Non-stationary Data.** 21756-21777
- Shikun Liu, Tianchun Li, Yongbin Feng, Nhan Tran, Han Zhao, Qiang Qiu, Pan Li:  
**Structural Re-weighting Improves Graph Domain Adaptation.** 21778-21793
- Yue Liu, Ke Liang, Jun Xia, Sihang Zhou, Xihong Yang, Xinwang Liu, Stan Z. Li:  
**Dink-Net: Neural Clustering on Large Graphs.** 21794-21812
- Shih-Yang Liu, Zechun Liu, Kwang-Ting Cheng:  
**Oscillation-free Quantization for Low-bit Vision Transformers.** 21813-21824
- Xuejie Liu, Anji Liu, Guy Van den Broeck, Yitao Liang:  
**Understanding the Distillation Process from Deep Generative Models to Tractable Probabilistic Circuits.** 21825-21838
- Risheng Liu, Yaohua Liu, Wei Yao, Shangzhi Zeng, Jin Zhang:  
**Averaged Method of Multipliers for Bi-Level Optimization without Lower-Level Strong Convexity.** 21839-21866
- Yongtuo Liu, Sara Magliacane, Miltiadis Kofinas, Efstratios Gavves:  
**Graph Switching Dynamical Systems.** 21867-21883



- Zijian Liu, Ta Duy Nguyen, Thien Hang Nguyen, Alina Ene , Huy Nguyen:  
**High Probability Convergence of Stochastic Gradient Methods.** 21884-21914
- Enshu Liu, Xuefei Ning, Zinan Lin, Huazhong Yang, Yu Wang:  
**OMS-DPM: Optimizing the Model Schedule for Diffusion Probabilistic Models.** 21915-21936
- Boyin Liu, Zhiqiang Pu, Yi Pan, Jianqiang Yi, Yanyan Liang, Du Zhang:  
**Lazy Agents: A New Perspective on Solving Sparse Reward Problem in Multi-agent Reinforcement Learning.** 21937-21950
- Zirui Liu, Shengyuan Chen, Kaixiong Zhou, Daochen Zha, Xiao Huang, Xia Hu:  
**RSC: Accelerate Graph Neural Networks Training via Randomized Sparse Computations.** 21951-21968
- Yuhang Liu , Ananda Theertha Suresh, Wennan Zhu, Peter Kairouz, Marco Gruteser:  
**Algorithms for bounding contribution for histogram estimation under user-level privacy.** 21969-21996
- Evan Zheran Liu, Sahaana Suri, Tong Mu, Allan Zhou, Chelsea Finn:  
**Simple Embodied Language Learning as a Byproduct of Meta-Reinforcement Learning.** 21997-22008
- Terrance Liu, Jingwu Tang, Giuseppe Vietri, Steven Wu:  
**Generating Private Synthetic Data with Genetic Algorithms.** 22009-22027
- Songtao Liu, Zhengkai Tu, Minkai Xu, Zuobai Zhang, Lu Lin, Rex Ying, Jian Tang, Peilin Zhao, Dinghao Wu:  
**FusionRetro: Molecule Representation Fusion via In-Context Learning for Retrosynthetic Planning.** 22028-22041
- Guan-Horng Liu, Arash Vahdat, De-An Huang, Evangelos A. Theodorou, Weili Nie, Anima Anandkumar:  
**I<sup>2</sup>SB: Image-to-Image Schrödinger Bridge.** 22042-22062
- Fanghui Liu, Luca Viano, Volkan Cevher:  
**What can online reinforcement learning with function approximation benefit from general coverage conditions?** 22063-22091
- Zhaoyan Liu, Noël Vouitsis, Satya Krishna Gorti, Jimmy Ba, Gabriel Loaiza-Ganem:  
**TRON: Translator Networks for 0-Shot Plug-and-Play Conditional Generation.** 22092-22112
- Chong Liu, Yu-Xiang Wang:  
**Global Optimization with Parametric Function Approximation.** 22113-22136
- Zichang Liu, Jue Wang, Tri Dao, Tianyi Zhou, Binhang Yuan, Zhao Song, Anshumali Shrivastava, Ce Zhang, Yuandong Tian, Christopher Ré, Beidi Chen:  
**Deja Vu: Contextual Sparsity for Efficient LLMs at Inference Time.** 22137-22176
- Hanwen Liu, Zhenyu Weng, Yuesheng Zhu, Yadong Mu:  
**Trapdoor Normalization with Irreversible Ownership Verification.** 22177-22187
- Hong Liu, Sang Michael Xie, Zhiyuan Li , Tengyu Ma:  
**Same Pre-training Loss, Better Downstream: Implicit Bias Matters for Language Models.** 22188-22214
- Tianyi Liu, Zihao Xu, Hao He, Guang-Yuan Hao, Guang-He Lee, Hao Wang:  
**Taxonomy-Structured Domain Adaptation.** 22215-22232
- Zhuang Liu, Zhiqiu Xu, Joseph Jin, Zhiqiang Shen, Trevor Darrell:  
**Dropout Reduces Underfitting.** 22233-22248
- Biao Liu, Ning Xu, Jiaqi Lv, Xin Geng:  
**Revisiting Pseudo-Label for Single-Positive Multi-Label Learning.** 22249-22265
- Guoqing Liu, Di Xue, Shufang Xie, Yingce Xia, Austin Tripp, Krzysztof Maziarz, Marwin H. S. Segler, Tao Qin, Zongzhang Zhang, Tie-Yan Liu:  
**Retrosynthetic Planning with Dual Value Networks.** 22266-22276
- Xin Liu, Zixian Yang, Lei Ying:  
**Online Nonstochastic Control with Adversarial and Static Constraints.** 22277-22288
- Tianci Liu, Tong Yang, Quan Zhang, Qi Lei:  
**Optimization for Amortized Inverse Problems.** 22289-22319
- Xuefeng Liu, Takuma Yoneda, Chaoqi Wang, Matthew R. Walter, Yuxin Chen:  
**Active Policy Improvement from Multiple Black-box Oracles.** 22320-22337
- Weitang Liu, Yi-Zhuang You, Ying-Wai Li, Jingbo Shang:  
**Gradient-based Wang-Landau Algorithm: A Novel Sampler for Output Distribution of Neural Networks over the Input Space.** 22338-22351
- Yicheng Liu, Tianyuan Yuan, Yue Wang, Yilun Wang, Hang Zhao:  
**VectorMapNet: End-to-end Vectorized HD Map Learning.** 22352-22369
- Xiangyu Liu, Kaiqing Zhang:  
**Partially Observable Multi-agent RL with (Quasi-)Efficiency: The Blessing of Information Sharing.** 22370-22419
- Zhuqing Liu, Xin Zhang, Prashant Khanduri, Songtao Lu, Jia Liu:  
**Prometheus: Taming Sample and Communication Complexities in Constrained Decentralized Stochastic Bilevel Learning.** 22420-22453
- Xuanzhou Liu, Lin Zhang, Jiaqi Sun, Yujiu Yang, Haiqin Yang:  
**D2Match: Leveraging Deep Learning and Degeneracy for Subgraph Matching.** 22454-22472



Zhuoran Liu, Zhengyu Zhao, Martha A. Larson:  
**Image Shortcut Squeezing: Countering Perturbative Availability Poisons with Compression.** 22473-22487

Mingzhou Liu, Xiangyu Zheng, Xinwei Sun, Fang Fang, Yizhou Wang:  
**Which Invariance Should We Transfer? A Causal Minimax Learning Approach.** 22488-22527

Zhenzhen Liu, Jin Peng Zhou, Yufan Wang, Kilian Q. Weinberger:  
**Unsupervised Out-of-Distribution Detection with Diffusion Inpainting.** 22528-22538

Zichuan Liu, Yuanyang Zhu, Chunlin Chen:  
**NA<sup>2</sup>Q: Neural Attention Additive Model for Interpretable Multi-Agent Q-Learning.** 22539-22558

Xutong Liu, Jinhang Zuo, Siwei Wang, John C. S. Lui, Mohammad Hajiesmaili, Adam Wierman, Wei Chen:  
**Contextual Combinatorial Bandits with Probabilistically Triggered Arms.** 22559-22593

Sam Lobel, Akhil Bagaria, George Konidaris:  
**Flipping Coins to Estimate Pseudocounts for Exploration in Reinforcement Learning.** 22594-22613

Charlotte Loh, Seungwook Han, Shivchander Sudalairaj, Rumen Dangovski, Kai Xu, Florian Wenzel, Marin Soljacic, Akash Srivastava:  
**Multi-Symmetry Ensembles: Improving Diversity and Generalization via Opposing Symmetries.** 22614-22630

Shayne Longpre, Le Hou, Tu Vu, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V. Le, Barret Zoph, Jason Wei, Adam Roberts:  
**The Flan Collection: Designing Data and Methods for Effective Instruction Tuning.** 22631-22648

Noel Loo, Ramin M. Hasani, Mathias Lechner, Daniela Rus:  
**Dataset Distillation with Convexified Implicit Gradients.** 22649-22674

Aaron Lou, Stefano Ermon:  
**Reflected Diffusion Models.** 22675-22701

David R. Lovell, Dimity Miller, Jaiden Capra, Andrew P. Bradley:  
**Never mind the metrics - what about the uncertainty? Visualising binary confusion matrix metric distributions to put performance in perspective.** 22702-22757

Songtao Lu:  
**Bilevel Optimization with Coupled Decision-Dependent Distributions.** 22758-22789

Yulong Lu:  
**Two-Scale Gradient Descent Ascent Dynamics Finds Mixed Nash Equilibria of Continuous Games: A Mean-Field Perspective.** 22790-22811

Yucheng Lu, Shivani Agrawal, Suvinay Subramanian, Oleg Rybakov, Christopher De Sa, Amir Yazdanbakhsh:  
**STEP: Learning N: M Structured Sparsity Masks from Scratch with Precondition.** 22812-22824

Cheng Lu, Huayu Chen, Jianfei Chen, Hang Su, Chongxuan Li, Jun Zhu:  
**Contrastive Energy Prediction for Exact Energy-Guided Diffusion Sampling in Offline Reinforcement Learning.** 22825-22855

Yiwei Lu, Gautam Kamath, Yaoliang Yu:  
**Exploring the Limits of Model-Targeted Indiscriminate Data Poisoning Attacks.** 22856-22879

Xudong Lu, Kaisen Pan, Ge Yan, Jiaming Shan, Wenjie Wu, Junchi Yan:  
**QAS-Bench: Rethinking Quantum Architecture Search and A Benchmark.** 22880-22898

Xuanchen Lu, Xiaolong Wang, Judith E. Fan:  
**Learning Dense Correspondences between Photos and Sketches.** 22899-22916

Chris Lu, Timon Willi, Alistair Letcher, Jakob Nicolaus Foerster:  
**Adversarial Cheap Talk.** 22917-22941

Charles Lu, Yaodong Yu, Sai Praneeth Karimireddy, Michael I. Jordan, Ramesh Raskar:  
**Federated Conformal Predictors for Distributed Uncertainty Quantification.** 22942-22964

Ekdeep Singh Lubana, Eric J. Bigelow, Robert P. Dick, David Scott Krueger, Hidenori Tanaka:  
**Mechanistic Mode Connectivity.** 22965-23004

Daniel Lundström, Meisam Razaviyayn:  
**A Unifying Framework to the Analysis of Interaction Methods using Synergy Functions.** 23005-23032

Huaishao Luo, Junwei Bao, Youzheng Wu, Xiaodong He, Tianrui Li:  
**SegCLIP: Patch Aggregation with Learnable Centers for Open-Vocabulary Semantic Segmentation.** 23033-23044

Ziwei Luo, Fredrik K. Gustafsson, Zheng Zhao, Jens Sjölund, Thomas B. Schön:  
**Image Restoration with Mean-Reverting Stochastic Differential Equations.** 23045-23066

Xinyu Luo, Christopher Musco, Cas Widdershoven:  
**Dimensionality Reduction for General KDE Mode Finding.** 23067-23082

Yuetian Luo, Zhimei Ren, Rina Barber:  
**Iterative Approximate Cross-Validation.** 23083-23102

Xu Luo, Hao Wu, Ji Zhang, Lianli Gao, Jing Xu, Jingkuan Song:  
**A Closer Look at Few-shot Classification Again.** 23103-23123



- Xiao Luo, Jingyang Yuan, Zijie Huang, Huiyu Jiang, Yifang Qin, Wei Ju, Ming Zhang, Yizhou Sun:  
**HOPE: High-order Graph ODE For Modeling Interacting Dynamics.** 23124-23139
- Tianjiao Luo, Ziyu Zhu, Jianfei Chen, Jun Zhu:  
**Stabilizing GANs' Training with Brownian Motion Controller.** 23140-23156
- Shahar Lutati, Lior Wolf:  
**OCD: Learning to Overfit with Conditional Diffusion Models.** 23157-23169
- Clare Lyle, Arash Mehrjou, Pascal Notin, Andrew Jesson, Stefan Bauer, Yarin Gal, Patrick Schwab:  
**DiscoBAX: Discovery of optimal intervention sets in genomic experiment design.** 23170-23189
- Clare Lyle, Zeyu Zheng, Evgenii Nikishin, Bernardo Ávila Pires, Razvan Pascanu, Will Dabney:  
**Understanding Plasticity in Neural Networks.** 23190-23211
- Lixing Lyu, Wang Chi Cheung:  
**Bandits with Knapsacks: Advice on Time-Varying Demands.** 23212-23238
- Boxiang Lyu, Zhe Feng, Zachary Robertson, Sanmi Koyejo:  
**Pairwise Ranking Losses of Click-Through Rates Prediction for Welfare Maximization in Ad Auctions.** 23239-23263
- Ivan Lyzhin, Aleksei Ustimenko, Andrey Gulin, Liudmila Prokhorenko:  
**Which Tricks are Important for Learning to Rank?** 23264-23278
- Pingchuan Ma, Peter Yichen Chen, Bolei Deng, Joshua B. Tenenbaum, Tao Du, Chuang Gan, Wojciech Matusik:  
**Learning Neural Constitutive Laws from Motion Observations for Generalizable PDE Dynamics.** 23279-23300
- Yecheng Jason Ma, Vikash Kumar, Amy Zhang, Osbert Bastani, Dinesh Jayaraman:  
**LIV: Language-Image Representations and Rewards for Robotic Control.** 23301-23320
- Liheng Ma, Chen Lin, Derek Lim, Adriana Romero-Soriano, Puneet K. Dokania, Mark Coates, Philip H. S. Torr, Ser-Nam Lim:  
**Graph Inductive Biases in Transformers without Message Passing.** 23321-23337
- Baorui Ma, Yu-Shen Liu, Zhizhong Han:  
**Learning Signed Distance Functions from Noisy 3D Point Clouds via Noise to Noise Mapping.** 23338-23357
- Mingwei Ma, Jizhou Liu, Samuel Sokota, Max Kleiman-Weiner, Jakob Nicolaus Foerster:  
**Learning Intuitive Policies Using Action Features.** 23358-23372
- Ziye Ma, Igor Molybog, Javad Lavaei, Somayeh Sojoudi:  
**Over-parametrization via Lifting for Low-rank Matrix Sensing: Conversion of Spurious Solutions to Strict Saddle Points.** 23373-23387
- Mingchen Ma, Christos Tzamos:  
**Buying Information for Stochastic Optimization.** 23388-23411
- Yihan Ma, Zhikun Zhang, Ning Yu, Xinlei He, Michael Backes, Yun Shen, Yang Zhang:  
**Generated Graph Detection.** 23412-23428
- Huan Ma, Qingyang Zhang, Changqing Zhang, Bingzhe Wu, Huazhu Fu, Joey Tianyi Zhou, Qinghua Hu:  
**Calibrating Multimodal Learning.** 23429-23450
- Alaa Maalouf, Murad Tukan, Vladimir Braverman, Daniela Rus:  
**AutoCoreset: An Automatic Practical Coreset Construction Framework.** 23451-23466
- Kanika Madan, Jarrid Rector-Brooks, Maksym Korablyov, Emmanuel Bengio, Moksh Jain, Andrei Cristian Nica, Tom Bosc, Yoshua Bengio, Nikolay Malkin:  
**Learning GFlowNets From Partial Episodes For Improved Convergence And Stability.** 23467-23483
- Jessica Maghakian, Russell Lee, Mohammad Hajiesmaili, Jian Li, Ramesh K. Sitaraman, Zhenhua Liu:  
**Applied Online Algorithms with Heterogeneous Predictors.** 23484-23497
- Gengchen Mai, Ni Lao, Yutong He, Jiaming Song, Stefano Ermon:  
**CSP: Self-Supervised Contrastive Spatial Pre-Training for Geospatial-Visual Representations.** 23498-23515
- Peihua Mai, Yan Pang:  
**Vertical Federated Graph Neural Network for Recommender System.** 23516-23535
- Pratyush Maini, Michael Curtis Mozer, Hanie Sedghi, Zachary Chase Lipton, J. Zico Kolter, Chiyuan Zhang:  
**Can Neural Network Memorization Be Localized?** 23536-23557
- Anirudha Majumdar:  
**Fundamental Tradeoffs in Learning with Prior Information.** 23558-23573
- Alan Malek, Virginia Aglietti, Silvia Chiappa:  
**Additive Causal Bandits with Unknown Graph.** 23574-23589
- Dhruv Malik, Conor Igoe, Yuanzhi Li, Aarti Singh:  
**Weighted Tallying Bandits: Overcoming Intractability via Repeated Exposure Optimality.** 23590-23609
- Sadhika Malladi, Alexander Wettig, Dingli Yu, Danqi Chen, Sanjeev Arora:  
**A Kernel-Based View of Language Model Fine-Tuning.** 23610-23641
- Debmalya Mandal, Stelios Triantafyllou, Goran Radanovic:  
**Performative Reinforcement Learning.** 23642-23680
- Paul Mangold, Michaël Perrot, Aurélien Bellet, Marc Tommasi:  
**Differential Privacy has Bounded Impact on Fairness in Classification.** 23681-23705



- Yishay Mansour, Richard Nock, Robert C. Williamson:  
**Random Classification Noise does not defeat All Convex Potential Boosters Irrespective of Model Choice.** 23706-23742
- Anqi Mao, Mehryar Mohri, Yutao Zhong:  
**H-Consistency Bounds for Pairwise Misranking Loss Surrogates.** 23743-23802
- Anqi Mao, Mehryar Mohri, Yutao Zhong:  
**Cross-Entropy Loss Functions: Theoretical Analysis and Applications.** 23803-23828
- Yixiu Mao, Hongchang Zhang, Chen Chen, Yi Xu, Xiangyang Ji:  
**Supported Trust Region Optimization for Offline Reinforcement Learning.** 23829-23851
- Chengzhi Mao, Lingyu Zhang, Abhishek Vaibhav Joshi, Junfeng Yang, Hao Wang, Carl Vondrick:  
**Robust Perception through Equivariance.** 23852-23870
- Anna C. Marbut, Katy McKinney-Bock, Travis J. Wheeler:  
**Reliable Measures of Spread in High Dimensional Latent Spaces.** 23871-23885
- Tanguy Marchand, Regis Loeb, Ulysse Marteau-Ferey, Jean Ogier du Terrail, Arthur Pignet:  
**SRATTA: Sample Re-ATTribution Attack of Secure Aggregation in Federated Learning.** 23886-23914
- Emanuele Marconato, Gianpaolo Bontempo, Elisa Ficarra, Simone Calderara, Andrea Passerini, Stefano Teso:  
**Neuro-Symbolic Continual Learning: Knowledge, Reasoning Shortcuts and Concept Rehearsal.** 23915-23936
- Adria Marcos-Morales, Matan Leibovich, Sreyas Mohan, Joshua Lawrence Vincent, Piyush Haluai, Mai Tan, Peter A. Crozier, Carlos Fernandez-Granda:  
**Evaluating Unsupervised Denoising Requires Unsupervised Metrics.** 23937-23957
- Étienne Marcotte, Valentina Zantedeschi, Alexandre Drouin, Nicolas Chapados:  
**Regions of Reliability in the Evaluation of Multivariate Probabilistic Forecasts.** 23958-24004
- Raja Marjeh, Ilia Sucholutsky, Thomas A. Langlois, Nori Jacoby, Thomas L. Griffiths:  
**Analyzing Diffusion as Serial Reproduction.** 24005-24019
- Ilia Markov, Adrian Vladu, Qi Guo, Dan Alistarh:  
**Quantized Distributed Training of Large Models with Convergence Guarantees.** 24020-24044
- Juan Maroñas, Daniel Hernández-Lobato:  
**Efficient Transformed Gaussian Processes for Non-Stationary Dependent Multi-class Classification.** 24045-24081
- Samuele Marro, Michele Lombardi:  
**Computational Asymmetries in Robust Classification.** 24082-24138
- Tanya Marwah, Zachary Chase Lipton, Jianfeng Lu, Andrej Risteski:  
**Neural Network Approximations of PDEs Beyond Linearity: A Representational Perspective.** 24139-24172
- Satvik Mehul Mashkaria, Siddarth Krishnamoorthy, Aditya Grover:  
**Generative Pretraining for Black-Box Optimization.** 24173-24197
- Aditya Mate, Bryan Wilder, Aparna Taneja, Milind Tambe:  
**Improved Policy Evaluation for Randomized Trials of Algorithmic Resource Allocation.** 24198-24213
- Aimee Maurais, Terrence Alsup, Benjamin Peherstorfer, Youssef M. Marzouk:  
**Multi-Fidelity Covariance Estimation in the Log-Euclidean Geometry.** 24214-24235
- Prathamesh Mayekar, Jonathan Scarlett, Vincent Y. F. Tan:  
**Communication-Constrained Bandits under Additive Gaussian Noise.** 24236-24250
- Alessio Mazzetto, Eli Upfal:  
**Nonparametric Density Estimation under Distribution Drift.** 24251-24270
- Sokhna Diarra Mbacke, Florence Clerc, Pascal Germain:  
**PAC-Bayesian Generalization Bounds for Adversarial Generative Models.** 24271-24290
- Brandon McKinzie, Vaishaal Shankar, Joseph Yitan Cheng, Yinfei Yang, Jonathon Shlens, Alexander T. Toshev:  
**Robustness in Multimodal Learning under Train-Test Modality Mismatch.** 24291-24303
- Mohammad Mehrabi, Ryan A. Rossi:  
**A Model-free Closeness-of-influence Test for Features in Supervised Learning.** 24304-24324
- Jincheng Mei, Zixin Zhong, Bo Dai, Alekh Agarwal, Csaba Szepesvári, Dale Schuurmans:  
**Stochastic Gradient Succeeds for Bandits.** 24325-24360
- Valentyn Melnychuk, Dennis Frauen, Stefan Feuerriegel:  
**Normalizing Flows for Interventional Density Estimation.** 24361-24397
- Igor Melnyk, Vijil Chenthamarakshan, Pin-Yu Chen, Payel Das, Amit Dhurandhar, Inkit Padhi, Devleena Das:  
**Reprogramming Pretrained Language Models for Antibody Sequence Infilling.** 24398-24419
- Omid Memarrast, Linh Vu, Brian D. Ziebart:  
**Superhuman Fairness.** 24420-24435
- Si Yi Meng, Robert M. Gower:  
**A Model-Based Method for Minimizing CVaR and Beyond.** 24436-24456
- Yu Meng, Martin Michalski, Jiaxin Huang, Yu Zhang, Tarek F. Abdelzaher, Jiawei Han:  
**Tuning Language Models as Training Data Generators for Augmentation-Enhanced Few-Shot Learning.** 24457-



24477

- Nadav Merlis, Hugo Richard, Flore Sentenac, Corentin Odic, Mathieu Molina, Vianney Perchet:  
**On Preemption and Learning in Stochastic Scheduling.** 24478-24516
- Thomas Mesnard, Wenqi Chen, Alaa Saade, Yunhao Tang, Mark Rowland, Theophane Weber, Clare Lyle, Audrunas Gruslys, Michal Valko, Will Dabney, Georg Ostrovski, Eric Moulines, Rémi Munos:  
**Quantile Credit Assignment.** 24517-24531
- Dmitry Metelev, Alexander Rogozin, Dmitry Kovalev, Alexander V. Gasnikov:  
**Is Consensus Acceleration Possible in Decentralized Optimization over Slowly Time-Varying Networks?** 24532-24554
- Alberto Maria Metelli, Filippo Lazzati, Marcello Restelli:  
**Towards Theoretical Understanding of Inverse Reinforcement Learning.** 24555-24591
- Nico Meyer , Daniel D. Scherer, Axel Plinge, Christopher Mutschler, Michael J. Hartmann:  
**Quantum Policy Gradient Algorithm with Optimized Action Decoding.** 24592-24613
- Lucas Thibaut Meyer, Marc Schouler , Robert Alexander Caulk, Alejandro Ribés, Bruno Raffin:  
**Training Deep Surrogate Models with Large Scale Online Learning.** 24614-24630
- David Henry Mguni, Haojun Chen, Taher Jafferjee, Jianhong Wang, Longfei Yue, Xidong Feng, Stephen Marcus McAleer, Feifei Tong, Jun Wang, Yaodong Yang:  
**MANSA: Learning Fast and Slow in Multi-Agent Systems.** 24631-24658
- Zakaria Mhammedi, Dylan J. Foster, Alexander Rakhlin:  
**Representation Learning with Multi-Step Inverse Kinematics: An Efficient and Optimal Approach to Rich-Observation RL.** 24659-24700
- Elissa Mhanna, Mohamad Assaad:  
**Single Point-Based Distributed Zeroth-Order Optimization with a Non-Convex Stochastic Objective Function.** 24701-24719
- Ning Miao, Tom Rainforth, Emile Mathieu, Yann Dubois, Yee Whye Teh, Adam Foster, Hyunjik Kim:  
**Learning Instance-Specific Augmentations by Capturing Local Invariances.** 24720-24736
- Gaspard Michel, Giannis Nikolentzos, Johannes F. Lutzeyer, Michalis Vazirgiannis:  
**Path Neural Networks: Expressive and Accurate Graph Neural Networks.** 24737-24755
- Thomas Miconi:  
**Learning to acquire novel cognitive tasks with evolution, plasticity and meta-meta-learning.** 24756-24774
- Eleni Miliotou, Panagiotis Kyriakis, Jason D. Hinman, Andrei Irimia, Paul Bogdan:  
**Generative Decoding of Visual Stimuli.** 24775-24784
- Yifei Min, Jiafan He, Tianhao Wang, Quanquan Gu:  
**Cooperative Multi-Agent Reinforcement Learning: Asynchronous Communication and Linear Function Approximation.** 24785-24811
- Ming Min, Ruimeng Hu, Tomoyuki Ichiba:  
**Directed Chain Generative Adversarial Networks.** 24812-24830
- Seungki Min, Daniel Russo:  
**An Information-Theoretic Analysis of Nonstationary Bandit Learning.** 24831-24849
- Hancheng Min, René Vidal, Enrique Mallada:  
**On the Convergence of Gradient Flow on Multi-layer Linear Models.** 24850-24887
- Aaron Mishkin, Mert Pilancı:  
**Optimal Sets and Solution Paths of ReLU Networks.** 24888-24924
- Gal Mishne , Zhengchao Wan, Yusu Wang, Sheng Yang:  
**The Numerical Stability of Hyperbolic Representation Learning.** 24925-24949
- Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D. Manning, Chelsea Finn:  
**DetectGPT: Zero-Shot Machine-Generated Text Detection using Probability Curvature.** 24950-24962
- Sarthak Mittal, Korbinian Abstreiter, Stefan Bauer, Bernhard Schölkopf, Arash Mehrjou:  
**Diffusion Based Representation Learning.** 24963-24982
- Yujie Mo, Yajie Lei, Jialie Shen, Xiaoshuang Shi, Heng Tao Shen, Xiaofeng Zhu:  
**Disentangled Multiplex Graph Representation Learning.** 24983-25005
- Shentong Mo, Pedro Morgado:  
**A Unified Audio-Visual Learning Framework for Localization, Separation, and Recognition.** 25006-25017
- Zhanfeng Mo, Haosen Shi, Sinno Jialin Pan:  
**Pruning via Sparsity-indexed ODE: a Continuous Sparsity Viewpoint.** 25018-25036
- Mazda Moayeri, Keivan Rezaei, Maziar Sanjabi, Soheil Feizi:  
**Text-To-Concept (and Back) via Cross-Model Alignment.** 25037-25060
- Mohamad Amin Mohamadi , Wonho Bae, Danica J. Sutherland:  
**A Fast, Well-Founded Approximation to the Empirical Neural Tangent Kernel.** 25061-25081
- Amirkeivan Mohtashami, Martin Jaggi, Sebastian U. Stich:  
**Special Properties of Gradient Descent with Large Learning Rates.** 25082-25104



- Roberto Molinaro, Yunan Yang , Björn Engquist, Siddhartha Mishra:  
**Neural Inverse Operators for Solving PDE Inverse Problems.** 25105-25139
- Paul Monchot, Loic Coquelin, Sébastien Julien Petit, Sébastien Marmin, Erwan Le Pennec, Nicolas Fischer:  
**Input uncertainty propagation through trained neural networks.** 25140-25173
- Andrea Montanari, Eric Weiner:  
**Compressing Tabular Data via Latent Variable Estimation.** 25174-25208
- Enea Monzio Compagnoni, Luca Biggio, Antonio Orvieto, Frank Norbert Proske, Hans Kersting, Aurélien Lucchi:  
**An SDE for Modeling SAM: Theory and Insights.** 25209-25253
- Terufumi Morishita, Gaku Morio, Atsuki Yamaguchi, Yasuhiro Sogawa:  
**Learning Deductive Reasoning from Synthetic Corpus based on Formal Logic.** 25254-25274
- Christopher Morris, Floris Geerts, Jan Tönshoff, Martin Grohe:  
**WL meet VC.** 25275-25302
- Ted Moskovitz, Brendan O'Donoghue, Vivek Veeriah, Sebastian Flennerhag, Satinder Singh, Tom Zahavy:  
**ReLOAD: Reinforcement Learning with Optimistic Ascent-Descent for Last-Iterate Convergence in Constrained MDPs.** 25303-25336
- Antoine Moulin, Gergely Neu:  
**Optimistic Planning by Regularized Dynamic Programming.** 25337-25357
- Nicola Muca Cirone, Maud Lemercier, Christopher Salvi:  
**Neural signature kernels as infinite-width-depth-limits of controlled ResNets.** 25358-25425
- Matthew J. Muckley, Alaaeldin El-Nouby, Karen Ullrich, Hervé Jégou, Jakob Verbeek:  
**Improving Statistical Fidelity for Neural Image Compression with Implicit Local Likelihood Models.** 25426-25443
- Samuel Müller, Matthias Feurer, Noah Hollmann, Frank Hutter:  
**PFNs4BO: In-Context Learning for Bayesian Optimization.** 25444-25470
- Johannes Müller, Marius Zeinhofer :  
**Achieving High Accuracy with PINNs via Energy Natural Gradient Descent.** 25471-25485
- Andreas Munk, Alexander Mead, Frank Wood:  
**Uncertain Evidence in Probabilistic Models and Stochastic Simulators.** 25486-25500
- Naoki Murata, Koichi Saito, Chieh-Hsin Lai, Yuhta Takida, Toshimitsu Uesaka, Yuki Mitsufuji, Stefano Ermon:  
**GibbsDDRM: A Partially Collapsed Gibbs Sampler for Solving Blind Inverse Problems with Denoising Diffusion Restoration.** 25501-25522
- Tomoya Murata, Taiji Suzuki:  
**DIFF2: Differential Private Optimization via Gradient Differences for Nonconvex Distributed Learning.** 25523-25548
- Michael Murphy, Stefanie Jegelka, Ernest Fraenkel, Tobias Kind, David Healey, Thomas Butler:  
**Efficiently predicting high resolution mass spectra with graph neural networks.** 25549-25562
- Marco Mussi, Alberto Maria Metelli, Marcello Restelli:  
**Dynamical Linear Bandits.** 25563-25587
- Ofir Nabati, Guy Tennenholtz, Shie Mannor:  
**Representation-Driven Reinforcement Learning.** 25588-25603
- Adel Nabli, Edouard Oyallon:  
**DADAO: Decoupled Accelerated Decentralized Asynchronous Optimization.** 25604-25626
- Dheeraj Mysore Nagaraj, Suhas S. Kowshik, Naman Agarwal, Praneeth Netrapalli, Prateek Jain:  
**Multi-User Reinforcement Learning with Low Rank Rewards.** 25627-25659
- Thomas Nagler :  
**Statistical Foundations of Prior-Data Fitted Networks.** 25660-25676
- Aaditya Naik, Yinjun Wu, Mayur Naik, Eric Wong :  
**Do Machine Learning Models Learn Statistical Rules Inferred from Data?** 25677-25693
- Ilan Naiman, Nimrod Berman, Omri Azencot:  
**Sample and Predict Your Latent: Modality-free Sequential Disentanglement via Contrastive Estimation.** 25694-25717
- Milad Nasr, Saeed Mahloujifar, Xinyu Tang, Prateek Mittal, Amir Houmansadr:  
**Effectively Using Public Data in Privacy Preserving Machine Learning.** 25718-25732
- Arash Nasr-Esfahany, Mohammad Alizadeh, Devavrat Shah:  
**Counterfactual Identifiability of Bijective Causal Models.** 25733-25754
- Somjit Nath, Gopeshh Raaj Subbaraj, Khimya Khetarpal, Samira Ebrahimi Kahou:  
**Discovering Object-Centric Generalized Value Functions From Pixels.** 25755-25768
- Michal Nauman, Marek Cygan:  
**On Many-Actions Policy Gradient.** 25769-25789
- Aviv Navon, Aviv Shamsian, Idan Achituv, Ethan Fetaya, Gal Chechik, Haggai Maron:  
**Equivariant Architectures for Learning in Deep Weight Spaces.** 25790-25816



Siddharth Nayak, Kenneth Choi, Wenqi Ding, Sydney Dolan, Karthik Gopalakrishnan, Hamsa Balakrishnan:  
**Scalable Multi-Agent Reinforcement Learning through Intelligent Information Aggregation.** 25817-25833

Philipp Nazari, Sebastian Damrich, Fred A. Hamprecht:  
**Geometric Autoencoders - What You See is What You Decode.** 25834-25857

Kirill Neklyudov , Rob Brekelmans, Daniel Severo, Alireza Makhzani:  
**Action Matching: Learning Stochastic Dynamics from Samples.** 25858-25889

Buddhika Nettasinghe, Samrat Chatterjee, Ramakrishna Tipireddy, Mahantesh M. Halappanavar:  
**Extending Conformal Prediction to Hidden Markov Models with Exact Validity via de Finetti's Theorem for Markov Chains.** 25890-25903

Tung Nguyen, Johannes Brandstetter, Ashish Kapoor, Jayesh K. Gupta, Aditya Grover:  
**ClimaX: A foundation model for weather and climate.** 25904-25938

Hoai-An Nguyen, Ching-An Cheng:  
**Provable Reset-free Reinforcement Learning by No-Regret Reduction.** 25939-25955

Khang Nguyen, Nong Minh Hieu, Vinh Duc Nguyen, Nhat Ho, Stanley J. Osher, Tan Minh Nguyen:  
**Revisiting Over-smoothing and Over-squashing Using Ollivier-Ricci Curvature.** 25956-25979

Tri Nguyen, Shahana Ibrahim, Xiao Fu:  
**Deep Clustering with Incomplete Noisy Pairwise Annotations: A Geometric Regularization Approach.** 25980-26007

Khai Nguyen, Dang Nguyen, Nhat Ho:  
**Self-Attention Amortized Distributional Projection Optimization for Sliced Wasserstein Point-Cloud Reconstruction.** 26008-26030

Xuan Son Nguyen, Shuo Yang:  
**Building Neural Networks on Matrix Manifolds: A Gyrovector Space Approach.** 26031-26062

Lilian Ngweta, Subha Maity, Alex Gittens, Yuekai Sun, Mikhail Yurochkin:  
**Simple Disentanglement of Style and Content in Visual Representations.** 26063-26086

Fei Ni, Jianye Hao, Yao Mu, Yifu Yuan, Yan Zheng, Bin Wang, Zhixuan Liang :  
**MetaDiffuser: Diffusion Model as Conditional Planner for Offline Meta-RL.** 26087-26105

Ansong Ni, Srini Iyer, Dragomir Radev, Veselin Stoyanov, Wen-Tau Yih, Sida I. Wang, Xi Victoria Lin:  
**LEVER: Learning to Verify Language-to-Code Generation with Execution.** 26106-26128

Zixuan Ni, Longhui Wei, Siliang Tang, Yueting Zhuang, Qi Tian:  
**Continual Vision-Language Representation Learning with Off-Diagonal Information.** 26129-26149

Guangyu Nie, Changhoon Kim, Yezhou Yang, Yi Ren:  
**Attributing Image Generative Models using Latent Fingerprints.** 26150-26165

Guanyu Nie, Yididiya Y. Nadew, Yanhui Zhu, Vaneet Aggarwal, Christopher John Quinn:  
**A Framework for Adapting Offline Algorithms to Solve Combinatorial Multi-Armed Bandit Problems with Bandit Feedback.** 26166-26198

Yaniv Nikankin, Niv Haim, Michal Irani:  
**SinFusion: Training Diffusion Models on a Single Image or Video.** 26199-26214

Mahdi Nikdan, Tommaso Pegolotti, Eugenia Iofinova, Eldar Kurtic, Dan Alistarh:  
**SparseProp: Efficient Sparse Backpropagation for Faster Training of Neural Networks at the Edge.** 26215-26227

Alexander Nikulin , Vladislav Kurenkov, Denis Tarasov, Sergey Kolesnikov:  
**Anti-Exploration by Random Network Distillation.** 26228-26244

Mang Ning, Enver Sangineto, Angelo Porrello, Simone Calderara, Rita Cucchiara:  
**Input Perturbation Reduces Exposure Bias in Diffusion Models.** 26245-26265

Atsushi Nitanda, Kazusato Oko, Denny Wu, Nobuhito Takenouchi, Taiji Suzuki:  
**Primal and Dual Analysis of Entropic Fictitious Play for Finite-sum Problems.** 26266-26282

Georgy Noarov, Aaron Roth:  
**The Statistical Scope of Multicalibration.** 26283-26310

Kolby Nottingham, Prithviraj Ammanabrolu, Alane Suhr, Yejin Choi, Hannaneh Hajishirzi, Sameer Singh, Roy Fox:  
**Do Embodied Agents Dream of Pixelated Sheep: Embodied Decision Making using Language Guided World Modelling.** 26311-26325

Azade Nova, Hanjun Dai, Dale Schuurmans:  
**Gradient-Free Structured Pruning with Unlabeled Data.** 26326-26341

Zachary Novack, Julian J. McAuley, Zachary Chase Lipton, Saurabh Garg:  
**CHiLS: Zero-Shot Image Classification with Hierarchical Label Sets.** 26342-26362

Georgii Sergeevich Novikov, Daniel Bershtsky, Julia Gusak, Alex Shonenkov, Denis Valerievich Dimitrov, Ivan V. Oseledets:  
**Few-bit Backward: Quantized Gradients of Activation Functions for Memory Footprint Reduction.** 26363-26381

Brendan O'Donoghue:  
**Efficient Exploration via Epistemic-Risk-Seeking Policy Optimization.** 26382-26402

Junsoo Oh, Chulhee Yun:  
**Provable Benefit of Mixup for Finding Optimal Decision Boundaries.** 26403-26450



- Ruben Ohana, Kimia Nadjahi, Alain Rakotomamonjy, Liva Ralaivola:  
**Shedding a PAC-Bayesian Light on Adaptive Sliced-Wasserstein Distances.** 26451-26473
- Guy Ohayon, Theo Joseph Adrai, Michael Elad, Tomer Michaeli:  
**Reasons for the Superiority of Stochastic Estimators over Deterministic Ones: Robustness, Consistency and Perceptual Quality.** 26474-26494
- Nastaran Okati, Stratis Tsirsis, Manuel Gomez Rodriguez:  
**On the Within-Group Fairness of Screening Classifiers.** 26495-26516
- Kazusato Oko, Shunta Akiyama, Taiji Suzuki:  
**Diffusion Models are Minimax Optimal Distribution Estimators.** 26517-26582
- Raphaël Olivier, Bhiksha Raj:  
**How Many Perturbations Break This Model? Evaluating Robustness Beyond Adversarial Accuracy.** 26583-26598
- Miruna Oprescu, Jacob Dorn, Marah Ghoummaid, Andrew Jesson, Nathan Kallus, Uri Shalit:  
**B-Learner: Quasi-Oracle Bounds on Heterogeneous Causal Effects Under Hidden Confounding.** 26599-26618
- Gabriel Orlanski, Kefan Xiao, Xavier Garcia, Jeffrey Hui, Joshua Howland, Jonathan Malmaud, Jacob Austin, Rishabh Singh, Michele Catasta:  
**Measuring the Impact of Programming Language Distribution.** 26619-26645
- Guillermo Ortiz-Jiménez, Mark Collier, Anant Nawalgaria, Alexander Nicholas D'Amour, Jesse Berent, Rodolphe Jenatton, Efi Kokopoulou:  
**When does Privileged information Explain Away Label Noise?** 26646-26669
- Antonio Orvieto, Samuel L. Smith, Albert Gu, Anushan Fernando, Çağlar Gülcühre, Razvan Pascanu, Soham De:  
**Resurrecting Recurrent Neural Networks for Long Sequences.** 26670-26698
- Yidong Ouyang, Liyan Xie, Guang Cheng:  
**Improving Adversarial Robustness Through the Contrastive-Guided Diffusion Process.** 26699-26723
- Samet Oymak, Ankit Singh Rawat, Mahdi Soltanolkotabi , Christos Thrampoulidis:  
**On the Role of Attention in Prompt-tuning.** 26724-26768
- Asuman E. Ozdaglar, Sarath Pattathil, Jiawei Zhang, Kaiqing Zhang:  
**Revisiting the Linear-Programming Framework for Offline RL with General Function Approximation.** 26769-26791
- Vishakh Padmakumar, Richard Yuanzhe Pang, He He, Ankur P. Parikh:  
**Extrapolative Controlled Sequence Generation via Iterative Refinement.** 26792-26808
- Avik Pal, Alan Edelman, Christopher Vincent Rackauckas:  
**Locally Regularized Neural Differential Equations: Some Black Boxes were meant to remain closed!** 26809-26819
- Sourav Pal, Zhanpeng Zeng, Sathya N. Ravi, Vikas Singh:  
**Controlled Differential Equations on Long Sequences via Non-standard Wavelets.** 26820-26836
- Alexander Pan, Jun Shern Chan, Andy Zou, Nathaniel Li, Steven Basart, Thomas Woodside, Hanlin Zhang, Scott Emmons, Dan Hendrycks:  
**Do the Rewards Justify the Means? Measuring Trade-Offs Between Rewards and Ethical Behavior in the Machiavelli Benchmark.** 26837-26867
- Erlin Pan, Zhao Kang:  
**Beyond Homophily: Reconstructing Structure for Graph-agnostic Clustering.** 26868-26877
- Ling Pan, Nikolay Malkin, Dinghuai Zhang, Yoshua Bengio:  
**Better Training of GFlowNets with Local Credit and Incomplete Trajectories.** 26878-26890
- Hongyi Pan, Xin Zhu, Salih Furkan Atici, Ahmet Enis Çetin:  
**A Hybrid Quantum-Classical Approach based on the Hadamard Transform for the Convolutional Layer.** 26891-26903
- Ioannis Panageas, Stratis Skoulakis , Luca Viano, Xiao Wang, Volkan Cevher:  
**Semi Bandit dynamics in Congestion Games: Convergence to Nash Equilibrium and No-Regret Guarantees.** 26904-26930
- Kunjal Panchal, Sunav Choudhary, Subrata Mitra, Koyel Mukherjee, Somdeb Sarkhel, Saayan Mitra, Hui Guan:  
**Flash: Concept Drift Adaptation in Federated Learning.** 26931-26962
- Deep Shankar Pandey, Qi Yu:  
**Learn to Accumulate Evidence from All Training Samples: Theory and Practice.** 26963-26989
- Qi Pang, Lun Wang, Shuai Wang, Wenting Zheng, Dawn Song:  
**Secure Federated Correlation Test and Entropy Estimation.** 26990-27010
- Abhishek Panigrahi, Nikunj Saunshi, Haoyu Zhao, Sanjeev Arora:  
**Task-Specific Skill Localization in Fine-tuned Language Models.** 27011-27033
- Junyoung Park, Jeongyoun Ahn, Cheolwoo Park:  
**Kernel Sufficient Dimension Reduction and Variable Selection for Compositional Data via Amalgamation.** 27034-27047
- Jin-Hwi Park, Jaesung Choe, Inhwon Bae, Hae-Gon Jeon:  
**Learning Affinity with Hyperbolic Representation for Spatial Propagation.** 27048-27073



- Sung Min Park, Kristian Georgiev, Andrew Ilyas, Guillaume Leclerc, Aleksander Madry:  
**TRAK: Attributing Model Behavior at Scale.** 27074-27113
- Jungwuk Park, Dong-Jun Han, Soyeong Kim, Jaekyun Moon:  
**Test-Time Style Shifting: Handling Arbitrary Styles in Domain Generalization.** 27114-27131
- Sejun Park, Kihun Hong, Ganguk Hwang:  
**Towards Understanding Ensemble Distillation in Federated Learning.** 27132-27187
- Seobin Park, Dongjin Kim, Sungyong Baik, Tae Hyun Kim:  
**Learning Controllable Degradation for Real-World Super-Resolution via Constrained Flows.** 27188-27203
- Jinseong Park , Hoki Kim, Yujin Choi, Jaewook Lee:  
**Differentially Private Sharpness-Aware Training.** 27204-27224
- Seohong Park, Kimin Lee, Youngwoon Lee, Pieter Abbeel:  
**Controllability-Aware Unsupervised Skill Discovery.** 27225-27245
- Seohong Park, Sergey Levine:  
**Predictable MDP Abstraction for Unsupervised Model-Based RL.** 27246-27268
- Sungwoo Park, Byoungwoo Park, Moontae Lee, Changhee Lee:  
**Neural Stochastic Differential Games for Time-series Analysis.** 27269-27293
- Jisun Park, Ernest K. Ryu:  
**Accelerated Infeasibility Detection of Constrained Optimization and Fixed-Point Iterations.** 27294-27345
- Paavo Parmas, Takuma Seno, Yuma Aoki:  
**Model-based Reinforcement Learning with Scalable Composite Policy Gradient Estimators.** 27346-27377
- Advait U. Parulekar, Karthikeyan Shanmugam , Sanjay Shakkottai:  
**PAC Generalization via Invariant Representations.** 27378-27400
- Farhad Pashakhanloo, Alexei Koulakov:  
**Stochastic Gradient Descent-Induced Drift of Representation in a Two-Layer Neural Network.** 27401-27419
- Saro Passaro, C. Lawrence Zitnick:  
**Reducing SO(3) Convolutions to SO(2) for Efficient Equivariant GNNs.** 27420-27438
- Kumar Kshitij Patel, Lingxiao Wang, Aadirupa Saha, Nathan Srebro:  
**Federated Online and Bandit Convex Optimization.** 27439-27460
- Edward Pearce-Crump:  
**Brauer's Group Equivariant Neural Networks.** 27461-27482
- Edward Pearce-Crump:  
**How Jellyfish Characterise Alternating Group Equivariant Neural Networks.** 27483-27495
- Kexin Pei, David Bieber, Kensen Shi, Charles Sutton, Pengcheng Yin:  
**Can Large Language Models Reason about Program Invariants?** 27496-27520
- Zhengqi Pei, Shuhui Wang:  
**Dynamics-inspired Neuromorphic Visual Representation Learning.** 27521-27541
- Peifeng Gao, Qianqian Xu, Peisong Wen, Zhiyong Yang, Huiyang Shao, Qingming Huang:  
**Feature Directions Matter: Long-Tailed Learning via Rotated Balanced Representation.** 27542-27563
- Jaakko Peltonen, Wen Xu, Timo Nummenmaa, Jyrki Nummenmaa:  
**Fair Neighbor Embedding.** 27564-27584
- Liangzu Peng, Paris Giampouras, René Vidal:  
**The Ideal Continual Learner: An Agent That Never Forgets.** 27585-27610
- Xingang Peng, Jiaqi Guan, Qiang Liu, Jianzhu Ma:  
**MolDiff: Addressing the Atom-Bond Inconsistency Problem in 3D Molecule Diffusion Generation.** 27611-27629
- Andi Peng, Aviv Netanyahu, Mark K. Ho, Tianmin Shu, Andreea Bobu , Julie Shah, Pulkit Agrawal:  
**Diagnosis, Feedback, Adaptation: A Human-in-the-Loop Framework for Test-Time Policy Adaptation.** 27630-27641
- Binyamin Perets, Mark Kozdoba, Shie Mannor:  
**Learning Hidden Markov Models When the Locations of Missing Observations are Unknown.** 27642-27667
- Lorenzo Perini, Paul-Christian Bürkner, Arto Klami:  
**Estimating the Contamination Factor's Distribution in Unsupervised Anomaly Detection.** 27668-27679
- Luca Pesce, Florent Krzakala , Bruno Loureiro, Ludovic Stephan:  
**Are Gaussian Data All You Need? The Extents and Limits of Universality in High-Dimensional Generalized Linear Estimation.** 27680-27708
- Aleksandar Petrov, Francisco Eiras, Amartya Sanyal, Philip H. S. Torr, Adel Bibi:  
**Certifying Ensembles: A General Certification Theory with S-Lipschitzness.** 27709-27736
- Daniel Pfrommer, Max Simchowitz, Tyler Westenbroek, Nikolai Matni, Stephen Tu:  
**The Power of Learned Locally Linear Models for Nonlinear Policy Optimization.** 27737-27821
- Chi Bach Pham , Wynita Griggs, James Saunderson:  
**A Scalable Frank-Wolfe-Based Algorithm for the Max-Cut SDP.** 27822-27839



Thomy Phan, Fabian Ritz, Philipp Altmann, Maximilian Zorn, Jonas Nüßlein, Michael Kölle, Thomas Gabor, Claudia Linnhoff-Popien:

**Attention-Based Recurrence for Multi-Agent Reinforcement Learning under Stochastic Partial Observability.**  
27840-27853



Jason Phang, Yi Mao, Pengcheng He, Weizhu Chen:  
**HyperTuning: Toward Adapting Large Language Models without Back-propagation.** 27854-27875



Hannah Pinson, Joeri Lenaerts, Vincent Ginis:  
**Linear CNNs Discover the Statistical Structure of the Dataset Using Only the Most Dominant Frequencies.** 27876-27906



Vincent Plassier, Mehdi Makni, Aleksandr Rubashevskii, Eric Moulines, Maxim Panov:  
**Conformal Prediction for Federated Uncertainty Quantification Under Label Shift.** 27907-27947



Lena Podina, Brydon Eastman, Mohammad Kohandel:  
**Universal Physics-Informed Neural Networks: Symbolic Differential Operator Discovery with Sparse Data.** 27948-27956



Aleksandr Podkopaev, Patrick Blöbaum, Shiva Prasad Kasiviswanathan, Aaditya Ramdas:  
**Sequential Kernelized Independence Testing.** 27957-27993



Riccardo Poiani, Alberto Maria Metelli, Marcello Restelli:  
**Truncating Trajectories in Monte Carlo Reinforcement Learning.** 27994-28042



Michael Poli, Stefano Massaroli, Eric Nguyen, Daniel Y. Fu, Tri Dao, Stephen Baccus, Yoshua Bengio, Stefano Ermon, Christopher Ré:  
**Hyena Hierarchy: Towards Larger Convolutional Language Models.** 28043-28078



Samuele Pollaci:  
**Spurious Valleys and Clustering Behavior of Neural Networks.** 28079-28099



Aram-Alexandre Pooladian, Heli Ben-Hamu, Carles Domingo-Enrich, Brandon Amos, Yaron Lipman, Ricky T. Q. Chen:  
**Multisample Flow Matching: Straightening Flows with Minibatch Couplings.** 28100-28127



Aram-Alexandre Pooladian, Vincent Divol, Jonathan Niles-Weed:  
**Minimax estimation of discontinuous optimal transport maps: The semi-discrete case.** 28128-28150



Mihir Prabhudesai, Anirudh Goyal, Sujoy Paul, Sjoerd van Steenkiste, Mehdi S. M. Sajjadi, Gaurav Aggarwal, Thomas Kipf, Deepak Pathak, Katerina Fragkiadaki:  
**Test-time Adaptation with Slot-Centric Models.** 28151-28166



Drew Prinster, Suchi Saria, Anqi Liu:  
**JAWS-X: Addressing Efficiency Bottlenecks of Conformal Prediction Under Standard and Feedback Covariate Shift.** 28167-28190



Omri Puny, Derek Lim, Bobak Toussi Kiani, Haggai Maron, Yaron Lipman:  
**Equivariant Polynomials for Graph Neural Networks.** 28191-28222



Zekun Qi, Runpei Dong, Guofan Fan, Zheng Ge, Xiangyu Zhang, Kaisheng Ma, Li Yi:  
**Contrast with Reconstruct: Contrastive 3D Representation Learning Guided by Generative Pretraining.** 28223-28243



Shiang Qi, Neeraj Kumar, Mahtab Farrokh, Weijie Sun, Li-Hao Kuan, Rajesh Ranganath, Ricardo Henao, Russell Greiner:  
**An Effective Meaningful Way to Evaluate Survival Models.** 28244-28276



Bo Qiang, Yuxuan Song, Minkai Xu, Jingjing Gong, Bowen Gao, Hao Zhou, Wei-Ying Ma, Yanyan Lan:  
**Coarse-to-Fine: a Hierarchical Diffusion Model for Molecule Generation in 3D.** 28277-28299



Rui Qiao, Xinyi Xu, Bryan Kian Hsiang Low:  
**Collaborative Causal Inference with Fair Incentives.** 28300-28320



Congyu Qiao, Ning Xu, Jiaqi Lv, Yi Ren, Xin Geng:  
**FREDIS: A Fusion Framework of Refinement and Disambiguation for Unreliable Partial Label Learning.** 28321-28336



Guanghui Qin, Benjamin Van Durme:  
**Nugget: Neural Agglomerative Embeddings of Text.** 28337-28350



Haotong Qin, Mingyuan Zhang, Yifu Ding, Aoyu Li, Zhongang Cai, Ziwei Liu, Fisher Yu, Xianglong Liu:  
**BiBench: Benchmarking and Analyzing Network Binarization.** 28351-28388



Zi-Hao Qiu, Quanqi Hu, Zhuoning Yuan, Denny Zhou, Lijun Zhang, Tianbao Yang:  
**Not All Semantics are Created Equal: Contrastive Self-supervised Learning with Automatic Temperature Individualization.** 28389-28421



Xin Qiu, Risto Miikkulainen:  
**Shortest Edit Path Crossover: A Theory-driven Solution to the Permutation Problem in Evolutionary Neural Architecture Search.** 28422-28447



Shikai Qiu, Andres Potapczynski, Pavel Izmailov, Andrew Gordon Wilson:  
**Simple and Fast Group Robustness by Automatic Feature Reweighting.** 28448-28467



Francesco Quinza, Ashkan Soleymani, Patrick Jaillet, Cristian R. Rojas, Stefan Bauer:  
**DRCFS: Doubly Robust Causal Feature Selection.** 28468-28491



- Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, Ilya Sutskever:  
**Robust Speech Recognition via Large-Scale Weak Supervision.** 28492-28518
- Matthew Raffel, Drew Penney, Lihong Chen:  
**Shiftable Context: Addressing Training-Inference Context Mismatch in Simultaneous Speech Translation.** 28519-28530
- Aniruddh Raghu, Payal Chandak, Ridwan Alam, John V. Guttag, Collin M. Stultz:  
**Sequential Multi-Dimensional Self-Supervised Learning for Clinical Time Series.** 28531-28548
- Arman Rahbar, Ashkan Panahi, Morteza Haghir Chehreghani, Devdatt P. Dubhashi, Hamid Krim:  
**Recovery Bounds on Class-Based Optimal Transport: A Sum-of-Norms Regularization Framework.** 28549-28577
- Anant Raj, Lingjiong Zhu, Mert Gürbüzbalaban, Umut Simsekli:  
**Algorithmic Stability of Heavy-Tailed SGD with General Loss Functions.** 28578-28597
- Sai Rajeswar, Pietro Mazzaglia, Tim Verbelen, Alexandre Piché, Bart Dhoedt, Aaron C. Courville, Alexandre Lacoste:  
**Mastering the Unsupervised Reinforcement Learning Benchmark from Pixels.** 28598-28617
- Santhosh Kumar Ramakrishnan, Ziad Al-Halah, Kristen Grauman:  
**SpotEM: Efficient Video Search for Episodic Memory.** 28618-28636
- Sameera Ramasinghe, Lachlan Ewen MacDonald, Moshiur R. Farazi, Hemanth Saratchandran, Simon Lucey:  
**How much does Initialization Affect Generalization?** 28637-28655
- Alexandre Ramé, Kartik Ahuja, Jianyu Zhang, Matthieu Cord, Léon Bottou, David Lopez-Paz:  
**Model Ratatouille: Recycling Diverse Models for Out-of-Distribution Generalization.** 28656-28679
- Rahul Ramesh, Jialin Mao, Itay Griniasty, Rubing Yang, Han Kheng Teoh, Mark K. Transtrum, James P. Sethna, Pratik Chaudhari:  
**A Picture of the Space of Typical Learnable Tasks.** 28680-28700
- Yuhang Ran, Yi-Chen Li, Fuxiang Zhang, Zongzhang Zhang, Yang Yu:  
**Policy Regularization with Dataset Constraint for Offline Reinforcement Learning.** 28701-28717
- Ran Ran, Xinwei Luo, Wei Wang, Tao Liu, Gang Quan, Xiaolin Xu, Caiwen Ding, Wujie Wen:  
**SpENCNN: Orchestrating Encoding and Sparsity for Fast Homomorphically Encrypted Neural Network Inference.** 28718-28728
- Akshay Rangamani, Marius Lindegaard, Tomer Galanti, Tomaso A. Poggio:  
**Feature learning in deep classifiers through Intermediate Neural Collapse.** 28729-28745
- Sarah Rathnam, Sonali Parbhoo, Weiwei Pan, Susan A. Murphy, Finale Doshi-Velez:  
**The Unintended Consequences of Discount Regularization: Improving Regularization in Certainty Equivalence Reinforcement Learning.** 28746-28767
- Jishnu Ray Chowdhury, Cornelia Caragea:  
**Beam Tree Recursive Cells.** 28768-28791
- Jishnu Ray Chowdhury, Cornelia Caragea:  
**Monotonic Location Attention for Length Generalization.** 28792-28808
- Ofir Razon, Yoav Harris, Shahar Gottlieb, Dan Carmon, Ofir David, Ido Kaminer:  
**Automated Search for Conjectures on Mathematical Constants using Analysis of Integer Sequences.** 28809-28842
- Maria Refinetti, Alessandro Ingrosso, Sebastian Goldt:  
**Neural networks trained with SGD learn distributions of increasing complexity.** 28843-28863
- Isaac Reid, Krzysztof Marcin Choromanski, Valerii Likhoshevstov, Adrian Weller:  
**Simplex Random Features.** 28864-28888
- Qihan Ren, Huiqi Deng, Yunuo Chen, Siyu Lou , Quanshi Zhang:  
**Bayesian Neural Networks Avoid Encoding Complex and Perturbation-Sensitive Concepts.** 28889-28913
- Zhaolin Ren, Yujie Tang, Na Li:  
**Escaping saddle points in zeroth-order optimization: the power of two-point estimators.** 28914-28975
- Jiaxiang Ren, Yang Zhou, Jiayin Jin, Lingjuan Lyu, Da Yan:  
**Dimension-independent Certified Neural Network Watermarks via Mollifier Smoothing.** 28976-29008
- Alex Daniel Reneau, Jerry Yao-Chieh Hu, Ammar Gilani, Han Liu:  
**Feature Programming for Multivariate Time Series Prediction.** 29009-29029
- Keivan Rezaei, Kiarash Banihashem, Atoosa Malemir Chegini, Soheil Feizi:  
**Run-off Election: Improved Provable Defense against Data Poisoning Attacks.** 29030-29050
- Spencer M. Richards, Jean-Jacques E. Slotine, Navid Azizan, Marco Pavone:  
**Learning Control-Oriented Dynamical Structure from Data.** 29051-29062
- Pierre Harvey Richemond, Allison C. Tam, Yunhao Tang, Florian Strub, Bilal Piot, Felix Hill:  
**The Edge of Orthogonality: A Simple View of What Makes BYOL Tick.** 29063-29081
- Alexandre Rio, Merwan Barlier, Igor Colin, Marta Soare:  
**Multi-Agent Best Arm Identification with Private Communications.** 29082-29102
- Nicholas Rittler, Kamalika Chaudhuri:  
**A Two-Stage Active Learning Algorithm for k-Nearest Neighbors.** 29103-29129



- Yeonju Ro, Zhangyang Wang, Vijay Chidambaram, Aditya Akella:  
**Lowering the Pre-training Tax for Gradient-based Subset Training: A Lightweight Distributed Pre-Training Toolkit.** 29130-29142
- Borja Rodríguez Gálvez, Arno Blaas, Pau Rodríguez, Adam Golinski, Xavier Suau, Jason Ramapuram, Dan Busbridge, Luca Zappella:  
**The Role of Entropy and Reconstruction in Multi-View Self-Supervised Learning.** 29143-29160
- Rafael Rodríguez-Sánchez, Benjamin Adin Spiegel, Jennifer Wang, Roma Patel, Stefanie Tellex, George Konidaris:  
**RLang: A Declarative Language for Describing Partial World Knowledge to Reinforcement Learning Agents.** 29161-29178
- Yuji Roh, Kangwook Lee, Steven Euijong Whang, Changho Suh:  
**Improving Fair Training under Correlation Shifts.** 29179-29209
- Mark Rowland, Yunhao Tang, Clare Lyle, Rémi Munos, Marc G. Bellemare, Will Dabney:  
**The Statistical Benefits of Quantile Temporal-Difference Learning for Value Estimation.** 29210-29231
- Haolin Ruan, Siyu Zhou, Zhi Chen, Chin Pang Ho:  
**Robust Satisficing MDPs.** 29232-29258
- Mark Rucker, Yinglun Zhu, Paul Mineiro:  
**Infinite Action Contextual Bandits with Reusable Data Exhaust.** 29259-29274
- Tim G. J. Rudner, Sanyam Kapoor, Shikai Qiu, Andrew Gordon Wilson:  
**Function-Space Regularization in Neural Networks: A Probabilistic Perspective.** 29275-29290
- David Rügamer:  
**A New PHO-rmula for Improved Performance of Semi-Structured Networks.** 29291-29305
- David Ruhe, Jayesh K. Gupta, Steven De Keninck, Max Welling, Johannes Brandstetter:  
**Geometric Clifford Algebra Networks.** 29306-29337
- Davor Runje , Sharath M. Shankaranarayana:  
**Constrained Monotonic Neural Networks.** 29338-29353
- Phillip Rust, Anders Søgaard:  
**Differential Privacy, Linguistic Fairness, and Training Data Influence: Impossibility and Possibility Theorems for Multilingual Language Models.** 29354-29387
- Raif M. Rustamov, Subhabrata Majumdar:  
**Intrinsic Sliced Wasserstein Distances for Comparing Collections of Probability Distributions on Manifolds and Graphs.** 29388-29415
- Max Ryabinin, Tim Dettmers, Michael Diskin, Alexander Borzunov:  
**SWARM Parallelism: Training Large Models Can Be Surprisingly Communication-Efficient.** 29416-29440
- Chaitanya Ryali, Yuan-Ting Hu, Daniel Bolya, Chen Wei, Haoqi Fan, Po-Yao Huang, Vaibhav Aggarwal, Arkabandhu Chowdhury, Omid Poursaeed, Judy Hoffman, Jitendra Malik, Yanghao Li, Christoph Feichtenhofer:  
**Hiera: A Hierarchical Vision Transformer without the Bells-and-Whistles.** 29441-29454
- Yves Rychener, Daniel Kuhn, Tobias Sutter:  
**End-to-End Learning for Stochastic Optimization: A Bayesian Perspective.** 29455-29472
- Feras Saad, Brian Patton, Matthew Douglas Hoffman, Rif A. Saurous, Vikash Mansinghka:  
**Sequential Monte Carlo Learning for Time Series Structure Discovery.** 29473-29489
- El Mehdi Saad, Nicolas Verzelen, Alexandra Carpentier:  
**Active Ranking of Experts Based on their Performances in Many Tasks.** 29490-29513
- Seyed Amir Hossein Saberi, Amir Najafi, Abolfazl S. Motahari, Babak H. Khalaj:  
**Sample Complexity Bounds for Learning High-dimensional Simplices in Noisy Regimes.** 29514-29541
- Vin Sachidananda, Ziyi Yang, Chenguang Zhu:  
**Global Selection of Contrastive Batches via Optimization on Sample Permutations.** 29542-29562
- Abdurakhmon Sadiev, Marina Danilova, Eduard Gorbunov, Samuel Horváth, Gauthier Gidel, Pavel E. Dvurechensky, Alexander V. Gasnikov, Peter Richtárik:  
**High-Probability Bounds for Stochastic Optimization and Variational Inequalities: the Case of Unbounded Variance.** 29563-29648
- Bishwajit Saha, Dmitry Krotov, Mohammed J. Zaki, Parikshit Ram:  
**End-to-end Differentiable Clustering with Associative Memories.** 29649-29670
- Eden Saig, Nir Rosenfeld:  
**Learning to Suggest Breaks: Sustainable Optimization of Long-Term User Engagement.** 29671-29696
- Shota Saito, Mark Herbster:  
**Multi-class Graph Clustering via Approximated Effective p-Resistance.** 29697-29733
- Yuta Saito, Qingyang Ren, Thorsten Joachims:  
**Off-Policy Evaluation for Large Action Spaces via Conjunct Effect Modeling.** 29734-29759
- Shinsaku Sakaue, Taihei Oki:  
**Rethinking Warm-Starts with Predictions: Learning Predictions Close to Sets of Optimal Solutions for Faster L-L<sup>1</sup>-Convex Function Minimization.** 29760-29776



- Otmane Sakhi, Pierre Alquier, Nicolas Chopin:  
**PAC-Bayesian Offline Contextual Bandits With Guarantees.** 29777-29799
- Sudeep Salgia:  
**Provably and Practically Efficient Neural Contextual Bandits.** 29800-29844
- Sudeep Salgia, Qing Zhao:  
**Distributed Linear Bandits under Communication Constraints.** 29845-29875
- David Salinas, Jacek Golebiowski, Aaron Klein, Matthias W. Seeger, Cédric Archambeau:  
**Optimizing Hyperparameters with Conformal Quantile Regression.** 29876-29893
- Hadi Salman, Alaa Khaddaj, Guillaume Leclerc, Andrew Ilyas, Aleksander Madry:  
**Raising the Cost of Malicious AI-Powered Image Editing.** 29894-29918
- Michael Eli Sander, Joan Puigcerver, Josip Djolonga, Gabriel Peyré, Mathieu Blondel:  
**Fast, Differentiable and Sparse Top-k: a Convex Analysis Perspective.** 29919-29936
- Tom Sander, Pierre Stock, Alexandre Sablayrolles:  
**TAN Without a Burn: Scaling Laws of DP-SGD.** 29937-29949
- Parth Vipul Sangani, Arjun Shashank Kashettiwar, Pritish Chakraborty, Bhuvan Reddy Gangula, Durga Sivasubramanian, Ganesh Ramakrishnan, Rishabh K. Iyer, Abir De:  
**Discrete Continuous Optimization Framework for Simultaneous Clustering and Training in Mixture Models.** 29950-29970
- Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, Tatsunori Hashimoto:  
**Whose Opinions Do Language Models Reflect?** 29971-30004
- Akanksha Saran, Safoora Yousefi, Akshay Krishnamurthy, John Langford, Jordan T. Ash:  
**Streaming Active Learning with Deep Neural Networks.** 30005-30021
- Felix Sarnthein, Gregor Bachmann, Sotiris Anagnostidis, Thomas Hofmann:  
**Random Teachers are Good Teachers.** 30022-30041
- Remo Sasso, Michelangelo Conserva, Paulo E. Rauber:  
**Posterior Sampling for Deep Reinforcement Learning.** 30042-30061
- Ryoma Sato:  
**Graph Neural Networks can Recover the Hidden Features Solely from the Graph Structure.** 30062-30079
- Naoki Sato, Hideaki Iiduka:  
**Existence and Estimation of Critical Batch Size for Training Generative Adversarial Networks with Two Time-Scale Update Rule.** 30080-30104
- Axel Sauer, Tero Karras, Samuli Laine, Andreas Geiger, Timo Aila:  
**StyleGAN-T: Unlocking the Power of GANs for Fast Large-Scale Text-to-Image Synthesis.** 30105-30118
- Andrey V. Savchenko:  
**Facial Expression Recognition with Adaptive Frame Rate based on Multiple Testing Correction.** 30119-30129
- Naman Saxena, Subhojoyoti Khastagir, Shishir Kolathaya, Shalabh Bhatnagar:  
**Off-Policy Average Reward Actor-Critic with Deterministic Policy Search.** 30130-30203
- Philip Schär, Michael Habeck, Daniel Rudolf:  
**Gibbsian Polar Slice Sampling.** 30204-30223
- Andreas Schlaginhaufen, Maryam Kamgarpour:  
**Identifiability and Generalizability in Constrained Inverse Reinforcement Learning.** 30224-30251
- Dominik Schnaus, Jongseok Lee, Daniel Cremers, Rudolph Triebel:  
**Learning Expressive Priors for Generalization and Uncertainty Estimation in Neural Networks.** 30252-30284
- Dominik Schröder, Hugo Cui, Daniil Dmitriev , Bruno Loureiro:  
**Deterministic equivalent and error universality of deep random features learning.** 30285-30320
- Luca M. Schulze Buschoff, Eric Schulz, Marcel Binz:  
**The Acquisition of Physical Knowledge in Generative Neural Networks.** 30321-30341
- Jonathan Richard Schwarz, Jihoon Tack, Yee Whye Teh, Jaeho Lee, Jinwoo Shin:  
**Modality-Agnostic Variational Compression of Implicit Neural Representations.** 30342-30364
- Max Schwarzer, Johan Samir Obando-Ceron, Aaron C. Courville, Marc G. Bellemare, Rishabh Agarwal, Pablo Samuel Castro:  
**Bigger, Better, Faster: Human-level Atari with human-level efficiency.** 30365-30380
- Antonio Sclocchi , Mario Geiger, Matthieu Wyart:  
**Dissecting the Effects of SGD Noise in Distinct Regimes of Deep Learning.** 30381-30405
- Michael Sedlmayer, Dang-Khoa Nguyen , Radu Ioan Bot:  
**A Fast Optimistic Method for Monotone Variational Inequalities.** 30406-30438
- José Ignacio Segovia-Martín, Santiago Mazuelas, Anqi Liu:  
**Double-Weighting for Covariate Shift Adaptation.** 30439-30457
- Philipp Seidl, Andreu Vall, Sepp Hochreiter, Günter Klambauer:  
**Enhancing Activity Prediction Models in Drug Discovery with the Ability to Understand Human Language.** 30458-30490



- Jacob H. Seidman, Georgios Kissas, George J. Pappas, Paris Perdikaris:  
**Variational Autoencoding Neural Operators.** 30491-30522
- Patrick Seifner, Ramsés J. Sánchez:  
**Neural Markov Jump Processes.** 30523-30552
- Jeremy Sellier, Petros Dellaportas:  
**Bayesian online change point detection with Hilbert space approximate Student-t process.** 30553-30569
- Mark Sellke:  
**Incentivizing Exploration with Linear Contexts and Combinatorial Actions.** 30570-30583
- Hugo Henri Joseph Sénétaire, Damien Garreau, Jes Frellsen, Pierre-Alexandre Mattei:  
**Explainability as statistical inference.** 30584-30612
- Younggyo Seo, Junsu Kim, Stephen James, Kimin Lee, Jinwoo Shin, Pieter Abbeel:  
**Multi-View Masked World Models for Visual Robotic Manipulation.** 30613-30632
- Daniel Severo, James Townsend, Ashish J. Khisti, Alireza Makhzani:  
**One-Shot Compression of Large Edge-Exchangeable Graphs using Bits-Back Coding.** 30633-30645
- Harshay Shah, Sung Min Park, Andrew Ilyas, Aleksander Madry:  
**ModelDiff: A Framework for Comparing Learning Algorithms.** 30646-30688
- Aviv Shamsian, Aviv Navon, Neta Glazer, Kenji Kawaguchi, Gal Chechik, Ethan Fetaya:  
**Auxiliary Learning as an Asymmetric Bargaining Game.** 30689-30705
- Zhihong Shao, Yeyun Gong, Yelong Shen, Minlie Huang, Nan Duan, Weizhu Chen:  
**Synthetic Prompting: Generating Chain-of-Thought Demonstrations for Large Language Models.** 30706-30775
- Jianzhun Shao, Hongchang Zhang, Yun Qu, Chang Liu, Shuncheng He, Yuhang Jiang, Xiangyang Ji:  
**Complementary Attention for Multi-Agent Reinforcement Learning.** 30776-30793
- Ron Shapira Weber, Oren Freifeld:  
**Regularization-free Diffeomorphic Temporal Alignment Nets.** 30794-30826
- Arsalan Sharifnassab, Richard S. Sutton:  
**Toward Efficient Gradient-Based Value Estimation.** 30827-30849
- Louis Sharrock, Christopher Nemeth:  
**Coin Sampling: Gradient-Based Bayesian Inference without Learning Rates.** 30850-30882
- Neta Shaul, Ricky T. Q. Chen, Maximilian Nickel, Matthew Le, Yaron Lipman:  
**On Kinetic Optimal Probability Paths for Generative Models.** 30883-30907
- Shubhangshu Shekhar, Aaditya Ramdas:  
**Sequential Changepoint Detection via Backward Confidence Sequences.** 30908-30930
- Alexander Shekhovtsov:  
**Cold Analysis of Rao-Blackwellized Straight-Through Gumbel-Softmax Gradient Estimator.** 30931-30955
- Max W. Shen, Emmanuel Bengio, Ehsan Hajiramezanali, Andreas Loukas, Kyunghyun Cho, Tommaso Biancalani:  
**Towards Understanding and Improving GFlowNet Training.** 30956-30975
- Maohao Shen, Yuheng Bu, Gregory W. Wornell:  
**On Balancing Bias and Variance in Unsupervised Multi-Source-Free Domain Adaptation.** 30976-30991
- Han Shen, Tianyi Chen:  
**On Penalty-based Bilevel Gradient Descent Method.** 30992-31015
- Lifeng Shen, James T. Kwok:  
**Non-autoregressive Conditional Diffusion Models for Time Series Prediction.** 31016-31029
- Junhong Shen, Liam Li, Lucio M. Dery, Corey Staten, Mikhail Khodak, Graham Neubig, Ameet Talwalkar:  
**Cross-Modal Fine-Tuning: Align then Refine.** 31030-31056
- Yu Shen, Xijun Wang, Peng Gao, Ming C. Lin:  
**Auxiliary Modality Learning with Generalized Curriculum Distillation.** 31057-31076
- Idan Shenfeld, Zhang-Wei Hong, Aviv Tamar, Pulkit Agrawal:  
**TGRL: An Algorithm for Teacher Guided Reinforcement Learning.** 31077-31093
- Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Beidi Chen, Percy Liang, Christopher Ré, Ion Stoica, Ce Zhang:  
**FlexGen: High-Throughput Generative Inference of Large Language Models with a Single GPU.** 31094-31116
- Uri Sherman, Tomer Koren, Yishay Mansour:  
**Improved Regret for Efficient Online Reinforcement Learning with Linear Function Approximation.** 31117-31150
- Aleksandr Shevchenko, Kevin Kögler, Hamed Hassani, Marco Mondelli:  
**Fundamental Limits of Two-layer Autoencoders, and Achieving Them with Gradient Methods.** 31151-31209
- Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Schärli, Denny Zhou:  
**Large Language Models Can Be Easily Distracted by Irrelevant Context.** 31210-31227
- Hui Shi, Yupeng Gu, Yitong Zhou, Bo Zhao, Sicun Gao, Jishen Zhao:  
**Everyone's Preference Changes Differently: A Weighted Multi-Interest Model For Retrieval.** 31228-31242



- Ming Shi, Yingbin Liang, Ness B. Shroff:  
**A Near-Optimal Algorithm for Safe Reinforcement Learning Under Instantaneous Hard Constraints.** 31243-31268
- Yifan Shi, Li Shen, Kang Wei, Yan Sun, Bo Yuan, Xueqian Wang, Dacheng Tao:  
**Improving the Model Consistency of Decentralized Federated Learning.** 31269-31291
- Dachuan Shi, Chaofan Tao, Ying Jin, Zhendong Yang, Chun Yuan, Jiaqi Wang:  
**UPop: Unified and Progressive Pruning for Compressing Vision-Language Transformers.** 31292-31311
- Jiaxin Shi, Ke Alexander Wang, Emily B. Fox:  
**Sequence Modeling with Multiresolution Convolutional Memory.** 31312-31327
- Lei Shi, Jingshen Wang, Tianhao Wu:  
**Statistical Inference on Multi-armed Bandits with Delayed Feedback.** 31328-31352
- Chengshuai Shi, Wei Xiong, Cong Shen, Jing Yang:  
**Provably Efficient Offline Reinforcement Learning with Perturbed Data Sources.** 31353-31388
- Yu-Zhe Shi, Manjie Xu, John E. Hopcroft, Kun He, Joshua B. Tenenbaum, Song-Chun Zhu, Ying Nian Wu, Wenjuan Han, Yixin Zhu:  
**On the Complexity of Bayesian Generalization.** 31389-31407
- Liangliang Shi, Gu Zhang, Haoyu Zhen, Jintao Fan, Junchi Yan:  
**Understanding and Generalizing Contrastive Learning from the Inverse Optimal Transport Perspective.** 31408-31421
- Andy Shih, Dorsa Sadigh, Stefano Ermon:  
**Long Horizon Temperature Scaling.** 31422-31434
- Alistair Shilton, Sunil Gupta, Santu Rana, Svetha Venkatesh:  
**Gradient Descent in Neural Networks as Sequential Learning in Reproducing Kernel Banach Space.** 31435-31488
- Dongseok Shim, Seungjae Lee, H. Jin Kim:  
**SNeRL: Semantic-aware Neural Radiance Fields for Reinforcement Learning.** 31489-31503
- Sungbin Shin, Yohan Jo, Sungsoo Ahn, Namhoon Lee:  
**A Closer Look at the Intervention Procedure of Concept Bottleneck Models.** 31504-31520
- Wooseok Shin, Byung Hoon Lee, Jin Sob Kim, Hyun Joon Park, Sung Won Han:  
**MetricGAN-OKD: Multi-Metric Optimization of MetricGAN via Online Knowledge Distillation for Speech Enhancement.** 31521-31538
- Yongho Shin, Changyeol Lee, Gukryeol Lee, Hyung-Chan An:  
**Improved Learning-Augmented Algorithms for the Multi-Option Ski Rental Problem via Best-Possible Competitive Analysis.** 31539-31561
- Sangwoo Shin, Daehee Lee, Minjong Yoo, Woo Kyung Kim, Honguk Woo:  
**One-shot Imitation in a Non-Stationary Environment via Multi-Modal Skill.** 31562-31578
- Yooju Shin, Susik Yoon, Hwanjun Song, Dongmin Park, Byunghyun Kim, Jae-Gil Lee, Byung Suk Lee:  
**Context Consistency Regularization for Label Sparsity in Time Series.** 31579-31595
- Shayan Shirahmad Gale Bagi, Zahra Gharaee, Oliver Schulte, Mark Crowley:  
**Generative Causal Representation Learning for Out-of-Distribution Motion Forecasting.** 31596-31612
- Hamed Shirzad, Ameya Velingker, Balaji Venkatachalam, Danica J. Sutherland, Ali Kemal Sinop:  
**Exphormer: Sparse Transformers for Graphs.** 31613-31632
- Alon Shoshan, Nadav Bhonker, Igor Kviatkovsky, Matan Fintz, Gérard G. Medioni:  
**Synthetic data for model selection.** 31633-31656
- Xiao Shou, Debarun Bhattacharjya, Tian Gao, Dharmashankar Subramanian, Oktie Hassanzadeh, Kristin P. Bennett:  
**Probabilistic Attention-to-Influence Neural Models for Event Sequences.** 31657-31674
- Madhumitha Shridharan, Garud Iyengar:  
**Causal Bounds in Quasi-Markovian Graphs.** 31675-31692
- Disha Srivastava, Hugo Larochelle, Daniel Tarlow:  
**Repository-Level Prompt Generation for Large Language Models of Code.** 31693-31715
- Yang Shu, Xingzhuo Guo, Jialong Wu, Ximei Wang, Jianmin Wang, Mingsheng Long:  
**CLIPood: Generalizing CLIP to Out-of-Distributions.** 31716-31731
- Phillip Si, Zeyi Chen, Subham Sekhar Sahoo, Yair Schiff, Volodymyr Kuleshov:  
**Semi-Autoregressive Energy Flows: Exploring Likelihood-Free Training of Normalizing Flows.** 31732-31753
- Ali Siahkoohi, Rudy Morel, Maarten V. de Hoop, Erwan Ally, Grégory Sainton, Taichi Kawamura:  
**Unearthing InSights into Mars: Unsupervised Source Separation with Limited Data.** 31754-31772
- Julian Sieber, Johann Gehringer:  
**Quantitative Universal Approximation Bounds for Deep Belief Networks.** 31773-31787
- David Simchi-Levi, Chonghuan Wang:  
**Pricing Experimental Design: Causal Effect, Expected Revenue and Tail Risk.** 31788-31799
- Max Simchowitz, Anurag Ajay, Pulkit Agrawal, Akshay Krishnamurthy:  
**Statistical Learning under Heterogenous Distribution Shift.** 31800-31851



- James B. Simon, Maksis Knutins, Ziyin Liu, Daniel Geisz, Abraham J. Fetterman, Joshua Albrecht:  
**On the Stepwise Nature of Self-Supervised Learning.** 31852-31876
- Sean R. Sinclair, Felipe Vieira Frujeri, Ching-An Cheng, Luke Marshall, Hugo de Oliveira Barbalho, Jingling Li, Jennifer Neville, Ishai Menache, Adith Swaminathan:  
**Hindsight Learning for MDPs with Exogenous Inputs.** 31877-31914
- Uriel Singer, Shelly Sheynin, Adam Polyak, Oron Ashual, Iurii Makarov, Filippos Kokkinos, Naman Goyal, Andrea Vedaldi, Devi Parikh, Justin Johnson, Yaniv Taigman:  
**Text-To-4D Dynamic Scene Generation.** 31915-31929
- Sidak Pal Singh, Thomas Hofmann, Bernhard Schölkopf:  
**The Hessian perspective into the Nature of Convolutional Neural Networks.** 31930-31968
- Harvineet Singh, Matthäus Kleindessner, Volkan Cevher, Rumi Chunara, Chris Russell:  
**When do Minimax-fair Learning and Empirical Risk Minimization Coincide?** 31969-31989
- Martin Sípka, Johannes C. B. Dietschreit , Lukás Grajciar, Rafael Gómez-Bombarelli:  
**Differentiable Simulations for Enhanced Sampling of Rare Events.** 31990-32007
- Chawin Sitawarin, Florian Tramèr, Nicholas Carlini:  
**Preprocessors Matter! Realistic Decision-Based Attacks on Machine Learning Systems.** 32008-32032
- Joar Max Viktor Skalse, Matthew Farrugia-Roberts, Stuart Russell, Alessandro Abate, Adam Gleave:  
**Invariance in Policy Optimisation and Partial Identifiability in Reward Learning.** 32033-32058
- Oliver Slumbers, David Henry Mguni, Stefano B. Blumberg, Stephen Marcus McAleer, Yaodong Yang, Jun Wang:  
**A Game-Theoretic Framework for Managing Risk in Multi-Agent Systems.** 32059-32087
- Paloma Sodhi, Felix Wu, Ethan R. Elenberg, Kilian Q. Weinberger, Ryan McDonald:  
**On the Effectiveness of Offline RL for Dialogue Response Generation.** 32088-32104
- Alexander Soen, Hisham Husain, Richard Nock:  
**Fair Densities via Boosting the Sufficient Statistics of Exponential Families.** 32105-32144
- Ghada Sokar, Rishabh Agarwal, Pablo Samuel Castro, Utku Evci:  
**The Dormant Neuron Phenomenon in Deep Reinforcement Learning.** 32145-32168
- Samuel Sokota, Ryan D'Orazio, Chun Kai Ling, David J. Wu, J. Zico Kolter, Noam Brown:  
**Abstracting Imperfect Information Away from Two-Player Zero-Sum Games.** 32169-32193
- Jiwoo Son, Minsu Kim, Hyeonah Kim, Jinkyoo Park:  
**Meta-SAGE: Scale Meta-Learning Scheduled Adaptation with Guided Exploration for Mitigating Scale Shift on Combinatorial Optimization.** 32194-32210
- Yang Song, Prafulla Dhariwal, Mark Chen, Ilya Sutskever:  
**Consistency Models.** 32211-32252
- Xujie Song, Jingliang Duan, Wenxuan Wang, Shengbo Eben Li, Chen Chen, Bo Cheng, Bo Zhang, Junqing Wei, Xiaoming Simon Wang:  
**LipsNet: A Smooth and Robust Neural Network with Adaptive Lipschitz Constant for High Accuracy Optimal Control.** 32253-32272
- Zifan Song, Xiao Gong, Guosheng Hu, Cairong Zhao:  
**Deep Perturbation Learning: Enhancing the Network Performance via Image Perturbations.** 32273-32287
- Yue Song, T. Anderson Keller, Nicu Sebe, Max Welling:  
**Latent Traversals in Generative Models as Potential Flows.** 32288-32303
- Bingqing Song, Prashant Khanduri, Xinwei Zhang, Jinfeng Yi, Mingyi Hong:  
**FedAvg Converges to Zero Training Loss Linearly for Overparameterized Multi-Layer Neural Networks.** 32304-32330
- Jaeyun Song, Sungyub Kim, Eunho Yang:  
**RGE: A Repulsive Graph Rectification for Node Classification via Influence.** 32331-32348
- Zhenqiao Song, Lei Li:  
**Importance Weighted Expectation-Maximization for Protein Sequence Design.** 32349-32364
- Zhao Song, Yitan Wang, Zheng Yu, Lichen Zhang:  
**Sketching for First Order Method: Efficient Algorithm for Low-Bandwidth Channel and Vulnerability.** 32365-32417
- Zhao Song, Xin Yang, Yuanyuan Yang, Lichen Zhang:  
**Sketching Meets Differential Privacy: Fast Algorithm for Dynamic Kronecker Projection Maintenance.** 32418-32462
- Zhao Song, Mingquan Ye, Junze Yin, Lichen Zhang:  
**A Nearly-Optimal Bound for Fast Regression with  $\ell_\infty$  Guarantee.** 32463-32482
- Jiaming Song, Qinsheng Zhang, Hongxu Yin, Morteza Mardani, Ming-Yu Liu, Jan Kautz, Yongxin Chen, Arash Vahdat:  
**Loss-Guided Diffusion Models for Plug-and-Play Controllable Generation.** 32483-32498
- Paul Soulos, Edward J. Hu, Kate McCurdy, Yunmo Chen, Roland Fernandez, Paul Smolensky, Jianfeng Gao:  
**Differentiable Tree Operations Promote Compositional Generalization.** 32499-32520
- Aude Sportisse, Hugo Schmutz, Olivier Humbert, Charles Bouveyron, Pierre-Alexandre Mattei:  
**Are labels informative in semi-supervised learning? Estimating and leveraging the missing-data mechanism.** 32521-32539



- Chandler Squires, Anna Seigal, Salil S. Bhate, Caroline Uhler:  
**Linear Causal Disentanglement via Interventions.** 32540-32560
- Megha Srivastava, Noah D. Goodman, Dorsa Sadigh:  
**Generating Language Corrections for Teaching Physical Control Tasks.** 32561-32574
- Guillaume Staerman, Cédric Allain, Alexandre Gramfort, Thomas Moreau:  
**FaDIn: Fast Discretized Inference for Hawkes Processes with General Parametric Kernels.** 32575-32597
- David Stein, Silvia Di Gregorio, Bjoern Andres:  
**Partial Optimality in Cubic Correlation Clustering.** 32598-32617
- Benoit Steiner, Mostafa Elhoushi, Jacob Kahn, James Hegarty:  
**MODeL: Memory Optimizations for Deep Learning.** 32618-32632
- Eleni Straitouri, Lequn Wang, Nastaran Okati, Manuel Gomez Rodriguez:  
**Improving Expert Predictions with Conformal Prediction.** 32633-32653
- Grant P. Strimel, Yi Xie, Brian John King, Martin Radfar, Ariya Rastrow, Athanasios Mouchtaris:  
**Lookahead When It Matters: Adaptive Non-causal Transformers for Streaming Neural Transducers.** 32654-32676
- Diego Stucchi, Paolo Rizzo, Nicolò Folloni, Giacomo Boracchi:  
**Kernel QuantTree.** 32677-32697
- Nico Stucki, Johannes C. Paetzold, Suprosanna Shit, Bjoern H. Menze, Ulrich Bauer:  
**Topologically Faithful Image Segmentation via Induced Matching of Persistence Barcodes.** 32698-32727
- Junwei Su, Difan Zou, Zijun Zhang, Chuan Wu:  
**Towards Robust Graph Incremental Learning on Evolving Graphs.** 32728-32748
- Xavier Suau, Federico Danieli, T. Anderson Keller, Arno Blaas, Chen Huang, Jason Ramapuram, Dan Busbridge, Luca Zappella:  
**DUET: 2D Structured and Approximately Equivariant Representations.** 32749-32769
- Min-Kook Suh, Seung-Woo Seo:  
**Long-Tailed Recognition by Mutual Information Maximization between Latent Features and Ground-Truth Labels.** 32770-32782
- Yang Sui, Yukun Huang, Hongtu Zhu, Fan Zhou:  
**Adversarial Learning of Distributional Reinforcement Learning.** 32783-32796
- Theodore R. Sumers, Kenneth Marino, Arun Ahuja, Rob Fergus, Ishita Dasgupta:  
**Distilling Internet-Scale Vision-Language Models into Embodied Agents.** 32797-32818
- Zhuo Sun, Alessandro Barp, François-Xavier Briol:  
**Vector-Valued Control Variates.** 32819-32846
- Wenfang Sun, Yingjun Du, Xiantong Zhen, Fan Wang, Ling Wang, Cees G. M. Snoek:  
**MetaModulation: Learning Variational Feature Hierarchies for Few-Shot Learning with Fewer Tasks.** 32847-32858
- Haoran Sun, Katayoon Goshvadi, Azade Nova, Dale Schuurmans, Hanjun Dai:  
**Revisiting Sampling for Combinatorial Optimization.** 32859-32874
- Zequn Sun, Jiacheng Huang, Xiaozhou Xu, Qijin Chen, Weijun Ren, Wei Hu:  
**What Makes Entities Similar? A Similarity Flooding Perspective for Multi-sourced Knowledge Graph Embeddings.** 32875-32885
- Chunlin Sun, Shang Liu, Xiaocheng Li:  
**Maximum Optimality Margin: A Unified Approach for Contextual Linear Programming and Inverse Linear Programming.** 32886-32912
- Hu Sun, Ward Manchester, Meng Jin, Yang Liu, Yang Chen:  
**Tensor Gaussian Process with Contraction for Multi-Channel Imaging Analysis.** 32913-32935
- Jennifer J. Sun, Markus Marks, Andrew Wesley Ulmer, Dipam Chakraborty, Brian Geuther, Edward Hayes, Heng Jia, Vivek Kumar, Sebastian Oleszko, Zachary Partridge, Milan Peelman, Alice Robie, Catherine E. Schretter, Keith Sheppard, Chao Sun, Param Uttarwar, Julian Morgan Wagner, Erik Werner, Joseph Parker, Pietro Perona, Yisong Yue, Kristin Branson, Ann Kennedy:  
**MABe22: A Multi-Species Multi-Task Benchmark for Learned Representations of Behavior.** 32936-32990
- Yan Sun, Li Shen, Shixiang Chen, Liang Ding, Dacheng Tao:  
**Dynamic Regularized Sharpness Aware Minimization in Federated Learning: Approaching Global Consistency and Smooth Landscape.** 32991-33013
- Yiyou Sun, Zhenmei Shi, Yingyu Liang, Yixuan Li:  
**When and How Does Known Class Help Discover Unknown Ones? Provable Understanding Through Spectral Analysis.** 33014-33043
- Wei Sun, Asterios Tsourvas:  
**Learning Prescriptive ReLU Networks.** 33044-33060
- Junshu Sun, Shuhui Wang, Xinzhe Han, Zhe Xue, Qingming Huang:  
**All in a Row: Compressed Convolution Networks for Graphs.** 33061-33076
- Tao Sun, Qingsong Wang, Dongsheng Li, Bao Wang:  
**Momentum Ensures Convergence of SIGNSGD under Weaker Assumptions.** 33077-33099



Jiachen Sun, Jiongxiao Wang, Weili Nie, Zhiding Yu, Zuoqing Mao, Chaowei Xiao:  
**A Critical Revisit of Adversarial Robustness in 3D Point Cloud Recognition with Diffusion-Driven Purification.** 33100-33114

Shikun Sun, Longhui Wei, Junliang Xing, Jia Jia, Qi Tian:  
**SDDM: Score-Decomposed Diffusion Models on Manifolds for Unpaired Image-to-Image Translation.** 33115-33134

Zhiqing Sun, Yiming Yang, Shinjae Yoo:  
**A Neural PDE Solver with Temporal Stencil Modeling.** 33135-33155

Jiaqi Sun, Lin Zhang, Guangyi Chen, Peng Xu, Kun Zhang, Yujiu Yang:  
**Feature Expansion for Graph Neural Networks.** 33156-33176

Yihao Sun, Jiaji Zhang, Chengxing Jia, Haoxin Lin, Junyin Ye, Yang Yu:  
**Model-Bellman Inconsistency for Model-based Offline Reinforcement Learning.** 33177-33194

Mukund Sundararajan, Walid Krichene:  
**Inflow, Outflow, and Reciprocity in Machine Learning.** 33195-33208

Vinith Menon Suriyakumar, Marzyeh Ghazemi, Berk Ustun :  
**When Personalization Harms Performance: Reconsidering the Use of Group Attributes in Prediction.** 33209-33228

André Susano Pinto, Alexander Kolesnikov, Yuge Shi, Lucas Beyer, Xiaohua Zhai:  
**Tuning Computer Vision Models With Task Rewards.** 33229-33239

Wesley A. Suttle, Amrit S. Bedi, Bhrij Patel, Brian M. Sadler, Alec Koppel, Dinesh Manocha:  
**Beyond Exponentially Fast Mixing in Average-Reward Reinforcement Learning via Multi-Level Monte Carlo Actor-Critic.** 33240-33267

Atsushi Suzuki, Atsushi Nitanda, Taiji Suzuki, Jing Wang, Feng Tian, Kenji Yamanishi:  
**Tight and fast generalization error bound of graph embedding in metric space.** 33268-33284

Erik Sverdrup, Yifan Cui:  
**Proximal Causal Learning of Conditional Average Treatment Effects.** 33285-33298

Gokul Swamy, David Wu, Sanjiban Choudhury, Drew Bagnell, Zhiwei Steven Wu:  
**Inverse Reinforcement Learning without Reinforcement Learning.** 33299-33318

Kirk Swanson, Jake Lawrence Williams, Eric M. Jonas:  
**Von Mises Mixture Distributions for Molecular Conformation Generation.** 33319-33342

Yasa Syed, Guanyang Wang:  
**Optimal randomized multilevel Monte Carlo for repeatedly nested expectations.** 33343-33364

Andrew Szot, Unnat Jain, Dhruv Batra, Zsolt Kira, Ruta Desai, Akshara Rai:  
**Adaptive Coordination in Social Embodied Rearrangement.** 33365-33380

Ali Taghribakhshi, Nicolas Nytka, Tareq Uz Zaman, Scott P. MacLachlan, Luke N. Olson, Matthew West:  
**MG-GNN: Multigrid Graph Neural Networks for Learning Multilevel Domain Decomposition Methods.** 33381-33395

Wai Ming Tai, Bryon Aragam:  
**Learning Mixtures of Gaussians with Censored Data.** 33396-33415

Shokichi Takakura, Taiji Suzuki:  
**Approximation and Estimation Ability of Transformers for Sequence-to-Sequence Functions with Infinite Dimensional Input.** 33416-33447

Makoto Takamoto, Francesco Alesiani, Mathias Niepert:  
**Learning Neural PDE Solvers with Parameter-Guided Channel Attention.** 33448-33467

Kei Takemura:  
**Contextual Conservative Interleaving Bandits.** 33468-33489

Shion Takeno, Yu Inatsu, Masayuki Karasuyama:  
**Randomized Gaussian Process Upper Confidence Bound with Tighter Bayesian Regret Bounds.** 33490-33515

Shion Takeno, Masahiro Nomura, Masayuki Karasuyama:  
**Towards Practical Preferential Bayesian Optimization with Skew Gaussian Processes.** 33516-33533

Zeren Tan, Yang Tian:  
**Robust Explanation for Free or At the Cost of Faithfulness.** 33534-33562

Xiaoyu Tan, Lin Yong, Shengyu Zhu, Chao Qu, Xihe Qiu, Yinghui Xu, Peng Cui, Yuan Qi:  
**Provably Invariant Learning without Domain Information.** 33563-33580

Yuxin Tang, Zhimin Ding, Dimitrije Jankov, Binhang Yuan, Daniel Bourgeois, Chris Jermaine:  
**Auto-Differentiation of Relational Computations for Very Large Scale Machine Learning.** 33581-33598

Xiaohang Tang, Le Cong Dinh, Stephen Marcus McAleer, Yaodong Yang:  
**Regret-Minimizing Double Oracle for Extensive-Form Games.** 33599-33615

Hao Tang, Kevin Ellis:  
**From Perception to Programs: Regularize, Overparameterize, and Amortize.** 33616-33631

Yunhao Tang, Zhaoan Daniel Guo, Pierre Harvey Richemond, Bernardo Ávila Pires, Yash Chandak, Rémi Munos, Mark Rowland, Mohammad Gheshlaghi Azar, Charline Le Lan, Clare Lyle, András György, Shantanu Thakoor, Will Dabney, Bilal Piot, Daniele Calandriello, Michal Valko:  
**Understanding Self-Predictive Learning for Reinforcement Learning.** 33632-33656



- Yunhao Tang, Tadashi Kozuno, Mark Rowland, Anna Harutyunyan, Rémi Munos, Bernardo Ávila Pires, Michal Valko:  
**DoMo-AC: Doubly Multi-step Off-policy Actor-Critic Algorithm.** 33657-33673
- Huayi Tang, Yong Liu:  
**Towards Understanding Generalization of Graph Neural Networks.** 33674-33719
- Yunhao Tang, Rémi Munos:  
**Towards a better understanding of representation dynamics under TD-learning.** 33720-33738
- Yunhao Tang, Rémi Munos, Mark Rowland, Michal Valko:  
**VA-learning as a more efficient alternative to Q-learning.** 33739-33757
- Ling Tang, Wen Shen, Zhanpeng Zhou , Yuefeng Chen, Quanshi Zhang:  
**Defects of Convolutional Decoder Networks in Frequency Representation.** 33758-33791
- Caizhi Tang, Huiyuan Wang , Xinyu Li, Qing Cui, Longfei Li, Jun Zhou:  
**Difference-in-Differences Meets Tree-based Methods: Heterogeneous Treatment Effects Estimation with Unmeasured Confounding.** 33792-33803
- Shohei Taniguchi, Masahiro Suzuki, Yusuke Iwasawa, Yutaka Matsuo:  
**End-to-end Training of Deep Boltzmann Machines by Unbiased Contrastive Divergence with Local Mode Initialization.** 33804-33815
- Korawat Tanwisuth, Shujian Zhang, Huangjie Zheng, Pengcheng He, Mingyuan Zhou:  
**POUF: Prompt-Oriented Unsupervised Fine-tuning for Large Pre-trained Models.** 33816-33832
- Linwei Tao, Minjing Dong, Chang Xu:  
**Dual Focal Loss for Calibration.** 33833-33849
- Stone Tao, Xiaochen Li, Tongzhou Mu, Zhiao Huang, Yuzhe Qin, Hao Su:  
**Abstract-to-Executable Trajectory Translation for One-Shot Task Generalization.** 33850-33882
- Rohan Taori, Tatsunori Hashimoto:  
**Data Feedback Loops: Model-driven Amplification of Dataset Biases.** 33883-33920
- Ayush Kumar Tarun, Vikram Singh Chundawat, Murari Mandal, Mohan S. Kankanhalli:  
**Deep Regression Unlearning.** 33921-33939
- Jacopo Teneggi, Matthew Tivnan, J. Webster Stayman, Jeremias Sulam:  
**How to Trust Your Diffusion Model: A Convex Optimization Approach to Conformal Risk Control.** 33940-33960
- Jay Tenenbaum, Haim Kaplan, Yishay Mansour, Uri Stemmer:  
**Concurrent Shuffle Differential Privacy Under Continual Observation.** 33961-33982
- Jiaye Teng, Bohang Zhang, Ruichen Li, Haowei He, Yequan Wang, Yan Tian, Yang Yuan:  
**Finding Generalization Measures by Contrasting Signal and Noise.** 33983-34010
- Guy Tennenholz, Nadav Merlis, Lior Shani, Martin Mladenov, Craig Boutilier:  
**Reinforcement Learning with History Dependent Dynamic Contexts.** 34011-34053
- Lucile Ter-Minassian, Oscar Clivio, Karla DiazOrdaz, Robin J. Evans, Christopher C. Holmes:  
**PWSHAP: A Path-Wise Explanation Model for Targeted Variables.** 34054-34089
- Ashutosh Tewari:  
**On the Estimation of Gaussian Mixture Copula Models.** 34090-34104
- Kowshik Thopalli, Rakshit Subramanyam, Pavan K. Turaga, Jayaraman J. Thiagarajan:  
**Target-Aware Generative Augmentations for Single-Shot Adaptation.** 34105-34119
- Qinglong Tian, Xin Zhang, Jiwei Zhao:  
**ELSA: Efficient Label Shift Adaptation through the Lens of Semiparametric Models.** 34120-34142
- Louis C. Tiao, Vincent Dutordoir, Victor Picheny:  
**Spherical Inducing Features for Orthogonally-Decoupled Gaussian Processes.** 34143-34160
- Daniil Tiapkin , Denis Belomestny, Daniele Calandriello, Eric Moulines, Rémi Munos, Alexey Naumov, Pierre Perrault, Yunhao Tang, Michal Valko, Pierre Ménard:  
**Fast Rates for Maximum Entropy Exploration.** 34161-34221
- Alexandru Tifrea, Jacob Clarysse, Fanny Yang:  
**Margin-based sampling in high dimensions: When being active is less efficient than staying passive.** 34222-34262
- Panagiotis Tigas, Yashas Annadani, Desi R. Ivanova, Andrew Jesson, Yarin Gal, Adam Foster, Stefan Bauer:  
**Differentiable Multi-Target Causal Bayesian Experimental Design.** 34263-34279
- Malik Tiomoko, Romain Couillet, Frédéric Pascal:  
**PCA-based Multi-Task Learning: a Random Matrix Approach.** 34280-34300
- Tom Tirer, Haoxiang Huang, Jonathan Niles-Weed:  
**Perturbation Analysis of Neural Collapse.** 34301-34329
- Rishabh Tiwari, Pradeep Shenoy:  
**Overcoming Simplicity Bias in Deep Networks using a Feature Sieve.** 34330-34343
- Christian Tomani, Futa Kai Waseda, Yuesong Shen, Daniel Cremers:  
**Beyond In-Domain Scenarios: Robust Density-Aware Calibration.** 34344-34368



- Peifeng Tong , Wu Su, He Li, Jialin Ding, Zhan Haoxiang, Song Xi Chen:  
**Distribution Free Domain Generalization.** 34369-34378
- Francesco Tonin, Alex Lambert, Panagiotis Patrinos, Johan A. K. Suykens:  
**Extending Kernel PCA through Dualization: Sparsity, Robustness and Fast Algorithms.** 34379-34393
- Francesco Tonolini, Nikolaos Aletras, Yunlong Jiao, Gabriella Kazai:  
**Robust Weak Supervision with Variational Auto-Encoders.** 34394-34408
- Ba-Hien Tran, Babak Shahbaba , Stephan Mandt, Maurizio Filippone:  
**Fully Bayesian Autoencoders with Latent Sparse Gaussian Processes.** 34409-34430
- Frederik Träuble, Anirudh Goyal, Nasim Rahaman, Michael Curtis Mozer, Kenji Kawaguchi, Yoshua Bengio, Bernhard Schölkopf:  
**Discrete Key-Value Bottleneck.** 34431-34455
- Asher Trockman, J. Zico Kolter:  
**Mimetic Initialization of Self-Attention Layers.** 34456-34468
- Che-Ping Tsai, Jiong Zhang, Hsiang-Fu Yu, Eli Chien, Cho-Jui Hsieh, Pradeep Kumar Ravikumar:  
**Representer Point Selection for Explaining Regularized High-dimensional Models.** 34469-34490
- Hanna Tseran, Guido Montúfar:  
**Expected Gradients of Maxout Networks and Consequences to Parameter Initialization.** 34491-34532
- Murad Tukan, Samson Zhou, Alaa Maalouf, Daniela Rus, Vladimir Braverman, Dan Feldman:  
**Provable Data Subset Selection For Efficient Neural Networks Training.** 34533-34555
- Ikechukwu Uchendu, Ted Xiao, Yao Lu, Banghua Zhu, Mengyuan Yan, Joséphine Simon, Matthew Bennice, Chuyuan Fu, Cong Ma, Jiantao Jiao, Sergey Levine, Karol Hausman:  
**Jump-Start Reinforcement Learning.** 34556-34583
- Rajan Udwani:  
**Submodular Order Functions and Assortment Optimization.** 34584-34614
- Masatoshi Uehara, Ayush Sekhari, Jason D. Lee, Nathan Kallus, Wen Sun:  
**Computationally Efficient PAC RL in POMDPs with Latent Determinism and Conditional Embeddings.** 34615-34641
- Enayat Ullah, Raman Arora:  
**From Adaptive Query Release to Machine Unlearning.** 34642-34667
- Enayat Ullah, Christopher A. Choquette-Choo, Peter Kairouz, Sewoong Oh:  
**Private Federated Learning with Autotuned Compression.** 34668-34708
- Théo Uscidda, Marco Cuturi:  
**The Monge Gap: A Regularizer to Learn All Transport Maps.** 34709-34733
- Adrien Vacher, François-Xavier Vialard:  
**Semi-Dual Unbalanced Quadratic Optimal Transport: fast statistical rates and convergent algorithm.** 34734-34758
- Arnaud Vadeboncoeur , Ieva Kazlauskaitė , Yanni Papandreou, Fehmi Cirak, Mark Girolami, Ömer Deniz Akyıldız:  
**Random Grid Neural Processes for Parametric Partial Differential Equations.** 34759-34778
- Sattar Vakili, Danyal Ahmed, Alberto Bernacchia, Ciara Pike-Burke:  
**Delayed Feedback in Kernel Bandits.** 34779-34792
- Boris van Breugel, Zhaozhi Qian, Mihaela van der Schaar:  
**Synthetic Data, Real Errors: How (Not) to Publish and Use Synthetic Data.** 34793-34808
- Dirk van der Hoeven, Ciara Pike-Burke, Hao Qiu, Nicolò Cesa-Bianchi:  
**Trading-Off Payments and Accuracy in Online Classification with Paid Stochastic Experts.** 34809-34830
- Lars van der Laan, Ernesto Ulloa-Pérez, Marco Carone, Alex Luedtke:  
**Causal Isotonic Calibration for Heterogeneous Treatment Effects.** 34831-34854
- Toon Vanderschueren, Alicia Curth, Wouter Verbeke, Mihaela van der Schaar:  
**Accounting For Informative Sampling When Learning to Forecast Treatment Outcomes Over Time.** 34855-34874
- Filippo Vannella, Alexandre Proutière, Jaeseong Jeong:  
**Best Arm Identification in Multi-Agent Multi-Armed Bandits.** 34875-34907
- Harshit Varma, Abhijeet Awasthi, Sunita Sarawagi:  
**Conditional Tree Matching for Inference-Time Adaptation of Tree Prediction Models.** 34908-34923
- Nate Veldt:  
**Optimal LP Rounding and Linear-Time Approximation Algorithms for Clustering Edge-Colored Hypergraphs.** 34924-34951
- Ameya Velingker, Maximilian Vötsch, David P. Woodruff, Samson Zhou:  
**Fast  $(1+\epsilon)$ -Approximation Algorithms for Binary Matrix Factorization.** 34952-34977
- Anirudh Vemula, Yuda Song, Aarti Singh, Drew Bagnell, Sanjiban Choudhury:  
**The Virtues of Laziness in Model-based RL: A Unified Objective and Algorithms.** 34978-35005
- Sara Venturini, Andrea Cristofari, Francesco Rinaldi, Francesco Tudisco:  
**Learning the Right Layers a Data-Driven Layer-Aggregation Strategy for Semi-Supervised Learning on Multilayer Graphs.** 35006-35023



- David Venuto, Sherry Yang, Pieter Abbeel, Doina Precup, Igor Mordatch, Ofir Nachum:  
**Multi-Environment Pretraining Enables Transfer to Action Limited Datasets.** 35024-35036
- Yogesh Verma, Markus Heinonen, Vikas Garg:  
**AbODE: Ab initio antibody design using conjoined ODEs.** 35037-35050
- Mark Vero, Mislav Balunovic, Dimitar Iliev Dimitrov, Martin T. Vechev:  
**TabLeak: Tabular Data Leakage in Federated Learning.** 35051-35083
- Paul Vicol:  
**Low-Variance Gradient Estimation in Unrolled Computation Graphs with ES-Single.** 35084-35119
- Luke Vilnis, Yury Zemlyanskiy, Patrick Murray, Alexandre Tachard Passos, Sumit Sanghai:  
**Arithmetic Sampling: Parallel Diverse Decoding for Large Language Models.** 35120-35136
- Cameron Voloshin, Abhinav Verma , Yisong Yue:  
**Eventual Discounting Temporal Logic Counterfactual Experience Replay.** 35137-35150
- Johannes von Oswald, Eyvind Niklasson, Ettore Randazzo, João Sacramento, Alexander Mordvintsev, Andrey Zhmoginov, Max Vladymyrov:  
**Transformers Learn In-Context by Gradient Descent.** 35151-35174
- Julius von Rohrscheidt, Bastian Rieck:  
**Topological Singularity Detection at Multiple Scales.** 35175-35197
- Václav Voráček, Matthias Hein:  
**Improving  $\ell_1$ -Certified Robustness via Randomized Smoothing by Leveraging Box Constraints.** 35198-35222
- Long Tung Vuong, Trung Le, He Zhao, Chuanxia Zheng, Mehrtash Harandi, Jianfei Cai, Dinh Q. Phung:  
**Vector Quantized Wasserstein Auto-Encoder.** 35223-35242
- Abhijeet Vyas, Brian Bullins, Kamyar Azizzadenesheli:  
**Competitive Gradient Optimization.** 35243-35276
- Nikhil Vyas, Sham M. Kakade, Boaz Barak:  
**On Provable Copyright Protection for Generative Models.** 35277-35299
- Andrew Wagenmaker, Aldo Pacchiano:  
**Leveraging Offline Data in Online Reinforcement Learning.** 35300-35338
- Tal Wagner, Yonatan Naamad, Nina Mishra:  
**Fast Private Kernel Density Estimation via Locality Sensitive Quantization.** 35339-35367
- Jacob C. Walker, Eszter Vértes, Yazhe Li, Gabriel Dulac-Arnold, Ankesh Anand, Theophane Weber, Jessica B. Hamrick:  
**Investigating the Role of Model-Based Learning in Exploration and Transfer.** 35368-35383
- Alvin Wan, Hanxiang Hao, Kaushik Patnaik, Yueyang Xu, Omer Hadad, David Güera, Zhile Ren, Qi Shan:  
**UPSCALE: Unconstrained Channel Pruning.** 35384-35412
- Alexander Wan, Eric Wallace, Sheng Shen, Dan Klein:  
**Poisoning Language Models During Instruction Tuning.** 35413-35425
- Shenghua Wan, Yucen Wang, Minghao Shao, Ruying Chen, De-Chuan Zhan:  
**SeMAIL: Eliminating Distractors in Visual Imitation via Separated Models.** 35426-35443
- Runzhe Wan, Haoyu Wei, Branislav Kveton, Rui Song:  
**Multiplier Bootstrap-based Exploration.** 35444-35490
- Zongqi Wan, Jialin Zhang, Wei Chen, Xiaoming Sun, Zhijie Zhang:  
**Bandit Multi-linear DR-Submodular Maximization and Its Applications on Adversarial Submodular Bandits.** 35491-35524
- Chen Wang:  
**Tight Regret Bounds for Single-pass Streaming Multi-armed Bandits.** 35525-35547
- Yiping Wang, Yifang Chen, Kevin Jamieson, Simon Shaolei Du:  
**Improved Active Multi-Task Representation Learning via Lasso.** 35548-35578
- Yingjie Wang, Hong Chen, Weifeng Liu, Fengxiang He, Teliang Gong, Youcheng Fu, Dacheng Tao:  
**Tilted Sparse Additive Models.** 35579-35604
- Yuxin Wang, Quan Gan, Xipeng Qiu, Xuanjing Huang, David Wipf:  
**From Hypergraph Energy Functions to Hypergraph Neural Networks.** 35605-35623
- Shaoru Wang, Jin Gao, Zeming Li, Xiaoqin Zhang, Weiming Hu:  
**A Closer Look at Self-Supervised Lightweight Vision Transformers.** 35624-35641
- Haibin Wang, Ce Ge , Hesen Chen, Xiuyu Sun:  
**PreNAS: Preferred One-Shot Learning Towards Efficient Neural Architecture Search.** 35642-35654
- Tony Tong Wang, Adam Gleave, Tom Tseng , Kellin Peltine, Nora Belrose, Joseph Miller, Michael D. Dennis, Yawen Duan, Viktor Pogrebniak, Sergey Levine, Stuart Russell:  
**Adversarial Policies Beat Superhuman Go Als.** 35655-35739
- Kaifu Wang, Hangfeng He, Tin D. Nguyen, Piyush Kumar, Dan Roth:  
**On Regularization and Inference with Label Constraints.** 35740-35762



- Qiuhan Wang, Chin Pang Ho, Marek Petrik:  
**Policy Gradient in Robust MDPs with Global Convergence Guarantee.** 35763-35797
- Ziming Wang, Runhao Jiang, Shuang Lian, Rui Yan, Huajin Tang:  
**Adaptive Smoothing Gradient Learning for Spiking Neural Networks.** 35798-35816
- Yansen Wang, Xinyang Jiang, Kan Ren, Caihua Shan, Xufang Luo, Dongqi Han, Kaitao Song, Yifei Shen, Dongsheng Li:  
**CircuitNet: A Generic Neural Network to Realize Universal Circuit Motif Modeling.** 35817-35835
- Xiaoyu Wang, Mikael Johansson, Tong Zhang:  
**Generalized Polyak Step Size for First Order Optimization with Momentum.** 35836-35863
- Kaiwen Wang, Nathan Kallus, Wen Sun:  
**Near-Minimax-Optimal Risk-Sensitive Reinforcement Learning with CVaR.** 35864-35907
- Zhen Wang, Weirui Kuang, Ce Zhang, Bolin Ding, Yaliang Li:  
**FedHPO-Bench: A Benchmark Suite for Federated Hyperparameter Optimization.** 35908-35948
- Jialu Wang, Ping Li, Feifang Hu:  
**A/B Testing in Network Data with Covariate-Adaptive Randomization.** 35949-35969
- Andrew Wang, Andrew C. Li, Toryn Q. Klassen, Rodrigo Toro Icarte, Sheila A. McIlraith:  
**Learning Belief Representations for Partially Observable Deep RL.** 35970-35988
- Hang Wang, Sen Lin, Junshan Zhang:  
**Warm-Start Actor-Critic: From Approximation Error to Sub-optimality Gap.** 35989-36019
- Yanbo Wang, Letao Liu, Justin Dauwels:  
**Slot-VAE: Object-Centric Scene Generation with Slot Attention.** 36020-36035
- Guanchu Wang, Zirui Liu, Zhimeng Jiang, Ninghao Liu, Na Zou, Xia Ben Hu:  
**DIVISION: Memory Efficient Training via Dual Activation Precision.** 36036-36057
- Jue Wang, Yucheng Lu, Binhang Yuan, Beidi Chen, Percy Liang, Christopher De Sa, Christopher Ré, Ce Zhang:  
**CocktailSGD: Fine-tuning Foundation Models over 500Mbps Networks.** 36058-36076
- Hongyu Wang, Shuming Ma, Shaohan Huang, Li Dong, Wenhui Wang, Zhiliang Peng, Yu Wu, Payal Bajaj, Saksham Singhal, Alon Benhaim, Barun Patra, Zhun Liu, Vishrav Chaudhary, Xia Song, Furu Wei:  
**Magneto: A Foundation Transformer.** 36077-36092
- Ruigang Wang, Ian R. Manchester:  
**Direct Parameterization of Lipschitz-Bounded Deep Networks.** 36093-36110
- Ziqiao Wang, Yongyi Mao:  
**Tighter Information-Theoretic Generalization Bounds from Supersamples.** 36111-36137
- Jianfeng Wang, Daniela Massiceti, Xiaolin Hu, Vladimir Pavlovic, Thomas Lukasiewicz:  
**NP-SemiSeg: When Neural Processes meet Semi-Supervised Semantic Segmentation.** 36138-36156
- Tianchun Wang, Farzaneh Mirzazadeh, Xiang Zhang, Jie Chen:  
**GC-Flow: A Graph-Based Flow Network for Effective Clustering.** 36157-36173
- Xin Wang, Zirui Pan, Yuwei Zhou, Hong Chen, Chendi Ge, Wenwu Zhu:  
**Curriculum Co-disentangled Representation Learning across Multiple Environments for Social Recommendation.** 36174-36192
- Peihao Wang, Rameswar Panda, Zhangyang Wang:  
**Data Efficient Neural Scaling Law via Model Reusing.** 36193-36204
- Dingrong Wang, Deep Shankar Pandey, Krishna Prasad Neupane, Zhiwei Yu, Ervine Zheng, Zhi Zheng , Qi Yu:  
**Deep Temporal Sets with Evidential Reinforced Attentions for Unique Behavioral Pattern Discovery.** 36205-36223
- Aoran Wang , Jun Pang:  
**Active Learning based Structural Inference.** 36224-36245
- Zekai Wang, Tianyu Pang, Chao Du, Min Lin, Weiwei Liu, Shuicheng Yan:  
**Better Diffusion Models Further Improve Adversarial Training.** 36246-36263
- Qingyang Wang, Michael Alan Powell, Eric W. Bridgeford, Ali Geisa, Joshua T. Vogelstein:  
**Polarity Is All You Need to Learn and Transfer Faster.** 36264-36284
- Jinxin Wang, Yuen-Man Pun, Xiaolu Wang, Peng Wang , Anthony Man-Cho So:  
**Projected Tensor Power Method for Hypergraph Community Recovery.** 36285-36307
- Tian-Zuo Wang, Tian Qin, Zhi-Hua Zhou:  
**Estimating Possible Causal Effects with Latent Variables via Adjustment.** 36308-36335
- Yingheng Wang, Yair Schiff, Aaron Gokaslan, Weishen Pan, Fei Wang, Christopher De Sa, Volodymyr Kuleshov:  
**InfoDiffusion: Representation Learning Using Information Maximizing Diffusion Models.** 36336-36354
- Jitao Wang, Chengchun Shi, Zhenke Wu:  
**A Robust Test for the Stationarity Assumption in Sequential Decision Making.** 36355-36379
- Hanjing Wang, Man-Kit Sit, Congjie He, Ying Wen, Weinan Zhang, Jun Wang, Yaodong Yang, Luo Mai:  
**GEAR: A GPU-Centric Experience Replay System for Large Reinforcement Learning Models.** 36380-36390
- Aoran Wang , Tsz Pan Tong, Jun Pang:  
**Effective and Efficient Structural Inference with Reservoir Computing.** 36391-36410



- Tongzhou Wang, Antonio Torralba, Phillip Isola, Amy Zhang:  
**Optimal Goal-Reaching Reinforcement Learning via Quasimetric Learning.** 36411-36430
- Yue Wang, Alvaro Velasquez, George K. Atia, Ashley Prater-Bennette, Shaofeng Zou:  
**Model-Free Robust Average-Reward Reinforcement Learning.** 36431-36469
- Xiyao Wang, Wichayaporn Wongkamjan, Ruonan Jia , Furong Huang:  
**Live in the Moment: Learning Dynamics Model Adapted to Evolving Policy.** 36470-36493
- Qian Wang , Zongjun Yang, Xiaotie Deng, Yuqing Kong:  
**Learning to Bid in Repeated First-Price Auctions with Budgets.** 36494-36513
- Xiaolu Wang, Chung-Yiu Yau, Hoi-To Wai:  
**Network Effects in Performative Prediction Games.** 36514-36540
- Puqian Wang, Nikos Zarifis, Ilias Diakonikolas, Jelena Diakonikolas:  
**Robustly Learning a Single Neuron via Sharpness.** 36541-36577
- Zifeng Wang, Zheng Zhan, Yifan Gong, Yucai Shao, Stratis Ioannidis, Yanzhi Wang, Jennifer G. Dy:  
**DualHSIC: HSIC-Bottleneck and Alignment for Continual Learning.** 36578-36592
- Yixuan Wang, Simon Sinong Zhan, Ruochen Jiao, Zhilu Wang, Wanxin Jin, Zhuoran Yang, Zhaoran Wang, Chao Huang, Qi Zhu:  
**Enforcing Hard Constraints with Soft Barriers: Safe Reinforcement Learning in Unknown Stochastic Environments.** 36593-36604
- Runzhong Wang, Yunhao Zhang, Ziao Guo, Tianyi Chen, Xiaokang Yang, Junchi Yan:  
**LinSATNet: The Positive Linear Satisfiability Neural Networks.** 36605-36625
- Jianhao Wang, Jin Zhang, Haozhe Jiang, Junyu Zhang, Liwei Wang, Chongjie Zhang:  
**Offline Meta Reinforcement Learning with In-Distribution Online Adaptation.** 36626-36669
- Kaixin Wang, Kuangqi Zhou, Jiashi Feng, Bryan Hooi, Xinchao Wang:  
**Reachability-Aware Laplacian Representation in Reinforcement Learning.** 36670-36693
- Kaixin Wang, Daquan Zhou, Jiashi Feng, Shie Mannor:  
**PPG Reloaded: An Empirical Study on What Matters in Phasic Policy Gradient.** 36694-36713
- Richard A. Watson, Hengrui Cai, Xinming An, Samuel A. McLean, Rui Song:  
**On Heterogeneous Treatment Effects in Heterogeneous Causal Graphs.** 36714-36747
- Ian Waudby-Smith, Zhiwei Steven Wu, Aaditya Ramdas:  
**Nonparametric Extensions of Randomized Response for Private Confidence Sets.** 36748-36789
- Melanie Weber, Suvrit Sra:  
**Global optimality for Euclidean CCCP under Riemannian convexity.** 36790-36803
- Zixi Wei, Lei Feng, Bo Han, Tongliang Liu, Gang Niu, Xiaofeng Zhu, Heng Tao Shen:  
**A Universal Unbiased Method for Classification from Aggregate Observations.** 36804-36820
- Tianxin Wei, Zeming Guo, Yifan Chen, Jingrui He:  
**NTK-approximating MLP Fusion for Efficient Language Model Fine-tuning.** 36821-36838
- Chunyu Wei, Yu Wang, Bing Bai, Kai Ni, David Brady, Lu Fang:  
**Boosting Graph Contrastive Learning via Graph Contrastive Saliency.** 36839-36855
- Wei Wei, Lijun Zhang, Lin Li, Huizhong Song, Jiye Liang:  
**Set-membership Belief State-based Reinforcement Learning for POMDPs.** 36856-36867
- Hongxin Wei, Huiping Zhuang, Renchunzi Xie, Lei Feng, Gang Niu, Bo An, Yixuan Li:  
**Mitigating Memorization of Noisy Labels by Clipping the Model Prediction.** 36868-36886
- Christian Dietrich Weilbach, William Harvey, Frank Wood:  
**Graphically Structured Diffusion Models.** 36887-36909
- Pascal Welke, Maximilian Thiessen, Fabian Jogl, Thomas Gärtner:  
**Expectation-Complete Graph Representations with Homomorphisms.** 36910-36925
- Jeffrey Wen, Rizwan Ahmad, Philip Schniter:  
**A Conditional Normalizing Flow for Accelerated Multi-Coil MR Imaging.** 36926-36939
- Haitao Wen, Haoyang Cheng, Heqian Qiu, Lanxiao Wang, Lili Pan, Hongliang Li:  
**Optimizing Mode Connectivity for Class Incremental Learning.** 36940-36957
- Yijia Weng, Kaichun Mo, Ruoxi Shi, Yanchao Yang, Leonidas J. Guibas:  
**Towards Learning Geometric Eigen-Lengths Crucial for Fitting Tasks.** 36958-36977
- Zejia Weng, Xitong Yang, Ang Li, Zuxuan Wu, Yu-Gang Jiang:  
**Open-VCLIP: Transforming CLIP to an Open-vocabulary Video Model via Interpolated Weight Optimization.** 36978-36989
- Justin Whitehouse, Aaditya Ramdas, Ryan Rogers, Steven Wu:  
**Fully-Adaptive Composition in Differential Privacy.** 36990-37007
- Jeffrey Willette, Seanie Lee, Bruno Andreis, Kenji Kawaguchi, Juho Lee, Sung Ju Hwang:  
**Scalable Set Encoding with Universal Mini-Batch Consistency and Unbiased Full Set Gradient Approximation.** 37008-37041



- Ezekiel Williams, Colin Bredenberg, Guillaume Lajoie:  
**Flexible Phase Dynamics for Bio-Plausible Contrastive Learning.** 37042-37065
- Daniel J. Williams, Song Liu:  
**Approximate Stein Classes for Truncated Density Estimation.** 37066-37090
- Rick Wilming , Leo Kieslich, Benedict Clark, Stefan Haufe:  
**Theoretical Behavior of XAI Methods in the Presence of Suppressor Variables.** 37091-37107
- David Wipf:  
**Marginalization is not Marginal: No Bad VAE Local Minima when Learning Optimal Sparse Representations.** 37108-37132
- Tom Wollschläger, Nicholas Gao, Bertrand Charpentier, Mohamed Amine Ketata, Stephan Günnemann:  
**Uncertainty Estimation for Molecules: Desiderata and Methods.** 37133-37156
- Jiin Woo, Gauri Joshi, Yuejie Chi:  
**The Blessing of Heterogeneity in Federated Q-Learning: Linear Speedup and Beyond.** 37157-37216
- Gerald Woo, Chenghao Liu, Doyen Sahoo, Akshat Kumar, Steven C. H. Hoi:  
**Learning Deep Time-index Models for Time Series Forecasting.** 37217-37237
- David P. Woodruff, Taisuke Yasuda:  
**Sharper Bounds for  $\ell_p$  Sensitivity Sampling.** 37238-37272
- Blake E. Woodworth, Konstantin Mishchenko, Francis R. Bach:  
**Two Losses Are Better Than One: Faster Optimization Using a Cheaper Proxy.** 37273-37292
- Junran Wu, Xueyuan Chen, Bowen Shi, Shangzhe Li, Ke Xu:  
**SEGA: Structural Entropy Guided Anchor View for Graph Contrastive Learning.** 37293-37312
- Zhengxuan Wu, Karel D'Oosterlinck, Atticus Geiger, Amir Zur, Christopher Potts:  
**Causal Proxy Models for Concept-based Model Explanations.** 37313-37334
- Xiaobao Wu, Xinshuai Dong, Thong Thanh Nguyen, Anh Tuan Luu:  
**Effective Neural Topic Modeling with Embedding Clustering Regularization.** 37335-37357
- Bin Wu , Jinyuan Fang, Xiangxiang Zeng, Shangsong Liang, Qiang Zhang:  
**Adaptive Compositional Continual Meta-Learning.** 37358-37378
- Feijie Wu, Song Guo, Zhihao Qu, Shiqi He, Ziming Liu, Jing Gao:  
**Anchor Sampling for Federated Learning with Partial Client Participation.** 37379-37416
- Haixu Wu, Tengge Hu, Huakun Luo, Jianmin Wang, Mingsheng Long:  
**Solving High-Dimensional PDEs with Latent Spectral Models.** 37417-37438
- Yihan Wu, Heng Huang, Hongyang Zhang:  
**A Law of Robustness beyond Isoperimetry.** 37439-37455
- Tong Wu, Feiran Jia, Xiangyu Qi, Jiachen T. Wang, Vikash Sehwag, Saeed Mahloujifar, Prateek Mittal:  
**Uncovering Adversarial Risks of Test-Time Adaptation.** 37456-37495
- Anpeng Wu, Kun Kuang, Ruoxuan Xiong, Bo Li, Fei Wu:  
**Stable Estimation of Heterogeneous Treatment Effects.** 37496-37510
- Fang Wu, Siyuan Li, Xurui Jin, Yinghui Jiang, Dragomir Radev, Zhangming Niu, Stan Z. Li:  
**Rethinking Explaining Graph Neural Networks via Non-parametric Subgraph Matching.** 37511-37523
- Xiaoxia Wu, Cheng Li, Reza Yazdani Aminabadi, Zhewei Yao, Yuxiong He:  
**Understanding Int4 Quantization for Language Models: Latency Speedup, Composability, and Failure Cases.** 37524-37539
- Guoqiang Wu, Chongxuan Li, Yilong Yin:  
**Towards Understanding Generalization of Macro-AUC in Multi-label Learning.** 37540-37570
- Lirong Wu, Haitao Lin, Yufei Huang, Stan Z. Li:  
**Quantifying the Knowledge in GNNs for Reliable Distillation into MLPs.** 37571-37581
- Xuyang Wu, Changxin Liu, Sindri Magnússon, Mikael Johansson:  
**Delay-agnostic Asynchronous Coordinate Update Algorithm.** 37582-37606
- Philipp Wu, Arjun Majumdar, Kevin Stone, Yixin Lin, Igor Mordatch, Pieter Abbeel, Aravind Rajeswaran:  
**Masked Trajectory Models for Prediction, Representation, and Control.** 37607-37623
- Dongxia Wu, Ruijia Niu, Matteo Chinazzi, Yi-An Ma, Rose Yu:  
**Disentangled Multi-Fidelity Deep Bayesian Active Learning.** 37624-37634
- Hao Wu, Olga Ohremenko, Anthony Wirth:  
**Tight Data Access Bounds for Private Top-k Selection.** 37635-37655
- Lei Wu, Weijie J. Su:  
**The Implicit Regularization of Dynamical Stability in Stochastic Gradient Descent.** 37656-37684
- Runzhe Wu, Masatoshi Uehara, Wen Sun:  
**Distributional Offline Policy Evaluation with Predictive Error Guarantees.** 37685-37712
- Chengyue Wu, Teng Wang, Yixiao Ge, Zeyu Lu, Ruisong Zhou, Ying Shan, Ping Luo:  
 **$\pi$ -Tuning: Transferring Multimodal Foundation Models with Optimal Multi-task Interpolation.** 37713-37727



- Changlong Wu, Yifan Wang, Ananth Grama, Wojciech Szpankowski:  
**Learning Functional Distributions with Private Labels.** 37728-37744
- Wenjie Wu, Ge Yan, Xudong Lu, Kaisen Pan, Junchi Yan:  
**QuantumDARTS: Differentiable Quantum Architecture Search for Variational Quantum Algorithms.** 37745-37764
- Shirley Wu, Mert Yüksekgönül, Linjun Zhang, James Zou:  
**Discover and Cure: Concept-aware Mitigation of Spurious Correlation.** 37765-37786
- David Xing Wu, Chulhee Yun, Suvrit Sra:  
**On the Training Instability of Shuffling SGD with Batch Normalization.** 37787-37845
- Jiawei Wu, Changqing Zhang, Zuoyong Li, Huazhu Fu, Xi Peng, Joey Tianyi Zhou:  
**dugMatting: Decomposed-Uncertainty-Guided Matting.** 37846-37859
- Yue Wu, Shuaicheng Zhang, Wenchao Yu, Yanchi Liu, Quanquan Gu, Dawei Zhou, Haifeng Chen, Wei Cheng:  
**Personalized Federated Learning under Mixture of Distributions.** 37860-37879
- Yulian Wu, Xingyu Zhou, Sayak Ray Chowdhury, Di Wang:  
**Differentially Private Episodic Reinforcement Learning with Heavy-tailed Rewards.** 37880-37918
- Jingfeng Wu, Difan Zou, Zixiang Chen, Vladimir Braverman, Quanquan Gu, Sham M. Kakade:  
**Finite-Sample Analysis of Learning High-Dimensional Single ReLU Neuron.** 37919-37951
- Xun Xian, Ganghua Wang, Jayanth Srinivasa, Ashish Kundu, Xuan Bi, Mingyi Hong, Jie Ding:  
**Understanding Backdoor Attacks through the Adaptability Hypothesis.** 37952-37976
- Ruicheng Xian, Lang Yin, Han Zhao:  
**Fair and Optimal Classification via Post-Processing.** 37977-38012
- Zhen Xiang, Zidi Xiong, Bo Li:  
**UMD: Unsupervised Model Detection for X2X Backdoor Attacks.** 38013-38038
- Jie Xiao, Xueyang Fu, Man Zhou, Hongjian Liu, Zheng-Jun Zha:  
**Random Shuffle Transformer for Image Restoration.** 38039-38058
- Peiyao Xiao, Kaiyi Ji:  
**Communication-Efficient Federated Hypergradient Computation via Aggregated Iterative Differentiation.** 38059-38086
- Guangxuan Xiao, Ji Lin, Mickaël Seznec, Hao Wu, Julien Demouth, Song Han:  
**SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models.** 38087-38099
- Wei Xiao, Tsun-Hsuan Wang, Ramin M. Hasani, Mathias Lechner, Yutong Ban, Chuang Gan, Daniela Rus:  
**On the Forward Invariance of Neural ODEs.** 38100-38124
- Jinqi Xiao , Miao Yin, Yu Gong, Xiao Zang, Jian Ren, Bo Yuan:  
**COMCAT: Towards Efficient Compression and Customization of Attention-Based Vision Models.** 38125-38136
- Pengtao Xie:  
**Improving Bi-level Optimization Based Methods with Inspiration from Humans' Classroom Study Techniques.** 38137-38186
- Zhihui Xie, Zichuan Lin, Deheng Ye, Qiang Fu, Yang Wei, Shuai Li:  
**Future-conditioned Unsupervised Pretraining for Decision Transformer.** 38187-38203
- Liangbin Xie, Xintao Wang, Xiangyu Chen, Gen Li, Ying Shan, Jiantao Zhou, Chao Dong:  
**DeSRA: Detect and Delete the Artifacts of GAN-based Real-World Super-Resolution Models.** 38204-38226
- Chuhan Xie, Wenhao Yang, Zhihua Zhang:  
**Semiparametrically Efficient Off-Policy Evaluation in Linear Markov Decision Processes.** 38227-38257
- Hanchen Xie, Jiageng Zhu, Mahyar Khayatkhoei, Jiazhi Li , Mohamed E. Hussein , Wael AbdAlmageed:  
**A Critical View of Vision-Based Long-Term Dynamics Prediction Under Environment Misalignment.** 38258-38271
- Dong Xing, Pengjie Gu, Qian Zheng, Xinrun Wang, Shanqi Liu, Longtao Zheng, Bo An, Gang Pan:  
**Controlling Type Confounding in Ad Hoc Teamwork with Instance-wise Teammate Feedback Rectification.** 38272-38285
- Zheng Xiong, Jacob Beck, Shimon Whiteson:  
**Universal Morphology Control via Contextual Modulation.** 38286-38300
- Ping Xiong, Thomas Schnake, Michael Gastegger, Grégoire Montavon, Klaus-Robert Müller, Shinichi Nakajima:  
**Relevant Walk Search for Explaining Graph Neural Networks.** 38301-38324
- Frank F. Xu, Uri Alon, Graham Neubig:  
**Why do Nearest Neighbor Language Models Work?** 38325-38341
- Zuheng Xu, Naitong Chen, Trevor Campbell:  
**MixFlows: principled variational inference via mixed flows.** 38342-38376
- Tongda Xu, Han Gao, Chenjian Gao, Yuanyuan Wang, Dailan He, Jinyong Pi, Jixiang Luo, Ziyu Zhu, Mao Ye, Hongwei Qin, Yan Wang, Jingjing Liu, Ya-Qin Zhang:  
**Bit Allocation using Optimization.** 38377-38399
- Wenhao Xu, Xuefeng Gao, Xuedong He:  
**Regret Bounds for Markov Decision Processes with Recursive Optimized Certainty Equivalents.** 38400-38427



- Han Xu, Pengfei He, Jie Ren, Yuxuan Wan, Zitao Liu, Hui Liu, Jiliang Tang:  
**Probabilistic Categorical Adversarial Attack and Adversarial Training.** 38428-38442
- Xiang Xu, Pradeep Kumar Jayaraman, Joseph George Lambourne, Karl D. D. Willis, Yasutaka Furukawa:  
**Hierarchical Neural Coding for Controllable CAD Model Generation.** 38443-38461
- Hainan Xu, Fei Jia, Somshubra Majumdar, He Huang, Shinji Watanabe, Boris Ginsburg:  
**Efficient Sequence Transduction by Jointly Predicting Tokens and Durations.** 38462-38484
- Wenjie Xu, Yuning Jiang, Bratislav Svetozarevic, Colin N. Jones:  
**Constrained Efficient Global Optimization of Expensive Black-box Functions.** 38485-38498
- Mengfan Xu, Diego Klabjan:  
**Pareto Regret Analyses in Multi-objective Multi-armed Bandit.** 38499-38517
- Yan Xu, Deqian Kong, Dehong Xu, Ziwei Ji, Bo Pang, Pascale Fung, Ying Nian Wu:  
**Diverse and Faithful Knowledge-Grounded Dialogue Generation via Sequential Posterior Inference.** 38518-38534
- Jing Xu, Haoxiong Liu:  
**Quantifying the Variability Collapse of Neural Networks.** 38535-38550
- Ning Xu, Biao Liu, Jiaqi Lv, Congyu Qiao, Xin Geng:  
**Progressive Purification for Instance-Dependent Partial Label Learning.** 38551-38565
- Yilun Xu, Ziming Liu, Yonglong Tian, Shangyuan Tong, Max Tegmark, Tommi S. Jaakkola:  
**PFGM++: Unlocking the Potential of Physics-Inspired Generative Models.** 38566-38591
- Minkai Xu, Alexander S. Powers, Ron O. Dror, Stefano Ermon, Jure Leskovec:  
**Geometric Latent Diffusion Models for 3D Molecule Generation.** 38592-38610
- Xingyu Xu, Yandi Shen, Yuejie Chi, Cong Ma:  
**The Power of Preconditioning in Overparameterized Low-Rank Matrix Sensing.** 38611-38654
- Hongzuo Xu, Yijie Wang, Juhui Wei, Songlei Jian, Yizhou Li, Ning Liu:  
**Fascinating Supervisory Signals and Where to Find Them: Deep Anomaly Detection with Scale Learning.** 38655-38673
- Renzhe Xu, Haotian Wang, Xingxuan Zhang, Bo Li, Peng Cui:  
**Competing for Shareable Arms in Multi-Player Multi-Armed Bandits.** 38674-38706
- Chen Xu, Yao Xie:  
**Sequential Predictive Conformal Inference for Time Series.** 38707-38727
- Haiyang Xu, Qinghao Ye, Ming Yan, Yaya Shi, Jiabo Ye, Yuanhong Xu, Chenliang Li, Bin Bi, Qi Qian, Wei Wang, Guohai Xu, Ji Zhang, Songfang Huang, Fei Huang, Jingren Zhou:  
**mPLUG-2: A Modularized Multi-modal Foundation Model Across Text, Image and Video.** 38728-38748
- Minghao Xu, Xinyu Yuan, Santiago Miret, Jian Tang:  
**ProtST: Multi-Modality Learning of Protein Sequences and Biomedical Texts.** 38749-38767
- Yunbei Xu, Assaf Zeevi:  
**Bayesian Design Principles for Frequentist Sequential Learning.** 38768-38800
- Hang Xu, Wenxuan Zhang, Jiawei Fei, Yuzhe Wu, Tingwen Xie, Jun Huang, Yuchen Xie, Mohamed Elhoseiny, Panos Kalnis:  
**SLAMB: Accelerated Large Batch Training with Sparse Communication.** 38801-38825
- Peng Xu, Lin Zhang, Xuanzhou Liu, Jiaqi Sun, Yue Zhao, Haiqin Yang, Bei Yu:  
**Do Not Train It: A Linear Neural Architecture Search of Graph Neural Networks.** 38826-38847
- Yang Xu, Jin Zhu, Chengchun Shi, Shikai Luo, Rui Song:  
**An Instrumental Variable Approach to Confounded Off-Policy Evaluation.** 38848-38880
- Yecheng Xue, Xiaoyu Chen, Tongyang Li, Shaofeng H.-C. Jiang:  
**Near-Optimal Quantum Coreset Construction Algorithms for Clustering.** 38881-38912
- Fuzhao Xue, Jianghai Chen, Aixin Sun, Xiaozhe Ren, Zangwei Zheng, Xiaoxin He, Yongming Chen, Xin Jiang, Yang You:  
**A Study on Transformer Configuration and Training Objective.** 38913-38925
- Rui Xue, Haoyu Han, MohamadAli Torkamani, Jian Pei, Xiaorui Liu:  
**LazyGNN: Large-Scale Graph Neural Networks via Lazy Propagation.** 38926-38937
- Yihao Xue, Siddharth Joshi, Eric Gan, Pin-Yu Chen, Baharan Mirzasoleiman:  
**Which Features are Learnt by Contrastive Learning? On the Role of Simplicity Bias in Class Collapse and Feature Suppression.** 38938-38970
- Fuzhao Xue, Valerii Likhoshevstov, Anurag Arnab, Neil Houlsby, Mostafa Dehghani, Yang You:  
**Adaptive Computation with Elastic Input Sequence.** 38971-38988
- Taku Yamagata, Ahmed Khalil, Raúl Santos-Rodríguez:  
**Q-learning Decision Transformer: Leveraging Dynamic Programming for Conditional Sequence Modelling in Offline RL.** 38989-39007
- Hayata Yamasaki, Sathyawageeswar Subramanian, Satoshi Hayakawa, Sho Sonoda:  
**Quantum Ridgelet Transform: Winning Lottery Ticket of Neural Networks with Quantum Computation.** 39008-39034
- Yonggui Yan, Jie Chen, Pin-Yu Chen, Xiaodong Cui, Songtao Lu, Yangyang Xu:  
**Compressed Decentralized Proximal Stochastic Gradient Method for Nonconvex Composite Problems with**

**Heterogeneous Data.** 39035-39061

Wilson Yan, Danijar Hafner, Stephen James, Pieter Abbeel:  
**Temporally Consistent Transformers for Video Generation.** 39062-39098

Zhiqiang Yan, Xiang Li, Kun Wang, Shuo Chen, Jun Li, Jian Yang:  
**Distortion and Uncertainty Aware Loss for Panoramic Depth Completion.** 39099-39109

Jingquan Yan, Hao Wang:  
**Self-Interpretable Time Series Prediction with Counterfactual Explanations.** 39110-39125

Ge Yan, Huaijin Wu, Junchi Yan:  
**Quantum 3D Graph Learning with Applications to Molecule Embedding.** 39126-39137

Yu-Hu Yan, Peng Zhao, Zhi-Hua Zhou:  
**Fast Rates in Time-Varying Strongly Monotone Games.** 39138-39164

Hiroki Yanagisawa:  
**Proper Scoring Rules for Survival Analysis.** 39165-39182

Rushuai Yang, Chenjia Bai, Hongyi Guo, Siyuan Li, Bin Zhao, Zhen Wang, Peng Liu, Xuelong Li:  
**Behavior Contrastive Learning for Unsupervised Skill Discovery.** 39183-39204

Junwen Yang, Yifan Feng:  
**Nested Elimination: A Simple Algorithm for Best-Item Identification From Choice-Based Feedback.** 39205-39233

Mingqi Yang, Wenjie Feng, Yanming Shen, Bryan Hooi:  
**Towards Better Graph Representation Learning with Parameterized Decomposition & Filtering.** 39234-39251

Shenghao Yang, Kimon Fountoulakis:  
**Weighted Flow Diffusion for Local Graph Clustering with Node Attributes: an Algorithm and Statistical Guarantees.** 39252-39276

Soojung Yang, Rafael Gómez-Bombarelli:  
**Chemically Transferable Generative Backmapping of Coarse-Grained Proteins.** 39277-39298

Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, Yang Zhang:  
**Data Poisoning Attacks Against Multimodal Encoders.** 39299-39313

Yu Yang, Hao Kang, Baharan Mirzasoleiman:  
**Towards Sustainable Learning: Coresets for Data-efficient Deep Learning.** 39314-39330

Dongyoon Yang, Insung Kong, Yongdai Kim:  
**Improving Adversarial Robustness by Putting More Regularizations on Less Robust Samples.** 39331-39348

Zonghan Yang, Peng Li, Tianyu Pang, Yang Liu:  
**Improving Adversarial Robustness of Deep Equilibrium Models with Explicit Regulations Along the Neural Dynamics.** 39349-39364

Yu Yang, Besmira Nushi, Hamid Palangi, Baharan Mirzasoleiman:  
**Mitigating Spurious Correlations in Multi-modal Models during Fine-tuning.** 39365-39379

Adam X. Yang, Maxime Robeyns, Edward Milsom, Ben Anson, Nandi Schoots, Laurence Aitchison:  
**A theory of representation learning gives a deep generalisation of kernel methods.** 39380-39415

Joonhyuk Yang, Dongpil Shin, Hye Won Chung:  
**Efficient Algorithms for Exact Graph Matching on Correlated Stochastic Block Models with Constant Correlation.** 39416-39452

Yongyi Yang, Jacob Steinhardt, Wei Hu:  
**Are Neurons Actually Collapsed? On the Fine-Grained Structure in Neural Representations.** 39453-39487

Jianke Yang, Robin Walters, Nima Dehmamy, Rose Yu:  
**Generative Adversarial Symmetry Discovery.** 39488-39508

Qisen Yang, Shenzhi Wang, Matthieu Gaetan Lin, Shiji Song, Gao Huang:  
**Boosting Offline Reinforcement Learning with Action Preference Query.** 39509-39523

Jianan Yang , Haobo Wang, Sai Wu, Gang Chen, Junbo Zhao:  
**Towards Controlled Data Augmentations for Active Learning.** 39524-39542

Rui Yang, Lin Yong, Xiaoteng Ma, Hao Hu, Chongjie Zhang, Tong Zhang:  
**What is Essential for Unseen Goal Generalization of Offline Goal-conditioned RL?** 39543-39571

Lingxiao Yang, Hongzhi You, Zonglei Zhen, Dahui Wang, Xiaohong Wan, Xiaohua Xie, Ru-Yuan Zhang:  
**Neural Prediction Errors enable Analogical Visual Reasoning in Human Standard Intelligence Tests.** 39572-39583

Yuzhe Yang, Haoran Zhang, Dina Katabi, Marzyeh Ghassemi:  
**Change is Hard: A Closer Look at Subpopulation Shift.** 39584-39622

Yijun Yang, Tianyi Zhou, Jing Jiang, Guodong Long, Yuhui Shi:  
**Continual Task Allocation in Meta-Policy Network via Sparse Prompting.** 39623-39638

Menglin Yang, Min Zhou, Rex Ying, Yankai Chen, Irwin King:  
**Hyperbolic Representation Learning: Revisiting and Advancing.** 39639-39659

Yu Yao, Mingming Gong, Yuxuan Du, Jun Yu, Bo Han, Kun Zhang, Tongliang Liu:  
**Which is Better for Learning with Noisy Labels: The Semi-supervised Method or Modeling Label Noise?** 39660-



39673

- Fan Yao, Chuanhao Li, Denis Nekipelov, Hongning Wang, Haifeng Xu:  
**How Bad is Top-K Recommendation under Competing Content Creators?** 39674-39701
- Jiachen Yao, Chang Su, Zhongkai Hao, Songming Liu, Hang Su, Jun Zhu:  
**MultiAdam: Parameter-wise Scale-invariant Optimizer for Multiscale Training of Physics-informed Neural Networks.** 39702-39721
- Batuhan Yardim, Semih Cayci, Matthieu Geist, Niao He:  
**Policy Mirror Ascent for Efficient and Independent Learning in Mean Field Games.** 39722-39754
- Michihiro Yasunaga, Armen Aghajanyan, Weijia Shi, Richard James, Jure Leskovec, Percy Liang, Mike Lewis, Luke Zettlemoyer, Wen-Tau Yih:  
**Retrieval-Augmented Multimodal Language Modeling.** 39755-39769
- Haotian Ye, Xiaoyu Chen, Liwei Wang, Simon Shaolei Du:  
**On the Power of Pre-training for Generalization in RL: Provable Benefits and Hardness.** 39770-39800
- Rui Ye, Zhenyang Ni, Fangzhao Wu, Siheng Chen, Yanfeng Wang:  
**Personalized Federated Learning with Inferred Collaboration Graphs.** 39801-39817
- Jiacheng Ye, Zhiyong Wu, Jiangtao Feng, Tao Yu, Lingpeng Kong:  
**Compositional Exemplars for In-context Learning.** 39818-39833
- Chenlu Ye, Wei Xiong, Quanquan Gu, Tong Zhang:  
**Corruption-Robust Algorithms with Uncertainty Weighting for Nonlinear Contextual Bandits and Markov Decision Processes.** 39834-39863
- Huigen Ye, Hua Xu, Hongyan Wang, Chengming Wang, Yu Jiang:  
**GNN&GBDT-Guided Fast Optimizing Framework for Large-scale Integer Programming.** 39864-39878
- Rui Ye, Mingkai Xu, Jianyu Wang, Chenxin Xu, Siheng Chen, Yanfeng Wang:  
**FedDisco: Federated Learning with Discrepancy-Aware Collaboration.** 39879-39902
- Xinyu Ye, Ge Yan, Junchi Yan:  
**Towards Quantum Machine Learning for Constrained Combinatorial Optimization: a Quantum QAP Solver.** 39903-39912
- Hugo Yèche, Alizée Pace, Gunnar Rätsch, Rita Kuznetsova:  
**Temporal Label Smoothing for Early Event Prediction.** 39913-39938
- Raanan Y. Yehezkel Rohekar, Shami Nisimov, Yaniv Gurwicz, Gal Novik:  
**From Temporal to Contemporaneous Iterative Causal Discovery in the Presence of Latent Confounders.** 39939-39950
- Jialin Yi, Milan Vojnovic:  
**Doubly Adversarial Federated Bandits.** 39951-39967
- Qi Yi, Rui Zhang, Shaohui Peng, Jiaming Guo, Yunkai Gao, Kaizhao Yuan, Ruizhi Chen, Siming Lan, Xing Hu, Zidong Du, Xishan Zhang, Qi Guo, Yunji Chen:  
**Online Prototype Alignment for Few-shot Policy Transfer.** 39968-39983
- Mingxuan Yi, Zhanxing Zhu, Song Liu:  
**MonoFlow: Rethinking Divergence GANs via the Perspective of Wasserstein Gradient Flows.** 39984-40000
- Jason Yim, Brian L. Trippe, Valentin De Bortoli, Emile Mathieu, Arnaud Doucet, Regina Barzilay, Tommi S. Jaakkola:  
**SE(3) diffusion model with application to protein backbone generation.** 40001-40039
- Nan Yin, Li Shen, Mengzhu Wang, Long Lan, Zeyu Ma, Chong Chen, Xian-Sheng Hua, Xiao Luo:  
**CoCo: A Coupled Contrastive Framework for Unsupervised Domain Adaptive Graph Classification.** 40040-40053
- Jiaxi Ying, José Vinícius de Miranda Cardoso, Daniel P. Palomar:  
**Adaptive Estimation of Graphical Models under Total Positivity.** 40054-40074
- Seungryong Yoo, Eunji Kim, Dahuin Jung, Jungbeom Lee, Sungroh Yoon:  
**Improving Visual Prompt Tuning for Self-supervised Vision Transformers.** 40075-40092
- Jaeyoung Yoo, Hojun Lee, Seunghyeon Seo, Inseop Chung, Nojun Kwak:  
**End-to-End Multi-Object Detection with a Regularized Mixture Model.** 40093-40110
- Ji Won Yoon, Sunghwan Ahn, Hyeonseung Lee, Minchan Kim, Seok Min Kim, Nam Soo Kim:  
**EM-Network: Oracle Guided Self-distillation for Sequence Learning.** 40111-40128
- Jaehong Yoon, Sung Ju Hwang, Yue Cao:  
**Continual Learners are Incremental Model Generalizers.** 40129-40146
- Jaesik Yoon, Yi-Fu Wu, Heechul Bae, Sungjin Ahn:  
**An Investigation into Pre-Training Object-Centric Representations for Reinforcement Learning.** 40147-40174
- Minji Yoon, Yue Wu, John Palowitch, Bryan Perozzi, Russ Salakhutdinov:  
**Graph Generative Model for Benchmarking Graph Neural Networks.** 40175-40198
- Xuchen You, Shouvanik Chakrabarti, Boyang Chen, Xiaodi Wu:  
**Analyzing Convergence in Quantum Neural Networks: Deviations from Neural Tangent Kernels.** 40199-40224
- Ali Younes, Simone Schaub-Meyer, Georgia Chalvatzaki:  
**Entropy-driven Unsupervised Keypoint Representation Learning in Videos.** 40225-40253



- Kenny John Young, Aditya A. Ramesh, Louis Kirsch, Jürgen Schmidhuber:  
**The Benefits of Model-Based Generalization in Reinforcement Learning.** 40254-40276
- Anlan Yu, Ning Lyu, Jieming Yin, Zhiyuan Yan, Wujie Wen:  
**COLA: Orchestrating Error Coding and Learning for Robust Neural Network Inference Against Hardware Defects.** 40277-40289
- Chenglin Yu, Xinsong Ma, Weiwei Liu:  
**Delving into Noisy Label Detection with Clean Data.** 40290-40305
- Weichen Yu, Tianyu Pang, Qian Liu, Chao Du, Bingyi Kang, Yan Huang, Min Lin, Shuicheng Yan:  
**Bag of Tricks for Training Data Extraction from Language Models.** 40306-40320
- Dayou Yu, Weishi Shi, Qi Yu:  
**Discover-Then-Rank Unlabeled Support Vectors in the Dual Space for Multi-Class Active Learning.** 40321-40338
- Jiashuo Yu, Yaohui Wang, Xinyuan Chen, Xiao Sun, Yu Qiao:  
**Long-Term Rhythmic Video Soundtracker.** 40339-40353
- Lijia Yu, Yihan Wang, Xiao-Shan Gao:  
**Adversarial Parameter Attack on Deep Neural Networks.** 40354-40372
- Zhiyuan Yu , Yuhao Wu, Ning Zhang, Chenguang Wang, Yevgeniy Vorobeychik, Chaowei Xiao:  
**CodeIPPrompt: Intellectual Property Infringement Assessment of Code Language Models.** 40373-40389
- Liren Yu, Jiaming Xu, Xiaojun Lin:  
**SeedGNN: Graph Neural Network for Supervised Seeded Graph Matching.** 40390-40411
- Haiyang Yu, Zhao Xu, Xiaofeng Qian , Xiaoning Qian, Shuiwang Ji:  
**Efficient and Equivariant Graph Networks for Predicting Quantum Hamiltonian.** 40412-40424
- Xian Yu, Lei Ying:  
**On the Global Convergence of Risk-Averse Policy Gradient Methods with Expected Conditional Risk Measures.** 40425-40451
- Zishun Yu, Xinhua Zhang:  
**Actor-Critic Alignment for Offline-to-Online Reinforcement Learning.** 40452-40474
- Zhongzhi Yu, Yang Zhang, Kaizhi Qian, Cheng Wan, Yonggan Fu, Yongan Zhang, Yingyan Celine Lin:  
**Master-ASR: Achieving Multilingual Scalability and Low-Resource Adaptation in ASR with Modular Learning.** 40475-40487
- Ganzhao Yuan:  
**Coordinate Descent Methods for Fractional Minimization.** 40488-40518
- Yang Yuan:  
**On the Power of Foundation Models.** 40519-40530
- Mingqi Yuan, Bo Li, Xin Jin, Wenjun Zeng:  
**Automatic Intrinsic Reward Shaping for Exploration in Deep Reinforcement Learning.** 40531-40554
- Eunggu Yun, Hyungi Lee, Giung Nam, Juho Lee:  
**Traversing Between Modes in Function Space for Fast Ensembling.** 40555-40577
- Margaux Zaffran, Aymeric Dieuleveut, Julie Josse, Yaniv Romano:  
**Conformal Prediction with Missing Values.** 40578-40604
- Amir Zandieh, Insu Han, Majid Daliri, Amin Karbasi:  
**KDEformer: Accelerating Transformers via Kernel Density Estimation.** 40605-40623
- Santiago Zanella Béguelin, Lukas Wutschitz, Shruti Tople, Ahmed Salem, Victor Rühle, Andrew Paverd, Mohammad Naseri, Boris Köpf, Daniel Jones:  
**Bayesian Estimation of Differential Privacy.** 40624-40636
- Andrea Zanette:  
**When is Realizability Sufficient for Off-Policy Reinforcement Learning?** 40637-40668
- Sepanta Zeighami, Cyrus Shahabi:  
**On Distribution Dependent Sub-Logarithmic Query Time of Learned Indexing.** 40669-40680
- Houssam Zenati, Eustache Diemert, Matthieu Martin, Julien Mairal, Pierre Gaillard:  
**Sequential Counterfactual Risk Minimization.** 40681-40706
- Zhanpeng Zeng, Michael Davies, Pranav Pulijala, Karthikeyan Sankaralingam, Vikas Singh:  
**LookupFFN: Making Transformers Compute-lite for CPU inference.** 40707-40718
- Shiwei Zeng, Jie Shen:  
**Attribute-Efficient PAC Learning of Low-Degree Polynomial Threshold Functions with Nasty Noise.** 40719-40748
- Zhichen Zeng, Ruike Zhu, Yinglong Xia, Hanqing Zeng, Hanghang Tong:  
**Generative Graph Dictionary Learning.** 40749-40769
- Shuangfei Zhai, Tatiana Likhomanenko, Eta Littwin, Dan Busbridge, Jason Ramapuram, Yizhe Zhang, Jiatao Gu, Joshua M. Susskind:  
**Stabilizing Transformer Training by Preventing Attention Entropy Collapse.** 40770-40803
- Yuheng Zhang, Yu Bai, Nan Jiang:  
**Offline Learning in Markov Games with General Function Approximation.** 40804-40829



- Jianyu Zhang, Léon Bottou:  
**Learning useful representations for shifting tasks and distributions.** 40830-40850
- Chen Zhang, Xiaofeng Cao, Weiyang Liu, Ivor W. Tsang , James T. Kwok:  
**Nonparametric Iterative Machine Teaching.** 40851-40870
- Cindy Y. Zhang, Sarah Huiyi Cen, Devavrat Shah:  
**Matrix Estimation for Individual Fairness.** 40871-40887
- Hangfan Zhang, Jinghui Chen, Lu Lin, Jinyuan Jia, Dinghao Wu:  
**Graph Contrastive Backdoor Attacks.** 40888-40910
- Zixuan Zhang, Minshuo Chen, Mengdi Wang, Wenjing Liao, Tuo Zhao:  
**Effective Minkowski Dimension of Deep Nonparametric Regression: Function Approximation and Statistical Theories.** 40911-40931
- Honghua Zhang, Meihua Dang, Nanyun Peng, Guy Van den Broeck:  
**Tractable Control for Autoregressive Language Generation.** 40932-40945
- Yuhua Zhang, Walter H. Dempsey:  
**CataBEM: Integrating Latent Interaction Categories in Node-wise Community Detection Models for Network Data.** 40946-40975
- Jiuling Zhang, Zhiming Ding:  
**Rethink DARTS Search Space and Renovate a New Benchmark.** 40976-40995
- Brian Hu Zhang, Gabriele Farina, Tuomas Sandholm:  
**Team Belief DAG: Generalizing the Sequence Form to Team Games for Fast Computation of Correlated Team Max-Min Equilibria via Regret Minimization.** 40996-41018
- Bohang Zhang, Guhao Feng, Yiheng Du, Di He, Liwei Wang:  
**A Complete Expressiveness Hierarchy for Subgraph GNNs via Subgraph Weisfeiler-Lehman Tests.** 41019-41077
- Ruofan Zhang, Jinjin Gu, Haoyu Chen, Chao Dong, Yulun Zhang, Wenming Yang:  
**Crafting Training Degradation Distribution for the Accuracy-Generalization Trade-off in Real-World Super-Resolution.** 41078-41091
- Biao Zhang, Barry Haddow, Alexandra Birch:  
**Prompting Large Language Model for Machine Translation: A Case Study.** 41092-41110
- Weitong Zhang, Jiafan He, Zhiyuan Fan, Quanquan Gu:  
**On the Interplay Between Misspecification and Sub-optimality Gap in Linear Contextual Bandits.** 41111-41132
- Chunhui Zhang, Chao Huang, Yijun Tian, Qianlong Wen, Zhongyu Ouyang, Youhuan Li, Yanfang Ye, Chuxu Zhang:  
**When Sparsity Meets Contrastive Models: Less Graph Data Can Bring Better Class-Balanced Representations.** 41133-41150
- Qianru Zhang, Chao Huang, Lianghao Xia, Zheng Wang, Siu Ming Yiu, Ruihua Han:  
**Spatial-Temporal Graph Learning with Adversarial Contrastive Adaptation.** 41151-41163
- Guanhua Zhang, Jiabao Ji, Yang Zhang, Mo Yu, Tommi S. Jaakkola, Shiyu Chang:  
**Towards Coherent Image Inpainting Using Denoising Diffusion Implicit Models.** 41164-41193
- Jun Zhang, Shuyang Jiang, Jianguo Feng, Lin Zheng, Lingpeng Kong:  
**CAB: Comprehensive Attention Benchmarking on Long Sequence Modeling.** 41194-41218
- Shenao Zhang, Wanxin Jin, Zhaoran Wang:  
**Adaptive Barrier Smoothing for First-Order Policy Gradient with Contact Dynamics.** 41219-41243
- Hang Zhang, Ping Li:  
**One-Step Estimator for Permuted Sparse Recovery.** 41244-41267
- Chenyi Zhang, Tongyang Li:  
**Quantum Lower Bounds for Finding Stationary Points of Nonconvex Functions.** 41268-41299
- Xinlu Zhang, Shiyang Li, Zhiyu Chen, Xifeng Yan, Linda Ruth Petzold:  
**Improving Medical Predictions by Irregular Multimodal Electronic Health Records Modeling.** 41300-41313
- Hao Zhang, Chenglin Li, Wenrui Dai, Junni Zou, Hongkai Xiong:  
**FedCR: Personalized Federated Learning Based on Across-Client Common Representation with Conditional Mutual Information Regularization.** 41314-41330
- Haobo Zhang, Yicheng Li, Weihao Lu, Qian Lin:  
**On the Optimality of Misspecified Kernel Ridge Regression.** 41331-41353
- Jianyi Zhang, Ang Li, Minxue Tang, Jingwei Sun, Xiang Chen, Fan Zhang, Changyou Chen, Yiran Chen, Hai Li:  
**Fed-CBS: A Heterogeneity-Aware Client Sampling Mechanism for Federated Learning via Class-Imbalance Reduction.** 41354-41381
- Zaixi Zhang, Qi Liu:  
**Learning Subpocket Prototypes for Generalizable Structure-based Drug Design.** 41382-41398
- Feilong Zhang , Xianming Liu, Shiyi Lin, Gang Wu, Xiong Zhou, Junjun Jiang, Xiangyang Ji:  
**No One Idles: Efficient Heterogeneous Federated Learning with Parallel Edge and Server Computation.** 41399-41413



- Tianjun Zhang, Fangchen Liu, Justin Wong, Pieter Abbeel, Joseph E. Gonzalez:  
**The Wisdom of Hindsight Makes Language Models Better Instruction Followers.** 41414-41428
- Shuhai Zhang, Feng Liu, Jiahao Yang, Yifan Yang, Changsheng Li, Bo Han, Mingkui Tan:  
**Detecting Adversarial Data by Probing Multiple Perturbations Using Expected Perturbation Score.** 41429-41451
- Shijun Zhang, Jianfeng Lu, Hongkai Zhao:  
**On Enhancing Expressive Power via Compositions of Single Fixed-Size ReLU Network.** 41452-41487
- Yuxin Zhang, Yiting Luo, Mingbao Lin, Yunshan Zhong, Jingjing Xie, Fei Chao, Rongrong Ji:  
**Bi-directional Masks for Efficient N: M Sparse Training.** 41488-41497
- Jie Zhang, Xiaosong Ma, Song Guo, Wenchao Xu:  
**Towards Unbiased Training in Federated Open-world Semi-supervised Learning.** 41498-41509
- Shengping Zhang, Quanling Meng, Qinglin Liu, Liqiang Nie, Bineng Zhong, Xiaopeng Fan, Rongrong Ji:  
**Interactive Object Placement with Reinforcement Learning.** 41510-41522
- Fangzhao Zhang, Mert Pilanci:  
**Optimal Shrinkage for Distributed Second-Order Optimization.** 41523-41549
- Haoran Zhang, Harvineet Singh, Marzyeh Ghassemi, Shalmali Joshi:  
**"Why did the Model Fail?": Attributing Model Performance Changes to Distribution Shifts.** 41550-41578
- Fengxue Zhang, Jialin Song, James C. Bowden, Alexander Ladd, Yisong Yue, Thomas Desautels, Yuxin Chen:  
**Learning Regions of Interest for Bayesian Optimization with Adaptive Level-Set Estimation.** 41579-41595
- Yivan Zhang, Masashi Sugiyama:  
**A Category-theoretical Meta-analysis of Definitions of Disentanglement.** 41596-41612
- Shangtong Zhang, Remi Tachet des Combes, Romain Laroche:  
**On the Convergence of SARSA with Linear Function Approximation.** 41613-41646
- Yifan Zhang, Xue Wang, Kexin Jin, Kun Yuan, Zhang Zhang, Liang Wang, Rong Jin, Tieniu Tan:  
**AdaNPC: Exploring Non-Parametric Classifier for Test-Time Adaptation.** 41647-41676
- Qi Zhang, Yifei Wang, Yisen Wang:  
**On the Generalization of Multi-modal Contrastive Learning.** 41677-41693
- Wang Zhang, Tsui-Wei Weng, Subhro Das, Alexandre Megretski, Luca Daniel, Lam M. Nguyen:  
**ConCerNet: A Contrastive Learning Based Framework for Automated Conservation Law Discovery and Trustworthy Dynamical System Prediction.** 41694-41714
- Wenbo Zhang, Tong Wu, Yunlong Wang, Yong Cai, Hengrui Cai:  
**Towards Trustworthy Explanation: On Causal Rationalization.** 41715-41736
- He Zhang, Bang Wu, Shuo Wang, Xiangwen Yang, Minhui Xue, Shirui Pan, Xingliang Yuan:  
**Demystifying Uneven Vulnerability of Link Stealing Attacks against Graph Neural Networks.** 41737-41752
- Qingyang Zhang, Haitao Wu, Changqing Zhang, Qinghua Hu, Huazhu Fu, Joey Tianyi Zhou, Xi Peng:  
**Provable Dynamic Fusion for Low-Quality Multimodal Data.** 41753-41769
- Kexun Zhang, Xianjun Yang, William Yang Wang, Lei Li:  
**ReDi: Efficient Learning-Free Diffusion Inference via Trajectory Retrieval.** 41770-41785
- Qixin Zhang, Wenbing Ye, Zaiyi Chen, Haoyuan Hu, Enhong Chen, Yu Yang:  
**Nearly Optimal Competitive Ratio for Online Allocation Problems with Two-sided Resource Constraints and Finite Requests.** 41786-41818
- Xiaohui Zhang, Jiangyan Yi, Jianhua Tao, Chenglong Wang, Chu Yuan Zhang:  
**Do You Remember? Overcoming Catastrophic Forgetting for Fake Audio Detection.** 41819-41831
- Tianyi Zhang, Tao Yu, Tatsunori Hashimoto, Mike Lewis, Wen-Tau Yih, Daniel Fried, Sida Wang:  
**Coder Reviewer Reranking for Code Generation.** 41832-41846
- Wanrong Zhang, Ruqi Zhang:  
**DP-Fast MH: Private, Fast, and Accurate Metropolis-Hastings for Large-Scale Bayesian Inference.** 41847-41860
- Yifan Zhang, Min-Ling Zhang:  
**Nearly-tight Bounds for Deep Kernel Learning.** 41861-41879
- Tianping Zhang, Zheyu Aqa Zhang, Zhiyuan Fan, Haoyan Luo, Fengyuan Liu, Qian Liu, Wei Cao, Li Jian:  
**OpenFE: Automated Feature Generation with Expert-level Performance.** 41880-41901
- Junkai Zhang, Weitong Zhang, Quanquan Gu:  
**Optimal Horizon-Free Reward-Free Exploration for Linear Mixture MDPs.** 41902-41930
- Yan Zhang, David W. Zhang, Simon Lacoste-Julien, Gertjan J. Burghouts, Cees G. M. Snoek:  
**Unlocking Slot Attention by Changing Optimal Transport Costs.** 41931-41951
- Hang Zhang, Kaifeng Zhang, Kai Ming Ting, Ye Zhu:  
**Towards a Persistence Diagram that is Robust to Noise and Varied Densities.** 41952-41972
- Jinpeng Zhang, Yufeng Zheng, Chuheng Zhang, Li Zhao, Lei Song, Yuan Zhou, Jiang Bian:  
**Robust Situational Reinforcement Learning in Face of Context Disturbances.** 41973-41989
- Shaofeng Zhang, Qiang Zhou, Zhibin Wang, Fan Wang, Junchi Yan:  
**Patch-level Contrastive Learning via Positional Query for Visual Pre-training.** 41990-41999



- Dora Zhao, Jerome Theodore Alexander Andrews, Alice Xiang:  
**Men Also Do Laundry: Multi-Attribute Bias Amplification.** 42000-42017
- Xunyi Zhao, Théotime Le Hellard, Lionel Eyraud-Dubois, Julia Gusak, Olivier Beaumont:  
**Rockmate: an Efficient, Fast, Automatic and Generic Tool for Re-materialization in PyTorch.** 42018-42045
- Yixiu Zhao, Scott W. Linderman:  
**Revisiting Structured Variational Autoencoders.** 42046-42057
- Hao Zhao, Yuejiang Liu, Alexandre Alahi, Tao Lin:  
**On Pitfalls of Test-Time Adaptation.** 42058-42080
- Boxin Zhao, Boxiang Lyu, Raul Castro Fernandez, Mladen Kolar:  
**Addressing Budget Allocation and Revenue Allocation in Data Market Environments Using an Adaptive Sampling Algorithm.** 42081-42097
- Hanqing Zhao, Dianmo Sheng, Jianmin Bao, Dongdong Chen, Dong Chen, Fang Wen, Lu Yuan, Ce Liu, Wenbo Zhou, Qi Chu, Weiming Zhang, Nenghai Yu:  
**X-Paste: Revisiting Scalable Copy-Paste for Instance Segmentation using CLIP and StableDiffusion.** 42098-42109
- Yao Zhao, Connor Stephens, Csaba Szepesvári, Kwang-Sung Jun:  
**Revisiting Simple Regret: Fast Rates for Returning a Good Arm.** 42110-42158
- He Zhao, Ke Sun, Amir Dezfouli, Edwin V. Bonilla:  
**Transformed Distribution Matching for Missing Value Imputation.** 42159-42186
- Xuandong Zhao, Yu-Xiang Wang, Lei Li:  
**Protecting Language Generation Models via Invisible Watermarking.** 42187-42199
- Yulai Zhao, Zhuoran Yang, Zhaoran Wang, Jason D. Lee:  
**Local Optimization Achieves Global Optimality in Multi-Agent Reinforcement Learning.** 42200-42226
- Yi Zhao, Wenshuai Zhao, Rinu Boney, Juho Kannala, Joni Pajarinen:  
**Simplified Temporal Consistency Reinforcement Learning.** 42227-42246
- Liming Zhao, Kecheng Zheng, Yun Zheng, Deli Zhao, Jingren Zhou:  
**RLEG: Vision-Language Representation Learning with Diffusion-based Embedding Generation.** 42247-42258
- Heyang Zhao, Dongruo Zhou, Jiafan He, Quanquan Gu:  
**Optimal Online Generalized Linear Regression with Stochastic Noise and Its Application to Heteroscedastic Bandits.** 42259-42279
- Haiyan Zhao, Tianyi Zhou, Guodong Long, Jing Jiang, Chengqi Zhang:  
**Does Continual Learning Equally Forget All Parameters?** 42280-42303
- Geng Zhao, Banghua Zhu, Jiantao Jiao, Michael I. Jordan:  
**Online Learning in Stackelberg Games with an Omnipotent Follower.** 42304-42316
- Zaixiang Zheng, Yifan Deng, Dongyu Xue, Yi Zhou, Fei Ye, Quanquan Gu:  
**Structure-informed Language Models Are Protein Designers.** 42317-42338
- Qinqing Zheng, Mikael Henaff, Brandon Amos, Aditya Grover:  
**Semi-Supervised Offline Reinforcement Learning with Action-Free Trajectories.** 42339-42362
- Kaiwen Zheng, Cheng Lu, Jianfei Chen, Jun Zhu:  
**Improved Techniques for Maximum Likelihood Estimation for Diffusion ODEs.** 42363-42389
- Hongkai Zheng, Weili Nie, Arash Vahdat, Kamyar Azizzadenesheli, Anima Anandkumar:  
**Fast Sampling of Diffusion Models via Operator Learning.** 42390-42402
- Wenqing Zheng, S. P. Sharan, Ajay Kumar Jaiswal, Kevin Wang, Yihan Xi, Dejia Xu, Zhangyang Wang:  
**Outline, Then Details: Syntactically Guided Coarse-To-Fine Code Generation.** 42403-42419
- Chenyu Zheng, Guoqiang Wu, Fan Bao, Yue Cao, Chongxuan Li, Jun Zhu:  
**Revisiting Discriminative vs. Generative Classifiers: Theory and Implications.** 42420-42477
- Ervine Zheng, Qi Yu:  
**Evidential Interactive Learning for Medical Image Captioning.** 42478-42491
- Yizhen Zheng, He Zhang, Vincent Cheng-Siong Lee, Yu Zheng, Xiao Wang, Shirui Pan:  
**Finding the Missing-half: Graph Complementary Learning for Homophily-prone and Heterophily-prone Graphs.** 42492-42505
- Shengwei Zhou, Rufan Bai, Xiaowei Wu:  
**Multi-agent Online Scheduling: MMS Allocations for Indivisible Items.** 42506-42516
- Dawei Zhou , Yukun Chen, Nannan Wang, Decheng Liu, Xinbo Gao, Tongliang Liu:  
**Eliminating Adversarial Noise via Information Discard and Robust Representation Restoration.** 42517-42530
- Yanqi Zhou, Nan Du, Yanping Huang, Daiyi Peng, Chang Lan, Da Huang, Siamak Shakeri, David R. So, Andrew M. Dai, Yifeng Lu, Zhifeng Chen, Quoc V. Le, Claire Cui, James Laudon, Jeff Dean:  
**Brainformers: Trading Simplicity for Efficiency.** 42531-42542
- Mo Zhou, Rong Ge:  
**Implicit Regularization Leads to Benign Overfitting for Sparse Linear Regression.** 42543-42573
- Zhi Zhou , Lan-Zhe Guo, Lin-Han Jia, Dingchu Zhang, Yufeng Li:  
**ODS: Test-Time Adaptation in the Presence of Open-World Data Shift.** 42574-42588



- Man Zhou, Jie Huang, Chun-Le Guo, Chongyi Li:  
**Fourmer: An Efficient Global Modeling Paradigm for Image Restoration.** 42589-42601
- Wangchunshu Zhou, Yuchen Eleanor Jiang, Ethan Wilcox, Ryan Cotterell , Mrinmaya Sachan:  
**Controlled Text Generation with Natural Language Instructions.** 42602-42613
- Tong Zhou, Yukui Luo, Shaolei Ren, Xiaolin Xu:  
**NNSplitter: An Active Defense Solution for DNN Model via Automated Weight Obfuscation.** 42614-42624
- Linqi Zhou, Michael Poli, Winnie Xu, Stefano Massaroli, Stefano Ermon:  
**Deep Latent State Space Models for Time-Series Generation.** 42625-42643
- Ziang Zhou, Jieming Shi, Renchi Yang , Yuanhang Zou, Qing Li:  
**SlotGAT: Slot-based Message Passing for Heterogeneous Graphs.** 42644-42657
- Baojian Zhou, Yifan Sun, Reza Babanezhad Harikandeh:  
**Fast Online Node Labeling for Very Large Graphs.** 42658-42697
- Runlong Zhou, Ruosong Wang, Simon Shaolei Du:  
**Horizon-Free and Variance-Dependent Reinforcement Learning for Latent Markov Decision Processes.** 42698-42723
- Dawei Zhou , Nannan Wang, Heng Yang, Xinbo Gao, Tongliang Liu:  
**Phase-aware Adversarial Defense for Improving Adversarial Robustness.** 42724-42741
- Cai Zhou, Xiyuan Wang, Muhan Zhang:  
**From Relational Pooling to Subgraph GNNs: A Universal Framework for More Expressive Graph Neural Networks.** 42742-42768
- Jianan Zhou, Yaxin Wu, Wen Song, Zhiguang Cao, Jie Zhang:  
**Towards Omni-generalizable Neural Methods for Vehicle Routing Problems.** 42769-42789
- Yefan Zhou, Yaoqing Yang, Arin Chang, Michael W. Mahoney:  
**A Three-regime Model of Network Pruning.** 42790-42809
- Zihan Zhou, Tianshu Yu:  
**Learning to Decouple Complex Systems.** 42810-42828
- Kaiwen Zhou, Kaizhi Zheng, Connor Pryor, Yilin Shen, Hongxia Jin, Lise Getoor, Xin Eric Wang:  
**ESC: Exploration with Soft Commonsense Constraints for Zero-shot Object Navigation.** 42829-42842
- Zhanke Zhou, Chenyu Zhou, Xuan Li, Jiangchao Yao, Quanming Yao, Bo Han:  
**On Strengthening and Defending Graph Reconstruction Attack with Markov Chain Approximation.** 42843-42877
- Runlong Zhou, Zihan Zhang, Simon Shaolei Du:  
**Sharp Variance-Dependent Bounds in Reinforcement Learning: Best of Both Worlds in Stochastic and Deterministic Environments.** 42878-42914
- Sicheng Zhu, Bang An, Furong Huang, Sanghyun Hong:  
**Learning Unforeseen Robustness from Out-of-distribution Data Using Equivariant Domain Translator.** 42915-42937
- Harrison Zhu, Carles Balsells Rodas, Yingzhen Li:  
**Markovian Gaussian Process Variational Autoencoders.** 42938-42961
- Yilun Zhu, Aaron Fjeldsted, Darren Holland, George Landon, Azaree Lintereur, Clayton Scott:  
**Mixture Proportion Estimation Beyond Irreducibility.** 42962-42982
- Jianing Zhu, Xiawei Guo, Jiangchao Yao, Chao Du, Li He, Shuo Yuan, Tongliang Liu, Liang Wang, Bo Han:  
**Exploring Model Dynamics for Accumulative Poisoning Discovery.** 42983-43004
- Tongtian Zhu, Fengxiang He, Kaixuan Chen, Mingli Song, Dacheng Tao:  
**Decentralized SGD and Average-direction SAM are Asymptotically Equivalent.** 43005-43036
- Banghua Zhu, Michael I. Jordan, Jiantao Jiao:  
**Principled Reinforcement Learning with Human Feedback from Pairwise or K-wise Comparisons.** 43037-43067
- Jianing Zhu, Hengzhuang Li, Jiangchao Yao, Tongliang Liu, Jianliang Xu, Bo Han:  
**Unleashing Mask: Explore the Intrinsic Out-of-Distribution Detection Capability.** 43068-43104
- Zhenyu Zhu, Fanghui Liu, Grigoris Chrysis, Francesco Locatello, Volkan Cevher:  
**Benign Overfitting in Deep Neural Networks under Lazy Training.** 43105-43128
- Jiacheng Zhu, Jielin Qiu, Aritra Guha, Zhuolin Yang, XuanLong Nguyen, Bo Li, Ding Zhao:  
**Interpolation for Robust Learning: Data Augmentation on Wasserstein Geodesics.** 43129-43157
- Chaoyi Zhu, Stefanie Roos, Lydia Y. Chen:  
**LeadFL: Client Self-Defense against Model Poisoning in Federated Learning.** 43158-43180
- Bingzhao Zhu, Xingjian Shi, Nick Erickson, Mu Li, George Karypis, Mahsa Shoaran:  
**XTab: Cross-table Pretraining for Tabular Transformers.** 43181-43204
- Dixian Zhu, Bokun Wang, Zhi Chen, Yaxing Wang, Milan Sonka, Xiaodong Wu, Tianbao Yang:  
**Provable Multi-instance Deep AUC Maximization with Stochastic Pooling.** 43205-43227
- Junyi Zhu, Ruicong Yao, Matthew B. Blaschko:  
**Surrogate Model Extension (SME): A Fast and Accurate Weight Update Attack on Federated Learning.** 43228-43257



Zhaowei Zhu, Yuanshun Yao, Jiankai Sun, Hang Li, Yang Liu:  
**Weak Proxies are Sufficient and Preferable for Fairness with Missing Sensitive Attributes.** 43258-43288

Dixian Zhu, Yiming Ying, Tianbao Yang:  
**Label Distributionally Robust Losses for Multi-class Classification: Consistency, Robustness and Adaptivity.** 43289-43325

Yubo Zhuang, Xiaohui Chen, Yun Yang:  
**Likelihood Adjusted Semidefinite Programs for Clustering Heterogeneous Data.** 43326-43346

Juliusz Krysztof Ziomek, Haitham Bou-Ammar:  
**Are Random Decompositions all we need in High Dimensional Bayesian Optimisation?** 43347-43368

Joshua P. Zitovsky, Daniel de Marchi, Rishabh Agarwal, Michael Rene Kosorok:  
**Revisiting Bellman Errors for Offline Model Selection.** 43369-43406

Ziyin Liu, Zihao Wang:  
**spred: Solving L1 Penalty with SGD.** 43407-43422

Difan Zou, Yuan Cao, Yuanzhi Li, Quanquan Gu:  
**The Benefits of Mixup for Feature Learning.** 43423-43479

---

last updated on 2024-12-11 20:31 CET by the dblp team

all metadata released as open data under CC0 1.0 license

see also: Terms of Use | Privacy Policy | Imprint

dblp was originally created in 1993 at:



since 2018, dblp has been operated and maintained by:



the dblp computer science bibliography is funded and supported by:

